

Mark3 Realtime Kernel

Generated by Doxygen 1.8.13

Contents

1	The Mark3 Realtime Kernel	1
2	License	3
2.1	License	3
3	Configuring The Mark3 Kernel	5
3.1	Overview	5
3.2	Timer Options	5
3.3	Blocking Objects	7
3.4	Inter-process/thread Communication	7
3.5	Debug Features	8
3.6	Enhancements, Security, Miscellaneous	8
4	Building Mark3	11
4.1	Source Layout	11
4.2	Toolchain Integration	11
4.3	Installing Dependencies	12
4.4	Building Mark3 Kernel and Libraries	12
4.5	Exporting the kernel source	13

5	Getting Started With The Mark3 API	15
5.1	Kernel Setup	15
5.2	Threads	16
5.2.1	Thread Setup	16
5.2.2	Entry Functions	17
5.3	Timers	18
5.4	Semaphores	18
5.5	Mutexes	19
5.6	Event Flags	20
5.7	Messages	20
5.7.1	Message Objects	21
5.7.2	Global Message Pool	21
5.7.3	Message Queues	22
5.7.4	Messaging Example	22
5.8	Mailboxes	22
5.8.1	Mailbox Example	23
5.9	Notification Objects	23
5.9.1	Notification Example	23
5.10	Condition Variables	24
5.10.1	Condition Variable Example	24
5.11	Reader-Write Locks	25
5.11.1	Reader-Write Lock Example	25
5.12	Sleep	25
5.13	Round-Robin Quantum	26
6	Why Mark3?	27

7	When should you use an RTOS?	29
7.1	The reality of system code	29
7.2	Superloops, and their limitations	30
7.2.1	Intro to Superloops	30
7.2.2	The simplest loop	31
7.2.3	Interrupt-Driven Super-loop	31
7.2.4	Cooperative multi-tasking	33
7.2.5	Hybrid cooperative/preemptive multi-tasking	34
7.3	Problems with superloops	35
7.3.1	Hidden Costs	35
7.3.2	Tightly-coupled code	35
7.3.3	No blocking Calls	35
7.3.4	Difficult to guarantee responsiveness	35
7.3.5	Limited preemption capability	36
8	Can you afford an RTOS?	37
8.1	Intro	37
8.2	Application description	38
8.3	Runtime Overhead	39
8.4	Analysis	40
9	Mark3 Design Goals	41
9.1	Overview	41
9.1.1	Services Provided by an RTOS Kernel	41
9.1.2	Guiding Principles of Mark3	41
9.1.3	Be feature competitive	41
9.1.4	Be highly configuration	41
9.1.5	No external dependencies, no new language features	42
9.1.6	Target the most popular hobbyist platforms available	42
9.1.7	Maximize determinism – but be pragmatic	42
9.1.8	Apply engineering principles – and that means discipline, measurement and verification	42
9.1.9	Use Virtualization For Verification	43

10 Mark3 Kernel Architecture	45
10.1 Overview	45
10.2 Threads and Scheduling	47
10.2.1 A Bit About Threads	48
10.2.2 Thread States and ThreadLists	49
10.2.3 Blocking and Unblocking	49
10.2.4 Blocking Objects	49
10.3 Inside the Mark3 Scheduler	50
10.3.1 Considerations for Round-Robin Scheduling	52
10.3.2 Context Switching	53
10.3.3 Putting It All Together	53
10.4 Timers	53
10.4.1 Tick-based Timers	55
10.4.2 Tickless Timers	55
10.4.3 Timer Processing Algorithm	55
10.5 Synchronization and IPC	56
10.6 Blocking Objects	56
10.6.1 Semaphores	57
10.6.2 Mutex	57
10.6.3 Event Flags	57
10.6.4 Notification Objects	58
10.7 Messages and Global Message Queue	58
10.7.1 Messages	58
10.7.2 Message Objects	58
10.7.3 Global Message Pool	59
10.7.4 Message Queues	59
10.7.5 Mailboxes	59
10.7.6 Atomic Operations	60
10.7.7 Atomic Operations	60
10.7.8 Drivers	61
10.8 Kernel Proper and Porting	63

11 Mark3C - C-language API bindings for the Mark3 Kernel.	71
11.1 API Conventions	71
11.2 Allocating Objects	72
12 Release Notes	73
12.1 R7 Release	73
12.2 R6 Release	73
12.3 R5 Release	74
12.4 R4 Release	74
12.5 R3 Release	75
12.6 R2	75
12.7 R1 - 2nd Release Candidate	75
12.8 R1 - 1st Release Candidate	75
13 Code Size Profiling	77
13.1 Information	77
13.2 Compiler Version	77
13.3 Profiling Results	77
14 Namespace Index	79
14.1 Namespace List	79
15 Hierarchical Index	81
15.1 Class Hierarchy	81
16 Class Index	83
16.1 Class List	83
17 File Index	85
17.1 File List	85

18 Namespace Documentation	89
18.1 Mark3 Namespace Reference	89
18.1.1 Detailed Description	91
18.1.2 Typedef Documentation	91
18.1.2.1 TimerCallback	91
18.1.3 Enumeration Type Documentation	91
18.1.3.1 EventFlagOperation	91
18.2 Mark3::Atomic Namespace Reference	92
18.2.1 Detailed Description	92
18.2.2 Function Documentation	92
18.2.2.1 Add()	92
18.2.2.2 Set()	93
18.2.2.3 Sub()	93
18.2.2.4 TestAndSet()	94
18.3 Mark3::KernelAware Namespace Reference	94
18.3.1 Detailed Description	95
18.3.2 Function Documentation	95
18.3.2.1 ExitSimulator()	95
18.3.2.2 IsSimulatorAware()	95
18.3.2.3 Print()	95
18.3.2.4 ProfileInit()	96
18.3.2.5 ProfileReport()	96
18.3.2.6 ProfileStart()	96
18.3.2.7 ProfileStop()	97
18.3.2.8 Trace() [1/3]	97
18.3.2.9 Trace() [2/3]	97
18.3.2.10 Trace() [3/3]	98

19 Class Documentation	99
19.1 Mark3::BlockingObject Class Reference	99
19.1.1 Detailed Description	100
19.1.2 Member Function Documentation	100
19.1.2.1 Block()	100
19.1.2.2 BlockPriority()	100
19.1.2.3 IsInitialized()	101
19.1.2.4 Unblock()	101
19.2 Mark3::CircularLinkedList Class Reference	102
19.2.1 Detailed Description	102
19.2.2 Member Function Documentation	102
19.2.2.1 Add()	102
19.2.2.2 InsertNodeBefore()	103
19.2.2.3 PivotBackward()	103
19.2.2.4 PivotForward()	103
19.2.2.5 Remove()	104
19.3 Mark3::ConditionVariable Class Reference	104
19.3.1 Detailed Description	104
19.3.2 Member Function Documentation	105
19.3.2.1 Init()	105
19.3.2.2 Wait() [1/2]	105
19.3.2.3 Wait() [2/2]	105
19.4 Mark3::DoubleLinkedList Class Reference	106
19.4.1 Detailed Description	106
19.4.2 Constructor & Destructor Documentation	106
19.4.2.1 DoubleLinkedList()	107
19.4.3 Member Function Documentation	107
19.4.3.1 Add()	107
19.4.3.2 Remove()	107
19.5 Mark3::EventFlag Class Reference	108

19.5.1 Detailed Description	109
19.5.2 Member Function Documentation	109
19.5.2.1 Clear()	109
19.5.2.2 GetMask()	109
19.5.2.3 Set()	110
19.5.2.4 Wait() [1/2]	110
19.5.2.5 Wait() [2/2]	110
19.5.2.6 Wait_i()	112
19.5.2.7 WakeMe()	112
19.6 Mark3::FakeThread_t Struct Reference	113
19.6.1 Detailed Description	113
19.7 Mark3::Kernel Class Reference	114
19.7.1 Detailed Description	115
19.7.2 Member Function Documentation	115
19.7.2.1 GetIdleThread()	115
19.7.2.2 GetThreadContextSwitchCallout()	115
19.7.2.3 GetThreadCreateCallout()	116
19.7.2.4 GetThreadExitCallout()	116
19.7.2.5 Init()	116
19.7.2.6 IsPanic()	117
19.7.2.7 IsStarted()	117
19.7.2.8 Panic()	117
19.7.2.9 SetIdleFunc()	117
19.7.2.10 SetPanic()	118
19.7.2.11 SetThreadContextSwitchCallout()	118
19.7.2.12 SetThreadCreateCallout()	119
19.7.2.13 SetThreadExitCallout()	119
19.7.2.14 Start()	120
19.8 Mark3::KernelTimer Class Reference	120
19.8.1 Detailed Description	121

19.8.2	Member Function Documentation	121
19.8.2.1	ClearExpiry()	121
19.8.2.2	Config()	122
19.8.2.3	DI()	122
19.8.2.4	EI()	122
19.8.2.5	GetOvertime()	122
19.8.2.6	Read()	123
19.8.2.7	RI()	123
19.8.2.8	SetExpiry()	123
19.8.2.9	Start()	124
19.8.2.10	Stop()	124
19.8.2.11	SubtractExpiry()	124
19.8.2.12	TimeToExpiry()	125
19.9	Mark3::LinkedList Class Reference	125
19.9.1	Detailed Description	126
19.9.2	Member Function Documentation	126
19.9.2.1	GetHead()	126
19.9.2.2	GetTail()	127
19.9.2.3	Init()	127
19.10	Mark3::LinkedListNode Class Reference	127
19.10.1	Detailed Description	128
19.10.2	Member Function Documentation	128
19.10.2.1	ClearNode()	128
19.10.2.2	GetNext()	128
19.10.2.3	GetPrev()	129
19.11	Mark3::LockGuard Class Reference	129
19.11.1	Detailed Description	129
19.11.2	Constructor & Destructor Documentation	129
19.11.2.1	LockGuard() [1/2]	129
19.11.2.2	LockGuard() [2/2]	130

19.11.3 Member Function Documentation	130
19.11.3.1 isAcquired()	130
19.12Mark3::Mailbox Class Reference	130
19.12.1 Detailed Description	132
19.12.2 Member Function Documentation	132
19.12.2.1 CopyData()	132
19.12.2.2 GetHeadPointer()	133
19.12.2.3 GetTailPointer()	133
19.12.2.4 Init() [1/2]	133
19.12.2.5 Init() [2/2]	134
19.12.2.6 MoveHeadBackward()	134
19.12.2.7 MoveHeadForward()	134
19.12.2.8 MoveTailBackward()	135
19.12.2.9 MoveTailForward()	135
19.12.2.10Receive() [1/2]	135
19.12.2.11Receive() [2/2]	135
19.12.2.12Receive_i()	136
19.12.2.13ReceiveTail() [1/2]	136
19.12.2.14ReceiveTail() [2/2]	137
19.12.2.15Send() [1/2]	137
19.12.2.16Send() [2/2]	138
19.12.2.17Send_i()	138
19.12.2.18SendTail() [1/2]	139
19.12.2.19SendTail() [2/2]	139
19.12.3 Member Data Documentation	140
19.12.3.1 m_clSendSem	140
19.13Mark3::Message Class Reference	140
19.13.1 Detailed Description	141
19.13.2 Member Function Documentation	141
19.13.2.1 GetCode()	141

19.13.2.2 GetData()	141
19.13.2.3 Init()	142
19.13.2.4 SetCode()	142
19.13.2.5 SetData()	142
19.14Mark3::MessagePool Class Reference	143
19.14.1 Detailed Description	143
19.14.2 Member Function Documentation	143
19.14.2.1 GetHead()	144
19.14.2.2 Init()	144
19.14.2.3 Pop()	144
19.14.2.4 Push()	144
19.15Mark3::MessageQueue Class Reference	145
19.15.1 Detailed Description	146
19.15.2 Member Function Documentation	146
19.15.2.1 GetCount()	146
19.15.2.2 Init()	146
19.15.2.3 Receive() [1/2]	146
19.15.2.4 Receive() [2/2]	146
19.15.2.5 Receive_i()	147
19.15.2.6 Send()	147
19.16Mark3::Mutex Class Reference	148
19.16.1 Detailed Description	149
19.16.2 Member Function Documentation	149
19.16.2.1 Claim() [1/2]	149
19.16.2.2 Claim() [2/2]	149
19.16.2.3 Claim_i()	150
19.16.2.4 Init()	150
19.16.2.5 Release()	151
19.16.2.6 WakeMe()	151
19.16.2.7 WakeNext()	151

19.17Mark3::Notify Class Reference	152
19.17.1 Detailed Description	152
19.17.2 Member Function Documentation	152
19.17.2.1 Init()	153
19.17.2.2 Signal()	153
19.17.2.3 Wait() [1/2]	153
19.17.2.4 Wait() [2/2]	153
19.17.2.5 WakeMe()	154
19.18Mark3::PriorityMap Class Reference	154
19.18.1 Detailed Description	155
19.18.2 Constructor & Destructor Documentation	155
19.18.2.1 PriorityMap()	155
19.18.3 Member Function Documentation	155
19.18.3.1 Clear()	155
19.18.3.2 HighestPriority()	155
19.18.3.3 Set()	156
19.19Mark3::ProfileTimer Class Reference	156
19.19.1 Detailed Description	157
19.19.2 Member Function Documentation	157
19.19.2.1 ComputeCurrentTicks()	157
19.19.2.2 GetAverage()	158
19.19.2.3 GetCurrent()	158
19.19.2.4 Init()	158
19.19.2.5 Start()	158
19.19.2.6 Stop()	159
19.20Mark3::Quantum Class Reference	159
19.20.1 Detailed Description	159
19.20.2 Member Function Documentation	159
19.20.2.1 AddThread()	160
19.20.2.2 ClearInTimer()	160

19.20.2.3 RemoveThread()	160
19.20.2.4 SetInTimer()	160
19.20.2.5 SetTimer()	160
19.20.2.6 UpdateTimer()	161
19.21Mark3::ReaderWriterLock Class Reference	161
19.21.1 Detailed Description	162
19.21.2 Member Function Documentation	162
19.21.2.1 AcquireReader() [1/2]	162
19.21.2.2 AcquireReader() [2/2]	162
19.21.2.3 AcquireReader_i()	163
19.21.2.4 AcquireWriter() [1/2]	163
19.21.2.5 AcquireWriter() [2/2]	163
19.21.2.6 AcquireWriter_i()	164
19.21.2.7 Init()	164
19.22Mark3::Scheduler Class Reference	164
19.22.1 Detailed Description	165
19.22.2 Member Function Documentation	165
19.22.2.1 Add()	166
19.22.2.2 GetCurrentThread()	167
19.22.2.3 GetNextThread()	167
19.22.2.4 GetStopList()	168
19.22.2.5 GetThreadList()	168
19.22.2.6 Init()	168
19.22.2.7 IsEnabled()	169
19.22.2.8 QueueScheduler()	169
19.22.2.9 Remove()	169
19.22.2.10Schedule()	169
19.22.2.11SetScheduler()	170
19.23Mark3::Semaphore Class Reference	170
19.23.1 Detailed Description	171

19.23.2 Member Function Documentation	171
19.23.2.1 GetCount()	172
19.23.2.2 Init()	172
19.23.2.3 Pend() [1/2]	172
19.23.2.4 Pend() [2/2]	173
19.23.2.5 Pend_i()	173
19.23.2.6 Post()	173
19.23.2.7 WakeMe()	174
19.23.2.8 WakeNext()	174
19.24 Mark3::Thread Class Reference	174
19.24.1 Detailed Description	177
19.24.2 Member Function Documentation	177
19.24.2.1 ContextSwitchSWI()	178
19.24.2.2 Exit()	178
19.24.2.3 GetCurPriority()	178
19.24.2.4 GetCurrent()	179
19.24.2.5 GetEventFlagMask()	179
19.24.2.6 GetEventFlagMode()	179
19.24.2.7 GetExpired()	180
19.24.2.8 GetExtendedContext()	180
19.24.2.9 GetID()	180
19.24.2.10 GetOwner()	181
19.24.2.11 GetPriority()	181
19.24.2.12 GetQuantum()	181
19.24.2.13 GetStack()	182
19.24.2.14 GetStackSize()	182
19.24.2.15 GetStackSlack()	182
19.24.2.16 GetState()	183
19.24.2.17 InheritPriority()	183
19.24.2.18 Init() [1/2]	183

19.24.2.19	init() [2/2]	184
19.24.2.20	initIdle()	184
19.24.2.21	setCurrent()	185
19.24.2.22	setEventFlagMask()	185
19.24.2.23	setEventFlagMode()	185
19.24.2.24	setExpired()	186
19.24.2.25	setExtendedContext()	186
19.24.2.26	setId()	186
19.24.2.27	setOwner()	187
19.24.2.28	setPriority()	187
19.24.2.29	setPriorityBase()	188
19.24.2.30	setQuantum()	188
19.24.2.31	setState()	188
19.24.2.32	sleep()	189
19.24.2.33	start()	189
19.24.2.34	stop()	189
19.24.2.35	uSleep()	190
19.24.2.36	yield()	190
19.25	Mark3::ThreadList Class Reference	190
19.25.1	Detailed Description	191
19.25.2	Constructor & Destructor Documentation	191
19.25.2.1	ThreadList()	191
19.25.3	Member Function Documentation	192
19.25.3.1	Add() [1/2]	192
19.25.3.2	Add() [2/2]	192
19.25.3.3	AddPriority()	192
19.25.3.4	HighestWaiter()	193
19.25.3.5	Remove()	193
19.25.3.6	SetMapPointer()	193
19.25.3.7	SetPriority()	194

19.26Mark3::ThreadPort Class Reference	194
19.26.1 Detailed Description	195
19.26.2 Member Function Documentation	195
19.26.2.1 InitStack()	195
19.26.2.2 StartThreads()	195
19.27Mark3::Timer Class Reference	196
19.27.1 Detailed Description	197
19.27.2 Constructor & Destructor Documentation	197
19.27.2.1 Timer()	198
19.27.3 Member Function Documentation	198
19.27.3.1 GetInterval()	198
19.27.3.2 Init()	198
19.27.3.3 IsInitialized()	198
19.27.3.4 SetCallback()	198
19.27.3.5 SetData()	199
19.27.3.6 SetFlags()	199
19.27.3.7 SetIntervalMSeconds()	199
19.27.3.8 SetIntervalSeconds()	200
19.27.3.9 SetIntervalTicks()	200
19.27.3.10SetIntervalUSeconds()	200
19.27.3.11SetOwner()	201
19.27.3.12SetTolerance()	201
19.27.3.13Start() [1/3]	201
19.27.3.14Start() [2/3]	203
19.27.3.15Start() [3/3]	203
19.27.3.16Stop()	204
19.28Mark3::TimerList Class Reference	204
19.28.1 Detailed Description	205
19.28.2 Member Function Documentation	205
19.28.2.1 Add()	205
19.28.2.2 Init()	205
19.28.2.3 Process()	205
19.28.2.4 Remove()	205
19.29Mark3::TimerScheduler Class Reference	206
19.29.1 Detailed Description	206
19.29.2 Member Function Documentation	206
19.29.2.1 Add()	206
19.29.2.2 Init()	207
19.29.2.3 Process()	207
19.29.2.4 Remove()	207

20 File Documentation	209
20.1 /home/moslevin/projects/github/m3-repo/kernel/libs/mark3c/src/public/fake_types.h File Reference	209
20.1.1 Detailed Description	209
20.2 fake_types.h	209
20.3 /home/moslevin/projects/github/m3-repo/kernel/libs/mark3c/src/public/mark3c.h File Reference	212
20.3.1 Detailed Description	213
20.3.2 Enumeration Type Documentation	214
20.3.2.1 event_flag_operation_t	214
20.3.3 Function Documentation	214
20.3.3.1 Kernel_Init()	214
20.3.3.2 Kernel_IsPanic()	214
20.3.3.3 Kernel_IsStarted()	215
20.3.3.4 Kernel_Panic()	215
20.3.3.5 Kernel_SetPanic()	215
20.3.3.6 Kernel_Start()	216
20.3.3.7 Scheduler_Enable()	216
20.3.3.8 Scheduler_GetCurrentThread()	216
20.3.3.9 Scheduler_IsEnabled()	217
20.3.3.10 Thread_GetCurPriority()	217
20.3.3.11 Thread_GetID()	218
20.3.3.12 Thread_GetPriority()	218
20.3.3.13 Thread_GetStackSlack()	219
20.3.3.14 Thread_GetState()	219
20.3.3.15 Thread_Init()	220
20.3.3.16 Thread_SetID()	220
20.3.3.17 Thread_SetPriority()	221
20.3.3.18 Thread_Start()	221
20.3.3.19 Thread_Stop()	221
20.3.3.20 Thread_Yield()	222
20.4 mark3c.h	222

20.5	/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/kernelprofile.cpp	
	File Reference	228
20.5.1	Detailed Description	228
20.6	kernelprofile.cpp	229
20.7	/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/kernelswi.cpp	File
	Reference	230
20.7.1	Detailed Description	230
20.8	kernelswi.cpp	230
20.9	/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/kerneltimer.cpp	File
	Reference	231
20.9.1	Detailed Description	231
20.10	kerneltimer.cpp	232
20.11	/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/kernelprofile.h	
	File Reference	234
20.11.1	Detailed Description	234
20.12	kernelprofile.h	235
20.13	/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/kernelswi.h	
	File Reference	235
20.13.1	Detailed Description	235
20.14	kernelswi.h	236
20.15	/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/kerneltimer.h	
	File Reference	236
20.15.1	Detailed Description	236
20.16	kerneltimer.h	237
20.17	/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/portcfg.h	File
	Reference	237
20.17.1	Detailed Description	238
20.17.2	Macro Definition Documentation	238
20.17.2.1	AVR	238
20.17.2.2	K_WORD	239
20.17.2.3	PORT_PRIO_TYPE	239
20.17.2.4	PORT_SYSTEM_FREQ	239
20.17.2.5	PORT_TIMER_COUNT_TYPE	239

20.17.2.6 PORT_TIMER_FREQ	240
20.18portcfg.h	240
20.19/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/threadport.h File Reference	240
20.19.1 Detailed Description	241
20.19.2 Macro Definition Documentation	241
20.19.2.1 CS_ENTER	242
20.20threadport.h	242
20.21/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/threadport.cpp File Reference	244
20.21.1 Detailed Description	245
20.22threadport.cpp	245
20.23/home/moslevin/projects/github/m3-repo/kernel/src/atomic.cpp File Reference	247
20.23.1 Detailed Description	247
20.24atomic.cpp	247
20.25/home/moslevin/projects/github/m3-repo/kernel/src/autoalloc.cpp File Reference	249
20.25.1 Detailed Description	249
20.26autoalloc.cpp	249
20.27/home/moslevin/projects/github/m3-repo/kernel/src/blocking.cpp File Reference	253
20.27.1 Detailed Description	253
20.28blocking.cpp	254
20.29/home/moslevin/projects/github/m3-repo/kernel/src/condvar.cpp File Reference	255
20.29.1 Detailed Description	255
20.30condvar.cpp	255
20.31/home/moslevin/projects/github/m3-repo/kernel/src/eventflag.cpp File Reference	256
20.31.1 Detailed Description	257
20.32eventflag.cpp	257
20.33/home/moslevin/projects/github/m3-repo/kernel/src/kernel.cpp File Reference	261
20.33.1 Detailed Description	261
20.34kernel.cpp	261
20.35/home/moslevin/projects/github/m3-repo/kernel/src/kernelaware.cpp File Reference	263

20.35.1 Detailed Description	263
20.36kernelaware.cpp	263
20.37/home/moslevin/projects/github/m3-repo/kernel/src/ksemaphore.cpp File Reference	265
20.37.1 Detailed Description	265
20.38ksemaphore.cpp	266
20.39/home/moslevin/projects/github/m3-repo/kernel/src/ll.cpp File Reference	269
20.39.1 Detailed Description	269
20.40ll.cpp	269
20.41/home/moslevin/projects/github/m3-repo/kernel/src/lockguard.cpp File Reference	271
20.41.1 Detailed Description	271
20.42lockguard.cpp	272
20.43/home/moslevin/projects/github/m3-repo/kernel/src/mailbox.cpp File Reference	272
20.43.1 Detailed Description	272
20.44mailbox.cpp	273
20.45/home/moslevin/projects/github/m3-repo/kernel/src/message.cpp File Reference	276
20.45.1 Detailed Description	276
20.46message.cpp	276
20.47/home/moslevin/projects/github/m3-repo/kernel/src/mutex.cpp File Reference	278
20.47.1 Detailed Description	278
20.48mutex.cpp	279
20.49/home/moslevin/projects/github/m3-repo/kernel/src/notify.cpp File Reference	282
20.49.1 Detailed Description	282
20.50notify.cpp	282
20.51/home/moslevin/projects/github/m3-repo/kernel/src/priomap.cpp File Reference	285
20.51.1 Detailed Description	285
20.52priomap.cpp	285
20.53/home/moslevin/projects/github/m3-repo/kernel/src/profile.cpp File Reference	287
20.53.1 Detailed Description	287
20.54profile.cpp	287
20.55/home/moslevin/projects/github/m3-repo/kernel/src/public/atomic.h File Reference	289

20.55.1 Detailed Description	289
20.56atomic.h	289
20.57/home/moslevin/projects/github/m3-repo/kernel/src/public/autoalloc.h File Reference	290
20.57.1 Detailed Description	290
20.58autoalloc.h	290
20.59/home/moslevin/projects/github/m3-repo/kernel/src/public/blocking.h File Reference	293
20.59.1 Detailed Description	293
20.60blocking.h	294
20.61/home/moslevin/projects/github/m3-repo/kernel/src/public/buffalogger.h File Reference	294
20.61.1 Detailed Description	294
20.62buffalogger.h	295
20.63/home/moslevin/projects/github/m3-repo/kernel/src/public/condvar.h File Reference	295
20.63.1 Detailed Description	295
20.64condvar.h	296
20.65/home/moslevin/projects/github/m3-repo/kernel/src/public/eventflag.h File Reference	296
20.65.1 Detailed Description	297
20.66eventflag.h	297
20.67/home/moslevin/projects/github/m3-repo/kernel/src/public/kernel.h File Reference	298
20.67.1 Detailed Description	298
20.68kernel.h	298
20.69/home/moslevin/projects/github/m3-repo/kernel/src/public/kernelaware.h File Reference	299
20.69.1 Detailed Description	300
20.70kernelaware.h	301
20.71/home/moslevin/projects/github/m3-repo/kernel/src/public/kerneldebug.h File Reference	301
20.71.1 Detailed Description	302
20.72kerneldebug.h	302
20.73/home/moslevin/projects/github/m3-repo/kernel/src/public/kerneltypes.h File Reference	307
20.73.1 Detailed Description	308
20.74kerneltypes.h	308
20.75/home/moslevin/projects/github/m3-repo/kernel/src/public/ksemaphore.h File Reference	309

20.75.1 Detailed Description	309
20.76ksemaphore.h	309
20.77/home/moslevin/projects/github/m3-repo/kernel/src/public/ll.h File Reference	310
20.77.1 Detailed Description	311
20.78ll.h	311
20.79/home/moslevin/projects/github/m3-repo/kernel/src/public/lockguard.h File Reference	312
20.79.1 Detailed Description	313
20.80lockguard.h	313
20.81/home/moslevin/projects/github/m3-repo/kernel/src/public/mailbox.h File Reference	313
20.81.1 Detailed Description	314
20.82mailbox.h	314
20.83/home/moslevin/projects/github/m3-repo/kernel/src/public/manual.h File Reference	316
20.83.1 Detailed Description	316
20.84manual.h	316
20.85/home/moslevin/projects/github/m3-repo/kernel/src/public/mark3.h File Reference	316
20.85.1 Detailed Description	317
20.86mark3.h	317
20.87/home/moslevin/projects/github/m3-repo/kernel/src/public/mark3cfg.h File Reference	318
20.87.1 Detailed Description	319
20.87.2 Macro Definition Documentation	319
20.87.2.1 GLOBAL_MESSAGE_POOL_SIZE	320
20.87.2.2 KERNEL_AWARE_SIMULATION	320
20.87.2.3 KERNEL_EXTRA_CHECKS	320
20.87.2.4 KERNEL_NUM_PRIORITIES	320
20.87.2.5 KERNEL_TIMERS_MINIMUM_DELAY_US	321
20.87.2.6 KERNEL_TIMERS_THREADED	321
20.87.2.7 KERNEL_TIMERS_TICKLESS	321
20.87.2.8 KERNEL_USE_ATOMIC	321
20.87.2.9 KERNEL_USE_AUTO_ALLOC	322
20.87.2.10KERNEL_USE_CONDVAR	322

20.87.2.11	KERNEL_USE_DYNAMIC_THREADS	322
20.87.2.12	KERNEL_USE_EVENTFLAG	322
20.87.2.13	KERNEL_USE_IDLE_FUNC	323
20.87.2.14	KERNEL_USE_MAILBOX	323
20.87.2.15	KERNEL_USE_MESSAGE	323
20.87.2.16	KERNEL_USE_PROFILER	323
20.87.2.17	KERNEL_USE_QUANTUM	324
20.87.2.18	KERNEL_USE_READERWRITER	324
20.87.2.19	KERNEL_USE_SEMAPHORE	324
20.87.2.20	KERNEL_USE_STACK_GUARD	324
20.87.2.21	KERNEL_USE_THREAD_CALLOUTS	325
20.87.2.22	KERNEL_USE_THREADNAME	325
20.87.2.23	KERNEL_USE_TIMEOUTS	325
20.87.2.24	KERNEL_USE_TIMERS	325
20.87.2.25	SAFE_UNLINK	326
20.87.2.26	THREAD_QUANTUM_DEFAULT	326
20.88	mark3cfg.h	326
20.89	/home/moslevin/projects/github/m3-repo/kernel/src/public/message.h File Reference	328
20.89.1	Detailed Description	328
20.89.2	using Messages, Queues, and the Global Message Pool	329
20.90	message.h	329
20.91	/home/moslevin/projects/github/m3-repo/kernel/src/public/mutex.h File Reference	330
20.91.1	Detailed Description	331
20.91.2	Initializing	331
20.91.3	Resource protection example	331
20.92	mutex.h	332
20.93	/home/moslevin/projects/github/m3-repo/kernel/src/public/notify.h File Reference	332
20.93.1	Detailed Description	333
20.94	notify.h	333
20.95	/home/moslevin/projects/github/m3-repo/kernel/src/public/paniccodes.h File Reference	334

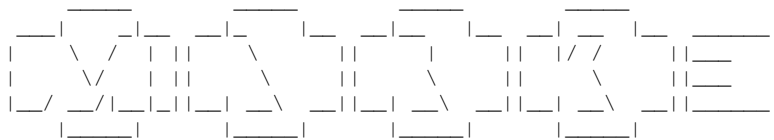
20.95.1 Detailed Description	334
20.96paniccodes.h	334
20.97/home/moslevin/projects/github/m3-repo/kernel/src/public/priomap.h File Reference	334
20.97.1 Detailed Description	335
20.98priomap.h	335
20.99/home/moslevin/projects/github/m3-repo/kernel/src/public/profile.h File Reference	336
20.99.1 Detailed Description	336
20.100profile.h	337
20.101/home/moslevin/projects/github/m3-repo/kernel/src/public/quantum.h File Reference	337
20.101.1 Detailed Description	338
20.102quantum.h	338
20.103/home/moslevin/projects/github/m3-repo/kernel/src/public/readerwriter.h File Reference	339
20.103.1 Detailed Description	339
20.104readerwriter.h	339
20.105/home/moslevin/projects/github/m3-repo/kernel/src/public/scheduler.h File Reference	340
20.105.1 Detailed Description	341
20.106scheduler.h	341
20.107/home/moslevin/projects/github/m3-repo/kernel/src/public/thread.h File Reference	342
20.107.1 Detailed Description	342
20.108thread.h	343
20.109/home/moslevin/projects/github/m3-repo/kernel/src/public/threadlist.h File Reference	345
20.109.1 Detailed Description	346
20.110threadlist.h	346
20.111/home/moslevin/projects/github/m3-repo/kernel/src/public/timer.h File Reference	346
20.111.1 Detailed Description	347
20.111.2 Macro Definition Documentation	347
20.111.2.1TIMERLIST_FLAG_EXPIRED	347
20.112timer.h	348
20.113/home/moslevin/projects/github/m3-repo/kernel/src/public/timerlist.h File Reference	349
20.113.1 Detailed Description	350

20.114	merlist.h	350
20.115	home/moslevin/projects/github/m3-repo/kernel/src/public/timerscheduler.h File Reference	351
20.115.1	Detailed Description	351
20.116	merscheduler.h	351
20.117	home/moslevin/projects/github/m3-repo/kernel/src/public/tracebuffer.h File Reference	352
20.117.1	Detailed Description	352
20.118	acebuffer.h	352
20.119	home/moslevin/projects/github/m3-repo/kernel/src/quantum.cpp File Reference	353
20.119.1	Detailed Description	353
20.120	quantum.cpp	353
20.121	home/moslevin/projects/github/m3-repo/kernel/src/readerwriter.cpp File Reference	355
20.121.1	Detailed Description	355
20.122	readerwriter.cpp	355
20.123	home/moslevin/projects/github/m3-repo/kernel/src/scheduler.cpp File Reference	357
20.123.1	Detailed Description	357
20.124	cheduler.cpp	358
20.125	home/moslevin/projects/github/m3-repo/kernel/src/thread.cpp File Reference	359
20.125.1	Detailed Description	359
20.126	hread.cpp	360
20.127	home/moslevin/projects/github/m3-repo/kernel/src/threadlist.cpp File Reference	366
20.127.1	Detailed Description	367
20.128	hreadlist.cpp	367
20.129	home/moslevin/projects/github/m3-repo/kernel/src/timer.cpp File Reference	368
20.129.1	Detailed Description	369
20.130	mer.cpp	369
20.131	home/moslevin/projects/github/m3-repo/kernel/src/timerlist.cpp File Reference	371
20.131.1	Detailed Description	371
20.132	merlist.cpp	372
20.133	home/moslevin/projects/github/m3-repo/kernel/src/tracebuffer.cpp File Reference	375
20.133.1	Detailed Description	375
20.134	acebuffer.cpp	375

21 Example Documentation	377
21.1 buffalogger/main.cpp	377
21.2 lab10_notifications/main.cpp	377
21.3 lab11_mailboxes/main.cpp	378
21.4 lab1_kernel_setup/main.cpp	380
21.5 lab2_idle_function/main.cpp	382
21.6 lab3_round_robin/main.cpp	383
21.7 lab4_semaphores/main.cpp	385
21.8 lab5_mutexes/main.cpp	386
21.9 lab6_timers/main.cpp	388
21.10lab7_events/main.cpp	390
21.11lab8_messages/main.cpp	392
21.12lab9_dynamic_threads/main.cpp	394
 Index	 399

Chapter 1

The Mark3 Realtime Kernel



--[Mark3 Realtime Platform]-----

Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
See license for more information

The [Mark3](#) Realtime Kernel is a completely free, open-source, real-time operating system aimed at bringing powerful, easy-to-use multitasking to microcontroller systems without MMUs.

It uses modern programming languages and concepts to minimize code duplication, and its object-oriented design enhances readability. The API is simple – in six function calls, you can set up the kernel, initialize two threads, and start the scheduler.

The source is fully-documented with example code provided to illustrate concepts. The result is a performant RTOS, which is easy to read, easy to understand, and easy to extend to fit your needs.

But [Mark3](#) is bigger than just a real-time kernel, it also contains a number of class-leading features:

- Native implementation in C++, with C-language bindings.
- Device driver HAL which provides a meaningful abstraction around device-specific peripherals.
- CMake-based build system which can be used to build all libraries, examples, tests, documentation, and user-projects for any number of targets from the command-line.
- Graphics and UI code designed to simplify the implementation of systems using displays, keypads, joysticks, and touchscreens
- Robust and deterministic dynamic memory management libraries
- A Variety of general-purpose libraries to speed up embedded app development
- Emulator-aware debugging via the fIAVR AVR emulator
- A bulletproof, well-documented bootloader for AVR microcontrollers Support for kernel-aware simulators, including Funkenstein's own fIAVR.

Chapter 2

License

2.1 License

Copyright (c) 2012 - 2018, m0slevin
All rights reserved.

Redistribution and use in source and binary forms, with or without modification, are permitted provided that the following conditions are met:

- Redistributions of source code must retain the above copyright notice, this list of conditions and the following disclaimer.
- Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution.
- Neither the name of m0slevin, nor the names of its contributors may be used to endorse or promote products derived from this software without specific prior written permission.

THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL FUNKENSTEIN SOFTWARE AND/OR ITS CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

Chapter 3

Configuring The Mark3 Kernel

3.1 Overview

The [Mark3](#) Kernel features a large number of compile-time options that can be set by the user. In this way, the user can build a custom OS kernel that provides only the necessary feature set required by the application, and reduce the code and data requirements of the kernel.

Care has been taken to ensure that all valid combinations of features can be enabled or disabled, barring direct dependencies.

When [Mark3](#) is built, the various compile-time definitions are used to alter how the kernel is compiled, and include or exclude various bits and pieces in order to satisfy the requirements of the selected features. As a result, the kernel must be rebuilt whenever changes are made to the configuration header.

Note that not all demos, libraries, and tests will build successfully if the prerequisite features are not included.

Kernel options are set by modifying [mark3cfg.h](#), located within the /kernel/public folder.

In the following sections, we will discuss the various configuration options, grouped by functionality.

3.2 Timer Options

KERNEL_USE_TIMERS

This option is related to all kernel time-tracking:

- Timers provide a way for events to be periodically triggered in a lightweight manner. These can be periodic, or one-shot.
- Thread Quantum (used for round-robin scheduling) is dependent on this module, as is Thread Sleep functionality.

Setting this option to 0 disables all timer-based functionality within the kernel.

KERNEL_TIMERS_TICKLESS

If you've opted to use the kernel timers module, you have an option as to which timer implementation to use: Tick-based or Tick-less.

Tick-based timers provide a "traditional" RTOS timer implementation based on a fixed-frequency timer interrupt. While this provides very accurate, reliable timing, it also means that the CPU is being interrupted far more often than may be necessary (as not all timer ticks result in "real work" being done).

Tick-less timerLs still rely on a hardware timer interrupt, but uses a dynamic expiry interval to ensure that the interrupt is only called when the next timer expires. This increases the complexity of the timer interrupt handler, but reduces the number and frequency.

Note that the CPU port ([kerneltimer.cpp](#)) must be implemented for the particular timer variant desired.

Set this option to 1 to use the tickless timer implementation, 0 to use the traditional tick-based approach. Tickless timers are a bit more heavy weight (larger code footprint), but can yield significant power savings as the CPU does not need to wake up at a fixed, high frequency.

KERNEL_USE_TIMEOUTS

By default, if you opt to enable kernel timers, you also get timeout- enabled versions of the blocking object APIs along with it. This support comes at a small cost to code size, but a slightly larger cost to realtime performance - as checking for the use of timers in the underlying internal code costs some cycles.

As a result, the option is given to the user here to manually disable these timeout-based APIs if desired by the user for performance and code-size reasons.

Set this option to 1 to enable timeout-based APIs for blocking calls.

KERNEL_USE_QUANTUM

Do you want to enable time quanta? This is useful when you want to have tasks in the same priority group share time in a controlled way. This allows equal tasks to use unequal amounts of the CPU, which is a great way to set up CPU budgets per thread in a round-robin scheduling system. If enabled, you can specify a number of ticks that serves as the default time period (quantum). Unless otherwise specified, every thread in a priority will get the default quantum.

Set this option to 1 to enable round-robin scheduling.

THREAD_QUANTUM_DEFAULT

This value defines the default thread quantum when `KERNEL_USE_QUANTUM` is enabled. The value defined is a time in milliseconds.

KERNEL_USE_SLEEP

This define enables the `Thread::Sleep()` API, which allows a thread to suspend its operation for a defined length of time, specified in ms.

3.3 Blocking Objects

KERNEL_USE_NOTIFY

This is a simple blocking object, where a thread (or threads) are guaranteed to block until an asynchronous event signals the object.

KERNEL_USE_SEMAPHORE

Do you want the ability to use counting/binary semaphores for thread synchronization? Enabling this feature provides fully-blocking semaphores and enables all API functions declared in `semaphore.h`. If you have to pick one blocking mechanism, this is the one to choose.

Note that all IPC mechanisms (mailboxes, messages) rely on semaphores, so keep in mind that this is a prerequisite for many other features in the kernel.

KERNEL_USE_MUTEX

Do you want the ability to use mutual exclusion semaphores (mutex) for resource/block protection? Enabling this feature provides mutexes, with priority inheritance, as declared in [mutex.h](#).

KERNEL_USE_EVENTFLAG

Provides additional event-flag based blocking. This relies on an additional per-thread flag-mask to be allocated, which adds 2 bytes to the size of each thread object.

KERNEL_USE_READERWRITER

Provides reader-writer locks. Allows current read access, or single write-access to a resource. Readers wait for the writer to release the lock, and writers wait for all readers to release the lock before acquiring the resource.

KERNEL_USE_CONDVAR

Provides condition variables. Allows a thread to wait for a specific condition to be true before proceeding, with a mutual-exclusion lock held.

3.4 Inter-process/thread Communication

KERNEL_USE_MESSAGE

Enable inter-thread messaging using message queues. This is the preferred mechanism for IPC for serious multi-threaded communications; generally anywhere a semaphore or event-flag is insufficient.

GLOBAL_MESSAGE_POOL_SIZE

If Messages are enabled, define the size of the default kernel message pool. Messages can be manually added to the message pool, but this mechanism is more convenient and automatic. All message queues can share their message objects from this global pool to maximize efficiency and simplify data management.

KERNEL_USE_MAILBOX

Enable inter-thread messaging using mailboxes. A mailbox manages a blob of data provided by the user, that is partitioned into fixed-size blocks called envelopes. The size of an envelope is set by the user when the mailbox is initialized. Any number of threads can read-from and write-to the mailbox. Envelopes can be sent-to or received-from the mailbox at the head or tail. In this way, mailboxes essentially act as a circular buffer that can be used as a blocking FIFO or LIFO queue.

3.5 Debug Features

KERNEL_USE_THREADNAME

Provide Thread method to allow the user to set a name for each thread in the system. Adds a const char* pointer to the size of the thread object.

KERNEL_USE_DEBUG

Provides extra logic for kernel debugging, and instruments the kernel with extra asserts, and kernel trace functionality.

KERNEL_ENABLE_LOGGING

Set this to 1 to enable very chatty kernel logging. Since most important things in the kernel emit logs, a large log-buffer and fast output are required in order to keep up. This is a pretty advanced power-user type feature, so it's disabled by default.

KERNEL_ENABLE_USER_LOGGING

This enables a set of logging macros similar to the kernel-logging macros; however, these can be enabled or disabled independently. This allows for user-code to benefit from the built-in kernel logging macros without having to account for the super-high-volume of logs generated by kernel code.

KERNEL_EXTRA_CHECKS

This option provides extra safety checks within the kernel APIs in order to minimize the potential for unsafe operations. This is especially helpful during development, and can help catch problems at development time, instead of in the field.

KERNEL_USE_STACK_GUARD

This feature, when enabled, tells the kernel to check whether any Thread's stack has been exhausted (or slack falls below a certain safety threshold) before executing each context switch. Enabling this is the most effective means to guard against stack corruption and stack overflow in the kernel, at the cost of increased context

3.6 Enhancements, Security, Miscellaneous

KERNEL_USE_DRIVER

Enabling device drivers provides a posix-like filesystem interface for peripheral device drivers.

KERNEL_USE_DYNAMIC_THREADS

Provide extra Thread methods to allow the application to create (and more importantly destroy) threads at runtime. useful for designs implementing worker threads, or threads that can be restarted after encountering error conditions.

KERNEL_USE_PROFILER

Provides extra classes for profiling the performance of code. useful for debugging and development, but uses an additional hardware timer.

KERNEL_USE_ATOMIC

Provides support for atomic operations, including addition, subtraction, set, and test-and-set. Add/Sub/Set contain 8, 16, and 32-bit variants.

SAFE_UNLINK

"Safe unlinking" performs extra checks on data to make sure that there are no inconsistencies when performing operations on linked lists. This goes beyond pointer checks, adding a layer of structural and metadata validation to help detect system corruption early.

KERNEL_AWARE_SIMULATION

Include support for kernel-aware simulation. Enabling this feature adds advanced profiling, trace, and environment-aware debugging and diagnostic functionality when Mark3-based applications are run on the fIAVR AVR simulator.

KERNEL_USE_IDLE_FUNC

Enabling this feature removes the necessity for the user to dedicate a complete thread for idle functionality. This saves a full thread stack, but also requires a bit extra static data. This also adds a slight overhead to the context switch and scheduler, as a special case has to be taken into account.

KERNEL_USE_AUTO_ALLOC

This feature enables an additional set of APIs that allow for objects to be created on-the-fly out of a special heap, without having to explicitly allocate them (from stack, heap, or static memory).

AUTO_ALLOC_SIZE

Size (in bytes) of the static pool of memory reserved from RAM for use by the auto allocator (if enabled).

KERNEL_USE_THREAD_CALLOUTS

This feature provides additional kernel APIs to register callout functions that are activated when threads are created or exited. This is useful for implementing low-level instrumentation based on information held in the threads.

KERNEL_USE_EXTENDED_CONTEXT

Allocate an extra pointer's worth of storage within a Thread object (and corresponding accessor methods) to provide the user with a means to implement arbitrary Thread-local storage.

Chapter 4

Building Mark3

4.1 Source Layout

One key aspect of [Mark3](#) is that system features are organized into their own separate modules. These modules are further grouped together into folders based on the type of features represented:

Root	Base folder, contains license info and build system configuration
arduino	Arduino-specific headers and API documentation files
bootloader	Mark3 Bootloader code for AVR microcontrollers
build	Device-specific toolchain configuraton files for various platforms
docs	Documentation (pdf + html)
drivers	Device driver code for various supported devices
example	Example applications
export	Platform specific output folder, used when running export.sh
fonts	Bitmap fonts converted from TTF, used by Mark3 graphics library
kbuild	Build output directory
kernel	Basic Mark3 Components (the focus of this manual)
cpu	CPU-specific porting code
scripts	Scripts used to simplify build, documentation, and profiling
libs	Utility code and services, extended system features
tests	Unit tests, written as C/C++ applications
util	Host utilities - including a ttf font converter and device programmer

4.2 Toolchain Integration

[Mark3](#) supports a variety of GCC ports out of the box - however, depending on your host OS and target processor, there may be some effort required to tie the toolchain into the build system.

After installing your toolchain of choice, you must make sure that the main toolchain binary paths are set in your systems PATH environment variable, ensuring that they are accessible directly from the command-line. Without this step, the build configuration step (cmake) will inevitably fail.

Depending on your toolchain, you may also be required to add toolchain-specific include directories to the build flags. These flags can be added to the cmake variables defined in `/build/<cpu>/<variant>/<toolchain>/platform.cmake` for your target architecture.

4.3 Installing Dependencies

The [Mark3](#) build system uses CMake (3.4.2 or above) for configuration management, and Ninja to execute the build steps. The combination of these two tools results in exceptionally fast builds - so fast that the previous makefile build system was scrapped in its favor.

These tools are readily available for most common host operating systems.

CMake can be found here: <https://cmake.org> Ninja can be found here: <https://ninja-build.org>

4.4 Building Mark3 Kernel and Libraries

Once a sane environment has been created, the kernel, libraries, examples and tests can be built by running `./scripts/build.sh` from the root directory. By default, [Mark3](#) builds for the `atmega328p` target, but the target can be selected by manually configuring the above environment variables, or by running the included `./scripts/set_target.sh` script as follows:

```
./scripts/set_target.sh <architecture> <variant> <toolchain>
```

Where:

```
<architecture> is the target CPU architecture(i.e. avr, msp430, cm0, cm3, cm4f)
<variant>      is the part name (i.e. atmega328p, msp430f2274, generic)
<toolchain>    is the build toolchain (i.e. gcc)
```

This script is a thin wrapper for the cmake configuration commands, and clears the `kbuild` output directory before re-initializing cmake for the selected target.

To build the [Mark3](#) kernel and middleware libraries for a generic ARM Cortex-M0 using a pre-configured `arm-none-eabi-gcc` toolchain, one would run the following commands:

```
./scripts/set_target.sh cm0 generic gcc
./scripts/build.sh
```

To perform an incremental build, go into the cmake build directory (`kbuild`) and simply run `'ninja'`.

Note that not all libraries/tests/examples will build in all kernel configurations. The default kernel configuration may need adjustment/tweaking to support a specific part. See `CMakeLists.txt` and [mark3cfg.h](#) respectively for more information

4.5 Exporting the kernel source

While the build system is flexible enough to adapt to any toolchain, it may be desirable to integrate the [Mark3](#) kernel and associated drivers/libraries into another build system.

[Mark3](#) provides a script (the aptly-named `export.sh`) which allow for the source for any supported port to be exported for this purpose. This script will also generate appropriate doxygen documentation, and package the whole of it together in a zip file. The files in the archive are placed in a "flat" hierarchy, and do not require any specific path structure to be maintained when imported into another build system.

As a special feature, if the "arduino" AVR target is specified, additional pre-processing is done on the source to turn the standard [Mark3](#) kernel into a library that can be imported directly into Arduino IDE. This is also how the official [Mark3](#) arduino-compatible releases are generated (hosted on [mark3os.com](#) and [sourceforge.net](#))

To exercise the build system, type the following from the main mark3 embedded source directory:

```
> ./scripts/export.sh <target>
```

Where:

Target is one of the following:

```
atmega328p
atmega644
atmega1280
atmega2560
atmega1284p
atxmega256a3
arduino
arduino2560
cortex_m0
cortex_m3
cortex_m4f
msp430f2274
```

If successful, the generated artifacts will be placed in an output folder under the `./export` directory.

Additionally, if doxygen is found on the host system's PATH, a copy of the manual (using the specific port's source code) will be generated and archived with the source release. If `pdflatex` is also found on the host's PATH, a PDF copy of the manual will be generated, tailored to the selected target.

Chapter 5

Getting Started With The Mark3 API

5.1 Kernel Setup

This section details the process of defining threads, initializing the kernel, and adding threads to the scheduler.

If you're at all familiar with real-time operating systems, then these setup and initialization steps should be familiar. I've tried very hard to ensure that as much of the heavy lifting is hidden from the user, so that only the bare minimum of calls are required to get things started.

The examples presented in this chapter are real, working examples taken from the ATmega328p port.

First, you'll need to create the necessary data structures and functions for the threads:

1. Create a Thread object for all of the "root" or "initial" tasks.
2. Allocate stacks for each of the Threads
3. Define an entry-point function for each Thread

This is shown in the example code below:

```
//-----  
#include "thread.h"  
#include "kernel.h"  
  
//1) Create a thread object for all of the "root" or "initial" tasks  
static Thread AppThread;  
static Thread IdleThread;  
  
//2) Allocate stacks for each thread  
#define STACK_SIZE_APP      (192)  
#define STACK_SIZE_IDLE     (128)  
  
static uint8_t aucAppStack[STACK_SIZE_APP];  
static uint8_t aucIdleStack[STACK_SIZE_IDLE];  
  
//3) Define entry point functions for each thread  
void AppThread(void);  
void IdleThread(void);
```

Next, we'll need to add the required kernel initialization code to main. This consists of running the Kernel's init routine, initializing all of the threads we defined, adding the threads to the scheduler, and finally calling `Kernel::Start()`, which transfers control of the system to the RTOS.

These steps are illustrated in the following example.

```

int main(void)
{
    //1) Initialize the kernel prior to use
    Kernel::Init();           // MUST be before other kernel ops

    //2) Initialize all of the threads we've defined
    AppThread.Init( aucAppStack,    // Pointer to the stack
                   STACK_SIZE_APP,  // Size of the stack
                   1,               // Thread priority
                   (void*)AppEntry, // Entry function
                   NULL );          // Entry function argument

    IdleThread.Init( aucIdleStack,   // Pointer to the stack
                   STACK_SIZE_IDLE,  // Size of the stack
                   0,               // Thread priority
                   (void*)IdleEntry, // Entry function
                   NULL );           // Entry function argument

    //3) Add the threads to the scheduler
    AppThread.Start();           // Actively schedule the threads
    IdleThread.Start();

    //4) Give control of the system to the kernel
    Kernel::Start();             // Start the kernel!
}

```

Not much to it, is there? There are a few noteworthy points in this code, though.

In order for the kernel to work properly, a system must always contain an idle thread; that is, a thread at priority level 0 that never blocks. This thread is responsible for performing any of the low-level power management on the CPU in order to maximize battery life in an embedded device. The idle thread must also never block, and it must never exit. Either of these operations will cause undefined behavior in the system.

The App thread is at a priority level greater-than 0. This ensures that as long as the App thread has something useful to do, it will be given control of the CPU. In this case, if the app thread blocks, control will be given back to the Idle thread, which will put the CPU into a power-saving mode until an interrupt occurs.

Stack sizes must be large enough to accommodate not only the requirements of the threads, but also the requirements of interrupts - up to the maximum interrupt-nesting level used. Stack overflows are super-easy to run into in an embedded system; if you encounter strange and unexplained behavior in your code, chances are good that one of your threads is blowing its stack.

5.2 Threads

[Mark3](#) Threads act as independent tasks in the system. While they share the same address-space, global data, device-drivers, and system peripherals, each thread has its own set of CPU registers and stack, collectively known as the thread's **context**. The context is what allows the RTOS kernel to rapidly switch between threads at a high rate, giving the illusion that multiple things are happening in a system, when really, only one thread is executing at a time.

5.2.1 Thread Setup

Each instance of the Thread class represents a thread, its stack, its CPU context, and all of the state and metadata maintained by the kernel. Before a Thread will be scheduled to run, it must first be initialized with the necessary configuration data.

The Init function gives the user the opportunity to set the stack, stack size, thread priority, entry-point function, entry-function argument, and round-robin time quantum:

Thread stacks are pointers to blobs of memory (usually char arrays) carved out of the system's address space. Each thread must have a stack defined that's large enough to handle not only the requirements of local variables in the thread's code path, but also the maximum depth of the ISR stack.

Priorities should be chosen carefully such that the shortest tasks with the most strict determinism requirements are executed first - and are thus located in the highest priorities. Tasks that take the longest to execute (and require the least degree of responsiveness) must occupy the lower thread priorities. The idle thread must be the only thread occupying the lowest priority level.

The thread quantum only applies when there are multiple threads in the ready queue at the same priority level. This interval is used to kick-off a timer that will cycle execution between the threads in the priority list so that they each get a fair chance to execute.

The entry function is the function that the kernel calls first when the thread instance is first started. Entry functions have at most one argument - a pointer to a data-object specified by the user during initialization.

An example thread initialization is shown below:

```
Thread clMyThread;
uint8_t aucStack[192];

void AppEntry(void)
{
    while(1)
    {
        // Do something!
    }
}

...
{
    clMyThread.Init(aucStack,    // Pointer to the stack to use by this thread
                   192,        // Size of the stack in bytes
                   1,          // Thread priority (0 = idle, 7 = max)
                   (void*)AppEntry, // Function where the thread starts executing
                   NULL );      // Argument passed into the entry function
}
```

Once a thread has been initialized, it can be added to the scheduler by calling:

```
clMyThread.Start();
```

The thread will be placed into the Scheduler's queue at the designated priority, where it will wait its turn for execution.

5.2.2 Entry Functions

Mark3 Threads should not run-to-completion - they should execute as infinite loops that perform a series of tasks, appropriately partitioned to provide the responsiveness characteristics desired in the system.

The most basic Thread loop is shown below:

```
void Thread( void *param )
{
    while(1)
    {
        // Do Something
    }
}
```

Threads can interact with each other in the system by means of synchronization objects (Semaphore), mutual-exclusion objects (Mutex), Inter-process messaging (MessageQueue), and timers (Timer).

Threads can suspend their own execution for a predetermined period of time by using the static Thread::Sleep() method. Calling this will block the Thread's execution until the amount of time specified has elapsed. Upon expiry, the thread will be placed back into the ready queue for its priority level, where it awaits its next turn to run.

5.3 Timers

Timer objects are used to trigger callback events periodic or on a one-shot (alarm) basis.

While extremely simple to use, they provide one of the most powerful execution contexts in the system. The timer callbacks execute from within the timer callback ISR in an interrupt-enabled context. As such, timer callbacks are considered higher-priority than any thread in the system, but lower priority than other interrupts. Care must be taken to ensure that timer callbacks execute as quickly as possible to minimize the impact of processing on the throughput of tasks in the system. Wherever possible, heavy-lifting should be deferred to the threads by way of semaphores or messages.

Below is an example showing how to start a periodic system timer which will trigger every second:

```
{
    Timer clTimer;
    clTimer.Init();

    clTimer.Start( 1000,
                  1,
                  MyCallback,
                  (void*)&my_data );

    ... // Keep doing work in the thread
}

// Callback function, executed from the timer-expiry context.
void MyCallBack( Thread *pclOwner_, void *pvData_ )
{
    LED.Flash(); // Flash an LED.
}
```

5.4 Semaphores

Semaphores are used to synchronized execution of threads based on the availability (and quantity) of application-specific resources in the system. They are extremely useful for solving producer-consumer problems, and are the method-of-choice for creating efficient, low latency systems, where ISRs post semaphores that are handled from within the context of individual threads. (Yes, Semaphores can be posted - but not pended - from the interrupt context).

The following is an example of the producer-consumer usage of a binary semaphore:

```
Semaphore clSemaphore; // Declare a semaphore shared between a producer and a consumer thread.

void Producer()
{
    clSemaphore.Init(0, 1);
    while(1)
    {
        // Do some work, create something to be consumed

        // Post a semaphore, allowing another thread to consume the data
        clSemaphore.Post();
    }
}

void Consumer()
{
    // Assumes semaphore initialized before use...
    While(1)
    {
        // Wait for new data from the producer thread
        clSemaphore.Pend();

        // Consume the data!
    }
}
```

And an example of using semaphores from the ISR context to perform event- driven processing.

```
Semaphore clSemaphore;

__interrupt__ MyISR()
{
    clSemaphore.Post(); // Post the interrupt. Lightweight when uncontested.
}

void MyThread()
{
    clSemaphore.Init(0, 1); // Ensure this is initialized before the MyISR interrupt is enabled.
    while(1)
    {
        // Wait until we get notification from the interrupt
        clSemaphore.Pend();

        // Interrupt has fired, do the necessary work in this thread's context
        HeavyLifting();
    }
}
```

5.5 Mutexes

Mutexes (Mutual exclusion objects) are provided as a means of creating "protected sections" around a particular resource, allowing for access of these objects to be serialized. Only one thread can hold the mutex at a time - other threads have to wait until the region is released by the owner thread before they can take their turn operating on the protected resource. Note that mutexes can only be owned by threads - they are not available to other contexts (i.e. interrupts). Calling the mutex APIs from an interrupt will cause catastrophic system failures.

Note that these objects are also not recursive- that is, the owner thread can not attempt to claim a mutex more than once.

Priority inheritance is provided with these objects as a means to avoid priority inversions. Whenever a thread at a priority than the mutex owner blocks on a mutex, the priority of the current thread is boosted to the highest-priority waiter to ensure that other tasks at intermediate priorities cannot artificially prevent progress from being made.

Mutex objects are very easy to use, as there are only three operations supported: Initialize, Claim and Release. An example is shown below.

```
Mutex clMutex; // Create a mutex globally.

void Init()
{
    // Initialize the mutex before use.
    clMutex.Init();
}

// Some function called from a thread
void Thread1Function()
{
    clMutex.Claim();

    // Once the mutex is owned, no other thread can
    // enter a block protect by the same mutex

    my_protected_resource.do_something();
    my_protected_resource.do_something_else();

    clMutex.Release();
}

// Some function called from another thread
void Thread2Function()
{
    clMutex.Claim();

    // Once the mutex is owned, no other thread can
    // enter a block protect by the same mutex

    my_protected_resource.do_something();
    my_protected_resource.do_different_things();

    clMutex.Release();
}
```

5.6 Event Flags

Event Flags are another synchronization object, conceptually similar to a semaphore.

Unlike a semaphore, however, the condition on which threads are unblocked is determined by a more complex set of rules. Each Event Flag object contains a 16-bit field, and threads block, waiting for combinations of bits within this field to become set.

A thread can wait on any pattern of bits from this field to be set, and any number of threads can wait on any number of different patterns. Threads can wait on a single bit, multiple bits, or bits from within a subset of bits within the field.

As a result, setting a single value in the flag can result in any number of threads becoming unblocked simultaneously. This mechanism is extremely powerful, allowing for all sorts of complex, yet efficient, thread synchronization schemes that can be created using a single shared object.

Note that Event Flags can be set from interrupts, but you cannot wait on an event flag from within an interrupt.

Examples demonstrating the use of event flags are shown below.

```
// Simple example showing a thread blocking on a multiple bits in the
// fields within an event flag.

EventFlag clEventFlag;

int main()
{
    ...
    clEventFlag.Init(); // Initialize event flag prior to use
    ...
}

void MyInterrupt()
{
    // Some interrupt corresponds to event 0x0020
    clEventFlag.Set(0x0020);
}

void MyThreadFunc()
{
    ...
    while(1)
    {
        ...
        uint16_t ul6WakeCondition;

        // Allow this thread to block on multiple flags
        ul6WakeCondition = clEventFlag.Wait(0x00FF, EventFlagOperation::Any_Set);

        // Clear the event condition that caused the thread to wake (in this case,
        // ul6WakeCondition will equal 0x20 when triggered from the interrupt above)
        clEventFlag.Clear(ul6WakeCondition);

        // <do something>
    }
}
```

5.7 Messages

Sending messages between threads is the key means of synchronizing access to data, and the primary mechanism to perform asynchronous data processing operations.

Sending a message consists of the following operations:

- Obtain a Message object from the global message pool

- Set the message data and event fields
- Send the message to the destination message queue

While receiving a message consists of the following steps:

- Wait for a messages in the destination message queue
- Process the message data
- Return the message back to the global message pool

These operations, and the various data objects involved are discussed in more detail in the following section.

5.7.1 Message Objects

Message objects are used to communicate arbitrary data between threads in a safe and synchronous way.

The message object consists of an event code field and a data field. The event code is used to provide context to the message object, while the data field (essentially a void * data pointer) is used to provide a payload of data corresponding to the particular event.

Access to these fields is marshalled by accessors - the transmitting thread uses the `SetData()` and `SetCode()` methods to seed the data, while the receiving thread uses the `GetData()` and `GetCode()` methods to retrieve it.

By providing the data as a void data pointer instead of a fixed-size message, we achieve an unprecedented measure of simplicity and flexibility. Data can be either statically or dynamically allocated, and sized appropriately for the event without having to format and reformat data by both sending and receiving threads. The choices here are left to the user - and the kernel doesn't get in the way of efficiency.

It is worth noting that you can send messages to message queues from within ISR context. This helps maintain consistency, since the same APIs can be used to provide event-driven programming facilities throughout the whole of the OS.

5.7.2 Global Message Pool

To maintain efficiency in the messaging system (and to prevent over-allocation of data), a global pool of message objects is provided. The size of this message pool is specified in the implementation, and can be adjusted depending on the requirements of the target application as a compile-time option.

Allocating a message from the message pool is as simple as calling the `GlobalMessagePool::Pop()` Method.

Messages are returned back to the `GlobalMessagePool::Push()` method once the message contents are no longer required.

One must be careful to ensure that discarded messages always are returned to the pool, otherwise a resource leak can occur, which may cripple the operating system's ability to pass data between threads.

5.7.3 Message Queues

Message objects specify data with context, but do not specify where the messages will be sent. For this purpose we have a `MessageQueue` object. Sending an object to a message queue involves calling the `MessageQueue::Send()` method, passing in a pointer to the `Message` object as an argument.

When a message is sent to the queue, the first thread blocked on the queue (as a result of calling the `MessageQueue::Receive()` method) will wake up, with a pointer to the `Message` object returned.

It's worth noting that multiple threads can block on the same message queue, providing a means for multiple threads to share work in parallel.

5.7.4 Messaging Example

```
// Message queue object shared between threads
MessageQueue cMsgQ;

// Function that initializes the shared message queue
void MsgQInit()
{
    cMsgQ.Init();
}

// Function called by one thread to send message data to
// another
void TxMessage()
{
    // Get a message, initialize its data
    Message *pclMesg = GlobalMessagePool::Pop();

    pclMesg->SetCode(0xAB);
    pclMesg->SetData((void*)some_data);

    // Send the data to the message queue
    cMsgQ.Send(pclMesg);
}

// Function called in the other thread to block until
// a message is received in the message queue.
void RxMessage()
{
    Message *pclMesg;

    // Block until we have a message in the queue
    pclMesg = cMsgQ.Receive();

    // Do something with the data once the message is received
    pclMesg->GetCode();

    // Free the message once we're done with it.
    GlobalMessagePool::Push(pclMesg);
}
```

5.8 Mailboxes

Another form of IPC is provided by [Mark3](#), in the form of Mailboxes and Envelopes.

Mailboxes are similar to message queues in that they provide a synchronized interface by which data can be transmitted between threads.

Where Message Queues rely on linked lists of lightweight message objects (containing only message code and a `void*` data-pointer), which are inherently abstract, Mailboxes use a dedicated blob of memory, which is carved up into fixed-size chunks called Envelopes (defined by the user), which are sent and received. Unlike message queues, mailbox data is copied to and from the mailboxes dedicated pool.

Mailboxes also differ in that they provide not only a blocking "receive" call, but also a blocking "send" call, providing the opportunity for threads to block on "mailbox full" as well as "mailbox empty" conditions.

All send/receive APIs support an optional timeout parameter if the `KERNEL_USE_TIMEOUTS` option has been configured in [mark3cfg.h](#)

5.8.1 Mailbox Example

```
// Create a mailbox object, and define a buffer that will be used to store the
// mailbox' envelopes.
static Mailbox clMbox;
static uint8_t aucMBoxBuffer[128];

...
void InitMailbox(void)
{
    // Initialize our mailbox, telling it to use our defined buffer for envelope
    // storage. Pass in the size of the buffer, and set the size of each
    // envelope to 16 bytes. This gives u16 a mailbox capacity of (128 / 16) = 8
    // envelopes.
    clMbox.Init((void*)aucMBoxBuffer, 128, 16);
}

...
void SendThread(void)
{
    // Define a buffer that we'll eventually send to the
    // mailbox. Note the size is the same as that of an
    // envelope.
    uint8_t aucTxBuf[16];

    while(1)
    {
        // Copy some data into aucTxBuf, a 16-byte buffer, the
        // same size as a mailbox envelope.
        ...

        // Deliver the envelope (our buffer) into the mailbox
        clMbox.Send((void*)aucTxBuf);
    }
}

...
void RecvThred(void)
{
    uint8_t aucRxBuf[16];

    while(1)
    {
        // Wait until there's a message in our mailbox. Once
        // there is a message, read it into our local buffer.
        cmMbox.Receive((void*)aucRxBuf);

        // Do something with the contents of aucRxBuf, which now
        // contains an envelope of data read from the mailbox.
        ...
    }
}
```

5.9 Notification Objects

Notification objects are the most lightweight of all blocking objects supplied by [Mark3](#).

using this blocking primitive, one or more threads wait for the notification object to be signalled by code elsewhere in the system (i.e. another thread or interrupt). Once the the notification has been signalled, all threads currently blocked on the object become unblocked.

5.9.1 Notification Example

```
static Notify clNotifier;

...
void MyThread(void *unused_)
{
    // Initialize our notification object before use
    clNotifier.Init();

    while (1)
```

```

    {
        // Wait until our thread has been notified that it
        // can wake up.
        clNotify.Wait();

        ...
        // Thread has woken up now -- do something!
    }
}

...
void SignalCallback(void)
{
    // Something in the system (interrupt, thread event, IPC,
    // etc.,) has called this function. As a result, we need
    // our other thread to wake up. Call the Notify object's
    // Signal() method to wake the thread up. Note that this
    // will have no effect if the thread is not presently
    // blocked.

    clNotify.Signal();
}

```

5.10 Condition Variables

Condition Variables, implemented in [Mark3](#) with the `ConditionVariable` class, provide an implementation of the classic Monitor pattern. This object allows a thread to wait for a specific condition to occur, claiming a shared lock once the condition is met. Threads may also choose to signal a single blocking thread to indicate a condition has changed, or broadcast condition changes to all waiting threads.

5.10.1 Condition Variable Example

```

// Define a condition variable object, a shared lock, and
// a piece of common data shared between threads to represent
// a condition.

// Assume Mutex and ConditionVariable are initialized
// prior to use.
static ConditionVariable clCondVar;
static Mutex clSharedLock;
static volatile int iCondition = 0;

...

void CondThread1(void *unused_)
{
    while (1)
    {
        // Wait until
        clCondVar.Wait(&clSharedLock);

        // Only act on a specific condition
        if (iCondition == 1337) {
            // Do something
        }

        clSharedLock.Release();
    }
}

void CondThread2(void *unused_)
{
    // Assume Mutex and ConditionVariable are initialized
    // prior to use.
    while (1)
    {
        // Wait until
        clCondVar.Wait(&clSharedLock);

        // Act on a *different* condition than the other thread
        if (iCondition == 5454) {
            // Do something
        }

        clSharedLock.Release();
    }
}

```

```

    }
}

void SignalThread(void* unused)
{
    while (1) {
        // Sleep for a while
        Thread::Sleep(100);

        // Update the condition in a thread-safe manner
        clSharedLock.Claim();
        iCondition = 1337;
        clSharedLock.Release();

        // Wake one thread to check for the updated condition
        clCondVar.Signal();

        // Sleep for a while
        Thread::Sleep(100);

        // Update the condition in a thread-safe manner
        clSharedLock.Claim();
        iCondition = 1337;
        clSharedLock.Release();

        // Wake all threads to check for the updated condition
        clCondVar.Broadcast();
    }
}

```

5.11 Reader-Write Locks

Reader-Writer locks are provided in [Mark3](#) to provide an efficient way for multiple threads to share concurrent, non-destructive access to a resource, while preventing concurrent destructive/non-destructive accesses. A single "writer" may hold the lock, or 1-or-more "readers" may hold the lock. In the case that readers hold the lock, writers will block until all readers have relinquished their access to the resource. In the case that a writer holds the lock, all other readers and writers must wait until the lock is relinquished.

5.11.1 Reader-Write Lock Example

```

void WriterTask(void* param)
{
    auto pclRWLock = static_cast<ReaderWriterLock*>(param);

    pclRWLock->AcquireWriter();
    // All other readers and writers will have to wait until
    // the lock is released.
    iNumWrites++;
    ...
    pclRWLock->ReleaseWriter();
}

void ReaderTask(void* param)
{
    auto pclRWLock = static_cast<ReaderWriterLock*>(param);

    pclRWLock->AcquireReader();
    // Any number of reader threads can also acquire the lock
    // without having to block, waiting for this task to release it.
    // Writers must block until all readers have released their references
    // to the lock.
    iNumReads++;
    ...
    pclRWLock->ReleaseReader();
}

```

5.12 Sleep

There are instances where it may be necessary for a thread to poll a resource, or wait a specific amount of time before proceeding to operate on a peripheral or volatile piece of data.

While the Timer object is generally a better choice for performing time-sensitive operations (and certainly a better choice for periodic operations), the Thread::Sleep() method provides a convenient (and efficient) mechanism that allows for a thread to suspend its execution for a specified interval.

Note that when a thread is sleeping it is blocked, during which other threads can operate, or the system can enter its idle state.

```
int GetPeripheralData();
{
    int value;
    // The hardware manual for a peripheral specifies that
    // the "foo()" method will result in data being generated
    // that can be captured using the "bar()" method.
    // However, the value only becomes valid after 10ms

    peripheral.foo();
    Thread::Sleep(10); // Wait 10ms for data to become valid
    value = peripheral.bar();
    return value;
}
```

5.13 Round-Robin Quantum

Threads at the same thread priority are scheduled using a round-robin scheme. Each thread is given a timeslice (which can be configured) of which it shares time amongst ready threads in the group. Once a thread's timeslice has expired, the next thread in the priority group is chosen to run until its quantum has expired - the cycle continues over and over so long as each thread has work to be done.

By default, the round-robin interval is set at 4ms.

This value can be overridden by calling the thread's SetQuantum() with a new interval specified in milliseconds.

Chapter 6

Why Mark3?

My first job after graduating from university in 2005 was with a small company that had a very old-school, low-budget philosophy when it came to software development.

Every make-or-buy decision ended with "make" when it came to tools. It was the kind of environment where vendors cost us money, but manpower was free. In retrospect, we didn't have a ton of business during the time that I worked there, and that may have had something to do with the fact that we were constantly short on ready cash for things we could code ourselves.

Early on, I asked why we didn't use industry-standard tools - like JTAG debuggers or IDEs. One senior engineer scoffed that debuggers were tools for wimps - and something that a good programmer should be able to do without. After all - we had serial ports, GPIOs, and a bi-color LED on our boards. Since these were built into the hardware, they didn't cost us a thing. We also had a single software "build" server that took 5 minutes to build a 32k binary on its best days, so when we had to debug code, it was a painful process of trial and error, with lots of Youtube between iterations. We complained that tens of thousands of dollars of productivity was being flushed away that could have been solved by implementing a proper build server - and while we eventually got our wish, it took far more time than it should have.

Needless to say, software development was painful at that company. We made life hard on ourselves purely out of pride, and for the right to say that we walked "up-hills both ways through 3 feet of snow, everyday". Our code was tied ever-so-tightly to our hardware platform, and the system code was indistinguishable from the application. While we didn't use an RTOS, we had effectively implemented a 3-priority threading scheme using a carefully designed interrupt nesting scheme with event flags and a while(1) superloop running as a background thread. Nothing was abstracted, and the code was always optimized for the platform, presumably in an effort to save on code size and wasted cycles. I asked why we didn't use an RTOS in any of our systems and received dismissive scoffs - the overhead from thread switching and maintaining multiple threads could not be tolerated in our systems according to our chief engineers. In retrospect, our ad-hoc system was likely as large as my smallest kernel, and had just as much context switching (although it was hidden by the compiler).

And every time a new iteration of our product was developed, the firmware took far too long to bring up, because the algorithms and data structures had to be re-tooled to work with the peripherals and sensors attached to the new boards. We worked very hard in an attempt to reinvent the wheel, all in the name of producing "efficient" code.

Regardless, I learned a lot about embedded software development.

Most important, I learned that good design is the key to good software; and good design doesn't have to come at a price. In all but the smallest of projects, the well-designed, well-abstracted code is not only more portable, but it's usually smaller, easier to read, and easier to reuse.

Also, since we had all the time in the world to invest in developing our own tools, I gained a lot of experience building them, and making use of good, free PC tools that could be used to develop and debug a large portion of our code. I ended up writing PC-based device and peripheral simulators, state-machine frameworks, and abstractions for our

horrible ad-hoc system code. At the end of the day, I had developed enough tools that I could solve a lot of our development problems without having to re-inventing the wheel at each turn. Gaining a background in how these tools worked gave me a better understanding of how to use them - making me more productive at the jobs that I've had since.

I am convinced that designing good software takes honest effort up-front, and that good application code cannot be written unless it is based on a solid framework. Just as the wise man builds his house on rocks, and not on sand, wise developers write applications based on a well-defined platforms. And while you can probably build a house using nothing but a hammer and sheer will, you can certainly build one a lot faster with all the right tools.

This conviction lead me to development my first RTOS kernel in 2009 - FunkOS. It is a small, yet surprisingly full-featured kernel. It has all the basics (semaphores, mutexes, round-robin and preemptive scheduling), and some pretty advanced features as well (device drivers and other middleware). However, it had two major problems - it doesn't scale well, and it doesn't support many devices.

While I had modest success with this kernel (it has been featured on some blogs, and still gets around 125 downloads a month), it was nothing like the success of other RTOS kernels like uC/OS-II and FreeRTOS. To be honest, as a one-man show, I just don't have the resources to support all of the devices, toolchains, and evaluation boards that a real vendor can. I had never expected my kernel to compete with the likes of them, and I don't expect [Mark3](#) to change the embedded landscape either.

My main goal with [Mark3](#) was to solve the technical shortfalls in the FunkOS kernel by applying my experience in kernel development. As a result, [Mark3](#) is better than FunkOS in almost every way; it scales better, has lower interrupt latency, and is generally more thoughtfully designed (all at a small cost to code size).

Another goal I had was to create something easy to understand, that could be documented and serve as a good introduction to RTOS kernel design. The end result of these goals is the kernel as presented in this book - a full source listing of a working OS kernel, with each module completely documented and explained in detail.

Finally, I wanted to prove that a kernel written entirely in C++ could perform just as well as one written in C. [Mark3](#) is fully benchmarked and profiled – you can see exactly how much it costs to call certain APIs or include various features in the kernel.

And in addition, the code is more readable and easier to understand as a result of making use of object-oriented concepts provided by C++. Applications are easier to write because common concepts are encapsulated into objects (Threads, Semaphores, Mutexes, etc.) with their own methods and data, as opposed to APIs which rely on lots of explicit pointer or handle-passing, type casting, and other operations that are typically considered "unsafe" or "advanced" topics in C.

Chapter 7

When should you use an RTOS?

7.1 The reality of system code

System code can be defined as the program logic required to manage, synchronize, and schedule all of the resources (CPU time, memory, peripherals, etc.) used by the application running on the CPU. And it's true that a significant portion of the code running on an embedded system will be system code. No matter how simple a system is, whether or not this logic is embedded directly into the application (bare-metal system), or included as part of a well-defined stack on which an application is written (RTOS-based); system code is still present, and it comes with a cost.

As an embedded systems is being designed, engineers have to decide which approach to take: Bare-metal, or RTOS. There are advantages and disadvantages to each – and a reasonable engineer should always perform a thorough analysis of the pros and cons of each - in the context of the given application - before choosing a path.

The following figure demonstrates the differences between the architecture of a bare-metal system and RTOS based system at a high level:

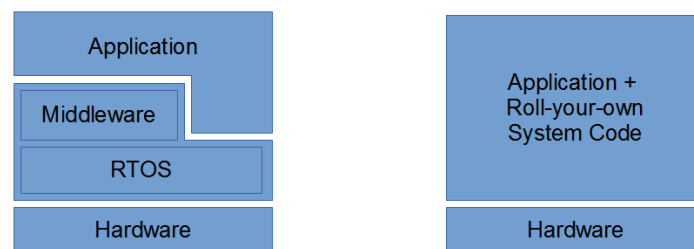


Figure 7.1 Arch

As can be seen, the RTOS (And associated middleware + libraries) captures a certain fixed size.

As a generalization, bare-metal systems typically have the advantage in that the system code overhead is small to start – but grows significantly as the application grows in complexity. At a certain point, it becomes extremely difficult and error-prone to add more functionality to an application running on such a system. There's a tipping point, where the cost of the code used to work-around the limitations of a bare-metal system outweigh the cost of a capable RTOS. Bare-metal systems also generally take longer to implement, because the system code has to be written from scratch (or derived from existing code) for the application. The resulting code also tends to be less portable, as it takes serious discipline to keep the system-specific elements of the code separated – in an RTOS-based system, once the kernel and drivers are ported, the application code is generally platform agnostic.

Conversely, an RTOS-based system incurs a slightly higher fixed cost up-front, but scales infinitely better than a bare-metal system as application's complexity increases. Using an RTOS for simple systems reduces application development time, but may cause an application not to fit into some extremely size-constrained microcontroller. An RTOS can also cause the size of an application to grow more slowly relative to a bare-metal system – especially as a result of applying synchronization mechanisms and judicious IPC. As a result, an RTOS makes it significantly easier to “go agile” with an application – iteratively adding features and functionality, without having to consider refactoring the underlying system at each turn.

Some of these factors may be more important than others. Requirements, specifications, schedules, chip-selection, and volume projections for a project should all be used to feed into the discussions to decide whether or to go bare-metal or RTOS as a result.

Consider the following questions when making that decision:

- What is the application?
- How efficient is efficient enough?
- How fast is fast enough?
- How small is small enough?
- How responsive is responsive enough?
- How much code space/RAM/etc is available on the target system?
- How much code space/RAM do I need for an RTOS?
- How much code space/RAM do I think I'll need for my application?
- How much time do I have to deliver my system?
- How many units do we plan to sell?

7.2 Superloops, and their limitations

7.2.1 Intro to Superloops

Before we start taking a look at designing a real-time operating system, it's worthwhile taking a look through one of the most-common design patterns that developers use to manage task execution in bare-metal embedded systems - Superloops.

Systems based on superloops favor the system control logic baked directly into the application code, usually under the guise of simplicity, or memory (code and RAM) efficiency. For simple systems, superloops can definitely get the job done. However, they have some serious limitations, and are not suitable for every kind of project. In a lot of cases you can squeak by using superloops - especially in extremely constrained systems, but in general they are not a solid basis for reusable, portable code.

Nonetheless, a variety of examples are presented here- from the extremely simple, to cooperative and limited-preemptive multitasking systems, all of which are examples are representative of real-world systems that I've either written the firmware for, or have seen in my experience.

7.2.2 The simplest loop

Let's start with the simplest embedded system design possible - an infinite loop that performs a single task repeatedly:

```
int main()
{
    while(1)
    {
        Do_Something();
    }
}
```

Here, the code inside the loop will run a single function forever and ever. Not much to it, is there? But you might be surprised at just how much embedded system firmware is implemented using essentially the same mechanism - there isn't anything wrong with that, but it's just not that interesting.

Despite its simplicity we can see the beginnings of some core OS concepts. Here, the `while(1)` statement can be logically seen as the operating system kernel - this one control statement determines what tasks can run in the system, and defines the constraints that could modify their execution. But at the end of the day, that's a big part of what a kernel is - a mechanism that controls the execution of application code.

The second concept here is the task. This is application code provided by the user to perform some useful purpose in a system. In this case `Do_something()` represents that task - it could be monitoring blood pressure, reading a sensor and writing its data to a terminal, or playing an MP3; anything you can think of for an embedded system to do. A simple round-robin multi-tasking system can be built off of this example by simply adding additional tasks in sequence in the main while-loop. Note that in this example the CPU is always busy running tasks - at no time is the CPU idle, meaning that it is likely burning a lot of power.

While we conceptually have two separate pieces of code involved here (an operating system kernel and a set of running tasks), they are not logically separate. The OS code is indistinguishable from the application. It's like a single-celled organism - everything is crammed together within the walls of an indivisible unit; and specialized to perform its given function relying solely on instinct.

7.2.3 Interrupt-Driven Super-loop

In the previous example, we had a system without any way to control the execution of the task- it just runs forever. There's no way to control when the task can (or more importantly can't) run, which greatly limits the usefulness of the system. Say you only want your task to run every 100 milliseconds - in the previous code, you have to add a hard-coded delay at the end of your task's execution to ensure your code runs only when it should.

Fortunately, there is a much more elegant way to do this. In this example, we introduce the concept of the synchronization object. A Synchronization object is some data structure which works within the bounds of the operating system to tell tasks when they can run, and in many cases includes special data unique to the synchronization event.

There are a whole family of synchronization objects, which we'll get into later. In this example, we make use of the simplest synchronization primitive

- the global flag.

With the addition of synchronization brings the addition of event-driven systems. If you're programming a microcontroller system, you generally have scores of peripherals available to you - timers, GPIOs, ADCs, UARTs, ethernet, USB, etc. All of which can be configured to provide a stimulus to your system by means of interrupts. This stimulus gives us the ability not only to program our micros to `do_something()`, but to `do_something()` if-and-only-if a corresponding trigger has occurred.

The following concepts are shown in the example below:

```
volatile K_BOOL something_to_do = false;

__interrupt__ My_Interrupt_Source(void)
{
    something_to_do = true;
}

int main()
{
    while (1)
    {
        if (something_to_do)
        {
            Do_something();
            something_to_do = false;
        }
        else
        {
            Idle();
        }
    }
}
```

So there you have it - an event driven system which uses a global variable to synchronize the execution of our task based on the occurrence of an interrupt. It's still just a bare-metal, OS-baked-into-the-application system, but it's introduced a whole bunch of added complexity (and control!) into the system.

The first thing to notice in the source is that the global variable, `something_to_do`, is used as a synchronization object. When an interrupt occurs from some external event, triggering the `My_Interrupt_Source()` ISR, program flow in `main()` is interrupted, the interrupt handler is run, and `something_to_do` is set to true, letting us know that when we get back to `main()`, that we should run our `Do_something()` task.

Another new concept at play here is that of the idle function. In general, when running an event driven system, there are times when the CPU has no application tasks to run. In order to minimize power consumption, CPUs usually contain instructions or registers that can be set up to disable non-essential subsets of the system when there's nothing to do. In general, the sleeping system can be re-activated quickly as a result of an interrupt or other external stimulus, allowing normal processing to resume.

Now, we could just call `Do_something()` from the interrupt itself - but that's generally not a great solution. In general, the more time we spend inside an interrupt, the more time we spend with at least some interrupts disabled. As a result, we end up with interrupt latency. Now, in this system, with only one interrupt source and only one task this might not be a big deal, but say that `Do_something()` takes several seconds to complete, and in that time several other interrupts occur from other sources. While executing in our long-running interrupt, no other interrupts can be processed - in many cases, if two interrupts of the same type occur before the first is processed, one of these interrupt events will be lost. This can be utterly disastrous in a real-time system and should be avoided at all costs. As a result, it's generally preferable to use synchronization objects whenever possible to defer processing outside of the ISR.

Another OS concept that is implicitly introduced in this example is that of task priority. When an interrupt occurs, the normal execution of code in `main()` is preempted: control is swapped over to the ISR (which runs to completion), and then control is given back to `main()` where it left off. The very fact that interrupts take precedence over what's running shows that `main` is conceptually a "low-priority" task, and that all ISRs are "high-priority" tasks. In this example, our "high-priority" task is setting a variable to tell our "low-priority" task that it can do something useful. We will investigate the concept of task priority further in the next example.

Preemption is another key principle in embedded systems. This is the notion that whatever the CPU is doing when an interrupt occurs, it should stop, cache its current state (referred to as its context), and allow the high-priority event to be processed. The context of the previous task is then restored its state before the interrupt, and resumes processing. We'll come back to preemption frequently, since the concept comes up frequently in RTOS-based systems.

7.2.4 Cooperative multi-tasking

Our next example takes the previous example one step further by introducing cooperative multi-tasking:

```
// Bitfield values used to represent three distinct tasks
#define TASK_1_EVENT (0x01)
#define TASK_2_EVENT (0x02)
#define TASK_3_EVENT (0x04)

volatile K_UCHAR event_flags = 0;

// Interrupt sources used to trigger event execution

__interrupt__ My_Interrupt_1(void)
{
    event_flags |= TASK_1_EVENT;
}

__interrupt__ My_Interrupt_2(void)
{
    event_flags |= TASK_2_EVENT;
}

__interrupt__ My_Interrupt_3(void)
{
    event_flags |= TASK_3_EVENT;
}

// Main tasks
int main(void)
{
    while(1)
    {
        while(event_flags)
        {
            if( event_flags & TASK_1_EVENT)
            {
                Do_Task_1();
                event_flags &= ~TASK_1_EVENT;
            } else if( event_flags & TASK_2_EVENT) {
                Do_Task_2();
                event_flags &= ~TASK_2_EVENT;
            } else if( event_flags & TASK_3_EVENT) {
                Do_Task_3();
                event_flags &= ~TASK_3_EVENT;
            }
        }
        Idle();
    }
}
```

This system is very similar to what we had before - however the differences are worth discussing. First, we have stimulus from multiple interrupt sources: each ISR is responsible for setting a single bit in our global event flag, which is then used to control execution of individual tasks from within main().

Next, we can see that tasks are explicitly given priorities inside the main loop based on the logic of the if/else if structure. As long as there is something set in the event flag, we will always try to execute Task1 first, and only when Task1 isn't set will we attempt to execute Task2, and then Task3. This added logic provides the notion of priority. However, because each of these tasks exist within the same context (they're just different functions called from our main control loop), we don't have the same notion of preemption that we have when dealing with interrupts.

That means that even through we may be running Task2 and an event flag for Task1 is set by an interrupt, the CPU still has to finish processing Task2 to completion before Task1 can be run. And that's why this kind of scheduling is referred to as cooperative multitasking: we can have as many tasks as we want, but unless they cooperate by means of returning back to main, the system can end up with high-priority tasks getting starved for CPU time by lower-priority, long-running tasks.

This is one of the more popular Os-baked-into-the-application approaches, and is widely used in a variety of real-time embedded systems.

7.2.5 Hybrid cooperative/preemptive multi-tasking

The final variation on the superloop design utilizes software-triggered interrupts to simulate a hybrid cooperative/preemptive multitasking system. Consider the example code below.

```
// Bitfields used to represent high-priority tasks. Tasks in this group
// can preempt tasks in the group below - but not eachother.
#define HP_TASK_1(0x01)
#define HP_TASK_2(0x02)

volatile K_UCHAR hp_tasks = 0;

// Bitfields used to represent low-priority tasks.
#define LP_TASK_1(0x01)
#define LP_TASK_2(0x02)

volatile K_UCHAR lp_tasks = 0;

// Interrupt sources, used to trigger both high and low priority tasks.
__interrupt__ System_Interrupt_1(void)
{
    // Set any of the other tasks from here...
    hp_tasks |= HP_TASK_1;
    // Trigger the SWI that calls the High_Priority_Tasks interrupt handler
    SWI();
}

__interrupt__ System_Interrupt_n...(void)
{
    // Set any of the other tasks from here...
}

// Interrupt handler that is used to implement the high-priority event context
__interrupt__ High_Priority_Tasks(void)
{
    // Enabled every interrupt except this one
    Disable_My_Interrupt();
    Enable_Interrupts();
    while( hp_tasks)
    {
        if( hp_tasks & HP_TASK_1)
        {
            HP_Task1();
            hp_tasks &= ~HP_TASK_1;
        }
        else if (hp_tasks & HP_TASK_2)
        {
            HP_Task2();
            hp_tasks &= ~HP_TASK_2;
        }
    }
    Restore_Interrupts();
    Enable_My_Interrupt();
}

// Main loop, used to implement the low-priority events
int main(void)
{
    // Set the function to run when a SWI is triggered
    Set_SWI(High_Priority_Tasks);

    // Run our super-loop
    while(1)
    {
        while (lp_tasks)
        {
            if (lp_tasks & LP_TASK_1)
            {
                LP_Task1();
                lp_tasks &= ~LP_TASK_1;
            }
            else if (lp_tasks & LP_TASK_2)
            {
                LP_Task2();
                lp_tasks &= ~LP_TASK_2;
            }
        }
        Idle();
    }
}
```

In this example, `High_Priority_Tasks()` can be triggered at any time as a result of a software interrupt (SWI). When a high-priority event is set, the code that sets the event calls the SWI as well, which instantly preempts whatever is happening in `main`, switching to the high-priority interrupt handler. If the CPU is executing in an interrupt handler already, the current ISR completes, at which point control is given to the high priority interrupt handler.

Once inside the HP ISR, all interrupts (except the software interrupt) are re-enabled, which allows this interrupt to be preempted by other interrupt sources, which is called interrupt nesting. As a result, we end up with two distinct execution contexts (`main` and `HighPriorityTasks()`), in which all tasks in the high-priority group are guaranteed to preempt `main()` tasks, and will run to completion before returning control back to tasks in `main()`. This is a very basic preemptive multitasking scenario, approximating a "real" RTOS system with two threads of different priorities.

7.3 Problems with superloops

As mentioned earlier, a lot of real-world systems are implemented using a superloop design; and while they are simple to understand due to the limited and obvious control logic involved, they are not without their problems.

7.3.1 Hidden Costs

It's difficult to calculate the overhead of the superloop and the code required to implement workarounds for blocking calls, scheduling, and preemption. There's a cost in both the logic used to implement workarounds (usually involving state machines), as well as a cost to maintainability that comes with breaking up into chunks based on execution time instead of logical operations. In moderate firmware systems, this size cost can exceed the overhead of a reasonably well-featured RTOS, and the deficit in maintainability is something that is measurable in terms of lost productivity through debugging and profiling.

7.3.2 Tightly-coupled code

Because the control logic is integrated so closely with the application logic, a lot of care must be taken not to compromise the separation between application and system code. The timing loops, state machines, and architecture-specific control mechanisms used to avoid (or simulate) preemption can all contribute to the problem. As a result, a lot of superloop code ends up being difficult to port without effectively simulating or replicating the underlying system for which the application was written. Abstraction layers can mitigate the risks, but a lot of care should be taken to fully decouple the application code from the system code.

7.3.3 No blocking Calls

In a super-loop environment, there's no such thing as a blocking call or blocking objects. Tasks cannot stop mid-execution for event-driven I/O from other contexts - they must always run to completion. If busy-waiting and polling are used as a substitute, it increases latency and wastes cycles. As a result, extra code complexity is often times necessary to work-around this lack of blocking objects, often times through implementing additional state machines. In a large enough system, the added overhead in code size and cycles can add up.

7.3.4 Difficult to guarantee responsiveness

Without multiple levels of priority, it may be difficult to guarantee a certain degree of real-time responsiveness without added profiling and tweaking. The latency of a given task in a priority-based cooperative multitasking system is the length of the longest task. Care must be taken to break tasks up into appropriate sized chunks in order to ensure that higher- priority tasks can run in a timely fashion - a manual process that must be repeated as new tasks are added in the system. Once again, this adds extra complexity that makes code larger, more difficult to understand and maintain due to the artificial subdivision of tasks into time-based components.

7.3.5 Limited preemption capability

As shown in the example code, the way to gain preemption in a superloop is through the use of nested interrupts. While this isn't unwieldy for two levels of priority, adding more levels beyond this is becomes complicated. In this case, it becomes necessary to track interrupt nesting manually, and separate sets of tasks that can run within given priority loops - and deadlock becomes more difficult to avoid.

Chapter 8

Can you afford an RTOS?

8.1 Intro

Of course, since you're reading the manual for an RTOS that I've been developing over the course of the several years, you can guess that the conclusion that I draw.

If your code is of any sort of non-trivial complexity (say, at least a few- thousand lines), then a more appropriate question would be "can you afford not* to use an RTOS in your system?".

In short, there are simply too many benefits of an RTOS to ignore, the most important being:

Threading, along with priority and time-based scheduling
Sophisticated synchronization objects and IPC
Flexible, powerful Software Timers
Ability to write more portable, decoupled code

Sure, these features have a cost in code space and RAM, but from my experience the cost of trying to code around a lack of these features will cost you as much - if not more. The results are often far less maintainable, error prone, and complex. And that simply adds time and cost. Real developers ship, and the RTOS is quickly becoming one of the standard tools that help keep developers shipping.

One of the main arguments against using an RTOS in an embedded project is that the overhead incurred is too great to be justified. Concerns over "wasted" RAM caused by using multiple stacks, added CPU utilization, and the "large" code footprint from the kernel cause a large number of developers to shun using a preemptive RTOS, instead favoring a non-preemptive, application-specific solution.

I believe that not only is the impact negligible in most cases, but that the benefits of writing an application with an RTOS can lead to savings around the board (code size, quality, reliability, and development time). While these other benefits provide the most compelling case for using an RTOS, they are far more challenging to demonstrate in a quantitative way, and are clearly documented in numerous industry-based case studies.

While there is some overhead associated with an RTOS, the typical arguments are largely unfounded when an RTOS is correctly implemented in a system. By measuring the true overhead of a preemptive RTOS in a typical application, we will demonstrate that the impact to code space, RAM, and CPU usage is minimal, and indeed acceptable for a wide range of CPU targets.

To illustrate just how little an RTOS impacts the size of an embedded software design we will look at a typical microcontroller project and analyze the various types of overhead associated with using a pre-emptive realtime kernel versus a similar non-preemptive event-based framework.

RTOS overhead can be broken into three distinct areas:

- Code space: The amount of code space eaten up by the kernel (static)
- Memory overhead: The RAM associated with running the kernel and application threads.
- Runtime overhead: The CPU cycles required for the kernel's functionality (primarily scheduling and thread switching)

While there are other notable reasons to include or avoid the use of an RTOS in certain applications (determinism, responsiveness, and interrupt latency among others), these are not considered in this discussion - as they are difficult to consider for the scope of our "canned" application.

8.2 Application description

For the purpose of this comparison, we first create an application using the standard preemptive [Mark3](#) kernel with 2 system threads running: A foreground thread and a background thread. This gives three total priority levels in the system - the interrupt level (high), and two application priority threads (medium and low), which is quite a common paradigm for microcontroller firmware designs. The foreground thread processes a variety of time-critical events at a fixed frequency, while the background thread processes lower priority, aperiodic events. When there are no background thread events to process, the processor enters its low-power mode until the next interrupt is acknowledged.

The contents of the threads themselves are unimportant for this comparison, but we can assume they perform a variety of realtime I/O functions. As a result, a number of device drivers are also implemented.

Code Space and Memory Overhead:

The application is compiled for an ATmega328p processor which contains 32kB of code space in flash, and 2kB of RAM, which is a lower-mid-range microcontroller in Atmel's 8-bit AVR line of microcontrollers. Using the AVR GCC compiler with -Os level optimizations, an executable is produced with the following code/RAM utilization:

```
Program: 27914 bytes
Data:    1313 bytes
```

An alternate version of this project is created using a custom "super-loop" kernel, which uses a single application thread and provides 2 levels of priority (interrupt and application). In this case, the event handler processes the different priority application events to completion from highest to lowest priority.

This approach leaves the application itself largely unchanged. Using the same optimization levels as the preemptive kernel, the code compiles as follows:

```
Program: 24886 bytes
Data:    750 bytes
```

At first glance, the difference in RAM utilization seems quite a lot higher for the preemptive mode version of the application, but the raw numbers don't tell the whole story.

The first issue is that the cooperative-mode total does not take into account the system stack - whereas these values are included in the totals for RTOS version of the project. As a result, some further analysis is required to determine how the stack sizes truly compare.

In cooperative mode, there is only one thread of execution - so considering that multiple event handlers are executed in turn, the stack requirements for cooperative mode is simply determined by those of the most stack-intensive event handler (ignoring stack use contributions due to interrupts).

In contrast, the preemptive kernel requires a separate stack for each active thread, and as a result the stack usage of the system is the sum of the stacks for all threads.

Since the application and idle events are the same for both preemptive and cooperative mode, we know that their (independent) stack requirements will be the same in both cases.

For cooperative mode, we see that the idle thread stack utilization is lower than that of the application thread, and so the application thread's determines the stack size requirement. Again, with the preemptive kernel the stack utilization is the sum of the stacks defined for both threads.

As a result, the difference in overhead between the two cases becomes the extra stack required for the idle thread - which in our case is (a somewhat generous) 128 bytes.

The numbers still don't add up completely, but looking into the linker output we see that the rest of the difference comes from the extra data structures used to manage the kernel in preemptive mode, and the kernel data itself.

Fixed kernel data costs:

```
--- 134 Bytes Kernel data
--- 26 Bytes Kernel Vtables
```

Application (Variable) data costs:

```
--- 24 Bytes Driver Vtables
--- 123 Bytes - statically-allocated kernel objects (semaphores, timers, etc.)
```

With this taken into account, the true memory cost of a 2-thread system ends up being around 428 bytes of RAM - which is about 20% of the total memory available on this particular microcontroller. Whether or not this is reasonable certainly depends on the application, but more importantly, it is not so unreasonable as to eliminate an RTOS-based solution from being considered. Also note that by using the “simulated idle” feature provided in [Mark3 R3](#) and onward, the idle thread (and its associated stack) can be eliminated altogether to reduce the cost in constrained devices.

The difference in code space overhead between the preemptive and cooperative mode solutions is less of an issue. Part of this reason is that both the preemptive and cooperative kernels are relatively small, and even an average target device (like the Atmega328 we've chosen) has plenty of room.

[Mark3](#) can be configured so that only features necessary for the application are included in the RTOS - you only pay for the parts of the system that you use. In this way, we can measure the overhead on a feature-by-feature basis, which is shown below for the kernel as configured for this application:

```
Kernel ..... 2563 Bytes
Synchronization Objects. 644 Bytes
Port ..... 974 Bytes
Features ..... 871 Bytes
```

The configuration tested in this comparison uses the thread/port module with timers, drivers, and semaphores, and mutexes, for a total kernel size of 5052 Bytes, with the rest of the code space occupied by the application.

As can be seen from the compiler's output, the difference in code space between the two versions of the application is 3028 bytes - or about 9% of the available code space on the selected processor. While nearly all of this comes from the added overhead of the kernel, the rest of the difference comes the changes to the application necessary to facilitate the different frameworks. This also demonstrates that the system-software code size in the cooperative case is about 2024 bytes.

8.3 Runtime Overhead

On the cooperative kernel, the overhead associated with running the thread is the time it takes the kernel to notice a pending event flag and launch the appropriate event handler, plus the timer interrupt execution time.

Similarly, on the preemptive kernel, the overhead is the time it takes to switch contexts to the application thread, plus the timer interrupt execution time.

The timer interrupt overhead is similar for both cases, so the overhead then becomes the difference between the following:

Preemptive mode:

- Posting the semaphore that wakes the high-priority thread
- Performing a context switch to the high-priority thread

Cooperative mode:

- Setting the event flag from the timer interrupt
- Acknowledging the event from the event loop

coop – 438 cycles preempt – 764 cycles

Using a cycle-accurate AVR simulator (fIAVR) running with a simulated speed of 16MHz, we find the end-to-end event sequence time to be 27us for the cooperative mode scheduler and 48us for the preemptive, and a raw difference of 20us.

With a fixed high-priority event frequency of 30Hz, we achieve a runtime overhead of 611us per second, or 0.06% of the total available CPU time. Now, obviously this value would expand at higher event frequencies and/or slower CPU frequencies, but for this typical application we find the difference in runtime overhead to be negligible for a preemptive system.

8.4 Analysis

For the selected test application and platform, including a preemptive RTOS is entirely reasonable, as the costs are low relative to a non-preemptive kernel solution. But these costs scale relative to the speed, memory and code space of the target processor. Because of these variables, there is no "magic bullet" environment suitable for every application, but [Mark3](#) attempts to provide a framework suitable for a wide range of targets.

On the one hand, if these tests had been performed on a higher-end microcontroller such as the ATmega1284p (containing 128kB of code space and 16kB of RAM), the overhead would be in the noise. For this type of resource-rich microcontroller, there would be no reason to avoid using the [Mark3](#) preemptive kernel.

Conversely, using a lower-end microcontroller like an ATmega88pa (which has only 8kB of code space and 1kB of RAM), the added overhead would likely be prohibitive for including a preemptive kernel. In this case, the cooperative-mode kernel would be a better choice.

As a rule of thumb, if one budgets 25% of a microcontroller's code space/RAM for system code, you should only require at minimum a microcontroller with 16k of code space and 2kB of RAM as a base platform for an RTOS. Unless there are serious constraints on the system that require much better latency or responsiveness than can be achieved with RTOS overhead, almost any modern platform is sufficient for hosting a kernel. In the event you find yourself with a microprocessor with external memory, there should be no reason to avoid using an RTOS at all.

Chapter 9

Mark3 Design Goals

9.1 Overview

9.1.1 Services Provided by an RTOS Kernel

At its lowest-levels, an operating system kernel is responsible for managing and scheduling resources within a system according to the application. In a typical thread-based RTOS, the resources involved is CPU time, and the kernel manages this by scheduling threads and timers. But capable RTOS kernels provide much more than just threading and timers.

In the following section, we discuss the [Mark3](#) kernel architecture, all of its features, and a thorough discussion of how the pieces all work together to make an awesome RTOS kernel.

9.1.2 Guiding Principles of Mark3

[Mark3](#) was designed with a number of over-arching principles, coming from years of experience designing, implementing, refining, and experimenting with RTOS kernels. Through that process I not only discovered what features I wanted in an RTOS, but how I wanted to build those features to look, work, and “feel”. With that understanding, I started with a clean slate and began designing a new RTOS. [Mark3](#) is the result of that process, and its design goals can be summarized in the following guiding principles.

9.1.3 Be feature competitive

To truly be taken seriously as more than just a toy or educational tool, an RTOS needs to have a certain feature suite. While [Mark3](#) isn't a clone of any existing RTOS, it should at least attempt parity with the most common software in its class.

Looking at its competitors, [Mark3](#) as a kernel supports most, if not all of the compelling features found in modern RTOS kernels, including dynamic threads, tickless timers, efficient message passing, and multiple types of synchronization primitives.

9.1.4 Be highly configuration

[Mark3](#) isn't a one-size-fits-all kernel – and as a result, it provides the means to build a custom kernel to suit your needs. By configuring the kernel at compile-time, [Mark3](#) can be built to contain the optimal feature set for a given application. And since features can be configured individually, you only pay the code/RAM footprint for the features you actually use.

9.1.5 No external dependencies, no new language features

To maximize portability and promote adoption to new platforms, [Mark3](#) is written in a widely supported subset of C++ that lends itself to embedded applications. It avoids RTTI, exceptions, templates, and libraries (C standard, STL, etc.), with all fundamental data structures and types implemented completely for use by the kernel. As a result, the portable parts of [Mark3](#) should compile for any capable C++ toolchain.

9.1.6 Target the most popular hobbyist platforms available

Realistically, this means supporting the various Arduino-compatible target CPUs, including AVR and ARM Cortex-M series microcontrollers. As a result, the current default target for [Mark3](#) is the atmega328p, which has 32KB of flash and 2KB of RAM. All decisions regarding default features, code size, and performance need to take that target system into account.

[Mark3](#) integrates cleanly as a library into the Arduino IDE to support atmega328-based targets. Other AVR and Cortex-M targets can be supported using the port code provided in the source package.

9.1.7 Maximize determinism – but be pragmatic

Guaranteeing deterministic and predictable behavior is tough to do in an embedded system, and often comes with a heavy price tag in either RAM or code-space. With [Mark3](#), we strive to keep the core kernel APIs and features as lightweight as possible, while avoiding algorithms that don't scale to large numbers of threads. We also achieve minimal latency by keeping interrupts enabled (operating out of the critical section) wherever possible.

In [Mark3](#), the most important parts of the kernel are fixed-time, including thread scheduling and context switching. Operations that are not fixed time can be characterized as a function of their dependent data. For instances, the Mutex and Semaphore APIs operate in fixed time in the uncontested case, and execute in linear time for the contested case – where the speed of execution is dependent on the number of threads currently waiting on that object.

The caveat here is that while we want to minimize latency and time spent in critical sections, that has to be balanced against increases in code size, and uncontested-case performance.

9.1.8 Apply engineering principles – and that means discipline, measurement and verification

My previous RTOS, FunkOS, was designed to be very ad-hoc. The usage instructions were along the lines of “drag and drop the source files into your IDE and compile”. There was no regression/unit testing, no code size/speed profiling, and all documentation was done manually. It worked, but the process was a bit of a mess, and resulted in a lot of re-spins of the software, and a lot of time spent stepping through emulators to measure parameters.

We take a different approach in [Mark3](#). Here, we've designed not only the kernel-code, but the build system, unit tests, profiling code, documentation and reporting that supports the kernel. Each release is built and tested using automation in order to ensure quality and correctness, with supporting documentation containing all critical metrics. Only code that passes testing is submitted to the repos and public forums for distribution. These metrics can be traced from build-to-build to ensure that performance remains consistent from one drop to the next, and that no regressions are introduced by new/refactored code.

And while the kernel code can still be exported into an IDE directly, that takes place with the knowledge that the kernel code has already been rigorously tested and profiled. Exporting source in [Mark3](#) is also supported by scripting to ensure reliable, reproducible results without the possibility for human-error.

9.1.9 Use Virtualization For Verification

Mark3 was designed to work with automated simulation tools as the primary means to validate changes to the kernel, due to the power and flexibility of automatic tests on virtual hardware. I was also intrigued by the thought of extending the virtual target to support functionality useful for a kernel, but not found on real hardware.

When the project was started, simavr was the tool of choice- however, its simulation was found to be incorrect compared to execution on a real MCU, and it did not provide the degree of extension that I desired for use with kernel development.

The flAVR AVR simulator was written to replace the dependency on that tool, and overcome those limitations. It also provides a GDB interface, as well as its own built-in debugger, profilers, and trace tools.

The example and test code relies heavily on flAVR kernel aware messaging, so it is recommended that you familiarize yourself with that tool if you intend to do any sort of customizations or extensions to the kernel.

flAVR is hosted on sourceforge at <http://www.sourceforge.net/projects/flavr/> . In its basic configuration, it builds with minimal external dependencies.

- On linux, it requires only pthreads.
- On Windows, it requires pthreads and ws2_32, both satisfied via MinGW.
- Optional SDL builds for both targets (featuring graphics and simulated joystick input) can be built, and rely on libSDL.

Chapter 10

Mark3 Kernel Architecture

10.1 Overview

At a high level, the [Mark3](#) RTOS is organized into the following features, and layered as shown below:

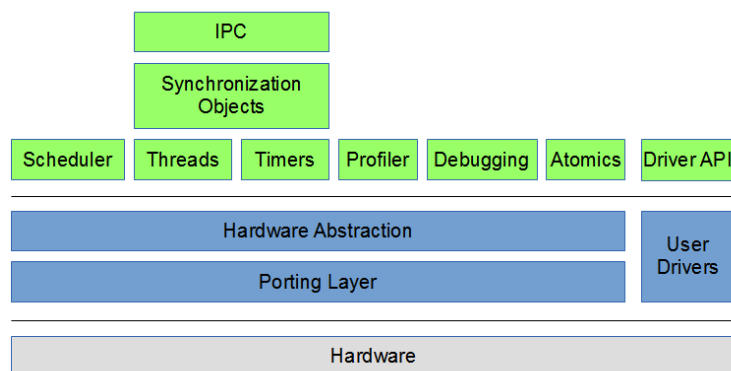


Figure 10.1 Overview

Everything in the “green” layer represents the [Mark3](#) public API and classes, beneath which lives all hardware abstraction and CPU-specific porting and driver code, which runs on a given target CPU.

The features and concepts introduced in this diagram can be described as follows:

Threads: The ability to multiplex the CPU between multiple tasks to give the perception that multiple programs are running simultaneously. Each thread runs in its own context with its own stack.

Scheduler: Algorithm which determines the thread that gets to run on the CPU at any given time. This algorithm takes into account the priorities (and other execution parameters) associated with the threads in the system.

IPC: Inter-process-communications. Message-passing and Mailbox interfaces used to communicate between threads synchronously or asynchronously.

Synchronization Objects: Ability to schedule thread execution relative to system conditions and events, allowing for sharing global data and resources safely and effectively.

Timers: High-resolution software timers that allow for actions to be triggered on a periodic or one-shot basis.

Profiler: Special timer used to measure the performance of arbitrary blocks of code.

Debugging: Realtime logging and trace functionality, facilitating simplified debugging of systems using the OS.

Atomics: Support for UN-interruptible arithmetic operations.

Driver API: Hardware abstraction interface allowing for device drivers to be written in a consistent, portable manner.

Hardware Abstraction Layer: Class interface definitions to represent threading, context-switching, and timers in a generic, abstracted manner.

Porting Layer: Class interface implementation to support threading, context-switching, and timers for a given CPU.

User Drivers: Code written by the user to implement device-specific peripheral drivers, built to make use of the [Mark3](#) driver API.

Each of these features will be described in more detail in the following sections of this chapter.

The concepts introduced in the above architecture are implemented in a variety of source modules, which are logically broken down into classes (or in some cases, groups of functions/macros). The relationship between objects in the [Mark3](#) kernel is shown below:

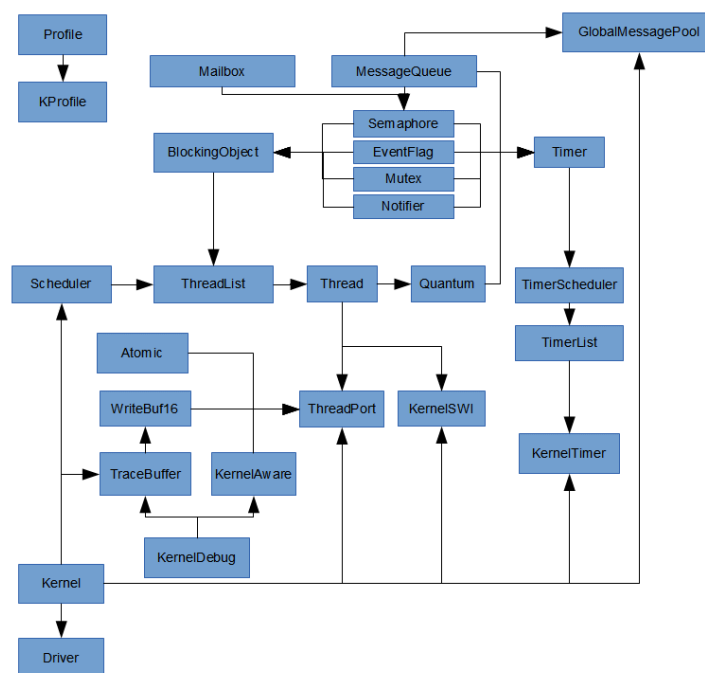


Figure 10.2 Overview

The objects shown in the preceding table can be grouped together by feature. In the table below, we group each feature by object, referencing the source module in which they can be found in the [Mark3](#) source tree.

Feature	Kernel Object	Source Files
Profiling	ProfileTimer	profile.cpp/.h
Threads + Scheduling	Thread	thread.cpp/.h
	Scheduler	scheduler.cpp/.h
	PriorityMap	priomap.cpp/.h
	Quantum	quantum.cpp/.h

Feature	Kernel Object	Source Files
	ThreadPort	threadport.cpp/.h **
	KernelSWI	kernelswi.cpp/.h **
Timers	Timer	timer.h/timer.cpp
	TimerScheduler	timerscheduler.h
	TimerList	timerlist.h/cpp
	KernelTimer	kerneltimer.cpp/.h **
Synchronization	BlockingObject	blocking.cpp/.h
	Semaphore	ksemaphore.cpp/.h
	EventFlag	eventflag.cpp/.h
	Mutex	mutex.cpp/.h
	Notify	notify.cpp/.h
	ConditionVariable	condvar.cpp/.h
	ReaderWriterLock	readerwriter.cpp/.h
IPC/Message-passing	Mailbox	mailbox.cpp/.h
	MessageQueue	message.cpp/.h
	GlobalMessagePool	message.cpp/.h
Debugging	Miscellaneous Macros	kerneldebug.h
	KernelAware	kernelaware.cpp/.h
	TraceBuffer	tracebuffer.cpp/.h
	Buffalogger	buffalogger.h
Atomic Operations	Atomic	atomic.cpp/.h
Kernel	Kernel	kernel.cpp/.h

** implementation is platform-dependent, and located under the kernel's
 ** /cpu/<arch>/<variant>/<toolchain> folder in the source tree

10.2 Threads and Scheduling

The classes involved in threading and scheduling in [Mark3](#) are highlighted in the following diagram, and are discussed in detail in this chapter:

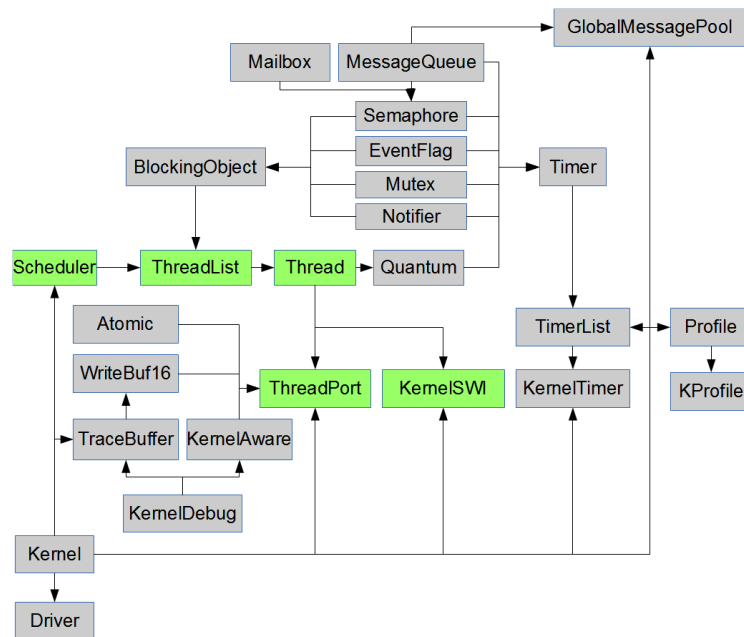


Figure 10.3 Threads and Scheduling

10.2.1 A Bit About Threads

Before we get started talking about the internals of the [Mark3](#) scheduler, it's necessary to go over some background material - starting with: what is a thread, anyway?

Let's look at a very basic CPU without any sort of special multi-threading hardware, and without interrupts. When the CPU is powered up, the program counter is loaded with some default location, at which point the processor core will start executing instructions sequentially - running forever and ever according to whatever has been loaded into program memory. This single instance of a simple program sequence is the only thing that runs on the processor, and the execution of the program can be predicted entirely by looking at the CPU's current register state, its program, and any affected system memory (the CPU's "context").

It's simple enough, and that's exactly the definition we have for a thread in an RTOS.

Each thread contains an instance of a CPU's register context, its own stack, and any other bookkeeping information necessary to define the minimum unique execution state of a system at runtime. It is the job of a RTOS to multiplex the execution of multiple threads on a single physical CPU, thereby creating the illusion that many programs are being executed simultaneously. In reality there can only ever be one thread truly executing at any given moment on a CPU core, so it's up to the scheduler to set and enforce rules about what thread gets to run when, for how long, and under what conditions. As mentioned earlier, any system without an RTOS executes as a single thread, so at least two threads are required for an RTOS to serve any useful purpose.

Note that all of this information is common to pretty well every RTOS in existence - the implementation details, including the scheduler rules, are all part of what differentiates one RTOS from another.

10.2.2 Thread States and ThreadLists

Since only one thread can run on a CPU at a time, the scheduler relies on thread information to make its decisions. [Mark3](#)'s scheduler relies on a variety of such information, including:

- The thread's current priority
- Round-Robin execution quanta
- Whether or not the thread is blocked on a synchronization object, such as a mutex or semaphore
- Whether or not the thread is currently suspended

The scheduler further uses this information to logically place each thread into 1 of 4 possible states:

```
- Ready - The thread is currently running
- Running - The thread is able to run
- Blocked - The thread cannot run until a system condition is met
- Stopped - The thread cannot run because its execution has been suspended
.
```

In order to determine a thread's state, threads are placed in "buckets" corresponding to these states. Ready and running threads exist in the scheduler's buckets, blocked threads exist in a bucket belonging to the object they're blocked on, and stopped threads exist in a separate bucket containing all stopped threads.

In reality, the various buckets are just doubly-linked lists of Thread objects - implemented in something called the ThreadList class. To facilitate this, the Thread class directly inherits from a LinkListNode class, which contains the node pointers required to implement a doubly-linked list. As a result, Threads may be effortlessly moved from one state to another using efficient linked-list operations built into the ThreadList class.

10.2.3 Blocking and Unblocking

While many developers new to the concept of an RTOS assume that all threads in a system are entirely separate from each other, the reality is that practical systems typically involve multiple threads working together, or at the very least sharing resources. In order to synchronize the execution of threads for that purpose, a number of synchronization primitives (blocking objects) are implemented to create specific sets of conditions under which threads can continue execution. The concept of "blocking" a thread until a specific condition is met is fundamental to understanding RTOS applications design, as well as any highly-multithreaded applications.

10.2.4 Blocking Objects

Blocking objects and primitives provided by [Mark3](#) include:

- Semaphores (binary and counting)
- Mutexes
- Event Flags
- Thread Notification Objects
- Thread Sleep
- Message Queues

- Mailboxes

The relationship between these objects in the system are shown below:

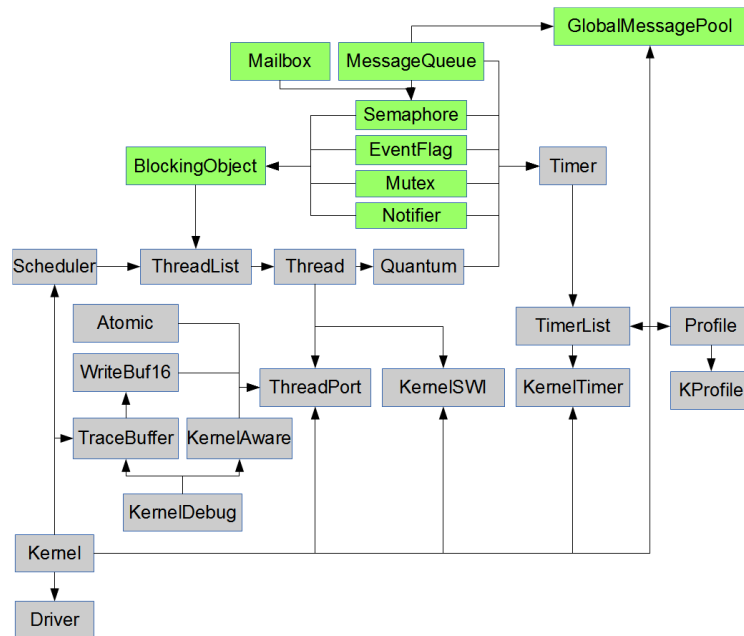


Figure 10.4 Blocking Objects

Each of these objects inherit from the BlockingObject class, which itself contains a ThreadList object. This class contains methods to Block() a thread (remove it from the Scheduler's "Ready" or "Running" ThreadLists), as well as Unblock() a thread (move a thread back to the "Ready" lists). This object handles transitioning threads from list-to-list (and state-to-state), as well as taking care of any other Scheduler bookkeeping required in the process. While each of the Blocking types implement a different condition, they are effectively variations on the same theme. Many simple Blocking objects are also used to build complex blocking objects - for instance, the Thread Sleep mechanism is essentially a binary semaphore and a timer object, while a message queue is a linked-list of message objects combined with a semaphore.

10.3 Inside the Mark3 Scheduler

At this point we've covered the following concepts:

- Threads
- Thread States and Thread Lists
- Blocking and Un-Blocking Threads

Thankfully, this is all the background required to understand how the [Mark3](#) Scheduler works. In technical terms, [Mark3](#) implements "strict priority scheduling, with round-robin scheduling among threads in each priority group". In plain English, this boils down to a scheduler which follows a few simple rules:

```
Find the highest-priority "Ready" list that has at least one Threads.
If the first thread in that bucket is not the current thread, select it
to run next. Otherwise, rotate the linked list, and choose the next
thread in the list to run
```

Since context switching is one of the most common and frequent operation performed by an RTOS, this needs to be as fast and deterministic as possible. While the logic is simple, a lot of care must be put into optimizing the scheduler to achieve those goals. In the section below we discuss the optimization approaches taken in [Mark3](#).

There are a number of ways to find the highest-priority thread. The naive approach would be to simply iterate through the scheduler's array of ThreadLists from highest to lowest, stopping when the first non-empty list is found, such as in the following block of code:

```
for (prio = num_prio - 1; prio >= 0; prio--)
{
    if (thread_list[prio].get_head() != NULL)
    {
        break;
    }
}
```

While that would certainly work and be sufficient for a variety of systems, it's a non-deterministic approach (complexity $O(n)$) whose cost varies substantially based on how many priorities have to be evaluated. It's simple to read and understand, but it's non-optimal.

Fortunately, a functionally-equivalent and more deterministic approach can be implemented with a few tricks.

In addition to maintaining an array of ThreadLists, [Mark3](#) also maintains a bitmap (one bit per priority level) that indicates which thread lists have ready threads. This bitmap is maintained automatically by the ThreadList class, and is updated every time a thread is moved to/from the Scheduler's ready lists.

By inspecting this bitmap using a technique to count the leading zero bits in the bitmap, we determine which threadlist to choose in fixed time.

Now, to implement the leading-zeros check, this can once again be performed iteratively using bitshifts and compares (which isn't any more efficient than the raw list traversal), but it can also be evaluated using either a lookup table, or via a special CPU instruction to count the leading zeros in a value. In [Mark3](#), we opt for the lookup-table approach since we have a limited number of priorities and not all supported CPU architectures support a count leading zero instruction. To achieve a balance between performance and memory use, we use a 4-bit lookup table (costing 16 bytes) to perform the lookup.

(As a sidenote - this is actually a very common approach in OS schedulers. It's actually part of the reason why modern ARM cores implement a dedicated count-leading-zeros [CLZ] instruction!)

With a 4-bit lookup table and an 8-bit priority-level bitmap, the priority check algorithm looks something like this:

```
// Check the highest 4 priority levels, represented in the
// upper 4 bits in the bitmap
priority = priority_lookup_table[(priority_bitmap >> 4)];

// priority is non-zero if we found something there
if( priority )
{
    // Add 4 because we were looking at the higher levels
    priority += 4;
}
else
{
    // Nothing in the upper 4, look at the lowest 4 priority levels
    // represented by the lowest 4 bits in the bitmap
    priority = priority_lookup_table[priority_bitmap & 0x0F];
}
```

Deconstructing this algorithm, you can see that the priority lookup will have an $O(1)$ complexity - and is extremely low-cost.

This operation is thus fully deterministic and time bound - no matter how many threads are scheduled, the operation will always be time-bound to the most expensive of these two code paths. Even with only 8 priority levels, this is still much faster than iteratively checking the thread lists manually, compared with the previous example implementation.

Once the priority level has been found, selecting the next thread to run is trivial, consisting of something like this:

```
next_thread = thread_list[prio].get_head();
```

In the case of the `get_head()` calls, this evaluates to an inline-load of the "head" pointer in the particular thread list.

One important thing to take away from this analysis is that the scheduler is only responsible for selecting the next-to-run thread. In fact, these two operations are totally decoupled - no context switching is performed by the scheduler, and the scheduler isn't called from the context switch. The scheduler simply produces new "next thread" values that are consumed from within the context switch code.

10.3.1 Considerations for Round-Robin Scheduling

One thing that isn't considered directly from the scheduler algorithm is the problem of dealing with multiple threads within a single priority group; all of the algorithms that have been explored above simply look at the first Thread in each group.

Mark3 addresses this issue indirectly, using a software timer to manage round-robin scheduling, as follows.

In some instances where the scheduler is run by the kernel directly (typically as a result of calling `Thread::Yield()`), the kernel will perform an additional check after running the Scheduler to determine whether or there are multiple ready Threads in the priority of the next ready thread.

If there are multiple threads within that priority, the kernel adds a one-shot software timer which is programmed to expire at the next Thread's configured quantum. When this timer expires, the timer's callback function executes to perform two simple operations:

"Pivot" the current Thread's priority list. Set a flag telling the kernel to trigger a Yield after exiting the main Timer ↔ Scheduler processing loop

Pivoting the thread list basically moves the head of a circular-linked-list to its next value, which in our case ensures that a new thread will be chosen the next time the scheduler is run (the scheduler only looks at the head node of the priority lists). And by calling `Yield`, the system forces the scheduler to run, a new round-robin software timer to be installed (if necessary), and triggers a context switch SWI to load the newly-chosen thread. Note that if the thread attached to the round-robin timer is pre-empted, the kernel will take steps to abort and invalidate that round-robin software timer, installing a new one tied to the next thread to run if necessary.

Because the round-robin software timer is dynamically installed when there are multiple ready threads at the highest ready priority level, there is no CPU overhead with this feature unless that condition is met. The cost of round-robin scheduling is also fixed - no matter how many threads there are, and the cost is identical to any other one-shot software timer in the system.

10.3.2 Context Switching

There's really not much to say about the actual context switch operation at a high level. Context switches are triggered whenever it has been determined that a new thread needs to be swapped into the CPU core when the scheduler is run. [Mark3](#) implements also context switches as a call to a software interrupt - on AVR platforms, we typically use INT0 or INT2 for this (although any pin-change GPIO interrupt can be used), and on ARM we achieve this by triggering a PendSV exception.

However, regardless of the architecture, the contex-switch ISR will perform the following three operations:

Save the current Thread's context to the current Thread stack
Make the "next to run" thread the "currently running" thread
Restore the context of the next Thread from the Thread stack

The code to implement the context switch is entirely architecture-specific, so it won't be discussed in detail here. It's almost always gory inline-assembly which is used to load and store various CPU registers, and is highly-optimized for speed. We dive into an example implementation for the ARM Cortex-M0 microcontroller in a later section of this book.

10.3.3 Putting It All Together

In short, we can say that the [Mark3](#) scheduler works as follows:

- The scheduler is run whenever a `Thread::Yield()` is called by a user, as part of blocking calls, or whenever a new thread is started
- The [Mark3](#) scheduler is deterministic, selecting the next thread to run in fixed-time
- The scheduler only chooses the next thread to run, the context switch SWI consumes that information to get that thread running
- Where there are multiple ready threads in the highest populated priority level, a software timer is used to manage round-robin scheduling

While we've covered a lot of ground in this section, there's not a whole lot of code involved. However, the code that performs these operations is nuanced and subtle. If you're interested in seeing how this all works in practice, I suggest reading through the [Mark3](#) source code (which is heavily annotated), and stepping through the code with a simulator/emulator.

10.4 Timers

[Mark3](#) implements one-shot and periodic software-timers via the `Timer` class. The user configures the timer for duration, repetition, and action, at which point the timer can be activated. When an active timer expires, the kernel calls a user-specified callback function, and then reloads the timer in the case of periodic timers. The same timer objects exposed to the user are also used within the kernel to implement round-robin scheduling, and timeout-based APIs for semaphores, mutexes, events, and messages.

Timers are implemented using the following components in the [Mark3](#) Kernel:

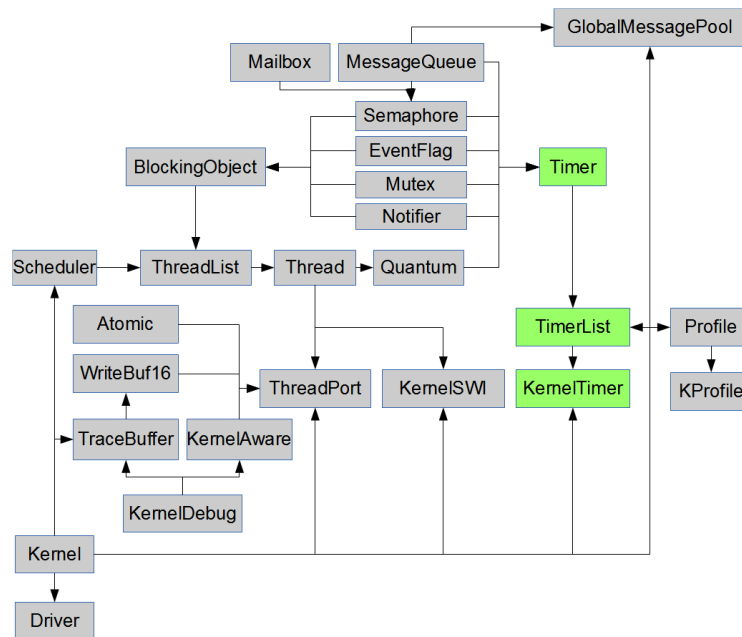


Figure 10.5 Timers

The `Timer` class provides the basic periodic and one-shot timer functionality used by application code, blocking objects, and IPC.

The `TimerList` class implements a doubly-linked list of `Timer` objects, and the logic required to implement a timer tick (tick-based kernel) or timer expiry (tickless kernel) event.

The `TimerScheduler` class contains a single `TimerList` object, implementing a single, system-wide list of `Timer` objects within the kernel. It also provides hooks for the hardware timer, such that when a timer tick or expiry event occurs, the `TimerList` expiry handler is run.

The `KernelTimer` class ([kerneltimer.cpp/.h](#)) implements the CPU specific hardware timer driver that is used by the kernel and the `TimerScheduler` to implement software timers.

While extremely simple to use, they provide one of the most powerful execution contexts in the system.

The software timers implemented in [Mark3](#) use interrupt-nesting within the kernel timer's interrupt handler. This context is be considered higher-priority than the highest priority user thread, but lower-priority than other interrupts in the system. As a result, this minimizes critical interrupt latency in the system, albeit at the expense of responsiveness of the user-threads.

For this reason, it's critical to ensure that all timer callback events are kept as short as possible to prevent adding thread-level latency. All heavy-lifting should be left to the threads, so the callback should only implement signalling via IPC or synchronization object.

The time spent in this interrupt context is also dependent on the number of active timers at any given time. However, [Mark3](#) also can be used to minimize the frequency of these interrupts wakeups, by using an optional "tolerance" parameter in the timer API calls. In this way, periodic tasks that have less rigorous real-time constraints can all be grouped together – executing as a group instead of one-after-another.

[Mark3](#) also contains two different timer implementations that can be configured at build-time, each with their own advantages.

10.4.1 Tick-based Timers

In a tick-based timing scheme, the kernel relies on a system-timer interrupt to fire at a relatively-high frequency, on which all kernel timer events are derived. On modern CPUs and microcontrollers, a 1kHz system tick is common, although quite often lower frequencies such as 60Hz, 100Hz, or 120Hz are used. The resolution of this timer also defines the maximum resolution of timer objects as a result. That is, if the timer frequency is 1kHz, a user cannot specify a timer resolution lower than 1ms.

The advantage of a tick-based timer is its sheer simplicity. It typically doesn't take much to set up a timer to trigger an interrupt at a fixed-interval, at which point, all system timer intervals are decremented by 1 count. When each system timer interval reaches zero, a callback is called for the event, and the events are either reset and restarted (repeated timers) or cleared (1-shot).

Unfortunately, that simplicity comes at a cost of increased interrupt count, which cause frequent CPU wakeups and utilization, and power consumption.

10.4.2 Tickless Timers

In a tickless system, the kernel timer only runs when there are active timers pending expiry, and even then, the timer module only generates interrupts when a timer expires, or a timer reaches its maximum count value. Additionally, when there are no active timer objects, the timer can be completely disabled – saving even more cycles, power, and CPU wakeups. These factors make the tickless timer approach a highly-optimal solution, suitable for a wide array of low-power applications.

Also, since tickless timers do not rely on a fixed, periodic clock, they can potentially be higher resolution. The only limitation in timer resolution is the precision of the underlying hardware timer as configured. For example, if a 32kHz hardware timer is being used to drive the timer scheduler, the resolution of timer objects would be in the $\sim 33\mu\text{s}$ range.

The only downside of the tickless timer system is an added complexity to the timer code, requiring more code space, and slightly longer execution of the timer routines when the timer interrupt is executed.

10.4.3 Timer Processing Algorithm

Timer interrupts occur at either a fixed-frequency (tick-based), or at the next timer expiry interval (tickless), at which point the timer processing algorithm runs. While the timer count is reset by the timer-interrupt, it is still allowed to accumulate ticks while this algorithm is executed in order to ensure that timer-accuracy is kept in real-time. It is also important to note that round-robin scheduling changes are disabled during the execution of this algorithm to prevent race conditions, as the round-robin code also relies on timer objects.

All active timer objects are stored in a doubly-linked list within the timer-scheduler, and this list is processed in two passes by the algorithm which runs from the timer-interrupt (with interrupt nesting enabled). The first pass determines which timers have expired and the next timer interval, while the second pass deals with executing the timer callbacks themselves. Both phases are discussed in more detail below.

In the first pass, the active timers are decremented by either 1 tick (tick-based), or by the duration of the last elapsed timer interval (tickless). Timers that have zero (or less-than-zero) time remaining have a “callback” flag set, telling the algorithm to call the timer's callback function in the second pass of the loop. In the event of a periodic timer, the timer's interval is reset to its starting value.

For the tickless case, the next timer interval is also computed in the first-pass by looking for the active timer with the least amount of time remaining in its interval. Note that this calculation is irrelevant in the tick-based timer code, as the timer interrupt fires at a fixed-frequency.

In the second pass, the algorithm loops through the list of active timers, looking for those with their “callback” flag set in the first pass. The callback function is then executed for each expired timer, and the “callback” flag cleared. In the event that a non-periodic (one-shot) timer expires, the timer is also removed from the timer scheduler at this time.

In a tickless system, once the second pass of the loop has been completed, the hardware timer is checked to see if the next timer interval has expired while processing the expired timer callbacks. In that event, the complete algorithm is re-run to ensure that no expired timers are missed. Once the algorithm has completed without the next timer expiring during processing, the expiry time is programmed into the hardware timer. Round-robin scheduling is re-enabled, and if a new thread has been scheduled as a result of action taken during a timer callback, a context switch takes place on return from the timer interrupt.

10.5 Synchronization and IPC

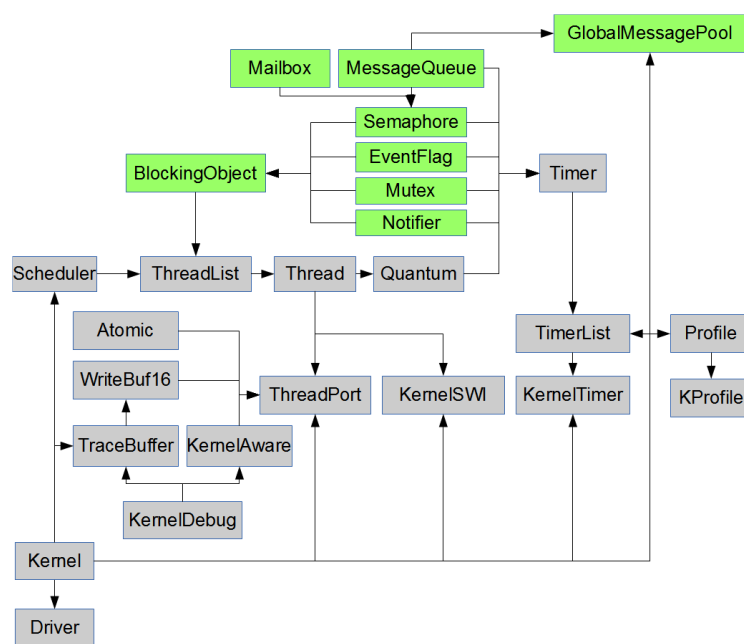


Figure 10.6 Synchronization and IPC

10.6 Blocking Objects

A Blocking object in [Mark3](#) is essentially a thread list. Any blocking object implementation (being a semaphore, mutex, event flag, etc.) can be built on top of this class, utilizing the provided functions to manipulate thread location within the Kernel.

Blocking a thread results in that thread becoming de-scheduled, placed in the blocking object's own private list of threads which are waiting on the object.

Unblocking a thread results in the reverse: The thread is moved back to its original location from the blocking list.

The only difference between a blocking object based on this class is the logic used to determine what constitutes a Block or Unblock condition.

For instance, a semaphore Pend operation may result in a call to the Block() method with the currently-executing thread in order to make that thread wait for a semaphore Post. That operation would then invoke the Unblock() method, removing the blocking thread from the semaphore's list, and back into the appropriate thread inside the scheduler.

Care must be taken when implementing blocking objects to ensure that critical sections are used judiciously, otherwise asynchronous events like timers and interrupts could result in non-deterministic and often catastrophic behavior.

[Mark3](#) implements a variety of blocking objects including semaphores, mutexes, event flags, and IPC mechanisms that all inherit from the basic Blocking-object class found in [blocking.h/cpp](#), ensuring consistency and a high degree of code-reuse between components.

10.6.1 Semaphores

Semaphores are used to synchronized execution of threads based on the availability (and quantity) of application-specific resources in the system. They are extremely useful for solving producer-consumer problems, and are the method-of-choice for creating efficient, low latency systems, where ISRs post semaphores that are handled from within the context of individual threads. Semaphores can also be posted (but not pended) from within the interrupt context.

10.6.2 Mutex

Mutexes (Mutual exclusion objects) are provided as a means of creating "protected sections" around a particular resource, allowing for access of these objects to be serialized. Only one thread can hold the mutex at a time

- other threads have to wait until the region is released by the owner thread before they can take their turn operating on the protected resource. Note that mutexes can only be owned by threads - they are not available to other contexts (i.e. interrupts). Calling the mutex APIs from an interrupt will cause catastrophic system failures.

Note that these objects are recursive in [Mark3](#) - that is, the owner thread can claim a mutex more than once. The caveat here is that a recursively-held mutex will not be released until a matching "release" call is made for each "claim" call.

Priority inheritance is provided with these objects as a means to avoid priority inversions. Whenever a thread at a priority than the mutex owner blocks on a mutex, the priority of the current thread is boosted to the highest-priority waiter to ensure that other tasks at intermediate priorities cannot artificially prevent progress from being made.

10.6.3 Event Flags

Event Flags are another synchronization object, conceptually similar to a semaphore.

Unlike a semaphore, however, the condition on which threads are unblocked is determined by a more complex set of rules. Each Event Flag object contains a 16-bit field, and threads block, waiting for combinations of bits within this field to become set.

A thread can wait on any pattern of bits from this field to be set, and any number of threads can wait on any number of different patterns. Threads can wait on a single bit, multiple bits, or bits from within a subset of bits within the field.

As a result, setting a single value in the flag can result in any number of threads becoming unblocked simultaneously. This mechanism is extremely powerful, allowing for all sorts of complex, yet efficient, thread synchronization schemes that can be created using a single shared object.

Note that Event Flags can be set from interrupts, but you cannot wait on an event flag from within an interrupt.

10.6.4 Notification Objects

Notification objects are the most lightweight of all blocking objects supplied by [Mark3](#).

using this blocking primitive, one or more threads wait for the notification object to be signalled by code elsewhere in the system (i.e. another thread or interrupt). Once the notification has been signalled, all threads currently blocked on the object become unblocked and moved into the ready list.

Signalling a notification object that has no actively-waiting threads has no effect.

10.7 Messages and Global Message Queue

10.7.1 Messages

Sending messages between threads is the key means of synchronizing access to data, and the primary mechanism to perform asynchronous data processing operations.

Sending a message consists of the following operations:

- Obtain a Message object from the global message pool
- Set the message data and event fields
- Send the message to the destination message queue

While receiving a message consists of the following steps:

- Wait for a messages in the destination message queue
- Process the message data
- Return the message back to the global message pool

These operations, and the various data objects involved are discussed in more detail in the following section.

10.7.2 Message Objects

Message objects are used to communicate arbitrary data between threads in a safe and synchronous way.

The message object consists of an event code field and a data field. The event code is used to provide context to the message object, while the data field (essentially a void * data pointer) is used to provide a payload of data corresponding to the particular event.

Access to these fields is marshalled by accessors - the transmitting thread uses the `SetData()` and `SetCode()` methods to seed the data, while the receiving thread uses the `GetData()` and `GetCode()` methods to retrieve it.

By providing the data as a void data pointer instead of a fixed-size message, we achieve an unprecedented measure of simplicity and flexibility. Data can be either statically or dynamically allocated, and sized appropriately for the event without having to format and reformat data by both sending and receiving threads. The choices here are left to the user - and the kernel doesn't get in the way of efficiency.

It is worth noting that you can send messages to message queues from within ISR context. This helps maintain consistency, since the same APIs can be used to provide event-driven programming facilities throughout the whole of the OS.

10.7.3 Global Message Pool

To maintain efficiency in the messaging system (and to prevent over-allocation of data), a global pool of message objects is provided. The size of this message pool is specified in the implementation, and can be adjusted depending on the requirements of the target application as a compile-time option.

Allocating a message from the message pool is as simple as calling the

`GlobalMessagePool::Pop()` Method.

Messages are returned back to the `GlobalMessagePool::Push()` method once the message contents are no longer required.

One must be careful to ensure that discarded messages always are returned to the pool, otherwise a resource leak will occur, which may cripple the operating system's ability to pass data between threads.

10.7.4 Message Queues

Message objects specify data with context, but do not specify where the messages will be sent. For this purpose we have a `MessageQueue` object. Sending an object to a message queue involves calling the `MessageQueue::Send()` method, passing in a pointer to the `Message` object as an argument.

When a message is sent to the queue, the first thread blocked on the queue (as a result of calling the `MessageQueue::Receive()` method) will wake up, with a pointer to the `Message` object returned.

It's worth noting that multiple threads can block on the same message queue, providing a means for multiple threads to share work in parallel.

10.7.5 Mailboxes

Another form of IPC is provided by [Mark3](#), in the form of Mailboxes and Envelopes. Mailboxes are similar to message queues in that they provide a synchronized interface by which data can be transmitted between threads.

Where Message Queues rely on linked lists of lightweight message objects (containing only message code and a `void*` data-pointer), which are inherently abstract, Mailboxes use a dedicated blob of memory, which is carved up into fixed-size chunks called Envelopes (defined by the user), which are sent and received. Unlike message queues, mailbox data is copied to and from the mailboxes dedicated pool.

Mailboxes also differ in that they provide not only a blocking "receive" call, but also a blocking "send" call, providing the opportunity for threads to block on "mailbox full" as well as "mailbox empty" conditions.

All send/receive APIs support an optional timeout parameter if the `KERNEL_USE_TIMEOUTS` option has been configured in [mark3cfg.h](#)

10.7.6 Atomic Operations

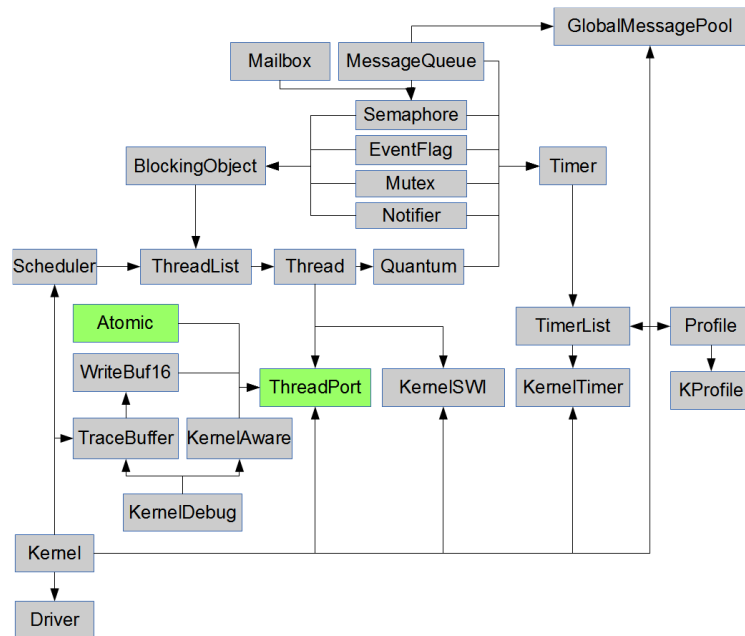


Figure 10.7 Atomic operations

This utility class provides primitives for atomic operations - that is, operations that are guaranteed to execute uninterrupted. Basic atomic primitives provided here include Set/Add/Delete for 8, 16, and 32-bit integer types, as well as an atomic test-and-set.

10.7.7 Atomic Operations

10.7.8 Drivers

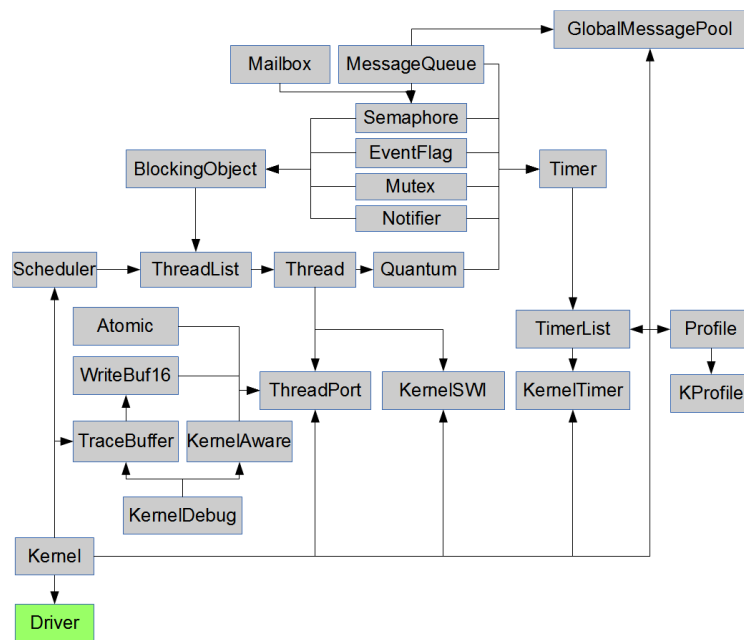


Figure 10.8 Drivers

This is the basis of the driver framework. In the context of [Mark3](#), drivers don't necessarily have to be based on physical hardware peripherals. They can be used to represent algorithms (such as random number generators), files, or protocol stacks. Unlike FunkOS, where driver IO is protected automatically by a mutex, we do not use this kind of protection - we leave it up to the driver implementor to do what's right in its own context. This also frees up the driver to implement all sorts of other neat stuff, like sending messages to threads associated with the driver. Drivers are implemented as character devices, with the standard array of posix-style accessor methods for reading, writing, and general driver control.

A global driver list is provided as a convenient and minimal "filesystem" structure, in which devices can be accessed by name.

Driver Design

A device driver needs to be able to perform the following operations:

- Initialize a peripheral
- Start/stop a peripheral
- Handle I/O control operations
- Perform various read/write operations

At the end of the day, that's pretty much all a device driver has to do, and all of the functionality that needs to be presented to the developer.

We abstract all device drivers using a base-class which implements the following methods:

- Start/Open
- Stop/Close
- Control
- Read
- Write

A basic driver framework and API can thus be implemented in five function calls - that's it! You could even reduce that further by handling the initialize, start, and stop operations inside the "control" operation.

Driver API

In C++, we can implement this as a class to abstract these event handlers, with virtual void functions in the base class overridden by the inherited objects.

To add and remove device drivers from the global table, we use the following methods:

```
void DriverList::Add( Driver *pclDriver_ );
void DriverList::Remove( Driver *pclDriver_ );
```

`DriverList::Add()/Remove()` takes a single argument - the pointer to the object to operate on.

Once a driver has been added to the table, drivers are opened by NAME using `DriverList::FindByName("/dev/name")`. This function returns a pointer to the specified driver if successful, or to a built in /dev/null device if the path name is invalid. After a driver is open, that pointer is used for all other driver access functions.

This abstraction is incredibly useful - any peripheral or service can be accessed through a consistent set of APIs, that make it easy to substitute implementations from one platform to another. Portability is ensured, the overhead is negligible, and it emphasizes the reuse of both driver and application code as separate entities.

Consider a system with drivers for I2C, SPI, and UART peripherals - under our driver framework, an application can initialize these peripherals and write a greeting to each using the same simple API functions for all drivers:

```
pclI2C = DriverList::FindByName("/dev/i2c");
pclUART = DriverList::FindByName("/dev/tty0");
pclSPI = DriverList::FindByName("/dev/spi");

pclI2C->Write(12, "Hello World!");
pclUART->Write(12, "Hello World!");
pclSPI->Write(12, "Hello World!");
```

10.8 Kernel Proper and Porting

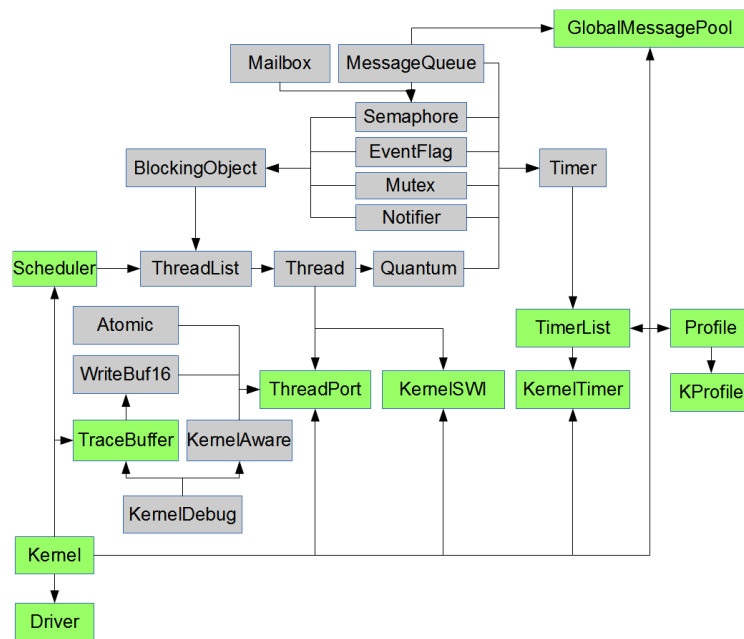


Figure 10.9 Kernel Proper and Porting

The Kernel class is a static class with methods to handle the initialization and startup of the RTOS, manage errors, and provide user-hooks for fatal error handling (functions called when `Kernel::Panic()` conditions are encountered), or when the Idle function is run.

Internally, `Kernel::Init()` calls the initialization routines for various kernel objects, providing a single interface by which all RTOS-related system initialization takes place.

`Kernel::Start()` is called to begin running OS functionality, and does not return. Control of the CPU is handed over to the scheduler, and the highest-priority ready thread begins execution in the RTOS environment.

Hardware Abstraction Layer

Almost all of the [Mark3](#) kernel (and middleware) is completely platform independent, and should compile cleanly on any platform with a modern C++ compiler. However, there are a few areas within [Mark3](#) that can only be implemented by touching hardware directly.

These interfaces generally cover four features:

- Thread initialization and context-switching logic
- Software interrupt control (used to generate context switches)
- Hardware timer control (support for time-based functionality, such as `Sleep()`)
- Code-execution profiling timer (not necessary to port if code-profiling is not compiled into the kernel)

The hardware abstraction layer in [Mark3](#) provides a consistent interface for each of these four features. [Mark3](#) is ported to new target architectures by providing an implementation for all of the interfaces declared in the abstraction layer. In the following section, we will explore how this was used to port the kernel to ARM Cortex-M0.

Real-world Porting Example – Cortex M0

This section serves as a real-world example of how [Mark3](#) can be ported to new architectures, how the [Mark3](#) abstraction layer works, and as a practical reference for using the RTOS support functionality baked in modern ARM Cortex-M series microcontrollers. Most of this documentation here is taken directly from the source code found in the `/kernel/cpu/cm0/ports` directory, with additional annotations to explain the port in more detail. Note that a familiarity with Cortex-M series parts will go a long way to understanding the subject matter presented, especially a basic understanding of the ARM CPU registers, exception models, and OS support features (PendSV, SysTick and SVC). If you're unfamiliar with ARM architecture, pay attention to the comments more than the source itself to illustrate the concepts.

Porting [Mark3](#) to a new architecture consists of a few basic pieces; for developers familiar with the target architecture and the porting process, it's not a tremendously onerous endeavour to get [Mark3](#) up-and-running somewhere new. For starters, all non-portable components are completely isolated in the source-tree under:

```
/embedded/kernel/CPU/VARIANT/TOOLCHAIN/
```

where CPU is the architecture, VARIANT is the vendor/part, and TOOLCHAIN is the compiler tool suite used to build the code.

From within the specific port folder, a developer needs only implement a few classes and headers that define the port-specific behavior of [Mark3](#):

- KernelSWI ([kernelswi.cpp/kernelswi.h](#)) - Provides a maskable software-triggered interrupt used to perform context switching.
- KernelTimer ([kerneltimer.cpp/kerneltimer.h](#)) - Provides either a fixed-frequency or programmable-interval timer, which triggers an interrupt on expiry. This is used for implementing round-robin scheduling, thread-sleeps, and generic software timers.
- Profiler ([kprofile.cpp/kprofile.h](#)) - Contains code for runtime code-profiling. This is optional and may be stubbed out if left unimplemented (we won't cover profiling timers here).
- ThreadPort ([threadport.cpp/threadport.h](#)) - The meat-and-potatoes of the port code lives here. This class contains architecture/part-specific code used to initialize threads, implement critical-sections, perform context-switching, and start the kernel. Most of the time spent in this article focuses on the code found here.

Summarizing the above, these modules provide the following list of functionality:

```
- Thread stack initialization
- Kernel startup and first thread entry
- Context switch and SWI
- Kernel timers
- Critical Sections
.
```

The implementation of each of these pieces will be analyzed in detail in the sections that follow.

Thread Stack Initialization

Before a thread can be used, its stack must first be initialized to its default state. This default state ensures that when the thread is scheduled for the first time and its context restored, that it will cause the CPU to jump to the user's specified entry-point function.

All of the platform independent thread setup is handled by the generic kernel code. However, since every CPU architecture has its own register set, and stacks different information as part of an interrupt/exception, we have to implement this thread setup code for each platform we want the kernel to support (Combination of Architecture + Variant + Toolchain).

In the ARM Cortex-M0 architecture, the stack frame consists of the following information:

a) Exception Stack Frame

Contains the 8 registers which the ARM Cortex-M0 CPU automatically pushes to the stack when entering an exception. The following registers are included (in stack'd order):

```
[ XPSR ] <-- Highest address in context
[ PC   ]
[ LR   ]
[ R12  ]
[ R3   ]
[ R2   ]
[ R1   ]
[ R0   ]
```

XPSR – This is the CPU's status register. We need to set this to 0x01000000 (the "T" bit), which indicates that the CPU is executing in "thumb" mode. Note that ARMv6m and ARMv7m processors only run thumb2 instructions, so an exception is liable to occur if this bit ever gets cleared.

PC – Program Counter. This should be set with the initial entry point (function pointer) for that the user wishes to start executing with this thread.

LR - The base link register. Normally, this register contains the return address of the calling function, which is where the CPU jumps when a function returns. However, our threads generally don't return (and if they do, they're placed into the stop state). As a result we can leave this as 0.

The other registers in the stack frame are generic working registers, and have no special meaning, with the exception that R0 will hold the user's argument value passed into the endpoint.

b) Complimentary CPU Register Context

```
[ R11  ]
...
[ R4   ] <-- Lowest address in context
```

These are the other general-purpose CPU registers that need to be backed up/ restored on a context switch, but aren't stacked by default on a Cortex-M0 exception. If there were any additional hardware registers to back up, then we'd also have to include them in this part of the context as well.

As a result, these registers all need to be manually pushed to the stack on stack creation, and will need to be explicitly pushed and pop as part of a normal context switch.

With this default exception state in mind, the following code is used to initialize a thread's stack for a Cortex-M0.

```

void ThreadPort::InitStack(Thread *pclThread_)
{
    K_ULONG *pulStack;
    K_ULONG *pulTemp;
    K_ULONG ulAddr;
    K_USHORT i;

    // Get the entrypoint for the thread
    ulAddr = (K_ULONG)(pclThread_>m_pfEntryPoint);

    // Get the top-of-stack pointer for the thread
    pulStack = (K_ULONG*)pclThread_>m_pwStackTop;

    // Initialize the stack to all FF's to aid in stack depth checking
    pulTemp = (K_ULONG*)pclThread_>m_pwStack;
    for (i = 0; i < pclThread_>m_usStackSize / sizeof(K_ULONG); i++)
    {
        pulTemp[i] = 0xFFFFFFFF;
    }

    PUSH_TO_STACK(pulStack, 0);          // Apply one word of padding

    //-- Simulated Exception Stack Frame --
    PUSH_TO_STACK(pulStack, 0x01000000); // XSPR;set "T" bit for thumb-mode
    PUSH_TO_STACK(pulStack, ulAddr);     // PC
    PUSH_TO_STACK(pulStack, 0);          // LR
    PUSH_TO_STACK(pulStack, 0x12);
    PUSH_TO_STACK(pulStack, 0x3);
    PUSH_TO_STACK(pulStack, 0x2);
    PUSH_TO_STACK(pulStack, 0x1);
    PUSH_TO_STACK(pulStack, (K_ULONG)pclThread_>m_pvArg); // R0 = argument

    //-- Simulated Manually-Stacked Registers --
    PUSH_TO_STACK(pulStack, 0x11);
    PUSH_TO_STACK(pulStack, 0x10);
    PUSH_TO_STACK(pulStack, 0x09);
    PUSH_TO_STACK(pulStack, 0x08);
    PUSH_TO_STACK(pulStack, 0x07);
    PUSH_TO_STACK(pulStack, 0x06);
    PUSH_TO_STACK(pulStack, 0x05);
    PUSH_TO_STACK(pulStack, 0x04);
    pulStack++;

    pclThread_>m_pwStackTop = pulStack;
}

```

Kernel Startup

The same general process applies to starting the kernel on an ARM Cortex-M0 as on other platforms. Here, we initialize and start the platform specific timer and software-interrupt modules, find the first thread to run, and then jump to that first thread.

Now, to perform that last step, we have two options:

- 1) Simulate a return from an exception manually to start the first thread, or..
- 2) Use a software interrupt to trigger the first "Context Restore/Return from Interrupt"

For 1), we basically have to restore the whole stack manually, not relying on the CPU to do any of this for us. That's certainly doable, but not all Cortex parts support this (other members of the family support privileged modes, etc.). That, and the code required to do this is generally more complex due to all of the exception-state simulation. So, we will opt for the second option instead.

To implement a software to start our first thread, we will use the SVC instruction to generate an exception. From that exception, we can then restore the context from our first thread, set the CPU up to use the right "process" stack, and return-from-exception back to our first thread. We'll explore the code for that later.

But, before we can call the SVC exception, we're going to do a couple of things.

First, we're going to reset the default MSP stack pointer to its original top-of-stack value. The rationale here is that we no longer care about the data on the MSP stack, since calling the SVC instruction triggers a chain of events from which we never return. The MSP is also used by all exception-handling, so regaining a few words of stack here can be useful. We'll also enable all maskable exceptions at this point, since this code results in the kernel being started with the CPU executing the RTOS threads, at which point a user would expect interrupts to be enabled.

Note, the default stack pointer location is stored at address 0x00000000 on all ARM Cortex M0 parts. That explains the code below...

```

void ThreadPort_StartFirstThread( void )
{
    asm(
        " ldr r1, [r0] \n" // Reset the MSP to the default base address
        " msr msp, r1 \n"
        " cpsie i \n"      // Enable interrupts
        " svc 0 \n"        // Jump to SVC Call
    );
}

```

First Thread Entry

This handler has the job of taking the first thread object's stack, and restoring the default state data in a way that ensures that the thread starts executing when returning from the call.

We also keep in mind that there's an 8-byte offset from the beginning of the thread object to the location of the thread stack pointer. This offset is a result of the thread object inheriting from the linked-list node class, which has 8-bytes of data. This is stored first in the object, before the first element of the class, which is the "stack top" pointer.

The following assembly code shows how the SVC call is implemented in [Mark3](#) for the purpose of starting the first thread.

```

get_thread_stack:
    ; Get the stack pointer for the current thread
    ldr r0, g_pstCurrent
    ldr r1, [r0]
    add r1, #8
    ldr r2, [r1]          ; r2 contains the current stack-top

load_manually_placed_context_r11_r8:
    ; Handle the bottom 32-bytes of the stack frame
    ; Start with r11-r8, because only r0-r7 can be used
    ; with ldmia on CM0.
    add r2, #16
    ldmia r2!, {r4-r7}
    mov r11, r7
    mov r10, r6
    mov r9, r5
    mov r8, r4

set_psp:
    ; Since r2 is coincidentally back to where the stack pointer should be,
    ; Set the program stack pointer such that returning from the exception handler
    msr psp, r2

load_manually_placed_context_r7_r4:
    ; Get back to the bottom of the manually stacked registers and pop.
    sub r2, #32
    ldmia r2!, {r4-r7} ; Register r4-r11 are restored.

set_thread_and_privilege_modes:
    ; Also modify the control register to force use of thread mode as well
    ; For CM3 forward-compatibility, also set user mode.
    mrs r0, control
    mov r1, #0x03
    orr r0, r1
    control, r0

set_lr:
    ; Set up the link register such that on return, the code operates
    ; in thread mode using the PSP. To do this, we or 0x0D to the value stored
    ; in the lr by the exception hardware EXC_RETURN. Alternately, we could
    ; just force lr to be 0xFFFFFFFED (we know that's what we want from the
    ; hardware, anyway)
    mov r0, #0x0D
    mov r1, lr
    orr r0, r1

exit_exception:
    ; Return from the exception handler.
    ; The CPU will automatically unstack R0-R3, R12, PC, LR, and xPSR
    ; for us. If all goes well, our thread will start execution at the
    ; entrypoint, with the us-specified argument.
    bx r0

```

On ARM Cortex parts, there's dedicated hardware that's used primarily to support RTOS (or RTOS-like) functionality. This functionality includes the SysTick timer, and the PendSV Exception. SysTick is used for a tick-based kernel timer, while the PendSV exception is used for performing context switches. In reality, it's a "special SVC" call that's designed to be lower-overhead, in that it isn't mux'd with a bunch of other system or application functionality.

So how do we go about actually implementing a context switch here? There are a lot of different parts involved, but it essentially comes down to 3 steps:

1) Saving the context.

Thread's top-of-stack value is stored, all registers are stacked. We're good to go!

2) Swap threads

We swap the Scheduler's "next" thread with the "current" thread.

3) Restore Context

This is more or less identical to what we did when restoring the first context. Some operations may be optimized for data already stored in registers.

The code used to implement these steps on Cortex-M0 is presented below:

```
void PendSV_Handler(void)
{
    ASM(
        // Thread_SaveContext()
        " ldr r1, CURR_ \n"
        " ldr r1, [r1] \n "
        " mov r3, r1 \n "
        " add r3, #8 \n "

        // Grab the psp and adjust it by 32 based on extra registers we're going
        // to be manually stacking.
        " mrs r2, psp \n "
        " sub r2, #32 \n "

        // While we're here, store the new top-of-stack value
        " str r2, [r3] \n "

        // And, while r2 is at the bottom of the stack frame, stack r7-r4
        " stmia r2!, {r4-r7} \n "

        // Stack r11-r8
        " mov r7, r11 \n "
        " mov r6, r10 \n "
        " mov r5, r9 \n "
        " mov r4, r8 \n "
        " stmia r2!, {r4-r7} \n "

        // Equivalent of Thread_Swap() - performs g_pstCurrent = g_pstNext
        " ldr r1, CURR_ \n"
        " ldr r0, NEXT_ \n"
        " ldr r0, [r0] \n"
        " str r0, [r1] \n"

        // Thread_RestoreContext()
        // Get the pointer to the next thread's stack
        " add r0, #8 \n "
        " ldr r2, [r0] \n "

        // Stack pointer is in r2, start loading registers from
        // the "manually-stacked" set
        // Start with r11-r8, since these can't be accessed directly.
        " add r2, #16 \n "
        " ldmba r2!, {r4-r7} \n "
        " mov r11, r7 \n "
        " mov r10, r6 \n "
        " mov r9, r5 \n "
        " mov r8, r4 \n "
```



```

// After subbing R2 #16 manually, and #16 through ldmia, our PSP is where it
// needs to be when we return from the exception handler
" msr psp, r2 \n "

// Pop manually-stacked R4-R7
" sub r2, #32 \n "
" ldmia r2!, {r4-r7} \n "

// lr contains the proper EXC_RETURN value
// we're done with the exception, so return back to newly-chosen thread
" bx lr \n "
" nop \n "

// Must be 4-byte aligned.
" NEXT_: .word g_pstNext \n"
" CURR_: .word g_pstCurrent \n"
);
}

```

Kernel Timers

ARM Cortex-M series microcontrollers each contain a SysTick timer, which was designed to facilitate a fixed-interval RTOS timer-tick. This timer is a precise 24-bit down-count timer, run at the main CPU clock frequency, that can be programmed to trigger an exception when the timer expires. The handler for this exception can thus be used to drive software timers throughout the system on a fixed interval.

Unfortunately, this hardware is extremely simple, and does not offer the flexibility of other timer hardware commonly implemented by MCU vendors - specifically a suitable timer prescaler that can be used to generate efficient, long-counting intervals. As a result, while the "generic" port of [Mark3](#) for Cortex-M0 leverages the common SysTick timer interface, it only supports the tick-based version of the kernel's timer (note that specific Cortex-M0 ports such as the Atmel SAMD20 do have tickless timers).

Setting up a tick-based `KernelTimer` class to use the SysTick timer is, however, extremely easy, as is illustrated below:

```

void KernelTimer::Start(void)
{
    SysTick_Config(PORT_SYSTEM_FREQ / 1000); // 1kHz fixed clock...
    NVIC_EnableIRQ(SysTick_IRQn);
}

In this instance, the call to SysTick_Config() generates a 1kHz system-tick
signal, and the NVIC_EnableIRQ() call ensures that a SysTick exception is
generated for each tick. All other functions in the Cortex version of the
KernelTimer class are essentially stubbed out (see the source for more details).

```

Note that the functions used in this call are part of the ARM Cortex Microcontroller Software Interface Standard (CMSIS), and are supplied by all parts vendors selling Cortex hardware. This greatly simplifies the design of our port-code, since we can be reasonably assured that these APIs will work the same on all devices.

The handler code called when a SysTick exception occurs is basically the same as on other platforms (such as [AVR](#)), except that we explicitly clear the "exception pending" bit before returning. This is implemented in the following code:

```

\code{.cpp}
void SysTick_Handler(void)
{
    #if KERNEL_USE_TIMERS
        TimerScheduler::Process();
    #endif
    #if KERNEL_USE_QUANTUM
        Quantum::UpdateTimer();
    #endif

    // Clear the systick interrupt pending bit.
    SCB->ICSR |= SCB_ICSR_PENDSTCLR_Msk;
}

```

Critical Sections

A "critical section" is a block of code whose execution cannot be interrupted by means of context switches or an interrupt. In a traditional single-core operating system, it is typically implemented as a block of code where the interrupts are disabled - this is also the approach taken by [Mark3](#). Given that every CPU has its own means of disabling/enabling interrupts, the implementation of the critical section APIs is also non-portable.

In the Cortex-M0 port, we implement the two critical section APIs ([CS_ENTER\(\)](#) and [CS_EXIT\(\)](#)) as function-like macros containing inline assembly. All uses of these calls are called in pairs within a function and must take place at the same level-of-scope. Also, as nesting may occur (critical section within a critical section), this must be taken into account in the code.

In general, [CS_ENTER\(\)](#) performs the following tasks:

- Cache the current interrupt-enabled state within a local variable in the thread's state
- Disable interrupts
- .

Conversely, [CS_EXIT\(\)](#) performs the following tasks:

- Read the original interrupt-enabled state from the cached value
- Restore interrupts to the original value
- .

On Cortex-M series microcontrollers, the PRIMASK special register contains a single status bit which can be used to enable/disable all maskable interrupts at once. This register can be read directly to examine or modify its state. For convenience, ARMv6m provides two instructions to enable/disable interrupts

- cpsid (disable interrupts) and cpsie (enable interrupts). [Mark3](#) Implements these steps according to the following code:

```
//-----
#define CS_ENTER() \
{ \
    K_ULONG __ulRegState; \
    asm ( \
        " mrs r0, PRIMASK \n" \
        " mov %[STATUS], r0 \n" \
        " cpsid i \n" \
        : [STATUS] "=r" (__ulRegState) \
        ); \
}

//-----
#define CS_EXIT() \
asm ( \
    " mov r0, %[STATUS] \n" \
    " msr primask, r0 \n" \
    : \
    : [STATUS] "r" (__ulRegState) \
    ); \
}
```

Summary

In this section we have investigated how the main non-portable areas of the [Mark3](#) RTOS are implemented on a Cortex-M0 microcontroller. [Mark3](#) leverages all of the hardware blocks designed to enable RTOS functionality on ARM Cortex-M series microcontrollers: the SVC call provides the mechanism by which we start the kernel, the PendSV exception provides the necessary software interrupt, and the SysTick timer provides an RTOS tick. As a result, [Mark3](#) is a perfect fit for these devices - and as a result of this approach, the same RTOS port code should work with little to no modification on all ARM Cortex-M parts.

We have discussed what functionality in the RTOS is not portable, and what interfaces must be implemented in order to complete a fully-functional port. The five specific areas which are non-portable (stack initialization, kernel startup/entry, kernel timers, context switching, and critical sections) have been discussed in detail, with the platform-specific source provided as a practical reference to ARM-specific OS features, as well as [Mark3](#)'s porting infrastructure. From this example (and the accompanying source), it should be possible for an experienced developers to create a port [Mark3](#) to other microcontroller targets.

Chapter 11

Mark3C - C-language API bindings for the Mark3 Kernel.

[Mark3](#) now includes an optional additional library with C language bindings for all core kernel APIs, known as Mark3C.

This library allows applications to be written in C, while still enjoying all of the benefits of the clean, modular design of the core RTOS kernel.

The C-language Mark3C APIs map directly to their [Mark3](#) counterparts using a simple set of conventions, documented below. As a result, explicit API documentation for Mark3C is not necessary, as the functions map 1-1 to their C++ counterparts.

11.1 API Conventions

1) Static Methods:

<code><ClassName>::<MethodName>()</code>	Becomes	<code><ClassName>_<MethodName>()</code>
i.e. <code>Kernel::Start()</code>	Becomes	<code>Kernel_Start()</code>

2) Kernel Object Methods:

In short, any class instance is represented using an object handle, and is always passed into the relevant APIs as the first argument. Further, any method that returns a pointer to an object in the C++ implementation now returns a handle to that object.

<code><Object>.<MethodName>(<args>)</code>	Becomes	<code><ClassName>_<MethodhooName>(<ObjectHandle>, <args>)</code>
i.e. <code>clAppThread.Start()</code>	Becomes	<code>Thread_Start(hAppThread)</code>

3) Overloaded Methods:

a) Methods overloaded with a Timeout parameter:

<code><Object>.<MethodName>(<args>)</code>	Becomes	<code><ClassName>_Timed<MethodName>(<ObjectHandle>, <args>)</code>
i.e. <code>clSemaphore.Wait(1000)</code>	Becomes	<code>Semaphore_Wait(hSemaphore, 1000)</code>

b) Methods overloaded based on number of arguments:

<Object>.<MethodName>()	Becomes	<ClassName>_<MethodName>(<ObjectHandle>)
<Object>.<MethodName>(<arg1>)	Becomes	<ClassName>_<MethodName>1(<ObjectHandle>, <arg1>)
<Object>.<MethodName>(<arg1>, <arg2>)	Becomes	<ClassName>_<MethodName>2(<ObjectHandle>, <arg1>, <arg2>)
<ClassName>::<MethodName>()	Becomes	<ClassName>_<MethodName>(<ObjectHandle>)
<ClassName>::<MethodName>(<arg1>)	Becomes	<ClassName>_<MethodName>1(<ObjectHandle>, <arg1>)
<ClassName>::<MethodName>(<arg1>, <arg2>)	Becomes	<ClassName>_<MethodName>2(<ObjectHandle>, <arg1>, <arg2>)

c) Methods overloaded base on parameter types:

<Object>.<MethodName>(<arg type_a>)	Becomes	<ClassName>_<MethodName><type_a>(<ObjectHandle>, <arg type_a>)
<Object>.<MethodName>(<arg type_b>)	Becomes	<ClassName>_<MethodName><type_b>(<ObjectHandle>, <arg type_b>)
<ClassName>::<MethodName>(<arg type_a>)	Becomes	<ClassName>_<MethodName><type_a>(<arg type_a>)
<ClassName>::<MethodName>(<arg type_b>)	Becomes	<ClassName>_<MethodName><type_b>(<arg type_b>)

d) Allocate-once memory allocation APIs

AutoAlloc::New<ObjectName>	Becomes	Alloc_<ObjectName>
AutoAlloc::Allocate(uint16_t u16Size_)	Becomes	AutoAlloc(uint16_t u16Size_)

11.2 Allocating Objects

Aside from the API name translations, the object allocation scheme is the major different between Mark3C and [Mark3](#). Instead of instantiating objects of the various kernel types, kernel objects must be declared using Declaration macros, which serve the purpose of reserving memory for the kernel object, and provide an opaque handle to that object memory. This is the case for statically-allocated objects, and objects allocated on the stack.

Example: Declaring a thread

```
#include "mark3c.h"

// Statically-allocated
DECLARE_THREAD(hMyThread1);
...

// On stack
int main(void)
{
    DECLARE_THREAD(hMyThread2);
    ...
}
```

Where:

hMyThread1 - is a handle to a statically-allocated thread
hMyThread2 - is a handle to a thread allocated from the main stack.

Alternatively, the AutoAlloc APIs can be used to dynamically allocate objects, as demonstrated in the following example.

```
void Allocate_Example(void)
{
    Thread_t hMyThread = AutoAlloc_Thread();

    Thread_Init(hMyThread, awMyStack, sizeof(awMyStack), 1, MyFunction, 0);
}
```

Note that the relevant kernel-object Init() function *must* be called prior to using any kernel object, whether or not they have been allocated statically, or dynamically.

Chapter 12

Release Notes

12.1 R7 Release

- Re-focusing project on kernel integrating with 3rd party code instead of 1st party middleware
- New: Refactored codebase to C++14 standard
- New: Moved non-kernel code, drivers, libs, and BSPs to separate repos from kernel
- New: Modular repository-based structure, managed via Android's Repo tool
- New: ConditionVariable kernel API
- New: ReaderWriterLock kernel API
- New: Slab memory allocator
- New: Bitmap object-allocator
- New: AutoAlloc redirects to user-defined allocators
- New: Global new() and delete() overrides redirect to AutoAlloc APIs
- Updated Mark3c for new APIs
- Moved driver layer out of the kernel

12.2 R6 Release

- New: Replace recursive-make build system with CMake and Ninja
- New: Transitioned version control to Git from Subversion.
- New: Socket library, implementing named "domain-socket" style IPC
- New: State Machine framework library
- New: Software I2C library completed, with demo app
- New: Kernel Timer loop can optionally be run within its own thread instead of a nested interrupt
- New: UART drivers are all now abstracted through UartDriver base class for portability
- Experimental: Process library, allowing for the creation of resource-isolated processes
- Removed: Bare-metal support for Atmel SAMD20 (generic port still works)
- Cleanup all compiler warnings on atmega328p
- Various Bugfixes and optimizations
- Various Script changes related to automating the build + release process

12.3 R5 Release

- New: Shell library for creating responsive CLIs for embedded applications (M3Shell)
- New: Stream library for creating thread-safe buffered streams (streamer)
- New: Blocking UART implementation for AVR (drvUARTplus)
- New: "Extended context" kernel feature, which is used to implement thread-local storage
- New: "Extra Checks" kernel feature, which enforces safe API usage under pain of Kernel Panic
- New: Realtime clock library
- New: Example application + bsp for the open-hardware Mark3no development board (mark3no)
- New: Kernel objects descoped/destroyed while still in active use will now cause kernel panic
- New: Kernel callouts for thread creation/destruction/context switching, used for time tracking
- New: Simple power management class
- New: WIP software-based I2C + SPI drivers
- Optimized thread scheduling via target-optimized "count-leading-zero" macros
- Expanded memutil library
- Various optimizations of ARM Cortex-M assembly code
- Various bugfixes to Timer code
- Improved stack overflow checking + warning (stack guard kernel feature)
- AVR bootloader now supports targets with more than 64K of flash
- Moved some port configuration out of platform.mak into header files in the kernel port code
- The usual minor bugfixes and "gentle refactoring"

12.4 R4 Release

- New: C-language bindings for [Mark3](#) kernel (mark3c library)
- New: Support for ARM Cortex-M3 and Cortex-M4 (floating point) targets
- New: Support for Atmel AVR atmega2560 and arduino pro mega
- New: Full-featured, lightweight heap implementation
- New: Mailbox IPC class
- New: Notification object class
- New: lightweight tracelogger/instrumentation implementation (buffalogger), with sample parser
- New: High-performance AVR Software UART implementation
- New: Allocate-once "AutoAlloc" memory allocator
- New: Fixed-time blocking/unblocking operations added to ThreadList/Blocking class
- Placement-new supported for all kernel objects
- Scheduler now supports up to 1024 levels of thread priority, up from 8 (configurable at build-time)

- Kernel now uses `stdint.h` types for standard integers (instead of `K_CHAR`, `K_ULONG`, etc.)
- Greatly expanded documentation, with many new examples covering all key kernel features
- Expanded unit test coverage on AVR
- Updated build system and scripts for easier kernel configuration
- Updated builds to only attempt to build tests for supported platforms

12.5 R3 Release

- New: Added support for MSP430 microcontrollers
- New: Added Kernel Idle-Function hook to eliminate the need for a dedicated idle-thread (where supported)
- New: Support for kernel-aware simulation and testing via fIAVR AVR simulator
- Updated AVR driver selection
- General bugfixes and maintenance
- Expanded documentation and test coverage

12.6 R2

- Experimental release, using a "kernel transaction queue" for serializing kernel calls
- Works as a proof-of-concept, but abandoned due to overhead of the transaction mechanism in the general case.

12.7 R1 - 2nd Release Candidate

- New: Added support for ARM Cortex-M0 targets
- New: Added support for various AVR targets
- New: Timers now support a "tolerance" parameter for grouping timers with close expiry times
- Expanded scripts and automation used in build/test
- Updated and expanded graphics APIs
- Large number of bugfixes

12.8 R1 - 1st Release Candidate

- Initial release, with support for AVR microcontrollers

Chapter 13

Code Size Profiling

The following report details the size of each module compiled into the kernel.

The size of each component is dependent on the flags specified in [mark3cfg.h](#) at compile time. Note that these sizes represent the maximum size of each module before dead code elimination and any additional link-time optimization, and represent the maximum possible size that any module can take.

The results below are for profiling on ARM Cortex-M3 qemu_lm3s6965evb-based targets using gcc. Results are not necessarily indicative of relative or absolute performance on other platforms or toolchains.

13.1 Information

Date Profiled: Mon Jul 2 20:49:30 EDT 2018

13.2 Compiler Version

arm-none-eabi-gcc (15:6.3.1+svn253039-1build1) 6.3.1 20170620 Copyright (C) 2016 Free Software Foundation, Inc. This is free software; see the source for copying conditions. There is NO warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.

13.3 Profiling Results

[Mark3](#) Module Size Report:

```
- Atomic Operations..... : 336 Bytes
- Allocate-once Heap..... : 724 Bytes
- Synchronization Objects - Base Class..... : 82 Bytes
- Condition Variables (Synchronization Object).... : 216 Bytes
- Synchronization Object - Event Flag..... : 494 Bytes
- Mark3 Kernel Base Class..... : 151 Bytes
- ARM Cortex-M3 - Kernel Aware Simulation Support..... : 266 Bytes
- Semaphore (Synchronization Object)..... : 366 Bytes
- Fundamental Kernel Linked-List Classes..... : 200 Bytes
- RAII Locking Support based on Mark3 Mutex class. : 58 Bytes
- Mailbox IPC Support..... : 572 Bytes
- Message-based IPC..... : 194 Bytes
- Mutex (Synchronization Object)..... : 464 Bytes
- Notification Blocking Object..... : 396 Bytes
```

- 2D Priority Map Object - Scheduler..... : 58 Bytes
- Performance-profiling timers..... : 240 Bytes
- Round-Robin Scheduling Support..... : 289 Bytes
- Reader-writer Locks (Synchronization Object).... : 158 Bytes
- Thread Scheduling..... : 414 Bytes
- Thread Implementation..... : 1277 Bytes
- Fundamental Kernel Thread-list Data Structures.. : 150 Bytes
- Software Timer Kernel Object..... : 236 Bytes
- Software Timer Management..... : 240 Bytes
- Runtime Kernel [Trace](#) Implementation..... : 0 Bytes
- ARM Cortex-M3 - Profiling Timer Implementation..... : 119 Bytes
- ARM Cortex-M3 - Kernel Interrupt Implementation..... : 66 Bytes
- ARM Cortex-M3 - Kernel Timer Implementation..... : 152 Bytes
- ARM Cortex-M3 - Basic Threading Support..... : 344 Bytes

[Mark3](#) Kernel Size Summary:

- Kernel : 2332 Bytes
- Synchronization Objects : 1950 Bytes
- Port : 1915 Bytes
- Features : 2065 Bytes
- Total Size : 8262 Bytes

Chapter 14

Namespace Index

14.1 Namespace List

Here is a list of all documented namespaces with brief descriptions:

Mark3	Class providing the software-interrupt required for context-switching in the kernel	89
Mark3::Atomic	The Atomic class	92
Mark3::KernelAware	The KernelAware class	94

Chapter 15

Hierarchical Index

15.1 Class Hierarchy

This inheritance list is sorted roughly, but not completely, alphabetically:

Mark3::BlockingObject	99
Mark3::EventFlag	108
Mark3::Mutex	148
Mark3::Notify	152
Mark3::Semaphore	170
Mark3::ConditionVariable	104
Mark3::FakeThread_t	113
Mark3::Kernel	114
Mark3::KernelTimer	120
Mark3::LinkList	125
Mark3::CircularLinkList	102
Mark3::ThreadList	190
Mark3::DoubleLinkList	106
Mark3::TimerList	204
Mark3::LinkListNode	127
Mark3::Message	140
Mark3::Thread	174
Mark3::Timer	196
Mark3::LockGuard	129
Mark3::Mailbox	130
Mark3::MessagePool	143
Mark3::MessageQueue	145
Mark3::PriorityMap	154
Mark3::ProfileTimer	156
Mark3::Quantum	159
Mark3::ReaderWriterLock	161
Mark3::Scheduler	164
Mark3::ThreadPort	194
Mark3::TimerScheduler	206

Chapter 16

Class Index

16.1 Class List

Here are the classes, structs, unions and interfaces with brief descriptions:

Mark3::BlockingObject	Class implementing thread-blocking primitives	99
Mark3::CircularLinkedList	Circular-linked-list data type, inherited from the base LinkedList type	102
Mark3::ConditionVariable	The ConditionVariable class This class implements a condition variable	104
Mark3::DoubleLinkedList	Doubly-linked-list data type, inherited from the base LinkedList type	106
Mark3::EventFlag	Blocking object, similar to a semaphore or mutex, commonly used for synchronizing thread execution based on events occurring within the system	108
Mark3::FakeThread_t	If the kernel is set up to use an idle function instead of an idle thread, we use a placeholder data structure to "simulate" the effect of having an idle thread in the system	113
Mark3::Kernel	Class that encapsulates all of the kernel startup functions	114
Mark3::KernelTimer	Hardware timer interface, used by all scheduling/timer subsystems	120
Mark3::LinkedList	Abstract-data-type from which all other linked-lists are derived	125
Mark3::LinkedListNode	Basic linked-list node data structure	127
Mark3::LockGuard	129
Mark3::Mailbox	Implements an IPC mechanism based on envelopes containing data of a fixed size (configured at initialization) that reside within a buffer of memory provided by the user	130
Mark3::Message	Class to provide message-based IPC services in the kernel	140
Mark3::MessagePool	Implements a list of message objects	143
Mark3::MessageQueue	List of messages, used as the channel for sending and receiving messages between threads	145
Mark3::Mutex	Mutual-exclusion locks, based on BlockingObject	148

Mark3::Notify	Blocking object type, that allows one or more threads to wait for an event to occur before resuming operation	152
Mark3::PriorityMap	The PriorityMap class	154
Mark3::ProfileTimer	Profiling timer	156
Mark3::Quantum	Static-class used to implement Thread quantum functionality, which is a key part of round-robin scheduling	159
Mark3::ReaderWriterLock	The ReaderWriterLock class This class implements an object that marshalls access to a resource based on the intended usage of the resource	161
Mark3::Scheduler	Priority-based round-robin Thread scheduling, using ThreadLists for housekeeping	164
Mark3::Semaphore	Binary & Counting semaphores, based on BlockingObject base class	170
Mark3::Thread	Object providing fundamental multitasking support in the kernel	174
Mark3::ThreadList	This class is used for building thread-management facilities, such as schedulers, and blocking objects	190
Mark3::ThreadPort	Class defining the architecture specific functions required by the kernel	194
Mark3::Timer	Kernel-managed software timers	196
Mark3::TimerList	TimerList class - a doubly-linked-list of timer objects	204
Mark3::TimerScheduler	"Static" Class used to interface a global TimerList with the rest of the kernel	206

Chapter 17

File Index

17.1 File List

Here is a list of all documented files with brief descriptions:

/home/moslevin/projects/github/m3-repo/kernel/libs/mark3c/src/ mark3c.cpp	??
/home/moslevin/projects/github/m3-repo/kernel/libs/mark3c/src/public/ fake_types.h	
C-struct definitions that mirror	209
/home/moslevin/projects/github/m3-repo/kernel/libs/mark3c/src/public/ mark3c.h	
Implementation of C-language wrappers for the Mark3 kernel	212
/home/moslevin/projects/github/m3-repo/kernel/src/ atomic.cpp	
Basic Atomic Operations	247
/home/moslevin/projects/github/m3-repo/kernel/src/ autoalloc.cpp	
Automatic memory allocation for kernel objects	249
/home/moslevin/projects/github/m3-repo/kernel/src/ blocking.cpp	
Implementation of base class for blocking objects	253
/home/moslevin/projects/github/m3-repo/kernel/src/ condvar.cpp	
Condition Variable implementation	255
/home/moslevin/projects/github/m3-repo/kernel/src/ eventflag.cpp	
Event Flag Blocking Object/IPC-Object implementation	256
/home/moslevin/projects/github/m3-repo/kernel/src/ kernel.cpp	
Kernel initialization and startup code	261
/home/moslevin/projects/github/m3-repo/kernel/src/ kernelaware.cpp	
Kernel aware simulation support	263
/home/moslevin/projects/github/m3-repo/kernel/src/ ksemaphore.cpp	
Semaphore Blocking-Object Implemenation	265
/home/moslevin/projects/github/m3-repo/kernel/src/ ll.cpp	
Core Linked-List implementation, from which all kernel objects are derived	269
/home/moslevin/projects/github/m3-repo/kernel/src/ lockguard.cpp	
Mutex RAIL helper class	271
/home/moslevin/projects/github/m3-repo/kernel/src/ mailbox.cpp	
Mailbox + Envelope IPC mechanism	272
/home/moslevin/projects/github/m3-repo/kernel/src/ message.cpp	
Inter-thread communications via message passing	276
/home/moslevin/projects/github/m3-repo/kernel/src/ mutex.cpp	
Mutual-exclusion object	278
/home/moslevin/projects/github/m3-repo/kernel/src/ notify.cpp	
Lightweight thread notification - blocking object	282
/home/moslevin/projects/github/m3-repo/kernel/src/ priomap.cpp	
Priority map data structure	285

/home/moslevin/projects/github/m3-repo/kernel/src/ profile.cpp	
Code profiling utilities	287
/home/moslevin/projects/github/m3-repo/kernel/src/ quantum.cpp	
Thread Quantum Implementation for Round-Robin Scheduling	353
/home/moslevin/projects/github/m3-repo/kernel/src/ readerwriter.cpp	
Reader-writer lock implementation	355
/home/moslevin/projects/github/m3-repo/kernel/src/ scheduler.cpp	
Strict-Priority + Round-Robin thread scheduler implementation	357
/home/moslevin/projects/github/m3-repo/kernel/src/ thread.cpp	
Platform-Independent thread class Definition	359
/home/moslevin/projects/github/m3-repo/kernel/src/ threadlist.cpp	
Thread linked-list definitions	366
/home/moslevin/projects/github/m3-repo/kernel/src/ timer.cpp	
Timer implementations	368
/home/moslevin/projects/github/m3-repo/kernel/src/ timerlist.cpp	
Implements timer list processing algorithms, responsible for all timer tick and expiry logic . . .	371
/home/moslevin/projects/github/m3-repo/kernel/src/ tracebuffer.cpp	
Kernel trace buffer class definition	375
/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/ kernelprofile.cpp	
ATMega328p Profiling timer implementation	228
/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/ kernelswi.cpp	
Kernel Software interrupt implementation for ATMega328p	230
/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/ kerneltimer.cpp	
Kernel Timer Implementation for ATMega328p	231
/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/ threadport.cpp	
ATMega1284p Multithreading	244
/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/ kernelprofile.h	
Profiling timer hardware interface	234
/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/ kernelswi.h	
Kernel Software interrupt declarations	235
/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/ kerneltimer.h	
Kernel Timer Class declaration	236
/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/ portcfg.h	
Mark3 Port Configuration	237
/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/ threadport.h	
ATMega328p Multithreading support	240
/home/moslevin/projects/github/m3-repo/kernel/src/public/ atomic.h	
Basic Atomic Operations	289
/home/moslevin/projects/github/m3-repo/kernel/src/public/ autoalloc.h	
Automatic memory allocation for kernel objects	290
/home/moslevin/projects/github/m3-repo/kernel/src/public/ blocking.h	
Blocking object base class declarations	293
/home/moslevin/projects/github/m3-repo/kernel/src/public/ buffalogger.h	
Super-efficient, super-secure logging routines	294
/home/moslevin/projects/github/m3-repo/kernel/src/public/ condvar.h	
Condition Variable implementation	295
/home/moslevin/projects/github/m3-repo/kernel/src/public/ dbg_file_list.h	??
/home/moslevin/projects/github/m3-repo/kernel/src/public/ eventflag.h	
Event Flag Blocking Object/IPC-Object definition	296
/home/moslevin/projects/github/m3-repo/kernel/src/public/ kernel.h	
Kernel initialization and startup class	298
/home/moslevin/projects/github/m3-repo/kernel/src/public/ kernelaware.h	
Kernel aware simulation support	299
/home/moslevin/projects/github/m3-repo/kernel/src/public/ kerneldebug.h	
Macros and functions used for assertions, kernel traces, etc	301
/home/moslevin/projects/github/m3-repo/kernel/src/public/ kerneltypes.h	
Basic data type primitives used throughout the OS	307

/home/moslevin/projects/github/m3-repo/kernel/src/public/ ksemaphore.h	
Semaphore Blocking Object class declarations	309
/home/moslevin/projects/github/m3-repo/kernel/src/public/ ll.h	
Core linked-list declarations, used by all kernel list types	310
/home/moslevin/projects/github/m3-repo/kernel/src/public/ lockguard.h	
Mutex RAIL helper class	312
/home/moslevin/projects/github/m3-repo/kernel/src/public/ mailbox.h	
Mailbox + Envelope IPC Mechanism	313
/home/moslevin/projects/github/m3-repo/kernel/src/public/ manual.h	
/brief Ascii-format documentation, used by doxygen to create various printable and viewable forms	316
/home/moslevin/projects/github/m3-repo/kernel/src/public/ mark3.h	
Single include file given to users of the Mark3 Kernel API	316
/home/moslevin/projects/github/m3-repo/kernel/src/public/ mark3cfg.h	
Mark3 Kernel Configuration	318
/home/moslevin/projects/github/m3-repo/kernel/src/public/ message.h	
Inter-thread communication via message-passing	328
/home/moslevin/projects/github/m3-repo/kernel/src/public/ mutex.h	
Mutual exclusion class declaration	330
/home/moslevin/projects/github/m3-repo/kernel/src/public/ notify.h	
Lightweight thread notification - blocking object	332
/home/moslevin/projects/github/m3-repo/kernel/src/public/ paniccodes.h	
Defines the reason codes thrown when a kernel panic occurs	334
/home/moslevin/projects/github/m3-repo/kernel/src/public/ priomap.h	
Priority map data structure	334
/home/moslevin/projects/github/m3-repo/kernel/src/public/ profile.h	
High-precision profiling timers	336
/home/moslevin/projects/github/m3-repo/kernel/src/public/ profiling_results.h	??
/home/moslevin/projects/github/m3-repo/kernel/src/public/ quantum.h	
Thread Quantum declarations for Round-Robin Scheduling	337
/home/moslevin/projects/github/m3-repo/kernel/src/public/ readerwriter.h	
Reader-Writer lock implementation	339
/home/moslevin/projects/github/m3-repo/kernel/src/public/ scheduler.h	
Thread scheduler function declarations	340
/home/moslevin/projects/github/m3-repo/kernel/src/public/ sizeprofile.h	??
/home/moslevin/projects/github/m3-repo/kernel/src/public/ thread.h	
Platform independent thread class declarations	342
/home/moslevin/projects/github/m3-repo/kernel/src/public/ threadlist.h	
Thread linked-list declarations	345
/home/moslevin/projects/github/m3-repo/kernel/src/public/ timer.h	
Timer object declarations	346
/home/moslevin/projects/github/m3-repo/kernel/src/public/ timerlist.h	
Timer list declarations	349
/home/moslevin/projects/github/m3-repo/kernel/src/public/ timerscheduler.h	
Timer scheduler declarations	351
/home/moslevin/projects/github/m3-repo/kernel/src/public/ tracebuffer.h	
Kernel trace buffer class declaration	352

Chapter 18

Namespace Documentation

18.1 Mark3 Namespace Reference

Class providing the software-interrupt required for context-switching in the kernel.

Namespaces

- [Atomic](#)
The [Atomic](#) class.
- [KernelAware](#)
The [KernelAware](#) class.

Classes

- class [BlockingObject](#)
Class implementing thread-blocking primitives.
- class [CircularLinkedList](#)
Circular-linked-list data type, inherited from the base [LinkedList](#) type.
- class [ConditionVariable](#)
The [ConditionVariable](#) class This class implements a condition variable.
- class [DoubleLinkedList](#)
Doubly-linked-list data type, inherited from the base [LinkedList](#) type.
- class [EventFlag](#)
The [EventFlag](#) class is a blocking object, similar to a semaphore or mutex, commonly used for synchronizing thread execution based on events occurring within the system.
- struct [FakeThread_t](#)
If the kernel is set up to use an idle function instead of an idle thread, we use a placeholder data structure to "simulate" the effect of having an idle thread in the system.
- class [Kernel](#)
Class that encapsulates all of the kernel startup functions.
- class [KernelTimer](#)
Hardware timer interface, used by all scheduling/timer subsystems.
- class [LinkedList](#)
Abstract-data-type from which all other linked-lists are derived.

- class [LinkListNode](#)
Basic linked-list node data structure.
- class [LockGuard](#)
- class [Mailbox](#)
The [Mailbox](#) class implements an IPC mechanism based on envelopes containing data of a fixed size (configured at initialization) that reside within a buffer of memory provided by the user.
- class [Message](#)
Class to provide message-based IPC services in the kernel.
- class [MessagePool](#)
Implements a list of message objects.
- class [MessageQueue](#)
List of messages, used as the channel for sending and receiving messages between threads.
- class [Mutex](#)
Mutual-exclusion locks, based on [BlockingObject](#).
- class [Notify](#)
The [Notify](#) class is a blocking object type, that allows one or more threads to wait for an event to occur before resuming operation.
- class [PriorityMap](#)
The [PriorityMap](#) class.
- class [ProfileTimer](#)
Profiling timer.
- class [Quantum](#)
Static-class used to implement [Thread](#) quantum functionality, which is a key part of round-robin scheduling.
- class [ReaderWriterLock](#)
The [ReaderWriterLock](#) class This class implements an object that marshalls access to a resource based on the intended usage of the resource.
- class [Scheduler](#)
Priority-based round-robin [Thread](#) scheduling, using ThreadLists for housekeeping.
- class [Semaphore](#)
Binary & Counting semaphores, based on [BlockingObject](#) base class.
- class [Thread](#)
Object providing fundamental multitasking support in the kernel.
- class [ThreadList](#)
This class is used for building thread-management facilities, such as schedulers, and blocking objects.
- class [ThreadPort](#)
Class defining the architecture specific functions required by the kernel.
- class [Timer](#)
Kernel-managed software timers.
- class [TimerList](#)
[TimerList](#) class - a doubly-linked-list of timer objects.
- class [TimerScheduler](#)
"Static" Class used to interface a global [TimerList](#) with the rest of the kernel.

Typedefs

- using [PanicFunc](#) = void (*)(uint16_t u16PanicCode_)
Function pointer type used to implement kernel-panic handlers.
- using [IdleFunc](#) = void (*)()
Function pointer type used to implement the idle function, where support for an idle function (as opposed to an idle thread) exists.
- using [ThreadEntryFunc](#) = void (*)(void *pvArg_)
Function pointer type used for thread entryptoint functions.
- using [TimerCallback](#) = void (*)([Thread](#) *pclOwner_, void *pvData_)
This type defines the callback function type for timer events.

Enumerations

- enum [EventFlagOperation](#) : uint8_t {
[EventFlagOperation::All_Set](#) = 0, [EventFlagOperation::Any_Set](#), [EventFlagOperation::All_Clear](#), [EventFlagOperation::Any_Clear](#),
[EventFlagOperation::Pending_Unblock](#) }

This enumeration describes the different operations supported by the event flag blocking object.

- enum [ThreadState](#) : uint8_t

Enumeration representing the different states a thread can exist in.

Functions

- [ISR](#) (INT2_vect) __attribute__((signal))
ISR(INT2_vect) SWI using INT2 - used to trigger a context switch.

18.1.1 Detailed Description

Class providing the software-interrupt required for context-switching in the kernel.

18.1.2 Typedef Documentation

18.1.2.1 TimerCallback

```
using Mark3::TimerCallback = typedef void (*)(Thread* pclOwner_, void* pvData_)
```

This type defines the callback function type for timer events.

Since these are called from an interrupt context, they do not operate from within a thread or object context directly – as a result, the context must be manually passed into the calls.

pclOwner_ is a pointer to the thread that owns the timer pvData_ is a pointer to some data or object that needs to know about the timer's expiry from within the timer interrupt context.

Definition at line 94 of file [timer.h](#).

18.1.3 Enumeration Type Documentation

18.1.3.1 EventFlagOperation

```
enum Mark3::EventFlagOperation : uint8_t [strong]
```

This enumeration describes the different operations supported by the event flag blocking object.

Enumerator

All_Set	Block until all bits in the specified bitmask are set.
Any_Set	Block until any bits in the specified bitmask are set.
All_Clear	Block until all bits in the specified bitmask are cleared.
Any_Clear	Block until any bits in the specified bitmask are cleared.
Pending_Unblock	Special code. Not used by user

Definition at line 50 of file [kerneltypes.h](#).

18.2 Mark3::Atomic Namespace Reference

The [Atomic](#) class.

Functions

- [uint8_t Set](#) (uint8_t *pu8Source_, uint8_t u8Val_)
Set Set a variable to a given value in an uninterruptable operation.
- [uint8_t Add](#) (uint8_t *pu8Source_, uint8_t u8Val_)
Add Add a value to a variable in an uninterruptable operation.
- [uint8_t Sub](#) (uint8_t *pu8Source_, uint8_t u8Val_)
Sub Subtract a value from a variable in an uninterruptable operation.
- [bool TestAndSet](#) (bool *pbLock)
TestAndSet Test to see if a variable is set, and set it if is not already set.

18.2.1 Detailed Description

The [Atomic](#) class.

This utility class provides primitives for atomic operations - that is, operations that are guaranteed to execute uninterrupted. Basic atomic primitives provided here include Set/Add/Delete for 8, 16, and 32-bit integer types, as well as an atomic test-and-set.

18.2.2 Function Documentation

18.2.2.1 Add()

```
uint8_t Mark3::Atomic::Add (
    uint8_t * pu8Source_,
    uint8_t u8Val_ )
```

Add Add a value to a variable in an uninterruptable operation.

Parameters

<i>pu8↵ Source_</i>	Pointer to a variable
<i>u8Val_</i>	Value to add to the variable

Returns

Previously-held value in *pu8Source_*

18.2.2.2 Set()

```
uint8_t Mark3::Atomic::Set (  
    uint8_t * pu8Source_,  
    uint8_t u8Val_ )
```

Set Set a variable to a given value in an uninterruptable operation.

Parameters

<i>pu8↵ Source_</i>	Pointer to a variable to set the value of
<i>u8Val_</i>	New value to set in the variable

Returns

Previously-set value

18.2.2.3 Sub()

```
uint8_t Mark3::Atomic::Sub (  
    uint8_t * pu8Source_,  
    uint8_t u8Val_ )
```

Sub Subtract a value from a variable in an uninterruptable operation.

Parameters

<i>pu8↵ Source_</i>	Pointer to a variable
<i>u8Val_</i>	Value to subtract from the variable

Returns

Previously-held value in pu8Source_

18.2.2.4 TestAndSet()

```
bool Mark3::Atomic::TestAndSet (
    bool * pbLock )
```

TestAndSet Test to see if a variable is set, and set it if is not already set.

This is an uninterruptable operation.

If the value is false, set the variable to true, and return the previously-held value.

If the value is already true, return true.

Parameters

<i>pbLock</i>	Pointer to a value to test against. This will always be set to "true" at the end of a call to TestAndSet.
---------------	---

Returns

true - Lock value was "true" on entry, false - Lock was set

18.3 Mark3::KernelAware Namespace Reference

The [KernelAware](#) class.

Functions

- void [ProfileInit](#) (const char *szStr_)
ProfileInit.
- void [ProfileStart](#) (void)
ProfileStart.
- void [ProfileStop](#) (void)
ProfileStop.
- void [ProfileReport](#) (void)
ProfileReport.
- void [ExitSimulator](#) (void)
ExitSimulator.
- void [Print](#) (const char *szStr_)
Print.
- void [Trace](#) (uint16_t u16File_, uint16_t u16Line_)
Trace.

- void [Trace](#) (uint16_t u16File_, uint16_t u16Line_, uint16_t u16Arg1_)
Trace.
- void [Trace](#) (uint16_t u16File_, uint16_t u16Line_, uint16_t u16Arg1_, uint16_t u16Arg2_)
Trace.
- bool [IsSimulatorAware](#) (void)
IsSimulatorAware.

18.3.1 Detailed Description

The [KernelAware](#) class.

This class contains functions that are used to trigger kernel-aware functionality within a supported simulation environment (i.e. fIAVR).

These static methods operate on a singleton set of global variables, which are monitored for changes from within the simulator. The simulator hooks into these variables by looking for the correctly-named symbols in an elf-formatted binary being run and registering callbacks that are called whenever the variables are changed. On each change of the command variable, the kernel-aware data is analyzed and interpreted appropriately.

If these methods are run in an unsupported simulator or on actual hardware the commands generally have no effect (except for the exit-on-reset command, which will result in a jump-to-0 reset).

18.3.2 Function Documentation

18.3.2.1 ExitSimulator()

```
void Mark3::KernelAware::ExitSimulator (
    void )
```

[ExitSimulator.](#)

Instruct the kernel-aware simulator to terminate (destroying the virtual CPU).

18.3.2.2 IsSimulatorAware()

```
bool Mark3::KernelAware::IsSimulatorAware (
    void )
```

[IsSimulatorAware.](#)

use this function to determine whether or not the code is running on a simulator that is aware of the kernel.

Returns

true - the application is being run in a kernel-aware simulator. false - otherwise.

18.3.2.3 Print()

```
void Mark3::KernelAware::Print (
    const char * szStr_ )
```

[Print.](#)

Instruct the kernel-aware simulator to print a char string

Parameters

<i>sz↔</i>	
<i>Str_</i>	

Examples:

[lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_idle_↔
function/main.cpp](#), [lab3_round_robin/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab5_mutexes/main.cpp](#), [lab6_↔
_timers/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

18.3.2.4 ProfileInit()

```
void Mark3::KernelAware::ProfileInit (
    const char * szStr_ )
```

ProfileInit.

Initializes the kernel-aware profiler. This function instructs the kernel-aware simulator to reset its accounting variables, and prepare to start counting profiling data tagged to the given string. How this is handled is the responsibility of the simulator.

Parameters

<i>sz↔</i>	String to use as a tag for the profiling session.
<i>Str_</i>	

18.3.2.5 ProfileReport()

```
void Mark3::KernelAware::ProfileReport (
    void )
```

ProfileReport.

Instruct the kernel-aware simulator to print a report for its current profiling data.

18.3.2.6 ProfileStart()

```
void Mark3::KernelAware::ProfileStart (
    void )
```

ProfileStart.

Instruct the kernel-aware simulator to begin counting cycles towards the current profiling counter.

18.3.2.7 ProfileStop()

```
void Mark3::KernelAware::ProfileStop (
    void )
```

ProfileStop.

Instruct the kernel-aware simulator to end counting cycles relative to the current profiling counter's iteration.

18.3.2.8 Trace() [1/3]

```
void Mark3::KernelAware::Trace (
    uint16_t u16File_,
    uint16_t u16Line_ )
```

Trace.

Insert a kernel trace statement into the kernel-aware simulator's debug data stream.

Parameters

<i>u16File_</i>	16-bit code representing the file
<i>u16Line_</i>	16-bit code representing the line in the file

Examples:

[lab11_mailboxes/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

18.3.2.9 Trace() [2/3]

```
void Mark3::KernelAware::Trace (
    uint16_t u16File_,
    uint16_t u16Line_,
    uint16_t u16Arg1_ )
```

Trace.

Insert a kernel trace statement into the kernel-aware simulator's debug data stream.

Parameters

<i>u16File_</i>	16-bit code representing the file
<i>u16Line_</i>	16-bit code representing the line in the file
<i>u16Arg1_</i>	16-bit argument to the format string.

18.3.2.10 Trace() [3/3]

```
void Mark3::KernelAware::Trace (
    uint16_t u16File_,
    uint16_t u16Line_,
    uint16_t u16Arg1_,
    uint16_t u16Arg2_ )
```

Trace.

Insert a kernel trace statement into the kernel-aware simulator's debug data stream.

Parameters

<i>u16File_</i> ↔ —	16-bit code representing the file
<i>u16_</i> ↔ <i>Line_</i>	16-bit code representing the line in the file
<i>u16_</i> ↔ <i>Arg1_</i>	16-bit argument to the format string.
<i>u16_</i> ↔ <i>Arg2_</i>	16-bit argument to the format string.

Chapter 19

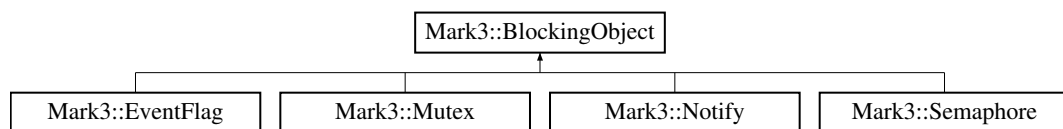
Class Documentation

19.1 Mark3::BlockingObject Class Reference

Class implementing thread-blocking primitives.

```
#include <blocking.h>
```

Inheritance diagram for Mark3::BlockingObject:



Protected Member Functions

- void **Block** (Thread *pciThread_)
Block.
- void **BlockPriority** (Thread *pciThread_)
BlockPriority.
- void **UnBlock** (Thread *pciThread_)
UnBlock.
- void **SetInitialized** (void)
SetInitialized.
- bool **IsInitialized** (void)
IsInitialized.

Protected Attributes

- ThreadList **m_clBlockList**
ThreadList which is used to hold the list of threads blocked on a given object.
- uint8_t **m_u8Initialized**
Token used to check whether or not the object has been initialized prior to use.

19.1.1 Detailed Description

Class implementing thread-blocking primitives.

used for implementing things like semaphores, mutexes, message queues, or anything else that could cause a thread to suspend execution on some external stimulus.

Definition at line 69 of file [blocking.h](#).

19.1.2 Member Function Documentation

19.1.2.1 Block()

```
void Mark3::BlockingObject::Block (  
    Thread * pclThread_ ) [protected]
```

Block.

Blocks a thread on this object. This is the fundamental operation performed by any sort of blocking operation in the operating system. All semaphores/mutexes/sleeping/messaging/etc ends up going through the blocking code at some point as part of the code that manages a transition from an "active" or "waiting" thread to a "blocked" thread.

The steps involved in blocking a thread (which are performed in the function itself) are as follows;

1) Remove the specified thread from the current owner's list (which is likely one of the scheduler's thread lists) 2) Add the thread to this object's thread list 3) Setting the thread's "current thread-list" point to reference this object's threadlist.

Parameters

<i>pclThread_</i>	Pointer to the thread object that will be blocked.
-------------------	--

Definition at line 42 of file [blocking.cpp](#).

19.1.2.2 BlockPriority()

```
void Mark3::BlockingObject::BlockPriority (  
    Thread * pclThread_ ) [protected]
```

BlockPriority.

Same as [Block\(\)](#), but ensures that threads are added to the block-list in priority-order, which optimizes the unblock procedure.

Parameters

<i>pcl</i> ↪ <i>Thread_</i>	Pointer to the Thread to Block.
--------------------------------	---

Definition at line 58 of file [blocking.cpp](#).

19.1.2.3 IsInitialized()

```
bool Mark3::BlockingObject::IsInitialized (
    void ) [inline], [protected]
```

IsInitialized.

Returns

Definition at line 145 of file [blocking.h](#).

19.1.2.4 Unblock()

```
void Mark3::BlockingObject::UnBlock (
    Thread * pclThread_ ) [protected]
```

UnBlock.

Unblock a thread that is already blocked on this object, returning it to the "ready" state by performing the following steps:

Parameters

<i>pcl</i> ↪ <i>Thread_</i>	Pointer to the thread to unblock.
--------------------------------	-----------------------------------

1) Removing the thread from this object's threadlist 2) Restoring the thread to its "original" owner's list

Definition at line 74 of file [blocking.cpp](#).

The documentation for this class was generated from the following files:

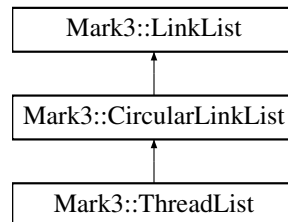
- [/home/moslevin/projects/github/m3-repo/kernel/src/public/blocking.h](#)
- [/home/moslevin/projects/github/m3-repo/kernel/src/blocking.cpp](#)

19.2 Mark3::CircularLinkedList Class Reference

Circular-linked-list data type, inherited from the base [LinkedList](#) type.

```
#include <ll.h>
```

Inheritance diagram for Mark3::CircularLinkedList:



Public Member Functions

- void [Add](#) ([LinkedListNode](#) *node_)
Add the linked list node to this linked list.
- void [Remove](#) ([LinkedListNode](#) *node_)
Remove.
- void [PivotForward](#) ()
PivotForward.
- void [PivotBackward](#) ()
PivotBackward.
- void [InsertNodeBefore](#) ([LinkedListNode](#) *node_, [LinkedListNode](#) *insert_)
InsertNodeBefore.

Additional Inherited Members

19.2.1 Detailed Description

Circular-linked-list data type, inherited from the base [LinkedList](#) type.

Definition at line [182](#) of file [ll.h](#).

19.2.2 Member Function Documentation

19.2.2.1 Add()

```
void Mark3::CircularLinkedList::Add (
    LinkedListNode * node_ )
```

Add the linked list node to this linked list.

Parameters

<i>node</i> ↔ —	Pointer to the node to add
--------------------	----------------------------

Definition at line 100 of file [ll.cpp](#).

19.2.2.2 InsertNodeBefore()

```
void Mark3::CircularLinkedList::InsertNodeBefore (
    LinkListNode * node_,
    LinkListNode * insert_ )
```

InsertNodeBefore.

Insert a linked-list node into the list before the specified insertion point.

Parameters

<i>node</i> ↔ —	Node to insert into the list
<i>insert</i> ↔ —	Insert point.

Definition at line 174 of file [ll.cpp](#).

19.2.2.3 PivotBackward()

```
void Mark3::CircularLinkedList::PivotBackward ( )
```

PivotBackward.

Pivot the head of the circularly linked list backward (Head = Head->prev, Tail = Tail->prev)

Definition at line 165 of file [ll.cpp](#).

19.2.2.4 PivotForward()

```
void Mark3::CircularLinkedList::PivotForward ( )
```

PivotForward.

Pivot the head of the circularly linked list forward (Head = Head->next, Tail = Tail->next)

Definition at line 156 of file [ll.cpp](#).

19.2.2.5 Remove()

```
void Mark3::CircularLinkedList::Remove (
    LinkListNode * node_ )
```

Remove.

Add the linked list node to this linked list

Parameters

<i>node_</i> ↔	Pointer to the node to remove
—	

Definition at line 122 of file [ll.cpp](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/github/m3-repo/kernel/src/public/ll.h](#)
- [/home/moslevin/projects/github/m3-repo/kernel/src/ll.cpp](#)

19.3 Mark3::ConditionVariable Class Reference

The [ConditionVariable](#) class This class implements a condition variable.

```
#include <condvar.h>
```

Public Member Functions

- void [Init](#) ()
Init Initialize the condition variable prior to use.
- void [Wait](#) ([Mutex](#) *pclMutex_)
Wait Block the current thread, and wait for the object to be signalled.
- bool [Wait](#) ([Mutex](#) *pclMutex_, uint32_t u32WaitTimeMS_)
Wait Block the current thread, and wait for the object to be signalled.
- void [Signal](#) ()
Signal Signal/Unblock the next thread currently blocked on this condition variable.
- void [Broadcast](#) ()
Broadcast Unblock all threads currently blocked on this condition variable.

19.3.1 Detailed Description

The [ConditionVariable](#) class This class implements a condition variable.

This is a synchronization object that allows multiple threads to block, each waiting for specific signals unique to them. Access to the specified condition is guarded by a mutex that is supplied by the caller. This object can permit multiple waiters that can be unblocked one-at-a-time via signalling, or unblocked all at once via broadcasting. This object is built upon lower-level primitives, and is somewhat more heavyweight than the primitive types supplied by the kernel.

Definition at line 38 of file [condvar.h](#).

19.3.2 Member Function Documentation

19.3.2.1 Init()

```
void Mark3::ConditionVariable::Init (
    void )
```

Init Initialize the condition variable prior to use.

Must be called before the object can be used

Definition at line 38 of file [condvar.cpp](#).

19.3.2.2 Wait() [1/2]

```
void Mark3::ConditionVariable::Wait (
    Mutex * pclMutex_ )
```

Wait Block the current thread, and wait for the object to be signalled.

The specified mutex will be locked when the thread returns.

Parameters

<i>pcl</i> ↔ <i>Mutex_</i>	Mutex to claim once the calling thread has access to the condvar
-------------------------------	--

Definition at line 46 of file [condvar.cpp](#).

19.3.2.3 Wait() [2/2]

```
bool Mark3::ConditionVariable::Wait (
    Mutex * pclMutex_,
    uint32_t u32WaitTimeMS_ )
```

Wait Block the current thread, and wait for the object to be signalled.

The specified mutex will be locked when the thread returns.

Parameters

<i>pclMutex_</i>	Mutex to claim once the calling thread has access to the condvar
<i>u32WaitTimeM</i> ↔ <i>S_</i>	Maximum time in ms to wait before abandoning the operation

Returns

true on success, false on timeout

The documentation for this class was generated from the following files:

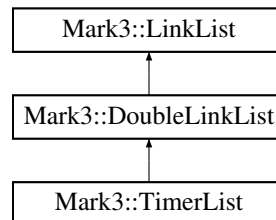
- </home/moslevin/projects/github/m3-repo/kernel/src/public/condvar.h>
- </home/moslevin/projects/github/m3-repo/kernel/src/condvar.cpp>

19.4 Mark3::DoubleLinkedList Class Reference

Doubly-linked-list data type, inherited from the base [LinkedList](#) type.

```
#include <ll.h>
```

Inheritance diagram for Mark3::DoubleLinkedList:



Public Member Functions

- [DoubleLinkedList](#) ()
DoubleLinkedList.
- void [Add](#) ([LinkedListNode](#) *node_)
Add.
- void [Remove](#) ([LinkedListNode](#) *node_)
Remove.

Additional Inherited Members

19.4.1 Detailed Description

Doubly-linked-list data type, inherited from the base [LinkedList](#) type.

Definition at line [144](#) of file [ll.h](#).

19.4.2 Constructor & Destructor Documentation

19.4.2.1 DoubleLinkedList()

```
Mark3::DoubleLinkedList::DoubleLinkedList ( ) [inline]
```

[DoubleLinkedList](#).

Default constructor - initializes the head/tail nodes to NULL

Definition at line [153](#) of file [ll.h](#).

19.4.3 Member Function Documentation

19.4.3.1 Add()

```
void Mark3::DoubleLinkedList::Add (
    LinkedListNode * node_ )
```

Add.

Add the linked list node to this linked list

Parameters

<i>node_</i> ↔	Pointer to the node to add
—	

Definition at line [49](#) of file [ll.cpp](#).

19.4.3.2 Remove()

```
void Mark3::DoubleLinkedList::Remove (
    LinkedListNode * node_ )
```

Remove.

Add the linked list node to this linked list

Parameters

<i>node_</i> ↔	Pointer to the node to remove
—	

Definition at line [70](#) of file [ll.cpp](#).

The documentation for this class was generated from the following files:

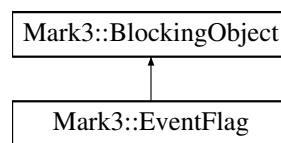
- </home/moslevin/projects/github/m3-repo/kernel/src/public/ll.h>
- </home/moslevin/projects/github/m3-repo/kernel/src/ll.cpp>

19.5 Mark3::EventFlag Class Reference

The [EventFlag](#) class is a blocking object, similar to a semaphore or mutex, commonly used for synchronizing thread execution based on events occurring within the system.

```
#include <eventflag.h>
```

Inheritance diagram for Mark3::EventFlag:



Public Member Functions

- void [Init](#) ()
Init Initializes the [EventFlag](#) object prior to use.
- uint16_t [Wait](#) (uint16_t u16Mask_, [EventFlagOperation](#) eMode_)
Wait - Block a thread on the specific flags in this event flag group.
- uint16_t [Wait](#) (uint16_t u16Mask_, [EventFlagOperation](#) eMode_, uint32_t u32TimeMS_)
Wait - Block a thread on the specific flags in this event flag group.
- void [WakeMe](#) ([Thread](#) *pclChosenOne_)
WakeMe.
- void [Set](#) (uint16_t u16Mask_)
Set - Set additional flags in this object (logical OR).
- void [Clear](#) (uint16_t u16Mask_)
ClearFlags - Clear a specific set of flags within this object, specific by bitmask.
- uint16_t [GetMask](#) ()
GetMask Returns the state of the 16-bit bitmask within this object.

Private Member Functions

- uint16_t [Wait_i](#) (uint16_t u16Mask_, [EventFlagOperation](#) eMode_, uint32_t u32TimeMS_)
Wait_i.

Private Attributes

- uint16_t [m_u16SetMask](#)
Event flags currently set in this object.

Additional Inherited Members

19.5.1 Detailed Description

The [EventFlag](#) class is a blocking object, similar to a semaphore or mutex, commonly used for synchronizing thread execution based on events occurring within the system.

Each [EventFlag](#) object contains a 16-bit bitmask, which is used to trigger events on associated threads. Threads wishing to block, waiting for a specific event to occur can wait on any pattern within this 16-bit bitmask to be set. Here, we provide the ability for a thread to block, waiting for ANY bits in a specified mask to be set, or for ALL bits within a specific mask to be set. Depending on how the object is configured, the bits that triggered the wakeup can be automatically cleared once a match has occurred.

Examples:

[lab7_events/main.cpp](#).

Definition at line 45 of file [eventflag.h](#).

19.5.2 Member Function Documentation

19.5.2.1 Clear()

```
void Mark3::EventFlag::Clear (
    uint16_t u16Mask_ )
```

ClearFlags - Clear a specific set of flags within this object, specific by bitmask.

Parameters

<i>u16Mask_</i>	- Bitmask of flags to clear
-----------------	-----------------------------

Examples:

[lab7_events/main.cpp](#).

19.5.2.2 GetMask()

```
uint16_t Mark3::EventFlag::GetMask ( )
```

GetMask Returns the state of the 16-bit bitmask within this object.

Returns

The state of the 16-bit bitmask

19.5.2.3 Set()

```
void Mark3::EventFlag::Set (
    uint16_t u16Mask_ )
```

Set - Set additional flags in this object (logical OR).

This API can potentially result in threads blocked on [Wait\(\)](#) to be unblocked.

Parameters

<i>u16Mask_</i>	- Bitmask of flags to set.
-----------------	----------------------------

Examples:

[lab7_events/main.cpp](#).

19.5.2.4 Wait() [1/2]

```
uint16_t Mark3::EventFlag::Wait (
    uint16_t u16Mask_,
    EventFlagOperation eMode_ )
```

Wait - Block a thread on the specific flags in this event flag group.

Parameters

<i>u16Mask_</i>	- 16-bit bitmask to block on
<i>eMode_</i>	- EventFlagOperation::Any_Set: Thread will block on any of the bits in the mask <ul style="list-style-type: none"> EventFlagOperation::All_Set: Thread will block on all of the bits in the mask

Returns

Bitmask condition that caused the thread to unblock, or 0 on error or timeout

Examples:

[lab7_events/main.cpp](#).

19.5.2.5 Wait() [2/2]

```
uint16_t Mark3::EventFlag::Wait (
    uint16_t u16Mask_,
```

```
EventFlagOperation eMode_,  
uint32_t u32TimeMS_ )
```

Wait - Block a thread on the specific flags in this event flag group.

Parameters

<i>u16Mask_</i>	- 16-bit bitmask to block on
<i>eMode_</i>	- EventFlagOperation::Any_Set : Thread will block on any of the bits in the mask <ul style="list-style-type: none"> • EventFlagOperation::All_Set: Thread will block on all of the bits in the mask
<i>u32TimeMS_</i>	- Time to block (in ms)

Returns

Bitmask condition that caused the thread to unblock, or 0 on error or timeout

19.5.2.6 Wait_i()

```
uint16_t Mark3::EventFlag::Wait_i (
    uint16_t u16Mask_,
    EventFlagOperation eMode_,
    uint32_t u32TimeMS_ ) [private]
```

Wait_i.

Internal abstraction used to manage both timed and untimed wait operations

Parameters

<i>u16Mask_</i>	- 16-bit bitmask to block on
<i>eMode_</i>	- EventFlagOperation::Any_Set : Thread will block on any of the bits in the mask <ul style="list-style-type: none"> • EventFlagOperation::All_Set: Thread will block on all of the bits in the mask
<i>u32TimeMS_</i>	- Time to block (in ms)

Returns

Bitmask condition that caused the thread to unblock, or 0 on error or timeout

19.5.2.7 WakeMe()

```
void Mark3::EventFlag::WakeMe (
    Thread * pClChosenOne_ )
```

WakeMe.

Wake the given thread, currently blocking on this object

Parameters

<i>pc</i> ↔ <i>Owner_</i>	Pointer to the owner thread to unblock.
------------------------------	---

The documentation for this class was generated from the following file:

- /home/moslevin/projects/github/m3-repo/kernel/src/public/[eventflag.h](#)

19.6 Mark3::FakeThread_t Struct Reference

If the kernel is set up to use an idle function instead of an idle thread, we use a placeholder data structure to "simulate" the effect of having an idle thread in the system.

```
#include <thread.h>
```

Public Attributes

- [K_WORD](#) * [m_pwStackTop](#)
Pointer to the top of the thread's stack.
- [K_WORD](#) * [m_pwStack](#)
Pointer to the thread's stack.
- [uint8_t](#) [m_u8ThreadID](#)
Thread ID.
- [PORT_PRIO_TYPE](#) [m_uXPriority](#)
Default priority of the thread.
- [PORT_PRIO_TYPE](#) [m_uXCurPriority](#)
Current priority of the thread (priority inheritance)
- [ThreadState](#) [m_eState](#)
Enum indicating the thread's current state.
- [void](#) * [m_pvExtendedContext](#)
Pointer provided to a [Thread](#) to implement thread-local storage.

19.6.1 Detailed Description

If the kernel is set up to use an idle function instead of an idle thread, we use a placeholder data structure to "simulate" the effect of having an idle thread in the system.

When cast to a [Thread](#), this data structure will still result in `GetPriority()` calls being valid, which is all that is needed to support the tick-based/tickless times – while saving a fairly decent chunk of RAM on a small micro.

Note that this struct must have the same memory layout as the [Thread](#) class up to the last item.

Definition at line [535](#) of file [thread.h](#).

The documentation for this struct was generated from the following file:

- /home/moslevin/projects/github/m3-repo/kernel/src/public/[thread.h](#)

19.7 Mark3::Kernel Class Reference

Class that encapsulates all of the kernel startup functions.

```
#include <kernel.h>
```

Static Public Member Functions

- static void [Init](#) (void)
Kernel Initialization Function, call before any other OS function.
- static void [Start](#) (void)
Start the operating system kernel - the current execution context is cancelled, all kernel services are started, and the processor resumes execution at the entrypoint for the highest-priority thread.
- static bool [IsStarted](#) ()
IsStarted.
- static void [SetPanic](#) ([PanicFunc](#) pfPanic_)
SetPanic Set a function to be called when a kernel panic occurs, giving the user to determine the behavior when a catastrophic failure is observed.
- static bool [IsPanic](#) ()
IsPanic Returns whether or not the kernel is in a panic state.
- static void [Panic](#) (uint16_t u16Cause_)
Panic Cause the kernel to enter its panic state.
- static void [SetIdleFunc](#) ([IdleFunc](#) pfIdle_)
SetIdleFunc Set the function to be called when no active threads are available to be scheduled by the scheduler.
- static void [Idle](#) (void)
IdleFunc Call the low-priority idle function when no active threads are available to be scheduled.
- static [Thread](#) * [GetIdleThread](#) (void)
GetIdleThread Return a pointer to the Kernel's idle thread object to the user.
- static void [SetThreadCreateCallout](#) ([ThreadCreateCallout](#) pfCreate_)
SetThreadCreateCallout.
- static void [SetThreadExitCallout](#) ([ThreadExitCallout](#) pfExit_)
SetThreadExitCallout.
- static void [SetThreadContextSwitchCallout](#) ([ThreadContextCallout](#) pfContext_)
SetThreadContextSwitchCallout.
- static [ThreadCreateCallout](#) [GetThreadCreateCallout](#) (void)
GetThreadCreateCallout.
- static [ThreadExitCallout](#) [GetThreadExitCallout](#) (void)
GetThreadExitCallout.
- static [ThreadContextCallout](#) [GetThreadContextSwitchCallout](#) (void)
GetThreadContextSwitchCallout.

Static Private Attributes

- static bool [m_bIsStarted](#)
true if kernel is running, false otherwise
- static bool [m_bIsPanic](#)
true if kernel is in panic state, false otherwise
- static [PanicFunc](#) [m_pfPanic](#)
set panic function
- static [IdleFunc](#) [m_pfIdle](#)

- set idle function*
- static [FakeThread_t m_clIdle](#)
Idle thread object (note: not a real thread)
- static ThreadCreateCallout [m_pfThreadCreateCallout](#)
Function to call on thread creation.
- static ThreadExitCallout [m_pfThreadExitCallout](#)
Function to call on thread exit.
- static ThreadContextCallout [m_pfThreadContextCallout](#)
Function to call on context switch.

19.7.1 Detailed Description

Class that encapsulates all of the kernel startup functions.

Definition at line 45 of file [kernel.h](#).

19.7.2 Member Function Documentation

19.7.2.1 GetIdleThread()

```
static Thread* Mark3::Kernel::GetIdleThread (
    void ) [inline], [static]
```

GetIdleThread Return a pointer to the [Kernel](#)'s idle thread object to the user.

Note that the [Thread](#) object involved is to be used for comparisons only – the thread itself is "virtual", and doesn't represent a unique execution context with its own stack.

Returns

Pointer to the [Kernel](#)'s idle thread object

Definition at line 124 of file [kernel.h](#).

19.7.2.2 GetThreadContextSwitchCallout()

```
static ThreadContextCallout Mark3::Kernel::GetThreadContextSwitchCallout (
    void ) [inline], [static]
```

GetThreadContextSwitchCallout.

Return the current function called on every [Thread::ContextSwitchSWI\(\)](#)

Returns

Pointer to the currently-installed callout function, or NULL if not set.

Definition at line 192 of file [kernel.h](#).

19.7.2.3 GetThreadCreateCallout()

```
static ThreadCreateCallout Mark3::Kernel::GetThreadCreateCallout (
    void ) [inline], [static]
```

GetThreadCreateCallout.

Return the current function called on every [Thread::Init\(\)](#);

Returns

Pointer to the currently-installed callout function, or NULL if not set.

Definition at line 174 of file [kernel.h](#).

19.7.2.4 GetThreadExitCallout()

```
static ThreadExitCallout Mark3::Kernel::GetThreadExitCallout (
    void ) [inline], [static]
```

GetThreadExitCallout.

Return the current function called on every [Thread::Exit\(\)](#);

Returns

Pointer to the currently-installed callout function, or NULL if not set.

Definition at line 183 of file [kernel.h](#).

19.7.2.5 Init()

```
void Mark3::Kernel::Init (
    void ) [static]
```

[Kernel](#) Initialization Function, call before any other OS function.

Initializes all global resources used by the operating system. This must be called before any other kernel function is invoked.

Examples:

[lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_idle_function/main.cpp](#), [lab3_round_robin/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab5_mutexes/main.cpp](#), [lab6_timers/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 75 of file [kernel.cpp](#).

19.7.2.6 IsPanic()

```
static bool Mark3::Kernel::IsPanic ( ) [inline], [static]
```

IsPanic Returns whether or not the kernel is in a panic state.

Returns

Whether or not the kernel is in a panic state

Definition at line 92 of file [kernel.h](#).

19.7.2.7 IsStarted()

```
static bool Mark3::Kernel::IsStarted ( ) [inline], [static]
```

IsStarted.

Returns

Whether or not the kernel has started - true = running, false = not started

Definition at line 79 of file [kernel.h](#).

19.7.2.8 Panic()

```
void Mark3::Kernel::Panic (
    uint16_t u16Cause_ ) [static]
```

Panic Cause the kernel to enter its panic state.

Parameters

<i>u16Cause_↵</i>	Reason for the kernel panic
-------------------	-----------------------------

Definition at line 109 of file [kernel.cpp](#).

19.7.2.9 SetIdleFunc()

```
static void Mark3::Kernel::SetIdleFunc (
    IdleFunc pfIdle_ ) [inline], [static]
```

SetIdleFunc Set the function to be called when no active threads are available to be scheduled by the scheduler.

Parameters

<i>pf</i> ↔	Pointer to the idle function
<i>Idle</i> ↔	
—	

Examples:

[lab2_idle_function/main.cpp](#).

Definition at line 105 of file [kernel.h](#).

19.7.2.10 SetPanic()

```
static void Mark3::Kernel::SetPanic (
    PanicFunc pfPanic_ ) [inline], [static]
```

SetPanic Set a function to be called when a kernel panic occurs, giving the user to determine the behavior when a catastrophic failure is observed.

Parameters

<i>pf</i> ↔	Panic function pointer
<i>Panic</i> ↔	
—	

Definition at line 87 of file [kernel.h](#).

19.7.2.11 SetThreadContextSwitchCallout()

```
static void Mark3::Kernel::SetThreadContextSwitchCallout (
    ThreadContextCallout pfContext_ ) [inline], [static]
```

SetThreadContextSwitchCallout.

Set a function to be called on each context switch.

A callout is only executed if this method has been called to set a valid handler function.

Parameters

<i>pf</i> ↔	Pointer to a function to call on context switch
<i>Context</i> ↔	
—	

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 161 of file [kernel.h](#).

19.7.2.12 SetThreadCreateCallout()

```
static void Mark3::Kernel::SetThreadCreateCallout (
    ThreadCreateCallout pfCreate_ ) [inline], [static]
```

SetThreadCreateCallout.

Set a function to be called on creation of a new thread. This callout is executed on the successful completion of a [Thread::Init\(\)](#) call. A callout is only executed if this method has been called to set a valid handler function.

Parameters

<i>pfCreate_</i>	Pointer to a function to call on thread creation
—	

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 138 of file [kernel.h](#).

19.7.2.13 SetThreadExitCallout()

```
static void Mark3::Kernel::SetThreadExitCallout (
    ThreadExitCallout pfExit_ ) [inline], [static]
```

SetThreadExitCallout.

Set a function to be called on thread exit. This callout is executed from the beginning of [Thread::Exit\(\)](#).

A callout is only executed if this method has been called to set a valid handler function.

Parameters

<i>pfExit_</i>	Pointer to a function to call on thread exit
—	

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 150 of file [kernel.h](#).

19.7.2.14 Start()

```
void Mark3::Kernel::Start (  
    void ) [static]
```

Start the operating system kernel - the current execution context is cancelled, all kernel services are started, and the processor resumes execution at the entrypoint for the highest-priority thread.

You must have at least one thread added to the kernel before calling this function, otherwise the behavior is undefined. The exception to this is if the system is configured to use the threadless idle hook, in which case the kernel is allowed to run without any ready threads.

Examples:

[lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_idle_↔
function/main.cpp](#), [lab3_round_robin/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab5_mutexes/main.cpp](#), [lab6_↔
_timers/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 100 of file [kernel.cpp](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/github/m3-repo/kernel/src/public/kernel.h](#)
- [/home/moslevin/projects/github/m3-repo/kernel/src/kernel.cpp](#)

19.8 Mark3::KernelTimer Class Reference

Hardware timer interface, used by all scheduling/timer subsystems.

```
#include <kerneltimer.h>
```

Static Public Member Functions

- static void [Config](#) (void)
Config.
- static void [Start](#) (void)
Start.
- static void [Stop](#) (void)
Stop.
- static uint8_t [DI](#) (void)
DI.
- static void [RI](#) (bool bEnable_)
RI.
- static void [EI](#) (void)
EI.
- static [PORT_TIMER_COUNT_TYPE SubtractExpiry](#) ([PORT_TIMER_COUNT_TYPE](#) ulInterval_)
SubtractExpiry.
- static [PORT_TIMER_COUNT_TYPE TimeToExpiry](#) (void)
TimeToExpiry.
- static [PORT_TIMER_COUNT_TYPE SetExpiry](#) (uint32_t u32Interval_)
SetExpiry.
- static [PORT_TIMER_COUNT_TYPE GetOvertime](#) (void)
GetOvertime.
- static void [ClearExpiry](#) (void)
ClearExpiry.
- static [PORT_TIMER_COUNT_TYPE Read](#) (void)
Read.

19.8.1 Detailed Description

Hardware timer interface, used by all scheduling/timer subsystems.

Definition at line 32 of file [kerneltimer.h](#).

19.8.2 Member Function Documentation

19.8.2.1 [ClearExpiry\(\)](#)

```
void Mark3::KernelTimer::ClearExpiry (
    void ) [static]
```

[ClearExpiry.](#)

Clear the hardware timer expiry register

Definition at line 179 of file [kerneltimer.cpp](#).

19.8.2.2 Config()

```
void Mark3::KernelTimer::Config (  
    void ) [static]
```

Config.

Initializes the kernel timer before use

Definition at line 66 of file [kerneltimer.cpp](#).

19.8.2.3 DI()

```
uint8_t Mark3::KernelTimer::DI (  
    void ) [static]
```

DI.

Disable the kernel timer's expiry interrupt

Definition at line 187 of file [kerneltimer.cpp](#).

19.8.2.4 EI()

```
void Mark3::KernelTimer::EI (  
    void ) [static]
```

EI.

Enable the kernel timer's expiry interrupt

Definition at line 200 of file [kerneltimer.cpp](#).

19.8.2.5 GetOvertime()

```
PORT\_TIMER\_COUNT\_TYPE Mark3::KernelTimer::GetOvertime (  
    void ) [static]
```

GetOvertime.

Return the number of ticks that have elapsed since the last expiry.

Returns

Number of ticks that have elapsed after last timer expiration

Definition at line 155 of file [kerneltimer.cpp](#).

19.8.2.6 Read()

```
PORT_TIMER_COUNT_TYPE Mark3::KernelTimer::Read (
    void ) [static]
```

Read.

Safely read the current value in the timer register

Returns

Value held in the timer register

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 109 of file [kerneltimer.cpp](#).

19.8.2.7 RI()

```
void Mark3::KernelTimer::RI (
    bool bEnable_ ) [static]
```

RI.

Retstore the state of the kernel timer's expiry interrupt.

Parameters

$b \leftrightarrow$ <i>Enable</i> \leftrightarrow —	1 enable, 0 disable
---	---------------------

Definition at line 206 of file [kerneltimer.cpp](#).

19.8.2.8 SetExpiry()

```
PORT_TIMER_COUNT_TYPE Mark3::KernelTimer::SetExpiry (
    uint32_t u32Interval_ ) [static]
```

SetExpiry.

Resets the kernel timer's expiry interval to the specified value

Parameters

<i>u32↔ Interval_</i>	Desired interval in ticks to set the timer for
---------------------------	--

Returns

Actual number of ticks set (may be less than desired)

Definition at line 161 of file [kerneltimer.cpp](#).

19.8.2.9 Start()

```
void Mark3::KernelTimer::Start (
    void ) [static]
```

Start.

Starts the kernel time (must be configured first)

Definition at line 82 of file [kerneltimer.cpp](#).

19.8.2.10 Stop()

```
void Mark3::KernelTimer::Stop (
    void ) [static]
```

Stop.

Shut down the kernel timer, used when no timers are scheduled

Definition at line 97 of file [kerneltimer.cpp](#).

19.8.2.11 SubtractExpiry()

```
PORT_TIMER_COUNT_TYPE Mark3::KernelTimer::SubtractExpiry (
    PORT_TIMER_COUNT_TYPE uInterval_ ) [static]
```

SubtractExpiry.

Subtract the specified number of ticks from the timer's expiry count register. Returns the new expiry value stored in the register.

Parameters

<i>u32↔ Interval_</i>	Time (in HW-specific) ticks to subtract
---------------------------	---

Returns

Value in ticks stored in the timer's expiry register

Definition at line 127 of file [kerneltimer.cpp](#).

19.8.2.12 TimeToExpiry()

```
PORT_TIMER_COUNT_TYPE Mark3::KernelTimer::TimeToExpiry (
    void ) [static]
```

TimeToExpiry.

Returns the number of ticks remaining before the next timer expiry.

Returns

Time before next expiry in platform-specific ticks

Definition at line 138 of file [kerneltimer.cpp](#).

The documentation for this class was generated from the following files:

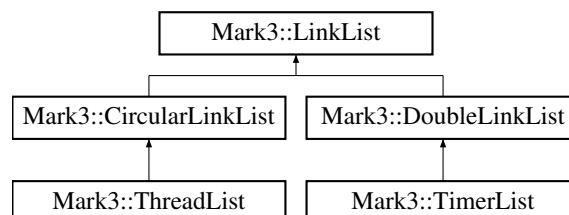
- [/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/kerneltimer.h](#)
- [/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/kerneltimer.cpp](#)

19.9 Mark3::LinkedList Class Reference

Abstract-data-type from which all other linked-lists are derived.

```
#include <ll.h>
```

Inheritance diagram for Mark3::LinkedList:



Public Member Functions

- void [Init](#) ()
Init.
- [LinkedListNode](#) * [GetHead](#) ()
GetHead.
- [LinkedListNode](#) * [GetTail](#) ()
GetTail.

Protected Attributes

- [LinkedListNode](#) * [m_pclHead](#)
Pointer to the head node in the list.
- [LinkedListNode](#) * [m_pclTail](#)
Pointer to the tail node in the list.

19.9.1 Detailed Description

Abstract-data-type from which all other linked-lists are derived.

Definition at line 104 of file [ll.h](#).

19.9.2 Member Function Documentation

19.9.2.1 [GetHead\(\)](#)

```
LinkedListNode* Mark3::LinkedList::GetHead ( ) [inline]
```

[GetHead.](#)

Get the head node in the linked list

Returns

Pointer to the head node in the list

Definition at line 129 of file [ll.h](#).

19.9.2.2 GetTail()

```
LinkedListNode* Mark3::LinkedList::GetTail ( ) [inline]
```

GetTail.

Get the tail node of the linked list

Returns

Pointer to the tail node in the list

Definition at line 137 of file ll.h.

19.9.2.3 Init()

```
void Mark3::LinkedList::Init (
    void ) [inline]
```

Init.

Clear the linked list.

Definition at line 116 of file ll.h.

The documentation for this class was generated from the following file:

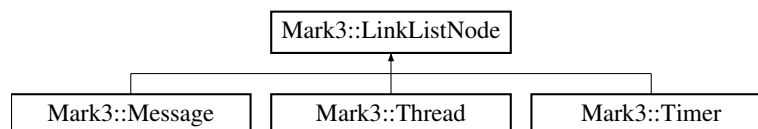
- /home/moslevin/projects/github/m3-repo/kernel/src/public/ll.h

19.10 Mark3::LinkedListNode Class Reference

Basic linked-list node data structure.

```
#include <ll.h>
```

Inheritance diagram for Mark3::LinkedListNode:



Public Member Functions

- [LinkedListNode](#) * [GetNext](#) (void)
GetNext.
- [LinkedListNode](#) * [GetPrev](#) (void)
GetPrev.

Protected Member Functions

- void [ClearNode](#) ()
ClearNode.

Protected Attributes

- [LinkListNode](#) * [next](#)
Pointer to the next node in the list.
- [LinkListNode](#) * [prev](#)
Pointer to the previous node in the list.

19.10.1 Detailed Description

Basic linked-list node data structure.

This data is managed by the linked-list class types, and can be used transparently between them.

Definition at line [63](#) of file [ll.h](#).

19.10.2 Member Function Documentation

19.10.2.1 ClearNode()

```
void Mark3::LinkListNode::ClearNode ( ) [protected]
```

[ClearNode](#).

Initialize the linked list node, clearing its next and previous node.

Definition at line [42](#) of file [ll.cpp](#).

19.10.2.2 GetNext()

```
LinkListNode\* Mark3::LinkListNode::GetNext (   
    void ) [inline]
```

[GetNext](#).

Returns a pointer to the next node in the list.

Returns

a pointer to the next node in the list.

Definition at line [85](#) of file [ll.h](#).

19.10.2.3 GetPrev()

```
LinkedListNode* Mark3::LinkedListNode::GetPrev (
    void ) [inline]
```

GetPrev.

Returns a pointer to the previous node in the list.

Returns

a pointer to the previous node in the list.

Definition at line 93 of file [ll.h](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/github/m3-repo/kernel/src/public/ll.h](#)
- [/home/moslevin/projects/github/m3-repo/kernel/src/ll.cpp](#)

19.11 Mark3::LockGuard Class Reference

```
#include <lockguard.h>
```

Public Member Functions

- [LockGuard \(Mutex *pclMutex\)](#)
- [LockGuard \(Mutex *pclMutex, uint32_t u32TimeoutMs_\)](#)
- [bool isAcquired \(\)](#)

19.11.1 Detailed Description

Implement RAII-style locks based on [Mark3's](#) kernel [Mutex](#) object. Note that [Mark3](#) does not support exceptions, so

Definition at line 32 of file [lockguard.h](#).

19.11.2 Constructor & Destructor Documentation

19.11.2.1 LockGuard() [1/2]

```
Mark3::LockGuard::LockGuard (
    Mutex * pclMutex )
```

Parameters

<i>pclMutex</i>	mutex to lock during construction
-----------------	-----------------------------------

Definition at line 23 of file [lockguard.cpp](#).

19.11.2.2 LockGuard() [2/2]

```
Mark3::LockGuard::LockGuard (
    Mutex * pclMutex,
    uint32_t u32TimeoutMs_ )
```

Parameters

<i>pclMutex</i>	mutex to lock during construction
<i>u32TimeoutMs_</i>	timeout (in ms) to wait before bailing

19.11.3 Member Function Documentation

19.11.3.1 isAcquired()

```
bool Mark3::LockGuard::isAcquired ( ) [inline]
```

Verify that lock was correctly initialized and locked during acquisition. This is used to provide error-checking for timed RAIL locks, where [Mark3](#) does not use exceptions, and a kernel-panic is too heavy-handed.

Returns

true if the lock was initialed correctly, false on error

Definition at line 56 of file [lockguard.h](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/github/m3-repo/kernel/src/public/lockguard.h](#)
- [/home/moslevin/projects/github/m3-repo/kernel/src/lockguard.cpp](#)

19.12 Mark3::Mailbox Class Reference

The [Mailbox](#) class implements an IPC mechnism based on envelopes containing data of a fixed size (configured at initialization) that reside within a buffer of memory provided by the user.

```
#include <mailbox.h>
```

Public Member Functions

- void [Init](#) (void *pvBuffer_, uint16_t u16BufferSize_, uint16_t u16ElementSize_)
Init.
- bool [Send](#) (void *pvData_)
Send.
- bool [SendTail](#) (void *pvData_)
SendTail.
- bool [Send](#) (void *pvData_, uint32_t u32TimeoutMS_)
Send.
- bool [SendTail](#) (void *pvData_, uint32_t u32TimeoutMS_)
SendTail.
- void [Receive](#) (void *pvData_)
Receive.
- void [ReceiveTail](#) (void *pvData_)
ReceiveTail.
- bool [Receive](#) (void *pvData_, uint32_t u32TimeoutMS_)
Receive.
- bool [ReceiveTail](#) (void *pvData_, uint32_t u32TimeoutMS_)
ReceiveTail.

Static Public Member Functions

- static [Mailbox](#) * [Init](#) (uint16_t u16BufferSize_, uint16_t u16ElementSize_)
Init.

Private Member Functions

- void * [GetHeadPointer](#) (void)
GetHeadPointer.
- void * [GetTailPointer](#) (void)
GetTailPointer.
- void [CopyData](#) (const void *src_, const void *dst_, uint16_t len_)
CopyData.
- void [MoveTailForward](#) (void)
MoveTailForward.
- void [MoveHeadForward](#) (void)
MoveHeadForward.
- void [MoveTailBackward](#) (void)
MoveTailBackward.
- void [MoveHeadBackward](#) (void)
MoveHeadBackward.
- bool [Send_i](#) (const void *pvData_, bool bTail_, uint32_t u32TimeoutMS_)
Send_i.
- bool [Receive_i](#) (const void *pvData_, bool bTail_, uint32_t u32WaitTimeMS_)
Receive_i.

Private Attributes

- `uint16_t m_u16Head`
Current head index.
- `uint16_t m_u16Tail`
Current tail index.
- `uint16_t m_u16Count`
Count of items in the mailbox.
- `volatile uint16_t m_u16Free`
Current number of free slots in the mailbox.
- `uint16_t m_u16ElementSize`
Size of the objects tracked in this mailbox.
- `const void * m_pvBuffer`
Pointer to the data-buffer managed by this mailbox.
- `Semaphore m_clRecvSem`
Counting semaphore used to synchronize threads on the object.
- `Semaphore m_clSendSem`
Binary semaphore for send-blocked threads.

19.12.1 Detailed Description

The `Mailbox` class implements an IPC mechanism based on envelopes containing data of a fixed size (configured at initialization) that reside within a buffer of memory provided by the user.

Examples:

[lab11_mailboxes/main.cpp](#).

Definition at line 35 of file [mailbox.h](#).

19.12.2 Member Function Documentation

19.12.2.1 CopyData()

```
void Mark3::Mailbox::CopyData (
    const void * src_,
    const void * dst_,
    uint16_t len_ ) [inline], [private]
```

CopyData.

Perform a direct byte-copy from a source to a destination object.

Parameters

<code>src_↔</code>	Pointer to an object to read from
<code>dst_↔</code>	Pointer to an object to write to
<code>len_↔</code>	Length to copy (in bytes)

Definition at line 238 of file [mailbox.h](#).

19.12.2.2 GetHeadPointer()

```
void* Mark3::Mailbox::GetHeadPointer (
    void ) [inline], [private]
```

GetHeadPointer.

Return a pointer to the current head of the mailbox's internal circular buffer.

Returns

pointer to the head element in the mailbox

Definition at line 207 of file [mailbox.h](#).

19.12.2.3 GetTailPointer()

```
void* Mark3::Mailbox::GetTailPointer (
    void ) [inline], [private]
```

GetTailPointer.

Return a pointer to the current tail of the mailbox's internal circular buffer.

Returns

pointer to the tail element in the mailbox

Definition at line 222 of file [mailbox.h](#).

19.12.2.4 Init() [1/2]

```
void Mark3::Mailbox::Init (
    void * pvBuffer_,
    uint16_t u16BufferSize_,
    uint16_t u16ElementSize_ )
```

Init.

Initialize the mailbox object prior to its use. This must be called before any calls can be made to the object.

Parameters

<i>pvBuffer_</i>	Pointer to the static buffer to use for the mailbox
<i>u16BufferSize_</i>	Size of the mailbox buffer, in bytes
<i>u16Element↔Size_</i>	Size of each envelope, in bytes

19.12.2.5 Init() [2/2]

```
static Mailbox* Mark3::Mailbox::Init (
    uint16_t u16BufferSize_,
    uint16_t u16ElementSize_ ) [static]
```

Init.

Create and initialize the mailbox object prior to its use. This must be called before any calls can be made to the object. This version of the API allocates the buffer space from the kernel's Auto-Allocation heap, which cannot be returned back. As a result, this is only suitable for cases where the mailbox will be created once on startup, and persist for the duration of the system.

Parameters

<i>u16BufferSize_</i>	Size of the mailbox buffer, in bytes
<i>u16Element↔Size_</i>	Size of each envelope, in bytes

19.12.2.6 MoveHeadBackward()

```
void Mark3::Mailbox::MoveHeadBackward (
    void ) [inline], [private]
```

MoveHeadBackward.

Move the head index backward one element

Definition at line 291 of file [mailbox.h](#).

19.12.2.7 MoveHeadForward()

```
void Mark3::Mailbox::MoveHeadForward (
    void ) [inline], [private]
```

MoveHeadForward.

Move the head index forward one element

Definition at line 265 of file [mailbox.h](#).

19.12.2.8 MoveTailBackward()

```
void Mark3::Mailbox::MoveTailBackward (
    void ) [inline], [private]
```

MoveTailBackward.

Move the tail index backward one element

Definition at line 278 of file [mailbox.h](#).

19.12.2.9 MoveTailForward()

```
void Mark3::Mailbox::MoveTailForward (
    void ) [inline], [private]
```

MoveTailForward.

Move the tail index forward one element

Definition at line 252 of file [mailbox.h](#).

19.12.2.10 Receive() [1/2]

```
void Mark3::Mailbox::Receive (
    void * pvData_ )
```

Receive.

Read one envelope from the head of the mailbox. If the mailbox is currently empty, the calling thread will block until an envelope is delivered.

Parameters

<i>pvData_</i>	Pointer to a buffer that will have the envelope's contents copied into upon delivery.
--------------------------------	---

Examples:

[lab11_mailboxes/main.cpp](#).

19.12.2.11 Receive() [2/2]

```
bool Mark3::Mailbox::Receive (
    void * pvData_,
    uint32_t u32TimeoutMS_ )
```

Receive.

Read one envelope from the head of the mailbox. If the mailbox is currently empty, the calling thread will block until an envelope is delivered, or the specified time has elapsed without delivery.

Parameters

<i>pvData_</i>	Pointer to a buffer that will have the envelope's contents copied into upon delivery.
<i>u32TimeoutMS_</i>	Maximum time to wait for delivery.

Returns

true - envelope was delivered, false - delivery timed out.

19.12.2.12 Receive_i()

```
bool Mark3::Mailbox::Receive_i (
    const void * pvData_,
    bool bTail_,
    uint32_t u32WaitTimeMS_ ) [private]
```

Receive_i.

Internal method which implements all Read() methods in the class.

Parameters

<i>pvData_</i>	Pointer to the envelope data
<i>bTail_</i>	true - read from tail, false - read from head
<i>u32WaitTimeMS_</i>	Time to wait before timeout (in ms).

Returns

true - read successfully, false - timeout.

19.12.2.13 ReceiveTail() [1/2]

```
void Mark3::Mailbox::ReceiveTail (
    void * pvData_ )
```

ReceiveTail.

Read one envelope from the tail of the mailbox. If the mailbox is currently empty, the calling thread will block until an envelope is delivered.

Parameters

<i>pvData_</i>	Pointer to a buffer that will have the envelope's contents copied into upon delivery.
----------------	---

19.12.2.14 ReceiveTail() [2/2]

```
bool Mark3::Mailbox::ReceiveTail (
    void * pvData_,
    uint32_t u32TimeoutMS_ )
```

ReceiveTail.

Read one envelope from the tail of the mailbox. If the mailbox is currently empty, the calling thread will block until an envelope is delivered, or the specified time has elapsed without delivery.

Parameters

<i>pvData_</i>	Pointer to a buffer that will have the envelope's contents copied into upon delivery.
<i>u32TimeoutMS_</i>	Maximum time to wait for delivery.

Returns

true - envelope was delivered, false - delivery timed out.

19.12.2.15 Send() [1/2]

```
bool Mark3::Mailbox::Send (
    void * pvData_ )
```

Send.

Send an envelope to the mailbox. This safely copies the data contents of the datastructure to the previously-initialized mailbox buffer. If there is a thread already blocking, awaiting delivery to the mailbox, it will be unblocked at this time.

This method delivers the envelope at the head of the mailbox.

Parameters

<i>pvData_</i>	Pointer to the data object to send to the mailbox.
----------------	--

Returns

true - envelope was delivered, false - mailbox is full.

Examples:

[lab11_mailboxes/main.cpp](#).

19.12.2.16 Send() [2/2]

```
bool Mark3::Mailbox::Send (
    void * pvData_,
    uint32_t u32TimeoutMS_ )
```

Send.

Send an envelope to the mailbox. This safely copies the data contents of the datastructure to the previously-initialized mailbox buffer. If there is a thread already blocking, awaiting delivery to the mailbox, it will be unblocked at this time.

This method delivers the envelope at the head of the mailbox.

Parameters

<i>pvData_</i>	Pointer to the data object to send to the mailbox.
<i>u32TimeoutMS_</i>	Maximum time to wait for a free transmit slot

Returns

true - envelope was delivered, false - mailbox is full.

19.12.2.17 Send_i()

```
bool Mark3::Mailbox::Send_i (
    const void * pvData_,
    bool bTail_,
    uint32_t u32TimeoutMS_ ) [private]
```

Send_i.

Internal method which implements all [Send\(\)](#) methods in the class.

Parameters

<i>pvData_</i>	Pointer to the envelope data
<i>bTail_</i>	true - write to tail, false - write to head
<i>u32WaitTimeMS_</i>	Time to wait before timeout (in ms).
<i>S_</i>	

Returns

true - data successfully written, false - buffer full

19.12.2.18 SendTail() [1/2]

```
bool Mark3::Mailbox::SendTail (
    void * pvData_ )
```

SendTail.

Send an envelope to the mailbox. This safely copies the data contents of the datastructure to the previously-initialized mailbox buffer. If there is a thread already blocking, awaiting delivery to the mailbox, it will be unblocked at this time.

This method delivers the envelope at the tail of the mailbox.

Parameters

<i>pvData_</i>	Pointer to the data object to send to the mailbox.
----------------	--

Returns

true - envelope was delivered, false - mailbox is full.

19.12.2.19 SendTail() [2/2]

```
bool Mark3::Mailbox::SendTail (
    void * pvData_,
    uint32_t u32TimeoutMS_ )
```

SendTail.

Send an envelope to the mailbox. This safely copies the data contents of the datastructure to the previously-initialized mailbox buffer. If there is a thread already blocking, awaiting delivery to the mailbox, it will be unblocked at this time.

This method delivers the envelope at the tail of the mailbox.

Parameters

<i>pvData_</i>	Pointer to the data object to send to the mailbox.
<i>u32TimeoutMS_</i>	Maximum time to wait for a free transmit slot

Returns

true - envelope was delivered, false - mailbox is full.

19.12.3 Member Data Documentation

19.12.3.1 m_clSendSem

`Semaphore` Mark3::Mailbox::m_clSendSem [private]

Binary semaphore for send-blocked threads.

Definition at line 360 of file [mailbox.h](#).

The documentation for this class was generated from the following file:

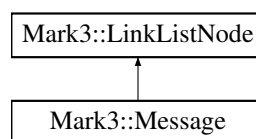
- [/home/moslevin/projects/github/m3-repo/kernel/src/public/mailbox.h](#)

19.13 Mark3::Message Class Reference

Class to provide message-based IPC services in the kernel.

```
#include <message.h>
```

Inheritance diagram for Mark3::Message:



Public Member Functions

- void [Init](#) ()
Init.
- void [SetData](#) (void *pvData_)
SetData.
- void * [GetData](#) ()
GetData.
- void [SetCode](#) (uint16_t u16Code_)
SetCode.
- uint16_t [GetCode](#) ()
GetCode.

Private Attributes

- void * [m_pvData](#)
Pointer to the message data.
- uint16_t [m_u16Code](#)
Message code, providing context for the message.

Additional Inherited Members

19.13.1 Detailed Description

Class to provide message-based IPC services in the kernel.

Examples:

[lab8_messages/main.cpp](#).

Definition at line 98 of file [message.h](#).

19.13.2 Member Function Documentation

19.13.2.1 GetCode()

```
uint16_t Mark3::Message::GetCode ( ) [inline]
```

GetCode.

Return the code set in the message upon receipt

Returns

user code set in the object

Definition at line 145 of file [message.h](#).

19.13.2.2 GetData()

```
void* Mark3::Message::GetData ( ) [inline]
```

GetData.

Get the data pointer stored in the message upon receipt

Returns

Pointer to the data set in the message object

Definition at line 129 of file [message.h](#).

19.13.2.3 Init()

```
void Mark3::Message::Init (
    void ) [inline]
```

Init.

Initialize the data and code in the message.

Definition at line 107 of file [message.h](#).

19.13.2.4 SetCode()

```
void Mark3::Message::SetCode (
    uint16_t u16Code_ ) [inline]
```

SetCode.

Set the code in the message before transmission

Parameters

<i>u16↵ Code_</i>	Data code to set in the object
-----------------------	--------------------------------

Examples:

[lab8_messages/main.cpp](#).

Definition at line 137 of file [message.h](#).

19.13.2.5 SetData()

```
void Mark3::Message::SetData (
    void * pvData_ ) [inline]
```

SetData.

Set the data pointer for the message before transmission.

Parameters

<i>pv↵ Data_</i>	Pointer to the data object to send in the message
----------------------	---

Definition at line 121 of file [message.h](#).

The documentation for this class was generated from the following file:

- [/home/moslevin/projects/github/m3-repo/kernel/src/public/message.h](#)

19.14 Mark3::MessagePool Class Reference

Implements a list of message objects.

```
#include <message.h>
```

Public Member Functions

- void [Init](#) ()
Init.
- void [Push](#) ([Message](#) *pclMessage_)
Push.
- [Message](#) * [Pop](#) ()
Pop.
- [Message](#) * [GetHead](#) ()
GetHead.

Private Attributes

- [DoubleLinkedList](#) [m_clList](#)
Linked list used to manage the [Message](#) objects.

19.14.1 Detailed Description

Implements a list of message objects.

Examples:

[lab8_messages/main.cpp](#).

Definition at line [158](#) of file [message.h](#).

19.14.2 Member Function Documentation

19.14.2.1 GetHead()

```
Message* Mark3::MessagePool::GetHead ( )
```

GetHead.

Return a pointer to the first element in the message list

Returns

19.14.2.2 Init()

```
void Mark3::MessagePool::Init ( )
```

Init.

Initialize the message queue prior to use

19.14.2.3 Pop()

```
Message* Mark3::MessagePool::Pop ( )
```

Pop.

Pop a message from the queue, returning it to the user to be popu32ated before sending by a transmitter.

Returns

Pointer to a [Message](#) object

Examples:

[lab8_messages/main.cpp](#).

19.14.2.4 Push()

```
void Mark3::MessagePool::Push (
    Message * pclMessage_ )
```

Push.

Return a previously-claimed message object back to the queue. used once the message has been processed by a receiver.

Parameters

<i>pcl</i> ↔ <i>Message_</i>	Pointer to the Message object to return back to the queue
---------------------------------	---

Examples:

[lab8_messages/main.cpp](#).

The documentation for this class was generated from the following file:

- [/home/moslevin/projects/github/m3-repo/kernel/src/public/message.h](#)

19.15 Mark3::MessageQueue Class Reference

List of messages, used as the channel for sending and receiving messages between threads.

```
#include <message.h>
```

Public Member Functions

- void [Init](#) ()
Init.
- [Message](#) * [Receive](#) ()
Receive.
- [Message](#) * [Receive](#) (uint32_t u32TimeWaitMS_)
Receive.
- void [Send](#) ([Message](#) *pclSrc_)
Send.
- uint16_t [GetCount](#) ()
GetCount.

Private Member Functions

- [Message](#) * [Receive_i](#) (uint32_t u32TimeWaitMS_)
Receive_i.

Private Attributes

- [Semaphore](#) [m_clSemaphore](#)
Counting semaphore used to manage thread blocking.
- [DoubleLinkedList](#) [m_clLinkList](#)
List object used to store messages.

19.15.1 Detailed Description

List of messages, used as the channel for sending and receiving messages between threads.

Examples:

[lab8_messages/main.cpp](#).

Definition at line 210 of file [message.h](#).

19.15.2 Member Function Documentation

19.15.2.1 GetCount()

```
uint16_t Mark3::MessageQueue::GetCount ( )
```

GetCount.

Return the number of messages pending in the "receive" queue.

Returns

Count of pending messages in the queue.

19.15.2.2 Init()

```
void Mark3::MessageQueue::Init ( )
```

Init.

Initialize the message queue prior to use.

19.15.2.3 Receive() [1/2]

```
Message* Mark3::MessageQueue::Receive ( )
```

Receive.

Receive a message from the message queue. If the message queue is empty, the thread will block until a message is available.

Returns

Pointer to a message object at the head of the queue

Examples:

[lab8_messages/main.cpp](#).

19.15.2.4 Receive() [2/2]

```
Message* Mark3::MessageQueue::Receive (
    uint32_t u32TimeWaitMS_ )
```

Receive.

Receive a message from the message queue. If the message queue is empty, the thread will block until a message is available for the duration specified. If no message arrives within that duration, the call will return with NULL.

Parameters

<i>u32TimeWaitMS_</i>	The amount of time in ms to wait for a message before timing out and unblocking the waiting thread.
-----------------------	---

Returns

Pointer to a message object at the head of the queue or NULL on timeout.

19.15.2.5 Receive_i()

```
Message* Mark3::MessageQueue::Receive_i (
    uint32_t u32TimeWaitMS_ ) [private]
```

Receive_i.

Internal function used to abstract timed and un-timed Receive calls.

Parameters

<i>u32TimeWaitMS_</i>	Time (in ms) to block, 0 for un-timed call.
-----------------------	---

Returns

Pointer to a message, or 0 on timeout.

19.15.2.6 Send()

```
void Mark3::MessageQueue::Send (
    Message * pclSrc_ )
```

Send.

Send a message object into this message queue. Will un-block the first waiting thread blocked on this queue if that occurs.

Parameters

<i>pclSrc_</i>	Pointer to the message object to add to the queue
----------------	---

Examples:

[lab8_messages/main.cpp](#).

The documentation for this class was generated from the following file:

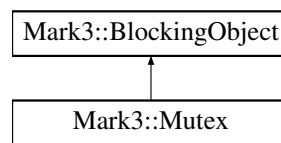
- /home/moslevin/projects/github/m3-repo/kernel/src/public/[message.h](#)

19.16 Mark3::Mutex Class Reference

Mutual-exclusion locks, based on [BlockingObject](#).

```
#include <mutex.h>
```

Inheritance diagram for Mark3::Mutex:



Public Member Functions

- void [Init](#) (bool bRecursive_=true)
Init.
- void [Claim](#) ()
Claim.
- bool [Claim](#) (uint32_t u32WaitTimeMS_)
Claim.
- void [WakeMe](#) ([Thread](#) *pclOwner_)
WakeMe.
- void [Release](#) ()
Release.

Private Member Functions

- uint8_t [WakeNext](#) ()
WakeNext.
- bool [Claim_i](#) (uint32_t u32WaitTimeMS_)
Claim_i.

Private Attributes

- uint8_t [m_u8Recurse](#)
The recursive lock-count when a mutex is claimed multiple times by the same owner.
- bool [m_bReady](#)
State of the mutex - true = ready, false = claimed.
- bool [m_bRecursive](#)
Whether or not the lock is recursive.
- uint8_t [m_u8MaxPri](#)
Maximum priority of thread in queue, used for priority inheritance.
- [Thread](#) * [m_pclOwner](#)
Pointer to the thread that owns the mutex (when claimed)

Additional Inherited Members

19.16.1 Detailed Description

Mutual-exclusion locks, based on [BlockingObject](#).

Examples:

[lab5_mutexes/main.cpp](#).

Definition at line 64 of file [mutex.h](#).

19.16.2 Member Function Documentation

19.16.2.1 Claim() [1/2]

```
void Mark3::Mutex::Claim ( )
```

Claim.

Claim the mutex. When the mutex is claimed, no other thread can claim a region protected by the object. If another [Thread](#) currently holds the [Mutex](#) when the Claim method is called, that [Thread](#) will block until the current owner of the mutex releases the [Mutex](#).

If the calling [Thread](#)'s priority is lower than that of a [Thread](#) that currently owns the [Mutex](#) object, then the priority of that [Thread](#) will be elevated to that of the highest-priority calling object until the [Mutex](#) is released. This property is known as "Priority Inheritance"

Note: A single thread can recursively claim a mutex up to a count of

1. Attempting to claim a mutex beyond that will cause a kernel panic.

Examples:

[lab5_mutexes/main.cpp](#).

19.16.2.2 Claim() [2/2]

```
bool Mark3::Mutex::Claim (
    uint32_t u32WaitTimeMS_ )
```

Claim.

Claim a mutex, with timeout.

Parameters

<i>u32WaitTimeMS_</i>	
-----------------------	--

Returns

true - mutex was claimed within the time period specified
false - mutex operation timed-out before the claim operation.

19.16.2.3 Claim_i()

```
bool Mark3::Mutex::Claim_i (
    uint32_t u32WaitTimeMS_ ) [private]
```

Claim_i.

Abstracts out timed/non-timed mutex claim operations.

Parameters

<i>u32WaitTimeMS_</i>	Time in MS to wait, 0 for infinite
-----------------------	------------------------------------

Returns

true on successful claim, false otherwise

19.16.2.4 Init()

```
void Mark3::Mutex::Init (
    bool bRecursive_ = true )
```

Init.

Initialize a mutex object for use - must call this function before using the object.

Parameters

<i>bRecursive_</i>	Whether or not the mutex can be recursively locked.
--------------------	---

19.16.2.5 Release()

```
void Mark3::Mutex::Release ( )
```

Release.

Release the mutex. When the mutex is released, another object can enter the mutex-protected region.

If there are Threads waiting for the [Mutex](#) to become available, then the highest priority [Thread](#) will be unblocked at this time and will claim the [Mutex](#) lock immediately - this may result in an immediate context switch, depending on relative priorities.

If the calling [Thread](#)'s priority was boosted as a result of priority inheritance, the [Thread](#)'s previous priority will also be restored at this time.

Note that if a [Mutex](#) is held recursively, it must be Release'd the same number of times that it was Claim'd before it will be available for use by another [Thread](#).

Examples:

[lab5_mutexes/main.cpp](#).

19.16.2.6 WakeMe()

```
void Mark3::Mutex::WakeMe (
    Thread * pOwner_ )
```

WakeMe.

Wake a thread blocked on the mutex. This is an internal function used for implementing timed mutexes relying on timer callbacks. Since these do not have access to the private data of the mutex and its base classes, we have to wrap this as a public method - do not use this for any other purposes.

Parameters

<i>pOwner_</i>	Thread to unblock from this object.
----------------	---

19.16.2.7 WakeNext()

```
uint8_t Mark3::Mutex::WakeNext ( ) [private]
```

WakeNext.

Wake the next thread waiting on the [Mutex](#).

The documentation for this class was generated from the following file:

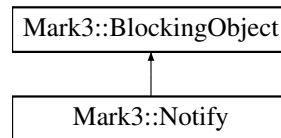
- [/home/moslevin/projects/github/m3-repo/kernel/src/public/mutex.h](#)

19.17 Mark3::Notify Class Reference

The [Notify](#) class is a blocking object type, that allows one or more threads to wait for an event to occur before resuming operation.

```
#include <notify.h>
```

Inheritance diagram for Mark3::Notify:



Public Member Functions

- void [Init](#) (void)
Init.
- void [Signal](#) (void)
Signal.
- void [Wait](#) (bool *pbFlag_)
Wait.
- bool [Wait](#) (uint32_t u32WaitTimeMS_, bool *pbFlag_)
Wait.
- void [WakeMe](#) ([Thread](#) *pclChosenOne_)
WakeMe.

Additional Inherited Members

19.17.1 Detailed Description

The [Notify](#) class is a blocking object type, that allows one or more threads to wait for an event to occur before resuming operation.

Examples:

[lab10_notifications/main.cpp](#).

Definition at line 33 of file [notify.h](#).

19.17.2 Member Function Documentation

19.17.2.1 Init()

```
void Mark3::Notify::Init (
    void )
```

Init.

Initialize the Notification object prior to use.

19.17.2.2 Signal()

```
void Mark3::Notify::Signal (
    void )
```

Signal.

Signal the notification object. This will cause the highest priority thread currently blocking on the object to wake. If no threads are currently blocked on the object, the call has no effect.

Examples:

[lab10_notifications/main.cpp](#).

19.17.2.3 Wait() [1/2]

```
void Mark3::Notify::Wait (
    bool * pbFlag_ )
```

Wait.

Block the current thread, waiting for a signal on the object.

Parameters

<i>pbFlag_</i>	Flag set to false on block, and true upon wakeup.
----------------	---

Examples:

[lab10_notifications/main.cpp](#).

19.17.2.4 Wait() [2/2]

```
bool Mark3::Notify::Wait (
    uint32_t u32WaitTimeMS_,
    bool * pbFlag_ )
```

Wait.

Block the current thread, waiting for a signal on the object.

Parameters

<i>u32WaitTimeM↔ S_</i>	Time to wait for the notification event.
<i>pbFlag_</i>	Flag set to false on block, and true upon wakeup.

Returns

true on notification, false on timeout

19.17.2.5 WakeMe()

```
void Mark3::Notify::WakeMe (
    Thread * pclChosenOne_ )
```

WakeMe.

Wake the specified thread from its current blocking queue. Note that this is only public in order to be accessible from a timer callack.

Parameters

<i>pclChosen↔ One_</i>	Thread to wake up
----------------------------	-------------------

The documentation for this class was generated from the following file:

- </home/moslevin/projects/github/m3-repo/kernel/src/public/notify.h>

19.18 Mark3::PriorityMap Class Reference

The [PriorityMap](#) class.

```
#include <priomap.h>
```

Public Member Functions

- [PriorityMap](#) ()
PriorityMap.
- void [Set](#) ([PORT_PRIO_TYPE](#) uXPrio_)
Set Set the priority map bitmap data, at all levels, for the given priority.
- void [Clear](#) ([PORT_PRIO_TYPE](#) uXPrio_)
Clear Clear the priority map bitmap data, at all levels, for the given priority.
- [PORT_PRIO_TYPE](#) [HighestPriority](#) (void)
HighestPriority.

19.18.1 Detailed Description

The [PriorityMap](#) class.

Definition at line 70 of file [priomap.h](#).

19.18.2 Constructor & Destructor Documentation

19.18.2.1 PriorityMap()

```
Mark3::PriorityMap::PriorityMap ( )
```

[PriorityMap](#).

Initialize the priority map object, clearing the bitmap data to all 0's.

Definition at line 54 of file [priomap.cpp](#).

19.18.3 Member Function Documentation

19.18.3.1 Clear()

```
void Mark3::PriorityMap::Clear (
    PORT_PRIO_TYPE uXPrio_ )
```

Clear Clear the priority map bitmap data, at all levels, for the given priority.

Parameters

<i>uX_← Prio_</i>	Priority level to clear the bitmap data for.
---------------------------------	--

Definition at line 81 of file [priomap.cpp](#).

19.18.3.2 HighestPriority()

```
PORT_PRIO_TYPE Mark3::PriorityMap::HighestPriority (
    void )
```

HighestPriority.

Computes the numeric priority of the highest-priority thread represented in the priority map.

Returns

Highest priority ready-thread's number.

Definition at line 97 of file [priomap.cpp](#).

19.18.3.3 Set()

```
void Mark3::PriorityMap::Set (
    PORT_PRIO_TYPE uXPrio_ )
```

Set Set the priority map bitmap data, at all levels, for the given priority.

Parameters

<i>uX_← Prio_</i>	Priority level to set the bitmap data for.
---------------------------------	--

Definition at line 67 of file [priomap.cpp](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/github/m3-repo/kernel/src/public/priomap.h](#)
- [/home/moslevin/projects/github/m3-repo/kernel/src/priomap.cpp](#)

19.19 Mark3::ProfileTimer Class Reference

Profiling timer.

```
#include <profile.h>
```

Public Member Functions

- void [Init](#) ()
Init.
- void [Start](#) ()
Start.
- void [Stop](#) ()
Stop.
- uint32_t [GetAverage](#) ()
GetAverage.
- uint32_t [GetCurrent](#) ()
GetCurrent.

Private Member Functions

- uint32_t [ComputeCurrentTicks](#) (uint16_t u16Current_, uint32_t u32Epoch_)
ComputeCurrentTicks.

Private Attributes

- uint32_t [m_u32Cumulative](#)
Cumulative tick-count for this timer.
- uint32_t [m_u32CurrentIteration](#)
Tick-count for the current iteration.
- uint16_t [m_u16Initial](#)
Initial count.
- uint32_t [m_u32InitialEpoch](#)
Initial Epoch.
- uint16_t [m_u16Iterations](#)
Number of iterations executed for this profiling timer.
- bool [m_bActive](#)
Whether or not the timer is active or stopped.

19.19.1 Detailed Description

Profiling timer.

This class is used to perform high-performance profiling of code to see how int32_t certain operations take. useful in instrumenting the performance of key algorithms and time-critical operations to ensure real-timer behavior.

Definition at line 69 of file [profile.h](#).

19.19.2 Member Function Documentation

19.19.2.1 ComputeCurrentTicks()

```
uint32_t Mark3::ProfileTimer::ComputeCurrentTicks (
    uint16_t u16Current_,
    uint32_t u32Epoch_ ) [private]
```

ComputeCurrentTicks.

Figure out how many ticks have elapsed in this iteration

Parameters

<i>u16↵</i> <i>Count_</i>	Current timer count
<i>u32↵</i> <i>Epoch_</i>	Current timer epoch

Returns

Current tick count

19.19.2.2 GetAverage()

```
uint32_t Mark3::ProfileTimer::GetAverage ( )
```

GetAverage.

Get the average time associated with this operation.

Returns

Average tick count normalized over all iterations

19.19.2.3 GetCurrent()

```
uint32_t Mark3::ProfileTimer::GetCurrent ( )
```

GetCurrent.

Return the current tick count held by the profiler. Valid for both active and stopped timers.

Returns

The currently held tick count.

19.19.2.4 Init()

```
void Mark3::ProfileTimer::Init ( )
```

Init.

Initialize the profiling timer prior to use. Can also be used to reset a timer that's been used previously.

19.19.2.5 Start()

```
void Mark3::ProfileTimer::Start ( )
```

Start.

Start a profiling session, if the timer is not already active. Has no effect if the timer is already active.

19.19.2.6 Stop()

```
void Mark3::ProfileTimer::Stop ( )
```

Stop.

Stop the current profiling session, adding to the cumulative time for this timer, and the total iteration count.

The documentation for this class was generated from the following file:

- </home/moslevin/projects/github/m3-repo/kernel/src/public/profile.h>

19.20 Mark3::Quantum Class Reference

Static-class used to implement [Thread](#) quantum functionality, which is a key part of round-robin scheduling.

```
#include <quantum.h>
```

Static Public Member Functions

- static void [UpdateTimer](#) ()
UpdateTimer.
- static void [AddThread](#) ([Thread](#) *pciThread_)
AddThread.
- static void [RemoveThread](#) ()
RemoveThread.
- static void [SetInTimer](#) (void)
SetInTimer.
- static void [ClearInTimer](#) (void)
ClearInTimer.

Static Private Member Functions

- static void [SetTimer](#) ([Thread](#) *pciThread_)
SetTimer.

19.20.1 Detailed Description

Static-class used to implement [Thread](#) quantum functionality, which is a key part of round-robin scheduling.

Definition at line [42](#) of file [quantum.h](#).

19.20.2 Member Function Documentation

19.20.2.1 AddThread()

```
static void Mark3::Quantum::AddThread (
    Thread * pClThread_ ) [static]
```

AddThread.

Add the thread to the quantum timer. Only one thread can own the quantum, since only one thread can be running on a core at a time.

19.20.2.2 ClearInTimer()

```
static void Mark3::Quantum::ClearInTimer (
    void ) [inline], [static]
```

ClearInTimer.

Clear the flag once the timer callback function has been completed.

Definition at line 84 of file [quantum.h](#).

19.20.2.3 RemoveThread()

```
static void Mark3::Quantum::RemoveThread ( ) [static]
```

RemoveThread.

Remove the thread from the quantum timer. This will cancel the timer.

19.20.2.4 SetInTimer()

```
static void Mark3::Quantum::SetInTimer (
    void ) [inline], [static]
```

SetInTimer.

Set a flag to indicate that the CPU is currently running within the timer-callback routine. This prevents the [Quantum](#) timer from being updated in the middle of a callback cycle, potentially resulting in the kernel timer becoming disabled.

Definition at line 78 of file [quantum.h](#).

19.20.2.5 SetTimer()

```
static void Mark3::Quantum::SetTimer (
    Thread * pClThread_ ) [static], [private]
```

SetTimer.

Set up the quantum timer in the timer scheduler. This creates a one-shot timer, which calls a static callback in [quantum.cpp](#) that on expiry will pivot the head of the threadlist for the thread's priority. This is the mechanism that provides round-robin scheduling in the system.

Parameters

<i>pcl</i> ↔ <i>Thread_</i>	Pointer to the thread to set the Quantum timer on
--------------------------------	---

19.20.2.6 UpdateTimer()

```
static void Mark3::Quantum::UpdateTimer ( ) [static]
```

UpdateTimer.

This function is called to update the thread quantum timer whenever something in the scheduler has changed. This can result in the timer being re-loaded or started. The timer is never stopped, but it may be ignored on expiry.

The documentation for this class was generated from the following file:

- </home/moslevin/projects/github/m3-repo/kernel/src/public/quantum.h>

19.21 Mark3::ReaderWriterLock Class Reference

The [ReaderWriterLock](#) class This class implements an object that marshalls access to a resource based on the intended usage of the resource.

```
#include <readerwriter.h>
```

Public Member Functions

- void [Init](#) ()
Init Initialize the reader-writer lock before use.
- void [AcquireReader](#) ()
AcquireReader Acquire the object's reader lock.
- bool [AcquireReader](#) (uint32_t u32TimeoutMs_)
AcquireReader Acquire the object's reader lock.
- void [ReleaseReader](#) ()
ReleaseReader Release a previously-held reader lock.
- void [AcquireWriter](#) ()
AcquireWriter Acquire the writer lock.
- bool [AcquireWriter](#) (uint32_t u32TimeoutMs_)
AcquireWriter Acquire the writer lock.
- void [ReleaseWriter](#) ()
ReleaseWriter Release the currently held writer, allowing other readers/writers to access the object.

Private Member Functions

- bool [AcquireReader_i](#) (uint32_t u32TimeoutMs_)
AcquireReader_i Internal helper function for AcquireReader.
- bool [AcquireWriter_i](#) (uint32_t u32TimeoutMs_)
AcquireWriter_i Internal helper function for AcquireWriter.

Private Attributes

- [Mutex m_clGlobalMutex](#)
Mutex used to lock the object against concurrent read + write.
- [Mutex m_clReaderMutex](#)
Mutex used to lock object for readers.
- [uint8_t m_u8ReadCount](#)
Number of concurrent readers.

19.21.1 Detailed Description

The [ReaderWriterLock](#) class This class implements an object that marshalls access to a resource based on the intended usage of the resource.

A reader-writer lock permits multiple concurrent read access, or single-writer access to a resource. If the object holds a write lock, other writers, and all readers will block until the writer is finished. If the object holds reader locks, all writers will block until all readers are finished before the first writer can take ownership of the resource. This is based upon lower-level synchronization primitives, and is somewhat more heavyweight than primitive synchronization types.

Definition at line 41 of file [readerwriter.h](#).

19.21.2 Member Function Documentation

19.21.2.1 AcquireReader() [1/2]

```
void Mark3::ReaderWriterLock::AcquireReader ( )
```

AcquireReader Acquire the object's reader lock.

Multiple concurrent readers are allowed. If the writer lock is currently held, the calling thread will wait until the writer lock is relinquished

19.21.2.2 AcquireReader() [2/2]

```
bool Mark3::ReaderWriterLock::AcquireReader (
    uint32_t u32TimeoutMs_ )
```

AcquireReader Acquire the object's reader lock.

Multiple concurrent readers are allowed. If the writer lock is currently held, the calling thread will wait until the writer lock is relinquished

Parameters

<i>u32TimeoutMs_</i>	Maximum time to wait (in ms) before the operation is aborted
----------------------	--

Returns

true on success, false on timeout

19.21.2.3 AcquireReader_i()

```
bool Mark3::ReaderWriterLock::AcquireReader_i (
    uint32_t u32TimeoutMs_ ) [private]
```

AcquireReader_i Internal helper function for AcquireReaer.

Parameters

<i>u32TimeoutMs_</i>	Maximum time to wait (in ms) before the operation is aborted
----------------------	--

Returns

true on success, false on timeout

19.21.2.4 AcquireWriter() [1/2]

```
void Mark3::ReaderWriterLock::AcquireWriter ( )
```

AcquireWriter Acquire the writer lock.

Only a single writer is allowed to access the object at a time. This will block the currently-running thread until all other readers/writers have released their locks.

19.21.2.5 AcquireWriter() [2/2]

```
bool Mark3::ReaderWriterLock::AcquireWriter (
    uint32_t u32TimeoutMs_ )
```

AcquireWriter Acquire the writer lock.

Only a single writer is allowed to access the object at a time. This will block the currently-running thread until all other readers/writers have released their locks.

Parameters

<i>u32TimeoutMs_</i>	Maximum time to wait (in ms) before the operation is aborted
----------------------	--

Returns

true on success, false on timeout

19.21.2.6 AcquireWriter_i()

```
bool Mark3::ReaderWriterLock::AcquireWriter_i (
    uint32_t u32TimeoutMs_ ) [private]
```

AcquireWriter_i Internal helper function for AcquireWriter.

Parameters

<i>u32TimeoutMs_</i>	Maximum time to wait (in ms) before the operation is aborted
----------------------	--

Returns

true on success, false on timeout

19.21.2.7 Init()

```
void Mark3::ReaderWriterLock::Init ( )
```

Init Initialize the reader-writer lock before use.

Must be called before attempting any other operations on the object.

The documentation for this class was generated from the following file:

- /home/moslevin/projects/github/m3-repo/kernel/src/public/[readerwriter.h](#)

19.22 Mark3::Scheduler Class Reference

Priority-based round-robin [Thread](#) scheduling, using ThreadLists for housekeeping.

```
#include <scheduler.h>
```


Static Public Member Functions

- static void [Init](#) ()
Init.
- static void [Schedule](#) ()
Schedule.
- static void [Add](#) ([Thread](#) *pciThread_)
Add.
- static void [Remove](#) ([Thread](#) *pciThread_)
Remove.
- static bool [SetScheduler](#) (bool bEnable_)
SetScheduler.
- static [Thread](#) * [GetCurrentThread](#) ()
GetCurrentThread.
- static volatile [Thread](#) * [GetNextThread](#) ()
GetNextThread.
- static [ThreadList](#) * [GetThreadList](#) ([PORT_PRIO_TYPE](#) uXPriority_)
GetThreadList.
- static [ThreadList](#) * [GetStopList](#) ()
GetStopList.
- static bool [IsEnabled](#) ()
IsEnabled.
- static void [QueueScheduler](#) ()
QueueScheduler.

Static Private Attributes

- static bool [m_bEnabled](#)
[Scheduler](#)'s state - enabled or disabled.
- static bool [m_bQueuedSchedule](#)
Variable representing whether or not there's a queued scheduler operation.
- static [ThreadList](#) [m_clStopList](#)
[ThreadList](#) for all stopped threads.
- static [ThreadList](#) [m_aclPriorities](#) [[KERNEL_NUM_PRIORITIES](#)]
[ThreadLists](#) for all threads at all priorities.
- static [PriorityMap](#) [m_clPrioMap](#)
Priority bitmap lookup structure, 1-bit per thread priority.

19.22.1 Detailed Description

Priority-based round-robin [Thread](#) scheduling, using [ThreadLists](#) for housekeeping.

Definition at line 63 of file [scheduler.h](#).

19.22.2 Member Function Documentation

19.22.2.1 Add()

```
void Mark3::Scheduler::Add (  
    Thread * pclThread_ ) [static]
```

Add.

Add a thread to the scheduler at its current priority level.

Parameters

<i>pthread_</i> <i>Thread_</i>	Pointer to the thread to add to the scheduler
-----------------------------------	---

Definition at line 86 of file [scheduler.cpp](#).

19.22.2.2 GetCurrentThread()

```
static Thread* Mark3::Scheduler::GetCurrentThread ( ) [inline], [static]
```

GetCurrentThread.

Return the pointer to the currently-running thread.

Returns

Pointer to the currently-running thread

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 123 of file [scheduler.h](#).

19.22.2.3 GetNextThread()

```
static volatile Thread* Mark3::Scheduler::GetNextThread ( ) [inline], [static]
```

GetNextThread.

Return the pointer to the thread that should run next, according to the last run of the scheduler.

Returns

Pointer to the next-running thread

Definition at line 132 of file [scheduler.h](#).

19.22.2.4 GetStopList()

```
static ThreadList* Mark3::Scheduler::GetStopList ( ) [inline], [static]
```

GetStopList.

Return the pointer to the list of threads that are in the scheduler's stopped state.

Returns

Pointer to the [ThreadList](#) containing the stopped threads

Definition at line 152 of file [scheduler.h](#).

19.22.2.5 GetThreadList()

```
static ThreadList* Mark3::Scheduler::GetThreadList (
    PORT_PRIO_TYPE uXPriority_ ) [inline], [static]
```

GetThreadList.

Return the pointer to the active list of threads that are at the given priority level in the scheduler.

Parameters

$uX \leftrightarrow$ <i>Priority_</i>	Priority level of the threadlist
--	----------------------------------

Returns

Pointer to the [ThreadList](#) for the given priority level

Definition at line 143 of file [scheduler.h](#).

19.22.2.6 Init()

```
void Mark3::Scheduler::Init (
    void ) [static]
```

Init.

Intiaillize the scheduler, must be called before use.

Definition at line 53 of file [scheduler.cpp](#).

19.22.2.7 IsEnabled()

```
static bool Mark3::Scheduler::IsEnabled ( ) [inline], [static]
```

IsEnabled.

Return the current state of the scheduler - whether or not scheduling is enabled or disabled.

Returns

true - scheduler enabled, false - disabled

Definition at line 161 of file [scheduler.h](#).

19.22.2.8 QueueScheduler()

```
static void Mark3::Scheduler::QueueScheduler ( ) [inline], [static]
```

QueueScheduler.

Tell the kernel to perform a scheduling operation as soon as the scheduler is re-enabled.

Definition at line 168 of file [scheduler.h](#).

19.22.2.9 Remove()

```
void Mark3::Scheduler::Remove (
    Thread * pclThread_ ) [static]
```

Remove.

Remove a thread from the scheduler at its current priority level.

Parameters

<i>pc</i> ↔ <i>Thread_</i>	Pointer to the thread to be removed from the scheduler
-------------------------------	--

Definition at line 92 of file [scheduler.cpp](#).

19.22.2.10 Schedule()

```
void Mark3::Scheduler::Schedule ( ) [static]
```

Schedule.

Run the scheduler, determines the next thread to run based on the current state of the threads. Note that the next-thread chosen from this function is only valid while in a critical section.

Definition at line 62 of file [scheduler.cpp](#).

19.22.2.11 SetScheduler()

```
bool Mark3::Scheduler::SetScheduler (
    bool bEnable_ ) [static]
```

SetScheduler.

Set the active state of the scheduler. When the scheduler is disabled, the *next thread* is never set; the currently running thread will run forever until the scheduler is enabled again. Care must be taken to ensure that we don't end up trying to block while the scheduler is disabled, otherwise the system ends up in an unusable state.

Parameters

<i>b</i> ↔ <i>Enable</i> ↔ —	true to enable, false to disable the scheduler
------------------------------------	--

Definition at line 98 of file [scheduler.cpp](#).

The documentation for this class was generated from the following files:

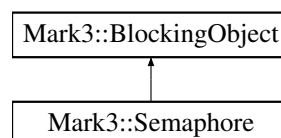
- [/home/moslevin/projects/github/m3-repo/kernel/src/public/scheduler.h](#)
- [/home/moslevin/projects/github/m3-repo/kernel/src/scheduler.cpp](#)

19.23 Mark3::Semaphore Class Reference

Binary & Counting semaphores, based on [BlockingObject](#) base class.

```
#include <ksemaphore.h>
```

Inheritance diagram for Mark3::Semaphore:



Public Member Functions

- void [Init](#) (uint16_t u16InitVal_, uint16_t u16MaxVal_)
Initialize a semaphore before use.
- bool [Post](#) ()
Increment the semaphore count.
- void [Pend](#) ()
Decrement the semaphore count.
- uint16_t [GetCount](#) ()
Return the current semaphore counter.
- bool [Pend](#) (uint32_t u32WaitTimeMS_)
Decrement the semaphore count.
- void [WakeMe](#) ([Thread](#) *pclChosenOne_)
Wake a thread blocked on the semaphore.

Private Member Functions

- uint8_t [WakeNext](#) ()
Wake the next thread waiting on the semaphore.
- bool [Pend_i](#) (uint32_t u32WaitTimeMS_)
Pend_i.

Private Attributes

- uint16_t [m_u16Value](#)
Current count held by the semaphore.
- uint16_t [m_u16MaxValue](#)
Maximum count that can be held by this semaphore.

Additional Inherited Members

19.23.1 Detailed Description

Binary & Counting semaphores, based on [BlockingObject](#) base class.

Examples:

[lab4_semaphores/main.cpp](#), [lab6_timers/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 36 of file [ksemaphore.h](#).

19.23.2 Member Function Documentation

19.23.2.1 GetCount()

```
uint16_t Mark3::Semaphore::GetCount ( )
```

Return the current semaphore counter.

This can be used by a thread to bypass blocking on a semaphore - allowing it to do other things until a non-zero count is returned, instead of blocking until the semaphore is posted.

Returns

The current semaphore counter value.

19.23.2.2 Init()

```
void Mark3::Semaphore::Init (
    uint16_t u16InitVal_,
    uint16_t u16MaxVal_ )
```

Initialize a semaphore before use.

Must be called before attempting post/pend operations on the object.

This initialization is required to configure the behavior of the semaphore with regards to the initial and maximum values held by the semaphore. By providing access to the raw initial and maximum count elements of the semaphore, these objects are able to be used as either counting or binary semaphores.

To configure a semaphore object for use as a binary semaphore, set values of 0 and 1 respectively for the initial/maximum value parameters.

Any other combination of values can be used to implement a counting semaphore.

Parameters

<i>u16InitVal_</i>	Initial value held by the semaphore
<i>u16MaxVal_</i>	Maximum value for the semaphore. Must be nonzero.

Examples:

[lab6_timers/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

19.23.2.3 Pend() [1/2]

```
void Mark3::Semaphore::Pend ( )
```


Decrement the semaphore count.

If the count is zero, the calling [Thread](#) will block until the semaphore is posted, and the [Thread](#)'s priority is higher than that of any other [Thread](#) blocked on the object.

Examples:

[lab4_semaphores/main.cpp](#).

19.23.2.4 `Pend()` [2/2]

```
bool Mark3::Semaphore::Pend (
    uint32_t u32WaitTimeMS_ )
```

Decrement the semaphore count.

If the count is zero, the thread will block until the semaphore is pended. If the specified interval expires before the thread is unblocked, then the status is returned back to the user.

Returns

true - semaphore was acquired before the timeout false - timeout occurred before the semaphore was claimed.

19.23.2.5 `Pend_i()`

```
bool Mark3::Semaphore::Pend_i (
    uint32_t u32WaitTimeMS_ ) [private]
```

`Pend_i`.

Internal function used to abstract timed and untimed semaphore pend operations.

Parameters

<code>u32WaitTimeMS_</code>	Time in MS to wait
-----------------------------	--------------------

Returns

true on success, false on failure.

19.23.2.6 `Post()`

```
bool Mark3::Semaphore::Post ( )
```

Increment the semaphore count.

If the semaphore count is zero at the time this is called, and there are threads blocked on the object, this will immediately unblock the highest-priority blocked [Thread](#).

Note that if the priority of that [Thread](#) is higher than the current thread's priority, a context switch will occur and control will be relinquished to that [Thread](#).

Returns

true if the semaphore was posted, false if the count is already maxed out.

Examples:

[lab4_semaphores/main.cpp](#), and [lab6_timers/main.cpp](#).

19.23.2.7 WakeMe()

```
void Mark3::Semaphore::WakeMe (
    Thread * pChosenOne_ )
```

Wake a thread blocked on the semaphore.

This is an internal function used for implementing timed semaphores relying on timer callbacks. Since these do not have access to the private data of the semaphore and its base classes, we have to wrap this as a public method - do not use this for any other purposes.

19.23.2.8 WakeNext()

```
uint8_t Mark3::Semaphore::WakeNext ( ) [private]
```

Wake the next thread waiting on the semaphore.

Used internally.

The documentation for this class was generated from the following file:

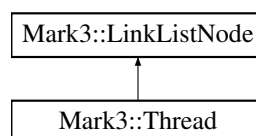
- [/home/moslevin/projects/github/m3-repo/kernel/src/public/ksemaphore.h](#)

19.24 Mark3::Thread Class Reference

Object providing fundamental multitasking support in the kernel.

```
#include <thread.h>
```

Inheritance diagram for Mark3::Thread:



Public Member Functions

- void [Init](#) ([K_WORD](#) *pwStack_, [uint16_t](#) u16StackSize_, [PORT_PRIO_TYPE](#) uXPriority_, [ThreadEntryFunc](#) pfEntryPoint_, void *pvArg_)
Init.
- void [Start](#) ()
Start.
- void [Stop](#) ()
Stop.
- [ThreadList](#) * [GetOwner](#) (void)
GetOwner.
- [ThreadList](#) * [GetCurrent](#) (void)
GetCurrent.
- [PORT_PRIO_TYPE](#) [GetPriority](#) (void)
GetPriority.
- [PORT_PRIO_TYPE](#) [GetCurPriority](#) (void)
GetCurPriority.
- void [SetQuantum](#) ([uint16_t](#) u16Quantum_)
SetQuantum.
- [uint16_t](#) [GetQuantum](#) (void)
GetQuantum.
- void [SetCurrent](#) ([ThreadList](#) *pclNewList_)
SetCurrent.
- void [SetOwner](#) ([ThreadList](#) *pclNewList_)
SetOwner.
- void [SetPriority](#) ([PORT_PRIO_TYPE](#) uXPriority_)
SetPriority.
- void [InheritPriority](#) ([PORT_PRIO_TYPE](#) uXPriority_)
InheritPriority.
- void [Exit](#) ()
Exit.
- void [SetID](#) ([uint8_t](#) u8ID_)
SetID.
- [uint8_t](#) [GetID](#) ()
GetID.
- [uint16_t](#) [GetStackSlack](#) ()
GetStackSlack.
- [uint16_t](#) [GetEventFlagMask](#) ()
GetEventFlagMask returns the thread's current event-flag mask, which is used in conjunction with the [EventFlag](#) blocking object type.
- void [SetEventFlagMask](#) ([uint16_t](#) u16Mask_)
SetEventFlagMask Sets the active event flag bitfield mask.
- void [SetEventFlagMode](#) ([EventFlagOperation](#) eMode_)
SetEventFlagMode Sets the active event flag operation mode.
- [EventFlagOperation](#) [GetEventFlagMode](#) ()
GetEventFlagMode Returns the thread's event flag's operating mode.
- [Timer](#) * [GetTimer](#) ()
Return a pointer to the thread's timer object.
- void [SetExpired](#) (bool bExpired_)
SetExpired.
- bool [GetExpired](#) ()

- *GetExpired.*
- void [InitIdle](#) ()
 - InitIdle Initialize this [Thread](#) object as the [Kernel](#)'s idle thread.*
- void * [GetExtendedContext](#) ()
 - GetExtendedContext.*
- void [SetExtendedContext](#) (void *pvData_)
 - SetExtendedContext.*
- [ThreadState](#) [GetState](#) ()
 - GetState Returns the current state of the thread to the caller.*
- void [SetState](#) ([ThreadState](#) eState_)
 - SetState Set the thread's state to a new value.*
- [K_WORD](#) * [GetStack](#) ()
 - GetStack.*
- [uint16_t](#) [GetStackSize](#) ()
 - GetStackSize.*

Static Public Member Functions

- static [Thread](#) * [Init](#) ([uint16_t](#) u16StackSize_, [PORT_PRIO_TYPE](#) uXPriority_, [ThreadEntryFunc](#) pfEntry↵
Point_, void *pvArg_)
 - Init.*
- static void [Sleep](#) ([uint32_t](#) u32TimeMs_)
 - Sleep.*
- static void [USleep](#) ([uint32_t](#) u32TimeUs_)
 - USleep.*
- static void [Yield](#) (void)
 - Yield.*

Private Member Functions

- void [SetPriorityBase](#) ([PORT_PRIO_TYPE](#) uXPriority_)
 - SetPriorityBase.*

Static Private Member Functions

- static void [ContextSwitchSWI](#) (void)
 - ContextSwitchSWI.*

Private Attributes

- [K_WORD](#) * [m_pwStackTop](#)
 - Pointer to the top of the thread's stack.*
- [K_WORD](#) * [m_pwStack](#)
 - Pointer to the thread's stack.*
- [uint8_t](#) [m_u8ThreadID](#)
 - [Thread](#) ID.*
- [PORT_PRIO_TYPE](#) [m_uXPriority](#)
 - Default priority of the thread.*

- [PORT_PRIO_TYPE m_uXCurPriority](#)
Current priority of the thread (priority inheritance)
- [ThreadState m_eState](#)
Enum indicating the thread's current state.
- `void * m_pvExtendedContext`
Pointer provided to a [Thread](#) to implement thread-local storage.
- `uint16_t m_u16StackSize`
Size of the stack (in bytes)
- `ThreadList * m_pclCurrent`
Pointer to the thread-list where the thread currently resides.
- `ThreadList * m_pclOwner`
Pointer to the thread-list where the thread resides when active.
- `ThreadEntryFunc m_pfEntryPoint`
The entry-point function called when the thread starts.
- `void * m_pvArg`
Pointer to the argument passed into the thread's entrypoint.
- `uint16_t m_u16Quantum`
[Thread](#) quantum (in milliseconds)
- `uint16_t m_u16FlagMask`
Event-flag mask.
- `EventFlagOperation m_eFlagMode`
Event-flag mode.
- `Timer m_clTimer`
[Timer](#) used for blocking-object timeouts.
- `bool m_bExpired`
Indicate whether or not a blocking-object timeout has occurred.

Additional Inherited Members

19.24.1 Detailed Description

Object providing fundamental multitasking support in the kernel.

Examples:

[lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_idle_↵
function/main.cpp](#), [lab3_round_robin/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab5_mutexes/main.cpp](#), [lab6_↵
timers/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 62 of file [thread.h](#).

19.24.2 Member Function Documentation

19.24.2.1 ContextSwitchSWI()

```
void Mark3::Thread::ContextSwitchSWI (
    void ) [static], [private]
```

ContextSwitchSWI.

This code is used to trigger the context switch interrupt. Called whenever the kernel decides that it is necessary to swap out the current thread for the "next" thread.

Definition at line 493 of file [thread.cpp](#).

19.24.2.2 Exit()

```
void Mark3::Thread::Exit ( )
```

Exit.

Remove the thread from being scheduled again. The thread is effectively destroyed when this occurs. This is extremely useful for cases where a thread encounters an unrecoverable error and needs to be restarted, or in the context of systems where threads need to be created and destroyed dynamically.

This must not be called on the idle thread.

Examples:

[lab9_dynamic_threads/main.cpp](#).

19.24.2.3 GetCurPriority()

```
PORT_PRIO_TYPE Mark3::Thread::GetCurPriority (
    void ) [inline]
```

GetCurPriority.

Return the priority of the current thread

Returns

Priority of the current thread

Definition at line 184 of file [thread.h](#).

19.24.2.4 GetCurrent()

```
ThreadList* Mark3::Thread::GetCurrent (
    void ) [inline]
```

GetCurrent.

Return the [ThreadList](#) where the thread is currently located

Returns

Pointer to the thread's current list

Definition at line 167 of file [thread.h](#).

19.24.2.5 GetEventFlagMask()

```
uint16_t Mark3::Thread::GetEventFlagMask ( ) [inline]
```

GetEventFlagMask returns the thread's current event-flag mask, which is used in conjunction with the [EventFlag](#) blocking object type.

Returns

A copy of the thread's event flag mask

Definition at line 329 of file [thread.h](#).

19.24.2.6 GetEventFlagMode()

```
EventFlagOperation Mark3::Thread::GetEventFlagMode ( ) [inline]
```

GetEventFlagMode Returns the thread's event flag's operating mode.

Returns

The thread's event flag mode.

Definition at line 345 of file [thread.h](#).

19.24.2.7 GetExpired()

```
bool Mark3::Thread::GetExpired ( )
```

GetExpired.

Return the status of the most-recent blocking call on the thread.

Returns

true - call expired, false - call did not expire

19.24.2.8 GetExtendedContext()

```
void* Mark3::Thread::GetExtendedContext ( ) [inline]
```

GetExtendedContext.

Return the [Thread](#) object's extended-context data pointer. Used by code implementing a user-defined thread-local storage model. Pointer exists only for the lifespan of the [Thread](#).

Returns

[Thread](#)'s extended context data pointer.

Definition at line [393](#) of file [thread.h](#).

19.24.2.9 GetID()

```
uint8_t Mark3::Thread::GetID ( ) [inline]
```

GetID.

Return the 8-bit ID corresponding to this thread.

Returns

[Thread](#)'s 8-bit ID, set by the user

Definition at line [306](#) of file [thread.h](#).

19.24.2.10 GetOwner()

```
ThreadList* Mark3::Thread::GetOwner (
    void ) [inline]
```

GetOwner.

Return the [ThreadList](#) where the thread belongs when it's in the active/ready state in the scheduler.

Returns

Pointer to the [Thread](#)'s owner list

Definition at line 159 of file [thread.h](#).

19.24.2.11 GetPriority()

```
PORT_PRIO_TYPE Mark3::Thread::GetPriority (
    void ) [inline]
```

GetPriority.

Return the priority of the current thread

Returns

Priority of the current thread

Definition at line 176 of file [thread.h](#).

19.24.2.12 GetQuantum()

```
uint16_t Mark3::Thread::GetQuantum (
    void ) [inline]
```

GetQuantum.

Get the thread's round-robin execution quantum.

Returns

The thread's quantum

Definition at line 201 of file [thread.h](#).

19.24.2.13 GetStack()

```
K_WORD* Mark3::Thread::GetStack ( ) [inline]
```

GetStack.

Returns

Pointer to the blob of memory used as the thread's stack

Definition at line 429 of file [thread.h](#).

19.24.2.14 GetStackSize()

```
uint16_t Mark3::Thread::GetStackSize ( ) [inline]
```

GetStackSize.

Returns

Size of the thread's stack in bytes

Definition at line 435 of file [thread.h](#).

19.24.2.15 GetStackSlack()

```
uint16_t Mark3::Thread::GetStackSlack ( )
```

GetStackSlack.

Performs a (somewhat lengthy) check on the thread stack to check the amount of stack margin (or "slack") remaining on the stack. If you're having problems with blowing your stack, you can run this function at points in your code during development to see what operations cause problems. Also useful during development as a tool to optimally size thread stacks.

Returns

The amount of slack (unused bytes) on the stack

ToDo : Reverse the logic for MCUs where stack grows UP instead of down

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 352 of file [thread.cpp](#).

19.24.2.16 GetState()

```
ThreadState Mark3::Thread::GetState ( ) [inline]
```

GetState Returns the current state of the thread to the caller.

Can be used to determine whether or not a thread is ready (or running), stopped, or terminated/exit'd.

Returns

ThreadState_t representing the thread's current state

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 415 of file [thread.h](#).

19.24.2.17 InheritPriority()

```
void Mark3::Thread::InheritPriority (
    PORT_PRIO_TYPE uXPriority_ )
```

InheritPriority.

Allow the thread to run at a different priority level (temporarily) for the purpose of avoiding priority inversions. This should only be called from within the implementation of blocking-objects.

Parameters

<i>uX_← Priority_</i>	New Priority to boost to.
-------------------------------------	---------------------------

Definition at line 482 of file [thread.cpp](#).

19.24.2.18 Init() [1/2]

```
void Mark3::Thread::Init (
    K_WORD * pwStack_,
    uint16_t u16StackSize_,
    PORT_PRIO_TYPE uXPriority_,
    ThreadEntryFunc pfEntryPoint_,
    void * pvArg_ )
```

Init.

Initialize a thread prior to its use. Initialized threads are placed in the stopped state, and are not scheduled until the thread's start method has been invoked first.

Parameters

<i>pwStack_</i>	Pointer to the stack to use for the thread
<i>u16StackSize_</i>	Size of the stack (in bytes)
<i>uXPriority_</i>	Priority of the thread (0 = idle, 7 = max)
<i>pfEntryPoint_</i> —	This is the function that gets called when the thread is started
<i>pvArg_</i>	Pointer to the argument passed into the thread's entrypoint function.

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 72 of file [thread.cpp](#).

19.24.2.19 Init() [2/2]

```
static Thread* Mark3::Thread::Init (
    uint16_t u16StackSize_,
    PORT_PRIO_TYPE uXPriority_,
    ThreadEntryFunc pfEntryPoint_,
    void * pvArg_ ) [static]
```

Init.

Create and initialize a new thread, using memory from the auto-allocated heap region to supply both the thread object and its stack. The thread returned can then be started using the [Start\(\)](#) method directly. Note that the memory used to create this thread cannot be reclaimed, and so this API is only suitable for threads that exist for the duration of runtime.

Parameters

<i>u16StackSize_</i>	Size of the stack (in bytes)
<i>uXPriority_</i>	Priority of the thread (0 = idle, 7 = max)
<i>pfEntryPoint_</i> —	This is the function that gets called when the thread is started
<i>pvArg_</i>	Pointer to the argument passed into the thread's entrypoint function.

Returns

Pointer to a newly-created thread.

19.24.2.20 InitIdle()

```
void Mark3::Thread::InitIdle ( )
```

InitIdle Initialize this [Thread](#) object as the [Kernel](#)'s idle thread.

There should only be one of these, maximum, in a given system.

19.24.2.21 SetCurrent()

```
void Mark3::Thread::SetCurrent (
    ThreadList * pclNewList_ ) [inline]
```

SetCurrent.

Set the thread's current to the specified thread list

Parameters

<i>pcNewList_</i>	Pointer to the threadlist to apply thread ownership
-------------------	---

Definition at line 211 of file [thread.h](#).

19.24.2.22 SetEventFlagMask()

```
void Mark3::Thread::SetEventFlagMask (
    uint16_t u16Mask_ ) [inline]
```

SetEventFlagMask Sets the active event flag bitfield mask.

Parameters

<i>u16Mask_</i>	
-----------------	--

Definition at line 334 of file [thread.h](#).

19.24.2.23 SetEventFlagMode()

```
void Mark3::Thread::SetEventFlagMode (
    EventFlagOperation eMode_ ) [inline]
```

SetEventFlagMode Sets the active event flag operation mode.

Parameters

<i>eMode_</i>	Event flag operation mode, defines the logical operator to apply to the event flag.
---------------	---

Definition at line 340 of file [thread.h](#).

19.24.2.24 SetExpired()

```
void Mark3::Thread::SetExpired (
    bool bExpired_ )
```

SetExpired.

Set the status of the current blocking call on the thread.

Parameters

<i>bExpired_</i>	true - call expired, false - call did not expire
—	

19.24.2.25 SetExtendedContext()

```
void Mark3::Thread::SetExtendedContext (
    void * pvData_ ) [inline]
```

SetExtendedContext.

Assign the [Thread](#) object's extended-context data pointer. Used by code implementing a user-defined thread-local storage model.

Object assigned to the context pointer should persist for the duration of the [Thread](#).

Parameters

<i>pvData_</i>	Object to assign to the extended data pointer.+
<i>Data_</i>	

Definition at line 406 of file [thread.h](#).

19.24.2.26 SetID()

```
void Mark3::Thread::SetID (
    uint8_t u8ID_ ) [inline]
```

SetID.

Set an 8-bit ID to uniquely identify this thread.

Parameters

<i>u8Id_</i>	8-bit Thread ID, set by the user
--------------	--

Definition at line 298 of file [thread.h](#).

19.24.2.27 SetOwner()

```
void Mark3::Thread::SetOwner (
    ThreadList * pclNewList_ ) [inline]
```

SetOwner.

Set the thread's owner to the specified thread list

Parameters

<i>pcNewList_</i>	Pointer to the threadlist to apply thread ownership
-------------------	---

Definition at line 219 of file [thread.h](#).

19.24.2.28 SetPriority()

```
void Mark3::Thread::SetPriority (
    PORT_PRIO_TYPE uXPriority_ )
```

SetPriority.

Set the priority of the [Thread](#) (running or otherwise) to a different level. This activity involves re-scheduling, and must be done so with due caution, as it may effect the determinism of the system.

This should *always* be called from within a critical section to prevent system issues.

Parameters

<i>uXPriority_</i>	New priority of the thread
--------------------	----------------------------

Definition at line 433 of file [thread.cpp](#).

19.24.2.29 SetPriorityBase()

```
void Mark3::Thread::SetPriorityBase (
    PORT_PRIO_TYPE uXPriority_ ) [private]
```

SetPriorityBase.

Parameters

<i>uX</i> ↔ <i>Priority_</i>	
---------------------------------	--

Definition at line 419 of file [thread.cpp](#).

19.24.2.30 SetQuantum()

```
void Mark3::Thread::SetQuantum (
    uint16_t u16Quantum_ ) [inline]
```

SetQuantum.

Set the thread's round-robin execution quantum.

Parameters

<i>u16</i> ↔ <i>Quantum_</i>	Thread 's execution quantum (in milliseconds)
---------------------------------	---

Definition at line 193 of file [thread.h](#).

19.24.2.31 SetState()

```
void Mark3::Thread::SetState (
    ThreadState eState_ ) [inline]
```

SetState Set the thread's state to a new value.

This is only to be used by code within the kernel, and is not intended for use by an end-user.

Parameters

<i>e</i> ↔ <i>State</i> ↔ _	New thread state to set.
-----------------------------------	--------------------------

Definition at line 423 of file [thread.h](#).

19.24.2.32 Sleep()

```
static void Mark3::Thread::Sleep (
    uint32_t u32TimeMs_ ) [static]
```

Sleep.

Put the thread to sleep for the specified time (in milliseconds). Actual time slept may be longer (but not less than) the interval specified.

Parameters

<i>u32Time</i> ↔ <i>Ms_</i>	Time to sleep (in ms)
--------------------------------	-----------------------

Examples:

[lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_idle_↔
function/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

19.24.2.33 Start()

```
void Mark3::Thread::Start (
    void )
```

Start.

Start the thread - remove it from the stopped list, add it to the scheduler's list of threads (at the thread's set priority), and continue along.

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line [147](#) of file [thread.cpp](#).

19.24.2.34 Stop()

```
void Mark3::Thread::Stop ( )
```

Stop.

Stop a thread that's actively scheduled without destroying its stacks. Stopped threads can be restarted using the [Start\(\)](#) API.

Definition at line [189](#) of file [thread.cpp](#).

19.24.2.35 USleep()

```
static void Mark3::Thread::USleep (
    uint32_t u32TimeUs_ ) [static]
```

USleep.

Put the thread to sleep for the specified time (in microseconds). Actual time slept may be longer (but not less than) the interval specified.

Parameters

<i>u32TimeUs_</i>	Time to sleep (in microseconds)
-------------------	---------------------------------

19.24.2.36 Yield()

```
void Mark3::Thread::Yield (
    void ) [static]
```

Yield.

Yield the thread - this forces the system to call the scheduler and determine what thread should run next. This is typically used when threads are moved in and out of the scheduler.

Definition at line 387 of file [thread.cpp](#).

The documentation for this class was generated from the following files:

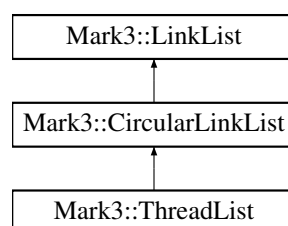
- [/home/moslevin/projects/github/m3-repo/kernel/src/public/thread.h](#)
- [/home/moslevin/projects/github/m3-repo/kernel/src/thread.cpp](#)

19.25 Mark3::ThreadList Class Reference

This class is used for building thread-management facilities, such as schedulers, and blocking objects.

```
#include <threadlist.h>
```

Inheritance diagram for Mark3::ThreadList:



Public Member Functions

- [ThreadList](#) ()
ThreadList.
- void [SetPriority](#) ([PORT_PRIO_TYPE](#) uXPriority_)
SetPriority.
- void [SetMapPointer](#) ([PriorityMap](#) *pclMap_)
SetMapPointer.
- void [Add](#) ([LinkListNode](#) *node_)
Add.
- void [Add](#) ([LinkListNode](#) *node_, [PriorityMap](#) *pclMap_, [PORT_PRIO_TYPE](#) uXPriority_)
Add.
- void [AddPriority](#) ([LinkListNode](#) *node_)
AddPriority.
- void [Remove](#) ([LinkListNode](#) *node_)
Remove.
- [Thread](#) * [HighestWaiter](#) ()
HighestWaiter.

Private Attributes

- [PORT_PRIO_TYPE](#) m_uXPriority
Priority of the threadlist.
- [PriorityMap](#) * m_pclMap
Pointer to the bitmap/flag to set when used for scheduling.

Additional Inherited Members

19.25.1 Detailed Description

This class is used for building thread-management facilities, such as schedulers, and blocking objects.

Definition at line 36 of file [threadlist.h](#).

19.25.2 Constructor & Destructor Documentation

19.25.2.1 ThreadList()

```
Mark3::ThreadList::ThreadList ( ) [inline]
```

[ThreadList](#).

Default constructor - zero-initializes the data.

Definition at line 45 of file [threadlist.h](#).

19.25.3 Member Function Documentation

19.25.3.1 Add() [1/2]

```
void Mark3::ThreadList::Add (
    LinkListNode * node_ )
```

Add.

Add a thread to the threadlist.

Parameters

<i>node_</i> ↔	Pointer to the thread (link list node) to add to the list
—	

Definition at line 53 of file [threadlist.cpp](#).

19.25.3.2 Add() [2/2]

```
void Mark3::ThreadList::Add (
    LinkListNode * node_,
    PriorityMap * pclMap_,
    PORT_PRIO_TYPE uXPriority_ )
```

Add.

Add a thread to the threadlist, specifying the flag and priority at the same time.

Parameters

<i>node_</i>	Pointer to the thread to add (link list node)
<i>pclMap_</i>	Pointer to the bitmap flag to set (if used in a scheduler context), or NULL for non-scheduler.
<i>uXPriority_</i> ↔	Priority of the threadlist

Definition at line 102 of file [threadlist.cpp](#).

19.25.3.3 AddPriority()

```
void Mark3::ThreadList::AddPriority (
    LinkListNode * node_ )
```

AddPriority.

Add a thread to the list such that threads are ordered from highest to lowest priority from the head of the list.

Parameters

<i>node</i> ↔	Pointer to a thread to add to the list.
—	

Definition at line 66 of file [threadlist.cpp](#).

19.25.3.4 HighestWaiter()

```
Thread * Mark3::ThreadList::HighestWaiter ( )
```

HighestWaiter.

Return a pointer to the highest-priority thread in the thread-list.

Returns

Pointer to the highest-priority thread

Definition at line 125 of file [threadlist.cpp](#).

19.25.3.5 Remove()

```
void Mark3::ThreadList::Remove (
    LinkListNode * node_ )
```

Remove.

Remove the specified thread from the threadlist

Parameters

<i>node</i> ↔	Pointer to the thread to remove
—	

Definition at line 112 of file [threadlist.cpp](#).

19.25.3.6 SetMapPointer()

```
void Mark3::ThreadList::SetMapPointer (
    PriorityMap * pclMap_ )
```

SetMapPointer.

Set the pointer to a bitmap to use for this threadlist. Once again, only needed when the threadlist is being used for scheduling purposes.

Parameters

<i>pcl↔</i> <i>Map_</i>	Pointer to the priority map object used to track this thread.
----------------------------	---

Definition at line 47 of file [threadlist.cpp](#).

19.25.3.7 SetPriority()

```
void Mark3::ThreadList::SetPriority (
    PORT_PRIO_TYPE uXPriority_ )
```

SetPriority.

Set the priority of this threadlist (if used for a scheduler).

Parameters

<i>uX↔</i> <i>Priority_</i>	Priority level of the thread list
--------------------------------	-----------------------------------

Definition at line 41 of file [threadlist.cpp](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/github/m3-repo/kernel/src/public/threadlist.h](#)
- [/home/moslevin/projects/github/m3-repo/kernel/src/threadlist.cpp](#)

19.26 Mark3::ThreadPort Class Reference

Class defining the architecture specific functions required by the kernel.

```
#include <threadport.h>
```

Static Public Member Functions

- static void [StartThreads](#) ()
StartThreads.

Static Private Member Functions

- static void [InitStack](#) ([Thread](#) *pstThread_)
InitStack.

19.26.1 Detailed Description

Class defining the architecture specific functions required by the kernel.

This is limited (at this point) to a function to start the scheduler, and a function to initialize the default stack-frame for a thread.

Definition at line 187 of file [threadport.h](#).

19.26.2 Member Function Documentation

19.26.2.1 InitStack()

```
void Mark3::ThreadPort::InitStack (
    Thread * pstThread_ ) [static], [private]
```

InitStack.

Initialize the thread's stack.

Parameters

<i>pstThread_</i>	Pointer to the thread to initialize
-------------------	-------------------------------------

Definition at line 40 of file [threadport.cpp](#).

19.26.2.2 StartThreads()

```
void Mark3::ThreadPort::StartThreads ( ) [static]
```

StartThreads.

Function to start the scheduler, initial threads, etc.

Definition at line 134 of file [threadport.cpp](#).

The documentation for this class was generated from the following files:

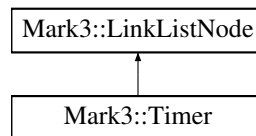
- [/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/threadport.h](#)
- [/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/threadport.cpp](#)

19.27 Mark3::Timer Class Reference

Kernel-managed software timers.

```
#include <timer.h>
```

Inheritance diagram for Mark3::Timer:



Public Member Functions

- [Timer](#) ()
Timer.
- void [Init](#) ()
Init.
- void [Start](#) (bool bRepeat_, uint32_t u32IntervalMs_, [TimerCallback](#) pfCallback_, void *pvData_)
Start.
- void [Start](#) (bool bRepeat_, uint32_t u32IntervalMs_, uint32_t u32ToleranceMs_, [TimerCallback](#) pfCallback_, void *pvData_)
Start.
- void [Start](#) ()
Start.
- void [Stop](#) ()
Stop.
- void [SetFlags](#) (uint8_t u8Flags_)
SetFlags.
- void [SetCallback](#) ([TimerCallback](#) pfCallback_)
SetCallback.
- void [SetData](#) (void *pvData_)
SetData.
- void [SetOwner](#) ([Thread](#) *pclOwner_)
SetOwner.
- void [SetIntervalTicks](#) (uint32_t u32Ticks_)
SetIntervalTicks.
- void [SetIntervalSeconds](#) (uint32_t u32Seconds_)
SetIntervalSeconds.
- uint32_t [GetInterval](#) ()
GetInterval.
- void [SetIntervalMSeconds](#) (uint32_t u32MSeconds_)
SetIntervalMSeconds.
- void [SetIntervalUSeconds](#) (uint32_t u32USeconds_)
SetIntervalUSeconds.
- void [SetTolerance](#) (uint32_t u32Ticks_)
SetTolerance.

Private Member Functions

- void [SetInitialized](#) ()
SetInitialized.
- bool [IsInitialized](#) (void)
IsInitialized.

Private Attributes

- uint8_t [m_u8Initialized](#)
Cookie used to determine whether or not the timer is initialized.
- uint8_t [m_u8Flags](#)
Flags for the timer, defining if the timer is one-shot or repeated.
- [TimerCallback](#) [m_pfCallback](#)
Pointer to the callback function.
- uint32_t [m_u32Interval](#)
Interval of the timer in timer ticks.
- uint32_t [m_u32TimeLeft](#)
Time remaining on the timer.
- uint32_t [m_u32TimerTolerance](#)
Maximum tolerance (used for timer harmonization)
- [Thread](#) * [m_pclOwner](#)
Pointer to the owner thread.
- void * [m_pvData](#)
Pointer to the callback data.

Additional Inherited Members

19.27.1 Detailed Description

Kernel-managed software timers.

Kernel-managed timers, used to provide high-precision high-resolution delays. Functionality is useful to both user-code, and is used extensively within the kernel and its blocking objects to implement round-robin scheduling, thread sleep, and timeouts. Relies on a single hardware timer, which is multiplexed through the kernel.

Examples:

[lab6_timers/main.cpp](#).

Definition at line 111 of file [timer.h](#).

19.27.2 Constructor & Destructor Documentation

19.27.2.1 Timer()

```
Mark3::Timer::Timer ( )
```

[Timer](#).

Default Constructor - Do nothing. Allow the init call to perform the necessary object initialization prior to use.

19.27.3 Member Function Documentation

19.27.3.1 GetInterval()

```
uint32_t Mark3::Timer::GetInterval ( ) [inline]
```

GetInterval.

Return the timer's configured interval in ticks

Returns

[Timer](#) interval in ticks.

Definition at line [239](#) of file [timer.h](#).

19.27.3.2 Init()

```
void Mark3::Timer::Init ( )
```

Init.

Re-initialize the [Timer](#) to default values.

19.27.3.3 IsInitialized()

```
bool Mark3::Timer::IsInitialized (
    void ) [inline], [private]
```

IsInitialized.

Returns

Definition at line [281](#) of file [timer.h](#).

19.27.3.4 SetCallback()

```
void Mark3::Timer::SetCallback (
    TimerCallback pfCallback_ ) [inline]
```

SetCallback.

Define the callback function to be executed on expiry of the timer

Parameters

<i>pf</i> ↔ <i>Callback</i> ↔ —	Pointer to the callback function to call
---------------------------------------	--

Definition at line 196 of file [timer.h](#).

19.27.3.5 SetData()

```
void Mark3::Timer::SetData (
    void * pvData_ ) [inline]
```

SetData.

Define a pointer to be sent to the timer callbacak on timer expiry

Parameters

<i>pv</i> ↔ <i>Data_</i>	Pointer to data to pass as argument into the callback
-----------------------------	---

Definition at line 204 of file [timer.h](#).

19.27.3.6 SetFlags()

```
void Mark3::Timer::SetFlags (
    uint8_t u8Flags_ ) [inline]
```

SetFlags.

Set the timer's flags based on the bits in the *u8Flags_* argument

Parameters

<i>u8</i> ↔ <i>Flags_</i>	Flags to assign to the timer object. <code>TIMERLIST_FLAG_ONE_SHOT</code> for a one-shot timer, 0 for a continuous timer.
------------------------------	---

Definition at line 188 of file [timer.h](#).

19.27.3.7 SetIntervalMSeconds()

```
void Mark3::Timer::SetIntervalMSeconds (
    uint32_t u32MSeconds_ )
```

SetIntervalMSeconds.

Set the timer expiry interval in milliseconds (platform agnostic)

Parameters

<i>u32M↔ Seconds_</i>	Time in milliseconds
---------------------------	----------------------

19.27.3.8 SetIntervalSeconds()

```
void Mark3::Timer::SetIntervalSeconds (
    uint32_t u32Seconds_ )
```

SetIntervalSeconds.

Set the timer expiry interval in seconds (platform agnostic)

Parameters

<i>u32↔ Seconds_</i>	Time in seconds
--------------------------	-----------------

19.27.3.9 SetIntervalTicks()

```
void Mark3::Timer::SetIntervalTicks (
    uint32_t u32Ticks_ )
```

SetIntervalTicks.

Set the timer expiry in system-ticks (platform specific!)

Parameters

<i>u32↔ Ticks_</i>	Time in ticks
------------------------	---------------

19.27.3.10 SetIntervalUSEconds()

```
void Mark3::Timer::SetIntervalUSEconds (
    uint32_t u32USEconds_ )
```

SetIntervalUSEconds.

Set the timer expiry interval in microseconds (platform agnostic)

Parameters

<i>u32U↔ Seconds_</i>	Time in microseconds
---------------------------	----------------------

19.27.3.11 SetOwner()

```
void Mark3::Timer::SetOwner (
    Thread * pclOwner_ ) [inline]
```

SetOwner.

Set the owner-thread of this timer object (all timers must be owned by a thread).

Parameters

<i>pcl↔ Owner_</i>	Owner thread of this timer object
------------------------	-----------------------------------

Definition at line 213 of file [timer.h](#).

19.27.3.12 SetTolerance()

```
void Mark3::Timer::SetTolerance (
    uint32_t u32Ticks_ )
```

SetTolerance.

Set the timer's maximum tolerance in order to synchronize timer processing with other timers in the system.

Parameters

<i>u32↔ Ticks_</i>	Maximum tolerance in ticks
------------------------	----------------------------

19.27.3.13 Start() [1/3]

```
void Mark3::Timer::Start (
    bool bRepeat_,
    uint32_t u32IntervalMs_,
    TimerCallback pfCallback_,
    void * pvData_ )
```

Start.

Start a timer using default ownership, using repeats as an option, and millisecond resolution.

Parameters

<i>bRepeat_</i>	0 - timer is one-shot. 1 - timer is repeating.
<i>u32IntervalMs_</i>	- Interval of the timer in milliseconds
<i>pfCallback_</i>	- Function to call on timer expiry
<i>pvData_</i>	- Data to pass into the callback function

Examples:

[lab6_timers/main.cpp](#).

19.27.3.14 Start() [2/3]

```
void Mark3::Timer::Start (
    bool bRepeat_,
    uint32_t u32IntervalMs_,
    uint32_t u32ToleranceMs_,
    TimerCallback pfCallback_,
    void * pvData_ )
```

Start.

Start a timer using default ownership, using repeats as an option, and millisecond resolution.

Parameters

<i>bRepeat_</i>	0 - timer is one-shot. 1 - timer is repeating.
<i>u32IntervalMs_</i>	- Interval of the timer in milliseconds
<i>u32ToleranceMs_</i>	- Allow the timer expiry to be delayed by an additional maximum time, in order to have as many timers expire at the same time as possible.
<i>pfCallback_</i>	- Function to call on timer expiry
<i>pvData_</i>	- Data to pass into the callback function

19.27.3.15 Start() [3/3]

```
void Mark3::Timer::Start ( )
```

Start.

Start or restart a timer using parameters previously configured via calls to Start(<with args>), or via the a-la-carte parameter setter methods. This is especially useful for retriggering one-shot timers that have previously expired, using the timer's previous configuration.

19.27.3.16 Stop()

```
void Mark3::Timer::Stop ( )
```

Stop.

Stop a timer already in progress. Has no effect on timers that have already been stopped.

The documentation for this class was generated from the following file:

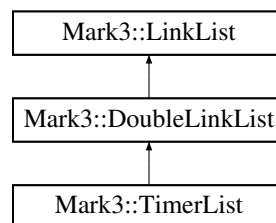
- </home/moslevin/projects/github/m3-repo/kernel/src/public/timer.h>

19.28 Mark3::TimerList Class Reference

[TimerList](#) class - a doubly-linked-list of timer objects.

```
#include <timerlist.h>
```

Inheritance diagram for Mark3::TimerList:



Public Member Functions

- void [Init](#) ()
Init.
- void [Add](#) ([Timer](#) *pclListNode_)
Add.
- void [Remove](#) ([Timer](#) *pclLinkedListNode_)
Remove.
- void [Process](#) ()
Process.

Private Attributes

- uint32_t [m_u32NextWakeup](#)
The time (in system clock ticks) of the next wakeup event.
- bool [m_bTimerActive](#)
Whether or not the timer is active.

Additional Inherited Members

19.28.1 Detailed Description

[TimerList](#) class - a doubly-linked-list of timer objects.

Definition at line 39 of file [timerlist.h](#).

19.28.2 Member Function Documentation

19.28.2.1 Add()

```
void Mark3::TimerList::Add (
    Timer * pclListNode_ )
```

Add.

Add a timer to the [TimerList](#).

Parameters

<i>pclList← Node_</i>	Pointer to the Timer to Add
---------------------------	---

19.28.2.2 Init()

```
void Mark3::TimerList::Init ( )
```

Init.

Initialize the [TimerList](#) object. Must be called before using the object.

19.28.2.3 Process()

```
void Mark3::TimerList::Process ( )
```

Process.

Process all timers in the timerlist as a result of the timer expiring. This will select a new timer epoch based on the next timer to expire.

19.28.2.4 Remove()

```
void Mark3::TimerList::Remove (
    Timer * pclLinkListNode_ )
```

Remove.

Remove a timer from the [TimerList](#), cancelling its expiry.

Parameters

<i>pcListNode_</i>	Pointer to the Timer to remove
--------------------	--

The documentation for this class was generated from the following file:

- </home/moslevin/projects/github/m3-repo/kernel/src/public/timerlist.h>

19.29 Mark3::TimerScheduler Class Reference

"Static" Class used to interface a global [TimerList](#) with the rest of the kernel.

```
#include <timerscheduler.h>
```

Static Public Member Functions

- static void [Init](#) ()
Init.
- static void [Add](#) ([Timer](#) *pcListNode_)
Add.
- static void [Remove](#) ([Timer](#) *pcListNode_)
Remove.
- static void [Process](#) ()
Process.

Static Private Attributes

- static [TimerList](#) m_clTimerList
TimerList object manipulated by the Timer Scheduler.

19.29.1 Detailed Description

"Static" Class used to interface a global [TimerList](#) with the rest of the kernel.

Definition at line [38](#) of file [timerscheduler.h](#).

19.29.2 Member Function Documentation

19.29.2.1 Add()

```
static void Mark3::TimerScheduler::Add (
    Timer * pcListNode_ ) [inline], [static]
```

Add.

Add a timer to the timer scheduler. Adding a timer implicitly starts the timer as well.

Parameters

<i>pclListNode_</i>	Pointer to the timer list node to add
---------------------	---------------------------------------

Definition at line 56 of file [timerscheduler.h](#).

19.29.2.2 Init()

```
static void Mark3::TimerScheduler::Init (  
    void ) [inline], [static]
```

Init.

Initialize the timer scheduler. Must be called before any timer, or timer-derived functions are used.

Definition at line 47 of file [timerscheduler.h](#).

19.29.2.3 Process()

```
static void Mark3::TimerScheduler::Process ( ) [inline], [static]
```

Process.

This function must be called on timer expiry (from the timer's ISR context). This will result in all timers being updated based on the epoch that just elapsed. The next timer epoch is set based on the next [Timer](#) object to expire.

Definition at line 74 of file [timerscheduler.h](#).

19.29.2.4 Remove()

```
static void Mark3::TimerScheduler::Remove (  
    Timer * pclListNode_ ) [inline], [static]
```

Remove.

Remove a timer from the timer scheduler. May implicitly stop the timer if this is the only active timer scheduled.

Parameters

<i>pclListNode_</i>	Pointer to the timer list node to remove
---------------------	--

Definition at line 65 of file [timerscheduler.h](#).

The documentation for this class was generated from the following file:

- </home/moslevin/projects/github/m3-repo/kernel/src/public/timerscheduler.h>

Chapter 20

File Documentation

20.1 /home/moslevin/projects/github/m3-repo/kernel/libs/mark3c/src/public/fake_types.h File Reference

C-struct definitions that mirror.

```
#include <stdint.h>
#include <stddef.h>
#include <stdbool.h>
#include "mark3cfg.h"
```

20.1.1 Detailed Description

C-struct definitions that mirror.

This header contains a set of "fake" structures that have the same memory layout as the kernel objects in C++ (taking into account inheritance, etc.). These are used for sizing the opaque data blobs that are declared in C, which then become instantiated as C++ kernel objects via the bindings provided.

Definition in file [fake_types.h](#).

20.2 fake_types.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00026 #include <stdint.h>
00027 #include <stddef.h>
00028 #include <stdbool.h>
00029 #include "mark3cfg.h"
```

```

00030
00031 #pragma once
00032
00033 #if defined(__cplusplus)
00034 extern "C" {
00035 #endif
00036
00037 //-----
00038 typedef struct {
00039     void* prev;
00040     void* next;
00041 } Fake_LinkedListNode;
00042
00043 //-----
00044 typedef struct {
00045     void* head;
00046     void* tail;
00047 } Fake_LinkedList;
00048
00049 //-----
00050 typedef struct {
00051     Fake_LinkedList fake_list;
00052     PORT_PRIO_TYPE m_uXPriority;
00053     void* m_pclMap;
00054 } Fake_ThreadList;
00055
00056 //-----
00057 typedef struct {
00058     Fake_LinkedListNode m_ll_node;
00059 #if KERNEL_EXTRA_CHECKS
00060     uint8_t m_u8Initialized;
00061 #endif
00062     uint8_t m_u8Flags;
00063     void* m_pfCallback;
00064     uint32_t m_u32Interval;
00065     uint32_t m_u32TimeLeft;
00066     uint32_t m_u32TimerTolerance;
00067     void* m_pclOwner;
00068     void* m_pvData;
00069 } Fake_Timer;
00070
00071 //-----
00072 typedef struct {
00073     Fake_LinkedListNode m_ll_node;
00074     K_WORD* m_pwStackTop;
00075     K_WORD* m_pwStack;
00076     uint8_t m_u8ThreadID;
00077     PORT_PRIO_TYPE m_uXPriority;
00078     PORT_PRIO_TYPE m_uXCurPriority;
00079     uint8_t m_eState;
00080 #if KERNEL_USE_EXTENDED_CONTEXT
00081     void* m_pvExtendedContext;
00082 #endif
00083 #if KERNEL_USE_THREADNAME
00084     const char* m_szName;
00085 #endif
00086     uint16_t m_u16StackSize;
00087     void* m_pclCurrent;
00088     void* m_pclOwner;
00089     void* m_pfEntryPoint;
00090     void* m_pvArg;
00091 #if KERNEL_USE_QUANTUM
00092     uint16_t m_u16Quantum;
00093 #endif
00094 #if KERNEL_USE_EVENTFLAG
00095     uint16_t m_u16FlagMask;
00096     uint8_t m_eFlagMode;
00097 #endif
00098 #if KERNEL_USE_TIMEOUTS || KERNEL_USE_SLEEP
00099     Fake_Timer m_clTimer;
00100 #endif
00101 #if KERNEL_USE_TIMEOUTS
00102     bool m_bExpired;
00103 #endif
00104 } Fake_Thread;
00105
00106 //-----
00107 typedef struct {
00108     Fake_ThreadList thread_list;
00109 #if KERNEL_EXTRA_CHECKS
00110     uint8_t m_u8Initialized;
00111 #endif
00112     uint16_t m_u16Value;
00113     uint16_t m_u16MaxValue;
00114 } Fake_Semaphore;
00115
00116 //-----

```

```

00117 typedef struct {
00118     Fake_ThreadList thread_list;
00119 #if KERNEL_EXTRA_CHECKS
00120     uint8_t          m_u8Initialized;
00121 #endif
00122     uint8_t          m_u8Recurse;
00123     bool             m_bReady;
00124     uint8_t          m_u8MaxPri;
00125     void*            m_pclOwner;
00126 } Fake_Mutex;
00127
00128 //-----
00129 typedef struct {
00130     Fake_LinkedListNode list_node;
00131     void*                m_pvData;
00132     uint16_t             m_ul6Code;
00133 } Fake_Message;
00134
00135 //-----
00136 typedef struct {
00137     Fake_Semaphore m_clSemaphore;
00138     Fake_LinkedList m_clLinkList;
00139 } Fake_MessageQueue;
00140
00141 //-----
00142 typedef struct {
00143     Fake_LinkedList m_clList;
00144 } Fake_MessagePool;
00145
00146 //-----
00147 typedef struct {
00148     uint16_t          m_ul6Head;
00149     uint16_t          m_ul6Tail;
00150     uint16_t          m_ul6Count;
00151     uint16_t          m_ul6Free;
00152     uint16_t          m_ul6ElementSize;
00153     void*             m_pvBuffer;
00154     Fake_Semaphore m_clRecvSem;
00155 #if KERNEL_USE_TIMEOUTS
00156     Fake_Semaphore m_clSendSem;
00157 #endif
00158 } Fake_Mailbox;
00159
00160 //-----
00161 typedef struct {
00162     Fake_ThreadList thread_list;
00163 #if KERNEL_EXTRA_CHECKS
00164     uint8_t          m_u8Initialized;
00165 #endif
00166     bool             m_bPending;
00167 } Fake_Notify;
00168
00169 //-----
00170 typedef struct {
00171     Fake_ThreadList thread_list;
00172 #if KERNEL_EXTRA_CHECKS
00173     uint8_t          m_u8Initialized;
00174 #endif
00175     uint16_t          m_ul6EventFlag;
00176 } Fake_EventFlag;
00177
00178 //-----
00179 typedef struct {
00180     Fake_Mutex m_clGlobalMutex;
00181     Fake_Mutex m_clReaderMutex;
00182     uint8_t m_u8ReadCount;
00183 } Fake_ConditionVariable;
00184
00185 //-----
00186 typedef struct {
00187     Fake_Mutex m_clMutex;
00188     Fake_Semaphore m_clSemaphore;
00189     uint8_t m_u8Waiters;
00190 } Fake_ReaderWriterLock;
00191
00192 #if defined(__cplusplus)
00193 }
00194 #endif

```

20.3 /home/moslevin/projects/github/m3-repo/kernel/libs/mark3c/src/public/mark3c.h File Reference

Implementation of C-language wrappers for the [Mark3](#) kernel.

```
#include "mark3cfg.h"
#include "fake_types.h"
#include <stdint.h>
#include <stdbool.h>
```

Typedefs

- typedef void * [EventFlag_t](#)
EventFlag opaque handle data type.
- typedef void * [Mailbox_t](#)
Mailbox opaque handle data type.
- typedef void * [Message_t](#)
Message opaque handle data type.
- typedef void * [MessagePool_t](#)
MessagePool opaque handle data type.
- typedef void * [MessageQueue_t](#)
MessageQueue opaque handle data type.
- typedef void * [Mutex_t](#)
Mutex opaque handle data type.
- typedef void * [Notify_t](#)
Notification object opaque handle data type.
- typedef void * [Semaphore_t](#)
Semaphore opaque handle data type.
- typedef void * [Thread_t](#)
Thread opaque handle data type.
- typedef void * [Timer_t](#)
Timer opaque handle data type.
- typedef void * [ConditionVariable_t](#)
Condition Variable opaque handle data type.
- typedef void * [ReaderWriterLock_t](#)
Reader-writer-lock opaque handle data type.

Enumerations

- enum [event_flag_operation_t](#) {
[EVENT_FLAG_ALL_SET](#), [EVENT_FLAG_ANY_SET](#), [EVENT_FLAG_ALL_CLEAR](#), [EVENT_FIAG_ANY_↵](#)
[CLEAR](#),
[EVENT_FLAG_PENDING_UNBLOCK](#) }

Functions

- void [Kernel_Init](#) (void)
Kernel_Init.
- void [Kernel_Start](#) (void)
Kernel_Start.
- bool [Kernel_IsStarted](#) (void)
Kernel_IsStarted.
- void [Kernel_SetPanic](#) (panic_func_t pfPanic_)
Kernel_SetPanic.
- bool [Kernel_IsPanic](#) (void)
Kernel_IsPanic.
- void [Kernel_Panic](#) (uint16_t u16Cause_)
Kernel_Panic.
- void [Scheduler_Enable](#) (bool bEnable_)
Scheduler_Enable.
- bool [Scheduler_IsEnabled](#) (void)
Scheduler_IsEnabled.
- [Thread_t](#) [Scheduler_GetCurrentThread](#) (void)
Scheduler_GetCurrentThread.
- void [Thread_Init](#) ([Thread_t](#) handle, [K_WORD](#) *pwStack_, uint16_t u16StackSize_, [PORT_PRIO_TYPE](#) uXPriority_, thread_entry_func_t pfEntryPoint_, void *pvArg_)
Thread_Init.
- void [Thread_Start](#) ([Thread_t](#) handle)
Thread_Start.
- void [Thread_Stop](#) ([Thread_t](#) handle)
Thread_Stop.
- [PORT_PRIO_TYPE](#) [Thread_GetPriority](#) ([Thread_t](#) handle)
Thread_GetPriority.
- [PORT_PRIO_TYPE](#) [Thread_GetCurPriority](#) ([Thread_t](#) handle)
Thread_GetCurPriority.
- void [Thread_SetPriority](#) ([Thread_t](#) handle, [PORT_PRIO_TYPE](#) uXPriority_)
Thread_SetPriority.
- void [Thread_Yield](#) (void)
Thread_Yield.
- void [Thread_SetID](#) ([Thread_t](#) handle, uint8_t u8ID_)
Thread_SetID.
- uint8_t [Thread_GetID](#) ([Thread_t](#) handle)
Thread_GetID.
- uint16_t [Thread_GetStackSlack](#) ([Thread_t](#) handle)
Thread_GetStackSlack.
- thread_state_t [Thread_GetState](#) ([Thread_t](#) handle)
Thread_GetState.

20.3.1 Detailed Description

Implementation of C-language wrappers for the [Mark3](#) kernel.

Header providing C-language API bindings for the [Mark3](#) kernel.

Definition in file [mark3c.h](#).

20.3.2 Enumeration Type Documentation

20.3.2.1 event_flag_operation_t

enum `event_flag_operation_t`

Enumerator

<code>EVENT_FLAG_ALL_SET</code>	Block until all bits in the specified bitmask are set.
<code>EVENT_FLAG_ANY_SET</code>	Block until any bits in the specified bitmask are set.
<code>EVENT_FLAG_ALL_CLEAR</code>	Block until all bits in the specified bitmask are cleared.
<code>EVENT_FIAG_ANY_CLEAR</code>	Block until any bits in the specified bitmask are cleared.
<code>EVENT_FLAG_PENDING_UNBLOCK</code>	Special code. Not used by user

Definition at line 72 of file [mark3c.h](#).

20.3.3 Function Documentation

20.3.3.1 Kernel_Init()

```
void Kernel_Init (
    void )
```

`Kernel_Init`.

See also

`void Kernel::Init()`

Definition at line 204 of file [mark3c.cpp](#).

20.3.3.2 Kernel_IsPanic()

```
bool Kernel_IsPanic (
    void )
```

`Kernel_IsPanic`.

See also

`bool Kernel::IsPanic()`

Returns

Whether or not the kernel is in a panic state

Definition at line 228 of file [mark3c.cpp](#).

20.3.3.3 Kernel_IsStarted()

```
bool Kernel_IsStarted (
    void )
```

Kernel_IsStarted.

See also

bool Kernel::IsStarted()

Returns

Whether or not the kernel has started - true = running, false = not started

Definition at line 216 of file [mark3c.cpp](#).

20.3.3.4 Kernel_Panic()

```
void Kernel_Panic (
    uint16_t u16Cause_ )
```

Kernel_Panic.

See also

void Kernel::Panic(uint16_t u16Cause_)

Parameters

<i>u16Cause_</i>	Reason for the kernel panic
------------------	-----------------------------

Definition at line 234 of file [mark3c.cpp](#).

20.3.3.5 Kernel_SetPanic()

```
void Kernel_SetPanic (
    panic_func_t pfPanic_ )
```

Kernel_SetPanic.

See also

void Kernel::SetPanic(PanicFunc_t pfPanic_)

Parameters

<i>pf</i> ↔	Panic function pointer
<i>Panic</i> ↔	
—	

Definition at line 222 of file [mark3c.cpp](#).

20.3.3.6 Kernel_Start()

```
void Kernel_Start (
    void )
```

Kernel_Start.

See also

`void Kernel::Start()`

Definition at line 210 of file [mark3c.cpp](#).

20.3.3.7 Scheduler_Enable()

```
void Scheduler_Enable (
    bool bEnable_ )
```

Scheduler_Enable.

See also

`void Scheduler::SetScheduler(bool bEnable_)`

Parameters

<i>bEnable_true</i>	to enable, false to disable the scheduler
---------------------	---

Definition at line 303 of file [mark3c.cpp](#).

20.3.3.8 Scheduler_GetCurrentThread()

```
Thread_t Scheduler_GetCurrentThread (
    void )
```

Scheduler_GetCurrentThread.

See also

Thread* Scheduler::GetCurrentThread()

Returns

Handle of the currently-running thread

Definition at line 315 of file [mark3c.cpp](#).

20.3.3.9 Scheduler_IsEnabled()

```
bool Scheduler_IsEnabled (
    void )
```

Scheduler_IsEnabled.

See also

bool Scheduler::IsEnabled()

Returns

true - scheduler enabled, false - disabled

Definition at line 309 of file [mark3c.cpp](#).

20.3.3.10 Thread_GetCurPriority()

```
PORT_PRIO_TYPE Thread_GetCurPriority (
    Thread_t handle )
```

Thread_GetCurPriority.

See also

[PORT_PRIO_TYPE](#) Thread::GetCurPriority()

Parameters

<i>handle</i>	Handle of the thread
---------------	----------------------

Returns

Current priority of the thread considering priority inheritance

Definition at line 370 of file [mark3c.cpp](#).

20.3.3.11 Thread_GetID()

```
uint8_t Thread_GetID (
    Thread_t handle )
```

Thread_GetID.

See also

`uint8_t Thread::GetID()`

Parameters

<i>handle</i>	Handle of the thread
---------------	----------------------

Returns

Return ID assigned to the thread

Definition at line 449 of file [mark3c.cpp](#).

20.3.3.12 Thread_GetPriority()

```
PORT_PRIO_TYPE Thread_GetPriority (
    Thread_t handle )
```

Thread_GetPriority.

See also

`PORT_PRIO_TYPE Thread::GetPriority()`

Parameters

<i>handle</i>	Handle of the thread
---------------	----------------------

Returns

Current priority of the thread not considering priority inheritance

Definition at line 364 of file [mark3c.cpp](#).

20.3.3.13 Thread_GetStackSlack()

```
uint16_t Thread_GetStackSlack (
    Thread_t handle )
```

Thread_GetStackSlack.

See also

uint16_t Thread::GetStackSlack()

Parameters

<i>handle</i>	Handle of the thread
---------------	----------------------

Returns

Return the amount of unused stack on the given thread

Definition at line 455 of file [mark3c.cpp](#).

20.3.3.14 Thread_GetState()

```
thread_state_t Thread_GetState (
    Thread_t handle )
```

Thread_GetState.

See also

ThreadState Thread::GetState()

Parameters

<i>handle</i>	Handle of the thread
---------------	----------------------

Returns

The thread's current execution state

Definition at line 461 of file [mark3c.cpp](#).

20.3.3.15 Thread_Init()

```
void Thread_Init (
    Thread_t handle,
    K_WORD * pwStack_,
    uint16_t u16StackSize_,
    PORT_PRIO_TYPE uXPriority_,
    thread_entry_func_t pfEntryPoint_,
    void * pvArg_ )
```

Thread_Init.

See also

```
void Thread::Init(K_WORD *pwStack_, uint16_t u16StackSize_, PORT_PRIO_TYPE uXPriority_, Thread_↵
Entry_t pfEntryPoint_, void *pvArg_)
```

Parameters

<i>handle</i>	Handle of the thread to initialize
<i>pwStack_</i>	Pointer to the stack to use for the thread
<i>u16Stack_↵ Size_</i>	Size of the stack (in bytes)
<i>uXPriority_</i>	Priority of the thread (0 = idle, 7 = max)
<i>pfEntryPoint_↵ _</i>	This is the function that gets called when the thread is started
<i>pvArg_</i>	Pointer to the argument passed into the thread's entryptoint function.

Definition at line 324 of file [mark3c.cpp](#).

20.3.3.16 Thread_SetID()

```
void Thread_SetID (
    Thread_t handle,
    uint8_t u8ID_ )
```

Thread_SetID.

See also

```
void Thread::SetID(uint8_t u8ID_)
```

Parameters

<i>handle</i>	Handle of the thread
<i>u8ID_↵ D_</i>	ID To assign to the thread

Definition at line 443 of file [mark3c.cpp](#).

20.3.3.17 Thread_SetPriority()

```
void Thread_SetPriority (
    Thread_t handle,
    PORT_PRIO_TYPE uXPriority_ )
```

Thread_SetPriority.

See also

void Thread::SetPriority(PORT_PRIO_TYPE uXPriority_)

Parameters

<i>handle</i>	Handle of the thread
<i>uX_↔</i> <i>Priority_</i>	New priority level

Definition at line 391 of file [mark3c.cpp](#).

20.3.3.18 Thread_Start()

```
void Thread_Start (
    Thread_t handle )
```

Thread_Start.

See also

void Thread::Start()

Parameters

<i>handle</i>	Handle of the thread to start
---------------	-------------------------------

Definition at line 336 of file [mark3c.cpp](#).

20.3.3.19 Thread_Stop()

```
void Thread_Stop (
    Thread_t handle )
```

Thread_Stop.

See also

`void Thread::Stop()`

Parameters

<i>handle</i>	Handle of the thread to stop
---------------	------------------------------

Definition at line 343 of file [mark3c.cpp](#).

20.3.3.20 Thread_Yield()

```
void Thread_Yield (
    void )
```

`Thread_Yield`.

See also

`void Thread::Yield()`

Definition at line 438 of file [mark3c.cpp](#).

20.4 mark3c.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00021 #pragma once
00022
00023 #include "mark3cfg.h"
00024 #include "fake_types.h"
00025
00026 #include <stdint.h>
00027 #include <stdbool.h>
00028
00029 #if defined(__cplusplus)
00030 extern "C" {
00031 #endif
00032
00033 //-----
00034 // Define a series of handle types to be used in place of the underlying classes
00035 // of Mark3.
00036 typedef void* EventFlag_t;
00037 typedef void* Mailbox_t;
00038 typedef void* Message_t;
00039 typedef void* MessagePool_t;
00040 typedef void* MessageQueue_t;
00041 typedef void* Mutex_t;
00042 typedef void* Notify_t;
00043 typedef void* Semaphore_t;
00044 typedef void* Thread_t;
00045 typedef void* Timer_t;
```

```

00046 typedef void* ConditionVariable_t;
00047 typedef void* ReaderWriterLock_t;
00048
00049 //-----
00050 // Function pointer types used by Kernel APIs
00051 typedef void (*thread_create_callout_t)(Thread_t hThread_);
00052 typedef void (*thread_exit_callout_t)(Thread_t hThread_);
00053 typedef void (*thread_context_callout_t)(Thread_t hThread_);
00054
00055 //-----
00056 // Use the sizes of the structs in fake_types.h to generate opaque object-blobs
00057 // that get instantiated as kernel objects (from the C++ code) later.
00058 #define THREAD_SIZE (sizeof(Fake_Thread))
00059 #define TIMER_SIZE (sizeof(Fake_Timer))
00060 #define SEMAPHORE_SIZE (sizeof(Fake_Semaphore))
00061 #define MUTEX_SIZE (sizeof(Fake_Mutex))
00062 #define MESSAGE_SIZE (sizeof(Fake_Message))
00063 #define MESSAGEQUEUE_SIZE (sizeof(Fake_MessageQueue))
00064 #define MAILBOX_SIZE (sizeof(Fake_Mailbox))
00065 #define NOTIFY_SIZE (sizeof(Fake_Notify))
00066 #define EVENTFLAG_SIZE (sizeof(Fake_EventFlag))
00067 #define MESSAGEPOOL_SIZE (sizeof(Fake_MessagePool))
00068 #define CONDITIONVARIABLE_SIZE (sizeof(Fake_ConditionVariable))
00069 #define READERWRITERLOCK_SIZE (sizeof(Fake_ReaderWriterLock))
00070
00071 //-----
00072 typedef enum {
00073     EVENT_FLAG_ALL_SET,
00074     EVENT_FLAG_ANY_SET,
00075     EVENT_FLAG_ALL_CLEAR,
00076     EVENT_FLAG_ANY_CLEAR,
00077     EVENT_FLAG_PENDING_UNBLOCK
00078 } event_flag_operation_t;
00079
00080 //-----
00081 typedef enum {
00082     THREAD_STATE_EXIT = 0,
00083     THREAD_STATE_READY,
00084     THREAD_STATE_BLOCKED,
00085     THREAD_STATE_STOP,
00086     THREAD_STATE_INVALID
00087 } thread_state_t;
00088 //-----
00089 // Macros for declaring opaque buffers of an appropriate size for the given
00090 // kernel objects
00091 #define TOKEN_1(x, y) x##y
00092 #define TOKEN_2(x, y) TOKEN_1(x, y)
00093
00094 // Ensure that opaque buffers are sized to the nearest word - which is
00095 // a platform-dependent value.
00096 #define WORD_ROUND(x) ((x) + (sizeof(K_WORD) - 1)) / sizeof(K_WORD)
00097
00098 #define DECLARE_THREAD(name)
00099     K_WORD    TOKEN_2(__thread_, name)[WORD_ROUND(THREAD_SIZE)];
00100     Thread_t  name = (Thread_t)TOKEN_2(__thread_, name);
00101
00102 #define DECLARE_TIMER(name)
00103     K_WORD    TOKEN_2(__timer_, name)[WORD_ROUND(TIMER_SIZE)];
00104     Timer_t   name = (Timer_t)TOKEN_2(__timer_, name);
00105
00106 #define DECLARE_SEMAPHORE(name)
00107     K_WORD    TOKEN_2(__semaphore_, name)[WORD_ROUND(SEMAPHORE_SIZE)];
00108     Semaphore_t name = (Semaphore_t)TOKEN_2(__semaphore_, name);
00109
00110 #define DECLARE_MUTEX(name)
00111     K_WORD    TOKEN_2(__mutex_, name)[WORD_ROUND(MUTEX_SIZE)];
00112     Mutex_t   name = (Mutex_t)TOKEN_2(__mutex_, name);
00113
00114 #define DECLARE_MESSAGE(name)
00115     K_WORD    TOKEN_2(__message_, name)[WORD_ROUND(MESSAGE_SIZE)];
00116     Message_t name = (Message_t)TOKEN_2(__message_, name);
00117
00118 #define DECLARE_MESSAGEPOOL(name)
00119     K_WORD    TOKEN_2(__messagepool_, name)[WORD_ROUND(MESSAGEPOOL_SIZE)];
00120     MessagePool_t name = (MessagePool_t)TOKEN_2(__messagepool_, name);

```

```

00121
00122 #define DECLARE_MESSAGEQUEUE(name)
00123     K_WORD      \      TOKEN_2(__messagequeue_, name) [WORD_ROUND(MESSAGEQUEUE_SIZE)];
00124     MessageQueue_t name = (MessageQueue_t)TOKEN_2(__messagequeue_, name);
00125
00126 #define DECLARE_MAILBOX(name)
00127     K_WORD      \      TOKEN_2(__mailbox_, name) [WORD_ROUND(MAILBOX_SIZE)];
00128     Mailbox_t name = (Mailbox_t)TOKEN_2(__mailbox_, name);
00129
00130 #define DECLARE_NOTIFY(name)
00131     K_WORD      \      TOKEN_2(__notify_, name) [WORD_ROUND(NOTIFY_SIZE)];
00132     Notify_t name = (Notify_t)TOKEN_2(__notify_, name);
00133
00134 #define DECLARE_EVENTFLAG(name)
00135     K_WORD      \      TOKEN_2(__eventflag_, name) [WORD_ROUND(EVENTFLAG_SIZE)];
00136     EventFlag_t name = (EventFlag_t)TOKEN_2(__eventflag_, name);
00137
00138 #define DECLARE_CONDITIONVARIABLE(name)
00139     K_WORD      \      TOKEN_2(__condvar_, name) [WORD_ROUND(EVENTFLAG_SIZE)];
00140     ConditionVariable_t name = (ConditionVariable_t)TOKEN_2(__condvar_, name);
00141
00142 #define DECLARE_READERWRITERLOCK(name)
00143     K_WORD      \      TOKEN_2(__readerwriterlock_, name) [WORD_ROUND(EVENTFLAG_SIZE)];
00144     ReaderWriterLock_t name = (ReaderWriterLock_t)TOKEN_2(__readerwriterlock_, name);
00145
00146 //-----
00147 // Allocate-once Memory managment APIs
00148 #if defined KERNEL_USE_AUTO_ALLOC
00149
00155 void* Alloc_Memory(size_t eSize_);
00156
00161 void Free_Memory(void* pvObject_);
00162
00163 #if KERNEL_USE_SEMAPHORE
00164
00169 Semaphore_t Alloc_Semaphore(void);
00170 void Free_Semaphore(Semaphore_t handle);
00171
00172 #endif
00173 #if KERNEL_USE_MUTEX
00174
00179 Mutex_t Alloc_Mutex(void);
00180 void Free_Mutex(Mutex_t handle);
00181
00182 #endif
00183 #if KERNEL_USE_EVENTFLAG
00184
00189 EventFlag_t Alloc_EventFlag(void);
00190 void Free_EventFlag(EventFlag_t handle);
00191
00192 #endif
00193 #if KERNEL_USE_MESSAGE
00194
00199 Message_t Alloc_Message(void);
00200 void Free_Message(Message_t handle);
00201
00207 MessageQueue_t Alloc_MessageQueue(void);
00208 void Free_MessageQueue(MessageQueue_t handle);
00209
00210 MessagePool_t Alloc_MessagePool(void);
00211 void Free_MessagePool(MessagePool_t handle);
00212
00213 #endif
00214 #endif
00215 #if KERNEL_USE_NOTIFY
00216
00221 Notify_t Alloc_Notify(void);
00222 void Free_Notify(Notify_t handle);
00223
00224 #endif
00225 #if KERNEL_USE_MAILBOX
00226
00231 Mailbox_t Alloc_Mailbox(void);
00232 void Free_Mailbox(Mailbox_t handle);
00233 #endif

```

```

00234
00239 Thread_t Alloc_Thread(void);
00240 void Free_Thread(Thread_t handle);
00241
00242 #if KERNEL_USE_TIMERS
00243
00248 Timer_t Alloc_Timer(void);
00249 void Free_Timer(Timer_t handle);
00250
00251 #endif
00252 #endif
00253
00254 //-----
00255 // Kernel APIs
00260 void Kernel_Init(void);
00265 void Kernel_Start(void);
00272 bool Kernel_IsStarted(void);
00273
00274 typedef void (*panic_func_t)(uint16_t ul6PanicCode_);
00280 void Kernel_SetPanic(panic_func_t pfPanic_);
00286 bool Kernel_IsPanic(void);
00292 void Kernel_Panic(uint16_t ul6Cause_);
00293 #if KERNEL_USE_IDLE_FUNC
00294 typedef void (*idle_func_t)(void);
00295
00301 void Kernel_SetIdleFunc(idle_func_t pfIdle_);
00302 #endif
00303
00304 #if KERNEL_USE_THREAD_CALLOUTS
00305
00310 void Kernel_SetThreadCreateCallout(thread_create_callout_t pfCreate_);
00316 void Kernel_SetThreadExitCallout(thread_exit_callout_t pfExit_);
00317
00323 void Kernel_SetThreadContextSwitchCallout(thread_context_callout_t pfContext_);
00324
00330 thread_create_callout_t Kernel_GetThreadCreateCallout(void);
00331
00337 thread_exit_callout_t Kernel_GetThreadExitCallout(void);
00338
00344 thread_context_callout_t Kernel_GetThreadContextSwitchCallout(void);
00345 #endif
00346
00347 #if KERNEL_USE_STACK_GUARD
00348
00354 void Kernel_SetStackGuardThreshold(uint16_t ul6Threshold_);
00355
00361 uint16_t Kernel_GetStackGuardThreshold(void);
00362 #endif
00363 //-----
00364 // Scheduler APIs
00370 void Scheduler_Enable(bool bEnable_);
00376 bool Scheduler_IsEnabled(void);
00382 Thread_t Scheduler_GetCurrentThread(void);
00383
00384 typedef void (*thread_entry_func_t)(void* pvArg_);
00385 //-----
00386 // Thread APIs
00400 void Thread_Init(Thread_t handle,
00401                  K_WORD* pwStack_,
00402                  uint16_t ul6StackSize_,
00403                  PORT_PRIO_TYPE uXPriority_,
00404                  thread_entry_func_t pfEntryPoint_,
00405                  void* pvArg_);
00411 void Thread_Start(Thread_t handle);
00417 void Thread_Stop(Thread_t handle);
00418 #if KERNEL_USE_THREADNAME
00419
00425 void Thread_SetName(Thread_t handle, const char* szName_);
00432 const char* Thread_GetName(Thread_t handle);
00433 #endif
00434
00440 PORT_PRIO_TYPE Thread_GetPriority(Thread_t handle);
00447 PORT_PRIO_TYPE Thread_GetCurPriority(Thread_t handle);
00448 #if KERNEL_USE_QUANTUM
00449
00455 void Thread_SetQuantum(Thread_t handle, uint16_t ul6Quantum_);
00462 uint16_t Thread_GetQuantum(Thread_t handle);
00463 #endif
00464
00470 void Thread_SetPriority(Thread_t handle, PORT_PRIO_TYPE uXPriority_);
00471 #if KERNEL_USE_DYNAMIC_THREADS
00472
00477 void Thread_Exit(Thread_t handle);
00478 #endif
00479 #if KERNEL_USE_SLEEP
00480
00485 void Thread_Sleep(uint32_t u32TimeMs_);

```

```

00491 void Thread_USleep(uint32_t u32TimeUs_);
00492 #endif
00493 #if KERNEL_USE_EXTENDED_CONTEXT
00494
00500 void* Thread_GetExtendedContext(Thread_t handle);
00501
00508 void Thread_SetExtendedContext(Thread_t handle, void* pvData_);
00509
00510 #endif
00511
00515 void Thread_Yield(void);
00522 void Thread_SetID(Thread_t handle, uint8_t u8ID_);
00529 uint8_t Thread_GetID(Thread_t handle);
00536 uint16_t Thread_GetStackSlack(Thread_t handle);
00543 thread_state_t Thread_GetState(Thread_t handle);
00544
00545 //-----
00546 // Timer APIs
00547 #if KERNEL_USE_TIMERS
00548 typedef void (*timer_callback_t)(Thread_t hOwner_, void* pvData_);
00554 void Timer_Init(Timer_t handle);
00566 void Timer_Start(Timer_t handle,
00567                  bool bRepeat_,
00568                  uint32_t u32IntervalMs_,
00569                  uint32_t u32ToleranceMs_,
00570                  timer_callback_t pfCallback_,
00571                  void* pvData_);
00572
00578 void Timer_Restart(Timer_t handle);
00579
00585 void Timer_Stop(Timer_t handle);
00586 #endif
00587
00588 //-----
00589 // Semaphore APIs
00590 #if KERNEL_USE_SEMAPHORE
00591
00598 void Semaphore_Init(Semaphore_t handle, uint16_t ul6InitVal_, uint16_t ul6MaxVal_);
00604 void Semaphore_Post(Semaphore_t handle);
00610 void Semaphore_Pend(Semaphore_t handle);
00611 #if KERNEL_USE_TIMEOUTS
00612
00619 bool Semaphore_TimedPend(Semaphore_t handle, uint32_t u32WaitTimeMS_);
00620 #endif
00621 #endif
00622
00623 //-----
00624 // Mutex APIs
00625 #if KERNEL_USE_MUTEX
00626
00631 void Mutex_Init(Mutex_t handle);
00637 void Mutex_Claim(Mutex_t handle);
00643 void Mutex_Release(Mutex_t handle);
00644 #if KERNEL_USE_TIMEOUTS
00645
00652 bool Mutex_TimedClaim(Mutex_t handle, uint32_t u32WaitTimeMS_);
00653 #endif
00654 #endif
00655
00656 //-----
00657 // EventFlag APIs
00658 #if KERNEL_USE_EVENTFLAG
00659
00664 void EventFlag_Init(EventFlag_t handle);
00673 uint16_t EventFlag_Wait(EventFlag_t handle, uint16_t ul6Mask_, event_flag_operation_t eMode_);
00674 #if KERNEL_USE_TIMEOUTS
00675
00684 uint16_t EventFlag_TimedWait(EventFlag_t handle, uint16_t ul6Mask_, event_flag_operation_t eMode_, uint32_t
    u32TimeMS_);
00685 #endif
00686
00692 void EventFlag_Set(EventFlag_t handle, uint16_t ul6Mask_);
00699 void EventFlag_Clear(EventFlag_t handle, uint16_t ul6Mask_);
00706 uint16_t EventFlag_GetMask(EventFlag_t handle);
00707 #endif
00708
00709 //-----
00710 // Notification APIs
00711 #if KERNEL_USE_NOTIFY
00712
00717 void Notify_Init(Notify_t handle);
00723 void Notify_Signal(Notify_t handle);
00730 void Notify_Wait(Notify_t handle, bool* pbFlag_);
00731 #if KERNEL_USE_TIMEOUTS
00732
00740 bool Notify_TimedWait(Notify_t handle, uint32_t u32WaitTimeMS_, bool* pbFlag_);
00741 #endif

```

```

00742 #endif
00743
00744 //-----
00745 // Atomic Functions
00746 #if KERNEL_USE_ATOMIC
00747
00754 uint8_t Atomic_Set8(uint8_t* pu8Source_, uint8_t u8Val_);
00762 uint16_t Atomic_Set16(uint16_t* pu16Source_, uint16_t u16Val_);
00770 uint32_t Atomic_Set32(uint32_t* pu32Source_, uint32_t u32Val_);
00778 uint8_t Atomic_Add8(uint8_t* pu8Source_, uint8_t u8Val_);
00786 uint16_t Atomic_Add16(uint16_t* pu16Source_, uint16_t u16Val_);
00794 uint32_t Atomic_Add32(uint32_t* pu32Source_, uint32_t u32Val_);
00802 uint8_t Atomic_Sub8(uint8_t* pu8Source_, uint8_t u8Val_);
00810 uint16_t Atomic_Sub16(uint16_t* pu16Source_, uint16_t u16Val_);
00818 uint32_t Atomic_Sub32(uint32_t* pu32Source_, uint32_t u32Val_);
00827 bool Atomic_TestAndSet(bool* pbLock);
00828 #endif
00829
00830 //-----
00831 // Message/Message Queue APIs
00832 #if KERNEL_USE_MESSAGE
00833
00838 void Message_Init(Message_t handle);
00845 void Message_SetData(Message_t handle, void* pvData_);
00852 void* Message_GetData(Message_t handle);
00859 void Message_SetCode(Message_t handle, uint16_t u16Code_);
00866 uint16_t Message_GetCode(Message_t handle);
00872 void MessageQueue_Init(MessageQueue_t handle);
00879 Message_t MessageQueue_Receive(MessageQueue_t handle);
00880 #if KERNEL_USE_TIMEOUTS
00881
00891 Message_t MessageQueue_TimedReceive(MessageQueue_t handle, uint32_t u32TimeWaitMS_);
00892 #endif
00893
00900 void MessageQueue_Send(MessageQueue_t handle, Message_t hMessage_);
00901
00907 uint16_t MessageQueue_GetCount(MessageQueue_t handle);
00908
00914 void MessagePool_Init(MessagePool_t handle);
00915
00922 void MessagePool_Push(MessagePool_t handle, Message_t msg);
00923
00930 Message_t MessagePool_Pop(MessagePool_t handle);
00931
00932 #endif
00933
00934 //-----
00935 // Mailbox APIs
00936 #if KERNEL_USE_MAILBOX
00937
00946 void Mailbox_Init(Mailbox_t handle, void* pvBuffer_, uint16_t u16BufferSize_, uint16_t u16ElementSize_);
00947
00955 bool Mailbox_Send(Mailbox_t handle, void* pvData_);
00956
00964 bool Mailbox_SendTail(Mailbox_t handle, void* pvData_);
00965
00974 bool Mailbox_TimedSend(Mailbox_t handle, void* pvData_, uint32_t u32TimeoutMS_);
00975
00984 bool Mailbox_TimedSendTail(Mailbox_t handle, void* pvData_, uint32_t u32TimeoutMS_);
00985
00993 void Mailbox_Receive(Mailbox_t handle, void* pvData_);
00994
01002 void Mailbox_ReceiveTail(Mailbox_t handle, void* pvData_);
01003 #if KERNEL_USE_TIMEOUTS
01004
01014 bool Mailbox_TimedReceive(Mailbox_t handle, void* pvData_, uint32_t u32TimeoutMS_);
01015
01025 bool Mailbox_TimedReceiveTail(Mailbox_t handle, void* pvData_, uint32_t u32TimeoutMS_);
01026
01033 uint16_t Mailbox_GetFreeSlots(Mailbox_t handle);
01034
01041 bool Mailbox_IsFull(Mailbox_t handle);
01042
01049 bool Mailbox_IsEmpty(Mailbox_t handle);
01050 #endif
01051 #endif
01052
01053 //-----
01054 // Condition Variables
01055 #if KERNEL_USE_CONDVAR
01056
01061 void ConditionVariable_Init(ConditionVariable_t handle);
01062
01069 void ConditionVariable_Wait(ConditionVariable_t handle, Mutex_t hMutex_);
01070
01076 void ConditionVariable_Signal(ConditionVariable_t handle);
01077

```

```

01083 void ConditionVariable_Broadcast(ConditionVariable_t handle);
01084 #if KERNEL_USE_TIMEOUTS
01085
01093 bool ConditionVariable_TimedWait(ConditionVariable_t handle, Mutex_t hMutex_, uint32_t u32WaitTimeMS_);
01094 #endif
01095 #endif
01096
01097 //-----
01098 // Reader-writer locks
01099 #if KERNEL_USE_READERWRITER
01100
01105 void ReaderWriterLock_Init(ReaderWriterLock_t handle);
01106
01112 void ReaderWriterLock_AcquireReader(ReaderWriterLock_t handle);
01113
01119 void ReaderWriterLock_ReleaseReader(ReaderWriterLock_t handle);
01120
01126 void ReaderWriterLock_AcquireWriter(ReaderWriterLock_t handle);
01127
01133 void ReaderWriterLock_ReleaseWriter(ReaderWriterLock_t handle);
01134 #if KERNEL_USE_TIMEOUTS
01135
01142 bool ReaderWriterLock_TimedAcquireWriter(ReaderWriterLock_t handle, uint32_t u32TimeoutMs_);
01143
01151 bool ReaderWriterLock_TimedAcquireReader(ReaderWriterLock_t handle, uint32_t u32TimeoutMs_);
01152 #endif
01153 #endif
01154
01155 //-----
01156 // Kernel-Aware Simulation APIs
01157 #if KERNEL_AWARE_SIMULATION
01158
01164 void KernelAware_ProfileInit(const char* szStr_);
01165
01170 void KernelAware_ProfileStart(void);
01171
01176 void KernelAware_ProfileStop(void);
01177
01182 void KernelAware_ProfileReport(void);
01183
01189 void KernelAware_ExitSimulator(void);
01190
01196 void KernelAware_Print(const char* szStr_);
01197
01204 void KernelAware_Trace(uint16_t u16File_, uint16_t u16Line_);
01205
01213 void KernelAware_Trace1(uint16_t u16File_, uint16_t u16Line_, uint16_t u16Arg1_);
01222 void KernelAware_Trace2(uint16_t u16File_, uint16_t u16Line_, uint16_t u16Arg1_, uint16_t u16Arg2_);
01232 bool KernelAware_IsSimulatorAware(void);
01233 #endif
01234
01235 #if defined(__cplusplus)
01236 }
01237 #endif

```

20.5 /home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/kernelprofile.cpp

File Reference

ATMega328p Profiling timer implementation.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "profile.h"
#include "kernelprofile.h"
#include "threadport.h"
#include <avr/io.h>
#include <avr/interrupt.h>

```

20.5.1 Detailed Description

ATMega328p Profiling timer implementation.

Definition in file [kernelprofile.cpp](#).

20.6 kernelprofile.cpp

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00020 #include "kerneltypes.h"
00021 #include "mark3cfg.h"
00022 #include "profile.h"
00023 #include "kernelprofile.h"
00024 #include "threadport.h"
00025 #include <avr/io.h>
00026 #include <avr/interrupt.h>
00027
00028 #if KERNEL_USE_PROFILER
00029 namespace Mark3
00030 {
00031     uint32_t Profiler::m_u32Epoch;
00032
00033     //-----
00034     void Profiler::Init()
00035     {
00036         TCCR0A = 0;
00037         TCCR0B = 0;
00038         TIFR0 = 0;
00039         TIMSK0 = 0;
00040         m_u32Epoch = 0;
00041     }
00042
00043     //-----
00044     void Profiler::Start()
00045     {
00046         TIFR0 = 0;
00047         TCNT0 = 0;
00048         TCCR0B |= (1 << CS01);
00049         TIMSK0 |= (1 << TOIE0);
00050     }
00051
00052     //-----
00053     void Profiler::Stop()
00054     {
00055         TIFR0 = 0;
00056         TCCR0B &= ~(1 << CS01);
00057         TIMSK0 &= ~(1 << TOIE0);
00058     }
00059     //-----
00060     uint16_t Profiler::Read()
00061     {
00062         uint16_t u16Ret;
00063         CS_ENTER();
00064         TCCR0B &= ~(1 << CS01);
00065         u16Ret = TCNT0;
00066         TCCR0B |= (1 << CS01);
00067         CS_EXIT();
00068         return u16Ret;
00069     }
00070
00071     //-----
00072     void Profiler::Process()
00073     {
00074         CS_ENTER();
00075         m_u32Epoch++;
00076         CS_EXIT();
00077     }
00078
00079     //-----
00080     ISR(TIMER0_OVF_vect)
00081     {
00082         Profiler::Process();
00083     }
00084 } //namespace Mark3
00085 #endif

```

20.7 /home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/kernelswi.cpp File Reference

Kernel Software interrupt implementation for ATmega328p.

```
#include "kerneltypes.h"
#include "kernelswi.h"
#include <avr/io.h>
#include <avr/interrupt.h>
```

Namespaces

- [Mark3](#)

Class providing the software-interrupt required for context-switching in the kernel.

20.7.1 Detailed Description

Kernel Software interrupt implementation for ATmega328p.

Definition in file [kernelswi.cpp](#).

20.8 kernelswi.cpp

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00022 #include "kerneltypes.h"
00023 #include "kernelswi.h"
00024
00025 #include <avr/io.h>
00026 #include <avr/interrupt.h>
00027
00028
00029 namespace Mark3
00030 {
00031 //-----
00032 void KernelSWI::Config(void)
00033 {
00034     PORTB &= ~0x04;           // Clear INT2
00035     DDRB |= 0x04;           // Set PortB, bit 2 (INT2) As Output
00036     EICRA |= (1 << ISC20) | (1 << ISC21); // Rising edge on INT2
00037 }
00038
00039 //-----
00040 void KernelSWI::Start(void)
00041 {
00042     EIFR &= ~(1 << INTF2); // Clear any pending interrupts on INT2
00043     EIMSK |= (1 << INT2); // Enable INT2 interrupt (as int32_t as I-bit is set)
00044 }
00045
00046 //-----
00047 void KernelSWI::Stop(void)
00048 {
```

```

00049     EIMSK &= ~(1 << INT2); // Disable INT0 interrupts
00050 }
00051
00052 //-----
00053 uint8_t KernelSWI::DI()
00054 {
00055     bool bEnabled = ((EIMSK & (1 << INT2)) != 0);
00056     EIMSK &= ~(1 << INT2);
00057     return bEnabled;
00058 }
00059
00060 //-----
00061 void KernelSWI::RI(bool bEnable_)
00062 {
00063     if (bEnable_) {
00064         EIMSK |= (1 << INT2);
00065     } else {
00066         EIMSK &= ~(1 << INT2);
00067     }
00068 }
00069
00070 //-----
00071 void KernelSWI::Clear(void)
00072 {
00073     EIFR &= ~(1 << INTF2); // Clear the interrupt flag for INT0
00074 }
00075
00076 //-----
00077 void KernelSWI::Trigger(void)
00078 {
00079     // if(Thread_IsSchedulerEnabled())
00080     {
00081         PORTB &= ~0x04;
00082         PORTB |= 0x04;
00083     }
00084 }
00085 } //namespace Mark3

```

20.9 /home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/kerneltimer.cpp File Reference

Kernel Timer Implementation for ATmega328p.

```

#include "kerneltypes.h"
#include "kerneltimer.h"
#include "mark3cfg.h"
#include "ksemaphore.h"
#include "thread.h"
#include <avr/common.h>
#include <avr/io.h>
#include <avr/interrupt.h>

```

Namespaces

- [Mark3](#)

Class providing the software-interrupt required for context-switching in the kernel.

20.9.1 Detailed Description

Kernel Timer Implementation for ATmega328p.

Definition in file [kerneltimer.cpp](#).


```

00092     TIFR1 &= ~TIMER_IFR;
00093     TIMSK1 |= TIMER_IMSK;
00094 }
00095
00096 //-----
00097 void KernelTimer::Stop(void)
00098 {
00099     #if KERNEL_TIMERS_TICKLESS
00100         TIFR1 &= ~TIMER_IFR;
00101         TIMSK1 &= ~TIMER_IMSK;
00102         TCCR1B &= ~(1 << CS12); // Disable count...
00103         TCNT1 = 0;
00104         OCR1A = 0;
00105     #endif
00106 }
00107
00108 //-----
00109 PORT_TIMER_COUNT_TYPE KernelTimer::Read(void)
00110 {
00111     #if KERNEL_TIMERS_TICKLESS
00112         volatile uint16_t u16Read1;
00113         volatile uint16_t u16Read2;
00114
00115         do {
00116             u16Read1 = TCNT1;
00117             u16Read2 = TCNT1;
00118         } while (u16Read1 != u16Read2);
00119
00120         return u16Read1;
00121     #else
00122         return 0;
00123     #endif
00124 }
00125
00126 //-----
00127 PORT_TIMER_COUNT_TYPE KernelTimer::SubtractExpiry(
    PORT_TIMER_COUNT_TYPE uInterval)
00128 {
00129     #if KERNEL_TIMERS_TICKLESS
00130         OCR1A -= uInterval;
00131         return OCR1A;
00132     #else
00133         return 0;
00134     #endif
00135 }
00136
00137 //-----
00138 PORT_TIMER_COUNT_TYPE KernelTimer::TimeToExpiry(void)
00139 {
00140     #if KERNEL_TIMERS_TICKLESS
00141         uint16_t u16Read = KernelTimer::Read();
00142         uint16_t u16OCR1A = OCR1A;
00143
00144         if (u16Read >= u16OCR1A) {
00145             return 0;
00146         } else {
00147             return (u16OCR1A - u16Read);
00148         }
00149     #else
00150         return 0;
00151     #endif
00152 }
00153
00154 //-----
00155 PORT_TIMER_COUNT_TYPE KernelTimer::GetOvertime(void)
00156 {
00157     return KernelTimer::Read();
00158 }
00159
00160 //-----
00161 PORT_TIMER_COUNT_TYPE KernelTimer::SetExpiry(uint32_t
    u32Interval_)
00162 {
00163     #if KERNEL_TIMERS_TICKLESS
00164         uint16_t u16SetInterval;
00165         if (u32Interval_ > 65535) {
00166             u16SetInterval = 65535;
00167         } else {
00168             u16SetInterval = static_cast<uint16_t>(u32Interval_);
00169         }
00170
00171         OCR1A = u16SetInterval;
00172         return u16SetInterval;
00173     #else
00174         return 0;
00175     #endif
00176 }

```

```

00177
00178 //-----
00179 void KernelTimer::ClearExpiry(void)
00180 {
00181     #if KERNEL_TIMERS_TICKLESS
00182         OCR1A = 65535; // Clear the compare value
00183     #endif
00184 }
00185
00186 //-----
00187 uint8_t KernelTimer::DI(void)
00188 {
00189     #if KERNEL_TIMERS_TICKLESS
00190         bool bEnabled = ((TIMSK1 & (TIMER_IMSK)) != 0);
00191         TIFR1 &= ~TIMER_IFR; // Clear interrupt flags
00192         TIMSK1 &= ~TIMER_IMSK; // Disable interrupt
00193         return bEnabled;
00194     #else
00195         return 0;
00196     #endif
00197 }
00198
00199 //-----
00200 void KernelTimer::EI(void)
00201 {
00202     KernelTimer::RI(0);
00203 }
00204
00205 //-----
00206 void KernelTimer::RI(bool bEnable_)
00207 {
00208     #if KERNEL_TIMERS_TICKLESS
00209         if (bEnable_) {
00210             TIMSK1 |= (1 << OCIE1A); // Enable interrupt
00211         } else {
00212             TIMSK1 &= ~(1 << OCIE1A);
00213         }
00214     #endif
00215 }
00216 } //namespace Mark3
00217
00218 //-----
00223 //-----
00224 using namespace Mark3;
00225 ISR(TIMER1_COMPA_vect)
00226 {
00227     #if KERNEL_TIMERS_THREADED
00228         KernelTimer::ClearExpiry();
00229         s_clTimerSemaphore.Post();
00230     #else
00231         #if KERNEL_USE_TIMERS
00232             TimerScheduler::Process();
00233         #endif
00234         #if KERNEL_USE_QUANTUM
00235             Quantum::UpdateTimer();
00236         #endif
00237     #endif
00238 }

```

20.11 /home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/kernelprofile.h File Reference

Profiling timer hardware interface.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ll.h"

```

20.11.1 Detailed Description

Profiling timer hardware interface.

Definition in file [kernelprofile.h](#).

20.14 kernelswi.h

```

00001  /*=====
00002
00003  00004  00005  00006  00007  00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  ===== */
00021  #pragma once
00022  #include "kerneltypes.h"
00023
00024  //-----
00029  namespace Mark3
00030  {
00031  class KernelSWI
00032  {
00033  public:
00040      static void Config(void);
00041
00047      static void Start(void);
00048
00054      static void Stop(void);
00055
00061      static void Clear(void);
00062
00069      static void Trigger(void);
00070
00078      static uint8_t DI();
00079
00087      static void RI(bool bEnable_);
00088  };
00089  } // namespace Mark3

```

20.15 /home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/kerneltimer.h

File Reference

Kernel Timer Class declaration.

```

#include "kerneltypes.h"
#include "mark3cfg.h"

```

Classes

- class [Mark3::KernelTimer](#)
Hardware timer interface, used by all scheduling/timer subsystems.

Namespaces

- [Mark3](#)
Class providing the software-interrupt required for context-switching in the kernel.

20.15.1 Detailed Description

Kernel Timer Class declaration.

Definition in file [kerneltimer.h](#).

20.16 kerneletimer.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00021 #pragma once
00022
00023 #include "kerneltypes.h"
00024 #include "mark3cfg.h"
00025
00026 namespace Mark3
00027 {
00028 //-----
00032 class KernelTimer
00033 {
00034 public:
00040     static void Config(void);
00041
00047     static void Start(void);
00048
00054     static void Stop(void);
00055
00061     static uint8_t DI(void);
00062
00070     static void RI(bool bEnable_);
00071
00077     static void EI(void);
00078
00089     static PORT_TIMER_COUNT_TYPE SubtractExpiry(
PORT_TIMER_COUNT_TYPE uInterval_);
00090
00099     static PORT_TIMER_COUNT_TYPE TimeToExpiry(void);
00100
00109     static PORT_TIMER_COUNT_TYPE SetExpiry(uint32_t u32Interval_);
00110
00119     static PORT_TIMER_COUNT_TYPE GetOvertime(void);
00120
00126     static void ClearExpiry(void);
00127
00135     static PORT_TIMER_COUNT_TYPE Read(void);
00136 };
00137 } //namespace Mark3

```

20.17 /home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/portcfg.h File Reference

[Mark3](#) Port Configuration.

Macros

- `#define AVR (1)`
Define a macro indicating the CPU architecture for which this port belongs.
- `#define K_WORD uint8_t`
Define types that map to the CPU Architecture's default data-word and address size.
- `#define K_ADDR uint16_t`
Size of an address (pointer size)
- `#define PORT_PRIO_TYPE uint8_t`
Set a base datatype used to represent each element of the scheduler's priority bitmap.

- `#define PORT_PRIO_MAP_WORD_SIZE (1)`
size of PORT_PRIO_TYPE in bytes
- `#define PORT_SYSTEM_FREQ ((uint32_t)16000000)`
Define the running CPU frequency.
- `#define PORT_TIMER_FREQ ((uint32_t)(PORT_SYSTEM_FREQ / 1000))`
Set the timer frequency.
- `#define PORT_KERNEL_DEFAULT_STACK_SIZE ((K_ADDR)256)`
Define the default/minimum size of a thread stack.
- `#define PORT_KERNEL_TIMERS_THREAD_STACK ((K_ADDR)256)`
Define the size of the kernel-timer thread stack (if one is configured)
- `#define PORT_TIMER_COUNT_TYPE uint16_t`
Define the native type corresponding to the kernel timer hardware's counter register.
- `#define PORT_MIN_TIMER_TICKS (0)`
Minimum number of timer ticks for any delay or sleep, required to ensure that a timer cannot be initialized to a negative value.

20.17.1 Detailed Description

[Mark3](#) Port Configuration.

This file is used to configure the kernel for your specific target CPU in order to provide the optimal set of features for a given use case.

!! NOTE: This file must ONLY be included from [mark3cfg.h](#)

Definition in file [portcfg.h](#).

20.17.2 Macro Definition Documentation

20.17.2.1 AVR

```
#define AVR (1)
```

Define a macro indicating the CPU architecture for which this port belongs.

This may also be set by the toolchain, but that's not guaranteed.

Definition at line 32 of file [portcfg.h](#).

20.17.2.2 K_WORD

```
#define K_WORD uint8_t
```

Define types that map to the CPU Architecture's default data-word and address size.

Size of a data word

Examples:

[lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_idle_function/main.cpp](#), [lab3_round_robin/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab5_mutexes/main.cpp](#), [lab6_timers/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 39 of file [portcfg.h](#).

20.17.2.3 PORT_PRIO_TYPE

```
#define PORT_PRIO_TYPE uint8_t
```

Set a base datatype used to represent each element of the scheduler's priority bitmap.

PORT_PRIO_MAP_WORD_SIZE should map to the *size* of an element of type PORT_PRIO_TYPE. Type used for bitmap in the PriorityMap class

Definition at line 49 of file [portcfg.h](#).

20.17.2.4 PORT_SYSTEM_FREQ

```
#define PORT_SYSTEM_FREQ ((uint32_t)16000000)
```

Define the running CPU frequency.

This may be an integer constant, or an alias for another variable which holds the CPU's current running frequency.↔
CPU Frequency in Hz

Definition at line 57 of file [portcfg.h](#).

20.17.2.5 PORT_TIMER_COUNT_TYPE

```
#define PORT_TIMER_COUNT_TYPE uint16_t
```

Define the native type corresponding to the kernel timer hardware's counter register.

Timer counter type

Definition at line 86 of file [portcfg.h](#).

20.17.2.6 PORT_TIMER_FREQ

```
#define PORT_TIMER_FREQ ((uint32_t)(PORT_SYSTEM_FREQ / 1000))
```

Set the timer frequency.

If running in tickless mode, this is simply the frequency at which the free-running kernel timer increments.

In tick-based mode, this is the frequency at which the fixed-frequency kernel tick interrupt occurs. Fixed timer interrupt frequency

Definition at line 70 of file [portcfg.h](#).

20.18 portcfg.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00024 #pragma once
00025
00031 #ifndef AVR
00032 # define AVR (1)
00033 #endif
00034
00039 #define K_WORD uint8_t
00040 #define K_ADDR uint16_t
00041
00042
00049 #define PORT_PRIO_TYPE uint8_t
00050 #define PORT_PRIO_MAP_WORD_SIZE (1)
00051
00052
00056 #if !defined(PORT_SYSTEM_FREQ)
00057 #define PORT_SYSTEM_FREQ ((uint32_t)16000000)
00058 #endif
00059
00067 #if KERNEL_TIMERS_TICKLESS
00068 #define PORT_TIMER_FREQ ((uint32_t)(PORT_SYSTEM_FREQ / 256))
00069 #else
00070 #define PORT_TIMER_FREQ ((uint32_t)(PORT_SYSTEM_FREQ / 1000))
00071 #endif
00072
00076 #define PORT_KERNEL_DEFAULT_STACK_SIZE ((K_ADDR)256)
00077
00081 #define PORT_KERNEL_TIMERS_THREAD_STACK ((K_ADDR)256)
00082
00086 #define PORT_TIMER_COUNT_TYPE uint16_t
00087
00088
00092 #define PORT_MIN_TIMER_TICKS (0)
```

20.19 /home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/threadport.

File Reference

ATMega328p Multithreading support.

```
#include "kerneltypes.h"
#include "thread.h"
#include <avr/io.h>
#include <avr/interrupt.h>
```

Classes

- class [Mark3::ThreadPort](#)

Class defining the architecture specific functions required by the kernel.

Namespaces

- [Mark3](#)

Class providing the software-interrupt required for context-switching in the kernel.

Macros

- #define [ASM](#)(x) asm volatile(x);
ASM Macro - simplify the use of ASM directive in C.
- #define [SR_](#) 0x3F
Status register define - map to 0x003F.
- #define [SPH_](#) 0x3E
Stack pointer define.
- #define [TOP_OF_STACK](#)(x, y) (reinterpret_cast<[K_WORD](#)*>(reinterpret_cast<[K_ADDR](#)>(x) + (static_cast<[K_ADDR](#)>(y) - 1)))
Macro to find the top of a stack given its size and top address.
- #define [PUSH_TO_STACK](#)(x, y) *x = y; x--;
Push a value y to the stack pointer x and decrement the stack pointer.
- #define [Thread_SaveContext](#)()
Save the context of the Thread.
- #define [Thread_RestoreContext](#)()
Restore the context of the Thread.
- #define [CS_ENTER](#)()
These macros must be used in pairs !
- #define [CS_EXIT](#)()
Exit critical section (restore status register)
- #define [ENABLE_INTS](#)() [ASM](#)("sei");
Initiate a contex switch without using the SWI.

20.19.1 Detailed Description

ATMega328p Multithreading support.

Definition in file [threadport.h](#).

20.19.2 Macro Definition Documentation

20.19.2.1 CS_ENTER

```
#define CS_ENTER( )
```

Value:

```
{ \
uint8_t __x = _SFR_IO8(SR_); \
ASM("cli");
```

These macros *must* be used in pairs !

Enter critical section (copy status register, disable interrupts)

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 163 of file [threadport.h](#).

20.20 threadport.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00020 #pragma once
00021
00022 #include "kerneltypes.h"
00023 #include "thread.h"
00024
00025 #include <avr/io.h>
00026 #include <avr/interrupt.h>
00027
00028 namespace Mark3 {
00029
00030 // clang-format off
00031 //-----
00033 #define ASM(x)      asm volatile(x);
00034 #define SR_         0x3F
00036 #define SPH_         0x3E
00038 #define SPL_         0x3D
00039
00040 //-----
00042 #define TOP_OF_STACK(x, y)      (reinterpret_cast<K_WORD*>(reinterpret_cast<K_ADDR>(x) +
                                (static_cast<K_ADDR>(y) - 1)))
00043 #define PUSH_TO_STACK(x, y)     *x = y; x--;
00045 #define STACK_GROWS_DOWN      (1)
00046
00047 //-----
00048 // Lookup table based count-leading zeros implementation.
00049 inline uint8_t __mark3_clz8(uint8_t in_)
00050 {
00051     static const uint8_t u8Lookup[] = {4, 3, 2, 2, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0};
00052     uint8_t hi = __builtin_avr_swap(in_) & 0x0F;
00053     if (hi) {
00054         return u8Lookup[hi];
00055     }
00056     return 4 + u8Lookup[in_];
00057 }
00058
```

```
00059 //-----
00060 #define HW_CLZ      (1)
00061 #define CLZ(x)      __mark3_clz8(x)
00062
00063 //-----
00065 #define Thread_SaveContext() \
00066 ASM("push r0"); \
00067 ASM("in r0, __SREG__"); \
00068 ASM("cli"); \
00069 ASM("push r0"); \
00070 ASM("push r1"); \
00071 ASM("clr r1"); \
00072 ASM("push r2"); \
00073 ASM("push r3"); \
00074 ASM("push r4"); \
00075 ASM("push r5"); \
00076 ASM("push r6"); \
00077 ASM("push r7"); \
00078 ASM("push r8"); \
00079 ASM("push r9"); \
00080 ASM("push r10"); \
00081 ASM("push r11"); \
00082 ASM("push r12"); \
00083 ASM("push r13"); \
00084 ASM("push r14"); \
00085 ASM("push r15"); \
00086 ASM("push r16"); \
00087 ASM("push r17"); \
00088 ASM("push r18"); \
00089 ASM("push r19"); \
00090 ASM("push r20"); \
00091 ASM("push r21"); \
00092 ASM("push r22"); \
00093 ASM("push r23"); \
00094 ASM("push r24"); \
00095 ASM("push r25"); \
00096 ASM("push r26"); \
00097 ASM("push r27"); \
00098 ASM("push r28"); \
00099 ASM("push r29"); \
00100 ASM("push r30"); \
00101 ASM("push r31"); \
00102 ASM("in r0, 0x3B"); \
00103 ASM("push r0"); \
00104 ASM("lds r26, g_pclCurrent"); \
00105 ASM("lds r27, g_pclCurrent + 1"); \
00106 ASM("adiw r26, 4"); \
00107 ASM("in r0, 0x3D"); \
00108 ASM("st x+, r0"); \
00109 ASM("in r0, 0x3E"); \
00110 ASM("st x+, r0");
00111
00112 //-----
00114 #define Thread_RestoreContext() \
00115 ASM("lds r26, g_pclCurrent"); \
00116 ASM("lds r27, g_pclCurrent + 1"); \
00117 ASM("adiw r26, 4"); \
00118 ASM("ld r28, x+"); \
00119 ASM("out 0x3D, r28"); \
00120 ASM("ld r29, x+"); \
00121 ASM("out 0x3E, r29"); \
00122 ASM("pop r0"); \
00123 ASM("out 0x3B, r0"); \
00124 ASM("pop r31"); \
00125 ASM("pop r30"); \
00126 ASM("pop r29"); \
00127 ASM("pop r28"); \
00128 ASM("pop r27"); \
00129 ASM("pop r26"); \
00130 ASM("pop r25"); \
00131 ASM("pop r24"); \
00132 ASM("pop r23"); \
00133 ASM("pop r22"); \
00134 ASM("pop r21"); \
00135 ASM("pop r20"); \
00136 ASM("pop r19"); \
00137 ASM("pop r18"); \
00138 ASM("pop r17"); \
00139 ASM("pop r16"); \
00140 ASM("pop r15"); \
00141 ASM("pop r14"); \
00142 ASM("pop r13"); \
00143 ASM("pop r12"); \
00144 ASM("pop r11"); \
00145 ASM("pop r10"); \
00146 ASM("pop r9"); \
00147 ASM("pop r8"); \
```

```

00148 ASM("pop r7"); \
00149 ASM("pop r6"); \
00150 ASM("pop r5"); \
00151 ASM("pop r4"); \
00152 ASM("pop r3"); \
00153 ASM("pop r2"); \
00154 ASM("pop r1"); \
00155 ASM("pop r0"); \
00156 ASM("out __SREG__, r0"); \
00157 ASM("pop r0");
00158
00159 //-----
00161 //-----
00163 #define CS_ENTER() \
00164 { \
00165     uint8_t __x = _SFR_IO8(SR); \
00166     ASM("cli");
00167 //-----
00169 #define CS_EXIT() \
00170     _SFR_IO8(SR) = __x; \
00171 }
00172
00173 //-----
00175 #define ENABLE_INTS()      ASM("sei");
00176 #define DISABLE_INTS()    ASM("cli");
00177
00178 //-----
00179 class Thread;
00187 class ThreadPort
00188 {
00189 public:
00195     static void StartThreads();
00196     friend class Thread;
00197 private:
00198
00206     static void InitStack(Thread *pstThread_);
00207 };
00208
00209 } // namespace Mark3

```

20.21 /home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/threadport.cpp

File Reference

ATMega1284p Multithreading.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "thread.h"
#include "threadport.h"
#include "kernelprofile.h"
#include "kernelswi.h"
#include "kerneltimer.h"
#include "timerlist.h"
#include "quantum.h"
#include "kernel.h"
#include "kernelaware.h"
#include <avr/io.h>
#include <avr/interrupt.h>

```

Namespaces

- [Mark3](#)

Class providing the software-interrupt required for context-switching in the kernel.

Functions

- **Mark3::ISR** (INT2_vect) __attribute__((signal))
ISR(INT2_vect) SWI using INT2 - used to trigger a context switch.

20.21.1 Detailed Description

ATMega1284p Multithreading.

Definition in file [threadport.cpp](#).

20.22 threadport.cpp

```

00001  /*=====
00002
00003  _____
00004  |   \   /   |   |   |   |   |   |   |   |   |
00005  |   \   /   |   |   |   |   |   |   |   |   |
00006  |   \   /   |   |   |   |   |   |   |   |   |
00007  |   \   /   |   |   |   |   |   |   |   |   |
00008  |   \   /   |   |   |   |   |   |   |   |   |
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  =====*/
00022  #include "kerneltypes.h"
00023  #include "mark3cfg.h"
00024  #include "thread.h"
00025  #include "threadport.h"
00026  #include "kernelprofile.h"
00027  #include "kernelswi.h"
00028  #include "kerneltimer.h"
00029  #include "timerlist.h"
00030  #include "quantum.h"
00031  #include "kernel.h"
00032  #include "kernelaware.h"
00033  #include <avr/io.h>
00034  #include <avr/interrupt.h>
00035
00036  //-----
00037  namespace Mark3
00038  {
00039  //-----
00040  void ThreadPort::InitStack(Thread* pclThread_)
00041  {
00042      // Initialize the stack for a Thread
00043      uint16_t ul6Addr;
00044      uint8_t* pu8Stack;
00045      uint16_t i;
00046
00047      // Get the address of the thread's entry function
00048      ul6Addr = (uint16_t)(pclThread_>m_pfEntryPoint);
00049
00050      // Start by finding the bottom of the stack
00051      pu8Stack = (uint8_t*)pclThread_>m_pwStackTop;
00052
00053      // clear the stack, and initialize it to a known-default value (easier
00054      // to debug when things go sour with stack corruption or overflow)
00055      for (i = 0; i < pclThread_>m_ul6StackSize; i++) {
00056          pclThread_>m_pwStack[i] = 0xFF;
00057      }
00058
00059      // Our context starts with the entry function
00060      PUSH_TO_STACK(pu8Stack, (uint8_t)(ul6Addr & 0x00FF));
00061      PUSH_TO_STACK(pu8Stack, (uint8_t)((ul6Addr >> 8) & 0x00FF));
00062
00063      // R0
00064      PUSH_TO_STACK(pu8Stack, 0x00); // R0
00065
00066      // Push status register and R1 (which is used as a constant zero)
00067      PUSH_TO_STACK(pu8Stack, 0x80); // SR
00068      PUSH_TO_STACK(pu8Stack, 0x00); // R1
00069

```

```

00070 // Push other registers
00071 for (i = 2; i <= 23; i++) // R2-R23
00072 {
00073     PUSH_TO_STACK(pu8Stack, i);
00074 }
00075
00076 // Assume that the argument is the only stack variable
00077 PUSH_TO_STACK(pu8Stack, (uint8_t)((uint16_t)(pclThread->
m_pvArg)) & 0x00FF)); // R24
00078 PUSH_TO_STACK(pu8Stack, (uint8_t)((uint16_t)(pclThread->
m_pvArg)) >> 8) & 0x00FF)); // R25
00079
00080 // Push the rest of the registers in the context
00081 for (i = 26; i <= 31; i++) {
00082     PUSH_TO_STACK(pu8Stack, i);
00083 }
00084
00085 PUSH_TO_STACK(pu8Stack, 0x00); // RAMPZ
00086 // Set the top of the stack.
00087 pclThread->m_pwStackTop = (uint8_t*)pu8Stack;
00088
00089 // That's it! the thread is ready to run now.
00090 }
00091
00092 //-----
00093 static void Thread_Switch(void)
00094 {
00095     #if KERNEL_USE_IDLE_FUNC
00096         // If there's no next-thread-to-run...
00097         if (g_pclNext == Kernel::GetIdleThread()) {
00098             g_pclCurrent = Kernel::GetIdleThread();
00099
00100             // Disable the SWI, and re-enable interrupts -- enter nested interrupt
00101             // mode.
00102             KernelSWI::DI();
00103
00104             g_pclCurrent = Kernel::GetIdleThread();
00105
00106             uint8_t u8SR = _SFR_IO8(SR_);
00107
00108             // So long as there's no "next-to-run" thread, keep executing the Idle
00109             // function to conclusion...
00110
00111             while (g_pclNext == Kernel::GetIdleThread()) {
00112                 // Ensure that we run this block in an interrupt enabled context (but
00113                 // with the rest of the checks being performed in an interrupt disabled
00114                 // context).
00115                 ASM("sei");
00116                 Kernel::Idle();
00117                 ASM("cli");
00118             }
00119
00120             // Progress has been achieved -- an interrupt-triggered event has caused
00121             // the scheduler to run, and choose a new thread. Since we've already
00122             // saved the context of the thread we've hijacked to run idle, we can
00123             // proceed to disable the nested interrupt context and switch to the
00124             // new thread.
00125
00126             _SFR_IO8(SR_) = u8SR;
00127             KernelSWI::RI(true);
00128         }
00129     #endif
00130     g_pclCurrent = (Thread*)g_pclNext;
00131 }
00132
00133 //-----
00134 void ThreadPort::StartThreads()
00135 {
00136     KernelSWI::Config(); // configure the task switch SWI
00137     KernelTimer::Config(); // configure the kernel timer
00138     #if KERNEL_USE_PROFILER
00139         Profiler::Init();
00140     #endif
00141     Scheduler::SetScheduler(1); // enable the scheduler
00142     Scheduler::Schedule(); // run the scheduler - determine the first thread to run
00143
00144     Thread_Switch(); // Set the next scheduled thread to the current thread
00145
00146     KernelTimer::Start(); // enable the kernel timer
00147     KernelSWI::Start(); // enable the task switch SWI
00148
00149     #if KERNEL_USE_QUANTUM
00150         // Restart the thread quantum timer, as any value held prior to starting
00151         // the kernel will be invalid. This fixes a bug where multiple threads
00152         // started with the highest priority before starting the kernel causes problems
00153         // until the running thread voluntarily blocks.
00154         Quantum::RemoveThread();

```

```

00155     Quantum::AddThread(g_pclCurrent);
00156 #endif
00157
00158     // Restore the context...
00159     Thread_RestoreContext(); // restore the context of the first running thread
00160     ASM("reti");           // return from interrupt - will return to the first scheduled thread
00161 }
00162
00163 //-----
00164 //-----
00165 ISR(INT2_vect) __attribute__((signal, naked));
00170 ISR(INT2_vect)
00171 {
00172     Thread_SaveContext(); // Push the context (registers) of the current task
00173     Thread_Switch();      // Switch to the next task
00174     Thread_RestoreContext(); // Pop the context (registers) of the next task
00175     ASM("reti");         // Return to the next task
00176 }
00177 } //namespace Mark3
00178

```

20.23 /home/moslevin/projects/github/m3-repo/kernel/src/atomic.cpp File Reference

Basic Atomic Operations.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "atomic.h"
#include "threadport.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
```

20.23.1 Detailed Description

Basic Atomic Operations.

Definition in file [atomic.cpp](#).

20.24 atomic.cpp

```
00001 /*-----
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*
00021 #include "kerneltypes.h"
00022 #include "mark3cfg.h"
00023 #include "atomic.h"
00024 #include "threadport.h"
00025
00026 #define _CAN_HAS_DEBUG
00027 //--[Autogenerated - Do Not Modify]-----
00028 #include "dbg_file_list.h"
00029 #include "buffalogger.h"
00030 #if defined(DBG_FILE)
00031 #error "Debug logging file token already defined!  Bailing."
00032 #else
```

```

00033 #define DBG_FILE _DBG___KERNEL_ATOMIC_CPP
00034 #endif
00035 //--[End Autogenerated content]-----
00036
00037 #if KERNEL_USE_ATOMIC
00038 namespace Mark3
00039 {
00040     //--
00041     uint8_t Atomic::Set(uint8_t* pu8Source_, uint8_t u8Val_)
00042     {
00043         uint8_t u8Ret;
00044         CS_ENTER();
00045         u8Ret = *pu8Source_;
00046         *pu8Source_ = u8Val_;
00047         CS_EXIT();
00048         return u8Ret;
00049     }
00050     //--
00051     uint16_t Atomic::Set(uint16_t* pul6Source_, uint16_t u16Val_)
00052     {
00053         uint16_t u16Ret;
00054         CS_ENTER();
00055         u16Ret = *pul6Source_;
00056         *pul6Source_ = u16Val_;
00057         CS_EXIT();
00058         return u16Ret;
00059     }
00060     //--
00061     uint32_t Atomic::Set(uint32_t* pu32Source_, uint32_t u32Val_)
00062     {
00063         uint32_t u32Ret;
00064         CS_ENTER();
00065         u32Ret = *pu32Source_;
00066         *pu32Source_ = u32Val_;
00067         CS_EXIT();
00068         return u32Ret;
00069     }
00070
00071     //--
00072     uint8_t Atomic::Add(uint8_t* pu8Source_, uint8_t u8Val_)
00073     {
00074         uint8_t u8Ret;
00075         CS_ENTER();
00076         u8Ret = *pu8Source_;
00077         *pu8Source_ += u8Val_;
00078         CS_EXIT();
00079         return u8Ret;
00080     }
00081
00082     //--
00083     uint16_t Atomic::Add(uint16_t* pul6Source_, uint16_t u16Val_)
00084     {
00085         uint16_t u16Ret;
00086         CS_ENTER();
00087         u16Ret = *pul6Source_;
00088         *pul6Source_ += u16Val_;
00089         CS_EXIT();
00090         return u16Ret;
00091     }
00092
00093     //--
00094     uint32_t Atomic::Add(uint32_t* pu32Source_, uint32_t u32Val_)
00095     {
00096         uint32_t u32Ret;
00097         CS_ENTER();
00098         u32Ret = *pu32Source_;
00099         *pu32Source_ += u32Val_;
00100         CS_EXIT();
00101         return u32Ret;
00102     }
00103
00104     //--
00105     uint8_t Atomic::Sub(uint8_t* pu8Source_, uint8_t u8Val_)
00106     {
00107         uint8_t u8Ret;
00108         CS_ENTER();
00109         u8Ret = *pu8Source_;
00110         *pu8Source_ -= u8Val_;
00111         CS_EXIT();
00112         return u8Ret;
00113     }
00114
00115     //--
00116     uint16_t Atomic::Sub(uint16_t* pul6Source_, uint16_t u16Val_)
00117     {
00118         uint16_t u16Ret;
00119         CS_ENTER();

```

20.25 /home/moslevin/projects/github/m3-repo/kernel/src/autoalloc.cpp File Reference

```
#include "mark3cfg.h"
#include "mark3.h"
#include "autoalloc.h"
#include "threadport.h"
#include "kernel.h"
#include <stdint.h>
```

Definition in file [autoalloc.cpp](#).[illegible]

```

00021 #include "mark3.h"
00022 #include "autoalloc.h"
00023 #include "threadport.h"
00024 #include "kernel.h"
00025
00026 #include <stdint.h>
00027
00028 #if KERNEL_USE_AUTO_ALLOC
00029 using namespace Mark3;
00030 //-----
00031 // Override new() and delete() using functions provided to AutoAlloc
00032 //-----
00033 void* operator new(size_t n)
00034 {
00035     return AutoAlloc::NewRawData(n);
00036 }
00037
00038 //-----
00039 void* operator new[](size_t n)
00040 {
00041     return AutoAlloc::NewRawData(n);
00042 }
00043
00044 //-----
00045 void operator delete(void * p)
00046 {
00047     AutoAlloc::DestroyRawData(p);
00048 }
00049
00050 //-----
00051 void operator delete[](void * p)
00052 {
00053     AutoAlloc::DestroyRawData(p);
00054 }
00055
00056 //-----
00057 namespace Mark3
00058 {
00059     AutoAllocAllocator_t AutoAlloc::m_pfAllocator;
00060     AutoAllocFree_t AutoAlloc::m_pfFree;
00061
00062     //-----
00063     void AutoAlloc::Init()
00064     {
00065         m_pfAllocator = NULL;
00066         m_pfFree = NULL;
00067     }
00068
00069     //-----
00070     void* AutoAlloc::Allocate(AutoAllocType eType_, size_t sSize_)
00071     {
00072         if (!m_pfAllocator) {
00073             return NULL;
00074         }
00075         return m_pfAllocator(eType_, sSize_);
00076     }
00077
00078     //-----
00079     void AutoAlloc::Free(AutoAllocType eType_, void* pvObj_)
00080     {
00081         if (!m_pfFree) {
00082             return;
00083         }
00084         m_pfFree(eType_, pvObj_);
00085     }
00086
00087 #if KERNEL_USE_SEMAPHORE
00088 //-----
00089 Semaphore* AutoAlloc::NewSemaphore(void)
00090 {
00091     void* pvObj = Allocate(AutoAllocType::Semaphore, sizeof(Semaphore));
00092     if (pvObj) {
00093         return new (pvObj) Semaphore();
00094     }
00095     return 0;
00096 }
00097 //-----
00098 void AutoAlloc::DestroySemaphore(Semaphore *pclSemaphore_)
00099 {
00100     pclSemaphore_->~Semaphore();
00101     Free(AutoAllocType::Semaphore, pclSemaphore_);
00102 }
00103 #endif
00104
00105 #if KERNEL_USE_MUTEX
00106 //-----
00107 Mutex* AutoAlloc::NewMutex(void)

```

```

00108 {
00109     void* pvObj = Allocate(AutoAllocType::Mutex, sizeof(Mutex));
00110     if (pvObj) {
00111         return new (pvObj) Mutex();
00112     }
00113     return 0;
00114 }
00115 //-----
00116 void AutoAlloc::DestroyMutex(Mutex *pclMutex_)
00117 {
00118     pclMutex_->~Mutex();
00119     Free(AutoAllocType::Mutex, pclMutex_);
00120 }
00121 #endif
00122
00123 #if KERNEL_USE_EVENTFLAG
00124 //-----
00125 EventFlag* AutoAlloc::NewEventFlag(void)
00126 {
00127     void* pvObj = Allocate(AutoAllocType::EventFlag, sizeof(EventFlag));
00128     if (pvObj) {
00129         return new (pvObj) EventFlag();
00130     }
00131     return 0;
00132 }
00133 //-----
00134 void AutoAlloc::DestroyEventFlag(EventFlag *pclEventFlag_)
00135 {
00136     pclEventFlag_->~EventFlag();
00137     Free(AutoAllocType::EventFlag, pclEventFlag_);
00138 }
00139 #endif
00140
00141 #if KERNEL_USE_MESSAGE
00142 //-----
00143 Message* AutoAlloc::NewMessage(void)
00144 {
00145     void* pvObj = Allocate(AutoAllocType::Message, sizeof(Message));
00146     if (pvObj) {
00147         return new (pvObj) Message();
00148     }
00149     return 0;
00150 }
00151 //-----
00152 void AutoAlloc::DestroyMessage(Message *pclMessage_)
00153 {
00154     pclMessage_->~Message();
00155     Free(AutoAllocType::Message, pclMessage_);
00156 }
00157 //-----
00158 MessagePool* AutoAlloc::NewMessagePool(void)
00159 {
00160     void* pvObj = Allocate(AutoAllocType::MessagePool, sizeof(MessagePool));
00161     if (pvObj) {
00162         return new (pvObj) MessagePool();
00163     }
00164     return 0;
00165 }
00166 //-----
00167 void AutoAlloc::DestroyMessagePool(MessagePool *pclMessagePool_)
00168 {
00169     pclMessagePool_->~MessagePool();
00170     Free(AutoAllocType::MessagePool, pclMessagePool_);
00171 }
00172 //-----
00173 MessageQueue* AutoAlloc::NewMessageQueue(void)
00174 {
00175     void* pvObj = Allocate(AutoAllocType::MessageQueue, sizeof(MessageQueue));
00176     if (pvObj) {
00177         return new (pvObj) MessageQueue();
00178     }
00179     return 0;
00180 }
00181 //-----
00182 void AutoAlloc::DestroyMessageQueue(MessageQueue *pclMessageQ_)
00183 {
00184     pclMessageQ_->~MessageQueue();
00185     Free(AutoAllocType::MessageQueue, pclMessageQ_);
00186 }
00187 #endif
00188
00189 #if KERNEL_USE_NOTIFY
00190 //-----
00191 Notify* AutoAlloc::NewNotify(void)
00192 {
00193     void* pvObj = Allocate(AutoAllocType::Notify, sizeof(Notify));

```

```

00195     if (pvObj) {
00196         return new (pvObj) Notify();
00197     }
00198     return 0;
00199 }
00200 //-----
00201 void AutoAlloc::DestroyNotify(Notify *pclNotify_)
00202 {
00203     pclNotify_->~Notify();
00204     Free(AutoAllocType::Notify, pclNotify_);
00205 }
00206 #endif
00207
00208 #if KERNEL_USE_MAILBOX
00209 //-----
00210 Mailbox* AutoAlloc::NewMailbox(void)
00211 {
00212     void* pvObj = Allocate(AutoAllocType::MailBox, sizeof(Mailbox));
00213     if (pvObj) {
00214         return new (pvObj) Mailbox();
00215     }
00216     return 0;
00217 }
00218 //-----
00219 void AutoAlloc::DestroyMailbox(Mailbox *pclMailbox_)
00220 {
00221     pclMailbox_->~Mailbox();
00222     Free(AutoAllocType::MailBox, pclMailbox_);
00223 }
00224 #endif
00225
00226 #if KERNEL_USE_CONDVAR
00227 //-----
00228 ConditionVariable* AutoAlloc::NewConditionVariable()
00229 {
00230     void* pvObj = Allocate(AutoAllocType::ConditionVariable, sizeof(
00231 ConditionVariable));
00232     if (pvObj) {
00233         return new (pvObj) ConditionVariable();
00234     }
00235     return 0;
00236 }
00237 //-----
00238 void AutoAlloc::DestroyConditionVariable(ConditionVariable* pclCondvar_)
00239 {
00240     pclCondvar_->~ConditionVariable();
00241     Free(AutoAllocType::ConditionVariable, pclCondvar_);
00242 }
00243 #endif
00244
00245 #if KERNEL_USE_READERWRITER
00246 //-----
00247 ReaderWriterLock* AutoAlloc::NewReaderWriterLock()
00248 {
00249     void* pvObj = Allocate(AutoAllocType::ReaderWriterLock, sizeof(
00250 ReaderWriterLock));
00251     if (pvObj) {
00252         return new (pvObj) ReaderWriterLock();
00253     }
00254     return 0;
00255 }
00256 //-----
00257 void AutoAlloc::DestroyReaderWriterLock(ReaderWriterLock *pclReaderWriterLock_)
00258 {
00259     pclReaderWriterLock_->~ReaderWriterLock();
00260     Free(AutoAllocType::ReaderWriterLock, pclReaderWriterLock_);
00261 }
00262 #endif
00263 //-----
00264 Thread* AutoAlloc::NewThread(void)
00265 {
00266     void* pvObj = Allocate(AutoAllocType::Thread, sizeof(Thread));
00267     if (pvObj) {
00268         return new (pvObj) Thread();
00269     }
00270     return 0;
00271 }
00272 //-----
00273 void AutoAlloc::DestroyThread(Thread *pclThread_)
00274 {
00275     pclThread_->~Thread();
00276     Free(AutoAllocType::Thread, pclThread_);
00277 }
00278 #if KERNEL_USE_TIMERS
00279 //-----

```



```

00280 Timer* AutoAlloc::NewTimer(void)
00281 {
00282     void* pObj = Allocate(AutoAllocType::Timer, sizeof(Timer));
00283     if (pObj) {
00284         return new (pObj) Timer();
00285     }
00286     return 0;
00287 }
00288 //-----
00289 void AutoAlloc::DestroyTimer(Timer *pclTimer_)
00290 {
00291     pclTimer_->~Timer();
00292     Free(AutoAllocType::Timer, pclTimer_);
00293 }
00294 #endif
00295 //-----
00296 void* AutoAlloc::NewUserTypeAllocation(uint8_t eType_)
00297 {
00298     return Allocate(static_cast<AutoAllocType>(eType_), 0);
00299 }
00300 //-----
00301 void AutoAlloc::DestroyUserTypeAllocation(uint8_t eUserType_, void *pObj_)
00302 {
00303     Free(AutoAllocType::User, pObj_);
00304 }
00305 //-----
00306 void* AutoAlloc::NewRawData(size_t sSize_)
00307 {
00308     return Allocate(AutoAllocType::Raw, sSize_);
00309 }
00310 //-----
00311 void AutoAlloc::DestroyRawData(void *pvData_)
00312 {
00313     Free(AutoAllocType::Raw, pvData_);
00314 }
00315 } //namespace Mark3
00316 #endif

```

20.27 /home/moslevin/projects/github/m3-repo/kernel/src/blocking.cpp File Reference

Implementation of base class for blocking objects.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "blocking.h"
#include "thread.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"

```

Namespaces

- [Mark3](#)

Class providing the software-interrupt required for context-switching in the kernel.

20.27.1 Detailed Description

Implementation of base class for blocking objects.

Definition in file [blocking.cpp](#).

20.28 blocking.cpp

```

00001  /*=====
00002
00003  _____
00004  |   /   \   |   /   \   |   /   \   |   /   \   |
00005  |  /     \  |  /     \  |  /     \  |  /     \  |
00006  | /       \ | /       \ | /       \ | /       \ |
00007  |_____|   |_____|   |_____|   |_____|   |
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  =====*/
00021  #include "kerneltypes.h"
00022  #include "mark3cfg.h"
00023
00024  #include "blocking.h"
00025  #include "thread.h"
00026
00027  #define _CAN_HAS_DEBUG
00028  //--[Autogenerated - Do Not Modify]-----
00029  #include "dbg_file_list.h"
00030  #include "buffalogger.h"
00031  #if defined(DBG_FILE)
00032  #error "Debug logging file token already defined! Bailing."
00033  #else
00034  #define DBG_FILE _DBG__KERNEL_BLOCKING_CPP
00035  #endif
00036  //--[End Autogenerated content]-----
00037  #include "kerneldebug.h"
00038
00039  namespace Mark3
00040  {
00041  //-----
00042  void BlockingObject::Block(Thread* pclThread_)
00043  {
00044      KERNEL_ASSERT(pclThread_);
00045      KERNEL_TRACE_1("Blocking Thread %d", static_cast<uint16_t>(pclThread_>
GetID()));
00046
00047      // Remove the thread from its current thread list (the "owner" list)
00048      // ... And add the thread to this object's block list
00049      Scheduler::Remove(pclThread_);
00050      m_clBlockList.Add(pclThread_);
00051
00052      // Set the "current" list location to the blocklist for this thread
00053      pclThread_>SetCurrent(&m_clBlockList);
00054      pclThread_>SetState(ThreadState::Blocked);
00055  }
00056
00057  //-----
00058  void BlockingObject::BlockPriority(Thread* pclThread_)
00059  {
00060      KERNEL_ASSERT(pclThread_);
00061      KERNEL_TRACE_1("Blocking Thread %d", static_cast<uint16_t>(pclThread_>
GetID()));
00062
00063      // Remove the thread from its current thread list (the "owner" list)
00064      // ... And add the thread to this object's block list
00065      Scheduler::Remove(pclThread_);
00066      m_clBlockList.AddPriority(pclThread_);
00067
00068      // Set the "current" list location to the blocklist for this thread
00069      pclThread_>SetCurrent(&m_clBlockList);
00070      pclThread_>SetState(ThreadState::Blocked);
00071  }
00072
00073  //-----
00074  void BlockingObject::UnBlock(Thread* pclThread_)
00075  {
00076      KERNEL_ASSERT(pclThread_);
00077      KERNEL_TRACE_1("Unblocking Thread %d", static_cast<uint16_t>(pclThread_>
GetID()));
00078
00079      // Remove the thread from its current thread list (the "owner" list)
00080      pclThread_>GetCurrent()->Remove(pclThread_);
00081
00082      // Put the thread back in its active owner's list. This is usually
00083      // the ready-queue at the thread's original priority.
00084      Scheduler::Add(pclThread_);
00085
00086      // Tag the thread's current list location to its owner
00087      pclThread_>SetCurrent(pclThread_>GetOwner());
00088      pclThread_>SetState(ThreadState::Ready);

```

```
00089 }
00090 } //namespace Mark3
```

20.29 /home/moslevin/projects/github/m3-repo/kernel/src/condvar.cpp File Reference

Condition Variable implementation.

```
#include "mark3cfg.h"
#include "condvar.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
```

Namespaces

- [Mark3](#)

Class providing the software-interrupt required for context-switching in the kernel.

20.29.1 Detailed Description

Condition Variable implementation.

Definition in file [condvar.cpp](#).

20.30 condvar.cpp

```
00001
00002 /*=====
00003
00004
00005
00006
00007
00008
00009
00010 --[Mark3 Realtime Platform]-----
00011
00012 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00013 See license.txt for more information
00014 =====*/
00021 #include "mark3cfg.h"
00022 #include "condvar.h"
00023
00024 #define _CAN_HAS_DEBUG
00025 //--[Autogenerated - Do Not Modify]-----
00026 #include "dbg_file_list.h"
00027 #include "buffalogger.h"
00028 #if defined(DBG_FILE)
00029 # error "Debug logging file token already defined! Bailing."
00030 #else
00031 # define DBG_FILE _DBG__KERNEL_CONDVAR_CPP
00032 #endif
00033 //--[End Autogenerated content]-----
00034
00035 namespace Mark3 {
00036
00037 //-----
00038 void ConditionVariable::Init()
00039
00040 {
00041     m_clMutex.Init();
00042     m_clSemaphore.Init(0, 255);
```

```

00043 }
00044
00045 //-----
00046 void ConditionVariable::Wait(Mutex* pclMutex_)
00047 {
00048     m_clMutex.Claim();
00049     pclMutex_>Release();
00050     m_u8Waiters++;
00051     m_clMutex.Release();
00052     m_clSemaphore.Pend();
00053     pclMutex_>Claim();
00054 }
00055
00056 #if KERNEL_USE_TIMEOUTS
00057 //-----
00058 bool ConditionVariable::Wait(Mutex* pclMutex_, uint32_t u32WaitTimeMS_)
00059 {
00060     m_clMutex.Claim();
00061     pclMutex_>Release();
00062     m_u8Waiters++;
00063     m_clMutex.Release();
00064     if (!m_clSemaphore.Pend(u32WaitTimeMS_)) {
00065         return false;
00066     }
00067     return pclMutex_>Claim(u32WaitTimeMS_);
00068 }
00069 #endif
00070
00071 //-----
00072 void ConditionVariable::Signal()
00073 {
00074     m_clMutex.Claim();
00075     if (m_u8Waiters) {
00076         m_u8Waiters--;
00077         m_clSemaphore.Post();
00078     }
00079     m_clMutex.Release();
00080 }
00081
00082 //-----
00083 void ConditionVariable::Broadcast()
00084 {
00085     m_clMutex.Claim();
00086     while (m_u8Waiters > 0) {
00087         m_u8Waiters--;
00088         m_clSemaphore.Post();
00089     }
00090     m_clMutex.Release();
00091 }
00092
00093 // namespace Mark3
00094 }

```

20.31 /home/moslevin/projects/github/m3-repo/kernel/src/eventflag.cpp File Reference

Event Flag Blocking Object/IPC-Object implementation.

```

#include "mark3cfg.h"
#include "blocking.h"
#include "kernel.h"
#include "thread.h"
#include "eventflag.h"
#include "kernelaware.h"
#include "kerneldebug.h"
#include "dbg_file_list.h"
#include "buffalogger.h"

```

20.31.1 Detailed Description

Event Flag Blocking Object/IPC-Object implementation.

Definition in file [eventflag.cpp](#).

20.32 eventflag.cpp

```

00001 /*
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====
00014 #include "mark3cfg.h"
00015 #include "blocking.h"
00016 #include "kernel.h"
00017 #include "thread.h"
00018 #include "eventflag.h"
00019 #include "kernelaware.h"
00020 #include "kerneldebug.h"
00021
00022 #define _CAN_HAS_DEBUG
00023 //--[Autogenerated - Do Not Modify]-----
00024 #include "dbg_file_list.h"
00025 #include "buffallogger.h"
00026 #if defined(DBG_FILE)
00027 #error "Debug logging file token already defined! Bailing."
00028 #else
00029 #define DBG_FILE _DBG__KERNEL_EVENTFLAG_CPP
00030 #endif
00031 //--[End Autogenerated content]-----
00032
00033 #if KERNEL_USE_EVENTFLAG
00034
00035 #if KERNEL_USE_TIMEOUTS
00036 #include "timerlist.h"
00037 namespace Mark3 {
00038 namespace {
00039
00040 void TimedEventFlag_Callback(Thread* pclOwner_, void* pvData_)
00041 {
00042     auto* pclEventFlag = static_cast<EventFlag*>(pvData_);
00043
00044     pclEventFlag->WakeMe(pclOwner_);
00045     pclOwner_->SetExpired(true);
00046     pclOwner_->SetEventFlagMask(0);
00047
00048     if (pclOwner_->GetCurPriority() >= Scheduler::GetCurrentThread()->
00049         GetCurPriority()) {
00050         Thread::Yield();
00051     }
00052 }
00053 } // anonymous namespace
00054 }
00055
00056 EventFlag::~EventFlag()
00057 {
00058     // If there are any threads waiting on this object when it goes out
00059     // of scope, set a kernel panic.
00060     if (m_clBlockList.HighestWaiter() != nullptr) {
00061         Kernel::Panic(PANIC_ACTIVE_EVENTFLAG_DESCOPED);
00062     }
00063 }
00064
00065 void EventFlag::Init()
00066 {
00067     #if KERNEL_EXTRA_CHECKS
00068         KERNEL_ASSERT(!m_clBlockList.GetHead());
00069     #endif
00070     m_ul6SetMask = 0;
00071     #if KERNEL_EXTRA_CHECKS

```

```

00087     SetInitialized();
00088 #endif
00089 }
00090
00091 //-----
00092 void EventFlag::WakeMe(Thread* pClChosenOne_)
00093 {
00094     #if KERNEL_EXTRA_CHECKS
00095         KERNEL_ASSERT(IsInitialized());
00096     #endif
00097
00098     Unblock(pClChosenOne_);
00099 }
00100 #endif
00101
00102 //-----
00103 #if KERNEL_USE_TIMEOUTS
00104 uint16_t EventFlag::Wait_i(uint16_t u16Mask_,
00105     EventFlagOperation eMode_, uint32_t u32TimeMS_)
00106 #else
00107 uint16_t EventFlag::Wait_i(uint16_t u16Mask_,
00108     EventFlagOperation eMode_)
00109 #endif
00110 {
00111     #if KERNEL_EXTRA_CHECKS
00112         KERNEL_ASSERT(IsInitialized());
00113     #endif
00114
00115     auto bThreadYield = false;
00116     auto bMatch = false;
00117
00118     #if KERNEL_USE_TIMEOUTS
00119         Timer clEventTimer;
00120         auto bUseTimer = false;
00121     #endif
00122
00123     // Ensure we're operating in a critical section while we determine
00124     // whether or not we need to block the current thread on this object.
00125     CS_ENTER();
00126
00127     // Check to see whether or not the current mask matches any of the
00128     // desired bits.
00129     g_pclCurrent->SetEventFlagMask(u16Mask_);
00130
00131     if ((eMode_ == EventFlagOperation::All_Set) || (eMode_ ==
00132         EventFlagOperation::All_Clear)) {
00133         // Check to see if the flags in their current state match all of
00134         // the set flags in the event flag group, with this mask.
00135         if ((m_u16SetMask & u16Mask_) == u16Mask_) {
00136             bMatch = true;
00137             g_pclCurrent->SetEventFlagMask(u16Mask_);
00138         }
00139     } else if ((eMode_ == EventFlagOperation::Any_Set) || (eMode_ ==
00140         EventFlagOperation::Any_Clear)) {
00141         // Check to see if the existing flags match any of the set flags in
00142         // the event flag group with this mask
00143         if ((m_u16SetMask & u16Mask_) != 0) {
00144             bMatch = true;
00145             g_pclCurrent->SetEventFlagMask(m_u16SetMask & u16Mask_);
00146         }
00147     }
00148
00149     // We're unable to match this pattern as-is, so we must block.
00150     if (!bMatch) {
00151         // Reset the current thread's event flag mask & mode
00152         g_pclCurrent->SetEventFlagMask(u16Mask_);
00153         g_pclCurrent->SetEventFlagMode(eMode_);
00154
00155         #if KERNEL_USE_TIMEOUTS
00156             if (u32TimeMS_ != 0u) {
00157                 g_pclCurrent->SetExpired(false);
00158                 clEventTimer.Init();
00159                 clEventTimer.Start(false, u32TimeMS_, TimedEventFlag_Callback, this);
00160                 bUseTimer = true;
00161             }
00162         #endif
00163
00164         // Add the thread to the object's block-list.
00165         BlockPriority(g_pclCurrent);
00166
00167         // Trigger that
00168         bThreadYield = true;
00169     }
00170
00171     // If bThreadYield is set, it means that we've blocked the current thread,
00172     // and must therefore rerun the scheduler to determine what thread to
00173     // switch to.

```

```

00170     if (bThreadYield) {
00171         // Switch threads immediately
00172         Thread::Yield();
00173     }
00174
00175     // Exit the critical section and return back to normal execution
00176     CS_EXIT();
00177
00182 #if KERNEL_USE_TIMEOUTS
00183     if (bUseTimer && bThreadYield) {
00184         clEventTimer.Stop();
00185     }
00186 #endif
00187
00188     return g_pclCurrent->GetEventFlagMask();
00189 }
00190
00191 //-----
00192 uint16_t EventFlag::Wait(uint16_t ul6Mask_, EventFlagOperation eMode_)
00193 {
00194     #if KERNEL_USE_TIMEOUTS
00195         return Wait_i(ul6Mask_, eMode_, 0);
00196     #else
00197         return Wait_i(ul6Mask_, eMode_);
00198     #endif
00199 }
00200
00201 #if KERNEL_USE_TIMEOUTS
00202 //-----
00203 uint16_t EventFlag::Wait(uint16_t ul6Mask_, EventFlagOperation eMode_,
00204                          uint32_t u32TimeMS_)
00205 {
00206     return Wait_i(ul6Mask_, eMode_, u32TimeMS_);
00207 }
00208 #endif
00209 //-----
00210 void EventFlag::Set(uint16_t ul6Mask_)
00211 {
00212     #if KERNEL_EXTRA_CHECKS
00213         KERNEL_ASSERT(IsInitialized());
00214     #endif
00215
00216     Thread* pclPrev;
00217     Thread* pclCurrent;
00218     auto bReschedule = false;
00219     uint16_t ul6NewMask;
00220
00221     CS_ENTER();
00222
00223     // Walk through the whole block list, checking to see whether or not
00224     // the current flag set now matches any/all of the masks and modes of
00225     // the threads involved.
00226
00227     m_ul6SetMask |= ul6Mask_;
00228     ul6NewMask = m_ul6SetMask;
00229
00230     // Start at the head of the list, and iterate through until we hit the
00231     // "head" element in the list again. Ensure that we handle the case where
00232     // we remove the first or last elements in the list, or if there's only
00233     // one element in the list.
00234     pclCurrent = static_cast<Thread*>(m_clBlockList.GetHead());
00235
00236     // Do nothing when there are no objects blocking.
00237     if (pclCurrent != nullptr) {
00238         // First loop - process every thread in the block-list and check to
00239         // see whether or not the current flags match the event-flag conditions
00240         // on the thread.
00241         do {
00242             pclPrev = pclCurrent;
00243             pclCurrent = static_cast<Thread*>(pclCurrent->GetNext());
00244
00245             // Read the thread's event mask/mode
00246             uint16_t ul6ThreadMask = pclPrev->GetEventFlagMask();
00247             auto eThreadMode = pclPrev->GetEventFlagMode();
00248
00249             // For the "any" mode - unblock the blocked threads if one or more bits
00250             // in the thread's bitmask match the object's bitmask
00251             if ((EventFlagOperation::Any_Set == eThreadMode) || (
00252                 EventFlagOperation::Any_Clear == eThreadMode)) {
00253                 if ((ul6ThreadMask & m_ul6SetMask) != 0) {
00254                     pclPrev->SetEventFlagMode(
00255                         EventFlagOperation::Pending_Unblock);
00256                     pclPrev->SetEventFlagMask(m_ul6SetMask & ul6ThreadMask);
00257                     bReschedule = true;
00258                 }
00259                 // If the "clear" variant is set, then clear the bits in the mask

```

```

00258             // that caused the thread to unblock.
00259             if (EventFlagOperation::Any_Clear == eThreadMode) {
00260                 ul6NewMask &= ~(ul6ThreadMask & ul6Mask_);
00261             }
00262         }
00263     }
00264     // For the "all" mode, every set bit in the thread's requested bitmask must
00265     // match the object's flag mask.
00266     else if ((EventFlagOperation::All_Set == eThreadMode) || (
EventFlagOperation::All_Clear == eThreadMode)) {
00267         if ((ul6ThreadMask & m_ul6SetMask) == ul6ThreadMask) {
00268             pclPrev->SetEventFlagMode(
EventFlagOperation::Pending_Unblock);
00269             pclPrev->SetEventFlagMask(ul6ThreadMask);
00270             bReschedule = true;
00271
00272             // If the "clear" variant is set, then clear the bits in the mask
00273             // that caused the thread to unblock.
00274             if (EventFlagOperation::All_Clear == eThreadMode) {
00275                 ul6NewMask &= ~(ul6ThreadMask & ul6Mask_);
00276             }
00277         }
00278     }
00279 }
00280 // To keep looping, ensure that there's something in the list, and
00281 // that the next item isn't the head of the list.
00282 while (pclPrev != m_clBlockList.GetTail());
00283
00284 // Second loop - go through and unblock all of the threads that
00285 // were tagged for unblocking.
00286 pclCurrent = static_cast<Thread*>(m_clBlockList.GetHead());
00287 auto bIsTail = false;
00288 do {
00289     pclPrev = pclCurrent;
00290     pclCurrent = static_cast<Thread*>(pclCurrent->GetNext());
00291
00292     // Check to see if this is the condition to terminate the loop
00293     if (pclPrev == m_clBlockList.GetTail()) {
00294         bIsTail = true;
00295     }
00296
00297     // If the first pass indicated that this thread should be
00298     // unblocked, then unblock the thread
00299     if (pclPrev->GetEventFlagMode() ==
EventFlagOperation::Pending_Unblock) {
00300         Unblock(pclPrev);
00301     }
00302 } while (!bIsTail);
00303 }
00304
00305 // If we awoke any threads, re-run the scheduler
00306 if (bReschedule) {
00307     Thread::Yield();
00308 }
00309
00310 // Update the bitmask based on any "clear" operations performed along
00311 // the way
00312 m_ul6SetMask = ul6NewMask;
00313
00314 // Restore interrupts - will potentially cause a context switch if a
00315 // thread is unblocked.
00316 CS_EXIT();
00317 }
00318
00319 //-----
00320 void EventFlag::Clear(uint16_t ul6Mask_)
00321 {
00322     #if KERNEL_EXTRA_CHECKS
00323         KERNEL_ASSERT(IsInitialized());
00324     #endif
00325
00326     // Just clear the bitfields in the local object.
00327     CS_ENTER();
00328     m_ul6SetMask &= ~ul6Mask_;
00329     CS_EXIT();
00330 }
00331
00332 //-----
00333 uint16_t EventFlag::GetMask()
00334 {
00335     #if KERNEL_EXTRA_CHECKS
00336         KERNEL_ASSERT(IsInitialized());
00337     #endif
00338
00339     // Return the presently held event flag values in this object. Ensure
00340     // we get this within a critical section to guarantee atomicity.
00341     uint16_t ul6Return;

```



```

00342     CS_ENTER();
00343     ul6Return = m_ul6SetMask;
00344     CS_EXIT();
00345     return ul6Return;
00346 }
00347 } // namespace Mark3
00348 #endif // KERNEL_USE_EVENTFLAG

```

20.33 /home/moslevin/projects/github/m3-repo/kernel/src/kernel.cpp File Reference

Kernel initialization and startup code.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "kernel.h"
#include "scheduler.h"
#include "thread.h"
#include "threadport.h"
#include "timerlist.h"
#include "message.h"
#include "profile.h"
#include "kernelprofile.h"
#include "autoalloc.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"
#include "tracebuffer.h"

```

Namespaces

- [Mark3](#)

Class providing the software-interrupt required for context-switching in the kernel.

20.33.1 Detailed Description

Kernel initialization and startup code.

Definition in file [kernel.cpp](#).

20.34 kernel.cpp

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00021 #include "kerneltypes.h"

```

```

00022 #include "mark3cfg.h"
00023
00024 #include "kernel.h"
00025 #include "scheduler.h"
00026 #include "thread.h"
00027 #include "threadport.h"
00028 #include "timerlist.h"
00029 #include "message.h"
00030 #include "profile.h"
00031 #include "kernelprofile.h"
00032 #include "autoalloc.h"
00033
00034 #define _CAN_HAS_DEBUG
00035 //--[Autogenerated - Do Not Modify]-----
00036 #include "dbg_file_list.h"
00037 #include "buffalogger.h"
00038 #if defined(DBG_FILE)
00039 #error "Debug logging file token already defined! Bailing."
00040 #else
00041 #define DBG_FILE _DBG__KERNEL_KERNEL_CPP
00042 #endif
00043 //--[End Autogenerated content]-----
00044 #include "kerneldebug.h"
00045 #include "tracebuffer.h"
00046
00047 extern "C" {
00048     int __cxa_guard_acquire(long long int* wha) { return 0; }
00049     void __cxa_guard_release(long long int* wha) {}
00050     void atexit(void) {}
00051 }
00052
00053 namespace Mark3
00054 {
00055
00056     bool Kernel::m_bIsStarted;
00057     bool Kernel::m_bIsPanic;
00058     PanicFunc Kernel::m_pfPanic;
00059     #if KERNEL_USE_IDLE_FUNC
00060     IdleFunc Kernel::m_pfIdle;
00061     FakeThread_t Kernel::m_clIdle;
00062     #endif
00063
00064     #if KERNEL_USE_THREAD_CALLOUTS
00065     ThreadCreateCallout Kernel::m_pfThreadCreateCallout;
00066     ThreadExitCallout Kernel::m_pfThreadExitCallout;
00067     ThreadContextCallout Kernel::m_pfThreadContextCallout;
00068     #endif
00069
00070     #if KERNEL_USE_STACK_GUARD
00071     uint16_t Kernel::m_ul6GuardThreshold;
00072     #endif
00073
00074     //--[Autogenerated content]-----
00075     void Kernel::Init(void)
00076     {
00077         #if KERNEL_USE_AUTO_ALLOC
00078             AutoAlloc::Init();
00079         #endif
00080         #if KERNEL_USE_IDLE_FUNC
00081             (reinterpret_cast<Thread*>(&m_clIdle))->InitIdle();
00082         #endif
00083         #if KERNEL_USE_DEBUG && !KERNEL_AWARE_SIMULATION
00084             TraceBuffer::Init();
00085         #endif
00086         KERNEL_TRACE("Initializing Mark3 Kernel");
00087
00088         // Initialize the global kernel data - scheduler, timer-scheduler, and
00089         // the global message pool.
00090         Scheduler::Init();
00091         #if KERNEL_USE_TIMERS
00092             TimerScheduler::Init();
00093         #endif
00094         #if KERNEL_USE_STACK_GUARD
00095             m_ul6GuardThreshold = KERNEL_STACK_GUARD_DEFAULT;
00096         #endif
00097     }
00098
00099     //--[Autogenerated content]-----
00100     void Kernel::Start(void)
00101     {
00102         KERNEL_TRACE("Starting Mark3 Scheduler");
00103         m_bIsStarted = true;
00104         ThreadPort::StartThreads();
00105         KERNEL_TRACE("Error starting Mark3 Scheduler");
00106     }
00107
00108     //--[Autogenerated content]-----

```

```

00109 void Kernel::Panic(uint16_t u16Cause_)
00110 {
00111     m_bIsPanic = true;
00112     if (m_pfPanic != nullptr) {
00113         m_pfPanic(u16Cause_);
00114     } else {
00115 #if KERNEL_AWARE_SIMULATION
00116         KernelAware::Print("Panic\n");
00117         KernelAware::Trace(0, 0, u16Cause_, g_pclCurrent->
GetID());
00118         KernelAware::ExitSimulator();
00119 #endif
00120         while (true) { }
00121     }
00122 }
00123 } //namespace Mark3
00124

```

20.35 /home/moslevin/projects/github/m3-repo/kernel/src/kernelaware.cpp File Reference

Kernel aware simulation support.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "kernelaware.h"
#include "threadport.h"
#include "dbg_file_list.h"
#include "buffalogger.h"

```

20.35.1 Detailed Description

Kernel aware simulation support.

Definition in file [kernelaware.cpp](#).

20.36 kernelaware.cpp

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00021 #include "kerneltypes.h"
00022 #include "mark3cfg.h"
00023 #include "kernelaware.h"
00024 #include "threadport.h"
00025
00026 #define _CAN_HAS_DEBUG
00027 //--[Autogenerated - Do Not Modify]-----
00028 #include "dbg_file_list.h"
00029 #include "buffalogger.h"
00030 #if defined(DBG_FILE)
00031 #error "Debug logging file token already defined! Bailing."
00032 #else
00033 #define DBG_FILE _DBG___KERNEL_KERNELAWARE_CPP

```

```

00034 #endif
00035 //--[End Autogenerated content]-----
00036
00037 #if KERNEL_AWARE_SIMULATION
00038
00039 //-----
00044 typedef enum {
00045     KA_COMMAND_IDLE = 0,
00046     KA_COMMAND_PROFILE_INIT,
00047     KA_COMMAND_PROFILE_START,
00048     KA_COMMAND_PROFILE_STOP,
00049     KA_COMMAND_PROFILE_REPORT,
00050     KA_COMMAND_EXIT_SIMULATOR,
00051     KA_COMMAND_TRACE_0,
00052     KA_COMMAND_TRACE_1,
00053     KA_COMMAND_TRACE_2,
00054     KA_COMMAND_PRINT
00055 } KernelAwareCommand_t;
00056
00057 //-----
00066 typedef union {
00067     volatile uint16_t au16Buffer[5];
00068
00072     struct {
00073         volatile const char* szName;
00074     } Profiler;
00079     struct {
00080         volatile uint16_t u16File;
00081         volatile uint16_t u16Line;
00082         volatile uint16_t u16Arg1;
00083         volatile uint16_t u16Arg2;
00084     } Trace;
00089     struct {
00090         volatile const char* szString;
00091     } Print;
00092 } KernelAwareData_t;
00093
00094 //-----
00095 // Must not live in Mark3 namespace
00096 using namespace Mark3;
00097 volatile bool g_bIsKernelAware;
00098 volatile uint8_t g_u8KACommand;
00099 KernelAwareData_t g_stKAData;
00100
00101 namespace Mark3
00102 {
00103 //-----
00104 void Trace_i(
00105     uint16_t u16File_, uint16_t u16Line_, uint16_t u16Arg1_, uint16_t u16Arg2_, KernelAwareCommand_t eCmd_)
00106 {
00107     CS_ENTER();
00108     g_stKAData.Trace.u16File = u16File_;
00109     g_stKAData.Trace.u16Line = u16Line_;
00110     g_stKAData.Trace.u16Arg1 = u16Arg1_;
00111     g_stKAData.Trace.u16Arg2 = u16Arg2_;
00112     g_u8KACommand = eCmd_;
00113     CS_EXIT();
00114 }
00115
00116 //-----
00117 void KernelAware::ProfileInit(const char* szStr_)
00118 {
00119     CS_ENTER();
00120     g_stKAData.Profiler.szName = szStr_;
00121     g_u8KACommand = KA_COMMAND_PROFILE_INIT;
00122     CS_EXIT();
00123 }
00124
00125 //-----
00126 void KernelAware::ProfileStart(void)
00127 {
00128     g_u8KACommand = KA_COMMAND_PROFILE_START;
00129 }
00130
00131 //-----
00132 void KernelAware::ProfileStop(void)
00133 {
00134     g_u8KACommand = KA_COMMAND_PROFILE_STOP;
00135 }
00136
00137 //-----
00138 void KernelAware::ProfileReport(void)
00139 {
00140     g_u8KACommand = KA_COMMAND_PROFILE_REPORT;
00141 }
00142
00143 //-----

```

```

00144 void KernelAware::ExitSimulator(void)
00145 {
00146     g_u8KACommand = KA_COMMAND_EXIT_SIMULATOR;
00147 }
00148
00149 //-----
00150 void KernelAware::Trace(uint16_t u16File_, uint16_t u16Line_)
00151 {
00152     Trace_i(u16File_, u16Line_, 0, 0, KA_COMMAND_TRACE_0);
00153 }
00154
00155 //-----
00156 void KernelAware::Trace(uint16_t u16File_, uint16_t u16Line_, uint16_t u16Arg1_)
00157 {
00158     Trace_i(u16File_, u16Line_, u16Arg1_, 0, KA_COMMAND_TRACE_1);
00159 }
00160 //-----
00161 void KernelAware::Trace(uint16_t u16File_, uint16_t u16Line_, uint16_t u16Arg1_, uint16_t
    u16Arg2_)
00162 {
00163     Trace_i(u16File_, u16Line_, u16Arg1_, u16Arg2_, KA_COMMAND_TRACE_2);
00164 }
00165
00166 //-----
00167 void KernelAware::Print(const char* szStr_)
00168 {
00169     CS_ENTER();
00170     g_stKADData.Print.szString = szStr_;
00171     g_u8KACommand = KA_COMMAND_PRINT;
00172     CS_EXIT();
00173 }
00174
00175 //-----
00176 bool KernelAware::IsSimulatorAware(void)
00177 {
00178     return g_bIsKernelAware;
00179 }
00180 } //namespace Mark3
00181 #endif

```

20.37 /home/moslevin/projects/github/m3-repo/kernel/src/ksemaphore.cpp File Reference

Semaphore Blocking-Object Implemenation.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ksemaphore.h"
#include "blocking.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"

```

20.37.1 Detailed Description

Semaphore Blocking-Object Implemenation.

Definition in file [ksemaphore.cpp](#).


```

00102
00103     // Remove from the semaphore waitlist and back to its ready list.
00104     Unblock(pclChosenOne);
00105
00106     // Call a task switch if higher or equal priority thread
00107     if (pclChosenOne->GetCurPriority() >= Scheduler::GetCurrentThread()->
GetCurPriority()) {
00108         return 1;
00109     }
00110     return 0;
00111 }
00112
00113 //-----
00114 void Semaphore::Init(uint16_t ul6InitVal_, uint16_t ul6MaxVal_)
00115 {
00116     #if KERNEL_EXTRA_CHECKS
00117         KERNEL_ASSERT(!m_clBlockList.GetHead());
00118     #endif
00119
00120     // Copy the paramters into the object - set the maximum value for this
00121     // semaphore to implement either binary or counting semaphores, and set
00122     // the initial count. Clear the wait list for this object.
00123     m_ul6Value = ul6InitVal_;
00124     m_ul6MaxValue = ul6MaxVal_;
00125
00126     #if KERNEL_EXTRA_CHECKS
00127         SetInitialized();
00128     #endif
00129 }
00130
00131
00132 //-----
00133 bool Semaphore::Post()
00134 {
00135     #if KERNEL_EXTRA_CHECKS
00136         KERNEL_ASSERT(IsInitialized());
00137     #endif
00138
00139     KERNEL_TRACE_1("Posting semaphore, Thread %d", static_cast<uint16_t>(g_pclCurrent->
GetID()));
00140
00141     auto bThreadWake = false;
00142     auto bBail = false;
00143     // Increment the semaphore count - we can mess with threads so ensure this
00144     // is in a critical section. We don't just disable the scheduler since
00145     // we want to be able to do this from within an interrupt context as well.
00146     CS_ENTER();
00147
00148     // If nothing is waiting for the semaphore
00149     if (m_clBlockList.GetHead() == NULL) {
00150         // Check so see if we've reached the maximum value in the semaphore
00151         if (m_ul6Value < m_ul6MaxValue) {
00152             // Increment the count value
00153             m_ul6Value++;
00154         } else {
00155             // Maximum value has been reached, bail out.
00156             bBail = true;
00157         }
00158     } else {
00159         // Otherwise, there are threads waiting for the semaphore to be
00160         // posted, so wake the next one (highest priority goes first).
00161         bThreadWake = (WakeNext() != 0u);
00162     }
00163
00164     CS_EXIT();
00165
00166     // If we weren't able to increment the semaphore count, fail out.
00167     if (bBail) {
00168         return false;
00169     }
00170
00171     // if bThreadWake was set, it means that a higher-priority thread was
00172     // woken. Trigger a context switch to ensure that this thread gets
00173     // to execute next.
00174     if (bThreadWake) {
00175         Thread::Yield();
00176     }
00177     return true;
00178 }
00179
00180 //-----
00181 #if KERNEL_USE_TIMEOUTS
00182 bool Semaphore::Pend_i(uint32_t u32WaitTimeMS_)
00183 #else
00184 void Semaphore::Pend_i(void)
00185 #endif
00186 {

```

```

00187 #if KERNEL_EXTRA_CHECKS
00188     KERNEL_ASSERT(IsInitialized());
00189 #endif
00190
00191     KERNEL_TRACE_1("Pending semaphore, Thread %d", static_cast<uint16_t>(g_pclCurrent->
00192         GetID()));
00193
00194 #if KERNEL_USE_TIMEOUTS
00195     Timer clSemTimer;
00196     auto bUseTimer = false;
00197 #endif
00198
00199     // Once again, messing with thread data - ensure
00200     // we're doing all of these operations from within a thread-safe context.
00201     CS_ENTER();
00202
00203     // Check to see if we need to take any action based on the semaphore count
00204     if (m_ul6Value != 0) {
00205         // The semaphore count is non-zero, we can just decrement the count
00206         // and go along our merry way.
00207         m_ul6Value--;
00208     } else {
00209         // The semaphore count is zero - we need to block the current thread
00210         // and wait until the semaphore is posted from elsewhere.
00211         #if KERNEL_USE_TIMEOUTS
00212             if (u32WaitTimeMS_ != 0u) {
00213                 g_pclCurrent->SetExpired(false);
00214                 clSemTimer.Init();
00215                 clSemTimer.Start(false, u32WaitTimeMS_, TimedSemaphore_Callback, this);
00216                 bUseTimer = true;
00217             }
00218         #endif
00219         BlockPriority(g_pclCurrent);
00220
00221         // Switch Threads immediately
00222         Thread::Yield();
00223     }
00224
00225     CS_EXIT();
00226
00227 #if KERNEL_USE_TIMEOUTS
00228     if (bUseTimer) {
00229         clSemTimer.Stop();
00230         return (static_cast<int>(g_pclCurrent->GetExpired()) == 0);
00231     }
00232     return true;
00233 #endif
00234 }
00235
00236 //-----
00237 void Semaphore::Pend()
00238 {
00239     #if KERNEL_USE_TIMEOUTS
00240         Pend_i(0);
00241     #else
00242         Pend_i();
00243     #endif
00244 }
00245
00246 #if KERNEL_USE_TIMEOUTS
00247 //-----
00248 bool Semaphore::Pend(uint32_t u32WaitTimeMS_)
00249 {
00250     return Pend_i(u32WaitTimeMS_);
00251 }
00252 #endif
00253
00254 //-----
00255 uint16_t Semaphore::GetCount()
00256 {
00257     #if KERNEL_EXTRA_CHECKS
00258         KERNEL_ASSERT(IsInitialized());
00259     #endif
00260     uint16_t ul6Ret;
00261     CS_ENTER();
00262     ul6Ret = m_ul6Value;
00263     CS_EXIT();
00264     return ul6Ret;
00265 }
00266 } //namespace Mark3
00267 #endif

```


20.39 /home/moslevin/projects/github/m3-repo/kernel/src/ll.cpp File Reference

Core Linked-List implementation, from which all kernel objects are derived.

```
#include "kerneltypes.h"
#include "kernel.h"
#include "ll.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"
```

Namespaces

- [Mark3](#)

Class providing the software-interrupt required for context-switching in the kernel.

20.39.1 Detailed Description

Core Linked-List implementation, from which all kernel objects are derived.

Definition in file [ll.cpp](#).

20.40 ll.cpp

```
00001  /*=====
00002
00003  _____
00004  |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |
00005  |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |
00006  |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |
00007  |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |
00008  |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  =====*/
00022  #include "kerneltypes.h"
00023  #include "kernel.h"
00024  #include "ll.h"
00025
00026  #define _CAN_HAS_DEBUG
00027  //--[Autogenerated - Do Not Modify]-----
00028  #include "dbg_file_list.h"
00029  #include "buffalogger.h"
00030  #if defined(DBG_FILE)
00031  #error "Debug logging file token already defined! Bailing."
00032  #else
00033  #define DBG_FILE _DBG__KERNEL_LL_CPP
00034  #endif
00035  //--[End Autogenerated content]-----
00036
00037  #include "kerneldebug.h"
00038
00039  namespace Mark3
00040  {
00041  //-----
00042  void LinkListNode::ClearNode()
00043  {
00044      next = NULL;
00045      prev = NULL;
00046  }
00047
00048  //-----
```

```

00049 void DoubleLinkedList::Add(LinkListNode* node_)
00050 {
00051     KERNEL_ASSERT(node_);
00052
00053     node_>prev = m_pclTail;
00054     node_>next = NULL;
00055
00056     // If the list is empty, initilize the head
00057     if (m_pclHead == nullptr) {
00058         m_pclHead = node_;
00059     }
00060     // Otherwise, adjust the tail's next pointer
00061     else {
00062         m_pclTail->next = node_;
00063     }
00064
00065     // Move the tail node, and assign it to the new node just passed in
00066     m_pclTail = node_;
00067 }
00068
00069 //-----
00070 void DoubleLinkedList::Remove(LinkListNode* node_)
00071 {
00072     KERNEL_ASSERT(node_);
00073
00074     if (node_>prev != nullptr) {
00075 #if SAFE_UNLINK
00076         if (node_>prev->next != node_) {
00077             Kernel::Panic(PANIC_LIST_UNLINK_FAILED);
00078         }
00079 #endif
00080         node_>prev->next = node_>next;
00081     }
00082     if (node_>next != nullptr) {
00083 #if SAFE_UNLINK
00084         if (node_>next->prev != node_) {
00085             Kernel::Panic(PANIC_LIST_UNLINK_FAILED);
00086         }
00087 #endif
00088         node_>next->prev = node_>prev;
00089     }
00090     if (node_ == m_pclHead) {
00091         m_pclHead = node_>next;
00092     }
00093     if (node_ == m_pclTail) {
00094         m_pclTail = node_>prev;
00095     }
00096     node_>ClearNode();
00097 }
00098
00099 //-----
00100 void CircularLinkedList::Add(LinkListNode* node_)
00101 {
00102     KERNEL_ASSERT(node_);
00103
00104     if (m_pclHead == nullptr) {
00105         // If the list is empty, initilize the nodes
00106         m_pclHead = node_;
00107         m_pclTail = node_;
00108     } else {
00109         // Move the tail node, and assign it to the new node just passed in
00110         m_pclTail->next = node_;
00111     }
00112
00113     // Add a node to the end of the linked list.
00114     node_>prev = m_pclTail;
00115     node_>next = m_pclHead;
00116
00117     m_pclTail = node_;
00118     m_pclHead->prev = node_;
00119 }
00120
00121 //-----
00122 void CircularLinkedList::Remove(LinkListNode* node_)
00123 {
00124     KERNEL_ASSERT(node_);
00125
00126     // Check to see if this is the head of the list...
00127     if ((node_ == m_pclHead) && (m_pclHead == m_pclTail)) {
00128         // Clear the head and tail pointers - nothing else left.
00129         m_pclHead = NULL;
00130         m_pclTail = NULL;
00131         return;
00132     }
00133
00134 #if SAFE_UNLINK
00135     // Verify that all nodes are properly connected

```

```

00136     if ((node_>prev->next != node_) || (node_>next->prev != node_)) {
00137         Kernel::Panic(PANIC_LIST_UNLINK_FAILED);
00138     }
00139 #endif
00140
00141     // This is a circularly linked list - no need to check for connection,
00142     // just remove the node.
00143     node_>next->prev = node_>prev;
00144     node_>prev->next = node_>next;
00145
00146     if (node_ == m_pclHead) {
00147         m_pclHead = m_pclHead->next;
00148     }
00149     if (node_ == m_pclTail) {
00150         m_pclTail = m_pclTail->prev;
00151     }
00152     node_>ClearNode();
00153 }
00154
00155 //-----
00156 void CircularLinkedList::PivotForward()
00157 {
00158     if (m_pclHead != nullptr) {
00159         m_pclHead = m_pclHead->next;
00160         m_pclTail = m_pclTail->next;
00161     }
00162 }
00163
00164 //-----
00165 void CircularLinkedList::PivotBackward()
00166 {
00167     if (m_pclHead != nullptr) {
00168         m_pclHead = m_pclHead->prev;
00169         m_pclTail = m_pclTail->prev;
00170     }
00171 }
00172
00173 //-----
00174 void CircularLinkedList::InsertNodeBefore(
    LinkListNode* node_, LinkListNode* insert_)
00175 {
00176     KERNEL_ASSERT(node_);
00177
00178     node_>next = insert_;
00179     node_>prev = insert_>prev;
00180
00181     if (insert_>prev != nullptr) {
00182         insert_>prev->next = node_;
00183     }
00184     insert_>prev = node_;
00185 }
00186 } //namespace Mark3

```

20.41 /home/moslevin/projects/github/m3-repo/kernel/src/lockguard.cpp File Reference

Mutex RAIL helper class.

```
#include "lockguard.h"
```

Namespaces

- [Mark3](#)

Class providing the software-interrupt required for context-switching in the kernel.

20.41.1 Detailed Description

Mutex RAIL helper class.

Definition in file [lockguard.cpp](#).

20.42 lockguard.cpp

```

00001  /*=====
00002
00003  00004  00005  00006  00007  00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2018 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  =====*/
00020  #include "lockguard.h"
00021
00022  namespace Mark3 {
00023  LockGuard::LockGuard(Mutex *pclMutex_)
00024  : m_bIsAcquired{true}
00025  , m_pclMutex{pclMutex_}
00026  {
00027      m_pclMutex->Claim();
00028  }
00029
00030  #if KERNEL_USE_TIMEOUTS
00031  LockGuard::LockGuard(Mutex* pclMutex_, uint32_t u32TimeoutMs_)
00032  : m_pclMutex{pclMutex_}
00033  {
00034      m_bIsAcquired = m_pclMutex->Claim(u32TimeoutMs_);
00035  }
00036  #endif
00037
00038  LockGuard::~LockGuard()
00039  {
00040      if (m_bIsAcquired) {
00041          m_pclMutex->Release();
00042      }
00043  }
00044
00045  } // namespace Mark3

```

20.43 /home/moslevin/projects/github/m3-repo/kernel/src/mailbox.cpp File Reference

Mailbox + Envelope IPC mechanism.

```

#include "mark3cfg.h"
#include "kerneltypes.h"
#include "ksemaphore.h"
#include "mailbox.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"

```

20.43.1 Detailed Description

Mailbox + Envelope IPC mechanism.

Definition in file [mailbox.cpp](#).

20.44 mailbox.cpp

```

00001  /*=====
00002
00003  _____
00004  |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |
00005  |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |
00006  |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |
00007  |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  =====*/
00021  #include "mark3cfg.h"
00022  #include "kerneltypes.h"
00023  #include "ksemaphore.h"
00024  #include "mailbox.h"
00025
00026  #define _CAN_HAS_DEBUG
00027  //--[Autogenerated - Do Not Modify]-----
00028  #include "dbg_file_list.h"
00029  #include "buffalogger.h"
00030  #if defined(DBG_FILE)
00031  #error "Debug logging file token already defined! Bailing."
00032  #else
00033  #define DBG_FILE _DBG__KERNEL_MAILBOX_CPP
00034  #endif
00035  //--[End Autogenerated content]-----
00036
00037  #include "kerneldebug.h"
00038
00039  #if KERNEL_USE_MAILBOX
00040  namespace Mark3
00041  {
00042  //-----
00043  Mailbox::~Mailbox()
00044  {
00045      // If the mailbox isn't empty on destruction, kernel panic.
00046      if (m_ul6Free != m_ul6Count) {
00047          Kernel::Panic(PANIC_ACTIVE_MAILBOX_DESCOPEDED);
00048      }
00049  }
00050
00051  //-----
00052  void Mailbox::Init(void* pvBuffer_, uint16_t ul6BufferSize_, uint16_t ul6ElementSize_)
00053  {
00054      KERNEL_ASSERT(ul6BufferSize_);
00055      KERNEL_ASSERT(ul6ElementSize_);
00056      KERNEL_ASSERT(pvBuffer_);
00057
00058      m_pvBuffer = pvBuffer_;
00059      m_ul6ElementSize = ul6ElementSize_;
00060
00061      m_ul6Count = (ul6BufferSize_ / ul6ElementSize_);
00062      m_ul6Free = m_ul6Count;
00063
00064      m_ul6Head = 0;
00065      m_ul6Tail = 0;
00066
00067      // We use the counting semaphore to implement blocking - with one element
00068      // in the mailbox corresponding to a post/pend operation in the semaphore.
00069      m_clRecvSem.Init(0, m_ul6Free);
00070
00071      #if KERNEL_USE_TIMEOUTS
00072      // Binary semaphore is used to track any threads that are blocked on a
00073      // "send" due to lack of free slots.
00074      m_clSendSem.Init(0, 1);
00075      #endif
00076  }
00077
00078  //-----
00079  #if KERNEL_USE_AUTO_ALLOC
00080  Mailbox* Mailbox::Init(uint16_t ul6BufferSize_, uint16_t ul6ElementSize_)
00081  {
00082      auto* pclNew = AutoAlloc::NewMailbox();
00083      auto* pvBuffer = AutoAlloc::NewRawData(ul6BufferSize_);
00084      pclNew->Init(pvBuffer, ul6BufferSize_, ul6ElementSize_);
00085      return pclNew;
00086  }
00087  #endif
00088
00089  //-----
00090  void Mailbox::Receive(void* pvData_)
00091  {

```

```

00092     KERNEL_ASSERT(pvData_);
00093
00094 #if KERNEL_USE_TIMEOUTS
00095     Receive_i(pvData_, false, 0);
00096 #else
00097     Receive_i(pvData_, false);
00098 #endif
00099 }
00100
00101 #if KERNEL_USE_TIMEOUTS
00102 //-----
00103 bool Mailbox::Receive(void* pvData_, uint32_t u32TimeoutMS_)
00104 {
00105     KERNEL_ASSERT(pvData_);
00106     return Receive_i(pvData_, false, u32TimeoutMS_);
00107 }
00108 #endif
00109
00110 //-----
00111 void Mailbox::ReceiveTail(void* pvData_)
00112 {
00113     KERNEL_ASSERT(pvData_);
00114
00115 #if KERNEL_USE_TIMEOUTS
00116     Receive_i(pvData_, true, 0);
00117 #else
00118     Receive_i(pvData_, true);
00119 #endif
00120 }
00121
00122 #if KERNEL_USE_TIMEOUTS
00123 //-----
00124 bool Mailbox::ReceiveTail(void* pvData_, uint32_t u32TimeoutMS_)
00125 {
00126     KERNEL_ASSERT(pvData_);
00127     return Receive_i(pvData_, true, u32TimeoutMS_);
00128 }
00129 #endif
00130
00131 //-----
00132 bool Mailbox::Send(void* pvData_)
00133 {
00134     KERNEL_ASSERT(pvData_);
00135
00136 #if KERNEL_USE_TIMEOUTS
00137     return Send_i(pvData_, false, 0);
00138 #else
00139     return Send_i(pvData_, false);
00140 #endif
00141 }
00142
00143 //-----
00144 bool Mailbox::SendTail(void* pvData_)
00145 {
00146     KERNEL_ASSERT(pvData_);
00147
00148 #if KERNEL_USE_TIMEOUTS
00149     return Send_i(pvData_, true, 0);
00150 #else
00151     return Send_i(pvData_, true);
00152 #endif
00153 }
00154
00155 #if KERNEL_USE_TIMEOUTS
00156 //-----
00157 bool Mailbox::Send(void* pvData_, uint32_t u32TimeoutMS_)
00158 {
00159     KERNEL_ASSERT(pvData_);
00160     return Send_i(pvData_, false, u32TimeoutMS_);
00161 }
00162
00163 //-----
00164 bool Mailbox::SendTail(void* pvData_, uint32_t u32TimeoutMS_)
00165 {
00166     KERNEL_ASSERT(pvData_);
00167     return Send_i(pvData_, true, u32TimeoutMS_);
00168 }
00169 #endif
00170
00171 //-----
00172 #if KERNEL_USE_TIMEOUTS
00173 bool Mailbox::Send_i(const void* pvData_, bool bTail_, uint32_t u32TimeoutMS_)
00174 #else
00175 bool Mailbox::Send_i(const void* pvData_, bool bTail_)
00176 #endif
00177 #endif

```

```

00179 {
00180     const void* pvDst = NULL;
00181
00182     auto bRet      = false;
00183     auto bSchedState = Scheduler::SetScheduler(false);
00184
00185 #if KERNEL_USE_TIMEOUTS
00186     auto bBlock = false;
00187     auto bDone  = false;
00188     while (!bDone) {
00189         // Try to claim a slot first before resorting to blocking.
00190         if (bBlock) {
00191             bDone = true;
00192             Scheduler::SetScheduler(bSchedState);
00193             m_clSendSem.Pend(u32TimeoutMS_);
00194             Scheduler::SetScheduler(false);
00195         }
00196 #endif
00197
00198         CS_ENTER();
00199         // Ensure we have a free slot before we attempt to write data
00200         if (m_ul6Free != 0u) {
00201             m_ul6Free--;
00202
00203             if (bTail_) {
00204                 pvDst = GetTailPointer();
00205                 MoveTailBackward();
00206             } else {
00207                 MoveHeadForward();
00208                 pvDst = GetHeadPointer();
00209             }
00210             bRet = true;
00211 #if KERNEL_USE_TIMEOUTS
00212             bDone = true;
00213 #endif
00214         }
00215 #if KERNEL_USE_TIMEOUTS
00216         else if (u32TimeoutMS_ != 0u) {
00217             bBlock = true;
00218         } else {
00219             bDone = true;
00220         }
00221 #endif
00222
00223         CS_EXIT();
00224
00225 #if KERNEL_USE_TIMEOUTS
00226     }
00227 #endif
00228
00229     // Copy data to the claimed slot, and post the counting semaphore
00230     if (bRet) {
00231         CopyData(pvData_, pvDst, m_ul6ElementSize);
00232     }
00233
00234     Scheduler::SetScheduler(bSchedState);
00235
00236     if (bRet) {
00237         m_clRecvSem.Post();
00238     }
00239
00240     return bRet;
00241 }
00242
00243 //-----
00244 #if KERNEL_USE_TIMEOUTS
00245 bool Mailbox::Receive_i(const void* pvData_, bool bTail_, uint32_t u32WaitTimeMS_)
00246 #else
00247 void Mailbox::Receive_i(const void* pvData_, bool bTail_)
00248 #endif
00249 {
00250     const void* pvSrc;
00251
00252 #if KERNEL_USE_TIMEOUTS
00253     if (!m_clRecvSem.Pend(u32WaitTimeMS_)) {
00254         // Failed to get the notification from the counting semaphore in the
00255         // time allotted. Bail.
00256         return false;
00257     }
00258 #else
00259     m_clRecvSem.Pend();
00260 #endif
00261
00262     // Disable the scheduler while we do this -- this ensures we don't have
00263     // multiple concurrent readers off the same queue, which could be problematic
00264     // if multiple writes occur during reads, etc.
00265     auto bSchedState = Scheduler::SetScheduler(false);

```



```

00023 #include "mark3cfg.h"
00024
00025 #include "message.h"
00026 #include "threadport.h"
00027
00028 #define _CAN_HAS_DEBUG
00029 //--[Autogenerated - Do Not Modify]-----
00030 #include "dbg_file_list.h"
00031 #include "buffalogger.h"
00032 #if defined(DBG_FILE)
00033 #error "Debug logging file token already defined! Bailing."
00034 #else
00035 #define DBG_FILE _DBG__KERNEL_MESSAGE_CPP
00036 #endif
00037 //--[End Autogenerated content]-----
00038 #include "kerneldebug.h"
00039
00040 #if KERNEL_USE_MESSAGE
00041
00042 #if KERNEL_USE_TIMEOUTS
00043 #include "timerlist.h"
00044 #endif
00045 namespace Mark3
00046 {
00047 //-----
00048 void MessagePool::Init()
00049 {
00050     m_clList.Init();
00051 }
00052
00053 //-----
00054 void MessagePool::Push(Message* pclMessage_)
00055 {
00056     KERNEL_ASSERT(pclMessage_);
00057
00058     CS_ENTER();
00059
00060     m_clList.Add(pclMessage_);
00061
00062     CS_EXIT();
00063 }
00064
00065 //-----
00066 Message* MessagePool::Pop()
00067 {
00068     Message* pclRet;
00069     CS_ENTER();
00070
00071     pclRet = static_cast<Message*>(m_clList.GetHead());
00072     if (0 != pclRet) {
00073         m_clList.Remove(static_cast<LinkListNode*>(pclRet));
00074     }
00075
00076     CS_EXIT();
00077     return pclRet;
00078 }
00079
00080 //-----
00081 Message* MessagePool::GetHead()
00082 {
00083     return static_cast<Message*>(m_clList.GetHead());
00084 }
00085
00086 //-----
00087 void MessageQueue::Init()
00088 {
00089     m_clSemaphore.Init(0, GLOBAL_MESSAGE_POOL_SIZE);
00090 }
00091
00092 //-----
00093 Message* MessageQueue::Receive()
00094 {
00095     #if KERNEL_USE_TIMEOUTS
00096         return Receive_i(0);
00097     #else
00098         return Receive_i();
00099     #endif
00100 }
00101
00102 //-----
00103 #if KERNEL_USE_TIMEOUTS
00104 Message* MessageQueue::Receive(uint32_t u32TimeWaitMS_)
00105 {
00106     return Receive_i(u32TimeWaitMS_);
00107 }
00108 #endif
00109

```

```

00110 //-----
00111 #if KERNEL_USE_TIMEOUTS
00112 Message* MessageQueue::Receive_i(uint32_t u32TimeWaitMS_)
00113 #else
00114 Message* MessageQueue::Receive_i(void)
00115 #endif
00116 {
00117     Message* pclRet;
00118
00119     // Block the current thread on the counting semaphore
00120     #if KERNEL_USE_TIMEOUTS
00121         if (!m_clSemaphore.Pend(u32TimeWaitMS_)) {
00122             return NULL;
00123         }
00124     #else
00125         m_clSemaphore.Pend();
00126     #endif
00127
00128     CS_ENTER();
00129
00130     // Pop the head of the message queue and return it
00131     pclRet = static_cast<Message*>(m_clLinkList.GetHead());
00132     m_clLinkList.Remove(static_cast<Message*>(pclRet));
00133
00134     CS_EXIT();
00135
00136     return pclRet;
00137 }
00138
00139 //-----
00140 void MessageQueue::Send(Message* pclSrc_)
00141 {
00142     KERNEL_ASSERT(pclSrc_);
00143
00144     CS_ENTER();
00145
00146     // Add the message to the head of the linked list
00147     m_clLinkList.Add(pclSrc_);
00148
00149     // Post the semaphore, waking the blocking thread for the queue.
00150     m_clSemaphore.Post();
00151
00152     CS_EXIT();
00153 }
00154
00155 //-----
00156 uint16_t MessageQueue::GetCount()
00157 {
00158     return m_clSemaphore.GetCount();
00159 }
00160 } //namespace Mark3
00161 #endif // KERNEL_USE_MESSAGE

```

20.47 /home/moslevin/projects/github/m3-repo/kernel/src/mutex.cpp File Reference

Mutual-exclusion object.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "blocking.h"
#include "mutex.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"

```

20.47.1 Detailed Description

Mutual-exclusion object.

Definition in file [mutex.cpp](#).

20.48 mutex.cpp

```

00001 /*-----
00002
00003
00004 |-----|-----|-----|-----|-----|-----|
00005 | \ / | \ / | \ / | \ / | \ / | \ / | \ / |
00006 | / \ | / \ | / \ | / \ | / \ | / \ | / \ |
00007 |-----|-----|-----|-----|-----|-----|
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00020 #include "kerneltypes.h"
00021 #include "mark3cfg.h"
00022
00023 #include "blocking.h"
00024 #include "mutex.h"
00025
00026 #define _CAN_HAS_DEBUG
00027 //--[Autogenerated - Do Not Modify]-----
00028 #include "dbg_file_list.h"
00029 #include "buffallogger.h"
00030 #if defined(DBG_FILE)
00031 #error "Debug logging file token already defined! Bailing."
00032 #else
00033 #define DBG_FILE _DBG__KERNEL_MUTEX_CPP
00034 #endif
00035 //--[End Autogenerated content]-----
00036
00037 #include "kerneldebug.h"
00038
00039 #if KERNEL_USE_MUTEX
00040
00041 #if KERNEL_USE_TIMEOUTS
00042 namespace Mark3
00043 {
00044 namespace
00045 {
00046 //-----
00057 void TimedMutex_Callback(Thread* pclOwner_, void* pvData_)
00058 {
00059     auto* pclMutex = static_cast<Mutex*>(pvData_);
00060
00061     // Indicate that the semaphore has expired on the thread
00062     pclOwner_->SetExpired(true);
00063
00064     // Wake up the thread that was blocked on this semaphore.
00065     pclMutex->WakeMe(pclOwner_);
00066
00067     if (pclOwner_->GetCurPriority() >= Scheduler::GetCurrentThread()->
GetCurPriority()) {
00068         Thread::Yield();
00069     }
00070 }
00071 } // anonymous namespace
00072
00073 //-----
00074 Mutex::~Mutex()
00075 {
00076     // If there are any threads waiting on this object when it goes out
00077     // of scope, set a kernel panic.
00078     if (m_clBlockList.GetHead() != nullptr) {
00079         Kernel::Panic(PANIC_ACTIVE_MUTEX_DESCOPE);
00080     }
00081 }
00082
00083 //-----
00084 void Mutex::WakeMe(Thread* pclOwner_)
00085 {
00086     // Remove from the semaphore waitlist and back to its ready list.
00087     Unblock(pclOwner_);
00088 }
00089
00090 #endif
00091
00092 //-----
00093 uint8_t Mutex::WakeNext()
00094 {
00095     // Get the highest priority waiter thread
00096     auto* pclChosenOne = m_clBlockList.HighestWaiter();
00097
00098     // Unblock the thread
00099     Unblock(pclChosenOne);

```

```

00100
00101     // The chosen one now owns the mutex
00102     m_pclOwner = pclChosenOne;
00103
00104     // Signal a context switch if it's a greater than or equal to the current priority
00105     if (pclChosenOne->GetCurPriority() >= Scheduler::GetCurrentThread()->
GetCurPriority()) {
00106         return 1;
00107     }
00108     return 0;
00109 }
00110
00111 //-----
00112 void Mutex::Init(bool bRecursive_)
00113 {
00114     // Cannot re-init a mutex which has threads blocked on it
00115     #if KERNEL_EXTRA_CHECKS
00116         KERNEL_ASSERT(!m_clBlockList.GetHead());
00117     #endif
00118
00119     // Reset the data in the mutex
00120     m_bReady = true; // The mutex is free.
00121     m_u8MaxPri = 0; // Set the maximum priority inheritance state
00122     m_pclOwner = NULL; // Clear the mutex owner
00123     m_u8Recurse = 0; // Reset recurse count
00124     m_bRecursive = bRecursive_;
00125     #if KERNEL_EXTRA_CHECKS
00126         SetInitialized();
00127     #endif
00128 }
00129
00130 //-----
00131 #if KERNEL_USE_TIMEOUTS
00132 bool Mutex::Claim_i(uint32_t u32WaitTimeMS_)
00133 #else
00134 void Mutex::Claim_i(void)
00135 #endif
00136 {
00137     #if KERNEL_EXTRA_CHECKS
00138         KERNEL_ASSERT(IsInitialized());
00139     #endif
00140
00141     KERNEL_TRACE_1("Claiming Mutex, Thread %d", static_cast<uint16_t>(g_pclCurrent->
GetID()));
00142
00143     #if KERNEL_USE_TIMEOUTS
00144         Timer clTimer;
00145         auto bUseTimer = false;
00146     #endif
00147
00148     // Disable the scheduler while claiming the mutex - we're dealing with all
00149     // sorts of private thread data, can't have a thread switch while messing
00150     // with internal data structures.
00151     Scheduler::SetScheduler(false);
00152
00153     // Check to see if the mutex is claimed or not
00154     if (static_cast<int>(m_bReady) != 0) {
00155         // Mutex isn't claimed, claim it.
00156         m_bReady = false;
00157         m_u8Recurse = 0;
00158         m_u8MaxPri = g_pclCurrent->GetPriority();
00159         m_pclOwner = g_pclCurrent;
00160
00161         Scheduler::SetScheduler(true);
00162
00163     #if KERNEL_USE_TIMEOUTS
00164         return true;
00165     #else
00166         return;
00167     #endif
00168     }
00169
00170     // If the mutex is already claimed, check to see if this is the owner thread,
00171     // since we allow the mutex to be claimed recursively.
00172     if (m_bRecursive && (g_pclCurrent == m_pclOwner)) {
00173         // Ensure that we haven't exceeded the maximum recursive-lock count
00174         KERNEL_ASSERT((m_u8Recurse < 255));
00175         m_u8Recurse++;
00176
00177         // Increment the lock count and bail
00178         Scheduler::SetScheduler(true);
00179     #if KERNEL_USE_TIMEOUTS
00180         return true;
00181     #else
00182         return;
00183     #endif
00184 }

```

```

00185
00186 // The mutex is claimed already - we have to block now. Move the
00187 // current thread to the list of threads waiting on the mutex.
00188 #if KERNEL_USE_TIMEOUTS
00189     if (u32WaitTimeMS_ != 0u) {
00190         g_pclCurrent->SetExpired(false);
00191         clTimer.Init();
00192         clTimer.Start(false, u32WaitTimeMS_, (TimerCallback)TimedMutex_Callback, this);
00193         bUseTimer = true;
00194     }
00195 #endif
00196 BlockPriority(g_pclCurrent);
00197
00198 // Check if priority inheritance is necessary. We do this in order
00199 // to ensure that we don't end up with priority inversions in case
00200 // multiple threads are waiting on the same resource.
00201 if (m_u8MaxPri <= g_pclCurrent->GetPriority()) {
00202     m_u8MaxPri = g_pclCurrent->GetPriority();
00203
00204     auto* pclTemp = static_cast<Thread*>(m_clBlockList.GetHead());
00205     while (pclTemp != nullptr) {
00206         pclTemp->InheritPriority(m_u8MaxPri);
00207         if (pclTemp == static_cast<Thread*>(m_clBlockList.
GetTail())) {
00208             break;
00209         }
00210         pclTemp = static_cast<Thread*>(pclTemp->GetNext());
00211     }
00212     m_pclOwner->InheritPriority(m_u8MaxPri);
00213 }
00214
00215 // Done with thread data -reenable the scheduler
00216 Scheduler::SetScheduler(true);
00217
00218 // Switch threads if this thread acquired the mutex
00219 Thread::Yield();
00220
00221 #if KERNEL_USE_TIMEOUTS
00222     if (bUseTimer) {
00223         clTimer.Stop();
00224         return (static_cast<int>(g_pclCurrent->GetExpired()) == 0);
00225     }
00226     return true;
00227 #endif
00228 }
00229
00230 //-----
00231 void Mutex::Claim(void)
00232 {
00233     #if KERNEL_USE_TIMEOUTS
00234         Claim_i(0);
00235     #else
00236         Claim_i();
00237     #endif
00238 }
00239
00240 //-----
00241 #if KERNEL_USE_TIMEOUTS
00242 bool Mutex::Claim(uint32_t u32WaitTimeMS_)
00243 {
00244     return Claim_i(u32WaitTimeMS_);
00245 }
00246 #endif
00247
00248 //-----
00249 void Mutex::Release()
00250 {
00251     #if KERNEL_EXTRA_CHECKS
00252         KERNEL_ASSERT(IsInitialized());
00253     #endif
00254
00255     KERNEL_TRACE_1("Releasing Mutex, Thread %d", static_cast<uint16_t>(g_pclCurrent->
GetID()));
00256
00257     auto bSchedule = false;
00258
00259     // Disable the scheduler while we deal with internal data structures.
00260     Scheduler::SetScheduler(false);
00261
00262     // This thread had better be the one that owns the mutex currently...
00263     KERNEL_ASSERT((g_pclCurrent == m_pclOwner));
00264
00265     // If the owner had claimed the lock multiple times, decrease the lock
00266     // count and return immediately.
00267     if (m_bRecursive && (m_u8Recurse != 0u)) {
00268         m_u8Recurse--;
00269         Scheduler::SetScheduler(true);

```



```

00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00022 #include "mark3cfg.h"
00023 #include "notify.h"
00024 #include "mark3.h"
00025 #include "kerneldebug.h"
00026
00027 #define _CAN_HAS_DEBUG
00028 //--[Autogenerated - Do Not Modify]-----
00029 #include "dbg_file_list.h"
00030 #include "buffalogger.h"
00031 #if defined(DBG_FILE)
00032 #error "Debug logging file token already defined! Bailing."
00033 #else
00034 #define DBG_FILE _DBG__KERNEL_NOTIFY_CPP
00035 #endif
00036 //--[End Autogenerated content]-----
00037
00038 #if KERNEL_USE_NOTIFY
00039
00040 #if KERNEL_USE_TIMEOUTS
00041 namespace Mark3
00042 {
00043 namespace
00044 {
00045 //-----
00046 void TimedNotify_Callback(Thread* pClOwner_, void* pvData_)
00047 {
00048     auto* pClNotify = static_cast<Notify*>(pvData_);
00049
00050     // Indicate that the semaphore has expired on the thread
00051     pClOwner_->SetExpired(true);
00052
00053     // Wake up the thread that was blocked on this semaphore.
00054     pClNotify->WakeMe(pClOwner_);
00055
00056     if (pClOwner_->GetCurPriority() >= Scheduler::GetCurrentThread()->
GetCurPriority()) {
00057         Thread::Yield();
00058     }
00059 }
00060 } // anonymous namespace
00061 #endif
00062 //-----
00063 Notify::~Notify()
00064 {
00065     // If there are any threads waiting on this object when it goes out
00066     // of scope, set a kernel panic.
00067     if (m_clBlockList.GetHead() != nullptr) {
00068         Kernel::Panic(PANIC_ACTIVE_NOTIFY_DESCOPE);
00069     }
00070 }
00071
00072 //-----
00073 void Notify::Init(void)
00074 {
00075 #if KERNEL_EXTRA_CHECKS
00076     KERNEL_ASSERT(!m_clBlockList.GetHead());
00077     SetInitialized();
00078 #endif
00079     m_bPending = false;
00080 }
00081
00082 //-----
00083 void Notify::Signal(void)
00084 {
00085 #if KERNEL_EXTRA_CHECKS
00086     KERNEL_ASSERT(IsInitialized());
00087 #endif
00088
00089     auto bReschedule = false;
00090
00091     CS_ENTER();
00092     auto* pClCurrent = static_cast<Thread*>(m_clBlockList.GetHead());
00093     if (pClCurrent == nullptr) {
00094         m_bPending = true;
00095     } else {
00096         while (pClCurrent != NULL) {
00097             Unblock(pClCurrent);
00098             if (!bReschedule && (pClCurrent->GetCurPriority() >=
Scheduler::GetCurrentThread()->GetCurPriority())) {
00099                 bReschedule = true;
00100             }
00101             pClCurrent = (Thread*)m_clBlockList.GetHead();
00102         }
00103         m_bPending = false;

```

```

00104     }
00105     CS_EXIT();
00106
00107     if (bReschedule) {
00108         Thread::Yield();
00109     }
00110 }
00111
00112 //-----
00113 void Notify::Wait(bool* pbFlag_)
00114 {
00115     #if KERNEL_EXTRA_CHECKS
00116         KERNEL_ASSERT(IsInitialized());
00117     #endif
00118
00119     auto bEarlyExit = false;
00120     CS_ENTER();
00121     if (!m_bPending) {
00122         Block(g_pclCurrent);
00123         if (pbFlag_ != nullptr) {
00124             *pbFlag_ = false;
00125         }
00126     } else {
00127         m_bPending = false;
00128         bEarlyExit = true;
00129     }
00130     CS_EXIT();
00131
00132     if (bEarlyExit) {
00133         return;
00134     }
00135
00136     Thread::Yield();
00137     if (pbFlag_ != nullptr) {
00138         *pbFlag_ = true;
00139     }
00140 }
00141
00142 //-----
00143 #if KERNEL_USE_TIMEOUTS
00144 bool Notify::Wait(uint32_t u32WaitTimeMS_, bool* pbFlag_)
00145 {
00146     #if KERNEL_EXTRA_CHECKS
00147         KERNEL_ASSERT(IsInitialized());
00148     #endif
00149
00150     auto bUseTimer = false;
00151     auto bEarlyExit = false;
00152     Timer clNotifyTimer;
00153
00154     CS_ENTER();
00155     if (!m_bPending) {
00156         if (u32WaitTimeMS_ != 0u) {
00157             bUseTimer = true;
00158             g_pclCurrent->SetExpired(false);
00159
00160             clNotifyTimer.Init();
00161             clNotifyTimer.Start(false, u32WaitTimeMS_, TimedNotify_Callback, this);
00162         }
00163
00164         Block(g_pclCurrent);
00165
00166         if (pbFlag_ != nullptr) {
00167             *pbFlag_ = false;
00168         }
00169     } else {
00170         m_bPending = false;
00171         bEarlyExit = true;
00172     }
00173     CS_EXIT();
00174
00175     if (bEarlyExit) {
00176         return true;
00177     }
00178
00179     Thread::Yield();
00180
00181     if (bUseTimer) {
00182         clNotifyTimer.Stop();
00183         return (static_cast<int>(g_pclCurrent->GetExpired()) == 0);
00184     }
00185
00186     if (pbFlag_ != nullptr) {
00187         *pbFlag_ = true;
00188     }
00189
00190     return true;
00191 }

```



```
00191 #endif
00192 //-----
00193 void Notify::WakeMe(Thread* pClChosenOne_)
00194 {
00195     if KERNEL_EXTRA_CHECKS
00196         KERNEL_ASSERT(IsInitialized());
00197     #endif
00198     Unblock(pClChosenOne_);
00199 }
00200 } //namespace Mark3
00201 #endif
```

20.51 /home/moslevin/projects/github/m3-repo/kernel/src/priomap.cpp File Reference

Priority map data structure.

```
#include "mark3.h"
#include "priomap.h"
#include "threadport.h"
#include <stdint.h>
#include <stdbool.h>
```

Namespaces

- Mark3

Class providing the software-interrupt required for context-switching in the kernel.

20.51.1 Detailed Description

Priority map data structure.

Definition in file [priomap.cpp](#).

20.52 priomap.cpp

```

00001 /*-----
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*
00014 #include "mark3.h"
00015 #include "priomap.h"
00016 #include "threadport.h"
00017
00018 #include <stdint.h>
00019 #include <stdbool.h>
00020 namespace Mark3
00021 {
00022 //-----
00023 namespace {
00024 inline uint8_t priority_from_bitmap(PORT_PRIO_TYPE uXPrio_)
00025 {

```

```

00031 #if HW_CLZ
00032 // Support hardware-accelerated Count-leading-zeros instruction
00033 uint8_t rc = PRIO_MAP_BITS - CLZ(uXPrio_);
00034 return rc;
00035 #else
00036 // Default un-optimized count-leading zeros operation
00037 PORT_PRIO_TYPE uXMask = (1 << (PRIO_MAP_BITS - 1));
00038 uint8_t u8Zeros = 0;
00039
00040 while (uXMask) {
00041     if (uXMask & uXPrio_) {
00042         return (PRIO_MAP_BITS - u8Zeros);
00043     }
00044     uXMask >>= 1;
00045     u8Zeros++;
00046 }
00047 return 0;
00048 #endif
00049 #endif
00050 }
00051 } // anonymous namespace
00052
00053 //-----
00054 PriorityMap::PriorityMap()
00055 {
00056     #if PRIO_MAP_MULTI_LEVEL
00057         m_uXPriorityMapL2 = 0;
00058         for (int i = 0; i < PRIO_MAP_NUM_WORDS; i++) {
00059             m_auXPriorityMap[i] = 0;
00060         }
00061     #else
00062         m_uXPriorityMap = 0;
00063     #endif
00064 }
00065
00066 //-----
00067 void PriorityMap::Set(PORT_PRIO_TYPE uXPrio_)
00068 {
00069     PORT_PRIO_TYPE uXPrioBit = PRIO_BIT(uXPrio_);
00070     #if PRIO_MAP_MULTI_LEVEL
00071     PORT_PRIO_TYPE uXWordIdx = PRIO_MAP_WORD_INDEX(uXPrio_);
00072     m_auXPriorityMap[uXWordIdx] |= (1 << uXPrioBit);
00073     m_uXPriorityMapL2 |= (1 << uXWordIdx);
00074     #else
00075     m_uXPriorityMap |= (1 << uXPrioBit);
00076     #endif
00077 }
00078
00079 //-----
00080 void PriorityMap::Clear(PORT_PRIO_TYPE uXPrio_)
00081 {
00082     PORT_PRIO_TYPE uXPrioBit = PRIO_BIT(uXPrio_);
00083     #if PRIO_MAP_MULTI_LEVEL
00084     PORT_PRIO_TYPE uXWordIdx = PRIO_MAP_WORD_INDEX(uXPrio_);
00085     m_auXPriorityMap[uXWordIdx] &= ~(1 << uXPrioBit);
00086     if (!m_auXPriorityMap[uXWordIdx]) {
00087         m_uXPriorityMapL2 &= ~(1 << uXWordIdx);
00088     }
00089     #else
00090     m_uXPriorityMap &= ~(1 << uXPrioBit);
00091     #endif
00092 }
00093
00094 //-----
00095 PORT_PRIO_TYPE PriorityMap::HighestPriority(void)
00096 {
00097     #if PRIO_MAP_MULTI_LEVEL
00098     PORT_PRIO_TYPE uXMapIdx = priority_from_bitmap(m_uXPriorityMapL2);
00099     if (!uXMapIdx) {
00100         return 0;
00101     }
00102     uXMapIdx--;
00103     PORT_PRIO_TYPE uXPrio = priority_from_bitmap(m_auXPriorityMap[uXMapIdx]);
00104     uXPrio += (uXMapIdx * PRIO_MAP_BITS);
00105     #else
00106     PORT_PRIO_TYPE uXPrio = priority_from_bitmap(m_uXPriorityMap);
00107     #endif
00108     return uXPrio;
00109 }
00110 } //namespace Mark3

```

20.53 /home/moslevin/projects/github/m3-repo/kernel/src/profile.cpp File Reference

Code profiling utilities.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "profile.h"
#include "kernelprofile.h"
#include "threadport.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"
```

20.53.1 Detailed Description

Code profiling utilities.

Definition in file [profile.cpp](#).

20.54 profile.cpp

```
00001 /*=====
00002
00003
00004 | | | | | | | | | | | | | | | | | |
00005 | | | | | | | | | | | | | | | | | |
00006 | | | | | | | | | | | | | | | | | |
00007 | | | | | | | | | | | | | | | | | |
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00021 #include "kerneltypes.h"
00022 #include "mark3cfg.h"
00023 #include "profile.h"
00024 #include "kernelprofile.h"
00025 #include "threadport.h"
00026
00027 #define _CAN_HAS_DEBUG
00028 //--[Autogenerated - Do Not Modify]-----
00029 #include "dbg_file_list.h"
00030 #include "buffalogger.h"
00031 #if defined(DBG_FILE)
00032 #error "Debug logging file token already defined! Bailing."
00033 #else
00034 #define DBG_FILE _DBG__KERNEL_PROFILE_CPP
00035 #endif
00036 //--[End Autogenerated content]-----
00037
00038 #include "kerneldebug.h"
00039
00040 #if KERNEL_USE_PROFILER
00041 namespace Mark3
00042 {
00043     //--
00044     void ProfileTimer::Init()
00045     {
00046         m_u32Cumulative      = 0;
00047         m_u32CurrentIteration = 0;
00048         m_u16Iterations      = 0;
00049         m_bActive            = false;
00050     }
00051
00052     //--
00053     void ProfileTimer::Start()
00054     {
```

```

00055     if (!m_bActive) {
00056         CS_ENTER();
00057         m_u32CurrentIteration = 0;
00058         m_u32InitialEpoch    = Profiler::GetEpoch();
00059         m_ul6Initial          = Profiler::Read();
00060         CS_EXIT();
00061         m_bActive = true;
00062     }
00063 }
00064
00065 //-----
00066 void ProfileTimer::Stop()
00067 {
00068     if (m_bActive) {
00069         uint16_t u16Final;
00070         uint32_t u32Epoch;
00071         CS_ENTER();
00072         u16Final = Profiler::Read();
00073         u32Epoch = Profiler::GetEpoch();
00074         // Compute total for current iteration...
00075         m_u32CurrentIteration = ComputeCurrentTicks(u16Final,
u32Epoch);
00076         m_u32Cumulative += m_u32CurrentIteration;
00077         m_ul6Iterations++;
00078         CS_EXIT();
00079         m_bActive = false;
00080     }
00081 }
00082
00083 //-----
00084 uint32_t ProfileTimer::GetAverage()
00085 {
00086     if (m_ul6Iterations != 0u) {
00087         return m_u32Cumulative / (uint32_t)m_ul6Iterations;
00088     }
00089     return 0;
00090 }
00091
00092 //-----
00093 uint32_t ProfileTimer::GetCurrent()
00094 {
00095     if (m_bActive) {
00096         uint16_t u16Current;
00097         uint32_t u32Epoch;
00098         CS_ENTER();
00099         u16Current = Profiler::Read();
00100         u32Epoch = Profiler::GetEpoch();
00101         CS_EXIT();
00102         return ComputeCurrentTicks(u16Current, u32Epoch);
00103     }
00104     return m_u32CurrentIteration;
00105 }
00106
00107 //-----
00108 uint32_t ProfileTimer::ComputeCurrentTicks(uint16_t u16Current_, uint32_t
u32Epoch_)
00109 {
00110     uint32_t u32Total;
00111     uint32_t u32Overflows;
00112
00113     u32Overflows = u32Epoch_ - m_u32InitialEpoch;
00114
00115     // More than one overflow...
00116     if (u32Overflows > 1) {
00117         u32Total = ((uint32_t)(u32Overflows - 1) * TICKS_PER_OVERFLOW) + (uint32_t)(TICKS_PER_OVERFLOW -
m_ul6Initial)
00118             + (uint32_t)u16Current_;
00119     }
00120     // Only one overflow, or one overflow that has yet to be processed
00121     else if ((u32Overflows != 0u) || (u16Current_ < m_ul6Initial)) {
00122         u32Total = (uint32_t)(TICKS_PER_OVERFLOW - m_ul6Initial) + (uint32_t)u16Current_;
00123     }
00124     // No overflows, none pending.
00125     else {
00126         u32Total = (uint32_t)(u16Current_ - m_ul6Initial);
00127     }
00128
00129     return u32Total;
00130 }
00131 } //namespace Mark3
00132 #endif

```



```

00029 {
00039 namespace Atomic
00040 {
00047     uint8_t Set(uint8_t* pu8Source_, uint8_t u8Val_);
00048     uint16_t Set(uint16_t* pu16Source_, uint16_t u16Val_);
00049     uint32_t Set(uint32_t* pu32Source_, uint32_t u32Val_);
00050
00057     uint8_t Add(uint8_t* pu8Source_, uint8_t u8Val_);
00058     uint16_t Add(uint16_t* pu16Source_, uint16_t u16Val_);
00059     uint32_t Add(uint32_t* pu32Source_, uint32_t u32Val_);
00060
00067     uint8_t Sub(uint8_t* pu8Source_, uint8_t u8Val_);
00068     uint16_t Sub(uint16_t* pu16Source_, uint16_t u16Val_);
00069     uint32_t Sub(uint32_t* pu32Source_, uint32_t u32Val_);
00070
00085     bool TestAndSet(bool* pbLock);
00086 } // namespace Atomic
00087 //namespace Mark3
00088 #endif // KERNEL_USE_ATOMIC

```

20.57 /home/moslevin/projects/github/m3-repo/kernel/src/public/autoalloc.h File Reference

Automatic memory allocation for kernel objects.

```
#include <stdint.h>
#include <stdbool.h>
#include "mark3cfg.h"
```

Namespaces

- Mark3

Class providing the software-interrupt required for context-switching in the kernel.

20.57.1 Detailed Description

Automatic memory allocation for kernel objects.

Definition in file [autoalloc.h](#).

20.58 autoalloc.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00020 #pragma once
00021
00022 #include <stdint.h>
00023 #include <stdbool.h>
00024 #include "mark3cfg.h"
```

```

00025
00026 #if KERNEL_USE_AUTO_ALLOC
00027 namespace Mark3
00028 {
00029
00030 //-----
00031 // Define function pointer types used for interfacing with an external heap.
00032
00033 //-----
00034 enum class AutoAllocType : uint8_t {
00035     //-- Kernel object types
00036     EventFlag,
00037     MailBox,
00038     Message,
00039     MessagePool,
00040     MessageQueue,
00041     Mutex,
00042     Notify,
00043     Semaphore,
00044     Thread,
00045     Timer,
00046     ConditionVariable,
00047     ReaderWriterLock,
00048     //-- Allow for users to define their own object types beginning with AutoAllocType_t::User
00049     User,
00050     //--
00051     Raw = 0xFF
00052 };
00053
00054 //-----
00055 using AutoAllocAllocator_t = void* (*)(AutoAllocType eType_, size_t sSize_);
00056 using AutoAllocFree_t = void* (*)(AutoAllocType eType_, void* pObj_);
00057
00058 //-----
00059 // Forward declaration of kernel objects that can be automatically allocated.
00060 #if KERNEL_USE_EVENTFLAG
00061 class EventFlag;
00062 #endif
00063
00064 #if KERNEL_USE_MAILBOX
00065 class Mailbox;
00066 #endif
00067
00068 #if KERNEL_USE_MESSAGE
00069 class Message;
00070 class MessagePool;
00071 class MessageQueue;
00072 #endif
00073
00074 #if KERNEL_USE_MUTEX
00075 class Mutex;
00076 #endif
00077
00078 #if KERNEL_USE_NOTIFY
00079 class Notify;
00080 #endif
00081
00082 #if KERNEL_USE_SEMAPHORE
00083 class Semaphore;
00084 #endif
00085
00086 class Thread;
00087
00088 #if KERNEL_USE_TIMERS
00089 class Timer;
00090 #endif
00091
00092 #if KERNEL_USE_READERWRITER
00093 class ReaderWriterLock;
00094 #endif
00095
00096 #if KERNEL_USE_CONDVAR
00097 class ConditionVariable;
00098 #endif
00099
00100 class AutoAlloc
00101 {
00102 public:
00103
00110     static void Init(void);
00111
00118     static void SetAllocatorFunctions(AutoAllocAllocator_t pfAllocator_, AutoAllocFree_t pfFree_)
00119     { m_pfAllocator = pfAllocator_; m_pfFree = pfFree_; }
00120
00121 #if KERNEL_USE_SEMAPHORE
00122
00127     static Semaphore* NewSemaphore();

```

```

00128
00134     static void DestroySemaphore(Semaphore* pclSemaphore_);
00135 #endif
00136 #if KERNEL_USE_MUTEX
00137
00142     static Mutex* NewMutex();
00143
00149     static void DestroyMutex(Mutex* pclMutex_);
00150 #endif
00151 #if KERNEL_USE_EVENTFLAG
00152
00157     static EventFlag* NewEventFlag();
00158
00164     static void DestroyEventFlag(EventFlag* pclEventFlag_);
00165 #endif
00166 #if KERNEL_USE_TIMERS
00167
00172     static Timer* NewTimer();
00173
00179     static void DestroyTimer(Timer* pclTimer_);
00180 #endif
00181 #if KERNEL_USE_NOTIFY
00182
00187     static Notify* NewNotify();
00188
00194     static void DestroyNotify(Notify* pclNotify_);
00195 #endif
00196 #if KERNEL_USE_MAILBOX
00197
00202     static Mailbox* NewMailbox();
00203
00209     static void DestroyMailbox(Mailbox* pclMailbox_);
00210 #endif
00211 #if KERNEL_USE_MESSAGE
00212
00217     static Message* NewMessage();
00218
00224     static MessagePool* NewMessagePool();
00225
00231     static MessageQueue* NewMessageQueue();
00232
00238     static void DestroyMessage(Message* pclMessage_);
00239
00245     static void DestroyMessagePool(MessagePool *pclMessagePool_);
00246
00252     static void DestroyMessageQueue(MessageQueue* pclMessageQ_);
00253 #endif
00254 #if KERNEL_USE_CONDVAR
00255
00260     static ConditionVariable* NewConditionVariable();
00261
00267     static void DestroyConditionVariable(ConditionVariable* pclConditionVariable_);
00268 #endif
00269 #if KERNEL_USE_READERWRITER
00270
00275     static ReaderWriterLock* NewReaderWriterLock();
00276
00282     static void DestroyReaderWriterLock(ReaderWriterLock* pclReaderWriterLock_);
00283 #endif
00284
00289     static Thread* NewThread();
00290
00296     static void DestroyThread(Thread* pclThread_);
00297
00304     static void* NewUserTypeAllocation(uint8_t eUserType_);
00305
00312     static void DestroyUserTypeAllocation(uint8_t eUserType_, void* pvObj_);
00313
00320     static void* NewRawData(size_t sSize_);
00321
00327     static void DestroyRawData(void* pvData_);
00328
00329 private:
00330
00331     static void* Allocate(AutoAllocType eType_, size_t sSize_);
00332     static void Free(AutoAllocType eType_, void* pvObj_);
00333
00334     static AutoAllocAllocator_t m_pfAllocator;
00335     static AutoAllocFree_t      m_pfFree;
00336 };
00337 } //namespace Mark3
00338 #endif

```


20.59 /home/moslevin/projects/github/m3-repo/kernel/src/public/blocking.h File Reference

Blocking object base class declarations.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ll.h"
#include "threadlist.h"
```

Classes

- class [Mark3::BlockingObject](#)
Class implementing thread-blocking primitives.

Namespaces

- [Mark3](#)
Class providing the software-interrupt required for context-switching in the kernel.

20.59.1 Detailed Description

Blocking object base class declarations.

A Blocking object in [Mark3](#) is essentially a thread list. Any blocking object implementation (being a semaphore, mutex, event flag, etc.) can be built on top of this class, utilizing the provided functions to manipulate thread location within the Kernel.

Blocking a thread results in that thread becoming de-scheduled, placed in the blocking object's own private list of threads which are waiting on the object.

Unblocking a thread results in the reverse: The thread is moved back to its original location from the blocking list.

The only difference between a blocking object based on this class is the logic used to determine what constitutes a Block or Unblock condition.

For instance, a semaphore Pend operation may result in a call to the Block() method with the currently-executing thread in order to make that thread wait for a semaphore Post. That operation would then invoke the Unblock() method, removing the blocking thread from the semaphore's list, and back into the appropriate thread inside the scheduler.

Care must be taken when implementing blocking objects to ensure that critical sections are used judiciously, otherwise asynchronous events like timers and interrupts could result in non-deterministic and often catastrophic behavior.

Definition in file [blocking.h](#).

20.64 condvar.h

```

00001 #pragma once
00002
00003 /*=====
00004
00005
00006
00007
00008
00009
00010
00011 --[Mark3 Realtime Platform]-----
00012
00013 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00014 See license.txt for more information
00015 =====*/
00022 #include "mark3cfg.h"
00023 #include "ksemaphore.h"
00024 #include "mutex.h"
00025
00026 namespace Mark3 {
00027
00038 class ConditionVariable {
00039 public:
00040     void* operator new(size_t sz, void*pv) { return (ConditionVariable*)pv; }
00041
00047     void Init();
00048
00055     void Wait(Mutex* pclMutex_);
00056
00057 #if KERNEL_USE_TIMEOUTS
00066     bool Wait(Mutex* pclMutex_, uint32_t u32WaitTimeMS_);
00067 #endif
00073     void Signal();
00074
00079     void Broadcast();
00080
00081 private:
00082     Mutex m_clMutex;
00083     Semaphore m_clSemaphore;
00084     uint8_t m_u8Waiters;
00085 };
00086
00087
00088 } // namespace Mark3

```

20.65 /home/moslevin/projects/github/m3-repo/kernel/src/public/eventflag.h File Reference

Event Flag Blocking Object/IPC-Object definition.

```

#include "mark3cfg.h"
#include "kernel.h"
#include "kerneltypes.h"
#include "blocking.h"
#include "thread.h"

```

Classes

- class [Mark3::EventFlag](#)

The [EventFlag](#) class is a blocking object, similar to a semaphore or mutex, commonly used for synchronizing thread execution based on events occurring within the system.


```

00040 {
00041 //-----
00045 class Kernel
00046 {
00047 public:
00048
00057     static void Init(void);
00058
00071     static void Start(void);
00072
00079     static bool IsStarted() { return m_bIsStarted; }
00087     static void SetPanic(PanicFunc pfPanic_) { m_pfPanic = pfPanic_; }
00092     static bool IsPanic() { return m_bIsPanic; }
00097     static void Panic(uint16_t u16Cause_);
00098
00099 #if KERNEL_USE_IDLE_FUNC
00100
00105     static void SetIdleFunc(IdleFunc pfIdle_) { m_pfIdle = pfIdle_; }
00110     static void Idle(void)
00111     {
00112         if (m_pfIdle != 0) {
00113             m_pfIdle();
00114         }
00115     }
00116
00124     static Thread* GetIdleThread(void) { return (Thread*)&
m_clIdle; }
00125 #endif
00126
00127 #if KERNEL_USE_THREAD_CALLOUTS
00128
00138     static void SetThreadCreateCallout(ThreadCreateCallout pfCreate_) {
m_pfThreadCreateCallout = pfCreate_; }
00150     static void SetThreadExitCallout(ThreadExitCallout pfExit_) {
m_pfThreadExitCallout = pfExit_; }
00161     static void SetThreadContextSwitchCallout(ThreadContextCallout pfContext_)
00162     {
00163         m_pfThreadContextCallout = pfContext_;
00164     }
00165
00174     static ThreadCreateCallout GetThreadCreateCallout(void) { return
m_pfThreadCreateCallout; }
00183     static ThreadExitCallout GetThreadExitCallout(void) { return
m_pfThreadExitCallout; }
00192     static ThreadContextCallout GetThreadContextSwitchCallout(void) { return
m_pfThreadContextCallout; }
00193 #endif
00194
00195 #if KERNEL_USE_STACK_GUARD
00196     static void SetStackGuardThreshold(uint16_t u16Threshold_) { m_u16GuardThreshold = u16Threshold_; }
00197     static uint16_t
GetStackGuardThreshold(void) { return m_u16GuardThreshold; }
00198 #endif
00199
00200 private:
00201     static bool m_bIsStarted;
00202     static bool m_bIsPanic;
00203     static PanicFunc m_pfPanic;
00204 #if KERNEL_USE_IDLE_FUNC
00205     static IdleFunc m_pfIdle;
00206     static FakeThread_t m_clIdle;
00207 #endif
00208
00209 #if KERNEL_USE_THREAD_CALLOUTS
00210     static ThreadCreateCallout m_pfThreadCreateCallout;
00211     static ThreadExitCallout m_pfThreadExitCallout;
00212     static ThreadContextCallout m_pfThreadContextCallout;
00213 #endif
00214
00215 #if KERNEL_USE_STACK_GUARD
00216     static uint16_t m_u16GuardThreshold;
00217 #endif
00218 };
00219
00220 } //namespace Mark3

```

20.69 /home/moslevin/projects/github/m3-repo/kernel/src/public/kernelaware.h File Reference

Kernel aware simulation support.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
```

Namespaces

- [Mark3](#)

Class providing the software-interrupt required for context-switching in the kernel.

- [Mark3::KernelAware](#)

The [KernelAware](#) class.

Functions

- void [Mark3::KernelAware::ProfileInit](#) (const char *szStr_)

ProfileInit.

- void [Mark3::KernelAware::ProfileStart](#) (void)

ProfileStart.

- void [Mark3::KernelAware::ProfileStop](#) (void)

ProfileStop.

- void [Mark3::KernelAware::ProfileReport](#) (void)

ProfileReport.

- void [Mark3::KernelAware::ExitSimulator](#) (void)

ExitSimulator.

- void [Mark3::KernelAware::Print](#) (const char *szStr_)

Print.

- void [Mark3::KernelAware::Trace](#) (uint16_t u16File_, uint16_t u16Line_)

Trace.

- void [Mark3::KernelAware::Trace](#) (uint16_t u16File_, uint16_t u16Line_, uint16_t u16Arg1_)

Trace.

- void [Mark3::KernelAware::Trace](#) (uint16_t u16File_, uint16_t u16Line_, uint16_t u16Arg1_, uint16_t u16Arg2_)

Trace.

- bool [Mark3::KernelAware::IsSimulatorAware](#) (void)

IsSimulatorAware.

20.69.1 Detailed Description

Kernel aware simulation support.

Definition in file [kernelaware.h](#).

20.70 kernelaware.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00020 #pragma once
00021
00022 #include "kerneltypes.h"
00023 #include "mark3cfg.h"
00024
00025 #if KERNEL_AWARE_SIMULATION
00026 namespace Mark3
00027 {
00028 //-----
00046 namespace KernelAware
00047 {
00048 //-----
00059 void ProfileInit(const char* szStr_);
00060
00061 //-----
00069 void ProfileStart(void);
00070
00071 //-----
00078 void ProfileStop(void);
00079
00080 //-----
00088 void ProfileReport(void);
00089
00090 //-----
00098 void ExitSimulator(void);
00099
00100 //-----
00108 void Print(const char* szStr_);
00109
00110 //-----
00120 void Trace(uint16_t ul6File_, uint16_t ul6Line_);
00121
00122 //-----
00133 void Trace(uint16_t ul6File_, uint16_t ul6Line_, uint16_t ul6Arg1_);
00134
00135 //-----
00147 void Trace(uint16_t ul6File_, uint16_t ul6Line_, uint16_t ul6Arg1_, uint16_t ul6Arg2_);
00148
00149 //-----
00159 bool IsSimulatorAware(void);
00160 } // namespace KernelAware
00161 } //namespace Mark3
00162 #endif

```

20.71 /home/moslevin/projects/github/m3-repo/kernel/src/public/kerneldebug.h File Reference

Macros and functions used for assertions, kernel traces, etc.

```

#include "mark3cfg.h"
#include "tracebuffer.h"
#include "kernelaware.h"
#include "paniccodes.h"
#include "kernel.h"
#include "buffalogger.h"
#include "dbg_file_list.h"

```



```

00033 //-----
00034 #define KERNEL_TRACE(x)
00035     \
00036 {
00037     \
00038     EMIT_DBG_STRING(x);
00039     \
00040     uint16_t aul6Msg__[4];
00041     \
00042     aul6Msg__[0] = 0xACDC;
00043     \
00044     aul6Msg__[1] = DBG_FILE;
00045     \
00046     aul6Msg__[2] = __LINE__;
00047     \
00048     aul6Msg__[3] = TraceBuffer::Increment();
00049     \
00050     TraceBuffer::Write(aul6Msg__, 4);
00051 }
00052 //-----
00053 #define KERNEL_TRACE_1(x, arg1)
00054     \
00055 {
00056     \
00057     EMIT_DBG_STRING(x);
00058     \
00059     uint16_t aul6Msg__[5];
00060     \
00061     aul6Msg__[0] = 0xACDC;
00062     \
00063     aul6Msg__[1] = DBG_FILE;
00064     \
00065     aul6Msg__[2] = __LINE__;
00066     \
00067     aul6Msg__[3] = TraceBuffer::Increment();
00068     \
00069     aul6Msg__[4] = arg1;
00070     \
00071     TraceBuffer::Write(aul6Msg__, 5);
00072 }
00073 //-----
00074 #define KERNEL_TRACE_2(x, arg1, arg2)
00075     \
00076 {
00077     \
00078     EMIT_DBG_STRING(x);
00079     \
00080     uint16_t aul6Msg__[6];
00081     \
00082     aul6Msg__[0] = 0xACDC;
00083     \
00084     aul6Msg__[1] = DBG_FILE;
00085     \
00086     aul6Msg__[2] = __LINE__;
00087     \
00088     aul6Msg__[3] = TraceBuffer::Increment();
00089     \
00090     aul6Msg__[4] = arg1;
00091     \
00092     aul6Msg__[5] = arg2;
00093     \
00094     TraceBuffer::Write(aul6Msg__, 6);
00095 }
00096 //-----
00097 #define KERNEL_ASSERT(x)
00098     \
00099 {
00100     \
00101     if ((x) == false) {
00102         \
00103         EMIT_DBG_STRING("ASSERT FAILED");
00104         \
00105         uint16_t aul6Msg__[4];
00106         \
00107         aul6Msg__[0] = 0xACDC;
00108         \
00109         aul6Msg__[1] = DBG_FILE;
00110         \
00111         aul6Msg__[2] = __LINE__;
00112         \
00113         aul6Msg__[3] = TraceBuffer::Increment();
00114         \
00115         TraceBuffer::Write(aul6Msg__, 4);
00116     }
00117 }

```

```

00085         aul6Msg__[0] = 0xACDC;
00086         \
00087         aul6Msg__[1] = DBG_FILE;
00088         \
00089         aul6Msg__[2] = __LINE__;
00090         \
00091         aul6Msg__[3] = TraceBuffer::Increment();
00092         \
00093         TraceBuffer::Write(aul6Msg__, 4);
00094         \
00095         Kernel::Panic(PANIC_ASSERT_FAILED);
00096     }
00097     \
00098     \
00099     \
00100     \
00101     \
00102     \
00103     \
00104     \
00105     \
00106     \
00107     \
00108     \
00109     \
00110     \
00111     \
00112     \
00113     \
00114     \
00115     \
00116     \
00117     \
00118     \
00119     \
00120     \
00121     \
00122     \
00123     \
00124     \
00125     \
00126     \
00127     \
00128     \
00129     \
00130     \
00131     \
00132     \
00133     \
00134     \
00135     \
00136     \
00137     \
00138     \
00139     \
00140     \
00141     \
00142     \
00143     \
00144     \
00145     \

```

```

00146     if ((x) == false) {
00147         Kernel::Panic(PANIC_ASSERT_FAILED);
00148     }
00149 }
00150 #else
00151 //-----
00152 // Note -- when kernel-debugging is disabled, we still have to define the
00153 // macros to ensure that the expressions compile (albeit, by elimination
00154 // during pre-processing).
00155 //-----
00156 #define KERNEL_TRACE(x)
00157 //-----
00158 #define KERNEL_TRACE_1(x, arg1)
00159 //-----
00160 #define KERNEL_TRACE_2(x, arg1, arg2)
00161 //-----
00162 #define KERNEL_ASSERT(x)
00163 #endif // KERNEL_USE_DEBUG
00164
00165 //-----
00166 #if (KERNEL_USE_DEBUG && !KERNEL_AWARE_SIMULATION && KERNEL_ENABLE_USER_LOGGING)
00167 //-----
00168 //-----
00169 //-----
00170 #define USER_TRACE(x)
00171     \
00172 {
00173     \
00174     EMIT_DBG_STRING(x);
00175     \
00176     uint16_t aul6Msg__[4];
00177     \
00178     aul6Msg__[0] = 0xACDC;
00179     \
00180     aul6Msg__[1] = DBG_FILE;
00181     \
00182     aul6Msg__[2] = __LINE__;
00183     \
00184     aul6Msg__[3] = TraceBuffer::Increment();
00185     \
00186     TraceBuffer::Write(aul6Msg__, 4);
00187 }
00188 \
00189 };
00190 //-----
00191 #define USER_TRACE_1(x, arg1)
00192     \
00193 {
00194     \
00195     EMIT_DBG_STRING(x);
00196     \
00197     uint16_t aul6Msg__[5];
00198     \
00199     aul6Msg__[0] = 0xACDC;
00200     \
00201     aul6Msg__[1] = DBG_FILE;
00202     \
00203     aul6Msg__[2] = __LINE__;
00204     \
00205     aul6Msg__[3] = TraceBuffer::Increment();
00206     \
00207     aul6Msg__[4] = arg1;
00208     \
00209     TraceBuffer::Write(aul6Msg__, 5);
00210 }
00211 \
00212 };
00213 //-----
00214 #define USER_TRACE_2(x, arg1, arg2)
00215     \
00216 {
00217     \
00218     EMIT_DBG_STRING(x);
00219     \
00220     uint16_t aul6Msg__[6];
00221     \
00222     aul6Msg__[0] = 0xACDC;

```

```

00205         aul6Msg__[1] = DBG_FILE;
00206         aul6Msg__[2] = __LINE__;
00207         aul6Msg__[3] = TraceBuffer::Increment();
00208         aul6Msg__[4] = arg1;
00209         aul6Msg__[5] = arg2;
00210         TraceBuffer::Write(aul6Msg__, 6);
00211     \
00212 }
00213
00214 //-----
00215 #define USER_ASSERT(x)
00216     \
00217 {
00218     \
00219     if ((x) == false) {
00220         \
00221         EMIT_DBG_STRING("ASSERT FAILED");
00222         \
00223         uint16_t aul6Msg__[4];
00224         \
00225         aul6Msg__[0] = 0xACDC;
00226         \
00227         aul6Msg__[1] = DBG_FILE;
00228         \
00229         aul6Msg__[2] = __LINE__;
00230         \
00231         aul6Msg__[3] = TraceBuffer::Increment();
00232         \
00233         TraceBuffer::Write(aul6Msg__, 4);
00234         \
00235         Kernel::Panic(PANIC_ASSERT_FAILED);
00236     }
00237     \
00238 }
00239
00240 #elif (KERNEL_USE_DEBUG && KERNEL_AWARE_SIMULATION && KERNEL_ENABLE_USER_LOGGING)
00241 //-----
00242 #define USER_TRACE(x)
00243     \
00244 {
00245     \
00246     EMIT_DBG_STRING(x);
00247     \
00248     KernelAware::Trace(DBG_FILE, __LINE__);
00249 }
00250
00251 //-----
00252 #define USER_TRACE_1(x, arg1)
00253     \
00254 {
00255     \
00256     EMIT_DBG_STRING(x);
00257     \
00258     KernelAware::Trace(DBG_FILE, __LINE__, arg1);
00259 }
00260
00261 //-----
00262 #define USER_TRACE_2(x, arg1, arg2)
00263     \
00264 {
00265     \
00266     EMIT_DBG_STRING(x);
00267     \
00268     KernelAware::Trace(DBG_FILE, __LINE__, arg1, arg2);
00269 }
00270
00271 //-----
00272 #define USER_ASSERT(x)
00273     \
00274 {
00275     \
00276     if ((x) == false) {
00277         \
00278         EMIT_DBG_STRING("ASSERT FAILED");
00279         \
00280         uint16_t aul6Msg__[4];
00281         \
00282         aul6Msg__[0] = 0xACDC;
00283         \
00284         aul6Msg__[1] = DBG_FILE;
00285         \
00286         aul6Msg__[2] = __LINE__;
00287         \
00288         aul6Msg__[3] = TraceBuffer::Increment();
00289         \
00290         TraceBuffer::Write(aul6Msg__, 4);
00291         \
00292         Kernel::Panic(PANIC_ASSERT_FAILED);
00293     }
00294 }

```

```

00261     \
00262 {
00263     \
00264     if ((x) == false) {
00265         \
00266         EMIT_DBG_STRING("ASSERT FAILED");
00267         \
00268         KernelAware::Trace(DBG_FILE, __LINE__);
00269         \
00270         Kernel::Panic(PANIC_ASSERT_FAILED);
00271     }
00272     \
00273 }
00274 #else
00275 //-----
00276 // Note -- when kernel-debugging is disabled, we still have to define the
00277 // macros to ensure that the expressions compile (albeit, by elimination
00278 // during pre-processing).
00279 //-----
00280 #define USER_TRACE(x)
00281 //-----
00282 #define USER_TRACE_1(x, arg1)
00283 //-----
00284 #define USER_TRACE_2(x, arg1, arg2)
00285 //-----
00286 #define USER_ASSERT(x)
00287 #endif // KERNEL_USE_DEBUG
00288 } //namespace Mark3
00289

```

20.73 /home/moslevin/projects/github/m3-repo/kernel/src/public/kerneltypes.h File Reference

Basic data type primitives used throughout the OS.

```

#include <stdint.h>
#include <stdbool.h>
#include <stddef.h>

```

Namespaces

- [Mark3](#)

Class providing the software-interrupt required for context-switching in the kernel.

Typedefs

- using [Mark3::PanicFunc](#) = void(*)(uint16_t u16PanicCode_)
Function pointer type used to implement kernel-panic handlers.
- using [Mark3::IdleFunc](#) = void(*)()
Function pointer type used to implement the idle function, where support for an idle function (as opposed to an idle thread) exists.
- using [Mark3::ThreadEntryFunc](#) = void(*)(void *pvArg_)
Function pointer type used for thread entrypoint functions.

Enumerations

- enum [Mark3::EventFlagOperation](#) : uint8_t {
[Mark3::EventFlagOperation::All_Set](#) = 0, [Mark3::EventFlagOperation::Any_Set](#), [Mark3::EventFlagOperation::All_Clear](#), [Mark3::EventFlagOperation::Any_Clear](#),
[Mark3::EventFlagOperation::Pending_Unblock](#) }

This enumeration describes the different operations supported by the event flag blocking object.

- enum [Mark3::ThreadState](#) : uint8_t

Enumeration representing the different states a thread can exist in.

20.73.1 Detailed Description

Basic data type primitives used throughout the OS.

Definition in file [kerneltypes.h](#).

20.74 kerneltypes.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00014
00015 #include <stdint.h>
00016 #include <stdbool.h>
00017 #include <stddef.h>
00018
00019 #pragma once
00020 namespace Mark3
00021 {
00022 //-----
00023 using PanicFunc = void (*)(uint16_t u16PanicCode_);
00024
00025 //-----
00026 using IdleFunc = void (*)();
00027
00028 //-----
00029 using ThreadEntryFunc = void (*)(void* pvArg_);
00030
00031 //-----
00032 enum class EventFlagOperation : uint8_t {
00033     All_Set = 0,
00034     Any_Set,
00035     All_Clear,
00036     Any_Clear,
00037     Pending_Unblock
00038 };
00039
00040 //-----
00041 enum class ThreadState : uint8_t {
00042     Exit = 0,
00043     Ready,
00044     Blocked,
00045     Stop,
00046     Invalid
00047 };
00048 } //namespace Mark3

```


Semaphore Blocking Object class declarations.

Classes

- Binary & Counting semaphores, based on [BlockingObject](#) base class.

Class providing the software-interrupt required for context-switching in the kernel.

Semaphore Blocking Object class declarations.

Definition in file [ksemaphore.h](#).

```
00001 /*=====
00002
00003 |_____|_____|_____|_____|_____|_____|_____|_____|_____|_____|
00004 | \   /| \   /| \   /| \   /| \   /| \   /| \   /| \   /|
00005 |  V   |  V   |  V   |  V   |  V   |  V   |  V   |  V   |
00006 |_/___/\_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/_/
00007 |_____||_____||_____||_____||_____||_____||_____||_____||
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0sleivn, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00021 #pragma once
00022
00023 #include "kerneltypes.h"
00024 #include "mark3cfg.h"
00025
00026 #include "blocking.h"
00027 #include "threadlist.h"
00028
00029 #if KERNEL_USE_SEMAPHORE
00030 namespace Mark3
00031 {
00032 //-----
00036 class Semaphore : public BlockingObject
00037 {
00038 public:
00039     void* operator new(size_t sz, void* pv) { return (Semaphore*)pv; };
00040     ~Semaphore();
00041 }
```

```

00063     void Init(uint16_t u16InitVal_, uint16_t u16MaxVal_);
00064
00079     bool Post();
00080
00088     void Pend();
00089
00101     uint16_t GetCount();
00102
00103 #if KERNEL_USE_TIMEOUTS
00104     bool Pend(uint32_t u32WaitTimeMS_);
00116
00127     void WakeMe(Thread* p1ChosenOne_);
00128 #endif
00129
00130 private:
00136     uint8_t WakeNext();
00137
00138 #if KERNEL_USE_TIMEOUTS
00139     bool Pend_i(uint32_t u32WaitTimeMS_);
00148 #else
00149     void Pend_i(void);
00156 #endif
00157
00158     uint16_t m_u16Value;
00159     uint16_t m_u16MaxValue;
00160 };
00161 } //namespace Mark3
00162 #endif // KERNEL_USE_SEMAPHORE

```

20.77 /home/moslevin/projects/github/m3-repo/kernel/src/public/ll.h File Reference

Core linked-list declarations, used by all kernel list types.

```
#include "kerneltypes.h"
```

Classes

- class [Mark3::LinkListNode](#)
Basic linked-list node data structure.
- class [Mark3::LinkList](#)
Abstract-data-type from which all other linked-lists are derived.
- class [Mark3::DoubleLinkList](#)
Doubly-linked-list data type, inherited from the base [LinkList](#) type.
- class [Mark3::CircularLinkList](#)
Circular-linked-list data type, inherited from the base [LinkList](#) type.

Namespaces

- [Mark3](#)
Class providing the software-interrupt required for context-switching in the kernel.

20.77.1 Detailed Description

Core linked-list declarations, used by all kernel list types.

At the heart of RTOS data structures are linked lists. Having a robust and efficient set of linked-list types that we can use as a foundation for building the rest of our kernel types allows u16 to keep our RTOS code efficient and logically-separated.

So what data types rely on these linked-list classes?

-Threads -ThreadLists -The Scheduler -Timers, -The Timer Scheduler -Blocking objects (Semaphores, Mutexes, etc...)

Pretty much everything in the kernel uses these linked lists. By having objects inherit from the base linked-list node type, we're able to leverage the double and circular linked-list classes to manager virtually every object type in the system without duplicating code. These functions are very efficient as well, allowing for very deterministic behavior in our code.

Definition in file [ll.h](#).

20.78 ll.h

```

00001  /*=====
00002
00003
00004  |-----|-----|-----|-----|-----|-----|
00005  | \ / | \ / | \ / | \ / | \ / | \ / | \ / |
00006  | / \ | / \ | / \ | / \ | / \ | / \ | / \ |
00007  |-----|-----|-----|-----|-----|-----|
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  ===== */
00042  #pragma once
00043  #include "kerneltypes.h"
00044
00045  namespace Mark3
00046  {
00047
00048  //-----
00054  class LinkedList;
00055  class DoubleLinkedList;
00056  class CircularLinkedList;
00057
00058  //-----
00063  class LinkedListNode
00064  {
00065  protected:
00066      LinkedListNode* next;
00067      LinkedListNode* prev;
00068
00069      LinkedListNode() {}
00075      void ClearNode();
00076
00077  public:
00085      LinkedListNode* GetNext(void) { return next; }
00093      LinkedListNode* GetPrev(void) { return prev; }
00094      friend class LinkedList;
00095      friend class DoubleLinkedList;
00096      friend class CircularLinkedList;
00097      friend class ThreadList;
00098  };
00099
00100  //-----
00104  class LinkedList
00105  {
00106  protected:
00107      LinkedListNode* m_pclHead;
00108      LinkedListNode* m_pclTail;
00109

```

```

00110 public:
00116     void Init()
00117     {
00118         m_pclHead = NULL;
00119         m_pclTail = NULL;
00120     }
00121
00129     LinkListNode* GetHead() { return m_pclHead; }
00137     LinkListNode* GetTail() { return m_pclTail; }
00138 };
00139
00140 //-----
00144 class DoubleLinkedList : public LinkList
00145 {
00146 public:
00147     void* operator new(size_t sz, void* pv) { return (DoubleLinkedList*)pv; };
00153     DoubleLinkedList()
00154     {
00155         m_pclHead = NULL;
00156         m_pclTail = NULL;
00157     }
00166     void Add(LinkListNode* node_);
00167
00175     void Remove(LinkListNode* node_);
00176 };
00177
00178 //-----
00182 class CircularLinkedList : public LinkList
00183 {
00184 public:
00185     void* operator new(size_t sz, void* pv) { return (CircularLinkedList*)pv; };
00186     CircularLinkedList()
00187     {
00188         m_pclHead = NULL;
00189         m_pclTail = NULL;
00190     }
00191
00199     void Add(LinkListNode* node_);
00200
00208     void Remove(LinkListNode* node_);
00209
00216     void PivotForward();
00217
00224     void PivotBackward();
00225
00235     void InsertNodeBefore(LinkListNode* node_, LinkListNode* insert_);
00236 };
00237 } //namespace Mark3

```

20.79 /home/moslevin/projects/github/m3-repo/kernel/src/public/lockguard.h File Reference

Mutex RAI helper class.

```
#include "mark3.h"
```

Classes

- class [Mark3::LockGuard](#)

Namespaces

- [Mark3](#)

Class providing the software-interrupt required for context-switching in the kernel.

20.79.1 Detailed Description

Mutex RAIL helper class.

Definition in file [lockguard.h](#).

20.80 lockguard.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00020 #pragma once
00021
00022 #include "mark3.h"
00023
00024 #if KERNEL_USE_MUTEX
00025 namespace Mark3 {
00026
00032 class LockGuard {
00033 public:
00037     LockGuard(Mutex* pclMutex);
00038
00039 #if KERNEL_USE_TIMEOUTS
00040
00045     LockGuard(Mutex* pclMutex, uint32_t u32TimeoutMs);
00046 #endif
00047
00048     ~LockGuard();
00049
00056     bool isAcquired() { return m_bIsAcquired; }
00057
00058 private:
00059     bool m_bIsAcquired;
00060     Mutex* m_pclMutex;
00061 };
00062 } // namespace Mark3
00063
00064 #endif // KERNEL_USE_MUTEX

```

20.81 /home/moslevin/projects/github/m3-repo/kernel/src/public/mailbox.h File Reference

Mailbox + Envelope IPC Mechanism.

```

#include "mark3cfg.h"
#include "kerneltypes.h"
#include "threadport.h"
#include "ksemaphore.h"

```

Classes

- class [Mark3::Mailbox](#)

The [Mailbox](#) class implements an IPC mechnism based on envelopes containing data of a fixed size (configured at initialization) that reside within a buffer of memory provided by the user.


```

00189         uint16_t rc;
00190         CS_ENTER();
00191         rc = m_ul6Free;
00192         CS_EXIT();
00193         return rc;
00194     }
00195
00196     bool IsFull(void) { return (GetFreeSlots() == 0); }
00197     bool IsEmpty(void) { return (GetFreeSlots() == m_ul6Count); }
00198 private:
00207     void* GetHeadPointer(void)
00208     {
00209         K_ADDR uAddr = (K_ADDR)m_pvBuffer;
00210         uAddr += (K_ADDR)(m_ul6ElementSize) * (K_ADDR)(
m_ul6Head);
00211         return (void*)uAddr;
00212     }
00213
00222     void* GetTailPointer(void)
00223     {
00224         K_ADDR uAddr = (K_ADDR)m_pvBuffer;
00225         uAddr += (K_ADDR)(m_ul6ElementSize) * (K_ADDR)(
m_ul6Tail);
00226         return (void*)uAddr;
00227     }
00228
00238     void CopyData(const void* src_, const void* dst_, uint16_t len_)
00239     {
00240         uint8_t* u8Src = (uint8_t*)src_;
00241         uint8_t* u8Dst = (uint8_t*)dst_;
00242         while (len_-- > 0) {
00243             *u8Dst++ = *u8Src++;
00244         }
00245     }
00246
00252     void MoveTailForward(void)
00253     {
00254         m_ul6Tail++;
00255         if (m_ul6Tail == m_ul6Count) {
00256             m_ul6Tail = 0;
00257         }
00258     }
00259
00265     void MoveHeadForward(void)
00266     {
00267         m_ul6Head++;
00268         if (m_ul6Head == m_ul6Count) {
00269             m_ul6Head = 0;
00270         }
00271     }
00272
00278     void MoveTailBackward(void)
00279     {
00280         if (m_ul6Tail == 0) {
00281             m_ul6Tail = m_ul6Count;
00282         }
00283         m_ul6Tail--;
00284     }
00285
00291     void MoveHeadBackward(void)
00292     {
00293         if (m_ul6Head == 0) {
00294             m_ul6Head = m_ul6Count;
00295         }
00296         m_ul6Head--;
00297     }
00298
00299 #if KERNEL_USE_TIMEOUTS
00300
00310     bool Send_i(const void* pvData_, bool bTail_, uint32_t u32TimeoutMS_);
00311 #else
00312
00321     bool Send_i(const void* pvData_, bool bTail_);
00322 #endif
00323
00324 #if KERNEL_USE_TIMEOUTS
00325
00335     bool Receive_i(const void* pvData_, bool bTail_, uint32_t u32WaitTimeMS_);
00336 #else
00337
00345     void Receive_i(const void* pvData_, bool bTail_);
00346 #endif
00347
00348     uint16_t m_ul6Head;
00349     uint16_t m_ul6Tail;
00350
00351     uint16_t m_ul6Count;

```



```
#include "eventflag.h"
#include "message.h"
#include "notify.h"
#include "mailbox.h"
#include "readerwriter.h"
#include "condvar.h"
#include "atomic.h"
#include "kernelaware.h"
#include "profile.h"
#include "autoalloc.h"
#include "priomap.h"
```

20.85.1 Detailed Description

Single include file given to users of the [Mark3](#) Kernel API.

Definition in file [mark3.h](#).

20.86 mark3.h

```
00001 /*=====
00002
00003  _____
00004  |  _ \  /  _ \  /  _ \  /  _ \  /  _ \  /  _ \  /  _ \  /  _ \
00005  | / \ \ \ / \ \ \ / \ \ \ / \ \ \ / \ \ \ / \ \ \ / \ \ \ / \ \ \
00006  |/_/ \/_/ \/_/ \/_/ \/_/ \/_/ \/_/ \/_/ \/_/ \/_/ \/_/ \/_/ \/_/
00007  |_____|
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  =====*/
00020  #pragma once
00021
00022  #include "mark3cfg.h"
00023  #include "kerneltypes.h"
00024
00025  #include "threadport.h"
00026  #include "kernelswi.h"
00027  #include "kerneltimer.h"
00028  #include "kernelprofile.h"
00029
00030  #include "kernel.h"
00031  #include "thread.h"
00032  #include "timerlist.h"
00033
00034  #include "ksemaphore.h"
00035  #include "mutex.h"
00036  #include "eventflag.h"
00037  #include "message.h"
00038  #include "notify.h"
00039  #include "mailbox.h"
00040  #include "readerwriter.h"
00041  #include "condvar.h"
00042
00043  #include "atomic.h"
00044  #include "kernelaware.h"
00045
00046  #include "profile.h"
00047  #include "autoalloc.h"
00048  #include "priomap.h"
```

20.87 /home/moslevin/projects/github/m3-repo/kernel/src/public/mark3cfg.h File Reference

[Mark3](#) Kernel Configuration.

```
#include "portcfg.h"
```

Macros

- `#define` [KERNEL_NUM_PRIORITIES](#) (8)
Define the number of thread priorities that the kernel's scheduler will support.
- `#define` [KERNEL_USE_TIMERS](#) (1)
The following options is related to all kernel time-tracking.
- `#define` [KERNEL_TIMERS_TICKLESS](#) (1)
If you've opted to use the kernel timers module, you have an option as to which timer implementation to use: Tick-based or Tick-less.
- `#define` [KERNEL_TIMERS_MINIMUM_DELAY_US](#) (25)
When using tickless timers, it is useful to define a minimum sleep value.
- `#define` [KERNEL_TIMERS_THREADED](#) (0)
When timers are enabled, configure whether or not a dedicated thread is used to service timer maintenance.
- `#define` [KERNEL_USE_TIMEOUTS](#) (1)
Set the priority of the timer thread, if the kernel is configured to use a dedicated timer thread.
- `#define` [KERNEL_USE_QUANTUM](#) (1)
Do you want to enable time quanta? This is useful when you want to have tasks in the same priority group share time in a controlled way.
- `#define` [THREAD_QUANTUM_DEFAULT](#) (4)
This value defines the default thread quantum when [KERNEL_USE_QUANTUM](#) is enabled.
- `#define` [KERNEL_USE_NOTIFY](#) (1)
This is a simple blocking object, where a thread (or threads) are guaranteed to block until an asynchronous event signals the object.
- `#define` [KERNEL_USE_SEMAPHORE](#) (1)
Do you want the ability to use counting/binary semaphores for thread synchronization? Enabling this features provides fully-blocking semaphores and enables all API functions declared in [semaphore.h](#).
- `#define` [KERNEL_USE_MUTEX](#) (1)
Do you want the ability to use mutual exclusion semaphores (mutex) for resource/block protection? Enabling this feature provides mutexes, with priority inheritance, as declared in [mutex.h](#).
- `#define` [KERNEL_USE_EVENTFLAG](#) (1)
Provides additional event-flag based blocking.
- `#define` [KERNEL_USE_READERWRITER](#) (1)
Provides reader-writer locks.
- `#define` [KERNEL_USE_CONDVAR](#) (1)
Provides condition variables.
- `#define` [KERNEL_USE_MESSAGE](#) (1)
Enable inter-thread messaging using message queues.
- `#define` [GLOBAL_MESSAGE_POOL_SIZE](#) (8)
If Messages are enabled, define the size of the default kernel message pool.
- `#define` [KERNEL_USE_MAILBOX](#) (1)
Enable inter-thread messaging using mailboxes.
- `#define` [KERNEL_USE_SLEEP](#) (1)
Do you want to be able to set threads to sleep for a specified time? This enables the `Thread::Sleep()` API.

- `#define KERNEL_USE_THREADNAME (0)`
Provide Thread method to allow the user to set a name for each thread in the system.
- `#define KERNEL_USE_EXTENDED_CONTEXT (1)`
Allocate an extra pointer's worth of storage within a Thread object (and corresponding accessor methods) to provide the user with a means to implement arbitrary Thread-local storage.
- `#define KERNEL_USE_DYNAMIC_THREADS (1)`
Provide extra Thread methods to allow the application to create (and more importantly destroy) threads at runtime.
- `#define KERNEL_USE_PROFILER (1)`
Provides extra classes for profiling the performance of code.
- `#define KERNEL_USE_DEBUG (0)`
Provides extra logic for kernel debugging, and instruments the kernel with extra asserts, and kernel trace functionality.
- `#define KERNEL_USE_ATOMIC (1)`
Provides support for atomic operations, including addition, subtraction, set, and test-and-set.
- `#define SAFE_UNLINK (0)`
"Safe unlinking" performs extra checks on data to make sure that there are no consistencies when performing operations on linked lists.
- `#define KERNEL_AWARE_SIMULATION (1)`
Include support for kernel-aware simulation.
- `#define KERNEL_USE_IDLE_FUNC (1)`
Enabling this feature removes the necessity for the user to dedicate a complete thread for idle functionality.
- `#define KERNEL_USE_AUTO_ALLOC (1)`
This feature enables an additional set of APIs that allow for objects to be created on-the-fly out of a special heap, without having to explicitly allocate them (from stack, heap, or static memory).
- `#define KERNEL_USE_THREAD_CALLOUTS (1)`
This feature provides additional kernel APIs to register callout functions that are activated when threads are created or exited.
- `#define KERNEL_USE_STACK_GUARD (1)`
This feature, when enabled, tells the kernel to check whether any Thread's stack has been exhausted (or slack falls below a certain safety threshold) before executing each context switch.
- `#define KERNEL_EXTRA_CHECKS (1)`
This option provides extra safety checks within the kernel APIs in order to minimize the potential for unsafe operations.

20.87.1 Detailed Description

[Mark3](#) Kernel Configuration.

This file is used to configure the kernel for your specific application in order to provide the optimal set of features for a given use case.

Since you only pay the price (code space/RAM) for the features you use, you can usually find a sweet spot between features and resource usage by picking and choosing features a-la-carte. This config file is written in an "interactive" way, in order to minimize confusion about what each option provides, and to make dependencies obvious.

Definition in file [mark3cfg.h](#).

20.87.2 Macro Definition Documentation

20.87.2.1 GLOBAL_MESSAGE_POOL_SIZE

```
#define GLOBAL_MESSAGE_POOL_SIZE (8)
```

If Messages are enabled, define the size of the default kernel message pool.

Messages can be manually added to the message pool, but this mechanism is more convenient and automatic. All message queues share their message objects from this global pool to maximize efficiency and simplify data management.

Definition at line 211 of file [mark3cfg.h](#).

20.87.2.2 KERNEL_AWARE_SIMULATION

```
#define KERNEL_AWARE_SIMULATION (1)
```

Include support for kernel-aware simulation.

Enabling this feature adds advanced profiling, trace, and environment-aware debugging and diagnostic functionality when Mark3-based applications are run on the flavr AVR simulator.

Definition at line 315 of file [mark3cfg.h](#).

20.87.2.3 KERNEL_EXTRA_CHECKS

```
#define KERNEL_EXTRA_CHECKS (1)
```

This option provides extra safety checks within the kernel APIs in order to minimize the potential for unsafe operations.

This is especially helpful during development, and can help catch problems at development time, instead of in the field. include CPU/Port specific configuration options

Definition at line 370 of file [mark3cfg.h](#).

20.87.2.4 KERNEL_NUM_PRIORITIES

```
#define KERNEL_NUM_PRIORITIES (8)
```

Define the number of thread priorities that the kernel's scheduler will support.

The number of thread priorities is limited only by the memory of the host CPU, as a ThreadList object is statically-allocated for each thread priority.

In practice, systems rarely need more than 32 priority levels, with the most complex having the capacity for 256.

Definition at line 39 of file [mark3cfg.h](#).

20.87.2.5 KERNEL_TIMERS_MINIMUM_DELAY_US

```
#define KERNEL_TIMERS_MINIMUM_DELAY_US (25)
```

When using tickless timers, it is useful to define a minimum sleep value.

In the event that a delay/sleep/timeout value lower than this is provided to a timer-based API, the minimum value will be substituted.

Definition at line 84 of file [mark3cfg.h](#).

20.87.2.6 KERNEL_TIMERS_THREADED

```
#define KERNEL_TIMERS_THREADED (0)
```

When timers are enabled, configure whether or not a dedicated thread is used to service timer maintenance.

If set to 0, timer handlers are executed from a nested interrupt context.

Definition at line 93 of file [mark3cfg.h](#).

20.87.2.7 KERNEL_TIMERS_TICKLESS

```
#define KERNEL_TIMERS_TICKLESS (1)
```

If you've opted to use the kernel timers module, you have an option as to which timer implementation to use: Tick-based or Tick-less.

Tick-based timers provide a "traditional" RTOS timer implementation based on a fixed-frequency timer interrupt. While this provides very accurate, reliable timing, it also means that the CPU is being interrupted far more often than may be necessary (as not all timer ticks result in "real work" being done).

Tick-less timers still rely on a hardware timer interrupt, but uses a dynamic expiry interval to ensure that the interrupt is only called when the next timer expires. This increases the complexity of the timer interrupt handler, but reduces the number and frequency.

Note that the CPU port ([kerneltimer.cpp](#)) must be implemented for the particular timer variant desired.

Definition at line 75 of file [mark3cfg.h](#).

20.87.2.8 KERNEL_USE_ATOMIC

```
#define KERNEL_USE_ATOMIC (1)
```

Provides support for atomic operations, including addition, subtraction, set, and test-and-set.

Add/Sub/Set contain 8, 16, and 32-bit variants.

Definition at line 299 of file [mark3cfg.h](#).

20.87.2.9 KERNEL_USE_AUTO_ALLOC

```
#define KERNEL_USE_AUTO_ALLOC (1)
```

This feature enables an additional set of APIs that allow for objects to be created on-the-fly out of a special heap, without having to explicitly allocate them (from stack, heap, or static memory).

Note that auto-alloc memory cannot be reclaimed.

Definition at line 336 of file [mark3cfg.h](#).

20.87.2.10 KERNEL_USE_CONDVAR

```
#define KERNEL_USE_CONDVAR (1)
```

Provides condition variables.

Allows a thread to wait for a specific condition to be true before proceeding.

Definition at line 187 of file [mark3cfg.h](#).

20.87.2.11 KERNEL_USE_DYNAMIC_THREADS

```
#define KERNEL_USE_DYNAMIC_THREADS (1)
```

Provide extra Thread methods to allow the application to create (and more importantly destroy) threads at runtime.

useful for designs implementing worker threads, or threads that can be restarted after encountering error conditions.

Definition at line 259 of file [mark3cfg.h](#).

20.87.2.12 KERNEL_USE_EVENTFLAG

```
#define KERNEL_USE_EVENTFLAG (1)
```

Provides additional event-flag based blocking.

This relies on an additional per-thread flag-mask to be allocated, which adds 2 bytes to the size of each thread object.

Definition at line 168 of file [mark3cfg.h](#).

20.87.2.13 KERNEL_USE_IDLE_FUNC

```
#define KERNEL_USE_IDLE_FUNC (1)
```

Enabling this feature removes the necessity for the user to dedicate a complete thread for idle functionality.

This saves a full thread stack, but also requires a bit extra static data. This also adds a slight overhead to the context switch and scheduler, as a special case has to be taken into account.

Definition at line 325 of file [mark3cfg.h](#).

20.87.2.14 KERNEL_USE_MAILBOX

```
#define KERNEL_USE_MAILBOX (1)
```

Enable inter-thread messaging using mailboxes.

A mailbox manages a blob of data provided by the user, that is partitioned into fixed-size blocks called envelopes. The size of an envelope is set by the user when the mailbox is initialized. Any number of threads can read-from and write-to the mailbox. Envelopes can be sent-to or received-from the mailbox at the head or tail. In this way, mailboxes essentially act as a circular buffer that can be used as a blocking FIFO or LIFO queue.

Definition at line 224 of file [mark3cfg.h](#).

20.87.2.15 KERNEL_USE_MESSAGE

```
#define KERNEL_USE_MESSAGE (1)
```

Enable inter-thread messaging using message queues.

This is the preferred mechanism for IPC for serious multi-threaded communications; generally anywhere a semaphore or event-flag is insufficient.

Definition at line 198 of file [mark3cfg.h](#).

20.87.2.16 KERNEL_USE_PROFILER

```
#define KERNEL_USE_PROFILER (1)
```

Provides extra classes for profiling the performance of code.

useful for debugging and development, but uses an additional hardware timer.

Definition at line 265 of file [mark3cfg.h](#).

20.87.2.17 KERNEL_USE_QUANTUM

```
#define KERNEL_USE_QUANTUM (1)
```

Do you want to enable time quanta? This is useful when you want to have tasks in the same priority group share time in a controlled way.

This allows equal tasks to use unequal amounts of the CPU, which is a great way to set up CPU budgets per thread in a round-robin scheduling system. If enabled, you can specify a number of ticks that serves as the default time period (quantum). Unless otherwise specified, every thread in a priority will get the default quantum.

Definition at line 131 of file [mark3cfg.h](#).

20.87.2.18 KERNEL_USE_READERWRITER

```
#define KERNEL_USE_READERWRITER (1)
```

Provides reader-writer locks.

Allows current read access, or single write-access to a resource. Readers wait for the writer to release the lock, and writers wait for all readers to release the lock before acquiring the resource.

Definition at line 177 of file [mark3cfg.h](#).

20.87.2.19 KERNEL_USE_SEMAPHORE

```
#define KERNEL_USE_SEMAPHORE (1)
```

Do you want the ability to use counting/binary semaphores for thread synchronization? Enabling this features provides fully-blocking semaphores and enables all API functions declared in semaphore.h.

If you have to pick one blocking mechanism, this is the one to choose.

Definition at line 154 of file [mark3cfg.h](#).

20.87.2.20 KERNEL_USE_STACK_GUARD

```
#define KERNEL_USE_STACK_GUARD (1)
```

This feature, when enabled, tells the kernel to check whether any Thread's stack has been exhausted (or slack falls below a certain safety threshold) before executing each context switch.

Enabling this is the most effective means to guard against stack corruption and stack overflow in the kernel, at the cost of increased context switch latency.

Definition at line 358 of file [mark3cfg.h](#).

20.87.2.21 KERNEL_USE_THREAD_CALLOUTS

```
#define KERNEL_USE_THREAD_CALLOUTS (1)
```

This feature provides additional kernel APIs to register callout functions that are activated when threads are created or exited.

This is useful for implementing low-level instrumentation based on information held in the threads.

Definition at line 348 of file [mark3cfg.h](#).

20.87.2.22 KERNEL_USE_THREADNAME

```
#define KERNEL_USE_THREADNAME (0)
```

Provide Thread method to allow the user to set a name for each thread in the system.

Adds a const char* pointer to the size of the thread object.

Definition at line 244 of file [mark3cfg.h](#).

20.87.2.23 KERNEL_USE_TIMEOUTS

```
#define KERNEL_USE_TIMEOUTS (1)
```

Set the priority of the timer thread, if the kernel is configured to use a dedicated timer thread.

By default, if you opt to enable kernel timers, you also get timeout- enabled versions of the blocking object APIs along with it. This support comes at a small cost to code size, but a slightly larger cost to realtime performance - as checking for the use of timers in the underlying internal code costs some cycles.

As a result, the option is given to the user here to manually disable these timeout-based APIs if desired by the user for performance and code-size reasons.

Definition at line 116 of file [mark3cfg.h](#).

20.87.2.24 KERNEL_USE_TIMERS

```
#define KERNEL_USE_TIMERS (1)
```

The following options is related to all kernel time-tracking.

-timers provide a way for events to be periodically triggered in a lightweight manner. These can be periodic, or one-shot.

-Thread Quantum (usedd for round-robin scheduling) is dependent on this module, as is Thread Sleep functionality.

Definition at line 54 of file [mark3cfg.h](#).

20.87.2.25 SAFE_UNLINK

```
#define SAFE_UNLINK (0)
```

"Safe unlinking" performs extra checks on data to make sure that there are no consistencies when performing operations on linked lists.

This goes beyond pointer checks, adding a layer of structural and metadata validation to help detect system corruption early.

Definition at line 307 of file [mark3cfg.h](#).

20.87.2.26 THREAD_QUANTUM_DEFAULT

```
#define THREAD_QUANTUM_DEFAULT (4)
```

This value defines the default thread quantum when `KERNEL_USE_QUANTUM` is enabled.

The thread quantum value is in milliseconds

Definition at line 140 of file [mark3cfg.h](#).

20.88 mark3cfg.h

```
00001 /*=====
00002
00003  _____
00004  |   /   \   |   /   \   |   /   \   |   /   \   |
00005  |  /     \  |  /     \  |  /     \  |  /     \  |
00006  | /       \ | /       \ | /       \ | /       \ |
00007  |_____|   |_____|   |_____|   |_____|   |
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00028 #pragma once
00029
00039 #define KERNEL_NUM_PRIORITIES (8)
00040
00041 #if KERNEL_NUM_PRIORITIES > 1024
00042 #error "Mark3 supports a maximum of 1024 priorities"
00043 #endif
00044
00054 #define KERNEL_USE_TIMERS (1)
00055
00074 #if KERNEL_USE_TIMERS && !defined(ARM)
00075 #define KERNEL_TIMERS_TICKLESS (1)
00076 #endif
00077
00078 #if KERNEL_TIMERS_TICKLESS
00079
00084 #define KERNEL_TIMERS_MINIMUM_DELAY_US (25)
00085 #endif
00086
00092 #if KERNEL_USE_TIMERS
00093 # define KERNEL_TIMERS_THREADED (0)
00094 #endif
00095
00100 #if KERNEL_TIMERS_THREADED
00101 # define KERNEL_TIMERS_THREAD_PRIORITY (KERNEL_NUM_PRIORITIES - 1)
00102 #endif
00103
00115 #if KERNEL_USE_TIMERS
00116 #define KERNEL_USE_TIMEOUTS (1)
```

```
00117 #else
00118 #define KERNEL_USE_TIMEOUTS (0)
00119 #endif
00120
00130 #if KERNEL_USE_TIMERS
00131 #define KERNEL_USE_QUANTUM (1)
00132 #else
00133 #define KERNEL_USE_QUANTUM (0)
00134 #endif
00135
00140 #define THREAD_QUANTUM_DEFAULT (4)
00141
00146 #define KERNEL_USE_NOTIFY (1)
00147
00154 #define KERNEL_USE_SEMAPHORE (1)
00155
00161 #define KERNEL_USE_MUTEX (1)
00162
00168 #define KERNEL_USE_EVENTFLAG (1)
00169
00176 #if KERNEL_USE_MUTEX
00177 # define KERNEL_USE_READERWRITER (1)
00178 #else
00179 # define KERNEL_USE_READERWRITER (0)
00180 #endif
00181
00186 #if KERNEL_USE_SEMAPHORE && KERNEL_USE_MUTEX
00187 # define KERNEL_USE_CONDVAR (1)
00188 #else
00189 # define KERNEL_USE_CONDVAR (0)
00190 #endif
00191
00197 #if KERNEL_USE_SEMAPHORE
00198 #define KERNEL_USE_MESSAGE (1)
00199 #else
00200 #define KERNEL_USE_MESSAGE (0)
00201 #endif
00202
00210 #if KERNEL_USE_MESSAGE
00211 #define GLOBAL_MESSAGE_POOL_SIZE (8)
00212 #endif
00213
00223 #if KERNEL_USE_SEMAPHORE
00224 #define KERNEL_USE_MAILBOX (1)
00225 #else
00226 #define KERNEL_USE_MAILBOX (0)
00227 #endif
00228
00233 #if KERNEL_USE_TIMERS && KERNEL_USE_SEMAPHORE
00234 #define KERNEL_USE_SLEEP (1)
00235 #else
00236 #define KERNEL_USE_SLEEP (0)
00237 #endif
00238
00244 #define KERNEL_USE_THREADNAME (0)
00245
00251 #define KERNEL_USE_EXTENDED_CONTEXT (1)
00252
00259 #define KERNEL_USE_DYNAMIC_THREADS (1)
00260
00265 #define KERNEL_USE_PROFILER (1)
00266
00271 #define KERNEL_USE_DEBUG (0)
00272
00273 #if KERNEL_USE_DEBUG
00274
00280 #define KERNEL_ENABLE_LOGGING (0)
00281
00289 #define KERNEL_ENABLE_USER_LOGGING (0)
00290 #else
00291 #define KERNEL_ENABLE_LOGGING (0)
00292 #define KERNEL_ENABLE_USER_LOGGING (0)
00293 #endif
00294
00299 #define KERNEL_USE_ATOMIC (1)
00300
00307 #define SAFE_UNLINK (0)
00308
00315 #define KERNEL_AWARE_SIMULATION (1)
00316
00324 #if !defined(ARM)
00325 #define KERNEL_USE_IDLE_FUNC (1) // Supported everywhere but ARM
00326 #else
00327 #define KERNEL_USE_IDLE_FUNC (0) // Not currently supported on ARM
00328 #endif
00329
00336 #define KERNEL_USE_AUTO_ALLOC (1)
```

```

00337
00338 #if KERNEL_USE_AUTO_ALLOC
00339 #define AUTO_ALLOC_SIZE (512)
00340 #endif
00341
00348 #define KERNEL_USE_THREAD_CALLOUTS (1)
00349
00358 #define KERNEL_USE_STACK_GUARD (1)
00359
00360 #if KERNEL_USE_STACK_GUARD
00361 #define KERNEL_STACK_GUARD_DEFAULT (32) // words
00362 #endif
00363
00370 #define KERNEL_EXTRA_CHECKS (1)
00371
00372 #include "portcfg.h"

```

20.89 /home/moslevin/projects/github/m3-repo/kernel/src/public/message.h File Reference

Inter-thread communication via message-passing.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ll.h"
#include "ksemaphore.h"
#include "timerlist.h"

```

Classes

- class [Mark3::Message](#)
Class to provide message-based IPC services in the kernel.
- class [Mark3::MessagePool](#)
Implements a list of message objects.
- class [Mark3::MessageQueue](#)
List of messages, used as the channel for sending and receiving messages between threads.

Namespaces

- [Mark3](#)
Class providing the software-interrupt required for context-switching in the kernel.

20.89.1 Detailed Description

Inter-thread communication via message-passing.

Embedded systems guru Jack Ganssle once said that without a robust form of interprocess communications (IPC), an RTOS is just a toy. [Mark3](#) implements a form of IPC to provide safe and flexible messaging between threads.

using kernel-managed IPC offers significant benefits over other forms of data sharing (i.e. Global variables) in that it avoids synchronization issues and race conditions common to the practice. using IPC also enforces a more disciplined coding style that keeps threads decoupled from one another and minimizes global data, preventing careless and hard-to-debug errors.

20.89.2 using Messages, Queues, and the Global Message Pool

```
// Declare a message queue shared between two threads
MessageQueue my_queue;

int main()
{
    ...
    // Initialize the message queue
    my_queue.init();
    ...
}

void Thread1()
{
    // Example TX thread - sends a message every 10ms
    while(1)
    {
        // Grab a message from the global message pool
        Message *tx_message = GlobalMessagePool::Pop();

        // Set the message data/parameters
        tx_message->SetCode( 1234 );
        tx_message->SetData( NULL );

        // Send the message on the queue.
        my_queue.Send( tx_message );
        Thread::Sleep(10);
    }
}

void Thread2()
{
    while()
    {
        // Blocking receive - wait until we have messages to process
        Message *rx_message = my_queue.Recv();

        // Do something with the message data...

        // Return back into the pool when done
        GlobalMessagePool::Push(rx_message);
    }
}
```

Definition in file [message.h](#).

20.90 message.h

```

00001 /*-----
00002
00003 |-----|-----|-----|-----|-----|-----|
00004 |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |
00005 |  / \  \   |  / \  \   |  / \  \   |  / \  \   |  / \  \   |
00006 | /  \ / \  | /  \ / \  | /  \ / \  | /  \ / \  | /  \ / \  |
00007 |-----|-----|-----|-----|-----|
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00014 #pragma once
00015
00016 #include "kerneltypes.h"
00017 #include "mark3cfg.h"
00018
00019 #include "ll.h"
00020 #include "ksemaphore.h"
00021
00022 #if KERNEL_USE_MESSAGE
00023
00024 #if KERNEL_USE_TIMEOUTS
00025 #include "timerlist.h"
00026 #endif
00027 namespace Mark3
00028 {
00029
00030 //-----

```

```

00098 class Message : public LinkListNode
00099 {
00100 public:
00101     void* operator new(size_t sz, void* pv) { return (Message*)pv; }
00107     void Init()
00108     {
00109         ClearNode();
00110         m_pvData = NULL;
00111         m_ul6Code = 0;
00112     }
00113
00121     void SetData(void* pvData_) { m_pvData = pvData_; }
00129     void* GetData() { return m_pvData; }
00137     void SetCode(uint16_t u16Code_) { m_ul6Code = u16Code_; }
00145     uint16_t GetCode() { return m_ul6Code; }
00146 private:
00148     void* m_pvData;
00149
00151     uint16_t m_ul6Code;
00152 };
00153
00154 //-----
00158 class MessagePool
00159 {
00160 public:
00161     void* operator new(size_t sz, void* pv) { return (MessagePool*)pv; }
00162     ~MessagePool() {}
00168     void Init();
00179     void Push(Message* pclMessage_);
00180
00189     Message* Pop();
00190
00198     Message* GetHead();
00199
00200 private:
00202     DoubleLinkList m_clList;
00203 };
00204
00205 //-----
00210 class MessageQueue
00211 {
00212 public:
00213     void* operator new(size_t sz, void* pv) { return (MessageQueue*)pv; }
00214     ~MessageQueue() {}
00215
00221     void Init();
00222
00231     Message* Receive();
00232
00233     #if KERNEL_USE_TIMEOUTS
00234
00248     Message* Receive(uint32_t u32TimeWaitMS_);
00249     #endif
00250
00259     void Send(Message* pclSrc_);
00260
00268     uint16_t GetCount();
00269
00270 private:
00271     #if KERNEL_USE_TIMEOUTS
00272
00281     Message* Receive_i(uint32_t u32TimeWaitMS_);
00282     #else
00283
00290     Message* Receive_i(void);
00291     #endif
00292
00294     Semaphore m_clSemaphore;
00295
00297     DoubleLinkList m_clLinkList;
00298 };
00299 } //namespace Mark3
00300 #endif // KERNEL_USE_MESSAGE

```

20.91 /home/moslevin/projects/github/m3-repo/kernel/src/public/mutex.h File Reference

Mutual exclusion class declaration.

```

#include "kerneltypes.h"
#include "mark3cfg.h"

```

```
#include "blocking.h"
```

Classes

- class [Mark3::Mutex](#)

Mutual-exclusion locks, based on [BlockingObject](#).

Namespaces

- [Mark3](#)

Class providing the software-interrupt required for context-switching in the kernel.

20.91.1 Detailed Description

Mutual exclusion class declaration.

Resource locks are implemented using mutual exclusion semaphores (`Mutex_t`). Protected blocks can be placed around any resource that may only be accessed by one thread at a time. If additional threads attempt to access the protected resource, they will be placed in a wait queue until the resource becomes available. When the resource becomes available, the thread with the highest original priority claims the resource and is activated. Priority inheritance is included in the implementation to prevent priority inversion. Always ensure that you claim and release your mutex objects consistently, otherwise you may end up with a deadlock scenario that's hard to debug.

20.91.2 Initializing

Initializing a mutex object by calling:

```
clMutex.Init();
```

20.91.3 Resource protection example

```
clMutex.Claim();  
...  
<resource protected block>  
...  
clMutex.Release();
```

Definition in file [mutex.h](#).

20.92 mutex.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00050 #pragma once
00051
00052 #include "kerneltypes.h"
00053 #include "mark3cfg.h"
00054
00055 #include "blocking.h"
00056
00057 #if KERNEL_USE_MUTEX
00058 namespace Mark3
00059 {
00060 //-----
00064 class Mutex : public BlockingObject
00065 {
00066 public:
00067     void* operator new(size_t sz, void* pv) { return (Mutex*)pv; };
00068     ~Mutex();
00069
00078     void Init(bool bRecursive_ = true);
00079
00097     void Claim();
00098
00099 #if KERNEL_USE_TIMEOUTS
00100
00111     bool Claim(uint32_t u32WaitTimeMS_);
00112
00125     void WakeMe(Thread* pclOwner_);
00126
00127 #endif
00128
00149     void Release();
00150
00151 private:
00157     uint8_t WakeNext();
00158
00159 #if KERNEL_USE_TIMEOUTS
00160
00168     bool Claim_i(uint32_t u32WaitTimeMS_);
00169 #else
00176     void Claim_i(void);
00177 #endif
00178
00179     uint8_t m_u8Recurse;
00180     bool m_bReady;
00181     bool m_bRecursive;
00182     uint8_t m_u8MaxPri;
00183     Thread* m_pclOwner;
00184 };
00185 } //namespace Mark3
00186 #endif // KERNEL_USE_MUTEX

```

20.93 /home/moslevin/projects/github/m3-repo/kernel/src/public/notify.h File Reference

Lightweight thread notification - blocking object.

```

#include "mark3cfg.h"
#include "blocking.h"

```


Classes

- class [Mark3::Notify](#)

The [Notify](#) class is a blocking object type, that allows one or more threads to wait for an event to occur before resuming operation.

Namespaces

- [Mark3](#)

Class providing the software-interrupt required for context-switching in the kernel.

20.93.1 Detailed Description

Lightweight thread notification - blocking object.

Definition in file [notify.h](#).

20.94 notify.h

```

00001 /*=====
00002
00003
00004 | | | | | | | | | | | | | | | | | |
00005 | | | | | | | | | | | | | | | | | |
00006 | | | | | | | | | | | | | | | | | |
00007 | | | | | | | | | | | | | | | | | |
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00021 #pragma once
00022
00023 #include "mark3cfg.h"
00024 #include "blocking.h"
00025
00026 #if KERNEL_USE_NOTIFY
00027 namespace Mark3
00028 {
00033 class Notify : public BlockingObject
00034 {
00035 public:
00036     void* operator new(size_t sz, void* pv) { return (Notify*)pv; };
00037     ~Notify();
00038
00044     void Init(void);
00045
00055     void Signal(void);
00056
00066     void Wait(bool* pbFlag_);
00067
00068 #if KERNEL_USE_TIMEOUTS
00069     bool Wait(uint32_t u32WaitTimeMS_, bool* pbFlag_);
00082 #endif
00093     void WakeMe(Thread* p1ChosenOne_);
00094
00095 private:
00096     bool m_bPending;
00097 };
00098 } //namespace Mark3
00100 #endif
00101

```

20.95 /home/moslevin/projects/github/m3-repo/kernel/src/public/paniccodes.h File Reference

Defines the reason codes thrown when a kernel panic occurs.

20.95.1 Detailed Description

Defines the reason codes thrown when a kernel panic occurs.

Definition in file [paniccodes.h](#).

20.96 paniccodes.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00020 #pragma once
00021
00022 #define PANIC_ASSERT_FAILED (1)
00023 #define PANIC_LIST_UNLINK_FAILED (2)
00024 #define PANIC_STACK_SLACK_VIOLATED (3)
00025 #define PANIC_AUTO_HEAP_EXHAUSTED (4)
00026 #define PANIC_POWERMAN_EXHAUSTED (5)
00027 #define PANIC_NO_READY_THREADS (6)
00028 #define PANIC_RUNNING_THREAD_DESCOPED (7)
00029 #define PANIC_ACTIVE_SEMAPHORE_DESCOPED (8)
00030 #define PANIC_ACTIVE_MUTEX_DESCOPED (9)
00031 #define PANIC_ACTIVE_EVENTFLAG_DESCOPED (10)
00032 #define PANIC_ACTIVE_NOTIFY_DESCOPED (11)
00033 #define PANIC_ACTIVE_MAILBOX_DESCOPED (12)
00034 #define PANIC_ACTIVE_TIMER_DESCOPED (13)

```

20.97 /home/moslevin/projects/github/m3-repo/kernel/src/public/priomap.h File Reference

Priority map data structure.

```

#include "kerneltypes.h"
#include "mark3cfg.h"

```

Classes

- class [Mark3::PriorityMap](#)
The *PriorityMap* class.

Namespaces

- [Mark3](#)

Class providing the software-interrupt required for context-switching in the kernel.

20.97.1 Detailed Description

Priority map data structure.

Definition in file [priomap.h](#).

20.98 priomap.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00019 #pragma once
00020
00021 #include "kerneltypes.h"
00022 #include "mark3cfg.h"
00023
00024 //-----
00025 // Define the type used to store the priority map based on the word size of
00026 // the underlying host architecture.
00027 #if !defined(PORT_PRIO_MAP_WORD_SIZE)
00028 #error "undefined PORT_PRIO_MAP_WORD_SIZE"
00029 #endif
00030
00031 #define PRIO_MAP_WORD_TYPE K_WORD
00032
00033 // Size of the map index type in bits
00034 #define PRIO_MAP_BITS (8 * PORT_PRIO_MAP_WORD_SIZE)
00035
00036 // # of bits in an integer used to represent the number of bits in the map.
00037 // Used for bitshifting the bit index away from the map index.
00038 // i.e. 3 == 8 bits, 4 == 16 bits, 5 == 32 bits, etc...
00039 #define PRIO_MAP_WORD_SHIFT (2 + PORT_PRIO_MAP_WORD_SIZE)
00040
00041 // Bitmask used to separate out the priorities first-level bitmap from its
00042 // second-level map index for a given priority
00043 #define PRIO_MAP_BIT_MASK ((1 << PRIO_MAP_WORD_SHIFT) - 1)
00044
00045 // Get the priority bit for a given thread
00046 #define PRIO_BIT(x) ((x)&PRIO_MAP_BIT_MASK)
00047
00048 // Macro used to get the map index for a given priroity
00049 #define PRIO_MAP_WORD_INDEX(prio) ((prio) >> PRIO_MAP_WORD_SHIFT)
00050
00051 // Required size of the bitmap array in words
00052 #define PRIO_MAP_NUM_WORDS ((KERNEL_NUM_PRIORITIES + (PRIO_MAP_BITS - 1)) / (PRIO_MAP_BITS))
00053
00054 //-----
00055 #if (PRIO_MAP_NUM_WORDS == 1)
00056 // If there is only 1 word required to store the priority information, we don't
00057 // need an array, or a secondary bitmap.
00058 #define PRIO_MAP_MULTI_LEVEL (0)
00059 #else
00060 // An array of bitmaps are required, and a secondary index is required to
00061 // efficiently track which priority levels are active.
00062 #define PRIO_MAP_MULTI_LEVEL (1)
00063 #endif
00064 namespace Mark3
00065 {
00066 //-----

```

```

00070 class PriorityMap
00071 {
00072 public:
00073     PriorityMap();
00074
00075     void Set(PORT_PRIO_TYPE uXPrio_);
00076
00077     void Clear(PORT_PRIO_TYPE uXPrio_);
00078
00079     PORT_PRIO_TYPE HighestPriority(void);
00080
00081 private:
00082 #if PRIO_MAP_MULTI_LEVEL
00083     PRIO_MAP_WORD_TYPE m_auxPriorityMap[PRIO_MAP_NUM_WORDS];
00084     PRIO_MAP_WORD_TYPE m_uXPriorityMapL2;
00085 #else
00086     PRIO_MAP_WORD_TYPE m_uXPriorityMap;
00087 #endif
00088 };
00089 //namespace Mark3

```

20.99 /home/moslevin/projects/github/m3-repo/kernel/src/public/profile.h File Reference

High-precision profiling timers.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ll.h"

```

Classes

- class [Mark3::ProfileTimer](#)
Profiling timer.

Namespaces

- [Mark3](#)
Class providing the software-interrupt required for context-switching in the kernel.

20.99.1 Detailed Description

High-precision profiling timers.

Enables the profiling and instrumentation of performance-critical code. Multiple timers can be used simultaneously to enable system-wide performance metrics to be computed in a lightweight manner.

Usage:

```

ProfileTimer clMyTimer;
int i;

clMyTimer.Init();

// Profile the same block of code ten times
for (i = 0; i < 10; i++)
{
    clMyTimer.Start();
    ...
    //Block of code to profile
    ...
    clMyTimer.Stop();
}

// Get the average execution time of all iterations
u32AverageTimer = clMyTimer.GetAverage();

// Get the execution time from the last iteration
u32LastTimer = clMyTimer.GetCurrent();

```

Definition in file [profile.h](#).

20.100 profile.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00052 #pragma once
00053
00054 #include "kerneltypes.h"
00055 #include "mark3cfg.h"
00056 #include "ll.h"
00057
00058 #if KERNEL_USE_PROFILER
00059 namespace Mark3
00060 {
00069 class ProfileTimer
00070 {
00071 public:
00078     void Init();
00079
00086     void Start();
00087
00094     void Stop();
00095
00103     uint32_t GetAverage();
00104
00113     uint32_t GetCurrent();
00114
00115 private:
00126     uint32_t ComputeCurrentTicks(uint16_t u16Current_, uint32_t u32Epoch_);
00127
00128     uint32_t m_u32Cumulative;
00129     uint32_t m_u32CurrentIteration;
00130     uint16_t m_u16Initial;
00131     uint32_t m_u32InitialEpoch;
00132     uint16_t m_u16Iterations;
00133     bool m_bActive;
00134 };
00135 } //namespace Mark3
00136 #endif // KERNEL_USE_PROFILE

```

20.101 /home/moslevin/projects/github/m3-repo/kernel/src/public/quantum.h File Reference

Thread Quantum declarations for Round-Robin Scheduling.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "thread.h"
#include "timer.h"
#include "timerlist.h"
#include "timerscheduler.h"

```

Classes

- class [Mark3::Quantum](#)

Static-class used to implement [Thread](#) quantum functionality, which is a key part of round-robin scheduling.


```

00061
00069     bool AcquireReader(uint32_t u32TimeoutMs_);
00070 #endif
00071
00076     void ReleaseReader();
00077
00084     void AcquireWriter();
00085
00086 #if KERNEL_USE_TIMEOUTS
00087
00095     bool AcquireWriter(uint32_t u32TimeoutMs_);
00096 #endif
00097
00102     void ReleaseWriter();
00103
00104 private:
00105 #if KERNEL_USE_TIMEOUTS
00106
00112     bool AcquireReader_i(uint32_t u32TimeoutMs_);
00113 #else
00114
00117     void AcquireReader_i();
00118 #endif
00119
00120 #if KERNEL_USE_TIMEOUTS
00121
00127     bool AcquireWriter_i(uint32_t u32TimeoutMs_);
00128 #endif
00129
00130     Mutex m_clGlobalMutex;
00131     Mutex m_clReaderMutex;
00132     uint8_t m_u8ReadCount;
00133 };
00134
00135 } // namespace Mark3
00136 #endif

```

20.105 /home/moslevin/projects/github/m3-repo/kernel/src/public/scheduler.h File Reference

Thread scheduler function declarations.

```

#include "kerneltypes.h"
#include "thread.h"
#include "threadport.h"
#include "priomap.h"

```

Classes

- class [Mark3::Scheduler](#)

Priority-based round-robin [Thread](#) scheduling, using [ThreadLists](#) for housekeeping.

Namespaces

- [Mark3](#)

Class providing the software-interrupt required for context-switching in the kernel.

20.105.1 Detailed Description

Thread scheduler function declarations.

This scheduler implements a very flexible type of scheduling, which has become the defacto industry standard when it comes to real-time operating systems. This scheduling mechanism is referred to as priority round- robin.

From the name, there are two concepts involved here:

1) Priority scheduling:

Threads are each assigned a priority, and the thread with the highest priority which is ready to run gets to execute.

2) Round-robin scheduling:

Where there are multiple ready threads at the highest-priority level, each thread in that group gets to share time, ensuring that progress is made.

The scheduler uses an array of ThreadList objects to provide the necessary housekeeping required to keep track of threads at the various priorities. As a result, the scheduler contains one ThreadList per priority, with an additional list to manage the storage of threads which are in the "stopped" state (either have been stopped, or have not been started yet).

Definition in file [scheduler.h](#).

20.106 scheduler.h

```

00001  /*-----
00002
00003
00004
00005
00006
00007
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  ===== */
00046  #pragma once
00047
00048  #include "kerneltypes.h"
00049  #include "thread.h"
00050  #include "threadport.h"
00051  #include "priomap.h"
00052
00053  extern volatile Mark3::Thread* g_pclNext;
00054  extern Mark3::Thread* g_pclCurrent;
00055
00056  namespace Mark3
00057  {
00058  //-----
00063  class Scheduler
00064  {
00065  public:
00066
00072      static void Init();
00073
00081      static void Schedule();
00082
00090      static void Add(Thread* pclThread_);
00091
00100      static void Remove(Thread* pclThread_);
00101
00114      static bool SetScheduler(bool bEnable_);
00115
00123      static Thread* GetCurrentThread() { return g_pclCurrent; }
00132      static volatile Thread* GetNextThread() { return g_pclNext; }
00143      static ThreadList* GetThreadList(PORT_PRIO_TYPE uXPriority_) {

```

```

        return &m_aclPriorities[uXPriority_]; }
00152     static ThreadList* GetStopList() { return &m_clStopList; }
00161     static bool IsEnabled() { return m_bEnabled; }
00168     static void QueueScheduler() { m_bQueuedSchedule = true; }
00169 private:
00171     static bool m_bEnabled;
00172
00174     static bool m_bQueuedSchedule;
00175
00177     static ThreadList m_clStopList;
00178
00180     static ThreadList m_aclPriorities[
        KERNEL_NUM_PRIORITIES];
00181
00183     static PriorityMap m_clPrioMap;
00184 };
00185 } //namespace Mark3

```

20.107 /home/moslevin/projects/github/m3-repo/kernel/src/public/thread.h File Reference

Platform independent thread class declarations.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ll.h"
#include "threadlist.h"
#include "scheduler.h"
#include "threadport.h"
#include "quantum.h"
#include "autoalloc.h"
#include "priomap.h"

```

Classes

- class [Mark3::Thread](#)
Object providing fundamental multitasking support in the kernel.
- struct [Mark3::FakeThread_t](#)
If the kernel is set up to use an idle function instead of an idle thread, we use a placeholder data structure to "simulate" the effect of having an idle thread in the system.

Namespaces

- [Mark3](#)
Class providing the software-interrupt required for context-switching in the kernel.

20.107.1 Detailed Description

Platform independent thread class declarations.

Threads are an atomic unit of execution, and each instance of the thread class represents an instance of a program running on the processor. The Thread is the fundamental user-facing object in the kernel - it is what makes multiprocessing possible from application code.

In [Mark3](#), threads each have their own context - consisting of a stack, and all of the registers required to multiplex a processor between multiple threads.

The Thread class inherits directly from the LinkListNode class to facilitate efficient thread management using Double, or Double-Circular linked lists.

Definition in file [thread.h](#).


```

00232     void SetPriority(PORT_PRIO_TYPE uXPriority_);
00233
00243     void InheritPriority(PORT_PRIO_TYPE uXPriority_);
00244
00245     #if KERNEL_USE_DYNAMIC_THREADS
00246
00257         void Exit();
00258     #endif
00259
00260     #if KERNEL_USE_SLEEP
00261
00269         static void Sleep(uint32_t u32TimeMs_);
00270
00279         static void USleep(uint32_t u32TimeUs_);
00280     #endif
00281
00289         static void Yield(void);
00290
00298     void SetID(uint8_t u8ID_) { m_u8ThreadID = u8ID_; }
00306     uint8_t GetID() { return m_u8ThreadID; }
00319     uint16_t GetStackSlack();
00320
00321     #if KERNEL_USE_EVENTFLAG
00322
00329     uint16_t GetEventFlagMask() { return m_u16FlagMask; }
00334     void SetEventFlagMask(uint16_t u16Mask_) { m_u16FlagMask = u16Mask_; }
00340     void SetEventFlagMode(EventFlagOperation eMode_) {
00345         m_eFlagMode = eMode_; }
00346     EventFlagOperation GetEventFlagMode() { return
00347         m_eFlagMode; }
00348     #endif
00349
00348     #if KERNEL_USE_TIMEOUTS || KERNEL_USE_SLEEP
00349
00352         Timer* GetTimer();
00353     #endif
00354     #if KERNEL_USE_TIMEOUTS
00355
00363     void SetExpired(bool bExpired_);
00364
00371     bool GetExpired();
00372     #endif
00373
00374     #if KERNEL_USE_IDLE_FUNC
00375
00380     void InitIdle();
00381     #endif
00382
00383     #if KERNEL_USE_EXTENDED_CONTEXT
00384
00393     void* GetExtendedContext() { return m_pvExtendedContext; }
00394
00406     void SetExtendedContext(void* pvData_) {
00407         m_pvExtendedContext = pvData_; }
00408     #endif
00409
00415     ThreadState GetState() { return m_eState; }
00423     void SetState(ThreadState eState_) { m_eState = eState_; }
00424
00429     K_WORD* GetStack() { return m_pwStack; }
00430
00435     uint16_t GetStackSize() { return m_u16StackSize; }
00436
00437     friend class ThreadPort;
00438 private:
00446     static void ContextSwitchSWI(void);
00447
00453     void SetPriorityBase(PORT_PRIO_TYPE uXPriority_);
00454
00456     K_WORD* m_pwStackTop;
00457
00459     K_WORD* m_pwStack;
00460
00462     uint8_t m_u8ThreadID;
00463
00465     PORT_PRIO_TYPE m_uXPriority;
00466
00468     PORT_PRIO_TYPE m_uXCurPriority;
00469
00471     ThreadState m_eState;
00472
00473     #if KERNEL_USE_EXTENDED_CONTEXT
00474         void* m_pvExtendedContext;
00475     #endif
00476
00477
00478     #if KERNEL_USE_THREADNAME
00479         const char* m_szName;

```

```

00481 #endif
00482
00484     uint16_t m_ul6StackSize;
00485
00487     ThreadList* m_pclCurrent;
00488
00490     ThreadList* m_pclOwner;
00491
00493     ThreadEntryFunc m_pfEntryPoint;
00494
00496     void* m_pvArg;
00497
00498 #if KERNEL_USE_QUANTUM
00499     uint16_t m_ul6Quantum;
00501 #endif
00502
00503 #if KERNEL_USE_EVENTFLAG
00504     uint16_t m_ul6FlagMask;
00506
00508     EventFlagOperation m_eFlagMode;
00509 #endif
00510
00511 #if KERNEL_USE_TIMEOUTS || KERNEL_USE_SLEEP
00512     Timer m_clTimer;
00514 #endif
00515
00516 #if KERNEL_USE_TIMEOUTS
00517     bool m_bExpired;
00519 #endif
00520 };
00521
00522 #if KERNEL_USE_IDLE_FUNC
00523 //-----
00535 typedef struct {
00536     LinkListNode* next;
00537     LinkListNode* prev;
00538
00540     K_WORD* m_pwStackTop;
00541
00543     K_WORD* m_pwStack;
00544
00546     uint8_t m_u8ThreadID;
00547
00549     PORT_PRIO_TYPE m_uXPriority;
00550
00552     PORT_PRIO_TYPE m_uXCurPriority;
00553
00555     ThreadState m_eState;
00556
00557 #if KERNEL_USE_EXTENDED_CONTEXT
00558     void* m_pvExtendedContext;
00560 #endif
00561
00562 #if KERNEL_USE_THREADNAME
00563     const char* m_szName;
00565 #endif
00566
00567 } FakeThread_t;
00568 #endif //KERNEL_USE_IDLE_FUNC
00569 } //namespace Mark3

```

20.109 /home/moslevin/projects/github/m3-repo/kernel/src/public/threadlist.h File Reference

Thread linked-list declarations.

```

#include "kerneltypes.h"
#include "priomap.h"
#include "ll.h"

```

Classes

- class [Mark3::ThreadList](#)

This class is used for building thread-management facilities, such as schedulers, and blocking objects.

Namespaces

- [Mark3](#)

Class providing the software-interrupt required for context-switching in the kernel.

20.109.1 Detailed Description

Thread linked-list declarations.

Definition in file [threadlist.h](#).

20.110 threadlist.h

```

00001 /*=====
00002
00003 |-----|-----|-----|-----|-----|-----|
00004 | \ / | \ / | \ / | \ / | \ / | \ / | \ / |
00005 |  V  |  V  |  V  |  V  |  V  |  V  |  V  |
00006 | / \  | / \  | / \  | / \  | / \  | / \  |
00007 |-----|-----|-----|-----|-----|
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00022 #pragma once
00023
00024 #include "kerneltypes.h"
00025 #include "priomap.h"
00026 #include "ll.h"
00027
00028 namespace Mark3
00029 {
00030     class Thread;
00031
00036     class ThreadList : public CircularLinkedList
00037     {
00038     public:
00039         void* operator new(size_t sz, void* pv) { return (ThreadList*)pv; };
00045         ThreadList() : m_uXPriority(0), m_pclMap(NULL)
00046         { }
00047
00055         void SetPriority(PORT_PRIO_TYPE uXPriority_);
00056
00066         void SetMapPointer(PriorityMap* pclMap_);
00067
00075         void Add(LinkListNode* node_);
00076
00088         void Add(LinkListNode* node_, PriorityMap* pclMap_,
00089 PORT_PRIO_TYPE uXPriority_);
00089
00098         void AddPriority(LinkListNode* node_);
00099
00107         void Remove(LinkListNode* node_);
00108
00116         Thread* HighestWaiter();
00117
00118     private:
00120         PORT_PRIO_TYPE m_uXPriority;
00121
00123         PriorityMap* m_pclMap;
00124     };
00125 } //namespace Mark3

```

20.111 /home/moslevin/projects/github/m3-repo/kernel/src/public/timer.h File Reference

Timer object declarations.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ll.h"

```

Classes

- class [Mark3::Timer](#)
Kernel-managed software timers.

Namespaces

- [Mark3](#)
Class providing the software-interrupt required for context-switching in the kernel.

Macros

- `#define` [TIMERLIST_FLAG_ONE_SHOT](#) (0x01)
Timer is one-shot.
- `#define` [TIMERLIST_FLAG_ACTIVE](#) (0x02)
Timer is currently active.
- `#define` [TIMERLIST_FLAG_CALLBACK](#) (0x04)
Timer is pending a callback.
- `#define` [TIMERLIST_FLAG_EXPIRED](#) (0x08)
Timer is actually expired.
- `#define` [MAX_TIMER_TICKS](#) (0x7FFFFFFF)
Maximum value to set.
- `#define` [MIN_TICKS](#) (3)
The minimum tick value to set.

Typedefs

- using [Mark3::TimerCallback](#) = void(*)(Thread *pclOwner_, void *pvData_)
This type defines the callback function type for timer events.

20.111.1 Detailed Description

Timer object declarations.

Definition in file [timer.h](#).

20.111.2 Macro Definition Documentation

20.111.2.1 [TIMERLIST_FLAG_EXPIRED](#)

```
#define TIMERLIST_FLAG_EXPIRED (0x08)
```

Timer is actually expired.

Definition at line 36 of file [timer.h](#).

20.112 timer.h

```

00001  /*=====
00002
00003  _____
00004  |   /   \   |   /   \   |   /   \   |   /   \   |
00005  |  /     \  |  /     \  |  /     \  |  /     \  |
00006  | /       \ | /       \ | /       \ | /       \ |
00007  |_____|   |_____|   |_____|   |_____|   |
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  ===== */
00021  #pragma once
00022  #include "kerneltypes.h"
00023  #include "mark3cfg.h"
00024
00025  #include "ll.h"
00026
00027  #if KERNEL_USE_TIMERS
00028  namespace Mark3
00029  {
00030  class Thread;
00031
00032  //-----
00033  #define TIMERLIST_FLAG_ONE_SHOT (0x01)
00034  #define TIMERLIST_FLAG_ACTIVE (0x02)
00035  #define TIMERLIST_FLAG_CALLBACK (0x04)
00036  #define TIMERLIST_FLAG_EXPIRED (0x08)
00037
00038  //-----
00039  #define TIMER_INVALID_COOKIE (0x3C)
00040  #define TIMER_INIT_COOKIE (0xC3)
00041
00042  //-----
00043  #define MAX_TIMER_TICKS (0x7FFFFFFF)
00044  #define TIMER_TICKS_INVALID (0x80000000)
00045  //-----
00046  #if KERNEL_TIMERS_TICKLESS
00047
00048  //-----
00049  /*
00050      Ugly macros to support a wide resolution of delays.
00051      Given a 16-bit timer @ 16MHz & 256 cycle prescaler, this gives ul6...
00052      Max time, SECONDS_TO_TICKS: 68719s
00053      Max time, MSECONDS_TO_TICKS: 6871.9s
00054      Max time, USECONDS_TO_TICKS: 6.8719s
00055
00056      ...With a 16us tick resolution.
00057
00058      Depending on the system frequency and timer resolution, you may want to
00059      customize these values to suit your system more appropriately.
00060  */
00061  //-----
00062  #define SECONDS_TO_TICKS(x) (((uint32_t)x) * PORT_TIMER_FREQ))
00063  #define MSECONDS_TO_TICKS(x) (((((uint32_t)x) * (PORT_TIMER_FREQ / 100)) + 5) / 10))
00064  #define USECONDS_TO_TICKS(x) (((((uint32_t)x) * PORT_TIMER_FREQ) + 50000) / 1000000))
00065
00066  //-----
00067  #define MIN_TICKS (3)
00068  //-----
00069
00070  #else
00071
00072  //-----
00073  // add time because we don't know how far in an epoch we are when a call is made.
00074  #define SECONDS_TO_TICKS(x) (((uint32_t)(x)*1000) + 1)
00075  #define MSECONDS_TO_TICKS(x) ((uint32_t)(x + 1))
00076  #define USECONDS_TO_TICKS(x) (((uint32_t)(x + 999)) / 1000)
00077
00078  //-----
00079  #define MIN_TICKS (1)
00080  //-----
00081
00082  #endif // KERNEL_TIMERS_TICKLESS
00083  //-----
00094  using TimerCallback = void (*)(Thread* pOwner_, void* pvData_);
00095
00096  //-----
00097  class TimerList;
00098  class TimerScheduler;
00099  class Quantum;
00100
00101  //-----

```



```

00111 class Timer : public LinkListNode
00112 {
00113 public:
00114     void* operator new(size_t sz, void* pv) { return (Timer*)pv; }
00115     ~Timer() {}
00116
00123     Timer();
00124
00130     void Init();
00131
00143     void Start(bool bRepeat_, uint32_t u32IntervalMs_, TimerCallback pfCallback_, void*
pvData_);
00144
00158     void
00159     Start(bool bRepeat_, uint32_t u32IntervalMs_, uint32_t u32ToleranceMs_,
TimerCallback pfCallback_, void* pvData_);
00160
00169     void Start();
00170
00177     void Stop();
00178
00188     void SetFlags(uint8_t u8Flags_) { m_u8Flags = u8Flags_; }
00196     void SetCallback(TimerCallback pfCallback_) {
m_pfCallback = pfCallback_; }
00204     void SetData(void* pvData_) { m_pvData = pvData_; }
00213     void SetOwner(Thread* pclOwner_) { m_pclOwner = pclOwner_; }
00221     void SetIntervalTicks(uint32_t u32Ticks_);
00222
00230     void SetIntervalSeconds(uint32_t u32Seconds_);
00231
00239     uint32_t GetInterval() { return m_u32Interval; }
00247     void SetIntervalMSeconds(uint32_t u32MSeconds_);
00248
00256     void SetIntervalUSeconds(uint32_t u32USeconds_);
00257
00266     void SetTolerance(uint32_t u32Ticks_);
00267
00268 private:
00269     friend class TimerList;
00270
00271 #if KERNEL_EXTRA_CHECKS
00272
00275     void SetInitialized() { m_u8Initialized = TIMER_INIT_COOKIE; }
00276
00281     bool IsInitialized(void) { return (m_u8Initialized == TIMER_INIT_COOKIE); }
00282
00284     uint8_t m_u8Initialized;
00285 #endif
00286
00288     uint8_t m_u8Flags;
00289
00291     TimerCallback m_pfCallback;
00292
00294     uint32_t m_u32Interval;
00295
00297     uint32_t m_u32TimeLeft;
00298
00300     uint32_t m_u32TimerTolerance;
00301
00303     Thread* m_pclOwner;
00304
00306     void* m_pvData;
00307 };
00308 } //namespace Mark3
00309 #endif // KERNEL_USE_TIMERS
00310

```

20.113 /home/moslevin/projects/github/m3-repo/kernel/src/public/timerlist.h File Reference

Timer list declarations.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "mutex.h"

```

Classes

- class [Mark3::TimerList](#)
TimerList class - a doubly-linked-list of timer objects.

Namespaces

- [Mark3](#)
Class providing the software-interrupt required for context-switching in the kernel.

20.113.1 Detailed Description

Timer list declarations.

These classes implements a linked list of timer objects attached to the global kernel timer scheduler.

Definition in file [timerlist.h](#).

20.114 timerlist.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00024 #pragma once
00025
00026 #include "kerneltypes.h"
00027 #include "mark3cfg.h"
00028
00029 #include "mutex.h"
00030 #if KERNEL_USE_TIMERS
00031 namespace Mark3
00032 {
00033 class Timer;
00034
00035 //-----
00039 class TimerList : public DoubleLinkedList
00040 {
00041 public:
00048     void Init();
00049
00057     void Add(Timer* pclListNode_);
00058
00066     void Remove(Timer* pclLinkListNode_);
00067
00074     void Process();
00075
00076 private:
00078     uint32_t m_u32NextWakeup;
00079
00081     bool m_bTimerActive;
00082
00083 #if KERNEL_TIMERS_THREADED
00084     Mutex m_clMutex;
00086 #endif
00087
00088 };
00089 } //namespace Mark3
00090 #endif // KERNEL_USE_TIMERS

```

20.115 /home/moslevin/projects/github/m3-repo/kernel/src/public/timerscheduler.h File Reference

Timer scheduler declarations.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ll.h"
#include "timer.h"
#include "timerlist.h"
```

Classes

- class [Mark3::TimerScheduler](#)
"Static" Class used to interface a global [TimerList](#) with the rest of the kernel.

Namespaces

- [Mark3](#)
Class providing the software-interrupt required for context-switching in the kernel.

20.115.1 Detailed Description

Timer scheduler declarations.

Definition in file [timerscheduler.h](#).

20.116 timerscheduler.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00021 #pragma once
00022
00023 #include "kerneltypes.h"
00024 #include "mark3cfg.h"
00025
00026 #include "ll.h"
00027 #include "timer.h"
00028 #include "timerlist.h"
00029
00030 #if KERNEL_USE_TIMERS
00031 namespace Mark3
00032 {
00033 //-----
00038 class TimerScheduler
00039 {
00040 public:
```



```

00049
00054     static uint16_t Increment(void) { return m_ul6SyncNumber++; }
00063     static void Write(uint16_t* pu16Data_, uint16_t ul6Size_);
00064
00073     static void SetCallback(TraceBufferCallback_t pfCallback_) { m_pfCallback = pfCallback_; }
00074 private:
00075     static TraceBufferCallback_t m_pfCallback;
00076     static uint16_t             m_ul6SyncNumber;
00077     static uint16_t             m_ul6Index;
00078     static uint16_t             m_aul6Buffer[(TRACE_BUFFER_SIZE / sizeof(uint16_t))];
00079 };
00080
00081 } //namespace Mark3
00082 #endif // KERNEL_USE_DEBUG

```

20.119 /home/moslevin/projects/github/m3-repo/kernel/src/quantum.cpp File Reference

Thread Quantum Implementation for Round-Robin Scheduling.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "thread.h"
#include "timerlist.h"
#include "quantum.h"
#include "kernelaware.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"

```

20.119.1 Detailed Description

Thread Quantum Implementation for Round-Robin Scheduling.

Definition in file [quantum.cpp](#).

20.120 quantum.cpp

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00022 #include "kerneltypes.h"
00023 #include "mark3cfg.h"
00024
00025 #include "thread.h"
00026 #include "timerlist.h"
00027 #include "quantum.h"
00028 #include "kernelaware.h"
00029
00030 #define _CAN_HAS_DEBUG
00031 //--[Autogenerated - Do Not Modify]-----
00032 #include "dbg_file_list.h"
00033 #include "buffalogger.h"
00034 #if defined(DBG_FILE)

```

```

00035 #error "Debug logging file token already defined!  Bailing."
00036 #else
00037 #define DBG_FILE_DBG___KERNEL_QUANTUM_CPP
00038 #endif
00039 //--[End Autogenerated content]-----
00040 #include "kerneldebug.h"
00041
00042 #if KERNEL_USE_QUANTUM
00043 namespace Mark3
00044 {
00045     namespace
00046     {
00047         volatile bool bAddQuantumTimer; // Indicates that a timer add is pending
00048     } // anonymous namespace
00049
00050 #if KERNEL_TIMERS_THREADED
00051 Thread* Quantum::m_pclTimerThread;
00052 #endif
00053
00054 Timer Quantum::m_clQuantumTimer;
00055 bool Quantum::m_bActive;
00056 bool Quantum::m_bInTimer;
00057
00058 //-----
00059 void Quantum::SetTimer(Thread* pclThread_)
00060 {
00061     auto lQuantumCallback = [](Thread* pclThread_, void* /*pvData_*/) {
00062         if (pclThread_>GetCurrent()->GetHead() != pclThread_>GetCurrent()->GetTail()) {
00063             bAddQuantumTimer = true;
00064             pclThread_>GetCurrent()->PivotForward();
00065         }
00066     };
00067
00068     m_clQuantumTimer.SetIntervalMSeconds(pclThread_>GetQuantum());
00069     m_clQuantumTimer.SetFlags(TIMERLIST_FLAG_ONE_SHOT);
00070     m_clQuantumTimer.SetData(NULL);
00071     m_clQuantumTimer.SetCallback(lQuantumCallback);
00072     m_clQuantumTimer.SetOwner(pclThread_);
00073 }
00074
00075 //-----
00076 void Quantum::AddThread(Thread* pclThread_)
00077 {
00078     if (m_bActive
00079 #if KERNEL_USE_IDLE_FUNC
00080         || (pclThread_ == Kernel::GetIdleThread())
00081 #endif
00082     ) {
00083         return;
00084     }
00085
00086     // If this is called from the timer callback, queue a timer add...
00087     if (m_bInTimer) {
00088         bAddQuantumTimer = true;
00089         return;
00090     }
00091
00092     // If this isn't the only thread in the list.
00093     if (pclThread_>GetCurrent()->GetHead() != pclThread_>GetCurrent()->GetTail()) {
00094 #if KERNEL_EXTRA_CHECKS
00095         m_clQuantumTimer.Init();
00096 #endif
00097         Quantum::SetTimer(pclThread_);
00098         TimerScheduler::Add(&m_clQuantumTimer);
00099         m_bActive = true;
00100     }
00101 }
00102
00103 //-----
00104 void Quantum::RemoveThread(void)
00105 {
00106     if (!m_bActive) {
00107         return;
00108     }
00109
00110     // Cancel the current timer
00111     TimerScheduler::Remove(&m_clQuantumTimer);
00112     m_bActive = false;
00113 }
00114
00115 //-----
00116 void Quantum::UpdateTimer(void)
00117 {
00118     // If we have to re-add the quantum timer (more than 2 threads at the
00119     // high-priority level...)
00120     if (bAddQuantumTimer) {
00121         // Trigger a thread yield - this will also re-schedule the

```

```
00122 // thread *and* reset the round-robin scheduler.
00123 Thread::Yield();
00124 bAddQuantumTimer = false;
00125 }
00126 }
00127
00128 //-----
00129 #if KERNEL_TIMERS_THREADED
00130 void Quantum::SetTimerThread(Thread* pClThread_)
00131 {
00132     m_pClTimerThread = pClThread_;
00133 }
00134
00135 Thread* Quantum::GetTimerThread()
00136 {
00137     return m_pClTimerThread;
00138 }
00139
00140 #endif // KERNEL_TIMERS_THREADED
00141 } //namespace Mark3
00142 #endif // KERNEL_USE_QUANTUM
```

20.121 /home/moslevin/projects/github/m3-repo/kernel/src/readerwriter.cpp File Reference

Reader-writer lock implementation.

```
#include "mark3.h"
#include "kerneldebug.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
```

20.121.1 Detailed Description

Reader-writer lock implementation.

Definition in file [readerwriter.cpp](#).

20.122 readerwriter.cpp

```
00001 /*=====
00002
00003 |_____|_____|_____|_____|_____|_____|_____|_____|_____|_____|
00004 |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |
00005 |   /   \   |   /   \   |   /   \   |   /   \   |   /   \   |   /   \   |   /   \   |
00006 |_/___\___|_/___\___|_/___\___|_/___\___|_/___\___|_/___\___|_/___\___|_/___\___|
00007 |_____||_____||_____||_____||_____||_____||_____||_____||_____||_____||
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*
00020 #include "mark3.h"
00021 #include "kerneldebug.h"
00022
00023 #define _CAN_HAS_DEBUG
00024 //--[Autogenerated - Do Not Modify]-----
00025 #include "dbg_file_list.h"
00026 #include "buffalogger.h"
00027 #if defined(DBG_FILE)
00028 # error "Debug logging file token already defined! Bailing."
00029 #else
00030 # define DBG_FILE _DBG__KERNEL_READERWRITER_CPP
00031 #endif
```

```

00032 //--[End Autogenerated content]-----
00033
00034 #if KERNEL_USE_READERWRITER
00035 namespace Mark3 {
00036
00037
00038 //-----
00039 void ReaderWriterLock::Init()
00040 {
00041     m_u8ReadCount = 0;
00042     m_clGlobalMutex.Init();
00043     m_clReaderMutex.Init();
00044 }
00045
00046 //-----
00047 void ReaderWriterLock::AcquireReader()
00048 {
00049     #if KERNEL_USE_TIMEOUTS
00050         AcquireReader_i(0);
00051     #else
00052         AcquireReader_i();
00053     #endif
00054 }
00055
00056 //-----
00057 #if KERNEL_USE_TIMEOUTS
00058 bool ReaderWriterLock::AcquireReader(uint32_t u32TimeoutMs_)
00059 {
00060     return AcquireReader_i(u32TimeoutMs_);
00061 }
00062 #endif
00063
00064 //-----
00065 void ReaderWriterLock::ReleaseReader()
00066 {
00067     m_clReaderMutex.Claim();
00068     m_u8ReadCount--;
00069     if (m_u8ReadCount == 0) {
00070         m_clGlobalMutex.Release();
00071     }
00072     m_clReaderMutex.Release();
00073 }
00074
00075 //-----
00076 void ReaderWriterLock::AcquireWriter()
00077 {
00078     #if KERNEL_USE_TIMEOUTS
00079         AcquireWriter_i(0);
00080     #else
00081         AcquireWriter_i();
00082     #endif
00083 }
00084
00085 //-----
00086 #if KERNEL_USE_TIMEOUTS
00087 bool ReaderWriterLock::AcquireWriter(uint32_t u32TimeoutMs_)
00088 {
00089     return AcquireWriter_i(u32TimeoutMs_);
00090 }
00091 #endif
00092
00093 //-----
00094 void ReaderWriterLock::ReleaseWriter()
00095 {
00096     m_clGlobalMutex.Release();
00097 }
00098
00099 //-----
00100 #if KERNEL_USE_TIMEOUTS
00101 bool ReaderWriterLock::AcquireReader_i(uint32_t u32TimeoutMs_)
00102 {
00103     auto rc = true;
00104     if (!m_clReaderMutex.Claim(u32TimeoutMs_)) {
00105         return false;
00106     }
00107
00108     m_u8ReadCount++;
00109     if (m_u8ReadCount == 1) {
00110         rc = m_clGlobalMutex.Claim(u32TimeoutMs_);
00111     }
00112
00113     m_clReaderMutex.Release();
00114     return rc;
00115 }
00116 #else
00117 void ReaderWriterLock::AcquireReader_i()
00118 {

```



```

00119     m_clReaderMutex.Claim();
00120
00121     m_u8ReadCount++;
00122     if (m_u8ReadCount == 1) {
00123         m_clGlobalMutex.Claim();
00124     }
00125
00126     m_clReaderMutex.Release();
00127 }
00128 #endif
00129
00130 //-----
00131 #if KERNEL_USE_TIMEOUTS
00132 bool ReaderWriterLock::AcquireWriter_i(uint32_t u32TimeoutMs_)
00133 {
00134     return m_clGlobalMutex.Claim(u32TimeoutMs_);
00135 }
00136 #else
00137 void ReaderWriterLock::AcquireWriter_i()
00138 {
00139     m_clGlobalmutex.Claim();
00140 }
00141 #endif
00142 } // namespace Mark3
00143
00144 #endif // KERNEL_USE_READERWRITER
00145

```

20.123 /home/moslevin/projects/github/m3-repo/kernel/src/scheduler.cpp File Reference

Strict-Priority + Round-Robin thread scheduler implementation.

```

#include "kerneltypes.h"
#include "ll.h"
#include "scheduler.h"
#include "thread.h"
#include "threadport.h"
#include "kernel.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"

```

Namespaces

- [Mark3](#)

Class providing the software-interrupt required for context-switching in the kernel.

20.123.1 Detailed Description

Strict-Priority + Round-Robin thread scheduler implementation.

Definition in file [scheduler.cpp](#).


```

00089 }
00090
00091 //-----
00092 void Scheduler::Remove(Thread* pclThread_)
00093 {
00094     m_aclPriorities[pclThread_->GetPriority()].Remove(pclThread_);
00095 }
00096
00097 //-----
00098 bool Scheduler::SetScheduler(bool bEnable_)
00099 {
00100     bool bRet;
00101     CS_ENTER();
00102     bRet = m_bEnabled;
00103     m_bEnabled = bEnable_;
00104     // If there was a queued scheduler event, dequeue and trigger an
00105     // immediate Yield
00106     if (m_bEnabled && m_bQueuedSchedule) {
00107         m_bQueuedSchedule = false;
00108         Thread::Yield();
00109     }
00110     CS_EXIT();
00111     return bRet;
00112 }
00113 } //namespace Mark3

```

20.125 /home/moslevin/projects/github/m3-repo/kernel/src/thread.cpp File Reference

Platform-Independent thread class Definition.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "thread.h"
#include "scheduler.h"
#include "kernelswi.h"
#include "timerlist.h"
#include "ksemaphore.h"
#include "quantum.h"
#include "kernel.h"
#include "priomap.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"

```

Namespaces

- [Mark3](#)

Class providing the software-interrupt required for context-switching in the kernel.

20.125.1 Detailed Description

Platform-Independent thread class Definition.

Definition in file [thread.cpp](#).


```

00091     KERNEL_TRACE_1("Thread Id: %d", m_u8ThreadID);
00092     KERNEL_TRACE_1("Entrypoint: %x", static_cast<K_ADDR>(pfEntryPoint_));
00093
00094     // Initialize the thread parameters to their initial values.
00095     m_pwStack = pwStack_;
00096     m_pwStackTop = TOP_OF_STACK(pwStack_, ul6StackSize_);
00097
00098     m_ul6StackSize = ul6StackSize_;
00099
00100 #if KERNEL_USE_QUANTUM
00101     m_ul6Quantum = THREAD_QUANTUM_DEFAULT;
00102 #endif
00103
00104     m_uXPriority = uXPriority_;
00105     m_uXCurPriority = m_uXPriority;
00106     m_pfEntryPoint = pfEntryPoint_;
00107     m_pvArg = pvArg_;
00108
00109 #if KERNEL_USE_THREADNAME
00110     m_szName = NULL;
00111 #endif
00112 #if KERNEL_USE_TIMERS
00113     m_clTimer.Init();
00114 #endif
00115
00116     // Call CPU-specific stack initialization
00117     ThreadPort::InitStack(this);
00118
00119     // Add to the global "stop" list.
00120     CS_ENTER();
00121     m_pclOwner = Scheduler::GetThreadList(
00122         m_uXPriority);
00123     m_pclCurrent = Scheduler::GetStopList();
00124     m_eState = ThreadState::Stop;
00125     m_pclCurrent->Add(this);
00126     CS_EXIT();
00127
00128 #if KERNEL_USE_THREAD_CALLOUTS
00129     ThreadCreateCallout pfCallout = Kernel::GetThreadCreateCallout();
00130     if (pfCallout != nullptr) {
00131         pfCallout(this);
00132     }
00133 #endif
00134
00135 #if KERNEL_USE_AUTO_ALLOC
00136 //-----
00137 Thread* Thread::Init(uint16_t ul6StackSize_, PORT_PRIO_TYPE uXPriority_,
00138     ThreadEntryFunc pfEntryPoint_, void* pvArg_)
00139 {
00140     auto* pclNew = AutoAlloc::NewThread();
00141     auto* pwStack = static_cast<K_WORD*>(AutoAlloc::NewRawData(ul6StackSize_));
00142     pclNew->Init(pwStack, ul6StackSize_, uXPriority_, pfEntryPoint_, pvArg_);
00143     return pclNew;
00144 }
00145 #endif
00146 //-----
00147 void Thread::Start(void)
00148 {
00149 #if KERNEL_EXTRA_CHECKS
00150     KERNEL_ASSERT(IsInitialized());
00151 #endif
00152
00153     // Remove the thread from the scheduler's "stopped" list, and add it
00154     // to the scheduler's ready list at the proper priority.
00155     KERNEL_TRACE_1("Starting Thread %d", m_u8ThreadID);
00156
00157     CS_ENTER();
00158     Scheduler::GetStopList()->Remove(this);
00159     Scheduler::Add(this);
00160     m_pclOwner = Scheduler::GetThreadList(
00161         m_uXPriority);
00162     m_pclCurrent = m_pclOwner;
00163     m_eState = ThreadState::Ready;
00164
00165 #if KERNEL_USE_QUANTUM
00166     if (Kernel::IsStarted()) {
00167         if (GetCurPriority() >= Scheduler::GetCurrentThread()->
00168             GetCurPriority()) {
00169             // Deal with the thread Quantum
00170             #if KERNEL_TIMERS_THREADED
00171                 if (Quantum::GetTimerThread() != this) {
00172                     Quantum::RemoveThread();
00173                     Quantum::AddThread(this);
00174                 }
00175             #endif
00176         }
00177     }
00178 #endif

```

```

00174         }
00175 #endif
00176     }
00177 }
00178 #endif
00179
00180     if (Kernel::IsStarted()) {
00181         if (GetCurPriority() >= Scheduler::GetCurrentThread()->
00182             GetCurPriority()) {
00183             Thread::Yield();
00184         }
00185     }
00186     CS_EXIT();
00187 }
00188 //-----
00189 void Thread::Stop()
00190 {
00191     #if KERNEL_EXTRA_CHECKS
00192         KERNEL_ASSERT(IsInitialized());
00193     #endif
00194
00195     auto bReschedule = false;
00196     if (m_eState == ThreadState::Stop) {
00197         return;
00198     }
00199
00200     CS_ENTER();
00201
00202     // If a thread is attempting to stop itself, ensure we call the scheduler
00203     if (this == Scheduler::GetCurrentThread()) {
00204         bReschedule = true;
00205     }
00206
00207     // Add this thread to the stop-list (removing it from active scheduling)
00208     // Remove the thread from scheduling
00209     if (m_eState == ThreadState::Ready) {
00210         Scheduler::Remove(this);
00211     } else if (m_eState == ThreadState::Blocked) {
00212         m_pclCurrent->Remove(this);
00213     }
00214
00215     m_pclOwner = Scheduler::GetStopList();
00216     m_pclCurrent = m_pclOwner;
00217     m_pclOwner->Add(this);
00218     m_eState = ThreadState::Stop;
00219
00220     #if KERNEL_USE_TIMERS
00221         // Just to be safe - attempt to remove the thread's timer
00222         // from the timer-scheduler (does no harm if it isn't
00223         // in the timer-list)
00224         TimerScheduler::Remove(&m_clTimer);
00225     #endif
00226
00227     CS_EXIT();
00228
00229     if (bReschedule) {
00230         Thread::Yield();
00231     }
00232 }
00233
00234 #if KERNEL_USE_DYNAMIC_THREADS
00235 //-----
00236 void Thread::Exit()
00237 {
00238     #if KERNEL_EXTRA_CHECKS
00239         KERNEL_ASSERT(IsInitialized());
00240     #endif
00241     bool bReschedule = false;
00242
00243     KERNEL_TRACE_1("Exit Thread %d", m_u8ThreadID);
00244     if (m_eState == ThreadState::Exit) {
00245         return;
00246     }
00247
00248     CS_ENTER();
00249
00250     // If this thread is the actively-running thread, make sure we run the
00251     // scheduler again.
00252     if (this == Scheduler::GetCurrentThread()) {
00253         bReschedule = true;
00254     }
00255
00256     // Remove the thread from scheduling
00257     if (m_eState == ThreadState::Ready) {
00258         Scheduler::Remove(this);
00259     } else if ((m_eState == ThreadState::Blocked) || (m_eState == ThreadState::Stop)) {

```

```

00260         m_pclCurrent->Remove(this);
00261     }
00262
00263     m_pclCurrent = 0;
00264     m_pclOwner   = 0;
00265     m_eState     = ThreadState::Exit;
00266
00267     // We've removed the thread from scheduling, but interrupts might
00268     // trigger checks against this thread's currently priority before
00269     // we get around to scheduling new threads. As a result, set the
00270     // priority to idle to ensure that we always wind up scheduling
00271     // new threads.
00272     m_uXCurPriority = 0;
00273     m_uXPriority     = 0;
00274
00275     #if KERNEL_USE_TIMERS
00276     // Just to be safe - attempt to remove the thread's timer
00277     // from the timer-scheduler (does no harm if it isn't
00278     // in the timer-list)
00279     TimerScheduler::Remove(&m_clTimer);
00280     #endif
00281     CS_EXIT();
00282
00283     #if KERNEL_USE_THREAD_CALLOUTS
00284     ThreadExitCallout pfCallout = Kernel::GetThreadExitCallout();
00285     if (pfCallout != nullptr) {
00286         pfCallout(this);
00287     }
00288     #endif
00289
00290     if (bReschedule) {
00291         // Choose a new "next" thread if we must
00292         Thread::Yield();
00293     }
00294 }
00295 #endif
00296
00297 #if KERNEL_USE_SLEEP
00298 //-----
00299 void Thread::Sleep(uint32_t u32TimeMs_)
00300 {
00301     Semaphore clSemaphore;
00302     auto* pclTimer = g_pclCurrent->GetTimer();
00303     auto lTimerCallback = [](Thread* /*pclOwner*/, void* pvData_) {
00304         auto* pclSemaphore = static_cast<Semaphore*>(pvData_);
00305         pclSemaphore->Post();
00306     };
00307
00308     // Create a semaphore that this thread will block on
00309     clSemaphore.Init(0, 1);
00310
00311     // Create a one-shot timer that will call a callback that posts the
00312     // semaphore, waking our thread.
00313     pclTimer->Init();
00314     pclTimer->SetIntervalMSeconds(u32TimeMs_);
00315     pclTimer->SetCallback(lTimerCallback);
00316     pclTimer->SetData((void*)&clSemaphore);
00317     pclTimer->SetFlags(TIMERLIST_FLAG_ONE_SHOT);
00318
00319     // Add the new timer to the timer scheduler, and block the thread
00320     TimerScheduler::Add(pclTimer);
00321     clSemaphore.Pend();
00322 }
00323
00324 //-----
00325 void Thread::USleep(uint32_t u32TimeUs_)
00326 {
00327     Semaphore clSemaphore;
00328     auto* pclTimer = g_pclCurrent->GetTimer();
00329     auto lTimerCallback = [](Thread* /*pclOwner*/, void* pvData_) {
00330         auto* pclSemaphore = static_cast<Semaphore*>(pvData_);
00331         pclSemaphore->Post();
00332     };
00333
00334     // Create a semaphore that this thread will block on
00335     clSemaphore.Init(0, 1);
00336
00337     // Create a one-shot timer that will call a callback that posts the
00338     // semaphore, waking our thread.
00339     pclTimer->Init();
00340     pclTimer->SetIntervalUSeconds(u32TimeUs_);
00341     pclTimer->SetCallback(lTimerCallback);
00342     pclTimer->SetData((void*)&clSemaphore);
00343     pclTimer->SetFlags(TIMERLIST_FLAG_ONE_SHOT);
00344
00345     // Add the new timer to the timer scheduler, and block the thread
00346     TimerScheduler::Add(pclTimer);

```

```

00347     clSemaphore.Pend();
00348 }
00349 #endif // KERNEL_USE_SLEEP
00350
00351 //-----
00352 uint16_t Thread::GetStackSlack()
00353 {
00354     #if KERNEL_EXTRA_CHECKS
00355         KERNEL_ASSERT(IsInitialized());
00356     #endif
00357
00358     K_ADDR wBottom = 0;
00359     auto wTop      = static_cast<K_ADDR>(m_ul6StackSize - 1);
00360     auto wMid      = ((wTop + wBottom) + 1) / 2;
00361
00362     CS_ENTER();
00363
00364     // Logarithmic bisection - find the point where the contents of the
00365     // stack go from 0xFF's to non 0xFF. Not Definitive, but accurate enough
00366     while ((wTop - wBottom) > 1) {
00367         #if STACK_GROWS_DOWN
00368             if (m_pwStack[wMid] != (K_WORD)(-1))
00369         #else
00370             if (m_pwStack[wMid] == (K_WORD)(-1))
00371         #endif
00372             {
00373                 wTop = wMid;
00374             } else {
00375                 wBottom = wMid;
00376             }
00377             wMid = (wTop + wBottom + 1) / 2;
00378         }
00379     }
00380
00381     CS_EXIT();
00382
00383     return wMid;
00384 }
00385
00386 //-----
00387 void Thread::Yield()
00388 {
00389
00390     CS_ENTER();
00391     // Run the scheduler
00392     if (Scheduler::IsEnabled()) {
00393         Scheduler::Schedule();
00394
00395         // Only switch contexts if the new task is different than the old task
00396         if (Scheduler::GetCurrentThread() !=
00397             Scheduler::GetNextThread()) {
00397             #if KERNEL_USE_QUANTUM
00398             #if KERNEL_TIMERS_THREADED
00399                 if (Quantum::GetTimerThread() != Scheduler::GetNextThread()) {
00400                     #endif
00401                         // new thread scheduled. Stop current quantum timer (if it exists),
00402                         // and restart it for the new thread (if required).
00403                         Quantum::RemoveThread();
00404                         Quantum::AddThread((Thread*)
00405                             Scheduler::GetNextThread());
00406                     #if KERNEL_TIMERS_THREADED
00407                 }
00408             #endif
00409             Thread::ContextSwitchSWI();
00410         } else {
00411             Scheduler::QueueScheduler();
00412         }
00413     }
00414     CS_EXIT();
00415 }
00416
00417 //-----
00418 void Thread::SetPriorityBase(PORT_PRIO_TYPE /*uXPriority*/)
00419 {
00420     #if KERNEL_EXTRA_CHECKS
00421         KERNEL_ASSERT(IsInitialized());
00422     #endif
00423
00424     GetCurrent()->Remove(this);
00425
00426     SetCurrent(Scheduler::GetThreadList(
00427         m_uXPriority));
00428
00429     GetCurrent()->Add(this);
00430 }
00431

```



```

00432 //-----
00433 void Thread::SetPriority(PORT_PRIO_TYPE uXPriority_)
00434 {
00435     #if KERNEL_EXTRA_CHECKS
00436         KERNEL_ASSERT(IsInitialized());
00437     #endif
00438
00439     auto bSchedule = false;
00440
00441     CS_ENTER();
00442     // If this is the currently running thread, it's a good idea to reschedule
00443     // Or, if the new priority is a higher priority than the current thread's.
00444     if ((g_pclCurrent == this) || (uXPriority_ > g_pclCurrent->GetPriority())) {
00445         bSchedule = true;
00446     }
00447     Scheduler::Remove(this);
00448     CS_EXIT();
00449
00450     m_uXCurPriority = uXPriority_;
00451     m_uXPriority     = uXPriority_;
00452
00453     CS_ENTER();
00454     Scheduler::Add(this);
00455     CS_EXIT();
00456
00457     if (bSchedule) {
00458         if (Scheduler::IsEnabled()) {
00459             CS_ENTER();
00460             Scheduler::Schedule();
00461         }
00462         #if KERNEL_USE_QUANTUM
00463         #if KERNEL_TIMERS_THREADED
00464             if (Quantum::GetTimerThread() != Scheduler::GetNextThread()) {
00465                 // new thread scheduled. Stop current quantum timer (if it exists),
00466                 // and restart it for the new thread (if required).
00467                 Quantum::RemoveThread();
00468                 Quantum::AddThread((Thread*)
00469                     Scheduler::GetNextThread());
00470             }
00471         #endif
00472         #endif
00473         CS_EXIT();
00474         Thread::ContextSwitchSWI();
00475     } else {
00476         Scheduler::QueueScheduler();
00477     }
00478 }
00479 }
00480
00481 //-----
00482 void Thread::InheritPriority(PORT_PRIO_TYPE uXPriority_)
00483 {
00484     #if KERNEL_EXTRA_CHECKS
00485         KERNEL_ASSERT(IsInitialized());
00486     #endif
00487
00488     SetOwner(Scheduler::GetThreadList(uXPriority_));
00489     m_uXCurPriority = uXPriority_;
00490 }
00491
00492 //-----
00493 void Thread::ContextSwitchSWI()
00494 {
00495     // Call the context switch interrupt if the scheduler is enabled.
00496     if (static_cast<int>(Scheduler::IsEnabled()) == 1) {
00497         KERNEL_TRACE_1("Context switch to Thread %d",
00498             Scheduler::GetNextThread()->GetID());
00499         #if KERNEL_USE_STACK_GUARD
00500         #if KERNEL_USE_IDLE_FUNC
00501             if ((g_pclCurrent != nullptr) && (g_pclCurrent->GetID() != 255)) {
00502                 if (g_pclCurrent->GetStackSlack() <= Kernel::GetStackGuardThreshold()) {
00503                     #if KERNEL_AWARE_SIMULATION
00504                         KernelAware::Trace(DBG_FILE, __LINE__, g_pclCurrent->
00505                             GetID(), g_pclCurrent->GetStackSlack());
00506                     #endif
00507                     Kernel::Panic(PANIC_STACK_SLACK_VIOLATED);
00508                 }
00509             }
00510         #endif
00511         #endif
00512         #if KERNEL_USE_THREAD_CALLOUTS
00513             ThreadContextCallout pfCallout = Kernel::GetThreadContextSwitchCallout
00514             ();

```

```

00515         if (pfCallout != nullptr) {
00516             pfCallout(g_pclCurrent);
00517         }
00518 #endif
00519         KernelSWI::Trigger();
00520     }
00521 }
00522
00523 #if KERNEL_USE_TIMEOUTS || KERNEL_USE_SLEEP
00524 //-----
00525 Timer* Thread::GetTimer()
00526 {
00527     #if KERNEL_EXTRA_CHECKS
00528         KERNEL_ASSERT(IsInitialized());
00529     #endif
00530
00531     return &m_clTimer;
00532 }
00533 #endif
00534 #if KERNEL_USE_TIMEOUTS
00535 //-----
00536 void Thread::SetExpired(bool bExpired_)
00537 {
00538     #if KERNEL_EXTRA_CHECKS
00539         KERNEL_ASSERT(IsInitialized());
00540     #endif
00541
00542     m_bExpired = bExpired_;
00543 }
00544
00545 //-----
00546 bool Thread::GetExpired()
00547 {
00548     #if KERNEL_EXTRA_CHECKS
00549         KERNEL_ASSERT(IsInitialized());
00550     #endif
00551
00552     return m_bExpired;
00553 }
00554 #endif
00555
00556 #if KERNEL_USE_IDLE_FUNC
00557 //-----
00558 void Thread::InitIdle(void)
00559 {
00560     m_eState      = ThreadState::Ready;
00561     ClearNode();
00562
00563     m_uXPriority   = 0;
00564     m_uXCurPriority = 0;
00565     m_pfEntryPoint = 0;
00566     m_pvArg       = 0;
00567     m_u8ThreadID   = 255;
00568     #if KERNEL_USE_THREADNAME
00569         m_szName = "IDLE";
00570     #endif
00571 }
00572 #endif
00573 } //namespace Mark3
00574

```

20.127 /home/moslevin/projects/github/m3-repo/kernel/src/threadlist.cpp File Reference

Thread linked-list definitions.

```

#include "kerneltypes.h"
#include "ll.h"
#include "threadlist.h"
#include "thread.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"

```

Namespaces

- [Mark3](#)

Class providing the software-interrupt required for context-switching in the kernel.

20.127.1 Detailed Description

Thread linked-list definitions.

Definition in file [threadlist.cpp](#).

20.128 threadlist.cpp

```

00001  /*=====
00002
00003  _____
00004  |   \   |   |   |   |   |   |   |   |
00005  |   \   |   |   |   |   |   |   |   |
00006  |   \   |   |   |   |   |   |   |   |
00007  |   \   |   |   |   |   |   |   |   |
00008  |   \   |   |   |   |   |   |   |   |
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  =====*/
00022  #include "kerneltypes.h"
00023  #include "ll.h"
00024  #include "threadlist.h"
00025  #include "thread.h"
00026
00027  #define _CAN_HAS_DEBUG
00028  //--[Autogenerated - Do Not Modify]-----
00029  #include "dbg_file_list.h"
00030  #include "buffallogger.h"
00031  #if defined(DBG_FILE)
00032  #error "Debug logging file token already defined! Bailing."
00033  #else
00034  #define DBG_FILE _DBG__KERNEL_THREADLIST_CPP
00035  #endif
00036  //--[End Autogenerated content]-----
00037  #include "kerneldebug.h"
00038  namespace Mark3
00039  {
00040  //--
00041  void ThreadList::SetPriority(PORT_PRIO_TYPE uXPriority_)
00042  {
00043      m_uXPriority = uXPriority_;
00044  }
00045
00046  //--
00047  void ThreadList::SetMapPointer(PriorityMap* pclMap_)
00048  {
00049      m_pclMap = pclMap_;
00050  }
00051
00052  //--
00053  void ThreadList::Add(LinkListNode* node_)
00054  {
00055      CircularLinkList::Add(node_);
00056      CircularLinkList::PivotForward();
00057
00058      // We've specified a bitmap for this threadlist
00059      if (m_pclMap != nullptr) {
00060          // Set the flag for this priority level
00061          m_pclMap->Set(m_uXPriority);
00062      }
00063  }
00064
00065  //--
00066  void ThreadList::AddPriority(LinkListNode* node_)
00067  {
00068      auto* pclCurr = static_cast<Thread*>(GetHead());
00069      if (pclCurr == nullptr) {

```

```

00070         Add(node_);
00071         return;
00072     }
00073     auto uXHeadPri = pclCurr->GetCurPriority();
00074
00075     auto* pclTail = static_cast<Thread*>(GetTail());
00076     auto* pclNode = static_cast<Thread*>(node_);
00077
00078     // Set the threadlist's priority level, flag pointer, and then add the
00079     // thread to the threadlist
00080     auto uXPriority = pclNode->GetCurPriority();
00081     do {
00082         if (uXPriority > pclCurr->GetCurPriority()) {
00083             break;
00084         }
00085         pclCurr = static_cast<Thread*>(pclCurr->GetNext());
00086     } while (pclCurr != pclTail);
00087
00088     // Insert pclNode before pclCurr in the linked list.
00089     InsertNodeBefore(pclNode, pclCurr);
00090
00091     // If the priority is greater than current head, reset
00092     // the head pointer.
00093     if (uXPriority > uXHeadPri) {
00094         m_pclHead = pclNode;
00095         m_pclTail = m_pclHead->prev;
00096     } else if (pclNode->GetNext() == m_pclHead) {
00097         m_pclTail = pclNode;
00098     }
00099 }
00100
00101 //-----
00102 void ThreadList::Add(LinkListNode* node_, PriorityMap* pclMap_,
00103     PORT_PRIO_TYPE uXPriority_)
00104 {
00105     // Set the threadlist's priority level, flag pointer, and then add the
00106     // thread to the threadlist
00107     SetPriority(uXPriority_);
00108     SetMapPointer(pclMap_);
00109     Add(node_);
00110 }
00111 //-----
00112 void ThreadList::Remove(LinkListNode* node_)
00113 {
00114     // Remove the thread from the list
00115     CircularLinkedList::Remove(node_);
00116
00117     // If the list is empty...
00118     if ((m_pclHead == nullptr) && (m_pclMap != nullptr)) {
00119         // Clear the bit in the bitmap at this priority level
00120         m_pclMap->Clear(m_uXPriority);
00121     }
00122 }
00123
00124 //-----
00125 Thread* ThreadList::HighestWaiter()
00126 {
00127     return static_cast<Thread*>(GetHead());
00128 }
00129 } //namespace Mark3

```

20.129 /home/moslevin/projects/github/m3-repo/kernel/src/timer.cpp File Reference

Timer implementations.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "timer.h"
#include "timerlist.h"
#include "timerscheduler.h"
#include "kerneltimer.h"
#include "threadport.h"
#include "quantum.h"
#include "dbg_file_list.h"
#include "buffalogger.h"

```

```
#include "kerneldebug.h"
```

20.129.1 Detailed Description

Timer implementations.

Definition in file [timer.cpp](#).

20.130 timer.cpp

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00022 #include "kerneltypes.h"
00023 #include "mark3cfg.h"
00024
00025 #include "timer.h"
00026 #include "timerlist.h"
00027 #include "timerscheduler.h"
00028 #include "kerneltimer.h"
00029 #include "threadport.h"
00030 #include "quantum.h"
00031
00032 #define _CAN_HAS_DEBUG
00033 //--[Autogenerated - Do Not Modify]-----
00034 #include "dbg_file_list.h"
00035 #include "buffalogger.h"
00036 #if defined(DBG_FILE)
00037 #error "Debug logging file token already defined! Bailing."
00038 #else
00039 #define DBG_FILE _DBG__KERNEL_TIMER_CPP
00040 #endif
00041 //--[End Autogenerated content]-----
00042
00043 #include "kerneldebug.h"
00044
00045 #if KERNEL_USE_TIMERS
00046 namespace Mark3
00047 {
00048 TimerList TimerScheduler::m_clTimerList;
00049
00050 //-----
00051 Timer::Timer()
00052 {
00053 #if KERNEL_EXTRA_CHECKS
00054     m_u8Initialized = TIMER_INVALID_COOKIE;
00055 #endif
00056     m_u8Flags = 0;
00057 }
00058
00059 //-----
00060 void Timer::Init()
00061 {
00062 #if KERNEL_EXTRA_CHECKS
00063     if (IsInitialized()) {
00064         KERNEL_ASSERT((m_u8Flags & TIMERLIST_FLAG_ACTIVE) == 0);
00065     }
00066 #endif
00067
00068     ClearNode();
00069     m_u32Interval = 0;
00070     m_u32TimerTolerance = 0;
00071     m_u32TimeLeft = 0;
00072     m_u8Flags = 0;
```

```

00073
00074 #if KERNEL_EXTRA_CHECKS
00075     SetInitialized();
00076 #endif
00077 }
00078
00079 //-----
00080 void Timer::Start(bool bRepeat_, uint32_t u32IntervalMs_,
00081                 TimerCallback pfCallback_, void* pvData_)
00082 {
00083     #if KERNEL_EXTRA_CHECKS
00084         KERNEL_ASSERT(IsInitialized());
00085     #endif
00086     if ((m_u8Flags & TIMERLIST_FLAG_ACTIVE) != 0) {
00087         return;
00088     }
00089     SetIntervalMSeconds(u32IntervalMs_);
00090     m_u32TimerTolerance = 0;
00091     m_pfCallback = pfCallback_;
00092     m_pvData = pvData_;
00093     if (!bRepeat_) {
00094         m_u8Flags = TIMERLIST_FLAG_ONE_SHOT;
00095     } else {
00096         m_u8Flags = 0;
00097     }
00098     Start();
00099 }
00100
00101 //-----
00102 void Timer::Start(
00103     bool bRepeat_, uint32_t u32IntervalMs_, uint32_t u32ToleranceMs_,
00104     TimerCallback pfCallback_, void* pvData_)
00105 {
00106     #if KERNEL_EXTRA_CHECKS
00107         KERNEL_ASSERT(IsInitialized());
00108     #endif
00109     if ((m_u8Flags & TIMERLIST_FLAG_ACTIVE) != 0) {
00110         return;
00111     }
00112     m_u32TimerTolerance = MSECNDSTO_TCKS(u32ToleranceMs_);
00113     Start(bRepeat_, u32IntervalMs_, pfCallback_, pvData_);
00114 }
00115
00116 //-----
00117 void Timer::Start()
00118 {
00119     #if KERNEL_EXTRA_CHECKS
00120         KERNEL_ASSERT(IsInitialized());
00121     #endif
00122     if ((m_u8Flags & TIMERLIST_FLAG_ACTIVE) != 0) {
00123         return;
00124     }
00125     m_pclOwner = Scheduler::GetCurrentThread();
00126     TimerScheduler::Add(this);
00127 }
00128
00129 //-----
00130 void Timer::Stop()
00131 {
00132     #if KERNEL_EXTRA_CHECKS
00133         KERNEL_ASSERT(IsInitialized());
00134     #endif
00135     if ((m_u8Flags & TIMERLIST_FLAG_ACTIVE) == 0) {
00136         return;
00137     }
00138     TimerScheduler::Remove(this);
00139 }
00140
00141 //-----
00142 void Timer::SetIntervalTicks(uint32_t u32Ticks_)
00143 {
00144     #if KERNEL_EXTRA_CHECKS
00145         KERNEL_ASSERT(IsInitialized());
00146     #endif
00147     m_u32Interval = u32Ticks_;
00148 }

```

```

00158 //-----
00160 //-----
00161 void Timer::SetIntervalSeconds(uint32_t u32Seconds_)
00162 {
00163     #if KERNEL_EXTRA_CHECKS
00164         KERNEL_ASSERT(IsInitialized());
00165     #endif
00166     m_u32Interval = SECONDS_TO_TICKS(u32Seconds_);
00168 }
00169
00170 //-----
00171 void Timer::SetIntervalMSeconds(uint32_t u32MSeconds_)
00172 {
00173     #if KERNEL_EXTRA_CHECKS
00174         KERNEL_ASSERT(IsInitialized());
00175     #endif
00176     m_u32Interval = MSECONDS_TO_TICKS(u32MSeconds_);
00178 }
00179
00180 //-----
00181 void Timer::SetIntervalUSeconds(uint32_t u32USeconds_)
00182 {
00183     #if KERNEL_EXTRA_CHECKS
00184         KERNEL_ASSERT(IsInitialized());
00185     #endif
00186     #if KERNEL_TIMERS_TICKLESS
00187         if (u32USeconds_ < KERNEL_TIMERS_MINIMUM_DELAY_US) {
00188             u32USeconds_ = KERNEL_TIMERS_MINIMUM_DELAY_US;
00189         }
00191     #endif
00192     m_u32Interval = USECONDS_TO_TICKS(u32USeconds_);
00194 }
00195
00196 //-----
00197 void Timer::SetTolerance(uint32_t u32Ticks_)
00198 {
00199     #if KERNEL_EXTRA_CHECKS
00200         KERNEL_ASSERT(IsInitialized());
00201     #endif
00202     m_u32TimerTolerance = u32Ticks_;
00203 }
00204 } //namespace Mark3
00205 #endif

```

20.131 /home/moslevin/projects/github/m3-repo/kernel/src/timerlist.cpp File Reference

Implements timer list processing algorithms, responsible for all timer tick and expiry logic.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "timerlist.h"
#include "kerneltimer.h"
#include "threadport.h"
#include "quantum.h"
#include "mutex.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"

```

20.131.1 Detailed Description

Implements timer list processing algorithms, responsible for all timer tick and expiry logic.

Definition in file [timerlist.cpp](#).


```

00094         // If the new interval is less than the amount of time remaining...
00095         lDelta = (int32_t)((uint32_t)KernelTimer::TimeToExpiry() - pclListNode->
m_u32Interval);
00096
00097         if (lDelta > 0) {
00098             // Set the new expiry time on the timer.
00099             m_u32NextWakeup = (uint32_t)
KernelTimer::SubtractExpiry((uint32_t)lDelta);
00100         }
00101     } else {
00102         m_u32NextWakeup = pclListNode->m_u32Interval;
00103         KernelTimer::SetExpiry(m_u32NextWakeup);
00104         KernelTimer::Start();
00105     }
00106 #endif
00107
00108     // Set the timer as active.
00109     pclListNode->m_u8Flags |= TIMERLIST_FLAG_ACTIVE;
00110
00111     TIMERLIST_UNLOCK();
00112 }
00113
00114 //-----
00115 void TimerList::Remove(Timer* pclLinkListNode_)
00116 {
00117     TIMERLIST_LOCK();
00118
00119     DoubleLinkedList::Remove(pclLinkListNode_);
00120     pclLinkListNode->m_u8Flags &= ~TIMERLIST_FLAG_ACTIVE;
00121
00122     #if KERNEL_TIMERS_TICKLESS
00123     if (this->GetHead() == NULL) {
00124         KernelTimer::Stop();
00125     }
00126     #endif
00127
00128     TIMERLIST_UNLOCK();
00129 }
00130
00131 //-----
00132 void TimerList::Process(void)
00133 {
00134     #if KERNEL_TIMERS_TICKLESS
00135     uint32_t u32NewExpiry;
00136     uint32_t u32Overtime;
00137     bool     bContinue;
00138     #endif
00139
00140     Timer* pclNode;
00141     Timer* pclPrev;
00142
00143     TIMERLIST_LOCK();
00144
00145     #if KERNEL_USE_QUANTUM
00146     Quantum::SetInTimer();
00147     #endif
00148     #if KERNEL_TIMERS_TICKLESS
00149     #if !KERNEL_TIMERS_THREADED
00150         // Clear the timer and its expiry time - keep it running though
00151         KernelTimer::ClearExpiry();
00152     #endif
00153     do {
00154     #endif
00155         pclNode = static_cast<Timer*>(GetHead());
00156         pclPrev = NULL;
00157
00158         #if KERNEL_TIMERS_TICKLESS
00159         bContinue = false;
00160         u32NewExpiry = MAX_TIMER_TICKS; // Used to indicate that no timers are pending
00161     #endif
00162
00163         // Subtract the elapsed time interval from each active timer.
00164         while (pclNode != 0) {
00165             // Active timers only...
00166             if ((pclNode->m_u8Flags & TIMERLIST_FLAG_ACTIVE) != 0) {
00167                 // Did the timer expire?
00168                 #if KERNEL_TIMERS_TICKLESS
00169                 if (pclNode->m_u32TimeLeft <= m_u32NextWakeup)
00170                 #else
00171                 if (0 == pclNode->m_u32TimeLeft)
00172                 #endif
00173                 {
00174                     // Yes - set the "callback" flag - we'll execute the callbacks later
00175                     pclNode->m_u8Flags |= TIMERLIST_FLAG_CALLBACK;
00176
00177                     if ((pclNode->m_u8Flags & TIMERLIST_FLAG_ONE_SHOT) != 0) {

```

```

00179             // If this was a one-shot timer, deactivate the timer.
00180             pclNode->m_u8Flags |= TIMERLIST_FLAG_EXPIRED;
00181             pclNode->m_u8Flags &= ~TIMERLIST_FLAG_ACTIVE;
00182         } else {
00183             // Reset the interval timer.
00184             // I think we're good though...
00185             pclNode->m_u32TimeLeft = pclNode->m_u32Interval;
00186         }
00187     }
00188     #if KERNEL_TIMERS_TICKLESS
00189         // If the time remaining (plus the length of the tolerance interval)
00190         // is less than the next expiry interval, set the next expiry interval.
00191         uint32_t u32Tmp = pclNode->m_u32TimeLeft + pclNode->m_u32TimerTolerance;
00192         if (u32Tmp < u32NewExpiry) {
00193             u32NewExpiry = u32Tmp;
00194         }
00195     #endif
00196 }
00197 }
00198 }
00199 #if KERNEL_TIMERS_TICKLESS
00200     else {
00201         // Not expiring, but determine how long to run the next timer interval for.
00202         pclNode->m_u32TimeLeft -= m_u32NextWakeup;
00203         if (pclNode->m_u32TimeLeft < u32NewExpiry) {
00204             u32NewExpiry = pclNode->m_u32TimeLeft;
00205         }
00206     }
00207 #endif
00208 }
00209 pclNode = static_cast<Timer*>(pclNode->GetNext());
00210 }
00211
00212 // Process the expired timers callbacks.
00213 pclNode = static_cast<Timer*>(GetHead());
00214 while (pclNode != 0) {
00215     pclPrev = pclNode;
00216     pclNode = static_cast<Timer*>(pclNode->GetNext());
00217
00218     // If the timer expired, run the callbacks now.
00219     if ((pclPrev->m_u8Flags & TIMERLIST_FLAG_CALLBACK) != 0) {
00220         bool bRemove = false;
00221         // If this was a one-shot timer, tag it for removal
00222         if ((pclPrev->m_u8Flags & TIMERLIST_FLAG_ONE_SHOT) != 0) {
00223             bRemove = true;
00224         }
00225
00226         // Run the callback. these callbacks must be very fast...
00227         pclPrev->m_pfCallback(pclPrev->m_pclOwner, pclPrev->m_pvData);
00228         pclPrev->m_u8Flags &= ~TIMERLIST_FLAG_CALLBACK;
00229
00230         // Remove one-shot-timers
00231         if (bRemove) {
00232             Remove(pclPrev);
00233         }
00234     }
00235 }
00236
00237 #if KERNEL_TIMERS_TICKLESS
00238     // Check to see how much time has elapsed since the time we
00239     // acknowledged the interrupt...
00240     u32Overtime = (uint32_t)KernelTimer::GetOvertime();
00241
00242     if (u32Overtime >= u32NewExpiry) {
00243         m_u32NextWakeup = u32Overtime;
00244         bContinue = true;
00245     }
00246
00247     // If it's taken longer to go through this loop than would take us to
00248     // the next expiry, re-run the timing loop
00249 } while (bContinue);
00250
00251 // This timer elapsed, but there's nothing more to do...
00252 // Turn the timer off.
00253 if (u32NewExpiry >= MAX_TIMER_TICKS) {
00254     KernelTimer::Stop();
00255 } else {
00256     // Update the timer with the new "Next Wakeup" value, plus whatever
00257     // overtime has accumulated since the last time we called this handler
00258     m_u32NextWakeup = (uint32_t)KernelTimer::SetExpiry(
00259         u32NewExpiry + u32Overtime);
00260 }
00261 #endif
00262 #if KERNEL_USE_QUANTUM
00263     Quantum::ClearInTimer();
00264 #endif
00265 #endif

```

```

00266
00267     TIMERLIST_UNLOCK();
00268 }
00269 } //namespace Mark3
00270 #endif // KERNEL_USE_TIMERS

```

20.133 /home/moslevin/projects/github/m3-repo/kernel/src/tracebuffer.cpp File Reference

Kernel trace buffer class definition.

```

#include "kerneltypes.h"
#include "tracebuffer.h"
#include "mark3cfg.h"
#include "dbg_file_list.h"
#include "buffalogger.h"
#include "kerneldebug.h"

```

20.133.1 Detailed Description

Kernel trace buffer class definition.

Definition in file [tracebuffer.cpp](#).

20.134 tracebuffer.cpp

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00019 #include "kerneltypes.h"
00020 #include "tracebuffer.h"
00021 #include "mark3cfg.h"
00022
00023 #define _CAN_HAS_DEBUG
00024 //--[Autogenerated - Do Not Modify]-----
00025 #include "dbg_file_list.h"
00026 #include "buffalogger.h"
00027 #if defined(DBG_FILE)
00028 #error "Debug logging file token already defined! Bailing."
00029 #else
00030 #define DBG_FILE _DBG__KERNEL_TRACEBUFFER_CPP
00031 #endif
00032
00033 #include "kerneldebug.h"
00034
00035 //--[End Autogenerated content]-----
00036
00037 #if KERNEL_USE_DEBUG && !KERNEL_AWARE_SIMULATION
00038 namespace Mark3
00039 {
00040     TraceBufferCallback_t TraceBuffer::m_pfCallback;
00041     uint16_t TraceBuffer::m_ul6SyncNumber;
00042     uint16_t TraceBuffer::m_ul6Index;
00043     uint16_t TraceBuffer::m_aul6Buffer[(TRACE_BUFFER_SIZE / sizeof(uint16_t))];

```

```

00044
00045 //-----
00046 void TraceBuffer::Init()
00047 {
00048 }
00049
00050 //-----
00051 void TraceBuffer::Write(uint16_t* pul6Data_, uint16_t ul6Size_)
00052 {
00053     // Pipe the data directly to the circular buffer
00054     uint16_t ul6Start;
00055
00056     // Update the circular buffer index in a critical section. The
00057     // rest of the operations can take place in any context.
00058     CS_ENTER();
00059     uint16_t ul6NextIndex;
00060     ul6Start = m_ul6Index;
00061     ul6NextIndex = m_ul6Index + ul6Size_;
00062     if (ul6NextIndex >= (sizeof(m_aul6Buffer) / sizeof(uint16_t))) {
00063         ul6NextIndex -= (sizeof(m_aul6Buffer) / sizeof(uint16_t));
00064     }
00065     m_ul6Index = ul6NextIndex;
00066     CS_EXIT();
00067
00068     // Write the data into the circular buffer.
00069     uint16_t i;
00070     bool bCallback = false;
00071     bool bPingPong = false;
00072     for (i = 0; i < ul6Size_; i++) {
00073         m_aul6Buffer[ul6Start++] = pul6Data_[i];
00074         if (ul6Start >= (sizeof(m_aul6Buffer) / sizeof(uint16_t))) {
00075             ul6Start = 0;
00076             bCallback = true;
00077         } else if (ul6Start == ((sizeof(m_aul6Buffer) / sizeof(uint16_t)) / 2)) {
00078             bPingPong = true;
00079             bCallback = true;
00080         }
00081     }
00082
00083     // Done writing - see if there's a 50% or rollover callback
00084     if (bCallback && m_pfCallback) {
00085         uint16_t ul6Size = (sizeof(m_aul6Buffer) / sizeof(uint16_t)) / 2;
00086         if (bPingPong) {
00087             m_pfCallback(m_aul6Buffer, ul6Size, bPingPong);
00088         } else {
00089             m_pfCallback(m_aul6Buffer + ul6Size, ul6Size, bPingPong);
00090         }
00091     }
00092 }
00093 } //namespace Mark3
00094 #endif

```

Example Documentation

This example demonstrates how low-overhead logging can be implemented using buffalogger.

This examples demonstrates how to use notification objects as a thread synchronization mechanism.

```

=====
/*
   \   \   \   \   \   \   \   \   \   \   \   \   \   \   \   \   \   \   \   \
  /   /   /   /   /   /   /   /   /   /   /   /   /   /   /   /   /   /   /   /
  \   \   \   \   \   \   \   \   \   \   \   \   \   \   \   \   \   \   \   \
   \   \   \   \   \   \   \   \   \   \   \   \   \   \   \   \   \   \   \   \
=====
--[Mark3 Realtime Platform]-----

Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
See license.txt for more information
=====*/
#include "mark3.h"

/*=====

Lab Example 10:  Thread Notifications

Lessons covered in this example include:
- Create a notification object, and use it to synchronize execution of Threads.

Takeaway:
- Notification objects are a lightweight mechanism to signal thread execution
  in situations where even a semaphore would be a heavier-weight option.

=====*/
#if !KERNEL_USE_IDLE_FUNC
#error "This demo requires KERNEL_USE_IDLE_FUNC"
#endif
extern "C" {
void __cxa_pure_virtual(void)
{
}
}

namespace {
using namespace Mark3;
//-----

```

```

Thread clApp1Thread;
K_WORD awApp1Stack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void App1Main(void* unused_);

//-----
Thread clApp2Thread;
K_WORD awApp2Stack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void App2Main(void* unused_);

//-----
// Notification object used in the example.
Notify clNotify;

//-----
void App1Main(void* unused_)
{
    while (1) {
        auto bNotified = false;
        // Block the thread until the notification object is signalled from
        // elsewhere.
        clNotify.Wait(&bNotified);

        KernelAware::Print("T1: Notified\n");
    }
}

//-----
void App2Main(void* unused_)
{
    while (1) {
        // Wait a while, then signal the notification object

        KernelAware::Print("T2: Wait 1s\n");
        Thread::Sleep(1000);

        KernelAware::Print("T2: Notify\n");
        clNotify.Signal();
    }
}
} // anonymous namespace

using namespace Mark3;
//-----
int main(void)
{
    // See the annotations in previous labs for details on init.
    Kernel::Init();

    // Initialize notifer and notify-ee threads
    clApp1Thread.Init(awApp1Stack, sizeof(awApp1Stack), 1, App1Main, 0);
    clApp1Thread.Start();

    clApp2Thread.Init(awApp2Stack, sizeof(awApp2Stack), 1, App2Main, 0);
    clApp2Thread.Start();

    // Initialize the Notify objects
    clNotify.Init();

    Kernel::Start();

    return 0;
}

```

21.3 lab11_mailboxes/main.cpp

This examples shows how to use mailboxes to deliver data between threads in a synchronized way.

```

/*-----
  _ _ _ _ _
 /   \   /   \   /   \   /   \   /   \
| _ _ | | _ _ | | _ _ | | _ _ | | _ _ |
 \   /   \   /   \   /   \   /   \   \
  _ _ _ _ _

--[Mark3 Realtime Platform]-----

Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
See license.txt for more information
=====*/

```

```

#include "mark3.h"

/*=====

Lab Example 11: Mailboxes

Lessons covered in this example include:
- Initialize a mailbox for use as an IPC mechanism.
- Create and use mailboxes to pass data between threads.

Takeaway:
- Mailboxes are a powerful IPC mechanism used to pass messages of a fixed-size
  between threads.

=====*/
#if !KERNEL_USE_IDLE_FUNC
#error "This demo requires KERNEL_USE_IDLE_FUNC"
#endif

extern "C" {
void __cxa_pure_virtual(void)
{
}
}

namespace {
using namespace Mark3;
//-----
Thread clApp1Thread;
K_WORD awApp1Stack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void App1Main(void* unused_);

//-----
Thread clApp2Thread;
K_WORD awApp2Stack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void App2Main(void* unused_);

//-----
Mailbox clMailbox;
uint8_t au8MBData[100];

typedef struct {
    uint8_t au8Buffer[10];
} MBType_t;

//-----
void App1Main(void* unused_)
{
    while (1) {
        MBType_t stMsg;

        // Wait until there is an envelope available in the shared mailbox, and
        // then log a trace message.
        clMailbox.Receive(&stMsg);
        KernelAware::Trace(0, __LINE__, stMsg.au8Buffer[0], stMsg.au8Buffer[9]);
    }
}

//-----
void App2Main(void* unused_)
{
    while (1) {
        MBType_t stMsg;

        // Place a bunch of envelopes in the mailbox, and then wait for a
        // while. Note that this thread has a higher priority than the other
        // thread, so it will keep pushing envelopes to the other thread until
        // it gets to the sleep, at which point the other thread will be allowed
        // to execute.

        KernelAware::Print("Messages Begin\n");

        for (uint8_t i = 0; i < 10; i++) {
            for (uint8_t j = 0; j < 10; j++) {
                stMsg.au8Buffer[j] = (i * 10) + j;
            }
            clMailbox.Send(&stMsg);
        }

        KernelAware::Print("Messages End\n");
        Thread::Sleep(2000);
    }
}
} // anonymous namespace

using namespace Mark3;
//-----

```



```

// This block declares the thread data for the main application thread. It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
Thread clAppThread;
K_WORD awAppStack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void AppMain(void* unused_);

//-----
// This block declares the thread data for the idle thread. It defines a
// thread object, stack (in word-array form), and the entry-point function
// used by the idle thread.
Thread clIdleThread;
K_WORD awIdleStack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void IdleMain(void* unused_);

//-----
void AppMain(void* unused_)
{
    // This function is run from within the application thread. Here, we
    // simply print a friendly greeting and allow the thread to sleep for a
    // while before repeating the message. Note that while the thread is
    // sleeping, CPU execution will transition to the Idle thread.

    while (1) {
        KernelAware::Print("Hello World!\n");
        Thread::Sleep(1000);
    }
}

//-----
void IdleMain(void* unused_)
{
    while (1) {
        // Low priority task + power management routines go here.
        // The actions taken in this context must *not* cause the thread
        // to block, as the kernel requires that at least one thread is
        // schedulable at all times when not using an idle thread.

        // Note that if you have no special power-management code or idle
        // tasks, an empty while(1){} loop is sufficient to guarantee that
        // condition.
    }
}
} // anonymous namespace

using namespace Mark3;
//-----
int main(void)
{
    // Before any Mark3 RTOS APIs can be called, the user must call Kernel::Init().
    // Note that if you have any hardware-specific init code, it can be called
    // before Kernel::Init, so long as it does not enable interrupts, or
    // rely on hardware peripherals (timer, software interrupt, etc.) used by the
    // kernel.
    Kernel::Init();

    // Once the kernel initialization has been complete, the user can add their
    // application thread(s) and idle thread. Threads added before the kernel
    // is started are referred to as the "static threads" in the system, as they
    // are the default working-set of threads that make up the application on
    // kernel startup.

    // Initialize the application thread to use a specified word-array as its stack.
    // The thread will run at priority level "1", and start execution the
    // "AppMain" function when it's started.
    clAppThread.Init(awAppStack, sizeof(awAppStack), 1, AppMain, 0);

    // Initialize the idle thread to use a specific word-array as its stack.
    // The thread will run at priority level "0", which is reserved for the idle
    // priority thread. IdleMain will be run when the thread is started.
    clIdleThread.Init(awIdleStack, sizeof(awIdleStack), 0, IdleMain, 0);

    // Once the static threads have been added, the user must then ensure that the
    // threads are ready to execute. By default, creating a thread is created
    // in a STOPPED state. All threads must manually be started using the
    // Start() API before they will be scheduled by the system. Here, we are
    // starting the application and idle threads before starting the kernel - and
    // that's OK. When the kernel is started, it will choose which thread to run
    // first from the pool of ready threads.

    clAppThread.Start();
    clIdleThread.Start();

    // All threads have been initialized and made ready. The kernel will now
    // select the first thread to run, enable the hardware required to run the
    // kernel (Timers, software interrupts, etc.), and then do whatever is
    // necessary to maneuver control of thread execution to the kernel. At this

```

```

// point, execution will transition to the highest-priority ready thread.
// This function will not return.

Kernel::Start();

// As Kernel::Start() results in the operating system being executed, control
// will not be relinquished back to main(). The "return 0" is simply to
// avoid warnings.

return 0;
}

```

21.5 lab2_idle_function/main.cpp

This example demonstrates how to use the idle function, instead of an idle thread to manage system inactivity.

```

/*=====
|
|-----|-----|-----|-----|-----|-----|-----|-----|
| \    /  | \    /  | \    /  | \    /  | \    /  | \    /  | \    /  |
|  \  /   |  \  /   |  \  /   |  \  /   |  \  /   |  \  /   |  \  /   |
|   \/    |   \/    |   \/    |   \/    |   \/    |   \/    |   \/    |
|  /  \   |  /  \   |  /  \   |  /  \   |  /  \   |  /  \   |  /  \   |
| /    \  | /    \  | /    \  | /    \  | /    \  | /    \  | /    \  |
|-----|-----|-----|-----|-----|-----|-----|-----|
--[Mark3 Realtime Platform]-----

Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
See license.txt for more information
=====*/
#include "mark3.h"

/*=====

Lab Example 2: Initializing the Mark3 RTOS kernel with one thread.

The following example code presents a working example of how to initialize
the Mark3 RTOS kernel, configured to use an application thread and the special
Kernel-Idle function. This example is functionally identical to lab1, although
it uses less memory as a result of only requiring one thread. This example also
uses the f1AVR kernel-aware module to print out messages when run through the
f1AVR AVR Simulator.

Lessons covered in this example include:

- usage of the Kernel::SetIdleFunc() API
- Changing an idle thread into an idle function
- You can save a thread and a stack by using an idle function instead of a
  dedicated idle thread.

Takeaway:

The Kernel-Idle context allows you to run the Mark3 RTOS without running
a dedicated idle thread (where supported). This results in a lower overall
memory footprint for the application, as you can avoid having to declare
a thread object and stack for Idle functionality.

=====*/
#ifdef !KERNEL_USE_IDLE_FUNC
#error "This demo requires KERNEL_USE_IDLE_FUNC"
#endif
extern "C" {
void __cxa_pure_virtual(void)
{
}
}

namespace {
using namespace Mark3;
//-----
// This block declares the thread data for the main application thread. It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
Thread clAppThread;
K_WORD awAppStack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void AppMain(void* unused_);

//-----
// This block declares the special function called from with the special
// Kernel-Idle context. We use the Kernel::SetIdleFunc() API to ensure that
// this function is called to provide our idle context.

```

```

void IdleMain(void);

//-----
void AppMain(void* unused_)
{
    // Same as in lab1.
    while (1) {
        KernelAware::Print("Hello World!\n");
        Thread::Sleep(1000);
    }
}

//-----
void IdleMain(void)
{
    // Low priority task + power management routines go here.
    // The actions taken in this context must *not* cause a blocking call,
    // similar to the requirements for an idle thread.

    // Note that unlike an idle thread, the idle function must run to
    // completion. As this is also called from a nested interrupt context,
    // it's worthwhile keeping this function brief, limited to absolutely
    // necessary functionality, and with minimal stack use.
}

} // anonymous namespace

using namespace Mark3;

//-----
int main(void)
{
    // See the annotations in lab1.
    Kernel::Init();

    // Initialize the main application thread, as in lab1. Note that even
    // though we're using an Idle function and not a dedicated thread, priority
    // level 0 is still reserved for idle functionality. Application threads
    // should never be scheduled at priority level 0 when the idle function is
    // used instead of an idle thread.
    clAppThread.Init(awAppStack, sizeof(awAppStack), 1, AppMain, 0);
    clAppThread.Start();

    // This function is used to install our specified idle function to be called
    // whenever there are no ready threads in the system. Note that if no
    // Idle function is specified, a default will be used. Note that this default
    // function is essentially a null operation.
    Kernel::SetIdleFunc(IdleMain);

    Kernel::Start();

    return 0;
}

```

21.6 lab3_round_robin/main.cpp

This example demonstrates how to use round-robin thread scheduling with multiple threads of the same priority.

[illegible]

Takeaway:

- CPU Scheduling can be achieved using not just strict Thread priority, but also with round-robin time-slicing between threads at the same priority.

```

=====*/
#if !KERNEL_USE_IDLE_FUNC
#error "This demo requires KERNEL_USE_IDLE_FUNC"
#endif

extern "C" {
void __cxa_pure_virtual(void)
{
}
}

namespace {
using namespace Mark3;
//-----
// This block declares the thread data for one main application thread. It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
Thread clApp1Thread;
K_WORD awApp1Stack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void App1Main(void* unused_);

//-----
// This block declares the thread data for one main application thread. It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
Thread clApp2Thread;
K_WORD awApp2Stack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void App2Main(void* unused_);

//-----
void App1Main(void* unused_)
{
    // Simple loop that increments a volatile counter to 1000000 then resets
    // it while printing a message.
    volatile uint32_t u32Counter = 0;
    while (1) {
        u32Counter++;
        if (u32Counter == 1000000) {
            u32Counter = 0;
            KernelAware::Print("Thread 1 - Did some work\n");
        }
    }
}

//-----
void App2Main(void* unused_)
{
    // Same as App1Main. However, as this thread gets twice as much CPU time
    // as Thread 1, you should see its message printed twice as often as the
    // above function.
    volatile uint32_t u32Counter = 0;
    while (1) {
        u32Counter++;
        if (u32Counter == 1000000) {
            u32Counter = 0;
            KernelAware::Print("Thread 2 - Did some work\n");
        }
    }
}
} // anonymous namespace

using namespace Mark3;
//-----
int main(void)
{
    // See the annotations in lab1.
    Kernel::Init();

    // In this exercise, we create two threads at the same priority level.
    // As a result, the CPU will automatically swap between these threads
    // at runtime to ensure that each get a chance to execute.

    clApp1Thread.Init(awApp1Stack, sizeof(awApp1Stack), 1, App1Main, 0);
    clApp2Thread.Init(awApp2Stack, sizeof(awApp2Stack), 1, App2Main, 0);

    // Set the threads up so that Thread 1 can get 4ms of CPU time uninterrupted,
    // but Thread 2 can get 8ms of CPU time uninterrupted. This means that
    // in an ideal situation, Thread 2 will get to do twice as much work as
    // Thread 1 - even though they share the same scheduling priority.

    // Note that if SetQuantum() isn't called on a thread, a default value

```

```
// is set such that each thread gets equal timeslicing in the same
// priority group by default.  You can play around with these values and
// observe how it affects the execution of both threads.

clApp1Thread.SetQuantum(4);
clApp2Thread.SetQuantum(8);

clApp1Thread.Start();
clApp2Thread.Start();

Kernel::Start();

return 0;
}
```

21.7 lab4_semaphores/main.cpp

This example demonstrates how to use semaphores for Thread synchronization.

```

/*
=====
--[Mark3 Realtime Platform]-----
Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
See license.txt for more information
=====*/
#include "mark3.h"
/*=====

Lab Example 4:  using binary semaphores

In this example, we implement two threads, synchronized using a semaphore to
model the classic producer-consumer pattern.  One thread does work, and then
posts the semaphore indicating that the other thread can consume that work.
The blocking thread just waits idly until there is data for it to consume.

Lessons covered in this example include:
-Use of a binary semaphore to implement the producer-consumer pattern
-Synchronization of threads (within a single priority, or otherwise)
  using a semaphore

Takeaway:

Semaphores can be used to control which threads execute at which time.  This
allows threads to work cooperatively to achieve a goal in the system.

=====*/
#error "This demo requires KERNEL_USE_IDLE_FUNC"
#endif

extern "C" {
void __cxa_pure_virtual(void)
{
}
}

namespace {
using namespace Mark3;
//-----
// This block declares the thread data for one main application thread.  It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
#define APP1_STACK_SIZE (PORT_KERNEL_DEFAULT_STACK_SIZE)
Thread clApp1Thread;
K_WORD awApp1Stack[APP1_STACK_SIZE];
void App1Main(void* unused_);

//-----
// This block declares the thread data for one main application thread.  It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
#define APP2_STACK_SIZE (PORT_KERNEL_DEFAULT_STACK_SIZE)

```

```

Thread clApp2Thread;
K_WORD awApp2Stack[APP2_STACK_SIZE];
void App2Main(void* unused_);

//-----
// This is the semaphore that we'll use to synchronize two threads in this
// demo application
Semaphore clMySem;

//-----
void App1Main(void* unused_)
{
    while (1) {
        // Wait until the semaphore is posted from the other thread
        KernelAware::Print("Wait\n");
        clMySem.Pend();

        // Producer thread has finished doing its work -- do something to
        // consume its output. Once again - a contrived example, but we
        // can imagine that printing out the message is "consuming" the output
        // from the other thread.
        KernelAware::Print("Triggered!\n");
    }
}

//-----
void App2Main(void* unused_)
{
    volatile uint32_t u32Counter = 0;

    while (1) {
        // Do some work. Once the work is complete, post the semaphore. This
        // will cause the other thread to wake up and then take some action.
        // It's a bit contrived, but imagine that the results of this process
        // are necessary to drive the work done by that other thread.
        u32Counter++;
        if (u32Counter == 1000000) {
            u32Counter = 0;
            KernelAware::Print("Posted\n");
            clMySem.Post();
        }
    }
}

// anonymous namespace
using namespace Mark3;

//-----
int main(void)
{
    // See the annotations in previous labs for details on init.
    Kernel::Init();

    // In this example we create two threads to illustrate the use of a
    // binary semaphore as a synchronization method between two threads.

    // Thread 1 is a "consumer" thread -- It waits, blocked on the semaphore
    // until thread 2 is done doing some work. Once the semaphore is posted,
    // the thread is unblocked, and does some work.

    // Thread 2 is thus the "producer" thread -- It does work, and once that
    // work is done, the semaphore is posted to indicate that the other thread
    // can use the producer's work product.

    clApp1Thread.Init(awApp1Stack, APP1_STACK_SIZE, 1, App1Main, 0);
    clApp2Thread.Init(awApp2Stack, APP2_STACK_SIZE, 1, App2Main, 0);

    clApp1Thread.Start();
    clApp2Thread.Start();

    // Initialize a binary semaphore (maximum value of one, initial value of
    // zero).
    clMySem.Init(0, 1);

    Kernel::Start();

    return 0;
}

```

21.8 lab5_mutexes/main.cpp

This example demonstrates how to use mutexes to protect against concurrent access to resources.

```

/*
-----[Mark3 Realtime Platform]-----

Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
See license.txt for more information
=====*/
#include "mark3.h"

/*=====

Lab Example 5: using Mutexes.

Lessons covered in this example include:
You can use mutexes to lock accesses to a shared resource

Takeaway:

=====*/
#if !KERNEL_USE_IDLE_FUNC
#error "This demo requires KERNEL_USE_IDLE_FUNC"
#endif

extern "C" {
void __cxa_pure_virtual(void)
{
}
}
namespace {
using namespace Mark3;

//-----
// This block declares the thread data for one main application thread. It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
#define APP1_STACK_SIZE (PORT_KERNEL_DEFAULT_STACK_SIZE)
Thread clApp1Thread;
K_WORD awApp1Stack[APP1_STACK_SIZE];
void App1Main(void* unused_);

//-----
// This block declares the thread data for one main application thread. It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
#define APP2_STACK_SIZE (PORT_KERNEL_DEFAULT_STACK_SIZE)
Thread clApp2Thread;
K_WORD awApp2Stack[APP2_STACK_SIZE];
void App2Main(void* unused_);

//-----
// This is the mutex that we'll use to synchronize two threads in this
// demo application.
Mutex clMyMutex;

// This counter variable is the "shared resource" in the example, protected
// by the mutex. Only one thread should be given access to the counter at
// any time.
volatile uint32_t u32Counter = 0;

//-----
void App1Main(void* unused_)
{
    while (1) {
        // Claim the mutex. This will prevent any other thread from claiming
        // this lock simultaneously. As a result, the other thread has to
        // wait until we're done before it can do its work. You will notice
        // that the Start/Done prints for the thread will come as a pair (i.e.
        // you won't see "Thread2: Start" then "Thread1: Start").

        clMyMutex.Claim();

        // Start our work (incrementing a counter). Notice that the Start and
        // Done prints wind up as a pair when simulated with f1AVR.

        KernelAware::Print("Thread1: Start\n");
        u32Counter++;
        while (u32Counter <= 1000000) {
            u32Counter++;
        }
        u32Counter = 0;
        KernelAware::Print("Thread1: Done\n");
    }
}
}

```


Lessons covered in this example include:

Takeaway:

```

=====*/
#if !KERNEL_USE_IDLE_FUNC
#error "This demo requires KERNEL_USE_IDLE_FUNC"
#endif

extern "C" {
void __cxa_pure_virtual(void)
{
}
}

namespace {
using namespace Mark3;
//-----
// This block declares the thread data for one main application thread. It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
#define APPL_STACK_SIZE (PORT_KERNEL_DEFAULT_STACK_SIZE)
Thread clApplThread;
K_WORD awApplStack[APPL_STACK_SIZE];
void ApplMain(void* unused_);

//-----
void PeriodicCallback(Thread* owner, void* pvData_)
{
    // Timer callback function used to post a semaphore. Posting the semaphore
    // will wake up a thread that's pending on that semaphore.
    auto* pclSem = static_cast<Semaphore*>(pvData_);
    pclSem->Post();
}

//-----
void OneShotCallback(Thread* owner, void* pvData_)
{
    KernelAware::Print("One-shot timer expired.\n");
}

//-----
void ApplMain(void* unused_)
{
    Timer clMyTimer; // Periodic timer object
    Timer clOneShot; // One-shot timer object

    Semaphore clMySem; // Semaphore used to wake this thread

    // Initialize a binary semaphore (maximum value of one, initial value of
    // zero).
    clMySem.Init(0, 1);

    // Start a timer that triggers every 500ms that will call PeriodicCallback.
    // This timer simulates an external stimulus or event that would require
    // an action to be taken by this thread, but would be serviced by an
    // interrupt or other high-priority context.

    // PeriodicCallback will post the semaphore which wakes the thread
    // up to perform an action. Here that action consists of a trivial message
    // print.
    clMyTimer.Start(true, 500, PeriodicCallback, (void*)&clMySem);

    // Set up a one-shot timer to print a message after 2.5 seconds, asynchronously
    // from the execution of this thread.
    clOneShot.Start(false, 2500, OneShotCallback, 0);

    while (1) {
        // Wait until the semaphore is posted from the timer expiry
        clMySem.Pend();

        // Take some action after the timer posts the semaphore to wake this
        // thread.
        KernelAware::Print("Thread Triggered.\n");
    }
}
} // anonymous namespace

using namespace Mark3;

//-----
int main(void)
{
    // See the annotations in previous labs for details on init.
    Kernel::Init();

    clApplThread.Init(awApplStack, sizeof(awApplStack), 1, ApplMain, 0);
}

```



```
//-----
void App1Main(void* unused_)
{
    while (1) {
        // Block this thread until any of the event flags have been set by
        // some outside force (here, we use Thread 2). As an exercise to the
        // user, try playing around with the event mask to see the effect it
        // has on which events get processed. Different threads can block on
        // different bitmasks - this allows events with different real-time
        // priorities to be handled in different threads, while still using
        // the same event-flag object.

        // Also note that EventFlagOperation::Any_Set indicates that the thread will be
        // unblocked whenever any of the flags in the mask are selected. If
        // you wanted to trigger an action that only takes place once multiple
        // bits are set, you could block the thread waiting for a specific
        // event bitmask with EventFlagOperation::All_Set specified.
        auto ul6Flags = clFlags.Wait(0xFFFF, EventFlagOperation::Any_Set);

        // Print a message indicating which bit was set this time.
        switch (ul6Flags) {
            case 0x0001: KernelAware::Print("Event1\n"); break;
            case 0x0002: KernelAware::Print("Event2\n"); break;
            case 0x0004: KernelAware::Print("Event3\n"); break;
            case 0x0008: KernelAware::Print("Event4\n"); break;
            case 0x0010: KernelAware::Print("Event5\n"); break;
            case 0x0020: KernelAware::Print("Event6\n"); break;
            case 0x0040: KernelAware::Print("Event7\n"); break;
            case 0x0080: KernelAware::Print("Event8\n"); break;
            case 0x0100: KernelAware::Print("Event9\n"); break;
            case 0x0200: KernelAware::Print("Event10\n"); break;
            case 0x0400: KernelAware::Print("Event11\n"); break;
            case 0x0800: KernelAware::Print("Event12\n"); break;
            case 0x1000: KernelAware::Print("Event13\n"); break;
            case 0x2000: KernelAware::Print("Event14\n"); break;
            case 0x4000: KernelAware::Print("Event15\n"); break;
            case 0x8000: KernelAware::Print("Event16\n"); break;
            default: break;
        }

        // Clear the event-flag that we just printed a message about. This
        // will allow ul6 to acknowledge further events in that bit in the future.
        clFlags.Clear(ul6Flags);
    }
}

//-----
void App2Main(void* unused_)
{
    uint16_t ul6Flag = 1;
    while (1) {
        Thread::Sleep(100);

        // Event flags essentially map events to bits in a bitmap. Here we
        // set one bit each 100ms. In this loop, we cycle through bits 0-15
        // repeatedly. Note that this will wake the other thread, which is
        // blocked, waiting for *any* of the flags in the bitmap to be set.
        clFlags.Set(ul6Flag);

        // Bitshift the flag value to the left. This will be the flag we set
        // the next time this thread runs through its loop.
        if (ul6Flag != 0x8000) {
            ul6Flag <<= 1;
        } else {
            ul6Flag = 1;
        }
    }
}

} // anonymous namespace

using namespace Mark3;
//-----
int main(void)
{
    // See the annotations in previous labs for details on init.
    Kernel::Init();

    clApp1Thread.Init(awApp1Stack, sizeof(awApp1Stack), 1, App1Main, 0);
    clApp2Thread.Init(awApp2Stack, sizeof(awApp2Stack), 1, App2Main, 0);

    clApp1Thread.Start();
    clApp2Thread.Start();

    clFlags.Init();

    Kernel::Start();
}
```



```

//-----
void App1Main(void* unused_)
{
    auto ul6Data = 0;
    while (1) {
        // This thread grabs a message from the global message pool, sets a
        // code-value and the message data pointer, then sends the message to
        // a message queue object. Another thread (Thread2) is blocked, waiting
        // for a message to arrive in the queue.

        // Get the message object
        auto* pclMsg = s_clMessagePool.Pop();

        // Set the message object's data (contrived in this example)
        pclMsg->SetCode(0x1337);
        ul6Data++;
        pclMsg->SetData(&ul6Data);

        // Send the message to the shared message queue
        clMsgQ.Send(pclMsg);

        // Wait before sending another message.
        Thread::Sleep(200);
    }
}

//-----
void App2Main(void* unused_)
{
    while (1) {
        // This thread waits until it receives a message on the shared global
        // message queue. When it gets the message, it prints out information
        // about the message's code and data, before returning the message object
        // back to the global message pool. In a more practical application,
        // the user would typically use the code to tell the receiving thread
        // what kind of message was sent, and what type of data to expect in the
        // data field.

        // Wait for a message to arrive on the specified queue. Note that once
        // this thread receives the message, it is "owned" by the thread, and
        // must be returned back to its source message pool when it is no longer
        // needed.
        auto* pclMsg = clMsgQ.Receive();

        // We received a message, now print out its information
        KernelAware::Print("Received Message\n");
        KernelAware::Trace(0, __LINE__, pclMsg->GetCode(), *((uint16_t*)pclMsg->GetData()
    ));

        // Done with the message, return it back to the global message queue.
        s_clMessagePool.Push(pclMsg);
    }
} // anonymous namespace

using namespace Mark3;
//-----
int main(void)
{
    // See the annotations in previous labs for details on init.
    Kernel::Init();

    clApp1Thread.Init(awApp1Stack, sizeof(awApp1Stack), 1, App1Main, 0);
    clApp2Thread.Init(awApp2Stack, sizeof(awApp2Stack), 1, App2Main, 0);

    clApp1Thread.Start();
    clApp2Thread.Start();

    clMsgQ.Init();

    s_clMessagePool.Init();
    for (int i = 0; i < MESSAGE_POOL_SIZE; i++) {
        s_clMessages[i].Init();
        s_clMessagePool.Push(&s_clMessages[i]);
    }

    Kernel::Start();

    return 0;
}

```



```

        MemUtil::DecimalToHex((K_ADDR) apclActiveThreads[i], szStr);
        KernelAware::Print(szStr);
        KernelAware::Print(" ");
        MemUtil::DecimalToString(u16Slack, szStr);
        KernelAware::Print(szStr);
        KernelAware::Print("\n");
    }
}

void PrintCPUUsage(void)
{
    KernelAware::Print("Cpu usage\n");
    for (int i = 0; i < MAX_THREADS; i++) {
        if (apclActiveThreads[i] != 0) {
            KernelAware::Trace(0, __LINE__, (K_ADDR) apclActiveThreads[i],
                               aul6ActiveTime[i]);
        }
    }
}

void ThreadCreate(Thread* pclThread_)
{
    KernelAware::Print("TC\n");
    CS_ENTER();
    for (uint8_t i = 0; i < MAX_THREADS; i++) {
        if (apclActiveThreads[i] == 0) {
            apclActiveThreads[i] = pclThread_;
            break;
        }
    }
    CS_EXIT();

    PrintThreadSlack();
    PrintCPUUsage();
}

void ThreadExit(Thread* pclThread_)
{
    KernelAware::Print("TX\n");
    CS_ENTER();
    for (uint8_t i = 0; i < MAX_THREADS; i++) {
        if (apclActiveThreads[i] == pclThread_) {
            apclActiveThreads[i] = 0;
            aul6ActiveTime[i] = 0;
            break;
        }
    }
    CS_EXIT();

    PrintThreadSlack();
    PrintCPUUsage();
}

void ThreadContextSwitch(Thread* pclThread_)
{
    KernelAware::Print("CS\n");
    static uint16_t u16LastTick = 0;
    auto u16Ticks = KernelTimer::Read();

    CS_ENTER();
    for (uint8_t i = 0; i < MAX_THREADS; i++) {
        if (apclActiveThreads[i] == pclThread_) {
            aul6ActiveTime[i] += u16Ticks - u16LastTick;
            break;
        }
    }
    CS_EXIT();

    u16LastTick = u16Ticks;
}

#endif

//-----
void WorkerMain1(void* arg_)
{
    auto* pclSem = static_cast<Semaphore*>(arg_);
    uint32_t u32Count = 0;

    // Do some work. Post a semaphore to notify the other thread that the
    // work has been completed.
    while (u32Count < 1000000) {
        u32Count++;
    }

    KernelAware::Print("Worker1 -- Done Work\n");
}

```

```

    pclSem->Post();

    // Work is completed, just spin now. Let another thread destroy u16.
    while (1) {
    }
}
//-----
void WorkerMain2(void* arg_)
{
    uint32_t u32Count = 0;
    while (u32Count < 1000000) {
        u32Count++;
    }

    KernelAware::Print("Worker2 -- Done Work\n");

    // A dynamic thread can self-terminate as well:
    Scheduler::GetCurrentThread()->Exit();
}

//-----
void App1Main(void* unused_)
{
    Thread    clMyThread;
    Semaphore clMySem;

    clMySem.Init(0, 1);
    while (1) {
        // Example 1 - create a worker thread at our current priority in order to
        // parallelize some work.
        clMyThread.Init(awApp2Stack, sizeof(awApp2Stack), 1, WorkerMain1, (void*)&clMySem);
        clMyThread.Start();

        // Do some work of our own in parallel, while the other thread works on its project.
        uint32_t u32Count = 0;
        while (u32Count < 100000) {
            u32Count++;
        }

        KernelAware::Print("Thread -- Done Work\n");

        PrintThreadSlack();

        // Wait for the other thread to finish its job.
        clMySem.Pend();

        // Once the thread has signalled u16, we can safely call "Exit" on the thread to
        // remove it from scheduling and recycle it later.
        clMyThread.Exit();

        // Spin the thread up again to do something else in parallel. This time, the thread
        // will run completely asynchronously to this thread.
        clMyThread.Init(awApp2Stack, sizeof(awApp2Stack), 1, WorkerMain2, 0);
        clMyThread.Start();

        u32Count = 0;
        while (u32Count < 1000000) {
            u32Count++;
        }

        KernelAware::Print("Thread -- Done Work\n");

        // Check that we're sure the worker thread has terminated before we try running the
        // test loop again.
        while (clMyThread.GetState() != ThreadState::Exit) {
        }

        KernelAware::Print(" Test Done\n");
        Thread::Sleep(1000);
        PrintThreadSlack();
    }
}
} // anonymous namespace

using namespace Mark3;
//-----
int main(void)
{
    // See the annotations in previous labs for details on init.
    Kernel::Init();

    Kernel::SetThreadCreateCallout(ThreadCreate);
    Kernel::SetThreadExitCallout(ThreadExit);
    Kernel::SetThreadContextSwitchCallout(ThreadContextSwitch);

    clApp1Thread.Init(awApp1Stack, sizeof(awApp1Stack), 1, App1Main, 0);
    clApp1Thread.Start();
}

```



```
    Kernel::Start();  
    return 0;  
}
```


Index

/home/moslevin/projects/github/m3-repo/kernel/libs/mark3c/src/public/make282
_types.h, 209 /home/moslevin/projects/github/m3-repo/kernel/src/priomap.↵
/home/moslevin/projects/github/m3-repo/kernel/libs/mark3c/src/public/mmap286.↵
h, 212, 222 /home/moslevin/projects/github/m3-repo/kernel/src/profile.↵
/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p267/kernelprofile.↵
cpp, 228, 229 /home/moslevin/projects/github/m3-repo/kernel/src/public/atomic.↵
/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p267/gcc/kernelswi.↵
cpp, 230 /home/moslevin/projects/github/m3-repo/kernel/src/public/autoalloc.↵
/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p267/gcc/kernelswi.↵
cpp, 231, 232 /home/moslevin/projects/github/m3-repo/kernel/src/public/blocking.↵
/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p267/public/kernelprofile.↵
h, 234, 235 /home/moslevin/projects/github/m3-repo/kernel/src/public/buffalogger.↵
/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p267/public/kernelswi.↵
h, 235, 236 /home/moslevin/projects/github/m3-repo/kernel/src/public/condvar.↵
/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p267/public/kernelswi.↵
h, 236, 237 /home/moslevin/projects/github/m3-repo/kernel/src/public/kernelswi.↵
/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p267/public/kernelswi.↵
h, 237, 240 /home/moslevin/projects/github/m3-repo/kernel/src/public/eventflag.↵
/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p267/public/portcfg.↵
h, 237, 240 /home/moslevin/projects/github/m3-repo/kernel/src/public/kernel.↵
/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p267/gcc/public/threadport.↵
h, 240, 242 /home/moslevin/projects/github/m3-repo/kernel/src/public/kernelaware.↵
/home/moslevin/projects/github/m3-repo/kernel/src/arch/avr/atmega1284p267/gcc/public/threadport.↵
cpp, 244, 245 /home/moslevin/projects/github/m3-repo/kernel/src/public/kerneldebug.↵
/home/moslevin/projects/github/m3-repo/kernel/src/atomic.↵ h, 301, 302
cpp, 247 /home/moslevin/projects/github/m3-repo/kernel/src/public/kerneltypes.↵
/home/moslevin/projects/github/m3-repo/kernel/src/autoalloc.↵ h, 307, 308
cpp, 249 /home/moslevin/projects/github/m3-repo/kernel/src/public/ksemaphore.↵
/home/moslevin/projects/github/m3-repo/kernel/src/blocking.↵ h, 309
cpp, 253, 254 /home/moslevin/projects/github/m3-repo/kernel/src/public/ll.↵
/home/moslevin/projects/github/m3-repo/kernel/src/condvar.↵ h, 310, 311
cpp, 255 /home/moslevin/projects/github/m3-repo/kernel/src/public/lockguard.↵
/home/moslevin/projects/github/m3-repo/kernel/src/eventflag.↵ h, 312, 313
cpp, 256, 257 /home/moslevin/projects/github/m3-repo/kernel/src/public/mailbox.↵
/home/moslevin/projects/github/m3-repo/kernel/src/kernel.↵ h, 313, 314
cpp, 261 /home/moslevin/projects/github/m3-repo/kernel/src/public/manual.↵
/home/moslevin/projects/github/m3-repo/kernel/src/kernelaware.↵ h, 316
cpp, 263 /home/moslevin/projects/github/m3-repo/kernel/src/public/mark3.↵
/home/moslevin/projects/github/m3-repo/kernel/src/ksemaphore.↵ h, 316, 317
cpp, 265, 266 /home/moslevin/projects/github/m3-repo/kernel/src/public/mark3cfg.↵
/home/moslevin/projects/github/m3-repo/kernel/src/ll.↵ h, 318, 326
cpp, 269 /home/moslevin/projects/github/m3-repo/kernel/src/public/message.↵
/home/moslevin/projects/github/m3-repo/kernel/src/lockguard.↵ h, 328, 329
cpp, 271, 272 /home/moslevin/projects/github/m3-repo/kernel/src/public/mutex.↵
/home/moslevin/projects/github/m3-repo/kernel/src/mailbox.↵ h, 330, 332
cpp, 272, 273 /home/moslevin/projects/github/m3-repo/kernel/src/public/notify.↵
/home/moslevin/projects/github/m3-repo/kernel/src/message.↵ h, 332, 333
cpp, 276 /home/moslevin/projects/github/m3-repo/kernel/src/public/paniccodes.↵
/home/moslevin/projects/github/m3-repo/kernel/src/mutex.↵ h, 334
cpp, 278, 279 /home/moslevin/projects/github/m3-repo/kernel/src/public/priomap.↵
/home/moslevin/projects/github/m3-repo/kernel/src/notify.↵ h, 334, 335

- /home/moslevin/projects/github/m3-repo/kernel/src/public/Block.h.↔
 - h, [336](#), [337](#)
- /home/moslevin/projects/github/m3-repo/kernel/src/public/BlockPriority
 - h, [337](#), [338](#)
- /home/moslevin/projects/github/m3-repo/kernel/src/public/readerwriter.h.↔
 - h, [339](#)
- /home/moslevin/projects/github/m3-repo/kernel/src/public/schedthreadport.h, [241](#)
 - Claim
- /home/moslevin/projects/github/m3-repo/kernel/src/public/thread.h
 - Mark3::Mutex, [149](#)
 - Claim_i
- /home/moslevin/projects/github/m3-repo/kernel/src/public/threadlist.h
 - Mark3::Mutex, [150](#)
 - Clear
- /home/moslevin/projects/github/m3-repo/kernel/src/public/timer.h
 - Mark3::EventFlag, [109](#)
 - Mark3::PriorityMap, [155](#)
- /home/moslevin/projects/github/m3-repo/kernel/src/public/timerexpiry
 - h, [349](#), [350](#)
- /home/moslevin/projects/github/m3-repo/kernel/src/public/timerlist.h.↔
 - h, [351](#)
- /home/moslevin/projects/github/m3-repo/kernel/src/public/timerthread
 - h, [352](#)
- /home/moslevin/projects/github/m3-repo/kernel/src/quantum/ComputeCurrentTicks
 - cpp, [353](#)
- /home/moslevin/projects/github/m3-repo/kernel/src/readerwriter
 - cpp, [355](#)
- /home/moslevin/projects/github/m3-repo/kernel/src/scheduler
 - ContextSwitchSWI
 - cpp, [357](#), [358](#)
- /home/moslevin/projects/github/m3-repo/kernel/src/thread.CopyData
 - cpp, [359](#), [360](#)
- /home/moslevin/projects/github/m3-repo/kernel/src/threadlist.h.↔
 - cpp, [366](#), [367](#)
- /home/moslevin/projects/github/m3-repo/kernel/src/timer.h.↔
 - cpp, [368](#), [369](#)
- /home/moslevin/projects/github/m3-repo/kernel/src/timerlist.h.↔
 - cpp, [371](#), [372](#)
- /home/moslevin/projects/github/m3-repo/kernel/src/tracebuffer.h.↔
 - cpp, [375](#)
- AVR
 - portcfg.h, [238](#)
- AcquireReader
 - Mark3::ReaderWriterLock, [162](#)
- AcquireReader_i
 - Mark3::ReaderWriterLock, [163](#)
- AcquireWriter
 - Mark3::ReaderWriterLock, [163](#)
- AcquireWriter_i
 - Mark3::ReaderWriterLock, [164](#)
- Add
 - Mark3::Atomic, [92](#)
 - Mark3::CircularLinkList, [102](#)
 - Mark3::DoubleLinkList, [107](#)
 - Mark3::Scheduler, [165](#)
 - Mark3::ThreadList, [192](#)
 - Mark3::TimerList, [205](#)
 - Mark3::TimerScheduler, [206](#)
- AddPriority
 - Mark3::ThreadList, [192](#)
- AddThread
 - Mark3::Quantum, [159](#)
- Block
 - Mark3::BlockingObject, [100](#)
- BlockPriority
 - Mark3::BlockingObject, [100](#)
- CS_ENTER
- ClearExpiry
 - Mark3::KernelTimer, [121](#)
- ClearTimer
 - Mark3::Quantum, [160](#)
- ClearNode
 - Mark3::LinkListNode, [128](#)
- ComputeCurrentTicks
 - Mark3::ProfileTimer, [157](#)
- ContextSwitchSWI
 - Mark3::Thread, [177](#)
- CopyData
 - Mark3::Mailbox, [132](#)
- DI
 - Mark3::KernelTimer, [122](#)
- DoubleLinkList
 - Mark3::DoubleLinkList, [106](#)
- EventFlagOperation
 - event_flag_operation_t
 - mark3c.h, [214](#)
- Exit
 - Mark3::Thread, [178](#)
- ExitSimulator
 - Mark3::KernelAware, [95](#)
- GLOBAL_MESSAGE_POOL_SIZE
 - mark3cfg.h, [319](#)
- GetAverage
 - Mark3::ProfileTimer, [158](#)
- GetCode
 - Mark3::Message, [141](#)
- GetCount
 - Mark3::MessageQueue, [146](#)
 - Mark3::Semaphore, [171](#)
- GetCurPriority
 - Mark3::Thread, [178](#)
- GetCurrent
 - Mark3::ProfileTimer, [158](#)
 - Mark3::Thread, [178](#)

- GetCurrentThread
 - Mark3::Scheduler, [167](#)
- GetData
 - Mark3::Message, [141](#)
- GetEventFlagMask
 - Mark3::Thread, [179](#)
- GetEventFlagMode
 - Mark3::Thread, [179](#)
- GetExpired
 - Mark3::Thread, [179](#)
- GetExtendedContext
 - Mark3::Thread, [180](#)
- GetHead
 - Mark3::LinkList, [126](#)
 - Mark3::MessagePool, [143](#)
- GetHeadPointer
 - Mark3::Mailbox, [133](#)
- GetID
 - Mark3::Thread, [180](#)
- GetIdleThread
 - Mark3::Kernel, [115](#)
- GetInterval
 - Mark3::Timer, [198](#)
- GetMask
 - Mark3::EventFlag, [109](#)
- GetNext
 - Mark3::LinkListNode, [128](#)
- GetNextThread
 - Mark3::Scheduler, [167](#)
- GetOvertime
 - Mark3::KernelTimer, [122](#)
- GetOwner
 - Mark3::Thread, [180](#)
- GetPrev
 - Mark3::LinkListNode, [128](#)
- GetPriority
 - Mark3::Thread, [181](#)
- GetQuantum
 - Mark3::Thread, [181](#)
- GetStack
 - Mark3::Thread, [181](#)
- GetStackSize
 - Mark3::Thread, [182](#)
- GetStackSlack
 - Mark3::Thread, [182](#)
- GetState
 - Mark3::Thread, [182](#)
- GetStopList
 - Mark3::Scheduler, [167](#)
- GetTail
 - Mark3::LinkList, [126](#)
- GetTailPointer
 - Mark3::Mailbox, [133](#)
- GetThreadContextSwitchCallout
 - Mark3::Kernel, [115](#)
- GetThreadCreateCallout
 - Mark3::Kernel, [115](#)
- GetThreadExitCallout
 - Mark3::Kernel, [116](#)
- GetThreadList
 - Mark3::Scheduler, [168](#)
- HighestPriority
 - Mark3::PriorityMap, [155](#)
- HighestWaiter
 - Mark3::ThreadList, [193](#)
- InheritPriority
 - Mark3::Thread, [183](#)
- Init
 - Mark3::ConditionVariable, [105](#)
 - Mark3::Kernel, [116](#)
 - Mark3::LinkList, [127](#)
 - Mark3::Mailbox, [133](#), [134](#)
 - Mark3::Message, [141](#)
 - Mark3::MessagePool, [144](#)
 - Mark3::MessageQueue, [146](#)
 - Mark3::Mutex, [150](#)
 - Mark3::Notify, [152](#)
 - Mark3::ProfileTimer, [158](#)
 - Mark3::ReaderWriterLock, [164](#)
 - Mark3::Scheduler, [168](#)
 - Mark3::Semaphore, [172](#)
 - Mark3::Thread, [183](#), [184](#)
 - Mark3::Timer, [198](#)
 - Mark3::TimerList, [205](#)
 - Mark3::TimerScheduler, [207](#)
- InitIdle
 - Mark3::Thread, [184](#)
- InitStack
 - Mark3::ThreadPort, [195](#)
- InsertNodeBefore
 - Mark3::CircularLinkList, [103](#)
- isAcquired
 - Mark3::LockGuard, [130](#)
- IsEnabled
 - Mark3::Scheduler, [168](#)
- IsInitialized
 - Mark3::BlockingObject, [101](#)
 - Mark3::Timer, [198](#)
- IsPanic
 - Mark3::Kernel, [116](#)
- IsSimulatorAware
 - Mark3::KernelAware, [95](#)
- IsStarted
 - Mark3::Kernel, [117](#)
- K_WORD
 - portcfg.h, [238](#)
- KERNEL_AWARE_SIMULATION
 - mark3cfg.h, [320](#)
- KERNEL_EXTRA_CHECKS
 - mark3cfg.h, [320](#)
- KERNEL_NUM_PRIORITIES
 - mark3cfg.h, [320](#)
- KERNEL_TIMERS_MINIMUM_DELAY_US
 - mark3cfg.h, [320](#)

- KERNEL_TIMERS_THREADED
 - mark3cfg.h, [321](#)
- KERNEL_TIMERS_TICKLESS
 - mark3cfg.h, [321](#)
- KERNEL_USE_ATOMIC
 - mark3cfg.h, [321](#)
- KERNEL_USE_AUTO_ALLOC
 - mark3cfg.h, [321](#)
- KERNEL_USE_CONDVAR
 - mark3cfg.h, [322](#)
- KERNEL_USE_DYNAMIC_THREADS
 - mark3cfg.h, [322](#)
- KERNEL_USE_EVENTFLAG
 - mark3cfg.h, [322](#)
- KERNEL_USE_IDLE_FUNC
 - mark3cfg.h, [322](#)
- KERNEL_USE_MAILBOX
 - mark3cfg.h, [323](#)
- KERNEL_USE_MESSAGE
 - mark3cfg.h, [323](#)
- KERNEL_USE_PROFILER
 - mark3cfg.h, [323](#)
- KERNEL_USE_QUANTUM
 - mark3cfg.h, [323](#)
- KERNEL_USE_READERWRITER
 - mark3cfg.h, [324](#)
- KERNEL_USE_SEMAPHORE
 - mark3cfg.h, [324](#)
- KERNEL_USE_STACK_GUARD
 - mark3cfg.h, [324](#)
- KERNEL_USE_THREAD_CALLOUTS
 - mark3cfg.h, [324](#)
- KERNEL_USE_THREADNAME
 - mark3cfg.h, [325](#)
- KERNEL_USE_TIMEOUTS
 - mark3cfg.h, [325](#)
- KERNEL_USE_TIMERS
 - mark3cfg.h, [325](#)
- Kernel_Init
 - mark3c.h, [214](#)
- Kernel_IsPanic
 - mark3c.h, [214](#)
- Kernel_IsStarted
 - mark3c.h, [214](#)
- Kernel_Panic
 - mark3c.h, [215](#)
- Kernel_SetPanic
 - mark3c.h, [215](#)
- Kernel_Start
 - mark3c.h, [216](#)
- LockGuard
 - Mark3::LockGuard, [129](#), [130](#)
- m_clSendSem
 - Mark3::Mailbox, [140](#)
- Mark3, [89](#)
 - EventFlagOperation, [91](#)
 - TimerCallback, [91](#)
- Mark3::Atomic, [92](#)
 - Add, [92](#)
 - Set, [93](#)
 - Sub, [93](#)
 - TestAndSet, [94](#)
- Mark3::BlockingObject, [99](#)
 - Block, [100](#)
 - BlockPriority, [100](#)
 - IsInitialized, [101](#)
 - UnBlock, [101](#)
- Mark3::CircularLinkList, [102](#)
 - Add, [102](#)
 - InsertNodeBefore, [103](#)
 - PivotBackward, [103](#)
 - PivotForward, [103](#)
 - Remove, [103](#)
- Mark3::ConditionVariable, [104](#)
 - Init, [105](#)
 - Wait, [105](#)
- Mark3::DoubleLinkList, [106](#)
 - Add, [107](#)
 - DoubleLinkList, [106](#)
 - Remove, [107](#)
- Mark3::EventFlag, [108](#)
 - Clear, [109](#)
 - GetMask, [109](#)
 - Set, [109](#)
 - Wait, [110](#)
 - Wait_i, [112](#)
 - WakeMe, [112](#)
- Mark3::FakeThread_t, [113](#)
- Mark3::Kernel, [114](#)
 - GetIdleThread, [115](#)
 - GetThreadContextSwitchCallout, [115](#)
 - GetThreadCreateCallout, [115](#)
 - GetThreadExitCallout, [116](#)
 - Init, [116](#)
 - IsPanic, [116](#)
 - IsStarted, [117](#)
 - Panic, [117](#)
 - SetIdleFunc, [117](#)
 - SetPanic, [118](#)
 - SetThreadContextSwitchCallout, [118](#)
 - SetThreadCreateCallout, [119](#)
 - SetThreadExitCallout, [119](#)
 - Start, [120](#)
- Mark3::KernelAware, [94](#)
 - ExitSimulator, [95](#)
 - IsSimulatorAware, [95](#)
 - Print, [95](#)
 - ProfileInit, [96](#)
 - ProfileReport, [96](#)
 - ProfileStart, [96](#)
 - ProfileStop, [96](#)
 - Trace, [97](#), [98](#)
- Mark3::KernelTimer, [120](#)
 - ClearExpiry, [121](#)
 - Config, [121](#)

- DI, [122](#)
- EI, [122](#)
- GetOvertime, [122](#)
- Read, [122](#)
- RI, [123](#)
- SetExpiry, [123](#)
- Start, [124](#)
- Stop, [124](#)
- SubtractExpiry, [124](#)
- TimeToExpiry, [125](#)
- Mark3::LinkList, [125](#)
 - GetHead, [126](#)
 - GetTail, [126](#)
 - Init, [127](#)
- Mark3::LinkListNode, [127](#)
 - ClearNode, [128](#)
 - GetNext, [128](#)
 - GetPrev, [128](#)
- Mark3::LockGuard, [129](#)
 - isAcquired, [130](#)
 - LockGuard, [129](#), [130](#)
- Mark3::Mailbox, [130](#)
 - CopyData, [132](#)
 - GetHeadPointer, [133](#)
 - GetTailPointer, [133](#)
 - Init, [133](#), [134](#)
 - m_clSendSem, [140](#)
 - MoveHeadBackward, [134](#)
 - MoveHeadForward, [134](#)
 - MoveTailBackward, [134](#)
 - MoveTailForward, [135](#)
 - Receive, [135](#)
 - Receive_i, [136](#)
 - ReceiveTail, [136](#), [137](#)
 - Send, [137](#), [138](#)
 - Send_i, [138](#)
 - SendTail, [139](#)
- Mark3::Message, [140](#)
 - GetCode, [141](#)
 - GetData, [141](#)
 - Init, [141](#)
 - SetCode, [142](#)
 - SetData, [142](#)
- Mark3::MessagePool, [143](#)
 - GetHead, [143](#)
 - Init, [144](#)
 - Pop, [144](#)
 - Push, [144](#)
- Mark3::MessageQueue, [145](#)
 - GetCount, [146](#)
 - Init, [146](#)
 - Receive, [146](#)
 - Receive_i, [147](#)
 - Send, [147](#)
- Mark3::Mutex, [148](#)
 - Claim, [149](#)
 - Claim_i, [150](#)
 - Init, [150](#)
 - Release, [150](#)
 - WakeMe, [151](#)
 - WakeNext, [151](#)
- Mark3::Notify, [152](#)
 - Init, [152](#)
 - Signal, [153](#)
 - Wait, [153](#)
 - WakeMe, [154](#)
- Mark3::PriorityMap, [154](#)
 - Clear, [155](#)
 - HighestPriority, [155](#)
 - PriorityMap, [155](#)
 - Set, [156](#)
- Mark3::ProfileTimer, [156](#)
 - ComputeCurrentTicks, [157](#)
 - GetAverage, [158](#)
 - GetCurrent, [158](#)
 - Init, [158](#)
 - Start, [158](#)
 - Stop, [158](#)
- Mark3::Quantum, [159](#)
 - AddThread, [159](#)
 - ClearInTimer, [160](#)
 - RemoveThread, [160](#)
 - SetInTimer, [160](#)
 - SetTimer, [160](#)
 - UpdateTimer, [161](#)
- Mark3::ReaderWriterLock, [161](#)
 - AcquireReader, [162](#)
 - AcquireReader_i, [163](#)
 - AcquireWriter, [163](#)
 - AcquireWriter_i, [164](#)
 - Init, [164](#)
- Mark3::Scheduler, [164](#)
 - Add, [165](#)
 - GetCurrentThread, [167](#)
 - GetNextThread, [167](#)
 - GetStopList, [167](#)
 - GetThreadList, [168](#)
 - Init, [168](#)
 - IsEnabled, [168](#)
 - QueueScheduler, [169](#)
 - Remove, [169](#)
 - Schedule, [169](#)
 - SetScheduler, [170](#)
- Mark3::Semaphore, [170](#)
 - GetCount, [171](#)
 - Init, [172](#)
 - Pend, [172](#), [173](#)
 - Pend_i, [173](#)
 - Post, [173](#)
 - WakeMe, [174](#)
 - WakeNext, [174](#)
- Mark3::Thread, [174](#)
 - ContextSwitchSWI, [177](#)
 - Exit, [178](#)
 - GetCurPriority, [178](#)
 - GetCurrent, [178](#)

- GetEventFlagMask, [179](#)
- GetEventFlagMode, [179](#)
- GetExpired, [179](#)
- GetExtendedContext, [180](#)
- GetID, [180](#)
- GetOwner, [180](#)
- GetPriority, [181](#)
- GetQuantum, [181](#)
- GetStack, [181](#)
- GetStackSize, [182](#)
- GetStackSlack, [182](#)
- GetState, [182](#)
- InheritPriority, [183](#)
- Init, [183](#), [184](#)
- InitIdle, [184](#)
- SetCurrent, [185](#)
- SetEventFlagMask, [185](#)
- SetEventFlagMode, [185](#)
- SetExpired, [186](#)
- SetExtendedContext, [186](#)
- SetID, [186](#)
- SetOwner, [187](#)
- SetPriority, [187](#)
- SetPriorityBase, [187](#)
- SetQuantum, [188](#)
- SetState, [188](#)
- Sleep, [189](#)
- Start, [189](#)
- Stop, [189](#)
- USleep, [189](#)
- Yield, [190](#)
- Mark3::ThreadList, [190](#)
 - Add, [192](#)
 - AddPriority, [192](#)
 - HighestWaiter, [193](#)
 - Remove, [193](#)
 - SetMapPointer, [193](#)
 - SetPriority, [194](#)
 - ThreadList, [191](#)
- Mark3::ThreadPort, [194](#)
 - InitStack, [195](#)
 - StartThreads, [195](#)
- Mark3::Timer, [196](#)
 - GetInterval, [198](#)
 - Init, [198](#)
 - IsInitialized, [198](#)
 - SetCallback, [198](#)
 - SetData, [199](#)
 - SetFlags, [199](#)
 - SetIntervalMSeconds, [199](#)
 - SetIntervalSeconds, [200](#)
 - SetIntervalTicks, [200](#)
 - SetIntervalUSeconds, [200](#)
 - SetOwner, [201](#)
 - SetTolerance, [201](#)
 - Start, [201](#), [203](#)
 - Stop, [203](#)
 - Timer, [197](#)
- Mark3::TimerList, [204](#)
 - Add, [205](#)
 - Init, [205](#)
 - Process, [205](#)
 - Remove, [205](#)
- Mark3::TimerScheduler, [206](#)
 - Add, [206](#)
 - Init, [207](#)
 - Process, [207](#)
 - Remove, [207](#)
- mark3c.h
 - event_flag_operation_t, [214](#)
 - Kernel_Init, [214](#)
 - Kernel_IsPanic, [214](#)
 - Kernel_IsStarted, [214](#)
 - Kernel_Panic, [215](#)
 - Kernel_SetPanic, [215](#)
 - Kernel_Start, [216](#)
 - Scheduler_Enable, [216](#)
 - Scheduler_GetCurrentThread, [216](#)
 - Scheduler_IsEnabled, [217](#)
 - Thread_GetCurPriority, [217](#)
 - Thread_GetID, [218](#)
 - Thread_GetPriority, [218](#)
 - Thread_GetStackSlack, [218](#)
 - Thread_GetState, [219](#)
 - Thread_Init, [219](#)
 - Thread_SetID, [220](#)
 - Thread_SetPriority, [221](#)
 - Thread_Start, [221](#)
 - Thread_Stop, [221](#)
 - Thread_Yield, [222](#)
- mark3cfg.h
 - GLOBAL_MESSAGE_POOL_SIZE, [319](#)
 - KERNEL_AWARE_SIMULATION, [320](#)
 - KERNEL_EXTRA_CHECKS, [320](#)
 - KERNEL_NUM_PRIORITIES, [320](#)
 - KERNEL_TIMERS_MINIMUM_DELAY_US, [320](#)
 - KERNEL_TIMERS_THREADED, [321](#)
 - KERNEL_TIMERS_TICKLESS, [321](#)
 - KERNEL_USE_ATOMIC, [321](#)
 - KERNEL_USE_AUTO_ALLOC, [321](#)
 - KERNEL_USE_CONDVAR, [322](#)
 - KERNEL_USE_DYNAMIC_THREADS, [322](#)
 - KERNEL_USE_EVENTFLAG, [322](#)
 - KERNEL_USE_IDLE_FUNC, [322](#)
 - KERNEL_USE_MAILBOX, [323](#)
 - KERNEL_USE_MESSAGE, [323](#)
 - KERNEL_USE_PROFILER, [323](#)
 - KERNEL_USE_QUANTUM, [323](#)
 - KERNEL_USE_READERWRITER, [324](#)
 - KERNEL_USE_SEMAPHORE, [324](#)
 - KERNEL_USE_STACK_GUARD, [324](#)
 - KERNEL_USE_THREAD_CALLOUTS, [324](#)
 - KERNEL_USE_THREADNAME, [325](#)
 - KERNEL_USE_TIMEOUTS, [325](#)
 - KERNEL_USE_TIMERS, [325](#)
 - SAFE_UNLINK, [325](#)

- THREAD_QUANTUM_DEFAULT, 326
- MoveHeadBackward
 - Mark3::Mailbox, 134
- MoveHeadForward
 - Mark3::Mailbox, 134
- MoveTailBackward
 - Mark3::Mailbox, 134
- MoveTailForward
 - Mark3::Mailbox, 135
- PORT_PRIO_TYPE
 - portcfg.h, 239
- PORT_SYSTEM_FREQ
 - portcfg.h, 239
- PORT_TIMER_COUNT_TYPE
 - portcfg.h, 239
- PORT_TIMER_FREQ
 - portcfg.h, 239
- Panic
 - Mark3::Kernel, 117
- Pend
 - Mark3::Semaphore, 172, 173
- Pend_i
 - Mark3::Semaphore, 173
- PivotBackward
 - Mark3::CircularLinkList, 103
- PivotForward
 - Mark3::CircularLinkList, 103
- Pop
 - Mark3::MessagePool, 144
- portcfg.h
 - AVR, 238
 - K_WORD, 238
 - PORT_PRIO_TYPE, 239
 - PORT_SYSTEM_FREQ, 239
 - PORT_TIMER_COUNT_TYPE, 239
 - PORT_TIMER_FREQ, 239
- Post
 - Mark3::Semaphore, 173
- Print
 - Mark3::KernelAware, 95
- PriorityMap
 - Mark3::PriorityMap, 155
- Process
 - Mark3::TimerList, 205
 - Mark3::TimerScheduler, 207
- ProfileInit
 - Mark3::KernelAware, 96
- ProfileReport
 - Mark3::KernelAware, 96
- ProfileStart
 - Mark3::KernelAware, 96
- ProfileStop
 - Mark3::KernelAware, 96
- Push
 - Mark3::MessagePool, 144
- QueueScheduler
 - Mark3::Scheduler, 169
- Read
 - Mark3::KernelTimer, 122
- Receive
 - Mark3::Mailbox, 135
 - Mark3::MessageQueue, 146
- Receive_i
 - Mark3::Mailbox, 136
 - Mark3::MessageQueue, 147
- ReceiveTail
 - Mark3::Mailbox, 136, 137
- Release
 - Mark3::Mutex, 150
- Remove
 - Mark3::CircularLinkList, 103
 - Mark3::DoubleLinkList, 107
 - Mark3::Scheduler, 169
 - Mark3::ThreadList, 193
 - Mark3::TimerList, 205
 - Mark3::TimerScheduler, 207
- RemoveThread
 - Mark3::Quantum, 160
- RI
 - Mark3::KernelTimer, 123
- SAFE_UNLINK
 - mark3cfg.h, 325
- Schedule
 - Mark3::Scheduler, 169
- Scheduler_Enable
 - mark3c.h, 216
- Scheduler_GetCurrentThread
 - mark3c.h, 216
- Scheduler_IsEnabled
 - mark3c.h, 217
- Send
 - Mark3::Mailbox, 137, 138
 - Mark3::MessageQueue, 147
- Send_i
 - Mark3::Mailbox, 138
- SendTail
 - Mark3::Mailbox, 139
- Set
 - Mark3::Atomic, 93
 - Mark3::EventFlag, 109
 - Mark3::PriorityMap, 156
- SetCallback
 - Mark3::Timer, 198
- SetCode
 - Mark3::Message, 142
- SetCurrent
 - Mark3::Thread, 185
- SetData
 - Mark3::Message, 142
 - Mark3::Timer, 199
- SetEventFlagMask
 - Mark3::Thread, 185
- SetEventFlagMode
 - Mark3::Thread, 185
- SetExpired

- Mark3::Thread, [186](#)
- SetExpiry
 - Mark3::KernelTimer, [123](#)
- SetExtendedContext
 - Mark3::Thread, [186](#)
- SetFlags
 - Mark3::Timer, [199](#)
- SetID
 - Mark3::Thread, [186](#)
- SetIdleFunc
 - Mark3::Kernel, [117](#)
- SetInTimer
 - Mark3::Quantum, [160](#)
- SetIntervalMSeconds
 - Mark3::Timer, [199](#)
- SetIntervalSeconds
 - Mark3::Timer, [200](#)
- SetIntervalTicks
 - Mark3::Timer, [200](#)
- SetIntervalUSeconds
 - Mark3::Timer, [200](#)
- SetMapPointer
 - Mark3::ThreadList, [193](#)
- SetOwner
 - Mark3::Thread, [187](#)
 - Mark3::Timer, [201](#)
- SetPanic
 - Mark3::Kernel, [118](#)
- SetPriority
 - Mark3::Thread, [187](#)
 - Mark3::ThreadList, [194](#)
- SetPriorityBase
 - Mark3::Thread, [187](#)
- SetQuantum
 - Mark3::Thread, [188](#)
- SetScheduler
 - Mark3::Scheduler, [170](#)
- SetState
 - Mark3::Thread, [188](#)
- SetThreadContextSwitchCallout
 - Mark3::Kernel, [118](#)
- SetThreadCreateCallout
 - Mark3::Kernel, [119](#)
- SetThreadExitCallout
 - Mark3::Kernel, [119](#)
- SetTimer
 - Mark3::Quantum, [160](#)
- SetTolerance
 - Mark3::Timer, [201](#)
- Signal
 - Mark3::Notify, [153](#)
- Sleep
 - Mark3::Thread, [189](#)
- Start
 - Mark3::Kernel, [120](#)
 - Mark3::KernelTimer, [124](#)
 - Mark3::ProfileTimer, [158](#)
 - Mark3::Thread, [189](#)
 - Mark3::Timer, [201](#), [203](#)
- StartThreads
 - Mark3::ThreadPort, [195](#)
- Stop
 - Mark3::KernelTimer, [124](#)
 - Mark3::ProfileTimer, [158](#)
 - Mark3::Thread, [189](#)
 - Mark3::Timer, [203](#)
- Sub
 - Mark3::Atomic, [93](#)
- SubtractExpiry
 - Mark3::KernelTimer, [124](#)
- THREAD_QUANTUM_DEFAULT
 - mark3cfg.h, [326](#)
- TIMERLIST_FLAG_EXPIRED
 - timer.h, [347](#)
- TestAndSet
 - Mark3::Atomic, [94](#)
- Thread_GetCurPriority
 - mark3c.h, [217](#)
- Thread_GetID
 - mark3c.h, [218](#)
- Thread_GetPriority
 - mark3c.h, [218](#)
- Thread_GetStackSlack
 - mark3c.h, [218](#)
- Thread_GetState
 - mark3c.h, [219](#)
- Thread_Init
 - mark3c.h, [219](#)
- Thread_SetID
 - mark3c.h, [220](#)
- Thread_SetPriority
 - mark3c.h, [221](#)
- Thread_Start
 - mark3c.h, [221](#)
- Thread_Stop
 - mark3c.h, [221](#)
- Thread_Yield
 - mark3c.h, [222](#)
- ThreadList
 - Mark3::ThreadList, [191](#)
- threadport.h
 - CS_ENTER, [241](#)
- TimeToExpiry
 - Mark3::KernelTimer, [125](#)
- Timer
 - Mark3::Timer, [197](#)
- timer.h
 - TIMERLIST_FLAG_EXPIRED, [347](#)
- TimerCallback
 - Mark3, [91](#)
- Trace
 - Mark3::KernelAware, [97](#), [98](#)
- USleep
 - Mark3::Thread, [189](#)
- UnBlock

- Mark3::BlockingObject, [101](#)
- UpdateTimer
 - Mark3::Quantum, [161](#)
- Wait
 - Mark3::ConditionVariable, [105](#)
 - Mark3::EventFlag, [110](#)
 - Mark3::Notify, [153](#)
- Wait_i
 - Mark3::EventFlag, [112](#)
- WakeMe
 - Mark3::EventFlag, [112](#)
 - Mark3::Mutex, [151](#)
 - Mark3::Notify, [154](#)
 - Mark3::Semaphore, [174](#)
- WakeNext
 - Mark3::Mutex, [151](#)
 - Mark3::Semaphore, [174](#)
- Yield
 - Mark3::Thread, [190](#)