

Mark3 Realtime Kernel

Generated by Doxygen 1.8.13

Contents

1	The Mark3 Realtime Kernel	1
2	License	3
2.1	License	3
3	Configuring The Kernel	5
3.1	Overview	5
3.1.1	Kernel Configuration Options	5
3.1.2	Port Configuration Options	6
4	Building The Kernel	9
4.1	Prerequisites	9
4.2	Building	9
4.3	Supported Targets	10
4.4	More Info	10
5	The Mark3 API	11
5.1	Kernel Setup	11
5.2	Threads	12
5.2.1	Thread Setup	12
5.2.2	Entry Functions	13
5.3	Timers	14
5.4	Semaphores	14
5.5	Mutexes	15
5.6	Event Flags	16

5.7	Messages	16
5.7.1	Message Objects	17
5.7.2	Message Queues	17
5.7.3	Messaging Example	18
5.8	Mailboxes	18
5.8.1	Mailbox Example	18
5.9	Notification Objects	19
5.9.1	Notification Example	19
5.10	Condition Variables	20
5.10.1	Condition Variable Example	20
5.11	Reader-Write Locks	21
5.11.1	Reader-Write Lock Example	21
5.12	Sleep	21
5.13	Round-Robin Quantum	22
5.14	Coroutines	22
5.15	Critical Guards	23
5.16	Lock Guards	23
5.17	Scheduler Guards	24
6	Why Mark3?	25
7	When to use an RTOS?	27
7.1	The reality of system code	27
7.2	Superloops, and their limitations	28
7.2.1	Intro to Superloops	28
7.2.2	The simplest loop	29
7.2.3	Interrupt-Driven Super-loop	29
7.2.4	Cooperative multi-tasking	31
7.2.5	Hybrid cooperative/preemptive multi-tasking	32
7.3	Problems with superloops	33
7.3.1	Hidden Costs	33
7.3.2	Tightly-coupled code	33
7.3.3	No blocking Calls	33
7.3.4	Difficult to guarantee responsiveness	33
7.3.5	Limited preemption capability	34

8	Can you afford an RTOS?	35
8.1	Intro	35
8.2	Application description	36
8.3	Runtime Overhead	37
8.4	Analysis	38
9	Mark3 Design Goals	39
9.1	Overview	39
9.1.1	Services Provided by an RTOS Kernel	39
9.1.2	Guiding Principles of Mark3	39
9.1.3	Be feature competitive	39
9.1.4	No external dependencies, no non-language features	39
9.1.5	Target the most popular hobbyist platforms available	40
9.1.6	Maximize determinism – but be pragmatic	40
9.1.7	Apply engineering principles – and that means discipline, measurement and verification	40
9.1.8	Use Virtualization For Verification	41
10	Mark3 Kernel Architecture	43
10.1	Overview	43
10.2	Threads and Scheduling	46
10.2.1	A Bit About Threads	46
10.2.2	Thread States and ThreadLists	47
10.2.3	Blocking and Unblocking	47
10.2.4	Blocking Objects	47
10.3	Inside the Mark3 Scheduler	48
10.3.1	Considerations for Round-Robin Scheduling	50
10.3.2	Context Switching	51
10.3.3	Putting It All Together	51
10.4	Timers	51
10.4.1	Tick-based Timers	53
10.4.2	Tickless Timers	53

10.4.3 Timer Processing Algorithm	53
10.5 Synchronization and IPC	54
10.6 Blocking Objects	54
10.6.1 Semaphores	55
10.6.2 Mutex	55
10.6.3 Event Flags	56
10.6.4 Notification Objects	56
10.7 Messages and Message Queues	56
10.7.1 Messages	56
10.7.2 Message Objects	57
10.7.3 Message Queues	57
10.7.4 Mailboxes	57
10.7.5 Atomic Operations	58
10.7.6 Drivers	59
10.8 Kernel Proper and Porting	61
11 C-language bindings	69
11.1 API Conventions	69
11.2 Allocating Objects	70
12 Release Notes	71
12.1 R10 Release	71
12.2 R9 Release	71
12.3 R8 Release	71
12.4 R7 (Full Throttle) Release	72
12.5 R6 Release	72
12.6 R5 Release	73
12.7 R4 Release	74
12.8 R3 Release	74
12.9 R2	74
12.10R1 - 2nd Release Candidate	75
12.11R1 - 1st Release Candidate	75

13 Code Size Profiling	77
13.1 Information	77
13.2 Compiler Version	77
13.3 Profiling Results	77
14 Namespace Index	79
14.1 Namespace List	79
15 Hierarchical Index	81
15.1 Class Hierarchy	81
16 Class Index	83
16.1 Class List	83
17 File Index	87
17.1 File List	87
18 Namespace Documentation	91
18.1 Mark3 Namespace Reference	91
18.1.1 Typedef Documentation	95
18.1.1.1 AutoAllocAllocator_t	95
18.1.1.2 AutoAllocFree_t	95
18.1.1.3 CoPrioMap	96
18.1.1.4 CoroutineHandler	96
18.1.1.5 IdleFunc	96
18.1.1.6 PanicFunc	96
18.1.1.7 PriorityMap	96
18.1.1.8 ThreadContextCallout	96
18.1.1.9 ThreadCreateCallout	97
18.1.1.10 ThreadEntryFunc	97
18.1.1.11 ThreadExitCallout	97
18.1.1.12 TimerCallback	97
18.1.2 Enumeration Type Documentation	97

18.1.2.1	AutoAllocType	97
18.1.2.2	EventFlagOperation	98
18.1.2.3	ThreadState	98
18.1.3	Function Documentation	99
18.1.3.1	ISR()	99
18.1.3.2	KernelTimer_Task()	99
18.1.3.3	PORT_CLZ()	99
18.1.3.4	PORT_CS_ENTER()	99
18.1.3.5	PORT_CS_EXIT()	99
18.1.3.6	PORT_CS_NESTING()	100
18.1.3.7	PORT_IRQ_DISABLE()	100
18.1.3.8	PORT_IRQ_ENABLE()	100
18.1.3.9	Thread_Switch()	100
18.1.4	Variable Documentation	100
18.1.4.1	g_kwCriticalCount	100
18.1.4.2	g_kwSFR	100
18.1.4.3	naked	101
18.1.4.4	SR_	101
18.1.4.5	uMaxTimerTicks	101
18.1.4.6	uTimerFlagActive	101
18.1.4.7	uTimerFlagCallback	101
18.1.4.8	uTimerFlagExpired	102
18.1.4.9	uTimerFlagOneShot	102
18.1.4.10	uTimerTicksInvalid	102
18.2	Mark3::Atomic Namespace Reference	102
18.2.1	Detailed Description	103
18.2.2	Function Documentation	103
18.2.2.1	Add()	103
18.2.2.2	Set()	103
18.2.2.3	Sub()	104
18.2.2.4	TestAndSet()	104

19 Class Documentation	105
19.1 Mark3::AutoAlloc Class Reference	105
19.1.1 Detailed Description	106
19.1.2 Member Function Documentation	106
19.1.2.1 Allocate()	106
19.1.2.2 DestroyObject()	106
19.1.2.3 DestroyRawData()	106
19.1.2.4 DestroyUserTypeAllocation()	107
19.1.2.5 Free()	107
19.1.2.6 Init()	107
19.1.2.7 NewObject()	108
19.1.2.8 NewRawData()	108
19.1.2.9 NewUserTypeAllocation()	108
19.1.2.10 SetAllocatorFunctions()	109
19.1.3 Member Data Documentation	109
19.1.3.1 m_pfAllocator	109
19.1.3.2 m_pfFree	109
19.2 Mark3::BlockingObject Class Reference	110
19.2.1 Detailed Description	111
19.2.2 Constructor & Destructor Documentation	111
19.2.2.1 BlockingObject()	111
19.2.2.2 ~BlockingObject()	111
19.2.3 Member Function Documentation	111
19.2.3.1 Block()	111
19.2.3.2 BlockPriority()	112
19.2.3.3 IsInitialized()	112
19.2.3.4 SetInitialized()	112
19.2.3.5 Unblock()	113
19.2.4 Member Data Documentation	114
19.2.4.1 m_clBlockList	114

19.2.4.2	<code>m_u8Initialized</code>	114
19.2.4.3	<code>m_uBlockingInitCookie</code>	114
19.2.4.4	<code>m_uBlockingInvalidCookie</code>	114
19.3	Mark3::CircularLinkedList Class Reference	115
19.3.1	Detailed Description	115
19.3.2	Constructor & Destructor Documentation	115
19.3.2.1	<code>CircularLinkedList()</code>	115
19.3.3	Member Function Documentation	116
19.3.3.1	<code>Add()</code>	116
19.3.3.2	<code>InsertNodeBefore()</code>	116
19.3.3.3	<code>operator new()</code>	116
19.3.3.4	<code>PivotBackward()</code>	117
19.3.3.5	<code>PivotForward()</code>	117
19.3.3.6	<code>Remove()</code>	117
19.4	Mark3::CoList Class Reference	117
19.4.1	Detailed Description	118
19.4.2	Member Function Documentation	118
19.4.2.1	<code>Add()</code>	118
19.4.2.2	<code>Remove()</code>	119
19.4.2.3	<code>SetPrioMap()</code>	119
19.4.2.4	<code>SetPriority()</code>	119
19.4.3	Member Data Documentation	120
19.4.3.1	<code>m_pclPrioMap</code>	120
19.4.3.2	<code>m_uPriority</code>	120
19.5	Mark3::ConditionVariable Class Reference	120
19.5.1	Detailed Description	121
19.5.2	Member Function Documentation	121
19.5.2.1	<code>Broadcast()</code>	121
19.5.2.2	<code>Init()</code>	122
19.5.2.3	<code>operator new()</code>	122

19.5.2.4	Signal()	122
19.5.2.5	Wait() [1/2]	122
19.5.2.6	Wait() [2/2]	123
19.5.3	Member Data Documentation	123
19.5.3.1	m_clMutex	123
19.5.3.2	m_clSemaphore	123
19.5.3.3	m_u8Waiters	123
19.6	Mark3::Coroutine Class Reference	124
19.6.1	Detailed Description	124
19.6.2	Constructor & Destructor Documentation	125
19.6.2.1	~Coroutine()	125
19.6.3	Member Function Documentation	125
19.6.3.1	Activate()	125
19.6.3.2	GetPriority()	125
19.6.3.3	Init()	125
19.6.3.4	Run()	126
19.6.3.5	SetPriority()	126
19.6.4	Member Data Documentation	126
19.6.4.1	m_bQueued	126
19.6.4.2	m_pclOwner	127
19.6.4.3	m_pfHandler	127
19.6.4.4	m_pvContext	127
19.6.4.5	m_uPriority	127
19.7	Mark3::CoScheduler Class Reference	127
19.7.1	Detailed Description	128
19.7.2	Member Function Documentation	128
19.7.2.1	GetCoList()	128
19.7.2.2	GetPrioMap()	129
19.7.2.3	GetStopList()	129
19.7.2.4	Init()	129

19.7.2.5	Schedule()	130
19.7.3	Member Data Documentation	130
19.7.3.1	m_aclPriorities	130
19.7.3.2	m_clPrioMap	130
19.7.3.3	m_clStopList	130
19.8	Mark3::CriticalGuard Class Reference	131
19.8.1	Detailed Description	131
19.8.2	Constructor & Destructor Documentation	131
19.8.2.1	CriticalGuard()	131
19.8.2.2	~CriticalGuard()	131
19.9	Mark3::CriticalSection Class Reference	132
19.9.1	Detailed Description	132
19.9.2	Member Function Documentation	132
19.9.2.1	Enter()	132
19.9.2.2	Exit()	133
19.9.2.3	NestingCount()	133
19.10	Mark3::DoubleLinkedList Class Reference	133
19.10.1	Detailed Description	134
19.10.2	Constructor & Destructor Documentation	134
19.10.2.1	DoubleLinkedList()	134
19.10.3	Member Function Documentation	134
19.10.3.1	Add()	134
19.10.3.2	operator new()	134
19.10.3.3	Remove()	135
19.11	Mark3::EventFlag Class Reference	135
19.11.1	Detailed Description	136
19.11.2	Constructor & Destructor Documentation	137
19.11.2.1	~EventFlag()	137
19.11.3	Member Function Documentation	137
19.11.3.1	Clear()	137

19.11.3.2 GetMask()	137
19.11.3.3 Init()	137
19.11.3.4 operator new()	138
19.11.3.5 Set()	138
19.11.3.6 Wait() [1/2]	138
19.11.3.7 Wait() [2/2]	139
19.11.3.8 Wait_i()	139
19.11.3.9 WakeMe()	140
19.11.4 Member Data Documentation	140
19.11.4.1 m_u16SetMask	140
19.12Mark3::Kernel Class Reference	140
19.12.1 Detailed Description	142
19.12.2 Member Function Documentation	142
19.12.2.1 CompleteStart()	142
19.12.2.2 DebugPrint()	142
19.12.2.3 GetStackGuardThreshold()	143
19.12.2.4 GetThreadContextSwitchCallout()	143
19.12.2.5 GetThreadCreateCallout()	143
19.12.2.6 GetThreadExitCallout()	143
19.12.2.7 GetTicks()	144
19.12.2.8 Init()	144
19.12.2.9 IsPanic()	144
19.12.2.10IsStarted()	144
19.12.2.11Panic()	144
19.12.2.12SetDebugPrintFunction()	145
19.12.2.13SetPanic()	145
19.12.2.14SetStackGuardThreshold()	146
19.12.2.15SetThreadContextSwitchCallout()	146
19.12.2.16SetThreadCreateCallout()	146
19.12.2.17SetThreadExitCallout()	147

19.12.2.18	Start()	147
19.12.2.19	Tick()	148
19.12.3	Member Data Documentation	148
19.12.3.1	m_blsPanic	148
19.12.3.2	m_blsStarted	148
19.12.3.3	m_pfDebugPrintFunction	148
19.12.3.4	m_pfPanic	148
19.12.3.5	m_pfThreadContextCallout	149
19.12.3.6	m_pfThreadCreateCallout	149
19.12.3.7	m_pfThreadExitCallout	149
19.12.3.8	m_u16GuardThreshold	149
19.12.3.9	m_u32Ticks	149
19.13	Mark3::KernelSWI Class Reference	150
19.13.1	Detailed Description	150
19.13.2	Member Function Documentation	150
19.13.2.1	Config()	150
19.13.2.2	Start()	150
19.13.2.3	Trigger()	151
19.14	Mark3::KernelTimer Class Reference	151
19.14.1	Detailed Description	151
19.14.2	Member Function Documentation	151
19.14.2.1	Config()	151
19.14.2.2	Start()	152
19.14.2.3	Stop()	152
19.15	Mark3::LinkList Class Reference	152
19.15.1	Detailed Description	153
19.15.2	Member Function Documentation	153
19.15.2.1	GetHead()	153
19.15.2.2	GetTail()	153
19.15.2.3	Init()	154

19.15.2.4 SetHead()	154
19.15.2.5 SetTail()	154
19.15.3 Member Data Documentation	154
19.15.3.1 m_pclHead	155
19.15.3.2 m_pclTail	155
19.16Mark3::LinkedList Class Reference	155
19.16.1 Detailed Description	156
19.16.2 Constructor & Destructor Documentation	156
19.16.2.1 LinkedList()	156
19.16.3 Member Function Documentation	156
19.16.3.1 ClearNode()	156
19.16.3.2 GetNext()	157
19.16.3.3 GetPrev()	157
19.16.4 Friends And Related Function Documentation	157
19.16.4.1 CircularLinkedList	157
19.16.4.2 DoubleLinkedList	157
19.16.4.3 LinkedList	158
19.16.5 Member Data Documentation	158
19.16.5.1 next	158
19.16.5.2 prev	158
19.17Mark3::LockGuard Class Reference	158
19.17.1 Detailed Description	159
19.17.2 Constructor & Destructor Documentation	159
19.17.2.1 LockGuard() [1/2]	159
19.17.2.2 LockGuard() [2/2]	159
19.17.2.3 ~LockGuard()	160
19.17.3 Member Function Documentation	160
19.17.3.1 isAcquired()	160
19.17.4 Member Data Documentation	160
19.17.4.1 m_bIsAcquired	160

19.17.4.2 m_pclMutex	160
19.18Mark3::Mailbox Class Reference	161
19.18.1 Detailed Description	162
19.18.2 Constructor & Destructor Documentation	163
19.18.2.1 ~Mailbox()	163
19.18.3 Member Function Documentation	163
19.18.3.1 CopyData()	163
19.18.3.2 GetFreeSlots()	163
19.18.3.3 GetHeadPointer()	164
19.18.3.4 GetTailPointer()	164
19.18.3.5 Init() [1/2]	164
19.18.3.6 Init() [2/2]	165
19.18.3.7 IsEmpty()	165
19.18.3.8 IsFull()	165
19.18.3.9 MoveHeadBackward()	165
19.18.3.10MoveHeadForward()	166
19.18.3.11MoveTailBackward()	166
19.18.3.12MoveTailForward()	166
19.18.3.13operator new()	166
19.18.3.14Receive() [1/2]	166
19.18.3.15Receive() [2/2]	167
19.18.3.16Receive_i()	167
19.18.3.17ReceiveTail() [1/2]	168
19.18.3.18ReceiveTail() [2/2]	168
19.18.3.19Send() [1/2]	169
19.18.3.20Send() [2/2]	169
19.18.3.21Send_i()	170
19.18.3.22SendTail() [1/2]	170
19.18.3.23SendTail() [2/2]	171
19.18.4 Member Data Documentation	171

19.18.4.1 m_clRecvSem	171
19.18.4.2 m_clSendSem	171
19.18.4.3 m_pvBuffer	172
19.18.4.4 m_u16Count	172
19.18.4.5 m_u16ElementSize	172
19.18.4.6 m_u16Free	172
19.18.4.7 m_u16Head	172
19.18.4.8 m_u16Tail	173
19.19Mark3::MemUtil Class Reference	173
19.19.1 Detailed Description	174
19.19.2 Member Function Documentation	174
19.19.2.1 Checksum16()	174
19.19.2.2 Checksum8()	175
19.19.2.3 CompareMemory()	175
19.19.2.4 CompareStrings() [1/2]	176
19.19.2.5 CompareStrings() [2/2]	176
19.19.2.6 CopyMemory()	176
19.19.2.7 CopyString()	177
19.19.2.8 DecimalToHex() [1/4]	177
19.19.2.9 DecimalToHex() [2/4]	178
19.19.2.10DecimalToHex() [3/4]	178
19.19.2.11DecimalToHex() [4/4]	178
19.19.2.12DecimalToString() [1/4]	178
19.19.2.13DecimalToString() [2/4]	179
19.19.2.14DecimalToString() [3/4]	179
19.19.2.15DecimalToString() [4/4]	179
19.19.2.16SetMemory()	179
19.19.2.17StringLength()	180
19.19.2.18StringSearch()	180
19.19.2.19StringToDecimal16()	181

19.19.2.20	StringToDecimal32()	181
19.19.2.21	StringToDecimal64()	181
19.19.2.22	StringToDecimal8()	181
19.19.2.23	Tokenize()	182
19.20	Mark3::Message Class Reference	182
19.20.1	Detailed Description	183
19.20.2	Member Function Documentation	183
19.20.2.1	GetCode()	183
19.20.2.2	GetData()	184
19.20.2.3	Init()	184
19.20.2.4	operator new()	184
19.20.2.5	SetCode()	184
19.20.2.6	SetData()	185
19.20.3	Member Data Documentation	185
19.20.3.1	m_pvData	185
19.20.3.2	m_u16Code	185
19.21	Mark3::MessagePool Class Reference	186
19.21.1	Detailed Description	186
19.21.2	Constructor & Destructor Documentation	186
19.21.2.1	~MessagePool()	186
19.21.3	Member Function Documentation	187
19.21.3.1	GetHead()	187
19.21.3.2	Init()	187
19.21.3.3	operator new()	187
19.21.3.4	Pop()	188
19.21.3.5	Push()	188
19.21.4	Member Data Documentation	188
19.21.4.1	m_cList	189
19.22	Mark3::MessageQueue Class Reference	189
19.22.1	Detailed Description	190

19.22.2 Constructor & Destructor Documentation	190
19.22.2.1 ~MessageQueue()	190
19.22.3 Member Function Documentation	190
19.22.3.1 GetCount()	190
19.22.3.2 Init()	190
19.22.3.3 operator new()	191
19.22.3.4 Receive() [1/2]	191
19.22.3.5 Receive() [2/2]	191
19.22.3.6 Receive_i()	192
19.22.3.7 Send()	192
19.22.4 Member Data Documentation	192
19.22.4.1 m_clLinkList	193
19.22.4.2 m_clSemaphore	193
19.23Mark3::Mutex Class Reference	193
19.23.1 Detailed Description	194
19.23.2 Constructor & Destructor Documentation	194
19.23.2.1 ~Mutex()	194
19.23.3 Member Function Documentation	194
19.23.3.1 Claim() [1/2]	195
19.23.3.2 Claim() [2/2]	195
19.23.3.3 Claim_i()	195
19.23.3.4 Init()	196
19.23.3.5 operator new()	196
19.23.3.6 Release()	196
19.23.3.7 WakeMe()	197
19.23.3.8 WakeNext()	197
19.23.4 Member Data Documentation	197
19.23.4.1 m_bReady	197
19.23.4.2 m_bRecursive	198
19.23.4.3 m_pclOwner	198

19.23.4.4 m_u8Recurse	198
19.23.4.5 m_uMaxPri	198
19.24Mark3::Notify Class Reference	199
19.24.1 Detailed Description	199
19.24.2 Constructor & Destructor Documentation	200
19.24.2.1 ~Notify()	200
19.24.3 Member Function Documentation	200
19.24.3.1 Init()	200
19.24.3.2 operator new()	200
19.24.3.3 Signal()	200
19.24.3.4 Wait() [1/2]	200
19.24.3.5 Wait() [2/2]	201
19.24.3.6 WakeMe()	201
19.24.4 Member Data Documentation	202
19.24.4.1 m_bPending	202
19.25Mark3::PriorityMapL1< T, C > Class Template Reference	202
19.25.1 Detailed Description	203
19.25.2 Constructor & Destructor Documentation	203
19.25.2.1 PriorityMapL1()	203
19.25.3 Member Function Documentation	203
19.25.3.1 Clear()	203
19.25.3.2 HighestPriority()	204
19.25.3.3 PrioBit()	204
19.25.3.4 PrioMapWordIndex()	204
19.25.3.5 PriorityFromBitmap()	204
19.25.3.6 Set()	204
19.25.4 Member Data Documentation	205
19.25.4.1 m_uXPrioMapBitMask	205
19.25.4.2 m_uXPrioMapBits	205
19.25.4.3 m_uXPrioMapShiftLUT	205

19.25.4.4 m_uXPrioMapWordShift	205
19.25.4.5 m_uXPriorityMap	206
19.26Mark3::PriorityMapL2< T, C > Class Template Reference	206
19.26.1 Detailed Description	207
19.26.2 Constructor & Destructor Documentation	207
19.26.2.1 PriorityMapL2()	207
19.26.3 Member Function Documentation	207
19.26.3.1 Clear()	207
19.26.3.2 HighestPriority()	208
19.26.3.3 PrioBit()	208
19.26.3.4 PrioMapWordIndex()	208
19.26.3.5 PriorityFromBitmap()	208
19.26.3.6 Set()	208
19.26.4 Member Data Documentation	209
19.26.4.1 m_auXPriorityMap	209
19.26.4.2 m_uXPrioMapBitMask	209
19.26.4.3 m_uXPrioMapBits	209
19.26.4.4 m_uXPrioMapNumWords	209
19.26.4.5 m_uXPrioMapShiftLUT	210
19.26.4.6 m_uXPrioMapWordShift	210
19.26.4.7 m_uXPriorityMapL2	210
19.27Mark3::ProfileTimer Class Reference	210
19.27.1 Detailed Description	211
19.27.2 Member Function Documentation	211
19.27.2.1 GetAverage()	211
19.27.2.2 GetCurrent()	211
19.27.2.3 Init()	212
19.27.2.4 Start()	212
19.27.2.5 Stop()	212
19.27.3 Member Data Documentation	212

19.27.3.1 m_bActive	212
19.27.3.2 m_u16Iterations	213
19.27.3.3 m_u32Cumulative	213
19.27.3.4 m_u32CurrentIteration	213
19.27.3.5 m_u32StartTicks	213
19.28Mark3::Quantum Class Reference	213
19.28.1 Detailed Description	214
19.28.2 Member Function Documentation	214
19.28.2.1 Cancel()	214
19.28.2.2 ClearInTimer()	214
19.28.2.3 Init()	215
19.28.2.4 SetInTimer()	215
19.28.2.5 SetTimerThread()	215
19.28.2.6 Update()	215
19.28.3 Member Data Documentation	216
19.28.3.1 m_bInTimer	216
19.28.3.2 m_pclActiveThread	216
19.28.3.3 m_pclTimerThread	216
19.28.3.4 m_u16TicksRemain	216
19.29Mark3::ReaderWriterLock Class Reference	216
19.29.1 Detailed Description	217
19.29.2 Member Function Documentation	218
19.29.2.1 AcquireReader() [1/2]	218
19.29.2.2 AcquireReader() [2/2]	218
19.29.2.3 AcquireReader_i()	218
19.29.2.4 AcquireWriter() [1/2]	219
19.29.2.5 AcquireWriter() [2/2]	219
19.29.2.6 AcquireWriter_i()	219
19.29.2.7 Init()	220
19.29.2.8 operator new()	220

19.29.2.9 ReleaseReader()	220
19.29.2.10ReleaseWriter()	220
19.29.3 Member Data Documentation	221
19.29.3.1 m_clGlobalMutex	221
19.29.3.2 m_clReaderMutex	221
19.29.3.3 m_u8ReadCount	221
19.30Mark3::Scheduler Class Reference	221
19.30.1 Detailed Description	222
19.30.2 Member Function Documentation	223
19.30.2.1 Add()	223
19.30.2.2 GetCurrentThread()	223
19.30.2.3 GetNextThread()	223
19.30.2.4 GetStopList()	224
19.30.2.5 GetThreadList()	224
19.30.2.6 Init()	224
19.30.2.7 IsEnabled()	225
19.30.2.8 QueueScheduler()	225
19.30.2.9 Remove()	225
19.30.2.10Schedule()	225
19.30.2.11SetScheduler()	226
19.30.3 Member Data Documentation	226
19.30.3.1 m_aclPriorities	226
19.30.3.2 m_bEnabled	226
19.30.3.3 m_bQueuedSchedule	226
19.30.3.4 m_clPrioMap	227
19.30.3.5 m_clStopList	227
19.30.3.6 m_uNumPriorities	227
19.31Mark3::SchedulerGuard Class Reference	227
19.31.1 Detailed Description	228
19.31.2 Constructor & Destructor Documentation	228

19.31.2.1 SchedulerGuard()	228
19.31.2.2 ~SchedulerGuard()	228
19.31.3 Member Data Documentation	228
19.31.3.1 m_bSchedState	228
19.32Mark3::Semaphore Class Reference	229
19.32.1 Detailed Description	230
19.32.2 Constructor & Destructor Documentation	230
19.32.2.1 ~Semaphore()	230
19.32.3 Member Function Documentation	230
19.32.3.1 GetCount()	230
19.32.3.2 Init()	230
19.32.3.3 operator new()	231
19.32.3.4 Pend() [1/2]	231
19.32.3.5 Pend() [2/2]	231
19.32.3.6 Pend_i()	232
19.32.3.7 Post()	233
19.32.3.8 WakeMe()	233
19.32.3.9 WakeNext()	234
19.32.4 Member Data Documentation	234
19.32.4.1 m_u16MaxValue	234
19.32.4.2 m_u16Value	234
19.33Mark3::Streamer Class Reference	234
19.33.1 Detailed Description	235
19.33.2 Member Function Documentation	236
19.33.2.1 CanRead()	236
19.33.2.2 CanWrite()	236
19.33.2.3 Claim()	236
19.33.2.4 GetAvailable()	237
19.33.2.5 Init()	237
19.33.2.6 IsEmpty()	237

19.33.2.7 Lock()	238
19.33.2.8 Read() [1/2]	238
19.33.2.9 Read() [2/2]	239
19.33.2.10Unlock()	239
19.33.2.11Write() [1/2]	239
19.33.2.12Write() [2/2]	240
19.33.3 Member Data Documentation	240
19.33.3.1 m_pau8Buffer	240
19.33.3.2 m_pu8LockAddr	240
19.33.3.3 m_u16Avail	241
19.33.3.4 m_u16Head	241
19.33.3.5 m_u16Size	241
19.33.3.6 m_u16Tail	241
19.34Mark3::Thread Class Reference	242
19.34.1 Detailed Description	245
19.34.2 Constructor & Destructor Documentation	245
19.34.2.1 ~Thread()	245
19.34.2.2 Thread()	246
19.34.3 Member Function Documentation	246
19.34.3.1 ContextSwitchSWI()	246
19.34.3.2 CoopYield()	246
19.34.3.3 Exit()	246
19.34.3.4 GetCurPriority()	247
19.34.3.5 GetCurrent()	247
19.34.3.6 GetEventFlagMask()	247
19.34.3.7 GetEventFlagMode()	248
19.34.3.8 GetExpired()	248
19.34.3.9 GetExtendedContext()	248
19.34.3.10GetID()	248
19.34.3.11GetName()	249

19.34.3.12	GetOwner()	249
19.34.3.13	GetPriority()	249
19.34.3.14	GetQuantum()	250
19.34.3.15	GetStack()	250
19.34.3.16	GetStackSize()	250
19.34.3.17	GetStackSlack()	251
19.34.3.18	GetState()	251
19.34.3.19	GetTimer()	251
19.34.3.20	InheritPriority()	251
19.34.3.21	Init() [1 / 2]	252
19.34.3.22	Init() [2 / 2]	252
19.34.3.23	IsInitialized()	253
19.34.3.24	operator new()	253
19.34.3.25	SetCurrent()	253
19.34.3.26	SetEventFlagMask()	254
19.34.3.27	SetEventFlagMode()	254
19.34.3.28	SetExpired()	254
19.34.3.29	SetExtendedContext()	255
19.34.3.30	SetID()	255
19.34.3.31	SetName()	255
19.34.3.32	SetOwner()	257
19.34.3.33	SetPriority()	257
19.34.3.34	SetPriorityBase()	257
19.34.3.35	SetQuantum()	258
19.34.3.36	SetState()	258
19.34.3.37	Sleep()	258
19.34.3.38	Start()	259
19.34.3.39	Stop()	259
19.34.3.40	Yield()	259
19.34.4	Friends And Related Function Documentation	260

19.34.4.1 ThreadPort	260
19.34.5 Member Data Documentation	260
19.34.5.1 m_bExpired	260
19.34.5.2 m_clTimer	260
19.34.5.3 m_eFlagMode	260
19.34.5.4 m_eState	261
19.34.5.5 m_pclCurrent	261
19.34.5.6 m_pclOwner	261
19.34.5.7 m_pfEntryPoint	261
19.34.5.8 m_pvArg	261
19.34.5.9 m_pvExtendedContext	262
19.34.5.10m_pwStack	262
19.34.5.11m_pwStackTop	262
19.34.5.12m_szName	262
19.34.5.13m_u16FlagMask	262
19.34.5.14m_u16Quantum	263
19.34.5.15m_u16StackSize	263
19.34.5.16m_u8ThreadID	263
19.34.5.17m_uXCurPriority	263
19.34.5.18m_uXPriority	263
19.35Mark3::ThreadList Class Reference	264
19.35.1 Detailed Description	265
19.35.2 Constructor & Destructor Documentation	265
19.35.2.1 ThreadList()	265
19.35.3 Member Function Documentation	265
19.35.3.1 Add() [1/2]	265
19.35.3.2 Add() [2/2]	265
19.35.3.3 AddPriority()	267
19.35.3.4 HighestWaiter()	267
19.35.3.5 operator new()	267

19.35.3.6 Remove()	268
19.35.3.7 SetMapPointer()	269
19.35.3.8 SetPriority()	269
19.35.4 Member Data Documentation	269
19.35.4.1 m_pciMap	269
19.35.4.2 m_uXPriority	270
19.36Mark3::ThreadListList Class Reference	270
19.36.1 Detailed Description	270
19.36.2 Member Function Documentation	270
19.36.2.1 Add()	270
19.36.2.2 GetHead()	271
19.36.2.3 Remove()	271
19.36.3 Member Data Documentation	271
19.36.3.1 m_clThreadListList	271
19.37Mark3::ThreadPort Class Reference	272
19.37.1 Detailed Description	272
19.37.2 Member Function Documentation	272
19.37.2.1 Init()	272
19.37.2.2 InitStack()	272
19.37.2.3 StartThreads()	273
19.37.3 Friends And Related Function Documentation	273
19.37.3.1 Thread	273
19.38Mark3::Timer Class Reference	273
19.38.1 Detailed Description	275
19.38.2 Constructor & Destructor Documentation	275
19.38.2.1 ~Timer()	275
19.38.2.2 Timer()	275
19.38.3 Member Function Documentation	276
19.38.3.1 Init()	276
19.38.3.2 IsInitialized()	276

19.38.3.3 MSecondsToTicks()	276
19.38.3.4 operator new()	276
19.38.3.5 SecondsToTicks()	277
19.38.3.6 SetInitialized()	277
19.38.3.7 Start() [1/2]	277
19.38.3.8 Start() [2/2]	277
19.38.3.9 Stop()	278
19.38.3.10 USecondsToTicks()	278
19.38.4 Friends And Related Function Documentation	278
19.38.4.1 TimerList	278
19.38.5 Member Data Documentation	278
19.38.5.1 m_pclOwner	278
19.38.5.2 m_pfCallback	279
19.38.5.3 m_pvData	279
19.38.5.4 m_u32Interval	279
19.38.5.5 m_u32TimeLeft	279
19.38.5.6 m_u8Flags	279
19.38.5.7 m_u8Initialized	280
19.38.5.8 m_uTimerInitCookie	280
19.38.5.9 m_uTimerInvalidCookie	280
19.39Mark3::TimerList Class Reference	280
19.39.1 Detailed Description	281
19.39.2 Member Function Documentation	281
19.39.2.1 Add()	281
19.39.2.2 Init()	282
19.39.2.3 Process()	282
19.39.2.4 Remove()	282
19.39.3 Member Data Documentation	282
19.39.3.1 m_bTimerActive	282
19.39.3.2 m_clMutex	283

19.39.3.3 m_u32NextWakeup	283
19.40Mark3::TimerScheduler Class Reference	283
19.40.1 Detailed Description	284
19.40.2 Member Function Documentation	284
19.40.2.1 Add()	284
19.40.2.2 Init()	284
19.40.2.3 Process()	284
19.40.2.4 Remove()	285
19.40.3 Member Data Documentation	285
19.40.3.1 m_clTimerList	285
19.41Mark3::Token_t Struct Reference	285
19.41.1 Detailed Description	286
19.41.2 Member Data Documentation	286
19.41.2.1 pcToken	286
19.41.2.2 u8Len	286
19.42Mark3::TypedCircularLinkedList< T > Class Template Reference	286
19.42.1 Detailed Description	287
19.42.2 Constructor & Destructor Documentation	287
19.42.2.1 TypedCircularLinkedList()	287
19.42.3 Member Function Documentation	287
19.42.3.1 Add()	287
19.42.3.2 GetHead()	288
19.42.3.3 GetTail()	288
19.42.3.4 InsertNodeBefore()	288
19.42.3.5 operator new()	289
19.42.3.6 Remove()	289
19.42.3.7 SetHead()	289
19.42.3.8 SetTail()	290
19.43Mark3::TypedDoubleLinkedList< T > Class Template Reference	290
19.43.1 Detailed Description	291

19.43.2 Constructor & Destructor Documentation	291
19.43.2.1 TypedDoubleLinkedList()	291
19.43.3 Member Function Documentation	291
19.43.3.1 Add()	291
19.43.3.2 GetHead()	292
19.43.3.3 GetTail()	292
19.43.3.4 operator new()	292
19.43.3.5 Remove()	292
19.43.3.6 SetHead()	293
19.43.3.7 SetTail()	293
19.44 Mark3::TypedLinkedListNode< T > Class Template Reference	293
19.44.1 Detailed Description	294
19.44.2 Member Function Documentation	294
19.44.2.1 GetNext()	294
19.44.2.2 GetPrev()	294
20 File Documentation	295
20.1 /home/moslevin/projects/m3-repo/kernel/lib/mark3c/src/mark3c.cpp File Reference	295
20.1.1 Detailed Description	299
20.1.2 Function Documentation	300
20.1.2.1 Alloc_ConditionVariable()	300
20.1.2.2 Alloc_EventFlag()	300
20.1.2.3 Alloc_Mailbox()	300
20.1.2.4 Alloc_Memory()	300
20.1.2.5 Alloc_Message()	301
20.1.2.6 Alloc_MessagePool()	301
20.1.2.7 Alloc_MessageQueue()	301
20.1.2.8 Alloc_Mutex()	302
20.1.2.9 Alloc_Notify()	302
20.1.2.10 Alloc_ReaderWriterLock()	302
20.1.2.11 Alloc_Semaphore()	303

20.1.2.12 Alloc_Thread()	303
20.1.2.13 Alloc_Timer()	303
20.1.2.14 Atomic_Add16()	304
20.1.2.15 Atomic_Add32()	304
20.1.2.16 Atomic_Add8()	305
20.1.2.17 Atomic_Set16()	305
20.1.2.18 Atomic_Set32()	306
20.1.2.19 Atomic_Set8()	306
20.1.2.20 Atomic_Sub16()	307
20.1.2.21 Atomic_Sub32()	307
20.1.2.22 Atomic_Sub8()	308
20.1.2.23 Atomic_TestAndSet()	308
20.1.2.24 ConditionVariable_Broadcast()	309
20.1.2.25 ConditionVariable_Init()	309
20.1.2.26 ConditionVariable_Signal()	309
20.1.2.27 ConditionVariable_TimedWait()	310
20.1.2.28 ConditionVariable_Wait()	310
20.1.2.29 Free_ConditionVariable()	311
20.1.2.30 Free_EventFlag()	311
20.1.2.31 Free_Mailbox()	311
20.1.2.32 Free_Memory()	311
20.1.2.33 Free_Message()	312
20.1.2.34 Free_MessagePool()	312
20.1.2.35 Free_MessageQueue()	312
20.1.2.36 Free_Mutex()	312
20.1.2.37 Free_Notify()	312
20.1.2.38 Free_ReaderWriterLock()	313
20.1.2.39 Free_Semaphore()	313
20.1.2.40 Free_Thread()	313
20.1.2.41 Free_Timer()	313

20.1.2.42 Kernel_DebugPrint()	313
20.1.2.43 Kernel_GetStackGuardThreshold()	314
20.1.2.44 Kernel_GetTicks()	314
20.1.2.45 Kernel_Init()	314
20.1.2.46 Kernel_IsPanic()	315
20.1.2.47 Kernel_IsStarted()	315
20.1.2.48 Kernel_Panic()	315
20.1.2.49 Kernel_SetDebugPrintFunction()	316
20.1.2.50 Kernel_SetPanic()	316
20.1.2.51 Kernel_SetStackGuardThreshold()	317
20.1.2.52 Kernel_Start()	317
20.1.2.53 Mailbox_GetFreeSlots()	317
20.1.2.54 Mailbox_Init()	318
20.1.2.55 Mailbox_IsEmpty()	318
20.1.2.56 Mailbox_IsFull()	319
20.1.2.57 Mailbox_Receive()	319
20.1.2.58 Mailbox_ReceiveTail()	319
20.1.2.59 Mailbox_Send()	320
20.1.2.60 Mailbox_SendTail()	320
20.1.2.61 Mailbox_TimedReceive()	321
20.1.2.62 Mailbox_TimedReceiveTail()	321
20.1.2.63 Mailbox_TimedSend()	322
20.1.2.64 Mailbox_TimedSendTail()	323
20.1.2.65 Message_GetCode()	323
20.1.2.66 Message_GetData()	324
20.1.2.67 Message_Init()	324
20.1.2.68 Message_SetCode()	324
20.1.2.69 Message_SetData()	325
20.1.2.70 MessagePool_Init()	325
20.1.2.71 MessagePool_Pop()	326

20.1.2.72 MessagePool_Push()	326
20.1.2.73 MessageQueue_GetCount()	327
20.1.2.74 MessageQueue_Init()	327
20.1.2.75 MessageQueue_Receive()	327
20.1.2.76 MessageQueue_Send()	328
20.1.2.77 MessageQueue_TimedReceive()	328
20.1.2.78 Mutex_Claim()	329
20.1.2.79 Mutex_Init()	329
20.1.2.80 Mutex_Release()	330
20.1.2.81 Mutex_TimedClaim()	330
20.1.2.82 Notify_Init()	331
20.1.2.83 Notify_Signal()	331
20.1.2.84 Notify_TimedWait()	331
20.1.2.85 Notify_Wait()	332
20.1.2.86 ReaderWriterLock_AcquireReader()	332
20.1.2.87 ReaderWriterLock_AcquireWriter()	333
20.1.2.88 ReaderWriterLock_Init()	333
20.1.2.89 ReaderWriterLock_ReleaseReader()	334
20.1.2.90 ReaderWriterLock_ReleaseWriter()	334
20.1.2.91 ReaderWriterLock_TimedAcquireReader()	334
20.1.2.92 ReaderWriterLock_TimedAcquireWriter()	335
20.1.2.93 Scheduler_Enable()	335
20.1.2.94 Scheduler_GetCurrentThread()	336
20.1.2.95 Scheduler_IsEnabled()	336
20.1.2.96 Semaphore_Init()	336
20.1.2.97 Semaphore_Pend()	337
20.1.2.98 Semaphore_Post()	337
20.1.2.99 Semaphore_TimedPend()	338
20.1.2.100 Thread_CoopYield()	338
20.1.2.101 Thread_Exit()	338

20.1.2.102Thread_GetCurPriority()	339
20.1.2.103Thread_GetID()	339
20.1.2.104Thread_GetName()	340
20.1.2.105Thread_GetPriority()	340
20.1.2.106Thread_GetQuantum()	340
20.1.2.107Thread_GetStackSlack()	341
20.1.2.108Thread_GetState()	341
20.1.2.109Thread_Init()	341
20.1.2.110Thread_SetID()	342
20.1.2.111Thread_SetName()	342
20.1.2.112Thread_SetPriority()	343
20.1.2.113Thread_SetQuantum()	343
20.1.2.114Thread_Sleep()	343
20.1.2.115Thread_Start()	344
20.1.2.116Thread_Stop()	344
20.1.2.117Thread_Yield()	344
20.1.2.118Timer_Init()	345
20.1.2.119Timer_Restart()	345
20.1.2.120Timer_Start()	345
20.1.2.121Timer_Stop()	346
20.2 mark3c.cpp	346
20.3 /home/moslevin/projects/m3-repo/kernel/lib/mark3c/src/public/fake_types.h File Reference	358
20.3.1 Detailed Description	358
20.4 fake_types.h	358
20.5 /home/moslevin/projects/m3-repo/kernel/lib/mark3c/src/public/mark3c.h File Reference	360
20.5.1 Detailed Description	366
20.5.2 Macro Definition Documentation	366
20.5.2.1 CONDITIONVARIABLE_SIZE	366
20.5.2.2 DECLARE_CONDITIONVARIABLE	366
20.5.2.3 DECLARE_MAILBOX	366

20.5.2.4	DECLARE_MESSAGE	367
20.5.2.5	DECLARE_MESSAGEPOOL	367
20.5.2.6	DECLARE_MESSAGEQUEUE	367
20.5.2.7	DECLARE_MUTEX	368
20.5.2.8	DECLARE_NOTIFY	368
20.5.2.9	DECLARE_READERWRITERLOCK	368
20.5.2.10	DECLARE_SEMAPHORE	369
20.5.2.11	DECLARE_THREAD	369
20.5.2.12	DECLARE_TIMER	369
20.5.2.13	MAILBOX_SIZE	369
20.5.2.14	MESSAGE_SIZE	370
20.5.2.15	MESSAGEPOOL_SIZE	370
20.5.2.16	MESSAGEQUEUE_SIZE	370
20.5.2.17	MUTEX_SIZE	370
20.5.2.18	NOTIFY_SIZE	370
20.5.2.19	READERWRITERLOCK_SIZE	370
20.5.2.20	SEMAPHORE_SIZE	371
20.5.2.21	THREAD_SIZE	371
20.5.2.22	TIMER_SIZE	371
20.5.2.23	TOKEN_1	371
20.5.2.24	TOKEN_2	371
20.5.2.25	WORD_ROUND	371
20.5.3	Typedef Documentation	372
20.5.3.1	ConditionVariable_t	372
20.5.3.2	kernel_debug_print_t	372
20.5.3.3	Mailbox_t	372
20.5.3.4	Message_t	372
20.5.3.5	MessagePool_t	372
20.5.3.6	MessageQueue_t	373
20.5.3.7	Mutex_t	373

20.5.3.8	Notify_t	373
20.5.3.9	panic_func_t	373
20.5.3.10	ReaderWriterLock_t	373
20.5.3.11	Semaphore_t	374
20.5.3.12	thread_entry_func_t	374
20.5.3.13	Thread_t	374
20.5.3.14	timer_callback_t	374
20.5.3.15	Timer_t	374
20.5.4	Enumeration Type Documentation	374
20.5.4.1	thread_state_t	374
20.5.5	Function Documentation	375
20.5.5.1	Alloc_Mailbox()	375
20.5.5.2	Alloc_Memory()	375
20.5.5.3	Alloc_Message()	376
20.5.5.4	Alloc_MessagePool()	376
20.5.5.5	Alloc_MessageQueue()	376
20.5.5.6	Alloc_Mutex()	377
20.5.5.7	Alloc_Notify()	377
20.5.5.8	Alloc_Semaphore()	377
20.5.5.9	Alloc_Thread()	378
20.5.5.10	Alloc_Timer()	378
20.5.5.11	Atomic_Add16()	378
20.5.5.12	Atomic_Add32()	379
20.5.5.13	Atomic_Add8()	379
20.5.5.14	Atomic_Set16()	380
20.5.5.15	Atomic_Set32()	380
20.5.5.16	Atomic_Set8()	381
20.5.5.17	Atomic_Sub16()	381
20.5.5.18	Atomic_Sub32()	382
20.5.5.19	Atomic_Sub8()	382

20.5.5.20 Atomic_TestAndSet()	383
20.5.5.21 ConditionVariable_Broadcast()	383
20.5.5.22 ConditionVariable_Init()	384
20.5.5.23 ConditionVariable_Signal()	384
20.5.5.24 ConditionVariable_TimedWait()	385
20.5.5.25 ConditionVariable_Wait()	385
20.5.5.26 Free_Mailbox()	386
20.5.5.27 Free_Memory()	386
20.5.5.28 Free_Message()	386
20.5.5.29 Free_MessagePool()	386
20.5.5.30 Free_MessageQueue()	386
20.5.5.31 Free_Mutex()	387
20.5.5.32 Free_Notify()	387
20.5.5.33 Free_Semaphore()	387
20.5.5.34 Free_Thread()	387
20.5.5.35 Free_Timer()	387
20.5.5.36 Kernel_DebugPrint()	387
20.5.5.37 Kernel_GetTicks()	388
20.5.5.38 Kernel_Init()	388
20.5.5.39 Kernel_IsPanic()	389
20.5.5.40 Kernel_IsStarted()	389
20.5.5.41 Kernel_Panic()	389
20.5.5.42 Kernel_SetDebugPrintFunction()	390
20.5.5.43 Kernel_SetPanic()	390
20.5.5.44 Kernel_Start()	391
20.5.5.45 Mailbox_GetFreeSlots()	391
20.5.5.46 Mailbox_Init()	391
20.5.5.47 Mailbox_IsEmpty()	392
20.5.5.48 Mailbox_IsFull()	392
20.5.5.49 Mailbox_Receive()	393

20.5.5.50 Mailbox_ReceiveTail()	393
20.5.5.51 Mailbox_Send()	394
20.5.5.52 Mailbox_SendTail()	394
20.5.5.53 Mailbox_TimedReceive()	395
20.5.5.54 Mailbox_TimedReceiveTail()	395
20.5.5.55 Mailbox_TimedSend()	396
20.5.5.56 Mailbox_TimedSendTail()	396
20.5.5.57 Message_GetCode()	397
20.5.5.58 Message_GetData()	397
20.5.5.59 Message_Init()	398
20.5.5.60 Message_SetCode()	398
20.5.5.61 Message_SetData()	399
20.5.5.62 MessagePool_Init()	399
20.5.5.63 MessagePool_Pop()	399
20.5.5.64 MessagePool_Push()	400
20.5.5.65 MessageQueue_GetCount()	400
20.5.5.66 MessageQueue_Init()	401
20.5.5.67 MessageQueue_Receive()	401
20.5.5.68 MessageQueue_Send()	402
20.5.5.69 MessageQueue_TimedReceive()	402
20.5.5.70 Mutex_Claim()	403
20.5.5.71 Mutex_Init()	403
20.5.5.72 Mutex_Release()	403
20.5.5.73 Mutex_TimedClaim()	404
20.5.5.74 Notify_Init()	404
20.5.5.75 Notify_Signal()	405
20.5.5.76 Notify_TimedWait()	405
20.5.5.77 Notify_Wait()	406
20.5.5.78 ReaderWriterLock_AcquireReader()	406
20.5.5.79 ReaderWriterLock_AcquireWriter()	406

20.5.5.80 ReaderWriterLock_Init()	407
20.5.5.81 ReaderWriterLock_ReleaseReader()	407
20.5.5.82 ReaderWriterLock_ReleaseWriter()	408
20.5.5.83 ReaderWriterLock_TimedAcquireReader()	408
20.5.5.84 ReaderWriterLock_TimedAcquireWriter()	408
20.5.5.85 Scheduler_Enable()	409
20.5.5.86 Scheduler_GetCurrentThread()	409
20.5.5.87 Scheduler_IsEnabled()	410
20.5.5.88 Semaphore_Init()	410
20.5.5.89 Semaphore_Pend()	411
20.5.5.90 Semaphore_Post()	411
20.5.5.91 Semaphore_TimedPend()	411
20.5.5.92 Thread_CoopYield()	412
20.5.5.93 Thread_Exit()	412
20.5.5.94 Thread_GetCurPriority()	413
20.5.5.95 Thread_GetID()	413
20.5.5.96 Thread_GetPriority()	413
20.5.5.97 Thread_GetState()	414
20.5.5.98 Thread_Init()	414
20.5.5.99 Thread_SetID()	415
20.5.5.100 Thread_SetPriority()	415
20.5.5.101 Thread_Sleep()	416
20.5.5.102 Thread_Start()	416
20.5.5.103 Thread_Stop()	417
20.5.5.104 Thread_Yield()	417
20.5.5.105 Timer_Init()	417
20.5.5.106 Timer_Restart()	418
20.5.5.107 Timer_Start()	418
20.5.5.108 Timer_Stop()	419
20.6 mark3c.h	419

20.7 /home/moslevin/projects/m3-repo/kernel/lib/memutil/memutil.cpp File Reference	424
20.7.1 Detailed Description	425
20.8 memutil.cpp	425
20.9 /home/moslevin/projects/m3-repo/kernel/lib/memutil/public/memutil.h File Reference	431
20.9.1 Detailed Description	432
20.10 memutil.h	432
20.11 /home/moslevin/projects/m3-repo/kernel/lib/streamer/public/streamer.h File Reference	433
20.11.1 Detailed Description	433
20.12 streamer.h	434
20.13 /home/moslevin/projects/m3-repo/kernel/lib/streamer/streamer.cpp File Reference	434
20.13.1 Detailed Description	435
20.14 streamer.cpp	435
20.15 /home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/kernelswi.cpp File Reference	438
20.15.1 Detailed Description	438
20.16 kernelswi.cpp	438
20.17 /home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/kerneltimer.cpp File Refer-	
ence	439
20.17.1 Detailed Description	439
20.17.2 Macro Definition Documentation	440
20.17.2.1 TCCR1B_INIT	440
20.17.2.2 TIMER_IFR	440
20.17.2.3 TIMER_IMSK	440
20.17.3 Function Documentation	440
20.17.3.1 ISR()	440
20.18 kerneltimer.cpp	441
20.19 /home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/portcfg.h File Refer-	
ence	442
20.19.1 Detailed Description	442
20.19.2 Macro Definition Documentation	443
20.19.2.1 AVR	443
20.19.2.2 K_ADDR	443

20.19.2.3 K_INT	443
20.19.2.4 K_WORD	443
20.19.2.5 KERNEL_NUM_PRIORITIES	444
20.19.2.6 KERNEL_STACK_GUARD_DEFAULT	444
20.19.2.7 KERNEL_TIMERS_THREAD_PRIORITY	444
20.19.2.8 PORT_COROUTINE_PRIORITIES	444
20.19.2.9 PORT_KERNEL_DEFAULT_STACK_SIZE	444
20.19.2.10PORT_KERNEL_TIMERS_THREAD_STACK	445
20.19.2.11PORT_MIN_TIMER_TICKS	445
20.19.2.12PORT_OVERLOAD_NEW	445
20.19.2.13PORT_PRIO_MAP_WORD_SIZE	445
20.19.2.14PORT_PRIO_TYPE	445
20.19.2.15PORT_STACK_GROWS_DOWN	446
20.19.2.16PORT_SYSTEM_FREQ	446
20.19.2.17PORT_TIMER_COUNT_TYPE	446
20.19.2.18PORT_TIMER_FREQ	446
20.19.2.19PORT_USE_HW_CLZ	447
20.19.2.20THREAD_QUANTUM_DEFAULT	447
20.20portcfg.h	447
20.21/home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/threadport.h	File
Reference	448
20.21.1 Detailed Description	448
20.21.2 Macro Definition Documentation	449
20.21.2.1 ASM	449
20.21.2.2 PORT_PUSH_TO_STACK	449
20.21.2.3 PORT_TOP_OF_STACK	449
20.21.2.4 Thread_RestoreContext	449
20.21.2.5 Thread_SaveContext	450
20.22threadport.h	450
20.23/home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/threadport.cpp	File
Reference	452

20.23.1 Detailed Description	453
20.23.2 Variable Documentation	453
20.23.2.1 g_kwCriticalCount	453
20.23.2.2 g_kwSFR	453
20.24threadport.cpp	453
20.25/home/moslevin/projects/m3-repo/kernel/src/atomic.cpp File Reference	455
20.25.1 Detailed Description	455
20.26atomic.cpp	455
20.27/home/moslevin/projects/m3-repo/kernel/src/autoalloc.cpp File Reference	456
20.27.1 Detailed Description	456
20.28autoalloc.cpp	456
20.29/home/moslevin/projects/m3-repo/kernel/src/blocking.cpp File Reference	457
20.29.1 Detailed Description	458
20.30blocking.cpp	458
20.31/home/moslevin/projects/m3-repo/kernel/src/colist.cpp File Reference	459
20.31.1 Detailed Description	459
20.32colist.cpp	459
20.33/home/moslevin/projects/m3-repo/kernel/src/condvar.cpp File Reference	460
20.33.1 Detailed Description	460
20.34condvar.cpp	460
20.35/home/moslevin/projects/m3-repo/kernel/src/coroutine.cpp File Reference	461
20.35.1 Detailed Description	461
20.36coroutine.cpp	462
20.37/home/moslevin/projects/m3-repo/kernel/src/cosched.cpp File Reference	463
20.37.1 Detailed Description	463
20.38cosched.cpp	463
20.39/home/moslevin/projects/m3-repo/kernel/src/eventflag.cpp File Reference	464
20.39.1 Detailed Description	464
20.40eventflag.cpp	464
20.41/home/moslevin/projects/m3-repo/kernel/src/kernel.cpp File Reference	468

20.41.1 Detailed Description	468
20.42kernel.cpp	468
20.43/home/moslevin/projects/m3-repo/kernel/src/ksemaphore.cpp File Reference	469
20.43.1 Detailed Description	470
20.44ksemaphore.cpp	470
20.45/home/moslevin/projects/m3-repo/kernel/src/ll.cpp File Reference	472
20.45.1 Detailed Description	472
20.46ll.cpp	473
20.47/home/moslevin/projects/m3-repo/kernel/src/lockguard.cpp File Reference	474
20.47.1 Detailed Description	475
20.48lockguard.cpp	475
20.49/home/moslevin/projects/m3-repo/kernel/src/mailbox.cpp File Reference	475
20.49.1 Detailed Description	476
20.50mailbox.cpp	476
20.51/home/moslevin/projects/m3-repo/kernel/src/message.cpp File Reference	478
20.51.1 Detailed Description	479
20.52message.cpp	479
20.53/home/moslevin/projects/m3-repo/kernel/src/mutex.cpp File Reference	480
20.53.1 Detailed Description	480
20.54mutex.cpp	481
20.55/home/moslevin/projects/m3-repo/kernel/src/notify.cpp File Reference	483
20.55.1 Detailed Description	484
20.56notify.cpp	484
20.57/home/moslevin/projects/m3-repo/kernel/src/profile.cpp File Reference	486
20.57.1 Detailed Description	486
20.58profile.cpp	486
20.59/home/moslevin/projects/m3-repo/kernel/src/public/atomic.h File Reference	487
20.59.1 Detailed Description	488
20.60atomic.h	488
20.61/home/moslevin/projects/m3-repo/kernel/src/public/autoalloc.h File Reference	489

20.61.1 Detailed Description	490
20.62autoalloc.h	490
20.63/home/moslevin/projects/m3-repo/kernel/src/public/blocking.h File Reference	491
20.63.1 Detailed Description	492
20.64blocking.h	492
20.65/home/moslevin/projects/m3-repo/kernel/src/public/colist.h File Reference	493
20.65.1 Detailed Description	493
20.66colist.h	493
20.67/home/moslevin/projects/m3-repo/kernel/src/public/condvar.h File Reference	494
20.67.1 Detailed Description	494
20.68condvar.h	495
20.69/home/moslevin/projects/m3-repo/kernel/src/public/coroutine.h File Reference	495
20.69.1 Detailed Description	496
20.70coroutine.h	496
20.71/home/moslevin/projects/m3-repo/kernel/src/public/cosched.h File Reference	497
20.71.1 Detailed Description	497
20.72cosched.h	497
20.73/home/moslevin/projects/m3-repo/kernel/src/public/criticalguard.h File Reference	498
20.73.1 Detailed Description	498
20.74criticalguard.h	498
20.75/home/moslevin/projects/m3-repo/kernel/src/public/criticalsection.h File Reference	499
20.75.1 Detailed Description	499
20.76criticalsection.h	499
20.77/home/moslevin/projects/m3-repo/kernel/src/public/eventflag.h File Reference	500
20.77.1 Detailed Description	500
20.78eventflag.h	500
20.79/home/moslevin/projects/m3-repo/kernel/src/public/ithreadport.h File Reference	501
20.79.1 Detailed Description	501
20.80ithreadport.h	502
20.81/home/moslevin/projects/m3-repo/kernel/src/public/kernel.h File Reference	502

20.81.1 Detailed Description	503
20.81.2 Typedef Documentation	503
20.81.2.1 DebugPrintFunction	503
20.82kernel.h	503
20.83/home/moslevin/projects/m3-repo/kernel/src/public/kerneldebug.h File Reference	504
20.83.1 Detailed Description	505
20.83.2 Macro Definition Documentation	505
20.83.2.1 KERNEL_ASSERT	505
20.84kerneldebug.h	505
20.85/home/moslevin/projects/m3-repo/kernel/src/public/kernelswi.h File Reference	506
20.85.1 Detailed Description	506
20.86kernelswi.h	506
20.87/home/moslevin/projects/m3-repo/kernel/src/public/kerneltimer.h File Reference	507
20.87.1 Detailed Description	507
20.88kerneltimer.h	507
20.89/home/moslevin/projects/m3-repo/kernel/src/public/kerneltypes.h File Reference	508
20.89.1 Detailed Description	508
20.90kerneltypes.h	508
20.91/home/moslevin/projects/m3-repo/kernel/src/public/ksemaphore.h File Reference	509
20.91.1 Detailed Description	509
20.92ksemaphore.h	510
20.93/home/moslevin/projects/m3-repo/kernel/src/public/ll.h File Reference	510
20.93.1 Detailed Description	511
20.94ll.h	511
20.95/home/moslevin/projects/m3-repo/kernel/src/public/lockguard.h File Reference	513
20.95.1 Detailed Description	514
20.96lockguard.h	514
20.97/home/moslevin/projects/m3-repo/kernel/src/public/mailbox.h File Reference	514
20.97.1 Detailed Description	515
20.98mailbox.h	515

20.99/home/moslevin/projects/m3-repo/kernel/src/public/manual.h File Reference	516
20.99.1 Detailed Description	517
20.100manual.h	517
20.101/home/moslevin/projects/m3-repo/kernel/src/public/mark3.h File Reference	517
20.101.1 Detailed Description	518
20.102mark3.h	518
20.103/home/moslevin/projects/m3-repo/kernel/src/public/mark3cfg.h File Reference	518
20.103.1 Detailed Description	519
20.103.2 Macro Definition Documentation	519
20.103.2.1 KERNEL_CONTEXT_SWITCH_CALLOUT	519
20.103.2.2 KERNEL_DEBUG	519
20.103.2.3 KERNEL_EVENT_FLAGS	520
20.103.2.4 KERNEL_EXTENDED_CONTEXT	520
20.103.2.5 KERNEL_NAMED_THREADS	520
20.103.2.6 KERNEL_ROUND_ROBIN	520
20.103.2.7 KERNEL_STACK_CHECK	521
20.103.2.8 KERNEL_THREAD_CREATE_CALLOUT	521
20.103.2.9 KERNEL_THREAD_EXIT_CALLOUT	521
20.104mark3cfg.h	521
20.105/home/moslevin/projects/m3-repo/kernel/src/public/message.h File Reference	522
20.105.1 Detailed Description	522
20.105.2 Using Messages, Queues, and the Global Message Pool	523
20.106message.h	523
20.107/home/moslevin/projects/m3-repo/kernel/src/public/mutex.h File Reference	524
20.107.1 Detailed Description	525
20.107.2 Initializing	525
20.107.3 Resource protection example	525
20.108mutex.h	525
20.109/home/moslevin/projects/m3-repo/kernel/src/public/notify.h File Reference	526
20.109.1 Detailed Description	526

20.110	otify.h	527
20.111	home/moslevin/projects/m3-repo/kernel/src/public/paniccodes.h File Reference	527
20.111.1	Detailed Description	528
20.111.2	Macro Definition Documentation	528
20.111.2.1	PANIC_ACTIVE_COROUTINE_DESCOPED	528
20.111.2.2	PANIC_ACTIVE_EVENTFLAG_DESCOPED	528
20.111.2.3	PANIC_ACTIVE_MAILBOX_DESCOPED	528
20.111.2.4	PANIC_ACTIVE_MUTEX_DESCOPED	528
20.111.2.5	PANIC_ACTIVE_NOTIFY_DESCOPED	528
20.111.2.6	PANIC_ACTIVE_SEMAPHORE_DESCOPED	529
20.111.2.7	PANIC_ACTIVE_TIMER_DESCOPED	529
20.111.2.8	PANIC_ASSERT_FAILED	529
20.111.2.9	PANIC_AUTO_HEAP_EXHAUSTED	529
20.111.2.10	PANIC_LIST_UNLINK_FAILED	529
20.111.2.11	PANIC_NO_READY_THREADS	529
20.111.2.12	PANIC_POWERMAN_EXHAUSTED	530
20.111.2.13	PANIC_RUNNING_THREAD_DESCOPED	530
20.111.2.14	PANIC_STACK_SLACK_VIOLATED	530
20.112	paniccodes.h	530
20.113	home/moslevin/projects/m3-repo/kernel/src/public/priomap.h File Reference	530
20.113.1	Detailed Description	531
20.114	priomap.h	531
20.115	home/moslevin/projects/m3-repo/kernel/src/public/priomapl1.h File Reference	531
20.115.1	Detailed Description	532
20.116	priomapl1.h	532
20.117	home/moslevin/projects/m3-repo/kernel/src/public/priomapl2.h File Reference	533
20.117.1	Detailed Description	533
20.118	priomapl2.h	534
20.119	home/moslevin/projects/m3-repo/kernel/src/public/profile.h File Reference	535
20.119.1	Detailed Description	535

20.12	profile.h	536
20.12	home/moslevin/projects/m3-repo/kernel/src/public/profiling_results.h File Reference	536
20.12	profiling_results.h	536
20.12	home/moslevin/projects/m3-repo/kernel/src/public/quantum.h File Reference	537
20.123	Detailed Description	537
20.12	quantum.h	537
20.12	home/moslevin/projects/m3-repo/kernel/src/public/readerwriter.h File Reference	538
20.125	Detailed Description	538
20.12	readerwriter.h	539
20.12	home/moslevin/projects/m3-repo/kernel/src/public/scheduler.h File Reference	539
20.127	Detailed Description	540
20.127.2	Variable Documentation	540
20.127.2.1	g_pclCurrent	540
20.127.2.2	g_pclNext	540
20.12	scheduler.h	541
20.12	home/moslevin/projects/m3-repo/kernel/src/public/schedulerguard.h File Reference	541
20.129	Detailed Description	542
20.13	schedulerguard.h	542
20.13	home/moslevin/projects/m3-repo/kernel/src/public/sizeprofile.h File Reference	542
20.13	sizeprofile.h	542
20.13	home/moslevin/projects/m3-repo/kernel/src/public/thread.h File Reference	543
20.133	Detailed Description	543
20.13	thread.h	544
20.13	home/moslevin/projects/m3-repo/kernel/src/public/threadlist.h File Reference	546
20.135	Detailed Description	546
20.13	threadlist.h	547
20.13	home/moslevin/projects/m3-repo/kernel/src/public/threadlistlist.h File Reference	547
20.137	Detailed Description	548
20.13	threadlistlist.h	548
20.13	home/moslevin/projects/m3-repo/kernel/src/public/timer.h File Reference	548

20.139.1 Detailed Description	549
20.140 timer.h	549
20.141 home/moslevin/projects/m3-repo/kernel/src/public/timerlist.h File Reference	550
20.141.1 Detailed Description	551
20.142 timerlist.h	551
20.143 home/moslevin/projects/m3-repo/kernel/src/public/timerscheduler.h File Reference	551
20.143.1 Detailed Description	552
20.144 timerscheduler.h	552
20.145 home/moslevin/projects/m3-repo/kernel/src/quantum.cpp File Reference	552
20.145.1 Detailed Description	553
20.146 quantum.cpp	553
20.147 home/moslevin/projects/m3-repo/kernel/src/readerwriter.cpp File Reference	554
20.147.1 Detailed Description	554
20.148 readerwriter.cpp	554
20.149 home/moslevin/projects/m3-repo/kernel/src/scheduler.cpp File Reference	555
20.149.1 Detailed Description	556
20.149.2 variable Documentation	556
20.149.2.1 g_pclCurrent	556
20.149.2.2 g_pclNext	556
20.150 scheduler.cpp	556
20.151 home/moslevin/projects/m3-repo/kernel/src/thread.cpp File Reference	557
20.151.1 Detailed Description	557
20.152 thread.cpp	558
20.153 home/moslevin/projects/m3-repo/kernel/src/threadlist.cpp File Reference	563
20.153.1 Detailed Description	563
20.154 threadlist.cpp	563
20.155 home/moslevin/projects/m3-repo/kernel/src/threadlistlist.cpp File Reference	565
20.155.1 Detailed Description	565
20.156 threadlistlist.cpp	565
20.157 home/moslevin/projects/m3-repo/kernel/src/timer.cpp File Reference	565
20.157.1 Detailed Description	566
20.158 timer.cpp	566
20.159 home/moslevin/projects/m3-repo/kernel/src/timerlist.cpp File Reference	567
20.159.1 Detailed Description	567
20.160 timerlist.cpp	567

21 Example Documentation	569
21.1 lab10_notifications/main.cpp	569
21.2 lab11_mailboxes/main.cpp	570
21.3 lab1_kernel_setup/main.cpp	572
21.4 lab2_coroutines/main.cpp	574
21.5 lab3_round_robin/main.cpp	576
21.6 lab4_semaphores/main.cpp	578
21.7 lab5_mutexes/main.cpp	579
21.8 lab6_timers/main.cpp	581
21.9 lab7_events/main.cpp	583
21.10lab8_messages/main.cpp	585
21.11lab9_dynamic_threads/main.cpp	587

Chapter 1

The Mark3 Realtime Kernel



The [Mark3](#) Realtime Kernel is a completely free, open-source, real-time operating system aimed at bringing powerful, easy-to-use multitasking to microcontroller systems without MMUs.



The RTOS is written using a super portable design that scales to many common processor architectures, including a variety of 8, 16, 32, and 64-bit targets. The flexible CMake-based build system facilitates compiling the kernel, tests, examples, and user-application code for any supported target with a consistent interface.

The API is rich and surprisingly simple – with six function calls, you can set up the kernel, initialize two threads, and start the scheduler.

Written in modern C++, [Mark3](#) makes use of modern language features that improve code quality, reduces duplication, and simplifies API usage. C++ isn't your thing? No problem! there's a complete set of C-language bindings available to facilitate the use of [Mark3](#) in a wider variety of environments.

The [Mark3](#) kernel releases contain zero compiler warnings, zero compiler errors, and have zero unit test failures. The build and test process can be automated through the Mark3-docker project, allowing for easy integration with continuous integration environments. The kernel is also run through static analysis tools, automated profiling and documentation tools.

The source is fully-documented with doxygen, and example code is provided to illustrate core concepts. The result is a performant RTOS, which is easy to read, easy to understand, and easy to extend to fit your needs.

Chapter 2

License

2.1 License

Copyright (c) 2012 - 2019, m0slevin
All rights reserved.

Redistribution and use in source and binary forms, with or without modification, are permitted provided that the following conditions are met:

- Redistributions of source code must retain the above copyright notice, this list of conditions and the following disclaimer.
- Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution.
- Neither the names m0slevin, Mark Slevinsky, nor the names of its contributors may be used to endorse or promote products derived from this software without specific prior written permission.

THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL FUNKENSTEIN SOFTWARE AND/OR ITS CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

Chapter 3

Configuring The Kernel

3.1 Overview

Configuration is done through setting options in [mark3cfg.h](#) , and [portcfg.h](#).

[mark3cfg.h](#) contains global kernel configuration options, which determine specific kernel behaviors, and enable certain kernel APIs independent of the any target architecture. Previous to the R7 release, all kernel configuration options were set from [mark3cfg.h](#), and there was an incredible amount of granularity in the configuration options.

One main motivating factor behind removing the granular configuration in [mark3cfg.h](#) is that it added a ton of `#ifdefs` throughout the code, which made it look less clean. It was also difficult to maintain since there were too many permutations and combinations of configuration options to reasonably test.

Another motivation for removing the vast array of configuration options from [mark3cfg.h](#) is that there's limited benefit to code size. With the advent of modern compiler optimizations such as section-based garbage collection and link-time optimizations, compilers to a remarkable job of optimizing out unused code. [Mark3](#) supports these optimizations, allowing for nearly the same level of performance benefit as feature specific `#ifdefs`. In short - you still only pay for what you use, without having to select the features you want ahead of time.

3.1.1 Kernel Configuration Options

Kernel configuration is performed by setting `#define`'s in [mark3cfg.h](#) to values of 0 or 1.

KERNEL_DEBUG

When enabled, assert statements within the kernel are checked at runtime. This is useful for tracking kernel-breaking changes, memory corruption, etc. in debug builds.

Should be disabled in release builds for performance reasons.

KERNEL_STACK_CHECK

Perform stack-depth checks on threads at each context switch, which is useful in detecting stack overflows / near overflows. Near-overflow detection uses thresholds defined in the target's [portcfg.h](#). Enabling this also adds the `Thread::GetStackSlack()` method, which allows a thread's stack to be profiled on-demand.

Note: When enabled, the additional stack checks result in a performance hit to context switches and thread initialization.

KERNEL_NAMED_THREADS

Enabling this provides the `Thread::SetName()` and `Thread::GetName()` methods, allowing for each thread to be named with a null-terminated `const char*` string.

Note: the string passed to `Thread::SetName()` must persist for the lifetime of the thread

KERNEL_EVENT_FLAGS

This flag enables the event-flags synchronization object. This feature allows threads to be blocked, waiting on specific condition bits to be set or cleared on an `EventFlag` object.

While other synchronization objects are enabled by default, this one is configurable because it impacts the `Thread` object's member data.

KERNEL_CONTEXT_SWITCH_CALLOUT

When enabled, this feature allows a user to define a callback to be executed whenever a context switch occurs. Enabling this provides a means for a user to track thread statistics, but it does result in additional overhead during a context switch.

KERNEL_THREAD_CREATE_CALLOUT

This feature provides a user-defined kernel callback that is executed whenever a thread is started.

KERNEL_THREAD_EXIT_CALLOUT

This feature provides a user-defined kernel callback that is executed whenever a thread is terminated.

KERNEL_ROUND_ROBIN

Enable round-robin scheduling within each priority level. When selected, this results in a small performance hit during context switching and in the system tick handler, as a special software timer is used to manage the running thread's quantum. Can be disabled to optimize performance if not required.

KERNEL_EXTENDED_CONTEXT

Provide a special data pointer in the thread object, which may be used to add additional context to a thread. Typically this would be used to implement thread-local-storage.

3.1.2 Port Configuration Options

The bulk of kernel configuration options reside in the target's `portcfg.h` file. These options determine various sizes, priorities, and default values related to registers, timers, and threads. Some ports may define their own configuration options used locally by its `kerneltimer/kernelswi/threadport` modules; these are not shown here. The common configuration options are described below.

KERNEL_NUM_PRIORITIES

Define the number of thread priorities that the kernel's scheduler will support. The number of thread priorities is limited only by the memory of the host CPU, as a `ThreadList` object is statically-allocated for each thread priority.

In practice, systems rarely need more than 32 priority levels, with the most complex having the capacity for 256.

KERNEL_TIMERS_THREAD_PRIORITY

Define the priority at which the kernel timer thread runs. Typically, this needs to be one of the highest

Note: Other threads at or above the timer thread's priority are not permitted to use the kernel's Timer API, as the thread relies on strict priority scheduling to manage the global timer list without requiring additional/excessive critical sections.

THREAD_QUANTUM_DEFAULT

Number of milliseconds to set as the default epoch for round-robin scheduling when multiple ready threads are active within the same priority.

KERNEL_STACK_GUARD_DEFAULT

Set the minimum number of words of margin that each thread's stack must maintain. If a thread's stack grows into this margin, a kernel assert will be generated in debug builds. This is useful for ensuring that threads are not in danger of overflowing their stacks during development/verification.

K_WORD

Define the size of a data word (in bytes) on the target

K_ADDR

Define the size of an address (in bytes) on the target. This typically only differs from K_WORD on Harvard-architecture CPUs.

K_INT

Define a type to be used as an integer by the kernel.

PORT_PRIO_TYPE

Set a base datatype used to represent each element of the scheduler's priority bitmap.

PORT_PRIO_MAP_WORD_SIZE

Size of PORT_PRIO_TYPE in bytes

PORT_SYSTEM_FREQ

Define the running CPU frequency. This may be an integer constant, or an alias for another variable which holds the CPU's current running frequency.

PORT_TIMER_FREQ

Set the timer tick frequency. This is the frequency at which the fixed-frequency kernel tick interrupt occurs.

PORT_KERNEL_DEFAULT_STACK_SIZE

Define the default thread stack size (in bytes)

PORT_KERNEL_TIMERS_THREAD_STACK

Define the Timer thread's stack size (in bytes)

PORT_TIMER_COUNT_TYPE

Define the native type corresponding to the target timer hardware's counter register.

PORT_MIN_TIMER_TICKS

Minimum number of timer ticks for any delay or sleep, required to ensure that a timer cannot be initialized to a negative value.

PORT_OVERLOAD_NEW

Set this to 1 to overload the system's New/Free functions with the kernel's allocator functions. A user must configure the Kernel's allocator functions to point to a real heap implementation backed with real memory in order to use dynamic object creation.

PORT_STACK_GROWS_DOWN

Set this to 1 if the stack grows down in the target architecture, or 0 if the stack grows up

PORT_USE_HW_CLZ

Set this to 1 if the target CPU/toolchain supports an optimized Count-leading-zeros instruction, or count-leading-zeros intrinsic. If such functionality is not available, a general-purpose implementation will be used.

See also

[portcfg.h](#)
[mark3cfg.h](#)

Chapter 4

Building The Kernel

[Mark3](#) contains its own build system and support scripts, based on CMake and Ninja.

4.1 Prerequisites

To build via CMake, a user requires a suitable, supported toolchain (i.e. gcc-avr, arm-none-eabi-gcc), CMake 3.4.2 or higher, and a backend supported by CMake (i.e. Ninja build).

For example, on debian-based distributions, such as Ubuntu, the avr toolchain can be installed using:

```
apt-get install avr-libc gcc-avr cmake ninja-build
```

Once a sane build environment has been created, the kernel, libraries, examples and tests can be built by running `ninja` from the target's build directory (`/out/<target>/`). By default, [Mark3](#) builds for the `atmega328p` target, but the target can be selected by manually configuring environment variables, or by running the `./build/set_target.sh` script as follows:

```
./build/set_target.sh <architecture> <variant> <toolchain>
```

Where:

<code><architecture></code>	is the target CPU architecture (i.e. <code>avr</code> , <code>mcp430</code> , <code>cm0</code> , <code>cm3</code> , <code>cm4f</code>)
<code><variant></code>	is the part name (i.e. <code>atmega328p</code> , <code>mcp430f2274</code> , <code>generic</code>)
<code><toolchain></code>	is the build toolchain (i.e. <code>gcc</code>)

This script is a thin wrapper for the cmake configuration commands, and clears the output directory before re-initializing cmake for the selected target.

4.2 Building

To build the [Mark3](#) kernel and middleware libraries for a generic ARM Cortex-M0 using a pre-configured `arm-none-eabi-gcc` toolchain, one would run the following commands:

```
./build/set_target.sh cm0 generic gcc
./build/build.sh
```

To perform an incremental build, go into the target's cmake build directory (`/out/<target>/`) and simply run `'ninja'`.

To perform a clean build, go into the target's cmake build directory (`/out/<target>/`) and run `'ninja clean'`, then `'ninja'`.

Note that not all libraries/tests/examples will build in all kernel configurations. The default kernel configuration may need adjustment/tweaking to support a specific part. See the documentation for details.

4.3 Supported Targets

Currently, [Mark3](#) supports GCC toolchains for the following targets:

```
atmega328p
atmega644
atmega1284p
atxmega256a3
atmega1280
atmega2560
msp430f2274
ARM Cortex-M0
ARM Cortex-M3 (Note: Also supports Cortex-M4)
ARM Cortex-M4F (Note: Also supports Cortex-M7)
```

4.4 More Info

The [Mark3](#) project also has a ready-made docker image to simplify the process. Please see the [Mark3 Docker](#) project here for more details:

<https://github.com/moslevin/mark3-docker>

Chapter 5

The Mark3 API

5.1 Kernel Setup

This section details the process of defining threads, initializing the kernel, and adding threads to the scheduler.

If you're at all familiar with real-time operating systems, then these setup and initialization steps should be familiar. I've tried very hard to ensure that as much of the heavy lifting is hidden from the user, so that only the bare minimum of calls are required to get things started.

The examples presented in this chapter are real, working examples taken from the ATmega328p port.

First, you'll need to create the necessary data structures and functions for the threads:

1. Create a Thread object for all of the "root" or "initial" tasks.
2. Allocate stacks for each of the Threads
3. Define an entry-point function for each Thread

This is shown in the example code below:

```
//-----  
#include "thread.h"  
#include "kernel.h"  
  
//1) Create a thread object for all of the "root" or "initial" tasks  
static Thread AppThread;  
static Thread IdleThread;  
  
//2) Allocate stacks for each thread - in bytes  
#define STACK_SIZE_APP      (192)  
#define STACK_SIZE_IDLE     (128)  
  
static K_WORD awAppStack[STACK_SIZE_APP / sizeof(K_WORD)];  
static K_WORD awIdleStack[STACK_SIZE_IDLE / sizeof(K_WORD)];  
  
//3) Define entry point functions for each thread  
void AppThread(void);  
void IdleThread(void);
```

Next, we'll need to add the required kernel initialization code to main. This consists of running the Kernel's init routine, initializing all of the threads we defined, adding the threads to the scheduler, and finally calling Kernel::Start(), which transfers control of the system to the RTOS.

These steps are illustrated in the following example.

```

int main(void)
{
    //1) Initialize the kernel prior to use
    Kernel::Init();           // MUST be before other kernel ops

    //2) Initialize all of the threads we've defined
    AppThread.Init( awAppStack,      // Pointer to the stack
                   sizeof(awAppStack), // Size of the stack
                   1,                // Thread priority
                   (void*)AppEntry,   // Entry function
                   nullptr );         // Entry function argument

    IdleThread.Init( awIdleStack,     // Pointer to the stack
                   sizeof(awIdleStack), // Size of the stack
                   0,                // Thread priority
                   (void*)IdleEntry,   // Entry function
                   nullptr );         // Entry function argument

    //3) Add the threads to the scheduler
    AppThread.Start();           // Actively schedule the threads
    IdleThread.Start();

    //4) Give control of the system to the kernel
    Kernel::Start();            // Start the kernel!
}

```

Not much to it, is there? There are a few noteworthy points in this code, though.

In order for the kernel to work properly, a system must always contain an idle thread; that is, a thread at priority level 0 that never blocks. This thread is responsible for performing any of the low-level power management on the CPU in order to maximize battery life in an embedded device. The idle thread must also never block, and it must never exit. Either of these operations will cause undefined behavior in the system.

The App thread is at a priority level greater-than 0. This ensures that as long as the App thread has something useful to do, it will be given control of the CPU. In this case, if the app thread blocks, control will be given back to the Idle thread, which will put the CPU into a power-saving mode until an interrupt occurs.

Stack sizes must be large enough to accommodate not only the requirements of the threads, but also the requirements of interrupts - up to the maximum interrupt-nesting level used. Stack overflows are super-easy to run into in an embedded system; if you encounter strange and unexplained behavior in your code, chances are good that one of your threads is blowing its stack.

5.2 Threads

[Mark3](#) Threads act as independent tasks in the system. While they share the same address-space, global data, device-drivers, and system peripherals, each thread has its own set of CPU registers and stack, collectively known as the thread's **context**. The context is what allows the RTOS kernel to rapidly switch between threads at a high rate, giving the illusion that multiple things are happening in a system, when really, only one thread is executing at a time.

5.2.1 Thread Setup

Each instance of the Thread class represents a thread, its stack, its CPU context, and all of the state and metadata maintained by the kernel. Before a Thread will be scheduled to run, it must first be initialized with the necessary configuration data.

The Init function gives the user the opportunity to set the stack, stack size, thread priority, entry-point function, entry-function argument, and round-robin time quantum:

Thread stacks are pointers to blobs of memory (usually char arrays) carved out of the system's address space. Each thread must have a stack defined that's large enough to handle not only the requirements of local variables in the thread's code path, but also the maximum depth of the ISR stack.

Priorities should be chosen carefully such that the shortest tasks with the most strict determinism requirements are executed first - and are thus located in the highest priorities. Tasks that take the longest to execute (and require the least degree of responsiveness) must occupy the lower thread priorities. The idle thread must be the only thread occupying the lowest priority level.

The thread quantum only applies when there are multiple threads in the ready queue at the same priority level. This interval is used to kick-off a timer that will cycle execution between the threads in the priority list so that they each get a fair chance to execute.

The entry function is the function that the kernel calls first when the thread instance is first started. Entry functions have at most one argument - a pointer to a data-object specified by the user during initialization.

An example thread initialization is shown below:

```
Thread clMyThread;
K_WORD awStack[192];

void AppEntry(void)
{
    while(1) {
        // Do something!
    }
}

...
{
    clMyThread.Init(awStack,           // Pointer to the stack to use by this thread
                   sizeof(awStack),    // Size of the stack in bytes
                   1,                  // Thread priority (0 = idle, 7 = max)
                   (void*)AppEntry,    // Function where the thread starts executing
                   nullptr);           // Argument passed into the entry function
}
```

Once a thread has been initialized, it can be added to the scheduler by calling:

```
clMyThread.Start();
```

The thread will be placed into the Scheduler's queue at the designated priority, where it will wait its turn for execution.

5.2.2 Entry Functions

Mark3 Threads should not run-to-completion - they should execute as infinite loops that perform a series of tasks, appropriately partitioned to provide the responsiveness characteristics desired in the system.

The most basic Thread loop is shown below:

```
void Thread( void *param )
{
    while(1) {
        // Do Something
    }
}
```

Threads can interact with each other in the system by means of synchronization objects (Semaphore), mutual-exclusion objects (Mutex), Inter-process messaging (MessageQueue), and timers (Timer).

Threads can suspend their own execution for a predetermined period of time by using the static Thread::Sleep() method. Calling this will block the Thread's execution until the amount of time specified has elapsed. Upon expiry, the thread will be placed back into the ready queue for its priority level, where it awaits its next turn to run.

5.3 Timers

Timer objects are used to trigger callback events periodic or on a one-shot (alarm) basis.

While extremely simple to use, they provide one of the most powerful execution contexts in the system. The timer callbacks execute from within a timer thread, as a result of a semaphore posted in a timer interrupt. Timer callbacks are executed from a high-priority thread – typically at the highest priority thread in the system. Care must be taken to ensure that timer callbacks execute as quickly as possible to minimize the impact of processing on the throughput of tasks in the system. Wherever possible, heavy-lifting should be deferred to lower-priority threads by way of semaphores or messages.

Below is an example showing how to start a periodic system timer which will trigger every second:

```
{
    Timer clTimer;
    clTimer.Init();

    clTimer.Start( 1000,
                  1,
                  MyCallback,
                  (void*)&my_data );

    ... // Keep doing work in the thread
}

// Callback function, executed from the timer-expiry context.
void MyCallBack( Thread *pclOwner_, void *pvData_ )
{
    LED.Flash(); // Flash an LED.
}
```

5.4 Semaphores

Semaphores are used to synchronized execution of threads based on the availability (and quantity) of application-specific resources in the system. They are extremely useful for solving producer-consumer problems, and are the method-of-choice for creating efficient, low latency systems, where ISRs post semaphores that are handled from within the context of individual threads. (Yes, Semaphores can be posted - but not pended - from the interrupt context).

The following is an example of the producer-consumer usage of a binary semaphore:

```
Semaphore clSemaphore; // Declare a semaphore shared between a producer and a consumer thread.

void Producer()
{
    clSemaphore.Init(0, 1);
    while(1) {
        // Do some work, create something to be consumed

        // Post a semaphore, allowing another thread to consume the data
        clSemaphore.Post();
    }
}

void Consumer()
{
    // Assumes semaphore initialized before use...
    While(1) {
        // Wait for new data from the producer thread
        clSemaphore.Pend();

        // Consume the data!
    }
}
```

And an example of using semaphores from the ISR context to perform event- driven processing.

```

Semaphore clSemaphore;

__interrupt__ MyISR()
{
    clSemaphore.Post(); // Post the interrupt. Lightweight when uncontested.
}

void MyThread()
{
    clSemaphore.Init(0, 1); // Ensure this is initialized before the MyISR interrupt is enabled.
    while(1) {
        // Wait until we get notification from the interrupt
        clSemaphore.Pend();

        // Interrupt has fired, do the necessary work in this thread's context
        HeavyLifting();
    }
}

```

5.5 Mutexes

Mutexes (Mutual exclusion objects) are provided as a means of creating "protected sections" around a particular resource, allowing for access of these objects to be serialized. Only one thread can hold the mutex at a time - other threads have to wait until the region is released by the owner thread before they can take their turn operating on the protected resource. Note that mutexes can only be owned by threads - they are not available to other contexts (i.e. interrupts). Calling the mutex APIs from an interrupt will cause catastrophic system failures.

Note that these objects are also not recursive- that is, the owner thread can not attempt to claim a mutex more than once.

Priority inheritance is provided with these objects as a means to avoid priority inversions. Whenever a thread at a priority than the mutex owner blocks on a mutex, the priority of the current thread is boosted to the highest-priority waiter to ensure that other tasks at intermediate priorities cannot artificially prevent progress from being made.

Mutex objects are very easy to use, as there are only three operations supported: Initialize, Claim and Release. An example is shown below.

```

Mutex clMutex; // Create a mutex globally.

void Init()
{
    // Initialize the mutex before use.
    clMutex.Init();
}

// Some function called from a thread
void Thread1Function()
{
    clMutex.Claim();

    // Once the mutex is owned, no other thread can
    // enter a block protect by the same mutex

    my_protected_resource.do_something();
    my_protected_resource.do_something_else();

    clMutex.Release();
}

// Some function called from another thread
void Thread2Function()
{
    clMutex.Claim();

    // Once the mutex is owned, no other thread can
    // enter a block protect by the same mutex

    my_protected_resource.do_something();
    my_protected_resource.do_different_things();

    clMutex.Release();
}

```

5.6 Event Flags

Event Flags are another synchronization object, conceptually similar to a semaphore.

Unlike a semaphore, however, the condition on which threads are unblocked is determined by a more complex set of rules. Each Event Flag object contains a 16-bit field, and threads block, waiting for combinations of bits within this field to become set.

A thread can wait on any pattern of bits from this field to be set, and any number of threads can wait on any number of different patterns. Threads can wait on a single bit, multiple bits, or bits from within a subset of bits within the field.

As a result, setting a single value in the flag can result in any number of threads becoming unblocked simultaneously. This mechanism is extremely powerful, allowing for all sorts of complex, yet efficient, thread synchronization schemes that can be created using a single shared object.

Note that Event Flags can be set from interrupts, but you cannot wait on an event flag from within an interrupt.

Examples demonstrating the use of event flags are shown below.

```
// Simple example showing a thread blocking on a multiple bits in the
// fields within an event flag.

EventFlag clEventFlag;

int main()
{
    ...
    clEventFlag.Init(); // Initialize event flag prior to use
    ...
}

void MyInterrupt()
{
    // Some interrupt corresponds to event 0x0020
    clEventFlag.Set(0x0020);
}

void MyThreadFunc()
{
    ...
    while(1) {
        ...
        uint16_t ul6WakeCondition;

        // Allow this thread to block on multiple flags
        ul6WakeCondition = clEventFlag.Wait(0x00FF, EventFlagOperation::Any_Set);

        // Clear the event condition that caused the thread to wake (in this case,
        // ul6WakeCondition will equal 0x20 when triggered from the interrupt above)
        clEventFlag.Clear(ul6WakeCondition);

        // <do something>
    }
}
```

5.7 Messages

Sending messages between threads is the key means of synchronizing access to data, and the primary mechanism to perform asynchronous data processing operations.

Sending a message consists of the following operations:

- Obtain a Message object from the global message pool
- Set the message data and event fields

- Send the message to the destination message queue

While receiving a message consists of the following steps:

- Wait for a messages in the destination message queue
- Process the message data
- Return the message back to the global message pool

These operations, and the various data objects involved are discussed in more detail in the following section.

5.7.1 Message Objects

Message objects are used to communicate arbitrary data between threads in a safe and synchronous way.

The message object consists of an event code field and a data field. The event code is used to provide context to the message object, while the data field (essentially a void * data pointer) is used to provide a payload of data corresponding to the particular event.

Access to these fields is marshalled by accessors - the transmitting thread uses the `SetData()` and `SetCode()` methods to seed the data, while the receiving thread uses the `GetData()` and `GetCode()` methods to retrieve it.

By providing the data as a void data pointer instead of a fixed-size message, we achieve an unprecedented measure of simplicity and flexibility. Data can be either statically or dynamically allocated, and sized appropriately for the event without having to format and reformat data by both sending and receiving threads. The choices here are left to the user - and the kernel doesn't get in the way of efficiency.

It is worth noting that you can send messages to message queues from within ISR context. This helps maintain consistency, since the same APIs can be used to provide event-driven programming facilities throughout the whole of the OS.

5.7.2 Message Queues

Message objects specify data with context, but do not specify where the messages will be sent. For this purpose we have a `MessageQueue` object. Sending an object to a message queue involves calling the `MessageQueue::Send()` method, passing in a pointer to the `Message` object as an argument.

When a message is sent to the queue, the first thread blocked on the queue (as a result of calling the `MessageQueue::Receive()` method) will wake up, with a pointer to the `Message` object returned.

It's worth noting that multiple threads can block on the same message queue, providing a means for multiple threads to share work in parallel.

5.7.3 Messaging Example

```
// Message queue object shared between threads
MessageQueue c1MsgQ;

// Function that initializes the shared message queue
void MsgQInit()
{
    c1MsgQ.Init();
}

// Function called by one thread to send message data to
// another
void TxMessage()
{
    // Get a message, initialize its data
    Message *pclMesg = MyMessagePool.Pop();

    pclMesg->SetCode(0xAB);
    pclMesg->SetData((void*)some_data);

    // Send the data to the message queue
    c1MsgQ.Send(pclMesg);
}

// Function called in the other thread to block until
// a message is received in the message queue.
void RxMessage()
{
    Message *pclMesg;

    // Block until we have a message in the queue
    pclMesg = c1MsgQ.Receive();

    // Do something with the data once the message is received
    pclMesg->GetCode();

    // Free the message once we're done with it.
    MyMessagePool.Push(pclMesg);
}
```

5.8 Mailboxes

Another form of IPC is provided by [Mark3](#), in the form of Mailboxes and Envelopes.

Mailboxes are similar to message queues in that they provide a synchronized interface by which data can be transmitted between threads.

Where Message Queues rely on linked lists of lightweight message objects (containing only message code and a void* data-pointer), which are inherently abstract, Mailboxes use a dedicated blob of memory, which is carved up into fixed-size chunks called Envelopes (defined by the user), which are sent and received. Unlike message queues, mailbox data is copied to and from the mailboxes dedicated pool.

Mailboxes also differ in that they provide not only a blocking "receive" call, but also a blocking "send" call, providing the opportunity for threads to block on "mailbox full" as well as "mailbox empty" conditions.

All send/receive APIs support an optional timeout parameter if the KERNEL_USE_TIMEOUTS option has been configured in [mark3cfg.h](#)

5.8.1 Mailbox Example

```
// Create a mailbox object, and define a buffer that will be used to store the
// mailbox' envelopes.
static Mailbox c1Mbox;
static uint8_t aucMBoxBuffer[128];

...
void InitMailbox(void)
{
```

```

    // Initialize our mailbox, telling it to use our defined buffer for envelope
    // storage. Pass in the size of the buffer, and set the size of each
    // envelope to 16 bytes. This gives u16 a mailbox capacity of (128 / 16) = 8
    // envelopes.
    clMbox.Init((void*)aucMBoxBuffer, 128, 16);
}

...
void SendThread(void)
{
    // Define a buffer that we'll eventually send to the
    // mailbox. Note the size is the same as that of an
    // envelope.
    uint8_t aucTxBuf[16];

    while(1) {
        // Copy some data into aucTxBuf, a 16-byte buffer, the
        // same size as a mailbox envelope.
        ...

        // Deliver the envelope (our buffer) into the mailbox
        clMbox.Send((void*)aucTxBuf);
    }
}

...
void RecvThred(void)
{
    uint8_t aucRxBuf[16];

    while(1) {
        // Wait until there's a message in our mailbox. Once
        // there is a message, read it into our local buffer.
        cmMbox.Receive((void*)aucRxBuf);

        // Do something with the contents of aucRxBuf, which now
        // contains an envelope of data read from the mailbox.
        ...
    }
}

```

5.9 Notification Objects

Notification objects are the most lightweight of all blocking objects supplied by [Mark3](#).

using this blocking primitive, one or more threads wait for the notification object to be signalled by code elsewhere in the system (i.e. another thread or interrupt). Once the the notification has been signalled, all threads currently blocked on the object become unblocked.

5.9.1 Notification Example

```

NOTIFIER

static Notify clNotifier;

...
void MyThread(void *unused_)
{
    // Initialize our notification object before use
    clNotifier.Init();

    while (1)
    {
        // Wait until our thread has been notified that it
        // can wake up.
        clNotify.Wait();

        ...
        // Thread has woken up now -- do something!
    }
}

...
void SignalCallback(void)

```

```

{
    // Something in the system (interrupt, thread event, IPC,
    // etc.,) has called this function. As a result, we need
    // our other thread to wake up. Call the Notify object's
    // Signal() method to wake the thread up. Note that this
    // will have no effect if the thread is not presently
    // blocked.

    clNotify.Signal();
}

```

5.10 Condition Variables

Condition Variables, implemented in [Mark3](#) with the `ConditionVariable` class, provide an implementation of the classic Monitor pattern. This object allows a thread to wait for a specific condition to occur, claiming a shared lock once the condition is met. Threads may also choose to signal a single blocking thread to indicate a condition has changed, or broadcast condition changes to all waiting threads.

5.10.1 Condition Variable Example

```

// Define a condition variable object, a shared lock, and
// a piece of common data shared between threads to represent
// a condition.

// Assume Mutex and ConditionVariable are initialized
// prior to use.
static ConditionVariable clCondVar;
static Mutex clSharedLock;
static volatile int iCondition = 0;

...

void CondThread1(void *unused_)
{
    while (1)
    {
        // Wait until
        clCondVar.Wait(&clSharedLock);

        // Only act on a specific condition
        if (iCondition == 1337) {
            // Do something
        }

        clSharedLock.Release();
    }
}

void CondThread2(void *unused_)
{
    // Assume Mutex and ConditionVariable are initialized
    // prior to use.
    while (1)
    {
        // Wait until
        clCondVar.Wait(&clSharedLock);

        // Act on a *different* condition than the other thread
        if (iCondition == 5454) {
            // Do something
        }

        clSharedLock.Release();
    }
}

void SignalThread(void* unused)
{
    while (1) {
        // Sleep for a while
        Thread::Sleep(100);

        // Update the condition in a thread-safe manner
        clSharedLock.Claim();
        iCondition = 1337;
    }
}

```



```

        clSharedLock.Release();

        // Wake one thread to check for the updated condition
        clCondVar.Signal();

        // Sleep for a while
        Thread.Sleep(100);

        // Update the condition in a thread-safe manner
        clSharedLock.Claim();
        iCondition = 1337;
        clSharedLock.Release();

        // Wake all threads to check for the updated condition
        clCondVar.Broadcast();
    }
}

```

5.11 Reader-Write Locks

Reader-Writer locks are provided in [Mark3](#) to provide an efficient way for multiple threads to share concurrent, non-destructive access to a resource, while preventing concurrent destructive/non-destructive accesses. A single "writer" may hold the lock, or 1-or-more "readers" may hold the lock. In the case that readers hold the lock, writers will block until all readers have relinquished their access to the resource. In the case that a writer holds the lock, all other readers and writers must wait until the lock is relinquished.

5.11.1 Reader-Write Lock Example

```

void WriterTask(void* param)
{
    auto pclRWLock = static_cast<ReaderWriterLock*>(param);

    pclRWLock->AcquireWriter();
    // All other readers and writers will have to wait until
    // the lock is released.
    iNumWrites++;
    ...
    pclRWLock->ReleaseWriter();
}

void ReaderTask(void* param)
{
    auto pclRWLock = static_cast<ReaderWriterLock*>(param);

    pclRWLock->AcquireReader();
    // Any number of reader threads can also acquire the lock
    // without having to block, waiting for this task to release it.
    // Writers must block until all readers have released their references
    // to the lock.
    iNumReads++;
    ...
    pclRWLock->ReleaseReader();
}

```

5.12 Sleep

There are instances where it may be necessary for a thread to poll a resource, or wait a specific amount of time before proceeding to operate on a peripheral or volatile piece of data.

While the Timer object is generally a better choice for performing time-sensitive operations (and certainly a better choice for periodic operations), the Thread::Sleep() method provides a convenient (and efficient) mechanism that allows for a thread to suspend its execution for a specified interval.

Note that when a thread is sleeping it is blocked, during which other threads can operate, or the system can enter its idle state.

```
int GetPeripheralData()
{
    int value;
    // The hardware manual for a peripheral specifies that
    // the "foo()" method will result in data being generated
    // that can be captured using the "bar()" method.
    // However, the value only becomes valid after 10ms

    peripheral.foo();
    Thread::Sleep(10); // Wait 10ms for data to become valid
    value = peripheral.bar();
    return value;
}
```

5.13 Round-Robin Quantum

Threads at the same thread priority are scheduled using a round-robin scheme. Each thread is given a timeslice (which can be configured) of which it shares time amongst ready threads in the group. Once a thread's timeslice has expired, the next thread in the priority group is chosen to run until its quantum has expired - the cycle continues over and over so long as each thread has work to be done.

By default, the round-robin interval is set at 4ms.

This value can be overridden by calling the thread's `SetQuantum()` with a new interval specified in milliseconds.

5.14 Coroutines

[Mark3](#) implements a coroutine scheduler, capable of managing a set of prioritized run-to-completion tasks. This is a simple and lightweight cooperative scheduling mechanism, that trades the preemption and synchronization capabilities of threads for simplicity. It is an ideal mechanism to use for background processes in a system, or for performing a coordinating a group of tasks where the relative priority of task execution is important, but the duration of individual tasks is less important.

Like the [Mark3](#) thread scheduler, the coroutine scheduler supports multiple priorities of tasks. Multiple coroutines activated at the same priority level are executed in first-in first-out order.

Coroutines are activated by interrupts, threads, or from within other co-routines. Once activated, the co-routine is able to be scheduled.

The coroutine scheduler is called by the application from a thread priority. So long as there are activated tasks to be scheduled, the scheduler will return a pointer to the highest priority active coroutine to be run.

Running a co-routine de-activates the co-routine, meaning that coroutines must be re-activated every time they are run.

```
Coroutine cr1;
Coroutine cr2;
Coroutine cr3;

void CoRoutineInit()
{
    CoScheduler::Init();

    cr1.Init(<priority>, <handler function>, <data passed to handler function>);
    cr2.Init(<priority>, <handler function>, <data passed to handler function>);
    cr3.Init(<priority>, <handler function>, <data passed to handler function>);
}

void AsyncEvent1()
{
    // Some event occurred requiring us to run cr1. Activating a task does not
    // cause the coroutine to be run immediately, but schedules it to be returned by
```

```

    // CoScheduler::Schedule() when there are no higher-priority active tasks to
    // run.
    cr1.Activate();
}

void AsyncEvent2()
{
    // Some event occurred requiring us to run cr2.
    cr2.Activate();
}

void AsyncEvent3()
{
    // Some event occurred requiring us to run cr3.
    cr3.Activate();
}

...
void CoRoutineSchedulerTask()
{
    while (1) {
        auto* coroutineToRun = CoScheduler::Schedule();
        if (coroutineToRun) {
            // Scheduled task is removed from the scheduler once run. It must
            // be reactivated to be run again
            coroutineToRun->Run();
        }
        // If run from idle thread, could go into sleep mode instead of running in
        // a tight loop.
    }
}

```

5.15 Critical Guards

Often times, it is useful in a real-time multi-threaded system to place a critical section around a block of code to protect it against concurrent access, or to protect global data from access from interrupts. In [Mark3](#) there are a few different ways of implementing critical sections. Historically, the `CS_ENTER()` and `CS_EXIT()` macros were used for this purpose; however, the usage of matching entry/exit macros can often-times be cumbersome and error-prone.

The `CriticalGuard` object allows a user to wrap a block of code in a critical section, where the critical section is entered when the critical guard object is declared, and the critical section is exited when the object goes out-of-scope.

It is essentially an RAIL-style critical section object, that provides the benefit of critical sections without the hassle of having to carefully match enter/exit statements.

```

void MyFunc()
{
    // operations outside of critical section
    {
        auto cg = CriticalGuard{};
        // operations protected by critical section
        // critical section ends when CriticalGuard object goes out of scope
    }
    // Operations outside of critical section
}

```

5.16 Lock Guards

The modern C++ standard library provides RAIL-style mutex locking. Unfortunately such an implementation is not usable in the context of [Mark3](#), because the STL implementation supplied for an embedded C++ toolchain would not be aware of [Mark3](#)'s threading model or synchronization primitives. [Mark3](#) provides its own RAIL mutex locking mechanism in the form of `LockGuard` objects. When a `LockGuard` object is declared (referencing a valid and initialized mutex object at construction), the lock is claimed upon declaration, and released when the object goes out-of-scope.

```
Mutex clMutex;
void MyFunc()
{
    // operations outside of mutex-locked context
    {
        auto lg = LockGuard{&clMutex};
        // operations inside mutex-locked context
        // mutex automatically unlocked when LockGuard goes out of scope
    }
    // Operations outside of mutex-locked context
}
```

5.17 Scheduler Guards

Similar to the LockGuard and CriticalGuard objects, the SchedulerGuard object provides a copped scheduler-disabled blocks. This essentially gives the executing thread exclusive control of the CPU - except for interrupts - for the duration of the block wrapped in the SchedulerGuard. The scheduler is disabled when the object is declared, and scheduler state is restored when the SchedulerGuard object goes out-of-scope. This is yet another form of RAII-based resource locking in [Mark3](#).

```
void MyFunc()
{
    // thread scheduler enabled
    {
        auto sg = SchedulerGuard{};
        // thread-safety guaranteed - scheduler disabled
        // thread scheduler re-enabled when SchedulerGuard object goes out of scope
    }
    // thread scheduler enabled
}
```

Chapter 6

Why Mark3?

My first job after graduating from university in 2005 was with a small company that had a very old-school, low-budget philosophy when it came to software development. Every make-or-buy decision ended with "make" when it came to tools. It was the kind of environment where vendors cost us money, but manpower was free. In retrospect, we didn't have a ton of business during the time that I worked there, and that may have had something to do with the fact that we were constantly short on ready cash for things we could code ourselves.

Early on, I asked why we didn't use industry-standard tools - like JTAG debuggers or IDEs. One senior engineer scoffed that debuggers were tools for wimps - and something that a good programmer should be able to do without. After all - we had serial ports, GPIOs, and a bi-color LED on our boards. Since these were built into the hardware, they didn't cost us a thing. We also had a single software "build" server that took 5 minutes to build a 32k binary on its best days, so when we had to debug code, it was a painful process of trial and error, with lots of Youtube between iterations. We complained that tens of thousands of dollars of productivity was being flushed away that could have been solved by implementing a proper build server - and while we eventually got our wish, it took far more time than it should have.

Needless to say, software development was painful at that company. We made life hard on ourselves purely out of pride, and for the right to say that we walked "up-hills both ways through 3 feet of snow, everyday". Our code was tied ever-so-tightly to our hardware platform, and the system code was indistinguishable from the application. While we didn't use an RTOS, we had effectively implemented a 3-priority threading scheme using a carefully designed interrupt nesting scheme with event flags and a while(1) superloop running as a background thread. Nothing was abstracted, and the code was always optimized for the platform, presumably in an effort to save on code size and wasted cycles. I asked why we didn't use an RTOS in any of our systems and received dismissive scoffs - the overhead from thread switching and maintaining multiple threads could not be tolerated in our systems according to our chief engineers. In retrospect, our ad-hoc system was likely as large as my smallest kernel, and had just as much context switching (although it was hidden by the compiler).

And every time a new iteration of our product was developed, the firmware took far too long to bring up, because the algorithms and data structures had to be re-tooled to work with the peripherals and sensors attached to the new boards. We worked very hard in an attempt to reinvent the wheel, all in the name of producing "efficient" code.

Regardless, I learned a lot about embedded software development.

Most important, I learned that good design is the key to good software; and good design doesn't have to come at a price. In all but the smallest of projects, the well-designed, well-abstracted code is not only more portable, but it's usually smaller, easier to read, and easier to reuse.

Also, since we had all the time in the world to invest in developing our own tools, I gained a lot of experience building them, and making use of good, free PC tools that could be used to develop and debug a large portion of our code. I ended up writing PC-based device and peripheral simulators, state-machine frameworks, and abstractions for our horrible ad-hoc system code. At the end of the day, I had developed enough tools that I could solve a lot of our development problems without having to re-inventing the wheel at each turn. Gaining a background in how these

tools worked gave me a better understanding of how to use them - making me more productive at the jobs that I've had since.

I am convinced that designing good software takes honest effort up-front, and that good application code cannot be written unless it is based on a solid framework. Just as the wise man builds his house on rocks, and not on sand, wise developers write applications based on a well-defined platforms. And while you can probably build a house using nothing but a hammer and sheer will, you can certainly build one a lot faster with all the right tools.

This conviction lead me to development my first RTOS kernel in 2009 - FunkOS. It is a small, yet surprisingly full-featured kernel. It has all the basics (semaphores, mutexes, round-robin and preemptive scheduling), and some pretty advanced features as well (device drivers and other middleware). However, it had two major problems - it doesn't scale well, and it doesn't support many devices.

While I had modest success with this kernel (it has been featured on some blogs, and still gets around 125 downloads a month), it was nothing like the success of other RTOS kernels like uC/OS-II and FreeRTOS. To be honest, as a one-man show, I just don't have the resources to support all of the devices, toolchains, and evaluation boards that a real vendor can. I had never expected my kernel to compete with the likes of them, and I don't expect [Mark3](#) to change the embedded landscape either.

My main goal with [Mark3](#) was to solve the technical shortfalls in the FunkOS kernel by applying my experience in kernel development. As a result, [Mark3](#) is better than FunkOS in almost every way; it scales better, has lower interrupt latency, and is generally more thoughtfully designed (all at a small cost to code size).

Another goal I had was to create something easy to understand, that could be documented and serve as a good introduction to RTOS kernel design. The end result of these goals is the kernel as presented in this book - a full source listing of a working OS kernel, with each module completely documented and explained in detail.

Finally, I wanted to prove that a kernel written entirely in C++ could perform just as well as one written in C. [Mark3](#) is fully benchmarked and profiled – you can see exactly how much it costs to call certain APIs or include various features in the kernel.

And in addition, the code is more readable and easier to understand as a result of making use of object-oriented concepts provided by C++. Applications are easier to write because common concepts are encapsulated into objects (Threads, Semaphores, Mutexes, etc.) with their own methods and data, as opposed to APIs which rely on lots of explicit pointer or handle-passing, type casting, and other operations that are typically considered "unsafe" or "advanced" topics in C.

Chapter 7

When to use an RTOS?

7.1 The reality of system code

System code can be defined as the program logic required to manage, synchronize, and schedule all of the resources (CPU time, memory, peripherals, etc.) used by the application running on the CPU. And it's true that a significant portion of the code running on an embedded system will be system code. No matter how simple a system is, whether or not this logic is embedded directly into the application (bare-metal system), or included as part of a well-defined stack on which an application is written (RTOS-based); system code is still present, and it comes with a cost.

As an embedded systems is being designed, engineers have to decide which approach to take: Bare-metal, or RTOS. There are advantages and disadvantages to each – and a reasonable engineer should always perform a thorough analysis of the pros and cons of each - in the context of the given application - before choosing a path.

The following figure demonstrates the differences between the architecture of a bare-metal system and RTOS based system at a high level:




Figure 7.1 Arch

As can be seen, the RTOS (And associated middleware + libraries) captures a certain fixed size.

As a generalization, bare-metal systems typically have the advantage in that the system code overhead is small to start – but grows significantly as the application grows in complexity. At a certain point, it becomes extremely difficult and error-prone to add more functionality to an application running on such a system. There's a tipping point, where the cost of the code used to work-around the limitations of a bare-metal system outweigh the cost of a capable RTOS. Bare-metal systems also generally take longer to implement, because the system code has to be written from scratch (or derived from existing code) for the application. The resulting code also tends to be less portable, as it takes serious discipline to keep the system-specific elements of the code separated – in an RTOS-based system, once the kernel and drivers are ported, the application code is generally platform agnostic.

Conversely, an RTOS-based system incurs a slightly higher fixed cost up-front, but scales infinitely better than a bare-metal system as application's complexity increases. Using an RTOS for simple systems reduces application development time, but may cause an application not to fit into some extremely size-constrained microcontroller. An RTOS can also cause the size of an application to grow more slowly relative to a bare-metal system – especially as a result of applying synchronization mechanisms and judicious IPC. As a result, an RTOS makes it significantly easier to “go agile” with an application – iteratively adding features and functionality, without having to consider refactoring the underlying system at each turn.

Some of these factors may be more important than others. Requirements, specifications, schedules, chip-selection, and volume projections for a project should all be used to feed into the discussions to decide whether or to go bare-metal or RTOS as a result.

Consider the following questions when making that decision:

- What is the application?
- How efficient is efficient enough?
- How fast is fast enough?
- How small is small enough?
- How responsive is responsive enough?
- How much code space/RAM/etc is available on the target system?
- How much code space/RAM do I need for an RTOS?
- How much code space/RAM do I think I'll need for my application?
- How much time do I have to deliver my system?
- How many units do we plan to sell?

7.2 Superloops, and their limitations

7.2.1 Intro to Superloops

Before we start taking a look at designing a real-time operating system, it's worthwhile taking a look through one of the most-common design patterns that developers use to manage task execution in bare-metal embedded systems - Superloops.

Systems based on superloops favor the system control logic baked directly into the application code, usually under the guise of simplicity, or memory (code and RAM) efficiency. For simple systems, superloops can definitely get the job done. However, they have some serious limitations, and are not suitable for every kind of project. In a lot of cases you can squeak by using superloops - especially in extremely constrained systems, but in general they are not a solid basis for reusable, portable code.

Nonetheless, a variety of examples are presented here- from the extremely simple, to cooperative and limited-preemptive multitasking systems, all of which are examples are representative of real-world systems that I've either written the firmware for, or have seen in my experience.

7.2.2 The simplest loop

Let's start with the simplest embedded system design possible - an infinite loop that performs a single task repeatedly:

```
int main()
{
    while(1)
    {
        Do_Something();
    }
}
```

Here, the code inside the loop will run a single function forever and ever. Not much to it, is there? But you might be surprised at just how much embedded system firmware is implemented using essentially the same mechanism - there isn't anything wrong with that, but it's just not that interesting.

Despite its simplicity we can see the beginnings of some core OS concepts. Here, the `while(1)` statement can be logically seen as the operating system kernel - this one control statement determines what tasks can run in the system, and defines the constraints that could modify their execution. But at the end of the day, that's a big part of what a kernel is - a mechanism that controls the execution of application code.

The second concept here is the task. This is application code provided by the user to perform some useful purpose in a system. In this case `Do_something()` represents that task - it could be monitoring blood pressure, reading a sensor and writing its data to a terminal, or playing an MP3; anything you can think of for an embedded system to do. A simple round-robin multi-tasking system can be built off of this example by simply adding additional tasks in sequence in the main while-loop. Note that in this example the CPU is always busy running tasks - at no time is the CPU idle, meaning that it is likely burning a lot of power.

While we conceptually have two separate pieces of code involved here (an operating system kernel and a set of running tasks), they are not logically separate. The OS code is indistinguishable from the application. It's like a single-celled organism - everything is crammed together within the walls of an indivisible unit; and specialized to perform its given function relying solely on instinct.

7.2.3 Interrupt-Driven Super-loop

In the previous example, we had a system without any way to control the execution of the task- it just runs forever. There's no way to control when the task can (or more importantly can't) run, which greatly limits the usefulness of the system. Say you only want your task to run every 100 milliseconds - in the previous code, you have to add a hard-coded delay at the end of your task's execution to ensure your code runs only when it should.

Fortunately, there is a much more elegant way to do this. In this example, we introduce the concept of the synchronization object. A Synchronization object is some data structure which works within the bounds of the operating system to tell tasks when they can run, and in many cases includes special data unique to the synchronization event.

There are a whole family of synchronization objects, which we'll get into later. In this example, we make use of the simplest synchronization primitive

- the global flag.

With the addition of synchronization brings the addition of event-driven systems. If you're programming a microcontroller system, you generally have scores of peripherals available to you - timers, GPIOs, ADCs, UARTs, ethernet, USB, etc. All of which can be configured to provide a stimulus to your system by means of interrupts. This stimulus gives us the ability not only to program our micros to `do_something()`, but to `do_something()` if-and-only-if a corresponding trigger has occurred.

The following concepts are shown in the example below:

```
volatile K_BOOL something_to_do = false;

__interrupt__ My_Interrupt_Source(void)
{
    something_to_do = true;
}

int main()
{
    while (1)
    {
        if (something_to_do)
        {
            Do_something();
            something_to_do = false;
        }
        else
        {
            Idle();
        }
    }
}
```

So there you have it - an event driven system which uses a global variable to synchronize the execution of our task based on the occurrence of an interrupt. It's still just a bare-metal, OS-baked-into-the-application system, but it's introduced a whole bunch of added complexity (and control!) into the system.

The first thing to notice in the source is that the global variable, `something_to_do`, is used as a synchronization object. When an interrupt occurs from some external event, triggering the `My_Interrupt_Source()` ISR, program flow in `main()` is interrupted, the interrupt handler is run, and `something_to_do` is set to true, letting us know that when we get back to `main()`, that we should run our `Do_something()` task.

Another new concept at play here is that of the idle function. In general, when running an event driven system, there are times when the CPU has no application tasks to run. In order to minimize power consumption, CPUs usually contain instructions or registers that can be set up to disable non-essential subsets of the system when there's nothing to do. In general, the sleeping system can be re-activated quickly as a result of an interrupt or other external stimulus, allowing normal processing to resume.

Now, we could just call `Do_something()` from the interrupt itself - but that's generally not a great solution. In general, the more time we spend inside an interrupt, the more time we spend with at least some interrupts disabled. As a result, we end up with interrupt latency. Now, in this system, with only one interrupt source and only one task this might not be a big deal, but say that `Do_something()` takes several seconds to complete, and in that time several other interrupts occur from other sources. While executing in our long-running interrupt, no other interrupts can be processed - in many cases, if two interrupts of the same type occur before the first is processed, one of these interrupt events will be lost. This can be utterly disastrous in a real-time system and should be avoided at all costs. As a result, it's generally preferable to use synchronization objects whenever possible to defer processing outside of the ISR.

Another OS concept that is implicitly introduced in this example is that of task priority. When an interrupt occurs, the normal execution of code in `main()` is preempted: control is swapped over to the ISR (which runs to completion), and then control is given back to `main()` where it left off. The very fact that interrupts take precedence over what's running shows that `main` is conceptually a "low-priority" task, and that all ISRs are "high-priority" tasks. In this example, our "high-priority" task is setting a variable to tell our "low-priority" task that it can do something useful. We will investigate the concept of task priority further in the next example.

Preemption is another key principle in embedded systems. This is the notion that whatever the CPU is doing when an interrupt occurs, it should stop, cache its current state (referred to as its context), and allow the high-priority event to be processed. The context of the previous task is then restored its state before the interrupt, and resumes processing. We'll come back to preemption frequently, since the concept comes up frequently in RTOS-based systems.

7.2.4 Cooperative multi-tasking

Our next example takes the previous example one step further by introducing cooperative multi-tasking:

```
// Bitfield values used to represent three distinct tasks
#define TASK_1_EVENT (0x01)
#define TASK_2_EVENT (0x02)
#define TASK_3_EVENT (0x04)

volatile K_UCHAR event_flags = 0;

// Interrupt sources used to trigger event execution

__interrupt__ My_Interrupt_1(void)
{
    event_flags |= TASK_1_EVENT;
}

__interrupt__ My_Interrupt_2(void)
{
    event_flags |= TASK_2_EVENT;
}

__interrupt__ My_Interrupt_3(void)
{
    event_flags |= TASK_3_EVENT;
}

// Main tasks
int main(void)
{
    while(1)
    {
        while(event_flags)
        {
            if( event_flags & TASK_1_EVENT)
            {
                Do_Task_1();
                event_flags &= ~TASK_1_EVENT;
            } else if( event_flags & TASK_2_EVENT) {
                Do_Task_2();
                event_flags &= ~TASK_2_EVENT;
            } else if( event_flags & TASK_3_EVENT) {
                Do_Task_3();
                event_flags &= ~TASK_3_EVENT;
            }
        }
        Idle();
    }
}
```

This system is very similar to what we had before - however the differences are worth discussing. First, we have stimulus from multiple interrupt sources: each ISR is responsible for setting a single bit in our global event flag, which is then used to control execution of individual tasks from within main().

Next, we can see that tasks are explicitly given priorities inside the main loop based on the logic of the if/else if structure. As long as there is something set in the event flag, we will always try to execute Task1 first, and only when Task1 isn't set will we attempt to execute Task2, and then Task3. This added logic provides the notion of priority. However, because each of these tasks exist within the same context (they're just different functions called from our main control loop), we don't have the same notion of preemption that we have when dealing with interrupts.

That means that even through we may be running Task2 and an event flag for Task1 is set by an interrupt, the CPU still has to finish processing Task2 to completion before Task1 can be run. And that's why this kind of scheduling is referred to as cooperative multitasking: we can have as many tasks as we want, but unless they cooperate by means of returning back to main, the system can end up with high-priority tasks getting starved for CPU time by lower-priority, long-running tasks.

This is one of the more popular Os-baked-into-the-application approaches, and is widely used in a variety of real-time embedded systems.

7.2.5 Hybrid cooperative/preemptive multi-tasking

The final variation on the superloop design utilizes software-triggered interrupts to simulate a hybrid cooperative/preemptive multitasking system. Consider the example code below.

```
// Bitfields used to represent high-priority tasks. Tasks in this group
// can preempt tasks in the group below - but not eachother.
#define HP_TASK_1(0x01)
#define HP_TASK_2(0x02)

volatile K_UCHAR hp_tasks = 0;

// Bitfields used to represent low-priority tasks.
#define LP_TASK_1(0x01)
#define LP_TASK_2(0x02)

volatile K_UCHAR lp_tasks = 0;

// Interrupt sources, used to trigger both high and low priority tasks.
__interrupt__ System_Interrupt_1(void)
{
    // Set any of the other tasks from here...
    hp_tasks |= HP_TASK_1;
    // Trigger the SWI that calls the High_Priority_Tasks interrupt handler
    SWI();
}

__interrupt__ System_Interrupt_n...(void)
{
    // Set any of the other tasks from here...
}

// Interrupt handler that is used to implement the high-priority event context
__interrupt__ High_Priority_Tasks(void)
{
    // Enabled every interrupt except this one
    Disable_My_Interrupt();
    Enable_Interrupts();
    while( hp_tasks)
    {
        if( hp_tasks & HP_TASK_1)
        {
            HP_Task1();
            hp_tasks &= ~HP_TASK_1;
        }
        else if (hp_tasks & HP_TASK_2)
        {
            HP_Task2();
            hp_tasks &= ~HP_TASK_2;
        }
    }
    Restore_Interrupts();
    Enable_My_Interrupt();
}

// Main loop, used to implement the low-priority events
int main(void)
{
    // Set the function to run when a SWI is triggered
    Set_SWI(High_Priority_Tasks);

    // Run our super-loop
    while(1)
    {
        while (lp_tasks)
        {
            if (lp_tasks & LP_TASK_1)
            {
                LP_Task1();
                lp_tasks &= ~LP_TASK_1;
            }
            else if (lp_tasks & LP_TASK_2)
            {
                LP_Task2();
                lp_tasks &= ~LP_TASK_2;
            }
        }
        Idle();
    }
}
```

In this example, `High_Priority_Tasks()` can be triggered at any time as a result of a software interrupt (SWI). When a high-priority event is set, the code that sets the event calls the SWI as well, which instantly preempts whatever is happening in `main`, switching to the high-priority interrupt handler. If the CPU is executing in an interrupt handler already, the current ISR completes, at which point control is given to the high priority interrupt handler.

Once inside the HP ISR, all interrupts (except the software interrupt) are re-enabled, which allows this interrupt to be preempted by other interrupt sources, which is called interrupt nesting. As a result, we end up with two distinct execution contexts (`main` and `HighPriorityTasks()`), in which all tasks in the high-priority group are guaranteed to preempt `main()` tasks, and will run to completion before returning control back to tasks in `main()`. This is a very basic preemptive multitasking scenario, approximating a "real" RTOS system with two threads of different priorities.

7.3 Problems with superloops

As mentioned earlier, a lot of real-world systems are implemented using a superloop design; and while they are simple to understand due to the limited and obvious control logic involved, they are not without their problems.

7.3.1 Hidden Costs

It's difficult to calculate the overhead of the superloop and the code required to implement workarounds for blocking calls, scheduling, and preemption. There's a cost in both the logic used to implement workarounds (usually involving state machines), as well as a cost to maintainability that comes with breaking up into chunks based on execution time instead of logical operations. In moderate firmware systems, this size cost can exceed the overhead of a reasonably well-featured RTOS, and the deficit in maintainability is something that is measurable in terms of lost productivity through debugging and profiling.

7.3.2 Tightly-coupled code

Because the control logic is integrated so closely with the application logic, a lot of care must be taken not to compromise the separation between application and system code. The timing loops, state machines, and architecture-specific control mechanisms used to avoid (or simulate) preemption can all contribute to the problem. As a result, a lot of superloop code ends up being difficult to port without effectively simulating or replicating the underlying system for which the application was written. Abstraction layers can mitigate the risks, but a lot of care should be taken to fully decouple the application code from the system code.

7.3.3 No blocking Calls

In a super-loop environment, there's no such thing as a blocking call or blocking objects. Tasks cannot stop mid-execution for event-driven I/O from other contexts - they must always run to completion. If busy-waiting and polling are used as a substitute, it increases latency and wastes cycles. As a result, extra code complexity is often times necessary to work-around this lack of blocking objects, often times through implementing additional state machines. In a large enough system, the added overhead in code size and cycles can add up.

7.3.4 Difficult to guarantee responsiveness

Without multiple levels of priority, it may be difficult to guarantee a certain degree of real-time responsiveness without added profiling and tweaking. The latency of a given task in a priority-based cooperative multitasking system is the length of the longest task. Care must be taken to break tasks up into appropriate sized chunks in order to ensure that higher- priority tasks can run in a timely fashion - a manual process that must be repeated as new tasks are added in the system. Once again, this adds extra complexity that makes code larger, more difficult to understand and maintain due to the artificial subdivision of tasks into time-based components.

7.3.5 Limited preemption capability

As shown in the example code, the way to gain preemption in a superloop is through the use of nested interrupts. While this isn't unwieldy for two levels of priority, adding more levels beyond this is becomes complicated. In this case, it becomes necessary to track interrupt nesting manually, and separate sets of tasks that can run within given priority loops - and deadlock becomes more difficult to avoid.

Chapter 8

Can you afford an RTOS?

8.1 Intro

Of course, since you're reading the manual for an RTOS that I've been developing over the course of the several years, you can guess that the conclusion that I draw.

If your code is of any sort of non-trivial complexity (say, at least a few- thousand lines), then a more appropriate question would be "can you afford not* to use an RTOS in your system?".

In short, there are simply too many benefits of an RTOS to ignore, the most important being:

Threading, along with priority and time-based scheduling
Sophisticated synchronization objects and IPC
Flexible, powerful Software Timers
Ability to write more portable, decoupled code

Sure, these features have a cost in code space and RAM, but from my experience the cost of trying to code around a lack of these features will cost you as much - if not more. The results are often far less maintainable, error prone, and complex. And that simply adds time and cost. Real developers ship, and the RTOS is quickly becoming one of the standard tools that help keep developers shipping.

One of the main arguments against using an RTOS in an embedded project is that the overhead incurred is too great to be justified. Concerns over "wasted" RAM caused by using multiple stacks, added CPU utilization, and the "large" code footprint from the kernel cause a large number of developers to shun using a preemptive RTOS, instead favoring a non-preemptive, application-specific solution.

I believe that not only is the impact negligible in most cases, but that the benefits of writing an application with an RTOS can lead to savings around the board (code size, quality, reliability, and development time). While these other benefits provide the most compelling case for using an RTOS, they are far more challenging to demonstrate in a quantitative way, and are clearly documented in numerous industry-based case studies.

While there is some overhead associated with an RTOS, the typical arguments are largely unfounded when an RTOS is correctly implemented in a system. By measuring the true overhead of a preemptive RTOS in a typical application, we will demonstrate that the impact to code space, RAM, and CPU usage is minimal, and indeed acceptable for a wide range of CPU targets.

To illustrate just how little an RTOS impacts the size of an embedded software design we will look at a typical microcontroller project and analyze the various types of overhead associated with using a pre-emptive realtime kernel versus a similar non-preemptive event-based framework.

RTOS overhead can be broken into three distinct areas:

- Code space: The amount of code space eaten up by the kernel (static)
- Memory overhead: The RAM associated with running the kernel and application threads.
- Runtime overhead: The CPU cycles required for the kernel's functionality (primarily scheduling and thread switching)

While there are other notable reasons to include or avoid the use of an RTOS in certain applications (determinism, responsiveness, and interrupt latency among others), these are not considered in this discussion - as they are difficult to consider for the scope of our "canned" application.

8.2 Application description

For the purpose of this comparison, we first create an application using the standard preemptive [Mark3](#) kernel with 2 system threads running: A foreground thread and a background thread. This gives three total priority levels in the system - the interrupt level (high), and two application priority threads (medium and low), which is quite a common paradigm for microcontroller firmware designs. The foreground thread processes a variety of time-critical events at a fixed frequency, while the background thread processes lower priority, aperiodic events. When there are no background thread events to process, the processor enters its low-power mode until the next interrupt is acknowledged.

The contents of the threads themselves are unimportant for this comparison, but we can assume they perform a variety of realtime I/O functions. As a result, a number of device drivers are also implemented.

Code Space and Memory Overhead:

The application is compiled for an ATMega328p processor which contains 32kB of code space in flash, and 2kB of RAM, which is a lower-mid-range microcontroller in Atmel's 8-bit AVR line of microcontrollers. Using the AVR GCC compiler with -Os level optimizations, an executable is produced with the following code/RAM utilization:

```
Program: 27914 bytes
Data:    1313 bytes
```

An alternate version of this project is created using a custom "super-loop" kernel, which uses a single application thread and provides 2 levels of priority (interrupt and application). In this case, the event handler processes the different priority application events to completion from highest to lowest priority.

This approach leaves the application itself largely unchanged. Using the same optimization levels as the preemptive kernel, the code compiles as follows:

```
Program: 24886 bytes
Data:    750 bytes
```

At first glance, the difference in RAM utilization seems quite a lot higher for the preemptive mode version of the application, but the raw numbers don't tell the whole story.

The first issue is that the cooperative-mode total does not take into account the system stack - whereas these values are included in the totals for RTOS version of the project. As a result, some further analysis is required to determine how the stack sizes truly compare.

In cooperative mode, there is only one thread of execution - so considering that multiple event handlers are executed in turn, the stack requirements for cooperative mode is simply determined by those of the most stack-intensive event handler (ignoring stack use contributions due to interrupts).

In contrast, the preemptive kernel requires a separate stack for each active thread, and as a result the stack usage of the system is the sum of the stacks for all threads.

Since the application and idle events are the same for both preemptive and cooperative mode, we know that their (independent) stack requirements will be the same in both cases.

For cooperative mode, we see that the idle thread stack utilization is lower than that of the application thread, and so the application thread's determines the stack size requirement. Again, with the preemptive kernel the stack utilization is the sum of the stacks defined for both threads.

As a result, the difference in overhead between the two cases becomes the extra stack required for the idle thread - which in our case is (a somewhat generous) 128 bytes.

The numbers still don't add up completely, but looking into the linker output we see that the rest of the difference comes from the extra data structures used to manage the kernel in preemptive mode, and the kernel data itself.

Fixed kernel data costs:


```
--- 134 Bytes Kernel data
--- 26 Bytes Kernel Vtables
```

Application (Variable) data costs:

```
--- 24 Bytes Driver Vtables
--- 123 Bytes - statically-allocated kernel objects (semaphores, timers, etc.)
```

With this taken into account, the true memory cost of a 2-thread system ends up being around 428 bytes of RAM - which is about 20% of the total memory available on this particular microcontroller. Whether or not this is reasonable certainly depends on the application, but more importantly, it is not so unreasonable as to eliminate an RTOS-based solution from being considered. Also note that by using the “simulated idle” feature provided in [Mark3 R3](#) and onward, the idle thread (and its associated stack) can be eliminated altogether to reduce the cost in constrained devices.

The difference in code space overhead between the preemptive and cooperative mode solutions is less of an issue. Part of this reason is that both the preemptive and cooperative kernels are relatively small, and even an average target device (like the Atmega328 we've chosen) has plenty of room.

[Mark3](#) can be configured so that only features necessary for the application are included in the RTOS - you only pay for the parts of the system that you use. In this way, we can measure the overhead on a feature-by-feature basis, which is shown below for the kernel as configured for this application:

```
Kernel ..... 2563 Bytes
Synchronization Objects. 644 Bytes
Port ..... 974 Bytes
Features ..... 871 Bytes
```

The configuration tested in this comparison uses the thread/port module with timers, drivers, and semaphores, and mutexes, for a total kernel size of 5052 Bytes, with the rest of the code space occupied by the application.

As can be seen from the compiler's output, the difference in code space between the two versions of the application is 3028 bytes - or about 9% of the available code space on the selected processor. While nearly all of this comes from the added overhead of the kernel, the rest of the difference comes the changes to the application necessary to facilitate the different frameworks. This also demonstrates that the system-software code size in the cooperative case is about 2024 bytes.

8.3 Runtime Overhead

On the cooperative kernel, the overhead associated with running the thread is the time it takes the kernel to notice a pending event flag and launch the appropriate event handler, plus the timer interrupt execution time.

Similarly, on the preemptive kernel, the overhead is the time it takes to switch contexts to the application thread, plus the timer interrupt execution time.

The timer interrupt overhead is similar for both cases, so the overhead then becomes the difference between the following:

Preemptive mode:

- Posting the semaphore that wakes the high-priority thread
- Performing a context switch to the high-priority thread

Cooperative mode:

- Setting the event flag from the timer interrupt
- Acknowledging the event from the event loop

coop – 438 cycles preempt – 764 cycles

Using a cycle-accurate AVR simulator (fIAVR) running with a simulated speed of 16MHz, we find the end-to-end event sequence time to be 27us for the cooperative mode scheduler and 48us for the preemptive, and a raw difference of 20us.

With a fixed high-priority event frequency of 30Hz, we achieve a runtime overhead of 611us per second, or 0.06% of the total available CPU time. Now, obviously this value would expand at higher event frequencies and/or slower CPU frequencies, but for this typical application we find the difference in runtime overhead to be negligible for a preemptive system.

8.4 Analysis

For the selected test application and platform, including a preemptive RTOS is entirely reasonable, as the costs are low relative to a non-preemptive kernel solution. But these costs scale relative to the speed, memory and code space of the target processor. Because of these variables, there is no "magic bullet" environment suitable for every application, but [Mark3](#) attempts to provide a framework suitable for a wide range of targets.

On the one hand, if these tests had been performed on a higher-end microcontroller such as the ATmega1284p (containing 128kB of code space and 16kB of RAM), the overhead would be in the noise. For this type of resource-rich microcontroller, there would be no reason to avoid using the [Mark3](#) preemptive kernel.

Conversely, using a lower-end microcontroller like an ATmega88pa (which has only 8kB of code space and 1kB of RAM), the added overhead would likely be prohibitive for including a preemptive kernel. In this case, the cooperative-mode kernel would be a better choice.

As a rule of thumb, if one budgets 25% of a microcontroller's code space/RAM for system code, you should only require at minimum a microcontroller with 16k of code space and 2kB of RAM as a base platform for an RTOS. Unless there are serious constraints on the system that require much better latency or responsiveness than can be achieved with RTOS overhead, almost any modern platform is sufficient for hosting a kernel. In the event you find yourself with a microprocessor with external memory, there should be no reason to avoid using an RTOS at all.

Chapter 9

Mark3 Design Goals

9.1 Overview

9.1.1 Services Provided by an RTOS Kernel

At its lowest-levels, an operating system kernel is responsible for managing and scheduling resources within a system according to the application. In a typical thread-based RTOS, the resources involved is CPU time, and the kernel manages this by scheduling threads and timers. But capable RTOS kernels provide much more than just threading and timers.

In the following section, we discuss the [Mark3](#) kernel architecture, all of its features, and a thorough discussion of how the pieces all work together to make an awesome RTOS kernel.

9.1.2 Guiding Principles of Mark3

[Mark3](#) was designed with a number of over-arching principles, coming from years of experience designing, implementing, refining, and experimenting with RTOS kernels. Through that process I not only discovered what features I wanted in an RTOS, but how I wanted to build those features to look, work, and “feel”. With that understanding, I started with a clean slate and began designing a new RTOS. [Mark3](#) is the result of that process, and its design goals can be summarized in the following guiding principles.

9.1.3 Be feature competitive

To truly be taken seriously as more than just a toy or educational tool, an RTOS needs to have a certain feature suite. While [Mark3](#) isn't a clone of any existing RTOS, it should at least attempt parity with the most common software in its class.

Looking at its competitors, [Mark3](#) as a kernel supports most, if not all of the compelling features found in modern RTOS kernels, including dynamic threads, dynamic timers, efficient message passing, and multiple types of synchronization primitives.

9.1.4 No external dependencies, no non-language features

To maximize portability and promote adoption to new platforms, [Mark3](#) is written in a widely supported subset of C++ that lends itself to embedded applications. It avoids RTTI, exceptions, and libraries (C standard, STL, boost, etc.), with all fundamental data structures and types implemented completely for use by the kernel. As a result, the portable parts of [Mark3](#) should compile for any capable C++ toolchain.

9.1.5 Target the most popular hobbyist platforms available

Realistically, this means supporting the various Arduino-compatible target CPUs, including AVR and ARM Cortex-M series microcontrollers. As a result, the current default target for [Mark3](#) is the atmega328p, which has 32KB of flash and 2KB of RAM. All decisions regarding default features, code size, and performance need to take that target system into account.

[Mark3](#) integrates cleanly as a library into the Arduino IDE to support atmega328-based targets. Other AVR and Cortex-M targets can be supported using the port code provided in the source package.

9.1.6 Maximize determinism – but be pragmatic

Guaranteeing deterministic and predictable behavior is tough to do in an embedded system, and often comes with a heavy price tag in either RAM or code-space. With [Mark3](#), we strive to keep the core kernel APIs and features as lightweight as possible, while avoiding algorithms that don't scale to large numbers of threads. We also achieve minimal latency by keeping interrupts enabled (operating out of the critical section) wherever possible.

In [Mark3](#), the most important parts of the kernel are fixed-time, including thread scheduling and context switching. Operations that are not fixed time can be characterized as a function of their dependent data. For instances, the Mutex and Semaphore APIs operate in fixed time in the uncontested case, and execute in linear time for the contested case – where the speed of execution is dependent on the number of threads currently waiting on that object.

The caveat here is that while we want to minimize latency and time spent in critical sections, that has to be balanced against increases in code size, and uncontested-case performance.

9.1.7 Apply engineering principles – and that means discipline, measurement and verification

My previous RTOS, FunkOS, was designed to be very ad-hoc. The usage instructions were along the lines of “drag and drop the source files into your IDE and compile”. There was no regression/unit testing, no code size/speed profiling, and all documentation was done manually. It worked, but the process was a bit of a mess, and resulted in a lot of re-spins of the software, and a lot of time spent stepping through emulators to measure parameters.

We take a different approach in [Mark3](#). Here, we've designed not only the kernel-code, but the build system, unit tests, profiling code, documentation and reporting that supports the kernel. Each release is built and tested using automation in order to ensure quality and correctness, with supporting documentation containing all critical metrics. Only code that passes testing is submitted to the repos and public forums for distribution. These metrics can be traced from build-to-build to ensure that performance remains consistent from one drop to the next, and that no regressions are introduced by new/refactored code.

And while the kernel code can still be exported into an IDE directly, that takes place with the knowledge that the kernel code has already been rigorously tested and profiled. Exporting source in [Mark3](#) is also supported by scripting to ensure reliable, reproducible results without the possibility for human-error.

9.1.8 Use Virtualization For Verification

[Mark3](#) was designed to work with automated simulation tools as the primary means to validate changes to the kernel, due to the power and flexibility of automatic tests on virtual hardware. I was also intrigued by the thought of extending the virtual target to support functionality useful for a kernel, but not found on real hardware.

When the project was started, simavr was the tool of choice- however, its simulation was found to be incorrect compared to execution on a real MCU, and it did not provide the degree of extension that I desired for use with kernel development.

The flAVR AVR simulator was written to replace the dependency on that tool, and overcome those limitations. It also provides a GDB interface, as well as its own built-in debugger, profilers, and trace tools.

flAVR is hosted on sourceforge at <http://www.sourceforge.net/projects/flavr/> . In its basic configuration, it builds with minimal external dependencies.

- On linux, it requires only pthreads.
- On Windows, it requires pthreads and ws2_32, both satisfied via MinGW.
- Optional SDL builds for both targets (featuring graphics and simulated joystick input) can be built, and rely on libSDL.

Chapter 10

Mark3 Kernel Architecture

10.1 Overview

At a high level, the [Mark3](#) RTOS is organized into the following features, and layered as shown below:

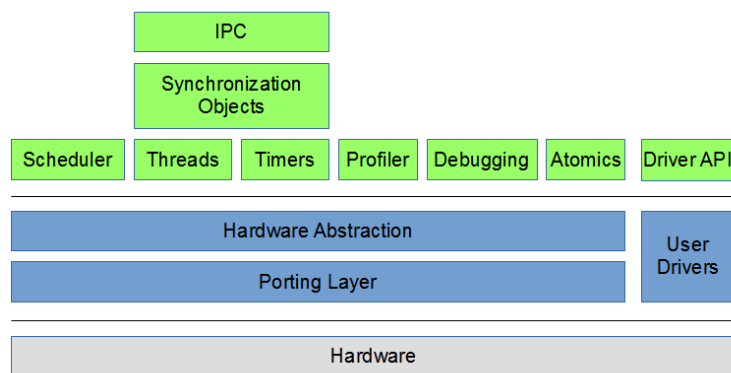


Figure 10.1 Overview

Everything in the “green” layer represents the [Mark3](#) public API and classes, beneath which lives all hardware abstraction and CPU-specific porting and driver code, which runs on a given target CPU.

The features and concepts introduced in this diagram can be described as follows:

Threads: The ability to multiplex the CPU between multiple tasks to give the perception that multiple programs are running simultaneously. Each thread runs in its own context with its own stack.

Scheduler: Algorithm which determines the thread that gets to run on the CPU at any given time. This algorithm takes into account the priorities (and other execution parameters) associated with the threads in the system.

IPC: Inter-process-communications. Message-passing and Mailbox interfaces used to communicate between threads synchronously or asynchronously.

Synchronization Objects: Ability to schedule thread execution relative to system conditions and events, allowing for sharing global data and resources safely and effectively.

Timers: High-resolution software timers that allow for actions to be triggered on a periodic or one-shot basis.

Profiler: Special timer used to measure the performance of arbitrary blocks of code.

Debugging: Realtime logging and trace functionality, facilitating simplified debugging of systems using the OS.

Atomics: Support for UN-interruptible arithmetic operations.

Driver API: Hardware abstraction interface allowing for device drivers to be written in a consistent, portable manner.

Hardware Abstraction Layer: Class interface definitions to represent threading, context-switching, and timers in a generic, abstracted manner.

Porting Layer: Class interface implementation to support threading, context-switching, and timers for a given CPU.

User Drivers: Code written by the user to implement device-specific peripheral drivers, built to make use of the [Mark3](#) driver API.

Each of these features will be described in more detail in the following sections of this chapter.

The concepts introduced in the above architecture are implemented in a variety of source modules, which are logically broken down into classes (or in some cases, groups of functions/macros). The relationship between objects in the [Mark3](#) kernel is shown below:

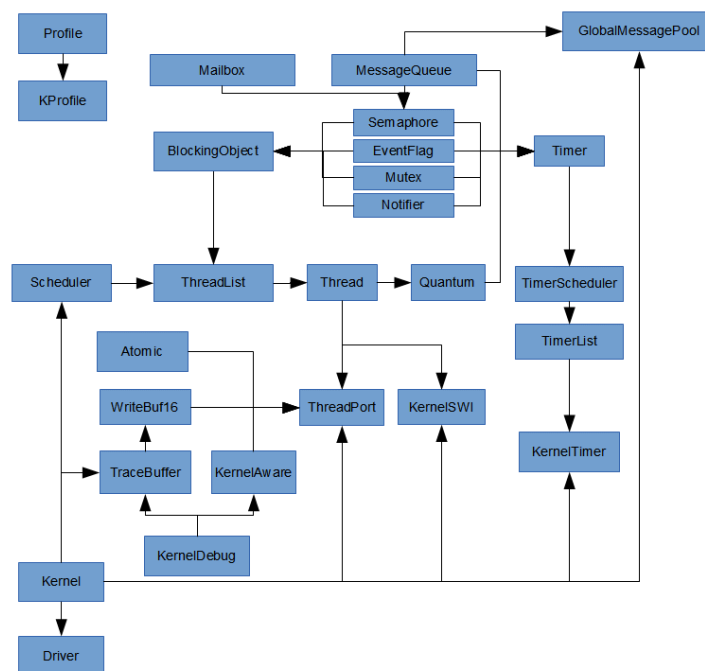


Figure 10.2 Overview

The objects shown in the preceding table can be grouped together by feature. In the table below, we group each feature by object, referencing the source module in which they can be found in the [Mark3](#) source tree.

Atomic Operations

Atomic - [atomic.cpp/.h](#)

Memory Allocators

AutoAlloc - [autoalloc.cpp/.h](#)

CoRoutines

CoScheduler - [cosched.cpp/.h](#)

CoList - [colist.cpp/.h](#)

Coroutine - [coroutine.cpp/.h](#)

Data Structures

LinkedList - [ll.cpp/.h](#)

PriorityMapL1 - [priomapl1.h](#)

PriorityMapL2 - [priomapl2.h](#)

Threads + Scheduling

Thread - [thread.cpp/.h](#)

Scheduler - [scheduler.cpp/.h](#)

Quantum - [quantum.cpp/.h](#)

ThreadPort - [threadport.cpp/.h](#) **

KernelSWI - [kernelswi.cpp/.h](#) **

ThreadList - [threadlist.cpp/.h](#)

ThreadListList - [threadlistlist.cpp/.h](#)

Profiling

ProfileTimer - [profile.cpp/.h](#)

Timers

Timer - [timer.h/timer.cpp](#)

TimerScheduler - [timerscheduler.h](#)

TimerList - [timerlist.h/cpp](#)

KernelTimer - [kerneltimer.cpp/.h](#) **

Synchronization

BlockingObject - [blocking.cpp/.h](#)

Semaphore - [ksemaphore.cpp/.h](#)

EventFlag - [eventflag.cpp/.h](#)

Mutex - [mutex.cpp/.h](#)

Notify - [notify.cpp/.h](#)

ConditionVariable - [condvar.cpp/.h](#)

ReaderWriterLock - [readerwriter.cpp/.h](#)

CriticalGuard - [criticalguard.h](#)

CriticalSection - [criticalsection.h](#)

LockGuard - [lockguard.h](#)

SchedGuard - [schedguard.h](#)

IPC/Message-passing

Mailbox - [mailbox.cpp/.h](#)

MessageQueue - [message.cpp/.h](#)

Debugging

Miscellaneous Macros - [kerneldebug.h](#)

Kernel

Kernel - [kernel.cpp/.h](#)

```
** implementation is platform-dependent, and located under the kernel's  
** /cpu/<arch>/<variant>/<toolchain> folder in the source tree
```

10.2 Threads and Scheduling

The classes involved in threading and scheduling in [Mark3](#) are highlighted in the following diagram, and are discussed in detail in this chapter:

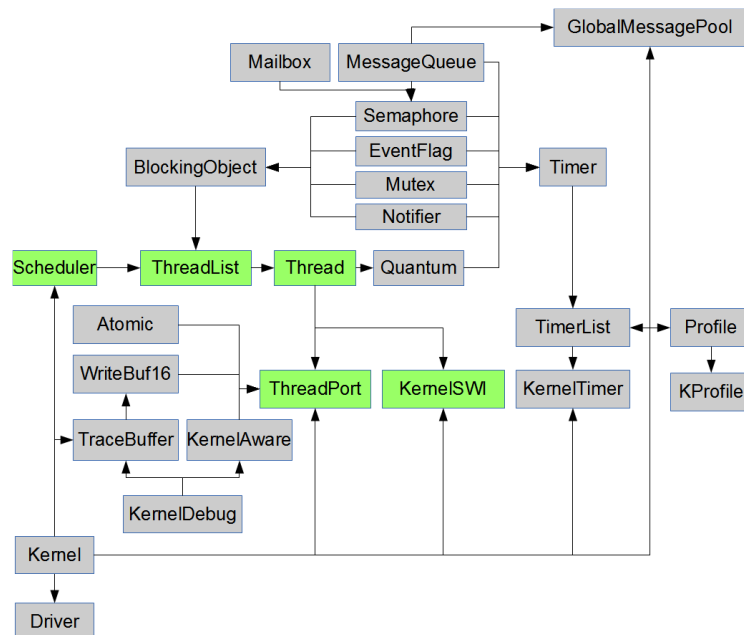


Figure 10.3 Threads and Scheduling

10.2.1 A Bit About Threads

Before we get started talking about the internals of the [Mark3](#) scheduler, it's necessary to go over some background material - starting with: what is a thread, anyway?

Let's look at a very basic CPU without any sort of special multi-threading hardware, and without interrupts. When the CPU is powered up, the program counter is loaded with some default location, at which point the processor core will start executing instructions sequentially - running forever and ever according to whatever has been loaded into program memory. This single instance of a simple program sequence is the only thing that runs on the processor, and the execution of the program can be predicted entirely by looking at the CPU's current register state, its program, and any affected system memory (the CPU's "context").

It's simple enough, and that's exactly the definition we have for a thread in an RTOS.

Each thread contains an instance of a CPU's register context, its own stack, and any other bookkeeping information necessary to define the minimum unique execution state of a system at runtime. It is the job of a RTOS to multiplex the execution of multiple threads on a single physical CPU, thereby creating the illusion that many programs are being executed simultaneously. In reality there can only ever be one thread truly executing at any given moment on a CPU core, so it's up to the scheduler to set and enforce rules about what thread gets to run when, for how long, and under what conditions. As mentioned earlier, any system without an RTOS executes as a single thread, so at least two threads are required for an RTOS to serve any useful purpose.

Note that all of this information is common to pretty well every RTOS in existence - the implementation details, including the scheduler rules, are all part of what differentiates one RTOS from another.

10.2.2 Thread States and ThreadLists

Since only one thread can run on a CPU at a time, the scheduler relies on thread information to make its decisions. [Mark3](#)'s scheduler relies on a variety of such information, including:

- The thread's current priority
- Round-Robin execution quanta
- Whether or not the thread is blocked on a synchronization object, such as a mutex or semaphore
- Whether or not the thread is currently suspended

The scheduler further uses this information to logically place each thread into 1 of 4 possible states:

```
- Ready - The thread is currently running
- Running - The thread is able to run
- Blocked - The thread cannot run until a system condition is met
- Stopped - The thread cannot run because its execution has been suspended
.
```

In order to determine a thread's state, threads are placed in "buckets" corresponding to these states. Ready and running threads exist in the scheduler's buckets, blocked threads exist in a bucket belonging to the object they're blocked on, and stopped threads exist in a separate bucket containing all stopped threads.

In reality, the various buckets are just doubly-linked lists of Thread objects - implemented in something called the ThreadList class. To facilitate this, the Thread class directly inherits from a LinkListNode class, which contains the node pointers required to implement a doubly-linked list. As a result, Threads may be effortlessly moved from one state to another using efficient linked-list operations built into the ThreadList class.

10.2.3 Blocking and Unblocking

While many developers new to the concept of an RTOS assume that all threads in a system are entirely separate from each other, the reality is that practical systems typically involve multiple threads working together, or at the very least sharing resources. In order to synchronize the execution of threads for that purpose, a number of synchronization primitives (blocking objects) are implemented to create specific sets of conditions under which threads can continue execution. The concept of "blocking" a thread until a specific condition is met is fundamental to understanding RTOS applications design, as well as any highly-multithreaded applications.

10.2.4 Blocking Objects

Blocking objects and primitives provided by [Mark3](#) include:

- Semaphores (binary and counting)
- Mutexes
- Event Flags
- Thread Notification Objects
- Thread Sleep
- Message Queues


```
Find the highest-priority "Ready" list that has at least one Thread.
Select the next thread to run as the first thread in that list
```

Since context switching is one of the most common and frequent operation performed by an RTOS, this needs to be as fast and deterministic as possible. While the logic is simple, a lot of care must be put into optimizing the scheduler to achieve those goals. In the section below we discuss the optimization approaches taken in [Mark3](#).

There are a number of ways to find the highest-priority thread. The naive approach would be to simply iterate through the scheduler's array of ThreadLists from highest to lowest, stopping when the first non-empty list is found, such as in the following block of code:

```
for (prio = num_prio - 1; prio >= 0; prio--)
{
    if (thread_list[prio].get_head() != nullptr)
    {
        break;
    }
}
```

While that would certainly work and be sufficient for a variety of systems, it's a non-deterministic approach (complexity $O(n)$) whose cost varies substantially based on how many priorities have to be evaluated. It's simple to read and understand, but it's non-optimal.

Fortunately, a functionally-equivalent and more deterministic approach can be implemented with a few tricks.

In addition to maintaining an array of ThreadLists, [Mark3](#) also maintains a bitmap (one bit per priority level) that indicates which thread lists have ready threads. This bitmap is maintained automatically by the ThreadList class, and is updated every time a thread is moved to/from the Scheduler's ready lists.

By inspecting this bitmap using a technique to count the leading zero bits in the bitmap, we determine which threadlist to choose in fixed time.

Now, to implement the leading-zeros check, this can once again be performed iteratively using bitshifts and compares (which isn't any more efficient than the raw list traversal), but it can also be evaluated using either a lookup table, or via a special CPU instruction to count the leading zeros in a value. In [Mark3](#), we use all approaches. In the event a target architecture or toolchain has intrinsic support for a count-leading-zeroes (PORT_CLZ) instruction, that implementation is used. Otherwise, a software-based implementation is provided – either using a lookup table, or a bitshift-and-compare algorithm.

(As a sidenote - this is actually a very common approach used in OS schedulers. It's actually part of the reason why modern ARM cores implement a dedicated count-leading-zeros [PORT_CLZ] instruction!)

For the lookup-table approach - a 4-bit lookup table can be used with an 8-bit priority-level bitmap would look something like this:

```
// Check the highest 4 priority levels, represented in the
// upper 4 bits in the bitmap
priority = priority_lookup_table[(priority_bitmap >> 4)];

// priority is non-zero if we found something there
if( priority )
{
    // Add 4 because we were looking at the higher levels
    priority += 4;
}
else
{
    // Nothing in the upper 4, look at the lowest 4 priority levels
    // represented by the lowest 4 bits in the bitmap
    priority = priority_lookup_table[priority_bitmap & 0x0F];
}
```

Deconstructing this algorithm, you can see that the priority lookup will have an $O(1)$ complexity - and is extremely low-cost.

This operation is thus fully deterministic and time bound - no matter how many threads are scheduled, the operation will always be time-bound to the most expensive of these two code paths. Even with only 8 priority levels, this is still much faster than iteratively checking the thread lists manually, compared with the previous example implementation.

Once the priority level has been found, selecting the next thread to run is trivial, consisting of something like this:

```
next_thread = thread_list[prio].get_head();
```

In the case of the `get_head()` calls, this evaluates to an inline-load of the "head" pointer in the particular thread list.

One important thing to take away from this analysis is that the scheduler is only responsible for selecting the next-to-run thread. In fact, these two operations are totally decoupled - no context switching is performed by the scheduler, and the scheduler isn't called from the context switch. The scheduler simply produces new "next thread" values that are consumed from within the context switch code.

10.3.1 Considerations for Round-Robin Scheduling

One thing that isn't considered directly from the scheduler algorithm is the problem of dealing with multiple threads within a single priority group; all of the algorithms that have been explored above simply look at the first Thread in each group.

[Mark3](#) addresses this issue indirectly, using an optimized software timer to manage round-robin scheduling, as follows.

In some instances where the scheduler is run by the kernel directly (typically as a result of calling `Thread::Yield()`), the kernel will perform an additional check after running the Scheduler to determine whether or there are multiple ready Threads in the priority of the next ready thread.

If there are multiple threads within that priority, the kernel starts a one-shot software timer which is programmed to expire at the next Thread's configured quantum. When this timer expires, a timer callback function executes to perform two simple operations:

"Pivot" the current Thread's priority list. Set a flag telling the kernel to trigger a Yield after exiting the main Timer↔ Scheduler processing loop

Pivoting the thread list basically moves the head of a circular-linked-list to its next value, which in our case ensures that a new thread will be chosen the next time the scheduler is run (the scheduler only looks at the head node of the priority lists). And by calling `Yield`, the system forces the scheduler to run, a new round-robin software timer to be installed (if necessary), and triggers a context switch SWI to load the newly-chosen thread. Note that if the thread attached to the round-robin timer is pre-empted, the kernel will take steps to abort and invalidate that round-robin software timer, installing a new one tied to the next thread to run if necessary.

Because the round-robin software timer is dynamically installed when there are multiple ready threads at the highest ready priority level, there is no CPU overhead with this feature unless that condition is met. The cost of round-robin scheduling is also fixed - no matter how many threads there are, and the cost is identical to any other one-shot software timer in the system.

10.3.2 Context Switching

There's really not much to say about the actual context switch operation at a high level. Context switches are triggered whenever it has been determined that a new thread needs to be swapped into the CPU core when the scheduler is run. [Mark3](#) implements also context switches as a call to a software interrupt - on AVR platforms, we typically use INT0 or INT2 for this (although any pin-change GPIO interrupt can be used), and on ARM we achieve this by triggering a PendSV exception.

However, regardless of the architecture, the contex-switch ISR will perform the following three operations:

Save the current Thread's context to the current Thread stack
Make the "next to run" thread the "currently running" thread
Restore the context of the next Thread from the Thread stack

The code to implement the context switch is entirely architecture-specific, so it won't be discussed in detail here. It's almost always gory inline-assembly which is used to load and store various CPU registers, and is highly-optimized for speed. We dive into an example implementation for the ARM Cortex-M0 microcontroller in a later section of this book.

10.3.3 Putting It All Together

In short, we can say that the [Mark3](#) scheduler works as follows:

- The scheduler is run whenever a `Thread::Yield()` is called by a user, as part of blocking calls, or whenever a new thread is started
- The [Mark3](#) scheduler is deterministic, selecting the next thread to run in fixed-time
- The scheduler only chooses the next thread to run, the context switch SWI consumes that information to get that thread running
- Where there are multiple ready threads in the highest populated priority level, a software timer is used to manage round-robin scheduling

While we've covered a lot of ground in this section, there's not a whole lot of code involved. However, the code that performs these operations is nuanced and subtle. If you're interested in seeing how this all works in practice, I suggest reading through the [Mark3](#) source code (which is heavily annotated), and stepping through the code with a simulator/emulator.

10.4 Timers

[Mark3](#) implements one-shot and periodic software-timers via the `Timer` class. The user configures the timer for duration, repetition, and action, at which point the timer can be activated. When an active timer expires, the kernel calls a user-specified callback function, and then reloads the timer in the case of periodic timers. The same timer objects exposed to the user are also used within the kernel to implement round-robin scheduling, and timeout-based APIs for semaphores, mutexes, events, and messages.

Timers are implemented using the following components in the [Mark3](#) Kernel:

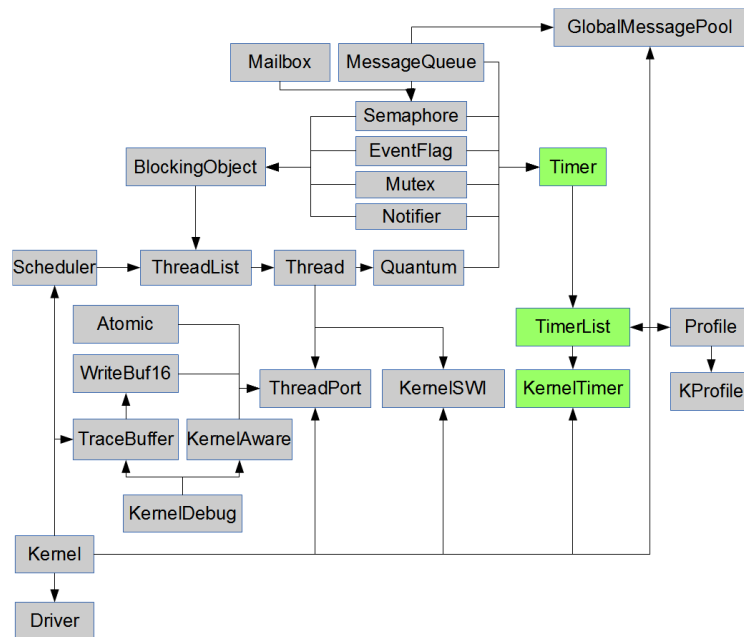


Figure 10.5 Timers

The `Timer` class provides the basic periodic and one-shot timer functionality used by application code, blocking objects, and IPC.

The `TimerList` class implements a doubly-linked list of `Timer` objects, and the logic required to implement a timer tick (tick-based kernel) or timer expiry (tickless kernel) event.

The `TimerScheduler` class contains a single `TimerList` object, implementing a single, system-wide list of `Timer` objects within the kernel. It also provides hooks for the hardware timer, such that when a timer tick or expiry event occurs, the `TimerList` expiry handler is run.

The `KernelTimer` class ([kerneltimer.cpp/.h](#)) implements the CPU specific hardware timer driver that is used by the kernel and the `TimerScheduler` to implement software timers.

While extremely simple to use, they provide one of the most powerful execution contexts in the system.

The software timers implemented in [Mark3](#) use interrupt-nesting within the kernel timer's interrupt handler. This context is be considered higher-priority than the highest priority user thread, but lower-priority than other interrupts in the system. As a result, this minimizes critical interrupt latency in the system, albeit at the expense of responsiveness of the user-threads.

For this reason, it's critical to ensure that all timer callback events are kept as short as possible to prevent adding thread-level latency. All heavy-lifting should be left to the threads, so the callback should only implement signalling via IPC or synchronization object.

The time spent in this interrupt context is also dependent on the number of active timers at any given time. However, [Mark3](#) also can be used to minimize the frequency of these interrupts wakeups, by using an optional "tolerance" parameter in the timer API calls. In this way, periodic tasks that have less rigorous real-time constraints can all be grouped together – executing as a group instead of one-after-another.

[Mark3](#) also contains two different timer implementations that can be configured at build-time, each with their own advantages.

10.4.1 Tick-based Timers

In a tick-based timing scheme, the kernel relies on a system-timer interrupt to fire at a relatively-high frequency, on which all kernel timer events are derived. On modern CPUs and microcontrollers, a 1kHz system tick is common, although quite often lower frequencies such as 60Hz, 100Hz, or 120Hz are used. The resolution of this timer also defines the maximum resolution of timer objects as a result. That is, if the timer frequency is 1kHz, a user cannot specify a timer resolution lower than 1ms.

The advantage of a tick-based timer is its sheer simplicity. It typically doesn't take much to set up a timer to trigger an interrupt at a fixed-interval, at which point, all system timer intervals are decremented by 1 count. When each system timer interval reaches zero, a callback is called for the event, and the events are either reset and restarted (repeated timers) or cleared (1-shot).

Unfortunately, that simplicity comes at a cost of increased interrupt count, which cause frequent CPU wakeups and utilization, and power consumption.

10.4.2 Tickless Timers

Note: Tickless timers are removed as of the R7 release. The below documentation is preserved for historical information only. The reason for removing tickless timers include the overhead associated with managing those timers (significantly more math and management is involved). In practice, there are few scenarios where purely tickless timers add benefit - beyond the most constrained devices. Also, it is entirely possible to disable software timers from the idle task when lower power/fewer interrupts are desired in most cases where a tickless timer would be of value.

In a tickless system, the kernel timer only runs when there are active timers pending expiry, and even then, the timer module only generates interrupts when a timer expires, or a timer reaches its maximum count value. Additionally, when there are no active timer objects, the timer can be completely disabled – saving even more cycles, power, and CPU wakeups. These factors make the tickless timer approach a highly-optimal solution, suitable for a wide array of low-power applications.

Also, since tickless timers do not rely on a fixed, periodic clock, they can potentially be higher resolution. The only limitation in timer resolution is the precision of the underlying hardware timer as configured. For example, if a 32kHz hardware timer is being used to drive the timer scheduler, the resolution of timer objects would be in the $\sim 33\mu\text{s}$ range.

The only downside of the tickless timer system is an added complexity to the timer code, requiring more code space, and slightly longer execution of the timer routines when the timer interrupt is executed.

10.4.3 Timer Processing Algorithm

Timer interrupts occur at either a fixed-frequency (tick-based), or at the next timer expiry interval (tickless), at which point the timer processing algorithm runs. While the timer count is reset by the timer-interrupt, it is still allowed to accumulate ticks while this algorithm is executed in order to ensure that timer-accuracy is kept in real-time. It is also important to note that round-robin scheduling changes are disabled during the execution of this algorithm to prevent race conditions, as the round-robin code also relies on timer objects.

All active timer objects are stored in a doubly-linked list within the timer-scheduler, and this list is processed in two passes by the algorithm which runs from the timer-interrupt (with interrupt nesting enabled). The first pass determines which timers have expired and the next timer interval, while the second pass deals with executing the timer callbacks themselves. Both phases are discussed in more detail below.

In the first pass, the active timers are decremented by either 1 tick (tick-based), or by the duration of the last elapsed timer interval (tickless). Timers that have zero (or less-than-zero) time remaining have a “callback” flag set, telling

Blocking a thread results in that thread becoming de-scheduled, placed in the blocking object's own private list of threads which are waiting on the object.

Unblocking a thread results in the reverse: The thread is moved back to its original location from the blocking list.

The only difference between a blocking object based on this class is the logic used to determine what constitutes a Block or Unblock condition.

For instance, a semaphore Pend operation may result in a call to the Block() method with the currently-executing thread in order to make that thread wait for a semaphore Post. That operation would then invoke the Unblock() method, removing the blocking thread from the semaphore's list, and back into the appropriate thread inside the scheduler.

Care must be taken when implementing blocking objects to ensure that critical sections are used judiciously, otherwise asynchronous events like timers and interrupts could result in non-deterministic and often catastrophic behavior.

[Mark3](#) implements a variety of blocking objects including semaphores, mutexes, event flags, and IPC mechanisms that all inherit from the basic Blocking-object class found in [blocking.h/cpp](#), ensuring consistency and a high degree of code-reuse between components.

10.6.1 Semaphores

Semaphores are used to synchronized execution of threads based on the availability (and quantity) of application-specific resources in the system. They are extremely useful for solving producer-consumer problems, and are the method-of-choice for creating efficient, low latency systems, where ISRs post semaphores that are handled from within the context of individual threads. Semaphores can also be posted (but not pended) from within the interrupt context.

10.6.2 Mutex

Mutexes (Mutual exclusion objects) are provided as a means of creating "protected sections" around a particular resource, allowing for access of these objects to be serialized. Only one thread can hold the mutex at a time

- other threads have to wait until the region is released by the owner thread before they can take their turn operating on the protected resource. Note that mutexes can only be owned by threads - they are not available to other contexts (i.e. interrupts). Calling the mutex APIs from an interrupt will cause catastrophic system failures.

Note that these objects are recursive in [Mark3](#) - that is, the owner thread can claim a mutex more than once. The caveat here is that a recursively-held mutex will not be released until a matching "release" call is made for each "claim" call.

Priority inheritance is provided with these objects as a means to avoid priority inversions. Whenever a thread at a priority than the mutex owner blocks on a mutex, the priority of the current thread is boosted to the highest-priority waiter to ensure that other tasks at intermediate priorities cannot artificially prevent progress from being made.

10.6.3 Event Flags

Event Flags are another synchronization object, conceptually similar to a semaphore.

Unlike a semaphore, however, the condition on which threads are unblocked is determined by a more complex set of rules. Each Event Flag object contains a 16-bit field, and threads block, waiting for combinations of bits within this field to become set.

A thread can wait on any pattern of bits from this field to be set, and any number of threads can wait on any number of different patterns. Threads can wait on a single bit, multiple bits, or bits from within a subset of bits within the field.

As a result, setting a single value in the flag can result in any number of threads becoming unblocked simultaneously. This mechanism is extremely powerful, allowing for all sorts of complex, yet efficient, thread synchronization schemes that can be created using a single shared object.

Note that Event Flags can be set from interrupts, but you cannot wait on an event flag from within an interrupt.

10.6.4 Notification Objects

Notification objects are the most lightweight of all blocking objects supplied by [Mark3](#).

using this blocking primitive, one or more threads wait for the notification object to be signalled by code elsewhere in the system (i.e. another thread or interrupt). Once the notification has been signalled, all threads currently blocked on the object become unblocked and moved into the ready list.

Signalling a notification object that has no actively-waiting threads has no effect.

10.7 Messages and Message Queues

10.7.1 Messages

Sending messages between threads is the key means of synchronizing access to data, and the primary mechanism to perform asynchronous data processing operations.

Sending a message consists of the following operations:

- Obtain a Message object from a source message pool
- Set the message data and event fields
- Send the message to the destination message queue

While receiving a message consists of the following steps:

- Wait for a messages in the destination message queue
- Process the message data
- Return the message back to the source message pool

These operations, and the various data objects involved are discussed in more detail in the following section.

10.7.2 Message Objects

Message objects are used to communicate arbitrary data between threads in a safe and synchronous way.

The message object consists of an event code field and a data field. The event code is used to provide context to the message object, while the data field (essentially a void * data pointer) is used to provide a payload of data corresponding to the particular event.

Access to these fields is marshalled by accessors - the transmitting thread uses the `SetData()` and `SetCode()` methods to seed the data, while the receiving thread uses the `GetData()` and `GetCode()` methods to retrieve it.

By providing the data as a void data pointer instead of a fixed-size message, we achieve an unprecedented measure of simplicity and flexibility. Data can be either statically or dynamically allocated, and sized appropriately for the event without having to format and reformat data by both sending and receiving threads. The choices here are left to the user - and the kernel doesn't get in the way of efficiency.

It is worth noting that you can send messages to message queues from within ISR context. This helps maintain consistency, since the same APIs can be used to provide event-driven programming facilities throughout the whole of the OS.

10.7.3 Message Queues

Message objects specify data with context, but do not specify where the messages will be sent. For this purpose we have a `MessageQueue` object. Sending an object to a message queue involves calling the `MessageQueue::Send()` method, passing in a pointer to the Message object as an argument.

When a message is sent to the queue, the first thread blocked on the queue (as a result of calling the `MessageQueue::Receive()` method) will wake up, with a pointer to the Message object returned.

It's worth noting that multiple threads can block on the same message queue, providing a means for multiple threads to share work in parallel.

10.7.4 Mailboxes

Another form of IPC is provided by [Mark3](#), in the form of Mailboxes and Envelopes. Mailboxes are similar to message queues in that they provide a synchronized interface by which data can be transmitted between threads.

Where Message Queues rely on linked lists of lightweight message objects (containing only message code and a void* data-pointer), which are inherently abstract, Mailboxes use a dedicated blob of memory, which is carved up into fixed-size chunks called Envelopes (defined by the user), which are sent and received. Unlike message queues, mailbox data is copied to and from the mailboxes dedicated pool.

Mailboxes also differ in that they provide not only a blocking "receive" call, but also a blocking "send" call, providing the opportunity for threads to block on "mailbox full" as well as "mailbox empty" conditions.

All send/receive APIs support an optional timeout parameter if the `KERNEL_USE_TIMEOUTS` option has been configured in [mark3cfg.h](#)

10.7.5 Atomic Operations

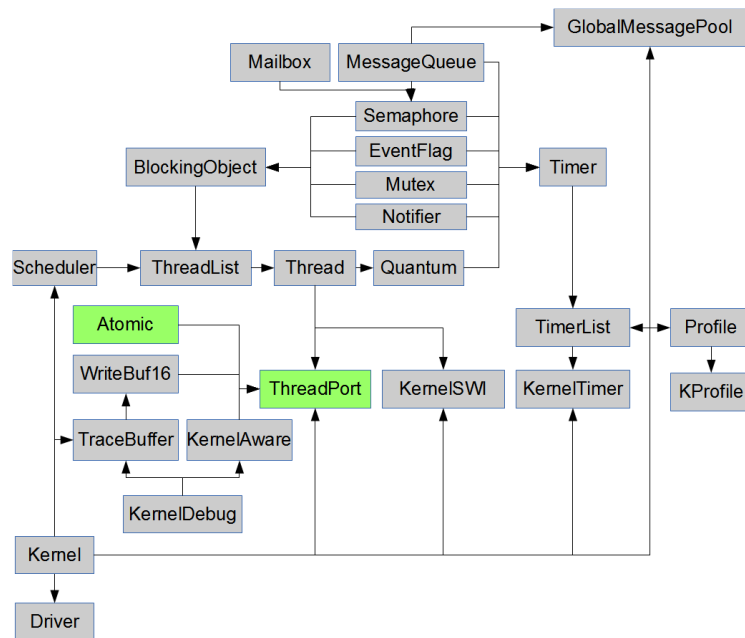


Figure 10.7 Atomic operations

This utility class provides primitives for atomic operations - that is, operations that are guaranteed to execute uninterrupted. Basic atomic primitives provided here include Set/Add/Delete for 8, 16, and 32-bit integer types, as well as an atomic test-and-set.

10.7.6 Drivers

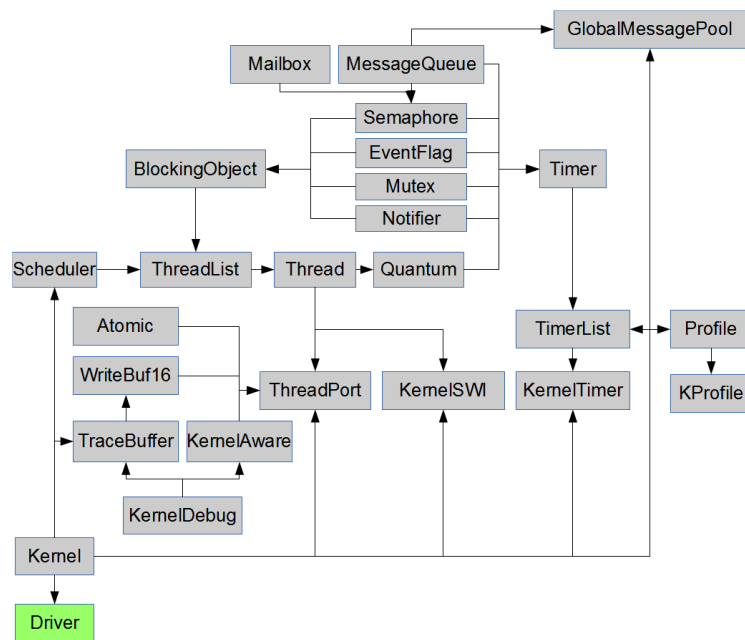


Figure 10.8 Drivers

This is the basis of the driver framework. In the context of [Mark3](#), drivers don't necessarily have to be based on physical hardware peripherals. They can be used to represent algorithms (such as random number generators), files, or protocol stacks. Unlike FunkOS, where driver IO is protected automatically by a mutex, we do not use this kind of protection - we leave it up to the driver implementor to do what's right in its own context. This also frees up the driver to implement all sorts of other neat stuff, like sending messages to threads associated with the driver. Drivers are implemented as character devices, with the standard array of posix-style accessor methods for reading, writing, and general driver control.

A global driver list is provided as a convenient and minimal "filesystem" structure, in which devices can be accessed by name.

Driver Design

A device driver needs to be able to perform the following operations:

- Initialize a peripheral
- Start/stop a peripheral
- Handle I/O control operations
- Perform various read/write operations

At the end of the day, that's pretty much all a device driver has to do, and all of the functionality that needs to be presented to the developer.

We abstract all device drivers using a base-class which implements the following methods:

- Start/Open
- Stop/Close
- Control
- Read
- Write

A basic driver framework and API can thus be implemented in five function calls - that's it! You could even reduce that further by handling the initialize, start, and stop operations inside the "control" operation.

Driver API

In C++, we can implement this as a class to abstract these event handlers, with virtual void functions in the base class overridden by the inherited objects.

To add and remove device drivers from the global table, we use the following methods:

```
void DriverList::Add( Driver *pclDriver_ );
void DriverList::Remove( Driver *pclDriver_ );
```

`DriverList::Add()/Remove()` takes a single argument - the pointer to the object to operate on.

Once a driver has been added to the table, drivers are opened by NAME using `DriverList::FindByName("/dev/name")`. This function returns a pointer to the specified driver if successful, or to a built in /dev/null device if the path name is invalid. After a driver is open, that pointer is used for all other driver access functions.

This abstraction is incredibly useful - any peripheral or service can be accessed through a consistent set of APIs, that make it easy to substitute implementations from one platform to another. Portability is ensured, the overhead is negligible, and it emphasizes the reuse of both driver and application code as separate entities.

Consider a system with drivers for I2C, SPI, and UART peripherals - under our driver framework, an application can initialize these peripherals and write a greeting to each using the same simple API functions for all drivers:

```
pclI2C = DriverList::FindByName("/dev/i2c");
pclUART = DriverList::FindByName("/dev/tty0");
pclSPI = DriverList::FindByName("/dev/spi");

pclI2C->Write(12, "Hello World!");
pclUART->Write(12, "Hello World!");
pclSPI->Write(12, "Hello World!");
```


10.8 Kernel Proper and Porting




Figure 10.9 Kernel Proper and Porting

The Kernel class is a static class with methods to handle the initialization and startup of the RTOS, manage errors, and provide user-hooks for fatal error handling (functions called when `Kernel::Panic()` conditions are encountered), or when the Idle function is run.

Internally, `Kernel::Init()` calls the initialization routines for various kernel objects, providing a single interface by which all RTOS-related system initialization takes place.

`Kernel::Start()` is called to begin running OS functionality, and does not return. Control of the CPU is handed over to the scheduler, and the highest-priority ready thread begins execution in the RTOS environment.

Hardware Abstraction Layer

Almost all of the [Mark3](#) kernel (and middleware) is completely platform independent, and should compile cleanly on any platform with a modern C++ compiler. However, there are a few areas within [Mark3](#) that can only be implemented by touching hardware directly.

These interfaces generally cover four features:

- Thread initialization and context-switching logic
- Software interrupt control (used to generate context switches)
- Hardware timer control (support for time-based functionality, such as `Sleep()`)
- Code-execution profiling timer (not necessary to port if code-profiling is not compiled into the kernel)

The hardware abstraction layer in [Mark3](#) provides a consistent interface for each of these four features. [Mark3](#) is ported to new target architectures by providing an implementation for all of the interfaces declared in the abstraction layer. In the following section, we will explore how this was used to port the kernel to ARM Cortex-M0.

Real-world Porting Example – Cortex M0

This section serves as a real-world example of how [Mark3](#) can be ported to new architectures, how the [Mark3](#) abstraction layer works, and as a practical reference for using the RTOS support functionality baked in modern ARM Cortex-M series microcontrollers. Most of this documentation here is taken directly from the source code found in the `/kernel/cpu/cm0/ports` directory, with additional annotations to explain the port in more detail. Note that a familiarity with Cortex-M series parts will go a long way to understanding the subject matter presented, especially a basic understanding of the ARM CPU registers, exception models, and OS support features (PendSV, SysTick and SVC). If you're unfamiliar with ARM architecture, pay attention to the comments more than the source itself to illustrate the concepts.

Porting [Mark3](#) to a new architecture consists of a few basic pieces; for developers familiar with the target architecture and the porting process, it's not a tremendously onerous endeavour to get [Mark3](#) up-and-running somewhere new. For starters, all non-portable components are completely isolated in the source-tree under:

```
/embedded/kernel/CPU/VARIANT/TOOLCHAIN/
```

where CPU is the architecture, VARIANT is the vendor/part, and TOOLCHAIN is the compiler tool suite used to build the code.

From within the specific port folder, a developer needs only implement a few classes and headers that define the port-specific behavior of [Mark3](#):

- KernelSWI ([kernelswi.cpp/kernelswi.h](#)) - Provides a maskable software-triggered interrupt used to perform context switching.
- KernelTimer ([kerneltimer.cpp/kerneltimer.h](#)) - Provides either a fixed-frequency or programmable-interval timer, which triggers an interrupt on expiry. This is used for implementing round-robin scheduling, thread-sleeps, and generic software timers.
- Profiler ([kprofile.cpp/kprofile.h](#)) - Contains code for runtime code-profiling. This is optional and may be stubbed out if left unimplemented (we won't cover profiling timers here).
- ThreadPort ([threadport.cpp/threadport.h](#)) - The meat-and-potatoes of the port code lives here. This class contains architecture/part-specific code used to initialize threads, implement critical-sections, perform context-switching, and start the kernel. Most of the time spent in this article focuses on the code found here.

Summarizing the above, these modules provide the following list of functionality:

```
- Thread stack initialization
- Kernel startup and first thread entry
- Context switch and SWI
- Kernel timers
- Critical Sections
.
```

The implementation of each of these pieces will be analyzed in detail in the sections that follow.

Thread Stack Initialization

Before a thread can be used, its stack must first be initialized to its default state. This default state ensures that when the thread is scheduled for the first time and its context restored, that it will cause the CPU to jump to the user's specified entry-point function.

All of the platform independent thread setup is handled by the generic kernel code. However, since every CPU architecture has its own register set, and stacks different information as part of an interrupt/exception, we have to implement this thread setup code for each platform we want the kernel to support (Combination of Architecture + Variant + Toolchain).

In the ARM Cortex-M0 architecture, the stack frame consists of the following information:

a) Exception Stack Frame

Contains the 8 registers which the ARM Cortex-M0 CPU automatically pushes to the stack when entering an exception. The following registers are included (in stack'd order):

```
[ XPSR ] <-- Highest address in context
[ PC   ]
[ LR   ]
[ R12  ]
[ R3   ]
[ R2   ]
[ R1   ]
[ R0   ]
```

XPSR – This is the CPU's status register. We need to set this to 0x01000000 (the "T" bit), which indicates that the CPU is executing in "thumb" mode. Note that ARMv6m and ARMv7m processors only run thumb2 instructions, so an exception is liable to occur if this bit ever gets cleared.

PC – Program Counter. This should be set with the initial entry point (function pointer) for that the user wishes to start executing with this thread.

LR - The base link register. Normally, this register contains the return address of the calling function, which is where the CPU jumps when a function returns. However, our threads generally don't return (and if they do, they're placed into the stop state). As a result we can leave this as 0.

The other registers in the stack frame are generic working registers, and have no special meaning, with the exception that R0 will hold the user's argument value passed into the entrypoint.

b) Complimentary CPU Register Context

```
[ R11  ]
...
[ R4   ] <-- Lowest address in context
```

These are the other general-purpose CPU registers that need to be backed up/ restored on a context switch, but aren't stacked by default on a Cortex-M0 exception. If there were any additional hardware registers to back up, then we'd also have to include them in this part of the context as well.

As a result, these registers all need to be manually pushed to the stack on stack creation, and will need to be explicitly pushed and pop as part of a normal context switch.

With this default exception state in mind, the following code is used to initialize a thread's stack for a Cortex-M0.

```

void ThreadPort::InitStack(Thread *pclThread_)
{
    K_ULONG *pulStack;
    K_ULONG *pulTemp;
    K_ULONG ulAddr;
    K_USHORT i;

    // Get the entrypoint for the thread
    ulAddr = (K_ULONG)(pclThread_>m_pfEntryPoint);

    // Get the top-of-stack pointer for the thread
    pulStack = (K_ULONG*)pclThread_>m_pwStackTop;

    // Initialize the stack to all FF's to aid in stack depth checking
    pulTemp = (K_ULONG*)pclThread_>m_pwStack;
    for (i = 0; i < pclThread_>m_usStackSize / sizeof(K_ULONG); i++)
    {
        pulTemp[i] = 0xFFFFFFFF;
    }

    PORT_PUSH_TO_STACK(pulStack, 0); // Apply one word of padding

    //-- Simulated Exception Stack Frame --
    PORT_PUSH_TO_STACK(pulStack, 0x01000000); // XSPR;set "T" bit for thumb-mode
    PORT_PUSH_TO_STACK(pulStack, ulAddr); // PC
    PORT_PUSH_TO_STACK(pulStack, 0); // LR
    PORT_PUSH_TO_STACK(pulStack, 0x12);
    PORT_PUSH_TO_STACK(pulStack, 0x3);
    PORT_PUSH_TO_STACK(pulStack, 0x2);
    PORT_PUSH_TO_STACK(pulStack, 0x1);
    PORT_PUSH_TO_STACK(pulStack, (K_ULONG)pclThread_>m_pvArg); // R0 = argument

    //-- Simulated Manually-Stacked Registers --
    PORT_PUSH_TO_STACK(pulStack, 0x11);
    PORT_PUSH_TO_STACK(pulStack, 0x10);
    PORT_PUSH_TO_STACK(pulStack, 0x09);
    PORT_PUSH_TO_STACK(pulStack, 0x08);
    PORT_PUSH_TO_STACK(pulStack, 0x07);
    PORT_PUSH_TO_STACK(pulStack, 0x06);
    PORT_PUSH_TO_STACK(pulStack, 0x05);
    PORT_PUSH_TO_STACK(pulStack, 0x04);
    pulStack++;

    pclThread_>m_pwStackTop = pulStack;
}

```

Kernel Startup

The same general process applies to starting the kernel on an ARM Cortex-M0 as on other platforms. Here, we initialize and start the platform specific timer and software-interrupt modules, find the first thread to run, and then jump to that first thread.

Now, to perform that last step, we have two options:

1) Simulate a return from an exception manually to start the first thread, or.. 2) Use a software interrupt to trigger the first "Context Restore/Return from Interrupt"

For 1), we basically have to restore the whole stack manually, not relying on the CPU to do any of this for us. That's certainly doable, but not all Cortex parts support this (other members of the family support privileged modes, etc.). That, and the code required to do this is generally more complex due to all of the exception-state simulation. So, we will opt for the second option instead.

To implement a software to start our first thread, we will use the SVC instruction to generate an exception. From that exception, we can then restore the context from our first thread, set the CPU up to use the right "process" stack, and return-from-exception back to our first thread. We'll explore the code for that later.

But, before we can call the SVC exception, we're going to do a couple of things.

First, we're going to reset the default MSP stack pointer to its original top-of-stack value. The rationale here is that we no longer care about the data on the MSP stack, since calling the SVC instruction triggers a chain of events from which we never return. The MSP is also used by all exception-handling, so regaining a few words of stack here can be useful. We'll also enable all maskable exceptions at this point, since this code results in the kernel being started with the CPU executing the RTOS threads, at which point a user would expect interrupts to be enabled.

Note, the default stack pointer location is stored at address 0x00000000 on all ARM Cortex M0 parts. That explains the code below...

```

void ThreadPort_StartFirstThread( void )
{
    asm(
        " ldr r1, [r0] \n" // Reset the MSP to the default base address
        " msr msp, r1 \n"
        " cpsie i \n"      // Enable interrupts
        " svc 0 \n"        // Jump to SVC Call
    );
}

```

First Thread Entry

This handler has the job of taking the first thread object's stack, and restoring the default state data in a way that ensures that the thread starts executing when returning from the call.

We also keep in mind that there's an 8-byte offset from the beginning of the thread object to the location of the thread stack pointer. This offset is a result of the thread object inheriting from the linked-list node class, which has 8-bytes of data. This is stored first in the object, before the first element of the class, which is the "stack top" pointer.

The following assembly code shows how the SVC call is implemented in [Mark3](#) for the purpose of starting the first thread.

```

get_thread_stack:
    ; Get the stack pointer for the current thread
    ldr r0, g_pstCurrent
    ldr r1, [r0]
    add r1, #8
    ldr r2, [r1]          ; r2 contains the current stack-top

load_manually_placed_context_r11_r8:
    ; Handle the bottom 32-bytes of the stack frame
    ; Start with r11-r8, because only r0-r7 can be used
    ; with ldmia on CM0.
    add r2, #16
    ldmia r2!, {r4-r7}
    mov r11, r7
    mov r10, r6
    mov r9, r5
    mov r8, r4

set_psp:
    ; Since r2 is coincidentally back to where the stack pointer should be,
    ; Set the program stack pointer such that returning from the exception handler
    msr psp, r2

load_manually_placed_context_r7_r4:
    ; Get back to the bottom of the manually stacked registers and pop.
    sub r2, #32
    ldmia r2!, {r4-r7} ; Register r4-r11 are restored.

set_thread_and_privilege_modes:
    ; Also modify the control register to force use of thread mode as well
    ; For CM3 forward-compatibility, also set user mode.
    mrs r0, control
    mov r1, #0x03
    orr r0, r1
    control, r0

set_lr:
    ; Set up the link register such that on return, the code operates
    ; in thread mode using the PSP. To do this, we or 0x0D to the value stored
    ; in the lr by the exception hardware EXC_RETURN. Alternately, we could
    ; just force lr to be 0xFFFFFFFED (we know that's what we want from the
    ; hardware, anyway)
    mov r0, #0x0D
    mov r1, lr
    orr r0, r1

exit_exception:
    ; Return from the exception handler.
    ; The CPU will automatically unstack R0-R3, R12, PC, LR, and xPSR
    ; for us. If all goes well, our thread will start execution at the
    ; entrypoint, with the us-specified argument.
    bx r0

```

On ARM Cortex parts, there's dedicated hardware that's used primarily to support RTOS (or RTOS-like) functionality. This functionality includes the SysTick timer, and the PendSV Exception. SysTick is used for a tick-based kernel timer, while the PendSV exception is used for performing context switches. In reality, it's a "special SVC" call that's designed to be lower-overhead, in that it isn't mux'd with a bunch of other system or application functionality.

So how do we go about actually implementing a context switch here? There are a lot of different parts involved, but it essentially comes down to 3 steps:

1) Saving the context.

Thread's top-of-stack value is stored, all registers are stacked. We're good to go!

2) Swap threads

We swap the Scheduler's "next" thread with the "current" thread.

3) Restore Context

This is more or less identical to what we did when restoring the first context. Some operations may be optimized for data already stored in registers.

The code used to implement these steps on Cortex-M0 is presented below:

```
void PendSV_Handler(void)
{
    ASM(
        // Thread_SaveContext()
        " ldr r1, CURR_ \n"
        " ldr r1, [r1] \n "
        " mov r3, r1 \n "
        " add r3, #8 \n "

        // Grab the psp and adjust it by 32 based on extra registers we're going
        // to be manually stacking.
        " mrs r2, psp \n "
        " sub r2, #32 \n "

        // While we're here, store the new top-of-stack value
        " str r2, [r3] \n "

        // And, while r2 is at the bottom of the stack frame, stack r7-r4
        " stmia r2!, {r4-r7} \n "

        // Stack r11-r8
        " mov r7, r11 \n "
        " mov r6, r10 \n "
        " mov r5, r9 \n "
        " mov r4, r8 \n "
        " stmia r2!, {r4-r7} \n "

        // Equivalent of Thread_Swap() - performs g_pstCurrent = g_pstNext
        " ldr r1, CURR_ \n"
        " ldr r0, NEXT_ \n"
        " ldr r0, [r0] \n"
        " str r0, [r1] \n"

        // Thread_RestoreContext()
        // Get the pointer to the next thread's stack
        " add r0, #8 \n "
        " ldr r2, [r0] \n "

        // Stack pointer is in r2, start loading registers from
        // the "manually-stacked" set
        // Start with r11-r8, since these can't be accessed directly.
        " add r2, #16 \n "
        " ldmba r2!, {r4-r7} \n "
        " mov r11, r7 \n "
        " mov r10, r6 \n "
        " mov r9, r5 \n "
        " mov r8, r4 \n "
```

```

// After subbing R2 #16 manually, and #16 through ldmia, our PSP is where it
// needs to be when we return from the exception handler
" msr psp, r2 \n "

// Pop manually-stacked R4-R7
" sub r2, #32 \n "
" ldmia r2!, {r4-r7} \n "

// lr contains the proper EXC_RETURN value
// we're done with the exception, so return back to newly-chosen thread
" bx lr \n "
" nop \n "

// Must be 4-byte aligned.
" NEXT_: .word g_pstNext \n"
" CURR_: .word g_pstCurrent \n"
);
}

```

Kernel Timers

ARM Cortex-M series microcontrollers each contain a SysTick timer, which was designed to facilitate a fixed-interval RTOS timer-tick. This timer is a precise 24-bit down-count timer, run at the main CPU clock frequency, that can be programmed to trigger an exception when the timer expires. The handler for this exception can thus be used to drive software timers throughout the system on a fixed interval.

Unfortunately, this hardware is extremely simple, and does not offer the flexibility of other timer hardware commonly implemented by MCU vendors - specifically a suitable timer prescaler that can be used to generate efficient, long-counting intervals. As a result, while the "generic" port of [Mark3](#) for Cortex-M0 leverages the common SysTick timer interface, it only supports the tick-based version of the kernel's timer (note that specific Cortex-M0 ports such as the Atmel SAMD20 do have tickless timers).

Setting up a tick-based `KernelTimer` class to use the SysTick timer is, however, extremely easy, as is illustrated below:

```

void KernelTimer::Start(void)
{
    SysTick_Config(PORT_SYSTEM_FREQ / 1000); // 1kHz fixed clock...
    NVIC_EnableIRQ(SysTick_IRQn);
}

In this instance, the call to SysTick_Config() generates a 1kHz system-tick
signal, and the NVIC_EnableIRQ() call ensures that a SysTick exception is
generated for each tick. All other functions in the Cortex version of the
KernelTimer class are essentially stubbed out (see the source for more details).

```

Note that the functions used in this call are part of the ARM Cortex Microcontroller Software Interface Standard (CMSIS), and are supplied by all parts vendors selling Cortex hardware. This greatly simplifies the design of our port-code, since we can be reasonably assured that these APIs will work the same on all devices.

The handler code called when a SysTick exception occurs is basically the same as on other platforms (such as [AVR](#)), except that we explicitly clear the "exception pending" bit before returning. This is implemented in the following code:

```

@code{.cpp}
void SysTick_Handler(void)
{
    #if KERNEL_USE_TIMERS
        TimerScheduler::Process();
    #endif
    #if KERNEL_USE_QUANTUM
        Quantum::UpdateTimer();
    #endif

    // Clear the systick interrupt pending bit.
    SCB->ICSR |= SCB_ICSR_PENDSTCLR_Msk;
}

```

Critical Sections

A "critical section" is a block of code whose execution cannot be interrupted by means of context switches or an interrupt. In a traditional single-core operating system, it is typically implemented as a block of code where the interrupts are disabled - this is also the approach taken by [Mark3](#). Given that every CPU has its own means of disabling/enabling interrupts, the implementation of the critical section APIs is also non-portable.

In the Cortex-M0 port, we implement the two critical section APIs (`CriticalSection::Enter()` and `CriticalSection::Exit()`) as function-like macros containing inline assembly. All uses of these calls are called in pairs within a function and must take place at the same level-of-scope. Also, as nesting may occur (critical section within a critical section), this must be taken into account in the code.

In general, `CriticalSection::Enter()` performs the following tasks:

- Cache the current interrupt-enabled state within a local variable in the thread's state
- Disable interrupts
- .

Conversely, `CriticalSection::Exit()` performs the following tasks:

- Read the original interrupt-enabled state from the cached value
- Restore interrupts to the original value
- .

On Cortex-M series microcontrollers, the PRIMASK special register contains a single status bit which can be used to enable/disable all maskable interrupts at once. This register can be read directly to examine or modify its state. For convenience, ARMv6m provides two instructions to enable/disable interrupts

- `cpsid` (disable interrupts) and `cpsie` (enable interrupts). [Mark3](#) Implements these steps according to the following code:

```
//-----
#define CriticalSection::Enter() \
{ \
    K_ULONG __ulRegState; \
    asm ( \
        " mrs r0, PRIMASK \n" \
        " mov %[STATUS], r0 \n" \
        " cpsid i \n " \
        : [STATUS] "=r" (__ulRegState) \
        ); \
}

//-----
#define CriticalSection::Exit() \
    asm ( \
        " mov r0, %[STATUS] \n" \
        " msr primask, r0 \n" \
        : \
        : [STATUS] "r" (__ulRegState) \
        ); \
}
```

Summary

In this section we have investigated how the main non-portable areas of the [Mark3](#) RTOS are implemented on a Cortex-M0 microcontroller. [Mark3](#) leverages all of the hardware blocks designed to enable RTOS functionality on ARM Cortex-M series microcontrollers: the SVC call provides the mechanism by which we start the kernel, the PendSV exception provides the necessary software interrupt, and the SysTick timer provides an RTOS tick. As a result, [Mark3](#) is a perfect fit for these devices - and as a result of this approach, the same RTOS port code should work with little to no modification on all ARM Cortex-M parts.

We have discussed what functionality in the RTOS is not portable, and what interfaces must be implemented in order to complete a fully-functional port. The five specific areas which are non-portable (stack initialization, kernel startup/entry, kernel timers, context switching, and critical sections) have been discussed in detail, with the platform-specific source provided as a practical reference to ARM-specific OS features, as well as [Mark3](#)'s porting infrastructure. From this example (and the accompanying source), it should be possible for an experienced developers to create a port [Mark3](#) to other microcontroller targets.

Chapter 11

C-language bindings

[Mark3](#) now includes an optional additional library with C language bindings for all core kernel APIs, known as Mark3C. This library allows applications to be written in C, while still enjoying all of the benefits of the clean, modular design of the core RTOS kernel.

The C-language Mark3C APIs map directly to their [Mark3](#) counterparts using a simple set of conventions, documented below. As a result, explicit API documentation for Mark3C is not necessary, as the functions map 1-1 to their C++ counterparts.

11.1 API Conventions

1) Static Methods:

<code><ClassName>::<MethodName>()</code>	Becomes	<code><ClassName>_<MethodName>()</code>
i.e. <code>Kernel::Start()</code>	Becomes	<code>Kernel_Start()</code>

2) Kernel Object Methods:

In short, any class instance is represented using an object handle, and is always passed into the relevant APIs as the first argument. Further, any method that returns a pointer to an object in the C++ implementation now returns a handle to that object.

<code><Object>.<MethodName>(<args>)</code>	Becomes	<code><ClassName>_<Methodname>(<ObjectHandle>, <args>)</code>
i.e. <code>clAppThread.Start()</code>	Becomes	<code>Thread_Start(hAppThread)</code>

3) Overloaded Methods:

a) Methods overloaded with a Timeout parameter:

<code><Object>.<MethodName>(<args>)</code>	Becomes	<code><ClassName>_Timed<MethodName>(<ObjectHandle>, <args>)</code>
i.e. <code>clSemaphore.Wait(1000)</code>	Becomes	<code>Semaphore_Wait(hSemaphore, 1000)</code>

b) Methods overloaded based on number of arguments:

<Object>.<MethodName>()	Becomes	<ClassName>_<MethodName>(<ObjectHandle>)
<Object>.<MethodName>(<arg1>)	Becomes	<ClassName>_<MethodName>1(<ObjectHandle>, <arg1>)
<Object>.<MethodName>(<arg1>, <arg2>)	Becomes	<ClassName>_<MethodName>2(<ObjectHandle>, <arg1>, <arg2>)
<ClassName>::<MethodName>()	Becomes	<ClassName>_<MethodName>(<ObjectHandle>)
<ClassName>::<MethodName>(<arg1>)	Becomes	<ClassName>_<MethodName>1(<ObjectHandle>, <arg1>)
<ClassName>::<MethodName>(<arg1>, <arg2>)	Becomes	<ClassName>_<MethodName>2(<ObjectHandle>, <arg1>, <arg2>)

c) Methods overloaded base on parameter types:

<Object>.<MethodName>(<arg type_a>)	Becomes	<ClassName>_<MethodName><type_a>(<ObjectHandle>, <arg type_a>)
<Object>.<MethodName>(<arg type_b>)	Becomes	<ClassName>_<MethodName><type_b>(<ObjectHandle>, <arg type_b>)
<ClassName>::<MethodName>(<arg type_a>)	Becomes	<ClassName>_<MethodName><type_a>(<arg type a>)
<ClassName>::<MethodName>(<arg type_b>)	Becomes	<ClassName>_<MethodName><type_b>(<arg type b>)

d) Allocate-once memory allocation APIs

AutoAlloc::New<ObjectName>	Becomes	Alloc_<ObjectName>
AutoAlloc::Allocate(uint16_t u16Size_)	Becomes	AutoAlloc(uint16_t u16Size_)

11.2 Allocating Objects

Aside from the API name translations, the object allocation scheme is the major different between Mark3C and [Mark3](#). Instead of instantiating objects of the various kernel types, kernel objects must be declared using Declaration macros, which serve the purpose of reserving memory for the kernel object, and provide an opaque handle to that object memory. This is the case for statically-allocated objects, and objects allocated on the stack.

Example: Declaring a thread

```
#include "mark3c.h"

// Statically-allocated
DECLARE_THREAD(hMyThread1);
...

// On stack
int main(void)
{
    DECLARE_THREAD(hMyThread2);
    ...
}
```

Where:

hMyThread1 - is a handle to a statically-allocated thread
hMyThread2 - is a handle to a thread allocated from the main stack.

Alternatively, the AutoAlloc APIs can be used to dynamically allocate objects, as demonstrated in the following example.

```
void Allocate_Example(void)
{
    Thread_t hMyThread = AutoAlloc_Thread();

    Thread_Init(hMyThread, awMyStack, sizeof(awMyStack), 1, MyFunction, 0);
}
```

Note that the relevant kernel-object Init() function *must* be called prior to using any kernel object, whether or not they have been allocated statically, or dynamically.

Chapter 12

Release Notes

12.1 R10 Release

- New: Coroutines + Cooperative scheduler
- New: Critical section APIs defined in kernel lib
- New: RAIL critical section (CriticalGuard object)
- New: RAIL scheduler-disabled context (SchedulerGuard object)
- Kernel code updated to use RAIL critical sections instead of CS_ENTER/CS_EXIT macros
- Updated documentation

12.2 R9 Release

- New: templated linked-lists to avoid explicit casting used in list traversal
- New: ThreadList class to efficiently track all threads in the system
- Remove use of C-style casts in kernel
- Remove use of 0 as nullptr in kernel
- Refactor code to use constexpr instead of C-style preprocessor defines where possible
- Refactor priority-map class as a set of template classes, reducing use of macros and defines
- Fix a "disappearing thread" bug where an inopportune context switch could cause a thread to get lost
- Docs no longer build by default

12.3 R8 Release

- Structural changes to separate the kernel from the rest of Mark3-repo
- Cleanup and reformatting

12.4 R7 (Full Throttle) Release

- Re-focusing project on kernel, integrating with 3rd party code instead of 1st party middleware
- Re-focusing on atmega1284p and cortex-m as default targets
- New: Refactored codebase to C++14 standard
- New: Moved non-kernel code, drivers, libs, and BSPs to separate repos from kernel
- New: Modular repository-based structure, managed via Android's Repo tool
- New: ConditionVariable kernel API
- New: ReaderWriterLock kernel API
- New: AutoAlloc redirects to user-defined allocators
- New: Global new() and delete() overrides redirect to AutoAlloc APIs
- New: RAII Mutex Locking APIs
- New: Support for cortex-a53 (aarch64) targets
- New: Doxygen builds as part of cmake process
- Updated Mark3c for new APIs
- Moved driver layer out of the kernel
- Moved all non-essential libraries out of the kernel (into other repos)
- Build system supports modular BSP architecture
- Removed fake idle-thread feature, since it doesn't support all targets
- Unit tests and examples will run on any target with a BSP
- Moved AVR-specific code out of the kernel (kernelaware debugging support)
- Remove most build-time configuration flags from [mark3cfg.h](#), and remove associated ifdefs throughout the code.
- Support qemu-system-arm's lm3s6965 evb target, with semihosting
- Support qemu-system-arm's rasbp3 evb target, with semihosting
- Increase test coverage
- Various bugfixes and improvements

12.5 R6 Release

- New: Replace recursive-make build system with CMake and Ninja
- New: Transitioned version control to Git from Subversion.
- New: Socket library, implementing named "domain-socket" style IPC
- New: State Machine framework library
- New: Software I2C library completed, with demo app
- New: Kernel Timer loop can optionally be run within its own thread instead of a nested interrupt

- New: UART drivers are all now abstracted through UartDriver base class for portability
- Experimental: Process library, allowing for the creation of resource-isolated processes
- Removed: Bare-metal support for Atmel SAMD20 (generic port still works)
- Cleanup all compiler warnings on atmega328p
- Various Bugfixes and optimizations
- Various Script changes related to automating the build + release process

12.6 R5 Release

- New: Shell library for creating responsive CLIs for embedded applications (M3Shell)
- New: Stream library for creating thread-safe buffered streams (streamer)
- New: Blocking UART implementation for AVR (drvUARTplus)
- New: "Extended context" kernel feature, which is used to implement thread-local storage
- New: "Extra Checks" kernel feature, which enforces safe API usage under pain of Kernel Panic
- New: Realtime clock library
- New: Example application + bsp for the open-hardware Mark3no development board (mark3no)
- New: Kernel objects descoped/destroyed while still in active use will now cause kernel panic
- New: Kernel callouts for thread creation/destruction/context switching, used for time tracking
- New: Simple power management class
- New: WIP software-based I2C + SPI drivers
- Optimized thread scheduling via target-optimized "count-leading-zero" macros
- Expanded memutil library
- Various optimizations of ARM Cortex-M assembly code
- Various bugfixes to Timer code
- Improved stack overflow checking + warning (stack guard kernel feature)
- AVR bootloader now supports targets with more than 64K of flash
- Moved some port configuration out of platform.mak into header files in the kernel port code
- The usual minor bugfixes and "gentle refactoring"

12.7 R4 Release

- New: C-language bindings for [Mark3](#) kernel (mark3c library)
- New: Support for ARM Cortex-M3 and Cortex-M4 (floating point) targets
- New: Support for Atmel AVR atmega2560 and arduino pro mega
- New: Full-featured, lightweight heap implementation
- New: Mailbox IPC class
- New: Notification object class
- New: lightweight tracelogger/instrumentation implementation (buffalogger), with sample parser
- New: High-performance AVR Software UART implementation
- New: Allocate-once "AutoAlloc" memory allocator
- New: Fixed-time blocking/unblocking operations added to ThreadList/Blocking class
- Placement-new supported for all kernel objects
- Scheduler now supports up to 1024 levels of thread priority, up from 8 (configurable at build-time)
- Kernel now uses stdint.h types for standard integers (instead of K_CHAR, K_ULONG, etc.)
- Greatly expanded documentation, with many new examples covering all key kernel features
- Expanded unit test coverage on AVR
- Updated build system and scripts for easier kernel configuration
- Updated builds to only attempt to build tests for supported platforms

12.8 R3 Release

- New: Added support for MSP430 microcontrollers
- New: Added Kernel Idle-Function hook to eliminate the need for a dedicated idle-thread (where supported)
- New: Support for kernel-aware simulation and testing via fIAVR AVR simulator
- Updated AVR driver selection
- General bugfixes and maintenance
- Expanded documentation and test coverage

12.9 R2

- Experimental release, using a "kernel transaction queue" for serializing kernel calls
- Works as a proof-of-concept, but abandoned due to overhead of the transaction mechanism in the general case.

12.10 R1 - 2nd Release Candidate

- New: Added support for ARM Cortex-M0 targets
- New: Added support for various AVR targets
- New: Timers now support a "tolerance" parameter for grouping timers with close expiry times
- Expanded scripts and automation used in build/test
- Updated and expanded graphics APIs
- Large number of bugfixes

12.11 R1 - 1st Release Candidate

- Initial release, with support for AVR microcontrollers

Chapter 13

Code Size Profiling

The following report details the size of each module compiled into the kernel. The size of each component is dependent on the flags specified in [mark3cfg.h](#) and [portcfg.h](#) at compile time. Note that these sizes represent the echo maximum size of each module before dead code elimination and any additional link-time optimization, and represent the maximum possible size that any module can take.

The results below are for profiling on Atmel AVR atmega1284p-based targets using gcc. Results are not necessarily indicative of relative or absolute performance on other platforms or toolchains.

13.1 Information

Date Profiled: Wed Jun 12 21:37:23 EDT 2019

13.2 Compiler Version

avr-gcc (GCC) 5.4.0 Copyright (C) 2015 Free Software Foundation, Inc. This is free software; see the source for copying conditions. There is NO warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.

13.3 Profiling Results

[Mark3](#) Module Size Report:

```
- Atomic Operations..... : 56 Bytes
- Allocate-once Heap..... : 120 Bytes
- Synchronization Objects - Base Class..... : 126 Bytes
- Condition Variables (Synchronization Object).... : 286 Bytes
- Coroutine task-list management..... : 122 Bytes
- Main coroutine task object..... : 336 Bytes
- Coroutine task scheduler..... : 318 Bytes
- Synchronization Object - Event Flag..... : 872 Bytes
- Mark3 Kernel Base Class..... : 142 Bytes
- Semaphore (Synchronization Object)..... : 568 Bytes
- Fundamental Kernel Linked-List Classes..... : 560 Bytes
- RAII Locking Support based on Mark3 Mutex class. : 62 Bytes
- Mailbox IPC Support..... : 1064 Bytes
- Message-based IPC..... : 288 Bytes
- Mutex (Synchronization Object)..... : 658 Bytes
```

- Notification Blocking Object..... : 626 Bytes
- Performance-profiling timers..... : 366 Bytes
- Round-Robin Scheduling Support..... : 265 Bytes
- Reader-writer Locks (Synchronization Object).... : 206 Bytes
- [Thread](#) Scheduling..... : 570 Bytes
- [Thread](#) Implementation..... : 1501 Bytes
- Fundamental Kernel [Thread](#)-list Data Structures.. : 440 Bytes
- ThreadListList Data Structures..... : 20 Bytes
- Software [Timer](#) Kernel Object..... : 241 Bytes
- Software [Timer](#) Management..... : 412 Bytes
- Atmel [AVR](#) - Kernel Interrupt Implemenation..... : 28 Bytes
- Atmel [AVR](#) - Kernel [Timer](#) Implementation..... : 810 Bytes
- Atmel [AVR](#) - Basic Threading Support..... : 506 Bytes

[Mark3](#) Kernel Size Summary:

- Kernel : 3359 Bytes
- Synchronization Objects : 4630 Bytes
- Port : 1344 Bytes
- Features : 1460 Bytes
- Coroutines : 776 Bytes
- Untracked Objects : 0 Bytes
- Total Size : 11569 Bytes

Chapter 14

Namespace Index

14.1 Namespace List

Here is a list of all namespaces with brief descriptions:

Mark3	91
Mark3::Atomic	
The Atomic namespace This utility module provides primitives for atomic operations - that is, operations that are guaranteed to execute uninterrupted. Basic atomic primitives provided here include Set/Add/Subtract, as well as an atomic test-and-set	102

Chapter 15

Hierarchical Index

15.1 Class Hierarchy

This inheritance list is sorted roughly, but not completely, alphabetically:

Mark3::AutoAlloc	105
Mark3::BlockingObject	110
Mark3::EventFlag	135
Mark3::Mutex	193
Mark3::Notify	199
Mark3::Semaphore	229
Mark3::ConditionVariable	120
Mark3::CoScheduler	127
Mark3::CriticalGuard	131
Mark3::CriticalSection	132
Mark3::Kernel	140
Mark3::KernelSWI	150
Mark3::KernelTimer	151
Mark3::LinkedList	152
Mark3::CircularLinkedList	115
Mark3::TypedCircularLinkedList< T >	286
Mark3::TypedCircularLinkedList< Coroutine >	286
Mark3::CoList	117
Mark3::TypedCircularLinkedList< Thread >	286
Mark3::ThreadList	264
Mark3::DoubleLinkedList	133
Mark3::TypedDoubleLinkedList< T >	290
Mark3::TypedDoubleLinkedList< Mark3::Message >	290
Mark3::TypedDoubleLinkedList< Mark3::ThreadList >	290
Mark3::TypedDoubleLinkedList< Timer >	290
Mark3::TimerList	280
Mark3::LinkedListNode	155
Mark3::ThreadList	264
Mark3::TypedLinkedListNode< T >	293
Mark3::TypedLinkedListNode< Coroutine >	293
Mark3::Coroutine	124
Mark3::TypedLinkedListNode< Message >	293
Mark3::Message	182
Mark3::TypedLinkedListNode< Thread >	293

Mark3::Thread	242
Mark3::TypedLinkedListNode< Timer >	293
Mark3::Timer	273
Mark3::LockGuard	158
Mark3::Mailbox	161
Mark3::MemUtil	173
Mark3::MessagePool	186
Mark3::MessageQueue	189
Mark3::PriorityMapL1< T, C >	202
Mark3::PriorityMapL2< T, C >	206
Mark3::ProfileTimer	210
Mark3::Quantum	213
Mark3::ReaderWriterLock	216
Mark3::Scheduler	221
Mark3::SchedulerGuard	227
Mark3::Streamer	234
Mark3::ThreadListList	270
Mark3::ThreadPort	272
Mark3::TimerScheduler	283
Mark3::Token_t	285

Chapter 16

Class Index

16.1 Class List

Here are the classes, structs, unions and interfaces with brief descriptions:

[Mark3::AutoAlloc](#)

The [AutoAlloc](#) class. This class provides an object-allocation interface for both kernel objects and user-defined types. This class supplies callouts for alloc/free that use object-type metadata to determine how objects may be allocated, allowing a user to create custom dynamic memory implementations for specific object types and sizes. As a result, the user-defined allocators can avoid the kinds of memory fragmentation and exhaustion issues that occur in typical embedded systems in which a single heap is used to satisfy all allocations in the application 105

[Mark3::BlockingObject](#)

The [BlockingObject](#) class. Class implementing thread-blocking primitives. used for implementing things like semaphores, mutexes, message queues, or anything else that could cause a thread to suspend execution on some external stimulus 110

[Mark3::CircularLinkedList](#)

The [CircularLinkedList](#) class Circular-linked-list data type, inherited from the base [LinkedList](#) type . . . 115

[Mark3::CoList](#)

The [CoList](#) class The [CoList](#) class implements a circular-linked-listed structure for coroutine objects. The intent of this object is to maintain a list of active coroutine objects with a specific priority or state, to ensure that a freshly-schedulable co-routine always exists at the head of the list 117

[Mark3::ConditionVariable](#)

The [ConditionVariable](#) class This class implements a condition variable. This is a synchronization object that allows multiple threads to block, each waiting for specific signals unique to them. Access to the specified condition is guarded by a mutex that is supplied by the caller. This object can permit multiple waiters that can be unblocked one-at-a-time via signalling, or unblocked all at once via broadcasting. This object is built upon lower-level primitives, and is somewhat more heavyweight than the primitive types supplied by the kernel 120

[Mark3::Coroutine](#)

The [Coroutine](#) class implements a lightweight, run-to-completion task that forms the basis for cooperative task scheduling in [Mark3](#). Coroutines are designed to be run from a singular context, and scheduled as a result of events occurring from threads, timers, interrupt sources, or other co-routines 124

[Mark3::CoScheduler](#)

The [CoScheduler](#) class. This class implements the coroutine scheduler. Similar to the [Mark3](#) thread scheduler, the highest-priority active object is scheduled / returned for execution. If no active co-routines are available to be scheduled, then the scheduler returns nullptr 127

Mark3::CriticalSectionGuard	
The CriticalSectionGuard class. This class provides an implementation of RAIL for critical sections. Object creation results in a critical section being invoked. The subsequent destructor call results in the critical section being released	131
Mark3::CriticalSection	
The CriticalSection class. This class implements a portable CriticalSection interface based on macros/inline functions that are implemented as part of each port	132
Mark3::DoubleLinkedList	
The DoubleLinkedList Class Doubly-linked-list data type, inherited from the base LinkedList type . .	133
Mark3::EventFlag	
The EventFlag class. This class implements a blocking object, similar to a semaphore or mutex, commonly used for synchronizing thread execution based on events occurring within the system. Each EventFlag object contains a 16-bit bitmask, which is used to trigger events on associated threads. Threads wishing to block, waiting for a specific event to occur can wait on any pattern within this 16-bit bitmask to be set. Here, we provide the ability for a thread to block, waiting for ANY bits in a specified mask to be set, or for ALL bits within a specific mask to be set. Depending on how the object is configured, the bits that triggered the wakeup can be automatically cleared once a match has occurred	135
Mark3::Kernel	
The Kernel Class encapsulates all of the kernel startup, configuration and management functions	140
Mark3::KernelSWI	
The KernelSWI Class provides the software-interrupt used to implement the context-switching interrupt used by the kernel. This interface must be implemented by target-specific code in the porting layer	150
Mark3::KernelTimer	
Timer interface used by all time-based scheduling/timer subsystems in the kernel. This interface must be implemented by target-specific code in the porting layer	151
Mark3::LinkedList	
The LinkedList Class Abstract-data-type from which all other linked-lists are derived	152
Mark3::LinkedListNode	
The LinkedListNode Class Basic linked-list node data structure. This data is managed by the linked-list class types, and can be used transparently between them	155
Mark3::LockGuard	
The LockGuard class. This class provides RAIL locks based on Mark3's kernel Mutex object. Note that Mark3 does not support exceptions, so care must be taken to ensure that this object is only used where that constraint can be met	158
Mark3::Mailbox	
The Mailbox class. This class implements an IPC mechanism based on sending/receiving envelopes containing data of a fixed size, configured at initialization) that reside within a buffer of memory provided by the user	161
Mark3::MemUtil	
String and Memory manipulation class	173
Mark3::Message	
Message class. This object provides threadsafe message-based IPC services based on exchange of objects containing a data pointer and minimal application-defined metadata. Messages are to be allocated/produced by the sender, and deallocated/consumed by the receiver .	182
Mark3::MessagePool	
The MessagePool Class The MessagePool class implements a simple allocator for message objects exchanged between threads. The sender allocates (pop's) messages, then sends them to the receiver. Upon receipt, it is the receiver's responsibility to deallocate (push) the message back to the pool	186
Mark3::MessageQueue	
The MessageQueue class. Implements a mechanism used to send/receive data between threads. Allows threads to block, waiting for messages to be sent from other contexts	189
Mark3::Mutex	
The Mutex Class. Class providing Mutual-exclusion locks, based on BlockingObject	193

[Mark3::Notify](#)

The [Notify](#) class. This class provides a blocking object type that allows one or more threads to wait for an event to occur before resuming operation 199

[Mark3::PriorityMapL1< T, C >](#)

The [PriorityMapL1](#) class This class implements a priority bitmap data structure. Each bit in the objects internal storage represents a priority. When a bit is set, it indicates that something is scheduled at the bit's corresponding priority, when a bit is clear it indicates that no entities are scheduled at that priority. This object provides the fundamental logic required to implement efficient priority-based scheduling for the thread + coroutine schedulers in the kernel 202

[Mark3::PriorityMapL2< T, C >](#)

The [PriorityMapL2](#) class This class implements a priority bitmap data structure. Each bit in the objects internal storage represents a priority. When a bit is set, it indicates that something is scheduled at the bit's corresponding priority, when a bit is clear it indicates that no entities are scheduled at that priority. This object provides the fundamental logic required to implement efficient priority-based scheduling for the thread + coroutine schedulers in the kernel 206

[Mark3::ProfileTimer](#)

Profiling timer. This class is used to perform high-performance profiling of code to see how `int32_t` certain operations take. useful in instrumenting the performance of key algorithms and time-critical operations to ensure real-timer behavior 210

[Mark3::Quantum](#)

The [Quantum](#) Class. Static-class used to implement [Thread](#) quantum functionality, which is fundamental to round-robin thread scheduling 213

[Mark3::ReaderWriterLock](#)

The [ReaderWriterLock](#) class. This class implements an object that marshalls access to a resource based on the intended usage of the resource. A reader-writer lock permits multiple concurrent read access, or single-writer access to a resource. If the object holds a write lock, other writers, and all readers will block until the writer is finished. If the object holds reader locks, all writers will block until all readers are finished before the first writer can take ownership of the resource. This is based upon lower-level synchronization primitives, and is somewhat more heavyweight than primitive synchronization types 216

[Mark3::Scheduler](#)

The [Scheduler](#) Class. This class provides priority-based round-robin [Thread](#) scheduling for all active threads managed by the kernel 221

[Mark3::SchedulerGuard](#)

The [SchedulerGuard](#) class This class implements RAIL-based control of the scheduler's global state. Upon object construction, the scheduler's state is cached locally and the scheduler is disabled (if not already disabled). Upon object destruction, the scheduler's previous state is restored. This object is interrupt-safe, although it has no effect when called from an interrupt given that interrupts are inherently higher-priority than threads 227

[Mark3::Semaphore](#)

[Semaphore](#) class provides Binary & Counting semaphore objects, based on [BlockingObject](#) base class 229

[Mark3::Streamer](#)

The [Streamer](#) class. This class implements a circular byte-buffer with thread and interrupt safe methods for writing-to and reading-from the buffer. Objects of this class type are designed to be shared between threads, or between threads and interrupts 234

[Mark3::Thread](#)

The [Thread](#) Class. This object providing the fundamental thread control data structures and functions that define a single thread of execution in the [Mark3](#) operating system. It is the fundamental data type used to provide multitasking support in the kernel 242

[Mark3::ThreadList](#)

The [ThreadList](#) Class. This class is used for building thread-management facilities, such as schedulers, and blocking objects 264

[Mark3::ThreadListList](#)

The [ThreadListList](#) class Class used to track all threadlists active in the OS kernel. At any point in time, the list can be traversed to get a complete view of all running, blocked, or stopped threads in the system 270

Mark3::ThreadPort	
The ThreadPort Class defines the target-specific functions required by the kernel for threading	272
Mark3::Timer	
The Timer Class. This class provides kernel-managed timers, used to provide high-precision delays. Functionality is useful to both user-code, and is used extensively within the kernel and its blocking objects to implement round-robin scheduling, thread sleep, and timeouts. Provides one-shot and periodic timers for use by application code. This object relies on a target-defined hardware timer implementation, which is multiplexed by the kernel's timer scheduler	273
Mark3::TimerList	
TimerList class. This class implements a doubly-linked-list of timer objects	280
Mark3::TimerScheduler	
The TimerScheduler Class. This implements a "Static" class used to manage a global list of timers used throughout the system	283
Mark3::Token_t	
Token descriptor struct format	285
Mark3::TypedCircularLinkedList< T >	
The TypedCircularLinkedList Class Circular-linked-list data type, inherited from the base LinkedList type, and templated for use with linked-list-node derived data-types	286
Mark3::TypedDoubleLinkedList< T >	
The TypedDoubleLinkedList Class Doubly-linked-list data type, inherited from the base LinkedList type, and templated for use with linked-list-node derived data-types	290
Mark3::TypedLinkedListNode< T >	
The TypedLinkedListNode class The TypedLinkedListNode class provides a linked-list node type for a specified object type. This can be used with typed link-list data structures to manage lists of objects without having to static-cast between the base type and the derived class	293

Chapter 17

File Index

17.1 File List

Here is a list of all files with brief descriptions:

/home/moslevin/projects/m3-repo/kernel/lib/mark3c/src/ mark3c.cpp	
Implementation of C-language wrappers for the Mark3 kernel	295
/home/moslevin/projects/m3-repo/kernel/lib/mark3c/src/public/ fake_types.h	
C-struct definitions that mirror	358
/home/moslevin/projects/m3-repo/kernel/lib/mark3c/src/public/ mark3c.h	
Header providing C-language API bindings for the Mark3 kernel	360
/home/moslevin/projects/m3-repo/kernel/lib/memutil/ memutil.cpp	
Implementation of memory, string, and conversion routines	424
/home/moslevin/projects/m3-repo/kernel/lib/memutil/public/ memutil.h	
Utility class containing memory, string, and conversion routines	431
/home/moslevin/projects/m3-repo/kernel/lib/streamer/ streamer.cpp	
Thread/Interrupt-safe byte-based data streaming	434
/home/moslevin/projects/m3-repo/kernel/lib/streamer/public/ streamer.h	
Thread/Interrupt-safe byte-based data streaming	433
/home/moslevin/projects/m3-repo/kernel/src/ atomic.cpp	
Basic Atomic Operations	455
/home/moslevin/projects/m3-repo/kernel/src/ autoalloc.cpp	
Automatic memory allocation for kernel objects	456
/home/moslevin/projects/m3-repo/kernel/src/ blocking.cpp	
Implementation of base class for blocking objects	457
/home/moslevin/projects/m3-repo/kernel/src/ colist.cpp	
CoRoutine List structure implementation	459
/home/moslevin/projects/m3-repo/kernel/src/ condvar.cpp	
Condition Variable implementation	460
/home/moslevin/projects/m3-repo/kernel/src/ coroutine.cpp	
Coroutine object implementation	461
/home/moslevin/projects/m3-repo/kernel/src/ cosched.cpp	
CoRoutine Scheduler implementation	463
/home/moslevin/projects/m3-repo/kernel/src/ eventflag.cpp	
Event Flag Blocking Object/IPC-Object implementation	464
/home/moslevin/projects/m3-repo/kernel/src/ kernel.cpp	
Kernel initialization and startup code	468
/home/moslevin/projects/m3-repo/kernel/src/ ksemaphore.cpp	
Semaphore Blocking-Object Implementation	469
/home/moslevin/projects/m3-repo/kernel/src/ ll.cpp	
Core Linked-List implementation, from which all kernel objects are derived	472

/home/moslevin/projects/m3-repo/kernel/src/ lockguard.cpp	
Mutex RAIL helper class	474
/home/moslevin/projects/m3-repo/kernel/src/ mailbox.cpp	
Mailbox + Envelope IPC mechanism	475
/home/moslevin/projects/m3-repo/kernel/src/ message.cpp	
Inter-thread communications via message passing	478
/home/moslevin/projects/m3-repo/kernel/src/ mutex.cpp	
Mutual-exclusion object	480
/home/moslevin/projects/m3-repo/kernel/src/ notify.cpp	
Lightweight thread notification - blocking object	483
/home/moslevin/projects/m3-repo/kernel/src/ profile.cpp	
Code profiling utilities	486
/home/moslevin/projects/m3-repo/kernel/src/ quantum.cpp	
Thread Quantum Implementation for Round-Robin Scheduling	552
/home/moslevin/projects/m3-repo/kernel/src/ readerwriter.cpp	
Reader-writer lock implementation	554
/home/moslevin/projects/m3-repo/kernel/src/ scheduler.cpp	
Strict-Priority + Round-Robin thread scheduler implementation	555
/home/moslevin/projects/m3-repo/kernel/src/ thread.cpp	
Platform-Independent thread class Definition	557
/home/moslevin/projects/m3-repo/kernel/src/ threadlist.cpp	
Thread linked-list definitions	563
/home/moslevin/projects/m3-repo/kernel/src/ threadlistlist.cpp	
Class implementing a doubly-linked list of thread lists	565
/home/moslevin/projects/m3-repo/kernel/src/ timer.cpp	
Timer implementations	565
/home/moslevin/projects/m3-repo/kernel/src/ timerlist.cpp	
Implements timer list processing algorithms, responsible for all timer tick and expiry logic	567
/home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/ kernelswi.cpp	
Kernel Software interrupt implementation for ATmega1284p	438
/home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/ kerneltimer.cpp	
Kernel Timer Implementation for ATmega1284p	439
/home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/ threadport.cpp	
ATmega1284p Multithreading	452
/home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/ portcfg.h	
Mark3 Port Configuration	442
/home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/ threadport.h	
ATmega1284p Multithreading support	448
/home/moslevin/projects/m3-repo/kernel/src/public/ atomic.h	
Basic Atomic Operations	487
/home/moslevin/projects/m3-repo/kernel/src/public/ autoalloc.h	
Automatic memory allocation for kernel objects	489
/home/moslevin/projects/m3-repo/kernel/src/public/ blocking.h	
Blocking object base class declarations	491
/home/moslevin/projects/m3-repo/kernel/src/public/ colist.h	
CoRoutine List structure implementation	493
/home/moslevin/projects/m3-repo/kernel/src/public/ condvar.h	
Condition Variable implementation	494
/home/moslevin/projects/m3-repo/kernel/src/public/ coroutine.h	
CoRoutine implementation	495
/home/moslevin/projects/m3-repo/kernel/src/public/ cosched.h	
CoRoutine Scheduler implementation	497
/home/moslevin/projects/m3-repo/kernel/src/public/ criticalguard.h	
RAIL Critical Section Implementation	498
/home/moslevin/projects/m3-repo/kernel/src/public/ criticalsection.h	
Critical Section Support	499
/home/moslevin/projects/m3-repo/kernel/src/public/ eventflag.h	
Event Flag Blocking Object/IPC-Object definition	500

/home/moslevin/projects/m3-repo/kernel/src/public/ithreadport.h	
Thread porting interface	501
/home/moslevin/projects/m3-repo/kernel/src/public/kernel.h	
Kernel initialization and startup class	502
/home/moslevin/projects/m3-repo/kernel/src/public/kerneldebug.h	
Macros and functions used for assertions, kernel traces, etc	504
/home/moslevin/projects/m3-repo/kernel/src/public/kernelswi.h	
Kernel Software interrupt declarations	506
/home/moslevin/projects/m3-repo/kernel/src/public/kerneltimer.h	
Kernel Timer Class declaration	507
/home/moslevin/projects/m3-repo/kernel/src/public/kerneltypes.h	
Basic data type primitives used throughout the OS	508
/home/moslevin/projects/m3-repo/kernel/src/public/ksemaphore.h	
Semaphore Blocking Object class declarations	509
/home/moslevin/projects/m3-repo/kernel/src/public/ll.h	
Core linked-list declarations, used by all kernel list types At the heart of RTOS data structures are linked lists. Having a robust and efficient set of linked-list types that we can use as a foundation for building the rest of our kernel types allows u16 to keep our RTOS code efficient and logically-separated	510
/home/moslevin/projects/m3-repo/kernel/src/public/lockguard.h	
Mutex RAIL helper class	513
/home/moslevin/projects/m3-repo/kernel/src/public/mailbox.h	
Mailbox + Envelope IPC Mechanism	514
/home/moslevin/projects/m3-repo/kernel/src/public/manual.h	
Single include file given to users of the Mark3 Kernel API	516
/home/moslevin/projects/m3-repo/kernel/src/public/mark3.h	
Mark3 Kernel Configuration This file is used to configure the kernel for your specific application in order to provide the optimal set of features for a given use case	517
/home/moslevin/projects/m3-repo/kernel/src/public/message.h	
Inter-thread communication via message-passing Embedded systems guru Jack Ganssle once said that without a robust form of interprocess communications (IPC), an RTOS is just a toy. Mark3 implements a form of IPC to provide safe and flexible messaging between threads . . .	518
/home/moslevin/projects/m3-repo/kernel/src/public/mutex.h	
Mutual exclusion class declaration Resource locks are implemented using mutual exclusion semaphores (Mutex_t). Protected blocks can be placed around any resource that may only be accessed by one thread at a time. If additional threads attempt to access the protected resource, they will be placed in a wait queue until the resource becomes available. When the resource becomes available, the thread with the highest original priority claims the resource and is activated. Priority inheritance is included in the implementation to prevent priority inversion. Always ensure that you claim and release your mutex objects consistently, otherwise you may end up with a deadlock scenario that's hard to debug	522
/home/moslevin/projects/m3-repo/kernel/src/public/notify.h	
Lightweight thread notification - blocking object	524
/home/moslevin/projects/m3-repo/kernel/src/public/paniccodes.h	
Defines the reason codes thrown when a kernel panic occurs	526
/home/moslevin/projects/m3-repo/kernel/src/public/priomap.h	
Priority map data structure	527
/home/moslevin/projects/m3-repo/kernel/src/public/priomap1.h	
1-Level bitmap allocator template-class used for scheduler implementation	530
/home/moslevin/projects/m3-repo/kernel/src/public/priomap2.h	
2-Level priority allocator template-class used for scheduler implementation	531
/home/moslevin/projects/m3-repo/kernel/src/public/profile.h	
High-precision profiling timers Enables the profiling and instrumentation of performance-critical code. Multiple timers can be used simultaneously to enable system-wide performance metrics to be computed in a lightweight manner	533
/home/moslevin/projects/m3-repo/kernel/src/public/profiling_results.h	
	535
	536

/home/moslevin/projects/m3-repo/kernel/src/public/ quantum.h	
Thread Quantum declarations for Round-Robin Scheduling	537
/home/moslevin/projects/m3-repo/kernel/src/public/ readerwriter.h	
Reader-Writer lock implementation	538
/home/moslevin/projects/m3-repo/kernel/src/public/ scheduler.h	
Thread scheduler function declarations This scheduler implements a very flexible type of scheduling, which has become the defacto industry standard when it comes to real-time operating systems. This scheduling mechanism is referred to as priority round- robin	539
/home/moslevin/projects/m3-repo/kernel/src/public/ schedulerguard.h	
RAII Scheduler Locking	541
/home/moslevin/projects/m3-repo/kernel/src/public/ sizeprofile.h	542
/home/moslevin/projects/m3-repo/kernel/src/public/ thread.h	
Platform independent thread class declarations Threads are an atomic unit of execution, and each instance of the thread class represents an instance of a program running of the processor. The Thread is the fundmanetal user-facing object in the kernel - it is what makes multiprocessing possible from application code	543
/home/moslevin/projects/m3-repo/kernel/src/public/ threadlist.h	
Thread linked-list declarations	546
/home/moslevin/projects/m3-repo/kernel/src/public/ threadlistlist.h	
Class implementing a doubly-linked list of thread lists	547
/home/moslevin/projects/m3-repo/kernel/src/public/ timer.h	
Timer object declarations	548
/home/moslevin/projects/m3-repo/kernel/src/public/ timerlist.h	
Timer list declarations	550
/home/moslevin/projects/m3-repo/kernel/src/public/ timerscheduler.h	
Timer scheduler declarations	551

Chapter 18

Namespace Documentation

18.1 Mark3 Namespace Reference

Namespaces

- [Atomic](#)

The [Atomic](#) namespace This utility module provides primitives for atomic operations - that is, operations that are guaranteed to execute uninterrupted. Basic atomic primitives provided here include Set/Add/Subtract, as well as an atomic test-and-set.

Classes

- class [AutoAlloc](#)

The [AutoAlloc](#) class. This class provides an object-allocation interface for both kernel objects and user-defined types. This class supplies callouts for alloc/free that use object-type metadata to determine how objects may be allocated, allowing a user to create custom dynamic memory implementations for specific object types and sizes. As a result, the user-defined allocators can avoid the kinds of memory fragmentation and exhaustion issues that occur in typical embedded systems in which a single heap is used to satisfy all allocations in the application.

- class [BlockingObject](#)

The [BlockingObject](#) class. Class implementing thread-blocking primitives. used for implementing things like semaphores, mutexes, message queues, or anything else that could cause a thread to suspend execution on some external stimulus.

- class [CircularLinkList](#)

The [CircularLinkList](#) class Circular-linked-list data type, inherited from the base [LinkList](#) type.

- class [CoList](#)

The [CoList](#) class The [CoList](#) class implements a circular-linked-listed structure for coroutine objects. The intent of this object is to maintain a list of active coroutine objects with a specific priority or state, to ensure that a freshly-schedulable co-routine always exists at the head of the list.

- class [ConditionVariable](#)

The [ConditionVariable](#) class This class implements a condition variable. This is a synchronization object that allows multiple threads to block, each waiting for specific signals unique to them. Access to the specified condition is guarded by a mutex that is supplied by the caller. This object can permit multiple waiters that can be unblocked one-at-a-time via signalling, or unblocked all at once via broadcasting. This object is built upon lower-level primitives, and is somewhat more heavyweight than the primitive types supplied by the kernel.

- class [Coroutine](#)

The [Coroutine](#) class implements a lightweight, run-to-completion task that forms the basis for co-operative task scheduling in [Mark3](#). Coroutines are designed to be run from a singular context, and scheduled as a result of events occurring from threads, timers, interrupt sources, or other co-routines.

- class [CoScheduler](#)
The [CoScheduler](#) class. This class implements the coroutine scheduler. Similar to the [Mark3](#) thread scheduler, the highest-priority active object is scheduled / returned for execution. If no active co-routines are available to be scheduled, then the scheduler returns nullptr.
- class [CriticalSectionGuard](#)
The [CriticalSectionGuard](#) class. This class provides an implementation of RAIL for critical sections. Object creation results in a critical section being invoked. The subsequent destructor call results in the critical section being released.
- class [CriticalSection](#)
The [CriticalSection](#) class. This class implements a portable [CriticalSection](#) interface based on macros/inline functions that are implemented as part of each port.
- class [DoubleLinkedList](#)
The [DoubleLinkedList](#) Class Doubly-linked-list data type, inherited from the base [LinkedList](#) type.
- class [EventFlag](#)
The [EventFlag](#) class. This class implements a blocking object, similar to a semaphore or mutex, commonly used for synchronizing thread execution based on events occurring within the system. Each [EventFlag](#) object contains a 16-bit bitmask, which is used to trigger events on associated threads. Threads wishing to block, waiting for a specific event to occur can wait on any pattern within this 16-bit bitmask to be set. Here, we provide the ability for a thread to block, waiting for ANY bits in a specified mask to be set, or for ALL bits within a specific mask to be set. Depending on how the object is configured, the bits that triggered the wakeup can be automatically cleared once a match has occurred.
- class [Kernel](#)
The [Kernel](#) Class encapsulates all of the kernel startup, configuration and management functions.
- class [KernelSWI](#)
The [KernelSWI](#) Class provides the software-interrupt used to implement the context-switching interrupt used by the kernel. This interface must be implemented by target-specific code in the porting layer.
- class [KernelTimer](#)
The [KernelTimer](#) class provides a timer interface used by all time-based scheduling/timer subsystems in the kernel. This interface must be implemented by target-specific code in the porting layer.
- class [LinkedList](#)
The [LinkedList](#) Class Abstract-data-type from which all other linked-lists are derived.
- class [LinkedListNode](#)
The [LinkedListNode](#) Class Basic linked-list node data structure. This data is managed by the linked-list class types, and can be used transparently between them.
- class [LockGuard](#)
The [LockGuard](#) class. This class provides RAIL locks based on [Mark3](#)'s kernel [Mutex](#) object. Note that [Mark3](#) does not support exceptions, so care must be taken to ensure that this object is only used where that constraint can be met.
- class [Mailbox](#)
The [Mailbox](#) class. This class implements an IPC mechanism based on sending/receiving envelopes containing data of a fixed size, configured at initialization) that reside within a buffer of memory provided by the user.
- class [MemUtil](#)
String and Memory manipulation class.
- class [Message](#)
the [Message](#) class. This object provides threadsafe message-based IPC services based on exchange of objects containing a data pointer and minimal application-defined metadata. Messages are to be allocated/produced by the sender, and deallocated/consumed by the receiver.
- class [MessagePool](#)
The [MessagePool](#) Class The [MessagePool](#) class implements a simple allocator for message objects exchanged between threads. The sender allocates (pop's) messages, then sends them to the receiver. Upon receipt, it is the receiver's responsibility to deallocate (push) the message back to the pool.
- class [MessageQueue](#)
The [MessageQueue](#) class. Implements a mechanism used to send/receive data between threads. Allows threads to block, waiting for messages to be sent from other contexts.
- class [Mutex](#)
The [Mutex](#) Class. Class providing Mutual-exclusion locks, based on [BlockingObject](#).
- class [Notify](#)

The [Notify](#) class. This class provides a blocking object type that allows one or more threads to wait for an event to occur before resuming operation.

- class [PriorityMapL1](#)

The [PriorityMapL1](#) class This class implements a priority bitmap data structure. Each bit in the objects internal storage represents a priority. When a bit is set, it indicates that something is scheduled at the bit's corresponding priority, when a bit is clear it indicates that no entities are scheduled at that priority. This object provides the fundamental logic required to implement efficient priority-based scheduling for the thread + coroutine schedulers in the kernel.

- class [PriorityMapL2](#)

The [PriorityMapL2](#) class This class implements a priority bitmap data structure. Each bit in the objects internal storage represents a priority. When a bit is set, it indicates that something is scheduled at the bit's corresponding priority, when a bit is clear it indicates that no entities are scheduled at that priority. This object provides the fundamental logic required to implement efficient priority-based scheduling for the thread + coroutine schedulers in the kernel.

- class [ProfileTimer](#)

Profiling timer. This class is used to perform high-performance profiling of code to see how `int32_t` certain operations take. useful in instrumenting the performance of key algorithms and time-critical operations to ensure real-timer behavior.

- class [Quantum](#)

The [Quantum](#) Class. Static-class used to implement [Thread](#) quantum functionality, which is fundamental to round-robin thread scheduling.

- class [ReaderWriterLock](#)

The [ReaderWriterLock](#) class. This class implements an object that marshalls access to a resource based on the intended usage of the resource. A reader-writer lock permits multiple concurrent read access, or single-writer access to a resource. If the object holds a write lock, other writers, and all readers will block until the writer is finished. If the object holds reader locks, all writers will block until all readers are finished before the first writer can take ownership of the resource. This is based upon lower-level synchronization primitives, and is somewhat more heavyweight than primitive synchronization types.

- class [Scheduler](#)

The [Scheduler](#) Class. This class provides priority-based round-robin [Thread](#) scheduling for all active threads managed by the kernel.

- class [SchedulerGuard](#)

The [SchedulerGuard](#) class This class implements RAIL-based control of the scheduler's global state. Upon object construction, the scheduler's state is cached locally and the scheduler is disabled (if not already disabled). Upon object destruction, the scheduler's previous state is restored. This object is interrupt-safe, although it has no effect when called from an interrupt given that interrupts are inherently higher-priority than threads.

- class [Semaphore](#)

the [Semaphore](#) class provides Binary & Counting semaphore objects, based on [BlockingObject](#) base class.

- class [Streamer](#)

The [Streamer](#) class. This class implements a circular byte-buffer with thread and interrupt safe methods for writing-to and reading-from the buffer. Objects of this class type are designed to be shared between threads, or between threads and interrupts.

- class [Thread](#)

The [Thread](#) Class. This object providing the fundamental thread control data structures and functions that define a single thread of execution in the [Mark3](#) operating system. It is the fundamental data type used to provide multitasking support in the kernel.

- class [ThreadList](#)

The [ThreadList](#) Class. This class is used for building thread-management facilities, such as schedulers, and blocking objects.

- class [ThreadListList](#)

The [ThreadListList](#) class Class used to track all threadlists active in the OS kernel. At any point in time, the list can be traversed to get a complete view of all running, blocked, or stopped threads in the system.

- class [ThreadPort](#)

The [ThreadPort](#) Class defines the target-specific functions required by the kernel for threading.

- class [Timer](#)

The [Timer](#) Class. This class provides kernel-managed timers, used to provide high-precision delays. Functionality is useful to both user-code, and is used extensively within the kernel and its blocking objects to implement round-robin scheduling, thread sleep, and timeouts. Provides one-shot and periodic timers for use by application code. This object relies on a target-defined hardware timer implementation, which is multiplexed by the kernel's timer scheduler.

- class [TimerList](#)
the [TimerList](#) class. This class implements a doubly-linked-list of timer objects.
- class [TimerScheduler](#)
The [TimerScheduler](#) Class. This implements a "Static" class used to manage a global list of timers used throughout the system.
- struct [Token_t](#)
Token descriptor struct format.
- class [TypedCircularLinkedList](#)
The [TypedCircularLinkedList](#) Class Circular-linked-list data type, inherited from the base [LinkedList](#) type, and templated for use with linked-list-node derived data-types.
- class [TypedDoubleLinkedList](#)
The [TypedDoubleLinkedList](#) Class Doubly-linked-list data type, inherited from the base [LinkedList](#) type, and templated for use with linked-list-node derived data-types.
- class [TypedLinkedListNode](#)
The [TypedLinkedListNode](#) class The [TypedLinkedListNode](#) class provides a linked-list node type for a specified object type. This can be used with typed link-list data structures to manage lists of objects without having to static-cast between the base type and the derived class.

Typedefs

- using [AutoAllocAllocator_t](#) = void (*)(AutoAllocType eType_, size_t sSize_)
- using [AutoAllocFree_t](#) = void (*)(AutoAllocType eType_, void *pvObj_)
- using [CoPrioMap](#) = PriorityMapL1 < PORT_PRIO_TYPE, PORT_COROUTINE_PRIORITIES >
- using [CoroutineHandler](#) = void (*)(Coroutine *pclCaller_, void *pvContext_)
- using [PanicFunc](#) = void (*)(uint16_t u16PanicCode_)
- using [IdleFunc](#) = void (*)()
- using [ThreadEntryFunc](#) = void (*)(void *pvArg_)
- using [PriorityMap](#) = PriorityMapL1 < PORT_PRIO_TYPE, KERNEL_NUM_PRIORITIES >
- using [ThreadCreateCallout](#) = void (*)(Thread *pclThread_)
- using [ThreadExitCallout](#) = void (*)(Thread *pclThread_)
- using [ThreadContextCallout](#) = void (*)(Thread *pclThread_)
- using [TimerCallback](#) = void (*)(Thread *pclOwner_, void *pvData_)

Enumerations

- enum [AutoAllocType](#) : uint8_t {
AutoAllocType::EventFlag, AutoAllocType::MailBox, AutoAllocType::Message, AutoAllocType::MessagePool,
AutoAllocType::MessageQueue, AutoAllocType::Mutex, AutoAllocType::Notify, AutoAllocType::Semaphore,
AutoAllocType::Thread, AutoAllocType::Timer, AutoAllocType::ConditionVariable, AutoAllocType::Reader↵
WriterLock,
AutoAllocType::User, AutoAllocType::Raw = 0xFF }
- enum [EventFlagOperation](#) : uint8_t {
EventFlagOperation::All_Set = 0, EventFlagOperation::Any_Set, EventFlagOperation::All_Clear, EventFlag↵
Operation::Any_Clear,
EventFlagOperation::Pending_Unblock }
- enum [ThreadState](#) : uint8_t {
ThreadState::Exit = 0, ThreadState::Ready, ThreadState::Blocked, ThreadState::Stop,
ThreadState::Invalid }

Functions

- static void [KernelTimer_Task](#) (void *unused)
- static void [Thread_Switch](#) (void)
- [ISR](#) (INT2_vect) [__attribute__\(\(signal](#)
ISR(INT2_vect) SWI using INT2 - used to trigger a context switch.
- uint8_t [PORT_CLZ](#) (uint8_t in_)
- void [PORT_IRQ_ENABLE](#) ()
- void [PORT_IRQ_DISABLE](#) ()
- void [PORT_CS_ENTER](#) ()
- void [PORT_CS_EXIT](#) ()
- [K_WORD](#) [PORT_CS_NESTING](#) ()

Variables

- static constexpr auto [uMaxTimerTicks](#) = uint32_t { 0x7FFFFFFF }
Maximum value to set.
- static constexpr auto [uTimerTicksInvalid](#) = uint32_t { 0 }
- static constexpr auto [uTimerFlagOneShot](#) = uint8_t { 0x01 }
Timer is one-shot.
- static constexpr auto [uTimerFlagActive](#) = uint8_t { 0x02 }
Timer is currently active.
- static constexpr auto [uTimerFlagCallback](#) = uint8_t { 0x04 }
Timer is pending a callback.
- static constexpr auto [uTimerFlagExpired](#) = uint8_t { 0x08 }
Timer is actually expired.
- [naked](#)
- static constexpr auto [SR_](#) = uint8_t{0x3F}
- [K_WORD](#) [g_kwSFR](#)
- [K_WORD](#) [g_kwCriticalCount](#)

18.1.1 Typedef Documentation

18.1.1.1 [AutoAllocAllocator_t](#)

using [Mark3::AutoAllocAllocator_t](#) = typedef void* (*) ([AutoAllocType](#) eType_, size_t sSize_)

Definition at line 53 of file [autoalloc.h](#).

18.1.1.2 [AutoAllocFree_t](#)

using [Mark3::AutoAllocFree_t](#) = typedef void (*) ([AutoAllocType](#) eType_, void* pObj_)

Definition at line 54 of file [autoalloc.h](#).

18.1.1.3 CoPrioMap

```
using Mark3::CoPrioMap = typedef PriorityMapL1<PORT_PRIO_TYPE, PORT_COROUTINE_PRIORITIES>
```

Definition at line 30 of file [coroutine.h](#).

18.1.1.4 CoroutineHandler

```
using Mark3::CoroutineHandler = typedef void (*) (Coroutine* pClCaller_, void* pvContext_)
```

Definition at line 40 of file [coroutine.h](#).

18.1.1.5 IdleFunc

```
using Mark3::IdleFunc = typedef void (*) ()
```

Function pointer type used to implement the idle function, where support for an idle function (as opposed to an idle thread) exists.

Definition at line 37 of file [kerneltypes.h](#).

18.1.1.6 PanicFunc

```
using Mark3::PanicFunc = typedef void (*) (uint16_t u16PanicCode_)
```

Function pointer type used to implement kernel-panic handlers.

Definition at line 30 of file [kerneltypes.h](#).

18.1.1.7 PriorityMap

```
using Mark3::PriorityMap = typedef PriorityMapL1<PORT_PRIO_TYPE, KERNEL_NUM_PRIORITIES>
```

Definition at line 29 of file [priomap.h](#).

18.1.1.8 ThreadContextCallout

```
using Mark3::ThreadContextCallout = typedef void (*) (Thread* pClThread_)
```

Definition at line 54 of file [thread.h](#).

18.1.1.9 ThreadCreateCallout

```
using Mark3::ThreadCreateCallout = typedef void (*) (Thread* pclThread_)
```

Definition at line 52 of file [thread.h](#).

18.1.1.10 ThreadEntryFunc

```
using Mark3::ThreadEntryFunc = typedef void (*) (void* pvArg_)
```

Function pointer type used for thread entrypoint functions

Definition at line 43 of file [kerneltypes.h](#).

18.1.1.11 ThreadExitCallout

```
using Mark3::ThreadExitCallout = typedef void (*) (Thread* pclThread_)
```

Definition at line 53 of file [thread.h](#).

18.1.1.12 TimerCallback

```
using Mark3::TimerCallback = typedef void (*) (Thread* pclOwner_, void* pvData_)
```

This type defines the callback function type for timer events. Since these are called from an interrupt context, they do not operate from within a thread or object context directly – as a result, the context must be manually passed into the calls.

`pclOwner_` is a pointer to the thread that owns the timer `pvData_` is a pointer to some data or object that needs to know about the timer's expiry from within the timer interrupt context.

Definition at line 50 of file [timer.h](#).

18.1.2 Enumeration Type Documentation

18.1.2.1 AutoAllocType

```
enum Mark3::AutoAllocType : uint8_t [strong]
```

Enumerator

EventFlag	
MailBox	
Message	
MessagePool	
MessageQueue	
Mutex	
Notify	
Semaphore	
Thread	
Timer	
ConditionVariable	
ReaderWriterLock	
User	
Raw	

Definition at line 32 of file [autoalloc.h](#).

18.1.2.2 EventFlagOperation

```
enum Mark3::EventFlagOperation : uint8_t [strong]
```

This enumeration describes the different operations supported by the event flag blocking object.

Enumerator

All_Set	Block until all bits in the specified bitmask are set.
Any_Set	Block until any bits in the specified bitmask are set.
All_Clear	Block until all bits in the specified bitmask are cleared.
Any_Clear	Block until any bits in the specified bitmask are cleared.
Pending_Unblock	Special code. Not used by user

Definition at line 50 of file [kerneltypes.h](#).

18.1.2.3 ThreadState

```
enum Mark3::ThreadState : uint8_t [strong]
```

Enumeration representing the different states a thread can exist in

Enumerator

Exit	
Ready	
Blocked	
Stop	
Invalid	

Definition at line 62 of file [kerneltypes.h](#).

18.1.3 Function Documentation

18.1.3.1 ISR()

```
Mark3::ISR (
    INT2_vect )
```

ISR(INT2_vect) SWI using INT2 - used to trigger a context switch.

Definition at line 137 of file [threadport.cpp](#).

18.1.3.2 KernelTimer_Task()

```
static void Mark3::KernelTimer_Task (
    void * unused ) [static]
```

Definition at line 44 of file [kerneltimer.cpp](#).

18.1.3.3 PORT_CLZ()

```
uint8_t Mark3::PORT_CLZ (
    uint8_t in_ ) [inline]
```

Definition at line 144 of file [threadport.h](#).

18.1.3.4 PORT_CS_ENTER()

```
void Mark3::PORT_CS_ENTER ( ) [inline]
```

Definition at line 167 of file [threadport.h](#).

18.1.3.5 PORT_CS_EXIT()

```
void Mark3::PORT_CS_EXIT ( ) [inline]
```

Definition at line 178 of file [threadport.h](#).

18.1.3.6 PORT_CS_NESTING()

`K_WORD Mark3::PORT_CS_NESTING () [inline]`

Definition at line 187 of file [threadport.h](#).

18.1.3.7 PORT_IRQ_DISABLE()

`void Mark3::PORT_IRQ_DISABLE () [inline]`

Definition at line 161 of file [threadport.h](#).

18.1.3.8 PORT_IRQ_ENABLE()

`void Mark3::PORT_IRQ_ENABLE () [inline]`

Definition at line 155 of file [threadport.h](#).

18.1.3.9 Thread_Switch()

`static void Mark3::Thread_Switch (
void) [static]`

Definition at line 94 of file [threadport.cpp](#).

18.1.4 Variable Documentation

18.1.4.1 g_kwCriticalCount

`K_WORD Mark3::g_kwCriticalCount`

Definition at line 36 of file [threadport.cpp](#).

18.1.4.2 g_kwSFR

`K_WORD Mark3::g_kwSFR`

Definition at line 35 of file [threadport.cpp](#).

18.1.4.3 naked

Mark3::naked

Definition at line 136 of file [threadport.cpp](#).

18.1.4.4 SR_

```
constexpr auto Mark3::SR_ = uint8_t{0x3F} [static]
```

Definition at line 137 of file [threadport.h](#).

18.1.4.5 uMaxTimerTicks

```
constexpr auto Mark3::uMaxTimerTicks = uint32_t { 0x7FFFFFFF } [static]
```

Maximum value to set.

Definition at line 32 of file [timer.h](#).

18.1.4.6 uTimerFlagActive

```
constexpr auto Mark3::uTimerFlagActive = uint8_t { 0x02 } [static]
```

[Timer](#) is currently active.

Definition at line 35 of file [timer.h](#).

18.1.4.7 uTimerFlagCallback

```
constexpr auto Mark3::uTimerFlagCallback = uint8_t { 0x04 } [static]
```

[Timer](#) is pending a callback.

Definition at line 36 of file [timer.h](#).

18.1.4.8 uTimerFlagExpired

```
constexpr auto Mark3::uTimerFlagExpired = uint8_t { 0x08 } [static]
```

[Timer](#) is actually expired.

Definition at line 37 of file [timer.h](#).

18.1.4.9 uTimerFlagOneShot

```
constexpr auto Mark3::uTimerFlagOneShot = uint8_t { 0x01 } [static]
```

[Timer](#) is one-shot.

Definition at line 34 of file [timer.h](#).

18.1.4.10 uTimerTicksInvalid

```
constexpr auto Mark3::uTimerTicksInvalid = uint32_t { 0 } [static]
```

Definition at line 33 of file [timer.h](#).

18.2 Mark3::Atomic Namespace Reference

The [Atomic](#) namespace This utility module provides primitives for atomic operations - that is, operations that are guaranteed to execute uninterrupted. Basic atomic primitives provided here include Set/Add/Subtract, as well as an atomic test-and-set.

Functions

- template<typename T >
T [Set](#) (T *pSource_, T val_)
Set Set a variable to a given value in an uninterruptable operation.
- template<typename T >
T [Add](#) (T *pSource_, T val_)
Add Add a value to a variable in an uninterruptable operation.
- template<typename T >
T [Sub](#) (T *pSource_, T val_)
Sub Subtract a value from a variable in an uninterruptable operation.
- bool [TestAndSet](#) (bool *pbLock)
TestAndSet Test to see if a variable is set, and set it if it is not already set. This is an uninterruptable operation.

18.2.1 Detailed Description

The [Atomic](#) namespace This utility module provides primitives for atomic operations - that is, operations that are guaranteed to execute uninterrupted. Basic atomic primitives provided here include Set/Add/Subtract, as well as an atomic test-and-set.

18.2.2 Function Documentation

18.2.2.1 Add()

```
template<typename T >
T Mark3::Atomic::Add (
    T * pSource_,
    T val_ )
```

Add Add a value to a variable in an uninterruptable operation.

Parameters

<i>p</i> _↔ <i>Source</i> _↔ _	Pointer to a variable
<i>val</i> _	Value to add to the variable

Returns

Previously-held value in pSource_

Definition at line 64 of file [atomic.h](#).

18.2.2.2 Set()

```
template<typename T >
T Mark3::Atomic::Set (
    T * pSource_,
    T val_ )
```

Set Set a variable to a given value in an uninterruptable operation.

Parameters

<i>p</i> _↔ <i>Source</i> _↔ _	Pointer to a variable to set the value of
<i>val</i> _	New value to set in the variable

Returns

Previously-set value

Definition at line 47 of file [atomic.h](#).

18.2.2.3 Sub()

```
template<typename T >
T Mark3::Atomic::Sub (
    T * pSource_,
    T val_ )
```

Sub Subtract a value from a variable in an uninterruptable operation.

Parameters

<i>p_↔</i> <i>Source_↔</i> _	Pointer to a variable
<i>val_</i>	Value to subtract from the variable

Returns

Previously-held value in pSource_

Definition at line 81 of file [atomic.h](#).

18.2.2.4 TestAndSet()

```
bool Mark3::Atomic::TestAndSet (
    bool * pbLock )
```

TestAndSet Test to see if a variable is set, and set it if is not already set. This is an uninterruptable operation.

If the value is false, set the variable to true, and return the previously-held value.

If the value is already true, return true.

Parameters

<i>pbLock</i>	Pointer to a value to test against. This will always be set to "true" at the end of a call to TestAndSet.
---------------	---

Returns

true - Lock value was "true" on entry, false - Lock was set

Definition at line 26 of file [atomic.cpp](#).

Chapter 19

Class Documentation

19.1 Mark3::AutoAlloc Class Reference

The [AutoAlloc](#) class. This class provides an object-allocation interface for both kernel objects and user-defined types. This class supplies callouts for alloc/free that use object-type metadata to determine how objects may be allocated, allowing a user to create custom dynamic memory implementations for specific object types and sizes. As a result, the user-defined allocators can avoid the kinds of memory fragmentation and exhaustion issues that occur in typical embedded systems in which a single heap is used to satisfy all allocations in the application.

```
#include <autoalloc.h>
```

Static Public Member Functions

- static void [Init](#) (void)
Init Initialize the AutoAllocator before use. Called by [Kernel::Init\(\)](#).
- static void [SetAllocatorFunctions](#) ([AutoAllocAllocator_t](#) pfAllocator_, [AutoAllocFree_t](#) pfFree_)
SetAllocatorFunctions Set the functions used by this class to allocate/free memory used in the kernel.
- template<typename T, AutoAllocType e>
static T * [NewObject](#) ()
- template<typename T, AutoAllocType e>
static void [DestroyObject](#) (T *pObj_)
- static void * [NewUserTypeAllocation](#) (uint8_t eUserType_)
NewUserTypeAllocation Attempt to allocate a user-defined object type from the heap.
- static void [DestroyUserTypeAllocation](#) (uint8_t eUserType_, void *pvObj_)
DestroyUserTypeAllocation Free a previously allocated user-defined object.
- static void * [NewRawData](#) (size_t sSize_)
NewRawData Attempt to allocate a blob of raw data from the heap.
- static void [DestroyRawData](#) (void *pvData_)
DestroyRawData Free a previously allocated blob of data allocated via [NewRawData\(\)](#)

Static Private Member Functions

- static void * [Allocate](#) ([AutoAllocType](#) eType_, size_t sSize_)
- static void [Free](#) ([AutoAllocType](#) eType_, void *pvObj_)

Static Private Attributes

- static [AutoAllocAllocator_t m_pfAllocator](#)
Function used to allocate objects.
- static [AutoAllocFree_t m_pfFree](#)
Function used to free objects.

19.1.1 Detailed Description

The [AutoAlloc](#) class. This class provides an object-allocation interface for both kernel objects and user-defined types. This class supplies callouts for alloc/free that use object-type metadata to determine how objects may be allocated, allowing a user to create custom dynamic memory implementations for specific object types and sizes. As a result, the user-defined allocators can avoid the kinds of memory fragmentation and exhaustion issues that occur in typical embedded systems in which a single heap is used to satisfy all allocations in the application.

Definition at line 82 of file [autoalloc.h](#).

19.1.2 Member Function Documentation

19.1.2.1 Allocate()

```
void * Mark3::AutoAlloc::Allocate (
    AutoAllocType eType_,
    size_t sSize_ ) [static], [private]
```

Definition at line 58 of file [autoalloc.cpp](#).

19.1.2.2 DestroyObject()

```
template<typename T , AutoAllocType e>
static void Mark3::AutoAlloc::DestroyObject (
    T * pObj_ ) [inline], [static]
```

Template function used to manage the destruction and de-allocation of predefined kernel objects

Definition at line 116 of file [autoalloc.h](#).

19.1.2.3 DestroyRawData()

```
void Mark3::AutoAlloc::DestroyRawData (
    void * pvData_ ) [static]
```

DestroyRawData Free a previously allocated blob of data allocated via [NewRawData\(\)](#)

Parameters

<i>pv</i> ↔ <i>Data_</i>	pointer to previously-created data object
-----------------------------	---

Definition at line 105 of file [autoalloc.cpp](#).

19.1.2.4 DestroyUserTypeAllocation()

```
void Mark3::AutoAlloc::DestroyUserTypeAllocation (
    uint8_t eUserType_,
    void * pvObj_ ) [static]
```

DestroyUserTypeAllocation Free a previously allocated user-defined object.

Parameters

<i>pvObj_</i>	Pointer to previously-allocated object, allocated through NewUserTypeAllocation ()
<i>eUser</i> ↔ <i>Type_</i>	User defined object type, interpreted by the allocator function

Definition at line 95 of file [autoalloc.cpp](#).

19.1.2.5 Free()

```
void Mark3::AutoAlloc::Free (
    AutoAllocType eType_,
    void * pvObj_ ) [static], [private]
```

Definition at line 67 of file [autoalloc.cpp](#).

19.1.2.6 Init()

```
void Mark3::AutoAlloc::Init (
    void ) [static]
```

Init Initialize the AutoAllocator before use. Called by [Kernel::Init\(\)](#).

Definition at line 83 of file [autoalloc.cpp](#).

19.1.2.7 NewObject()

```
template<typename T , AutoAllocType e>
static T* Mark3::AutoAlloc::NewObject ( ) [inline], [static]
```

Template function used to manage the allocation of predefined kernel object types

Definition at line 103 of file [autoalloc.h](#).

19.1.2.8 NewRawData()

```
void * Mark3::AutoAlloc::NewRawData (
    size_t sSize_ ) [static]
```

NewRawData Attempt to allocate a blob of raw data from the heap.

Parameters

<i>sSize_</i>	Size of the data blob (in bytes)
---------------	----------------------------------

Returns

pointer to newly-allocated blob of data, or nullptr on error.

Definition at line 100 of file [autoalloc.cpp](#).

19.1.2.9 NewUserTypeAllocation()

```
void * Mark3::AutoAlloc::NewUserTypeAllocation (
    uint8_t eUserType_ ) [static]
```

NewUserTypeAllocation Attempt to allocate a user-defined object type from the heap.

Parameters

<i>eUserType_</i>	User defined object type, interpreted by the allocator function
-------------------	---

Returns

pointer to a newly-created object, or nullptr on error.

Definition at line 90 of file [autoalloc.cpp](#).

19.1.2.10 SetAllocatorFunctions()

```
void Mark3::AutoAlloc::SetAllocatorFunctions (
    AutoAllocAllocator_t pfAllocator_,
    AutoAllocFree_t pfFree_ ) [static]
```

SetAllocatorFunctions Set the functions used by this class to allocate/free memory used in the kernel.

Parameters

<i>pf</i> ↔ <i>Allocator</i> ↔ _	Function to allocate an object based on its type and/or size
<i>pfFree</i> _	Function to free a previously-allocated object

Definition at line 76 of file [autoalloc.cpp](#).

19.1.3 Member Data Documentation

19.1.3.1 m_pfAllocator

```
AutoAllocAllocator_t Mark3::AutoAlloc::m_pfAllocator [static], [private]
```

Function used to allocate objects.

Definition at line 157 of file [autoalloc.h](#).

19.1.3.2 m_pfFree

```
AutoAllocFree_t Mark3::AutoAlloc::m_pfFree [static], [private]
```

Function used to free objects.

Function used to free objects.

Definition at line 158 of file [autoalloc.h](#).

The documentation for this class was generated from the following files:


- [/home/moslevin/projects/m3-repo/kernel/src/public/autoalloc.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/autoalloc.cpp](#)

19.2 Mark3::BlockingObject Class Reference

The [BlockingObject](#) class. Class implementing thread-blocking primitives. used for implementing things like semaphores, mutexes, message queues, or anything else that could cause a thread to suspend execution on some external stimulus.

```
#include <blocking.h>
```

Inheritance diagram for Mark3::BlockingObject:



Public Member Functions

- [BlockingObject](#) ()
- [~BlockingObject](#) ()

Protected Member Functions

- void [Block](#) ([Thread](#) *pclThread_)
- Block Blocks a thread on this object. This is the fundamental operation performed by any sort of blocking operation in the operating system. All semaphores/mutexes/sleeping/messaging/etc ends up going through the blocking code at some point as part of the code that manages a transition from an "active" or "waiting" thread to a "blocked" thread.*
- void [BlockPriority](#) ([Thread](#) *pclThread_)
- BlockPriority Same as [Block\(\)](#), but ensures that threads are added to the block-list in priority-order, which optimizes the unblock procedure.*
- void [UnBlock](#) ([Thread](#) *pclThread_)
- UnBlock Unblock a thread that is already blocked on this object, returning it to the "ready" state by performing the following steps:*
- void [SetInitialized](#) (void)
- SetInitialized.*
- bool [IsInitialized](#) (void)
- IsInitialized.*

Protected Attributes

- [ThreadList](#) [m_clBlockList](#)
- [uint8_t](#) [m_u8Initialized](#)

Static Protected Attributes

- static constexpr auto [m_uBlockingInvalidCookie](#) = [uint8_t](#) { 0x3C }
- static constexpr auto [m_uBlockingInitCookie](#) = [uint8_t](#) { 0xC3 }

19.2.1 Detailed Description

The [BlockingObject](#) class. Class implementing thread-blocking primitives. used for implementing things like semaphores, mutexes, message queues, or anything else that could cause a thread to suspend execution on some external stimulus.

Definition at line 65 of file [blocking.h](#).

19.2.2 Constructor & Destructor Documentation

19.2.2.1 BlockingObject()

```
Mark3::BlockingObject::BlockingObject ( ) [inline]
```

Definition at line 68 of file [blocking.h](#).

19.2.2.2 ~BlockingObject()

```
Mark3::BlockingObject::~~BlockingObject ( ) [inline]
```

Definition at line 69 of file [blocking.h](#).

19.2.3 Member Function Documentation

19.2.3.1 Block()

```
void Mark3::BlockingObject::Block (
    Thread * pclThread_ ) [protected]
```

Block Blocks a thread on this object. This is the fundamental operation performed by any sort of blocking operation in the operating system. All semaphores/mutexes/sleeping/messaging/etc ends up going through the blocking code at some point as part of the code that manages a transition from an "active" or "waiting" thread to a "blocked" thread.

The steps involved in blocking a thread (which are performed in the function itself) are as follows;

- 1) Remove the specified thread from the current owner's list (which is likely one of the scheduler's thread lists)
- 2) Add the thread to this object's thread list
- 3) Setting the thread's "current thread-list" point to reference this object's threadlist.

Parameters

<i>pcl↔ Thread_</i>	Pointer to the thread object that will be blocked.
-------------------------	--

Definition at line 26 of file [blocking.cpp](#).

19.2.3.2 BlockPriority()

```
void Mark3::BlockingObject::BlockPriority (
    Thread * pclThread_ ) [protected]
```

BlockPriority Same as [Block\(\)](#), but ensures that threads are added to the block-list in priority-order, which optimizes the unblock procedure.

Parameters

<i>pcl↔ Thread_</i>	Pointer to the Thread to Block.
-------------------------	---

Definition at line 41 of file [blocking.cpp](#).

19.2.3.3 IsInitialized()

```
bool Mark3::BlockingObject::IsInitialized (
    void ) [inline], [protected]
```

IsInitialized.

Returns

true if initialized, false if object uninitialized

Definition at line 123 of file [blocking.h](#).

19.2.3.4 SetInitialized()

```
void Mark3::BlockingObject::SetInitialized (
    void ) [inline], [protected]
```

SetInitialized.

Definition at line 117 of file [blocking.h](#).

19.2.3.5 Unblock()

```
void Mark3::BlockingObject::UnBlock (
    Thread * pclThread_ ) [protected]
```

UnBlock Unblock a thread that is already blocked on this object, returning it to the "ready" state by performing the following steps:

Parameters

<i>pthread_t</i> <i>Thread_</i>	Pointer to the thread to unblock.
------------------------------------	-----------------------------------

1) Removing the thread from this object's threadlist 2) Restoring the thread to its "original" owner's list

Definition at line 56 of file [blocking.cpp](#).

19.2.4 Member Data Documentation

19.2.4.1 m_clBlockList

```
ThreadList Mark3::BlockingObject::m_clBlockList [protected]
```

[ThreadList](#) which is used to hold the list of threads blocked on a given object.

Definition at line 133 of file [blocking.h](#).

19.2.4.2 m_u8Initialized

```
uint8_t Mark3::BlockingObject::m_u8Initialized [protected]
```

Token used to check whether or not the object has been initialized prior to use.

Definition at line 139 of file [blocking.h](#).

19.2.4.3 m_uBlockingInitCookie

```
constexpr auto Mark3::BlockingObject::m_uBlockingInitCookie = uint8_t { 0xC3 } [static],  
[protected]
```

Definition at line 127 of file [blocking.h](#).

19.2.4.4 m_uBlockingInvalidCookie

```
constexpr auto Mark3::BlockingObject::m_uBlockingInvalidCookie = uint8_t { 0xC3 } [static],  
[protected]
```

Definition at line 126 of file [blocking.h](#).

The documentation for this class was generated from the following files:


- [/home/moslevin/projects/m3-repo/kernel/src/public/blocking.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/blocking.cpp](#)

19.3 Mark3::CircularLinkedList Class Reference

The [CircularLinkedList](#) class Circular-linked-list data type, inherited from the base [LinkedList](#) type.

```
#include <ll.h>
```

Inheritance diagram for Mark3::CircularLinkedList:



Public Member Functions

- `void * operator new (size_t sz, void *pv)`
- `CircularLinkedList ()`
- `void Add (LinkedListNode *node_)`
Add the linked list node to this linked list.
- `void Remove (LinkedListNode *node_)`
Remove Add the linked list node to this linked list.
- `void PivotForward ()`
PivotForward Pivot the head of the circularly linked list forward (Head = Head->next, Tail = Tail->next)
- `void PivotBackward ()`
PivotBackward Pivot the head of the circularly linked list backward (Head = Head->prev, Tail = Tail->prev)
- `void InsertNodeBefore (LinkedListNode *node_, LinkedListNode *insert_)`
InsertNodeBefore Insert a linked-list node into the list before the specified insertion point.

Additional Inherited Members

19.3.1 Detailed Description

The [CircularLinkedList](#) class Circular-linked-list data type, inherited from the base [LinkedList](#) type.

Definition at line 218 of file [ll.h](#).

19.3.2 Constructor & Destructor Documentation

19.3.2.1 CircularLinkedList()

```
Mark3::CircularLinkedList::CircularLinkedList ( ) [inline]
```

Definition at line 222 of file [ll.h](#).

19.3.3 Member Function Documentation

19.3.3.1 Add()

```
void Mark3::CircularLinkedList::Add (
    LinkListNode * node_ )
```

Add the linked list node to this linked list.

Parameters

<i>node_</i>	Pointer to the node to add
—	

Definition at line 81 of file [ll.cpp](#).

19.3.3.2 InsertNodeBefore()

```
void Mark3::CircularLinkedList::InsertNodeBefore (
    LinkListNode * node_,
    LinkListNode * insert_ )
```

InsertNodeBefore Insert a linked-list node into the list before the specified insertion point.

Parameters

<i>node_</i>	Node to insert into the list
—	
<i>insert_</i>	Insert point.
—	

Definition at line 153 of file [ll.cpp](#).

19.3.3.3 operator new()

```
void* Mark3::CircularLinkedList::operator new (
    size_t sz,
    void * pv ) [inline]
```

Definition at line 221 of file [ll.h](#).

19.3.3.4 PivotBackward()

```
void Mark3::CircularLinkedList::PivotBackward ( )
```

PivotBackward Pivot the head of the circularly linked list backward (Head = Head->prev, Tail = Tail->prev)

Definition at line 144 of file [ll.cpp](#).

19.3.3.5 PivotForward()

```
void Mark3::CircularLinkedList::PivotForward ( )
```

PivotForward Pivot the head of the circularly linked list forward (Head = Head->next, Tail = Tail->next)

Definition at line 135 of file [ll.cpp](#).

19.3.3.6 Remove()

```
void Mark3::CircularLinkedList::Remove (
    LinkListNode * node_ )
```

Remove Add the linked list node to this linked list.

Parameters

<i>node_</i>	Pointer to the node to remove
—	

Definition at line 103 of file [ll.cpp](#).

The documentation for this class was generated from the following files:

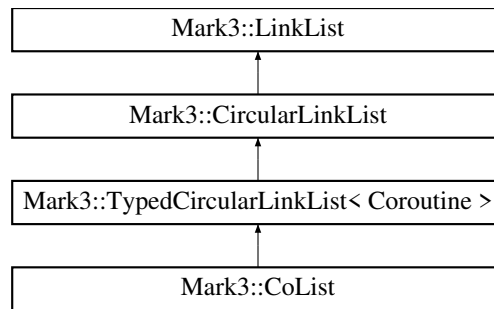
- [/home/moslevin/projects/m3-repo/kernel/src/public/ll.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/ll.cpp](#)

19.4 Mark3::CoList Class Reference

The [CoList](#) class The [CoList](#) class implements a circular-linked-listed structure for coroutine objects. The intent of this object is to maintain a list of active coroutine objects with a specific priority or state, to ensure that a freshly-schedulable co-routine always exists at the head of the list.

```
#include <colist.h>
```

Inheritance diagram for Mark3::CoList:



Public Member Functions

- void `SetPrioMap` (`CoPrioMap` *pclPrioMap_)
SetPrioMap Assign a priority map object to this co-routine list.
- void `SetPriority` (`PORT_PRIO_TYPE` uPriority_)
SetPriority Set the scheduling priority of this coroutine list; has no effect unless a SetPrioMap has been called with a valid coroutine priority map object.
- void `Add` (`Coroutine` *pclCoroutine_)
Add Add a coroutine object to this list.
- void `Remove` (`Coroutine` *pclCoroutine_)
Remove Remove a given coroutine object from this list.

Private Attributes

- `CoPrioMap` * m_pclPrioMap
- `uint8_t` m_uPriority

Additional Inherited Members

19.4.1 Detailed Description

The `CoList` class The `CoList` class implements a circular-linked-list structure for coroutine objects. The intent of this object is to maintain a list of active coroutine objects with a specific priority or state, to ensure that a freshly-schedulable co-routine always exists at the head of the list.

Definition at line 35 of file `colist.h`.

19.4.2 Member Function Documentation

19.4.2.1 Add()

```
void Mark3::CoList::Add (
    Coroutine * pclCoroutine_ )
```

Add Add a coroutine object to this list.

Parameters

<i>pcl</i> ↔ <i>Coroutine_</i>	Pointer to the coroutine object to add
-----------------------------------	--

Definition at line 36 of file [colist.cpp](#).

19.4.2.2 Remove()

```
void Mark3::CoList::Remove (
    Coroutine * pclCoroutine_ )
```

Remove Remove a given coroutine object from this list.

Parameters

<i>pcl</i> ↔ <i>Coroutine_</i>	Pointer to the coroutine object to remove
-----------------------------------	---

Definition at line 46 of file [colist.cpp](#).

19.4.2.3 SetPrioMap()

```
void Mark3::CoList::SetPrioMap (
    CoPrioMap * pclPrioMap_ )
```

SetPrioMap Assign a priority map object to this co-routine list.

Parameters

<i>pclPrio</i> ↔ <i>Map_</i>	priority map object to assign
---------------------------------	-------------------------------

Definition at line 24 of file [colist.cpp](#).

19.4.2.4 SetPriority()

```
void Mark3::CoList::SetPriority (
    PORT_PRIO_TYPE uPriority_ )
```

SetPriority Set the scheduling priority of this coroutine liss; has no effect unless a SetPrioMap has been called with a valid coroutine priority map object.

Parameters

<i>u</i> ↔ <i>Priority</i> ↔ —	Priority of coroutines associated with this list
--------------------------------------	--

Definition at line 30 of file [colist.cpp](#).

19.4.3 Member Data Documentation

19.4.3.1 m_pclPrioMap

```
CoPrioMap* Mark3::CoList::m_pclPrioMap [private]
```

Definition at line 73 of file [colist.h](#).

19.4.3.2 m_uPriority

```
uint8_t Mark3::CoList::m_uPriority [private]
```

Definition at line 74 of file [colist.h](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/m3-repo/kernel/src/public/colist.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/colist.cpp](#)

19.5 Mark3::ConditionVariable Class Reference

The [ConditionVariable](#) class This class implements a condition variable. This is a synchronization object that allows multiple threads to block, each waiting for specific signals unique to them. Access to the specified condition is guarded by a mutex that is supplied by the caller. This object can permit multiple waiters that can be unblocked one-at-a-time via signalling, or unblocked all at once via broadcasting. This object is built upon lower-level primitives, and is somewhat more heavyweight than the primitive types supplied by the kernel.

```
#include <condvar.h>
```

Public Member Functions

- void * [operator new](#) (size_t sz, void *pv)
- void [Init](#) ()

Init Initialize the condition variable prior to use. Must be called before the object can be used.
- void [Wait](#) ([Mutex](#) *pclMutex_)

Wait Block the current thread, and wait for the object to be signalled. The specified mutex will be locked when the thread returns.
- bool [Wait](#) ([Mutex](#) *pclMutex_, uint32_t u32WaitTimeMS_)

Wait Block the current thread, and wait for the object to be signalled. The specified mutex will be locked when the thread returns.
- void [Signal](#) ()

Signal Signal/Unblock the next thread currently blocked on this condition variable.
- void [Broadcast](#) ()

Broadcast Unblock all threads currently blocked on this condition variable.

Private Attributes

- [Mutex](#) m_clMutex
- [Semaphore](#) m_clSemaphore
- uint8_t m_u8Waiters

19.5.1 Detailed Description

The [ConditionVariable](#) class This class implements a condition variable. This is a synchronization object that allows multiple threads to block, each waiting for specific signals unique to them. Access to the specified condition is guarded by a mutex that is supplied by the caller. This object can permit multiple waiters that can be unblocked one-at-a-time via signalling, or unblocked all at once via broadcasting. This object is built upon lower-level primitives, and is somewhat more heavyweight than the primitive types supplied by the kernel.

Definition at line 39 of file [condvar.h](#).

19.5.2 Member Function Documentation

19.5.2.1 Broadcast()

```
void Mark3::ConditionVariable::Broadcast ( )
```

Broadcast Unblock all threads currently blocked on this condition variable.

Definition at line 77 of file [condvar.cpp](#).

19.5.2.2 Init()

```
void Mark3::ConditionVariable::Init (
    void )
```

Init Initialize the condition variable prior to use. Must be called before the object can be used.

Definition at line 25 of file [condvar.cpp](#).

19.5.2.3 operator new()

```
void* Mark3::ConditionVariable::operator new (
    size_t sz,
    void * pv ) [inline]
```

Definition at line 42 of file [condvar.h](#).

19.5.2.4 Signal()

```
void Mark3::ConditionVariable::Signal ( )
```

Signal Signal/Unblock the next thread currently blocked on this condition variable.

Definition at line 66 of file [condvar.cpp](#).

19.5.2.5 Wait() [1/2]

```
void Mark3::ConditionVariable::Wait (
    Mutex * pclMutex_ )
```

Wait Block the current thread, and wait for the object to be signalled. The specified mutex will be locked when the thread returns.

Parameters

<i>pcl</i> ↔ <i>Mutex_</i>	Mutex to claim once the calling thread has access to the condvar
-------------------------------	--

Definition at line 32 of file [condvar.cpp](#).

19.5.2.6 Wait() [2/2]

```
bool Mark3::ConditionVariable::Wait (
    Mutex * pclMutex_,
    uint32_t u32WaitTimeMS_ )
```

Wait Block the current thread, and wait for the object to be signalled. The specified mutex will be locked when the thread returns.

Parameters

<i>pclMutex_</i>	Mutex to claim once the calling thread has access to the condvar
<i>u32WaitTimeMS_</i>	Maximum time in ms to wait before abandoning the operation

Returns

true on success, false on timeout

Definition at line 48 of file [condvar.cpp](#).

19.5.3 Member Data Documentation

19.5.3.1 m_clMutex

```
Mutex Mark3::ConditionVariable::m_clMutex [private]
```

Definition at line 82 of file [condvar.h](#).

19.5.3.2 m_clSemaphore

```
Semaphore Mark3::ConditionVariable::m_clSemaphore [private]
```

Definition at line 83 of file [condvar.h](#).

19.5.3.3 m_u8Waiters

```
uint8_t Mark3::ConditionVariable::m_u8Waiters [private]
```

Definition at line 84 of file [condvar.h](#).

The documentation for this class was generated from the following files:


- [/home/moslevin/projects/m3-repo/kernel/src/public/condvar.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/condvar.cpp](#)

19.6 Mark3::Coroutine Class Reference

The [Coroutine](#) class implements a lightweight, run-to-completion task that forms the basis for co-operative task scheduling in [Mark3](#). Coroutines are designed to be run from a singular context, and scheduled as a result of events occurring from threads, timers, interrupt sources, or other co-routines.

```
#include <coroutine.h>
```

Inheritance diagram for Mark3::Coroutine:



Public Member Functions

- [~Coroutine](#) ()
- void [Init](#) ([PORT_PRIO_TYPE](#) uPriority_, [CoroutineHandler](#) pfHandler_, void *pvContext_)
Init Initialize the coroutine object prior to use. Must be called before using the other methods in the class.
- void [Run](#) ()
Run Clear the co-routine's pending execution flag, and execute the coroutine's handler function.
- void [Activate](#) ()
Activate Tag the co-routine as pending execution. Has no effect if the co-routine is already pending execution.
- void [SetPriority](#) ([PORT_PRIO_TYPE](#) uPriority_)
SetPriority Update the scheduling priority of the co-routine. Can be called from within the co-routine, or from any other context aware of the co-routine object.
- [PORT_PRIO_TYPE](#) [GetPriority](#) ()
GetPriority Retrieve the current scheduling priority of the co-routine.

Private Attributes

- [CoList](#) * [m_pclOwner](#)
- [CoroutineHandler](#) [m_pfHandler](#)
- void * [m_pvContext](#)
- [PORT_PRIO_TYPE](#) [m_uPriority](#)
- bool [m_bQueued](#)

Additional Inherited Members

19.6.1 Detailed Description

The [Coroutine](#) class implements a lightweight, run-to-completion task that forms the basis for co-operative task scheduling in [Mark3](#). Coroutines are designed to be run from a singular context, and scheduled as a result of events occurring from threads, timers, interrupt sources, or other co-routines.

Co-routines differ from Threads in that they cannot block, and must run to completion before other (potentially higher-priority) co-routines block.

Examples:

[lab2_coroutines/main.cpp](#).

Definition at line 53 of file [coroutine.h](#).

19.6.2 Constructor & Destructor Documentation

19.6.2.1 ~Coroutine()

```
Mark3::Coroutine::~~Coroutine ( )
```

Definition at line 28 of file [coroutine.cpp](#).

19.6.3 Member Function Documentation

19.6.3.1 Activate()

```
void Mark3::Coroutine::Activate ( )
```

Activate Tag the co-routine as pending execution. Has no effect if the co-routine is already pending execution.

Examples:

[lab2_coroutines/main.cpp](#).

Definition at line 67 of file [coroutine.cpp](#).

19.6.3.2 GetPriority()

```
PORT_PRIO_TYPE Mark3::Coroutine::GetPriority ( )
```

GetPriority Retrieve the current scheduling priority of the co-routine.

Returns

current scheduling priority of the co-routine.

Definition at line 97 of file [coroutine.cpp](#).

19.6.3.3 Init()

```
void Mark3::Coroutine::Init (
    PORT_PRIO_TYPE uPriority_,
    CoroutineHandler pfHandler_,
    void * pvContext_ )
```

Init Initialize the coroutine object prior to use. Must be called before using the other methods in the class.

Parameters

<i>uPriority</i> ↔ —	The scheduling priority of this coroutine as configured by the port.
<i>pf</i> ↔ <i>Handler</i> ↔ —	Function to run when the coroutine is executed
<i>pv</i> ↔ <i>Context</i> —	User-defined value passed into the handler function on execution

Definition at line 39 of file [coroutine.cpp](#).

19.6.3.4 Run()

```
void Mark3::Coroutine::Run ( )
```

Run Clear the co-routine's pending execution flag, and execute the coroutine's handler function.

Definition at line 53 of file [coroutine.cpp](#).

19.6.3.5 SetPriority()

```
void Mark3::Coroutine::SetPriority (
    PORT_PRIO_TYPE uPriority_ )
```

SetPriority Update the scheduling priority of the co-routine. Can be called from within the co-routine, or from any other context aware of the co-routine object.

Parameters

<i>u</i> ↔ <i>Priority</i> ↔ —	New priority of the co-routine
--------------------------------------	--------------------------------

Definition at line 82 of file [coroutine.cpp](#).

19.6.4 Member Data Documentation

19.6.4.1 m_bQueued

```
bool Mark3::Coroutine::m_bQueued [private]
```

Definition at line 108 of file [coroutine.h](#).

19.6.4.2 m_pclOwner

```
CoList* Mark3::Coroutine::m_pclOwner [private]
```

Definition at line 104 of file [coroutine.h](#).

19.6.4.3 m_pfHandler

```
CoroutineHandler Mark3::Coroutine::m_pfHandler [private]
```

Definition at line 105 of file [coroutine.h](#).

19.6.4.4 m_pvContext

```
void* Mark3::Coroutine::m_pvContext [private]
```

Definition at line 106 of file [coroutine.h](#).

19.6.4.5 m_uPriority

```
PORT_PRIO_TYPE Mark3::Coroutine::m_uPriority [private]
```

Definition at line 107 of file [coroutine.h](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/m3-repo/kernel/src/public/coroutine.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/coroutine.cpp](#)

19.7 Mark3::CoScheduler Class Reference

The [CoScheduler](#) class. This class implements the coroutine scheduler. Similar to the [Mark3](#) thread scheduler, the highest-priority active object is scheduled / returned for execution. If no active co-routines are available to be scheduled, then the scheduler returns nullptr.

```
#include <cosched.h>
```

Static Public Member Functions

- static void [Init](#) ()
Init Initialize the coroutine scheduler prior to use. Must be called prior to using any other functions in the coroutine scheduler.
- static [CoPrioMap](#) * [GetPrioMap](#) ()
GetPrioMap Get the pointer to the priority map object used by the scheduler.
- static [CoList](#) * [GetStopList](#) ()
GetStopList Get the pointer to the coroutine list managing initialized coroutines that are not awaiting execution.
- static [CoList](#) * [GetCoList](#) ([PORT_PRIO_TYPE](#) uPriority_)
GetCoList Retrieve the coroutine list associated with a given priority.
- static [Coroutine](#) * [Schedule](#) ()
Schedule Determine what coroutine (if any) is the next to be executed.

Static Private Attributes

- static [CoList](#) m_aclPriorities [[PORT_COROUTINE_PRIORITIES](#)]
- static [CoList](#) m_clStopList
- static [CoPrioMap](#) m_clPrioMap

19.7.1 Detailed Description

The [CoScheduler](#) class. This class implements the coroutine scheduler. Similar to the [Mark3](#) thread scheduler, the highest-priority active object is scheduled / returned for execution. If no active co-routines are available to be scheduled, then the scheduler returns nullptr.

Definition at line 33 of file [cosched.h](#).

19.7.2 Member Function Documentation

19.7.2.1 [GetCoList\(\)](#)

```
CoList * Mark3::CoScheduler::GetCoList (  
    PORT\_PRIO\_TYPE uPriority_ ) [static]
```

[GetCoList](#) Retrieve the coroutine list associated with a given priority.

Parameters

$u \leftrightarrow$ <i>Priority</i> \leftrightarrow —	Priority to get the coroutine list of.
---	--

Returns

coroutine list pointer or nullptr on invalid priority.

Definition at line 51 of file [cosched.cpp](#).

19.7.2.2 GetPrioMap()

```
CoPrioMap * Mark3::CoScheduler::GetPrioMap ( ) [static]
```

GetPrioMap Get the pointer to the priority map object used by the scheduler.

Returns

Return the priority map object owned by the schedule

Definition at line 39 of file [cosched.cpp](#).

19.7.2.3 GetStopList()

```
CoList * Mark3::CoScheduler::GetStopList ( ) [static]
```

GetStopList Get the pointer to the coroutine list managing initialized coroutines that are not awaiting execution.

Returns

Pointer to the coroutine stop list

Definition at line 45 of file [cosched.cpp](#).

19.7.2.4 Init()

```
void Mark3::CoScheduler::Init (
    void ) [static]
```

Init Initialize the coroutine scheduler prior to use. Must be called prior to using any other functions in the coroutine scheduler.

Examples:

[lab2_coroutines/main.cpp](#).

Definition at line 29 of file [cosched.cpp](#).

19.7.2.5 Schedule()

```
Coroutine * Mark3::CoScheduler::Schedule ( ) [static]
```

Schedule Determine what coroutine (if any) is the next to be executed.

Returns

next coroutine to execute, or nullptr if no coroutines are ready to be scheduled.

Examples:

[lab2_coroutines/main.cpp](#).

Definition at line 60 of file [cosched.cpp](#).

19.7.3 Member Data Documentation

19.7.3.1 m_aclPriorities

```
CoList Mark3::CoScheduler::m_aclPriorities [static], [private]
```

Definition at line 80 of file [cosched.h](#).

19.7.3.2 m_clPrioMap

```
CoPrioMap Mark3::CoScheduler::m_clPrioMap [static], [private]
```

Definition at line 82 of file [cosched.h](#).

19.7.3.3 m_clStopList

```
CoList Mark3::CoScheduler::m_clStopList [static], [private]
```

Definition at line 81 of file [cosched.h](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/m3-repo/kernel/src/public/cosched.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/cosched.cpp](#)

19.8 Mark3::CriticalGuard Class Reference

The [CriticalGuard](#) class. This class provides an implementation of RAII for critical sections. Object creation results in a critical section being invoked. The subsequent destructor call results in the critical section being released.

```
#include <criticalguard.h>
```

Public Member Functions

- [CriticalGuard](#) ()
- [~CriticalGuard](#) ()

19.8.1 Detailed Description

The [CriticalGuard](#) class. This class provides an implementation of RAII for critical sections. Object creation results in a critical section being invoked. The subsequent destructor call results in the critical section being released.

This is similar to the [LockGuard](#) class, except that class operates on a single [Mutex](#), and this class operates on a global interrupt-disabled lock.

Definition at line 38 of file [criticalguard.h](#).

19.8.2 Constructor & Destructor Documentation

19.8.2.1 CriticalGuard()

```
Mark3::CriticalGuard::CriticalGuard ( ) [inline]
```

Definition at line 40 of file [criticalguard.h](#).

19.8.2.2 ~CriticalGuard()

```
Mark3::CriticalGuard::~~CriticalGuard ( ) [inline]
```

Definition at line 44 of file [criticalguard.h](#).

The documentation for this class was generated from the following file:

- [/home/moslevin/projects/m3-repo/kernel/src/public/criticalguard.h](#)

19.9 Mark3::CriticalSection Class Reference

The [CriticalSection](#) class. This class implements a portable [CriticalSection](#) interface based on macros/inline functions that are implemented as part of each port.

```
#include <criticalsection.h>
```

Static Public Member Functions

- static void [Enter](#) ()
Enter Enter a critical section, disabling all kernel-aware interrupts, and giving exclusive control of the CPU to the curreninttly running task.
- static void [Exit](#) ()
Exit Exit a critical section, re-enabling kernel-aware interrupts, and releasing exclusive control of the CPU.
- static [K_WORD](#) [NestingCount](#) ()
NestingCount.

19.9.1 Detailed Description

The [CriticalSection](#) class. This class implements a portable [CriticalSection](#) interface based on macros/inline functions that are implemented as part of each port.

Critical sections *can* be safely nested, that is, multiple calls to [Enter\(\)](#) may be called before matching calls to [Exit\(\)](#) are called. In such cases, only the *final* call to [Exit\(\)](#) will cause the caller to relinquish control of the CPU.

Critical sections can be safely used within interrupts; although they have no effect in the general case. The exception is in the case where a system has multiple levels of interrupt nesting; at which point calling a critical section from an interrupt temporarily disables subsequent levels of nesting.

Care, however, must be taken to ensure that the currently executing thread does *not* block when a critical section is active. This condition is guaranteed to break the system.

Definition at line 48 of file [criticalsection.h](#).

19.9.2 Member Function Documentation

19.9.2.1 Enter()

```
static void Mark3::CriticalSection::Enter ( ) [inline], [static]
```

Enter Enter a critical section, disabling all kernel-aware interrupts, and giving exclusive control of the CPU to the curreninttly running task.

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 56 of file [criticalsection.h](#).

19.9.2.2 Exit()

```
static void Mark3::CriticalSection::Exit ( ) [inline], [static]
```

Exit Exit a critical section, re-enabling kernel-aware interrupts, and releasing exclusive control of the CPU.

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 65 of file [criticalsection.h](#).

19.9.2.3 NestingCount()

```
static K_WORD Mark3::CriticalSection::NestingCount ( ) [inline], [static]
```

NestingCount.

Returns

Number of [Enter\(\)](#) calls awaiting an [Exit\(\)](#) call.

Definition at line 73 of file [criticalsection.h](#).

The documentation for this class was generated from the following file:


- [/home/moslevin/projects/m3-repo/kernel/src/public/criticalsection.h](#)

19.10 Mark3::DoubleLinkedList Class Reference

The [DoubleLinkedList](#) Class Doubly-linked-list data type, inherited from the base [LinkedList](#) type.

```
#include <ll.h>
```

Inheritance diagram for Mark3::DoubleLinkedList:



Public Member Functions

- void * [operator new](#) (size_t sz, void *pv)
- [DoubleLinkedList](#) ()
DoubleLinkedList.
- void [Add](#) ([LinkedListNode](#) *node_)
Add.
- void [Remove](#) ([LinkedListNode](#) *node_)
Remove.

Additional Inherited Members

19.10.1 Detailed Description

The [DoubleLinkedList](#) Class Doubly-linked-list data type, inherited from the base [LinkedList](#) type.

Definition at line [179](#) of file [ll.h](#).

19.10.2 Constructor & Destructor Documentation

19.10.2.1 DoubleLinkedList()

```
Mark3::DoubleLinkedList::DoubleLinkedList ( ) [inline]
```

[DoubleLinkedList](#).

Default constructor - initializes the head/tail nodes to nullptr

Definition at line [188](#) of file [ll.h](#).

19.10.3 Member Function Documentation

19.10.3.1 Add()

```
void Mark3::DoubleLinkedList::Add (
    LinkedListNode * node_ )
```

Add.

Add the linked list node to this linked list

Parameters

<i>node_</i>	Pointer to the node to add
—	

Definition at line [34](#) of file [ll.cpp](#).

19.10.3.2 operator new()

```
void* Mark3::DoubleLinkedList::operator new (
```

```

    size_t sz,
    void * pv ) [inline]

```

Definition at line 182 of file [ll.h](#).

19.10.3.3 Remove()

```

void Mark3::DoubleLinkedList::Remove (
    LinkListNode * node_ )

```

Remove.

Add the linked list node to this linked list

Parameters

<i>node</i> ↔	Pointer to the node to remove
—	

Definition at line 55 of file [ll.cpp](#).

The documentation for this class was generated from the following files:


- [/home/moslevin/projects/m3-repo/kernel/src/public/ll.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/ll.cpp](#)

19.11 Mark3::EventFlag Class Reference

The [EventFlag](#) class. This class implements a blocking object, similar to a semaphore or mutex, commonly used for synchronizing thread execution based on events occurring within the system. Each [EventFlag](#) object contains a 16-bit bitmask, which is used to trigger events on associated threads. Threads wishing to block, waiting for a specific event to occur can wait on any pattern within this 16-bit bitmask to be set. Here, we provide the ability for a thread to block, waiting for ANY bits in a specified mask to be set, or for ALL bits within a specific mask to be set. Depending on how the object is configured, the bits that triggered the wakeup can be automatically cleared once a match has occurred.

```
#include <eventflag.h>
```

Inheritance diagram for Mark3::EventFlag:



Public Member Functions

- void * [operator new](#) (size_t sz, void *pv)
- [~EventFlag](#) ()
- void [Init](#) ()
Init Initializes the [EventFlag](#) object prior to use.
- uint16_t [Wait](#) (uint16_t u16Mask_, [EventFlagOperation](#) eMode_)
Wait Block a thread on the specific flags in this event flag group.
- uint16_t [Wait](#) (uint16_t u16Mask_, [EventFlagOperation](#) eMode_, uint32_t u32TimeMS_)
Wait Block a thread on the specific flags in this event flag group.
- void [WakeMe](#) ([Thread](#) *pclChosenOne_)
WakeMe Wake the given thread, currently blocking on this object.
- void [Set](#) (uint16_t u16Mask_)
Set Set additional flags in this object (logical OR). This API can potentially result in threads blocked on [Wait\(\)](#) to be unblocked.
- void [Clear](#) (uint16_t u16Mask_)
ClearFlags - Clear a specific set of flags within this object, specific by bitmask.
- uint16_t [GetMask](#) ()
GetMask Returns the state of the 16-bit bitmask within this object.

Private Member Functions

- uint16_t [Wait_i](#) (uint16_t u16Mask_, [EventFlagOperation](#) eMode_, uint32_t u32TimeMS_)
Wait_i Internal abstraction used to manage both timed and untimed wait operations.

Private Attributes

- uint16_t [m_u16SetMask](#)
Event flags currently set in this object.

Additional Inherited Members

19.11.1 Detailed Description

The [EventFlag](#) class. This class implements a blocking object, similar to a semaphore or mutex, commonly used for synchronizing thread execution based on events occurring within the system. Each [EventFlag](#) object contains a 16-bit bitmask, which is used to trigger events on associated threads. Threads wishing to block, waiting for a specific event to occur can wait on any pattern within this 16-bit bitmask to be set. Here, we provide the ability for a thread to block, waiting for ANY bits in a specified mask to be set, or for ALL bits within a specific mask to be set. Depending on how the object is configured, the bits that triggered the wakeup can be automatically cleared once a match has occurred.

Examples:

[lab7_events/main.cpp](#).

Definition at line 45 of file [eventflag.h](#).

19.11.2 Constructor & Destructor Documentation

19.11.2.1 ~EventFlag()

```
Mark3::EventFlag::~~EventFlag ( )
```

19.11.3 Member Function Documentation

19.11.3.1 Clear()

```
void Mark3::EventFlag::Clear (
    uint16_t u16Mask_ )
```

ClearFlags - Clear a specific set of flags within this object, specific by bitmask.

Parameters

<i>u16Mask_</i>	- Bitmask of flags to clear
-----------------	-----------------------------

Examples:

[lab7_events/main.cpp](#).

19.11.3.2 GetMask()

```
uint16_t Mark3::EventFlag::GetMask ( )
```

GetMask Returns the state of the 16-bit bitmask within this object.

Returns

The state of the 16-bit bitmask

19.11.3.3 Init()

```
void Mark3::EventFlag::Init ( )
```

Init Initializes the [EventFlag](#) object prior to use.

19.11.3.4 operator new()

```
void* Mark3::EventFlag::operator new (
    size_t sz,
    void * pv ) [inline]
```

Definition at line 48 of file [eventflag.h](#).

19.11.3.5 Set()

```
void Mark3::EventFlag::Set (
    uint16_t u16Mask_ )
```

Set Set additional flags in this object (logical OR). This API can potentially result in threads blocked on [Wait\(\)](#) to be unblocked.

Parameters

<i>u16↔ Mask_</i>	- Bitmask of flags to set.
-----------------------	----------------------------

Examples:

[lab7_events/main.cpp](#).

19.11.3.6 Wait() [1/2]

```
uint16_t Mark3::EventFlag::Wait (
    uint16_t u16Mask_,
    EventFlagOperation eMode_ )
```

Wait Block a thread on the specific flags in this event flag group.

Parameters

<i>u16↔ Mask_</i>	- 16-bit bitmask to block on
<i>eMode_</i>	- EventFlagOperation::Any_Set : Thread will block on any of the bits in the mask <ul style="list-style-type: none"> EventFlagOperation::All_Set: Thread will block on all of the bits in the mask

Returns

Bitmask condition that caused the thread to unblock, or 0 on error or timeout

Examples:

[lab7_events/main.cpp](#).

19.11.3.7 Wait() [2/2]

```
uint16_t Mark3::EventFlag::Wait (
    uint16_t u16Mask_,
    EventFlagOperation eMode_,
    uint32_t u32TimeMS_ )
```

Wait Block a thread on the specific flags in this event flag group.

Parameters

<i>u16Mask_</i>	- 16-bit bitmask to block on
<i>eMode_</i>	- EventFlagOperation::Any_Set : Thread will block on any of the bits in the mask <ul style="list-style-type: none"> EventFlagOperation::All_Set: Thread will block on all of the bits in the mask
<i>u32TimeMS_</i>	- Time to block (in ms)

Returns

Bitmask condition that caused the thread to unblock, or 0 on error or timeout

19.11.3.8 Wait_i()

```
uint16_t Mark3::EventFlag::Wait_i (
    uint16_t u16Mask_,
    EventFlagOperation eMode_,
    uint32_t u32TimeMS_ ) [private]
```

Wait_i Internal abstraction used to manage both timed and untimed wait operations.

Parameters

<i>u16Mask_</i>	- 16-bit bitmask to block on
<i>eMode_</i>	- EventFlagOperation::Any_Set : Thread will block on any of the bits in the mask <ul style="list-style-type: none"> EventFlagOperation::All_Set: Thread will block on all of the bits in the mask
<i>u32TimeMS_</i>	- Time to block (in ms)

Returns

Bitmask condition that caused the thread to unblock, or 0 on error or timeout

19.11.3.9 WakeMe()

```
void Mark3::EventFlag::WakeMe (
    Thread * pclChosenOne_ )
```

WakeMe Wake the given thread, currently blocking on this object.

Parameters

<i>pclChosenOne_</i>	Pointer to the owner thread to unblock.
----------------------	---

19.11.4 Member Data Documentation**19.11.4.1 m_u16SetMask**

```
uint16_t Mark3::EventFlag::m_u16SetMask [private]
```

Event flags currently set in this object.

Definition at line 119 of file [eventflag.h](#).

The documentation for this class was generated from the following file:

- [/home/moslevin/projects/m3-repo/kernel/src/public/eventflag.h](#)

19.12 Mark3::Kernel Class Reference

The [Kernel](#) Class encapsulates all of the kernel startup, configuration and management functions.

```
#include <kernel.h>
```


Static Public Member Functions

- static void [Init](#) ()
Kernel Initialization Function, call before any other OS function.
- static void [Start](#) ()
Start the operating system kernel - the current execution context is cancelled, all kernel services are started, and the processor resumes execution at the entrypoint for the highest-priority thread.
- static void [CompleteStart](#) ()
CompleteStart Call this from the thread initialization code at the point that the scheduler is to be run for the first time.
- static bool [IsStarted](#) ()
IsStarted.
- static void [SetPanic](#) ([PanicFunc](#) pfPanic_)
SetPanic Set a function to be called when a kernel panic occurs, giving the user to determine the behavior when a catastrophic failure is observed.
- static bool [IsPanic](#) ()
IsPanic Returns whether or not the kernel is in a panic state.
- static void [Panic](#) (uint16_t u16Cause_)
Panic Cause the kernel to enter its panic state.
- static void [SetThreadCreateCallout](#) ([ThreadCreateCallout](#) pfCreate_)
SetThreadCreateCallout Set a function to be called on creation of a new thread. This callout is executed on the successful completion of a [Thread::Init\(\)](#) call. A callout is only executed if this method has been called to set a valid handler function.
- static void [SetThreadExitCallout](#) ([ThreadExitCallout](#) pfExit_)
SetThreadExitCallout Set a function to be called on thread exit. This callout is executed from the beginning of [Thread::Exit\(\)](#).
- static void [SetThreadContextSwitchCallout](#) ([ThreadContextCallout](#) pfContext_)
SetThreadContextSwitchCallout Set a function to be called on each context switch.
- static void [SetDebugPrintFunction](#) ([DebugPrintFunction](#) pfPrintFunction_)
SetDebugPrintFunction Set the function to be used when printing kernel debug information.
- static void [DebugPrint](#) (const char *szString_)
DebugPrint Print a string to the configured output interface. Has no effect if [Kernel::SetDebugPrintFunction\(\)](#) has not been called with a valid print handler.
- static [ThreadCreateCallout](#) [GetThreadCreateCallout](#) ()
GetThreadCreateCallout Return the current function called on every [Thread::Init\(\)](#).
- static [ThreadExitCallout](#) [GetThreadExitCallout](#) ()
GetThreadExitCallout Return the current function called on every [Thread::Exit\(\)](#).
- static [ThreadContextCallout](#) [GetThreadContextSwitchCallout](#) ()
GetThreadContextSwitchCallout Return the current function called on every [Thread::ContextSwitchSWI\(\)](#)
- static void [SetStackGuardThreshold](#) (uint16_t u16Threshold_)
- static uint16_t [GetStackGuardThreshold](#) ()
- static void [Tick](#) ()
- static uint32_t [GetTicks](#) ()

Static Private Attributes

- static bool [m_bIsStarted](#)
true if kernel is running, false otherwise
- static bool [m_bIsPanic](#)
true if kernel is in panic state, false otherwise
- static [PanicFunc](#) [m_pfPanic](#)
set panic function
- static [ThreadCreateCallout](#) [m_pfThreadCreateCallout](#)

Function to call on thread creation.

- static [ThreadExitCallout m_pfThreadExitCallout](#)

Function to call on thread exit.

- static [ThreadContextCallout m_pfThreadContextCallout](#)

Function to call on context switch.

- static [DebugPrintFunction m_pfDebugPrintFunction](#)

Function to call to print debug info.

- static uint16_t [m_u16GuardThreshold](#)
- static uint32_t [m_u32Ticks](#)

19.12.1 Detailed Description

The [Kernel](#) Class encapsulates all of the kernel startup, configuration and management functions.

Definition at line 48 of file [kernel.h](#).

19.12.2 Member Function Documentation

19.12.2.1 CompleteStart()

```
void Mark3::Kernel::CompleteStart ( ) [static]
```

CompleteStart Call this from the thread initialization code at the point that the scheduler is to be run for the first time.

Definition at line 64 of file [kernel.cpp](#).

19.12.2.2 DebugPrint()

```
void Mark3::Kernel::DebugPrint (
    const char * szString_ ) [static]
```

DebugPrint Print a string to the configured output interface. Has no effect if [Kernel::SetDebugPrintFunction\(\)](#) has not been called with a valid print handler.

Parameters

<i>szString_</i>	string to print
------------------	-----------------

Examples:

[lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_coroutines/main.cpp](#), [lab3_round_robin/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab5_mutexes/main.cpp](#), [lab6_timers/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 81 of file [kernel.cpp](#).

19.12.2.3 GetStackGuardThreshold()

```
static uint16_t Mark3::Kernel::GetStackGuardThreshold ( ) [inline], [static]
```

Definition at line 202 of file [kernel.h](#).

19.12.2.4 GetThreadContextSwitchCallout()

```
static ThreadContextCallout Mark3::Kernel::GetThreadContextSwitchCallout ( ) [inline], [static]
```

GetThreadContextSwitchCallout Return the current function called on every [Thread::ContextSwitchSWI\(\)](#)

Returns

Pointer to the currently-installed callout function, or nullptr if not set.

Definition at line 198 of file [kernel.h](#).

19.12.2.5 GetThreadCreateCallout()

```
static ThreadCreateCallout Mark3::Kernel::GetThreadCreateCallout ( ) [inline], [static]
```

GetThreadCreateCallout Return the current function called on every [Thread::Init\(\)](#);

Returns

Pointer to the currently-installed callout function, or nullptr if not set.

Definition at line 178 of file [kernel.h](#).

19.12.2.6 GetThreadExitCallout()

```
static ThreadExitCallout Mark3::Kernel::GetThreadExitCallout ( ) [inline], [static]
```

GetThreadExitCallout Return the current function called on every [Thread::Exit\(\)](#);

Returns

Pointer to the currently-installed callout function, or nullptr if not set.

Definition at line 188 of file [kernel.h](#).

19.12.2.7 GetTicks()

```
uint32_t Mark3::Kernel::GetTicks ( ) [static]
```

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 90 of file [kernel.cpp](#).

19.12.2.8 Init()

```
void Mark3::Kernel::Init (
    void ) [static]
```

[Kernel](#) Initialization Function, call before any other OS function.

Initializes all global resources used by the operating system. This must be called before any other kernel function is invoked.

Examples:

[lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_coroutines/main.cpp](#), [lab3_round_robin/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab5_mutexes/main.cpp](#), [lab6_timers/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 44 of file [kernel.cpp](#).

19.12.2.9 IsPanic()

```
static bool Mark3::Kernel::IsPanic ( ) [inline], [static]
```

IsPanic Returns whether or not the kernel is in a panic state.

Returns

Whether or not the kernel is in a panic state

Definition at line 99 of file [kernel.h](#).

19.12.2.10 IsStarted()

```
static bool Mark3::Kernel::IsStarted ( ) [inline], [static]
```

IsStarted.

Returns

Whether or not the kernel has started - true = running, false = not started

Definition at line 86 of file [kernel.h](#).

19.12.2.11 Panic()

```
void Mark3::Kernel::Panic (
    uint16_t u16Cause_ ) [static]
```

Panic Cause the kernel to enter its panic state.

Parameters

<i>u16</i> ↔ <i>Cause_</i>	Reason for the kernel panic
-------------------------------	-----------------------------

Definition at line 70 of file [kernel.cpp](#).

19.12.2.12 SetDebugPrintFunction()

```
static void Mark3::Kernel::SetDebugPrintFunction (
    DebugPrintFunction pfPrintFunction_ ) [inline], [static]
```

SetDebugPrintFunction Set the function to be used when printing kernel debug information.

Parameters

<i>pfPrint</i> ↔ <i>Function_</i>	Function used to print kernel debug message strings
--------------------------------------	---

Examples:

[lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_coroutines/main.cpp](#), [lab3_round_robin/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab5_mutexes/main.cpp](#), [lab6_timers/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 155 of file [kernel.h](#).

19.12.2.13 SetPanic()

```
static void Mark3::Kernel::SetPanic (
    PanicFunc pfPanic_ ) [inline], [static]
```

SetPanic Set a function to be called when a kernel panic occurs, giving the user to determine the behavior when a catastrophic failure is observed.

Parameters

<i>pf</i> ↔ <i>Panic</i> ↔ _	Panic function pointer
------------------------------------	------------------------

Definition at line 94 of file [kernel.h](#).

19.12.2.14 SetStackGuardThreshold()

```
static void Mark3::Kernel::SetStackGuardThreshold (
    uint16_t u16Threshold_ ) [inline], [static]
```

Definition at line 201 of file [kernel.h](#).

19.12.2.15 SetThreadContextSwitchCallout()

```
static void Mark3::Kernel::SetThreadContextSwitchCallout (
    ThreadContextCallout pfContext_ ) [inline], [static]
```

SetThreadContextSwitchCallout Set a function to be called on each context switch.

A callout is only executed if this method has been called to set a valid handler function.

Parameters

<i>pfContext_</i>	Pointer to a function to call on context switch
—	

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 143 of file [kernel.h](#).

19.12.2.16 SetThreadCreateCallout()

```
static void Mark3::Kernel::SetThreadCreateCallout (
    ThreadCreateCallout pfCreate_ ) [inline], [static]
```

SetThreadCreateCallout Set a function to be called on creation of a new thread. This callout is executed on the successful completion of a [Thread::Init\(\)](#) call. A callout is only executed if this method has been called to set a valid handler function.

Parameters

<i>pfCreate_</i>	Pointer to a function to call on thread creation
—	

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 116 of file [kernel.h](#).

19.12.2.17 SetThreadExitCallout()

```
static void Mark3::Kernel::SetThreadExitCallout (
    ThreadExitCallout pFExit_ ) [inline], [static]
```

SetThreadExitCallout Set a function to be called on thread exit. This callout is executed from the beginning of [Thread::Exit\(\)](#).

A callout is only executed if this method has been called to set a valid handler function.

Parameters

<i>pfExit_</i>	Pointer to a function to call on thread exit
—	

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 130 of file [kernel.h](#).

19.12.2.18 Start()

```
void Mark3::Kernel::Start ( ) [static]
```

Start the operating system kernel - the current execution context is cancelled, all kernel services are started, and the processor resumes execution at the entrypoint for the highest-priority thread.

You must have at least one thread added to the kernel before calling this function, otherwise the behavior is undefined. The exception to this is if the system is configured to use the threadless idle hook, in which case the kernel is allowed to run without any ready threads.

Examples:

[lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_coroutines/main.cpp](#), [lab3_round_robin/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab5_mutexes/main.cpp](#), [lab6_timers/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 58 of file [kernel.cpp](#).

19.12.2.19 Tick()

```
static void Mark3::Kernel::Tick ( ) [inline], [static]
```

Definition at line 205 of file [kernel.h](#).

19.12.3 Member Data Documentation

19.12.3.1 m_bIsPanic

```
bool Mark3::Kernel::m_bIsPanic [static], [private]
```

true if kernel is in panic state, false otherwise

Definition at line 210 of file [kernel.h](#).

19.12.3.2 m_bIsStarted

```
bool Mark3::Kernel::m_bIsStarted [static], [private]
```

true if kernel is running, false otherwise

Definition at line 209 of file [kernel.h](#).

19.12.3.3 m_pfDebugPrintFunction

```
DebugPrintFunction Mark3::Kernel::m_pfDebugPrintFunction [static], [private]
```

Function to call to print debug info.

Function to call when printing debug info.

Definition at line 222 of file [kernel.h](#).

19.12.3.4 m_pfPanic

```
PanicFunc Mark3::Kernel::m_pfPanic [static], [private]
```

set panic function

Definition at line 211 of file [kernel.h](#).

19.12.3.5 m_pfThreadContextCallout

`ThreadContextCallout` Mark3::Kernel::m_pfThreadContextCallout [static], [private]

Function to call on context switch.

Definition at line 220 of file [kernel.h](#).

19.12.3.6 m_pfThreadCreateCallout

`ThreadCreateCallout` Mark3::Kernel::m_pfThreadCreateCallout [static], [private]

Function to call on thread creation.

Definition at line 214 of file [kernel.h](#).

19.12.3.7 m_pfThreadExitCallout

`ThreadExitCallout` Mark3::Kernel::m_pfThreadExitCallout [static], [private]

Function to call on thread exit.

Definition at line 217 of file [kernel.h](#).

19.12.3.8 m_u16GuardThreshold

`uint16_t` Mark3::Kernel::m_u16GuardThreshold [static], [private]

Definition at line 224 of file [kernel.h](#).

19.12.3.9 m_u32Ticks

`uint32_t` Mark3::Kernel::m_u32Ticks [static], [private]

Definition at line 226 of file [kernel.h](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/m3-repo/kernel/src/public/kernel.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/kernel.cpp](#)

19.13 Mark3::KernelSWI Class Reference

The [KernelSWI](#) Class provides the software-interrupt used to implement the context-switching interrupt used by the kernel. This interface must be implemented by target-specific code in the porting layer.

```
#include <kernelswi.h>
```

Static Public Member Functions

- static void [Config](#) (void)
Config Configure the software interrupt - must be called before any other software interrupt functions are called.
- static void [Start](#) (void)
Start Enable ("Start") the software interrupt functionality.
- static void [Trigger](#) (void)
Trigger Call the software interrupt.

19.13.1 Detailed Description

The [KernelSWI](#) Class provides the software-interrupt used to implement the context-switching interrupt used by the kernel. This interface must be implemented by target-specific code in the porting layer.

Definition at line 32 of file [kernelswi.h](#).

19.13.2 Member Function Documentation

19.13.2.1 Config()

```
void Mark3::KernelSWI::Config (  
    void ) [static]
```

Config Configure the software interrupt - must be called before any other software interrupt functions are called.

Definition at line 31 of file [kernelswi.cpp](#).

19.13.2.2 Start()

```
void Mark3::KernelSWI::Start (  
    void ) [static]
```

Start Enable ("Start") the software interrupt functionality.

Definition at line 39 of file [kernelswi.cpp](#).

19.13.2.3 Trigger()

```
void Mark3::KernelSWI::Trigger (
    void ) [static]
```

Trigger Call the software interrupt.

Definition at line 46 of file [kernelswi.cpp](#).

The documentation for this class was generated from the following files:

- /home/moslevin/projects/m3-repo/kernel/src/public/[kernelswi.h](#)
- /home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/[kernelswi.cpp](#)

19.14 Mark3::KernelTimer Class Reference

The [KernelTimer](#) class provides a timer interface used by all time-based scheduling/timer subsystems in the kernel. This interface must be implemented by target-specific code in the porting layer.

```
#include <kerneltimer.h>
```

Static Public Member Functions

- static void [Config](#) (void)
Config Initializes the kernel timer before use.
- static void [Start](#) (void)
Start Starts the kernel time (must be configured first)
- static void [Stop](#) (void)
Stop Shut down the kernel timer, used when no timers are scheduled.

19.14.1 Detailed Description

The [KernelTimer](#) class provides a timer interface used by all time-based scheduling/timer subsystems in the kernel. This interface must be implemented by target-specific code in the porting layer.

Definition at line 34 of file [kerneltimer.h](#).

19.14.2 Member Function Documentation

19.14.2.1 Config()

```
void Mark3::KernelTimer::Config (
    void ) [static]
```

Config Initializes the kernel timer before use.

Definition at line 60 of file [kerneltimer.cpp](#).

19.14.2.2 Start()

```
void Mark3::KernelTimer::Start (
    void ) [static]
```

Start Starts the kernel time (must be configured first)

Definition at line 76 of file [kerneltimer.cpp](#).

19.14.2.3 Stop()

```
void Mark3::KernelTimer::Stop (
    void ) [static]
```

Stop Shut down the kernel timer, used when no timers are scheduled.

Definition at line 86 of file [kerneltimer.cpp](#).

The documentation for this class was generated from the following files:


- [/home/moslevin/projects/m3-repo/kernel/src/public/kerneltimer.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/kerneltimer.cpp](#)

19.15 Mark3::LinkedList Class Reference

The [LinkedList](#) Class Abstract-data-type from which all other linked-lists are derived.

```
#include <ll.h>
```

Inheritance diagram for Mark3::LinkedList:



Public Member Functions

- void [Init](#) ()
Init.
- [LinkedListNode](#) * [GetHead](#) ()
GetHead.
- void [SetHead](#) ([LinkedListNode](#) *pclNode_)
SetHead.
- [LinkedListNode](#) * [GetTail](#) ()
GetTail.
- void [SetTail](#) ([LinkedListNode](#) *pclNode_)
SetTail.

Protected Attributes

- `LinkedListNode * m_pclHead`
Pointer to the head node in the list.
- `LinkedListNode * m_pclTail`
Pointer to the tail node in the list.

19.15.1 Detailed Description

The `LinkedList` Class Abstract-data-type from which all other linked-lists are derived.

Definition at line 119 of file `ll.h`.

19.15.2 Member Function Documentation

19.15.2.1 GetHead()

```
LinkedListNode* Mark3::LinkedList::GetHead ( ) [inline]
```

GetHead.

Get the head node in the linked list

Returns

Pointer to the head node in the list

Definition at line 144 of file `ll.h`.

19.15.2.2 GetTail()

```
LinkedListNode* Mark3::LinkedList::GetTail ( ) [inline]
```

GetTail.

Get the tail node of the linked list

Returns

Pointer to the tail node in the list

Definition at line 162 of file `ll.h`.

19.15.2.3 Init()

```
void Mark3::LinkedList::Init (
    void ) [inline]
```

Init.

Clear the linked list.

Definition at line 131 of file [ll.h](#).

19.15.2.4 SetHead()

```
void Mark3::LinkedList::SetHead (
    LinkedListNode * pclNode_ ) [inline]
```

SetHead.

Set the head node of a linked list

Parameters

<i>pcl</i> ↔ <i>Node_</i>	Pointer to node to set as the head of the linked list
------------------------------	---

Definition at line 153 of file [ll.h](#).

19.15.2.5 SetTail()

```
void Mark3::LinkedList::SetTail (
    LinkedListNode * pclNode_ ) [inline]
```

SetTail.

Set the tail node of the linked list

Parameters

<i>pcl</i> ↔ <i>Node_</i>	Pointer to the node to set as the tail of the linked list
------------------------------	---

Definition at line 171 of file [ll.h](#).

19.15.3 Member Data Documentation

19.15.3.1 m_pclHead

```
LinkedListNode* Mark3::LinkedList::m_pclHead [protected]
```

Pointer to the head node in the list.

Definition at line 122 of file [ll.h](#).

19.15.3.2 m_pclTail

```
LinkedListNode* Mark3::LinkedList::m_pclTail [protected]
```

Pointer to the tail node in the list.

Definition at line 123 of file [ll.h](#).

The documentation for this class was generated from the following file:


- [/home/moslevin/projects/m3-repo/kernel/src/public/ll.h](#)

19.16 Mark3::LinkedListNode Class Reference

The [LinkedListNode](#) Class Basic linked-list node data structure. This data is managed by the linked-list class types, and can be used transparently between them.

```
#include <ll.h>
```

Inheritance diagram for Mark3::LinkedListNode:



Public Member Functions

- [LinkedListNode](#) * [GetNext](#) (void)
GetNext.
- [LinkedListNode](#) * [GetPrev](#) (void)
GetPrev.

Protected Member Functions

- [LinkedListNode](#) ()
- void [ClearNode](#) ()
ClearNode.

Protected Attributes

- [LinkedListNode](#) * [next](#)
Pointer to the next node in the list.
- [LinkedListNode](#) * [prev](#)
Pointer to the previous node in the list.

Friends

- class [LinkedList](#)
- class [DoubleLinkedList](#)
- class [CircularLinkedList](#)

19.16.1 Detailed Description

The [LinkedListNode](#) Class Basic linked-list node data structure. This data is managed by the linked-list class types, and can be used transparently between them.

Definition at line 62 of file [ll.h](#).

19.16.2 Constructor & Destructor Documentation

19.16.2.1 [LinkedListNode\(\)](#)

```
Mark3::LinkedListNode::LinkedListNode ( ) [inline], [protected]
```

Definition at line 68 of file [ll.h](#).

19.16.3 Member Function Documentation

19.16.3.1 [ClearNode\(\)](#)

```
void Mark3::LinkedListNode::ClearNode ( ) [protected]
```

[ClearNode](#).

Initialize the linked list node, clearing its next and previous node.

Definition at line 27 of file [ll.cpp](#).

19.16.3.2 GetNext()

```
LinkedListNode* Mark3::LinkedListNode::GetNext (
    void ) [inline]
```

GetNext.

Returns a pointer to the next node in the list.

Returns

a pointer to the next node in the list.

Definition at line 85 of file [ll.h](#).

19.16.3.3 GetPrev()

```
LinkedListNode* Mark3::LinkedListNode::GetPrev (
    void ) [inline]
```

GetPrev.

Returns a pointer to the previous node in the list.

Returns

a pointer to the previous node in the list.

Definition at line 93 of file [ll.h](#).

19.16.4 Friends And Related Function Documentation

19.16.4.1 CircularLinkedList

```
friend class CircularLinkedList [friend]
```

Definition at line 96 of file [ll.h](#).

19.16.4.2 DoubleLinkedList

```
friend class DoubleLinkedList [friend]
```

Definition at line 95 of file [ll.h](#).

19.16.4.3 LinkList

```
friend class LinkList [friend]
```

Definition at line 94 of file [ll.h](#).

19.16.5 Member Data Documentation

19.16.5.1 next

```
LinkListNode* Mark3::LinkListNode::next [protected]
```

Pointer to the next node in the list.

Definition at line 65 of file [ll.h](#).

19.16.5.2 prev

```
LinkListNode* Mark3::LinkListNode::prev [protected]
```

Pointer to the previous node in the list.

Definition at line 66 of file [ll.h](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/m3-repo/kernel/src/public/ll.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/ll.cpp](#)

19.17 Mark3::LockGuard Class Reference

The [LockGuard](#) class. This class provides RAIL locks based on [Mark3](#)'s kernel [Mutex](#) object. Note that [Mark3](#) does not support exceptions, so care must be taken to ensure that this object is only used where that constraint can be met.

```
#include <lockguard.h>
```

Public Member Functions

- [LockGuard](#) ([Mutex](#) *pciMutex)
- [LockGuard](#) ([Mutex](#) *pciMutex, uint32_t u32TimeoutMs_)
- [~LockGuard](#) ()
- bool [isAcquired](#) ()

Private Attributes

- bool [m_bIsAcquired](#)
- [Mutex](#) * [m_pclMutex](#)

19.17.1 Detailed Description

The [LockGuard](#) class. This class provides RAII locks based on [Mark3](#)'s kernel [Mutex](#) object. Note that [Mark3](#) does not support exceptions, so care must be taken to ensure that this object is only used where that constraint can be met.

Definition at line 32 of file [lockguard.h](#).

19.17.2 Constructor & Destructor Documentation

19.17.2.1 LockGuard() [1/2]

```
Mark3::LockGuard::LockGuard (
    Mutex * pclMutex )
```

Parameters

<i>pclMutex</i>	mutex to lock during construction
-----------------	-----------------------------------

Definition at line 25 of file [lockguard.cpp](#).

19.17.2.2 LockGuard() [2/2]

```
Mark3::LockGuard::LockGuard (
    Mutex * pclMutex,
    uint32_t u32TimeoutMs_ )
```

Parameters

<i>pclMutex</i>	mutex to lock during construction
<i>u32TimeoutMs_</i>	timeout (in ms) to wait before bailing

Definition at line 34 of file [lockguard.cpp](#).

19.17.2.3 ~LockGuard()

```
Mark3::LockGuard::~~LockGuard ( )
```

Definition at line 42 of file [lockguard.cpp](#).

19.17.3 Member Function Documentation

19.17.3.1 isAcquired()

```
bool Mark3::LockGuard::isAcquired ( ) [inline]
```

Verify that lock was correctly initialized and locked during acquisition. This is used to provide error-checking for timed RAIL locks, where [Mark3](#) does not use exceptions, and a kernel-panic is too heavy-handed.

Returns

true if the lock was initialed correctly, false on error

Definition at line 54 of file [lockguard.h](#).

19.17.4 Member Data Documentation

19.17.4.1 m_bIsAcquired

```
bool Mark3::LockGuard::m_bIsAcquired [private]
```

Definition at line 57 of file [lockguard.h](#).

19.17.4.2 m_pclMutex

```
Mutex* Mark3::LockGuard::m_pclMutex [private]
```

Definition at line 58 of file [lockguard.h](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/m3-repo/kernel/src/public/lockguard.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/lockguard.cpp](#)

19.18 Mark3::Mailbox Class Reference

The [Mailbox](#) class. This class implements an IPC mechanism based on sending/receiving envelopes containing data of a fixed size, configured at initialization) that reside within a buffer of memory provided by the user.

```
#include <mailbox.h>
```

Public Member Functions

- void * [operator new](#) (size_t sz, void *pv)
- [~Mailbox](#) ()
- void [Init](#) (void *pvBuffer_, uint16_t u16BufferSize_, uint16_t u16ElementSize_)

Init Initialize the mailbox object prior to its use. This must be called before any calls can be made to the object.
- bool [Send](#) (void *pvData_)

Send Send an envelope to the mailbox. This safely copies the data contents of the datastructure to the previously-initialized mailbox buffer. If there is a thread already blocking, awaiting delivery to the mailbox, it will be unblocked at this time.
- bool [SendTail](#) (void *pvData_)

SendTail Send an envelope to the mailbox. This safely copies the data contents of the datastructure to the previously-initialized mailbox buffer. If there is a thread already blocking, awaiting delivery to the mailbox, it will be unblocked at this time.
- bool [Send](#) (void *pvData_, uint32_t u32TimeoutMS_)

Send Send an envelope to the mailbox. This safely copies the data contents of the datastructure to the previously-initialized mailbox buffer. If there is a thread already blocking, awaiting delivery to the mailbox, it will be unblocked at this time.
- bool [SendTail](#) (void *pvData_, uint32_t u32TimeoutMS_)

SendTail Send an envelope to the mailbox. This safely copies the data contents of the datastructure to the previously-initialized mailbox buffer. If there is a thread already blocking, awaiting delivery to the mailbox, it will be unblocked at this time.
- void [Receive](#) (void *pvData_)

Receive Read one envelope from the head of the mailbox. If the mailbox is currently empty, the calling thread will block until an envelope is delivered.
- void [ReceiveTail](#) (void *pvData_)

ReceiveTail Read one envelope from the tail of the mailbox. If the mailbox is currently empty, the calling thread will block until an envelope is delivered.
- bool [Receive](#) (void *pvData_, uint32_t u32TimeoutMS_)

Receive Read one envelope from the head of the mailbox. If the mailbox is currently empty, the calling thread will block until an envelope is delivered, or the specified time has elapsed without delivery.
- bool [ReceiveTail](#) (void *pvData_, uint32_t u32TimeoutMS_)

ReceiveTail Read one envelope from the tail of the mailbox. If the mailbox is currently empty, the calling thread will block until an envelope is delivered, or the specified time has elapsed without delivery.
- uint16_t [GetFreeSlots](#) (void)
- bool [IsFull](#) (void)
- bool [IsEmpty](#) (void)

Static Public Member Functions

- static [Mailbox](#) * [Init](#) (uint16_t u16BufferSize_, uint16_t u16ElementSize_)

Init Create and initialize the mailbox object prior to its use. This must be called before any calls can be made to the object. This version of the API allocates the buffer space from the kernel's Auto-Allocation heap, which cannot be returned back. As a result, this is only suitable for cases where the mailbox will be created once on startup, and persist for the duration of the system.

Private Member Functions

- void * [GetHeadPointer](#) (void)
GetHeadPointer Return a pointer to the current head of the mailbox's internal circular buffer.
- void * [GetTailPointer](#) (void)
GetTailPointer Return a pointer to the current tail of the mailbox's internal circular buffer.
- void [CopyData](#) (const void *src_, void *dst_, uint16_t len_)
CopyData Perform a direct byte-copy from a source to a destination object.
- void [MoveTailForward](#) (void)
MoveTailForward Move the tail index forward one element.
- void [MoveHeadForward](#) (void)
MoveHeadForward Move the head index forward one element.
- void [MoveTailBackward](#) (void)
MoveTailBackward Move the tail index backward one element.
- void [MoveHeadBackward](#) (void)
MoveHeadBackward Move the head index backward one element.
- bool [Send_i](#) (const void *pvData_, bool bTail_, uint32_t u32TimeoutMS_)
Send_i Internal method which implements all [Send\(\)](#) methods in the class.
- bool [Receive_i](#) (void *pvData_, bool bTail_, uint32_t u32WaitTimeMS_)
Receive_i Internal method which implements all [Read\(\)](#) methods in the class.

Private Attributes

- uint16_t [m_u16Head](#)
Current head index.
- uint16_t [m_u16Tail](#)
Current tail index.
- uint16_t [m_u16Count](#)
Count of items in the mailbox.
- volatile uint16_t [m_u16Free](#)
Current number of free slots in the mailbox.
- uint16_t [m_u16ElementSize](#)
Size of the objects tracked in this mailbox.
- const void * [m_pvBuffer](#)
Pointer to the data-buffer managed by this mailbox.
- Semaphore [m_clRecvSem](#)
Counting semaphore used to synchronize threads on the object.
- Semaphore [m_clSendSem](#)
Binary semaphore for send-blocked threads.

19.18.1 Detailed Description

The [Mailbox](#) class. This class implements an IPC mechanism based on sending/receiving envelopes containing data of a fixed size, configured at initialization) that reside within a buffer of memory provided by the user.

Examples:

[lab11_mailboxes/main.cpp](#).

Definition at line 35 of file [mailbox.h](#).

19.18.2 Constructor & Destructor Documentation

19.18.2.1 ~Mailbox()

```
Mark3::Mailbox::~~Mailbox ( )
```

Definition at line 25 of file [mailbox.cpp](#).

19.18.3 Member Function Documentation

19.18.3.1 CopyData()

```
void Mark3::Mailbox::CopyData (
    const void * src_,
    void * dst_,
    uint16_t len_ ) [inline], [private]
```

CopyData Perform a direct byte-copy from a source to a destination object.

Parameters

<i>src</i> ↔ —	Pointer to an object to read from
<i>dst</i> ↔ —	Pointer to an object to write to
<i>len</i> ↔ —	Length to copy (in bytes)

Definition at line 216 of file [mailbox.h](#).

19.18.3.2 GetFreeSlots()

```
uint16_t Mark3::Mailbox::GetFreeSlots (
    void ) [inline]
```

Definition at line 170 of file [mailbox.h](#).

19.18.3.3 GetHeadPointer()

```
void* Mark3::Mailbox::GetHeadPointer (
    void ) [inline], [private]
```

GetHeadPointer Return a pointer to the current head of the mailbox's internal circular buffer.

Returns

pointer to the head element in the mailbox

Definition at line 187 of file [mailbox.h](#).

19.18.3.4 GetTailPointer()

```
void* Mark3::Mailbox::GetTailPointer (
    void ) [inline], [private]
```

GetTailPointer Return a pointer to the current tail of the mailbox's internal circular buffer.

Returns

pointer to the tail element in the mailbox

Definition at line 201 of file [mailbox.h](#).

19.18.3.5 Init() [1/2]

```
void Mark3::Mailbox::Init (
    void * pvBuffer_,
    uint16_t u16BufferSize_,
    uint16_t u16ElementSize_ )
```

Init Initialize the mailbox object prior to its use. This must be called before any calls can be made to the object.

Parameters

<i>pvBuffer_</i>	Pointer to the static buffer to use for the mailbox
<i>u16BufferSize_</i>	Size of the mailbox buffer, in bytes
<i>u16ElementSize_</i>	Size of each envelope, in bytes

Definition at line 34 of file [mailbox.cpp](#).

19.18.3.6 Init() [2/2]

```
Mailbox * Mark3::Mailbox::Init (
    uint16_t u16BufferSize_,
    uint16_t u16ElementSize_ ) [static]
```

Init Create and initialize the mailbox object prior to its use. This must be called before any calls can be made to the object. This version of the API allocates the buffer space from the kernel's Auto-Allocation heap, which cannot be returned back. As a result, this is only suitable for cases where the mailbox will be created once on startup, and persist for the duration of the system.

Parameters

<i>u16BufferSize_</i>	Size of the mailbox buffer, in bytes
<i>u16ElementSize_</i>	Size of each envelope, in bytes

Definition at line 59 of file [mailbox.cpp](#).

19.18.3.7 IsEmpty()

```
bool Mark3::Mailbox::IsEmpty (
    void ) [inline]
```

Definition at line 177 of file [mailbox.h](#).

19.18.3.8 IsFull()

```
bool Mark3::Mailbox::IsFull (
    void ) [inline]
```

Definition at line 176 of file [mailbox.h](#).

19.18.3.9 MoveHeadBackward()

```
void Mark3::Mailbox::MoveHeadBackward (
    void ) [inline], [private]
```

MoveHeadBackward Move the head index backward one element.

Definition at line 263 of file [mailbox.h](#).

19.18.3.10 MoveHeadForward()

```
void Mark3::Mailbox::MoveHeadForward (
    void ) [inline], [private]
```

MoveHeadForward Move the head index forward one element.

Definition at line 239 of file [mailbox.h](#).

19.18.3.11 MoveTailBackward()

```
void Mark3::Mailbox::MoveTailBackward (
    void ) [inline], [private]
```

MoveTailBackward Move the tail index backward one element.

Definition at line 251 of file [mailbox.h](#).

19.18.3.12 MoveTailForward()

```
void Mark3::Mailbox::MoveTailForward (
    void ) [inline], [private]
```

MoveTailForward Move the tail index forward one element.

Definition at line 227 of file [mailbox.h](#).

19.18.3.13 operator new()

```
void* Mark3::Mailbox::operator new (
    size_t sz,
    void * pv ) [inline]
```

Definition at line 38 of file [mailbox.h](#).

19.18.3.14 Receive() [1/2]

```
void Mark3::Mailbox::Receive (
    void * pvData_ )
```

Receive Read one envelope from the head of the mailbox. If the mailbox is currently empty, the calling thread will block until an envelope is delivered.

Parameters

<i>pvData_</i>	Pointer to a buffer that will have the envelope's contents copied into upon delivery.
----------------	---

Examples:

[lab11_mailboxes/main.cpp](#).

Definition at line 83 of file [mailbox.cpp](#).

19.18.3.15 Receive() [2/2]

```
bool Mark3::Mailbox::Receive (
    void * pvData_,
    uint32_t u32TimeoutMS_ )
```

Receive Read one envelope from the head of the mailbox. If the mailbox is currently empty, the calling thread will block until an envelope is delivered, or the specified time has elapsed without delivery.

Parameters

<i>pvData_</i>	Pointer to a buffer that will have the envelope's contents copied into upon delivery.
<i>u32TimeoutMS_</i>	Maximum time to wait for delivery.

Returns

true - envelope was delivered, false - delivery timed out.

Definition at line 90 of file [mailbox.cpp](#).

19.18.3.16 Receive_i()

```
bool Mark3::Mailbox::Receive_i (
    void * pvData_,
    bool bTail_,
    uint32_t u32WaitTimeMS_ ) [private]
```

Receive_i Internal method which implements all Read() methods in the class.

Parameters

<i>pvData_</i>	Pointer to the envelope data
<i>bTail_</i>	true - read from tail, false - read from head
<i>u32WaitTimeMS_</i>	Time to wait before timeout (in ms).

Returns

true - read successfully, false - timeout.

Definition at line 197 of file [mailbox.cpp](#).

19.18.3.17 ReceiveTail() [1/2]

```
void Mark3::Mailbox::ReceiveTail (
    void * pvData_ )
```

ReceiveTail Read one envelope from the tail of the mailbox. If the mailbox is currently empty, the calling thread will block until an envelope is delivered.

Parameters

<i>pvData_</i>	Pointer to a buffer that will have the envelope's contents copied into upon delivery.
----------------	---

Definition at line 97 of file [mailbox.cpp](#).

19.18.3.18 ReceiveTail() [2/2]

```
bool Mark3::Mailbox::ReceiveTail (
    void * pvData_,
    uint32_t u32TimeoutMS_ )
```

ReceiveTail Read one envelope from the tail of the mailbox. If the mailbox is currently empty, the calling thread will block until an envelope is delivered, or the specified time has elapsed without delivery.

Parameters

<i>pvData_</i>	Pointer to a buffer that will have the envelope's contents copied into upon delivery.
<i>u32TimeoutMS_</i>	Maximum time to wait for delivery.

Returns

true - envelope was delivered, false - delivery timed out.

Definition at line 104 of file [mailbox.cpp](#).

19.18.3.19 Send() [1/2]

```
bool Mark3::Mailbox::Send (
    void * pvData_ )
```

Send Send an envelope to the mailbox. This safely copies the data contents of the datastructure to the previously-initialized mailbox buffer. If there is a thread already blocking, awaiting delivery to the mailbox, it will be unblocked at this time.

This method delivers the envelope at the head of the mailbox.

Parameters

<i>pvData_</i>	Pointer to the data object to send to the mailbox.
----------------	--

Returns

true - envelope was delivered, false - mailbox is full.

Examples:

[lab11_mailboxes/main.cpp](#).

Definition at line 111 of file [mailbox.cpp](#).

19.18.3.20 Send() [2/2]

```
bool Mark3::Mailbox::Send (
    void * pvData_,
    uint32_t u32TimeoutMS_ )
```

Send Send an envelope to the mailbox. This safely copies the data contents of the datastructure to the previously-initialized mailbox buffer. If there is a thread already blocking, awaiting delivery to the mailbox, it will be unblocked at this time.

This method delivers the envelope at the head of the mailbox.

Parameters

<i>pvData_</i>	Pointer to the data object to send to the mailbox.
<i>u32TimeoutMS_</i>	Maximum time to wait for a free transmit slot

Returns

true - envelope was delivered, false - mailbox is full.

Definition at line 125 of file [mailbox.cpp](#).

19.18.3.21 Send_i()

```
bool Mark3::Mailbox::Send_i (
    const void * pvData_,
    bool bTail_,
    uint32_t u32TimeoutMS_ ) [private]
```

Send_i Internal method which implements all [Send\(\)](#) methods in the class.

Parameters

<i>pvData_</i>	Pointer to the envelope data
<i>bTail_</i>	true - write to tail, false - write to head
<i>u32TimeoutMS_</i>	Time to wait before timeout (in ms).

Returns

true - data successfully written, false - buffer full

Definition at line [139](#) of file [mailbox.cpp](#).

19.18.3.22 SendTail() [[1/2](#)]

```
bool Mark3::Mailbox::SendTail (
    void * pvData_ )
```

SendTail Send an envelope to the mailbox. This safely copies the data contents of the datastructure to the previously-initialized mailbox buffer. If there is a thread already blocking, awaiting delivery to the mailbox, it will be unblocked at this time.

This method delivers the envelope at the tail of the mailbox.

Parameters

<i>pvData_</i>	Pointer to the data object to send to the mailbox.
----------------	--

Returns

true - envelope was delivered, false - mailbox is full.

Definition at line [118](#) of file [mailbox.cpp](#).

19.18.3.23 SendTail() [2/2]

```
bool Mark3::Mailbox::SendTail (
    void * pvData_,
    uint32_t u32TimeoutMS_ )
```

SendTail Send an envelope to the mailbox. This safely copies the data contents of the datastructure to the previously-initialized mailbox buffer. If there is a thread already blocking, awaiting delivery to the mailbox, it will be unblocked at this time.

This method delivers the envelope at the tail of the mailbox.

Parameters

<i>pvData_</i>	Pointer to the data object to send to the mailbox.
<i>u32TimeoutMS_</i>	Maximum time to wait for a free transmit slot

Returns

true - envelope was delivered, false - mailbox is full.

Definition at line 132 of file [mailbox.cpp](#).

19.18.4 Member Data Documentation

19.18.4.1 m_clRecvSem

[Semaphore](#) Mark3::Mailbox::m_clRecvSem [private]

Counting semaphore used to synchronize threads on the object.

Definition at line 302 of file [mailbox.h](#).

19.18.4.2 m_clSendSem

[Semaphore](#) Mark3::Mailbox::m_clSendSem [private]

Binary semaphore for send-blocked threads.

Definition at line 303 of file [mailbox.h](#).

19.18.4.3 m_pvBuffer

```
const void* Mark3::Mailbox::m_pvBuffer [private]
```

Pointer to the data-buffer managed by this mailbox.

Definition at line 300 of file [mailbox.h](#).

19.18.4.4 m_u16Count

```
uint16_t Mark3::Mailbox::m_u16Count [private]
```

Count of items in the mailbox.

Definition at line 296 of file [mailbox.h](#).

19.18.4.5 m_u16ElementSize

```
uint16_t Mark3::Mailbox::m_u16ElementSize [private]
```

Size of the objects tracked in this mailbox.

Definition at line 299 of file [mailbox.h](#).

19.18.4.6 m_u16Free

```
volatile uint16_t Mark3::Mailbox::m_u16Free [private]
```

Current number of free slots in the mailbox.

Definition at line 297 of file [mailbox.h](#).

19.18.4.7 m_u16Head

```
uint16_t Mark3::Mailbox::m_u16Head [private]
```

Current head index.

Definition at line 293 of file [mailbox.h](#).

19.18.4.8 m_u16Tail

```
uint16_t Mark3::Mailbox::m_u16Tail [private]
```

Current tail index.

Definition at line 294 of file [mailbox.h](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/m3-repo/kernel/src/public/mailbox.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/mailbox.cpp](#)

19.19 Mark3::MemUtil Class Reference

String and Memory manipulation class.

```
#include <memutil.h>
```

Static Public Member Functions

- static void [DecimalToHex](#) (uint8_t u8Data_, char *szText_)
DecimalToHex.
- static void [DecimalToHex](#) (uint16_t u16Data_, char *szText_)
- static void [DecimalToHex](#) (uint32_t u32Data_, char *szText_)
- static void [DecimalToHex](#) (uint64_t u64Data_, char *szText_)
- static void [DecimalToString](#) (uint8_t u8Data_, char *szText_)
DecimalToString.
- static void [DecimalToString](#) (uint16_t u16Data_, char *szText_)
- static void [DecimalToString](#) (uint32_t u32Data_, char *szText_)
- static void [DecimalToString](#) (uint64_t u64Data_, char *szText_)
- static bool [StringToDecimal8](#) (const char *szText_, uint8_t *pu8Out_)
StringToDecimial8.
- static bool [StringToDecimal16](#) (const char *szText_, uint16_t *pu16Out_)
- static bool [StringToDecimal32](#) (const char *szText_, uint32_t *pu32Out_)
- static bool [StringToDecimal64](#) (const char *szText_, uint64_t *pu64Out_)
- static uint8_t [Checksum8](#) (const void *pvSrc_, uint16_t u16Len_)
Checksum8.
- static uint16_t [Checksum16](#) (const void *pvSrc_, uint16_t u16Len_)
Checksum16.
- static uint16_t [StringLength](#) (const char *szStr_)
StringLength.
- static bool [CompareStrings](#) (const char *szStr1_, const char *szStr2_)
CompareStrings.
- static bool [CompareStrings](#) (const char *szStr1_, const char *szStr2_, uint16_t u16Length_)
- static void [CopyMemory](#) (void *pvDst_, const void *pvSrc_, uint16_t u16Len_)
CopyMemory.
- static void [CopyString](#) (char *szDst_, const char *szSrc_)
CopyString.
- static int16_t [StringSearch](#) (const char *szBuffer_, const char *szPattern_)

StringSearch.

- static bool [CompareMemory](#) (const void *pvMem1_, const void *pvMem2_, uint16_t u16Len_)

CompareMemory.

- static void [SetMemory](#) (void *pvDst_, uint8_t u8Val_, uint16_t u16Len_)

SetMemory.

- static uint8_t [Tokenize](#) (const char *szBuffer_, [Token_t](#) *pastTokens_, uint8_t u8MaxTokens_)

Tokenize Function to tokenize a string based on a space delimiter. This is a non-destructive function, which populates a [Token_t](#) descriptor array.

19.19.1 Detailed Description

String and Memory manipulation class.

Utility method class implementing common memory and string manipulation functions, without relying on an external standard library implementation which might not be available on some toolchains, may be closed source, or may not be thread-safe.

Definition at line 46 of file [memutil.h](#).

19.19.2 Member Function Documentation

19.19.2.1 Checksum16()

```
uint16_t Mark3::MemUtil::Checksum16 (
    const void * pvSrc_,
    uint16_t u16Len_ ) [static]
```

Checksum16.

Compute the 16-bit additive checksum of a memory buffer.

Parameters

<i>pvSrc_</i>	Memory buffer to compute a 16-bit checksum of.
<i>u16Len_</i>	Length of the buffer in bytes.

Returns

16-bit checksum of the memory block.

Definition at line 346 of file [memutil.cpp](#).

19.19.2.2 Checksum8()

```
uint8_t Mark3::MemUtil::Checksum8 (
    const void * pvSrc_,
    uint16_t u16Len_ ) [static]
```

Checksum8.

Compute the 8-bit additive checksum of a memory buffer.

Parameters

<i>pvSrc_</i>	Memory buffer to compute a 8-bit checksum of.
<i>u16Len_</i>	Length of the buffer in bytes.

Returns

8-bit checksum of the memory block.

Definition at line 333 of file [memutil.cpp](#).

19.19.2.3 CompareMemory()

```
bool Mark3::MemUtil::CompareMemory (
    const void * pvMem1_,
    const void * pvMem2_,
    uint16_t u16Len_ ) [static]
```

CompareMemory.

Compare the contents of two memory buffers to eachother

Parameters

<i>pvMem1_</i>	First buffer to compare
<i>pvMem2_</i>	Second buffer to compare
<i>u16Len_</i>	Length of buffer (in bytes) to compare

Returns

true if the buffers match, false if they do not.

Definition at line 467 of file [memutil.cpp](#).

19.19.2.4 CompareStrings() [1/2]

```
bool Mark3::MemUtil::CompareStrings (
    const char * szStr1_,
    const char * szStr2_ ) [static]
```

CompareStrings.

Compare the contents of two zero-terminated string buffers to eachother.

Parameters

<i>szStr1_↔</i>	First string to compare
<i>szStr2_↔</i>	Second string to compare

Returns

true if strings match, false otherwise.

Definition at line 372 of file [memutil.cpp](#).

19.19.2.5 CompareStrings() [2/2]

```
bool Mark3::MemUtil::CompareStrings (
    const char * szStr1_,
    const char * szStr2_,
    uint16_t u16Length_ ) [static]
```

Definition at line 391 of file [memutil.cpp](#).

19.19.2.6 CopyMemory()

```
void Mark3::MemUtil::CopyMemory (
    void * pvDst_,
    const void * pvSrc_,
    uint16_t u16Len_ ) [static]
```

CopyMemory.

Copy one buffer in memory into another.

Parameters

<i>pvDst_↔</i>	Pointer to the destination buffer
<i>pvSrc_↔</i>	Pointer to the source buffer
<i>u16Len_↔</i>	Number of bytes to copy from source to destination

Definition at line 408 of file [memutil.cpp](#).

19.19.2.7 CopyString()

```
void Mark3::MemUtil::CopyString (
    char * szDst_,
    const char * szSrc_ ) [static]
```

CopyString.

Copy a string from one buffer into another.

Parameters

<i>sz↔ Dst_</i>	Pointer to the buffer to copy into
<i>sz↔ Src_</i>	Pointer to the buffer to copy data from

Definition at line 422 of file [memutil.cpp](#).

19.19.2.8 DecimalToHex() [1/4]

```
void Mark3::MemUtil::DecimalToHex (
    uint8_t u8Data_,
    char * szText_ ) [static]
```

DecimalToHex.

Convert an 8-bit unsigned binary value as a hexadecimal string.

Parameters

<i>u8↔ Data_</i>	Value to convert into a string
<i>sz↔ Text_</i>	Destination string buffer (3 bytes minimum)

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 29 of file [memutil.cpp](#).

19.19.2.9 DecimalToHex() [2/4]

```
void Mark3::MemUtil::DecimalToHex (
    uint16_t u16Data_,
    char * szText_ ) [static]
```

Definition at line 55 of file [memutil.cpp](#).

19.19.2.10 DecimalToHex() [3/4]

```
void Mark3::MemUtil::DecimalToHex (
    uint32_t u32Data_,
    char * szText_ ) [static]
```

Definition at line 81 of file [memutil.cpp](#).

19.19.2.11 DecimalToHex() [4/4]

```
void Mark3::MemUtil::DecimalToHex (
    uint64_t u64Data_,
    char * szText_ ) [static]
```

Definition at line 107 of file [memutil.cpp](#).

19.19.2.12 DecimalToString() [1/4]

```
void Mark3::MemUtil::DecimalToString (
    uint8_t u8Data_,
    char * szText_ ) [static]
```

DecimalToString.

Convert an 8-bit unsigned binary value as a decimal string.

Parameters

<i>u8↔ Data_</i>	Value to convert into a string
<i>sz↔ Text_</i>	Destination string buffer (4 bytes minimum)

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 133 of file [memutil.cpp](#).

19.19.2.13 DecimalToString() [2/4]

```
void Mark3::MemUtil::DecimalToString (
    uint16_t u16Data_,
    char * szText_ ) [static]
```

Definition at line 157 of file [memutil.cpp](#).

19.19.2.14 DecimalToString() [3/4]

```
void Mark3::MemUtil::DecimalToString (
    uint32_t u32Data_,
    char * szText_ ) [static]
```

Definition at line 178 of file [memutil.cpp](#).

19.19.2.15 DecimalToString() [4/4]

```
void Mark3::MemUtil::DecimalToString (
    uint64_t u64Data_,
    char * szText_ ) [static]
```

Definition at line 199 of file [memutil.cpp](#).

19.19.2.16 SetMemory()

```
void Mark3::MemUtil::SetMemory (
    void * pvDst_,
    uint8_t u8Val_,
    uint16_t u16Len_ ) [static]
```

SetMemory.

Initialize a buffer of memory to a specified 8-bit pattern

Parameters

<i>pvDst</i> ↔ —	Destination buffer to set
<i>u8Val</i> ↔ —	8-bit pattern to initialize each byte of destination with
<i>u16Len</i> ↔ Generated by Doxygen Len_	Length of the buffer (in bytes) to initialize

Definition at line 486 of file [memutil.cpp](#).

19.19.2.17 StringLength()

```
uint16_t Mark3::MemUtil::StringLength (
    const char * szStr_ ) [static]
```

StringLength.

Compute the length of a string in bytes.

Parameters

<i>sz↔ Str_</i>	Pointer to the zero-terminated string to calculate the length of
---------------------	--

Returns

length of the string (in bytes), not including the 0-terminator.

Definition at line 360 of file [memutil.cpp](#).

19.19.2.18 StringSearch()

```
int16_t Mark3::MemUtil::StringSearch (
    const char * szBuffer_,
    const char * szPattern_ ) [static]
```

StringSearch.

Search for the presence of one string as a substring within another.

Parameters

<i>szBuffer↔ _</i>	Buffer to search for pattern within
<i>sz↔ Pattern_</i>	Pattern to search for in the buffer

Returns

Index of the first instance of the pattern in the buffer, or -1 on no match.

Definition at line 436 of file [memutil.cpp](#).

19.19.2.19 StringToDecimal16()

```
bool Mark3::MemUtil::StringToDecimal16 (
    const char * szText_,
    uint16_t * pu16Out_ ) [static]
```

Definition at line 248 of file [memutil.cpp](#).

19.19.2.20 StringToDecimal32()

```
bool Mark3::MemUtil::StringToDecimal32 (
    const char * szText_,
    uint32_t * pu32Out_ ) [static]
```

Definition at line 276 of file [memutil.cpp](#).

19.19.2.21 StringToDecimal64()

```
bool Mark3::MemUtil::StringToDecimal64 (
    const char * szText_,
    uint64_t * pu64Out_ ) [static]
```

Definition at line 304 of file [memutil.cpp](#).

19.19.2.22 StringToDecimal8()

```
bool Mark3::MemUtil::StringToDecimal8 (
    const char * szText_,
    uint8_t * pu8Out_ ) [static]
```

StringToDecimial8.

Convert a string to an unsigned integer value.

Parameters

<i>szText</i> ↔ —	String to convert
<i>pu8</i> ↔ <i>Out</i> —	Pointer to a uint8_t that will contain the result

Returns

true on success, false on invalid parameters or failure to convert the input string to an unsigned integer value

Definition at line 220 of file [memutil.cpp](#).

19.19.2.23 Tokenize()

```
uint8_t Mark3::MemUtil::Tokenize (
    const char * szBuffer_,
    Token_t * pastTokens_,
    uint8_t u8MaxTokens_ ) [static]
```

Tokenize Function to tokenize a string based on a space delimiter. This is a non-destructive function, which populates a [Token_t](#) descriptor array.

Parameters

<i>szBuffer_</i>	String to tokenize
<i>pastTokens_</i>	Pointer to the array of token descriptors
<i>u8MaxTokens_</i>	Maximum number of tokens to parse (i.e. size of <i>pastTokens_</i>)

Returns

Count of tokens parsed

Definition at line 496 of file [memutil.cpp](#).

The documentation for this class was generated from the following files:


- [/home/moslevin/projects/m3-repo/kernel/lib/memutil/public/memutil.h](#)
- [/home/moslevin/projects/m3-repo/kernel/lib/memutil/memutil.cpp](#)

19.20 Mark3::Message Class Reference

the [Message](#) class. This object provides threadsafe message-based IPC services based on exchange of objects containing a data pointer and minimal application-defined metadata. Messages are to be allocated/produced by the sender, and deallocated/consumed by the receiver.

```
#include <message.h>
```

Inheritance diagram for Mark3::Message:



Public Member Functions

- void * [operator new](#) (size_t sz, void *pv)
Init Initialize the data and code in the message.
- void [Init](#) ()
Init Initialize the data and code in the message.
- void [SetData](#) (void *pvData_)
SetData Set the data pointer for the message before transmission.
- void * [GetData](#) ()
GetData Get the data pointer stored in the message upon receipt.
- void [SetCode](#) (uint16_t u16Code_)
SetCode Set the code in the message before transmission.
- uint16_t [GetCode](#) ()
GetCode Return the code set in the message upon receipt.

Private Attributes

- void * [m_pvData](#)
Pointer to the message data.
- uint16_t [m_u16Code](#)
Message code, providing context for the message.

Additional Inherited Members

19.20.1 Detailed Description

the [Message](#) class. This object provides threadsafe message-based IPC services based on exchange of objects containing a data pointer and minimal application-defined metadata. Messages are to be allocated/produced by the sender, and deallocated/consumed by the receiver.

Examples:

[lab8_messages/main.cpp](#).

Definition at line 97 of file [message.h](#).

19.20.2 Member Function Documentation

19.20.2.1 GetCode()

```
uint16_t Mark3::Message::GetCode ( ) [inline]
```

GetCode Return the code set in the message upon receipt.

Returns

user code set in the object

Definition at line 139 of file [message.h](#).

19.20.2.2 GetData()

```
void* Mark3::Message::GetData ( ) [inline]
```

GetData Get the data pointer stored in the message upon receipt.

Returns

Pointer to the data set in the message object

Definition at line 125 of file [message.h](#).

19.20.2.3 Init()

```
void Mark3::Message::Init (
    void ) [inline]
```

Init Initialize the data and code in the message.

Definition at line 105 of file [message.h](#).

19.20.2.4 operator new()

```
void* Mark3::Message::operator new (
    size_t sz,
    void * pv ) [inline]
```

Definition at line 100 of file [message.h](#).

19.20.2.5 SetCode()

```
void Mark3::Message::SetCode (
    uint16_t u16Code_ ) [inline]
```

SetCode Set the code in the message before transmission.

Parameters

<i>u16Code_</i>	Data code to set in the object
-----------------	--------------------------------

Examples:

[lab8_messages/main.cpp](#).

Definition at line 132 of file [message.h](#).

19.20.2.6 SetData()

```
void Mark3::Message::SetData (
    void * pvData_ ) [inline]
```

SetData Set the data pointer for the message before transmission.

Parameters

<i>pvData_</i>	Pointer to the data object to send in the message
----------------	---

Definition at line 118 of file [message.h](#).

19.20.3 Member Data Documentation

19.20.3.1 m_pvData

```
void* Mark3::Message::m_pvData [private]
```

Pointer to the message data.

Definition at line 143 of file [message.h](#).

19.20.3.2 m_u16Code

```
uint16_t Mark3::Message::m_u16Code [private]
```

[Message](#) code, providing context for the message.

Definition at line 146 of file [message.h](#).

The documentation for this class was generated from the following file:

- [/home/moslevin/projects/m3-repo/kernel/src/public/message.h](#)

19.21 Mark3::MessagePool Class Reference

The [MessagePool](#) Class The [MessagePool](#) class implements a simple allocator for message objects exchanged between threads. The sender allocates (pop's) messages, then sends them to the receiver. Upon receipt, it is the receiver's responsibility to deallocate (push) the message back to the pool.

```
#include <message.h>
```

Public Member Functions

- void * [operator new](#) (size_t sz, void *pv)
- [~MessagePool](#) ()
- void [Init](#) ()
Init.
- void [Push](#) ([Message](#) *pclMessage_)
Push.
- [Message](#) * [Pop](#) ()
Pop.
- [Message](#) * [GetHead](#) ()
GetHead.

Private Attributes

- [TypedDoubleLinkedList](#)< [Message](#) > [m_clList](#)
Linked list used to manage the [Message](#) objects.

19.21.1 Detailed Description

The [MessagePool](#) Class The [MessagePool](#) class implements a simple allocator for message objects exchanged between threads. The sender allocates (pop's) messages, then sends them to the receiver. Upon receipt, it is the receiver's responsibility to deallocate (push) the message back to the pool.

Examples:

[lab8_messages/main.cpp](#).

Definition at line 157 of file [message.h](#).

19.21.2 Constructor & Destructor Documentation

19.21.2.1 ~MessagePool()

```
Mark3::MessagePool::~~MessagePool ( ) [inline]
```

Definition at line 161 of file [message.h](#).

19.21.3 Member Function Documentation

19.21.3.1 GetHead()

```
Message * Mark3::MessagePool::GetHead ( )
```

GetHead.

Return a pointer to the first element in the message list

Returns

Definition at line 52 of file [message.cpp](#).

19.21.3.2 Init()

```
void Mark3::MessagePool::Init (
    void )
```

Init.

Initialize the message queue prior to use

Definition at line 26 of file [message.cpp](#).

19.21.3.3 operator new()

```
void* Mark3::MessagePool::operator new (
    size_t sz,
    void * pv ) [inline]
```

Definition at line 160 of file [message.h](#).

19.21.3.4 Pop()

```
Message * Mark3::MessagePool::Pop ( )
```

Pop.

Pop a message from the queue, returning it to the user to be popu32ated before sending by a transmitter.

Returns

Pointer to a [Message](#) object

Examples:

[lab8_messages/main.cpp](#).

Definition at line 41 of file [message.cpp](#).

19.21.3.5 Push()

```
void Mark3::MessagePool::Push (
    Message * pclMessage_ )
```

Push.

Return a previously-claimed message object back to the queue. used once the message has been processed by a receiver.

Parameters

<i>pcl</i> ↔ <i>Message_</i>	Pointer to the Message object to return back to the queue
---------------------------------	---

Examples:

[lab8_messages/main.cpp](#).

Definition at line 32 of file [message.cpp](#).

19.21.4 Member Data Documentation

19.21.4.1 m_clList

```
TypedDoubleLinkedList<Message> Mark3::MessagePool::m_clList [private]
```

Linked list used to manage the [Message](#) objects.

Definition at line 201 of file [message.h](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/m3-repo/kernel/src/public/message.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/message.cpp](#)

19.22 Mark3::MessageQueue Class Reference

The [MessageQueue](#) class. Implements a mechanism used to send/receive data between threads. Allows threads to block, waiting for messages to be sent from other contexts.

```
#include <message.h>
```

Public Member Functions

- `void * operator new (size_t sz, void *pv)`
- `~MessageQueue ()`
- `void Init ()`
Init.
- `Message * Receive ()`
Receive.
- `Message * Receive (uint32_t u32TimeWaitMS_)`
Receive.
- `void Send (Message *pclSrc_)`
Send.
- `uint16_t GetCount ()`
GetCount.

Private Member Functions

- `Message * Receive_i (uint32_t u32TimeWaitMS_)`
Receive_i.

Private Attributes

- `Semaphore m_clSemaphore`
Counting semaphore used to manage thread blocking.
- `TypedDoubleLinkedList< Message > m_clLinkList`
List object used to store messages.

19.22.1 Detailed Description

The [MessageQueue](#) class. Implements a mechanism used to send/receive data between threads. Allows threads to block, waiting for messages to be sent from other contexts.

Examples:

[lab8_messages/main.cpp](#).

Definition at line 210 of file [message.h](#).

19.22.2 Constructor & Destructor Documentation

19.22.2.1 ~MessageQueue()

```
Mark3::MessageQueue::~MessageQueue ( ) [inline]
```

Definition at line 214 of file [message.h](#).

19.22.3 Member Function Documentation

19.22.3.1 GetCount()

```
uint16_t Mark3::MessageQueue::GetCount ( )
```

GetCount.

Return the number of messages pending in the "receive" queue.

Returns

Count of pending messages in the queue.

Definition at line 108 of file [message.cpp](#).

19.22.3.2 Init()

```
void Mark3::MessageQueue::Init (
    void )
```

Init.

Initialize the message queue prior to use.

Definition at line 58 of file [message.cpp](#).

19.22.3.3 operator new()

```
void* Mark3::MessageQueue::operator new (
    size_t sz,
    void * pv ) [inline]
```

Definition at line 213 of file [message.h](#).

19.22.3.4 Receive() [1/2]

```
Message * Mark3::MessageQueue::Receive ( )
```

Receive.

Receive a message from the message queue. If the message queue is empty, the thread will block until a message is available.

Returns

Pointer to a message object at the head of the queue

Examples:

[lab8_messages/main.cpp](#).

Definition at line 64 of file [message.cpp](#).

19.22.3.5 Receive() [2/2]

```
Message * Mark3::MessageQueue::Receive (
    uint32_t u32TimeWaitMS_ )
```

Receive.

Receive a message from the message queue. If the message queue is empty, the thread will block until a message is available for the duration specified. If no message arrives within that duration, the call will return with nullptr.

Parameters

<i>u32TimeWaitMS_</i>	The amount of time in ms to wait for a message before timing out and unblocking the waiting thread.
-----------------------	---

Returns

Pointer to a message object at the head of the queue or nullptr on timeout.

Definition at line 70 of file [message.cpp](#).

19.22.3.6 Receive_i()

```
Message * Mark3::MessageQueue::Receive_i (
    uint32_t u32TimeWaitMS_ ) [private]
```

Receive_i.

Internal function used to abstract timed and un-timed Receive calls.

Parameters

<i>u32TimeWaitMS_</i>	Time (in ms) to block, 0 for un-timed call.
-----------------------	---

Returns

Pointer to a message, or 0 on timeout.

Definition at line 76 of file [message.cpp](#).

19.22.3.7 Send()

```
void Mark3::MessageQueue::Send (
    Message * pclSrc_ )
```

Send.

Send a message object into this message queue. Will un-block the first waiting thread blocked on this queue if that occurs.

Parameters

<i>pclSrc_</i>	Pointer to the message object to add to the queue
----------------	---

Examples:

[lab8_messages/main.cpp](#).

Definition at line 92 of file [message.cpp](#).

19.22.4 Member Data Documentation

19.22.4.1 m_clLinkList

```
TypedDoubleLinkedList<Message> Mark3::MessageQueue::m_clLinkList [private]
```

List object used to store messages.

Definition at line 284 of file [message.h](#).

19.22.4.2 m_clSemaphore

```
Semaphore Mark3::MessageQueue::m_clSemaphore [private]
```

Counting semaphore used to manage thread blocking.

Definition at line 281 of file [message.h](#).

The documentation for this class was generated from the following files:


- [/home/moslevin/projects/m3-repo/kernel/src/public/message.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/message.cpp](#)

19.23 Mark3::Mutex Class Reference

The [Mutex](#) Class. Class providing Mutual-exclusion locks, based on [BlockingObject](#).

```
#include <mutex.h>
```

Inheritance diagram for Mark3::Mutex:



Public Member Functions

- void * [operator new](#) (size_t sz, void *pv)
- [~Mutex](#) ()
- void [Init](#) (bool bRecursive_=true)

Init Initialize a mutex object for use - must call this function before using the object.
- void [Claim](#) ()

Claim Claim the mutex. When the mutex is claimed, no other thread can claim a region protected by the object. If another [Thread](#) currently holds the [Mutex](#) when the Claim method is called, that [Thread](#) will block until the current owner of the mutex releases the [Mutex](#).
- bool [Claim](#) (uint32_t u32WaitTimeMS_)

Claim Claim a mutex, with timeout.
- void [WakeMe](#) ([Thread](#) *pclOwner_)

WakeMe Wake a thread blocked on the mutex. This is an internal function used for implementing timed mutexes relying on timer callbacks. Since these do not have access to the private data of the mutex and its base classes, we have to wrap this as a public method - do not use this for any other purposes.
- void [Release](#) ()

Release Release the mutex. When the mutex is released, another object can enter the mutex-protected region.

Private Member Functions

- `uint8_t WakeNext ()`
WakeNext.
- `bool Claim_i (uint32_t u32WaitTimeMS_)`
Claim_i Abstracts out timed/non-timed mutex claim operations.

Private Attributes

- `uint8_t m_u8Recurse`
The recursive lock-count when a mutex is claimed multiple times by the same owner.
- `bool m_bReady`
State of the mutex - true = ready, false = claimed.
- `bool m_bRecursive`
Whether or not the lock is recursive.
- `PORT_PRIO_TYPE m_uMaxPri`
Maximum priority of thread in queue, used for priority inheritance.
- `Thread * m_pclOwner`
Pointer to the thread that owns the mutex (when claimed)

Additional Inherited Members

19.23.1 Detailed Description

The [Mutex](#) Class. Class providing Mutual-exclusion locks, based on [BlockingObject](#).

Examples:

[lab5_mutexes/main.cpp](#).

Definition at line 63 of file [mutex.h](#).

19.23.2 Constructor & Destructor Documentation

19.23.2.1 ~Mutex()

```
Mark3::Mutex::~Mutex ( )
```

Definition at line 55 of file [mutex.cpp](#).

19.23.3 Member Function Documentation

19.23.3.1 Claim() [1/2]

```
void Mark3::Mutex::Claim (
    void )
```

Claim Claim the mutex. When the mutex is claimed, no other thread can claim a region protected by the object. If another [Thread](#) currently holds the [Mutex](#) when the Claim method is called, that [Thread](#) will block until the current owner of the mutex releases the [Mutex](#).

If the calling [Thread](#)'s priority is lower than that of a [Thread](#) that currently owns the [Mutex](#) object, then the priority of that [Thread](#) will be elevated to that of the highest-priority calling object until the [Mutex](#) is released. This property is known as "Priority Inheritance"

Note: A single thread can recursively claim a mutex up to a count of

1. Attempting to claim a mutex beyond that will cause a kernel panic.

Examples:

[lab5_mutexes/main.cpp](#).

Definition at line [185](#) of file [mutex.cpp](#).

19.23.3.2 Claim() [2/2]

```
bool Mark3::Mutex::Claim (
    uint32_t u32WaitTimeMS_ )
```

Claim Claim a mutex, with timeout.

Parameters

<i>u32WaitTimeMS_</i>	↔
S_	

Returns

true - mutex was claimed within the time period specified false - mutex operation timed-out before the claim operation.

Definition at line [191](#) of file [mutex.cpp](#).

19.23.3.3 Claim_i()

```
bool Mark3::Mutex::Claim_i (
    uint32_t u32WaitTimeMS_ ) [private]
```

[Claim_i](#) Abstracts out timed/non-timed mutex claim operations.

Parameters

<code>u32WaitTimeMS_</code>	Time in MS to wait, 0 for infinite
-----------------------------	------------------------------------

Returns

true on successful claim, false otherwise

Definition at line 108 of file [mutex.cpp](#).

19.23.3.4 Init()

```
void Mark3::Mutex::Init (
    bool bRecursive_ = true )
```

Init Initialize a mutex object for use - must call this function before using the object.

Parameters

<code>bRecursive_</code>	Whether or not the mutex can be recursively locked.
--------------------------	---

Definition at line 93 of file [mutex.cpp](#).

19.23.3.5 operator new()

```
void* Mark3::Mutex::operator new (
    size_t sz,
    void * pv ) [inline]
```

Definition at line 66 of file [mutex.h](#).

19.23.3.6 Release()

```
void Mark3::Mutex::Release ( )
```

Release Release the mutex. When the mutex is released, another object can enter the mutex-protected region.

If there are Threads waiting for the [Mutex](#) to become available, then the highest priority [Thread](#) will be unblocked at this time and will claim the [Mutex](#) lock immediately - this may result in an immediate context switch, depending on relative priorities.

If the calling [Thread](#)'s priority was boosted as a result of priority inheritance, the [Thread](#)'s previous priority will also be restored at this time.

Note that if a [Mutex](#) is held recursively, it must be Release'd the same number of times that it was Claim'd before it will be available for use by another [Thread](#).

Examples:

[lab5_mutexes/main.cpp](#).

Definition at line 197 of file [mutex.cpp](#).

19.23.3.7 WakeMe()

```
void Mark3::Mutex::WakeMe (
    Thread * pclOwner_ )
```

WakeMe Wake a thread blocked on the mutex. This is an internal function used for implementing timed mutexes relying on timer callbacks. Since these do not have access to the private data of the mutex and its base classes, we have to wrap this as a public method - do not use this for any other purposes.

Parameters

<i>pclOwner_</i>	Thread to unblock from this object.
------------------	---

Definition at line 65 of file [mutex.cpp](#).

19.23.3.8 WakeNext()

```
uint8_t Mark3::Mutex::WakeNext ( ) [private]
```

WakeNext.

Wake the next thread waiting on the [Mutex](#).

Definition at line 73 of file [mutex.cpp](#).

19.23.4 Member Data Documentation

19.23.4.1 m_bReady

```
bool Mark3::Mutex::m_bReady [private]
```

State of the mutex - true = ready, false = claimed.

Definition at line 159 of file [mutex.h](#).

19.23.4.2 m_bRecursive

```
bool Mark3::Mutex::m_bRecursive [private]
```

Whether or not the lock is recursive.

Definition at line 160 of file [mutex.h](#).

19.23.4.3 m_pclOwner

```
Thread* Mark3::Mutex::m_pclOwner [private]
```

Pointer to the thread that owns the mutex (when claimed)

Definition at line 162 of file [mutex.h](#).

19.23.4.4 m_u8Recurse

```
uint8_t Mark3::Mutex::m_u8Recurse [private]
```

The recursive lock-count when a mutex is claimed multiple times by the same owner.

Definition at line 158 of file [mutex.h](#).

19.23.4.5 m_uMaxPri

```
PORT_PRIO_TYPE Mark3::Mutex::m_uMaxPri [private]
```

Maximum priority of thread in queue, used for priority inheritance.

Definition at line 161 of file [mutex.h](#).

The documentation for this class was generated from the following files:

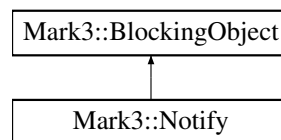
- [/home/moslevin/projects/m3-repo/kernel/src/public/mutex.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/mutex.cpp](#)

19.24 Mark3::Notify Class Reference

The [Notify](#) class. This class provides a blocking object type that allows one or more threads to wait for an event to occur before resuming operation.

```
#include <notify.h>
```

Inheritance diagram for Mark3::Notify:



Public Member Functions

- void * [operator new](#) (size_t sz, void *pv)
- [~Notify](#) ()
- void [Init](#) (void)
Init Initialize the Notification object prior to use.
- void [Signal](#) (void)
Signal Signal the notification object. This will cause the highest priority thread currently blocking on the object to wake. If no threads are currently blocked on the object, the call has no effect.
- void [Wait](#) (bool *pbFlag_)
Wait Block the current thread, waiting for a signal on the object.
- bool [Wait](#) (uint32_t u32WaitTimeMS_, bool *pbFlag_)
Wait Block the current thread, waiting for a signal on the object.
- void [WakeMe](#) ([Thread](#) *pclChosenOne_)
WakeMe Wake the specified thread from its current blocking queue. Note that this is only public in order to be accessible from a timer callack.

Private Attributes

- bool [m_bPending](#)

Additional Inherited Members

19.24.1 Detailed Description

The [Notify](#) class. This class provides a blocking object type that allows one or more threads to wait for an event to occur before resuming operation.

Examples:

[lab10_notifications/main.cpp](#).

Definition at line 33 of file [notify.h](#).

19.24.2 Constructor & Destructor Documentation

19.24.2.1 ~Notify()

```
Mark3::Notify::~~Notify ( )
```

Definition at line 47 of file [notify.cpp](#).

19.24.3 Member Function Documentation

19.24.3.1 Init()

```
void Mark3::Notify::Init (
    void )
```

Init Initialize the Notification object prior to use.

Definition at line 57 of file [notify.cpp](#).

19.24.3.2 operator new()

```
void* Mark3::Notify::operator new (
    size_t sz,
    void * pv ) [inline]
```

Definition at line 36 of file [notify.h](#).

19.24.3.3 Signal()

```
void Mark3::Notify::Signal (
    void )
```

Signal Signal the notification object. This will cause the highest priority thread currently blocking on the object to wake. If no threads are currently blocked on the object, the call has no effect.

Examples:

[lab10_notifications/main.cpp](#).

Definition at line 66 of file [notify.cpp](#).

19.24.3.4 Wait() [1/2]

```
void Mark3::Notify::Wait (
    bool * pbFlag_ )
```

Wait Block the current thread, waiting for a signal on the object.

Parameters

<i>pbFlag_</i>	Flag set to false on block, and true upon wakeup.
----------------	---

Examples:

[lab10_notifications/main.cpp](#).

Definition at line 95 of file [notify.cpp](#).

19.24.3.5 Wait() [2/2]

```
bool Mark3::Notify::Wait (  
    uint32_t u32WaitTimeMS_  
    bool * pbFlag_ )
```

Wait Block the current thread, waiting for a signal on the object.

Parameters

<i>u32WaitTimeMS_</i>	Time to wait for the notification event.
<i>pbFlag_</i>	Flag set to false on block, and true upon wakeup.

Returns

true on notification, false on timeout

Definition at line 125 of file [notify.cpp](#).

19.24.3.6 WakeMe()

```
void Mark3::Notify::WakeMe (  
    Thread * pclChosenOne_ )
```

WakeMe Wake the specified thread from its current blocking queue. Note that this is only public in order to be accessible from a timer callack.

Parameters

<i>pclChosenOne_</i>	Thread to wake up
----------------------	-----------------------------------

Definition at line 175 of file [notify.cpp](#).

19.24.4 Member Data Documentation

19.24.4.1 m_bPending

```
bool Mark3::Notify::m_bPending [private]
```

Definition at line 89 of file [notify.h](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/m3-repo/kernel/src/public/notify.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/notify.cpp](#)

19.25 Mark3::PriorityMapL1< T, C > Class Template Reference

The [PriorityMapL1](#) class This class implements a priority bitmap data structure. Each bit in the objects internal storage represents a priority. When a bit is set, it indicates that something is scheduled at the bit's corresponding priority, when a bit is clear it indicates that no entities are scheduled at that priority. This object provides the fundamental logic required to implement efficient priority-based scheduling for the thread + coroutine schedulers in the kernel.

```
#include <priomap11.h>
```

Public Member Functions

- [PriorityMapL1](#) ()
PriorityMap Initialize the priority map object, clearing the bitamp data to all 0's.
- void [Set](#) (T uXPrio_)
Set Set the priority map bitmap data, at all levels, for the given priority.
- void [Clear](#) (T uXPrio_)
Clear Clear the priority map bitmap data, at all levels, for the given priority.
- T [HighestPriority](#) (void)
HighestPriority Computes the numeric priority of the highest-priority thread represented in the priority map.

Static Private Member Functions

- static T [PrioBit](#) (T prio)
- static T [PrioMapWordIndex](#) (T prio)
- static T [PriorityFromBitmap](#) (T uXPrio_)

Private Attributes

- T [m_uXPriorityMap](#)

Static Private Attributes

- static constexpr size_t `m_uXPrioMapShiftLUT` [9] = {0, 3, 4, 0, 5, 0, 0, 0, 6}
- static constexpr auto `m_uXPrioMapWordShift` = T { `m_uXPrioMapShiftLUT`[sizeof(T)] }
- static constexpr auto `m_uXPrioMapBits` = T { 8 * sizeof(T) }
- static constexpr auto `m_uXPrioMapBitMask` = T { (1 << `m_uXPrioMapWordShift`) - 1 }

19.25.1 Detailed Description

```
template<typename T, size_t C>
class Mark3::PriorityMapL1< T, C >
```

The [PriorityMapL1](#) class This class implements a priority bitmap data structure. Each bit in the objects internal storage represents a priority. When a bit is set, it indicates that something is scheduled at the bit's corresponding priority, when a bit is clear it indicates that no entities are scheduled at that priority. This object provides the fundamental logic required to implement efficient priority-based scheduling for the thread + coroutine schedulers in the kernel.

The L1 version of the datastructure uses a single unsigned integer (specified by the templated typename "T"), to support "C" priorities. Thus, the maximum number of priorities supported in the map is the number of bits in the "T" parameter.

Definition at line 45 of file [priomapl1.h](#).

19.25.2 Constructor & Destructor Documentation

19.25.2.1 PriorityMapL1()

```
template<typename T , size_t C>
Mark3::PriorityMapL1< T, C >::PriorityMapL1 ( ) [inline]
```

PriorityMap Initialize the priority map object, clearing the bitamp data to all 0's.

Definition at line 52 of file [priomapl1.h](#).

19.25.3 Member Function Documentation

19.25.3.1 Clear()

```
template<typename T , size_t C>
void Mark3::PriorityMapL1< T, C >::Clear (
    T uXPrio_ ) [inline]
```

Clear Clear the priority map bitmap data, at all levels, for the given priority.

Parameters

<i>uX</i> ↔ <i>Prio_</i>	Priority level to clear the bitmap data for.
-----------------------------	--

Definition at line 73 of file [priomapl1.h](#).

19.25.3.2 HighestPriority()

```
template<typename T , size_t C>
T Mark3::PriorityMapL1< T, C >::HighestPriority (
    void ) [inline]
```

HighestPriority Computes the numeric priority of the highest-priority thread represented in the priority map.

Returns

Highest priority ready-thread's number.

Definition at line 86 of file [priomapl1.h](#).

19.25.3.3 PrioBit()

```
template<typename T , size_t C>
static T Mark3::PriorityMapL1< T, C >::PrioBit (
    T prio ) [inline], [static], [private]
```

Definition at line 93 of file [priomapl1.h](#).

19.25.3.4 PrioMapWordIndex()

```
template<typename T , size_t C>
static T Mark3::PriorityMapL1< T, C >::PrioMapWordIndex (
    T prio ) [inline], [static], [private]
```

Definition at line 95 of file [priomapl1.h](#).

19.25.3.5 PriorityFromBitmap()

```
template<typename T , size_t C>
static T Mark3::PriorityMapL1< T, C >::PriorityFromBitmap (
    T uXPrio_ ) [inline], [static], [private]
```

Definition at line 97 of file [priomapl1.h](#).

19.25.3.6 Set()

```
template<typename T , size_t C>
void Mark3::PriorityMapL1< T, C >::Set (
    T uXPrio_ ) [inline]
```

Set Set the priority map bitmap data, at all levels, for the given priority.

Parameters

<i>uX</i> ↔ <i>Prio_</i>	Priority level to set the bitmap data for.
-----------------------------	--

Definition at line 62 of file [priomapl1.h](#).

19.25.4 Member Data Documentation

19.25.4.1 m_uXPrioMapBitMask

```
template<typename T , size_t C>
constexpr auto Mark3::PriorityMapL1< T, C >::m_uXPrioMapBitMask = T { (1 << m_uXPrioMapWord↔
Shift) - 1 } [static], [private]
```

Definition at line 122 of file [priomapl1.h](#).

19.25.4.2 m_uXPrioMapBits

```
template<typename T , size_t C>
constexpr auto Mark3::PriorityMapL1< T, C >::m_uXPrioMapBits = T { 8 * sizeof(T) } [static],
[private]
```

Definition at line 121 of file [priomapl1.h](#).

19.25.4.3 m_uXPrioMapShiftLUT

```
template<typename T , size_t C>
constexpr size_t Mark3::PriorityMapL1< T, C >::m_uXPrioMapShiftLUT[9] = {0, 3, 4, 0, 5, 0, 0,
0, 6} [static], [private]
```

Definition at line 119 of file [priomapl1.h](#).

19.25.4.4 m_uXPrioMapWordShift

```
template<typename T , size_t C>
constexpr auto Mark3::PriorityMapL1< T, C >::m_uXPrioMapWordShift = T { m_uXPrioMapShiftL↔
UT[sizeof(T)] } [static], [private]
```

Definition at line 120 of file [priomapl1.h](#).

19.25.4.5 m_uXPriorityMap

```
template<typename T , size_t C>
T Mark3::PriorityMapL1< T, C >::m_uXPriorityMap [private]
```

Definition at line 124 of file [priomapl1.h](#).

The documentation for this class was generated from the following file:

- [/home/moslevin/projects/m3-repo/kernel/src/public/priomapl1.h](#)

19.26 Mark3::PriorityMapL2< T, C > Class Template Reference

The [PriorityMapL2](#) class This class implements a priority bitmap data structure. Each bit in the objects internal storage represents a priority. When a bit is set, it indicates that something is scheduled at the bit's corresponding priority, when a bit is clear it indicates that no entities are scheduled at that priority. This object provides the fundamental logic required to implement efficient priority-based scheduling for the thread + coroutine schedulers in the kernel.

```
#include <priomapl2.h>
```

Public Member Functions

- [PriorityMapL2](#) ()
PriorityMap Initialize the priority map object, clearing the bitamp data to all 0's.
- void [Set](#) (T uXPrio_)
Set Set the priority map bitmap data, at all levels, for the given priority.
- void [Clear](#) (T uXPrio_)
Clear Clear the priority map bitmap data, at all levels, for the given priority.
- T [HighestPriority](#) (void)
HighestPriority Computes the numeric priority of the highest-priority thread represented in the priority map.

Static Private Member Functions

- static T [PrioBit](#) (T prio)
- static T [PrioMapWordIndex](#) (T prio)
- static T [PriorityFromBitmap](#) (T uXPrio_)

Private Attributes

- T [m_auXPriorityMap](#) [m_uXPrioMapNumWords]
- T [m_uXPriorityMapL2](#)

Static Private Attributes

- static constexpr size_t [m_uXPrioMapShiftLUT](#) [9] = {0, 3, 4, 0, 5, 0, 0, 0, 6}
- static constexpr auto [m_uXPrioMapWordShift](#) = T { [m_uXPrioMapShiftLUT](#)[sizeof(T)] }
- static constexpr auto [m_uXPrioMapBits](#) = T { 8 * sizeof(T) }
- static constexpr auto [m_uXPrioMapBitMask](#) = T { (1 << [m_uXPrioMapWordShift](#)) - 1 }
- static constexpr auto [m_uXPrioMapNumWords](#) = T { (C + ([m_uXPrioMapBits](#) - 1)) / [m_uXPrioMapBits](#) }

19.26.1 Detailed Description

```
template<typename T, size_t C>
class Mark3::PriorityMapL2< T, C >
```

The [PriorityMapL2](#) class This class implements a priority bitmap data structure. Each bit in the objects internal storage represents a priority. When a bit is set, it indicates that something is scheduled at the bit's corresponding priority, when a bit is clear it indicates that no entities are scheduled at that priority. This object provides the fundamental logic required to implement efficient priority-based scheduling for the thread + coroutine schedulers in the kernel.

The L2 version of the datastructure uses a two-layer nested map structure, where a first layer bitmap contains a single unsigned integer of type "T", where each bit corresponds to an array entry in the second layer table. The second layer consists of an array of unsigned integers of type "T", where each bit in each element corresponds to a level of priority supported by the map structure. As a result, the maximum number of priorities ("C") supported by the object is n*n, where n is the number of bits in the "T" parameter. i.e., if an 8-bit unsigned int is used, this object supports up to 64 priorities; and if a 32-bit unsigned int is used, the object supports up to 1024 priorities.

Definition at line 50 of file [priomapl2.h](#).

19.26.2 Constructor & Destructor Documentation

19.26.2.1 PriorityMapL2()

```
template<typename T , size_t C>
Mark3::PriorityMapL2< T, C >::PriorityMapL2 ( ) [inline]
```

PriorityMap Initialize the priority map object, clearing the bitamp data to all 0's.

Definition at line 57 of file [priomapl2.h](#).

19.26.3 Member Function Documentation

19.26.3.1 Clear()

```
template<typename T , size_t C>
void Mark3::PriorityMapL2< T, C >::Clear (
    T uXPrio_ ) [inline]
```

Clear Clear the priority map bitmap data, at all levels, for the given priority.

Parameters

$uX \leftarrow$ <i>Prio_</i>	Priority level to clear the bitmap data for.
---------------------------------	--

Definition at line 82 of file [priomapl2.h](#).

19.26.3.2 HighestPriority()

```
template<typename T , size_t C>
T Mark3::PriorityMapL2< T, C >::HighestPriority (
    void ) [inline]
```

HighestPriority Computes the numeric priority of the highest-priority thread represented in the priority map.

Returns

Highest priority ready-thread's number.

Definition at line 100 of file [priomapl2.h](#).

19.26.3.3 PrioBit()

```
template<typename T , size_t C>
static T Mark3::PriorityMapL2< T, C >::PrioBit (
    T prio ) [inline], [static], [private]
```

Definition at line 113 of file [priomapl2.h](#).

19.26.3.4 PrioMapWordIndex()

```
template<typename T , size_t C>
static T Mark3::PriorityMapL2< T, C >::PrioMapWordIndex (
    T prio ) [inline], [static], [private]
```

Definition at line 115 of file [priomapl2.h](#).

19.26.3.5 PriorityFromBitmap()

```
template<typename T , size_t C>
static T Mark3::PriorityMapL2< T, C >::PriorityFromBitmap (
    T uXPrio_ ) [inline], [static], [private]
```

Definition at line 117 of file [priomapl2.h](#).

19.26.3.6 Set()

```
template<typename T , size_t C>
void Mark3::PriorityMapL2< T, C >::Set (
    T uXPrio_ ) [inline]
```

Set Set the priority map bitmap data, at all levels, for the given priority.

Parameters

<i>uX</i> ↔ <i>Prio_</i>	Priority level to set the bitmap data for.
-----------------------------	--

Definition at line 68 of file [priomapl2.h](#).

19.26.4 Member Data Documentation

19.26.4.1 m_auXPriorityMap

```
template<typename T , size_t C>
T Mark3::PriorityMapL2< T, C >::m_auXPriorityMap[m_uXPrioMapNumWords] [private]
```

Definition at line 148 of file [priomapl2.h](#).

19.26.4.2 m_uXPrioMapBitMask

```
template<typename T , size_t C>
constexpr auto Mark3::PriorityMapL2< T, C >::m_uXPrioMapBitMask = T { (1 << m_uXPrioMapWord↔
Shift) - 1 } [static], [private]
```

Definition at line 142 of file [priomapl2.h](#).

19.26.4.3 m_uXPrioMapBits

```
template<typename T , size_t C>
constexpr auto Mark3::PriorityMapL2< T, C >::m_uXPrioMapBits = T { 8 * sizeof(T) } [static],
[private]
```

Definition at line 141 of file [priomapl2.h](#).

19.26.4.4 m_uXPrioMapNumWords

```
template<typename T , size_t C>
constexpr auto Mark3::PriorityMapL2< T, C >::m_uXPrioMapNumWords = T { (C + (m_uXPrioMapBits -
1)) / m_uXPrioMapBits } [static], [private]
```

Definition at line 146 of file [priomapl2.h](#).

19.26.4.5 m_uXPrioMapShiftLUT

```
template<typename T , size_t C>
constexpr size_t Mark3::PriorityMapL2< T, C >::m_uXPrioMapShiftLUT[9] = {0, 3, 4, 0, 5, 0, 0,
0, 6} [static], [private]
```

Definition at line 139 of file [priomapl2.h](#).

19.26.4.6 m_uXPrioMapWordShift

```
template<typename T , size_t C>
constexpr auto Mark3::PriorityMapL2< T, C >::m_uXPrioMapWordShift = T { m_uXPrioMapShiftLUT[  
sizeof(T)] } [static], [private]
```

Definition at line 140 of file [priomapl2.h](#).

19.26.4.7 m_uXPriorityMapL2

```
template<typename T , size_t C>
T Mark3::PriorityMapL2< T, C >::m_uXPriorityMapL2 [private]
```

Definition at line 149 of file [priomapl2.h](#).

The documentation for this class was generated from the following file:

- [/home/moslevin/projects/m3-repo/kernel/src/public/priomapl2.h](#)

19.27 Mark3::ProfileTimer Class Reference

Profiling timer. This class is used to perform high-performance profiling of code to see how `int32_t` certain operations take. useful in instrumenting the performance of key algorithms and time-critical operations to ensure real-timer behavior.

```
#include <profile.h>
```

Public Member Functions

- void [Init](#) ()
Init Initialize the profiling timer prior to use. Can also be used to reset a timer that's been used previously.
- void [Start](#) ()
Start Start a profiling session, if the timer is not already active. Has no effect if the timer is already active.
- void [Stop](#) ()
Stop Stop the current profiling session, adding to the cumulative time for this timer, and the total iteration count.
- `uint32_t` [GetAverage](#) ()
GetAverage Get the average time associated with this operation.
- `uint32_t` [GetCurrent](#) ()
GetCurrent Return the current tick count held by the profiler. Valid for both active and stopped timers.

Private Attributes

- `uint32_t m_u32StartTicks`
Cumulative tick-count for this timer.
- `uint32_t m_u32CurrentIteration`
Tick count for current iteration.
- `uint32_t m_u32Cumulative`
Cumulative ticks tracked.
- `uint16_t m_u16Iterations`
Number of iterations executed for this profiling timer.
- `bool m_bActive`
Whether or not the timer is active or stopped.

19.27.1 Detailed Description

Profiling timer. This class is used to perform high-performance profiling of code to see how `uint32_t` certain operations take. useful in instrumenting the performance of key algorithms and time-critical operations to ensure real-timer behavior.

Definition at line 67 of file [profile.h](#).

19.27.2 Member Function Documentation

19.27.2.1 GetAverage()

```
uint32_t Mark3::ProfileTimer::GetAverage ( )
```

GetAverage Get the average time associated with this operation.

Returns

Average tick count normalized over all iterations

Definition at line 63 of file [profile.cpp](#).

19.27.2.2 GetCurrent()

```
uint32_t Mark3::ProfileTimer::GetCurrent ( )
```

GetCurrent Return the current tick count held by the profiler. Valid for both active and stopped timers.

Returns

The currently held tick count.

Definition at line 72 of file [profile.cpp](#).

19.27.2.3 Init()

```
void Mark3::ProfileTimer::Init (  
    void )
```

Init Initialize the profiling timer prior to use. Can also be used to reset a timer that's been used previously.

Definition at line 25 of file [profile.cpp](#).

19.27.2.4 Start()

```
void Mark3::ProfileTimer::Start (  
    void )
```

Start Start a profiling session, if the timer is not already active. Has no effect if the timer is already active.

Definition at line 34 of file [profile.cpp](#).

19.27.2.5 Stop()

```
void Mark3::ProfileTimer::Stop (  
    void )
```

Stop Stop the current profiling session, adding to the cumulative time for this timer, and the total iteration count.

Definition at line 46 of file [profile.cpp](#).

19.27.3 Member Data Documentation

19.27.3.1 m_bActive

```
bool Mark3::ProfileTimer::m_bActive [private]
```

Wheter or not the timer is active or stopped.

Definition at line 113 of file [profile.h](#).

19.27.3.2 m_u16Iterations

```
uint16_t Mark3::ProfileTimer::m_u16Iterations [private]
```

Number of iterations executed for this profiling timer.

Definition at line 112 of file [profile.h](#).

19.27.3.3 m_u32Cumulative

```
uint32_t Mark3::ProfileTimer::m_u32Cumulative [private]
```

Cumulative ticks tracked.

Definition at line 111 of file [profile.h](#).

19.27.3.4 m_u32CurrentIteration

```
uint32_t Mark3::ProfileTimer::m_u32CurrentIteration [private]
```

Tick count for current iteration.

Definition at line 110 of file [profile.h](#).

19.27.3.5 m_u32StartTicks

```
uint32_t Mark3::ProfileTimer::m_u32StartTicks [private]
```

Cumulative tick-count for this timer.

Definition at line 109 of file [profile.h](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/m3-repo/kernel/src/public/profile.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/profile.cpp](#)

19.28 Mark3::Quantum Class Reference

The [Quantum](#) Class. Static-class used to implement [Thread](#) quantum functionality, which is fundamental to round-robin thread scheduling.

```
#include <quantum.h>
```

Static Public Member Functions

- static void [Init](#) ()
- static void [SetInTimer](#) ()
SetInTimer Set a flag to indicate that the CPU is currently running within the timer-callback routine. This prevents the [Quantum](#) timer from being updated in the middle of a callback cycle, potentially resulting in the kernel timer becoming disabled.
- static void [ClearInTimer](#) ()
ClearInTimer Clear the flag once the timer callback function has been completed.
- static void [Update](#) ([Thread](#) *pclTargetThread_)
Update Update the current thread being tracked for round-robin scheduling. Note - this has no effect if called from the [Timer](#) thread, or if the [Timer](#) thread is active.
- static void [SetTimerThread](#) ([Thread](#) *pclTimerThread_)
SetTimerThread Pass the timer thread's [Thread](#) pointer to the [Quantum](#) module to track against requests to update the round-robin timer.
- static void [Cancel](#) ()

Static Private Attributes

- static [Thread](#) * [m_pclActiveThread](#)
- static [Thread](#) * [m_pclTimerThread](#)
- static uint16_t [m_u16TicksRemain](#)
- static bool [m_bInTimer](#)

19.28.1 Detailed Description

The [Quantum](#) Class. Static-class used to implement [Thread](#) quantum functionality, which is fundamental to round-robin thread scheduling.

Definition at line 42 of file [quantum.h](#).

19.28.2 Member Function Documentation

19.28.2.1 [Cancel\(\)](#)

```
static void Mark3::Quantum::Cancel ( ) [static]
```

Cancel the round-robin timer.

19.28.2.2 [ClearInTimer\(\)](#)

```
static void Mark3::Quantum::ClearInTimer ( ) [static]
```

[ClearInTimer](#) Clear the flag once the timer callback function has been completed.

19.28.2.3 Init()

```
static void Mark3::Quantum::Init ( ) [static]
```

19.28.2.4 SetInTimer()

```
static void Mark3::Quantum::SetInTimer ( ) [static]
```

SetInTimer Set a flag to indicate that the CPU is currently running within the timer-callback routine. This prevents the [Quantum](#) timer from being updated in the middle of a callback cycle, potentially resulting in the kernel timer becoming disabled.

19.28.2.5 SetTimerThread()

```
static void Mark3::Quantum::SetTimerThread (
    Thread * pclTimerThread_ ) [inline], [static]
```

SetTimerThread Pass the timer thread's [Thread](#) pointer to the [Quantum](#) module to track against requests to update the round-robin timer.

Parameters

<i>pclTimer↔ Thread_</i>	Pointer to the Timer thread's Thread object.
------------------------------	--

Definition at line [79](#) of file [quantum.h](#).

19.28.2.6 Update()

```
static void Mark3::Quantum::Update (
    Thread * pclTargetThread_ ) [static]
```

Update Update the current thread being tracked for round-robin scheduling. Note - this has no effect if called from the [Timer](#) thread, or if the [Timer](#) thread is active.

Parameters

<i>pclTarget↔ Thread_</i>	New thread to track.
-------------------------------	----------------------

19.28.3 Member Data Documentation

19.28.3.1 m_bInTimer

```
bool Mark3::Quantum::m_bInTimer [static], [private]
```

Definition at line 90 of file [quantum.h](#).

19.28.3.2 m_pclActiveThread

```
Thread* Mark3::Quantum::m_pclActiveThread [static], [private]
```

Definition at line 87 of file [quantum.h](#).

19.28.3.3 m_pclTimerThread

```
Thread* Mark3::Quantum::m_pclTimerThread [static], [private]
```

Definition at line 88 of file [quantum.h](#).

19.28.3.4 m_u16TicksRemain

```
uint16_t Mark3::Quantum::m_u16TicksRemain [static], [private]
```

Definition at line 89 of file [quantum.h](#).

The documentation for this class was generated from the following file:

- [/home/moslevin/projects/m3-repo/kernel/src/public/quantum.h](#)

19.29 Mark3::ReaderWriterLock Class Reference

The [ReaderWriterLock](#) class. This class implements an object that marshalls access to a resource based on the intended usage of the resource. A reader-writer lock permits multiple concurrent read access, or single-writer access to a resource. If the object holds a write lock, other writers, and all readers will block until the writer is finished. If the object holds reader locks, all writers will block until all readers are finished before the first writer can take ownership of the resource. This is based upon lower-level synchronization primitives, and is somewhat more heavyweight than primitive synchronization types.

```
#include <readerwriter.h>
```

Public Member Functions

- void * [operator new](#) (size_t sz, void *pv)
- void [Init](#) ()

Init Initialize the reader-writer lock before use. Must be called before attempting any other operations on the object.
- void [AcquireReader](#) ()

AcquireReader Acquire the object's reader lock. Multiple concurrent readers are allowed. If the writer lock is currently held, the calling thread will wait until the writer lock is relinquished.
- bool [AcquireReader](#) (uint32_t u32TimeoutMs_)

AcquireReader Acquire the object's reader lock. Multiple concurrent readers are allowed. If the writer lock is currently held, the calling thread will wait until the writer lock is relinquished.
- void [ReleaseReader](#) ()

ReleaseReader Release a previously-held reader lock.
- void [AcquireWriter](#) ()

AcquireWriter Acquire the writer lock. Only a single writer is allowed to access the object at a time. This will block the currently-running thread until all other readers/writers have released their locks.
- bool [AcquireWriter](#) (uint32_t u32TimeoutMs_)

AcquireWriter Acquire the writer lock. Only a single writer is allowed to access the object at a time. This will block the currently-running thread until all other readers/writers have released their locks.
- void [ReleaseWriter](#) ()

ReleaseWriter Release the currently held writer, allowing other readers/writers to access the object.

Private Member Functions

- bool [AcquireReader_i](#) (uint32_t u32TimeoutMs_)

AcquireReader_i Internal helper function for AcquireReader.
- bool [AcquireWriter_i](#) (uint32_t u32TimeoutMs_)

AcquireWriter_i Internal helper function for AcquireWriter.

Private Attributes

- [Mutex m_clGlobalMutex](#)

Mutex used to lock the object against concurrent read + write.
- [Mutex m_clReaderMutex](#)

Mutex used to lock object for readers.
- uint8_t [m_u8ReadCount](#)

Number of concurrent readers.

19.29.1 Detailed Description

The [ReaderWriterLock](#) class. This class implements an object that marshalls access to a resource based on the intended usage of the resource. A reader-writer lock permits multiple concurrent read access, or single-writer access to a resource. If the object holds a write lock, other writers, and all readers will block until the writer is finished. If the object holds reader locks, all writers will block until all readers are finished before the first writer can take ownership of the resource. This is based upon lower-level synchronization primitives, and is somewhat more heavyweight than primitive synchronization types.

Definition at line 40 of file [readerwriter.h](#).

19.29.2 Member Function Documentation

19.29.2.1 AcquireReader() [1/2]

```
void Mark3::ReaderWriterLock::AcquireReader ( )
```

AcquireReader Acquire the object's reader lock. Multiple concurrent readers are allowed. If the writer lock is currently held, the calling thread will wait until the writer lock is relinquished.

Definition at line 32 of file [readerwriter.cpp](#).

19.29.2.2 AcquireReader() [2/2]

```
bool Mark3::ReaderWriterLock::AcquireReader (
    uint32_t u32TimeoutMs_ )
```

AcquireReader Acquire the object's reader lock. Multiple concurrent readers are allowed. If the writer lock is currently held, the calling thread will wait until the writer lock is relinquished.

Parameters

<i>u32Timeout↵ Ms_</i>	Maximum time to wait (in ms) before the operation is aborted
----------------------------	--

Returns

true on success, false on timeout

Definition at line 38 of file [readerwriter.cpp](#).

19.29.2.3 AcquireReader_i()

```
bool Mark3::ReaderWriterLock::AcquireReader_i (
    uint32_t u32TimeoutMs_ ) [private]
```

AcquireReader_i Internal helper function for AcquireReaer.

Parameters

<i>u32Timeout↵ Ms_</i>	Maximum time to wait (in ms) before the operation is aborted
----------------------------	--

Returns

true on success, false on timeout

Definition at line 73 of file [readerwriter.cpp](#).

19.29.2.4 AcquireWriter() [1/2]

```
void Mark3::ReaderWriterLock::AcquireWriter ( )
```

AcquireWriter Acquire the writer lock. Only a single writer is allowed to access the object at a time. This will block the currently-runign thread until all other readers/writers have released their locks.

Definition at line 55 of file [readerwriter.cpp](#).

19.29.2.5 AcquireWriter() [2/2]

```
bool Mark3::ReaderWriterLock::AcquireWriter (
    uint32_t u32TimeoutMs_ )
```

AcquireWriter Acquire the writer lock. Only a single writer is allowed to access the object at a time. This will block the currently-runign thread until all other readers/writers have released their locks.

Parameters

<i>u32TimeoutMs_</i>	Maximum time to wait (in ms) before the operation is aborted
----------------------	--

Returns

true on success, false on timeout

Definition at line 61 of file [readerwriter.cpp](#).

19.29.2.6 AcquireWriter_i()

```
bool Mark3::ReaderWriterLock::AcquireWriter_i (
    uint32_t u32TimeoutMs_ ) [private]
```

AcquireWriter_i Internal helper function for **AcquireWriter**.

Parameters

<code>u32Timeout↵ Ms_</code>	Maximum time to wait (in ms) before the operation is aborted
----------------------------------	--

Returns

true on success, false on timeout

Definition at line 90 of file [readerwriter.cpp](#).

19.29.2.7 Init()

```
void Mark3::ReaderWriterLock::Init (
    void )
```

Init Initialize the reader-writer lock before use. Must be called before attempting any other operations on the object.

Definition at line 24 of file [readerwriter.cpp](#).

19.29.2.8 operator new()

```
void* Mark3::ReaderWriterLock::operator new (
    size_t sz,
    void * pv ) [inline]
```

Definition at line 43 of file [readerwriter.h](#).

19.29.2.9 ReleaseReader()

```
void Mark3::ReaderWriterLock::ReleaseReader ( )
```

ReleaseReader Release a previously-held reader lock.

Definition at line 44 of file [readerwriter.cpp](#).

19.29.2.10 ReleaseWriter()

```
void Mark3::ReaderWriterLock::ReleaseWriter ( )
```

ReleaseWriter Release the currently held writer, allowing other readers/writers to access the object.

Definition at line 67 of file [readerwriter.cpp](#).

19.29.3 Member Data Documentation

19.29.3.1 m_clGlobalMutex

`Mutex Mark3::ReaderWriterLock::m_clGlobalMutex [private]`

`Mutex` used to lock the object against concurrent read + write.

Definition at line 116 of file [readerwriter.h](#).

19.29.3.2 m_clReaderMutex

`Mutex Mark3::ReaderWriterLock::m_clReaderMutex [private]`

`Mutex` used to lock object for readers.

Definition at line 117 of file [readerwriter.h](#).

19.29.3.3 m_u8ReadCount

`uint8_t Mark3::ReaderWriterLock::m_u8ReadCount [private]`

Number of concurrent readers.

Definition at line 118 of file [readerwriter.h](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/m3-repo/kernel/src/public/readerwriter.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/readerwriter.cpp](#)

19.30 Mark3::Scheduler Class Reference

The `Scheduler` Class. This class provides priority-based round-robin `Thread` scheduling for all active threads managed by the kernel.

```
#include <scheduler.h>
```

Static Public Member Functions

- static void [Init](#) ()
Init Initialize the scheduler, must be called before use.
- static void [Schedule](#) ()
Schedule Run the scheduler, determines the next thread to run based on the current state of the threads. Note that the next-thread chosen from this function is only valid while in a critical section.
- static void [Add](#) ([Thread](#) *pciThread_)
Add Add a thread to the scheduler at its current priority level.
- static void [Remove](#) ([Thread](#) *pciThread_)
Remove Remove a thread from the scheduler at its current priority level.
- static bool [SetScheduler](#) (bool bEnable_)
SetScheduler Set the active state of the scheduler. When the scheduler is disabled, the next thread is never set; the currently running thread will run forever until the scheduler is enabled again. Care must be taken to ensure that we don't end up trying to block while the scheduler is disabled, otherwise the system ends up in an unusable state.
- static [Thread](#) * [GetCurrentThread](#) ()
GetCurrentThread Return the pointer to the currently-running thread.
- static volatile [Thread](#) * [GetNextThread](#) ()
GetNextThread Return the pointer to the thread that should run next, according to the last run of the scheduler.
- static [ThreadList](#) * [GetThreadList](#) ([PORT_PRIO_TYPE](#) uXPriority_)
GetThreadList Return the pointer to the active list of threads that are at the given priority level in the scheduler.
- static [ThreadList](#) * [GetStopList](#) ()
GetStopList Return the pointer to the list of threads that are in the scheduler's stopped state.
- static bool [IsEnabled](#) ()
IsEnabled Return the current state of the scheduler - whether or not scheudling is enabled or disabled.
- static void [QueueScheduler](#) ()
QueueScheduler Tell the kernel to perform a scheduling operation as soon as the scheduler is re-enabled.

Static Private Attributes

- static constexpr auto [m_uNumPriorities](#) = size_t { [KERNEL_NUM_PRIORITIES](#) }
- static bool [m_bEnabled](#)
Scheduler's state - enabled or disabled.
- static bool [m_bQueuedSchedule](#)
Variable representing whether or not there's a queued scheduler operation.
- static [ThreadList](#) [m_clStopList](#)
ThreadList for all stopped threads.
- static [ThreadList](#) [m_aclPriorities](#) [[m_uNumPriorities](#)]
ThreadLists for all threads at all priorities.
- static [PriorityMap](#) [m_clPrioMap](#)
Priority bitmap lookup structure, 1-bit per thread priority.

19.30.1 Detailed Description

The [Scheduler](#) Class. This class provides priority-based round-robin [Thread](#) scheduling for all active threads managed by the kernel.

Definition at line 63 of file [scheduler.h](#).

19.30.2 Member Function Documentation

19.30.2.1 Add()

```
void Mark3::Scheduler::Add (
    Thread * pclThread_ ) [static]
```

Add Add a thread to the scheduler at its current priority level.

Parameters

<i>pclThread_</i>	Pointer to the thread to add to the scheduler
-------------------	---

Definition at line 59 of file [scheduler.cpp](#).

19.30.2.2 GetCurrentThread()

```
static Thread* Mark3::Scheduler::GetCurrentThread ( ) [inline], [static]
```

GetCurrentThread Return the pointer to the currently-running thread.

Returns

Pointer to the currently-running thread

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 116 of file [scheduler.h](#).

19.30.2.3 GetNextThread()

```
static volatile Thread* Mark3::Scheduler::GetNextThread ( ) [inline], [static]
```

GetNextThread Return the pointer to the thread that should run next, according to the last run of the scheduler.

Returns

Pointer to the next-running thread

Definition at line 124 of file [scheduler.h](#).

19.30.2.4 GetStopList()

```
static ThreadList* Mark3::Scheduler::GetStopList ( ) [inline], [static]
```

GetStopList Return the pointer to the list of threads that are in the scheduler's stopped state.

Returns

Pointer to the [ThreadList](#) containing the stopped threads

Definition at line 142 of file [scheduler.h](#).

19.30.2.5 GetThreadList()

```
static ThreadList* Mark3::Scheduler::GetThreadList (
    PORT_PRIO_TYPE uXPriority_ ) [inline], [static]
```

GetThreadList Return the pointer to the active list of threads that are at the given priority level in the scheduler.

Parameters

<i>uX_↔</i> <i>Priority_</i>	Priority level of the threadlist
---	----------------------------------

Returns

Pointer to the [ThreadList](#) for the given priority level

Definition at line 134 of file [scheduler.h](#).

19.30.2.6 Init()

```
void Mark3::Scheduler::Init (
    void ) [static]
```

Init Initailize the scheduler, must be called before use.

Definition at line 36 of file [scheduler.cpp](#).

19.30.2.7 IsEnabled()

```
static bool Mark3::Scheduler::IsEnabled ( ) [inline], [static]
```

IsEnabled Return the current state of the scheduler - whether or not scheudling is enabled or disabled.

Returns

true - scheduler enabled, false - disabled

Definition at line 150 of file [scheduler.h](#).

19.30.2.8 QueueScheduler()

```
static void Mark3::Scheduler::QueueScheduler ( ) [inline], [static]
```

QueueScheduler Tell the kernel to perform a scheduling operation as soon as the scheduler is re-enabled.

Definition at line 156 of file [scheduler.h](#).

19.30.2.9 Remove()

```
void Mark3::Scheduler::Remove (
    Thread * pclThread_ ) [static]
```

Remove Remove a thread from the scheduler at its current priority level.

Parameters

<i>pcl</i> ↔ <i>Thread_</i>	Pointer to the thread to be removed from the scheduler
--------------------------------	--

Definition at line 67 of file [scheduler.cpp](#).

19.30.2.10 Schedule()

```
void Mark3::Scheduler::Schedule ( ) [static]
```

Schedule Run the scheduler, determines the next thread to run based on the current state of the threads. Note that the next-thread chosen from this function is only valid while in a critical section.

Definition at line 45 of file [scheduler.cpp](#).

19.30.2.11 SetScheduler()

```
bool Mark3::Scheduler::SetScheduler (
    bool bEnable_ ) [static]
```

SetScheduler Set the active state of the scheduler. When the scheduler is disabled, the *next thread* is never set; the currently running thread will run forever until the scheduler is enabled again. Care must be taken to ensure that we don't end up trying to block while the scheduler is disabled, otherwise the system ends up in an unusable state.

Parameters

<i>b↔ Enable↔ _</i>	true to enable, false to disable the scheduler
-----------------------------	--

Definition at line 75 of file [scheduler.cpp](#).

19.30.3 Member Data Documentation

19.30.3.1 m_aclPriorities

```
ThreadList Mark3::Scheduler::m_aclPriorities [static], [private]
```

ThreadLists for all threads at all priorities.

Definition at line 171 of file [scheduler.h](#).

19.30.3.2 m_bEnabled

```
bool Mark3::Scheduler::m_bEnabled [static], [private]
```

[Scheduler](#)'s state - enabled or disabled.

Definition at line 162 of file [scheduler.h](#).

19.30.3.3 m_bQueuedSchedule

```
bool Mark3::Scheduler::m_bQueuedSchedule [static], [private]
```

Variable representing whether or not there's a queued scheduler operation.

Definition at line 165 of file [scheduler.h](#).

19.30.3.4 m_clPrioMap

`PriorityMap` Mark3::Scheduler::m_clPrioMap [static], [private]

Priority bitmap lookup structure, 1-bit per thread priority.

Definition at line 174 of file [scheduler.h](#).

19.30.3.5 m_clStopList

`ThreadList` Mark3::Scheduler::m_clStopList [static], [private]

`ThreadList` for all stopped threads.

Definition at line 168 of file [scheduler.h](#).

19.30.3.6 m_uNumPriorities

```
constexpr auto Mark3::Scheduler::m_uNumPriorities = size_t { KERNEL_NUM_PRIORITIES } [static],
[private]
```

Definition at line 159 of file [scheduler.h](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/m3-repo/kernel/src/public/scheduler.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/scheduler.cpp](#)

19.31 Mark3::SchedulerGuard Class Reference

The [SchedulerGuard](#) class This class implements RAIL-based control of the scheduler's global state. Upon object construction, the scheduler's state is cached locally and the scheduler is disabled (if not already disabled). Upon object destruction, the scheduler's previous state is restored. This object is interrupt-safe, although it has no effect when called from an interrupt given that interrupts are inherently higher-priority than threads.

```
#include <schedulerguard.h>
```

Public Member Functions

- [SchedulerGuard](#) ()
- [~SchedulerGuard](#) ()

Private Attributes

- bool [m_bSchedState](#)

19.31.1 Detailed Description

The [SchedulerGuard](#) class This class implements RAII-based control of the scheduler's global state. Upon object construction, the scheduler's state is cached locally and the scheduler is disabled (if not already disabled). Upon object destruction, the scheduler's previous state is restored. This object is interrupt-safe, although it has no effect when called from an interrupt given that interrupts are inherently higher-priority than threads.

Definition at line 37 of file [schedulerguard.h](#).

19.31.2 Constructor & Destructor Documentation

19.31.2.1 SchedulerGuard()

```
Mark3::SchedulerGuard::SchedulerGuard ( ) [inline]
```

Definition at line 39 of file [schedulerguard.h](#).

19.31.2.2 ~SchedulerGuard()

```
Mark3::SchedulerGuard::~~SchedulerGuard ( ) [inline]
```

Definition at line 44 of file [schedulerguard.h](#).

19.31.3 Member Data Documentation

19.31.3.1 m_bSchedState

```
bool Mark3::SchedulerGuard::m_bSchedState [private]
```

Definition at line 50 of file [schedulerguard.h](#).

The documentation for this class was generated from the following file:


- [/home/moslevin/projects/m3-repo/kernel/src/public/schedulerguard.h](#)

19.32 Mark3::Semaphore Class Reference

the [Semaphore](#) class provides Binary & Counting semaphore objects, based on [BlockingObject](#) base class.

```
#include <ksemaphore.h>
```

Inheritance diagram for Mark3::Semaphore:



Public Member Functions

- void * [operator new](#) (size_t sz, void *pv)
- ~[Semaphore](#) ()
- void [Init](#) (uint16_t u16InitVal_, uint16_t u16MaxVal_)

Initialize a semaphore before use. Must be called before attempting post/pend operations on the object.
- bool [Post](#) ()

Increment the semaphore count. If the semaphore count is zero at the time this is called, and there are threads blocked on the object, this will immediately unblock the highest-priority blocked [Thread](#).
- void [Pend](#) ()

Decrement the semaphore count. If the count is zero, the calling [Thread](#) will block until the semaphore is posted, and the [Thread](#)'s priority is higher than that of any other [Thread](#) blocked on the object.
- uint16_t [GetCount](#) ()

Return the current semaphore counter. This can be used by a thread to bypass blocking on a semaphore - allowing it to do other things until a non-zero count is returned, instead of blocking until the semaphore is posted.
- bool [Pend](#) (uint32_t u32WaitTimeMS_)

Decrement the semaphore count. If the count is zero, the thread will block until the semaphore is pended. If the specified interval expires before the thread is unblocked, then the status is returned back to the user.
- void [WakeMe](#) ([Thread](#) *pclChosenOne_)

Wake a thread blocked on the semaphore. This is an internal function used for implementing timed semaphores relying on timer callbacks. Since these do not have access to the private data of the semaphore and its base classes, we have to wrap this as a public method - do not use this for any other purposes.

Private Member Functions

- uint8_t [WakeNext](#) ()

Wake the next thread waiting on the semaphore. Used internally.
- bool [Pend_i](#) (uint32_t u32WaitTimeMS_)

Pend_i.

Private Attributes

- uint16_t [m_u16Value](#)

Current count held by the semaphore.
- uint16_t [m_u16MaxValue](#)

Maximum count that can be held by this semaphore.

Additional Inherited Members

19.32.1 Detailed Description

the [Semaphore](#) class provides Binary & Counting semaphore objects, based on [BlockingObject](#) base class.

Examples:

[lab4_semaphores/main.cpp](#), [lab6_timers/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 36 of file [ksemaphore.h](#).

19.32.2 Constructor & Destructor Documentation

19.32.2.1 ~Semaphore()

```
Mark3::Semaphore::~~Semaphore ( )
```

Definition at line 58 of file [ksemaphore.cpp](#).

19.32.3 Member Function Documentation

19.32.3.1 GetCount()

```
uint16_t Mark3::Semaphore::GetCount ( )
```

Return the current semaphore counter. This can be usedd by a thread to bypass blocking on a semaphore - allowing it to do other things until a non-zero count is returned, instead of blocking until the semaphore is posted.

Returns

The current semaphore counter value.

Definition at line 206 of file [ksemaphore.cpp](#).

19.32.3.2 Init()

```
void Mark3::Semaphore::Init (
    uint16_t u16InitVal_,
    uint16_t u16MaxVal_ )
```

Initialize a semaphore before use. Must be called before attempting post/pend operations on the object.

This initialization is required to configure the behavior of the semaphore with regards to the initial and maximum values held by the semaphore. By providing access to the raw initial and maximum count elements of the semaphore, these objects are able to be used as either counting or binary semaphores.

To configure a semaphore object for use as a binary semaphore, set values of 0 and 1 respectively for the initial/maximum value parameters.

Any other combination of values can be used to implement a counting semaphore.

Parameters

<i>u16InitVal</i> ↔ —	Initial value held by the semaphore
<i>u16MaxVal</i> ↔ Val_	Maximum value for the semaphore. Must be nonzero.

Examples:

[lab6_timers/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 94 of file [ksemaphore.cpp](#).

19.32.3.3 operator new()

```
void* Mark3::Semaphore::operator new (
    size_t sz,
    void * pv ) [inline]
```

Definition at line 39 of file [ksemaphore.h](#).

19.32.3.4 Pend() [1/2]

```
void Mark3::Semaphore::Pend ( )
```

Decrement the semaphore count. If the count is zero, the calling [Thread](#) will block until the semaphore is posted, and the [Thread](#)'s priority is higher than that of any other [Thread](#) blocked on the object.

Examples:

[lab4_semaphores/main.cpp](#).

Definition at line 194 of file [ksemaphore.cpp](#).

19.32.3.5 Pend() [2/2]

```
bool Mark3::Semaphore::Pend (
    uint32_t u32WaitTimeMS_ )
```

Decrement the semaphore count. If the count is zero, the thread will block until the semaphore is pended. If the specified interval expires before the thread is unblocked, then the status is returned back to the user.

Returns

true - semaphore was acquired before the timeout
false - timeout occurred before the semaphore was claimed.

Definition at line 200 of file [ksemaphore.cpp](#).

19.32.3.6 Pend_i()

```
bool Mark3::Semaphore::Pend_i (
    uint32_t u32WaitTimeMS_ ) [private]
```

Pend_i.

Internal function used to abstract timed and untimed semaphore pend operations.

Parameters

<code>u32WaitTimeMS_</code>	Time in MS to wait
-----------------------------	--------------------

Returns

true on success, false on failure.

Definition at line 152 of file [ksemaphore.cpp](#).

19.32.3.7 Post()

```
bool Mark3::Semaphore::Post ( )
```

Increment the semaphore count. If the semaphore count is zero at the time this is called, and there are threads blocked on the object, this will immediately unblock the highest-priority blocked [Thread](#).

Note that if the priority of that [Thread](#) is higher than the current thread's priority, a context switch will occur and control will be relinquished to that [Thread](#).

Returns

true if the semaphore was posted, false if the count is already maxed out.

Examples:

[lab4_semaphores/main.cpp](#), and [lab6_timers/main.cpp](#).

Definition at line 108 of file [ksemaphore.cpp](#).

19.32.3.8 WakeMe()

```
void Mark3::Semaphore::WakeMe (
    Thread * pChosenOne_ )
```

Wake a thread blocked on the semaphore. This is an internal function used for implementing timed semaphores relying on timer callbacks. Since these do not have access to the private data of the semaphore and its base classes, we have to wrap this as a public method - do not use this for any other purposes.

Definition at line 68 of file [ksemaphore.cpp](#).

19.32.3.9 WakeNext()

```
uint8_t Mark3::Semaphore::WakeNext ( ) [private]
```

Wake the next thread waiting on the semaphore. Used internally.

Definition at line 78 of file [ksemaphore.cpp](#).

19.32.4 Member Data Documentation

19.32.4.1 m_u16MaxValue

```
uint16_t Mark3::Semaphore::m_u16MaxValue [private]
```

Maximum count that can be held by this semaphore.

Definition at line 140 of file [ksemaphore.h](#).

19.32.4.2 m_u16Value

```
uint16_t Mark3::Semaphore::m_u16Value [private]
```

Current count held by the semaphore.

Definition at line 139 of file [ksemaphore.h](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/m3-repo/kernel/src/public/ksemaphore.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/ksemaphore.cpp](#)

19.33 Mark3::Streamer Class Reference

The [Streamer](#) class. This class implements a circular byte-buffer with thread and interrupt safe methods for writing-to and reading-from the buffer. Objects of this class type are designed to be shared between threads, or between threads and interrupts.

```
#include <streamer.h>
```

Public Member Functions

- void **Init** (uint8_t *pau8Buffer_, uint16_t u16Size_)
Init. Initialize the [Streamer](#) object prior to its use, providing a blob of memory for the object to manage.
- bool **Read** (uint8_t *pu8Data_)
Read. Read a byte of data from the stream, if available.
- uint16_t **Read** (uint8_t *pu8Data_, uint16_t u16Len_)
Read. Read from the stream until a specified number of bytes have been read, or the stream is exhausted.
- bool **Write** (uint8_t u8Data_)
Write. Write a byte of data into the stream.
- uint16_t **Write** (uint8_t *pu8Data_, uint16_t u16Len_)
Write. Write a maximum number of bytes to the stream.
- bool **Claim** (uint8_t **pu8Addr_)
Claim. Claim a byte of data for writing, without actually writing into it. When the writer is ready to write into the data byte as a result of another operation, it can then populate the byte.
- void **Lock** (uint8_t *pu8LockAddr_)
Lock. When the lock is set, a client can neither read from, or write to the buffer at the index specified. This is used to in conjunction with Claim to safely reserve data from the buffer, while preventing race conditions occurring as a result of a consumer acting on the data before it is ready.
- void **Unlock** (void)
Unlock. Reset the lock pointer in the object, allowing a consumer to read any previously unavailable data that might still be in the stream.
- uint16_t **GetAvailable** (void)
GetAvailable.
- bool **CanRead** (void)
CanRead.
- bool **CanWrite** (void)
CanWrite.
- bool **IsEmpty** (void)
IsEmpty.

Private Attributes

- uint8_t * **m_pau8Buffer**
Pointer to the buffer managed in this object.
- uint8_t * **m_pu8LockAddr**
Address of the lock point in the stream.
- uint16_t **m_u16Size**
Size of the stream's circular buffer (in bytes)
- uint16_t **m_u16Avail**
Number of bytes free in the stream.
- uint16_t **m_u16Head**
Current head index (write to) of the stream.
- uint16_t **m_u16Tail**
Current tail index (read from) of the stream.

19.33.1 Detailed Description

The [Streamer](#) class. This class implements a circular byte-buffer with thread and interrupt safe methods for writing-to and reading-from the buffer. Objects of this class type are designed to be shared between threads, or between threads and interrupts.

Definition at line 35 of file [streamer.h](#).

19.33.2 Member Function Documentation

19.33.2.1 CanRead()

```
bool Mark3::Streamer::CanRead (
    void )
```

CanRead.

Returns

true if the stream has any unread data

Definition at line 200 of file [streamer.cpp](#).

19.33.2.2 CanWrite()

```
bool Mark3::Streamer::CanWrite (
    void )
```

CanWrite.

Returns

true if the stream has any free space

Definition at line 211 of file [streamer.cpp](#).

19.33.2.3 Claim()

```
bool Mark3::Streamer::Claim (
    uint8_t ** pu8Addr_ )
```

Claim. Claim a byte of data for writing, without actually writing into it. When the writer is ready to write into the data byte as a result of another operation, it can then populate the byte.

This method is useful when encoding data from a raw format into a packet-based format, where one byte of input may result in multiple bytes of output being generated. Especially in cases where a user wants to write to a stream while a peripheral reads from it to transmit its contents asynchronously (i.e. an interrupt-driven UART consuming a data packet, where the packet is being framed and consumed simultaneously).

This should be used in conjunction with the Lock method and judicious use of critical sections to prevent race conditions.

Parameters

<i>pU8_</i> <i>Addr_</i>	[out] Pointer to a byte pointer that will contain the address of the "claimed" byte on success.
-----------------------------	---

Returns

true if successful, false if buffer full or locked.

Definition at line 224 of file [streamer.cpp](#).

19.33.2.4 GetAvailable()

```
uint16_t Mark3::Streamer::GetAvailable (
    void ) [inline]
```

GetAvailable.

Returns

The current number of bytes available for write in the streams

Definition at line 134 of file [streamer.h](#).

19.33.2.5 Init()

```
void Mark3::Streamer::Init (
    uint8_t * pau8Buffer_,
    uint16_t u16Size_ )
```

Init. Initialize the [Streamer](#) object prior to its use, providing a blob of memory for the object to manage.

Parameters

<i>pau8_</i> <i>Buffer_</i>	Blob of memory to use as a circular buffer
<i>u16Size_</i>	Size of the supplied buffer in bytes

Definition at line 27 of file [streamer.cpp](#).

19.33.2.6 IsEmpty()

```
bool Mark3::Streamer::IsEmpty (
    void )
```

IsEmpty.

Returns

true if the stream is empty

Definition at line 264 of file [streamer.cpp](#).

19.33.2.7 Lock()

```
void Mark3::Streamer::Lock (
    uint8_t * pu8LockAddr_ )
```

Lock. When the lock is set, a client can neither read from, or write to the buffer at the index specified. This is used to in conjunction with Claim to safely reserve data from the buffer, while preventing race conditions occurring as a result of a consumer acting on the data before it is ready.

Parameters

<i>pu8Lock↵ Addr_</i>	Address (within the stream) to set as the lockpoint.
---------------------------	--

Definition at line 250 of file [streamer.cpp](#).

19.33.2.8 Read() [1/2]

```
bool Mark3::Streamer::Read (
    uint8_t * pu8Data_ )
```

Read. Read a byte of data from the stream, if available.

Parameters

<i>pu8↵ Data_</i>	Pointer to read data into from the stream
-----------------------	---

Returns

true if data was read, false if data unavailable or buffer locked

Definition at line 38 of file [streamer.cpp](#).

19.33.2.9 Read() [2/2]

```
uint16_t Mark3::Streamer::Read (
    uint8_t * pu8Data_,
    uint16_t u16Len_ )
```

Read. Read from the stream until a specified number of bytes have been read, or the stream is exhausted.

Parameters

<i>pu8Data_</i>	pointer to an array of data read into
<i>u16Len_</i>	maximum number of bytes to read

Returns

number of bytes read

Definition at line 64 of file [streamer.cpp](#).

19.33.2.10 Unlock()

```
void Mark3::Streamer::Unlock (
    void )
```

Unlock. Reset the lock pointer in the object, allowing a consumer to read any previously unavailable data that might still be in the stream.

Definition at line 257 of file [streamer.cpp](#).

19.33.2.11 Write() [1/2]

```
bool Mark3::Streamer::Write (
    uint8_t u8Data_ )
```

Write. Write a byte of data into the stream.

Parameters

<i>u8Data_</i>	Data byte to be written into the stream
----------------	---

Returns

true if byte was written, false on buffer full or buffer locked at index.

Definition at line 120 of file [streamer.cpp](#).

19.33.2.12 Write() [2/2]

```
uint16_t Mark3::Streamer::Write (
    uint8_t * pu8Data_,
    uint16_t u16Len_ )
```

Write. Write a maximum number of bytes to the stream.

Parameters

<i>pu8Data_</i>	pointer to an array of bytes to write out to the stream
<i>u16Len_</i>	Length of data held in the array

Returns

number of bytes written to the stream

Definition at line 144 of file [streamer.cpp](#).

19.33.3 Member Data Documentation

19.33.3.1 m_pau8Buffer

```
uint8_t* Mark3::Streamer::m_pau8Buffer [private]
```

Pointer to the buffer managed in this object.

Definition at line 154 of file [streamer.h](#).

19.33.3.2 m_pu8LockAddr

```
uint8_t* Mark3::Streamer::m_pu8LockAddr [private]
```

Address of the lock point in the stream.

Definition at line 155 of file [streamer.h](#).

19.33.3.3 m_u16Avail

```
uint16_t Mark3::Streamer::m_u16Avail [private]
```

Number of bytes free in the stream.

Definition at line 158 of file [streamer.h](#).

19.33.3.4 m_u16Head

```
uint16_t Mark3::Streamer::m_u16Head [private]
```

Current head index (write to) of the stream.

Definition at line 159 of file [streamer.h](#).

19.33.3.5 m_u16Size

```
uint16_t Mark3::Streamer::m_u16Size [private]
```

Size of the stream's circular buffer (in bytes)

Definition at line 157 of file [streamer.h](#).

19.33.3.6 m_u16Tail

```
uint16_t Mark3::Streamer::m_u16Tail [private]
```

Current tail index (read from) of the stream.

Definition at line 160 of file [streamer.h](#).

The documentation for this class was generated from the following files:


- [/home/moslevin/projects/m3-repo/kernel/lib/streamer/public/streamer.h](#)
- [/home/moslevin/projects/m3-repo/kernel/lib/streamer/streamer.cpp](#)

19.34 Mark3::Thread Class Reference

The [Thread](#) Class. This object providing the fundamental thread control data structures and functions that define a single thread of execution in the [Mark3](#) operating system. It is the fundamental data type used to provide multitasking support in the kernel.

```
#include <thread.h>
```

Inheritance diagram for Mark3::Thread:



Public Member Functions

- void * [operator new](#) (size_t sz, void *pv)
- [~Thread](#) ()
- [Thread](#) ()
- bool [IsInitialized](#) ()
 - IsInitialized* Used to check whether or not a thread has been initialized prior to use.
- void [Init](#) (K_WORD *pwStack_, uint16_t u16StackSize_, PORT_PRIO_TYPE uXPriority_, [ThreadEntryFunc](#) pfEntryPoint_, void *pvArg_)
 - Init* Initialize a thread prior to its use. Initialized threads are placed in the stopped state, and are not scheduled until the thread's start method has been invoked first.
- void [Start](#) ()
 - Start* Start the thread - remove it from the stopped list, add it to the scheduler's list of threads (at the thread's set priority), and continue along.
- void [Stop](#) ()
 - Stop* Stop a thread that's actively scheduled without destroying its stacks. Stopped threads can be restarted using the [Start\(\)](#) API.
- void [SetName](#) (const char *szName_)
 - SetName* Set the name of the thread - this is purely optional, but can be useful when identifying issues that come along when multiple threads are at play in a system.
- const char * [GetName](#) ()
 - GetName.*
- [ThreadList](#) * [GetOwner](#) (void)
 - GetOwner* Return the [ThreadList](#) where the thread belongs when it's in the active/ready state in the scheduler.
- [ThreadList](#) * [GetCurrent](#) (void)
 - GetCurrent* Return the [ThreadList](#) where the thread is currently located.
- PORT_PRIO_TYPE [GetPriority](#) (void)
 - GetPriority* Return the priority of the current thread.
- PORT_PRIO_TYPE [GetCurPriority](#) (void)
 - GetCurPriority* Return the priority of the current thread.
- void [SetQuantum](#) (uint16_t u16Quantum_)
 - SetQuantum* Set the thread's round-robin execution quantum.
- uint16_t [GetQuantum](#) (void)
 - GetQuantum* Get the thread's round-robin execution quantum.

- void [SetCurrent](#) ([ThreadList](#) *pclNewList_)
SetCurrent. Set the thread's current to the specified thread list.
- void [SetOwner](#) ([ThreadList](#) *pclNewList_)
SetOwner. Set the thread's owner to the specified thread list.
- void [SetPriority](#) ([PORT_PRIO_TYPE](#) uXPriority_)
SetPriority. Set the priority of the [Thread](#) (running or otherwise) to a different level. This activity involves re-scheduling, and must be done so with due caution, as it may effect the determinism of the system.
- void [InheritPriority](#) ([PORT_PRIO_TYPE](#) uXPriority_)
InheritPriority Allow the thread to run at a different priority level (temporarily) for the purpose of avoiding priority inversions. This should only be called from within the implementation of blocking-objects.
- void [Exit](#) ()
Exit. Remove the thread from being scheduled again. The thread is effectively destroyed when this occurs. This is extremely useful for cases where a thread encounters an unrecoverable error and needs to be restarted, or in the context of systems where threads need to be created and destroyed dynamically.
- void [SetID](#) (uint8_t u8ID_)
SetID Set an arbitrary 8-bit ID to uniquely identify this thread.
- uint8_t [GetID](#) ()
GetID Return the thread's integer ID. Note that this ID is not guaranteed to be unique when dynamic threading is used in the system, or there are more than 255 threads. Also not guaranteed to be unique if the [SetID](#) function is called by the user.
- uint16_t [GetStackSlack](#) ()
GetStackSlack Performs a (somewhat lengthy) check on the thread stack to check the amount of stack margin (or "slack") remaining on the stack. If you're having problems with blowing your stack, you can run this function at points in your code during development to see what operations cause problems. Also useful during development as a tool to optimally size thread stacks.
- uint16_t [GetEventFlagMask](#) ()
GetEventFlagMask returns the thread's current event-flag mask, which is used in conjunction with the [EventFlag](#) blocking object type.
- void [SetEventFlagMask](#) (uint16_t u16Mask_)
SetEventFlagMask Sets the active event flag bitfield mask.
- void [SetEventFlagMode](#) ([EventFlagOperation](#) eMode_)
SetEventFlagMode Sets the active event flag operation mode.
- [EventFlagOperation](#) [GetEventFlagMode](#) ()
GetEventFlagMode Returns the thread's event flag's operating mode.
- [Timer](#) * [GetTimer](#) ()
- void [SetExpired](#) (bool bExpired_)
SetExpired Set the status of the current blocking call on the thread.
- bool [GetExpired](#) ()
GetExpired Return the status of the most-recent blocking call on the thread.
- void * [GetExtendedContext](#) ()
GetExtendedContext Return the [Thread](#) object's extended-context data pointer. Used by code implementing a user-defined thread-local storage model. Pointer exists only for the lifespan of the [Thread](#).
- void [SetExtendedContext](#) (void *pvData_)
SetExtendedContext Assign the [Thread](#) object's extended-context data pointer. Used by code implementing a user-defined thread-local storage model.
- [ThreadState](#) [GetState](#) ()
GetState Returns the current state of the thread to the caller. Can be used to determine whether or not a thread is ready (or running), stopped, or terminated/exit'd.
- void [SetState](#) ([ThreadState](#) eState_)
SetState Set the thread's state to a new value. This is only to be used by code within the kernel, and is not intended for use by an end-user.
- [K_WORD](#) * [GetStack](#) ()
GetStack.
- uint16_t [GetStackSize](#) ()
GetStackSize.

Static Public Member Functions

- static [Thread](#) * [Init](#) (uint16_t u16StackSize_, [PORT_PRIO_TYPE](#) uXPriority_, [ThreadEntryFunc](#) pfEntry↵
Point_, void *pvArg_)

Init Create and initialize a new thread, using memory from the auto-allocated heap region to supply both the thread object and its stack. The thread returned can then be started using the [Start\(\)](#) method directly. Note that the memory used to create this thread cannot be reclaimed, and so this API is only suitable for threads that exist for the duration of runtime.
- static void [Sleep](#) (uint32_t u32TimeMs_)

Sleep Put the thread to sleep for the specified time (in milliseconds). Actual time slept may be longer (but not less than) the interval specified.
- static void [Yield](#) (void)

Yield Yield the thread - this forces the system to call the scheduler and determine what thread should run next. This is typically used when threads are moved in and out of the scheduler.
- static void [CoopYield](#) (void)

CoopYield Cooperative yield - This forces the system to not only call the scheduler, but also move the currently executing thread to the back of the current thread list, allowing other same-priority threads the opportunity to run. This is used primarily for cooperative scheduling between threads in the same priority level.

Private Member Functions

- void [SetPriorityBase](#) ([PORT_PRIO_TYPE](#) uXPriority_)

SetPriorityBase.

Static Private Member Functions

- static void [ContextSwitchSWI](#) (void)

ContextSwitchSWI This code is used to trigger the context switch interrupt. Called whenever the kernel decides that it is necessary to swap out the current thread for the "next" thread.

Private Attributes

- [K_WORD](#) * [m_pwStackTop](#)

Pointer to the top of the thread's stack.
- [K_WORD](#) * [m_pwStack](#)

Pointer to the thread's stack.
- [uint8_t](#) [m_u8ThreadID](#)

Thread ID.
- [PORT_PRIO_TYPE](#) [m_uXPriority](#)

Default priority of the thread.
- [PORT_PRIO_TYPE](#) [m_uXCurPriority](#)

Current priority of the thread (priority inheritance)
- [ThreadState](#) [m_eState](#)

Enum indicating the thread's current state.
- void * [m_pvExtendedContext](#)

Pointer provided to a [Thread](#) to implement thread-local storage.
- const char * [m_szName](#)

Thread name.
- [uint16_t](#) [m_u16StackSize](#)

Size of the stack (in bytes)
- [ThreadList](#) * [m_pclCurrent](#)

- Pointer to the thread-list where the thread currently resides.*
- [ThreadList * m_pclOwner](#)
Pointer to the thread-list where the thread resides when active.
- [ThreadEntryFunc m_pfEntryPoint](#)
The entry-point function called when the thread starts.
- `void * m_pvArg`
Pointer to the argument passed into the thread's entryptoint.
- `uint16_t m_u16Quantum`
Thread quantum (in milliseconds)
- `uint16_t m_u16FlagMask`
Event-flag mask.
- [EventFlagOperation m_eFlagMode](#)
Event-flag mode.
- [Timer m_clTimer](#)
Timer used for blocking-object timeouts.
- `bool m_bExpired`
Indicate whether or not a blocking-object timeout has occurred.

Friends

- class [ThreadPort](#)

Additional Inherited Members

19.34.1 Detailed Description

The [Thread](#) Class. This object providing the fundamental thread control data structures and functions that define a single thread of execution in the [Mark3](#) operating system. It is the fundamental data type used to provide multitasking support in the kernel.

Examples:

[lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_coroutines/main.cpp](#), [lab3_round_robin/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab5_mutexes/main.cpp](#), [lab6_timers/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 64 of file [thread.h](#).

19.34.2 Constructor & Destructor Documentation

19.34.2.1 ~Thread()

```
Mark3::Thread::~~Thread ( )
```

Definition at line 30 of file [thread.cpp](#).

19.34.2.2 Thread()

```
Mark3::Thread::Thread ( ) [inline]
```

Definition at line 70 of file [thread.h](#).

19.34.3 Member Function Documentation

19.34.3.1 ContextSwitchSWI()

```
void Mark3::Thread::ContextSwitchSWI (
    void ) [static], [private]
```

ContextSwitchSWI This code is used to trigger the context switch interrupt. Called whenever the kernel decides that it is necessary to swap out the current thread for the "next" thread.

Definition at line 392 of file [thread.cpp](#).

19.34.3.2 CoopYield()

```
void Mark3::Thread::CoopYield (
    void ) [static]
```

CoopYield Cooperative yield - This forces the system to not only call the scheduler, but also move the currently executing thread to the back of the current thread list, allowing other same-priority threads the opportunity to run. This is used primarily for cooperative scheduling between threads in the same priority level.

Definition at line 325 of file [thread.cpp](#).

19.34.3.3 Exit()

```
void Mark3::Thread::Exit ( )
```

Exit. Remove the thread from being scheduled again. The thread is effectively destroyed when this occurs. This is extremely useful for cases where a thread encounters an unrecoverable error and needs to be restarted, or in the context of systems where threads need to be created and destroyed dynamically.

This must not be called on the idle thread.

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 187 of file [thread.cpp](#).

19.34.3.4 GetCurPriority()

```
PORT_PRIO_TYPE Mark3::Thread::GetCurPriority (
    void ) [inline]
```

GetCurPriority Return the priority of the current thread.

Returns

Priority of the current thread

Definition at line 181 of file [thread.h](#).

19.34.3.5 GetCurrent()

```
ThreadList* Mark3::Thread::GetCurrent (
    void ) [inline]
```

GetCurrent Return the [ThreadList](#) where the thread is currently located.

Returns

Pointer to the thread's current list

Definition at line 166 of file [thread.h](#).

19.34.3.6 GetEventFlagMask()

```
uint16_t Mark3::Thread::GetEventFlagMask ( ) [inline]
```

GetEventFlagMask returns the thread's current event-flag mask, which is used in conjunction with the [EventFlag](#) blocking object type.

Returns

A copy of the thread's event flag mask

Definition at line 317 of file [thread.h](#).

19.34.3.7 GetEventFlagMode()

```
EventFlagOperation Mark3::Thread::GetEventFlagMode ( ) [inline]
```

GetEventFlagMode Returns the thread's event flag's operating mode.

Returns

The thread's event flag mode.

Definition at line 336 of file [thread.h](#).

19.34.3.8 GetExpired()

```
bool Mark3::Thread::GetExpired ( )
```

GetExpired Return the status of the most-recent blocking call on the thread.

Returns

true - call expired, false - call did not expire

Definition at line 425 of file [thread.cpp](#).

19.34.3.9 GetExtendedContext()

```
void* Mark3::Thread::GetExtendedContext ( ) [inline]
```

GetExtendedContext Return the [Thread](#) object's extended-context data pointer. Used by code implementing a user-defined thread-local storage model. Pointer exists only for the lifespan of the [Thread](#).

Returns

[Thread](#)'s extended context data pointer.

Definition at line 368 of file [thread.h](#).

19.34.3.10 GetID()

```
uint8_t Mark3::Thread::GetID ( ) [inline]
```

GetID Return the thread's integer ID. Note that this ID is not guaranteed to be unique when dynamic threading is used in the system, or there are more than 255 threads. Also not guaranteed to be unique if the SetID function is called by the user.

Returns

[Thread](#)'s 8-bit ID, set by the user

Definition at line 292 of file [thread.h](#).

19.34.3.11 GetName()

```
const char* Mark3::Thread::GetName ( ) [inline]
```

GetName.

Returns

Pointer to the name of the thread. If this is not set, will be nullptr.

Definition at line 149 of file [thread.h](#).

19.34.3.12 GetOwner()

```
ThreadList* Mark3::Thread::GetOwner (
    void ) [inline]
```

GetOwner Return the [ThreadList](#) where the thread belongs when it's in the active/ready state in the scheduler.

Returns

Pointer to the [Thread](#)'s owner list

Definition at line 159 of file [thread.h](#).

19.34.3.13 GetPriority()

```
PORT_PRIO_TYPE Mark3::Thread::GetPriority (
    void ) [inline]
```

GetPriority Return the priority of the current thread.

Returns

Priority of the current thread

Definition at line 174 of file [thread.h](#).

19.34.3.14 GetQuantum()

```
uint16_t Mark3::Thread::GetQuantum (
    void ) [inline]
```

GetQuantum Get the thread's round-robin execution quantum.

Returns

The thread's quantum

Definition at line 197 of file [thread.h](#).

19.34.3.15 GetStack()

```
K_WORD* Mark3::Thread::GetStack ( ) [inline]
```

GetStack.

Returns

Pointer to the blob of memory used as the thread's stack

Definition at line 403 of file [thread.h](#).

19.34.3.16 GetStackSize()

```
uint16_t Mark3::Thread::GetStackSize ( ) [inline]
```

GetStackSize.

Returns

Size of the thread's stack in bytes

Definition at line 409 of file [thread.h](#).

19.34.3.17 GetStackSlack()

```
uint16_t Mark3::Thread::GetStackSlack ( )
```

GetStackSlack Performs a (somewhat lengthy) check on the thread stack to check the amount of stack margin (or "slack") remaining on the stack. If you're having problems with blowing your stack, you can run this function at points in your code during development to see what operations cause problems. Also useful during development as a tool to optimally size thread stacks.

Returns

The amount of slack (unused bytes) on the stack

Examples:

[lab9_dynamic_threads/main.cpp](#).

19.34.3.18 GetState()

```
ThreadState Mark3::Thread::GetState ( ) [inline]
```

GetState Returns the current state of the thread to the caller. Can be used to determine whether or not a thread is ready (or running), stopped, or terminated/exit'd.

Returns

ThreadState_t representing the thread's current state

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 389 of file [thread.h](#).

19.34.3.19 GetTimer()

```
Timer * Mark3::Thread::GetTimer ( )
```

Return a pointer to the thread's timer object

Definition at line 412 of file [thread.cpp](#).

19.34.3.20 InheritPriority()

```
void Mark3::Thread::InheritPriority (
    PORT_PRIO_TYPE uXPriority_ )
```

InheritPriority Allow the thread to run at a different priority level (temporarily) for the purpose of avoiding priority inversions. This should only be called from within the implementation of blocking-objects.

Parameters

<i>uX_↔</i> <i>Priority_</i>	New Priority to boost to.
---	---------------------------

Definition at line 383 of file [thread.cpp](#).

19.34.3.21 Init() [1/2]

```
void Mark3::Thread::Init (
    K_WORD * pwStack_,
    uint16_t u16StackSize_,
    PORT_PRIO_TYPE uXPriority_,
    ThreadEntryFunc pfEntryPoint_,
    void * pvArg_ )
```

Init Initialize a thread prior to its use. Initialized threads are placed in the stopped state, and are not scheduled until the thread's start method has been invoked first.

Parameters

<i>pwStack_</i>	Pointer to the stack to use for the thread
<i>u16Stack_↔</i> <i>Size_</i>	Size of the stack (in bytes)
<i>uXPriority_</i>	Priority of the thread (0 = idle, 7 = max)
<i>pfEntryPoint_↔</i> —	This is the function that gets called when the thread is started
<i>pvArg_</i>	Pointer to the argument passed into the thread's entryptoint function.

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 49 of file [thread.cpp](#).

19.34.3.22 Init() [2/2]

```
Thread * Mark3::Thread::Init (
    uint16_t u16StackSize_,
    PORT_PRIO_TYPE uXPriority_,
    ThreadEntryFunc pfEntryPoint_,
    void * pvArg_ ) [static]
```

Init Create and initialize a new thread, using memory from the auto-allocated heap region to supply both the thread object and its stack. The thread returned can then be started using the [Start\(\)](#) method directly. Note that the memory used to create this thread cannot be reclaimed, and so this API is only suitable for threads that exist for the duration of runtime.

Parameters

<i>u16StackSize_</i>	Size of the stack (in bytes)
<i>uXPriority_</i>	Priority of the thread (0 = idle, 7 = max)
<i>pfEntryPoint_</i>	This is the function that gets called when the thread is started
<i>pvArg_</i>	Pointer to the argument passed into the thread's entryptoint function.

Returns

Pointer to a newly-created thread.

Definition at line 101 of file [thread.cpp](#).

19.34.3.23 IsInitialized()

```
bool Mark3::Thread::IsInitialized (
    void ) [inline]
```

IsInitialized Used to check whether or not a thread has been initialized prior to use.

Returns

return true if the thread object has been initialized, false otherwise.

Definition at line 77 of file [thread.h](#).

19.34.3.24 operator new()

```
void* Mark3::Thread::operator new (
    size_t sz,
    void * pv ) [inline]
```

Definition at line 67 of file [thread.h](#).

19.34.3.25 SetCurrent()

```
void Mark3::Thread::SetCurrent (
    ThreadList * pclNewList_ ) [inline]
```

SetCurrent. Set the thread's current to the specified thread list.

Parameters

<i>pcNew↔ List_</i>	Pointer to the threadlist to apply thread ownership
-------------------------	---

Definition at line 206 of file [thread.h](#).

19.34.3.26 SetEventFlagMask()

```
void Mark3::Thread::SetEventFlagMask (
    uint16_t u16Mask_ ) [inline]
```

SetEventFlagMask Sets the active event flag bitfield mask.

Parameters

<i>u16↔ Mask_</i>	
-----------------------	--

Definition at line 323 of file [thread.h](#).

19.34.3.27 SetEventFlagMode()

```
void Mark3::Thread::SetEventFlagMode (
    EventFlagOperation eMode_ ) [inline]
```

SetEventFlagMode Sets the active event flag operation mode.

Parameters

<i>e↔ Mode↔ _</i>	Event flag operation mode, defines the logical operator to apply to the event flag.
---------------------------	---

Definition at line 330 of file [thread.h](#).

19.34.3.28 SetExpired()

```
void Mark3::Thread::SetExpired (
    bool bExpired_ )
```

SetExpired Set the status of the current blocking call on the thread.

Parameters

<i>b↔ Expired↔ _</i>	true - call expired, false - call did not expire
------------------------------	--

Definition at line 418 of file [thread.cpp](#).

19.34.3.29 SetExtendedContext()

```
void Mark3::Thread::SetExtendedContext (
    void * pvData_ ) [inline]
```

SetExtendedContext Assign the [Thread](#) object's extended-context data pointer. Used by code implementing a user-defined thread-local storage model.

Object assigned to the context pointer should persist for the duration of the [Thread](#).

Parameters

<i>pv↔ Data_</i>	Object to assign to the extended data pointer.+
----------------------	---

Definition at line 380 of file [thread.h](#).

19.34.3.30 SetID()

```
void Mark3::Thread::SetID (
    uint8_t u8ID_ ) [inline]
```

SetID Set an arbitrary 8-bit ID to uniquely identify this thread.

Parameters

<i>u8l↔ D_</i>	8-bit Thread ID, set by the user
--------------------	--

Definition at line 282 of file [thread.h](#).

19.34.3.31 SetName()

```
void Mark3::Thread::SetName (
    const char * szName_ ) [inline]
```

SetName Set the name of the thread - this is purely optional, but can be useful when identifying issues that come along when multiple threads are at play in a system.

Parameters

<i>sz↔ Name_</i>	Char string containing the thread name
----------------------	--

Definition at line 143 of file [thread.h](#).

19.34.3.32 SetOwner()

```
void Mark3::Thread::SetOwner (
    ThreadList * pclNewList_ ) [inline]
```

SetOwner. Set the thread's owner to the specified thread list.

Parameters

<i>pc↔ NewList_</i>	Pointer to the threadlist to apply thread ownership
-------------------------	---

Definition at line 213 of file [thread.h](#).

19.34.3.33 SetPriority()

```
void Mark3::Thread::SetPriority (
    PORT_PRIO_TYPE uXPriority_ )
```

SetPriority. Set the priority of the [Thread](#) (running or otherwise) to a different level. This activity involves re-scheduling, and must be done so with due caution, as it may effect the determinism of the system.

This should *a*lways be called from within a critical section to prevent system issues.

Parameters

<i>uX↔ Priority_</i>	New priority of the thread
--------------------------	----------------------------

Definition at line 342 of file [thread.cpp](#).

19.34.3.34 SetPriorityBase()

```
void Mark3::Thread::SetPriorityBase (
    PORT_PRIO_TYPE uXPriority_ ) [private]
```

SetPriorityBase.

Parameters

<i>uX_↔</i> <i>Priority_</i>	
---	--

Definition at line 332 of file [thread.cpp](#).

19.34.3.35 SetQuantum()

```
void Mark3::Thread::SetQuantum (
    uint16_t u16Quantum_ ) [inline]
```

SetQuantum Set the thread's round-robin execution quantum.

Parameters

<i>u16_↔</i> <i>Quantum_</i>	Thread's execution quantum (in milliseconds)
---	--

Definition at line 190 of file [thread.h](#).

19.34.3.36 SetState()

```
void Mark3::Thread::SetState (
    ThreadState eState_ ) [inline]
```

SetState Set the thread's state to a new value. This is only to be used by code within the kernel, and is not intended for use by an end-user.

Parameters

<i>e_↔</i> <i>State_↔</i> _	New thread state to set.
---	--------------------------

Definition at line 397 of file [thread.h](#).

19.34.3.37 Sleep()

```
void Mark3::Thread::Sleep (
    uint32_t u32TimeMs_ ) [static]
```

Sleep Put the thread to sleep for the specified time (in milliseconds). Actual time slept may be longer (but not less than) the interval specified.

Parameters

<code>u32Time↔ Ms_</code>	Time to sleep (in ms)
-------------------------------	-----------------------

Examples:

[lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_coroutines/main.↔
cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line [249](#) of file [thread.cpp](#).

19.34.3.38 Start()

```
void Mark3::Thread::Start (
    void )
```

Start Start the thread - remove it from the stopped list, add it to the scheduler's list of threads (at the thread's set priority), and continue along.

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line [110](#) of file [thread.cpp](#).

19.34.3.39 Stop()

```
void Mark3::Thread::Stop (
    void )
```

Stop Stop a thread that's actively scheduled without destroying its stacks. Stopped threads can be restarted using the [Start\(\)](#) API.

Definition at line [141](#) of file [thread.cpp](#).

19.34.3.40 Yield()

```
void Mark3::Thread::Yield (
    void ) [static]
```

Yield Yield the thread - this forces the system to call the scheduler and determine what thread should run next. This is typically used when threads are moved in and out of the scheduler.

Definition at line [305](#) of file [thread.cpp](#).

19.34.4 Friends And Related Function Documentation

19.34.4.1 ThreadPort

```
friend class ThreadPort [friend]
```

Definition at line 411 of file [thread.h](#).

19.34.5 Member Data Documentation

19.34.5.1 m_bExpired

```
bool Mark3::Thread::m_bExpired [private]
```

Indicate whether or not a blocking-object timeout has occurred.

Definition at line 489 of file [thread.h](#).

19.34.5.2 m_clTimer

```
Timer Mark3::Thread::m_clTimer [private]
```

[Timer](#) used for blocking-object timeouts.

Definition at line 486 of file [thread.h](#).

19.34.5.3 m_eFlagMode

```
EventFlagOperation Mark3::Thread::m_eFlagMode [private]
```

Event-flag mode.

Definition at line 482 of file [thread.h](#).

19.34.5.4 m_eState

```
ThreadState Mark3::Thread::m_eState [private]
```

Enum indicating the thread's current state.

Definition at line 445 of file [thread.h](#).

19.34.5.5 m_pclCurrent

```
ThreadList* Mark3::Thread::m_pclCurrent [private]
```

Pointer to the thread-list where the thread currently resides.

Definition at line 461 of file [thread.h](#).

19.34.5.6 m_pclOwner

```
ThreadList* Mark3::Thread::m_pclOwner [private]
```

Pointer to the thread-list where the thread resides when active.

Definition at line 464 of file [thread.h](#).

19.34.5.7 m_pfEntryPoint

```
ThreadEntryFunc Mark3::Thread::m_pfEntryPoint [private]
```

The entry-point function called when the thread starts.

Definition at line 467 of file [thread.h](#).

19.34.5.8 m_pvArg

```
void* Mark3::Thread::m_pvArg [private]
```

Pointer to the argument passed into the thread's entypoint.

Definition at line 470 of file [thread.h](#).

19.34.5.9 m_pvExtendedContext

```
void* Mark3::Thread::m_pvExtendedContext [private]
```

Pointer provided to a [Thread](#) to implement thread-local storage.

Definition at line 449 of file [thread.h](#).

19.34.5.10 m_pwStack

```
K_WORD* Mark3::Thread::m_pwStack [private]
```

Pointer to the thread's stack.

Definition at line 433 of file [thread.h](#).

19.34.5.11 m_pwStackTop

```
K_WORD* Mark3::Thread::m_pwStackTop [private]
```

Pointer to the top of the thread's stack.

Definition at line 430 of file [thread.h](#).

19.34.5.12 m_szName

```
const char* Mark3::Thread::m_szName [private]
```

[Thread](#) name.

Definition at line 454 of file [thread.h](#).

19.34.5.13 m_u16FlagMask

```
uint16_t Mark3::Thread::m_u16FlagMask [private]
```

Event-flag mask.

Definition at line 479 of file [thread.h](#).

19.34.5.14 m_u16Quantum

```
uint16_t Mark3::Thread::m_u16Quantum [private]
```

[Thread](#) quantum (in milliseconds)

Definition at line 474 of file [thread.h](#).

19.34.5.15 m_u16StackSize

```
uint16_t Mark3::Thread::m_u16StackSize [private]
```

Size of the stack (in bytes)

Definition at line 458 of file [thread.h](#).

19.34.5.16 m_u8ThreadID

```
uint8_t Mark3::Thread::m_u8ThreadID [private]
```

[Thread](#) ID.

Definition at line 436 of file [thread.h](#).

19.34.5.17 m_uXCurPriority

```
PORT_PRIO_TYPE Mark3::Thread::m_uXCurPriority [private]
```

Current priority of the thread (priority inheritance)

Definition at line 442 of file [thread.h](#).

19.34.5.18 m_uXPriority

```
PORT_PRIO_TYPE Mark3::Thread::m_uXPriority [private]
```

Default priority of the thread.

Definition at line 439 of file [thread.h](#).

The documentation for this class was generated from the following files:

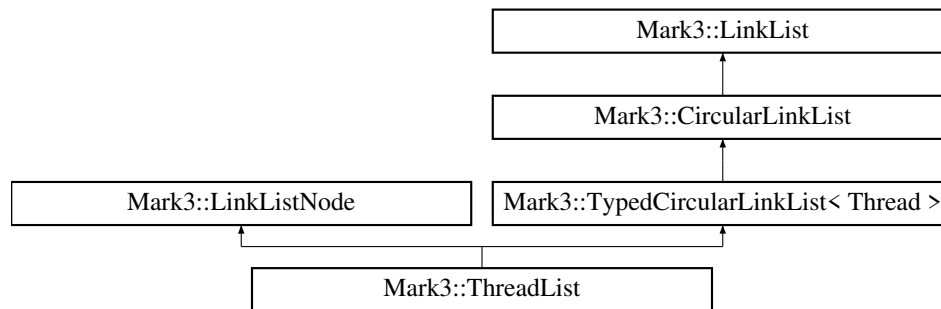
- [/home/moslevin/projects/m3-repo/kernel/src/public/thread.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/thread.cpp](#)

19.35 Mark3::ThreadList Class Reference

The [ThreadList](#) Class. This class is used for building thread-management facilities, such as schedulers, and blocking objects.

```
#include <threadlist.h>
```

Inheritance diagram for Mark3::ThreadList:



Public Member Functions

- void * [operator new](#) (size_t sz, void *pv)
- [ThreadList](#) ()
ThreadList Default constructor - zero-initializes the data.
- void [SetPriority](#) (PORT_PRIO_TYPE uXPriority_)
SetPriority Set the priority of this threadlist (if used for a scheduler).
- void [SetMapPointer](#) (PriorityMap *pclMap_)
SetMapPointer Set the pointer to a bitmap to use for this threadlist. Once again, only needed when the threadlist is being used for scheduling purposes.
- void [Add](#) (Thread *node_)
Add Add a thread to the threadlist.
- void [Add](#) (Thread *node_, PriorityMap *pclMap_, PORT_PRIO_TYPE uXPriority_)
Add Add a thread to the threadlist, specifying the flag and priority at the same time.
- void [AddPriority](#) (Thread *node_)
AddPriority Add a thread to the list such that threads are ordered from highest to lowest priority from the head of the list.
- void [Remove](#) (Thread *node_)
Remove Remove the specified thread from the threadlist.
- Thread * [HighestWaiter](#) ()
HighestWaiter Return a pointer to the highest-priority thread in the thread-list.

Private Attributes

- PORT_PRIO_TYPE m_uXPriority
Priority of the threadlist.
- PriorityMap * m_pclMap
Pointer to the bitmap/flag to set when used for scheduling.

Additional Inherited Members

19.35.1 Detailed Description

The [ThreadList](#) Class. This class is used for building thread-management facilities, such as schedulers, and blocking objects.

Definition at line 38 of file [threadlist.h](#).

19.35.2 Constructor & Destructor Documentation

19.35.2.1 ThreadList()

```
Mark3::ThreadList::ThreadList ( )
```

[ThreadList](#) Default constructor - zero-initializes the data.

Definition at line 27 of file [threadlist.cpp](#).

19.35.3 Member Function Documentation

19.35.3.1 Add() [1/2]

```
void Mark3::ThreadList::Add (
    Thread * node_ )
```

Add Add a thread to the threadlist.

Parameters

<i>node_↔</i> —	Pointer to the thread (link list node) to add to the list
--------------------	---

Definition at line 47 of file [threadlist.cpp](#).

19.35.3.2 Add() [2/2]

```
void Mark3::ThreadList::Add (
    Thread * node_,
```

```
PriorityMap * pclMap_,  
PORT_PRIO_TYPE uXPriority_ )
```

Add Add a thread to the threadlist, specifying the flag and priority at the same time.

Parameters

<i>node_</i>	Pointer to the thread to add (link list node)
<i>pclMap_</i>	Pointer to the bitmap flag to set (if used in a scheduler context), or nullptr for non-scheduler.
<i>uX_</i> <i>Priority_</i>	Priority of the threadlist

Definition at line 101 of file [threadlist.cpp](#).

19.35.3.3 AddPriority()

```
void Mark3::ThreadList::AddPriority (
    Thread * node_ )
```

AddPriority Add a thread to the list such that threads are ordered from highest to lowest priority from the head of the list.

Parameters

<i>node_</i> _	Pointer to a thread to add to the list.
-------------------	---

Definition at line 66 of file [threadlist.cpp](#).

19.35.3.4 HighestWaiter()

```
Thread * Mark3::ThreadList::HighestWaiter ( )
```

HighestWaiter Return a pointer to the highest-priority thread in the thread-list.

Returns

Pointer to the highest-priority thread

Definition at line 128 of file [threadlist.cpp](#).

19.35.3.5 operator new()

```
void* Mark3::ThreadList::operator new (
    size_t sz,
    void * pv ) [inline]
```

Definition at line 41 of file [threadlist.h](#).

19.35.3.6 Remove()

```
void Mark3::ThreadList::Remove (
    Thread * node_ )
```

Remove Remove the specified thread from the threadlist.

Parameters

<i>node</i> ↔ _	Pointer to the thread to remove
--------------------	---------------------------------

Definition at line 111 of file [threadlist.cpp](#).

19.35.3.7 SetMapPointer()

```
void Mark3::ThreadList::SetMapPointer (
    PriorityMap * pclMap_ )
```

SetMapPointer Set the pointer to a bitmap to use for this threadlist. Once again, only needed when the threadlist is being used for scheduling purposes.

Parameters

<i>pcl</i> ↔ <i>Map_</i>	Pointer to the priority map object used to track this thread.
-----------------------------	---

Definition at line 40 of file [threadlist.cpp](#).

19.35.3.8 SetPriority()

```
void Mark3::ThreadList::SetPriority (
    PORT_PRIO_TYPE uXPriority_ )
```

SetPriority Set the priority of this threadlist (if used for a scheduler).

Parameters

<i>uX</i> ↔ <i>Priority_</i>	Priority level of the thread list
---------------------------------	-----------------------------------

Definition at line 34 of file [threadlist.cpp](#).

19.35.4 Member Data Documentation

19.35.4.1 m_pclMap

```
PriorityMap* Mark3::ThreadList::m_pclMap [private]
```

Pointer to the bitmap/flag to set when used for scheduling.

Definition at line 116 of file [threadlist.h](#).

19.35.4.2 m_uXPriority

```
PORT_PRIO_TYPE Mark3::ThreadList::m_uXPriority [private]
```

Priority of the threadlist.

Definition at line 113 of file [threadlist.h](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/m3-repo/kernel/src/public/threadlist.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/threadlist.cpp](#)

19.36 Mark3::ThreadListList Class Reference

The [ThreadListList](#) class Class used to track all threadlists active in the OS kernel. At any point in time, the list can be traversed to get a complete view of all running, blocked, or stopped threads in the system.

```
#include <threadlistlist.h>
```

Static Public Member Functions

- static void [Add](#) ([ThreadList](#) *pclThreadList_)
Add Add a [ThreadList](#) to the list for tracking.
- static void [Remove](#) ([ThreadList](#) *pclThreadList_)
Remove Remove a threadlist from tracking.
- static [ThreadList](#) * [GetHead](#) ()
GetHead.

Static Private Attributes

- static [TypedDoubleLinkList](#)< [ThreadList](#) > [m_clThreadListList](#)

19.36.1 Detailed Description

The [ThreadListList](#) class Class used to track all threadlists active in the OS kernel. At any point in time, the list can be traversed to get a complete view of all running, blocked, or stopped threads in the system.

Definition at line 36 of file [threadlistlist.h](#).

19.36.2 Member Function Documentation

19.36.2.1 Add()

```
static void Mark3::ThreadListList::Add (  
    ThreadList * pclThreadList_ ) [inline], [static]
```

Add Add a [ThreadList](#) to the list for tracking.

Parameters

<i>pclThreadList_</i>	threadlist to add for tracking
-----------------------	--------------------------------

Definition at line 43 of file [threadlistlist.h](#).

19.36.2.2 GetHead()

```
static ThreadList* Mark3::ThreadListList::GetHead ( ) [inline], [static]
```

GetHead.

Returns

The threadlist at the beginning of the list

Definition at line 62 of file [threadlistlist.h](#).

19.36.2.3 Remove()

```
static void Mark3::ThreadListList::Remove (
    ThreadList * pclThreadList_ ) [inline], [static]
```

Remove Remove a threadlist from tracking.

Parameters

<i>pclThreadList_</i>	threadlist to remove from tracking
-----------------------	------------------------------------

Definition at line 53 of file [threadlistlist.h](#).

19.36.3 Member Data Documentation

19.36.3.1 m_clThreadListList

```
TypedDoubleLinkedList< ThreadList > Mark3::ThreadListList::m_clThreadListList [static], [private]
```

Definition at line 68 of file [threadlistlist.h](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/m3-repo/kernel/src/public/threadlistlist.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/threadlistlist.cpp](#)

19.37 Mark3::ThreadPort Class Reference

The [ThreadPort](#) Class defines the target-specific functions required by the kernel for threading.

```
#include <ithreadport.h>
```

Static Public Member Functions

- static void [Init](#) ()
Init Function to perform early init of the target environment prior to using OS primitives.
- static void [StartThreads](#) ()
StartThreads Function to start the scheduler, initial threads, etc.

Static Private Member Functions

- static void [InitStack](#) ([Thread](#) *pstThread_)
InitStack Initialize the thread's stack.

Friends

- class [Thread](#)

19.37.1 Detailed Description

The [ThreadPort](#) Class defines the target-specific functions required by the kernel for threading.

This is limited (at this point) to a function to start the scheduler, and a function to initialize the default stack-frame for a thread.

Definition at line [35](#) of file [ithreadport.h](#).

19.37.2 Member Function Documentation

19.37.2.1 Init()

```
static void Mark3::ThreadPort::Init (
    void ) [inline], [static]
```

Init Function to perform early init of the target environment prior to using OS primitives.

Definition at line [43](#) of file [ithreadport.h](#).

19.37.2.2 InitStack()

```
void Mark3::ThreadPort::InitStack (
    Thread * pstThread_ ) [static], [private]
```

InitStack Initialize the thread's stack.

Parameters

<i>pst</i> ↔ <i>Thread_</i>	Pointer to the thread to initialize
--------------------------------	-------------------------------------

Definition at line 43 of file [threadport.cpp](#).

19.37.2.3 StartThreads()

```
void Mark3::ThreadPort::StartThreads ( ) [static]
```

StartThreads Function to start the scheduler, initial threads, etc.

Definition at line 100 of file [threadport.cpp](#).

19.37.3 Friends And Related Function Documentation

19.37.3.1 Thread

```
friend class Thread [friend]
```

Definition at line 50 of file [ithreadport.h](#).

The documentation for this class was generated from the following files:


- [/home/moslevin/projects/m3-repo/kernel/src/public/ithreadport.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/threadport.cpp](#)

19.38 Mark3::Timer Class Reference

The [Timer](#) Class. This class provides kernel-managed timers, used to provide high-precision delays. Functionality is useful to both user-code, and is used extensively within the kernel and its blocking objects to implement round-robin scheduling, thread sleep, and timeouts. Provides one-shot and periodic timers for use by application code. This object relies on a target-defined hardware timer implementation, which is multiplexed by the kernel's timer scheduler.

```
#include <timer.h>
```

Inheritance diagram for Mark3::Timer:



Public Member Functions

- void * [operator new](#) (size_t sz, void *pv)
- [~Timer](#) ()
- [Timer](#) ()

[Timer](#) Default Constructor - Do nothing. Allow the init call to perform the necessary object initialization prior to use.
- void [Init](#) ()

Init Re-initialize the [Timer](#) to default values.
- void [Start](#) (bool bRepeat_, uint32_t u32IntervalMs_, [TimerCallback](#) pfCallback_, void *pvData_)

Start Start a timer using default ownership, using repeats as an option, and millisecond resolution.
- void [Start](#) ()

Start Start or restart a timer using parameters previously configured via calls to Start(<with args>), or via the a-la-carte parameter setter methods. This is especially useful for retriggering one-shot timers that have previously expired, using the timer's previous configuration.
- void [Stop](#) ()

Stop Stop a timer already in progress. Has no effect on timers that have already been stopped.

Private Member Functions

- void [SetInitialized](#) ()

SetInitialized.
- bool [IsInitialized](#) (void)

IsInitialized.

Static Private Member Functions

- static uint32_t [SecondsToTicks](#) (uint32_t x)
- static uint32_t [MSecondsToTicks](#) (uint32_t x)
- static uint32_t [USecondsToTicks](#) (uint32_t x)

Private Attributes

- uint8_t [m_u8Initialized](#)

Cookie used to determine whether or not the timer is initialized.
- uint8_t [m_u8Flags](#)

Flags for the timer, defining if the timer is one-shot or repeated.
- [TimerCallback](#) [m_pfCallback](#)

Pointer to the callback function.
- uint32_t [m_u32Interval](#)

Interval of the timer in timer ticks.
- uint32_t [m_u32TimeLeft](#)

Time remaining on the timer.
- [Thread](#) * [m_pclOwner](#)

Pointer to the owner thread.
- void * [m_pvData](#)

Pointer to the callback data.

Static Private Attributes

- static constexpr auto [m_uTimerInvalidCookie](#) = uint8_t { 0x3C }
- static constexpr auto [m_uTimerInitCookie](#) = uint8_t { 0xC3 }

Friends

- class [TimerList](#)

Additional Inherited Members

19.38.1 Detailed Description

The [Timer](#) Class. This class provides kernel-managed timers, used to provide high-precision delays. Functionality is useful to both user-code, and is used extensively within the kernel and its blocking objects to implement round-robin scheduling, thread sleep, and timeouts. Provides one-shot and periodic timers for use by application code. This object relies on a target-defined hardware timer implementation, which is multiplexed by the kernel's timer scheduler.

Examples:

[lab6_timers/main.cpp](#).

Definition at line 68 of file [timer.h](#).

19.38.2 Constructor & Destructor Documentation

19.38.2.1 ~Timer()

```
Mark3::Timer::~~Timer ( ) [inline]
```

Definition at line 72 of file [timer.h](#).

19.38.2.2 Timer()

```
Mark3::Timer::Timer ( )
```

[Timer](#) Default Constructor - Do nothing. Allow the init call to perform the necessary object initialization prior to use.

Definition at line 29 of file [timer.cpp](#).

19.38.3 Member Function Documentation

19.38.3.1 Init()

```
void Mark3::Timer::Init (  
    void )
```

Init Re-initialize the [Timer](#) to default values.

Definition at line [36](#) of file [timer.cpp](#).

19.38.3.2 IsInitialized()

```
bool Mark3::Timer::IsInitialized (  
    void ) [inline], [private]
```

IsInitialized.

Returns

true if initialized, false if not initialized.

Definition at line [127](#) of file [timer.h](#).

19.38.3.3 MSecondsToTicks()

```
static uint32_t Mark3::Timer::MSecondsToTicks (  
    uint32_t x ) [inline], [static], [private]
```

Definition at line [130](#) of file [timer.h](#).

19.38.3.4 operator new()

```
void* Mark3::Timer::operator new (  
    size_t sz,  
    void * pv ) [inline]
```

Definition at line [71](#) of file [timer.h](#).

19.38.3.5 SecondsToTicks()

```
static uint32_t Mark3::Timer::SecondsToTicks (
    uint32_t x ) [inline], [static], [private]
```

Definition at line 129 of file [timer.h](#).

19.38.3.6 SetInitialized()

```
void Mark3::Timer::SetInitialized (
    void ) [inline], [private]
```

SetInitialized.

Definition at line 121 of file [timer.h](#).

19.38.3.7 Start() [1/2]

```
void Mark3::Timer::Start (
    bool bRepeat_,
    uint32_t u32IntervalMs_,
    TimerCallback pfCallback_,
    void * pvData_ )
```

Start Start a timer using default ownership, using repeats as an option, and millisecond resolution.

Parameters

<i>bRepeat_</i>	0 - timer is one-shot. 1 - timer is repeating.
<i>u32IntervalMs_</i>	- Interval of the timer in milliseconds
<i>pfCallback_</i>	- Function to call on timer expiry
<i>pvData_</i>	- Data to pass into the callback function

Examples:

[lab6_timers/main.cpp](#).

Definition at line 51 of file [timer.cpp](#).

19.38.3.8 Start() [2/2]

```
void Mark3::Timer::Start (
    void )
```

Start or restart a timer using parameters previously configured via calls to `Start(<with args>)`, or via the a-la-carte parameter setter methods. This is especially useful for retriggering one-shot timers that have previously expired, using the timer's previous configuration.

Definition at line 73 of file [timer.cpp](#).

19.38.3.9 Stop()

```
void Mark3::Timer::Stop (  
    void )
```

Stop a timer already in progress. Has no effect on timers that have already been stopped.

Definition at line 86 of file [timer.cpp](#).

19.38.3.10 USecondsToTicks()

```
static uint32_t Mark3::Timer::USecondsToTicks (  
    uint32_t x ) [inline], [static], [private]
```

Definition at line 131 of file [timer.h](#).

19.38.4 Friends And Related Function Documentation

19.38.4.1 TimerList

```
friend class TimerList [friend]
```

Definition at line 116 of file [timer.h](#).

19.38.5 Member Data Documentation

19.38.5.1 m_pclOwner

```
Thread* Mark3::Timer::m_pclOwner [private]
```

Pointer to the owner thread.

Definition at line 152 of file [timer.h](#).

19.38.5.2 m_pfCallback

```
TimerCallback Mark3::Timer::m_pfCallback [private]
```

Pointer to the callback function.

Definition at line 143 of file [timer.h](#).

19.38.5.3 m_pvData

```
void* Mark3::Timer::m_pvData [private]
```

Pointer to the callback data.

Definition at line 155 of file [timer.h](#).

19.38.5.4 m_u32Interval

```
uint32_t Mark3::Timer::m_u32Interval [private]
```

Interval of the timer in timer ticks.

Definition at line 146 of file [timer.h](#).

19.38.5.5 m_u32TimeLeft

```
uint32_t Mark3::Timer::m_u32TimeLeft [private]
```

Time remaining on the timer.

Definition at line 149 of file [timer.h](#).

19.38.5.6 m_u8Flags

```
uint8_t Mark3::Timer::m_u8Flags [private]
```

Flags for the timer, defining if the timer is one-shot or repeated.

Definition at line 140 of file [timer.h](#).

19.38.5.7 m_u8Initialized

```
uint8_t Mark3::Timer::m_u8Initialized [private]
```

Cookie used to determine whether or not the timer is initialized.

Definition at line 137 of file [timer.h](#).

19.38.5.8 m_uTimerInitCookie

```
constexpr auto Mark3::Timer::m_uTimerInitCookie = uint8_t { 0xC3 } [static], [private]
```

Definition at line 134 of file [timer.h](#).

19.38.5.9 m_uTimerInvalidCookie

```
constexpr auto Mark3::Timer::m_uTimerInvalidCookie = uint8_t { 0x3C } [static], [private]
```

Definition at line 133 of file [timer.h](#).

The documentation for this class was generated from the following files:


- [/home/moslevin/projects/m3-repo/kernel/src/public/timer.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/timer.cpp](#)

19.39 Mark3::TimerList Class Reference

the [TimerList](#) class. This class implements a doubly-linked-list of timer objects.

```
#include <timerlist.h>
```

Inheritance diagram for Mark3::TimerList:



Public Member Functions

- void [Init](#) ()
Init Initialize the [TimerList](#) object. Must be called before using the object.
- void [Add](#) ([Timer](#) *pclListNode_)
Add Add a timer to the [TimerList](#).
- void [Remove](#) ([Timer](#) *pclLinkListNode_)
Remove Remove a timer from the [TimerList](#), cancelling its expiry.
- void [Process](#) ()
Process Process all timers in the timerlist as a result of the timer expiring. This will select a new timer epoch based on the next timer to expire.

Private Attributes

- uint32_t [m_u32NextWakeup](#)
The time (in system clock ticks) of the next wakeup event.
- bool [m_bTimerActive](#)
Whether or not the timer is active.
- [Mutex](#) [m_clMutex](#)
Guards against concurrent access to the timer list - Only needed when running threaded.

Additional Inherited Members

19.39.1 Detailed Description

the [TimerList](#) class. This class implements a doubly-linked-list of timer objects.

Definition at line 40 of file [timerlist.h](#).

19.39.2 Member Function Documentation

19.39.2.1 Add()

```
void Mark3::TimerList::Add (
    Timer * pclListNode_ )
```

Add Add a timer to the [TimerList](#).

Parameters

pclList ↔ Node_	Pointer to the Timer to Add
--	---

Definition at line 35 of file [timerlist.cpp](#).

19.39.2.2 Init()

```
void Mark3::TimerList::Init (
    void )
```

Init Initialize the [TimerList](#) object. Must be called before using the object.

Definition at line 27 of file [timerlist.cpp](#).

19.39.2.3 Process()

```
void Mark3::TimerList::Process (
    void )
```

Process Process all timers in the timerlist as a result of the timer expiring. This will select a new timer epoch based on the next timer to expire.

Definition at line 61 of file [timerlist.cpp](#).

19.39.2.4 Remove()

```
void Mark3::TimerList::Remove (
    Timer * pclLinkedListNode_ )
```

Remove Remove a timer from the [TimerList](#), cancelling its expiry.

Parameters

<i>pclLinkedListNode_</i>	Pointer to the Timer to remove
---	--

Definition at line 51 of file [timerlist.cpp](#).

19.39.3 Member Data Documentation

19.39.3.1 m_bTimerActive

```
bool Mark3::TimerList::m_bTimerActive [private]
```

Whether or not the timer is active.

Definition at line 78 of file [timerlist.h](#).

19.39.3.2 m_clMutex

```
Mutex Mark3::TimerList::m_clMutex [private]
```

Guards against concurrent access to the timer list - Only needed when running threaded.

Definition at line 81 of file [timerlist.h](#).

19.39.3.3 m_u32NextWakeup

```
uint32_t Mark3::TimerList::m_u32NextWakeup [private]
```

The time (in system clock ticks) of the next wakeup event.

Definition at line 75 of file [timerlist.h](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/m3-repo/kernel/src/public/timerlist.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/timerlist.cpp](#)

19.40 Mark3::TimerScheduler Class Reference

The [TimerScheduler](#) Class. This implements a "Static" class used to manage a global list of timers used throughout the system.

```
#include <timerscheduler.h>
```

Static Public Member Functions

- static void [Init](#) ()
Init Initialize the timer scheduler. Must be called before any timer, or timer-derived functions are used.
- static void [Add](#) ([Timer](#) *pclListNode_)
Add Add a timer to the timer scheduler. Adding a timer implicitly starts the timer as well.
- static void [Remove](#) ([Timer](#) *pclListNode_)
Remove Remove a timer from the timer scheduler. May implicitly stop the timer if this is the only active timer scheduled.
- static void [Process](#) ()
Process This function must be called on timer expiry (from the timer's ISR context). This will result in all timers being updated based on the epoch that just elapsed. The next timer epoch is set based on the next [Timer](#) object to expire.

Static Private Attributes

- static [TimerList](#) [m_clTimerList](#)
TimerList object manipulated by the [Timer Scheduler](#).

19.40.1 Detailed Description

The [TimerScheduler](#) Class. This implements a "Static" class used to manage a global list of timers used throughout the system.

Definition at line 38 of file [timerscheduler.h](#).

19.40.2 Member Function Documentation

19.40.2.1 Add()

```
static void Mark3::TimerScheduler::Add (
    Timer * pclListNode_ ) [inline], [static]
```

Add Add a timer to the timer scheduler. Adding a timer implicitly starts the timer as well.

Parameters

<i>pcListNode_</i>	Pointer to the timer list node to add
--------------------	---------------------------------------

Definition at line 54 of file [timerscheduler.h](#).

19.40.2.2 Init()

```
static void Mark3::TimerScheduler::Init (
    void ) [inline], [static]
```

Init Initialize the timer scheduler. Must be called before any timer, or timer-derived functions are used.

Definition at line 46 of file [timerscheduler.h](#).

19.40.2.3 Process()

```
static void Mark3::TimerScheduler::Process (
    void ) [inline], [static]
```

Process This function must be called on timer expiry (from the timer's ISR context). This will result in all timers being updated based on the epoch that just elapsed. The next timer epoch is set based on the next [Timer](#) object to expire.

Definition at line 70 of file [timerscheduler.h](#).

19.40.2.4 Remove()

```
static void Mark3::TimerScheduler::Remove (
    Timer * pclListNode_ ) [inline], [static]
```

Remove Remove a timer from the timer scheduler. May implicitly stop the timer if this is the only active timer scheduled.

Parameters

<i>pclList↵ Node_</i>	Pointer to the timer list node to remove
---------------------------	--

Definition at line 62 of file [timerscheduler.h](#).

19.40.3 Member Data Documentation

19.40.3.1 m_clTimerList

```
TimerList Mark3::TimerScheduler::m_clTimerList [static], [private]
```

[TimerList](#) object manipulated by the [Timer Scheduler](#).

Definition at line 74 of file [timerscheduler.h](#).

The documentation for this class was generated from the following files:

- [/home/moslevin/projects/m3-repo/kernel/src/public/timerscheduler.h](#)
- [/home/moslevin/projects/m3-repo/kernel/src/timer.cpp](#)

19.41 Mark3::Token_t Struct Reference

Token descriptor struct format.

```
#include <memutil.h>
```

Public Attributes

- const char * [pcToken](#)
Pointer to the beginning of the token string.
- uint8_t [u8Len](#)
Length of the token (in bytes)

19.41.1 Detailed Description

Token descriptor struct format.

Definition at line 32 of file [memutil.h](#).

19.41.2 Member Data Documentation

19.41.2.1 pcToken

```
const char* Mark3::Token_t::pcToken
```

Pointer to the beginning of the token string.

Definition at line 33 of file [memutil.h](#).

19.41.2.2 u8Len

```
uint8_t Mark3::Token_t::u8Len
```

Length of the token (in bytes)

Definition at line 34 of file [memutil.h](#).

The documentation for this struct was generated from the following file:


- [/home/moslevin/projects/m3-repo/kernel/lib/memutil/public/memutil.h](#)

19.42 Mark3::TypedCircularLinkedList< T > Class Template Reference

The [TypedCircularLinkedList](#) Class Circular-linked-list data type, inherited from the base [LinkedList](#) type, and templated for use with linked-list-node derived data-types.

```
#include <ll.h>
```

Inheritance diagram for Mark3::TypedCircularLinkedList< T >:



Public Member Functions

- void * [operator new](#) (size_t sz, void *pv)
- [TypedCircularLinkedList](#) ()
- T * [GetHead](#) ()
GetHead Get the head node in the linked list.
- void [SetHead](#) (T *pclNode_)
SetHead Set the head node of a linked list.
- T * [GetTail](#) ()
GetTail Get the tail node of the linked list.
- void [SetTail](#) (T *pclNode_)
SetTail Set the tail node of the linked list.
- void [Add](#) (T *pNode_)
Add Add the linked list node to this linked list.
- void [Remove](#) (T *pNode_)
Remove Remove Add the linked list node to this linked list.
- void [InsertNodeBefore](#) (T *pNode_, T *pInsert_)
InsertNodeBefore Insert a linked-list node into the list before the specified insertion point.

Additional Inherited Members

19.42.1 Detailed Description

```
template<typename T>
class Mark3::TypedCircularLinkedList< T >
```

The [TypedCircularLinkedList](#) Class Circular-linked-list data type, inherited from the base [LinkedList](#) type, and templated for use with linked-list-node derived data-types.

Definition at line 347 of file [ll.h](#).

19.42.2 Constructor & Destructor Documentation

19.42.2.1 TypedCircularLinkedList()

```
template<typename T>
Mark3::TypedCircularLinkedList< T >::TypedCircularLinkedList ( ) [inline]
```

Definition at line 352 of file [ll.h](#).

19.42.3 Member Function Documentation

19.42.3.1 Add()

```
template<typename T>
void Mark3::TypedCircularLinkedList< T >::Add (
    T * pNode_ ) [inline]
```

Add the linked list node to this linked list.

Parameters

<i>node</i> ↔	Pointer to the node to add
—	

Definition at line 394 of file [ll.h](#).

19.42.3.2 GetHead()

```
template<typename T>
T* Mark3::TypedCircularLinkedList< T >::GetHead ( ) [inline]
```

GetHead Get the head node in the linked list.

Returns

Pointer to the head node in the list

Definition at line 362 of file [ll.h](#).

19.42.3.3 GetTail()

```
template<typename T>
T* Mark3::TypedCircularLinkedList< T >::GetTail ( ) [inline]
```

GetTail Get the tail node of the linked list.

Returns

Pointer to the tail node in the list

Definition at line 378 of file [ll.h](#).

19.42.3.4 InsertNodeBefore()

```
template<typename T>
void Mark3::TypedCircularLinkedList< T >::InsertNodeBefore (
    T * pNode_,
    T * pInsert_ ) [inline]
```

InsertNodeBefore Insert a linked-list node into the list before the specified insertion point.

Parameters

<i>node</i> ↔ —	Node to insert into the list
<i>insert</i> ↔ —	Insert point.

Definition at line 418 of file [ll.h](#).

19.42.3.5 operator new()

```
template<typename T>
void* Mark3::TypedCircularLinkedList< T >::operator new (
    size_t sz,
    void * pv ) [inline]
```

Definition at line 350 of file [ll.h](#).

19.42.3.6 Remove()

```
template<typename T>
void Mark3::TypedCircularLinkedList< T >::Remove (
    T * pNode_ ) [inline]
```

Remove Add the linked list node to this linked list.

Parameters

<i>node</i> ↔ —	Pointer to the node to remove
--------------------	-------------------------------

Definition at line 405 of file [ll.h](#).

19.42.3.7 SetHead()

```
template<typename T>
void Mark3::TypedCircularLinkedList< T >::SetHead (
    T * pNode_ ) [inline]
```

SetHead Set the head node of a linked list.

Parameters

<i>pcNode_</i>	Pointer to node to set as the head of the linked list
----------------	---

Definition at line 370 of file [ll.h](#).

19.42.3.8 SetTail()

```
template<typename T>
void Mark3::TypedCircularLinkedList< T >::SetTail (
    T * pcNode_ ) [inline]
```

SetTail Set the tail node of the linked list.

Parameters

<i>pcNode_</i>	Pointer to the node to set as the tail of the linked list
----------------	---

Definition at line 386 of file [ll.h](#).

The documentation for this class was generated from the following file:


- [/home/moslevin/projects/m3-repo/kernel/src/public/ll.h](#)

19.43 Mark3::TypedDoubleLinkedList< T > Class Template Reference

The [TypedDoubleLinkedList](#) Class Doubly-linked-list data type, inherited from the base [LinkedList](#) type, and templated for use with linked-list-node derived data-types.

```
#include <ll.h>
```

Inheritance diagram for Mark3::TypedDoubleLinkedList< T >:



Public Member Functions

- void * [operator new](#) (size_t sz, void *pv)
- [TypedDoubleLinkedList](#) ()
- T * [GetHead](#) ()
GetHead Get the head node in the linked list.
- void [SetHead](#) (T *pclNode_)
SetHead Set the head node of a linked list.
- T * [GetTail](#) ()
GetTail Get the tail node of the linked list.
- void [SetTail](#) (T *pclNode_)
SetTail Set the tail node of the linked list.
- void [Add](#) (T *pNode_)
Add Add the linked list node to this linked list.
- void [Remove](#) (T *pNode_)
Remove Add the linked list node to this linked list.

Additional Inherited Members

19.43.1 Detailed Description

```
template<typename T>
class Mark3::TypedDoubleLinkedList< T >
```

The [TypedDoubleLinkedList](#) Class Doubly-linked-list data type, inherited from the base [LinkedList](#) type, and templated for use with linked-list-node derived data-types.

Definition at line 276 of file [ll.h](#).

19.43.2 Constructor & Destructor Documentation

19.43.2.1 TypedDoubleLinkedList()

```
template<typename T>
Mark3::TypedDoubleLinkedList< T >::TypedDoubleLinkedList ( ) [inline]
```

Definition at line 281 of file [ll.h](#).

19.43.3 Member Function Documentation

19.43.3.1 Add()

```
template<typename T>
void Mark3::TypedDoubleLinkedList< T >::Add (
    T * pNode_ ) [inline]
```

Add Add the linked list node to this linked list.

Parameters

<i>node</i> ↔	Pointer to the node to add
—	

Definition at line 323 of file [ll.h](#).

19.43.3.2 GetHead()

```
template<typename T>
T* Mark3::TypedDoubleLinkedList< T >::GetHead ( ) [inline]
```

GetHead Get the head node in the linked list.

Returns

Pointer to the head node in the list

Definition at line 291 of file [ll.h](#).

19.43.3.3 GetTail()

```
template<typename T>
T* Mark3::TypedDoubleLinkedList< T >::GetTail ( ) [inline]
```

GetTail Get the tail node of the linked list.

Returns

Pointer to the tail node in the list

Definition at line 307 of file [ll.h](#).

19.43.3.4 operator new()

```
template<typename T>
void* Mark3::TypedDoubleLinkedList< T >::operator new (
    size_t sz,
    void * pv ) [inline]
```

Definition at line 279 of file [ll.h](#).

19.43.3.5 Remove()

```
template<typename T>
void Mark3::TypedDoubleLinkedList< T >::Remove (
    T * pNode_ ) [inline]
```

Remove Add the linked list node to this linked list.

Parameters

<i>node</i> ↔ —	Pointer to the node to remove
--------------------	-------------------------------

Definition at line 334 of file [ll.h](#).

19.43.3.6 SetHead()

```
template<typename T>
void Mark3::TypedDoubleLinkedList< T >::SetHead (
    T * pclNode_ ) [inline]
```

SetHead Set the head node of a linked list.

Parameters

<i>pcl</i> ↔ <i>Node_</i>	Pointer to node to set as the head of the linked list
------------------------------	---

Definition at line 299 of file [ll.h](#).

19.43.3.7 SetTail()

```
template<typename T>
void Mark3::TypedDoubleLinkedList< T >::SetTail (
    T * pclNode_ ) [inline]
```

SetTail Set the tail node of the linked list.

Parameters

<i>pcl</i> ↔ <i>Node_</i>	Pointer to the node to set as the tail of the linked list
------------------------------	---

Definition at line 315 of file [ll.h](#).

The documentation for this class was generated from the following file:


- [/home/moslevin/projects/m3-repo/kernel/src/public/ll.h](#)

19.44 Mark3::TypedLinkedListNode< T > Class Template Reference

The [TypedLinkedListNode](#) class The [TypedLinkedListNode](#) class provides a linked-list node type for a specified object type. This can be used with typed link-list data structures to manage lists of objects without having to static-cast between the base type and the derived class.

```
#include <ll.h>
```

Inheritance diagram for `Mark3::TypedLinkedListNode< T >`:



Public Member Functions

- `T * GetNext ()`
- `T * GetPrev ()`

Additional Inherited Members

19.44.1 Detailed Description

```
template<typename T>
class Mark3::TypedLinkedListNode< T >
```

The [TypedLinkedListNode](#) class The [TypedLinkedListNode](#) class provides a linked-list node type for a specified object type. This can be used with typed link-list data structures to manage lists of objects without having to static-cast between the base type and the derived class.

Definition at line [108](#) of file [ll.h](#).

19.44.2 Member Function Documentation

19.44.2.1 [GetNext\(\)](#)

```
template<typename T>
T* Mark3::TypedLinkedListNode< T >::GetNext (
    void ) [inline]
```

Definition at line [111](#) of file [ll.h](#).

19.44.2.2 [GetPrev\(\)](#)

```
template<typename T>
T* Mark3::TypedLinkedListNode< T >::GetPrev (
    void ) [inline]
```

Definition at line [112](#) of file [ll.h](#).

The documentation for this class was generated from the following file:

- [/home/moslevin/projects/m3-repo/kernel/src/public/ll.h](#)

Chapter 20

File Documentation

20.1 /home/moslevin/projects/m3-repo/kernel/lib/mark3c/src/mark3c.cpp File Reference

Implementation of C-language wrappers for the [Mark3](#) kernel.

```
#include "mark3c.h"
#include "mark3.h"
```

Functions

- void * [Alloc_Memory](#) (uint16_t u16Size_)
- void [Free_Memory](#) (void *pvObject_)
Free_Memory.
- [Semaphore_t](#) [Alloc_Semaphore](#) (void)
Alloc_Semaphore.
- void [Free_Semaphore](#) ([Semaphore_t](#) handle)
- [Mutex_t](#) [Alloc_Mutex](#) (void)
Alloc_Mutex.
- void [Free_Mutex](#) ([Mutex_t](#) handle)
- [EventFlag_t](#) [Alloc_EventFlag](#) (void)
- void [Free_EventFlag](#) ([EventFlag_t](#) handle)
- [Message_t](#) [Alloc_Message](#) (void)
Alloc_Message.
- void [Free_Message](#) ([Message_t](#) handle)
- [MessageQueue_t](#) [Alloc_MessageQueue](#) (void)
Alloc_MessageQueue.
- void [Free_MessageQueue](#) ([MessageQueue_t](#) handle)
- [MessagePool_t](#) [Alloc_MessagePool](#) (void)
- void [Free_MessagePool](#) ([MessagePool_t](#) handle)
- [Notify_t](#) [Alloc_Notify](#) (void)
Alloc_Notify.
- void [Free_Notify](#) ([Notify_t](#) handle)
- [Mailbox_t](#) [Alloc_Mailbox](#) (void)
Alloc_Mailbox.
- void [Free_Mailbox](#) ([Mailbox_t](#) handle)

- [ConditionVariable_t Alloc_ConditionVariable](#) (void)
- void [Free_ConditionVariable](#) ([ConditionVariable_t](#) handle)
- [ReaderWriterLock_t Alloc_ReaderWriterLock](#) (void)
- void [Free_ReaderWriterLock](#) ([ReaderWriterLock_t](#) handle)
- [Thread_t Alloc_Thread](#) (void)
 - Alloc_Thread.*
- void [Free_Thread](#) ([Thread_t](#) handle)
- [Timer_t Alloc_Timer](#) (void)
 - Alloc_Timer.*
- void [Free_Timer](#) ([Timer_t](#) handle)
- void [Kernel_Init](#) (void)
 - Kernel_Init.*
- void [Kernel_Start](#) (void)
 - Kernel_Start.*
- bool [Kernel_IsStarted](#) (void)
 - Kernel_IsStarted.*
- void [Kernel_SetPanic](#) ([PanicFunc](#) pfPanic_)
 - Kernel_SetPanic.*
- bool [Kernel_IsPanic](#) (void)
 - Kernel_IsPanic.*
- void [Kernel_Panic](#) (uint16_t u16Cause_)
 - Kernel_Panic.*
- uint32_t [Kernel_GetTicks](#) (void)
 - Kernel_GetTicks.*
- void [Kernel_SetStackGuardThreshold](#) (uint16_t u16Threshold_)
- uint16_t [Kernel_GetStackGuardThreshold](#) (void)
- void [Kernel_SetDebugPrintFunction](#) ([kernel_debug_print_t](#) pfPrintFunction_)
 - Kernel_SetDebugPrintFunction.*
- void [Kernel_DebugPrint](#) (const char *szString_)
 - KernelDebug_DebugPrint.*
- void [Scheduler_Enable](#) (bool bEnable_)
 - Scheduler_Enable.*
- bool [Scheduler_IsEnabled](#) (void)
 - Scheduler_IsEnabled.*
- [Thread_t Scheduler_GetCurrentThread](#) (void)
 - Scheduler_GetCurrentThread.*
- void [Thread_Init](#) ([Thread_t](#) handle, [K_WORD](#) *pwStack_, uint16_t u16StackSize_, [PORT_PRIO_TYPE](#) uXPriority_, [ThreadEntryFunc](#) pfEntryPoint_, void *pvArg_)
 - Thread_Init.*
- void [Thread_Start](#) ([Thread_t](#) handle)
 - Thread_Start.*
- void [Thread_Stop](#) ([Thread_t](#) handle)
 - Thread_Stop.*
- void [Thread_SetName](#) ([Thread_t](#) handle, const char *szName_)
- const char * [Thread_GetName](#) ([Thread_t](#) handle)
- [PORT_PRIO_TYPE](#) [Thread_GetPriority](#) ([Thread_t](#) handle)
 - Thread_GetPriority.*
- [PORT_PRIO_TYPE](#) [Thread_GetCurPriority](#) ([Thread_t](#) handle)
 - Thread_GetCurPriority.*
- void [Thread_SetQuantum](#) ([Thread_t](#) handle, uint16_t u16Quantum_)
- uint16_t [Thread_GetQuantum](#) ([Thread_t](#) handle)
- void [Thread_SetPriority](#) ([Thread_t](#) handle, [PORT_PRIO_TYPE](#) uXPriority_)

- Thread_SetPriority.*
- void [Thread_Exit](#) ([Thread_t](#) handle)
- Thread_Exit.*
- void [Thread_Sleep](#) ([uint32_t](#) u32TimeMs_)
- Thread_Sleep.*
- void [Thread_Yield](#) (void)
- Thread_Yield.*
- void [Thread_CoopYield](#) (void)
- Thread_CoopYield.*
- void [Thread_SetID](#) ([Thread_t](#) handle, [uint8_t](#) u8ID_)
- Thread_SetID.*
- [uint8_t](#) [Thread_GetID](#) ([Thread_t](#) handle)
- Thread_GetID.*
- [uint16_t](#) [Thread_GetStackSlack](#) ([Thread_t](#) handle)
- [thread_state_t](#) [Thread_GetState](#) ([Thread_t](#) handle)
- Thread_GetState.*
- void [Timer_Init](#) ([Timer_t](#) handle)
- Timer_Init.*
- void [Timer_Start](#) ([Timer_t](#) handle, bool bRepeat_, [uint32_t](#) u32IntervalMs_, [timer_callback_t](#) pfCallback_, void *pvData_)
- Timer_Start.*
- void [Timer_Stop](#) ([Timer_t](#) handle)
- Timer_Stop.*
- void [Timer_Restart](#) ([Timer_t](#) handle)
- Timer_Restart.*
- void [Semaphore_Init](#) ([Semaphore_t](#) handle, [uint16_t](#) u16InitVal_, [uint16_t](#) u16MaxVal_)
- Semaphore_Init.*
- void [Semaphore_Post](#) ([Semaphore_t](#) handle)
- Semaphore_Post.*
- void [Semaphore_Pend](#) ([Semaphore_t](#) handle)
- Semaphore_Pend.*
- bool [Semaphore_TimedPend](#) ([Semaphore_t](#) handle, [uint32_t](#) u32WaitTimeMS_)
- Semaphore_TimedPend.*
- void [Mutex_Init](#) ([Mutex_t](#) handle)
- Mutex_Init.*
- void [Mutex_Claim](#) ([Mutex_t](#) handle)
- Mutex_Claim.*
- void [Mutex_Release](#) ([Mutex_t](#) handle)
- Mutex_Release.*
- bool [Mutex_TimedClaim](#) ([Mutex_t](#) handle, [uint32_t](#) u32WaitTimeMS_)
- Mutex_TimedClaim.*
- void [Notify_Init](#) ([Notify_t](#) handle)
- Notify_Init.*
- void [Notify_Signal](#) ([Notify_t](#) handle)
- Notify_Signal.*
- void [Notify_Wait](#) ([Notify_t](#) handle, bool *pbFlag_)
- Notify_Wait.*
- bool [Notify_TimedWait](#) ([Notify_t](#) handle, [uint32_t](#) u32WaitTimeMS_, bool *pbFlag_)
- Notify_TimedWait.*
- [uint8_t](#) [Atomic_Set8](#) ([uint8_t](#) *pu8Source_, [uint8_t](#) u8Val_)
- Atomic_Set8.*

- uint16_t [Atomic_Set16](#) (uint16_t *pu16Source_, uint16_t u16Val_)
Atomic_Set16.
- uint32_t [Atomic_Set32](#) (uint32_t *pu32Source_, uint32_t u32Val_)
Atomic_Set32.
- uint8_t [Atomic_Add8](#) (uint8_t *pu8Source_, uint8_t u8Val_)
Atomic_Add8.
- uint16_t [Atomic_Add16](#) (uint16_t *pu16Source_, uint16_t u16Val_)
Atomic_Add16.
- uint32_t [Atomic_Add32](#) (uint32_t *pu32Source_, uint32_t u32Val_)
Atomic_Add32.
- uint8_t [Atomic_Sub8](#) (uint8_t *pu8Source_, uint8_t u8Val_)
Atomic_Sub8.
- uint16_t [Atomic_Sub16](#) (uint16_t *pu16Source_, uint16_t u16Val_)
Atomic_Sub16.
- uint32_t [Atomic_Sub32](#) (uint32_t *pu32Source_, uint32_t u32Val_)
Atomic_Sub32.
- bool [Atomic_TestAndSet](#) (bool *pbLock)
Atomic_TestAndSet.
- void [Message_Init](#) (Message_t handle)
Message_Init.
- void [Message_SetData](#) (Message_t handle, void *pvData_)
Message_SetData.
- void * [Message_GetData](#) (Message_t handle)
Message_GetData.
- void [Message_SetCode](#) (Message_t handle, uint16_t u16Code_)
Message_SetCode.
- uint16_t [Message_GetCode](#) (Message_t handle)
Message_GetCode.
- void [MessageQueue_Init](#) (MessageQueue_t handle)
MessageQueue_Init.
- Message_t [MessageQueue_Receive](#) (MessageQueue_t handle)
MessageQueue_Receive.
- void [MessagePool_Init](#) (MessagePool_t handle)
MessagePool_Init.
- void [MessagePool_Push](#) (MessagePool_t handle, Message_t msg)
MessagePool_Push.
- Message_t [MessagePool_Pop](#) (MessagePool_t handle)
MessagePool_Pop.
- Message_t [MessageQueue_TimedReceive](#) (MessageQueue_t handle, uint32_t u32TimeWaitMS_)
MessageQueue_TimedReceive.
- void [MessageQueue_Send](#) (MessageQueue_t handle, Message_t hMessage_)
MessageQueue_Send.
- uint16_t [MessageQueue_GetCount](#) (MessageQueue_t handle)
MessageQueue_GetCount.
- void [Mailbox_Init](#) (Mailbox_t handle, void *pvBuffer_, uint16_t u16BufferSize_, uint16_t u16ElementSize_)
Mailbox_Init.
- bool [Mailbox_Send](#) (Mailbox_t handle, void *pvData_)
Mailbox_Send.
- bool [Mailbox_SendTail](#) (Mailbox_t handle, void *pvData_)
Mailbox_SendTail.
- bool [Mailbox_TimedSend](#) (Mailbox_t handle, void *pvData_, uint32_t u32TimeoutMS_)

- Mailbox_TimedSend.*
- bool [Mailbox_TimedSendTail](#) ([Mailbox_t](#) handle, void *pvData_, uint32_t u32TimeoutMS_)
- Mailbox_TimedSendTail.*
- void [Mailbox_Receive](#) ([Mailbox_t](#) handle, void *pvData_)
- Mailbox_Receive.*
- void [Mailbox_ReceiveTail](#) ([Mailbox_t](#) handle, void *pvData_)
- Mailbox_ReceiveTail.*
- bool [Mailbox_TimedReceive](#) ([Mailbox_t](#) handle, void *pvData_, uint32_t u32TimeoutMS_)
- Mailbox_TimedReceive.*
- bool [Mailbox_TimedReceiveTail](#) ([Mailbox_t](#) handle, void *pvData_, uint32_t u32TimeoutMS_)
- Mailbox_TimedReceiveTail.*
- uint16_t [Mailbox_GetFreeSlots](#) ([Mailbox_t](#) handle)
- Mailbox_GetFreeSlots.*
- bool [Mailbox_IsFull](#) ([Mailbox_t](#) handle)
- Mailbox_IsFull.*
- bool [Mailbox_IsEmpty](#) ([Mailbox_t](#) handle)
- Mailbox_IsEmpty.*
- void [ConditionVariable_Init](#) ([ConditionVariable_t](#) handle)
- ConditionVariable_Init.*
- void [ConditionVariable_Wait](#) ([ConditionVariable_t](#) handle, [Mutex_t](#) hMutex_)
- ConditionVariable_Wait.*
- void [ConditionVariable_Signal](#) ([ConditionVariable_t](#) handle)
- ConditionVariable_Signal.*
- void [ConditionVariable_Broadcast](#) ([ConditionVariable_t](#) handle)
- ConditionVariable_Broadcast.*
- bool [ConditionVariable_TimedWait](#) ([ConditionVariable_t](#) handle, [Mutex_t](#) hMutex_, uint32_t u32WaitTimeMS_)
- ConditionVariable_TimedWait.*
- void [ReaderWriterLock_Init](#) ([ReaderWriterLock_t](#) handle)
- ReaderWriterLock_Init.*
- void [ReaderWriterLock_AcquireReader](#) ([ReaderWriterLock_t](#) handle)
- ReaderWriterLock_AcquireReader.*
- void [ReaderWriterLock_ReleaseReader](#) ([ReaderWriterLock_t](#) handle)
- ReaderWriterLock_ReleaseReader.*
- void [ReaderWriterLock_AcquireWriter](#) ([ReaderWriterLock_t](#) handle)
- ReaderWriterLock_AcquireWriter.*
- void [ReaderWriterLock_ReleaseWriter](#) ([ReaderWriterLock_t](#) handle)
- ReaderWriterLock_ReleaseWriter.*
- bool [ReaderWriterLock_TimedAcquireWriter](#) ([ReaderWriterLock_t](#) handle, uint32_t u32TimeoutMs_)
- ReaderWriterLock_TimedAcquireWriter.*
- bool [ReaderWriterLock_TimedAcquireReader](#) ([ReaderWriterLock_t](#) handle, uint32_t u32TimeoutMs_)
- ReaderWriterLock_TimedAcquireReader.*

20.1.1 Detailed Description

Implementation of C-language wrappers for the [Mark3](#) kernel.

Definition in file [mark3c.cpp](#).

20.1.2 Function Documentation

20.1.2.1 Alloc_ConditionVariable()

```
ConditionVariable_t Alloc_ConditionVariable (
    void )
```

Definition at line 134 of file [mark3c.cpp](#).

20.1.2.2 Alloc_EventFlag()

```
EventFlag_t Alloc_EventFlag (
    void )
```

Definition at line 64 of file [mark3c.cpp](#).

20.1.2.3 Alloc_Mailbox()

```
Mailbox_t Alloc_Mailbox (
    void )
```

Alloc_Mailbox.

See also

Mailbox* AutoAlloc::NewMailbox()

Returns

Handle to an allocated object, or nullptr if heap exhausted

Definition at line 123 of file [mark3c.cpp](#).

20.1.2.4 Alloc_Memory()

```
void* Alloc_Memory (
    uint16_t u16Size_ )
```

Definition at line 29 of file [mark3c.cpp](#).

20.1.2.5 Alloc_Message()

```
Message_t Alloc_Message (
    void )
```

Alloc_Message.

See also

AutoAlloc::NewMessage()

Returns

Handle to an allocated object, or nullptr if heap exhausted

Definition at line 76 of file [mark3c.cpp](#).

20.1.2.6 Alloc_MessagePool()

```
MessagePool_t Alloc_MessagePool (
    void )
```

Definition at line 100 of file [mark3c.cpp](#).

20.1.2.7 Alloc_MessageQueue()

```
MessageQueue_t Alloc_MessageQueue (
    void )
```

Alloc_MessageQueue.

See also

MessageQueue* AutoAlloc::NewMessageQueue()

Returns

Handle to an allocated object, or nullptr if heap exhausted

Definition at line 88 of file [mark3c.cpp](#).

20.1.2.8 Alloc_Mutex()

```
Mutex_t Alloc_Mutex (
    void )
```

Alloc_Mutex.

See also

Mutex* AutoAlloc::NewMutex()

Returns

Handle to an allocated object, or nullptr if heap exhausted

Definition at line 52 of file [mark3c.cpp](#).

20.1.2.9 Alloc_Notify()

```
Notify_t Alloc_Notify (
    void )
```

Alloc_Notify.

See also

Notify* AutoAlloc::NewNotify()

Returns

Handle to an allocated object, or nullptr if heap exhausted

Definition at line 112 of file [mark3c.cpp](#).

20.1.2.10 Alloc_ReaderWriterLock()

```
ReaderWriterLock_t Alloc_ReaderWriterLock (
    void )
```

Definition at line 145 of file [mark3c.cpp](#).

20.1.2.11 Alloc_Semaphore()

```
Semaphore_t Alloc_Semaphore (  
    void )
```

Alloc_Semaphore.

See also

Semaphore* AutoAlloc::NewSemaphore()

Returns

Handle to an allocated object, or nullptr if heap exhausted

Definition at line 40 of file [mark3c.cpp](#).

20.1.2.12 Alloc_Thread()

```
Thread_t Alloc_Thread (  
    void )
```

Alloc_Thread.

See also

Thread* AutoAlloc::NewThread()

Returns

Handle to an allocated object, or nullptr if heap exhausted

Definition at line 156 of file [mark3c.cpp](#).

20.1.2.13 Alloc_Timer()

```
Timer_t Alloc_Timer (  
    void )
```

Alloc_Timer.

See also

Timer* AutoAlloc::NewTimer()

Returns

Handle to an allocated object, or nullptr if heap exhausted

Definition at line 167 of file [mark3c.cpp](#).

20.1.2.14 Atomic_Add16()

```
uint16_t Atomic_Add16 (
    uint16_t * pu16Source_,
    uint16_t u16Val_ )
```

Atomic_Add16.

See also

`uint16_t Atomic::Add(uint16_t *pu16Source_, uint16_t u16Val_)`

Parameters

<i>pu16Source_</i>	Pointer to a variable
<i>u16Val_</i>	Value to add to the variable

Returns

Previously-held value in pu16Source_

Definition at line 638 of file [mark3c.cpp](#).

20.1.2.15 Atomic_Add32()

```
uint32_t Atomic_Add32 (
    uint32_t * pu32Source_,
    uint32_t u32Val_ )
```

Atomic_Add32.

See also

`uint32_t Atomic::Add(uint32_t *pu32Source_, uint32_t u32Val_)`

Parameters

<i>pu32Source_</i>	Pointer to a variable
<i>u32Val_</i>	Value to add to the variable

Returns

Previously-held value in pu32Source_

Definition at line 644 of file [mark3c.cpp](#).

20.1.2.16 Atomic_Add8()

```
uint8_t Atomic_Add8 (
    uint8_t * pu8Source_,
    uint8_t u8Val_ )
```

Atomic_Add8.

See also

`uint8_t Atomic::Add(uint8_t *pu8Source_, uint8_t u8Val_)`

Parameters

<i>pu8Source_</i>	Pointer to a variable
<i>u8Val_</i>	Value to add to the variable

Returns

Previously-held value in *pu8Source_*

Definition at line [632](#) of file [mark3c.cpp](#).

20.1.2.17 Atomic_Set16()

```
uint16_t Atomic_Set16 (
    uint16_t * pu16Source_,
    uint16_t u16Val_ )
```

Atomic_Set16.

See also

`uint16_t Atomic::Set(uint16_t *pu16Source_, uint16_t u16Val_)`

Parameters

<i>pu16Source_</i>	Pointer to a variable to set the value of
<i>u16Val_</i>	New value to set in the variable

Returns

Previously-set value

Definition at line [620](#) of file [mark3c.cpp](#).

20.1.2.18 Atomic_Set32()

```
uint32_t Atomic_Set32 (
    uint32_t * pu32Source_,
    uint32_t u32Val_ )
```

Atomic_Set32.

See also

`uint32_t Atomic::Set(uint32_t *pu32Source_, uint32_t u32Val_)`

Parameters

<i>pu32↵ Source_</i>	Pointer to a variable to set the value of
<i>u32Val_</i>	New value to set in the variable

Returns

Previously-set value

Definition at line [626](#) of file [mark3c.cpp](#).

20.1.2.19 Atomic_Set8()

```
uint8_t Atomic_Set8 (
    uint8_t * pu8Source_,
    uint8_t u8Val_ )
```

Atomic_Set8.

See also

`uint8_t Atomic::Set(uint8_t *pu8Source_, uint8_t u8Val_)`

Parameters

<i>pu8↵ Source_</i>	Pointer to a variable to set the value of
<i>u8Val_</i>	New value to set in the variable

Returns

Previously-set value

Definition at line [614](#) of file [mark3c.cpp](#).

20.1.2.20 Atomic_Sub16()

```
uint16_t Atomic_Sub16 (
    uint16_t * pu16Source_,
    uint16_t u16Val_ )
```

Atomic_Sub16.

See also

`uint16_t Atomic::Sub(uint16_t *pu16Source_, uint16_t u16Val_)`

Parameters

<i>pu16Source_</i>	Pointer to a variable
<i>u16Val_</i>	Value to subtract from the variable

Returns

Previously-held value in *pu16Source_*

Definition at line [656](#) of file [mark3c.cpp](#).

20.1.2.21 Atomic_Sub32()

```
uint32_t Atomic_Sub32 (
    uint32_t * pu32Source_,
    uint32_t u32Val_ )
```

Atomic_Sub32.

See also

`uint32_t Atomic::Sub(uint32_t *pu32Source_, uint32_t u32Val_)`

Parameters

<i>pu32Source_</i>	Pointer to a variable
<i>u32Val_</i>	Value to subtract from the variable

Returns

Previously-held value in *pu32Source_*

Definition at line [662](#) of file [mark3c.cpp](#).

20.1.2.22 Atomic_Sub8()

```
uint8_t Atomic_Sub8 (
    uint8_t * pu8Source_,
    uint8_t u8Val_ )
```

Atomic_Sub8.

See also

`uint8_t Atomic::Sub(uint8_t *pu8Source_, uint8_t u8Val_)`

Parameters

<i>pu8Source_</i>	Pointer to a variable
<i>u8Val_</i>	Value to subtract from the variable

Returns

Previously-held value in pu8Source_

Definition at line [650](#) of file [mark3c.cpp](#).

20.1.2.23 Atomic_TestAndSet()

```
bool Atomic_TestAndSet (
    bool * pbLock )
```

Atomic_TestAndSet.

See also

`bool Atomic::TestAndSet(bool *pbLock)`

Parameters

<i>pbLock</i>	Pointer to a value to test against. This will always be set to "true" at the end of a call to TestAndSet.
---------------	---

Returns

true - Lock value was "true" on entry, false - Lock was set

Definition at line [668](#) of file [mark3c.cpp](#).

20.1.2.24 ConditionVariable_Broadcast()

```
void ConditionVariable_Broadcast (
    ConditionVariable_t handle )
```

ConditionVariable_Broadcast.

See also

void ConditionVariable::Broadcast()

Parameters

<i>handle</i>	Handle of the condition variable object
---------------	---

Definition at line 874 of file [mark3c.cpp](#).

20.1.2.25 ConditionVariable_Init()

```
void ConditionVariable_Init (
    ConditionVariable_t handle )
```

ConditionVariable_Init.

See also

void ConditionVariable::Init()

Parameters

<i>handle</i>	Handle of the condition variable object
---------------	---

Definition at line 855 of file [mark3c.cpp](#).

20.1.2.26 ConditionVariable_Signal()

```
void ConditionVariable_Signal (
    ConditionVariable_t handle )
```

ConditionVariable_Signal.

See also

void ConditionVariable::Signal()

Parameters

<i>handle</i>	Handle of the condition variable object
---------------	---

Definition at line 868 of file [mark3c.cpp](#).

20.1.2.27 ConditionVariable_TimedWait()

```
bool ConditionVariable_TimedWait (
    ConditionVariable_t handle,
    Mutex_t hMutex_,
    uint32_t u32WaitTimeMS_ )
```

ConditionVariable_TimedWait.

See also

`bool ConditionVariable::Wait(Mutex* pclMutex_, uint32_t u32WaitTimeMS_)`

Parameters

<i>handle</i>	Handle of the condition variable object
<i>hMutex_</i>	Handle of the mutex to lock on acquisition of the condition variable
<i>u32WaitTimeMS_</i>	Maximum time to wait for object

Returns

true on success, false on timeout

Definition at line 880 of file [mark3c.cpp](#).

20.1.2.28 ConditionVariable_Wait()

```
void ConditionVariable_Wait (
    ConditionVariable_t handle,
    Mutex_t hMutex_ )
```

ConditionVariable_Wait.

See also

`void ConditionVariable::Wait(Mutex* pclMutex_)`

Parameters

<i>handle</i>	Handle of the condition variable object
<i>h</i> ↔ <i>Mutex</i> ↔ —	Handle of the mutex to lock on acquisition of the condition variable

Definition at line 861 of file [mark3c.cpp](#).

20.1.2.29 Free_ConditionVariable()

```
void Free_ConditionVariable (
    ConditionVariable_t handle )
```

Definition at line 139 of file [mark3c.cpp](#).

20.1.2.30 Free_EventFlag()

```
void Free_EventFlag (
    EventFlag_t handle )
```

Definition at line 70 of file [mark3c.cpp](#).

20.1.2.31 Free_Mailbox()

```
void Free_Mailbox (
    Mailbox_t handle )
```

Definition at line 128 of file [mark3c.cpp](#).

20.1.2.32 Free_Memory()

```
void Free_Memory (
    void * pvObject_ )
```

Free_Memory.

Parameters

<i>pv</i> ↔ <i>Object_</i>	Pointer to previously allocated block of memory
-------------------------------	---

Definition at line 34 of file [mark3c.cpp](#).

20.1.2.33 Free_Message()

```
void Free_Message (
    Message_t handle )
```

Definition at line 82 of file [mark3c.cpp](#).

20.1.2.34 Free_MessagePool()

```
void Free_MessagePool (
    MessagePool_t handle )
```

Definition at line 106 of file [mark3c.cpp](#).

20.1.2.35 Free_MessageQueue()

```
void Free_MessageQueue (
    MessageQueue_t handle )
```

Definition at line 94 of file [mark3c.cpp](#).

20.1.2.36 Free_Mutex()

```
void Free_Mutex (
    Mutex_t handle )
```

Definition at line 58 of file [mark3c.cpp](#).

20.1.2.37 Free_Notify()

```
void Free_Notify (
    Notify_t handle )
```

Definition at line 118 of file [mark3c.cpp](#).

20.1.2.38 Free_ReaderWriterLock()

```
void Free_ReaderWriterLock (
    ReaderWriterLock_t handle )
```

Definition at line 150 of file [mark3c.cpp](#).

20.1.2.39 Free_Semaphore()

```
void Free_Semaphore (
    Semaphore_t handle )
```

Definition at line 46 of file [mark3c.cpp](#).

20.1.2.40 Free_Thread()

```
void Free_Thread (
    Thread_t handle )
```

Definition at line 161 of file [mark3c.cpp](#).

20.1.2.41 Free_Timer()

```
void Free_Timer (
    Timer_t handle )
```

Definition at line 172 of file [mark3c.cpp](#).

20.1.2.42 Kernel_DebugPrint()

```
void Kernel_DebugPrint (
    const char * szString_ )
```

KernelDebug_DebugPrint.

See also

`void DebugPrint(const char* szString_)`

Parameters

<i>sz↔ String_</i>	String to print to debug interface
------------------------	------------------------------------

Definition at line 282 of file [mark3c.cpp](#).

20.1.2.43 `Kernel_GetStackGuardThreshold()`

```
uint16_t Kernel_GetStackGuardThreshold (
    void )
```

Definition at line 270 of file [mark3c.cpp](#).

20.1.2.44 `Kernel_GetTicks()`

```
uint32_t Kernel_GetTicks (
    void )
```

`Kernel_GetTicks`.

See also

`Kernel::GetTicks()`

Returns

Number of kernel ticks that have elapsed since boot

Definition at line 216 of file [mark3c.cpp](#).

20.1.2.45 `Kernel_Init()`

```
void Kernel_Init (
    void )
```

`Kernel_Init`.

See also

`void Kernel::Init()`

Definition at line 180 of file [mark3c.cpp](#).

20.1.2.46 Kernel_IsPanic()

```
bool Kernel_IsPanic (  
    void )
```

Kernel_IsPanic.

See also

bool Kernel::IsPanic()

Returns

Whether or not the kernel is in a panic state

Definition at line [204](#) of file [mark3c.cpp](#).

20.1.2.47 Kernel_IsStarted()

```
bool Kernel_IsStarted (  
    void )
```

Kernel_IsStarted.

See also

bool Kernel::IsStarted()

Returns

Whether or not the kernel has started - true = running, false = not started

Definition at line [192](#) of file [mark3c.cpp](#).

20.1.2.48 Kernel_Panic()

```
void Kernel_Panic (  
    uint16_t u16Cause_ )
```

Kernel_Panic.

See also

void Kernel::Panic(uint16_t u16Cause_)

Parameters

<i>u16↔ Cause_</i>	Reason for the kernel panic
------------------------	-----------------------------

Definition at line 210 of file [mark3c.cpp](#).

20.1.2.49 Kernel_SetDebugPrintFunction()

```
void Kernel_SetDebugPrintFunction (
    kernel_debug_print_t pfPrintFunction_ )
```

Kernel_SetDebugPrintFunction.

See also

void Kernel::SetDebugPrintFunction()

Parameters

<i>pfPrint↔ Function_</i>	Function to use to print debug information from the kernel
-------------------------------	--

Definition at line 276 of file [mark3c.cpp](#).

20.1.2.50 Kernel_SetPanic()

```
void Kernel_SetPanic (
    panic_func_t pfPanic_ )
```

Kernel_SetPanic.

See also

void Kernel::SetPanic(PanicFunc_t pfPanic_)

Parameters

<i>pf↔ Panic↔ _</i>	Panic function pointer
-----------------------------	------------------------

Definition at line 198 of file [mark3c.cpp](#).

20.1.2.51 Kernel_SetStackGuardThreshold()

```
void Kernel_SetStackGuardThreshold (
    uint16_t u16Threshold_ )
```

Definition at line 264 of file [mark3c.cpp](#).

20.1.2.52 Kernel_Start()

```
void Kernel_Start (
    void )
```

Kernel_Start.

See also

`void Kernel::Start()`

Definition at line 186 of file [mark3c.cpp](#).

20.1.2.53 Mailbox_GetFreeSlots()

```
uint16_t Mailbox_GetFreeSlots (
    Mailbox_t handle )
```

Mailbox_GetFreeSlots.

See also

`uint16_t Mailbox::GetFreeSlots()`

Parameters

<i>handle</i>	Handle of the mailbox object
---------------	------------------------------

Returns

Number of free slots in the mailbox

Definition at line 832 of file [mark3c.cpp](#).

20.1.2.54 Mailbox_Init()

```
void Mailbox_Init (
    Mailbox_t handle,
    void * pvBuffer_,
    uint16_t u16BufferSize_,
    uint16_t u16ElementSize_ )
```

Mailbox_Init.

See also

`void Mailbox::Init(void *pvBuffer_, uint16_t u16BufferSize_, uint16_t u16ElementSize_)`

Parameters

<i>handle</i>	Handle of the mailbox object
<i>pvBuffer_</i>	Pointer to the static buffer to use for the mailbox
<i>u16BufferSize_</i>	Size of the mailbox buffer, in bytes
<i>u16ElementSize_</i>	Size of each envelope, in bytes

Definition at line 769 of file [mark3c.cpp](#).

20.1.2.55 Mailbox_IsEmpty()

```
bool Mailbox_IsEmpty (
    Mailbox_t handle )
```

Mailbox_IsEmpty.

See also

`bool Mailbox::IsEmpty()`

Parameters

<i>handle</i>	Handle of the mailbox object
---------------	------------------------------

Returns

true if the mailbox is empty, false otherwise

Definition at line 846 of file [mark3c.cpp](#).

20.1.2.56 Mailbox_IsFull()

```
bool Mailbox_IsFull (
    Mailbox_t handle )
```

Mailbox_IsFull.

See also

bool Mailbox::IsFull()

Parameters

<i>handle</i>	Handle of the mailbox object
---------------	------------------------------

Returns

true if the mailbox is full, false otherwise

Definition at line 839 of file [mark3c.cpp](#).

20.1.2.57 Mailbox_Receive()

```
void Mailbox_Receive (
    Mailbox_t handle,
    void * pvData_ )
```

Mailbox_Receive.

See also

void Mailbox::Receive(void *pvData_)

Parameters

<i>handle</i>	Handle of the mailbox object
<i>pv↔ Data_</i>	Pointer to a buffer that will have the envelope's contents copied into upon delivery.

Definition at line 804 of file [mark3c.cpp](#).

20.1.2.58 Mailbox_ReceiveTail()

```
void Mailbox_ReceiveTail (
    Mailbox_t handle,
    void * pvData_ )
```

Mailbox_ReceiveTail.

See also

void Mailbox::ReceiveTail(void *pvData_)

Parameters

<i>handle</i>	Handle of the mailbox object
<i>pvData_</i>	Pointer to a buffer that will have the envelope's contents copied into upon delivery.

Definition at line 811 of file [mark3c.cpp](#).

20.1.2.59 Mailbox_Send()

```
bool Mailbox_Send (
    Mailbox_t handle,
    void * pvData_ )
```

Mailbox_Send.

See also

bool Mailbox::Send(void *pvData_)

Parameters

<i>handle</i>	Handle of the mailbox object
<i>pvData_</i>	Pointer to the data object to send to the mailbox.

Returns

true - envelope was delivered, false - mailbox is full.

Definition at line 776 of file [mark3c.cpp](#).

20.1.2.60 Mailbox_SendTail()

```
bool Mailbox_SendTail (
    Mailbox_t handle,
    void * pvData_ )
```

Mailbox_SendTail.

See also

bool Mailbox::SendTail(void *pvData_)

Parameters

<i>handle</i>	Handle of the mailbox object
<i>pvData_</i>	Pointer to the data object to send to the mailbox.

Returns

true - envelope was delivered, false - mailbox is full.

Definition at line 783 of file [mark3c.cpp](#).

20.1.2.61 Mailbox_TimedReceive()

```
bool Mailbox_TimedReceive (  
    Mailbox_t handle,  
    void * pvData_,  
    uint32_t u32TimeoutMS_ )
```

Mailbox_TimedReceive.

See also

bool Mailbox::Receive(void *pvData_, uint32_t u32TimeoutMS_)

Parameters

<i>handle</i>	Handle of the mailbox object
<i>pvData_</i>	Pointer to a buffer that will have the envelope's contents copied into upon delivery.
<i>u32TimeoutMS_</i>	Maximum time to wait for delivery.

Returns

true - envelope was delivered, false - delivery timed out.

Definition at line 818 of file [mark3c.cpp](#).

20.1.2.62 Mailbox_TimedReceiveTail()

```
bool Mailbox_TimedReceiveTail (  
    Mailbox_t handle,  
    void * pvData_,  
    uint32_t u32TimeoutMS_ )
```

Mailbox_TimedReceiveTail.

See also

```
bool Mailbox::ReceiveTail(void *pvData_, uint32_t u32TimeoutMS_)
```

Parameters

<i>handle</i>	Handle of the mailbox object
<i>pvData_</i>	Pointer to a buffer that will have the envelope's contents copied into upon delivery.
<i>u32TimeoutMS_</i>	Maximum time to wait for delivery.

Returns

true - envelope was delivered, false - delivery timed out.

Definition at line 825 of file [mark3c.cpp](#).

20.1.2.63 Mailbox_TimedSend()

```
bool Mailbox_TimedSend (  
    Mailbox_t handle,  
    void * pvData_,  
    uint32_t u32TimeoutMS_ )
```

Mailbox_TimedSend.

See also

```
bool Mailbox::Send(void *pvData_, uint32_t u32TimeoutMS_)
```

Parameters

<i>handle</i>	Handle of the mailbox object
<i>pvData_</i>	Pointer to the data object to send to the mailbox.
<i>u32TimeoutMS_</i>	Maximum time to wait for a free transmit slot

Returns

true - envelope was delivered, false - mailbox is full.

Definition at line 790 of file [mark3c.cpp](#).

20.1.2.64 Mailbox_TimedSendTail()

```
bool Mailbox_TimedSendTail (
    Mailbox_t handle,
    void * pvData_,
    uint32_t u32TimeoutMS_ )
```

Mailbox_TimedSendTail.

See also

`bool Mailbox::Send(void *pvData_, uint32_t u32TimeoutMS_)`

Parameters

<i>handle</i>	Handle of the mailbox object
<i>pvData_</i>	Pointer to the data object to send to the mailbox.
<i>u32TimeoutMS_</i>	Maximum time to wait for a free transmit slot

Returns

true - envelope was delivered, false - mailbox is full.

Definition at line 797 of file [mark3c.cpp](#).

20.1.2.65 Message_GetCode()

```
uint16_t Message_GetCode (
    Message_t handle )
```

Message_GetCode.

See also

`uint16_t Message::GetCode()`

Parameters

<i>handle</i>	Handle of the message object
---------------	------------------------------

Returns

user code set in the object

Definition at line 704 of file [mark3c.cpp](#).

20.1.2.66 Message_GetData()

```
void* Message_GetData (
    Message_t handle )
```

Message_GetData.

See also

void* Message::GetData()

Parameters

<i>handle</i>	Handle of the message object
---------------	------------------------------

Returns

Pointer to the data set in the message object

Definition at line 690 of file [mark3c.cpp](#).

20.1.2.67 Message_Init()

```
void Message_Init (
    Message_t handle )
```

Message_Init.

See also

void Message::Init()

Parameters

<i>handle</i>	Handle of the message object
---------------	------------------------------

Definition at line 676 of file [mark3c.cpp](#).

20.1.2.68 Message_SetCode()

```
void Message_SetCode (
    Message_t handle,
    uint16_t u16Code_ )
```

Message_SetCode.

See also

```
void Message::SetCode(uint16_t u16Code_)
```

Parameters

<i>handle</i>	Handle of the message object
<i>u16↔ Code_</i>	Data code to set in the object

Definition at line 697 of file [mark3c.cpp](#).

20.1.2.69 Message_SetData()

```
void Message_SetData (
    Message_t handle,
    void * pvData_ )
```

Message_SetData.

See also

```
void Message::SetData(void *pvData_)
```

Parameters

<i>handle</i>	Handle of the message object
<i>pv↔ Data_</i>	Pointer to the data object to send in the message

Definition at line 683 of file [mark3c.cpp](#).

20.1.2.70 MessagePool_Init()

```
void MessagePool_Init (
    MessagePool_t handle )
```

MessagePool_Init.

See also

```
void MessagePool::Init()
```

Parameters

<i>handle</i>	Handle of the message pool object
---------------	-----------------------------------

Definition at line 725 of file [mark3c.cpp](#).

20.1.2.71 MessagePool_Pop()

```
Message_t MessagePool_Pop (  
    MessagePool_t handle )
```

MessagePool_Pop.

See also

Message* MessagePool::Pop()

Parameters

<i>handle</i>	Handle of the message pool object
---------------	-----------------------------------

Returns

Handle to a Message object, or nullptr on allocation error

Definition at line 739 of file [mark3c.cpp](#).

20.1.2.72 MessagePool_Push()

```
void MessagePool_Push (  
    MessagePool_t handle,  
    Message_t msg )
```

MessagePool_Push.

See also

void MessagePool::Push(Message* pclMessage_)

Parameters

<i>handle</i>	Handle of the message pool object
<i>msg</i>	Message object to return back to the pool

Definition at line 732 of file [mark3c.cpp](#).

20.1.2.73 MessageQueue_GetCount()

```
uint16_t MessageQueue_GetCount (
    MessageQueue_t handle )
```

MessageQueue_GetCount.

See also

`uint16_t MessageQueue::GetCount()`

Returns

Count of pending messages in the queue.

Definition at line 760 of file [mark3c.cpp](#).

20.1.2.74 MessageQueue_Init()

```
void MessageQueue_Init (
    MessageQueue_t handle )
```

MessageQueue_Init.

See also

`void MessageQueue::Init()`

Parameters

<i>handle</i>	Handle to the message queue to initialize
---------------	---

Definition at line 711 of file [mark3c.cpp](#).

20.1.2.75 MessageQueue_Receive()

```
Message_t MessageQueue_Receive (
    MessageQueue_t handle )
```

MessageQueue_Receive.

See also

[Message_t](#) MessageQueue::Receive()

Parameters

<i>handle</i>	Handle of the message queue object
---------------	------------------------------------

Returns

Pointer to a message object at the head of the queue

Definition at line 718 of file [mark3c.cpp](#).

20.1.2.76 MessageQueue_Send()

```
void MessageQueue_Send (
    MessageQueue_t handle,
    Message_t hMessage_ )
```

MessageQueue_Send.

See also

void MessageQueue::Send(Message *pciMessage_)

Parameters

<i>handle</i>	Handle of the message queue object
<i>h</i> ↔ <i>Message</i> ↔ —	Handle to the message to send to the given queue

Definition at line 753 of file [mark3c.cpp](#).

20.1.2.77 MessageQueue_TimedReceive()

```
Message_t MessageQueue_TimedReceive (
    MessageQueue_t handle,
    uint32_t u32TimeWaitMS_ )
```

MessageQueue_TimedReceive.

See also

[Message_t](#) MessageQueue::TimedReceive(uint32_t u32TimeWaitMS_)

Parameters

<i>handle</i>	Handle of the message queue object
<i>u32TimeWaitMS_</i>	The amount of time in ms to wait for a message before timing out and unblocking the waiting thread.

Returns

Pointer to a message object at the head of the queue or nullptr on timeout.

Definition at line 746 of file [mark3c.cpp](#).

20.1.2.78 Mutex_Claim()

```
void Mutex_Claim (
    Mutex_t handle )
```

Mutex_Claim.

See also

void Mutex::Claim()

Parameters

<i>handle</i>	Handle of the mutex
---------------	---------------------

Definition at line 515 of file [mark3c.cpp](#).

20.1.2.79 Mutex_Init()

```
void Mutex_Init (
    Mutex_t handle )
```

Mutex_Init.

See also

void Mutex::Init()

Parameters

<i>handle</i>	Handle of the mutex
---------------	---------------------

Definition at line 508 of file [mark3c.cpp](#).

20.1.2.80 Mutex_Release()

```
void Mutex_Release (
    Mutex_t handle )
```

Mutex_Release.

See also

`void Mutex::Release()`

Parameters

<i>handle</i>	Handle of the mutex
---------------	---------------------

Definition at line 522 of file [mark3c.cpp](#).

20.1.2.81 Mutex_TimedClaim()

```
bool Mutex_TimedClaim (
    Mutex_t handle,
    uint32_t u32WaitTimeMS_ )
```

Mutex_TimedClaim.

See also

`bool Mutex::Claim(uint32_t u32WaitTimeMS_)`

Parameters

<i>handle</i>	Handle of the mutex
<i>u32WaitTimeMS__S</i>	Time to wait before aborting

Returns

true if mutex was claimed, false on timeout

Definition at line 529 of file [mark3c.cpp](#).

20.1.2.82 Notify_Init()

```
void Notify_Init (
    Notify_t handle )
```

Notify_Init.

See also

void Notify::Init()

Parameters

<i>handle</i>	Handle of the notification object
---------------	-----------------------------------

Definition at line 584 of file [mark3c.cpp](#).

20.1.2.83 Notify_Signal()

```
void Notify_Signal (
    Notify_t handle )
```

Notify_Signal.

See also

void Notify::Signal()

Parameters

<i>handle</i>	Handle of the notification object
---------------	-----------------------------------

Definition at line 591 of file [mark3c.cpp](#).

20.1.2.84 Notify_TimedWait()

```
bool Notify_TimedWait (
    Notify_t handle,
    uint32_t u32WaitTimeMS_,
    bool * pbFlag_ )
```

Notify_TimedWait.

See also

bool Notify::Wait(uint32_t u32WaitTimeMS_, bool *pbFlag_)

Parameters

<i>handle</i>	Handle of the notification object
<i>u32WaitTimeMS_</i>	Maximum time to wait for notification in ms
<i>pbFlag_</i>	Flag to set to true on notification

Returns

true on unblock, false on timeout

Definition at line 605 of file [mark3c.cpp](#).

20.1.2.85 Notify_Wait()

```
void Notify_Wait (
    Notify_t handle,
    bool * pbFlag_ )
```

Notify_Wait.

See also

`void Notify::Wait(bool *pbFlag_)`

Parameters

<i>handle</i>	Handle of the notification object
<i>pbFlag_</i>	Flag to set to true on notification

Definition at line 598 of file [mark3c.cpp](#).

20.1.2.86 ReaderWriterLock_AcquireReader()

```
void ReaderWriterLock_AcquireReader (
    ReaderWriterLock_t handle )
```

ReaderWriterLock_AcquireReader.

See also

`void ReaderWriterLock::AcquireReader()`

Parameters

<i>handle</i>	Handle of the reader-writer object
---------------	------------------------------------

Definition at line 897 of file [mark3c.cpp](#).

20.1.2.87 ReaderWriterLock_AcquireWriter()

```
void ReaderWriterLock_AcquireWriter (  
    ReaderWriterLock_t handle )
```

ReaderWriterLock_AcquireWriter.

See also

`void ReaderWriterLock::AcquireWriter()`

Parameters

<i>handle</i>	Handle of the reader-writer object
---------------	------------------------------------

Definition at line 911 of file [mark3c.cpp](#).

20.1.2.88 ReaderWriterLock_Init()

```
void ReaderWriterLock_Init (  
    ReaderWriterLock_t handle )
```

ReaderWriterLock_Init.

See also

`void ReaderWriterLock::Init()`

Parameters

<i>handle</i>	Handle of the reader-writer object
---------------	------------------------------------

Definition at line 890 of file [mark3c.cpp](#).

20.1.2.89 ReaderWriterLock_ReleaseReader()

```
void ReaderWriterLock_ReleaseReader (
    ReaderWriterLock_t handle )
```

ReaderWriterLock_ReleaseReader.

See also

void ReaderWriterLock::ReleaseReader()

Parameters

<i>handle</i>	Handle of the reader-writer object
---------------	------------------------------------

Definition at line 904 of file [mark3c.cpp](#).

20.1.2.90 ReaderWriterLock_ReleaseWriter()

```
void ReaderWriterLock_ReleaseWriter (
    ReaderWriterLock_t handle )
```

ReaderWriterLock_ReleaseWriter.

See also

void ReaderWriterLock::ReleaseWriter()

Parameters

<i>handle</i>	Handle of the reader-writer object
---------------	------------------------------------

Definition at line 918 of file [mark3c.cpp](#).

20.1.2.91 ReaderWriterLock_TimedAcquireReader()

```
bool ReaderWriterLock_TimedAcquireReader (
    ReaderWriterLock_t handle,
    uint32_t u32TimeoutMs_ )
```

ReaderWriterLock_TimedAcquireReader.

See also

bool ReaderWriterLock::AcquireReader(uint32_t u32TimeoutMs_)

Parameters

<i>handle</i>	Handle of the reader-writer object
<i>u32Timeout</i> ↔ <i>Ms_</i>	Maximum time to wait for the reader lock before bailing

Returns

true on success, false on timeout

Definition at line 932 of file [mark3c.cpp](#).

20.1.2.92 ReaderWriterLock_TimedAcquireWriter()

```
bool ReaderWriterLock_TimedAcquireWriter (
    ReaderWriterLock_t handle,
    uint32_t u32TimeoutMs_ )
```

ReaderWriterLock_TimedAcquireWriter.

See also

`bool ReaderWriterLock::AcquireWriter(uint32_t u32TimeoutMs_)`

Parameters

<i>handle</i>	Handle of the reader-writer object
<i>u32Timeout</i> ↔ <i>Ms_</i>	Maximum time to wait for the writer lock before bailing

Returns

true on success, false on timeout

Definition at line 925 of file [mark3c.cpp](#).

20.1.2.93 Scheduler_Enable()

```
void Scheduler_Enable (
    bool bEnable_ )
```

Scheduler_Enable.

See also

`void Scheduler::SetScheduler(bool bEnable_)`

Parameters

<i>b</i> ↔ <i>Enable</i> ↔ —	true to enable, false to disable the scheduler
------------------------------------	--

Definition at line 289 of file [mark3c.cpp](#).

20.1.2.94 Scheduler_GetCurrentThread()

```
Thread_t Scheduler_GetCurrentThread (
    void )
```

Scheduler_GetCurrentThread.

See also

Thread* Scheduler::GetCurrentThread()

Returns

Handle of the currently-running thread

Definition at line 301 of file [mark3c.cpp](#).

20.1.2.95 Scheduler_IsEnabled()

```
bool Scheduler_IsEnabled (
    void )
```

Scheduler_IsEnabled.

See also

bool Scheduler::IsEnabled()

Returns

true - scheduler enabled, false - disabled

Definition at line 295 of file [mark3c.cpp](#).

20.1.2.96 Semaphore_Init()

```
void Semaphore_Init (
    Semaphore_t handle,
    uint16_t u16InitVal_,
    uint16_t u16MaxVal_ )
```

Semaphore_Init.

See also

void Semaphore::Init(uint16_t u16InitVal_, uint16_t u16MaxVal_)

Parameters

<i>handle</i>	Handle of the semaphore
<i>u16InitVal</i> \leftrightarrow —	Initial value of the semaphore
<i>u16Max</i> \leftrightarrow <i>Val_</i>	Maximum value that can be held for a semaphore

Definition at line 478 of file [mark3c.cpp](#).

20.1.2.97 Semaphore_Pend()

```
void Semaphore_Pend (
    Semaphore_t handle )
```

Semaphore_Pend.

See also

`void Semaphore::Pend()`

Parameters

<i>handle</i>	Handle of the semaphore
---------------	-------------------------

Definition at line 492 of file [mark3c.cpp](#).

20.1.2.98 Semaphore_Post()

```
void Semaphore_Post (
    Semaphore_t handle )
```

Semaphore_Post.

See also

`void Semaphore::Post()`

Parameters

<i>handle</i>	Handle of the semaphore
---------------	-------------------------

Definition at line 485 of file [mark3c.cpp](#).

20.1.2.99 Semaphore_TimedPend()

```
bool Semaphore_TimedPend (
    Semaphore_t handle,
    uint32_t u32WaitTimeMS_ )
```

Semaphore_TimedPend.

See also

bool Semaphore::Pend(uint32_t u32WaitTimeMS_)

Parameters

<i>handle</i>	Handle of the semaphore
<i>u32WaitTimeMS_</i>	Time in ms to wait

Returns

true if semaphore was acquired, false on timeout

Definition at line 499 of file [mark3c.cpp](#).

20.1.2.100 Thread_CoopYield()

```
void Thread_CoopYield (
    void )
```

Thread_CoopYield.

See also

void Thread::CoopYield()

Definition at line 415 of file [mark3c.cpp](#).

20.1.2.101 Thread_Exit()

```
void Thread_Exit (
    Thread_t handle )
```

Thread_Exit.

See also

void Thread::Exit()

Parameters

<i>handle</i>	Handle of the thread
---------------	----------------------

Definition at line 381 of file [mark3c.cpp](#).

20.1.2.102 Thread_GetCurPriority()

```
PORT_PRIO_TYPE Thread_GetCurPriority (
    Thread_t handle )
```

Thread_GetCurPriority.

See also

[PORT_PRIO_TYPE](#) Thread::GetCurPriority()

Parameters

<i>handle</i>	Handle of the thread
---------------	----------------------

Returns

Current priority of the thread considering priority inheritance

Definition at line 354 of file [mark3c.cpp](#).

20.1.2.103 Thread_GetID()

```
uint8_t Thread_GetID (
    Thread_t handle )
```

Thread_GetID.

See also

[uint8_t](#) Thread::GetID()

Parameters

<i>handle</i>	Handle of the thread
---------------	----------------------

Returns

Return ID assigned to the thread

Definition at line 426 of file [mark3c.cpp](#).

20.1.2.104 Thread_GetName()

```
const char* Thread_GetName (  
    Thread_t handle )
```

Definition at line 341 of file [mark3c.cpp](#).

20.1.2.105 Thread_GetPriority()

```
PORT_PRIO_TYPE Thread_GetPriority (  
    Thread_t handle )
```

Thread_GetPriority.

See also

[PORT_PRIO_TYPE Thread::GetPriority\(\)](#)

Parameters

<i>handle</i>	Handle of the thread
---------------	----------------------

Returns

Current priority of the thread not considering priority inheritance

Definition at line 348 of file [mark3c.cpp](#).

20.1.2.106 Thread_GetQuantum()

```
uint16_t Thread_GetQuantum (  
    Thread_t handle )
```

Definition at line 367 of file [mark3c.cpp](#).

20.1.2.107 Thread_GetStackSlack()

```
uint16_t Thread_GetStackSlack (
    Thread_t handle )
```

Definition at line 432 of file [mark3c.cpp](#).

20.1.2.108 Thread_GetState()

```
thread_state_t Thread_GetState (
    Thread_t handle )
```

Thread_GetState.

See also

ThreadState Thread::GetState()

Parameters

<i>handle</i>	Handle of the thread
---------------	----------------------

Returns

The thread's current execution state

Definition at line 438 of file [mark3c.cpp](#).

20.1.2.109 Thread_Init()

```
void Thread_Init (
    Thread_t handle,
    K_WORD * pwStack_,
    uint16_t u16StackSize_,
    PORT_PRIO_TYPE uXPriority_,
    thread_entry_func_t pfEntryPoint_,
    void * pvArg_ )
```

Thread_Init.

See also

void Thread::Init(K_WORD *pwStack_, uint16_t u16StackSize_, PORT_PRIO_TYPE uXPriority_, Thread↵
Entry_t pfEntryPoint_, void *pvArg_)

Parameters

<i>handle</i>	Handle of the thread to initialize
<i>pwStack_</i>	Pointer to the stack to use for the thread
<i>u16Stack↔ Size_</i>	Size of the stack (in bytes)
<i>uXPriority_</i>	Priority of the thread (0 = idle, 7 = max)
<i>pfEntryPoint↔ _</i>	This is the function that gets called when the thread is started
<i>pvArg_</i>	Pointer to the argument passed into the thread's entypoint function.

Definition at line 310 of file [mark3c.cpp](#).

20.1.2.110 Thread_SetID()

```
void Thread_SetID (
    Thread_t handle,
    uint8_t u8ID_ )
```

Thread_SetID.

See also

```
void Thread::SetID(uint8_t u8ID_)
```

Parameters

<i>handle</i>	Handle of the thread
<i>u8l↔ D_</i>	ID To assign to the thread

Definition at line 420 of file [mark3c.cpp](#).

20.1.2.111 Thread_SetName()

```
void Thread_SetName (
    Thread_t handle,
    const char * szName_ )
```

Definition at line 335 of file [mark3c.cpp](#).

20.1.2.112 Thread_SetPriority()

```
void Thread_SetPriority (
    Thread_t handle,
    PORT_PRIO_TYPE uXPriority_ )
```

Thread_SetPriority.

See also

void Thread::SetPriority(PORT_PRIO_TYPE uXPriority_)

Parameters

<i>handle</i>	Handle of the thread
<i>uX\leftrightarrowPriority_</i>	New priority level

Definition at line 374 of file [mark3c.cpp](#).

20.1.2.113 Thread_SetQuantum()

```
void Thread_SetQuantum (
    Thread_t handle,
    uint16_t u16Quantum_ )
```

Definition at line 361 of file [mark3c.cpp](#).

20.1.2.114 Thread_Sleep()

```
void Thread_Sleep (
    uint32_t u32TimeMs_ )
```

Thread_Sleep.

See also

void Thread::Sleep(uint32_t u32TimeMs_)

Parameters

<i>u32Time\leftrightarrowMs_</i>	Time in ms to block the thread for
---	------------------------------------

Definition at line 388 of file [mark3c.cpp](#).

20.1.2.115 Thread_Start()

```
void Thread_Start (
    Thread_t handle )
```

Thread_Start.

See also

void Thread::Start()

Parameters

<i>handle</i>	Handle of the thread to start
---------------	-------------------------------

Definition at line 322 of file [mark3c.cpp](#).

20.1.2.116 Thread_Stop()

```
void Thread_Stop (
    Thread_t handle )
```

Thread_Stop.

See also

void Thread::Stop()

Parameters

<i>handle</i>	Handle of the thread to stop
---------------	------------------------------

Definition at line 329 of file [mark3c.cpp](#).

20.1.2.117 Thread_Yield()

```
void Thread_Yield (
    void )
```

Thread_Yield.

See also

`void Thread::Yield()`

Definition at line 410 of file [mark3c.cpp](#).

20.1.2.118 `Timer_Init()`

```
void Timer_Init (
    Timer\_t handle )
```

`Timer_Init`.

See also

`void Timer::Init()`

Parameters

<i>handle</i>	Handle of the timer
---------------	---------------------

Definition at line 448 of file [mark3c.cpp](#).

20.1.2.119 `Timer_Restart()`

```
void Timer_Restart (
    Timer\_t handle )
```

`Timer_Restart`.

See also

`void Timer::Start()`

Parameters

<i>handle</i>	Handle of the timer to restart.
---------------	---------------------------------

Definition at line 469 of file [mark3c.cpp](#).

20.1.2.120 `Timer_Start()`

```
void Timer_Start (
    Timer\_t handle,
```



```

00007      |_____|      |_____|      |_____|      |_____|
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  =====*/
00021  #include "mark3c.h"
00022  #include "mark3.h"
00023
00024  using namespace Mark3;
00025
00026  //-----
00027  // Kernel Memory managment APIs
00028  //-----
00029  void* Alloc_Memory(uint16_t ul6Size_)
00030  {
00031      return AutoAlloc::NewRawData(ul6Size_);
00032  }
00033  //-----
00034  void Free_Memory(void* pvObject_)
00035  {
00036      AutoAlloc::DestroyRawData(pvObject_);
00037  }
00038
00039  //-----
00040  Semaphore_t Alloc_Semaphore(void)
00041  {
00042      return (Semaphore_t)AutoAlloc::NewObject<Semaphore, AutoAllocType::Semaphore>();
00043  }
00044
00045  //-----
00046  void Free_Semaphore(Semaphore_t handle)
00047  {
00048      AutoAlloc::DestroyObject<Semaphore, AutoAllocType::Semaphore>((Semaphore*)handle);
00049  }
00050
00051  //-----
00052  Mutex_t Alloc_Mutex(void)
00053  {
00054      return (Mutex_t)AutoAlloc::NewObject<Mutex, AutoAllocType::Mutex>();
00055  }
00056
00057  //-----
00058  void Free_Mutex(Mutex_t handle)
00059  {
00060      AutoAlloc::DestroyObject<Mutex, AutoAllocType::Mutex>((Mutex*)handle);
00061  }
00062
00063  //-----
00064  EventFlag_t Alloc_EventFlag(void)
00065  {
00066      return (EventFlag_t)AutoAlloc::NewObject<EventFlag, AutoAllocType::EventFlag>();
00067  }
00068
00069  //-----
00070  void Free_EventFlag(EventFlag_t handle)
00071  {
00072      AutoAlloc::DestroyObject<EventFlag, AutoAllocType::EventFlag>((EventFlag*)handle);
00073  }
00074
00075  //-----
00076  Message_t Alloc_Message(void)
00077  {
00078      return (Message_t)AutoAlloc::NewObject<Message, AutoAllocType::Message>();
00079  }
00080
00081  //-----
00082  void Free_Message(Message_t handle)
00083  {
00084      AutoAlloc::DestroyObject<Message, AutoAllocType::Message>((Message*)handle);
00085  }
00086
00087  //-----
00088  MessageQueue_t Alloc_MessageQueue(void)
00089  {
00090      return (MessageQueue_t)AutoAlloc::NewObject<MessageQueue, AutoAllocType::MessageQueue>();
00091  }
00092
00093  //-----
00094  void Free_MessageQueue(MessageQueue_t handle)
00095  {
00096      AutoAlloc::DestroyObject<MessageQueue, AutoAllocType::MessageQueue>((
00097          MessageQueue*)handle);
00098  }
00099  //-----

```

```

00100 MessagePool_t Alloc_MessagePool(void)
00101 {
00102     return (MessagePool_t)AutoAlloc::NewObject<MessagePool, AutoAllocType::MessagePool>();
00103 }
00104
00105 //-----
00106 void Free_MessagePool(MessagePool_t handle)
00107 {
00108     AutoAlloc::DestroyObject<MessagePool, AutoAllocType::MessagePool>((
MessagePool*)handle);
00109 }
00110
00111 //-----
00112 Notify_t Alloc_Notify(void)
00113 {
00114     return (Notify_t)AutoAlloc::NewObject<Notify, AutoAllocType::Notify>();
00115 }
00116
00117 //-----
00118 void Free_Notify(Notify_t handle)
00119 {
00120     AutoAlloc::DestroyObject<Notify, AutoAllocType::Notify>((Notify*)handle);
00121 }
00122 //-----
00123 Mailbox_t Alloc_Mailbox(void)
00124 {
00125     return (Mailbox_t)AutoAlloc::NewObject<Mailbox, AutoAllocType::MailBox>();
00126 }
00127 //-----
00128 void Free_Mailbox(Mailbox_t handle)
00129 {
00130     AutoAlloc::DestroyObject<Mailbox, AutoAllocType::MailBox>((Mailbox*)handle);
00131 }
00132
00133 //-----
00134 ConditionVariable_t Alloc_ConditionVariable(void)
00135 {
00136     return (ConditionVariable_t)AutoAlloc::NewObject<ConditionVariable,
AutoAllocType::ConditionVariable>();
00137 }
00138 //-----
00139 void Free_ConditionVariable(ConditionVariable_t handle)
00140 {
00141     AutoAlloc::DestroyObject<ConditionVariable, AutoAllocType::ConditionVariable>((
ConditionVariable*)handle);
00142 }
00143
00144 //-----
00145 ReaderWriterLock_t Alloc_ReaderWriterLock(void)
00146 {
00147     return (ReaderWriterLock_t)AutoAlloc::NewObject<ReaderWriterLock,
AutoAllocType::ReaderWriterLock>();
00148 }
00149 //-----
00150 void Free_ReaderWriterLock(ReaderWriterLock_t handle)
00151 {
00152     AutoAlloc::DestroyObject<ReaderWriterLock, AutoAllocType::ReaderWriterLock>((
ReaderWriterLock*)handle);
00153 }
00154
00155 //-----
00156 Thread_t Alloc_Thread(void)
00157 {
00158     return (Thread_t)AutoAlloc::NewObject<Thread, AutoAllocType::Thread>();
00159 }
00160 //-----
00161 void Free_Thread(Thread_t handle)
00162 {
00163     AutoAlloc::DestroyObject<Thread, AutoAllocType::Thread>((Thread*)handle);
00164 }
00165
00166 //-----
00167 Timer_t Alloc_Timer(void)
00168 {
00169     return (Thread_t)AutoAlloc::NewObject<Timer, AutoAllocType::Timer>();
00170 }
00171 //-----
00172 void Free_Timer(Timer_t handle)
00173 {
00174     AutoAlloc::DestroyObject<Timer, AutoAllocType::Timer>((Timer*)handle);
00175 }
00176
00177 //-----
00178 // Kernel APIs
00179 //-----
00180 void Kernel_Init(void)
00181 {

```



```

00182     Kernel::Init();
00183 }
00184
00185 //-----
00186 void Kernel_Start(void)
00187 {
00188     Kernel::Start();
00189 }
00190
00191 //-----
00192 bool Kernel_IsStarted(void)
00193 {
00194     return Kernel::IsStarted();
00195 }
00196
00197 //-----
00198 void Kernel_SetPanic(PanicFunc pfPanic_)
00199 {
00200     Kernel::SetPanic(pfPanic_);
00201 }
00202
00203 //-----
00204 bool Kernel_IsPanic(void)
00205 {
00206     return Kernel::IsPanic();
00207 }
00208
00209 //-----
00210 void Kernel_Panic(uint16_t ul6Cause_)
00211 {
00212     Kernel::Panic(ul6Cause_);
00213 }
00214
00215 //-----
00216 uint32_t Kernel_GetTicks(void)
00217 {
00218     return Kernel::GetTicks();
00219 }
00220
00221 #if KERNEL_THREAD_CREATE_CALLOUT
00222 //-----
00223 void Kernel_SetThreadCreateCallout(thread_create_callout_t pfCreate_)
00224 {
00225     Kernel::SetThreadCreateCallout((
00226         ThreadCreateCallout)pfCreate_);
00227 }
00228 #endif // #if KERNEL_THREAD_CREATE_CALLOUT
00229 //-----
00229 #if KERNEL_THREAD_EXIT_CALLOUT
00230 void Kernel_SetThreadExitCallout(thread_exit_callout_t pfExit_)
00231 {
00232     Kernel::SetThreadExitCallout((ThreadExitCallout)pfExit_);
00233 }
00234 #endif // #if KERNEL_THREAD_EXIT_CALLOUT
00235 #if KERNEL_CONTEXT_SWITCH_CALLOUT
00236 //-----
00237 void Kernel_SetThreadContextSwitchCallout(thread_context_callout_t pfContext_)
00238 {
00239     Kernel::SetThreadContextSwitchCallout((
00240         ThreadContextCallout)pfContext_);
00241 }
00242 #endif // #if KERNEL_CONTEXT_SWITCH_CALLOUT
00243 #if KERNEL_THREAD_CREATE_CALLOUT
00244 thread_create_callout_t Kernel_GetThreadCreateCallout(void)
00245 {
00246     return (thread_create_callout_t)Kernel::GetThreadCreateCallout();
00247 }
00248 #endif // #if KERNEL_THREAD_CREATE_CALLOUT
00249 #if KERNEL_THREAD_EXIT_CALLOUT
00250 //-----
00251 thread_exit_callout_t Kernel_GetThreadExitCallout(void)
00252 {
00253     return (thread_exit_callout_t)Kernel::GetThreadExitCallout();
00254 }
00255 #endif // #if KERNEL_THREAD_EXIT_CALLOUT
00256 #if KERNEL_CONTEXT_SWITCH_CALLOUT
00257 //-----
00258 thread_context_callout_t Kernel_GetThreadContextSwitchCallout(void)
00259 {
00260     return (thread_context_callout_t)Kernel::GetThreadContextSwitchCallout
00261         ();
00262 }
00263 #endif // #if KERNEL_CONTEXT_SWITCH_CALLOUT
00264 //-----
00264 void Kernel_SetStackGuardThreshold(uint16_t ul6Threshold_)
00265 {

```

```

00266     Kernel::SetStackGuardThreshold(u16Threshold_);
00267 }
00268
00269 //-----
00270 uint16_t Kernel_GetStackGuardThreshold(void)
00271 {
00272     return Kernel::GetStackGuardThreshold();
00273 }
00274
00275 //-----
00276 void Kernel_SetDebugPrintFunction(
    kernel_debug_print_t pfPrintFunction_)
00277 {
00278     Kernel::SetDebugPrintFunction(pfPrintFunction_);
00279 }
00280
00281 //-----
00282 void Kernel_DebugPrint(const char* szString_)
00283 {
00284     Kernel::DebugPrint(szString_);
00285 }
00286
00287 //-----
00288 // Scheduler APIs
00289 void Scheduler_Enable(bool bEnable_)
00290 {
00291     Scheduler::SetScheduler(bEnable_);
00292 }
00293
00294 //-----
00295 bool Scheduler_IsEnabled(void)
00296 {
00297     return Scheduler::IsEnabled();
00298 }
00299
00300 //-----
00301 Thread_t Scheduler_GetCurrentThread(void)
00302 {
00303     return (Thread_t)Scheduler::GetCurrentThread();
00304 }
00305
00306 //-----
00307 // Thread APIs
00308 //-----
00309
00310 void Thread_Init(Thread_t      handle,
00311                  K_WORD*      pwStack_,
00312                  uint16_t      u16StackSize_,
00313                  PORT_PRIO_TYPE uXPriority_,
00314                  ThreadEntryFunc pfEntryPoint_,
00315                  void*          pvArg_)
00316 {
00317     Thread* pclThread = new ((void*)handle) Thread();
00318     pclThread->Init(pwStack_, u16StackSize_, uXPriority_, pfEntryPoint_, pvArg_);
00319 }
00320
00321 //-----
00322 void Thread_Start(Thread_t handle)
00323 {
00324     Thread* pclThread = (Thread*)handle;
00325     pclThread->Start();
00326 }
00327
00328 //-----
00329 void Thread_Stop(Thread_t handle)
00330 {
00331     Thread* pclThread = (Thread*)handle;
00332     pclThread->Stop();
00333 }
00334 //-----
00335 void Thread_SetName(Thread_t handle, const char* szName_)
00336 {
00337     Thread* pclThread = (Thread*)handle;
00338     pclThread->SetName(szName_);
00339 }
00340 //-----
00341 const char* Thread_GetName(Thread_t handle)
00342 {
00343     Thread* pclThread = (Thread*)handle;
00344     return pclThread->GetName();
00345 }
00346
00347 //-----
00348 PORT_PRIO_TYPE Thread_GetPriority(Thread_t handle)
00349 {
00350     Thread* pclThread = (Thread*)handle;
00351     return pclThread->GetPriority();

```

```

00352 }
00353 //-----
00354 PORT_PRIO_TYPE Thread_GetCurPriority(Thread_t handle)
00355 {
00356     Thread* pclThread = (Thread*)handle;
00357     return pclThread->GetCurPriority();
00358 }
00359
00360 //-----
00361 void Thread_SetQuantum(Thread_t handle, uint16_t u16Quantum_)
00362 {
00363     Thread* pclThread = (Thread*)handle;
00364     pclThread->SetQuantum(u16Quantum_);
00365 }
00366 //-----
00367 uint16_t Thread_GetQuantum(Thread_t handle)
00368 {
00369     Thread* pclThread = (Thread*)handle;
00370     return pclThread->GetQuantum();
00371 }
00372
00373 //-----
00374 void Thread_SetPriority(Thread_t handle,
00375     PORT_PRIO_TYPE uXPriority_)
00376 {
00377     Thread* pclThread = (Thread*)handle;
00378     pclThread->SetPriority(uXPriority_);
00379 }
00380 //-----
00381 void Thread_Exit(Thread_t handle)
00382 {
00383     Thread* pclThread = (Thread*)handle;
00384     pclThread->Exit();
00385 }
00386
00387 //-----
00388 void Thread_Sleep(uint32_t u32TimeMs_)
00389 {
00390     Thread::Sleep(u32TimeMs_);
00391 }
00392
00393 #if KERNEL_EXTENDED_CONTEXT
00394 //-----
00395 void* Thread_GetExtendedContext(Thread_t handle)
00396 {
00397     Thread* pclThread = (Thread*)handle;
00398     return pclThread->GetExtendedContext();
00399 }
00400
00401 //-----
00402 void Thread_SetExtendedContext(Thread_t handle, void* pvData_)
00403 {
00404     Thread* pclThread = (Thread*)handle;
00405     pclThread->SetExtendedContext(pvData_);
00406 }
00407 #endif // #if KERNEL_EXTENDED_CONTEXT
00408
00409 //-----
00410 void Thread_Yield(void)
00411 {
00412     Thread::Yield();
00413 }
00414 //-----
00415 void Thread_CoopYield(void)
00416 {
00417     Thread::CoopYield();
00418 }
00419 //-----
00420 void Thread_SetID(Thread_t handle, uint8_t u8ID_)
00421 {
00422     Thread* pclThread = (Thread*)handle;
00423     pclThread->SetID(u8ID_);
00424 }
00425 //-----
00426 uint8_t Thread_GetID(Thread_t handle)
00427 {
00428     Thread* pclThread = (Thread*)handle;
00429     return pclThread->GetID();
00430 }
00431 //-----
00432 uint16_t Thread_GetStackSlack(Thread_t handle)
00433 {
00434     Thread* pclThread = (Thread*)handle;
00435     return pclThread->GetStackSlack();
00436 }
00437 //-----

```

```

00438 thread_state_t Thread_GetState(Thread_t handle)
00439 {
00440     Thread* pclThread = (Thread*)handle;
00441     return (thread_state_t)pclThread->GetState();
00442 }
00443 //-----
00444 // Timer APIs
00445 //-----
00446
00447 //-----
00448 void Timer_Init(Timer_t handle)
00449 {
00450     Timer* pclTimer = new ((void*)handle) Timer();
00451     pclTimer->Init();
00452 }
00453
00454 //-----
00455 void Timer_Start(Timer_t handle, bool bRepeat_, uint32_t u32IntervalMs_,
00456                 timer_callback_t pfCallback_, void* pvData_)
00457 {
00458     Timer* pclTimer = (Timer*)handle;
00459     pclTimer->Start(bRepeat_, u32IntervalMs_, (TimerCallback)pfCallback_, pvData_);
00460 }
00461 //-----
00462 void Timer_Stop(Timer_t handle)
00463 {
00464     Timer* pclTimer = (Timer*)handle;
00465     pclTimer->Stop();
00466 }
00467
00468 //-----
00469 void Timer_Restart(Timer_t handle)
00470 {
00471     Timer* pclTimer = (Timer*)handle;
00472     pclTimer->Start();
00473 }
00474
00475 //-----
00476 // Semaphore APIs
00477 //-----
00478 void Semaphore_Init(Semaphore_t handle, uint16_t u16InitVal_, uint16_t u16MaxVal_)
00479 {
00480     Semaphore* pclSemaphore = new ((void*)handle) Semaphore();
00481     pclSemaphore->Init(u16InitVal_, u16MaxVal_);
00482 }
00483
00484 //-----
00485 void Semaphore_Post(Semaphore_t handle)
00486 {
00487     Semaphore* pclSemaphore = (Semaphore*)handle;
00488     pclSemaphore->Post();
00489 }
00490
00491 //-----
00492 void Semaphore_Pend(Semaphore_t handle)
00493 {
00494     Semaphore* pclSemaphore = (Semaphore*)handle;
00495     pclSemaphore->Pend();
00496 }
00497
00498 //-----
00499 bool Semaphore_TimedPend(Semaphore_t handle, uint32_t u32WaitTimeMS_)
00500 {
00501     Semaphore* pclSemaphore = (Semaphore*)handle;
00502     return pclSemaphore->Pend(u32WaitTimeMS_);
00503 }
00504
00505 //-----
00506 // Mutex APIs
00507 //-----
00508 void Mutex_Init(Mutex_t handle)
00509 {
00510     Mutex* pclMutex = new ((void*)handle) Mutex();
00511     pclMutex->Init();
00512 }
00513
00514 //-----
00515 void Mutex_Claim(Mutex_t handle)
00516 {
00517     Mutex* pclMutex = (Mutex*)handle;
00518     pclMutex->Claim();
00519 }
00520
00521 //-----
00522 void Mutex_Release(Mutex_t handle)
00523 {

```

```

00524     Mutex* pclMutex = (Mutex*)handle;
00525     pclMutex->Release();
00526 }
00527
00528 //-----
00529 bool Mutex_TimedClaim(Mutex_t handle, uint32_t u32WaitTimeMS_)
00530 {
00531     Mutex* pclMutex = (Mutex*)handle;
00532     return pclMutex->Claim(u32WaitTimeMS_);
00533 }
00534
00535 #if KERNEL_EVENT_FLAGS
00536 //-----
00537 // EventFlag APIs
00538 //-----
00539 void EventFlag_Init(EventFlag_t handle)
00540 {
00541     EventFlag* pclFlag = new ((void*)handle) EventFlag();
00542     pclFlag->Init();
00543 }
00544
00545 //-----
00546 uint16_t EventFlag_Wait(EventFlag_t handle, uint16_t ul6Mask_, event_flag_operation_t eMode_)
00547 {
00548     EventFlag* pclFlag = (EventFlag*)handle;
00549     return pclFlag->Wait(ul6Mask_, (EventFlagOperation)eMode_);
00550 }
00551
00552 //-----
00553 uint16_t EventFlag_TimedWait(EventFlag_t handle, uint16_t ul6Mask_, event_flag_operation_t eMode_, uint32_t
    u32TimeMS_)
00554 {
00555     EventFlag* pclFlag = (EventFlag*)handle;
00556     return pclFlag->Wait(ul6Mask_, (EventFlagOperation)eMode_, u32TimeMS_);
00557 }
00558
00559 //-----
00560 void EventFlag_Set(EventFlag_t handle, uint16_t ul6Mask_)
00561 {
00562     EventFlag* pclFlag = (EventFlag*)handle;
00563     pclFlag->Set(ul6Mask_);
00564 }
00565
00566 //-----
00567 void EventFlag_Clear(EventFlag_t handle, uint16_t ul6Mask_)
00568 {
00569     EventFlag* pclFlag = (EventFlag*)handle;
00570     pclFlag->Clear(ul6Mask_);
00571 }
00572
00573 //-----
00574 uint16_t EventFlag_GetMask(EventFlag_t handle)
00575 {
00576     EventFlag* pclFlag = (EventFlag*)handle;
00577     return pclFlag->GetMask();
00578 }
00579 #endif // #if #if KERNEL_EVENT_FLAGS
00580
00581 //-----
00582 // Notification APIs
00583 //-----
00584 void Notify_Init(Notify_t handle)
00585 {
00586     Notify* pclNotify = new ((void*)handle) Notify();
00587     pclNotify->Init();
00588 }
00589
00590 //-----
00591 void Notify_Signal(Notify_t handle)
00592 {
00593     Notify* pclNotify = (Notify*)handle;
00594     pclNotify->Signal();
00595 }
00596
00597 //-----
00598 void Notify_Wait(Notify_t handle, bool* pbFlag_)
00599 {
00600     Notify* pclNotify = (Notify*)handle;
00601     pclNotify->Wait(pbFlag_);
00602 }
00603
00604 //-----
00605 bool Notify_TimedWait(Notify_t handle, uint32_t u32WaitTimeMS_, bool* pbFlag_)
00606 {
00607     Notify* pclNotify = (Notify*)handle;
00608     return pclNotify->Wait(u32WaitTimeMS_, pbFlag_);
00609 }

```

```

00610
00611 //-----
00612 // Atomic Functions
00613 //-----
00614 uint8_t Atomic_Set8(uint8_t* pu8Source_, uint8_t u8Val_)
00615 {
00616     return Atomic::Set(pu8Source_, u8Val_);
00617 }
00618
00619 //-----
00620 uint16_t Atomic_Set16(uint16_t* pu16Source_, uint16_t u16Val_)
00621 {
00622     return Atomic::Set(pu16Source_, u16Val_);
00623 }
00624
00625 //-----
00626 uint32_t Atomic_Set32(uint32_t* pu32Source_, uint32_t u32Val_)
00627 {
00628     return Atomic::Set(pu32Source_, u32Val_);
00629 }
00630
00631 //-----
00632 uint8_t Atomic_Add8(uint8_t* pu8Source_, uint8_t u8Val_)
00633 {
00634     return Atomic::Add(pu8Source_, u8Val_);
00635 }
00636
00637 //-----
00638 uint16_t Atomic_Add16(uint16_t* pu16Source_, uint16_t u16Val_)
00639 {
00640     return Atomic::Add(pu16Source_, u16Val_);
00641 }
00642
00643 //-----
00644 uint32_t Atomic_Add32(uint32_t* pu32Source_, uint32_t u32Val_)
00645 {
00646     return Atomic::Add(pu32Source_, u32Val_);
00647 }
00648
00649 //-----
00650 uint8_t Atomic_Sub8(uint8_t* pu8Source_, uint8_t u8Val_)
00651 {
00652     return Atomic::Sub(pu8Source_, u8Val_);
00653 }
00654
00655 //-----
00656 uint16_t Atomic_Sub16(uint16_t* pu16Source_, uint16_t u16Val_)
00657 {
00658     return Atomic::Sub(pu16Source_, u16Val_);
00659 }
00660
00661 //-----
00662 uint32_t Atomic_Sub32(uint32_t* pu32Source_, uint32_t u32Val_)
00663 {
00664     return Atomic::Sub(pu32Source_, u32Val_);
00665 }
00666
00667 //-----
00668 bool Atomic_TestAndSet(bool* pbLock)
00669 {
00670     return Atomic::TestAndSet(pbLock);
00671 }
00672
00673 //-----
00674 // Message/Message Queue APIs
00675 //-----
00676 void Message_Init(Message_t handle)
00677 {
00678     Message* pclMessage = new ((void*)handle) Message();
00679     return pclMessage->Init();
00680 }
00681
00682 //-----
00683 void Message_SetData(Message_t handle, void* pvData_)
00684 {
00685     Message* pclMessage = (Message*)handle;
00686     pclMessage->SetData(pvData_);
00687 }
00688
00689 //-----
00690 void* Message_GetData(Message_t handle)
00691 {
00692     Message* pclMessage = (Message*)handle;
00693     return pclMessage->GetData();
00694 }
00695
00696 //-----

```

```

00697 void Message_SetCode(Message_t handle, uint16_t ul6Code_)
00698 {
00699     Message* pclMessage = (Message*)handle;
00700     pclMessage->SetCode(ul6Code_);
00701 }
00702
00703 //-----
00704 uint16_t Message_GetCode(Message_t handle)
00705 {
00706     Message* pclMessage = (Message*)handle;
00707     return pclMessage->GetCode();
00708 }
00709
00710 //-----
00711 void MessageQueue_Init(MessageQueue_t handle)
00712 {
00713     MessageQueue* pclMsgQ = new ((void*)handle) MessageQueue();
00714     pclMsgQ->Init();
00715 }
00716
00717 //-----
00718 Message_t MessageQueue_Receive(MessageQueue_t handle)
00719 {
00720     MessageQueue* pclMsgQ = (MessageQueue*)handle;
00721     return pclMsgQ->Receive();
00722 }
00723
00724 //-----
00725 void MessagePool_Init(MessagePool_t handle)
00726 {
00727     MessagePool* pclMsgPool = new ((void*)handle) MessagePool();
00728     pclMsgPool->Init();
00729 }
00730
00731 //-----
00732 void MessagePool_Push(MessagePool_t handle,
00733     Message_t msg)
00734 {
00735     MessagePool* pclMsgPool = (MessagePool*)handle;
00736     pclMsgPool->Push((Message*)msg);
00737 }
00738
00739 //-----
00739 Message_t MessagePool_Pop(MessagePool_t handle)
00740 {
00741     MessagePool* pclMsgPool = (MessagePool*)handle;
00742     return (Message_t)pclMsgPool->Pop();
00743 }
00744
00745 //-----
00746 Message_t MessageQueue_TimedReceive(
00747     MessageQueue_t handle, uint32_t u32TimeWaitMS_)
00748 {
00749     MessageQueue* pclMsgQ = (MessageQueue*)handle;
00750     return (Message_t)pclMsgQ->Receive(u32TimeWaitMS_);
00751 }
00752
00753 //-----
00753 void MessageQueue_Send(MessageQueue_t handle,
00754     Message_t hMessage_)
00755 {
00756     MessageQueue* pclMsgQ = (MessageQueue*)handle;
00757     pclMsgQ->Send((Message*)hMessage_);
00758 }
00759
00759 //-----
00760 uint16_t MessageQueue_GetCount(MessageQueue_t handle)
00761 {
00762     MessageQueue* pclMsgQ = (MessageQueue*)handle;
00763     return pclMsgQ->GetCount();
00764 }
00765
00766 //-----
00767 // Mailbox APIs
00768 //-----
00769 void Mailbox_Init(Mailbox_t handle, void* pvBuffer_, uint16_t ul6BufferSize_, uint16_t
00770     ul6ElementSize_)
00771 {
00772     Mailbox* pclMBox = new ((void*)handle) Mailbox();
00773     pclMBox->Init(pvBuffer_, ul6BufferSize_, ul6ElementSize_);
00774 }
00775
00775 //-----
00776 bool Mailbox_Send(Mailbox_t handle, void* pvData_)
00777 {
00778     Mailbox* pclMBox = (Mailbox*)handle;
00779     return pclMBox->Send(pvData_);

```

```

00780 }
00781
00782 //-----
00783 bool Mailbox_SendTail(Mailbox_t handle, void* pvData_)
00784 {
00785     Mailbox* pclMBox = (Mailbox*)handle;
00786     return pclMBox->SendTail(pvData_);
00787 }
00788
00789 //-----
00790 bool Mailbox_TimedSend(Mailbox_t handle, void* pvData_, uint32_t u32TimeoutMS_)
00791 {
00792     Mailbox* pclMBox = (Mailbox*)handle;
00793     return pclMBox->Send(pvData_, u32TimeoutMS_);
00794 }
00795
00796 //-----
00797 bool Mailbox_TimedSendTail(Mailbox_t handle, void* pvData_, uint32_t
u32TimeoutMS_)
00798 {
00799     Mailbox* pclMBox = (Mailbox*)handle;
00800     return pclMBox->SendTail(pvData_, u32TimeoutMS_);
00801 }
00802
00803 //-----
00804 void Mailbox_Receive(Mailbox_t handle, void* pvData_)
00805 {
00806     Mailbox* pclMBox = (Mailbox*)handle;
00807     pclMBox->Receive(pvData_);
00808 }
00809
00810 //-----
00811 void Mailbox_ReceiveTail(Mailbox_t handle, void* pvData_)
00812 {
00813     Mailbox* pclMBox = (Mailbox*)handle;
00814     pclMBox->ReceiveTail(pvData_);
00815 }
00816
00817 //-----
00818 bool Mailbox_TimedReceive(Mailbox_t handle, void* pvData_, uint32_t
u32TimeoutMS_)
00819 {
00820     Mailbox* pclMBox = (Mailbox*)handle;
00821     return pclMBox->Receive(pvData_, u32TimeoutMS_);
00822 }
00823
00824 //-----
00825 bool Mailbox_TimedReceiveTail(Mailbox_t handle, void* pvData_, uint32_t
u32TimeoutMS_)
00826 {
00827     Mailbox* pclMBox = (Mailbox*)handle;
00828     return pclMBox->ReceiveTail(pvData_, u32TimeoutMS_);
00829 }
00830
00831 //-----
00832 uint16_t Mailbox_GetFreeSlots(Mailbox_t handle)
00833 {
00834     Mailbox* pclMBox = (Mailbox*)handle;
00835     return pclMBox->GetFreeSlots();
00836 }
00837
00838 //-----
00839 bool Mailbox_IsFull(Mailbox_t handle)
00840 {
00841     Mailbox* pclMBox = (Mailbox*)handle;
00842     return pclMBox->IsFull();
00843 }
00844
00845 //-----
00846 bool Mailbox_IsEmpty(Mailbox_t handle)
00847 {
00848     Mailbox* pclMBox = (Mailbox*)handle;
00849     return pclMBox->IsEmpty();
00850 }
00851
00852 //-----
00853 // Condition Variables
00854 //-----
00855 void ConditionVariable_Init(ConditionVariable_t handle)
00856 {
00857     ConditionVariable* pclCondvar = new ((void*)handle)
ConditionVariable();
00858     pclCondvar->Init();
00859 }
00860 //-----
00861 void ConditionVariable_Wait(ConditionVariable_t handle,
Mutex_t hMutex_)

```



```

00862 {
00863     ConditionVariable* pclCondvar = (ConditionVariable*)handle;
00864     Mutex*          pclMutex    = (Mutex*)hMutex_;
00865     pclCondvar->Wait(pclMutex);
00866 }
00867 //-----
00868 void ConditionVariable_Signal(ConditionVariable_t handle)
00869 {
00870     ConditionVariable* pclCondvar = (ConditionVariable*)handle;
00871     pclCondvar->Signal();
00872 }
00873 //-----
00874 void ConditionVariable_Broadcast(ConditionVariable_t handle)
00875 {
00876     ConditionVariable* pclCondvar = (ConditionVariable*)handle;
00877     pclCondvar->Broadcast();
00878 }
00879 //-----
00880 bool ConditionVariable_TimedWait(ConditionVariable_t handle,
00881                                 Mutex_t hMutex_, uint32_t u32WaitTimeMS_)
00882 {
00883     ConditionVariable* pclCondvar = (ConditionVariable*)handle;
00884     Mutex*          pclMutex    = (Mutex*)hMutex_;
00885     return pclCondvar->Wait(pclMutex, u32WaitTimeMS_);
00886 }
00887 //-----
00888 // Reader-writer locks
00889 //-----
00890 void ReaderWriterLock_Init(ReaderWriterLock_t handle)
00891 {
00892     ReaderWriterLock* pclReaderWriter = new ((void*)handle)
00893     ReaderWriterLock();
00894     pclReaderWriter->Init();
00895 }
00896 //-----
00897 void ReaderWriterLock_AcquireReader(
00898     ReaderWriterLock_t handle)
00899 {
00900     ReaderWriterLock* pclReaderWriter = (ReaderWriterLock*)handle;
00901     pclReaderWriter->AcquireReader();
00902 }
00903 //-----
00904 void ReaderWriterLock_ReleaseReader(
00905     ReaderWriterLock_t handle)
00906 {
00907     ReaderWriterLock* pclReaderWriter = (ReaderWriterLock*)handle;
00908     pclReaderWriter->ReleaseReader();
00909 }
00910 //-----
00911 void ReaderWriterLock_AcquireWriter(
00912     ReaderWriterLock_t handle)
00913 {
00914     ReaderWriterLock* pclReaderWriter = (ReaderWriterLock*)handle;
00915     pclReaderWriter->AcquireWriter();
00916 }
00917 //-----
00918 void ReaderWriterLock_ReleaseWriter(
00919     ReaderWriterLock_t handle)
00920 {
00921     ReaderWriterLock* pclReaderWriter = (ReaderWriterLock*)handle;
00922     pclReaderWriter->ReleaseWriter();
00923 }
00924 //-----
00925 bool ReaderWriterLock_TimedAcquireWriter(
00926     ReaderWriterLock_t handle, uint32_t u32TimeoutMs_)
00927 {
00928     ReaderWriterLock* pclReaderWriter = (ReaderWriterLock*)handle;
00929     return pclReaderWriter->AcquireWriter(u32TimeoutMs_);
00930 }
00931 //-----
00932 bool ReaderWriterLock_TimedAcquireReader(
00933     ReaderWriterLock_t handle, uint32_t u32TimeoutMs_)
00934 {
00935     ReaderWriterLock* pclReaderWriter = (ReaderWriterLock*)handle;
00936     return pclReaderWriter->AcquireReader(u32TimeoutMs_);
00937 }

```

20.3 /home/moslevin/projects/m3-repo/kernel/lib/mark3c/src/public/fake_types.h File Reference

C-struct definitions that mirror.

```
#include <stdint.h>
#include <stddef.h>
#include <stdbool.h>
#include "mark3cfg.h"
```

20.3.1 Detailed Description

C-struct definitions that mirror.

This header contains a set of "fake" structures that have the same memory layout as the kernel objects in C++ (taking into account inheritance, etc.). These are used for sizing the opaque data blobs that are declared in C, which then become instantiated as C++ kernel objects via the bindings provided.

Definition in file [fake_types.h](#).

20.4 fake_types.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00026 #include <stdint.h>
00027 #include <stddef.h>
00028 #include <stdbool.h>
00029 #include "mark3cfg.h"
00030
00031 #pragma once
00032 #if defined(__cplusplus)
00034 extern "C" {
00035 #endif
00036
00037 //-----
00038 typedef struct {
00039     void* prev;
00040     void* next;
00041 } Fake_LinkedListNode;
00042
00043 //-----
00044 typedef struct {
00045     void* head;
00046     void* tail;
00047 } Fake_LinkedList;
00048
00049 //-----
00050 typedef struct {
00051     Fake_LinkedListNode fake_node;
00052     Fake_LinkedList fake_list;
00053     PORT_PRIO_TYPE m_uXPriority;
00054     void* m_pclMap;
00055 } Fake_ThreadList;
00056
00057 //-----
```

```

00058 typedef struct {
00059     Fake_LinkedListNode m_ll_node;
00060     uint8_t             m_u8Initialized;
00061     uint8_t             m_u8Flags;
00062     void*               m_pfCallback;
00063     uint32_t            m_u32Interval;
00064     uint32_t            m_u32TimeLeft;
00065     void*               m_pclOwner;
00066     void*               m_pvData;
00067 } Fake_Timer;
00068
00069 //-----
00070 typedef struct {
00071     Fake_LinkedListNode m_ll_node;
00072     K_WORD*             m_pwStackTop;
00073     K_WORD*             m_pwStack;
00074     uint8_t             m_u8ThreadID;
00075     PORT_PRIO_TYPE      m_uXPriority;
00076     PORT_PRIO_TYPE      m_uXCurPriority;
00077     uint8_t             m_eState;
00078 #if KERNEL_EXTENDED_CONTEXT
00079     void* m_pvExtendedContext;
00080 #endif // #if KERNEL_EXTENDED_CONTEXT
00081 #if KERNEL_NAMED_THREADS
00082     const char* m_szName;
00083 #endif // #if KERNEL_NAMED_THREADS
00084     uint16_t m_ul6StackSize;
00085     void* m_pclCurrent;
00086     void* m_pclOwner;
00087     void* m_pfEntryPoint;
00088     void* m_pvArg;
00089 #if KERNEL_ROUND_ROBIN
00090     uint16_t m_ul6Quantum;
00091 #endif // #if KERNEL_ROUND_ROBIN
00092 #if KERNEL_EVENT_FLAGS
00093     uint16_t m_ul6FlagMask;
00094     uint8_t m_eFlagMode;
00095 #endif // #if KERNEL_EVENT_FLAGS
00096     Fake_Timer m_clTimer;
00097     bool m_bExpired;
00098 } Fake_Thread;
00099
00100 //-----
00101 typedef struct {
00102     Fake_ThreadList thread_list;
00103     uint8_t m_u8Initialized;
00104     uint16_t m_ul6Value;
00105     uint16_t m_ul6MaxValue;
00106 } Fake_Semaphore;
00107
00108 //-----
00109 typedef struct {
00110     Fake_ThreadList thread_list;
00111     uint8_t m_u8Initialized;
00112     uint8_t m_u8Recurse;
00113     bool m_bReady;
00114     bool m_bRecursive;
00115     uint8_t m_u8MaxPri;
00116     void* m_pclOwner;
00117 } Fake_Mutex;
00118
00119 //-----
00120 typedef struct {
00121     Fake_LinkedListNode list_node;
00122     void* m_pvData;
00123     uint16_t m_ul6Code;
00124 } Fake_Message;
00125
00126 //-----
00127 typedef struct {
00128     Fake_Semaphore m_clSemaphore;
00129     Fake_LinkedList m_clLinkList;
00130 } Fake_MessageQueue;
00131
00132 //-----
00133 typedef struct {
00134     Fake_LinkedList m_clList;
00135 } Fake_MessagePool;
00136
00137 //-----
00138 typedef struct {
00139     uint16_t m_ul6Head;
00140     uint16_t m_ul6Tail;
00141     uint16_t m_ul6Count;
00142     uint16_t m_ul6Free;
00143     uint16_t m_ul6ElementSize;
00144     void* m_pvBuffer;

```

```

00145     Fake_Semaphore m_clRecvSem;
00146     Fake_Semaphore m_clSendSem;
00147 } Fake_Mailbox;
00148
00149 //-----
00150 typedef struct {
00151     Fake_ThreadList thread_list;
00152     uint8_t          m_u8Initialized;
00153     bool             m_bPending;
00154 } Fake_Notify;
00155
00156 //-----
00157 typedef struct {
00158     Fake_ThreadList thread_list;
00159     uint8_t          m_u8Initialized;
00160     uint16_t         m_u16EventFlag;
00161 } Fake_EventFlag;
00162
00163 //-----
00164 typedef struct {
00165     Fake_Mutex m_clGlobalMutex;
00166     Fake_Mutex m_clReaderMutex;
00167     uint8_t    m_u8ReadCount;
00168 } Fake_ReaderWriterLock;
00169
00170 //-----
00171 typedef struct {
00172     Fake_Mutex    m_clMutex;
00173     Fake_Semaphore m_clSemaphore;
00174     uint8_t       m_u8Waiters;
00175 } Fake_ConditionVariable;
00176
00177 #if defined(__cplusplus)
00178 }
00179 #endif
00180

```

20.5 /home/moslevin/projects/m3-repo/kernel/lib/mark3c/src/public/mark3c.h File Reference

Header providing C-language API bindings for the [Mark3](#) kernel.

```

#include "mark3cfg.h"
#include "fake_types.h"
#include <stdint.h>
#include <stdbool.h>

```

Macros

- `#define THREAD_SIZE (sizeof(Fake_Thread))`
- `#define TIMER_SIZE (sizeof(Fake_Timer))`
- `#define SEMAPHORE_SIZE (sizeof(Fake_Semaphore))`
- `#define MUTEX_SIZE (sizeof(Fake_Mutex))`
- `#define MESSAGE_SIZE (sizeof(Fake_Message))`
- `#define MESSAGEQUEUE_SIZE (sizeof(Fake_MessageQueue))`
- `#define MAILBOX_SIZE (sizeof(Fake_Mailbox))`
- `#define NOTIFY_SIZE (sizeof(Fake_Notify))`
- `#define MESSAGEPOOL_SIZE (sizeof(Fake_MessagePool))`
- `#define CONDITIONVARIABLE_SIZE (sizeof(Fake_ConditionVariable))`
- `#define READERWRITERLOCK_SIZE (sizeof(Fake_ReaderWriterLock))`
- `#define TOKEN_1(x, y) x##y`
- `#define TOKEN_2(x, y) TOKEN_1(x, y)`
- `#define WORD_ROUND(x) (((x) + (sizeof(K_WORD) - 1)) / sizeof(K_WORD))`

- #define [DECLARE_THREAD](#)(name)
- #define [DECLARE_TIMER](#)(name)
- #define [DECLARE_SEMAPHORE](#)(name)
- #define [DECLARE_MUTEX](#)(name)
- #define [DECLARE_MESSAGE](#)(name)
- #define [DECLARE_MESSAGEPOOL](#)(name)
- #define [DECLARE_MESSAGEQUEUE](#)(name)
- #define [DECLARE_MAILBOX](#)(name)
- #define [DECLARE_NOTIFY](#)(name)
- #define [DECLARE_CONDITIONVARIABLE](#)(name)
- #define [DECLARE_READERWRITERLOCK](#)(name)

Typedefs

- typedef void * [Mailbox_t](#)
Mailbox opaque handle data type.
- typedef void * [Message_t](#)
Message opaque handle data type.
- typedef void * [MessagePool_t](#)
MessagePool opaque handle data type.
- typedef void * [MessageQueue_t](#)
MessageQueue opaque handle data type.
- typedef void * [Mutex_t](#)
Mutex opaque handle data type.
- typedef void * [Notify_t](#)
Notification object opaque handle data type.
- typedef void * [Semaphore_t](#)
Semaphore opaque handle data type.
- typedef void * [Thread_t](#)
Thread opaque handle data type.
- typedef void * [Timer_t](#)
Timer opaque handle data type.
- typedef void * [ConditionVariable_t](#)
Condition Variable opaque handle data type.
- typedef void * [ReaderWriterLock_t](#)
Reader-writer-lock opaque handle data type.
- typedef void(* [kernel_debug_print_t](#)) (const char *szString_)
- typedef void(* [panic_func_t](#)) (uint16_t u16PanicCode_)
- typedef void(* [thread_entry_func_t](#)) (void *pvArg_)
- typedef void(* [timer_callback_t](#)) ([Thread_t](#) hOwner_, void *pvData_)

Enumerations

- enum [thread_state_t](#) {
[THREAD_STATE_EXIT](#) = 0, [THREAD_STATE_READY](#), [THREAD_STATE_BLOCKED](#), [THREAD_STATE_STOP](#),
[THREAD_STATE_INVALID](#) }

Functions

- void * [Alloc_Memory](#) (size_t eSize_)
Alloc_Memory.
- void [Free_Memory](#) (void *pvObject_)
Free_Memory.
- [Semaphore_t Alloc_Semaphore](#) (void)
Alloc_Semaphore.
- void [Free_Semaphore](#) ([Semaphore_t](#) handle)
- [Mutex_t Alloc_Mutex](#) (void)
Alloc_Mutex.
- void [Free_Mutex](#) ([Mutex_t](#) handle)
- [Message_t Alloc_Message](#) (void)
Alloc_Message.
- void [Free_Message](#) ([Message_t](#) handle)
- [MessageQueue_t Alloc_MessageQueue](#) (void)
Alloc_MessageQueue.
- void [Free_MessageQueue](#) ([MessageQueue_t](#) handle)
- [MessagePool_t Alloc_MessagePool](#) (void)
- void [Free_MessagePool](#) ([MessagePool_t](#) handle)
- [Notify_t Alloc_Notify](#) (void)
Alloc_Notify.
- void [Free_Notify](#) ([Notify_t](#) handle)
- [Mailbox_t Alloc_Mailbox](#) (void)
Alloc_Mailbox.
- void [Free_Mailbox](#) ([Mailbox_t](#) handle)
- [Thread_t Alloc_Thread](#) (void)
Alloc_Thread.
- void [Free_Thread](#) ([Thread_t](#) handle)
- [Timer_t Alloc_Timer](#) (void)
Alloc_Timer.
- void [Free_Timer](#) ([Timer_t](#) handle)
- void [Kernel_Init](#) (void)
Kernel_Init.
- void [Kernel_Start](#) (void)
Kernel_Start.
- bool [Kernel_IsStarted](#) (void)
Kernel_IsStarted.
- void [Kernel_SetPanic](#) ([panic_func_t](#) pfPanic_)
Kernel_SetPanic.
- bool [Kernel_IsPanic](#) (void)
Kernel_IsPanic.
- void [Kernel_Panic](#) (uint16_t u16Cause_)
Kernel_Panic.
- uint32_t [Kernel_GetTicks](#) (void)
Kernel_GetTicks.
- void [Scheduler_Enable](#) (bool bEnable_)
Scheduler_Enable.
- bool [Scheduler_IsEnabled](#) (void)
Scheduler_IsEnabled.
- [Thread_t Scheduler_GetCurrentThread](#) (void)
Scheduler_GetCurrentThread.

- void [Thread_Init](#) ([Thread_t](#) handle, [K_WORD](#) *pwStack_, [uint16_t](#) u16StackSize_, [PORT_PRIO_TYPE](#) uXPriority_, [thread_entry_func_t](#) pfEntryPoint_, void *pvArg_)
Thread_Init.
- void [Thread_Start](#) ([Thread_t](#) handle)
Thread_Start.
- void [Thread_Stop](#) ([Thread_t](#) handle)
Thread_Stop.
- [PORT_PRIO_TYPE](#) [Thread_GetPriority](#) ([Thread_t](#) handle)
Thread_GetPriority.
- [PORT_PRIO_TYPE](#) [Thread_GetCurPriority](#) ([Thread_t](#) handle)
Thread_GetCurPriority.
- void [Thread_SetPriority](#) ([Thread_t](#) handle, [PORT_PRIO_TYPE](#) uXPriority_)
Thread_SetPriority.
- void [Thread_Exit](#) ([Thread_t](#) handle)
Thread_Exit.
- void [Thread_Sleep](#) ([uint32_t](#) u32TimeMs_)
Thread_Sleep.
- void [Thread_Yield](#) (void)
Thread_Yield.
- void [Thread_CoopYield](#) (void)
Thread_CoopYield.
- void [Thread_SetID](#) ([Thread_t](#) handle, [uint8_t](#) u8ID_)
Thread_SetID.
- [uint8_t](#) [Thread_GetID](#) ([Thread_t](#) handle)
Thread_GetID.
- [thread_state_t](#) [Thread_GetState](#) ([Thread_t](#) handle)
Thread_GetState.
- void [Timer_Init](#) ([Timer_t](#) handle)
Timer_Init.
- void [Timer_Start](#) ([Timer_t](#) handle, bool bRepeat_, [uint32_t](#) u32IntervalMs_, [timer_callback_t](#) pfCallback_, void *pvData_)
Timer_Start.
- void [Timer_Restart](#) ([Timer_t](#) handle)
Timer_Restart.
- void [Timer_Stop](#) ([Timer_t](#) handle)
Timer_Stop.
- void [Semaphore_Init](#) ([Semaphore_t](#) handle, [uint16_t](#) u16InitVal_, [uint16_t](#) u16MaxVal_)
Semaphore_Init.
- void [Semaphore_Post](#) ([Semaphore_t](#) handle)
Semaphore_Post.
- void [Semaphore_Pend](#) ([Semaphore_t](#) handle)
Semaphore_Pend.
- bool [Semaphore_TimedPend](#) ([Semaphore_t](#) handle, [uint32_t](#) u32WaitTimeMS_)
Semaphore_TimedPend.
- void [Mutex_Init](#) ([Mutex_t](#) handle)
Mutex_Init.
- void [Mutex_Claim](#) ([Mutex_t](#) handle)
Mutex_Claim.
- void [Mutex_Release](#) ([Mutex_t](#) handle)
Mutex_Release.
- bool [Mutex_TimedClaim](#) ([Mutex_t](#) handle, [uint32_t](#) u32WaitTimeMS_)

- Mutex_TimedClaim.*
- void [Notify_Init](#) ([Notify_t](#) handle)
 - Notify_Init.*
- void [Notify_Signal](#) ([Notify_t](#) handle)
 - Notify_Signal.*
- void [Notify_Wait](#) ([Notify_t](#) handle, bool *pbFlag_)
 - Notify_Wait.*
- bool [Notify_TimedWait](#) ([Notify_t](#) handle, uint32_t u32WaitTimeMS_, bool *pbFlag_)
 - Notify_TimedWait.*
- uint8_t [Atomic_Set8](#) (uint8_t *pu8Source_, uint8_t u8Val_)
 - Atomic_Set8.*
- uint16_t [Atomic_Set16](#) (uint16_t *pu16Source_, uint16_t u16Val_)
 - Atomic_Set16.*
- uint32_t [Atomic_Set32](#) (uint32_t *pu32Source_, uint32_t u32Val_)
 - Atomic_Set32.*
- uint8_t [Atomic_Add8](#) (uint8_t *pu8Source_, uint8_t u8Val_)
 - Atomic_Add8.*
- uint16_t [Atomic_Add16](#) (uint16_t *pu16Source_, uint16_t u16Val_)
 - Atomic_Add16.*
- uint32_t [Atomic_Add32](#) (uint32_t *pu32Source_, uint32_t u32Val_)
 - Atomic_Add32.*
- uint8_t [Atomic_Sub8](#) (uint8_t *pu8Source_, uint8_t u8Val_)
 - Atomic_Sub8.*
- uint16_t [Atomic_Sub16](#) (uint16_t *pu16Source_, uint16_t u16Val_)
 - Atomic_Sub16.*
- uint32_t [Atomic_Sub32](#) (uint32_t *pu32Source_, uint32_t u32Val_)
 - Atomic_Sub32.*
- bool [Atomic_TestAndSet](#) (bool *pbLock)
 - Atomic_TestAndSet.*
- void [Message_Init](#) ([Message_t](#) handle)
 - Message_Init.*
- void [Message_SetData](#) ([Message_t](#) handle, void *pvData_)
 - Message_SetData.*
- void * [Message_GetData](#) ([Message_t](#) handle)
 - Message_GetData.*
- void [Message_SetCode](#) ([Message_t](#) handle, uint16_t u16Code_)
 - Message_SetCode.*
- uint16_t [Message_GetCode](#) ([Message_t](#) handle)
 - Message_GetCode.*
- void [MessageQueue_Init](#) ([MessageQueue_t](#) handle)
 - MessageQueue_Init.*
- [Message_t](#) [MessageQueue_Receive](#) ([MessageQueue_t](#) handle)
 - MessageQueue_Receive.*
- [Message_t](#) [MessageQueue_TimedReceive](#) ([MessageQueue_t](#) handle, uint32_t u32TimeWaitMS_)
 - MessageQueue_TimedReceive.*
- void [MessageQueue_Send](#) ([MessageQueue_t](#) handle, [Message_t](#) hMessage_)
 - MessageQueue_Send.*
- uint16_t [MessageQueue_GetCount](#) ([MessageQueue_t](#) handle)
 - MessageQueue_GetCount.*
- void [MessagePool_Init](#) ([MessagePool_t](#) handle)
 - MessagePool_Init.*

- void [MessagePool_Push](#) ([MessagePool_t](#) handle, [Message_t](#) msg)
MessagePool_Push.
- [Message_t](#) [MessagePool_Pop](#) ([MessagePool_t](#) handle)
MessagePool_Pop.
- void [Mailbox_Init](#) ([Mailbox_t](#) handle, void *pvBuffer_, uint16_t u16BufferSize_, uint16_t u16ElementSize_)
Mailbox_Init.
- bool [Mailbox_Send](#) ([Mailbox_t](#) handle, void *pvData_)
Mailbox_Send.
- bool [Mailbox_SendTail](#) ([Mailbox_t](#) handle, void *pvData_)
Mailbox_SendTail.
- bool [Mailbox_TimedSend](#) ([Mailbox_t](#) handle, void *pvData_, uint32_t u32TimeoutMS_)
Mailbox_TimedSend.
- bool [Mailbox_TimedSendTail](#) ([Mailbox_t](#) handle, void *pvData_, uint32_t u32TimeoutMS_)
Mailbox_TimedSendTail.
- void [Mailbox_Receive](#) ([Mailbox_t](#) handle, void *pvData_)
Mailbox_Receive.
- void [Mailbox_ReceiveTail](#) ([Mailbox_t](#) handle, void *pvData_)
Mailbox_ReceiveTail.
- bool [Mailbox_TimedReceive](#) ([Mailbox_t](#) handle, void *pvData_, uint32_t u32TimeoutMS_)
Mailbox_TimedReceive.
- bool [Mailbox_TimedReceiveTail](#) ([Mailbox_t](#) handle, void *pvData_, uint32_t u32TimeoutMS_)
Mailbox_TimedReceiveTail.
- uint16_t [Mailbox_GetFreeSlots](#) ([Mailbox_t](#) handle)
Mailbox_GetFreeSlots.
- bool [Mailbox_IsFull](#) ([Mailbox_t](#) handle)
Mailbox_IsFull.
- bool [Mailbox_IsEmpty](#) ([Mailbox_t](#) handle)
Mailbox_IsEmpty.
- void [ConditionVariable_Init](#) ([ConditionVariable_t](#) handle)
ConditionVariable_Init.
- void [ConditionVariable_Wait](#) ([ConditionVariable_t](#) handle, [Mutex_t](#) hMutex_)
ConditionVariable_Wait.
- void [ConditionVariable_Signal](#) ([ConditionVariable_t](#) handle)
ConditionVariable_Signal.
- void [ConditionVariable_Broadcast](#) ([ConditionVariable_t](#) handle)
ConditionVariable_Broadcast.
- bool [ConditionVariable_TimedWait](#) ([ConditionVariable_t](#) handle, [Mutex_t](#) hMutex_, uint32_t u32WaitTimeMS_)
ConditionVariable_TimedWait.
- void [ReaderWriterLock_Init](#) ([ReaderWriterLock_t](#) handle)
ReaderWriterLock_Init.
- void [ReaderWriterLock_AcquireReader](#) ([ReaderWriterLock_t](#) handle)
ReaderWriterLock_AcquireReader.
- void [ReaderWriterLock_ReleaseReader](#) ([ReaderWriterLock_t](#) handle)
ReaderWriterLock_ReleaseReader.
- void [ReaderWriterLock_AcquireWriter](#) ([ReaderWriterLock_t](#) handle)
ReaderWriterLock_AcquireWriter.
- void [ReaderWriterLock_ReleaseWriter](#) ([ReaderWriterLock_t](#) handle)
ReaderWriterLock_ReleaseWriter.
- bool [ReaderWriterLock_TimedAcquireWriter](#) ([ReaderWriterLock_t](#) handle, uint32_t u32TimeoutMs_)
ReaderWriterLock_TimedAcquireWriter.

- bool [ReaderWriterLock_TimedAcquireReader](#) ([ReaderWriterLock_t](#) handle, uint32_t u32TimeoutMs_)
ReaderWriterLock_TimedAcquireReader.
- void [Kernel_SetDebugPrintFunction](#) ([kernel_debug_print_t](#) pfPrintFunction_)
Kernel_SetDebugPrintFunction.
- void [Kernel_DebugPrint](#) (const char *szString_)
KernelDebug_DebugPrint.

20.5.1 Detailed Description

Header providing C-language API bindings for the [Mark3](#) kernel.

Definition in file [mark3c.h](#).

20.5.2 Macro Definition Documentation

20.5.2.1 CONDITIONVARIABLE_SIZE

```
#define CONDITIONVARIABLE_SIZE (sizeof(Fake_ConditionVariable))
```

Definition at line 79 of file [mark3c.h](#).

20.5.2.2 DECLARE_CONDITIONVARIABLE

```
#define DECLARE_CONDITIONVARIABLE(  
    name )
```

Value:

```
K_WORD          TOKEN_2(__condvar_, name) [WORD_ROUND(EVENTFLAG_SIZE)];  
                \br/>ConditionVariable_t name = (ConditionVariable_t)TOKEN_2(__condvar_, name);
```

Definition at line 161 of file [mark3c.h](#).

20.5.2.3 DECLARE_MAILBOX

```
#define DECLARE_MAILBOX(  
    name )
```

Value:

```
K_WORD          TOKEN_2(__mailbox_, name) [WORD_ROUND(  
MAILBOX_SIZE)];  
                \br/>Mailbox_t name = (Mailbox_t)TOKEN_2(__mailbox_, name);
```

Definition at line 147 of file [mark3c.h](#).

20.5.2.4 DECLARE_MESSAGE

```
#define DECLARE_MESSAGE(  
    name )
```

Value:

```
K_WORD    TOKEN_2(__message_, name) [WORD_ROUND(  
    MESSAGE_SIZE)];  
Message_t name = (Message_t) TOKEN_2(__message_, name);
```

Definition at line 135 of file [mark3c.h](#).

20.5.2.5 DECLARE_MESSAGEPOOL

```
#define DECLARE_MESSAGEPOOL(  
    name )
```

Value:

```
K_WORD    TOKEN_2(__messagepool_, name) [WORD_ROUND(  
    MESSAGEPOOL_SIZE)];  
MessagePool_t name = (MessagePool_t) TOKEN_2(__messagepool_, name);
```

Definition at line 139 of file [mark3c.h](#).

20.5.2.6 DECLARE_MESSAGEQUEUE

```
#define DECLARE_MESSAGEQUEUE(  
    name )
```

Value:

```
K_WORD    TOKEN_2(__messagequeue_, name) [WORD_ROUND(  
    MESSAGEQUEUE_SIZE)];  
MessageQueue_t name = (MessageQueue_t) TOKEN_2(__messagequeue_, name);
```

Definition at line 143 of file [mark3c.h](#).

20.5.2.7 DECLARE_MUTEX

```
#define DECLARE_MUTEX(  
    name )
```

Value:

```
K_WORD    TOKEN_2(__mutex_, name) [WORD_ROUND (MUTEX_SIZE)];  
    \  
    Mutex_t name = (Mutex_t) TOKEN_2 (__mutex_, name);
```

Definition at line 131 of file [mark3c.h](#).

20.5.2.8 DECLARE_NOTIFY

```
#define DECLARE_NOTIFY(  
    name )
```

Value:

```
K_WORD    TOKEN_2(__notify_, name) [WORD_ROUND (NOTIFY_SIZE)];  
    \  
    Notify_t name = (Notify_t) TOKEN_2 (__notify_, name);
```

Definition at line 151 of file [mark3c.h](#).

20.5.2.9 DECLARE_READERWRITERLOCK

```
#define DECLARE_READERWRITERLOCK(  
    name )
```

Value:

```
K_WORD    TOKEN_2(__readerwriterlock_, name) [WORD_ROUND (EVENTFLAG_SIZE)];  
    \  
    ReaderWriterLock_t name = (ReaderWriterLock_t) TOKEN_2 (__readerwriterlock_,  
    name);
```

Definition at line 165 of file [mark3c.h](#).

20.5.2.10 DECLARE_SEMAPHORE

```
#define DECLARE_SEMAPHORE(  
    name )
```

Value:

```
K_WORD    TOKEN_2(__semaphore_, name) [WORD_ROUND(  
    SEMAPHORE_SIZE)];  
Semaphore_t name = (Semaphore_t)TOKEN_2(__semaphore_, name); \
```

Definition at line 127 of file [mark3c.h](#).

20.5.2.11 DECLARE_THREAD

```
#define DECLARE_THREAD(  
    name )
```

Value:

```
K_WORD    TOKEN_2(__thread_, name) [WORD_ROUND(THREAD_SIZE)];  
Thread_t name = (Thread_t)TOKEN_2(__thread_, name); \
```

Definition at line 119 of file [mark3c.h](#).

20.5.2.12 DECLARE_TIMER

```
#define DECLARE_TIMER(  
    name )
```

Value:

```
K_WORD    TOKEN_2(__timer_, name) [WORD_ROUND(TIMER_SIZE)];  
Timer_t name = (Timer_t)TOKEN_2(__timer_, name); \
```

Definition at line 123 of file [mark3c.h](#).

20.5.2.13 MAILBOX_SIZE

```
#define MAILBOX_SIZE (sizeof(Fake_Mailbox))
```

Definition at line 73 of file [mark3c.h](#).

20.5.2.14 MESSAGE_SIZE

```
#define MESSAGE_SIZE (sizeof(Fake_Message))
```

Definition at line 71 of file [mark3c.h](#).

20.5.2.15 MESSAGEPOOL_SIZE

```
#define MESSAGEPOOL_SIZE (sizeof(Fake_MessagePool))
```

Definition at line 78 of file [mark3c.h](#).

20.5.2.16 MESSAGEQUEUE_SIZE

```
#define MESSAGEQUEUE_SIZE (sizeof(Fake_MessageQueue))
```

Definition at line 72 of file [mark3c.h](#).

20.5.2.17 MUTEX_SIZE

```
#define MUTEX_SIZE (sizeof(Fake_Mutex))
```

Definition at line 70 of file [mark3c.h](#).

20.5.2.18 NOTIFY_SIZE

```
#define NOTIFY_SIZE (sizeof(Fake_Notify))
```

Definition at line 74 of file [mark3c.h](#).

20.5.2.19 READERWRITERLOCK_SIZE

```
#define READERWRITERLOCK_SIZE (sizeof(Fake_ReaderWriterLock))
```

Definition at line 80 of file [mark3c.h](#).

20.5.2.20 SEMAPHORE_SIZE

```
#define SEMAPHORE_SIZE (sizeof(Fake_Semaphore))
```

Definition at line 69 of file [mark3c.h](#).

20.5.2.21 THREAD_SIZE

```
#define THREAD_SIZE (sizeof(Fake_Thread))
```

Definition at line 67 of file [mark3c.h](#).

20.5.2.22 TIMER_SIZE

```
#define TIMER_SIZE (sizeof(Fake_Timer))
```

Definition at line 68 of file [mark3c.h](#).

20.5.2.23 TOKEN_1

```
#define TOKEN_1(  
    x,  
    y ) x##y
```

Definition at line 112 of file [mark3c.h](#).

20.5.2.24 TOKEN_2

```
#define TOKEN_2(  
    x,  
    y ) TOKEN_1(x, y)
```

Definition at line 113 of file [mark3c.h](#).

20.5.2.25 WORD_ROUND

```
#define WORD_ROUND(  
    x ) (((x) + (sizeof(K_WORD) - 1)) / sizeof(K_WORD))
```

Definition at line 117 of file [mark3c.h](#).

20.5.3 Typedef Documentation

20.5.3.1 ConditionVariable_t

```
typedef void* ConditionVariable_t
```

Condition Variable opaque handle data type.

Definition at line 48 of file [mark3c.h](#).

20.5.3.2 kernel_debug_print_t

```
typedef void(* kernel_debug_print_t) (const char *szString_)
```

Definition at line 62 of file [mark3c.h](#).

20.5.3.3 Mailbox_t

```
typedef void* Mailbox_t
```

Mailbox opaque handle data type.

Definition at line 39 of file [mark3c.h](#).

20.5.3.4 Message_t

```
typedef void* Message_t
```

Message opaque handle data type.

Definition at line 40 of file [mark3c.h](#).

20.5.3.5 MessagePool_t

```
typedef void* MessagePool_t
```

MessagePool opaque handle data type.

Definition at line 41 of file [mark3c.h](#).

20.5.3.6 MessageQueue_t

```
typedef void* MessageQueue_t
```

MessageQueue opaque handle data type.

Definition at line 42 of file [mark3c.h](#).

20.5.3.7 Mutex_t

```
typedef void* Mutex_t
```

Mutex opaque handle data type.

Definition at line 43 of file [mark3c.h](#).

20.5.3.8 Notify_t

```
typedef void* Notify_t
```

Notification object opaque handle data type.

Definition at line 44 of file [mark3c.h](#).

20.5.3.9 panic_func_t

```
typedef void(* panic_func_t) (uint16_t u16PanicCode_)
```

Definition at line 282 of file [mark3c.h](#).

20.5.3.10 ReaderWriterLock_t

```
typedef void* ReaderWriterLock_t
```

Reader-writer-lock opaque handle data type.

Definition at line 49 of file [mark3c.h](#).

20.5.3.11 Semaphore_t

```
typedef void* Semaphore_t
```

Semaphore opaque handle data type.

Definition at line 45 of file [mark3c.h](#).

20.5.3.12 thread_entry_func_t

```
typedef void(* thread_entry_func_t) (void *pvArg_)
```

Definition at line 401 of file [mark3c.h](#).

20.5.3.13 Thread_t

```
typedef void* Thread_t
```

Thread opaque handle data type.

Definition at line 46 of file [mark3c.h](#).

20.5.3.14 timer_callback_t

```
typedef void(* timer_callback_t) (Thread_t hOwner_, void *pvData_)
```

Definition at line 570 of file [mark3c.h](#).

20.5.3.15 Timer_t

```
typedef void* Timer_t
```

Timer opaque handle data type.

Definition at line 47 of file [mark3c.h](#).

20.5.4 Enumeration Type Documentation

20.5.4.1 thread_state_t

```
enum thread_state_t
```

Define the various states that a thread can be in

Enumerator

THREAD_STATE_EXIT	
THREAD_STATE_READY	!< Thread has terminated via exit path
THREAD_STATE_BLOCKED	!< Thread is ready to run
THREAD_STATE_STOP	!< Thread is blocked on a blocking call
THREAD_STATE_INVALID	!< Thread has been manually stopped !< Invalid thread state

Definition at line 101 of file [mark3c.h](#).

20.5.5 Function Documentation

20.5.5.1 Alloc_Mailbox()

```
Mailbox_t Alloc_Mailbox (  
    void )
```

Alloc_Mailbox.

See also

Mailbox* AutoAlloc::NewMailbox()

Returns

Handle to an allocated object, or nullptr if heap exhausted

Definition at line 123 of file [mark3c.cpp](#).

20.5.5.2 Alloc_Memory()

```
void* Alloc_Memory (  
    size_t eSize_ )
```

Alloc_Memory.

See also

void* AutoAlloc::NewRawData(size_t sSize_)

Parameters

<i>e↔</i> <i>Size↔</i> —	Size in bytes to allocate from the one-time-allocate heap
--------------------------------	---

Returns

Pointer to an allocated blob of memory, or nullptr if heap exhausted

20.5.5.3 Alloc_Message()

```
Message_t Alloc_Message (
    void )
```

Alloc_Message.

See also

AutoAlloc::NewMessage()

Returns

Handle to an allocated object, or nullptr if heap exhausted

Definition at line 76 of file [mark3c.cpp](#).

20.5.5.4 Alloc_MessagePool()

```
MessagePool_t Alloc_MessagePool (
    void )
```

Definition at line 100 of file [mark3c.cpp](#).

20.5.5.5 Alloc_MessageQueue()

```
MessageQueue_t Alloc_MessageQueue (
    void )
```

Alloc_MessageQueue.

See also

MessageQueue* AutoAlloc::NewMessageQueue()

Returns

Handle to an allocated object, or nullptr if heap exhausted

Definition at line 88 of file [mark3c.cpp](#).

20.5.5.6 Alloc_Mutex()

```
Mutex_t Alloc_Mutex (
    void )
```

Alloc_Mutex.

See also

Mutex* AutoAlloc::NewMutex()

Returns

Handle to an allocated object, or nullptr if heap exhausted

Definition at line 52 of file [mark3c.cpp](#).

20.5.5.7 Alloc_Notify()

```
Notify_t Alloc_Notify (
    void )
```

Alloc_Notify.

See also

Notify* AutoAlloc::NewNotify()

Returns

Handle to an allocated object, or nullptr if heap exhausted

Definition at line 112 of file [mark3c.cpp](#).

20.5.5.8 Alloc_Semaphore()

```
Semaphore_t Alloc_Semaphore (
    void )
```

Alloc_Semaphore.

See also

Semaphore* AutoAlloc::NewSemaphore()

Returns

Handle to an allocated object, or nullptr if heap exhausted

Definition at line 40 of file [mark3c.cpp](#).

20.5.5.9 Alloc_Thread()

```
Thread_t Alloc_Thread (
    void )
```

Alloc_Thread.

See also

Thread* AutoAlloc::NewThread()

Returns

Handle to an allocated object, or nullptr if heap exhausted

Definition at line 156 of file [mark3c.cpp](#).

20.5.5.10 Alloc_Timer()

```
Timer_t Alloc_Timer (
    void )
```

Alloc_Timer.

See also

Timer* AutoAlloc::NewTimer()

Returns

Handle to an allocated object, or nullptr if heap exhausted

Definition at line 167 of file [mark3c.cpp](#).

20.5.5.11 Atomic_Add16()

```
uint16_t Atomic_Add16 (
    uint16_t * pu16Source_,
    uint16_t u16Val_ )
```

Atomic_Add16.

See also

uint16_t Atomic::Add(uint16_t *pu16Source_, uint16_t u16Val_)

Parameters

<i>pu16Source_</i>	Pointer to a variable
<i>u16Val_</i>	Value to add to the variable

Returns

Previously-held value in *pu16Source_*

Definition at line 638 of file [mark3c.cpp](#).

20.5.5.12 Atomic_Add32()

```
uint32_t Atomic_Add32 (  
    uint32_t * pu32Source_,  
    uint32_t u32Val_ )
```

Atomic_Add32.

See also

`uint32_t Atomic::Add(uint32_t *pu32Source_, uint32_t u32Val_)`

Parameters

<i>pu32Source_</i>	Pointer to a variable
<i>u32Val_</i>	Value to add to the variable

Returns

Previously-held value in *pu32Source_*

Definition at line 644 of file [mark3c.cpp](#).

20.5.5.13 Atomic_Add8()

```
uint8_t Atomic_Add8 (  
    uint8_t * pu8Source_,  
    uint8_t u8Val_ )
```

Atomic_Add8.

See also

`uint8_t Atomic::Add(uint8_t *pu8Source_, uint8_t u8Val_)`

Parameters

<i>pu8↔ Source_</i>	Pointer to a variable
<i>u8Val_</i>	Value to add to the variable

Returns

Previously-held value in pu8Source_

Definition at line 632 of file [mark3c.cpp](#).

20.5.5.14 Atomic_Set16()

```
uint16_t Atomic_Set16 (  
    uint16_t * pu16Source_,  
    uint16_t u16Val_ )
```

Atomic_Set16.

See also

`uint16_t Atomic::Set(uint16_t *pu16Source_, uint16_t u16Val_)`

Parameters

<i>pu16↔ Source_</i>	Pointer to a variable to set the value of
<i>u16Val_</i>	New value to set in the variable

Returns

Previously-set value

Definition at line 620 of file [mark3c.cpp](#).

20.5.5.15 Atomic_Set32()

```
uint32_t Atomic_Set32 (  
    uint32_t * pu32Source_,  
    uint32_t u32Val_ )
```

Atomic_Set32.

See also

`uint32_t Atomic::Set(uint32_t *pu32Source_, uint32_t u32Val_)`

Parameters

<i>pu32↔ Source_</i>	Pointer to a variable to set the value of
<i>u32Val_</i>	New value to set in the variable

Returns

Previously-set value

Definition at line 626 of file [mark3c.cpp](#).

20.5.5.16 Atomic_Set8()

```
uint8_t Atomic_Set8 (  
    uint8_t * pu8Source_,  
    uint8_t u8Val_ )
```

Atomic_Set8.

See also

`uint8_t Atomic::Set(uint8_t *pu8Source_, uint8_t u8Val_)`

Parameters

<i>pu8↔ Source_</i>	Pointer to a variable to set the value of
<i>u8Val_</i>	New value to set in the variable

Returns

Previously-set value

Definition at line 614 of file [mark3c.cpp](#).

20.5.5.17 Atomic_Sub16()

```
uint16_t Atomic_Sub16 (  
    uint16_t * pu16Source_,  
    uint16_t u16Val_ )
```

Atomic_Sub16.

See also

`uint16_t Atomic::Sub(uint16_t *pu16Source_, uint16_t u16Val_)`

Parameters

<i>pu16Source_</i>	Pointer to a variable
<i>u16Val_</i>	Value to subtract from the variable

Returns

Previously-held value in *pu16Source_*

Definition at line 656 of file [mark3c.cpp](#).

20.5.5.18 Atomic_Sub32()

```
uint32_t Atomic_Sub32 (
    uint32_t * pu32Source_,
    uint32_t u32Val_ )
```

Atomic_Sub32.

See also

`uint32_t Atomic::Sub(uint32_t *pu32Source_, uint32_t u32Val_)`

Parameters

<i>pu32Source_</i>	Pointer to a variable
<i>u32Val_</i>	Value to subtract from the variable

Returns

Previously-held value in *pu32Source_*

Definition at line 662 of file [mark3c.cpp](#).

20.5.5.19 Atomic_Sub8()

```
uint8_t Atomic_Sub8 (
    uint8_t * pu8Source_,
    uint8_t u8Val_ )
```

Atomic_Sub8.

See also

`uint8_t Atomic::Sub(uint8_t *pu8Source_, uint8_t u8Val_)`

Parameters

<i>pu8Source_</i>	Pointer to a variable
<i>u8Val_</i>	Value to subtract from the variable

Returns

Previously-held value in *pu8Source_*

Definition at line 650 of file [mark3c.cpp](#).

20.5.5.20 Atomic_TestAndSet()

```
bool Atomic_TestAndSet (
    bool * pbLock )
```

Atomic_TestAndSet.

See also

`bool Atomic::TestAndSet(bool *pbLock)`

Parameters

<i>pbLock</i>	Pointer to a value to test against. This will always be set to "true" at the end of a call to <i>TestAndSet</i> .
---------------	---

Returns

true - Lock value was "true" on entry, false - Lock was set

Definition at line 668 of file [mark3c.cpp](#).

20.5.5.21 ConditionVariable_Broadcast()

```
void ConditionVariable_Broadcast (
    ConditionVariable\_t handle )
```

ConditionVariable_Broadcast.

See also

`void ConditionVariable::Broadcast()`

Parameters

<i>handle</i>	Handle of the condition variable object
---------------	---

Definition at line 874 of file [mark3c.cpp](#).

20.5.5.22 ConditionVariable_Init()

```
void ConditionVariable_Init (  
    ConditionVariable_t handle )
```

ConditionVariable_Init.

See also

void ConditionVariable::Init()

Parameters

<i>handle</i>	Handle of the condition variable object
---------------	---

Definition at line 855 of file [mark3c.cpp](#).

20.5.5.23 ConditionVariable_Signal()

```
void ConditionVariable_Signal (  
    ConditionVariable_t handle )
```

ConditionVariable_Signal.

See also

void ConditionVariable::Signal()

Parameters

<i>handle</i>	Handle of the condition variable object
---------------	---

Definition at line 868 of file [mark3c.cpp](#).

20.5.5.24 ConditionVariable_TimedWait()

```
bool ConditionVariable_TimedWait (
    ConditionVariable_t handle,
    Mutex_t hMutex_,
    uint32_t u32WaitTimeMS_ )
```

ConditionVariable_TimedWait.

See also

bool ConditionVariable::Wait(Mutex* pclMutex_, uint32_t u32WaitTimeMS_)

Parameters

<i>handle</i>	Handle of the condition variable object
<i>hMutex_</i>	Handle of the mutex to lock on acquisition of the condition variable
<i>u32WaitTimeMS_</i>	Maximum time to wait for object

Returns

true on success, false on timeout

Definition at line 880 of file [mark3c.cpp](#).

20.5.5.25 ConditionVariable_Wait()

```
void ConditionVariable_Wait (
    ConditionVariable_t handle,
    Mutex_t hMutex_ )
```

ConditionVariable_Wait.

See also

void ConditionVariable::Wait(Mutex* pclMutex_)

Parameters

<i>handle</i>	Handle of the condition variable object
<i>hMutex_</i>	Handle of the mutex to lock on acquisition of the condition variable

Definition at line 861 of file [mark3c.cpp](#).

20.5.5.26 Free_Mailbox()

```
void Free_Mailbox (
    Mailbox_t handle )
```

Definition at line 128 of file [mark3c.cpp](#).

20.5.5.27 Free_Memory()

```
void Free_Memory (
    void * pvObject_ )
```

Free_Memory.

Parameters

<i>pvObject_</i>	Pointer to previously allocated block of memory
------------------	---

Definition at line 34 of file [mark3c.cpp](#).

20.5.5.28 Free_Message()

```
void Free_Message (
    Message_t handle )
```

Definition at line 82 of file [mark3c.cpp](#).

20.5.5.29 Free_MessagePool()

```
void Free_MessagePool (
    MessagePool_t handle )
```

Definition at line 106 of file [mark3c.cpp](#).

20.5.5.30 Free_MessageQueue()

```
void Free_MessageQueue (
    MessageQueue_t handle )
```

Definition at line 94 of file [mark3c.cpp](#).

20.5.5.31 Free_Mutex()

```
void Free_Mutex (
    Mutex_t handle )
```

Definition at line 58 of file [mark3c.cpp](#).

20.5.5.32 Free_Notify()

```
void Free_Notify (
    Notify_t handle )
```

Definition at line 118 of file [mark3c.cpp](#).

20.5.5.33 Free_Semaphore()

```
void Free_Semaphore (
    Semaphore_t handle )
```

Definition at line 46 of file [mark3c.cpp](#).

20.5.5.34 Free_Thread()

```
void Free_Thread (
    Thread_t handle )
```

Definition at line 161 of file [mark3c.cpp](#).

20.5.5.35 Free_Timer()

```
void Free_Timer (
    Timer_t handle )
```

Definition at line 172 of file [mark3c.cpp](#).

20.5.5.36 Kernel_DebugPrint()

```
void Kernel_DebugPrint (
    const char * szString_ )
```

KernelDebug_DebugPrint.

See also

`void DebugPrint(const char* szString_)`

Parameters

<i>sz↔ String_</i>	String to print to debug interface
------------------------	------------------------------------

Definition at line 282 of file [mark3c.cpp](#).

20.5.5.37 Kernel_GetTicks()

```
uint32_t Kernel_GetTicks (  
    void )
```

Kernel_GetTicks.

See also

Kernel::GetTicks()

Returns

Number of kernel ticks that have elapsed since boot

Definition at line 216 of file [mark3c.cpp](#).

20.5.5.38 Kernel_Init()

```
void Kernel_Init (  
    void )
```

Kernel_Init.

See also

void Kernel::Init()

Definition at line 180 of file [mark3c.cpp](#).

20.5.5.39 Kernel_IsPanic()

```
bool Kernel_IsPanic (  
    void )
```

Kernel_IsPanic.

See also

bool Kernel::IsPanic()

Returns

Whether or not the kernel is in a panic state

Definition at line [204](#) of file [mark3c.cpp](#).

20.5.5.40 Kernel_IsStarted()

```
bool Kernel_IsStarted (  
    void )
```

Kernel_IsStarted.

See also

bool Kernel::IsStarted()

Returns

Whether or not the kernel has started - true = running, false = not started

Definition at line [192](#) of file [mark3c.cpp](#).

20.5.5.41 Kernel_Panic()

```
void Kernel_Panic (  
    uint16_t u16Cause_ )
```

Kernel_Panic.

See also

void Kernel::Panic(uint16_t u16Cause_)

Parameters

<i>u16↔ Cause_</i>	Reason for the kernel panic
------------------------	-----------------------------

Definition at line 210 of file [mark3c.cpp](#).

20.5.5.42 Kernel_SetDebugPrintFunction()

```
void Kernel_SetDebugPrintFunction (
    kernel_debug_print_t pfPrintFunction_ )
```

Kernel_SetDebugPrintFunction.

See also

void Kernel::SetDebugPrintFunction()

Parameters

<i>pfPrint↔ Function_</i>	Function to use to print debug information from the kernel
-------------------------------	--

Definition at line 276 of file [mark3c.cpp](#).

20.5.5.43 Kernel_SetPanic()

```
void Kernel_SetPanic (
    panic_func_t pfPanic_ )
```

Kernel_SetPanic.

See also

void Kernel::SetPanic(PanicFunc_t pfPanic_)

Parameters

<i>pf↔ Panic↔ _</i>	Panic function pointer
-----------------------------	------------------------

Definition at line 198 of file [mark3c.cpp](#).

20.5.5.44 Kernel_Start()

```
void Kernel_Start (
    void )
```

Kernel_Start.

See also

`void Kernel::Start()`

Definition at line 186 of file [mark3c.cpp](#).

20.5.5.45 Mailbox_GetFreeSlots()

```
uint16_t Mailbox_GetFreeSlots (
    Mailbox_t handle )
```

Mailbox_GetFreeSlots.

See also

`uint16_t Mailbox::GetFreeSlots()`

Parameters

<i>handle</i>	Handle of the mailbox object
---------------	------------------------------

Returns

Number of free slots in the mailbox

Definition at line 832 of file [mark3c.cpp](#).

20.5.5.46 Mailbox_Init()

```
void Mailbox_Init (
    Mailbox_t handle,
    void * pvBuffer_,
    uint16_t u16BufferSize_,
    uint16_t u16ElementSize_ )
```

Mailbox_Init.

See also

`void Mailbox::Init(void *pvBuffer_, uint16_t u16BufferSize_, uint16_t u16ElementSize_)`

Parameters

<i>handle</i>	Handle of the mailbox object
<i>pvBuffer_</i>	Pointer to the static buffer to use for the mailbox
<i>u16BufferSize_</i>	Size of the mailbox buffer, in bytes
<i>u16Element↔ Size_</i>	Size of each envelope, in bytes

Definition at line 769 of file [mark3c.cpp](#).

20.5.5.47 Mailbox_IsEmpty()

```
bool Mailbox_IsEmpty (
    Mailbox_t handle )
```

Mailbox_IsEmpty.

See also

bool Mailbox::IsEmpty()

Parameters

<i>handle</i>	Handle of the mailbox object
---------------	------------------------------

Returns

true if the mailbox is empty, false otherwise

Definition at line 846 of file [mark3c.cpp](#).

20.5.5.48 Mailbox_IsFull()

```
bool Mailbox_IsFull (
    Mailbox_t handle )
```

Mailbox_IsFull.

See also

bool Mailbox::IsFull()

Parameters

<i>handle</i>	Handle of the mailbox object
---------------	------------------------------

Returns

true if the mailbox is full, false otherwise

Definition at line 839 of file [mark3c.cpp](#).

20.5.5.49 Mailbox_Receive()

```
void Mailbox_Receive (
    Mailbox_t handle,
    void * pvData_ )
```

Mailbox_Receive.

See also

void Mailbox::Receive(void *pvData_)

Parameters

<i>handle</i>	Handle of the mailbox object
<i>pvData_</i>	Pointer to a buffer that will have the envelope's contents copied into upon delivery.

Definition at line 804 of file [mark3c.cpp](#).

20.5.5.50 Mailbox_ReceiveTail()

```
void Mailbox_ReceiveTail (
    Mailbox_t handle,
    void * pvData_ )
```

Mailbox_ReceiveTail.

See also

void Mailbox::ReceiveTail(void *pvData_)

Parameters

<i>handle</i>	Handle of the mailbox object
<i>pvData_</i>	Pointer to a buffer that will have the envelope's contents copied into upon delivery.

Definition at line 811 of file [mark3c.cpp](#).

20.5.5.51 Mailbox_Send()

```
bool Mailbox_Send (
    Mailbox_t handle,
    void * pvData_ )
```

Mailbox_Send.

See also

`bool Mailbox::Send(void *pvData_)`

Parameters

<i>handle</i>	Handle of the mailbox object
<i>pvData_</i>	Pointer to the data object to send to the mailbox.

Returns

true - envelope was delivered, false - mailbox is full.

Definition at line 776 of file [mark3c.cpp](#).

20.5.5.52 Mailbox_SendTail()

```
bool Mailbox_SendTail (
    Mailbox_t handle,
    void * pvData_ )
```

Mailbox_SendTail.

See also

`bool Mailbox::SendTail(void *pvData_)`

Parameters

<i>handle</i>	Handle of the mailbox object
<i>pvData_</i>	Pointer to the data object to send to the mailbox.

Returns

true - envelope was delivered, false - mailbox is full.

Definition at line 783 of file [mark3c.cpp](#).

20.5.5.53 Mailbox_TimedReceive()

```
bool Mailbox_TimedReceive (
    Mailbox_t handle,
    void * pvData_,
    uint32_t u32TimeoutMS_ )
```

Mailbox_TimedReceive.

See also

bool Mailbox::Receive(void *pvData_, uint32_t u32TimeoutMS_)

Parameters

<i>handle</i>	Handle of the mailbox object
<i>pvData_</i>	Pointer to a buffer that will have the envelope's contents copied into upon delivery.
<i>u32TimeoutMS_</i>	Maximum time to wait for delivery.

Returns

true - envelope was delivered, false - delivery timed out.

Definition at line 818 of file [mark3c.cpp](#).

20.5.5.54 Mailbox_TimedReceiveTail()

```
bool Mailbox_TimedReceiveTail (
    Mailbox_t handle,
    void * pvData_,
    uint32_t u32TimeoutMS_ )
```

Mailbox_TimedReceiveTail.

See also

bool Mailbox::ReceiveTail(void *pvData_, uint32_t u32TimeoutMS_)

Parameters

<i>handle</i>	Handle of the mailbox object
<i>pvData_</i>	Pointer to a buffer that will have the envelope's contents copied into upon delivery.
<i>u32TimeoutMS_</i>	Maximum time to wait for delivery.

Returns

true - envelope was delivered, false - delivery timed out.

Definition at line 825 of file [mark3c.cpp](#).

20.5.5.55 Mailbox_TimedSend()

```
bool Mailbox_TimedSend (
    Mailbox_t handle,
    void * pvData_,
    uint32_t u32TimeoutMS_ )
```

Mailbox_TimedSend.

See also

bool Mailbox::Send(void *pvData_, uint32_t u32TimeoutMS_)

Parameters

<i>handle</i>	Handle of the mailbox object
<i>pvData_</i>	Pointer to the data object to send to the mailbox.
<i>u32TimeoutMS_</i>	Maximum time to wait for a free transmit slot

Returns

true - envelope was delivered, false - mailbox is full.

Definition at line 790 of file [mark3c.cpp](#).

20.5.5.56 Mailbox_TimedSendTail()

```
bool Mailbox_TimedSendTail (
    Mailbox_t handle,
    void * pvData_,
    uint32_t u32TimeoutMS_ )
```

Mailbox_TimedSendTail.

See also

`bool Mailbox::Send(void *pvData_, uint32_t u32TimeoutMS_)`

Parameters

<i>handle</i>	Handle of the mailbox object
<i>pvData_</i>	Pointer to the data object to send to the mailbox.
<i>u32TimeoutMS_</i>	Maximum time to wait for a free transmit slot

Returns

true - envelope was delivered, false - mailbox is full.

Definition at line 797 of file [mark3c.cpp](#).

20.5.5.57 Message_GetCode()

```
uint16_t Message_GetCode (  
    Message_t handle )
```

Message_GetCode.

See also

`uint16_t Message::GetCode()`

Parameters

<i>handle</i>	Handle of the message object
---------------	------------------------------

Returns

user code set in the object

Definition at line 704 of file [mark3c.cpp](#).

20.5.5.58 Message_GetData()

```
void* Message_GetData (  
    Message_t handle )
```

Message_GetData.

See also

`void* Message::GetData()`

Parameters

<i>handle</i>	Handle of the message object
---------------	------------------------------

Returns

Pointer to the data set in the message object

Definition at line 690 of file [mark3c.cpp](#).

20.5.5.59 Message_Init()

```
void Message_Init (
    Message_t handle )
```

Message_Init.

See also

void Message::Init()

Parameters

<i>handle</i>	Handle of the message object
---------------	------------------------------

Definition at line 676 of file [mark3c.cpp](#).

20.5.5.60 Message_SetCode()

```
void Message_SetCode (
    Message_t handle,
    uint16_t u16Code_ )
```

Message_SetCode.

See also

void Message::SetCode(uint16_t u16Code_)

Parameters

<i>handle</i>	Handle of the message object
<i>u16Code_</i>	Data code to set in the object

Definition at line 697 of file [mark3c.cpp](#).

20.5.5.61 Message_SetData()

```
void Message_SetData (
    Message_t handle,
    void * pvData_ )
```

Message_SetData.

See also

`void Message::SetData(void *pvData_)`

Parameters

<i>handle</i>	Handle of the message object
<i>pvData_</i>	Pointer to the data object to send in the message

Definition at line 683 of file [mark3c.cpp](#).

20.5.5.62 MessagePool_Init()

```
void MessagePool_Init (
    MessagePool_t handle )
```

MessagePool_Init.

See also

`void MessagePool::Init()`

Parameters

<i>handle</i>	Handle of the message pool object
---------------	-----------------------------------

Definition at line 725 of file [mark3c.cpp](#).

20.5.5.63 MessagePool_Pop()

```
Message_t MessagePool_Pop (
    MessagePool_t handle )
```

MessagePool_Pop.

See also

Message* MessagePool::Pop()

Parameters

<i>handle</i>	Handle of the message pool object
---------------	-----------------------------------

Returns

Handle to a Message object, or nullptr on allocation error

Definition at line 739 of file [mark3c.cpp](#).

20.5.5.64 MessagePool_Push()

```
void MessagePool_Push (
    MessagePool_t handle,
    Message_t msg )
```

MessagePool_Push.

See also

void MessagePool::Push(Message* pclMessage_)

Parameters

<i>handle</i>	Handle of the message pool object
<i>msg</i>	Message object to return back to the pool

Definition at line 732 of file [mark3c.cpp](#).

20.5.5.65 MessageQueue_GetCount()

```
uint16_t MessageQueue_GetCount (
    MessageQueue_t handle )
```

MessageQueue_GetCount.

See also

uint16_t MessageQueue::GetCount()

Returns

Count of pending messages in the queue.

Definition at line 760 of file [mark3c.cpp](#).

20.5.5.66 MessageQueue_Init()

```
void MessageQueue_Init (
    MessageQueue_t handle )
```

MessageQueue_Init.

See also

`void MessageQueue::Init()`

Parameters

<i>handle</i>	Handle to the message queue to initialize
---------------	---

Definition at line 711 of file [mark3c.cpp](#).

20.5.5.67 MessageQueue_Receive()

```
Message_t MessageQueue_Receive (
    MessageQueue_t handle )
```

MessageQueue_Receive.

See also

[Message_t](#) `MessageQueue::Receive()`

Parameters

<i>handle</i>	Handle of the message queue object
---------------	------------------------------------

Returns

Pointer to a message object at the head of the queue

Definition at line 718 of file [mark3c.cpp](#).

20.5.5.68 MessageQueue_Send()

```
void MessageQueue_Send (
    MessageQueue_t handle,
    Message_t hMessage_ )
```

MessageQueue_Send.

See also

`void MessageQueue::Send(Message *pclMessage_)`

Parameters

<i>handle</i>	Handle of the message queue object
<i>h↔ Message↔ _</i>	Handle to the message to send to the given queue

Definition at line 753 of file [mark3c.cpp](#).

20.5.5.69 MessageQueue_TimedReceive()

```
Message_t MessageQueue_TimedReceive (
    MessageQueue_t handle,
    uint32_t u32TimeWaitMS_ )
```

MessageQueue_TimedReceive.

See also

`Message_t MessageQueue::TimedReceive(uint32_t u32TimeWaitMS_)`

Parameters

<i>handle</i>	Handle of the message queue object
<i>u32TimeWaitM↔ S_</i>	The amount of time in ms to wait for a message before timing out and unblocking the waiting thread.

Returns

Pointer to a message object at the head of the queue or nullptr on timeout.

Definition at line 746 of file [mark3c.cpp](#).

20.5.5.70 Mutex_Claim()

```
void Mutex_Claim (
    Mutex\_t handle )
```

Mutex_Claim.

See also

`void Mutex::Claim()`

Parameters

<i>handle</i>	Handle of the mutex
---------------	---------------------

Definition at line [515](#) of file [mark3c.cpp](#).

20.5.5.71 Mutex_Init()

```
void Mutex_Init (
    Mutex\_t handle )
```

Mutex_Init.

See also

`void Mutex::Init()`

Parameters

<i>handle</i>	Handle of the mutex
---------------	---------------------

Definition at line [508](#) of file [mark3c.cpp](#).

20.5.5.72 Mutex_Release()

```
void Mutex_Release (
    Mutex\_t handle )
```

Mutex_Release.

See also

`void Mutex::Release()`

Parameters

<i>handle</i>	Handle of the mutex
---------------	---------------------

Definition at line 522 of file [mark3c.cpp](#).

20.5.5.73 Mutex_TimedClaim()

```
bool Mutex_TimedClaim (
    Mutex_t handle,
    uint32_t u32WaitTimeMS_ )
```

Mutex_TimedClaim.

See also

`bool Mutex::Claim(uint32_t u32WaitTimeMS_)`

Parameters

<i>handle</i>	Handle of the mutex
<i>u32WaitTimeMS_</i>	Time to wait before aborting

Returns

true if mutex was claimed, false on timeout

Definition at line 529 of file [mark3c.cpp](#).

20.5.5.74 Notify_Init()

```
void Notify_Init (
    Notify_t handle )
```

Notify_Init.

See also

`void Notify::Init()`

Parameters

<i>handle</i>	Handle of the notification object
---------------	-----------------------------------

Definition at line 584 of file [mark3c.cpp](#).

20.5.5.75 Notify_Signal()

```
void Notify_Signal (
    Notify_t handle )
```

Notify_Signal.

See also

`void Notify::Signal()`

Parameters

<i>handle</i>	Handle of the notification object
---------------	-----------------------------------

Definition at line 591 of file [mark3c.cpp](#).

20.5.5.76 Notify_TimedWait()

```
bool Notify_TimedWait (
    Notify_t handle,
    uint32_t u32WaitTimeMS_,
    bool * pbFlag_ )
```

Notify_TimedWait.

See also

`bool Notify::Wait(uint32_t u32WaitTimeMS_, bool *pbFlag_)`

Parameters

<i>handle</i>	Handle of the notification object
<i>u32WaitTimeMS_</i>	Maximum time to wait for notification in ms
<i>pbFlag_</i>	Flag to set to true on notification

Returns

true on unblock, false on timeout

Definition at line 605 of file [mark3c.cpp](#).

20.5.5.77 Notify_Wait()

```
void Notify_Wait (
    Notify_t handle,
    bool * pbFlag_ )
```

Notify_Wait.

See also

void Notify::Wait(bool *pbFlag_)

Parameters

<i>handle</i>	Handle of the notification object
<i>pbFlag_</i>	Flag to set to true on notification

Definition at line 598 of file [mark3c.cpp](#).

20.5.5.78 ReaderWriterLock_AcquireReader()

```
void ReaderWriterLock_AcquireReader (
    ReaderWriterLock_t handle )
```

ReaderWriterLock_AcquireReader.

See also

void ReaderWriterLock::AcquireReader()

Parameters

<i>handle</i>	Handle of the reader-writer object
---------------	------------------------------------

Definition at line 897 of file [mark3c.cpp](#).

20.5.5.79 ReaderWriterLock_AcquireWriter()

```
void ReaderWriterLock_AcquireWriter (
    ReaderWriterLock_t handle )
```

ReaderWriterLock_AcquireWriter.

See also

void ReaderWriterLock::AcquireWriter()

Parameters

<i>handle</i>	Handle of the reader-writer object
---------------	------------------------------------

Definition at line 911 of file [mark3c.cpp](#).

20.5.5.80 ReaderWriterLock_Init()

```
void ReaderWriterLock_Init (
    ReaderWriterLock_t handle )
```

ReaderWriterLock_Init.

See also

void ReaderWriterLock::Init()

Parameters

<i>handle</i>	Handle of the reader-writer object
---------------	------------------------------------

Definition at line 890 of file [mark3c.cpp](#).

20.5.5.81 ReaderWriterLock_ReleaseReader()

```
void ReaderWriterLock_ReleaseReader (
    ReaderWriterLock_t handle )
```

ReaderWriterLock_ReleaseReader.

See also

void ReaderWriterLock::ReleaseReader()

Parameters

<i>handle</i>	Handle of the reader-writer object
---------------	------------------------------------

Definition at line 904 of file [mark3c.cpp](#).

20.5.5.82 ReaderWriterLock_ReleaseWriter()

```
void ReaderWriterLock_ReleaseWriter (
    ReaderWriterLock_t handle )
```

ReaderWriterLock_ReleaseWriter.

See also

`void ReaderWriterLock::ReleaseWriter()`

Parameters

<i>handle</i>	Handle of the reader-writer object
---------------	------------------------------------

Definition at line 918 of file [mark3c.cpp](#).

20.5.5.83 ReaderWriterLock_TimedAcquireReader()

```
bool ReaderWriterLock_TimedAcquireReader (
    ReaderWriterLock_t handle,
    uint32_t u32TimeoutMs_ )
```

ReaderWriterLock_TimedAcquireReader.

See also

`bool ReaderWriterLock::AcquireReader(uint32_t u32TimeoutMs_)`

Parameters

<i>handle</i>	Handle of the reader-writer object
<i>u32TimeoutMs_</i>	Maximum time to wait for the reader lock before bailing

Returns

true on success, false on timeout

Definition at line 932 of file [mark3c.cpp](#).

20.5.5.84 ReaderWriterLock_TimedAcquireWriter()

```
bool ReaderWriterLock_TimedAcquireWriter (
    ReaderWriterLock_t handle,
    uint32_t u32TimeoutMs_ )
```

ReaderWriterLock_TimedAcquireWriter.

See also

`bool ReaderWriterLock::AcquireWriter(uint32_t u32TimeoutMs_)`

Parameters

<i>handle</i>	Handle of the reader-writer object
<i>u32Timeout↔ Ms_</i>	Maximum time to wait for the writer lock before bailing

Returns

true on success, false on timeout

Definition at line 925 of file [mark3c.cpp](#).

20.5.5.85 Scheduler_Enable()

```
void Scheduler_Enable (
    bool bEnable_ )
```

Scheduler_Enable.

See also

`void Scheduler::SetScheduler(bool bEnable_)`

Parameters

<i>b↔ Enable↔ _</i>	true to enable, false to disable the scheduler
-----------------------------	--

Definition at line 289 of file [mark3c.cpp](#).

20.5.5.86 Scheduler_GetCurrentThread()

```
Thread_t Scheduler_GetCurrentThread (
    void )
```

Scheduler_GetCurrentThread.

See also

`Thread* Scheduler::GetCurrentThread()`

Returns

Handle of the currently-running thread

Definition at line 301 of file [mark3c.cpp](#).

20.5.5.87 Scheduler_IsEnabled()

```
bool Scheduler_IsEnabled (
    void )
```

Scheduler_IsEnabled.

See also

bool Scheduler::IsEnabled()

Returns

true - scheduler enabled, false - disabled

Definition at line 295 of file [mark3c.cpp](#).

20.5.5.88 Semaphore_Init()

```
void Semaphore_Init (
    Semaphore_t handle,
    uint16_t u16InitVal_,
    uint16_t u16MaxVal_ )
```

Semaphore_Init.

See also

void Semaphore::Init(uint16_t u16InitVal_, uint16_t u16MaxVal_)

Parameters

<i>handle</i>	Handle of the semaphore
<i>u16InitVal_↔</i> —	Initial value of the semaphore
<i>u16Max↔</i> <i>Val_</i>	Maximum value that can be held for a semaphore

Definition at line 478 of file [mark3c.cpp](#).

20.5.5.89 Semaphore_Pend()

```
void Semaphore_Pend (
    Semaphore_t handle )
```

Semaphore_Pend.

See also

void Semaphore::Pend()

Parameters

<i>handle</i>	Handle of the semaphore
---------------	-------------------------

Definition at line 492 of file [mark3c.cpp](#).

20.5.5.90 Semaphore_Post()

```
void Semaphore_Post (
    Semaphore_t handle )
```

Semaphore_Post.

See also

void Semaphore::Post()

Parameters

<i>handle</i>	Handle of the semaphore
---------------	-------------------------

Definition at line 485 of file [mark3c.cpp](#).

20.5.5.91 Semaphore_TimedPend()

```
bool Semaphore_TimedPend (
    Semaphore_t handle,
    uint32_t u32WaitTimeMS_ )
```

Semaphore_TimedPend.

See also

bool Semaphore::Pend(uint32_t u32WaitTimeMS_)

Parameters

<i>handle</i>	Handle of the semaphore
<i>u32WaitTimeM↔S_</i>	Time in ms to wait

Returns

true if semaphore was acquired, false on timeout

Definition at line 499 of file [mark3c.cpp](#).

20.5.5.92 Thread_CoopYield()

```
void Thread_CoopYield (
    void )
```

Thread_CoopYield.

See also

void Thread::CoopYield()

Definition at line 415 of file [mark3c.cpp](#).

20.5.5.93 Thread_Exit()

```
void Thread_Exit (
    Thread_t handle )
```

Thread_Exit.

See also

void Thread::Exit()

Parameters

<i>handle</i>	Handle of the thread
---------------	----------------------

Definition at line 381 of file [mark3c.cpp](#).

20.5.5.94 Thread_GetCurPriority()

```
PORT_PRIO_TYPE Thread_GetCurPriority (
    Thread_t handle )
```

Thread_GetCurPriority.

See also

[PORT_PRIO_TYPE](#) Thread::GetCurPriority()

Parameters

<i>handle</i>	Handle of the thread
---------------	----------------------

Returns

Current priority of the thread considering priority inheritance

Definition at line [354](#) of file [mark3c.cpp](#).

20.5.5.95 Thread_GetID()

```
uint8_t Thread_GetID (
    Thread_t handle )
```

Thread_GetID.

See also

[uint8_t](#) Thread::GetID()

Parameters

<i>handle</i>	Handle of the thread
---------------	----------------------

Returns

Return ID assigned to the thread

Definition at line [426](#) of file [mark3c.cpp](#).

20.5.5.96 Thread_GetPriority()

```
PORT_PRIO_TYPE Thread_GetPriority (
    Thread_t handle )
```

Thread_GetPriority.

See also

[PORT_PRIO_TYPE](#) Thread::GetPriority()

Parameters

<i>handle</i>	Handle of the thread
---------------	----------------------

Returns

Current priority of the thread not considering priority inheritance

Definition at line [348](#) of file [mark3c.cpp](#).

20.5.5.97 Thread_GetState()

```
thread_state_t Thread_GetState (
    Thread_t handle )
```

Thread_GetState.

See also

ThreadState Thread::GetState()

Parameters

<i>handle</i>	Handle of the thread
---------------	----------------------

Returns

The thread's current execution state

Definition at line [438](#) of file [mark3c.cpp](#).

20.5.5.98 Thread_Init()

```
void Thread_Init (
    Thread_t handle,
    K_WORD * pwStack_,
    uint16_t ul6StackSize_,
    PORT_PRIO_TYPE uXPriority_,
    thread_entry_func_t pfEntryPoint_,
    void * pvArg_ )
```

Thread_Init.

See also

```
void Thread::Init(K_WORD *pwStack_, uint16_t u16StackSize_, PORT_PRIO_TYPE uXPriority_, Thread↵
Entry_t pfEntryPoint_, void *pvArg_)
```

Parameters

<i>handle</i>	Handle of the thread to initialize
<i>pwStack_</i>	Pointer to the stack to use for the thread
<i>u16StackSize_</i>	Size of the stack (in bytes)
<i>uXPriority_</i>	Priority of the thread (0 = idle, 7 = max)
<i>pfEntryPoint_</i>	This is the function that gets called when the thread is started
<i>pvArg_</i>	Pointer to the argument passed into the thread's entryptoint function.

Definition at line 310 of file [mark3c.cpp](#).

20.5.5.99 Thread_SetID()

```
void Thread_SetID (
    Thread_t handle,
    uint8_t u8ID_ )
```

Thread_SetID.

See also

```
void Thread::SetID(uint8_t u8ID_)
```

Parameters

<i>handle</i>	Handle of the thread
<i>u8ID_</i>	ID To assign to the thread

Definition at line 420 of file [mark3c.cpp](#).

20.5.5.100 Thread_SetPriority()

```
void Thread_SetPriority (
    Thread_t handle,
    PORT_PRIO_TYPE uXPriority_ )
```

Thread_SetPriority.

See also

```
void Thread::SetPriority(PORT_PRIO_TYPE uXPriority_)
```

Parameters

<i>handle</i>	Handle of the thread
<i>uX↔ Priority_</i>	New priority level

Definition at line 374 of file [mark3c.cpp](#).

20.5.5.101 Thread_Sleep()

```
void Thread_Sleep (
    uint32_t u32TimeMs_ )
```

Thread_Sleep.

See also

`void Thread::Sleep(uint32_t u32TimeMs_)`

Parameters

<i>u32Time↔ Ms_</i>	Time in ms to block the thread for
-------------------------	------------------------------------

Definition at line 388 of file [mark3c.cpp](#).

20.5.5.102 Thread_Start()

```
void Thread_Start (
    Thread_t handle )
```

Thread_Start.

See also

`void Thread::Start()`

Parameters

<i>handle</i>	Handle of the thread to start
---------------	-------------------------------

Definition at line 322 of file [mark3c.cpp](#).

20.5.5.103 Thread_Stop()

```
void Thread_Stop (
    Thread_t handle )
```

Thread_Stop.

See also

void Thread::Stop()

Parameters

<i>handle</i>	Handle of the thread to stop
---------------	------------------------------

Definition at line 329 of file [mark3c.cpp](#).

20.5.5.104 Thread_Yield()

```
void Thread_Yield (
    void )
```

Thread_Yield.

See also

void Thread::Yield()

Definition at line 410 of file [mark3c.cpp](#).

20.5.5.105 Timer_Init()

```
void Timer_Init (
    Timer_t handle )
```

Timer_Init.

See also

void Timer::Init()

Parameters

<i>handle</i>	Handle of the timer
---------------	---------------------

Definition at line 448 of file [mark3c.cpp](#).

20.5.5.106 Timer_Restart()

```
void Timer_Restart (
    Timer_t handle )
```

Timer_Restart.

See also

`void Timer::Start()`

Parameters

<i>handle</i>	Handle of the timer to restart.
---------------	---------------------------------

Definition at line 469 of file [mark3c.cpp](#).

20.5.5.107 Timer_Start()

```
void Timer_Start (
    Timer_t handle,
    bool bRepeat_,
    uint32_t u32IntervalMs_,
    timer_callback_t pfCallback_,
    void * pvData_ )
```

Timer_Start.

See also

`void Timer::Start(bool bRepeat_, uint32_t u32IntervalMs_, uint32_t u32ToleranceMs_, TimerCallbackC_t pf↵
Callback_, void *pvData_)`

Parameters

<i>handle</i>	Handle of the timer
<i>bRepeat_</i>	Restart the timer continuously on expiry
<i>u32Interval↵ Ms_</i>	Time in ms to expiry
<i>pfCallback_</i>	Callback to run on timer expiry
<i>pvData_</i>	Data to pass to the callback on expiry

Definition at line 455 of file [mark3c.cpp](#).

20.5.5.108 Timer_Stop()

```
void Timer_Stop (
    Timer_t handle )
```

Timer_Stop.

See also

void Timer::Stop()

Parameters

<i>handle</i>	Handle of the timer
---------------	---------------------

Definition at line 462 of file [mark3c.cpp](#).

20.6 mark3c.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00021 #pragma once
00022
00023 #include "mark3cfg.h"
00024 #include "fake_types.h"
00025
00026 #include <stdint.h>
00027 #include <stdbool.h>
00028
00029 #if defined(__cplusplus)
00030 extern "C" {
00031 #endif
00032
00033 //-----
00034 // Define a series of handle types to be used in place of the underlying classes
00035 // of Mark3.
00036 #if KERNEL_EVENT_FLAGS
00037 typedef void* EventFlag_t;
00038 #endif // #if KERNEL_EVENT_FLAGS
00039 typedef void* Mailbox_t;
00040 typedef void* Message_t;
00041 typedef void* MessagePool_t;
00042 typedef void* MessageQueue_t;
00043 typedef void* Mutex_t;
00044 typedef void* Notify_t;
00045 typedef void* Semaphore_t;
00046 typedef void* Thread_t;
00047 typedef void* Timer_t;
00048 typedef void* ConditionVariable_t;
00049 typedef void* ReaderWriterLock_t;
00050
00051 //-----
00052 // Function pointer types used by Kernel APIs
00053 #if KERNEL_THREAD_CREATE_CALLOUT
```

```

00054 typedef void (*thread_create_callout_t)(Thread_t hThread_);
00055 #endif // #if KERNEL_THREAD_CREATE_CALLOUT
00056 #if KERNEL_THREAD_EXIT_CALLOUT
00057 typedef void (*thread_exit_callout_t)(Thread_t hThread_);
00058 #endif // #if KERNEL_THREAD_EXIT_CALLOUT
00059 #if KERNEL_CONTEXT_SWITCH_CALLOUT
00060 typedef void (*thread_context_callout_t)(Thread_t hThread_);
00061 #endif // #if KERNEL_CONTEXT_SWITCH_CALLOUT
00062 typedef void (*kernel_debug_print_t)(const char* szString_);
00063
00064 //-----
00065 // Use the sizes of the structs in fake_types.h to generate opaque object-blobs
00066 // that get instantiated as kernel objects (from the C++ code) later.
00067 #define THREAD_SIZE (sizeof(Fake_Thread))
00068 #define TIMER_SIZE (sizeof(Fake_Timer))
00069 #define SEMAPHORE_SIZE (sizeof(Fake_Semaphore))
00070 #define MUTEX_SIZE (sizeof(Fake_Mutex))
00071 #define MESSAGE_SIZE (sizeof(Fake_Message))
00072 #define MESSAGEQUEUE_SIZE (sizeof(Fake_MessageQueue))
00073 #define MAILBOX_SIZE (sizeof(Fake_Mailbox))
00074 #define NOTIFY_SIZE (sizeof(Fake_Notify))
00075 #if KERNEL_EVENT_FLAGS
00076 #define EVENTFLAG_SIZE (sizeof(Fake_EventFlag))
00077 #endif // #if KERNEL_EVENT_FLAGS
00078 #define MESSAGEPOOL_SIZE (sizeof(Fake_MessagePool))
00079 #define CONDITIONVARIABLE_SIZE (sizeof(Fake_ConditionVariable))
00080 #define READERWRITERLOCK_SIZE (sizeof(Fake_ReaderWriterLock))
00081
00082 #if KERNEL_EVENT_FLAGS
00083 //-----
00084 typedef enum {
00085     EVENT_FLAG_ALL_SET,
00086     EVENT_FLAG_ANY_SET,
00087     EVENT_FLAG_ALL_CLEAR,
00088     EVENT_FLAG_ANY_CLEAR,
00089     EVENT_FLAG_PENDING_UNBLOCK
00090 } event_flag_operation_t;
00091 #endif // #if KERNEL_EVENT_FLAGS
00092
00093 //-----
00094 typedef enum {
00095     THREAD_STATE_EXIT = 0,
00096     THREAD_STATE_READY,
00097     THREAD_STATE_BLOCKED,
00098     THREAD_STATE_STOP,
00099     THREAD_STATE_INVALID
00100 } thread_state_t;
00101
00102 //-----
00103 // Macros for declaring opaque buffers of an appropriate size for the given
00104 // kernel objects
00105 #define TOKEN_1(x, y) x##y
00106 #define TOKEN_2(x, y) TOKEN_1(x, y)
00107
00108 // Ensure that opaque buffers are sized to the nearest word - which is
00109 // a platform-dependent value.
00110 #define WORD_ROUND(x) ((x) + (sizeof(K_WORD) - 1)) / sizeof(K_WORD)
00111
00112 #define DECLARE_THREAD(name)
00113     K_WORD TOKEN_2(__thread_, name)[WORD_ROUND(THREAD_SIZE)];
00114     Thread_t name = (Thread_t)TOKEN_2(__thread_, name);
00115
00116 #define DECLARE_TIMER(name)
00117     K_WORD TOKEN_2(__timer_, name)[WORD_ROUND(TIMER_SIZE)];
00118     Timer_t name = (Timer_t)TOKEN_2(__timer_, name);
00119
00120 #define DECLARE_SEMAPHORE(name)
00121     K_WORD TOKEN_2(__semaphore_, name)[WORD_ROUND(SEMAPHORE_SIZE)];
00122     Semaphore_t name = (Semaphore_t)TOKEN_2(__semaphore_, name);
00123
00124 #define DECLARE_MUTEX(name)
00125     K_WORD TOKEN_2(__mutex_, name)[WORD_ROUND(MUTEX_SIZE)];
00126     Mutex_t name = (Mutex_t)TOKEN_2(__mutex_, name);
00127
00128 #define DECLARE_MESSAGE(name)
00129     K_WORD TOKEN_2(__message_, name)[WORD_ROUND(MESSAGE_SIZE)];
00130     Message_t name = (Message_t)TOKEN_2(__message_, name);

```



```

00138
00139 #define DECLARE_MESSAGEPOOL(name)
00140     K_WORD      TOKEN_2(__messagepool_, name) [WORD_ROUND(MESSAGEPOOL_SIZE)];
00141     MessagePool_t name = (MessagePool_t)TOKEN_2(__messagepool_, name);
00142
00143 #define DECLARE_MESSAGEQUEUE(name)
00144     K_WORD      TOKEN_2(__messagequeue_, name) [WORD_ROUND(MESSAGEQUEUE_SIZE)];
00145     MessageQueue_t name = (MessageQueue_t)TOKEN_2(__messagequeue_, name);
00146
00147 #define DECLARE_MAILBOX(name)
00148     K_WORD      TOKEN_2(__mailbox_, name) [WORD_ROUND(MAILBOX_SIZE)];
00149     Mailbox_t name = (Mailbox_t)TOKEN_2(__mailbox_, name);
00150
00151 #define DECLARE_NOTIFY(name)
00152     K_WORD      TOKEN_2(__notify_, name) [WORD_ROUND(NOTIFY_SIZE)];
00153     Notify_t name = (Notify_t)TOKEN_2(__notify_, name);
00154
00155 #if KERNEL_EVENT_FLAGS
00156 #define DECLARE_EVENTFLAG(name)
00157     K_WORD      TOKEN_2(__eventflag_, name) [WORD_ROUND(EVENTFLAG_SIZE)];
00158     EventFlag_t name = (EventFlag_t)TOKEN_2(__eventflag_, name);
00159 #endif // #if KERNEL_EVENT_FLAGS
00160
00161 #define DECLARE_CONDITIONVARIABLE(name)
00162     K_WORD      TOKEN_2(__condvar_, name) [WORD_ROUND(EVENTFLAG_SIZE)];
00163     ConditionVariable_t name = (ConditionVariable_t)TOKEN_2(__condvar_, name);
00164
00165 #define DECLARE_READERWRITERLOCK(name)
00166     K_WORD      TOKEN_2(__readerwriterlock_, name) [WORD_ROUND(EVENTFLAG_SIZE)];
00167     ReaderWriterLock_t name = (ReaderWriterLock_t)TOKEN_2(__readerwriterlock_, name);
00168
00169 //-----
00170 // Allocate-once Memory managment APIs
00171 void* Alloc_Memory(size_t eSize_);
00172
00173 void Free_Memory(void* pvObject_);
00174
00175 Semaphore_t Alloc_Semaphore(void);
00176 void Free_Semaphore(Semaphore_t handle);
00177
00178 Mutex_t Alloc_Mutex(void);
00179 void Free_Mutex(Mutex_t handle);
00180
00181 #if KERNEL_EVENT_FLAGS
00182 EventFlag_t Alloc_EventFlag(void);
00183 void Free_EventFlag(EventFlag_t handle);
00184 #endif // #if KERNEL_EVENT_FLAGS
00185
00186 Message_t Alloc_Message(void);
00187 void Free_Message(Message_t handle);
00188
00189 MessageQueue_t Alloc_MessageQueue(void);
00190 void Free_MessageQueue(MessageQueue_t handle);
00191
00192 MessagePool_t Alloc_MessagePool(void);
00193 void Free_MessagePool(MessagePool_t handle);
00194
00195 Notify_t Alloc_Notify(void);
00196 void Free_Notify(Notify_t handle);
00197
00198 Mailbox_t Alloc_Mailbox(void);
00199 void Free_Mailbox(Mailbox_t handle);
00200
00201 Thread_t Alloc_Thread(void);
00202 void Free_Thread(Thread_t handle);
00203
00204 Timer_t Alloc_Timer(void);
00205 void Free_Timer(Timer_t handle);
00206
00207 //-----
00208 // Kernel APIs
00209 void Kernel_Init(void);

```

```

00273 void Kernel_Start(void);
00280 bool Kernel_IsStarted(void);
00281
00282 typedef void (*panic_func_t)(uint16_t ul6PanicCode_);
00288 void Kernel_SetPanic(panic_func_t pfPanic_);
00294 bool Kernel_IsPanic(void);
00300 void Kernel_Panic(uint16_t ul6Cause_);
00301
00302 #if KERNEL_THREAD_CREATE_CALLOUT
00303
00308 void Kernel_SetThreadCreateCallout(thread_create_callout_t pfCreate_);
00309 #endif // #if KERNEL_THREAD_CREATE_CALLOUT
00310
00311 #if KERNEL_THREAD_EXIT_CALLOUT
00312
00317 void Kernel_SetThreadExitCallout(thread_exit_callout_t pfExit_);
00318 #endif // #if KERNEL_THREAD_EXIT_CALLOUT
00319
00320 #if KERNEL_CONTEXT_SWITCH_CALLOUT
00321
00326 void Kernel_SetThreadContextSwitchCallout(thread_context_callout_t pfContext_);
00327 #endif // #if KERNEL_CONTEXT_SWITCH_CALLOUT
00328
00329 #if KERNEL_THREAD_CREATE_CALLOUT
00330
00335 thread_create_callout_t Kernel_GetThreadCreateCallout(void);
00336 #endif // #if KERNEL_THREAD_CREATE_CALLOUT
00337
00338 #if KERNEL_THREAD_EXIT_CALLOUT
00339
00344 thread_exit_callout_t Kernel_GetThreadExitCallout(void);
00345 #endif // #if KERNEL_THREAD_EXIT_CALLOUT
00346
00347 #if KERNEL_CONTEXT_SWITCH_CALLOUT
00348
00353 thread_context_callout_t Kernel_GetThreadContextSwitchCallout(void);
00354 #endif // #if KERNEL_CONTEXT_SWITCH_CALLOUT
00355
00356 #if KERNEL_STACK_CHECK
00357
00363 void Kernel_SetStackGuardThreshold(uint16_t ul6Threshold_);
00364
00370 uint16_t Kernel_GetStackGuardThreshold(void);
00371 #endif // #if KERNEL_STACK_CHECK
00372
00378 uint32_t Kernel_GetTicks(void);
00379
00380 //-----
00381 // Scheduler APIs
00387 void Scheduler_Enable(bool bEnable_);
00393 bool Scheduler_IsEnabled(void);
00399 Thread_t Scheduler_GetCurrentThread(void);
00400
00401 typedef void (*thread_entry_func_t)(void* pvArg_);
00402 //-----
00403 // Thread APIs
00417 void Thread_Init(Thread_t handle,
00418                  K_WORD* pwStack_,
00419                  uint16_t ul6StackSize_,
00420                  PORT_PRIO_TYPE uXPriority_,
00421                  thread_entry_func_t pfEntryPoint_,
00422                  void* pvArg_);
00428 void Thread_Start(Thread_t handle);
00434 void Thread_Stop(Thread_t handle);
00435
00436 #if KERNEL_NAMED_THREADS
00437
00443 void Thread_SetName(Thread_t handle, const char* szName_);
00450 const char* Thread_GetName(Thread_t handle);
00451 #endif // #if KERNEL_NAMED_THREADS
00452
00459 PORT_PRIO_TYPE Thread_GetPriority(Thread_t handle);
00466 PORT_PRIO_TYPE Thread_GetCurPriority(Thread_t handle);
00467
00468 #if KERNEL_ROUND_ROBIN
00469
00475 void Thread_SetQuantum(Thread_t handle, uint16_t ul6Quantum_);
00482 uint16_t Thread_GetQuantum(Thread_t handle);
00483 #endif // #if KERNEL_ROUND_ROBIN
00484
00491 void Thread_SetPriority(Thread_t handle, PORT_PRIO_TYPE uXPriority_);
00497 void Thread_Exit(Thread_t handle);
00503 void Thread_Sleep(uint32_t u32TimeMs_);
00504
00505 #if KERNEL_EXTENDED_CONTEXT
00506
00512 void* Thread_GetExtendedContext(Thread_t handle);

```

```

00513
00520 void Thread_SetExtendedContext(Thread_t handle, void* pvData_);
00521 #endif // #if KERNEL_EXTENDED_CONTEXT
00522
00527 void Thread_Yield(void);
00528
00533 void Thread_CoopYield(void);
00534
00541 void Thread_SetID(Thread_t handle, uint8_t u8ID_);
00548 uint8_t Thread_GetID(Thread_t handle);
00549
00550 #if KERNEL_STACK_CHECK
00551
00557 uint16_t Thread_GetStackSlack(Thread_t handle);
00558 #endif // #if KERNEL_STACK_CHECK
00559
00566 thread_state_t Thread_GetState(Thread_t handle);
00567
00568 //-----
00569 // Timer APIs
00570 typedef void (*timer_callback_t)(Thread_t hOwner_, void* pvData_);
00576 void Timer_Init(Timer_t handle);
00587 void Timer_Start(Timer_t handle, bool bRepeat_, uint32_t u32IntervalMs_,
    timer_callback_t pfCallback_, void* pvData_);
00588
00594 void Timer_Restart(Timer_t handle);
00595
00601 void Timer_Stop(Timer_t handle);
00602
00603 //-----
00604 // Semaphore APIs
00612 void Semaphore_Init(Semaphore_t handle, uint16_t u16InitVal_, uint16_t u16MaxVal_);
00618 void Semaphore_Post(Semaphore_t handle);
00624 void Semaphore_Pend(Semaphore_t handle);
00632 bool Semaphore_TimedPend(Semaphore_t handle, uint32_t u32WaitTimeMS_);
00633
00634 //-----
00635 // Mutex APIs
00641 void Mutex_Init(Mutex_t handle);
00647 void Mutex_Claim(Mutex_t handle);
00653 void Mutex_Release(Mutex_t handle);
00661 bool Mutex_TimedClaim(Mutex_t handle, uint32_t u32WaitTimeMS_);
00662
00663 #if KERNEL_EVENT_FLAGS
00664 //-----
00665 // EventFlag APIs
00671 void EventFlag_Init(EventFlag_t handle);
00680 uint16_t EventFlag_Wait(EventFlag_t handle, uint16_t u16Mask_, event_flag_operation_t eMode_);
00690 uint16_t EventFlag_TimedWait(EventFlag_t handle, uint16_t u16Mask_, event_flag_operation_t eMode_, uint32_t
    u32TimeMS_);
00697 void EventFlag_Set(EventFlag_t handle, uint16_t u16Mask_);
00704 void EventFlag_Clear(EventFlag_t handle, uint16_t u16Mask_);
00711 uint16_t EventFlag_GetMask(EventFlag_t handle);
00712 #endif // #if KERNEL_EVENT_FLAGS
00713
00714 //-----
00715 // Notification APIs
00721 void Notify_Init(Notify_t handle);
00727 void Notify_Signal(Notify_t handle);
00734 void Notify_Wait(Notify_t handle, bool* pbFlag_);
00743 bool Notify_TimedWait(Notify_t handle, uint32_t u32WaitTimeMS_, bool* pbFlag_);
00744
00745 //-----
00746 // Atomic Functions
00754 uint8_t Atomic_Set8(uint8_t* pu8Source_, uint8_t u8Val_);
00762 uint16_t Atomic_Set16(uint16_t* pu16Source_, uint16_t u16Val_);
00770 uint32_t Atomic_Set32(uint32_t* pu32Source_, uint32_t u32Val_);
00778 uint8_t Atomic_Add8(uint8_t* pu8Source_, uint8_t u8Val_);
00786 uint16_t Atomic_Add16(uint16_t* pu16Source_, uint16_t u16Val_);
00794 uint32_t Atomic_Add32(uint32_t* pu32Source_, uint32_t u32Val_);
00802 uint8_t Atomic_Sub8(uint8_t* pu8Source_, uint8_t u8Val_);
00810 uint16_t Atomic_Sub16(uint16_t* pu16Source_, uint16_t u16Val_);
00818 uint32_t Atomic_Sub32(uint32_t* pu32Source_, uint32_t u32Val_);
00827 bool Atomic_TestAndSet(bool* pbLock_);
00828
00829 //-----
00830 // Message/Message Queue APIs
00836 void Message_Init(Message_t handle);
00843 void Message_SetData(Message_t handle, void* pvData_);
00850 void* Message_GetData(Message_t handle);
00857 void Message_SetCode(Message_t handle, uint16_t u16Code_);
00864 uint16_t Message_GetCode(Message_t handle);
00870 void MessageQueue_Init(MessageQueue_t handle);
00877 Message_t MessageQueue_Receive(MessageQueue_t handle);
00888 Message_t MessageQueue_TimedReceive(MessageQueue_t handle, uint32_t u32TimeWaitMS_
    );
00889

```

```

00896 void MessageQueue_Send(MessageQueue_t handle, Message_t hMessage_);
00897
00903 uint16_t MessageQueue_GetCount(MessageQueue_t handle);
00904
00910 void MessagePool_Init(MessagePool_t handle);
00911
00918 void MessagePool_Push(MessagePool_t handle, Message_t msg);
00919
00926 Message_t MessagePool_Pop(MessagePool_t handle);
00927
00928 //-----
00929 // Mailbox APIs
00930
00939 void Mailbox_Init(Mailbox_t handle, void* pvBuffer_, uint16_t u16BufferSize_, uint16_t
u16ElementSize_);
00940
00948 bool Mailbox_Send(Mailbox_t handle, void* pvData_);
00949
00957 bool Mailbox_SendTail(Mailbox_t handle, void* pvData_);
00958
00967 bool Mailbox_TimedSend(Mailbox_t handle, void* pvData_, uint32_t u32TimeoutMS_);
00968
00977 bool Mailbox_TimedSendTail(Mailbox_t handle, void* pvData_, uint32_t u32TimeoutMS_);
00978
00986 void Mailbox_Receive(Mailbox_t handle, void* pvData_);
00987
00995 void Mailbox_ReceiveTail(Mailbox_t handle, void* pvData_);
00996
01006 bool Mailbox_TimedReceive(Mailbox_t handle, void* pvData_, uint32_t u32TimeoutMS_);
01007
01017 bool Mailbox_TimedReceiveTail(Mailbox_t handle, void* pvData_, uint32_t
u32TimeoutMS_);
01018
01025 uint16_t Mailbox_GetFreeSlots(Mailbox_t handle);
01026
01033 bool Mailbox_IsFull(Mailbox_t handle);
01034
01041 bool Mailbox_IsEmpty(Mailbox_t handle);
01042
01043 //-----
01044 // Condition Variables
01050 void ConditionVariable_Init(ConditionVariable_t handle);
01051
01058 void ConditionVariable_Wait(ConditionVariable_t handle, Mutex_t hMutex_);
01059
01065 void ConditionVariable_Signal(ConditionVariable_t handle);
01066
01072 void ConditionVariable_Broadcast(ConditionVariable_t handle);
01073
01082 bool ConditionVariable_TimedWait(ConditionVariable_t handle, Mutex_t hMutex_,
uint32_t u32WaitTimeMS_);
01083
01084 //-----
01085 // Reader-writer locks
01091 void ReaderWriterLock_Init(ReaderWriterLock_t handle);
01092
01098 void ReaderWriterLock_AcquireReader(ReaderWriterLock_t handle);
01099
01105 void ReaderWriterLock_ReleaseReader(ReaderWriterLock_t handle);
01106
01112 void ReaderWriterLock_AcquireWriter(ReaderWriterLock_t handle);
01113
01119 void ReaderWriterLock_ReleaseWriter(ReaderWriterLock_t handle);
01120
01128 bool ReaderWriterLock_TimedAcquireWriter(ReaderWriterLock_t handle,
uint32_t u32TimeoutMs_);
01129
01137 bool ReaderWriterLock_TimedAcquireReader(ReaderWriterLock_t handle,
uint32_t u32TimeoutMs_);
01138
01144 void Kernel_SetDebugPrintFunction(
kernel_debug_print_t pfPrintFunction_);
01145
01151 void Kernel_DebugPrint(const char* szString_);
01152
01153 #if defined(__cplusplus)
01154 }
01155 #endif

```

20.7 /home/moslevin/projects/m3-repo/kernel/lib/memutil/memutil.cpp File Reference

Implementation of memory, string, and conversion routines.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "memutil.h"
#include "kerneldebug.h"
```

Namespaces

- [Mark3](#)

20.7.1 Detailed Description

Implementation of memory, string, and conversion routines.

Definition in file [memutil.cpp](#).

20.8 memutil.cpp

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00021 #include "kerneltypes.h"
00022 #include "mark3cfg.h"
00023 #include "memutil.h"
00024 #include "kerneldebug.h"
00025
00026 namespace Mark3
00027 {
00028 //-----
00029 void MemUtil::DecimalToHex(uint8_t u8Data_, char* szText_)
00030 {
00031     uint8_t u8Tmp = u8Data_;
00032     uint8_t u8Max;
00033
00034     KERNEL_ASSERT(szText_);
00035
00036     if (u8Tmp >= 0x10) {
00037         u8Max = 2;
00038     } else {
00039         u8Max = 1;
00040     }
00041
00042     u8Tmp = u8Data_;
00043     szText_[u8Max] = 0;
00044     while ((u8Max--) != 0u) {
00045         if ((u8Tmp & 0x0F) <= 9) {
00046             szText_[u8Max] = '0' + (u8Tmp & 0x0F);
00047         } else {
00048             szText_[u8Max] = 'A' + ((u8Tmp & 0x0F) - 10);
00049         }
00050         u8Tmp >>= 4;
00051     }
00052 }
00053
00054 //-----
00055 void MemUtil::DecimalToHex(uint16_t u16Data_, char* szText_)
00056 {
00057     uint16_t u16Tmp = u16Data_;
00058     uint16_t u16Max = 1;
00059     uint16_t u16Compare = 0x0010;
```

```

00060
00061     KERNEL_ASSERT(szText_);
00062
00063     while (u16Data_ > u16Compare && u16Max < 4) {
00064         u16Max++;
00065         u16Compare <= 4;
00066     }
00067
00068     u16Tmp      = u16Data_;
00069     szText_[u16Max] = 0;
00070     while ((u16Max--) != 0u) {
00071         if ((u16Tmp & 0x0F) <= 9) {
00072             szText_[u16Max] = '0' + (u16Tmp & 0x0F);
00073         } else {
00074             szText_[u16Max] = 'A' + ((u16Tmp & 0x0F) - 10);
00075         }
00076         u16Tmp >>= 4;
00077     }
00078 }
00079
00080 //-----
00081 void MemUtil::DecimalToHex(uint32_t u32Data_, char* szText_)
00082 {
00083     uint32_t u32Tmp      = u32Data_;
00084     uint32_t u32Max      = 1;
00085     uint32_t u32Compare = 0x0010;
00086
00087     KERNEL_ASSERT(szText_);
00088
00089     while (u32Data_ > u32Compare && u32Max < 8) {
00090         u32Max++;
00091         u32Compare <= 4;
00092     }
00093
00094     u32Tmp      = u32Data_;
00095     szText_[u32Max] = 0;
00096     while ((u32Max--) != 0u) {
00097         if ((u32Tmp & 0x0F) <= 9) {
00098             szText_[u32Max] = '0' + (u32Tmp & 0x0F);
00099         } else {
00100             szText_[u32Max] = 'A' + ((u32Tmp & 0x0F) - 10);
00101         }
00102         u32Tmp >>= 4;
00103     }
00104 }
00105
00106 //-----
00107 void MemUtil::DecimalToHex(uint64_t u64Data_, char* szText_)
00108 {
00109     uint64_t u64Tmp      = u64Data_;
00110     uint64_t u64Max      = 1;
00111     uint64_t u64Compare = 0x0010;
00112
00113     KERNEL_ASSERT(szText_);
00114
00115     while (u64Data_ > u64Compare && u64Max < 8) {
00116         u64Max++;
00117         u64Compare <= 4;
00118     }
00119
00120     u64Tmp      = u64Data_;
00121     szText_[u64Max] = 0;
00122     while ((u64Max--) != 0u) {
00123         if ((u64Tmp & 0x0F) <= 9) {
00124             szText_[u64Max] = '0' + (u64Tmp & 0x0F);
00125         } else {
00126             szText_[u64Max] = 'A' + ((u64Tmp & 0x0F) - 10);
00127         }
00128         u64Tmp >>= 4;
00129     }
00130 }
00131
00132 //-----
00133 void MemUtil::DecimalToString(uint8_t u8Data_, char* szText_)
00134 {
00135     uint8_t u8Tmp = u8Data_;
00136     uint8_t u8Max;
00137
00138     KERNEL_ASSERT(szText_);
00139
00140     // Find max index to print...
00141     if (u8Data_ >= 100) {
00142         u8Max = 3;
00143     } else if (u8Data_ >= 10) {
00144         u8Max = 2;
00145     } else {
00146         u8Max = 1;

```

```

00147     }
00148
00149     szText_[u8Max] = 0;
00150     while ((u8Max--) != 0u) {
00151         szText_[u8Max] = '0' + (u8Tmp % 10);
00152         u8Tmp /= 10;
00153     }
00154 }
00155
00156 //-----
00157 void MemUtil::DecimalToString(uint16_t u16Data_, char* szText_)
00158 {
00159     uint16_t u16Tmp    = u16Data_;
00160     uint16_t u16Max    = 1;
00161     uint16_t u16Compare = 10;
00162
00163     KERNEL_ASSERT(szText_);
00164
00165     while (u16Data_ >= u16Compare && u16Max < 5) {
00166         u16Compare *= 10;
00167         u16Max++;
00168     }
00169
00170     szText_[u16Max] = 0;
00171     while ((u16Max--) != 0u) {
00172         szText_[u16Max] = '0' + (u16Tmp % 10);
00173         u16Tmp /= 10;
00174     }
00175 }
00176
00177 //-----
00178 void MemUtil::DecimalToString(uint32_t u32Data_, char* szText_)
00179 {
00180     uint32_t u32Tmp    = u32Data_;
00181     uint32_t u32Max    = 1;
00182     uint32_t u32Compare = 10;
00183
00184     KERNEL_ASSERT(szText_);
00185
00186     while (u32Data_ >= u32Compare && u32Max < 12) {
00187         u32Compare *= 10;
00188         u32Max++;
00189     }
00190
00191     szText_[u32Max] = 0;
00192     while ((u32Max--) != 0u) {
00193         szText_[u32Max] = '0' + (u32Tmp % 10);
00194         u32Tmp /= 10;
00195     }
00196 }
00197
00198 //-----
00199 void MemUtil::DecimalToString(uint64_t u64Data_, char* szText_)
00200 {
00201     uint64_t u64Tmp    = u64Data_;
00202     uint64_t u64Max    = 1;
00203     uint64_t u64Compare = 10;
00204
00205     KERNEL_ASSERT(szText_);
00206
00207     while (u64Data_ >= u64Compare && u64Max < 12) {
00208         u64Compare *= 10;
00209         u64Max++;
00210     }
00211
00212     szText_[u64Max] = 0;
00213     while ((u64Max--) != 0u) {
00214         szText_[u64Max] = '0' + (u64Tmp % 10);
00215         u64Tmp /= 10;
00216     }
00217 }
00218
00219 //-----
00220 bool MemUtil::StringToDecimal8(const char* szText_, uint8_t* pu8Out_)
00221 {
00222     uint8_t u8Tmp = 0;
00223     uint8_t u8Len = 0;
00224
00225     for (uint8_t i = 0; i < 4; i++) {
00226         if (szText_[i] == 0) {
00227             if (i == 0) {
00228                 return false;
00229             }
00230             u8Len = i;
00231             break;
00232         }
00233     }

```

```

00234
00235     for (uint8_t i = 0; i < u8Len; i++) {
00236         if ((szText_[i] < '0') || (szText_[i] > '9')) {
00237             return false;
00238         }
00239         u8Tmp *= 10;
00240         u8Tmp += szText_[i] - '0';
00241     }
00242     *pu8Out_ = u8Tmp;
00243
00244     return true;
00245 }
00246
00247 //-----
00248 bool MemUtil::StringToDecimal16(const char* szText_, uint16_t* pu16Out_)
00249 {
00250     uint16_t u16Tmp = 0;
00251     uint16_t u16Len = 0;
00252
00253     for (uint8_t i = 0; i < 6; i++) {
00254         if (szText_[i] == 0) {
00255             if (i == 0) {
00256                 return false;
00257             }
00258             u16Len = i;
00259             break;
00260         }
00261     }
00262
00263     for (uint8_t i = 0; i < u16Len; i++) {
00264         if ((szText_[i] < '0') || (szText_[i] > '9')) {
00265             return false;
00266         }
00267         u16Tmp *= 10;
00268         u16Tmp += szText_[i] - '0';
00269     }
00270     *pu16Out_ = u16Tmp;
00271
00272     return true;
00273 }
00274
00275 //-----
00276 bool MemUtil::StringToDecimal32(const char* szText_, uint32_t* pu32Out_)
00277 {
00278     uint32_t u32Tmp = 0;
00279     uint32_t u32Len = 0;
00280
00281     for (uint8_t i = 0; i < 11; i++) {
00282         if (szText_[i] == 0) {
00283             if (i == 0) {
00284                 return false;
00285             }
00286             u32Len = i;
00287             break;
00288         }
00289     }
00290
00291     for (uint8_t i = 0; i < u32Len; i++) {
00292         if ((szText_[i] < '0') || (szText_[i] > '9')) {
00293             return false;
00294         }
00295         u32Tmp *= 10;
00296         u32Tmp += szText_[i] - '0';
00297     }
00298     *pu32Out_ = u32Tmp;
00299
00300     return true;
00301 }
00302
00303 //-----
00304 bool MemUtil::StringToDecimal64(const char* szText_, uint64_t* pu64Out_)
00305 {
00306     uint64_t u64Tmp = 0;
00307     uint64_t u64Len = 0;
00308
00309     for (uint8_t i = 0; i < 21; i++) {
00310         if (szText_[i] == 0) {
00311             if (i == 0) {
00312                 return false;
00313             }
00314             u64Len = i;
00315             break;
00316         }
00317     }
00318
00319     for (uint8_t i = 0; i < u64Len; i++) {
00320         if ((szText_[i] < '0') || (szText_[i] > '9')) {

```



```

00321         return false;
00322     }
00323     u64Tmp *= 10;
00324     u64Tmp += szText_[i] - '0';
00325 }
00326 *pu64Out_ = u64Tmp;
00327
00328 return true;
00329 }
00330
00331 //-----
00332 // Basic checksum routines
00333 uint8_t MemUtil::Checksum8(const void* pvSrc_, uint16_t ul6Len_)
00334 {
00335     uint8_t u8Ret = 0;
00336     uint8_t* pcData = (uint8_t*)pvSrc_;
00337     KERNEL_ASSERT(pvSrc_);
00338
00339     // 8-bit CRC, computed byte at a time
00340     while ((ul6Len_--) != 0u) { u8Ret += *pcData++; }
00341     return u8Ret;
00342 }
00343
00344 //-----
00345 //-----
00346 uint16_t MemUtil::Checksum16(const void* pvSrc_, uint16_t ul6Len_)
00347 {
00348     uint16_t ul6Ret = 0;
00349     uint8_t* pcData = (uint8_t*)pvSrc_;
00350
00351     KERNEL_ASSERT(pvSrc_);
00352
00353     // 16-bit CRC, computed byte at a time
00354     while ((ul6Len_--) != 0u) { ul6Ret += *pcData++; }
00355     return ul6Ret;
00356 }
00357
00358 //-----
00359 // Basic string routines
00360 uint16_t MemUtil::StringLength(const char* szStr_)
00361 {
00362     uint8_t* pcData = (uint8_t*)szStr_;
00363     uint16_t ul6Len = 0;
00364
00365     KERNEL_ASSERT(szStr_);
00366
00367     while (*pcData++ != 0u) { ul6Len++; }
00368     return ul6Len;
00369 }
00370
00371 //-----
00372 bool MemUtil::CompareStrings(const char* szStr1_, const char* szStr2_)
00373 {
00374     char* szTmp1 = (char*)szStr1_;
00375     char* szTmp2 = (char*)szStr2_;
00376
00377     KERNEL_ASSERT(szStr1_);
00378     KERNEL_ASSERT(szStr2_);
00379
00380     while ((*szTmp1 != 0) && (*szTmp2 != 0)) {
00381         if (*szTmp1++ != *szTmp2++) {
00382             return false;
00383         }
00384     }
00385
00386     // Both terminate at the same length
00387     return ((*szTmp1) == 0) && ((*szTmp2) == 0);
00388 }
00389
00390 //-----
00391 bool MemUtil::CompareStrings(const char* szStr1_, const char* szStr2_, uint16_t
ul6Length_)
00392 {
00393     char* szTmp1 = (char*)szStr1_;
00394     char* szTmp2 = (char*)szStr2_;
00395
00396     while ((*szTmp1 != 0) && (*szTmp2 != 0) && (ul6Length_ != 0u)) {
00397         if (*szTmp1++ != *szTmp2++) {
00398             return false;
00399         }
00400         ul6Length_--;
00401     }
00402
00403     // Both terminate at the same length
00404     return ((*szTmp1) == 0) && ((*szTmp2) == 0) || (ul6Length_ == 0u);
00405 }
00406

```

```

00407 //-----
00408 void MemUtil::CopyMemory(void* pvDst_, const void* pvSrc_, uint16_t ul6Len_)
00409 {
00410     char* szDst = (char*)pvDst_;
00411     char* szSrc = (char*)pvSrc_;
00412
00413     KERNEL_ASSERT(pvDst_);
00414     KERNEL_ASSERT(pvSrc_);
00415
00416     // Run through the strings verifying that each character matches
00417     // and the lengths are the same.
00418     while ((ul6Len_--) != 0u) { *szDst++ = *szSrc++; }
00419 }
00420
00421 //-----
00422 void MemUtil::CopyString(char* szDst_, const char* szSrc_)
00423 {
00424     char* szDst = (char*)szDst_;
00425     char* szSrc = (char*)szSrc_;
00426
00427     KERNEL_ASSERT(szDst_);
00428     KERNEL_ASSERT(szSrc_);
00429
00430     // Run through the strings verifying that each character matches
00431     // and the lengths are the same.
00432     while (*szSrc != 0) { *szDst++ = *szSrc++; }
00433 }
00434
00435 //-----
00436 int16_t MemUtil::StringSearch(const char* szBuffer_, const char* szPattern_)
00437 {
00438     char* szTmpPat = (char*)szPattern_;
00439     int16_t il6Idx = 0;
00440     int16_t il6Start;
00441     KERNEL_ASSERT(szBuffer_);
00442     KERNEL_ASSERT(szPattern_);
00443
00444     // Run through the big buffer looking for a match of the pattern
00445     while (szBuffer_[il6Idx] != 0) {
00446         // Reload the pattern
00447         il6Start = il6Idx;
00448         szTmpPat = (char*)szPattern_;
00449         while ((*szTmpPat != 0) && (szBuffer_[il6Idx] != 0)) {
00450             if (*szTmpPat != szBuffer_[il6Idx]) {
00451                 break;
00452             }
00453             szTmpPat++;
00454             il6Idx++;
00455         }
00456         // Made it to the end of the pattern, it's a match.
00457         if (*szTmpPat == '\0') {
00458             return il6Start;
00459         }
00460         il6Idx++;
00461     }
00462
00463     return -1;
00464 }
00465
00466 //-----
00467 bool MemUtil::CompareMemory(const void* pvMem1_, const void* pvMem2_, uint16_t
ul6Len_)
00468 {
00469     char* szTmp1 = (char*)pvMem1_;
00470     char* szTmp2 = (char*)pvMem2_;
00471
00472     KERNEL_ASSERT(pvMem1_);
00473     KERNEL_ASSERT(pvMem2_);
00474
00475     // Run through the strings verifying that each character matches
00476     // and the lengths are the same.
00477     while ((ul6Len_--) != 0u) {
00478         if (*szTmp1++ != *szTmp2++) {
00479             return false;
00480         }
00481     }
00482     return true;
00483 }
00484
00485 //-----
00486 void MemUtil::SetMemory(void* pvDst_, uint8_t u8Val_, uint16_t ul6Len_)
00487 {
00488     char* szDst = (char*)pvDst_;
00489
00490     KERNEL_ASSERT(pvDst_);
00491
00492     while ((ul6Len_--) != 0u) { *szDst++ = u8Val_; }

```

```

00493 }
00494
00495 //-----
00496 uint8_t MemUtil::Tokenize(const char* szBuffer_, Token_t* pastTokens_, uint8_t
    u8MaxTokens_)
00497 {
00498     uint8_t u8CurrArg = 0;
00499     uint8_t u8LastArg = 0;
00500     uint8_t i = 0;
00501
00502     bool bEscape = false;
00503
00504     KERNEL_ASSERT(szBuffer_);
00505     KERNEL_ASSERT(pastTokens_);
00506
00507     while (szBuffer_[i] != 0) {
00508         //-- Handle unescaped quotes
00509         if (szBuffer_[i] == '"' || szBuffer_[i] == '\\') {
00510             bEscape = !bEscape;
00511             i++;
00512             continue;
00513         }
00514
00515         //-- Handle all escaped chars - by ignoring them
00516         if (szBuffer_[i] == '\\') {
00517             i++;
00518             if (szBuffer_[i] != 0) {
00519                 i++;
00520             }
00521             continue;
00522         }
00523
00524         //-- Process chars based on current escape characters
00525         if (bEscape) {
00526             // Everything within the quote is treated as literal, but escaped chars are still treated the
same
00527             i++;
00528             continue;
00529         }
00530
00531         //-- Non-escaped case
00532         if (szBuffer_[i] != ' ' || szBuffer_[i] != '\0') {
00533             i++;
00534             continue;
00535         }
00536
00537         pastTokens_[u8CurrArg].pcToken = &(szBuffer_[u8LastArg]);
00538         pastTokens_[u8CurrArg].u8Len = i - u8LastArg;
00539         u8CurrArg++;
00540         if (u8CurrArg >= u8MaxTokens_) {
00541             return u8MaxTokens_;
00542         }
00543
00544         i++;
00545         while (szBuffer_[i] == ' ' || szBuffer_[i] == '\0') { i++; }
00546
00547         u8LastArg = i;
00548     }
00549     if ((i != 0) && (szBuffer_[i] == 0) && ((i - u8LastArg) != 0)) {
00550         pastTokens_[u8CurrArg].pcToken = &(szBuffer_[u8LastArg]);
00551         pastTokens_[u8CurrArg].u8Len = i - u8LastArg;
00552         u8CurrArg++;
00553     }
00554     return u8CurrArg;
00555 }
00556 } // namespace Mark3

```

20.9 /home/moslevin/projects/m3-repo/kernel/lib/memutil/public/memutil.h File Reference

Utility class containing memory, string, and conversion routines.

```

#include "kerneltypes.h"
#include "mark3cfg.h"

```

Classes

- struct [Mark3::Token_t](#)
Token descriptor struct format.
- class [Mark3::MemUtil](#)
String and Memory manipulation class.

Namespaces

- [Mark3](#)

20.9.1 Detailed Description

Utility class containing memory, string, and conversion routines.

Definition in file [memutil.h](#).

20.10 memutil.h

```

00001  /*=====
00002
00003  |-----|-----|-----|-----|-----|-----|
00004  |  \  /  |  \  /  |  \  /  |  \  /  |  \  /  |
00005  | /  \  | /  \  | /  \  | /  \  | /  \  |
00006  |-----|-----|-----|-----|-----|
00007  |-----|-----|-----|-----|
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  ===== */
00021  #pragma once
00022
00023  #include "kerneltypes.h"
00024  #include "mark3cfg.h"
00025
00026  namespace Mark3
00027  {
00028  //-----
00032  typedef struct {
00033      const char* pcToken;
00034      uint8_t      u8Len;
00035  } Token_t;
00036
00037  //-----
00046  class MemUtil
00047  {
00048  public:
00049      //-----
00059      static void DecimalToHex(uint8_t u8Data_, char* szText_);
00060      static void DecimalToHex(uint16_t u16Data_, char* szText_);
00061      static void DecimalToHex(uint32_t u32Data_, char* szText_);
00062      static void DecimalToHex(uint64_t u64Data_, char* szText_);
00063
00064      //-----
00073      static void DecimalToString(uint8_t u8Data_, char* szText_);
00074      static void DecimalToString(uint16_t u16Data_, char* szText_);
00075      static void DecimalToString(uint32_t u32Data_, char* szText_);
00076      static void DecimalToString(uint64_t u64Data_, char* szText_);
00077
00078      //-----
00089      static bool StringToDecimal8(const char* szText_, uint8_t* pu8Out_);
00090      static bool StringToDecimal16(const char* szText_, uint16_t* pu16Out_);
00091      static bool StringToDecimal32(const char* szText_, uint32_t* pu32Out_);
00092      static bool StringToDecimal64(const char* szText_, uint64_t* pu64Out_);
00093
00094      //-----
00104      static uint8_t Checksum8(const void* pvSrc_, uint16_t u16Len_);

```

```

00105
00106 //-----
00116 static uint16_t Checksum16(const void* pvSrc_, uint16_t ul6Len_);
00117
00118 //-----
00128 static uint16_t StringLength(const char* szStr_);
00129
00130 //-----
00140 static bool CompareStrings(const char* szStr1_, const char* szStr2_);
00141 static bool CompareStrings(const char* szStr1_, const char* szStr2_, uint16_t ul6Length_)
00142 ;
00143 //-----
00153 static void CopyMemory(void* pvDst_, const void* pvSrc_, uint16_t ul6Len_);
00154
00155 //-----
00164 static void CopyString(char* szDst_, const char* szSrc_);
00165
00166 //-----
00176 static int16_t StringSearch(const char* szBuffer_, const char* szPattern_);
00177
00178 //-----
00190 static bool CompareMemory(const void* pvMem1_, const void* pvMem2_, uint16_t ul6Len_);
00191
00192 //-----
00202 static void SetMemory(void* pvDst_, uint8_t u8Val_, uint16_t ul6Len_);
00203
00204 //-----
00214 static uint8_t Tokenize(const char* szBuffer_, Token_t* pastTokens_, uint8_t
    u8MaxTokens_);
00215 };
00216 } // namespace Mark3

```

20.11 /home/moslevin/projects/m3-repo/kernel/lib/streamer/public/streamer.h File Reference

Thread/Interrupt-safe byte-based data streaming.

```

#include "kerneltypes.h"
#include "mark3.h"

```

Classes

- class [Mark3::Streamer](#)

The *Streamer* class. This class implements a circular byte-buffer with thread and interrupt safe methods for writing-to and reading-from the buffer. Objects of this class type are designed to be shared between threads, or between threads and interrupts.

Namespaces

- [Mark3](#)

20.11.1 Detailed Description

Thread/Interrupt-safe byte-based data streaming.

Definition in file [streamer.h](#).

20.12 streamer.h

```

00001  /*=====
00002
00003  _____
00004  | \ / | | | \ / | | | \ / | | | \ / | | |
00005  | \ / | | | \ / | | | \ / | | | \ / | | |
00006  | \ / | | | \ / | | | \ / | | | \ / | | |
00007  | \ / | | | \ / | | | \ / | | | \ / | | |
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  ===== */
00020  #include "kerneltypes.h"
00021  #include "mark3.h"
00022
00023  #pragma once
00024
00025  namespace Mark3
00026  {
00027  //-----
00035  class Streamer
00036  {
00037  public:
00046      void Init(uint8_t* pau8Buffer_, uint16_t ul6Size_);
00047
00056      bool Read(uint8_t* pu8Data_);
00057
00067      uint16_t Read(uint8_t* pu8Data_, uint16_t ul6Len_);
00068
00077      bool Write(uint8_t u8Data_);
00078
00087      uint16_t Write(uint8_t* pu8Data_, uint16_t ul6Len_);
00088
00109      bool Claim(uint8_t** pu8Addr_);
00110
00120      void Lock(uint8_t* pu8LockAddr_);
00121
00128      void Unlock(void);
00129
00134      uint16_t GetAvailable(void) { return m_ul6Size; }
00139      bool CanRead(void);
00140
00145      bool CanWrite(void);
00146
00151      bool IsEmpty(void);
00152
00153  private:
00154      uint8_t* m_pau8Buffer;
00155      uint8_t* m_pu8LockAddr;
00156
00157      uint16_t m_ul6Size;
00158      uint16_t m_ul6Avail;
00159      uint16_t m_ul6Head;
00160      uint16_t m_ul6Tail;
00161  };
00162  } // namespace Mark3

```

20.13 /home/moslevin/projects/m3-repo/kernel/lib/streamer/streamer.cpp File Reference

Thread/Interrupt-safe byte-based data streaming.

```

#include "kerneltypes.h"
#include "mark3.h"
#include "streamer.h"

```

Namespaces

- [Mark3](#)


```

00078     const auto cs = CriticalGuard{};
00079     ul6Allocated = m_ul6Size - m_ul6Avail;
00080
00081     if (ul6Allocated > ul6Len_) {
00082         ul6ToRead = ul6Len_;
00083     } else {
00084         ul6ToRead = ul6Allocated;
00085     }
00086
00087     ul6PreWrap = m_ul6Size - m_ul6Tail;
00088
00089     pu8Src = &m_pau8Buffer[m_ul6Tail];
00090     pu8Dst = pu8Data_;
00091
00092     Lock(pu8Src);
00093 } // end critical section
00094
00095 if (ul6Allocated != 0u) {
00096     if (ul6PreWrap >= ul6ToRead) {
00097         for (uint16_t i = 0; i < ul6ToRead; i++) { *pu8Dst++ = *pu8Src++; }
00098     } else {
00099         for (uint16_t i = 0; i < ul6PreWrap; i++) { *pu8Dst++ = *pu8Src++; }
00100         pu8Src = m_pau8Buffer;
00101         for (uint16_t i = ul6PreWrap; i < ul6ToRead; i++) { *pu8Dst++ = *pu8Src++; }
00102     }
00103 }
00104
00105 { // Begin critical section
00106     const auto cs = CriticalGuard{};
00107     m_ul6Avail += ul6ToRead;
00108     if (ul6PreWrap >= ul6ToRead) {
00109         m_ul6Tail += ul6ToRead;
00110     } else {
00111         m_ul6Tail += ul6ToRead - m_ul6Size;
00112     }
00113
00114     Unlock();
00115 } // end critical section
00116 return ul6ToRead;
00117 }
00118
00119 //-----
00120 bool Streamer::Write(uint8_t u8Data_)
00121 {
00122     auto rc = true;
00123
00124     const auto cs = CriticalGuard{};
00125     if (m_ul6Avail == 0u) {
00126         rc = false;
00127     } else {
00128         if (m_pu8LockAddr == &m_pau8Buffer[m_ul6Head]) {
00129             rc = false;
00130         } else {
00131             m_pau8Buffer[m_ul6Head] = u8Data_;
00132             m_ul6Head++;
00133             if (m_ul6Head >= m_ul6Size) {
00134                 m_ul6Head = 0;
00135             }
00136             m_ul6Avail--;
00137         }
00138     }
00139
00140     return rc;
00141 }
00142
00143 //-----
00144 uint16_t Streamer::Write(uint8_t* pu8Data_, uint16_t ul6Len_)
00145 {
00146     uint16_t ul6ToWrite;
00147
00148     // Bail if the buffer is currently locked.
00149     if (m_pu8LockAddr != 0) {
00150         return 0;
00151     }
00152
00153     // Update the buffer metadata in a critical section, and lock it so that
00154     // we can safely write to it with interrupts enabled.
00155
00156     uint16_t ul6PreWrap;
00157     uint8_t* pu8Src;
00158     uint8_t* pu8Dst;
00159
00160     { // Begin critical section
00161         const auto cs = CriticalGuard{};
00162         if (m_ul6Avail > ul6Len_) {
00163             ul6ToWrite = ul6Len_;
00164         } else {

```



```

00165         ul6ToWrite = m_ul6Avail;
00166     }
00167
00168     ul6PreWrap = m_ul6Size - m_ul6Head;
00169
00170     pu8Src = pu8Data_;
00171     pu8Dst = &m_pau8Buffer[m_ul6Head];
00172
00173     m_ul6Avail -= ul6ToWrite;
00174
00175     if (ul6PreWrap >= ul6ToWrite) {
00176         m_ul6Head += ul6ToWrite;
00177     } else {
00178         m_ul6Head += ul6ToWrite - m_ul6Size;
00179     }
00180
00181     Lock(pu8Dst);
00182 } // End critical section
00183
00184 // Perform the buffer writes with interrupts enabled, buffers locked.
00185 if (ul6ToWrite != 0u) {
00186     if (ul6PreWrap >= ul6ToWrite) {
00187         for (uint16_t i = 0; i < ul6ToWrite; i++) { *pu8Dst++ = *pu8Src++; }
00188     } else {
00189         for (uint16_t i = 0; i < ul6PreWrap; i++) { *pu8Dst++ = *pu8Src++; }
00190         pu8Dst = m_pau8Buffer;
00191         for (uint16_t i = ul6PreWrap; i < ul6ToWrite; i++) { *pu8Dst++ = *pu8Src++; }
00192     }
00193 }
00194
00195 Unlock();
00196 return ul6ToWrite;
00197 }
00198
00199 //-----
00200 bool Streamer::CanRead(void)
00201 {
00202     auto bRc = true;
00203     const auto cs = CriticalGuard{};
00204     if (m_ul6Avail == m_ul6Size) {
00205         bRc = false;
00206     }
00207     return bRc;
00208 }
00209
00210 //-----
00211 bool Streamer::CanWrite(void)
00212 {
00213     auto bRc = false;
00214
00215     const auto cs = CriticalGuard{};
00216     if (m_ul6Avail != 0u) {
00217         bRc = true;
00218     }
00219
00220     return bRc;
00221 }
00222
00223 //-----
00224 bool Streamer::Claim(uint8_t** pu8Addr_)
00225 {
00226     auto rc = true;
00227
00228     const auto cs = CriticalGuard{};
00229     if (m_ul6Avail == 0u) {
00230         rc = false;
00231     } else {
00232         if (m_pu8LockAddr == &m_pau8Buffer[m_ul6Head]) {
00233             rc = false;
00234         } else {
00235             *pu8Addr_ = &m_pau8Buffer[m_ul6Head];
00236             if (m_pu8LockAddr == 0) {
00237                 m_pu8LockAddr = &m_pau8Buffer[m_ul6Head];
00238             }
00239             m_ul6Head++;
00240             if (m_ul6Head >= m_ul6Size) {
00241                 m_ul6Head = 0;
00242             }
00243             m_ul6Avail--;
00244         }
00245     }
00246     return rc;
00247 }
00248
00249 //-----
00250 void Streamer::Lock(uint8_t* pu8LockAddr_)
00251 {

```



```

00029 {
00030 //-----
00031 void KernelSWI::Config(void)
00032 {
00033     PORTB &= ~0x04; // Clear INT2
00034     DDRB |= 0x04; // Set PortB, bit 2 (INT2) As Output
00035     EICRA |= (1 << ISC20) | (1 << ISC21); // Rising edge on INT2
00036 }
00037
00038 //-----
00039 void KernelSWI::Start(void)
00040 {
00041     EIFR &= ~(1 << INTF2); // Clear any pending interrupts on INT2
00042     EIMSK |= (1 << INT2); // Enable INT2 interrupt (as int32_t as I-bit is set)
00043 }
00044
00045 //-----
00046 void KernelSWI::Trigger(void)
00047 {
00048     // if(Thread_IsSchedulerEnabled())
00049     {
00050         PORTB &= ~0x04;
00051         PORTB |= 0x04;
00052     }
00053 }
00054 } // namespace Mark3

```

20.17 /home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/kerneltimer.cpp File Reference

Kernel Timer Implementation for ATmega1284p.

```

#include "mark3.h"
#include <avr/common.h>
#include <avr/io.h>
#include <avr/interrupt.h>

```

Namespaces

- [Mark3](#)

Macros

- `#define TCCR1B_INIT ((1 << WGM12) | (1 << CS12))`
- `#define TIMER_IMSK (1 << OCIE1A)`
- `#define TIMER_IFR (1 << OCF1A)`

Functions

- static void [Mark3::KernelTimer_Task](#) (void *unused)
- [ISR](#) (TIMER1_COMPA_vect)

20.17.1 Detailed Description

Kernel Timer Implementation for ATmega1284p.

Definition in file [kerneltimer.cpp](#).

20.17.2 Macro Definition Documentation

20.17.2.1 TCCR1B_INIT

```
#define TCCR1B_INIT ((1 << WGM12) | (1 << CS12))
```

Definition at line 27 of file [kerneltimer.cpp](#).

20.17.2.2 TIMER_IFR

```
#define TIMER_IFR (1 << OCF1A)
```

Definition at line 29 of file [kerneltimer.cpp](#).

20.17.2.3 TIMER_IMSK

```
#define TIMER_IMSK (1 << OCIE1A)
```

Definition at line 28 of file [kerneltimer.cpp](#).

20.17.3 Function Documentation

20.17.3.1 ISR()

```
ISR (
    TIMER1_COMPA_vect )
```

Definition at line 103 of file [kerneltimer.cpp](#).

```

00001  /*-----
00002  |-----|-----|-----|-----|-----|-----|
00003  |-----|-----|-----|-----|-----|-----|
00004  |-----|-----|-----|-----|-----|-----|
00005  |-----|-----|-----|-----|-----|-----|
00006  |-----|-----|-----|-----|-----|-----|
00007  |-----|-----|-----|-----|-----|-----|
00008  |-----|-----|-----|-----|-----|-----|
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  =====*/
00021  #include "mark3.h"
00022
00023  #include <avr/common.h>
00024  #include <avr/io.h>
00025  #include <avr/interrupt.h>
00026
00027  #define TCCR1B_INIT ((1 << WGM12) | (1 << CS12))
00028  #define TIMER_IMSK (1 << OCIE1A)
00029  #define TIMER_IFR (1 << OCF1A)
00030
00031  namespace
00032  {
00033      using namespace Mark3;
00034      //-----
00035      // Static objects implementing the timer thread and its synchronization objects
00036      Thread      s_clTimerThread;
00037      K_WORD      s_clTimerThreadStack[PORT_KERNEL_TIMERS_THREAD_STACK];
00038      Semaphore    s_clTimerSemaphore;
00039  }
00040
00041  namespace Mark3
00042  {
00043      //-----
00044      static void KernelTimer_Task(void* unused)
00045      {
00046          (void)unused;
00047          while (1) {
00048              s_clTimerSemaphore.Pend();
00049              #if KERNEL_ROUND_ROBIN
00050                  Quantum::SetInTimer();
00051              #endif // #if KERNEL_ROUND_ROBIN
00052                  TimerScheduler::Process();
00053              #if KERNEL_ROUND_ROBIN
00054                  Quantum::ClearInTimer();
00055              #endif // #if KERNEL_ROUND_ROBIN
00056          }
00057      }
00058
00059      //-----
00060      void KernelTimer::Config(void)
00061      {
00062          TCCR1B = TCCR1B_INIT;
00063          s_clTimerSemaphore.Init(0, 1);
00064          s_clTimerThread.Init(s_clTimerThreadStack,
00065                              sizeof(s_clTimerThreadStack) / sizeof(K_WORD),
00066                              KERNEL_TIMERS_THREAD_PRIORITY,
00067                              KernelTimer_Task,
00068                              0);
00069          #if KERNEL_ROUND_ROBIN
00070              Quantum::SetTimerThread(&s_clTimerThread);
00071          #endif // #if KERNEL_ROUND_ROBIN
00072          s_clTimerThread.Start();
00073      }
00074
00075      //-----
00076      void KernelTimer::Start(void)
00077      {
00078          TCCR1B = ((1 << WGM12) | (1 << CS11) | (1 << CS10));
00079          OCR1A = ((PORT_SYSTEM_FREQ / 1000) / 64);
00080          TCNT1 = 0;
00081          TIFR1 &= ~TIMER_IFR;
00082          TIMSK1 |= TIMER_IMSK;
00083      }
00084
00085      //-----
00086      void KernelTimer::Stop(void)
00087      {
00088          TIFR1 &= ~TIMER_IFR;
00089          TIMSK1 &= ~TIMER_IMSK;
00090          TCCR1B &= ~(1 << CS12); // Disable count...
00091          TCNT1 = 0;

```

```

00092     OCR1A = 0;
00093 }
00094 } // namespace Mark3
00095
00096 //-----
00101 //-----
00102 using namespace Mark3;
00103 ISR(TIMER1_COMPA_vect)
00104 {
00105     Kernel::Tick();
00106     s_clTimerSemaphore.Post();
00107 }

```

20.19 /home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/portcfg.h

File Reference

[Mark3](#) Port Configuration.

Macros

- `#define KERNEL_NUM_PRIORITIES (16)`
- `#define KERNEL_TIMERS_THREAD_PRIORITY (KERNEL_NUM_PRIORITIES - 1)`
- `#define THREAD_QUANTUM_DEFAULT (4)`
- `#define KERNEL_STACK_GUARD_DEFAULT (32)`
- `#define PORT_COROUTINE_PRIORITIES (8)`
- `#define AVR (1)`
- `#define K_WORD uint8_t`
Size of a data word.
- `#define K_ADDR uint16_t`
Size of an address (pointer size)
- `#define K_INT int32_t`
- `#define PORT_PRIO_TYPE uint8_t`
Type used for bitmap in the PriorityMap class.
- `#define PORT_PRIO_MAP_WORD_SIZE (1)`
size of PORT_PRIO_TYPE in bytes
- `#define PORT_SYSTEM_FREQ ((uint32_t)16000000)`
CPU Frequency in Hz.
- `#define PORT_TIMER_FREQ ((uint32_t)(PORT_SYSTEM_FREQ / 1000))`
Fixed timer interrupt frequency.
- `#define PORT_KERNEL_DEFAULT_STACK_SIZE ((K_ADDR)384)`
- `#define PORT_KERNEL_TIMERS_THREAD_STACK ((K_ADDR)384)`
- `#define PORT_TIMER_COUNT_TYPE uint16_t`
Timer counter type.
- `#define PORT_MIN_TIMER_TICKS (0)`
- `#define PORT_OVERLOAD_NEW (1)`
- `#define PORT_STACK_GROWS_DOWN (1)`
- `#define PORT_USE_HW_CLZ (1)`

20.19.1 Detailed Description

[Mark3](#) Port Configuration.

This file is used to configure the kernel for your specific target CPU in order to provide the optimal set of features for a given use case.

!! NOTE: This file must ONLY be included from [mark3cfg.h](#)

Definition in file [portcfg.h](#).

20.19.2 Macro Definition Documentation

20.19.2.1 AVR

```
#define AVR (1)
```

Define a macro indicating the CPU architecture for which this port belongs.

This may also be set by the toolchain, but that's not guaranteed.

Definition at line 55 of file [portcfg.h](#).

20.19.2.2 K_ADDR

```
#define K_ADDR uint16_t
```

Size of an address (pointer size)

Examples:

[lab9_dynamic_threads/main.cpp](#).

Definition at line 63 of file [portcfg.h](#).

20.19.2.3 K_INT

```
#define K_INT int32_t
```

Definition at line 64 of file [portcfg.h](#).

20.19.2.4 K_WORD

```
#define K_WORD uint8_t
```

Size of a data word.

Define types that map to the CPU Architecture's default data-word and address size.

Examples:

[lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_coroutines/main.cpp](#), [lab3_round_robin/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab5_mutexes/main.cpp](#), [lab6_timers/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 62 of file [portcfg.h](#).

20.19.2.5 KERNEL_NUM_PRIORITIES

```
#define KERNEL_NUM_PRIORITIES (16)
```

Define the number of thread priorities that the kernel's scheduler will support. The number of thread priorities is limited only by the memory of the host CPU, as a ThreadList object is statically-allocated for each thread priority.

In practice, systems rarely need more than 32 priority levels, with the most complex having the capacity for 256.

Definition at line 35 of file [portcfg.h](#).

20.19.2.6 KERNEL_STACK_GUARD_DEFAULT

```
#define KERNEL_STACK_GUARD_DEFAULT (32)
```

Definition at line 41 of file [portcfg.h](#).

20.19.2.7 KERNEL_TIMERS_THREAD_PRIORITY

```
#define KERNEL_TIMERS_THREAD_PRIORITY (KERNEL_NUM_PRIORITIES - 1)
```

Definition at line 37 of file [portcfg.h](#).

20.19.2.8 PORT_COROUTINE_PRIORITIES

```
#define PORT_COROUTINE_PRIORITIES (8)
```

Set the number of priorities supported by the coroutine scheduler. The number of coroutine priorities is limited by the memory of the host CPU.

Definition at line 47 of file [portcfg.h](#).

20.19.2.9 PORT_KERNEL_DEFAULT_STACK_SIZE

```
#define PORT_KERNEL_DEFAULT_STACK_SIZE ((K_ADDR) 384)
```

Define the default/minimum size of a thread stack

Examples:

[lab10_notifications/main.cpp](#), [lab11_mailboxes/main.cpp](#), [lab1_kernel_setup/main.cpp](#), [lab2_coroutines/main.cpp](#), [lab3_round_robin/main.cpp](#), [lab4_semaphores/main.cpp](#), [lab5_mutexes/main.cpp](#), [lab6_timers/main.cpp](#), [lab7_events/main.cpp](#), [lab8_messages/main.cpp](#), and [lab9_dynamic_threads/main.cpp](#).

Definition at line 96 of file [portcfg.h](#).

20.19.2.10 PORT_KERNEL_TIMERS_THREAD_STACK

```
#define PORT_KERNEL_TIMERS_THREAD_STACK ((K_ADDR) 384)
```

Define the size of the kernel-timer thread stack (if one is configured)

Definition at line 101 of file [portcfg.h](#).

20.19.2.11 PORT_MIN_TIMER_TICKS

```
#define PORT_MIN_TIMER_TICKS (0)
```

Minimum number of timer ticks for any delay or sleep, required to ensure that a timer cannot be initialized to a negative value.

Definition at line 112 of file [portcfg.h](#).

20.19.2.12 PORT_OVERLOAD_NEW

```
#define PORT_OVERLOAD_NEW (1)
```

Set this to 1 to overload the system's New/Free functions with the kernel's allocator functions. A user must configure the Kernel's allocator functions to point to a real heap implementation backed with real memory in order to use dynamic object creation.

Definition at line 119 of file [portcfg.h](#).

20.19.2.13 PORT_PRIO_MAP_WORD_SIZE

```
#define PORT_PRIO_MAP_WORD_SIZE (1)
```

size of PORT_PRIO_TYPE in bytes

Definition at line 74 of file [portcfg.h](#).

20.19.2.14 PORT_PRIO_TYPE

```
#define PORT_PRIO_TYPE uint8_t
```

Type used for bitmap in the PriorityMap class.

Set a base datatype used to represent each element of the scheduler's priority bitmap.

PORT_PRIO_MAP_WORD_SIZE should map to the *size* of an element of type PORT_PROI_TYPE.

Definition at line 73 of file [portcfg.h](#).

20.19.2.15 PORT_STACK_GROWS_DOWN

```
#define PORT_STACK_GROWS_DOWN (1)
```

Set this to 1 if the stack grows down in the target architecture, or 0 if the stack grows up

Definition at line 124 of file [portcfg.h](#).

20.19.2.16 PORT_SYSTEM_FREQ

```
#define PORT_SYSTEM_FREQ ((uint32_t)16000000)
```

CPU Frequency in Hz.

Define the running CPU frequency. This may be an integer constant, or an alias for another variable which holds the CPU's current running frequency.

Definition at line 81 of file [portcfg.h](#).

20.19.2.17 PORT_TIMER_COUNT_TYPE

```
#define PORT_TIMER_COUNT_TYPE uint16_t
```

Timer counter type.

Define the native type corresponding to the kernel timer hardware's counter register.

Definition at line 106 of file [portcfg.h](#).

20.19.2.18 PORT_TIMER_FREQ

```
#define PORT_TIMER_FREQ ((uint32_t)(PORT_SYSTEM_FREQ / 1000))
```

Fixed timer interrupt frequency.

Set the timer frequency. If running in tickless mode, this is simply the frequency at which the free-running kernel timer increments.

In tick-based mode, this is the frequency at which the fixed-frequency kernel tick interrupt occurs.

Definition at line 91 of file [portcfg.h](#).

20.19.2.19 PORT_USE_HW_CLZ

```
#define PORT_USE_HW_CLZ (1)
```

Set this to 1 if the target CPU/toolchain supports an optimized Count-leading-zeros instruction, or count-leading-zeros intrinsic. If such functionality is not available, a general-purpose implementation will be used.

Definition at line 131 of file [portcfg.h](#).

20.19.2.20 THREAD_QUANTUM_DEFAULT

```
#define THREAD_QUANTUM_DEFAULT (4)
```

Definition at line 39 of file [portcfg.h](#).

20.20 portcfg.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00024 #pragma once
00025
00035 #define KERNEL_NUM_PRIORITIES (16)
00036
00037 #define KERNEL_TIMERS_THREAD_PRIORITY (KERNEL_NUM_PRIORITIES - 1)
00038
00039 #define THREAD_QUANTUM_DEFAULT (4)
00040
00041 #define KERNEL_STACK_GUARD_DEFAULT (32) // words
00042
00047 #define PORT_COROUTINE_PRIORITIES (8)
00048
00054 #ifndef AVR
00055 #define AVR (1)
00056 #endif
00057
00062 #define K_WORD uint8_t
00063 #define K_ADDR uint16_t
00064 #define K_INT int32_t
00065
00073 #define PORT_PRIO_TYPE uint8_t
00074 #define PORT_PRIO_MAP_WORD_SIZE (1)
00075
00076
00080 #if !defined(PORT_SYSTEM_FREQ)
00081 #define PORT_SYSTEM_FREQ ((uint32_t)16000000)
00082 #endif
00083
00091 #define PORT_TIMER_FREQ ((uint32_t)(PORT_SYSTEM_FREQ / 1000))
00092
00093
00096 #define PORT_KERNEL_DEFAULT_STACK_SIZE ((K_ADDR)384)
00097
00101 #define PORT_KERNEL_TIMERS_THREAD_STACK ((K_ADDR)384)
00102
00106 #define PORT_TIMER_COUNT_TYPE uint16_t
00107
00108
00112 #define PORT_MIN_TIMER_TICKS (0)
00113
00119 #define PORT_OVERLOAD_NEW (1)
00120
00124 #define PORT_STACK_GROWS_DOWN (1)
00125
00131 #define PORT_USE_HW_CLZ (1)
```

20.21 /home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/public/threadport.h File Reference

ATMega1284p Multithreading support.

```
#include "portcfg.h"
#include "kerneltypes.h"
#include <avr/io.h>
#include <avr/interrupt.h>
```

Namespaces

- [Mark3](#)

Macros

- `#define ASM(x)` asm volatile(x);
- `#define PORT_TOP_OF_STACK(x, y)` (reinterpret_cast<K_WORD*>(reinterpret_cast<K_ADDR>(x) + (static_cast<K_ADDR>(y) - 1)))
Macro to find the top of a stack given its size and top address.
- `#define PORT_PUSH_TO_STACK(x, y)` *x = y; x--;
Push a value y to the stack pointer x and decrement the stack pointer.
- `#define Thread_SaveContext()`
Save the context of the Thread.
- `#define Thread_RestoreContext()`
Restore the context of the Thread.

Functions

- `uint8_t Mark3::PORT_CLZ` (uint8_t in_)
- `void Mark3::PORT_IRQ_ENABLE` ()
- `void Mark3::PORT_IRQ_DISABLE` ()
- `void Mark3::PORT_CS_ENTER` ()
- `void Mark3::PORT_CS_EXIT` ()
- `K_WORD Mark3::PORT_CS_NESTING` ()

Variables

- `static constexpr auto Mark3::SR_` = uint8_t{0x3F}
- `K_WORD Mark3::g_kwSFR`
- `K_WORD Mark3::g_kwCriticalCount`

20.21.1 Detailed Description

ATMega1284p Multithreading support.

Definition in file [threadport.h](#).

20.21.2 Macro Definition Documentation

20.21.2.1 ASM

```
#define ASM(  
    x ) asm volatile(x);
```

Definition at line 32 of file [threadport.h](#).

20.21.2.2 PORT_PUSH_TO_STACK

```
#define PORT_PUSH_TO_STACK(  
    x,  
    y ) *x = y; x--;
```

Push a value y to the stack pointer x and decrement the stack pointer.

Definition at line 38 of file [threadport.h](#).

20.21.2.3 PORT_TOP_OF_STACK

```
#define PORT_TOP_OF_STACK(  
    x,  
    y ) (reinterpret_cast<K_WORD*>(reinterpret_cast<K_ADDR>(x) + (static_cast<K_ADDR>  
DDR>(y) - 1)))
```

Macro to find the top of a stack given its size and top address.

Definition at line 36 of file [threadport.h](#).

20.21.2.4 Thread_RestoreContext

```
#define Thread_RestoreContext( )
```

Restore the context of the Thread.

Definition at line 91 of file [threadport.h](#).

20.21.2.5 Thread_SaveContext

```
#define Thread_SaveContext ( )
```

Save the context of the Thread.

Definition at line 42 of file [threadport.h](#).

20.22 threadport.h

```
00001 /*=====
00002
00003
00004 | | | | | | | | | | | | | | | | | |
00005 | | | | | | | | | | | | | | | | | |
00006 | | | | | | | | | | | | | | | | | |
00007 | | | | | | | | | | | | | | | | | |
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00020 #pragma once
00021
00022 #include "portcfg.h"
00023 #include "kerneltypes.h"
00024
00025 #include <avr/io.h>
00026 #include <avr/interrupt.h>
00027
00028 namespace Mark3
00029 {
00030 // clang-format off
00031 //-----
00032 #define ASM(x) asm volatile(x);
00033
00034 //-----
00036 #define PORT_TOP_OF_STACK(x, y) (reinterpret_cast<K_WORD*>(reinterpret_cast<K_ADDR>(x) +
00037 (static_cast<K_ADDR>(y) - 1)))
00038 #define PORT_PUSH_TO_STACK(x, y) *x = y; x--;
00039
00040 //-----
00042 #define Thread_SaveContext() \
00043 ASM("push r0"); \
00044 ASM("in r0, __SREG__"); \
00045 ASM("cli"); \
00046 ASM("push r0"); \
00047 ASM("push r1"); \
00048 ASM("clr r1"); \
00049 ASM("push r2"); \
00050 ASM("push r3"); \
00051 ASM("push r4"); \
00052 ASM("push r5"); \
00053 ASM("push r6"); \
00054 ASM("push r7"); \
00055 ASM("push r8"); \
00056 ASM("push r9"); \
00057 ASM("push r10"); \
00058 ASM("push r11"); \
00059 ASM("push r12"); \
00060 ASM("push r13"); \
00061 ASM("push r14"); \
00062 ASM("push r15"); \
00063 ASM("push r16"); \
00064 ASM("push r17"); \
00065 ASM("push r18"); \
00066 ASM("push r19"); \
00067 ASM("push r20"); \
00068 ASM("push r21"); \
00069 ASM("push r22"); \
00070 ASM("push r23"); \
00071 ASM("push r24"); \
00072 ASM("push r25"); \
00073 ASM("push r26"); \
00074 ASM("push r27"); \
00075 ASM("push r28"); \
00076 ASM("push r29"); \
```

```

00077 ASM("push r30"); \
00078 ASM("push r31"); \
00079 ASM("in    r0, 0x3B"); \
00080 ASM("push r0"); \
00081 ASM("lds r26, g_pclCurrent"); \
00082 ASM("lds r27, g_pclCurrent + 1"); \
00083 ASM("adiw r26, 4"); \
00084 ASM("in    r0, 0x3D"); \
00085 ASM("st    x+, r0"); \
00086 ASM("in    r0, 0x3E"); \
00087 ASM("st    x+, r0");
00088
00089 //-----
00091 #define Thread_RestoreContext() \
00092 ASM("lds r26, g_pclCurrent"); \
00093 ASM("lds r27, g_pclCurrent + 1"); \
00094 ASM("adiw r26, 4"); \
00095 ASM("ld     r28, x+"); \
00096 ASM("out 0x3D, r28"); \
00097 ASM("ld     r29, x+"); \
00098 ASM("out 0x3E, r29"); \
00099 ASM("pop r0"); \
00100 ASM("out 0x3B, r0"); \
00101 ASM("pop r31"); \
00102 ASM("pop r30"); \
00103 ASM("pop r29"); \
00104 ASM("pop r28"); \
00105 ASM("pop r27"); \
00106 ASM("pop r26"); \
00107 ASM("pop r25"); \
00108 ASM("pop r24"); \
00109 ASM("pop r23"); \
00110 ASM("pop r22"); \
00111 ASM("pop r21"); \
00112 ASM("pop r20"); \
00113 ASM("pop r19"); \
00114 ASM("pop r18"); \
00115 ASM("pop r17"); \
00116 ASM("pop r16"); \
00117 ASM("pop r15"); \
00118 ASM("pop r14"); \
00119 ASM("pop r13"); \
00120 ASM("pop r12"); \
00121 ASM("pop r11"); \
00122 ASM("pop r10"); \
00123 ASM("pop r9"); \
00124 ASM("pop r8"); \
00125 ASM("pop r7"); \
00126 ASM("pop r6"); \
00127 ASM("pop r5"); \
00128 ASM("pop r4"); \
00129 ASM("pop r3"); \
00130 ASM("pop r2"); \
00131 ASM("pop r1"); \
00132 ASM("pop r0"); \
00133 ASM("out __SREG__, r0"); \
00134 ASM("pop r0");
00135
00136 //-----
00137 static constexpr auto SR_ = uint8_t{0x3F};
00138 extern "C" {
00139     extern K_WORD g_kwSFR;
00140     extern K_WORD g_kwCriticalCount;
00141 }
00142
00143 //-----
00144 inline uint8_t PORT_CLZ(uint8_t in_)
00145 {
00146     static const uint8_t u8Lookup[] = {4, 3, 2, 2, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0};
00147     uint8_t hi = __builtin_avr_swap(in_) & 0x0F;
00148     if (hi) {
00149         return u8Lookup[hi];
00150     }
00151     return 4 + u8Lookup[in_];
00152 }
00153
00154 //-----
00155 inline void PORT_IRQ_ENABLE()
00156 {
00157     ASM("sei");
00158 }
00159
00160 //-----
00161 inline void PORT_IRQ_DISABLE()
00162 {
00163     ASM("cli");
00164 }

```

```

00165
00166 //-----
00167 inline void PORT_CS_ENTER()
00168 {
00169     auto u8SFR = _SFR_IO8(SR_);
00170     ASM("cli");
00171     if (!g_kwCriticalCount) {
00172         g_kwSFR = u8SFR;
00173     }
00174     g_kwCriticalCount++;
00175 }
00176
00177 //-----
00178 inline void PORT_CS_EXIT()
00179 {
00180     g_kwCriticalCount--;
00181     if (!g_kwCriticalCount) {
00182         _SFR_IO8(SR_) = g_kwSFR;
00183     }
00184 }
00185
00186 //-----
00187 inline K_WORD PORT_CS_NESTING()
00188 {
00189     return g_kwCriticalCount;
00190 }
00191 } // namespace Mark3

```

20.23 /home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/threadport.cpp

File Reference

ATMega1284p Multithreading.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "thread.h"
#include "threadport.h"
#include "kernelswi.h"
#include "kerneltimer.h"
#include "timerlist.h"
#include "quantum.h"
#include "kernel.h"
#include <avr/io.h>
#include <avr/interrupt.h>

```

Namespaces

- [Mark3](#)

Functions

- static void [Mark3::Thread_Switch](#) (void)
- [Mark3::ISR](#) (INT2_vect) `__attribute__((signal`
ISR(INT2_vect) SWI using INT2 - used to trigger a context switch.

Variables

- `K_WORD` [g_kwSFR](#) = 0
- `K_WORD` [g_kwCriticalCount](#) = 0
- [Mark3::naked](#)

20.23.1 Detailed Description

ATMega1284p Multithreading.

Definition in file [threadport.cpp](#).

20.23.2 Variable Documentation

20.23.2.1 g_kwCriticalCount

```
K_WORD g_kwCriticalCount = 0
```

Definition at line 36 of file [threadport.cpp](#).

20.23.2.2 g_kwSFR

```
K_WORD g_kwSFR = 0
```

Definition at line 35 of file [threadport.cpp](#).

20.24 threadport.cpp

```
00001 /*=====
00002
00003 | | | | | | | | | | | | | | | | | |
00004 | | | | | | | | | | | | | | | | | |
00005 | | | | | | | | | | | | | | | | | |
00006 | | | | | | | | | | | | | | | | | |
00007 | | | | | | | | | | | | | | | | | |
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00022 #include "kerneltypes.h"
00023 #include "mark3cfg.h"
00024 #include "thread.h"
00025 #include "threadport.h"
00026 #include "kernelswi.h"
00027 #include "kerneltimer.h"
00028 #include "timerlist.h"
00029 #include "quantum.h"
00030 #include "kernel.h"
00031 #include <avr/io.h>
00032 #include <avr/interrupt.h>
00033
00034 extern "C" {
00035 K_WORD g_kwSFR = 0;
00036 K_WORD g_kwCriticalCount = 0;
00037 } // extern "C"
00038
00039 //-----
00040 namespace Mark3
00041 {
00042 //-----
00043 void ThreadPort::InitStack(Thread* pclThread_)
00044 {
```

```

00045 // Initialize the stack for a Thread
00046 uint16_t ul6Addr;
00047 uint8_t* pu8Stack;
00048 uint16_t i;
00049
00050 // Get the address of the thread's entry function
00051 ul6Addr = (uint16_t)(pclThread->m_pfEntryPoint);
00052
00053 // Start by finding the bottom of the stack
00054 pu8Stack = (uint8_t*)pclThread->m_pwStackTop;
00055
00056 #if KERNEL_STACK_CHECK
00057 // clear the stack, and initialize it to a known-default value (easier
00058 // to debug when things go sour with stack corruption or overflow)
00059 for (i = 0; i < pclThread->m_ul6StackSize; i++) { pclThread->
    m_pwStack[i] = 0xFF; }
00060 #endif // #if KERNEL_STACK_CHECK
00061
00062 // Our context starts with the entry function
00063 PORT_PUSH_TO_STACK(pu8Stack, (uint8_t)(ul6Addr & 0x00FF));
00064 PORT_PUSH_TO_STACK(pu8Stack, (uint8_t)((ul6Addr >> 8) & 0x00FF));
00065
00066 // R0
00067 PORT_PUSH_TO_STACK(pu8Stack, 0x00); // R0
00068
00069 // Push status register and R1 (which is used as a constant zero)
00070 PORT_PUSH_TO_STACK(pu8Stack, 0x80); // SR
00071 PORT_PUSH_TO_STACK(pu8Stack, 0x00); // R1
00072
00073 // Push other registers
00074 for (i = 2; i <= 23; i++) // R2-R23
00075 {
00076     PORT_PUSH_TO_STACK(pu8Stack, i);
00077 }
00078
00079 // Assume that the argument is the only stack variable
00080 PORT_PUSH_TO_STACK(pu8Stack, (uint8_t)((uint16_t)(pclThread->
    m_pvArg) & 0x00FF)); // R24
00081 PORT_PUSH_TO_STACK(pu8Stack, (uint8_t)((((uint16_t)(pclThread->
    m_pvArg) >> 8) & 0x00FF)); // R25
00082
00083 // Push the rest of the registers in the context
00084 for (i = 26; i <= 31; i++) { PORT_PUSH_TO_STACK(pu8Stack, i); }
00085
00086 PORT_PUSH_TO_STACK(pu8Stack, 0x00); // RAMPZ
00087 // Set the top of the stack.
00088 pclThread->m_pwStackTop = (uint8_t*)pu8Stack;
00089
00090 // That's it! the thread is ready to run now.
00091 }
00092
00093 //-----
00094 static void Thread_Switch(void)
00095 {
00096     g_pclCurrent = (Thread*)g_pclNext;
00097 }
00098
00099 //-----
00100 void ThreadPort::StartThreads()
00101 {
00102     KernelSWI::Config(); // configure the task switch SWI
00103     KernelTimer::Config(); // configure the kernel timer
00104
00105     // Tell the kernel that we're ready to start scheduling threads
00106     // for the first time.
00107     Kernel::CompleteStart();
00108
00109     Scheduler::SetScheduler(1); // enable the scheduler
00110     Scheduler::Schedule(); // run the scheduler - determine the first thread to run
00111
00112     Thread_Switch(); // Set the next scheduled thread to the current thread
00113
00114     KernelTimer::Start(); // enable the kernel timer
00115     KernelSWI::Start(); // enable the task switch SWI
00116
00117 #if KERNEL_ROUND_ROBIN
00118     // Restart the thread quantum timer, as any value held prior to starting
00119     // the kernel will be invalid. This fixes a bug where multiple threads
00120     // started with the highest priority before starting the kernel causes problems
00121     // until the running thread voluntarily blocks.
00122     Quantum::Update(g_pclCurrent);
00123 #endif // #if KERNEL_ROUND_ROBIN
00124
00125     // Restore the context...
00126     Thread_RestoreContext(); // restore the context of the first running thread
00127     ASM("reti"); // return from interrupt - will return to the first scheduled thread
00128 }

```

```

00129
00130 //-----
00131 //-----
00132
00133 ISR(INT2_vect) __attribute__((signal, naked));
00134 ISR(INT2_vect)
00135 {
00136     Thread_SaveContext();    // Push the context (registers) of the current task
00137     Thread_Switch();         // Switch to the next task
00138     Thread_RestoreContext(); // Pop the context (registers) of the next task
00139     ASM("reti");            // Return to the next task
00140 }
00141 } // namespace Mark3

```

20.25 /home/moslevin/projects/m3-repo/kernel/src/atomic.cpp File Reference

Basic Atomic Operations.

```
#include "mark3.h"
```

Namespaces

- Mark3

20.25.1 Detailed Description

Basic Atomic Operations.

Definition in file [atomic.cpp](#).

20.26 atomic.cpp

```

00001  /*=====
00002
00003
00004  |-----|-----|-----|-----|
00005  | \ / | | \ / | | \ / | | \ / | |
00006  | / \ | | / \ | | / \ | | / \ | |
00007  |-----|-----|-----|-----|
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2019 m0sleVIN, all rights reserved.
00012  See license.txt for more information
00013  =====*/
00021  #include "mark3.h"
00022
00023  namespace Mark3
00024  {
00025  //-----
00026  bool Atomic::TestAndSet(bool* pbLock_)
00027  {
00028      KERNEL_ASSERT(nullptr != pbLock_);
00029
00030      auto cs = CriticalGuard{};
00031      auto bRet = *pbLock_;
00032      if (!bRet) {
00033          *pbLock_ = 1;
00034      }
00035      return bRet;
00036  }
00037  } // namespace Mark3

```

20.27 /home/moslevin/projects/m3-repo/kernel/src/autoalloc.cpp File Reference

Automatic memory allocation for kernel objects.

```
#include "mark3.h"
#include <stdint.h>
```

Namespaces

- [Mark3](#)

20.27.1 Detailed Description

Automatic memory allocation for kernel objects.

Definition in file [autoalloc.cpp](#).

20.28 autoalloc.cpp

```
00001 /*=====
00002
00003 | | | | | | | | | | | | | | | | | | | | | |
00004 | | | | | | | | | | | | | | | | | | | | | |
00005 | | | | | | | | | | | | | | | | | | | | | |
00006 | | | | | | | | | | | | | | | | | | | | | |
00007 | | | | | | | | | | | | | | | | | | | | | |
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00019 #include "mark3.h"
00020
00021 #include <stdint.h>
00022
00023 //-----
00024 // Override new() and delete() using functions provided to AutoAlloc
00025 //-----
00026 #if PORT_OVERLOAD_NEW
00027 using namespace Mark3;
00028 void* operator new(size_t n)
00029 {
00030     return AutoAlloc::NewRawData(n);
00031 }
00032
00033 //-----
00034 void* operator new[](size_t n)
00035 {
00036     return AutoAlloc::NewRawData(n);
00037 }
00038
00039 //-----
00040 void operator delete(void* p)
00041 {
00042     AutoAlloc::DestroyRawData(p);
00043 }
00044
00045 //-----
00046 void operator delete[](void* p)
00047 {
00048     AutoAlloc::DestroyRawData(p);
00049 }
00050 #endif
00051 namespace Mark3
00052 {
00053
```

```

00054 AutoAllocAllocator_t AutoAlloc::m_pfAllocator;
00055 AutoAllocFree_t      AutoAlloc::m_pfFree;
00056
00057 //-----
00058 void* AutoAlloc::Allocate(AutoAllocType eType_, size_t sSize_)
00059 {
00060     if (!m_pfAllocator) {
00061         return nullptr;
00062     }
00063     return m_pfAllocator(eType_, sSize_);
00064 }
00065
00066 //-----
00067 void AutoAlloc::Free(AutoAllocType eType_, void* pObj_)
00068 {
00069     if (!m_pfFree) {
00070         return;
00071     }
00072     m_pfFree(eType_, pObj_);
00073 }
00074
00075 //-----
00076 void AutoAlloc::SetAllocatorFunctions(
    AutoAllocAllocator_t pfAllocator_, AutoAllocFree_t pfFree_)
00077 {
00078     m_pfAllocator = pfAllocator_;
00079     m_pfFree      = pfFree_;
00080 }
00081
00082 //-----
00083 void AutoAlloc::Init()
00084 {
00085     m_pfAllocator = nullptr;
00086     m_pfFree      = nullptr;
00087 }
00088
00089 //-----
00090 void* AutoAlloc::NewUserTypeAllocation(uint8_t eType_)
00091 {
00092     return Allocate(static_cast<AutoAllocType>(eType_), 0);
00093 }
00094 //-----
00095 void AutoAlloc::DestroyUserTypeAllocation(uint8_t eUserType_, void*
    pObj_)
00096 {
00097     Free(static_cast<AutoAllocType>(eUserType_), pObj_);
00098 }
00099 //-----
00100 void* AutoAlloc::NewRawData(size_t sSize_)
00101 {
00102     return Allocate(AutoAllocType::Raw, sSize_);
00103 }
00104 //-----
00105 void AutoAlloc::DestroyRawData(void* pData_)
00106 {
00107     Free(AutoAllocType::Raw, pData_);
00108 }
00109
00110 } // namespace Mark3

```

20.29 /home/moslevin/projects/m3-repo/kernel/src/blocking.cpp File Reference

Implementation of base class for blocking objects.

```
#include "mark3.h"
```

Namespaces

- [Mark3](#)

20.29.1 Detailed Description

Implementation of base class for blocking objects.

Definition in file [blocking.cpp](#).

20.30 blocking.cpp

```

00001  /*=====
00002
00003  _____
00004  /         \
00005  |         |
00006  |         |
00007  |         |
00008  |         |
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  =====*/
00021  #include "mark3.h"
00022
00023  namespace Mark3
00024  {
00025  //-----
00026  void BlockingObject::Block(Thread* pclThread_)
00027  {
00028      KERNEL_ASSERT(nullptr != pclThread_);
00029
00030      // Remove the thread from its current thread list (the "owner" list)
00031      // ... And add the thread to this object's block list
00032      Scheduler::Remove(pclThread_);
00033      m_clBlockList.Add(pclThread_);
00034
00035      // Set the "current" list location to the blocklist for this thread
00036      pclThread_>SetCurrent(&m_clBlockList);
00037      pclThread_>SetState(ThreadState::Blocked);
00038  }
00039
00040  //-----
00041  void BlockingObject::BlockPriority(Thread* pclThread_)
00042  {
00043      KERNEL_ASSERT(nullptr != pclThread_);
00044
00045      // Remove the thread from its current thread list (the "owner" list)
00046      // ... And add the thread to this object's block list
00047      Scheduler::Remove(pclThread_);
00048      m_clBlockList.AddPriority(pclThread_);
00049
00050      // Set the "current" list location to the blocklist for this thread
00051      pclThread_>SetCurrent(&m_clBlockList);
00052      pclThread_>SetState(ThreadState::Blocked);
00053  }
00054
00055  //-----
00056  void BlockingObject::UnBlock(Thread* pclThread_)
00057  {
00058      KERNEL_ASSERT(nullptr != pclThread_);
00059
00060      // Remove the thread from its current thread list (the "owner" list)
00061      pclThread_>GetCurrent()->Remove(pclThread_);
00062
00063      // Put the thread back in its active owner's list. This is usually
00064      // the ready-queue at the thread's original priority.
00065      Scheduler::Add(pclThread_);
00066
00067      // Tag the thread's current list location to its owner
00068      pclThread_>SetCurrent(pclThread_>GetOwner());
00069      pclThread_>SetState(ThreadState::Ready);
00070  }
00071  } // namespace Mark3

```


20.33 /home/moslevin/projects/m3-repo/kernel/src/condvar.cpp File Reference

Condition Variable implementation.

```
#include "mark3.h"
```

Namespaces

- [Mark3](#)

20.33.1 Detailed Description

Condition Variable implementation.

Definition in file [condvar.cpp](#).

20.34 condvar.cpp

```
00001  /*=====
00002
00003  _____
00004  | \ / | | | | \ / | | | \ / | | | \ / | | |
00005  | \ / | | | | \ / | | | \ / | | | \ / | | |
00006  | \ / | | | | \ / | | | \ / | | | \ / | | |
00007  | \ / | | | | \ / | | | \ / | | | \ / | | |
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2018 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  =====*/
00020  #include "mark3.h"
00021
00022  namespace Mark3
00023  {
00024  //-----
00025  void ConditionVariable::Init()
00026  {
00027      m_clMutex.Init();
00028      m_clSemaphore.Init(0, 255);
00029  }
00030
00031  //-----
00032  void ConditionVariable::Wait(Mutex* pclMutex_)
00033  {
00034      KERNEL_ASSERT(nullptr != pclMutex_);
00035
00036      m_clMutex.Claim();
00037
00038      pclMutex_>Release();
00039      m_u8Waiters++;
00040
00041      m_clMutex.Release();
00042
00043      m_clSemaphore.Pend();
00044      pclMutex_>Claim();
00045  }
00046
00047  //-----
00048  bool ConditionVariable::Wait(Mutex* pclMutex_, uint32_t u32WaitTimeMS_)
00049  {
00050      KERNEL_ASSERT(nullptr != pclMutex_);
00051
00052      m_clMutex.Claim();
00053
00054      pclMutex_>Release();
00055      m_u8Waiters++;
```



```

00056
00057     m_clMutex.Release();
00058
00059     if (!m_clSemaphore.Pend(u32WaitTimeMS_)) {
00060         return false;
00061     }
00062     return pclMutex_>Claim(u32WaitTimeMS_);
00063 }
00064
00065 //-----
00066 void ConditionVariable::Signal()
00067 {
00068     m_clMutex.Claim();
00069     if (m_u8Waiters) {
00070         m_u8Waiters--;
00071         m_clSemaphore.Post();
00072     }
00073     m_clMutex.Release();
00074 }
00075
00076 //-----
00077 void ConditionVariable::Broadcast()
00078 {
00079     m_clMutex.Claim();
00080
00081     while (m_u8Waiters > 0) {
00082         m_u8Waiters--;
00083         m_clSemaphore.Post();
00084     }
00085
00086     m_clMutex.Release();
00087 }
00088
00089 } // namespace Mark3

```

20.35 /home/moslevin/projects/m3-repo/kernel/src/coroutine.cpp File Reference

Coroutine object implementation.

```

#include "coroutine.h"
#include "cosched.h"
#include "criticalguard.h"
#include "kernel.h"

```

Namespaces

- [Mark3](#)

20.35.1 Detailed Description

Coroutine object implementation.

Definition in file [coroutine.cpp](#).


```
00089         m_pclOwner = CoScheduler::GetCoList (
00090             m_uPriority);
00091     } else {
00092         m_pclOwner = CoScheduler::GetStopList ();
00093     }
00094     m_pclOwner->Add(this);
00095 }
00096 //-----
00097 PORT_PRIO_TYPE Coroutine::GetPriority()
00098 {
00099     return m_uPriority;
00100 }
00101 } // namespace Mark3
```

20.37 /home/moslevin/projects/m3-repo/kernel/src/cosched.cpp File Reference

CoRoutine Scheduler implementation.

```
#include "cosched.h"
#include "criticalguard.h"
```

Namespaces

- Mark3

20.37.1 Detailed Description

CoRoutine Scheduler implementation.

Definition in file [cosched.cpp](#).

20.38 cosched.cpp

[illegible]


```

00021 #if KERNEL_EVENT_FLAGS
00022
00023 namespace Mark3
00024 {
00025     namespace
00026     {
00027         //-----
00028         void TimedEventFlag_Callback(Thread* pclOwner_, void* pvData_)
00029         {
00030             KERNEL_ASSERT(nullptr != pclOwner_);
00031             KERNEL_ASSERT(nullptr != pvData_);
00032
00033             auto* pclEventFlag = static_cast<EventFlag*>(pvData_);
00034
00035             pclOwner_->SetExpired(true);
00036             pclOwner_->SetEventFlagMask(0);
00037
00038             pclEventFlag->WakeMe(pclOwner_);
00039             if (pclOwner_->GetCurPriority() >= Scheduler::GetCurrentThread()->
GetCurPriority()) {
00040                 Thread::Yield();
00041             }
00042         }
00043     } // anonymous namespace
00044 //-----
00045 EventFlag::~EventFlag()
00046 {
00047     // If there are any threads waiting on this object when it goes out
00048     // of scope, set a kernel panic.
00049     if (nullptr != m_clBlockList.HighestWaiter()) {
00050         Kernel::Panic(PANIC_ACTIVE_EVENTFLAG_DESCOPED);
00051     }
00052 }
00053
00054 //-----
00055 void EventFlag::Init()
00056 {
00057     KERNEL_ASSERT(!m_clBlockList.GetHead());
00058     m_ul6SetMask = 0;
00059     SetInitialized();
00060 }
00061
00062 //-----
00063 void EventFlag::WakeMe(Thread* pclChosenOne_)
00064 {
00065     KERNEL_ASSERT(IsInitialized());
00066     KERNEL_ASSERT(nullptr != pclChosenOne_);
00067
00068     Unblock(pclChosenOne_);
00069 }
00070
00071 //-----
00072 uint16_t EventFlag::Wait_i(uint16_t ul6Mask_,
EventFlagOperation eMode_, uint32_t u32TimeMS_)
00073 {
00074     KERNEL_ASSERT(eMode_ <= EventFlagOperation::Pending_Unblock
);
00075     KERNEL_ASSERT(IsInitialized());
00076
00077     auto bThreadYield = false;
00078     auto bMatch = false;
00079
00080     auto clEventTimer = Timer {};
00081     auto bUseTimer = false;
00082
00083     // Ensure we're operating in a critical section while we determine
00084     // whether or not we need to block the current thread on this object.
00085
00086     { // Begin critical section
00087         const auto cs = CriticalGuard{};
00088
00089         // Check to see whether or not the current mask matches any of the
00090         // desired bits.
00091         g_pclCurrent->SetEventFlagMask(ul6Mask_);
00092
00093         if ((EventFlagOperation::All_Set == eMode_) || (
EventFlagOperation::All_Clear == eMode_)) {
00094             // Check to see if the flags in their current state match all of
00095             // the set flags in the event flag group, with this mask.
00096             if ((m_ul6SetMask & ul6Mask_) == ul6Mask_) {
00097                 bMatch = true;
00098                 g_pclCurrent->SetEventFlagMask(ul6Mask_);
00099
00100                 if (EventFlagOperation::All_Clear == eMode_) {
00101                     m_ul6SetMask &= ~ul6Mask_;
00102                     g_pclCurrent->SetExpired(false);
00103                 }
00104             }
00105         }
00106     }
00107 }

```

```

00114         }
00115     } else if ((EventFlagOperation::Any_Set == eMode_) || (
00116         EventFlagOperation::Any_Clear == eMode_)) {
00117         // Check to see if the existing flags match any of the set flags in
00118         // the event flag group with this mask
00119         if ((m_ul6SetMask & ul6Mask_) != 0) {
00120             bMatch = true;
00121             g_pclCurrent->SetEventFlagMask(
00122                 m_ul6SetMask & ul6Mask_);
00123             if (EventFlagOperation::Any_Clear == eMode_) {
00124                 m_ul6SetMask &= ~ul6Mask_;
00125                 g_pclCurrent->SetExpired(false);
00126             }
00127         }
00128
00129         // We're unable to match this pattern as-is, so we must block.
00130         if (!bMatch) {
00131             // Reset the current thread's event flag mask & mode
00132             g_pclCurrent->SetEventFlagMask(ul6Mask_);
00133             g_pclCurrent->SetEventFlagMode(eMode_);
00134
00135             if (0u != u32TimeMS_) {
00136                 g_pclCurrent->SetExpired(false);
00137                 clEventTimer.Init();
00138                 clEventTimer.Start(false, u32TimeMS_, TimedEventFlag_Callback, this);
00139                 bUseTimer = true;
00140             }
00141
00142             // Add the thread to the object's block-list.
00143             BlockPriority(g_pclCurrent);
00144
00145             // Trigger that
00146             bThreadYield = true;
00147         }
00148
00149         // If bThreadYield is set, it means that we've blocked the current thread,
00150         // and must therefore rerun the scheduler to determine what thread to
00151         // switch to.
00152         if (bThreadYield) {
00153             // Switch threads immediately
00154             Thread::Yield();
00155         }
00156     } // end critical section
00157
00162     if (bUseTimer && bThreadYield) {
00163         clEventTimer.Stop();
00164     }
00165
00166     return g_pclCurrent->GetEventFlagMask();
00167 }
00168
00169 //-----
00170 uint16_t EventFlag::Wait(uint16_t ul6Mask_, EventFlagOperation eMode_)
00171 {
00172     KERNEL_ASSERT(eMode_ <= EventFlagOperation::Pending_Unblock
00173 );
00174     return Wait_i(ul6Mask_, eMode_, 0);
00175 }
00176 //-----
00177 uint16_t EventFlag::Wait(uint16_t ul6Mask_, EventFlagOperation eMode_,
00178     uint32_t u32TimeMS_)
00179 {
00180     KERNEL_ASSERT(eMode_ <= EventFlagOperation::Pending_Unblock
00181 );
00182     return Wait_i(ul6Mask_, eMode_, u32TimeMS_);
00183 }
00184 //-----
00185 void EventFlag::Set(uint16_t ul6Mask_)
00186 {
00187     KERNEL_ASSERT(IsInitialized());
00188     auto bReschedule = false;
00189
00190     const auto cs = CriticalGuard{};
00191     // Walk through the whole block list, checking to see whether or not
00192     // the current flag set now matches any/all of the masks and modes of
00193     // the threads involved.
00194
00195     m_ul6SetMask |= ul6Mask_;
00196     auto ul6NewMask = m_ul6SetMask;
00197
00198     // Start at the head of the list, and iterate through until we hit the
00199     // "head" element in the list again. Ensure that we handle the case where

```

```

00200 // we remove the first or last elements in the list, or if there's only
00201 // one element in the list.
00202
00203 auto* pclCurrent = m_clBlockList.GetHead();
00204 // Do nothing when there are no objects blocking.
00205 if (nullptr != pclCurrent) {
00206     // First loop - process every thread in the block-list and check to
00207     // see whether or not the current flags match the event-flag conditions
00208     // on the thread.
00209     auto* pclPrev = (Thread*){};
00210     do {
00211         pclPrev = pclCurrent;
00212         pclCurrent = pclCurrent->GetNext();
00213
00214         // Read the thread's event mask/mode
00215         auto ul6ThreadMask = pclPrev->GetEventFlagMask();
00216         auto eThreadMode = pclPrev->GetEventFlagMode();
00217
00218         // For the "any" mode - unblock the blocked threads if one or more bits
00219         // in the thread's bitmask match the object's bitmask
00220         if ((EventFlagOperation::Any_Set == eThreadMode) || (
EventFlagOperation::Any_Clear == eThreadMode)) {
00221             if ((ul6ThreadMask & m_ul6SetMask) != 0) {
00222                 pclPrev->SetEventFlagMode(
EventFlagOperation::Pending_Unblock);
00223                 pclPrev->SetEventFlagMask(m_ul6SetMask & ul6ThreadMask);
00224                 bReschedule = true;
00225
00226                 // If the "clear" variant is set, then clear the bits in the mask
00227                 // that caused the thread to unblock.
00228                 if (EventFlagOperation::Any_Clear == eThreadMode) {
00229                     ul6NewMask &= ~(ul6ThreadMask & ul6Mask_);
00230                 }
00231             }
00232         }
00233         // For the "all" mode, every set bit in the thread's requested bitmask must
00234         // match the object's flag mask.
00235         else if ((EventFlagOperation::All_Set == eThreadMode) || (
EventFlagOperation::All_Clear == eThreadMode)) {
00236             if ((ul6ThreadMask & m_ul6SetMask) == ul6ThreadMask) {
00237                 pclPrev->SetEventFlagMode(
EventFlagOperation::Pending_Unblock);
00238                 pclPrev->SetEventFlagMask(ul6ThreadMask);
00239                 bReschedule = true;
00240
00241                 // If the "clear" variant is set, then clear the bits in the mask
00242                 // that caused the thread to unblock.
00243                 if (EventFlagOperation::All_Clear == eThreadMode) {
00244                     ul6NewMask &= ~(ul6ThreadMask & ul6Mask_);
00245                 }
00246             }
00247         }
00248     }
00249     // To keep looping, ensure that there's something in the list, and
00250     // that the next item isn't the head of the list.
00251     while (pclPrev != m_clBlockList.GetTail());
00252
00253     // Second loop - go through and unblock all of the threads that
00254     // were tagged for unblocking.
00255     pclCurrent = m_clBlockList.GetHead();
00256     auto bIsTail = false;
00257     do {
00258         pclPrev = pclCurrent;
00259         pclCurrent = pclCurrent->GetNext();
00260
00261         // Check to see if this is the condition to terminate the loop
00262         if (pclPrev == m_clBlockList.GetTail()) {
00263             bIsTail = true;
00264         }
00265
00266         // If the first pass indicated that this thread should be
00267         // unblocked, then unblock the thread
00268         if (pclPrev->GetEventFlagMode() ==
EventFlagOperation::Pending_Unblock) {
00269             Unblock(pclPrev);
00270         }
00271     } while (!bIsTail);
00272 }
00273
00274 // If we awoke any threads, re-run the scheduler
00275 if (bReschedule) {
00276     Thread::Yield();
00277 }
00278
00279 // Update the bitmask based on any "clear" operations performed along
00280 // the way
00281 m_ul6SetMask = ul6NewMask;

```



```

00027
00028 #if KERNEL_THREAD_CREATE_CALLOUT
00029 ThreadCreateCallout Kernel::m_pfThreadCreateCallout;
00030 #endif // #if KERNEL_THREAD_CREATE_CALLOUT
00031 #if KERNEL_THREAD_EXIT_CALLOUT
00032 ThreadExitCallout Kernel::m_pfThreadExitCallout;
00033 #endif // #if KERNEL_THREAD_EXIT_CALLOUT
00034 #if KERNEL_CONTEXT_SWITCH_CALLOUT
00035 ThreadContextCallout Kernel::m_pfThreadContextCallout;
00036 #endif // #if KERNEL_CONTEXT_SWITCH_CALLOUT
00037 DebugPrintFunction Kernel::m_pfDebugPrintFunction;
00038 #if KERNEL_STACK_CHECK
00039 uint16_t Kernel::m_ul6GuardThreshold;
00040 #endif // #if KERNEL_STACK_CHECK
00041 uint32_t Kernel::m_u32Ticks;
00042
00043 //-----
00044 void Kernel::Init()
00045 {
00046     // Call port-specific early init function
00047     ThreadPort::Init();
00048     AutoAlloc::Init();
00049     // Initialize the global kernel data - thread-scheduler, and timer-scheduler.
00050     Scheduler::Init();
00051     TimerScheduler::Init();
00052 #if KERNEL_STACK_CHECK
00053     m_ul6GuardThreshold = KERNEL_STACK_GUARD_DEFAULT;
00054 #endif // #if KERNEL_STACK_CHECK
00055 }
00056
00057 //-----
00058 void Kernel::Start()
00059 {
00060     ThreadPort::StartThreads();
00061 }
00062
00063 //-----
00064 void Kernel::CompleteStart()
00065 {
00066     m_bIsStarted = true;
00067 }
00068
00069 //-----
00070 void Kernel::Panic(uint16_t ul6Cause_)
00071 {
00072     m_bIsPanic = true;
00073     if (nullptr != m_pfPanic) {
00074         m_pfPanic(ul6Cause_);
00075     } else {
00076         while (true) {}
00077     }
00078 }
00079
00080 //-----
00081 void Kernel::DebugPrint(const char* szString_)
00082 {
00083     KERNEL_ASSERT(nullptr != szString_);
00084     if (nullptr != m_pfDebugPrintFunction) {
00085         m_pfDebugPrintFunction(szString_);
00086     }
00087 }
00088
00089 //-----
00090 uint32_t Kernel::GetTicks()
00091 {
00092     auto cs = CriticalGuard{};
00093     return m_u32Ticks;
00094 }
00095
00096 } // namespace Mark3

```

20.43 /home/moslevin/projects/m3-repo/kernel/src/ksemaphore.cpp File Reference

Semaphore Blocking-Object Implemenation.

```
#include "mark3.h"
```

Namespaces

- [Mark3](#)

20.43.1 Detailed Description

Semaphore Blocking-Object Implemenation.

Definition in file [ksemaphore.cpp](#).

20.44 ksemaphore.cpp

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00022 #include "mark3.h"
00023
00024 namespace Mark3
00025 {
00026     namespace
00027     {
00028         //-----
00038         void TimedSemaphore_Callback(Thread* pclOwner_, void* pvData_)
00039         {
00040             KERNEL_ASSERT(nullptr != pclOwner_);
00041             KERNEL_ASSERT(nullptr != pvData_);
00042
00043             auto* pclSemaphore = static_cast<Semaphore*>(pvData_);
00044
00045             // Indicate that the semaphore has expired on the thread
00046             pclOwner_>SetExpired(true);
00047
00048             // Wake up the thread that was blocked on this semaphore.
00049             pclSemaphore->WakeMe(pclOwner_);
00050
00051             if (pclOwner_>GetCurPriority() >= Scheduler::GetCurrentThread()->
GetCurPriority()) {
00052                 Thread::Yield();
00053             }
00054         }
00055     } // anonymous namespace
00056
00057 //-----
00058 Semaphore::~Semaphore()
00059 {
00060     // If there are any threads waiting on this object when it goes out
00061     // of scope, set a kernel panic.
00062     if (nullptr != m_clBlockList.GetHead()) {
00063         Kernel::Panic(PANIC_ACTIVE_SEMAPHORE_DESCOPEd);
00064     }
00065 }
00066
00067 //-----
00068 void Semaphore::WakeMe(Thread* pclChosenOne_)
00069 {
00070     KERNEL_ASSERT(pclChosenOne_);
00071     KERNEL_ASSERT(IsInitialized());
00072
00073     // Remove from the semaphore waitlist and back to its ready list.
00074     Unblock(pclChosenOne_);
00075 }
00076
00077 //-----
00078 uint8_t Semaphore::WakeNext()
00079 {

```

```

00080     auto* pclChosenOne = m_clBlockList.HighestWaiter();
00081     KERNEL_ASSERT(pclChosenOne);
00082
00083     // Remove from the semaphore waitlist and back to its ready list.
00084     Unblock(pclChosenOne);
00085
00086     // Call a task switch if higher or equal priority thread
00087     if (pclChosenOne->GetCurPriority() >= Scheduler::GetCurrentThread()->
GetCurPriority()) {
00088         return 1;
00089     }
00090     return 0;
00091 }
00092
00093 //-----
00094 void Semaphore::Init(uint16_t ul6InitVal_, uint16_t ul6MaxVal_)
00095 {
00096     KERNEL_ASSERT(!m_clBlockList.GetHead());
00097
00098     // Copy the paramters into the object - set the maximum value for this
00099     // semaphore to implement either binary or counting semaphores, and set
00100     // the initial count. Clear the wait list for this object.
00101     m_ul6Value = ul6InitVal_;
00102     m_ul6MaxValue = ul6MaxVal_;
00103
00104     SetInitialized();
00105 }
00106
00107 //-----
00108 bool Semaphore::Post()
00109 {
00110     KERNEL_ASSERT(IsInitialized());
00111
00112     auto bThreadWake = false;
00113     auto bBail = false;
00114     // Increment the semaphore count - we can mess with threads so ensure this
00115     // is in a critical section. We don't just disable the scheduler since
00116     // we want to be able to do this from within an interrupt context as well.
00117
00118     { // Begin critical section
00119         const auto cs = CriticalGuard{};
00120         // If nothing is waiting for the semaphore
00121         if (nullptr == m_clBlockList.GetHead()) {
00122             // Check to see if we've reached the maximum value in the semaphore
00123             if (m_ul6Value < m_ul6MaxValue) {
00124                 // Increment the count value
00125                 m_ul6Value++;
00126             } else {
00127                 // Maximum value has been reached, bail out.
00128                 bBail = true;
00129             }
00130         } else {
00131             // Otherwise, there are threads waiting for the semaphore to be
00132             // posted, so wake the next one (highest priority goes first).
00133             bThreadWake = (WakeNext() != 0u);
00134         }
00135     } // end critical section
00136
00137     // If we weren't able to increment the semaphore count, fail out.
00138     if (bBail) {
00139         return false;
00140     }
00141
00142     // if bThreadWake was set, it means that a higher-priority thread was
00143     // woken. Trigger a context switch to ensure that this thread gets
00144     // to execute next.
00145     if (bThreadWake) {
00146         Thread::Yield();
00147     }
00148     return true;
00149 }
00150
00151 //-----
00152 bool Semaphore::Pend_i(uint32_t u32WaitTimeMS_)
00153 {
00154     KERNEL_ASSERT(IsInitialized());
00155
00156     auto clSemTimer = Timer {};
00157     auto bUseTimer = false;
00158
00159     // Once again, messing with thread data - ensure
00160     // we're doing all of these operations from within a thread-safe context.
00161
00162     { // Begin critical section
00163         const auto cs = CriticalGuard{};
00164         // Check to see if we need to take any action based on the semaphore count
00165         if (0 != m_ul6Value) {

```

```

00166         // The semaphore count is non-zero, we can just decrement the count
00167         // and go along our merry way.
00168         m_ul6Value--;
00169     } else {
00170         // The semaphore count is zero - we need to block the current thread
00171         // and wait until the semaphore is posted from elsewhere.
00172         if (0u != u32WaitTimeMS_) {
00173             g_pclCurrent->SetExpired(false);
00174             clSemTimer.Init();
00175             clSemTimer.Start(false, u32WaitTimeMS_, TimedSemaphore_Callback, this);
00176             bUseTimer = true;
00177         }
00178         BlockPriority(g_pclCurrent);
00179
00180         // Switch Threads immediately
00181         Thread::Yield();
00182     }
00183 } // End critical section
00184
00185 if (bUseTimer) {
00186     clSemTimer.Stop();
00187     return (g_pclCurrent->GetExpired() == false);
00188 }
00189 return true;
00190 }
00191
00192 //-----
00193 // Redirect the untimed pend API to the timed pend, with a null timeout.
00194 void Semaphore::Pend()
00195 {
00196     Pend_i(0);
00197 }
00198
00199 //-----
00200 bool Semaphore::Pend(uint32_t u32WaitTimeMS_)
00201 {
00202     return Pend_i(u32WaitTimeMS_);
00203 }
00204
00205 //-----
00206 uint16_t Semaphore::GetCount()
00207 {
00208     KERNEL_ASSERT(IsInitialized());
00209
00210     auto cs = CriticalGuard{};
00211     return m_ul6Value;
00212 }
00213 } // namespace Mark3

```

20.45 /home/moslevin/projects/m3-repo/kernel/src/ll.cpp File Reference

Core Linked-List implementation, from which all kernel objects are derived.

```
#include "mark3.h"
```

Namespaces

- [Mark3](#)

20.45.1 Detailed Description

Core Linked-List implementation, from which all kernel objects are derived.

Definition in file [ll.cpp](#).


```

00093
00094 // Add a node to the end of the linked list.
00095 node_>prev = m_pclTail;
00096 node_>next = m_pclHead;
00097
00098 m_pclTail = node_;
00099 m_pclHead->prev = node_;
00100 }
00101
00102 //-----
00103 void CircularLinkedList::Remove(LinkListNode* node_)
00104 {
00105     KERNEL_ASSERT(nullptr != node_);
00106
00107     // Check to see if this is the head of the list...
00108     if ((node_ == m_pclHead) && (m_pclHead == m_pclTail)) {
00109         // Clear the head and tail pointers - nothing else left.
00110         m_pclHead = nullptr;
00111         m_pclTail = nullptr;
00112         return;
00113     }
00114
00115     // Verify that all nodes are properly connected
00116     if ((node_>prev->next != node_) || (node_>next->prev != node_)) {
00117         Kernel::Panic(PANIC_LIST_UNLINK_FAILED);
00118     }
00119
00120     // This is a circularly linked list - no need to check for connection,
00121     // just remove the node.
00122     node_>next->prev = node_>prev;
00123     node_>prev->next = node_>next;
00124
00125     if (node_ == m_pclHead) {
00126         m_pclHead = m_pclHead->next;
00127     }
00128     if (node_ == m_pclTail) {
00129         m_pclTail = m_pclTail->prev;
00130     }
00131     node_>ClearNode();
00132 }
00133
00134 //-----
00135 void CircularLinkedList::PivotForward()
00136 {
00137     if (nullptr != m_pclHead) {
00138         m_pclHead = m_pclHead->next;
00139         m_pclTail = m_pclTail->next;
00140     }
00141 }
00142
00143 //-----
00144 void CircularLinkedList::PivotBackward()
00145 {
00146     if (nullptr != m_pclHead) {
00147         m_pclHead = m_pclHead->prev;
00148         m_pclTail = m_pclTail->prev;
00149     }
00150 }
00151
00152 //-----
00153 void CircularLinkedList::InsertNodeBefore(
    LinkListNode* node_, LinkListNode* insert_)
00154 {
00155     KERNEL_ASSERT(nullptr != node_);
00156     KERNEL_ASSERT(nullptr != insert_);
00157
00158     node_>next = insert_;
00159     node_>prev = insert_>prev;
00160
00161     if (nullptr != insert_>prev) {
00162         insert_>prev->next = node_;
00163     }
00164     insert_>prev = node_;
00165 }
00166 } // namespace Mark3

```

20.47 /home/moslevin/projects/m3-repo/kernel/src/lockguard.cpp File Reference

Mutex RAIL helper class.

```
#include "mark3.h"
```



```

00079     return pclNew;
00080 }
00081
00082 //-----
00083 void Mailbox::Receive(void* pvData_)
00084 {
00085     KERNEL_ASSERT(nullptr != pvData_);
00086     Receive_i(pvData_, false, 0);
00087 }
00088
00089 //-----
00090 bool Mailbox::Receive(void* pvData_, uint32_t u32TimeoutMS_)
00091 {
00092     KERNEL_ASSERT(nullptr != pvData_);
00093     return Receive_i(pvData_, false, u32TimeoutMS_);
00094 }
00095
00096 //-----
00097 void Mailbox::ReceiveTail(void* pvData_)
00098 {
00099     KERNEL_ASSERT(nullptr != pvData_);
00100     Receive_i(pvData_, true, 0);
00101 }
00102
00103 //-----
00104 bool Mailbox::ReceiveTail(void* pvData_, uint32_t u32TimeoutMS_)
00105 {
00106     KERNEL_ASSERT(nullptr != pvData_);
00107     return Receive_i(pvData_, true, u32TimeoutMS_);
00108 }
00109
00110 //-----
00111 bool Mailbox::Send(void* pvData_)
00112 {
00113     KERNEL_ASSERT(nullptr != pvData_);
00114     return Send_i(pvData_, false, 0);
00115 }
00116
00117 //-----
00118 bool Mailbox::SendTail(void* pvData_)
00119 {
00120     KERNEL_ASSERT(nullptr != pvData_);
00121     return Send_i(pvData_, true, 0);
00122 }
00123
00124 //-----
00125 bool Mailbox::Send(void* pvData_, uint32_t u32TimeoutMS_)
00126 {
00127     KERNEL_ASSERT(nullptr != pvData_);
00128     return Send_i(pvData_, false, u32TimeoutMS_);
00129 }
00130
00131 //-----
00132 bool Mailbox::SendTail(void* pvData_, uint32_t u32TimeoutMS_)
00133 {
00134     KERNEL_ASSERT(nullptr != pvData_);
00135     return Send_i(pvData_, true, u32TimeoutMS_);
00136 }
00137
00138 //-----
00139 bool Mailbox::Send_i(const void* pvData_, bool bTail_, uint32_t u32TimeoutMS_)
00140 {
00141     KERNEL_ASSERT(nullptr != pvData_);
00142
00143     void* pvDst = nullptr;
00144
00145     auto bRet      = false;
00146     auto bSchedState = Scheduler::SetScheduler(false);
00147     auto bBlock     = false;
00148     auto bDone      = false;
00149
00150     while (!bDone) {
00151         // Try to claim a slot first before resorting to blocking.
00152         if (bBlock) {
00153             bDone = true;
00154             Scheduler::SetScheduler(bSchedState);
00155             m_clSendSem.Pend(u32TimeoutMS_);
00156             Scheduler::SetScheduler(false);
00157         }
00158
00159         { // Begin critical section
00160             const auto cs = CriticalGuard{};
00161             // Ensure we have a free slot before we attempt to write data
00162             if (0u != m_ul6Free) {
00163                 m_ul6Free--;
00164
00165                 if (bTail_) {

```

```

00166         pvDst = GetTailPointer();
00167         MoveTailBackward();
00168     } else {
00169         MoveHeadForward();
00170         pvDst = GetHeadPointer();
00171     }
00172     bRet = true;
00173     bDone = true;
00174     } else if (0u != u32TimeoutMS_) {
00175         bBlock = true;
00176     } else {
00177         bDone = true;
00178     }
00179     } // End critical section
00180 }
00181
00182 // Copy data to the claimed slot, and post the counting semaphore
00183 if (bRet) {
00184     CopyData(pvData_, pvDst, m_ul6ElementSize);
00185 }
00186
00187 Scheduler::SetScheduler(bSchedState);
00188
00189 if (bRet) {
00190     m_clRecvSem.Post();
00191 }
00192
00193 return bRet;
00194 }
00195
00196 //-----
00197 bool Mailbox::Receive_i(void* pvData_, bool bTail_, uint32_t u32WaitTimeMS_)
00198 {
00199     KERNEL_ASSERT(nullptr != pvData_);
00200     auto* pvSrc = (const void*){};
00201
00202     if (!m_clRecvSem.Pend(u32WaitTimeMS_)) {
00203         // Failed to get the notification from the counting semaphore in the
00204         // time allotted. Bail.
00205         return false;
00206     }
00207
00208     // Disable the scheduler while we do this -- this ensures we don't have
00209     // multiple concurrent readers off the same queue, which could be problematic
00210     // if multiple writes occur during reads, etc.
00211     auto bSchedState = Scheduler::SetScheduler(false);
00212
00213     // Update the head/tail indexes, and get the associated data pointer for
00214     // the read operation.
00215
00216     { // Begin critical section
00217         const auto cs = CriticalGuard{};
00218         m_ul6Free++;
00219         if (bTail_) {
00220             MoveTailForward();
00221             pvSrc = GetTailPointer();
00222         } else {
00223             pvSrc = GetHeadPointer();
00224             MoveHeadBackward();
00225         }
00226     } // end critical section
00227
00228     KERNEL_ASSERT(pvSrc);
00229     CopyData(pvSrc, pvData_, m_ul6ElementSize);
00230
00231     Scheduler::SetScheduler(bSchedState);
00232
00233     // Unblock a thread waiting for a free slot to send to
00234     m_clSendSem.Post();
00235
00236     return true;
00237 }
00238 } // namespace Mark3

```

20.51 /home/moslevin/projects/m3-repo/kernel/src/message.cpp File Reference

Inter-thread communications via message passing.

```
#include "mark3.h"
```



```

00072     return Receive_i(u32TimeWaitMS_);
00073 }
00074
00075 //-----
00076 Message* MessageQueue::Receive_i(uint32_t u32TimeWaitMS_)
00077 {
00078     // Block the current thread on the counting semaphore
00079     if (!m_clSemaphore.Pend(u32TimeWaitMS_)) {
00080         return nullptr;
00081     }
00082
00083     const auto cs = CriticalGuard{};
00084     // Pop the head of the message queue and return it
00085     auto* pclRet = m_clLinkList.GetHead();
00086     m_clLinkList.Remove(pclRet);
00087
00088     return pclRet;
00089 }
00090
00091 //-----
00092 void MessageQueue::Send(Message* pclSrc_)
00093 {
00094     KERNEL_ASSERT(pclSrc_);
00095
00096     CriticalSection::Enter();
00097
00098     // Add the message to the head of the linked list
00099     m_clLinkList.Add(pclSrc_);
00100
00101     CriticalSection::Exit();
00102
00103     // Post the semaphore, waking the blocking thread for the queue.
00104     m_clSemaphore.Post();
00105 }
00106
00107 //-----
00108 uint16_t MessageQueue::GetCount()
00109 {
00110     return m_clSemaphore.GetCount();
00111 }
00112 } // namespace Mark3

```

20.53 /home/moslevin/projects/m3-repo/kernel/src/mutex.cpp File Reference

Mutual-exclusion object.

```
#include "mark3.h"
```

Namespaces

- [Mark3](#)

20.53.1 Detailed Description

Mutual-exclusion object.

Definition in file [mutex.cpp](#).

20.54 mutex.cpp

```

00001  /*-----
00002
00003
00004
00005
00006
00007
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2019 m0slewin, all rights reserved.
00012  See license.txt for more information
00013  =====*/
00020  #include "mark3.h"
00021  namespace Mark3
00022  {
00023  namespace
00024  {
00025      //-----
00035      void TimedMutex_Callback(Thread* pclOwner_, void* pvData_)
00036      {
00037          KERNEL_ASSERT(nullptr != pclOwner_);
00038          KERNEL_ASSERT(nullptr != pvData_);
00039
00040          auto* pclMutex = static_cast<Mutex*>(pvData_);
00041
00042          // Indicate that the semaphore has expired on the thread
00043          pclOwner_->SetExpired(true);
00044
00045          // Wake up the thread that was blocked on this semaphore.
00046          pclMutex->WakeMe(pclOwner_);
00047
00048          if (pclOwner_->GetCurPriority() >= Scheduler::GetCurrentThread()->
GetCurPriority()) {
00049              Thread::Yield();
00050          }
00051      }
00052  } // anonymous namespace
00053
00054  //-----
00055  Mutex::~Mutex()
00056  {
00057      // If there are any threads waiting on this object when it goes out
00058      // of scope, set a kernel panic.
00059      if (nullptr != m_clBlockList.GetHead()) {
00060          Kernel::Panic(PANIC_ACTIVE_MUTEX_DESCOPED);
00061      }
00062  }
00063
00064  //-----
00065  void Mutex::WakeMe(Thread* pclOwner_)
00066  {
00067      KERNEL_ASSERT(nullptr != pclOwner_);
00068      // Remove from the semaphore waitlist and back to its ready list.
00069      Unblock(pclOwner_);
00070  }
00071
00072  //-----
00073  uint8_t Mutex::WakeNext()
00074  {
00075      // Get the highest priority waiter thread
00076      auto* pclChosenOne = m_clBlockList.HighestWaiter();
00077      KERNEL_ASSERT(pclChosenOne);
00078
00079      // Unblock the thread
00080      Unblock(pclChosenOne);
00081
00082      // The chosen one now owns the mutex
00083      m_pclOwner = pclChosenOne;
00084
00085      // Signal a context switch if it's a greater than or equal to the current priority
00086      if (pclChosenOne->GetCurPriority() >= Scheduler::GetCurrentThread()->
GetCurPriority()) {
00087          return 1;
00088      }
00089      return 0;
00090  }
00091
00092  //-----
00093  void Mutex::Init(bool bRecursive_)
00094  {
00095      // Cannot re-init a mutex which has threads blocked on it
00096      KERNEL_ASSERT(!m_clBlockList.GetHead());
00097

```

```

00098 // Reset the data in the mutex
00099 m_bReady = true; // The mutex is free.
00100 m_uMaxPri = 0; // Set the maximum priority inheritance state
00101 m_pclOwner = nullptr; // Clear the mutex owner
00102 m_u8Recurse = 0; // Reset recurse count
00103 m_bRecursive = bRecursive_;
00104 SetInitialized();
00105 }
00106
00107 //-----
00108 bool Mutex::Claim_i(uint32_t u32WaitTimeMS_)
00109 {
00110     KERNEL_ASSERT(IsInitialized());
00111
00112     auto clTimer = Timer {};
00113     auto bUseTimer = false;
00114
00115     // Disable the scheduler while claiming the mutex - we're dealing with all
00116     // sorts of private thread data, can't have a thread switch while messing
00117     // with internal data structures.
00118     Scheduler::SetScheduler(false);
00119
00120     // Check to see if the mutex is claimed or not
00121     if (false != m_bReady) {
00122         // Mutex isn't claimed, claim it.
00123         m_bReady = false;
00124         m_u8Recurse = 0;
00125         m_uMaxPri = g_pclCurrent->GetPriority();
00126         m_pclOwner = g_pclCurrent;
00127
00128         Scheduler::SetScheduler(true);
00129         return true;
00130     }
00131
00132     // If the mutex is already claimed, check to see if this is the owner thread,
00133     // since we allow the mutex to be claimed recursively.
00134     if (m_bRecursive && (g_pclCurrent == m_pclOwner)) {
00135         // Ensure that we haven't exceeded the maximum recursive-lock count
00136         KERNEL_ASSERT((m_u8Recurse < 255));
00137         m_u8Recurse++;
00138
00139         // Increment the lock count and bail
00140         Scheduler::SetScheduler(true);
00141         return true;
00142     }
00143
00144     // The mutex is claimed already - we have to block now. Move the
00145     // current thread to the list of threads waiting on the mutex.
00146     if (0u != u32WaitTimeMS_) {
00147         g_pclCurrent->SetExpired(false);
00148         clTimer.Init();
00149         clTimer.Start(false, u32WaitTimeMS_, TimedMutex_Callback, this);
00150         bUseTimer = true;
00151     }
00152     BlockPriority(g_pclCurrent);
00153
00154     // Check if priority inheritance is necessary. We do this in order
00155     // to ensure that we don't end up with priority inversions in case
00156     // multiple threads are waiting on the same resource.
00157     if (m_uMaxPri <= g_pclCurrent->GetPriority()) {
00158         m_uMaxPri = g_pclCurrent->GetPriority();
00159
00160         auto* pclTemp = m_clBlockList.GetHead();
00161         while (nullptr != pclTemp) {
00162             pclTemp->InheritPriority(m_uMaxPri);
00163             if (m_clBlockList.GetTail() == pclTemp) {
00164                 break;
00165             }
00166             pclTemp = pclTemp->GetNext();
00167         }
00168         m_pclOwner->InheritPriority(m_uMaxPri);
00169     }
00170
00171     // Done with thread data -reenable the scheduler
00172     Scheduler::SetScheduler(true);
00173
00174     // Switch threads if this thread acquired the mutex
00175     Thread::Yield();
00176
00177     if (bUseTimer) {
00178         clTimer.Stop();
00179         return (false == g_pclCurrent->GetExpired());
00180     }
00181     return true;
00182 }
00183
00184 //-----

```

```

00185 void Mutex::Claim(void)
00186 {
00187     Claim_i(0);
00188 }
00189
00190 //-----
00191 bool Mutex::Claim(uint32_t u32WaitTimeMS_)
00192 {
00193     return Claim_i(u32WaitTimeMS_);
00194 }
00195
00196 //-----
00197 void Mutex::Release()
00198 {
00199     KERNEL_ASSERT(IsInitialized());
00200
00201     auto bSchedule = false;
00202
00203     // Disable the scheduler while we deal with internal data structures.
00204     Scheduler::SetScheduler(false);
00205
00206     // This thread had better be the one that owns the mutex currently...
00207     KERNEL_ASSERT((g_pclCurrent == m_pclOwner));
00208
00209     // If the owner had claimed the lock multiple times, decrease the lock
00210     // count and return immediately.
00211     if (m_bRecursive && (0u != m_u8Recurse)) {
00212         m_u8Recurse--;
00213         Scheduler::SetScheduler(true);
00214         return;
00215     }
00216
00217     // Restore the thread's original priority
00218     if (g_pclCurrent->GetCurPriority() != g_pclCurrent->
GetPriority()) {
00219         g_pclCurrent->SetPriority(g_pclCurrent->
GetPriority());
00220
00221         // In this case, we want to reschedule
00222         bSchedule = true;
00223     }
00224
00225     // No threads are waiting on this semaphore?
00226     if (nullptr == m_clBlockList.GetHead()) {
00227         // Re-initialize the mutex to its default values
00228         m_bReady = true;
00229         m_uMaxPri = 0;
00230         m_pclOwner = nullptr;
00231     } else {
00232         // Wake the highest priority Thread pending on the mutex
00233         if (0u != WakeNext()) {
00234             // Switch threads if it's higher or equal priority than the current thread
00235             bSchedule = true;
00236         }
00237     }
00238
00239     // Must enable the scheduler again in order to switch threads.
00240     Scheduler::SetScheduler(true);
00241     if (bSchedule) {
00242         // Switch threads if a higher-priority thread was woken
00243         Thread::Yield();
00244     }
00245 }
00246 } // namespace Mark3

```

20.55 /home/moslevin/projects/m3-repo/kernel/src/notify.cpp File Reference

Lightweight thread notification - blocking object.

```
#include "mark3.h"
```

Namespaces

- [Mark3](#)


```

00078         while (nullptr != pclCurrent) {
00079             Unblock(pclCurrent);
00080             if (!bReschedule && (pclCurrent->GetCurPriority() >=
Scheduler::GetCurrentThread()->GetCurPriority())) {
00081                 bReschedule = true;
00082             }
00083             pclCurrent = m_clBlockList.GetHead();
00084         }
00085         m_bPending = false;
00086     }
00087 } // end critical section
00088
00089 if (bReschedule) {
00090     Thread::Yield();
00091 }
00092 }
00093
00094 //-----
00095 void Notify::Wait(bool* pbFlag_)
00096 {
00097     KERNEL_ASSERT(nullptr != pbFlag_);
00098     KERNEL_ASSERT(IsInitialized());
00099
00100     auto bEarlyExit = false;
00101     { // Begin critical section
00102         const auto cs = CriticalGuard{};
00103         if (!m_bPending) {
00104             Block(g_pclCurrent);
00105             if (nullptr != pbFlag_) {
00106                 *pbFlag_ = false;
00107             }
00108         } else {
00109             m_bPending = false;
00110             bEarlyExit = true;
00111         }
00112     } // End critical section
00113
00114     if (bEarlyExit) {
00115         return;
00116     }
00117
00118     Thread::Yield();
00119     if (nullptr != pbFlag_) {
00120         *pbFlag_ = true;
00121     }
00122 }
00123
00124 //-----
00125 bool Notify::Wait(uint32_t u32WaitTimeMS_, bool* pbFlag_)
00126 {
00127     KERNEL_ASSERT(nullptr != pbFlag_);
00128     KERNEL_ASSERT(IsInitialized());
00129
00130     auto bUseTimer = false;
00131     auto bEarlyExit = false;
00132     auto clNotifyTimer = Timer {};
00133
00134     { // Begin critical section
00135         const auto cs = CriticalGuard{};
00136         if (!m_bPending) {
00137             if (0u != u32WaitTimeMS_) {
00138                 bUseTimer = true;
00139                 g_pclCurrent->SetExpired(false);
00140
00141                 clNotifyTimer.Init();
00142                 clNotifyTimer.Start(false, u32WaitTimeMS_, TimedNotify_Callback, this);
00143             }
00144
00145             Block(g_pclCurrent);
00146
00147             if (nullptr != pbFlag_) {
00148                 *pbFlag_ = false;
00149             }
00150         } else {
00151             m_bPending = false;
00152             bEarlyExit = true;
00153         }
00154     } // end critical section
00155
00156     if (bEarlyExit) {
00157         return true;
00158     }
00159
00160     Thread::Yield();
00161
00162     if (bUseTimer) {
00163         clNotifyTimer.Stop();

```



```

00033 //-----
00034 void ProfileTimer::Start()
00035 {
00036     if (!m_bActive) {
00037         { // Begin critical section
00038             const auto cs = CriticalGuard{};
00039             m_u32StartTicks = Kernel::GetTicks();
00040         } // End Critical Section
00041         m_bActive = true;
00042     }
00043 }
00044
00045 //-----
00046 void ProfileTimer::Stop()
00047 {
00048     if (m_bActive) {
00049         uint32_t u32Final;
00050         { // Begin critical section
00051             const auto cs = CriticalGuard{};
00052             u32Final = Kernel::GetTicks();
00053             // Compute total for current iteration...
00054             m_u32CurrentIteration = u32Final -
m_u32StartTicks;
00055             m_u32Cumulative += m_u32CurrentIteration;
00056             m_ul6Iterations++;
00057         } // End critical section
00058         m_bActive = false;
00059     }
00060 }
00061
00062 //-----
00063 uint32_t ProfileTimer::GetAverage()
00064 {
00065     if (0u != m_ul6Iterations) {
00066         return (m_u32Cumulative + static_cast<uint32_t>(
m_ul6Iterations / 2)) / static_cast<uint32_t>(m_ul6Iterations);
00067     }
00068     return 0;
00069 }
00070
00071 //-----
00072 uint32_t ProfileTimer::GetCurrent()
00073 {
00074     if (m_bActive) {
00075         uint32_t u32Current;
00076         { // Begin critical section
00077             const auto cs = CriticalGuard{};
00078             u32Current = Kernel::GetTicks() - m_u32StartTicks;
00079         } // End critical section
00080         return u32Current;
00081     }
00082     return m_u32CurrentIteration;
00083 }
00084 } // namespace Mark3

```

20.59 /home/moslevin/projects/m3-repo/kernel/src/public/atomic.h File Reference

Basic Atomic Operations.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ithreadport.h"
#include "kerneldebug.h"

```

Namespaces

- [Mark3](#)
- [Mark3::Atomic](#)

The [Atomic](#) namespace This utility module provides primitives for atomic operations - that is, operations that are guaranteed to execute uninterrupted. Basic atomic primitives provided here include Set/Add/Subtract, as well as an atomic test-and-set.

Functions

- `template<typename T>`
`T Mark3::Atomic::Set (T *pSource_, T val_)`
Set Set a variable to a given value in an uninterruptable operation.
- `template<typename T>`
`T Mark3::Atomic::Add (T *pSource_, T val_)`
Add Add a value to a variable in an uninterruptable operation.
- `template<typename T>`
`T Mark3::Atomic::Sub (T *pSource_, T val_)`
Sub Subtract a value from a variable in an uninterruptable operation.
- `bool Mark3::Atomic::TestAndSet (bool *pbLock)`
TestAndSet Test to see if a variable is set, and set it if it is not already set. This is an uninterruptable operation.

20.59.1 Detailed Description

Basic Atomic Operations.

Definition in file [atomic.h](#).

20.60 atomic.h

```

00001  /*=====
00002
00003  _____
00004  /         \
00005  /         \
00006  /         \
00007  /         \
00008  /         \
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  =====*/
00021  #pragma once
00022
00023  #include "kerneltypes.h"
00024  #include "mark3cfg.h"
00025  #include "ithreadport.h"
00026  #include "kerneldebug.h"
00027
00028  namespace Mark3
00029  {
00038  namespace Atomic
00039  {
00047      template <typename T> T Set(T* pSource_, T val_)
00048      {
00049          KERNEL_ASSERT(nullptr != pSource_);
00050
00051          const auto cs = CriticalGuard{};
00052          auto ret      = *pSource_;
00053          *pSource_ = val_;
00054          return ret;
00055      }
00056
00064      template <typename T> T Add(T* pSource_, T val_)
00065      {
00066          KERNEL_ASSERT(nullptr != pSource_);
00067
00068          const auto cs = CriticalGuard{};
00069          auto ret = *pSource_;
00070          *pSource_ += val_;
00071          return ret;
00072      }
00073
00081      template <typename T> T Sub(T* pSource_, T val_)
00082      {

```

```

00083         KERNEL_ASSERT(nullptr != pSource_);
00084
00085         const auto cs = CriticalGuard{};
00086         auto ret = *pSource_;
00087         *pSource_ -= val_;
00088         return ret;
00089     }
00090
00106     bool TestAndSet(bool* pBlock);
00107 } // namespace Atomic
00108 } // namespace Mark3

```

20.61 /home/moslevin/projects/m3-repo/kernel/src/public/autoalloc.h File Reference

Automatic memory allocation for kernel objects.

```

#include <stddef.h>
#include <stdint.h>
#include <stdbool.h>
#include "mark3cfg.h"

```

Classes

- class [Mark3::AutoAlloc](#)

The [AutoAlloc](#) class. This class provides an object-allocation interface for both kernel objects and user-defined types. This class supplies callouts for alloc/free that use object-type metadata to determine how objects may be allocated, allowing a user to create custom dynamic memory implementations for specific object types and sizes. As a result, the user-defined allocators can avoid the kinds of memory fragmentation and exhaustion issues that occur in typical embedded systems in which a single heap is used to satisfy all allocations in the application.

Namespaces

- [Mark3](#)

Typedefs

- using [Mark3::AutoAllocAllocator_t](#) = void (*)(AutoAllocType eType_, size_t sSize_)
- using [Mark3::AutoAllocFree_t](#) = void (*)(AutoAllocType eType_, void *pvObj_)

Enumerations

- enum [Mark3::AutoAllocType](#) : uint8_t {
[Mark3::AutoAllocType::EventFlag](#), [Mark3::AutoAllocType::MailBox](#), [Mark3::AutoAllocType::Message](#),
[Mark3::AutoAllocType::MessagePool](#),
[Mark3::AutoAllocType::MessageQueue](#), [Mark3::AutoAllocType::Mutex](#), [Mark3::AutoAllocType::Notify](#),
[Mark3::AutoAllocType::Semaphore](#),
[Mark3::AutoAllocType::Thread](#), [Mark3::AutoAllocType::Timer](#), [Mark3::AutoAllocType::ConditionVariable](#),
[Mark3::AutoAllocType::ReaderWriterLock](#),
[Mark3::AutoAllocType::User](#), [Mark3::AutoAllocType::Raw](#) = 0xFF }


```

00099     AutoAllocFree_t pFree_);
00103     template <typename T, AutoAllocType e> static T* NewObject()
00104     {
00105         auto* pObj = Allocate(e, sizeof(T));
00106         if (pObj) {
00107             return new (pObj) T();
00108         }
00109         return 0;
00110     }
00111
00116     template <typename T, AutoAllocType e> static void DestroyObject(T* pObj_)
00117     {
00118         pObj_->~T();
00119         Free(e, pObj_);
00120     }
00121
00128     static void* NewUserTypeAllocation(uint8_t eUserType_);
00129
00136     static void DestroyUserTypeAllocation(uint8_t eUserType_, void* pObj_);
00137
00144     static void* NewRawData(size_t sSize_);
00145
00151     static void DestroyRawData(void* pData_);
00152
00153 private:
00154     static void* Allocate(AutoAllocType eType_, size_t sSize_);
00155     static void Free(AutoAllocType eType_, void* pObj_);
00156
00157     static AutoAllocAllocator_t m_pfAllocator;
00158     static AutoAllocFree_t m_pfFree;
00159 };
00160 } // namespace Mark3

```

20.63 /home/moslevin/projects/m3-repo/kernel/src/public/blocking.h File Reference

Blocking object base class declarations.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ll.h"
#include "threadlist.h"

```

Classes

- class [Mark3::BlockingObject](#)

The [BlockingObject](#) class. Class implementing thread-blocking primitives. used for implementing things like semaphores, mutexes, message queues, or anything else that could cause a thread to suspend execution on some external stimulus.

Namespaces

- [Mark3](#)

20.63.1 Detailed Description

Blocking object base class declarations.

A Blocking object in [Mark3](#) is essentially a thread list. Any blocking object implementation (being a semaphore, mutex, event flag, etc.) can be built on top of this class, utilizing the provided functions to manipulate thread location within the Kernel.

Blocking a thread results in that thread becoming de-scheduled, placed in the blocking object's own private list of threads which are waiting on the object.

Unblocking a thread results in the reverse: The thread is moved back to its original location from the blocking list.

The only difference between a blocking object based on this class is the logic used to determine what constitutes a Block or Unblock condition.

For instance, a semaphore Pend operation may result in a call to the Block() method with the currently-executing thread in order to make that thread wait for a semaphore Post. That operation would then invoke the Unblock() method, removing the blocking thread from the semaphore's list, and back into the appropriate thread inside the scheduler.

Care must be taken when implementing blocking objects to ensure that critical sections are used judiciously, otherwise asynchronous events like timers and interrupts could result in non-deterministic and often catastrophic behavior.

Definition in file [blocking.h](#).

20.64 blocking.h

```

00001  /*=====
00002
00003  00004  00005  00006  00007
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  ===== */
00046  #pragma once
00047
00048  #include "kerneltypes.h"
00049  #include "mark3cfg.h"
00050
00051  #include "ll.h"
00052  #include "threadlist.h"
00053
00054  namespace Mark3
00055  {
00056  class Thread;
00057
00058  //-----
00065  class BlockingObject
00066  {
00067  public:
00068      BlockingObject() { m_u8Initialized =
00069          m_uBlockingInvalidCookie; }
00069      ~BlockingObject() { m_u8Initialized =
00070          m_uBlockingInvalidCookie; }
00070
00071  protected:
00091      void Block(Thread* pclThread_);
00092
00100      void BlockPriority(Thread* pclThread_);
00101
00112      void Unblock(Thread* pclThread_);

```


20.65 /home/moslevin/projects/m3-repo/kernel/src/public/colist.h File Reference

```

00020
00021 #include "mark3cfg.h"
00022 #include "coroutine.h"
00023
00024 namespace Mark3 {
00025 class Coroutine;
00026
00035 class CoList : public TypedCircularLinkedList<Coroutine>
00036 {
00037 public:
00038
00045     void SetPrioMap(CoPrioMap* pClPrioMap_);
00046
00054     void SetPriority(PORT_PRIO_TYPE uPriority_);
00055
00062     void Add(Coroutine* pClCoroutine_);
00063
00070     void Remove(Coroutine* pClCoroutine_);
00071
00072 private:
00073     CoPrioMap* m_pClPrioMap;
00074     uint8_t m_uPriority;
00075 };
00076 } // namespace Mark3

```

20.67 /home/moslevin/projects/m3-repo/kernel/src/public/condvar.h File Reference

Condition Variable implementation.

```

#include "mark3cfg.h"
#include "ksemaphore.h"
#include "mutex.h"
#include <stddef.h>

```

Classes

- class [Mark3::ConditionVariable](#)

The [ConditionVariable](#) class This class implements a condition variable. This is a synchronization object that allows multiple threads to block, each waiting for specific signals unique to them. Access to the specified condition is guarded by a mutex that is supplied by the caller. This object can permit multiple waiters that can be unblocked one-at-a-time via signalling, or unblocked all at once via broadcasting. This object is built upon lower-level primitives, and is somewhat more heavyweight than the primitive types supplied by the kernel.

Namespaces

- [Mark3](#)

20.67.1 Detailed Description

Condition Variable implementation.

Definition in file [condvar.h](#).

Typedefs

- using [Mark3::CoPrioMap](#) = PriorityMapL1< [PORT_PRIO_TYPE](#), [PORT_COROUTINE_PRIORITIES](#) >
- using [Mark3::CoroutineHandler](#) = void(*) (Coroutine *pclCaller_, void *pvContext_)

20.69.1 Detailed Description

CoRoutine implementation.

Definition in file [coroutine.h](#).

20.70 [coroutine.h](#)

```

00001  /*=====
00002
00003
00004  |-----|-----|-----|-----|-----|-----|-----|-----|
00005  | \ / | \ / | \ / | \ / | \ / | \ / | \ / | \ / | \ / | \ / |
00006  | / \ | / \ | / \ | / \ | / \ | / \ | / \ | / \ | / \ |
00007  |-----|-----|-----|-----|-----|-----|-----|-----|
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  ===== */
00019  #pragma once
00020
00021  #include "mark3cfg.h"
00022  #include "ll.h"
00023  #include "priomap1.h"
00024  #include "priomap2.h"
00025
00026  namespace Mark3 {
00027
00028  // Priority map type declaration, based on port configuration
00029  #if PORT_COROUTINE_PRIORITIES <= (PORT_PRIO_MAP_WORD_SIZE * 8u)
00030  using CoPrioMap =
00031      PriorityMapL1<PORT_PRIO_TYPE, PORT_COROUTINE_PRIORITIES>
00032  ;
00033  #else
00034  using CoPrioMap =
00035      PriorityMapL2<PORT_PRIO_TYPE, PORT_COROUTINE_PRIORITIES>
00036  ;
00037  #endif
00038
00039  // Forward declarations
00040  class CoList;
00041  class Coroutine;
00042
00043  // CoRoutine functino handler type definition
00044  using CoroutineHandler = void (*) (Coroutine* pclCaller_, void* pvContext_);
00045
00046  class Coroutine : public TypedLinkedListNode<Coroutine>
00047  {
00048  public:
00049      ~Coroutine();
00050
00051      void Init(PORT_PRIO_TYPE uPriority_, CoroutineHandler pfHandler_,
00052          void* pvContext_);
00053
00054      void Run();
00055
00056      void Activate();
00057
00058      void SetPriority(PORT_PRIO_TYPE uPriority_);
00059
00060      PORT_PRIO_TYPE GetPriority();
00061
00062  private:
00063      CoList* m_pclOwner;
00064      CoroutineHandler m_pfHandler;
00065      void* m_pvContext;
00066      PORT_PRIO_TYPE m_uPriority;
00067      bool m_bQueued;
00068  };
00069  } // namespace Mark3

```

20.71 /home/moslevin/projects/m3-repo/kernel/src/public/cosched.h File Reference

CoRoutine Scheduler implementation.

```
#include "mark3cfg.h"
#include "coroutine.h"
#include "colist.h"
```

Classes

- class [Mark3::CoScheduler](#)

The [CoScheduler](#) class. This class implements the coroutine scheduler. Similar to the [Mark3](#) thread scheduler, the highest-priority active object is scheduled / returned for execution. If no active co-routines are available to be scheduled, then the scheduler returns nullptr.

Namespaces

- [Mark3](#)

20.71.1 Detailed Description

CoRoutine Scheduler implementation.

Definition in file [cosched.h](#).

20.72 cosched.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00019 #pragma once
00020
00021 #include "mark3cfg.h"
00022 #include "coroutine.h"
00023 #include "colist.h"
00024
00025 namespace Mark3 {
00033 class CoScheduler
00034 {
00035 public:
00036
00042     static void Init();
00043
00050     static CoPrioMap* GetPrioMap();
00051
00059     static CoList* GetStopList();
00060
00068     static CoList* GetCoList(PORT_PRIO_TYPE uPriority_);
00069
00077     static Coroutine* Schedule();
00078
00079 private:
00080     static CoList m_aclPriorities[PORT_COROUTINE_PRIORITIES];
00081     static CoList m_clStopList;
00082     static CoPrioMap m_clPrioMap;
00083 };
00084 } // namespace Mark3
```


20.75 /home/moslevin/projects/m3-repo/kernel/src/public/criticalsection.h File Reference

Critical Section Support.

```
#include "mark3cfg.h"
#include "threadport.h"
```

Classes

- class [Mark3::CriticalSection](#)

The [CriticalSection](#) class. This class implements a portable [CriticalSection](#) interface based on macros/inline functions that are implemented as part of each port.

Namespaces

- [Mark3](#)

20.75.1 Detailed Description

Critical Section Support.

Definition in file [criticalsection.h](#).

20.76 criticalsection.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00020 #pragma once
00021
00022 #include "mark3cfg.h"
00023 #include "threadport.h"
00024
00025 namespace Mark3 {
00026
00027 //-----
00048 class CriticalSection {
00049 public:
00050
00056     static inline void Enter() {
00057         PORT_CS_ENTER();
00058     }
00059
00065     static inline void Exit() {
00066         PORT_CS_EXIT();
00067     }
00068
00073     static inline K_WORD NestingCount() {
00074         return PORT_CS_NESTING();
00075     }
00076 };
00077
00078 } // namespace Mark3
```

20.77 /home/moslevin/projects/m3-repo/kernel/src/public/eventflag.h File Reference

Event Flag Blocking Object/IPC-Object definition.

```
#include "mark3cfg.h"
#include "kernel.h"
#include "kerneltypes.h"
#include "blocking.h"
#include "thread.h"
```

Classes

- class [Mark3::EventFlag](#)

The [EventFlag](#) class. This class implements a blocking object, similar to a semaphore or mutex, commonly used for synchronizing thread execution based on events occurring within the system. Each [EventFlag](#) object contains a 16-bit bitmask, which is used to trigger events on associated threads. Threads wishing to block, waiting for a specific event to occur can wait on any pattern within this 16-bit bitmask to be set. Here, we provide the ability for a thread to block, waiting for ANY bits in a specified mask to be set, or for ALL bits within a specific mask to be set. Depending on how the object is configured, the bits that triggered the wakeup can be automatically cleared once a match has occurred.

Namespaces

- [Mark3](#)

20.77.1 Detailed Description

Event Flag Blocking Object/IPC-Object definition.

Definition in file [eventflag.h](#).

20.78 eventflag.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00014 #pragma once
00015
00016 #include "mark3cfg.h"
00017 #include "kernel.h"
00018 #include "kerneltypes.h"
00019 #include "blocking.h"
00020 #include "thread.h"
00021
00022 #if KERNEL_EVENT_FLAGS
00023 namespace Mark3
00024 {
00025 //-----
00045 class EventFlag : public BlockingObject
```



```

00046 {
00047 public:
00048     void* operator new(size_t sz, void* pv) { return reinterpret_cast<EventFlag*>(pv); };
00049     ~EventFlag();
00050
00054     void Init();
00055
00064     uint16_t Wait(uint16_t ul6Mask_, EventFlagOperation eMode_);
00065
00075     uint16_t Wait(uint16_t ul6Mask_, EventFlagOperation eMode_, uint32_t u32TimeMS_);
00076
00083     void WakeMe(Thread* pClChosenOne_);
00084
00091     void Set(uint16_t ul6Mask_);
00092
00097     void Clear(uint16_t ul6Mask_);
00098
00103     uint16_t GetMask();
00104
00105 private:
00117     uint16_t Wait_i(uint16_t ul6Mask_, EventFlagOperation eMode_, uint32_t
        u32TimeMS_);
00118
00119     uint16_t m_ul6SetMask;
00120 };
00121 } // namespace Mark3
00122 #endif // #if KERNEL_EVENT_FLAGS

```

20.79 /home/moslevin/projects/m3-repo/kernel/src/public/ithreadport.h File Reference

Thread porting interface.

```
#include <stdint.h>
```

Classes

- class [Mark3::ThreadPort](#)

The [ThreadPort](#) Class defines the target-specific functions required by the kernel for threading.

Namespaces

- [Mark3](#)

20.79.1 Detailed Description

Thread porting interface.

Definition in file [ithreadport.h](#).

20.81.1 Detailed Description

Kernel initialization and startup class.

The Kernel namespace provides functions related to initializing and starting up the kernel.

The Kernel::Init() function must be called before any of the other functions in the kernel can be used.

Once the initial kernel configuration has been completed (i.e. first threads have been added to the scheduler), the Kernel::Start() function can then be called, which will transition code execution from the "main()" context to the threads in the scheduler.

Definition in file [kernel.h](#).

20.81.2 Typedef Documentation

20.81.2.1 DebugPrintFunction

using [DebugPrintFunction](#) = void (*)(const char* szString_)

Definition at line 39 of file [kernel.h](#).

20.82 kernel.h

```

00001  /*=====
00002
00003  _____
00004  /         \
00005  /         \
00006  /         \
00007  /         \
00008  /         \
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  =====*/
00032  #pragma once
00033
00034  #include "mark3cfg.h"
00035  #include "kerneltypes.h"
00036  #include "paniccodes.h"
00037  #include "thread.h"
00038
00039  using DebugPrintFunction = void (*)(const char* szString_);
00040
00041  namespace Mark3
00042  {
00043  //-----
00048  class Kernel
00049  {
00050  public:
00058      static void Init();
00059
00071      static void Start();
00072
00078      static void CompleteStart();
00079
00086      static bool IsStarted() { return m_bIsStarted; }
00094      static void SetPanic(PanicFunc pfPanic_) { m_pfPanic = pfPanic_; }
00099      static bool IsPanic() { return m_bIsPanic; }
00104      static void Panic(uint16_t u16Cause_);
00105

```

```

00106 #if KERNEL_THREAD_CREATE_CALLOUT
00107
00116     static void SetThreadCreateCallout(ThreadCreateCallout
pfCreate_) { m_pfThreadCreateCallout = pfCreate_; }
00117 #endif // #if KERNEL_THREAD_CREATE_HOOK
00118
00119 #if KERNEL_THREAD_EXIT_CALLOUT
00120
00130     static void SetThreadExitCallout(ThreadExitCallout pfExit_) {
m_pfThreadExitCallout = pfExit_; }
00131 #endif // #if KERNEL_THREAD_EXIT_CALLOUT
00132
00133 #if KERNEL_CONTEXT_SWITCH_CALLOUT
00134
00143     static void SetThreadContextSwitchCallout(
ThreadContextCallout pfContext_)
00144     {
00145         m_pfThreadContextCallout = pfContext_;
00146     }
00147 #endif // KERNEL_CONTEXT_SWITCH_CALLOUT
00148
00155     static void SetDebugPrintFunction(DebugPrintFunction
pfPrintFunction_)
00156     {
00157         m_pfDebugPrintFunction = pfPrintFunction_;
00158     }
00159
00168     static void DebugPrint(const char* szString_);
00169
00170 #if KERNEL_THREAD_CREATE_CALLOUT
00171
00178     static ThreadCreateCallout GetThreadCreateCallout() { return
m_pfThreadCreateCallout; }
00179 #endif // #if KERNEL_THREAD_CREATE_HOOK
00180 #if KERNEL_THREAD_EXIT_CALLOUT
00181
00188     static ThreadExitCallout GetThreadExitCallout() { return
m_pfThreadExitCallout; }
00189 #endif // #if KERNEL_THREAD_EXIT_HOOK
00190 #if KERNEL_CONTEXT_SWITCH_CALLOUT
00191
00198     static ThreadContextCallout
GetThreadContextSwitchCallout() { return
m_pfThreadContextCallout; }
00199 #endif // #if KERNEL_CONTEXT_SWITCH_CALLOUT
00200 #if KERNEL_STACK_CHECK
00201     static void SetStackGuardThreshold(uint16_t ul6Threshold_) {
m_ul6GuardThreshold = ul6Threshold_; }
00202     static uint16_t GetStackGuardThreshold() { return
m_ul6GuardThreshold; }
00203 #endif // #if KERNEL_STACK_CHECK
00204
00205     static void Tick() { m_u32Ticks++; }
00206     static uint32_t GetTicks();
00207
00208 private:
00209     static bool m_bIsStarted;
00210     static bool m_bIsPanic;
00211     static PanicFunc m_pfPanic;
00212
00213 #if KERNEL_THREAD_CREATE_CALLOUT
00214     static ThreadCreateCallout m_pfThreadCreateCallout;
00215 #endif // #if KERNEL_THREAD_CREATE_HOOK
00216 #if KERNEL_THREAD_EXIT_CALLOUT
00217     static ThreadExitCallout m_pfThreadExitCallout;
00218 #endif // #if KERNEL_THREAD_EXIT_HOOK
00219 #if KERNEL_CONTEXT_SWITCH_CALLOUT
00220     static ThreadContextCallout m_pfThreadContextCallout;
00221 #endif // #if KERNEL_CONTEXT_SWITCH_CALLOUT
00222     static DebugPrintFunction m_pfDebugPrintFunction;
00223 #if KERNEL_STACK_CHECK
00224     static uint16_t m_ul6GuardThreshold;
00225 #endif // #if KERNEL_STACK_CHECK
00226     static uint32_t m_u32Ticks;
00227 };
00228
00229 } // namespace Mark3

```

20.83 /home/moslevin/projects/m3-repo/kernel/src/public/kerneldebug.h File Reference

Macros and functions used for assertions, kernel traces, etc.

```
#include "mark3cfg.h"
#include "paniccodes.h"
#include "kernel.h"
```

Namespaces

- [Mark3](#)

Macros

- `#define` [KERNEL_ASSERT\(x\)](#)

20.83.1 Detailed Description

Macros and functions used for assertions, kernel traces, etc.

Definition in file [kerneldebug.h](#).

20.83.2 Macro Definition Documentation

20.83.2.1 KERNEL_ASSERT

```
#define KERNEL_ASSERT(  
    x )
```

Definition at line 36 of file [kerneldebug.h](#).

20.84 kerneldebug.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00019 #pragma once
00020
00021 #include "mark3cfg.h"
00022 #include "paniccodes.h"
00023 #include "kernel.h"
00024
00025 //-----
00026 namespace Mark3
00027 {
00028 #if KERNEL_DEBUG
```



```

00025 namespace Mark3
00026 {
00032 class KernelSWI
00033 {
00034 public:
00040     static void Config(void);
00041
00046     static void Start(void);
00047
00053     static void Trigger(void);
00054 };
00055 } // namespace Mark3

```

20.87 /home/moslevin/projects/m3-repo/kernel/src/public/kerneltimer.h File Reference

Kernel Timer Class declaration.

```

#include "kerneltypes.h"
#include "mark3cfg.h"

```

Classes

- class [Mark3::KernelTimer](#)

The [KernelTimer](#) class provides a timer interface used by all time-based scheduling/timer subsystems in the kernel. This interface must be implemented by target-specific code in the porting layer.

Namespaces

- [Mark3](#)

20.87.1 Detailed Description

Kernel Timer Class declaration.

Definition in file [kerneltimer.h](#).

20.88 kerneltimer.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00021 #pragma once
00022
00023 #include "kerneltypes.h"
00024 #include "mark3cfg.h"
00025
00026 namespace Mark3
00027 {
00028 //-----
00034 class KernelTimer
00035 {
00036 public:
00041     static void Config(void);
00042
00047     static void Start(void);
00048
00053     static void Stop(void);
00054 };
00055 } // namespace Mark3

```



```

00022
00023 #pragma once
00024 namespace Mark3
00025 {
00026 //-----
00030 using PanicFunc = void (*) (uint16_t u16PanicCode_);
00031
00032 //-----
00037 using IdleFunc = void (*) ();
00038
00039 //-----
00043 using ThreadEntryFunc = void (*) (void* pvArg_);
00044
00045 //-----
00050 enum class EventFlagOperation : uint8_t {
00051     All_Set = 0,
00052     Any_Set,
00053     All_Clear,
00054     Any_Clear,
00055     Pending_Unblock
00056 };
00057
00058 //-----
00062 enum class ThreadState : uint8_t { Exit = 0, Ready, Blocked,
    Stop, Invalid };
00063
00064 } // namespace Mark3

```

20.91 /home/moslevin/projects/m3-repo/kernel/src/public/ksemaphore.h File Reference

Semaphore Blocking Object class declarations.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "blocking.h"
#include "threadlist.h"

```

Classes

- class [Mark3::Semaphore](#)
the [Semaphore](#) class provides Binary & Counting semaphore objects, based on [BlockingObject](#) base class.

Namespaces

- [Mark3](#)

20.91.1 Detailed Description

Semaphore Blocking Object class declarations.

Definition in file [ksemaphore.h](#).

20.92 ksemaphore.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00021 #pragma once
00022
00023 #include "kerneltypes.h"
00024 #include "mark3cfg.h"
00025
00026 #include "blocking.h"
00027 #include "threadlist.h"
00028
00029 namespace Mark3
00030 {
00031 //-----
00036 class Semaphore : public BlockingObject
00037 {
00038 public:
00039     void* operator new(size_t sz, void* pv) { return reinterpret_cast<Semaphore*>(pv); };
00040     ~Semaphore();
00041
00042     void Init(uint16_t u16InitVal_, uint16_t u16MaxVal_);
00043
00044     bool Post();
00045
00046     void Pend();
00047
00048     uint16_t GetCount();
00049
00050     bool Pend(uint32_t u32WaitTimeMS_);
00051
00052     void WakeMe(Thread* pclChosenOne_);
00053
00054 private:
00055     uint8_t WakeNext();
00056
00057     bool Pend_i(uint32_t u32WaitTimeMS_);
00058
00059     uint16_t m_u16Value;
00060     uint16_t m_u16MaxValue;
00061 };
00062 } // namespace Mark3

```

20.93 /home/moslevin/projects/m3-repo/kernel/src/public/ll.h File Reference

Core linked-list declarations, used by all kernel list types At the heart of RTOS data structures are linked lists. Having a robust and efficient set of linked-list types that we can use as a foundation for building the rest of our kernel types allows u16 to keep our RTOS code efficient and logically-separated.

```
#include "kerneltypes.h"
```

Classes

- class [Mark3::LinkListNode](#)

The [LinkListNode](#) Class Basic linked-list node data structure. This data is managed by the linked-list class types, and can be used transparently between them.

- class [Mark3::TypedLinkListNode< T >](#)

The `TypedLinkedListNode` class provides a linked-list node type for a specified object type. This can be used with typed link-list data structures to manage lists of objects without having to static-cast between the base type and the derived class.

- class `Mark3::LinkList`

The [LinkedList](#) Class Abstract-data-type from which all other linked-lists are derived.

- class `Mark3::DoubleLinkedList`

The [DoubleLinkedList](#) Class Doubly-linked-list data type, inherited from the base [LinkedList](#) type.

- class **Mark3::CircularLinkList**

The `CircularLinkList` class Circular-linked-list data type, inherited from the base `LinkList` type.

- class `Mark3::TypedDoubleLinkList< T >`

The `TypedDoubleLinkedList` Class Doubly-linked-list data type, inherited from the base `LinkedList` type, and templated for use with linked-list-node derived data-types.

- class `Mark3::TypedCircularLinkList< T >`

The `TypedCircularLinkedList` Class Circular-linked-list data type, inherited from the base `LinkedList` type, and templated for use with linked-list-node derived data-types.

Namespaces

- Mark3

20.93.1 Detailed Description

Core linked-list declarations, used by all kernel list types At the heart of RTOS data structures are linked lists. Having a robust and efficient set of linked-list types that we can use as a foundation for building the rest of our kernel types allows u16 to keep our RTOS code efficient and logically-separated.

So what data types rely on these linked-list classes?

- Threads
- ThreadLists
- The Scheduler
- Timers, -The Timer Scheduler
- Blocking objects (Semaphores, Mutexes, etc...)

Pretty much everything in the kernel uses these linked lists. By having objects inherit from the base linked-list node type, we're able to leverage the double and circular linked-list classes to manager virtually every object type in the system without duplicating code. These functions are very efficient as well, allowing for very deterministic behavior in our code.

Definition in file [ll.h](#).

20.94 11.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00041 #pragma once
00042 #include "kerneltypes.h"
00043
00044 namespace Mark3

```

```

00045 {
00046 //-----
00052 class LinkList;
00053 class DoubleLinkList;
00054 class CircularLinkList;
00055
00056 //-----
00062 class LinkListNode
00063 {
00064 protected:
00065     LinkListNode* next;
00066     LinkListNode* prev;
00067
00068     LinkListNode() {}
00069
00075     void ClearNode();
00076
00077 public:
00085     LinkListNode* GetNext(void) { return next; }
00093     LinkListNode* GetPrev(void) { return prev; }
00094     friend class LinkList;
00095     friend class DoubleLinkList;
00096     friend class CircularLinkList;
00097 };
00098
00099 //-----
00107 template <typename T>
00108 class TypedLinkListNode : public LinkListNode
00109 {
00110 public:
00111     T* GetNext() { return static_cast<T*>(LinkListNode::GetNext()); }
00112     T* GetPrev() { return static_cast<T*>(LinkListNode::GetPrev()); }
00113 };
00114 //-----
00119 class LinkList
00120 {
00121 protected:
00122     LinkListNode* m_pclHead;
00123     LinkListNode* m_pclTail;
00124
00125 public:
00131     void Init()
00132     {
00133         m_pclHead = nullptr;
00134         m_pclTail = nullptr;
00135     }
00136
00144     LinkListNode* GetHead() { return m_pclHead; }
00145
00153     void SetHead(LinkListNode* pclNode_) { m_pclHead = pclNode_; }
00154
00162     LinkListNode* GetTail() { return m_pclTail; }
00163
00171     void SetTail(LinkListNode* pclNode_) { m_pclTail = pclNode_; }
00172 };
00173
00174 //-----
00179 class DoubleLinkList : public LinkList
00180 {
00181 public:
00182     void* operator new(size_t sz, void* pv) { return reinterpret_cast<
DoubleLinkList*>(pv); };
00188     DoubleLinkList()
00189     {
00190         m_pclHead = nullptr;
00191         m_pclTail = nullptr;
00192     }
00193
00201     void Add(LinkListNode* node_);
00202
00210     void Remove(LinkListNode* node_);
00211 };
00212
00213 //-----
00218 class CircularLinkList : public LinkList
00219 {
00220 public:
00221     void* operator new(size_t sz, void* pv) { return (CircularLinkList*)pv; };
00222     CircularLinkList()
00223     {
00224         m_pclHead = nullptr;
00225         m_pclTail = nullptr;
00226     }
00227
00234     void Add(LinkListNode* node_);
00235
00242     void Remove(LinkListNode* node_);

```

```

00243
00249     void PivotForward();
00250
00256     void PivotBackward();
00257
00266     void InsertNodeBefore(LinkListNode* node_, LinkListNode* insert_);
00267 };
00268
00269 //-----
00275 template <typename T>
00276 class TypedDoubleLinkedList : public DoubleLinkedList
00277 {
00278 public:
00279     void* operator new(size_t sz, void* pv) { return reinterpret_cast<
TypedDoubleLinkedList<T>*>(pv); }
00280
00281     TypedDoubleLinkedList<T>() {
00282         DoubleLinkedList::Init();
00283     }
00284
00291     T* GetHead() { return static_cast<T*>(DoubleLinkedList::GetHead()); }
00292
00299     void SetHead(T* pclNode_) { DoubleLinkedList::SetHead(pclNode_); }
00300
00307     T* GetTail() { return static_cast<T*>(DoubleLinkedList::GetTail()); }
00308
00315     void SetTail(T* pclNode_) { DoubleLinkedList::SetTail(pclNode_); }
00316
00323     void Add(T* pNode_)
00324     {
00325         DoubleLinkedList::Add(pNode_);
00326     }
00327
00334     void Remove(T* pNode_)
00335     {
00336         DoubleLinkedList::Remove(pNode_);
00337     }
00338 };
00339
00340 //-----
00346 template <typename T>
00347 class TypedCircularLinkedList : public CircularLinkedList
00348 {
00349 public:
00350     void* operator new(size_t sz, void* pv) { return reinterpret_cast<
TypedCircularLinkedList<T>*>(pv); }
00351
00352     TypedCircularLinkedList<T>() {
00353         CircularLinkedList::Init();
00354     }
00355
00362     T* GetHead() { return static_cast<T*>(CircularLinkedList::GetHead()); }
00363
00370     void SetHead(T* pclNode_) { CircularLinkedList::SetHead(pclNode_); }
00371
00378     T* GetTail() { return static_cast<T*>(CircularLinkedList::GetTail()); }
00379
00386     void SetTail(T* pclNode_) { CircularLinkedList::SetTail(pclNode_); }
00387
00394     void Add(T* pNode_)
00395     {
00396         CircularLinkedList::Add(pNode_);
00397     }
00398
00405     void Remove(T* pNode_)
00406     {
00407         CircularLinkedList::Remove(pNode_);
00408     }
00409
00418     void InsertNodeBefore(T* pNode_, T* pInsert_)
00419     {
00420         CircularLinkedList::InsertNodeBefore(pNode_, pInsert_);
00421     }
00422 };
00423 } // namespace Mark3

```

20.95 /home/moslevin/projects/m3-repo/kernel/src/public/lockguard.h File Reference

Mutex RAIL helper class.

```
#include "mark3.h"
```



```

00174     }
00175
00176     bool IsFull(void) { return (GetFreeSlots() == 0); }
00177     bool IsEmpty(void) { return (GetFreeSlots() == m_ul6Count); }
00178
00179 private:
00180     void* GetHeadPointer(void)
00181     {
00182         auto uAddr = reinterpret_cast<K_ADDR>(m_pvBuffer);
00183         uAddr += static_cast<K_ADDR>(m_ul6ElementSize) * static_cast<K_ADDR>(
00184             m_ul6Head);
00185         return reinterpret_cast<void*>(uAddr);
00186     }
00187
00188     void* GetTailPointer(void)
00189     {
00190         auto uAddr = reinterpret_cast<K_ADDR>(m_pvBuffer);
00191         uAddr += static_cast<K_ADDR>(m_ul6ElementSize) * static_cast<K_ADDR>(
00192             m_ul6Tail);
00193         return reinterpret_cast<void*>(uAddr);
00194     }
00195
00196     void CopyData(const void* src_, void* dst_, uint16_t len_)
00197     {
00198         auto* u8Src = reinterpret_cast<const uint8_t*>(src_);
00199         auto* u8Dst = reinterpret_cast<uint8_t*>(dst_);
00200         while (len_-- > 0) { *u8Dst++ = *u8Src++; }
00201     }
00202
00203     void MoveTailForward(void)
00204     {
00205         m_ul6Tail++;
00206         if (m_ul6Tail == m_ul6Count) {
00207             m_ul6Tail = 0;
00208         }
00209     }
00210
00211     void MoveHeadForward(void)
00212     {
00213         m_ul6Head++;
00214         if (m_ul6Head == m_ul6Count) {
00215             m_ul6Head = 0;
00216         }
00217     }
00218
00219     void MoveTailBackward(void)
00220     {
00221         if (m_ul6Tail == 0) {
00222             m_ul6Tail = m_ul6Count;
00223         }
00224         m_ul6Tail--;
00225     }
00226
00227     void MoveHeadBackward(void)
00228     {
00229         if (m_ul6Head == 0) {
00230             m_ul6Head = m_ul6Count;
00231         }
00232         m_ul6Head--;
00233     }
00234
00235     bool Send_i(const void* pvData_, bool bTail_, uint32_t u32TimeoutMS_);
00236
00237     bool Receive_i(void* pvData_, bool bTail_, uint32_t u32WaitTimeMS_);
00238
00239     uint16_t m_ul6Head;
00240     uint16_t m_ul6Tail;
00241
00242     uint16_t m_ul6Count;
00243     volatile uint16_t m_ul6Free;
00244
00245     uint16_t m_ul6ElementSize;
00246     const void* m_pvBuffer;
00247
00248     Semaphore m_clRecvSem;
00249     Semaphore m_clSendSem;
00250 };
00251 // namespace Mark3

```

20.99 /home/moslevin/projects/m3-repo/kernel/src/public/manual.h File Reference

20.101.1 Detailed Description

Single include file given to users of the [Mark3](#) Kernel API.

Definition in file [mark3.h](#).

20.102 mark3.h

```

00001  /*=====
00002
00003  _____
00004  /         \  /         \  /         \  /         \  /         \
00005  |         |  |         |  |         |  |         |  |         |
00006  \         /  \         /  \         /  \         /  \         /
00007  _____
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  =====*/
00020  #pragma once
00021
00022  #include "mark3cfg.h"
00023  #include "threadport.h"
00024  #include "criticalsection.h"
00025  #include "criticalguard.h"
00026
00027  #include "kerneltypes.h"
00028  #include "kerneldebug.h"
00029
00030  #include "ithreadport.h"
00031  #include "kernelswi.h"
00032  #include "kerneltimer.h"
00033
00034  #include "kernel.h"
00035  #include "thread.h"
00036  #include "timerlist.h"
00037
00038  #include "ksemaphore.h"
00039  #include "mutex.h"
00040  #include "lockguard.h"
00041  #include "eventflag.h"
00042  #include "message.h"
00043  #include "notify.h"
00044  #include "mailbox.h"
00045  #include "readerwriter.h"
00046  #include "condvar.h"
00047
00048  #include "atomic.h"
00049
00050  #include "profile.h"
00051  #include "autoalloc.h"
00052  #include "priomap.h"
00053
00054  #include "threadlist.h"
00055  #include "threadlistlist.h"
00056
00057  #include "schedulerguard.h"
00058
00059  #include "coroutine.h"
00060  #include "colist.h"
00061  #include "cosched.h"
00062

```

20.103 /home/moslevin/projects/m3-repo/kernel/src/public/mark3cfg.h File Reference

[Mark3](#) Kernel Configuration This file is used to configure the kernel for your specific application in order to provide the optimal set of features for a given use case.

```
#include "portcfg.h"
```

Macros

- `#define` [KERNEL_DEBUG](#) (0)
- `#define` [KERNEL_STACK_CHECK](#) (1)
- `#define` [KERNEL_NAMED_THREADS](#) (1)
- `#define` [KERNEL_EVENT_FLAGS](#) (1)
- `#define` [KERNEL_CONTEXT_SWITCH_CALLOUT](#) (1)
- `#define` [KERNEL_THREAD_CREATE_CALLOUT](#) (1)
- `#define` [KERNEL_THREAD_EXIT_CALLOUT](#) (1)
- `#define` [KERNEL_ROUND_ROBIN](#) (1)
- `#define` [KERNEL_EXTENDED_CONTEXT](#) (1)

include CPU/Port specific configuration options

20.103.1 Detailed Description

[Mark3](#) Kernel Configuration This file is used to configure the kernel for your specific application in order to provide the optimal set of features for a given use case.

Note: in the R7 and beyond version of the kernel, all options are enabled by default. As a result, the only configuration options presented are now located within the architecture-specific "portcfg.h".

Definition in file [mark3cfg.h](#).

20.103.2 Macro Definition Documentation

20.103.2.1 KERNEL_CONTEXT_SWITCH_CALLOUT

```
#define KERNEL_CONTEXT_SWITCH_CALLOUT (1)
```

When enabled, this feature allows a user to define a callback to be executed whenever a context switch occurs. Enabling this provides a means for a user to track thread statistics, but it does result in additional overhead during a context switch.

Definition at line 70 of file [mark3cfg.h](#).

20.103.2.2 KERNEL_DEBUG

```
#define KERNEL_DEBUG (0)
```

Enable kernel asserts at runtime.

Definition at line 30 of file [mark3cfg.h](#).

20.103.2.3 KERNEL_EVENT_FLAGS

```
#define KERNEL_EVENT_FLAGS (1)
```

This flag enables the event-flags synchronization object. This feature allows threads to be blocked, waiting on specific condition bits to be set or cleared on an EventFlag object.

While other synchronization objects are enabled by default, this one is configurable because it impacts the Thread object's member data.

Definition at line 62 of file [mark3cfg.h](#).

20.103.2.4 KERNEL_EXTENDED_CONTEXT

```
#define KERNEL_EXTENDED_CONTEXT (1)
```

include CPU/Port specific configuration options

Provide a special data pointer in the thread object, which may be used to add additional context to a thread. Typically this would be used to implement thread-local-storage.

Definition at line 97 of file [mark3cfg.h](#).

20.103.2.5 KERNEL_NAMED_THREADS

```
#define KERNEL_NAMED_THREADS (1)
```

Enabling this provides the Thread::SetName() and Thread::GetName() methods, allowing for each thread to be named with a null-terminated const char* string.

Note: the string passed to Thread::SetName() must persist for the lifetime of the thread

Definition at line 52 of file [mark3cfg.h](#).

20.103.2.6 KERNEL_ROUND_ROBIN

```
#define KERNEL_ROUND_ROBIN (1)
```

Enable round-robin scheduling within each priority level. When selected, this results in a small performance hit during context switching and in the system tick handler, as a special software timer is used to manage the running thread's quantum. Can be disabled to optimize performance if not required.

Definition at line 90 of file [mark3cfg.h](#).

20.105 /home/moslevin/projects/m3-repo/kernel/src/public/message.h File Reference

Inter-thread communication via message-passing Embedded systems guru Jack Ganssle once said that without a robust form of interprocess communications (IPC), an RTOS is just a toy. [Mark3](#) implements a form of IPC to provide safe and flexible messaging between threads.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ll.h"
#include "ksemaphore.h"
#include "timerlist.h"
```

Classes

- class [Mark3::Message](#)

the [Message](#) class. This object provides threadsafe message-based IPC services based on exchange of objects containing a data pointer and minimal application-defined metadata. Messages are to be allocated/produced by the sender, and deallocated/consumed by the receiver.

- class [Mark3::MessagePool](#)

The [MessagePool](#) Class The [MessagePool](#) class implements a simple allocator for message objects exchanged between threads. The sender allocates (pop's) messages, then sends them to the receiver. Upon receipt, it is the receiver's responsibility to deallocate (push) the message back to the pool.

- class [Mark3::MessageQueue](#)

The [MessageQueue](#) class. Implements a mechanism used to send/receive data between threads. Allows threads to block, waiting for messages to be sent from other contexts.

Namespaces

- [Mark3](#)

20.105.1 Detailed Description

Inter-thread communication via message-passing Embedded systems guru Jack Ganssle once said that without a robust form of interprocess communications (IPC), an RTOS is just a toy. [Mark3](#) implements a form of IPC to provide safe and flexible messaging between threads.

using kernel-managed IPC offers significant benefits over other forms of data sharing (i.e. Global variables) in that it avoids synchronization issues and race conditions common to the practice. using IPC also enforces a more disciplined coding style that keeps threads decoupled from one another and minimizes global data, preventing careless and hard-to-debug errors.

20.105.2 using Messages, Queues, and the Global Message Pool

```
// Declare a message queue shared between two threads
MessageQueue my_queue;

int main()
{
    ...
    // Initialize the message queue
    my_queue.init();
    ...
}

void Thread1()
{
    // Example TX thread - sends a message every 10ms
    while(1)
    {
        // Grab a message from the global message pool
        Message *tx_message = GlobalMessagePool::Pop();

        // Set the message data/parameters
        tx_message->SetCode( 1234 );
        tx_message->SetData( nullptr );

        // Send the message on the queue.
        my_queue.Send( tx_message );
        Thread::Sleep(10);
    }
}

void Thread2()
{
    while()
    {
        // Blocking receive - wait until we have messages to process
        Message *rx_message = my_queue.Recv();

        // Do something with the message data...

        // Return back into the pool when done
        GlobalMessagePool::Push(rx_message);
    }
}
```

Definition in file [message.h](#).

20.106 message.h

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slewin, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00078 #pragma once
00079
00080 #include "kerneltypes.h"
00081 #include "mark3cfg.h"
00082
00083 #include "ll.h"
00084 #include "ksemaphore.h"
00085 #include "timerlist.h"
00086
00087 namespace Mark3
00088 {
00089 //-----
00097 class Message : public TypedLinkedListNode<Message>
00098 {
00099 public:
00100     void* operator new(size_t sz, void* pv) { return reinterpret_cast<Message*>(pv); }
```

```

00105     void Init()
00106     {
00107         ClearNode();
00108         m_pvData = nullptr;
00109         m_ul6Code = 0;
00110     }
00111
00112     void SetData(void* pvData_) { m_pvData = pvData_; }
00113     void* GetData() { return m_pvData; }
00114     void SetCode(uint16_t u16Code_) { m_ul6Code = u16Code_; }
00115     uint16_t GetCode() { return m_ul6Code; }
00116
00117 private:
00118     void* m_pvData;
00119     uint16_t m_ul6Code;
00120 };
00121
00122 //-----
00123 class MessagePool
00124 {
00125 public:
00126     void* operator new(size_t sz, void* pv) { return (MessagePool*)pv; }
00127     ~MessagePool() {}
00128     void Init();
00129
00130     void Push(Message* pclMessage_);
00131
00132     Message* Pop();
00133
00134     Message* GetHead();
00135 private:
00136     TypedDoubleLinkedList<Message> m_clList;
00137 };
00138
00139 //-----
00140 class MessageQueue
00141 {
00142 public:
00143     void* operator new(size_t sz, void* pv) { return (MessageQueue*)pv; }
00144     ~MessageQueue() {}
00145
00146     void Init();
00147
00148     Message* Receive();
00149
00150     Message* Receive(uint32_t u32TimeWaitMS_);
00151
00152     void Send(Message* pclSrc_);
00153
00154     uint16_t GetCount();
00155 private:
00156     Message* Receive_i(uint32_t u32TimeWaitMS_);
00157     Semaphore m_clSemaphore;
00158     TypedDoubleLinkedList<Message> m_clLinkList;
00159 };
00160
00161 } // namespace Mark3

```

20.107 /home/moslevin/projects/m3-repo/kernel/src/public/mutex.h File Reference

Mutual exclusion class declaration Resource locks are implemented using mutual exclusion semaphores (Mutex_t). Protected blocks can be placed around any resource that may only be accessed by one thread at a time. If additional threads attempt to access the protected resource, they will be placed in a wait queue until the resource becomes available. When the resource becomes available, the thread with the highest original priority claims the resource and is activated. Priority inheritance is included in the implementation to prevent priority inversion. Always ensure that you claim and release your mutex objects consistently, otherwise you may end up with a deadlock scenario that's hard to debug.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "blocking.h"

```


Classes

- class `Mark3::Mutex`

The **Mutex** Class. Class providing Mutual-exclusion locks, based on **BlockingObject**.

Namespaces

- Mark3

20.107.1 Detailed Description

Mutual exclusion class declaration Resource locks are implemented using mutual exclusion semaphores (Mutex_t). Protected blocks can be placed around any resource that may only be accessed by one thread at a time. If additional threads attempt to access the protected resource, they will be placed in a wait queue until the resource becomes available. When the resource becomes available, the thread with the highest original priority claims the resource and is activated. Priority inheritance is included in the implementation to prevent priority inversion. Always ensure that you claim and release your mutex objects consistently, otherwise you may end up with a deadlock scenario that's hard to debug.

20.107.2 Initializing

Initializing a mutex object by calling:

```
clMutex.Init();
```

20.107.3 Resource protection example

```
clMutex.Claim();
...
<resource protected block>
...
clMutex.Release();
```

Definition in file [mutex.h](#).

20.108 mutex.h

```
00001 /*=====
00002
00003
00004 |-----|-----|-----|-----|
00005 | \ / | \ / | \ / | \ / | \ / |
00006 | / \ | / \ | / \ | / \ | / \ |
00007 |_____|_____|_____|_____|_____*/
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00049 #pragma once
00050
00051 #include "kerneltypes.h"
00052 #include "mark3cfg.h"
00053
```

```

00054 #include "blocking.h"
00055
00056 namespace Mark3
00057 {
00058 //-----
00063 class Mutex : public BlockingObject
00064 {
00065 public:
00066     void* operator new(size_t sz, void* pv) { return reinterpret_cast<Mutex*>(pv); };
00067     ~Mutex();
00068
00076     void Init(bool bRecursive_ = true);
00077
00094     void Claim();
00095
00105     bool Claim(uint32_t u32WaitTimeMS_);
00106
00118     void WakeMe(Thread* pclOwner_);
00119
00139     void Release();
00140
00141 private:
00147     uint8_t WakeNext();
00148
00156     bool Claim_i(uint32_t u32WaitTimeMS_);
00157
00158     uint8_t      m_u8Recurse;
00159     bool         m_bReady;
00160     bool         m_bRecursive;
00161     PORT_PRIO_TYPE m_uMaxPri;
00162     Thread*      m_pclOwner;
00163 };
00164 } // namespace Mark3

```

20.109 /home/moslevin/projects/m3-repo/kernel/src/public/notify.h File Reference

Lightweight thread notification - blocking object.

```

#include "mark3cfg.h"
#include "blocking.h"

```

Classes

- class [Mark3::Notify](#)

The [Notify](#) class. This class provides a blocking object type that allows one or more threads to wait for an event to occur before resuming operation.

Namespaces

- [Mark3](#)

20.109.1 Detailed Description

Lightweight thread notification - blocking object.

Definition in file [notify.h](#).

```

00001  /*=====
00002
00003
00004  |-----|-----|-----|-----|-----|
00005  | \      / | \      / | \      / | \      / | \      / |
00006  |  \    /  |  \    /  |  \    /  |  \    /  |  \    /  |
00007  |   \  /   |   \  /   |   \  /   |   \  /   |   \  /   |
00008  |_____|_____|_____|_____|_____|
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  =====*/
00021  #pragma once
00022
00023  #include "mark3cfg.h"
00024  #include "blocking.h"
00025
00026  namespace Mark3
00027  {
00033  class Notify : public BlockingObject
00034  {
00035  public:
00036      void* operator new(size_t sz, void* pv) { return reinterpret_cast<Notify*>(pv); };
00037      ~Notify();
00038
00043      void Init(void);
00044
00053      void Signal(void);
00054
00063      void Wait(bool* pbFlag_);
00064
00076      bool Wait(uint32_t u32WaitTimeMS_, bool* pbFlag_);
00077
00086      void WakeMe(Thread* pclChosenOne_);
00087
00088  private:
00089      bool m_bPending;
00090  };
00091  } // namespace Mark3

```

Defines the reason codes thrown when a kernel panic occurs.

- #define PANIC_ASSERT_FAILED (1)
- #define PANIC_LIST_UNLINK_FAILED (2)
- #define PANIC_STACK_SLACK_VIOLATED (3)
- #define PANIC_AUTO_HEAP_EXHAUSTED (4)
- #define PANIC_POWERMAN_EXHAUSTED (5)
- #define PANIC_NO_READY_THREADS (6)
- #define PANIC_RUNNING_THREAD_DESCOPED (7)
- #define PANIC_ACTIVE_SEMAPHORE_DESCOPED (8)
- #define PANIC_ACTIVE_MUTEX_DESCOPED (9)
- #define PANIC_ACTIVE_EVENTFLAG_DESCOPED (10)
- #define PANIC_ACTIVE_NOTIFY_DESCOPED (11)
- #define PANIC_ACTIVE_MAILBOX_DESCOPED (12)
- #define PANIC_ACTIVE_TIMER_DESCOPED (13)
- #define PANIC_ACTIVE_COROUTINE_DESCOPED (14)

20.111.1 Detailed Description

Defines the reason codes thrown when a kernel panic occurs.

Definition in file [paniccodes.h](#).

20.111.2 Macro Definition Documentation

20.111.2.1 PANIC_ACTIVE_COROUTINE_DESCOPE

```
#define PANIC_ACTIVE_COROUTINE_DESCOPE (14)
```

Definition at line 35 of file [paniccodes.h](#).

20.111.2.2 PANIC_ACTIVE_EVENTFLAG_DESCOPE

```
#define PANIC_ACTIVE_EVENTFLAG_DESCOPE (10)
```

Definition at line 31 of file [paniccodes.h](#).

20.111.2.3 PANIC_ACTIVE_MAILBOX_DESCOPE

```
#define PANIC_ACTIVE_MAILBOX_DESCOPE (12)
```

Definition at line 33 of file [paniccodes.h](#).

20.111.2.4 PANIC_ACTIVE_MUTEX_DESCOPE

```
#define PANIC_ACTIVE_MUTEX_DESCOPE (9)
```

Definition at line 30 of file [paniccodes.h](#).

20.111.2.5 PANIC_ACTIVE_NOTIFY_DESCOPE

```
#define PANIC_ACTIVE_NOTIFY_DESCOPE (11)
```

Definition at line 32 of file [paniccodes.h](#).

20.111.2.6 PANIC_ACTIVE_SEMAPHORE_DESCOPE

```
#define PANIC_ACTIVE_SEMAPHORE_DESCOPE (8)
```

Definition at line 29 of file [paniccodes.h](#).

20.111.2.7 PANIC_ACTIVE_TIMER_DESCOPE

```
#define PANIC_ACTIVE_TIMER_DESCOPE (13)
```

Definition at line 34 of file [paniccodes.h](#).

20.111.2.8 PANIC_ASSERT_FAILED

```
#define PANIC_ASSERT_FAILED (1)
```

Definition at line 22 of file [paniccodes.h](#).

20.111.2.9 PANIC_AUTO_HEAP_EXHAUSTED

```
#define PANIC_AUTO_HEAP_EXHAUSTED (4)
```

Definition at line 25 of file [paniccodes.h](#).

20.111.2.10 PANIC_LIST_UNLINK_FAILED

```
#define PANIC_LIST_UNLINK_FAILED (2)
```

Definition at line 23 of file [paniccodes.h](#).

20.111.2.11 PANIC_NO_READY_THREADS

```
#define PANIC_NO_READY_THREADS (6)
```

Definition at line 27 of file [paniccodes.h](#).

20.111.2.12 PANIC_POWERMAN_EXHAUSTED

```
#define PANIC_POWERMAN_EXHAUSTED (5)
```

Definition at line 26 of file [paniccodes.h](#).

20.111.2.13 PANIC_RUNNING_THREAD_DESCOPE

```
#define PANIC_RUNNING_THREAD_DESCOPE (7)
```

Definition at line 28 of file [paniccodes.h](#).

20.111.2.14 PANIC_STACK_SLACK_VIOLATED

```
#define PANIC_STACK_SLACK_VIOLATED (3)
```

Definition at line 24 of file [paniccodes.h](#).

20.112 paniccodes.h

```
00001 /*=====
00002
00003
00004 | | | | | | | | | | | | | | | | | |
00005 | | | | | | | | | | | | | | | | | |
00006 | | | | | | | | | | | | | | | | | |
00007 | | | | | | | | | | | | | | | | | |
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00020 #pragma once
00021
00022 #define PANIC_ASSERT_FAILED (1)
00023 #define PANIC_LIST_UNLINK_FAILED (2)
00024 #define PANIC_STACK_SLACK_VIOLATED (3)
00025 #define PANIC_AUTO_HEAP_EXHAUSTED (4)
00026 #define PANIC_POWERMAN_EXHAUSTED (5)
00027 #define PANIC_NO_READY_THREADS (6)
00028 #define PANIC_RUNNING_THREAD_DESCOPE (7)
00029 #define PANIC_ACTIVE_SEMAPHORE_DESCOPE (8)
00030 #define PANIC_ACTIVE_MUTEX_DESCOPE (9)
00031 #define PANIC_ACTIVE_EVENTFLAG_DESCOPE (10)
00032 #define PANIC_ACTIVE_NOTIFY_DESCOPE (11)
00033 #define PANIC_ACTIVE_MAILBOX_DESCOPE (12)
00034 #define PANIC_ACTIVE_TIMER_DESCOPE (13)
00035 #define PANIC_ACTIVE_COROUTINE_DESCOPE (14)
```

20.113 /home/moslevin/projects/m3-repo/kernel/src/public/priomap.h File Reference

Priority map data structure.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "priomapl1.h"
#include "priomapl2.h"
```

Namespaces

- [Mark3](#)

Typedefs

- using [Mark3::PriorityMap](#) = PriorityMapL1< [PORT_PRIO_TYPE](#), [KERNEL_NUM_PRIORITIES](#) >

20.113.1 Detailed Description

Priority map data structure.

Definition in file [priomap.h](#).

20.114 priomap.h

```

00001 /*=====
00002
00003 | | | | | | | | | | | | | | | | | |
00004 | | | | | | | | | | | | | | | | | |
00005 | | | | | | | | | | | | | | | | | |
00006 | | | | | | | | | | | | | | | | | |
00007 | | | | | | | | | | | | | | | | | |
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00019 #pragma once
00020
00021 #include "kerneltypes.h"
00022 #include "mark3cfg.h"
00023
00024 #include "priomapl1.h"
00025 #include "priomapl2.h"
00026
00027 namespace Mark3 {
00028     #if KERNEL_NUM_PRIORITIES <= (PORT_PRIO_MAP_WORD_SIZE * 8u)
00029     using PriorityMap =
00030         PriorityMapL1<PORT_PRIO_TYPE, KERNEL_NUM_PRIORITIES>;
00030     #else
00031     using PriorityMap =
00032         PriorityMapL2<PORT_PRIO_TYPE, KERNEL_NUM_PRIORITIES>;
00032     #endif
00033 } // namespace Mark3

```

20.115 /home/moslevin/projects/m3-repo/kernel/src/public/priomapl1.h File Reference

1-Level bitmap allocator template-class used for scheduler implementation

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "threadport.h"

```



```

00092 private:
00093     static inline T PrioBit(T prio) { return prio & m_uXPrioMapBitMask; }
00094
00095     static inline T PrioMapWordIndex(T prio) { return prio >>
m_uXPrioMapWordShift; }
00096
00097     static inline T PriorityFromBitmap(T uXPrio_)
00098     {
00099 #if PORT_USE_HW_CLZ
00100         // Support hardware-accelerated Count-leading-zeros instruction
00101         return m_uXPrioMapBits - PORT_CLZ(uXPrio_);
00102 #else
00103         // Default un-optimized count-leading zeros operation
00104         T uXMask = 1 << (m_uXPrioMapBits - 1);
00105         auto      u8Zeros = T { 0 };
00106
00107         while (uXMask) {
00108             if (uXMask & uXPrio_) {
00109                 return (m_uXPrioMapBits - u8Zeros);
00110             }
00111             uXMask >>= 1;
00112             u8Zeros++;
00113         }
00114         return 0;
00115 #endif
00116     }
00117
00118     static constexpr size_t m_uXPrioMapShiftLUT[9] = {0, 3, 4, 0, 5, 0, 0, 0, 6};
00119     static constexpr auto m_uXPrioMapWordShift = T { m_uXPrioMapShiftLUT[sizeof(T)] };
00120     static constexpr auto m_uXPrioMapBits      = T { 8 * sizeof(T) };
00121     static constexpr auto m_uXPrioMapBitMask = T { (1 <<
m_uXPrioMapWordShift) - 1 };
00122
00123     T m_uXPriorityMap;
00124 };
00125
00126 } // namespace Mark3

```

20.117 /home/moslevin/projects/m3-repo/kernel/src/public/priomapl2.h File Reference

2-Level priority allocator template-class used for scheduler implementation

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "threadport.h"

```

Classes

- class [Mark3::PriorityMapL2< T, C >](#)

The *PriorityMapL2* class This class implements a priority bitmap data structure. Each bit in the objects internal storage represents a priority. When a bit is set, it indicates that something is scheduled at the bit's corresponding priority, when a bit is clear it indicates that no entities are scheduled at that priority. This object provides the fundamental logic required to implement efficient priority-based scheduling for the thread + coroutine schedulers in the kernel.

Namespaces

- [Mark3](#)

20.117.1 Detailed Description

2-Level priority allocator template-class used for scheduler implementation

Definition in file [priomapl2.h](#).

20.118 priomapl2.h

```

00001  /*=====
00002
00003  00004  00005  00006  00007  00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  =====*/
00020  #pragma once
00021
00022  #include "kerneltypes.h"
00023  #include "mark3cfg.h"
00024  #include "threadport.h"
00025
00026  namespace Mark3
00027  {
00028  //-----
00049  template <typename T, size_t C>
00050  class PriorityMapL2
00051  {
00052  public:
00057      PriorityMapL2()
00058      {
00059          m_uXPriorityMapL2 = 0;
00060          for (auto i = PORT_PRIO_TYPE{0}; i < m_uXPrioMapNumWords; i++) {
00061              m_auXPriorityMap[i] = 0; }
00062
00068      void Set(T uXPrio_)
00069      {
00070          auto uXPrioBit = PrioBit(uXPrio_);
00071          auto uXWordIdx = PrioMapWordIndex(uXPrio_);
00072
00073          m_auXPriorityMap[uXWordIdx] |= (1 << uXPrioBit);
00074          m_uXPriorityMapL2 |= (1 << uXWordIdx);
00075      }
00076
00082      void Clear(T uXPrio_)
00083      {
00084          auto uXPrioBit = PrioBit(uXPrio_);
00085          auto uXWordIdx = PrioMapWordIndex(uXPrio_);
00086
00087          m_auXPriorityMap[uXWordIdx] &= ~(1 << uXPrioBit);
00088          if (!m_auXPriorityMap[uXWordIdx]) {
00089              m_uXPriorityMapL2 &= ~(1 << uXWordIdx);
00090          }
00091      }
00092
00100      T HighestPriority(void)
00101      {
00102          auto uXMapIdx = PriorityFromBitmap(m_uXPriorityMapL2);
00103          if (!uXMapIdx) {
00104              return 0;
00105          }
00106          uXMapIdx--;
00107          auto uXPrio = PriorityFromBitmap(m_auXPriorityMap[uXMapIdx]);
00108          uXPrio += (uXMapIdx * m_uXPrioMapBits);
00109          return uXPrio;
00110      }
00111
00112  private:
00113      static inline T PrioBit(T prio) { return prio & m_uXPrioMapBitMask; }
00114
00115      static inline T PrioMapWordIndex(T prio) { return prio >>
00116          m_uXPrioMapWordShift; }
00117
00117      static inline T PriorityFromBitmap(T uXPrio_)
00118      {
00119      #if PORT_USE_HW_CLZ
00120          // Support hardware-accelerated Count-leading-zeros instruction
00121          return m_uXPrioMapBits - PORT_CLZ(uXPrio_);
00122      #else
00123          // Default un-optimized count-leading zeros operation
00124          T uXMask = 1 << (m_uXPrioMapBits - 1);
00125          auto u8Zeros = T{ 0 };
00126
00127          while (uXMask) {
00128              if (uXMask & uXPrio_) {
00129                  return (m_uXPrioMapBits - u8Zeros);

```

```

00130         }
00131
00132         uXMask >>= 1;
00133         u8Zeros++;
00134     }
00135     return 0;
00136 #endif
00137 }
00138
00139 static constexpr size_t m_uXPrioMapShiftLUT[9] = {0, 3, 4, 0, 5, 0, 0, 0, 6};
00140 static constexpr auto m_uXPrioMapWordShift = T { m_uXPrioMapShiftLUT[sizeof(T)] };
00141 static constexpr auto m_uXPrioMapBits = T { 8 * sizeof(T) };
00142 static constexpr auto m_uXPrioMapBitMask = T { (1 <<
m_uXPrioMapWordShift) - 1 };
00143
00144 // Required size of the bitmap array in words
00145 static constexpr auto m_uXPrioMapNumWords
00146     = T { (C + (m_uXPrioMapBits - 1)) / m_uXPrioMapBits };
00147
00148 T m_uXPriorityMap[m_uXPrioMapNumWords];
00149 T m_uXPriorityMapL2;
00150 };
00151 } // namespace Mark3

```

20.119 /home/moslevin/projects/m3-repo/kernel/src/public/profile.h File Reference

High-precision profiling timers Enables the profiling and instrumentation of performance-critical code. Multiple timers can be used simultaneously to enable system-wide performance metrics to be computed in a lightweight manner.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ll.h"

```

Classes

- class [Mark3::ProfileTimer](#)

Profiling timer. This class is used to perform high-performance profiling of code to see how `int32_t` certain operations take. useful in instrumenting the performance of key algorithms and time-critical operations to ensure real-timer behavior.

Namespaces

- [Mark3](#)

20.119.1 Detailed Description

High-precision profiling timers Enables the profiling and instrumentation of performance-critical code. Multiple timers can be used simultaneously to enable system-wide performance metrics to be computed in a lightweight manner.

Usage:

```

ProfileTimer clMyTimer;
int i;

clMyTimer.Init();

// Profile the same block of code ten times
for (i = 0; i < 10; i++)
{
    clMyTimer.Start();
    ...
    //Block of code to profile
    ...
    clMyTimer.Stop();
}

// Get the average execution time of all iterations
u32AverageTimer = clMyTimer.GetAverage();

// Get the execution time from the last iteration
u32LastTimer = clMyTimer.GetCurrent();

```

Definition in file [profile.h](#).

20.120 profile.h

```

00001 /*=====
00002
00003  _ _ _ _ _
00004  / V \ / V \ / V \ / V \ / V \
00005  | / \ | / \ | / \ | / \ | / \ |
00006  | / \ | / \ | / \ | / \ | / \ |
00007  _ _ _ _ _
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00051 #pragma once
00052
00053 #include "kerneltypes.h"
00054 #include "mark3cfg.h"
00055 #include "ll.h"
00056
00057 namespace Mark3
00058 {
00067 class ProfileTimer
00068 {
00069 public:
00075     void Init();
00076
00082     void Start();
00083
00089     void Stop();
00090
00097     uint32_t GetAverage();
00098
00106     uint32_t GetCurrent();
00107
00108 private:
00109     uint32_t m_u32StartTicks;
00110     uint32_t m_u32CurrentIteration;
00111     uint32_t m_u32Cumulative;
00112     uint16_t m_u16Iterations;
00113     bool m_bActive;
00114 };
00115 } // namespace Mark3

```

20.121 /home/moslevin/projects/m3-repo/kernel/src/public/profiling_results.h File Reference

20.122 profiling_results.h


```

00044 public:
00045     static void Init();
00046
00054     static void SetInTimer();
00055
00060     static void ClearInTimer();
00061
00070     static void Update(Thread* pclTargetThread_);
00071
00079     static void SetTimerThread(Thread* pclTimerThread_) {
        m_pclTimerThread = pclTimerThread_; }
00080
00084     static void Cancel();
00085
00086 private:
00087     static Thread* m_pclActiveThread;
00088     static Thread* m_pclTimerThread;
00089     static uint16_t m_ul6TicksRemain;
00090     static bool m_bInTimer;
00091 };
00092 } // namespace Mark3
00093 #endif // #if KERNEL_ROUND_ROBIN

```

20.125 /home/moslevin/projects/m3-repo/kernel/src/public/readerwriter.h File Reference

Reader-Writer lock implementation.

```

#include "mark3cfg.h"
#include "blocking.h"
#include "mutex.h"

```

Classes

- class [Mark3::ReaderWriterLock](#)

The *ReaderWriterLock* class. This class implements an object that marshalls access to a resource based on the intended usage of the resource. A reader-writer lock permits multiple concurrent read access, or single-writer access to a resource. If the object holds a write lock, other writers, and all readers will block until the writer is finished. If the object holds reader locks, all writers will block until all readers are finished before the first writer can take ownership of the resource. This is based upon lower-level synchronization primitives, and is somewhat more heavyweight than primitive synchronization types.

Namespaces

- [Mark3](#)

20.125.1 Detailed Description

Reader-Writer lock implementation.

Definition in file [readerwriter.h](#).

20.126 readerwriter.h

```

00001  /*=====
00002
00003  _____
00004  |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |
00005  |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |
00006  |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |
00007  |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  =====*/
00021  #pragma once
00022
00023  #include "mark3cfg.h"
00024  #include "blocking.h"
00025  #include "mutex.h"
00026
00027  namespace Mark3
00028  {
00040  class ReaderWriterLock
00041  {
00042  public:
00043      void* operator new(size_t sz, void* pv) { return reinterpret_cast<
ReaderWriterLock*>(pv); }
00044
00050      void Init();
00051
00058      void AcquireReader();
00059
00068      bool AcquireReader(uint32_t u32TimeoutMs_);
00069
00074      void ReleaseReader();
00075
00082      void AcquireWriter();
00083
00092      bool AcquireWriter(uint32_t u32TimeoutMs_);
00093
00098      void ReleaseWriter();
00099
00100  private:
00107      bool AcquireReader_i(uint32_t u32TimeoutMs_);
00114      bool AcquireWriter_i(uint32_t u32TimeoutMs_);
00115
00116      Mutex    m_clGlobalMutex;
00117      Mutex    m_clReaderMutex;
00118      uint8_t  m_u8ReadCount;
00119  };
00120
00121  } // namespace Mark3

```

20.127 /home/moslevin/projects/m3-repo/kernel/src/public/scheduler.h File Reference

Thread scheduler function declarations This scheduler implements a very flexible type of scheduling, which has become the defacto industry standard when it comes to real-time operating systems. This scheduling mechanism is referred to as priority round- robin.

```

#include "kerneltypes.h"
#include "thread.h"
#include "ithreadport.h"
#include "priomap.h"

```

Classes

- class [Mark3::Scheduler](#)

The [Scheduler](#) Class. This class provides priority-based round-robin [Thread](#) scheduling for all active threads managed by the kernel.

Namespaces

- [Mark3](#)

Variables

- [Mark3::Thread](#) * [g_pclNext](#)
- [Mark3::Thread](#) * [g_pclCurrent](#)

20.127.1 Detailed Description

Thread scheduler function declarations This scheduler implements a very flexible type of scheduling, which has become the defacto industry standard when it comes to real-time operating systems. This scheduling mechanism is referred to as priority round- robin.

From the name, there are two concepts involved here:

1) Priority scheduling:

Threads are each assigned a priority, and the thread with the highest priority which is ready to run gets to execute.

2) Round-robin scheduling:

Where there are multiple ready threads at the highest-priority level, each thread in that group gets to share time, ensuring that progress is made.

The scheduler uses an array of ThreadList objects to provide the necessary housekeeping required to keep track of threads at the various priorities. As s result, the scheduler contains one ThreadList per priority, with an additional list to manage the storage of threads which are in the "stopped" state (either have been stopped, or have not been started yet).

Definition in file [scheduler.h](#).

20.127.2 Variable Documentation

20.127.2.1 g_pclCurrent

```
Mark3::Thread* g_pclCurrent
```

Definition at line 25 of file [scheduler.cpp](#).

20.127.2.2 g_pclNext

```
Mark3::Thread* g_pclNext
```

Definition at line 24 of file [scheduler.cpp](#).

20.128 scheduler.h

```

00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 ===== */
00045 #pragma once
00046
00047 #include "kerneltypes.h"
00048 #include "thread.h"
00049 #include "ithreadport.h"
00050 #include "priomap.h"
00051
00052 extern Mark3::Thread* g_pclNext;
00053 extern Mark3::Thread* g_pclCurrent;
00054
00055 namespace Mark3
00056 {
00057 //-----
00063 class Scheduler
00064 {
00065 public:
00070     static void Init();
00071
00078     static void Schedule();
00079
00086     static void Add(Thread* pclThread_);
00087
00095     static void Remove(Thread* pclThread_);
00096
00108     static bool SetScheduler(bool bEnable_);
00109
00116     static Thread* GetCurrentThread() { return
g_pclCurrent; }
00124     static volatile Thread* GetNextThread() { return
g_pclNext; }
00134     static ThreadList* GetThreadList(PORT_PRIO_TYPE uXPriority_) {
return &m_aclPriorities[uXPriority_]; }
00142     static ThreadList* GetStopList() { return &m_clStopList; }
00150     static bool IsEnabled() { return m_bEnabled; }
00156     static void QueueScheduler() { m_bQueuedSchedule = true; }
00157
00158 private:
00159     static constexpr auto m_uNumPriorities = size_t {
KERNEL_NUM_PRIORITIES };
00160
00162     static bool m_bEnabled;
00163
00165     static bool m_bQueuedSchedule;
00166
00168     static ThreadList m_clStopList;
00169
00171     static ThreadList m_aclPriorities[m_uNumPriorities];
00172
00174     static PriorityMap m_clPrioMap;
00175 };
00176 } // namespace Mark3

```

20.129 /home/moslevin/projects/m3-repo/kernel/src/public/schedulerguard.h File Reference

RAII Scheduler Locking.

```

#include "mark3cfg.h"
#include "scheduler.h"

```


20.133 /home/moslevin/projects/m3-repo/kernel/src/public/thread.h File Reference

Platform independent thread class declarations Threads are an atomic unit of execution, and each instance of the thread class represents an instance of a program running on the processor. The Thread is the fundamental user-facing object in the kernel - it is what makes multiprocessing possible from application code.

```
#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ll.h"
#include "threadlist.h"
#include "scheduler.h"
#include "ithreadport.h"
#include "quantum.h"
#include "autoalloc.h"
#include "priomap.h"
```

Classes

- class [Mark3::Thread](#)

The [Thread](#) Class. This object provides the fundamental thread control data structures and functions that define a single thread of execution in the [Mark3](#) operating system. It is the fundamental data type used to provide multitasking support in the kernel.

Namespaces

- [Mark3](#)

Typedefs

- using [Mark3::ThreadCreateCallout](#) = void(*) (Thread *pclThread_)
- using [Mark3::ThreadExitCallout](#) = void(*) (Thread *pclThread_)
- using [Mark3::ThreadContextCallout](#) = void(*) (Thread *pclThread_)

20.133.1 Detailed Description

Platform independent thread class declarations Threads are an atomic unit of execution, and each instance of the thread class represents an instance of a program running on the processor. The Thread is the fundamental user-facing object in the kernel - it is what makes multiprocessing possible from application code.

In [Mark3](#), threads each have their own context - consisting of a stack, and all of the registers required to multiplex a processor between multiple threads.

The Thread class inherits directly from the LinkListNode class to facilitate efficient thread management using Double, or Double-Circular linked lists.

Definition in file [thread.h](#).

20.134 thread.h

```

00001  /*=====
00002
00003  00004  00005  00006  00007  00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  ===== */
00034  #pragma once
00035
00036  #include "kerneltypes.h"
00037  #include "mark3cfg.h"
00038
00039  #include "ll.h"
00040  #include "threadlist.h"
00041  #include "scheduler.h"
00042  #include "ithreadport.h"
00043  #include "quantum.h"
00044  #include "autoalloc.h"
00045  #include "priomap.h"
00046
00047  namespace Mark3
00048  {
00049  class Thread;
00050
00051  //-----
00052  using ThreadCreateCallout = void (*) (Thread* pclThread_);
00053  using ThreadExitCallout = void (*) (Thread* pclThread_);
00054  using ThreadContextCallout = void (*) (Thread* pclThread_);
00055
00056  //-----
00064  class Thread : public TypedLinkedListNode<Thread>
00065  {
00066  public:
00067      void* operator new(size_t sz, void* pv) { return reinterpret_cast<Thread*>(pv); };
00068      ~Thread();
00069
00070      Thread() { m_eState = ThreadState::Invalid; }
00071
00077      bool IsInitialized() { return (m_eState !=
ThreadState::Invalid); }
00078
00093      void Init(K_WORD* pwStack_,
00094               uint16_t ul6StackSize_,
00095               PORT_PRIO_TYPE uXPriority_,
00096               ThreadEntryFunc pfEntryPoint_,
00097               void* pvArg_);
00098
00116      static Thread*
00117      Init(uint16_t ul6StackSize_, PORT_PRIO_TYPE uXPriority_,
ThreadEntryFunc pfEntryPoint_, void* pvArg_);
00118
00125      void Start();
00126
00132      void Stop();
00133
00134  #if KERNEL_NAMED_THREADS
00135
00143      void SetName(const char* szName_) { m_szName = szName_; }
00149      const char* GetName() { return m_szName; }
00150  #endif // #if KERNEL_NAMED_THREADS
00151
00159      ThreadList* GetOwner(void) { return m_pclOwner; }
00166      inline ThreadList* GetCurrent(void) { return
m_pclCurrent; }
00174      PORT_PRIO_TYPE GetPriority(void) { return
m_uXPriority; }
00181      PORT_PRIO_TYPE GetCurPriority(void) { return
m_uXCurPriority; }
00182
00183  #if KERNEL_ROUND_ROBIN
00184
00190      void SetQuantum(uint16_t ul6Quantum_) { m_ul6Quantum = ul6Quantum_; }
00197      uint16_t GetQuantum(void) { return m_ul6Quantum; }
00198  #endif // #if KERNEL_ROUND_ROBIN
00199
00206      void SetCurrent(ThreadList* pclNewList_) { m_pclCurrent = pclNewList_;
}
00213      void SetOwner(ThreadList* pclNewList_) { m_pclOwner = pclNewList_; }

```

```

00225     void SetPriority(PORT_PRIO_TYPE uXPriority_);
00226
00235     void InheritPriority(PORT_PRIO_TYPE uXPriority_);
00236
00247     void Exit();
00248
00256     static void Sleep(uint32_t u32TimeMs_);
00257
00264     static void Yield(void);
00265
00274     static void CoopYield(void);
00275
00282     void SetID(uint8_t u8ID_) { m_u8ThreadID = u8ID_; }
00292     uint8_t GetID() { return m_u8ThreadID; }
00293
00294 #if KERNEL_STACK_CHECK
00295     uint16_t GetStackSlack();
00307 #endif // #if KERNEL_STACK_CHECK
00308
00309 #if KERNEL_EVENT_FLAGS
00310     uint16_t GetEventFlagMask() { return m_u16FlagMask; }
00318
00323     void SetEventFlagMask(uint16_t u16Mask_) { m_u16FlagMask = u16Mask_; }
00324
00330     void SetEventFlagMode(EventFlagOperation eMode_) {
00331     m_eFlagMode = eMode_; }
00336     EventFlagOperation GetEventFlagMode() { return
00337     m_eFlagMode; }
00337 #endif // #if KERNEL_EVENT_FLAGS
00338
00342     Timer* GetTimer();
00343
00350     void SetExpired(bool bExpired_);
00351
00357     bool GetExpired();
00358
00359 #if KERNEL_EXTENDED_CONTEXT
00360     void* GetExtendedContext() { return m_pvExtendedContext; }
00369
00380     void SetExtendedContext(void* pvData_) {
00381     m_pvExtendedContext = pvData_; }
00381 #endif // #if KERNEL_EXTENDED_CONTEXT
00382
00389     ThreadState GetState() { return m_eState; }
00397     void SetState(ThreadState eState_) { m_eState = eState_; }
00398
00403     K_WORD* GetStack() { return m_pwStack; }
00404
00409     uint16_t GetStackSize() { return m_u16StackSize; }
00410
00411     friend class ThreadPort;
00412
00413 private:
00420     static void ContextSwitchSWI(void);
00421
00427     void SetPriorityBase(PORT_PRIO_TYPE uXPriority_);
00428
00430     K_WORD* m_pwStackTop;
00431
00433     K_WORD* m_pwStack;
00434
00436     uint8_t m_u8ThreadID;
00437
00439     PORT_PRIO_TYPE m_uXPriority;
00440
00442     PORT_PRIO_TYPE m_uXCurPriority;
00443
00445     ThreadState m_eState;
00446
00447 #if KERNEL_EXTENDED_CONTEXT
00448     void* m_pvExtendedContext;
00450 #endif // #if KERNEL_EXTENDED_CONTEXT
00451
00452 #if KERNEL_NAMED_THREADS
00453     const char* m_szName;
00455 #endif // #if KERNEL_NAMED_THREADS
00456
00458     uint16_t m_u16StackSize;
00459
00461     ThreadList* m_pclCurrent;
00462
00464     ThreadList* m_pclOwner;
00465

```

```

00467     ThreadEntryFunc m_pfEntryPoint;
00468
00470     void* m_pvArg;
00471
00472     #if KERNEL_ROUND_ROBIN
00473         uint16_t m_ul6Quantum;
00475     #endif // #if KERNEL_ROUND_ROBIN
00476
00477     #if KERNEL_EVENT_FLAGS
00478         uint16_t m_ul6FlagMask;
00480
00482         EventFlagOperation m_eFlagMode;
00483     #endif // #if KERNEL_EVENT_FLAGS
00484
00486         Timer m_clTimer;
00487
00489         bool m_bExpired;
00490 };
00491
00492 } // namespace Mark3

```

20.135 /home/moslevin/projects/m3-repo/kernel/src/public/threadlist.h File Reference

Thread linked-list declarations.

```

#include "kerneltypes.h"
#include "priomap.h"
#include "ll.h"

```

Classes

- class [Mark3::ThreadList](#)

The *ThreadList* Class. This class is used for building thread-management facilities, such as schedulers, and blocking objects.

Namespaces

- [Mark3](#)

20.135.1 Detailed Description

Thread linked-list declarations.

Definition in file [threadlist.h](#).

20.136 threadlist.h

```

00001  /*=====
00002
00003
00004  |-----|-----|-----|-----|-----|
00005  |   \   /   |   \   /   |   \   /   |   \   /   |   \   /   |
00006  |  / \  / \  |  / \  / \  |  / \  / \  |  / \  / \  |  / \  / \  |
00007  |_____|_____|_____|_____|_____|
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  ===== */
00022  #pragma once
00023
00024  #include "kerneltypes.h"
00025  #include "priomap.h"
00026  #include "ll.h"
00027
00028  //-----
00029  namespace Mark3
00030  {
00031  class Thread;
00032
00038  class ThreadList : public LinkListNode, public
    TypedCircularLinkList<Thread>
00039  {
00040  public:
00041      void* operator new(size_t sz, void* pv) { return reinterpret_cast<ThreadList*>(pv); };
00042      ThreadList();
00043
00044
00054      void SetPriority(PORT_PRIO_TYPE uXPriority_);
00055
00064      void SetMapPointer(PriorityMap* pclMap_);
00065
00072      void Add(Thread* node_);
00073
00084      void Add(Thread* node_, PriorityMap* pclMap_,
        PORT_PRIO_TYPE uXPriority_);
00085
00093      void AddPriority(Thread* node_);
00094
00101      void Remove(Thread* node_);
00102
00109      Thread* HighestWaiter();
00110
00111  private:
00113      PORT_PRIO_TYPE m_uXPriority;
00114
00116      PriorityMap* m_pclMap;
00117  };
00118  } // namespace Mark3

```

20.137 /home/moslevin/projects/m3-repo/kernel/src/public/threadlistlist.h File Reference

Class implementing a doubly-linked list of thread lists.

```
#include "mark3.h"
```

Classes

- class `Mark3::ThreadListList`

The `ThreadListList` class is used to track all threadlists active in the OS kernel. At any point in time, the list can be traversed to get a complete view of all running, blocked, or stopped threads in the system.

Namespaces

- Mark3

20.137.1 Detailed Description

Class implementing a doubly-linked list of thread lists.

Definition in file [threadlistlist.h](#).

20.138 threadlistlist.h

```

00001  /*=====
00002
00003  _____
00004  /         \
00005  /         \
00006  /         \
00007  /         \
00008  /         \
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  ===== */
00021  #pragma once
00022
00023  #include "mark3.h"
00024
00025  //-----
00026  namespace Mark3 {
00027  class ThreadList;
00028
00029  //-----
00036  class ThreadListList {
00037  public:
00043      static void Add(ThreadList* pclThreadList_)
00044      {
00045          m_clThreadListList.Add(pclThreadList_);
00046      }
00047
00053      static void Remove(ThreadList* pclThreadList_)
00054      {
00055          m_clThreadListList.Remove(pclThreadList_);
00056      }
00057
00062      static ThreadList* GetHead()
00063      {
00064          return m_clThreadListList.GetHead();
00065      }
00066
00067  private:
00068      static TypedDoubleLinkedList<ThreadList>
00069          m_clThreadListList;
00070  } // namespace Mark3

```

20.139 /home/moslevin/projects/m3-repo/kernel/src/public/timer.h File Reference

Timer object declarations.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "ll.h"

```

Classes

- class [Mark3::Timer](#)

The *Timer* Class. This class provides kernel-managed timers, used to provide high-precision delays. Functionality is useful to both user-code, and is used extensively within the kernel and its blocking objects to implement round-robin scheduling, thread sleep, and timeouts. Provides one-shot and periodic timers for use by application code. This object relies on a target-defined hardware timer implementation, which is multiplexed by the kernel's timer scheduler.


```

00038
00039 //-----
00050 using TimerCallback = void (*)(Thread* pclOwner_, void* pvData_);
00051
00052 //-----
00053 class TimerList;
00054 class TimerScheduler;
00055 class Quantum;
00056
00057 //-----
00068 class Timer : public TypedLinkedListNode<Timer>
00069 {
00070 public:
00071     void* operator new(size_t sz, void* pv) { return reinterpret_cast<Timer*>(pv); }
00072     ~Timer() {}
00073
00074     Timer();
00075
00076     void Init();
00077
00078     void Start(bool bRepeat_, uint32_t u32IntervalMs_, TimerCallback pfCallback_, void*
pvData_);
00079
00080     void Start();
00081
00082     void Stop();
00083
00084 private:
00085     friend class TimerList;
00086
00087     void SetInitialized() { m_u8Initialized =
m_uTimerInitCookie; }
00088
00089     bool IsInitialized(void) { return (m_u8Initialized ==
m_uTimerInitCookie); }
00090
00091     static inline uint32_t SecondsToTicks(uint32_t x) { return (x) * 1000; }
00092     static inline uint32_t MSecondsToTicks(uint32_t x) { return (x); }
00093     static inline uint32_t USecondsToTicks(uint32_t x) { return ((x + 999) / 1000); }
00094
00095     static constexpr auto m_uTimerInvalidCookie = uint8_t { 0xC3 };
00096     static constexpr auto m_uTimerInitCookie = uint8_t { 0xC3 };
00097
00098     uint8_t m_u8Initialized;
00099
00100     uint8_t m_u8Flags;
00101
00102     TimerCallback m_pfCallback;
00103
00104     uint32_t m_u32Interval;
00105
00106     uint32_t m_u32TimeLeft;
00107
00108     Thread* m_pclOwner;
00109
00110     void* m_pvData;
00111 };
00112 } // namespace Mark3

```

20.141 /home/moslevin/projects/m3-repo/kernel/src/public/timerlist.h File Reference

Timer list declarations.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "mutex.h"

```

Classes

- class [Mark3::TimerList](#)

the *TimerList* class. This class implements a doubly-linked-list of timer objects.

20.145.1 Detailed Description

Thread Quantum Implementation for Round-Robin Scheduling.

Definition in file [quantum.cpp](#).

20.146 quantum.cpp

```

00001 /*-----
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00022 #include "mark3.h"
00023
00024 #if KERNEL_ROUND_ROBIN
00025 namespace Mark3
00026 {
00027 //-----
00028 uint16_t Quantum::m_ul6TicksRemain;
00029 Thread* Quantum::m_pclActiveThread;
00030 Thread* Quantum::m_pclTimerThread;
00031 bool Quantum::m_bInTimer;
00032
00033 //-----
00034 void Quantum::SetInTimer()
00035 {
00036     const auto cs = CriticalGuard{};
00037     m_bInTimer = true;
00038
00039     // Timer is active
00040     if (m_ul6TicksRemain) {
00041         m_ul6TicksRemain--;
00042     }
00043 }
00044
00045 //-----
00046 void Quantum::ClearInTimer()
00047 {
00048     const auto cs = CriticalGuard{};
00049     m_bInTimer = false;
00050
00051     // Timer expired - Pivot the thread list.
00052     if (m_pclActiveThread && (!m_ul6TicksRemain)) {
00053         auto* pclThreadList = m_pclActiveThread->GetCurrent();
00054         if (pclThreadList->GetHead() != pclThreadList->GetTail()) {
00055             pclThreadList->PivotForward();
00056         }
00057         m_pclActiveThread = nullptr;
00058     }
00059 }
00060
00061 //-----
00062 void Quantum::Update(Thread* pclTargetThread_)
00063 {
00064     // Don't cancel the current RR interval if we're being interrupted by
00065     // the timer thread, or are in the middle of running the timer thread.
00066     // OR if the thread list only has one thread
00067     auto* pclThreadList = pclTargetThread_->GetCurrent();
00068     if ((pclThreadList->GetHead() == pclThreadList->GetTail()) || (pclTargetThread_ ==
m_pclTimerThread)
00069         || (pclTargetThread_ == m_pclActiveThread) ||
m_bInTimer) {
00070         return;
00071     }
00072
00073     // Update with a new thread and timeout.
00074     m_pclActiveThread = pclTargetThread_;
00075     m_ul6TicksRemain = pclTargetThread_->GetQuantum();
00076 }
00077

```



```

00043 //-----
00044 void ReaderWriterLock::ReleaseReader()
00045 {
00046     m_clReaderMutex.Claim();
00047     m_u8ReadCount--;
00048     if (0 == m_u8ReadCount) {
00049         m_clGlobalMutex.Release();
00050     }
00051     m_clReaderMutex.Release();
00052 }
00053
00054 //-----
00055 void ReaderWriterLock::AcquireWriter()
00056 {
00057     AcquireWriter_i(0);
00058 }
00059
00060 //-----
00061 bool ReaderWriterLock::AcquireWriter(uint32_t u32TimeoutMs_)
00062 {
00063     return AcquireWriter_i(u32TimeoutMs_);
00064 }
00065
00066 //-----
00067 void ReaderWriterLock::ReleaseWriter()
00068 {
00069     m_clGlobalMutex.Release();
00070 }
00071
00072 //-----
00073 bool ReaderWriterLock::AcquireReader_i(uint32_t u32TimeoutMs_)
00074 {
00075     auto rc = true;
00076     if (!m_clReaderMutex.Claim(u32TimeoutMs_)) {
00077         return false;
00078     }
00079
00080     m_u8ReadCount++;
00081     if (1 == m_u8ReadCount) {
00082         rc = m_clGlobalMutex.Claim(u32TimeoutMs_);
00083     }
00084
00085     m_clReaderMutex.Release();
00086     return rc;
00087 }
00088
00089 //-----
00090 bool ReaderWriterLock::AcquireWriter_i(uint32_t u32TimeoutMs_)
00091 {
00092     return m_clGlobalMutex.Claim(u32TimeoutMs_);
00093 }
00094 } // namespace Mark3

```

20.149 /home/moslevin/projects/m3-repo/kernel/src/scheduler.cpp File Reference

Strict-Priority + Round-Robin thread scheduler implementation.

```
#include "mark3.h"
```

Namespaces

- [Mark3](#)

Variables

- [Mark3::Thread * g_pclNext](#)
- [Mark3::Thread * g_pclCurrent](#)


```

00043
00044 //-----
00045 void Scheduler::Schedule()
00046 {
00047     auto uXPrio = m_clPrioMap.HighestPriority();
00048     if (0 == uXPrio) {
00049         Kernel::Panic(PANIC_NO_READY_THREADS);
00050     }
00051     // Priorities are one-indexed
00052     uXPrio--;
00053
00054     // Get the thread node at this priority.
00055     g_pclNext = m_aclPriorities[uXPrio].GetHead();
00056 }
00057
00058 //-----
00059 void Scheduler::Add(Thread* pclThread_)
00060 {
00061     KERNEL_ASSERT(pclThread_ != nullptr);
00062
00063     m_aclPriorities[pclThread_>GetPriority()].Add(pclThread_);
00064 }
00065
00066 //-----
00067 void Scheduler::Remove(Thread* pclThread_)
00068 {
00069     KERNEL_ASSERT(pclThread_ != nullptr);
00070
00071     m_aclPriorities[pclThread_>GetPriority()].Remove(pclThread_);
00072 }
00073
00074 //-----
00075 bool Scheduler::SetScheduler(bool bEnable_)
00076 {
00077     const auto cs = CriticalGuard{};
00078     auto bRet     = m_bEnabled;
00079     m_bEnabled = bEnable_;
00080     // If there was a queued scheduler event, dequeue and trigger an
00081     // immediate Yield
00082     if (m_bEnabled && m_bQueuedSchedule) {
00083         m_bQueuedSchedule = false;
00084         Thread::Yield();
00085     }
00086     return bRet;
00087 }
00088 } // namespace Mark3

```

20.151 /home/moslevin/projects/m3-repo/kernel/src/thread.cpp File Reference

Platform-Independent thread class Definition.

```

#include "kerneltypes.h"
#include "mark3cfg.h"
#include "mark3.h"

```

Namespaces

- [Mark3](#)

20.151.1 Detailed Description

Platform-Independent thread class Definition.

Definition in file [thread.cpp](#).


```

00091
00092 #if KERNEL_THREAD_CREATE_CALLOUT
00093     ThreadCreateCallout pfCallout =
00094         Kernel::GetThreadCreateCallout();
00095     if (nullptr != pfCallout) {
00096         pfCallout(this);
00097     }
00098 #endif
00099 }
00100 //-----
00101 Thread* Thread::Init(uint16_t ul6StackSize_, PORT_PRIO_TYPE uXPriority_,
00102     ThreadEntryFunc pfEntryPoint_, void* pvArg_)
00103 {
00104     auto* pclNew = AutoAlloc::NewObject<Thread, AutoAllocType::Thread>();
00105     auto* pwStack = static_cast<K_WORD*>(AutoAlloc::NewRawData(ul6StackSize_));
00106     pclNew->Init(pwStack, ul6StackSize_, uXPriority_, pfEntryPoint_, pvArg_);
00107     return pclNew;
00108 }
00109 //-----
00110 void Thread::Start(void)
00111 {
00112     KERNEL_ASSERT(IsInitialized());
00113
00114     // Remove the thread from the scheduler's "stopped" list, and add it
00115     // to the scheduler's ready list at the proper priority.
00116
00117     const auto cs = CriticalGuard{};
00118     Scheduler::GetStopList()->Remove(this);
00119     Scheduler::Add(this);
00120     m_pclOwner = Scheduler::GetThreadList(
00121         m_uXPriority);
00122     m_pclCurrent = m_pclOwner;
00123     m_eState = ThreadState::Ready;
00124 #if KERNEL_ROUND_ROBIN
00125     if (Kernel::IsStarted()) {
00126         if (GetCurPriority() >= Scheduler::GetCurrentThread()->
00127             GetCurPriority()) {
00128             // Deal with the thread Quantum
00129             Quantum::Update(this);
00130         }
00131     }
00132 #endif
00133     if (Kernel::IsStarted()) {
00134         if (GetCurPriority() >= Scheduler::GetCurrentThread()->
00135             GetCurPriority()) {
00136             Thread::Yield();
00137         }
00138     }
00139 //-----
00140 void Thread::Stop()
00141 {
00142     KERNEL_ASSERT(IsInitialized());
00143
00144     auto bReschedule = false;
00145     if (ThreadState::Stop == m_eState) {
00146         return;
00147     }
00148
00149     { // Begin critical section
00150         const auto cs = CriticalGuard{};
00151
00152         // If a thread is attempting to stop itself, ensure we call the scheduler
00153         if (this == Scheduler::GetCurrentThread()) {
00154             bReschedule = true;
00155         }
00156         #if KERNEL_ROUND_ROBIN
00157             // Cancel RR scheduling
00158             Quantum::Cancel();
00159         #endif
00160     }
00161
00162     // Add this thread to the stop-list (removing it from active scheduling)
00163     // Remove the thread from scheduling
00164     if (ThreadState::Ready == m_eState) {
00165         Scheduler::Remove(this);
00166     } else if (ThreadState::Blocked == m_eState) {
00167         m_pclCurrent->Remove(this);
00168     }
00169
00170     m_pclOwner = Scheduler::GetStopList();
00171     m_pclCurrent = m_pclOwner;
00172     m_pclOwner->Add(this);

```

```

00173         m_eState = ThreadState::Stop;
00174
00175         // Just to be safe - attempt to remove the thread's timer
00176         // from the timer-scheduler (does no harm if it isn't
00177         // in the timer-list)
00178         TimerScheduler::Remove(&m_clTimer);
00179     } // End Critical Section
00180
00181     if (bReschedule) {
00182         Thread::Yield();
00183     }
00184 }
00185
00186 //-----
00187 void Thread::Exit()
00188 {
00189     KERNEL_ASSERT(IsInitialized());
00190
00191     auto bReschedule = false;
00192
00193     if (ThreadState::Exit == m_eState) {
00194         return;
00195     }
00196
00197     { // Begin critical section
00198         const auto cs = CriticalGuard{};
00199
00200         // If this thread is the actively-running thread, make sure we run the
00201         // scheduler again.
00202         if (this == Scheduler::GetCurrentThread()) {
00203             bReschedule = true;
00204 #if KERNEL_ROUND_ROBIN
00205             // Cancel RR scheduling
00206             Quantum::Cancel();
00207 #endif
00208         }
00209
00210         // Remove the thread from scheduling
00211         if (ThreadState::Ready == m_eState) {
00212             Scheduler::Remove(this);
00213         } else if ((ThreadState::Blocked == m_eState) || (
ThreadState::Stop == m_eState)) {
00214             m_pclCurrent->Remove(this);
00215         }
00216
00217         m_pclCurrent = nullptr;
00218         m_pclOwner = nullptr;
00219         m_eState = ThreadState::Exit;
00220
00221         // We've removed the thread from scheduling, but interrupts might
00222         // trigger checks against this thread's currently priority before
00223         // we get around to scheduling new threads. As a result, set the
00224         // priority to idle to ensure that we always wind up scheduling
00225         // new threads.
00226         m_uXCurPriority = 0;
00227         m_uXPriority = 0;
00228
00229         // Just to be safe - attempt to remove the thread's timer
00230         // from the timer-scheduler (does no harm if it isn't
00231         // in the timer-list)
00232         TimerScheduler::Remove(&m_clTimer);
00233     } // End Critical Section
00234
00235 #if KERNEL_THREAD_EXIT_CALLOUT
00236     ThreadExitCallout pfCallout = Kernel::GetThreadExitCallout
00237     ();
00238     if (nullptr != pfCallout) {
00239         pfCallout(this);
00240     }
00241 #endif
00242
00243     if (bReschedule) {
00244         // Choose a new "next" thread if we must
00245         Thread::Yield();
00246     }
00247 }
00248 //-----
00249 void Thread::Sleep(uint32_t u32TimeMs_)
00250 {
00251     auto clSemaphore = Semaphore {};
00252     auto* pclTimer = g_pclCurrent->GetTimer();
00253     auto lTimerCallback = [](Thread* /*pclOwner*/, void* pvData_) {
00254         auto* pclSemaphore = static_cast<Semaphore*>(pvData_);
00255         pclSemaphore->Post();
00256     };
00257

```

```

00258     // Create a semaphore that this thread will block on
00259     clSemaphore.Init(0, 1);
00260
00261     // Create a one-shot timer that will call a callback that posts the
00262     // semaphore, waking our thread.
00263     pclTimer->Init();
00264     pclTimer->Start(false, u32TimeMs_, lTimerCallback, &clSemaphore);
00265
00266     clSemaphore.Pend();
00267 }
00268
00269 #if KERNEL_STACK_CHECK
00270 //-----
00271 uint16_t Thread::GetStackSlack()
00272 {
00273     KERNEL_ASSERT(IsInitialized());
00274
00275     auto wBottom = uint16_t { 0 };
00276     auto wTop = static_cast<uint16_t>((m_ul6StackSize - 1) / sizeof(
K_ADDR));
00277     auto wMid = static_cast<uint16_t>((wTop + wBottom) + 1) / 2;
00278
00279     { // Begin critical section
00280         const auto cs = CriticalGuard{};
00281
00282         // Logarithmic bisection - find the point where the contents of the
00283         // stack go from 0xFF's to non 0xFF. Not Definitive, but accurate enough
00284         while ((wTop - wBottom) > 1) {
00285             #if PORT_STACK_GROWS_DOWN
00286                 if (m_pwStack[wMid] != static_cast<K_WORD>(-1))
00287             #else
00288                 if (m_pwStack[wMid] == static_cast<K_WORD>(-1))
00289             #endif
00290             {
00291                 wTop = wMid;
00292             } else {
00293                 wBottom = wMid;
00294             }
00295             wMid = (wTop + wBottom + 1) / 2;
00296         }
00297     } // End Critical Section
00298
00299     return wMid * sizeof(K_ADDR);
00300 }
00301 #endif
00302
00303 //-----
00304 void Thread::Yield()
00305 {
00306     const auto cs = CriticalGuard{};
00307     // Run the scheduler
00308     if (Scheduler::IsEnabled()) {
00309         Scheduler::Schedule();
00310
00311         // Only switch contexts if the new task is different than the old task
00312         if (g_pclCurrent != g_pclNext) {
00313             #if KERNEL_ROUND_ROBIN
00314                 Quantum::Update(g_pclNext);
00315             #endif
00316             Thread::ContextSwitchSWI();
00317         }
00318     } else {
00319         Scheduler::QueueScheduler();
00320     }
00321 }
00322
00323 //-----
00324 void Thread::CoopYield(void)
00325 {
00326     g_pclCurrent->GetCurrent()->PivotForward();
00327     Yield();
00328 }
00329
00330 //-----
00331 void Thread::SetPriorityBase(PORT_PRIO_TYPE /*uXPriority_*/)
00332 {
00333     KERNEL_ASSERT(IsInitialized());
00334
00335     GetCurrent()->Remove(this);
00336     SetCurrent(Scheduler::GetThreadList(
m_uXPriority));
00337     GetCurrent()->Add(this);
00338 }
00339
00340 //-----
00341 void Thread::SetPriority(PORT_PRIO_TYPE uXPriority_)
00342 {
00343

```

```

00344     KERNEL_ASSERT(IsInitialized());
00345     auto bSchedule = false;
00346
00347     { // Begin critical section
00348         const auto cs = CriticalGuard{};
00349
00350         // If this is the currently running thread, it's a good idea to reschedule
00351         // Or, if the new priority is a higher priority than the current thread's.
00352         if ((this == g_pclCurrent) || (uXPriority_ > g_pclCurrent->
GetPriority())) {
00353             bSchedule = true;
00354             #if KERNEL_ROUND_ROBIN
00355                 Quantum::Cancel();
00356             #endif
00357             Scheduler::Remove(this);
00358
00359             m_uXCurPriority = uXPriority_;
00360             m_uXPriority    = uXPriority_;
00361
00362             Scheduler::Add(this);
00363         } // End critical section
00364
00365         if (bSchedule) {
00366             if (Scheduler::IsEnabled()) {
00367                 { // Begin critical section
00368                     const auto cs = CriticalGuard{};
00369                     Scheduler::Schedule();
00370                 }
00371                 #if KERNEL_ROUND_ROBIN
00372                     Quantum::Update(g_pclNext);
00373                 #endif
00374             } // End critical section
00375             Thread::ContextSwitchSWI();
00376         } else {
00377             Scheduler::QueueScheduler();
00378         }
00379     }
00380 }
00381
00382 //-----
00383 void Thread::InheritPriority(PORT_PRIO_TYPE uXPriority_)
00384 {
00385     KERNEL_ASSERT(IsInitialized());
00386
00387     SetOwner(Scheduler::GetThreadList(uXPriority_));
00388     m_uXCurPriority = uXPriority_;
00389 }
00390
00391 //-----
00392 void Thread::ContextSwitchSWI()
00393 {
00394     // Call the context switch interrupt if the scheduler is enabled.
00395     if (Scheduler::IsEnabled()) {
00396         #if KERNEL_STACK_CHECK
00397             if (g_pclCurrent && (g_pclCurrent->GetStackSlack() <=
Kernel::GetStackGuardThreshold())) {
00398                 Kernel::Panic(PANIC_STACK_SLACK_VIOLATED);
00399             }
00400         #endif
00401         #if KERNEL_CONTEXT_SWITCH_CALLOUT
00402             auto pfCallout = Kernel::GetThreadContextSwitchCallout();
00403             if (nullptr != pfCallout) {
00404                 pfCallout(g_pclCurrent);
00405             }
00406         #endif
00407         KernelSWI::Trigger();
00408     }
00409 }
00410
00411 //-----
00412 Timer* Thread::GetTimer()
00413 {
00414     KERNEL_ASSERT(IsInitialized());
00415     return &m_clTimer;
00416 }
00417 //-----
00418 void Thread::SetExpired(bool bExpired_)
00419 {
00420     KERNEL_ASSERT(IsInitialized());
00421     m_bExpired = bExpired_;
00422 }
00423
00424 //-----
00425 bool Thread::GetExpired()
00426 {
00427     KERNEL_ASSERT(IsInitialized());
00428     return m_bExpired;

```

```
00429 }
00430 } // namespace Mark3
```

20.153 /home/moslevin/projects/m3-repo/kernel/src/threadlist.cpp File Reference

Thread linked-list definitions.

```
#include "mark3.h"
#include "threadlistlist.h"
```

Namespaces

- [Mark3](#)

20.153.1 Detailed Description

Thread linked-list definitions.

Definition in file [threadlist.cpp](#).

20.154 threadlist.cpp

```
00001 /*=====
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*/
00022 #include "mark3.h"
00023 #include "threadlistlist.h"
00024 namespace Mark3
00025 {
00026 //-----
00027 ThreadList::ThreadList()
00028     : m_uXPriority(0)
00029     , m_pclMap(nullptr)
00030 {
00031 }
00032
00033 //-----
00034 void ThreadList::SetPriority(PORT_PRIO_TYPE uXPriority_)
00035 {
00036     m_uXPriority = uXPriority_;
00037 }
00038
00039 //-----
00040 void ThreadList::SetMapPointer(PriorityMap* pclMap_)
00041 {
00042     KERNEL_ASSERT(pclMap_ != nullptr);
00043     m_pclMap = pclMap_;
00044 }
00045
00046 //-----
00047 void ThreadList::Add(Thread* pclThread_)
00048 {
00049     KERNEL_ASSERT(pclThread_ != nullptr);
```

```

00050
00051 // If list was empty, add the object for global threadlist tracking
00052 if (!GetHead()) {
00053     ThreadListList::Add(this);
00054 }
00055 TypedCircularLinkedList<Thread>::Add(pclThread_);
00056 PivotForward();
00057
00058 // We've specified a bitmap for this threadlist
00059 if (nullptr != m_pclMap) {
00060     // Set the flag for this priority level
00061     m_pclMap->Set(m_uXPriority);
00062 }
00063 }
00064
00065 //-----
00066 void ThreadList::AddPriority(Thread* pclThread_)
00067 {
00068     KERNEL_ASSERT(node_ != nullptr);
00069     auto* pclCurr = GetHead();
00070     if (nullptr == pclCurr) {
00071         Add(pclThread_);
00072         return;
00073     }
00074     auto uXHeadPri = pclCurr->GetCurPriority();
00075     auto* pclTail = GetTail();
00076
00077     // Set the threadlist's priority level, flag pointer, and then add the
00078     // thread to the threadlist
00079     auto uXPriority = pclThread_->GetCurPriority();
00080     do {
00081         if (uXPriority > pclCurr->GetCurPriority()) {
00082             break;
00083         }
00084         pclCurr = pclCurr->GetNext();
00085     } while (pclCurr != pclTail);
00086
00087     // Insert pclNode before pclCurr in the linked list.
00088     InsertNodeBefore(pclThread_, pclCurr);
00089
00090     // If the priority is greater than current head, reset
00091     // the head pointer.
00092     if (uXPriority > uXHeadPri) {
00093         SetHead(pclThread_);
00094         SetTail(GetHead()->GetPrev());
00095     } else if (pclThread_->GetNext() == GetHead()) {
00096         SetTail(pclThread_);
00097     }
00098 }
00099
00100 //-----
00101 void ThreadList::Add(Thread* node_, PriorityMap* pclMap_,
00102     PORT_PRIO_TYPE uXPriority_)
00103 {
00104     // Set the threadlist's priority level, flag pointer, and then add the
00105     // thread to the threadlist
00106     SetPriority(uXPriority_);
00107     SetMapPointer(pclMap_);
00108     Add(node_);
00109 }
00110 //-----
00111 void ThreadList::Remove(Thread* node_)
00112 {
00113     // Remove the thread from the list
00114     TypedCircularLinkedList<Thread>::Remove(node_);
00115
00116     // If the list is empty...
00117     if (nullptr == GetHead()) {
00118         // No more threads - remove this object from global threadlist tracking
00119         ThreadListList::Remove(this);
00120         if (nullptr != m_pclMap) {
00121             // Clear the bit in the bitmap at this priority level
00122             m_pclMap->Clear(m_uXPriority);
00123         }
00124     }
00125 }
00126
00127 //-----
00128 Thread* ThreadList::HighestWaiter()
00129 {
00130     return GetHead();
00131 }
00132 } // namespace Mark3

```


20.155 /home/moslevin/projects/m3-repo/kernel/src/threadlistlist.cpp File Reference

Class implementing a doubly-linked list of thread lists.

```
#include "mark3.h"
#include "threadlistlist.h"
```

Namespaces

- [Mark3](#)

20.155.1 Detailed Description

Class implementing a doubly-linked list of thread lists.

Definition in file [threadlistlist.cpp](#).

20.156 threadlistlist.cpp

```
00001  /*=====
00002
00003  |-----|-----|-----|-----|-----|-----|-----|-----|
00004  | \ / | \ / | \ / | \ / | \ / | \ / | \ / | \ / | \ / | \ / |
00005  | / \ | / \ | / \ | / \ | / \ | / \ | / \ | / \ | / \ |
00006  |-----|-----|-----|-----|-----|-----|-----|-----|
00007
00008
00009  --[Mark3 Realtime Platform]-----
00010
00011  Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012  See license.txt for more information
00013  ===== */
00021  #include "mark3.h"
00022  #include "threadlistlist.h"
00023
00024  namespace Mark3 {
00025  TypedDoubleLinkedList<ThreadList> ThreadListList::m_clThreadListList;
00026  } // namespace Mark3
```

20.157 /home/moslevin/projects/m3-repo/kernel/src/timer.cpp File Reference

Timer implementations.

```
#include "mark3.h"
```

Namespaces

- [Mark3](#)


```
00079     }
00080
00081     m_pclOwner = Scheduler::GetCurrentThread();
00082     TimerScheduler::Add(this);
00083 }
00084
00085 //-----
00086 void Timer::Stop()
00087 {
00088     KERNEL_ASSERT(IsInitialized());
00089     if ((m_u8Flags & uTimerFlagActive) == 0) {
00090         return;
00091     }
00092     TimerScheduler::Remove(this);
00093 }
00094 } // namespace Mark3
```

20.159 /home/moslevin/projects/m3-repo/kernel/src/timerlist.cpp File Reference

Implements timer list processing algorithms, responsible for all timer tick and expiry logic.

```
#include "mark3.h"
```

Namespaces

- Mark3

20.159.1 Detailed Description

Implements timer list processing algorithms, responsible for all timer tick and expiry logic.

Definition in file [timerlist.cpp](#).

20.160 timerlist.cpp

```

00001 /*-----
00002
00003
00004
00005
00006
00007
00008
00009 --[Mark3 Realtime Platform]-----
00010
00011 Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
00012 See license.txt for more information
00013 =====*
00023 #include "mark3.h"
00024 namespace Mark3
00025 {
00026 //-----
00027 void TimerList::Init(void)
00028 {
00029     m_bTimerActive = false;
00030     m_u32NextWakeUp = 0;
00031     m_clMutex.Init();
00032 }
00033
00034 //-----
00035 void TimerList::Add(Timer* pclListNode_)
00036 {
00037     KERNEL_ASSERT(nullptr != pclListNode_);

```

```

00038     auto lock = LockGuard { &m_clMutex };
00039
00040     pclListNode->ClearNode();
00041     TypedDoubleLinkedList<Timer>::Add(pclListNode_);
00042
00043     // Set the initial timer value
00044     pclListNode->m_u32TimeLeft = pclListNode->m_u32Interval;
00045
00046     // Set the timer as active.
00047     pclListNode->m_u8Flags |= uTimerFlagActive;
00048 }
00049
00050 //-----
00051 void TimerList::Remove(Timer* pclLinkListNode_)
00052 {
00053     KERNEL_ASSERT(nullptr != pclLinkListNode_);
00054     auto lock = LockGuard { &m_clMutex };
00055
00056     TypedDoubleLinkedList<Timer>::Remove(pclLinkListNode_);
00057     pclLinkListNode->m_u8Flags &= ~uTimerFlagActive;
00058 }
00059
00060 //-----
00061 void TimerList::Process(void)
00062 {
00063     auto lock = LockGuard { &m_clMutex };
00064
00065     auto* pclCurr = GetHead();
00066     // Subtract the elapsed time interval from each active timer.
00067     while (nullptr != pclCurr) {
00068         auto* pclNext = pclCurr->GetNext();
00069
00070         // Active timers only...
00071         if ((pclCurr->m_u8Flags & uTimerFlagActive) != 0) {
00072             pclCurr->m_u32TimeLeft--;
00073             if (0 == pclCurr->m_u32TimeLeft) {
00074                 // Expired -- run the callback. these callbacks must be very fast...
00075                 if (nullptr != pclCurr->m_pfCallback) {
00076                     pclCurr->m_pfCallback(pclCurr->m_pclOwner, pclCurr->m_pvData);
00077                 }
00078                 if ((pclCurr->m_u8Flags & uTimerFlagOneShot) != 0) {
00079                     // If this was a one-shot timer, deactivate the timer + remove
00080                     pclCurr->m_u8Flags |= uTimerFlagExpired;
00081                     pclCurr->m_u8Flags &= ~uTimerFlagActive;
00082                     Remove(pclCurr);
00083                 } else {
00084                     // Reset the interval timer.
00085                     pclCurr->m_u32TimeLeft = pclCurr->m_u32Interval;
00086                 }
00087             }
00088         }
00089         pclCurr = pclNext;
00090     }
00091 }
00092
00093 } // namespace Mark3

```

Example Documentation

21.1 lab10_notifications/main.cpp

This examples demonstrates how to use notification objects as a thread synchronization mechanism.

```

/*=====
\      /      \      /      \      /      \      /
|      /      |      /      |      /      |      /
|     /       |     /       |     /       |     /
|    /        |    /        |    /        |    /
|   /         |   /         |   /         |   /
|  /          |  /          |  /          |  /
| /           | /           | /           | /
|/            |/            |/            |/
-----
--[Mark3 Realtime Platform]-----

Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
See license.txt for more information
=====*/
#include "mark3.h"

/*=====

Lab Example 10:  Thread Notifications

Lessons covered in this example include:
- Create a notification object, and use it to synchronize execution of Threads.

Takeaway:
- Notification objects are a lightweight mechanism to signal thread execution
  in situations where even a semaphore would be a heavier-weight option.

=====*/
extern "C" {
void __cxa_pure_virtual(void) {}
void DebugPrint(const char* szString_);
}

namespace
{
using namespace Mark3;
//-----
Thread clApp1Thread;
K_WORD awApp1Stack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void App1Main(void* unused_);

//-----
Thread clApp2Thread;
K_WORD awApp2Stack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void App2Main(void* unused_);

//-----
// idle thread -- do nothing
Thread clIdleThread;
K_WORD awIdleStack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void IdleMain(void* /*unused_*/)
{
    while (1) {}
}

```


Lab Example 11: Mailboxes

Lessons covered in this example include:

- Initialize a mailbox for use as an IPC mechanism.
- Create and use mailboxes to pass data between threads.

Takeaway:

- Mailboxes are a powerful IPC mechanism used to pass messages of a fixed-size between threads.

```

=====*/
extern "C" {
void __cxa_pure_virtual(void) {}
void DebugPrint(const char* szString_);
}

namespace
{
using namespace Mark3;
//-----
Thread clApp1Thread;
K_WORD awApp1Stack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void App1Main(void* unused_);

//-----
Thread clApp2Thread;
K_WORD awApp2Stack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void App2Main(void* unused_);

//-----
// idle thread -- do nothing
Thread clIdleThread;
K_WORD awIdleStack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void IdleMain(void* /*unused_*/)
{
    while (1) {}
}

//-----
Mailbox clMailbox;
uint8_t au8MBData[100];

typedef struct {
    uint8_t au8Buffer[10];
} MBType_t;

//-----
void App1Main(void* unused_)
{
    while (1) {
        MBType_t stMsg;

        // Wait until there is an envelope available in the shared mailbox, and
        // then log a trace message.
        clMailbox.Receive(&stMsg);
        // KernelAware::Trace(0, __LINE__, stMsg.au8Buffer[0], stMsg.au8Buffer[9]);
    }
}

//-----
void App2Main(void* unused_)
{
    while (1) {
        MBType_t stMsg;

        // Place a bunch of envelopes in the mailbox, and then wait for a
        // while. Note that this thread has a higher priority than the other
        // thread, so it will keep pushing envelopes to the other thread until
        // it gets to the sleep, at which point the other thread will be allowed
        // to execute.

        Kernel::DebugPrint("Messages Begin\n");

        for (uint8_t i = 0; i < 10; i++) {
            for (uint8_t j = 0; j < 10; j++) { stMsg.au8Buffer[j] = (i * 10) + j; }
            clMailbox.Send(&stMsg);
        }

        Kernel::DebugPrint("Messages End\n");
        Thread::Sleep(2000);
    }
} // anonymous namespace

using namespace Mark3;
//-----

```



```

{
using namespace Mark3;
//-----
// This block declares the thread data for the main application thread. It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
Thread clAppThread;
K_WORD awAppStack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void AppMain(void* unused_);

//-----
// This block declares the thread data for the idle thread. It defines a
// thread object, stack (in word-array form), and the entry-point function
// used by the idle thread.
Thread clIdleThread;
K_WORD awIdleStack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void IdleMain(void* unused_);

//-----
void AppMain(void* unused_)
{
    // This function is run from within the application thread. Here, we
    // simply print a friendly greeting and allow the thread to sleep for a
    // while before repeating the message. Note that while the thread is
    // sleeping, CPU execution will transition to the Idle thread.

    while (1) {
        Kernel::DebugPrint("Hello World!\n");
        Thread::Sleep(1000);
    }
}

//-----
void IdleMain(void* unused_)
{
    while (1) {
        // Low priority task + power management routines go here.
        // The actions taken in this context must *not* cause the thread
        // to block, as the kernel requires that at least one thread is
        // schedulable at all times when not using an idle thread.

        // Note that if you have no special power-management code or idle
        // tasks, an empty while(1){} loop is sufficient to guarantee that
        // condition.
    }
}
} // anonymous namespace

using namespace Mark3;
//-----
int main(void)
{
    // Before any Mark3 RTOS APIs can be called, the user must call Kernel::Init().
    // Note that if you have any hardware-specific init code, it can be called
    // before Kernel::Init, so long as it does not enable interrupts, or
    // rely on hardware peripherals (timer, software interrupt, etc.) used by the
    // kernel.
    Kernel::Init();
    Kernel::SetDebugPrintFunction(DebugPrint);

    // Once the kernel initialization has been complete, the user can add their
    // application thread(s) and idle thread. Threads added before the kernel
    // is started are referred to as the "static threads" in the system, as they
    // are the default working-set of threads that make up the application on
    // kernel startup.

    // Initialize the application thread to use a specified word-array as its stack.
    // The thread will run at priority level "1", and start execution the
    // "AppMain" function when it's started.
    clAppThread.Init(awAppStack, sizeof(awAppStack), 1, AppMain, 0);

    // Initialize the idle thread to use a specific word-array as its stack.
    // The thread will run at priority level "0", which is reserved for the idle
    // priority thread. IdleMain will be run when the thread is started.
    clIdleThread.Init(awIdleStack, sizeof(awIdleStack), 0, IdleMain, 0);

    // Once the static threads have been added, the user must then ensure that the
    // threads are ready to execute. By default, creating a thread is created
    // in a STOPPED state. All threads must manually be started using the
    // Start() API before they will be scheduled by the system. Here, we are
    // starting the application and idle threads before starting the kernel - and
    // that's OK. When the kernel is started, it will choose which thread to run
    // first from the pool of ready threads.

    clAppThread.Start();
    clIdleThread.Start();
}

```



```

// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
Thread clAppThread;
K_WORD awAppStack[PORT_KERNEL_DEFAULT_STACK_SIZE / sizeof(
    K_WORD)];
void AppMain(void* unused_);

//-----
// This block declares the thread data for the idle thread. It defines a
// thread object, stack (in word-array form), and the entry-point function
// used by the idle thread.
Thread clIdleThread;
K_WORD awIdleStack[PORT_KERNEL_DEFAULT_STACK_SIZE / sizeof(
    K_WORD)];
void IdleMain(void* unused_);

//-----
// Declare 3 co-routines, each with a variable which will be passed into its
// handle function.
Coroutine clCoroutine1;
int iCounter1;

Coroutine clCoroutine2;
int iCounter2;

Coroutine clCoroutine3;
int iCounter3;

//-----
// Declare handler functions used by the three coroutines. In this example,
// the coroutine tasks simply increments a counter and print their identity.
void CoroutineTask1(Coroutine* pclCaller_, void* pvArg_)
{
    Kernel::DebugPrint("Task1\n");
    int* piCounter = static_cast<int*>(pvArg_);
    (*piCounter)++;
}

//-----
void CoroutineTask2(Coroutine* pclCaller_, void* pvArg_)
{
    Kernel::DebugPrint("Task2\n");
    int* piCounter = static_cast<int*>(pvArg_);
    (*piCounter)++;
}

//-----
void CoroutineTask3(Coroutine* pclCaller_, void* pvArg_)
{
    Kernel::DebugPrint("Task3\n");
    int* piCounter = static_cast<int*>(pvArg_);
    (*piCounter)++;
}

//-----
void AppMain(void* unused_)
{
    // In this example, the application task is responsible for activating the
    // various coroutines in a round-robin fashion.
    while (1) {
        Thread::Sleep(100);
        clCoroutine1.Activate();
        Thread::Sleep(100);
        clCoroutine2.Activate();
        Thread::Sleep(100);
        clCoroutine3.Activate();
    }
}

//-----
void IdleMain(void* unused_)
{
    while (1) {
        // In this example, use the Idle context to run coroutines as they
        // are scheduled
        auto* pclCoRoutine = CoScheduler::Schedule();
        if (pclCoRoutine != nullptr) {
            pclCoRoutine->Run();
        }
    }
}

} // anonymous namespace

using namespace Mark3;
//-----
int main(void)
{

```



```

void App2Main(void* unused_);

//-----
// idle thread -- do nothing
Thread clIdleThread;
K_WORD awIdleStack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void IdleMain(void* /*unused*/)
{
    while (1) {}
}

//-----
void App1Main(void* unused_)
{
    // Simple loop that increments a volatile counter to 1000000 then resets
    // it while printing a message.
    volatile uint32_t u32Counter = 0;
    while (1) {
        u32Counter++;
        if (u32Counter == 1000000) {
            u32Counter = 0;
            Kernel::DebugPrint("Thread 1 - Did some work\n");
        }
    }
}

//-----
void App2Main(void* unused_)
{
    // Same as App1Main. However, as this thread gets twice as much CPU time
    // as Thread 1, you should see its message printed twice as often as the
    // above function.
    volatile uint32_t u32Counter = 0;
    while (1) {
        u32Counter++;
        if (u32Counter == 1000000) {
            u32Counter = 0;
            Kernel::DebugPrint("Thread 2 - Did some work\n");
        }
    }
}

} // anonymous namespace

using namespace Mark3;
//-----
int main(void)
{
    // See the annotations in lab1.
    Kernel::Init();
    Kernel::SetDebugPrintFunction(DebugPrint);

    clIdleThread.Init(awIdleStack, sizeof(awIdleStack), 0, IdleMain, 0);
    clIdleThread.Start();

    // In this exercise, we create two threads at the same priority level.
    // As a result, the CPU will automatically swap between these threads
    // at runtime to ensure that each get a chance to execute.

    clApp1Thread.Init(awApp1Stack, sizeof(awApp1Stack), 1, App1Main, 0);
    clApp2Thread.Init(awApp2Stack, sizeof(awApp2Stack), 1, App2Main, 0);

    // Set the threads up so that Thread 1 can get 4ms of CPU time uninterrupted,
    // but Thread 2 can get 8ms of CPU time uninterrupted. This means that
    // in an ideal situation, Thread 2 will get to do twice as much work as
    // Thread 1 - even though they share the same scheduling priority.

    // Note that if SetQuantum() isn't called on a thread, a default value
    // is set such that each thread gets equal timeslicing in the same
    // priority group by default. You can play around with these values and
    // observe how it affects the execution of both threads.

    clApp1Thread.SetQuantum(4);
    clApp2Thread.SetQuantum(8);

    clApp1Thread.Start();
    clApp2Thread.Start();

    Kernel::Start();

    return 0;
}

```



```

// Wait until the semaphore is posted from the other thread
Kernel::DebugPrint("Wait\n");
clMySem.Pend();

// Producer thread has finished doing its work -- do something to
// consume its output. Once again - a contrived example, but we
// can imagine that printing out the message is "consuming" the output
// from the other thread.
Kernel::DebugPrint("Triggered!\n");
}
}

//-----
void App2Main(void* unused_)
{
    volatile uint32_t u32Counter = 0;

    while (1) {
        // Do some work. Once the work is complete, post the semaphore. This
        // will cause the other thread to wake up and then take some action.
        // It's a bit contrived, but imagine that the results of this process
        // are necessary to drive the work done by that other thread.
        u32Counter++;
        if (u32Counter == 1000000) {
            u32Counter = 0;
            Kernel::DebugPrint("Posted\n");
            clMySem.Post();
        }
    }
} // anonymous namespace

using namespace Mark3;
//-----
int main(void)
{
    // See the annotations in previous labs for details on init.
    Kernel::Init();
    Kernel::SetDebugPrintFunction(DebugPrint);

    clIdleThread.Init(awIdleStack, sizeof(awIdleStack), 0, IdleMain, 0);
    clIdleThread.Start();

    // In this example we create two threads to illustrate the use of a
    // binary semaphore as a synchronization method between two threads.

    // Thread 1 is a "consumer" thread -- It waits, blocked on the semaphore
    // until thread 2 is done doing some work. Once the semaphore is posted,
    // the thread is unblocked, and does some work.

    // Thread 2 is thus the "producer" thread -- It does work, and once that
    // work is done, the semaphore is posted to indicate that the other thread
    // can use the producer's work product.

    clApp1Thread.Init(awApp1Stack, APP1_STACK_SIZE, 1, App1Main, 0);
    clApp2Thread.Init(awApp2Stack, APP2_STACK_SIZE, 1, App2Main, 0);

    clApp1Thread.Start();
    clApp2Thread.Start();

    // Initialize a binary semaphore (maximum value of one, initial value of
    // zero).
    clMySem.Init(0, 1);

    Kernel::Start();

    return 0;
}

```

21.7 lab5_mutexes/main.cpp

This example demonstrates how to use mutexes to protect against concurrent access to resources.

[illegible]

```

--[Mark3 Realtime Platform]-----
Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
See license.txt for more information
=====*/
#include "mark3.h"

/*=====

Lab Example 5:  using Mutexes.

Lessons covered in this example include:
-You can use mutexes to lock accesses to a shared resource

Takeaway:

=====*/
extern "C" {
void __cxa_pure_virtual(void) {}
void DebugPrint(const char* szString_);
}
namespace
{
using namespace Mark3;

//-----
// This block declares the thread data for one main application thread.  It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
#define APP1_STACK_SIZE (PORT_KERNEL_DEFAULT_STACK_SIZE)
Thread clApp1Thread;
K_WORD awApp1Stack[APP1_STACK_SIZE];
void App1Main(void* unused_);

//-----
// This block declares the thread data for one main application thread.  It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
#define APP2_STACK_SIZE (PORT_KERNEL_DEFAULT_STACK_SIZE)
Thread clApp2Thread;
K_WORD awApp2Stack[APP2_STACK_SIZE];
void App2Main(void* unused_);

//-----
// idle thread -- do nothing
Thread clIdleThread;
K_WORD awIdleStack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void IdleMain(void* /*unused*/)
{
    while (1) {}
}

//-----
// This is the mutex that we'll use to synchronize two threads in this
// demo application.
Mutex clMyMutex;

// This counter variable is the "shared resource" in the example, protected
// by the mutex.  Only one thread should be given access to the counter at
// any time.
volatile uint32_t u32Counter = 0;

//-----
void App1Main(void* unused_)
{
    while (1) {
        // Claim the mutex.  This will prevent any other thread from claiming
        // this lock simulatenously.  As a result, the other thread has to
        // wait until we're done before it can do its work.  You will notice
        // that the Start/Done prints for the thread will come as a pair (i.e.
        // you won't see "Thread2: Start" then "Thread1: Start").

        clMyMutex.Claim();

        // Start our work (incrementing a counter).  Notice that the Start and
        // Done prints wind up as a pair when simuated with f1AVR.

        Kernel::DebugPrint("Thread1: Start\n");
        u32Counter++;
        while (u32Counter <= 1000000) { u32Counter++; }
        u32Counter = 0;
        Kernel::DebugPrint("Thread1: Done\n");

        // Release the lock, allowing the other thread to do its thing.
        clMyMutex.Release();
    }
}

```



```
//-----
void App2Main(void* unused_)
{
    while (1) {
        // Claim the mutex. This will prevent any other thread from claiming
        // this lock simulatenously. As a result, the other thread has to
        // wait until we're done before it can do its work. You will notice
        // that the Start/Done prints for the thread will come as a pair (i.e.
        // you won't see "Thread2: Start" then "Thread1: Start").

        clMyMutex.Claim();

        // Start our work (incrementing a counter). Notice that the Start and
        // Done prints wind up as a pair when simulated with flAVR.

        Kernel::DebugPrint("Thread2: Start\n");
        u32Counter++;
        while (u32Counter <= 1000000) { u32Counter++; }
        u32Counter = 0;
        Kernel::DebugPrint("Thread2: Done\n");

        // Release the lock, allowing the other thread to do its thing.
        clMyMutex.Release();
    }
}
} // anonymous namespace

using namespace Mark3;

//-----
int main(void)
{
    // See the annotations in previous labs for details on init.
    Kernel::Init();
    Kernel::SetDebugPrintFunction(DebugPrint);

    clIdleThread.Init(awIdleStack, sizeof(awIdleStack), 0, IdleMain, 0);
    clIdleThread.Start();

    clApp1Thread.Init(awApp1Stack, sizeof(awApp1Stack), 1, App1Main, 0);
    clApp2Thread.Init(awApp2Stack, sizeof(awApp2Stack), 1, App2Main, 0);

    clApp1Thread.Start();
    clApp2Thread.Start();

    // Initialize the mutex used in this example.
    clMyMutex.Init();

    Kernel::Start();

    return 0;
}
```

21.8 lab6_timers/main.cpp

This example demonstrates how to create and use software timers.

[illegible]

Demonstration of the periodic and one-shot timer APIs provided by Mark3

Takeaway:

Mark3 can be used to provide flexible one-shot and periodic timers.

```

=====*/
extern "C" {
void __cxa_pure_virtual(void) {}
void DebugPrint(const char* szString_);
}

namespace
{
using namespace Mark3;
//-----
// This block declares the thread data for one main application thread. It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
#define APP1_STACK_SIZE (PORT_KERNEL_DEFAULT_STACK_SIZE)
Thread clApp1Thread;
K_WORD awApp1Stack[APP1_STACK_SIZE];
void App1Main(void* unused_);

//-----
// idle thread -- do nothing
Thread clIdleThread;
K_WORD awIdleStack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void IdleMain(void* /*unused_*/)
{
    while (1) {}
}

//-----
void PeriodicCallback(Thread* owner, void* pvData_)
{
    // Timer callback function used to post a semaphore. Posting the semaphore
    // will wake up a thread that's pending on that semaphore.
    auto* pclSem = static_cast<Semaphore*>(pvData_);
    pclSem->Post();
}

//-----
void OneShotCallback(Thread* owner, void* pvData_)
{
    Kernel::DebugPrint("One-shot timer expired.\n");
}

//-----
void App1Main(void* unused_)
{
    Timer clMyTimer; // Periodic timer object
    Timer clOneShot; // One-shot timer object

    Semaphore clMySem; // Semaphore used to wake this thread

    // Initialize a binary semaphore (maximum value of one, initial value of
    // zero).
    clMySem.Init(0, 1);

    // Start a timer that triggers every 500ms that will call PeriodicCallback.
    // This timer simulates an external stimulus or event that would require
    // an action to be taken by this thread, but would be serviced by an
    // interrupt or other high-priority context.

    // PeriodicCallback will post the semaphore which wakes the thread
    // up to perform an action. Here that action consists of a trivial message
    // print.
    clMyTimer.Start(true, 500, PeriodicCallback, (void*)&clMySem);

    // Set up a one-shot timer to print a message after 2.5 seconds, asynchronously
    // from the execution of this thread.
    clOneShot.Start(false, 2500, OneShotCallback, 0);

    while (1) {
        // Wait until the semaphore is posted from the timer expiry
        clMySem.Pend();

        // Take some action after the timer posts the semaphore to wake this
        // thread.
        Kernel::DebugPrint("Thread Triggered.\n");
    }
} // anonymous namespace

using namespace Mark3;

```

```
//-----
int main(void)
{
    // See the annotations in previous labs for details on init.
    Kernel::Init();
    Kernel::SetDebugPrintFunction(DebugPrint);

    clIdleThread.Init(awIdleStack, sizeof(awIdleStack), 0, IdleMain, 0);
    clIdleThread.Start();

    clApp1Thread.Init(awApp1Stack, sizeof(awApp1Stack), 1, App1Main, 0);
    clApp1Thread.Start();

    Kernel::Start();

    return 0;
}
```

21.9 lab7_events/main.cpp

This example demonstrates how to create and use event groups

```

=====
[-----] [-----] [-----] [-----] [-----] [-----]
|         |         |         |         |         |         | | | | | | |
| \       /        \|      /\     \|    /   \|   //    ||
|  \     /         \|     /\    \|   /    \|  //     ||
|   \   /          \|    /\   \|  /    \| //      ||
|    \ /           \|   /\   \| /    \|//       ||
|     V            \|  /\   \|/    \|//        ||
|_____|_____||_____|_____||_____|_____||_____|_____||
--[Mark3 Realtime Platform]-----

Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
See license.txt for more information
=====*/
#include "mark3.h"

/*=====

Lab Example 7: using Event Flags

Lessons covered in this example include:
-Using the EventFlag Class to synchronize thread execution
-Explore the behavior of the EventFlagOperation::Any_Set and EventFlagOperation::All_Set, and the event-mask bitfield.

Takeaway:

Like Semaphores and Mutexes, EventFlag objects can be used to synchronize
the execution of threads in a system. The EventFlag class allows for many
threads to share the same object, blocking on different event combinations.
This provides an efficient, robust way for threads to process asynchronous
system events that occur with a unified interface.

=====*/

extern "C" {
void __cxa_pure_virtual(void) {}
void DebugPrint(const char* szString_);
}

namespace
{
using namespace Mark3;

//-----
// This block declares the thread data for one main application thread. It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
#define APP1_STACK_SIZE (PORT_KERNEL_DEFAULT_STACK_SIZE)
Thread clApplThread;
K_WORD awApplStack[APP1_STACK_SIZE];
void ApplMain(void* unused_);

//-----
// This block declares the thread data for one main application thread. It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
#define APP2_STACK_SIZE (PORT_KERNEL_DEFAULT_STACK_SIZE)

```

```

Thread clApp2Thread;
K_WORD awApp2Stack[APP2_STACK_SIZE];
void App2Main(void* unused_)

//-----
// idle thread -- do nothing
Thread clIdleThread;
K_WORD awIdleStack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void IdleMain(void* /*unused_*/)
{
    while (1) {}
}

//-----
EventFlag clFlags;

//-----
void App1Main(void* unused_)
{
    while (1) {
        // Block this thread until any of the event flags have been set by
        // some outside force (here, we use Thread 2). As an exercise to the
        // user, try playing around with the event mask to see the effect it
        // has on which events get processed. Different threads can block on
        // different bitmasks - this allows events with different real-time
        // priorities to be handled in different threads, while still using
        // the same event-flag object.

        // Also note that EventFlagOperation::Any_Set indicates that the thread will be
        // unblocked whenever any of the flags in the mask are selected. If
        // you wanted to trigger an action that only takes place once multiple
        // bits are set, you could block the thread waiting for a specific
        // event bitmask with EventFlagOperation::All_Set specified.
        auto ul6Flags = clFlags.Wait(0xFFFF, EventFlagOperation::Any_Set);

        // Print a message indicating which bit was set this time.
        switch (ul6Flags) {
            case 0x0001: Kernel::DebugPrint("Event1\n"); break;
            case 0x0002: Kernel::DebugPrint("Event2\n"); break;
            case 0x0004: Kernel::DebugPrint("Event3\n"); break;
            case 0x0008: Kernel::DebugPrint("Event4\n"); break;
            case 0x0010: Kernel::DebugPrint("Event5\n"); break;
            case 0x0020: Kernel::DebugPrint("Event6\n"); break;
            case 0x0040: Kernel::DebugPrint("Event7\n"); break;
            case 0x0080: Kernel::DebugPrint("Event8\n"); break;
            case 0x0100: Kernel::DebugPrint("Event9\n"); break;
            case 0x0200: Kernel::DebugPrint("Event10\n"); break;
            case 0x0400: Kernel::DebugPrint("Event11\n"); break;
            case 0x0800: Kernel::DebugPrint("Event12\n"); break;
            case 0x1000: Kernel::DebugPrint("Event13\n"); break;
            case 0x2000: Kernel::DebugPrint("Event14\n"); break;
            case 0x4000: Kernel::DebugPrint("Event15\n"); break;
            case 0x8000: Kernel::DebugPrint("Event16\n"); break;
            default: break;
        }

        // Clear the event-flag that we just printed a message about. This
        // will allow ul6 to acknowledge further events in that bit in the future.
        clFlags.Clear(ul6Flags);
    }
}

//-----
void App2Main(void* unused_)
{
    uint16_t ul6Flag = 1;
    while (1) {
        Thread::Sleep(100);

        // Event flags essentially map events to bits in a bitmap. Here we
        // set one bit each 100ms. In this loop, we cycle through bits 0-15
        // repeatedly. Note that this will wake the other thread, which is
        // blocked, waiting for *any* of the flags in the bitmap to be set.
        clFlags.Set(ul6Flag);

        // Bitshift the flag value to the left. This will be the flag we set
        // the next time this thread runs through its loop.
        if (ul6Flag != 0x8000) {
            ul6Flag <<= 1;
        } else {
            ul6Flag = 1;
        }
    }
}
} // anonymous namespace

using namespace Mark3;

```

```
//
int main(void)
{
    // See the annotations in previous labs for details on init.
    Kernel::Init();
    Kernel::SetDebugPrintFunction(DebugPrint);

    clIdleThread.Init(awIdleStack, sizeof(awIdleStack), 0, IdleMain, 0);
    clIdleThread.Start();

    clApp1Thread.Init(awApp1Stack, sizeof(awApp1Stack), 1, App1Main, 0);
    clApp2Thread.Init(awApp2Stack, sizeof(awApp2Stack), 1, App2Main, 0);

    clApp1Thread.Start();
    clApp2Thread.Start();

    clFlags.Init();

    Kernel::Start();

    return 0;
}
```

21.10 lab8_messages/main.cpp

This example demonstrates how to pass data between threads using message passing.

```

/*
=====
--[Mark3 Realtime Platform]-----
Copyright (c) 2012 - 2019 m0slevin, all rights reserved.
See license.txt for more information
=====*/
#include "mark3.h"
/*=====

Lab Example 8:  using messages for IPC.

In this example, we present a typical asynchronous producer/consumer pattern
using Mark3's message-driven IPC.

Lessons covered in this example include:
- use of Message and MessageQueue objects to send data between threads
- use of GlobalMessagePool to allocate and free message objects

Takeaway:

Unlike cases presented in previous examples that relied on semaphores or
event flags, messages carry substantial context, specified in its "code" and
"data" members.  This mechanism can be used to pass data between threads
extremely efficiently, with a simple and flexible API.  Any number of threads
can write to/block on a single message queue, which give this method of
IPC even more flexibility.

=====*/
extern "C" {
void __cxa_pure_virtual(void) {}
void DebugPrint(const char* szString_);
}

namespace
{
using namespace Mark3;
//-----
// This block declares the thread data for one main application thread.  It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
#define APP1_STACK_SIZE (PORT_KERNEL_DEFAULT_STACK_SIZE)
Thread clApp1Thread;
K_WORD awApp1Stack[APP1_STACK_SIZE];
void App1Main(void* unused_);
}

```

```

//-----
// This block declares the thread data for one main application thread. It
// defines a thread object, stack (in word-array form), and the entry-point
// function used by the application thread.
#define APP2_STACK_SIZE (PORT_KERNEL_DEFAULT_STACK_SIZE)
Thread clApp2Thread;
K_WORD awApp2Stack[APP2_STACK_SIZE];
void App2Main(void* unused_);

//-----
// idle thread -- do nothing
Thread clIdleThread;
K_WORD awIdleStack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void IdleMain(void* /*unused_*/)
{
    while (1) {}
}

//-----
MessageQueue clMsgQ;

#define MESSAGE_POOL_SIZE (3)
MessagePool s_clMessagePool;
Message s_clMessages[MESSAGE_POOL_SIZE];

//-----
void App1Main(void* unused_)
{
    auto ul6Data = 0;
    while (1) {
        // This thread grabs a message from the global message pool, sets a
        // code-value and the message data pointer, then sends the message to
        // a message queue object. Another thread (Thread2) is blocked, waiting
        // for a message to arrive in the queue.

        // Get the message object
        auto* pclMsg = s_clMessagePool.Pop();

        // Set the message object's data (contrived in this example)
        pclMsg->SetCode(0x1337);
        ul6Data++;
        pclMsg->SetData(&ul6Data);

        // Send the message to the shared message queue
        clMsgQ.Send(pclMsg);

        // Wait before sending another message.
        Thread::Sleep(200);
    }
}

//-----
void App2Main(void* unused_)
{
    while (1) {
        // This thread waits until it receives a message on the shared global
        // message queue. When it gets the message, it prints out information
        // about the message's code and data, before returning the message object
        // back to the global message pool. In a more practical application,
        // the user would typically use the code to tell the receiving thread
        // what kind of message was sent, and what type of data to expect in the
        // data field.

        // Wait for a message to arrive on the specified queue. Note that once
        // this thread receives the message, it is "owned" by the thread, and
        // must be returned back to its source message pool when it is no longer
        // needed.
        auto* pclMsg = clMsgQ.Receive();

        // We received a message, now print out its information
        Kernel::DebugPrint("Received Message\n");
        // KernelAware::Trace(0, __LINE__, pclMsg->GetCode(), *((uint16_t*)pclMsg->GetData()));

        // Done with the message, return it back to the global message queue.
        s_clMessagePool.Push(pclMsg);
    }
}

} // anonymous namespace

using namespace Mark3;
//-----
int main(void)
{
    // See the annotations in previous labs for details on init.
    Kernel::Init();
    Kernel::SetDebugPrintFunction(DebugPrint);
}

```

```
clIdleThread.Init(awIdleStack, sizeof(awIdleStack), 0, IdleMain, 0);
clIdleThread.Start();

clApp1Thread.Init(awApp1Stack, sizeof(awApp1Stack), 1, App1Main, 0);
clApp2Thread.Init(awApp2Stack, sizeof(awApp2Stack), 1, App2Main, 0);

clApp1Thread.Start();
clApp2Thread.Start();

clMsgQ.Init();

s_clMessagePool.Init();
for (int i = 0; i < MESSAGE_POOL_SIZE; i++) {
    s_clMessages[i].Init();
    s_clMessagePool.Push(&s_clMessages[i]);
}

Kernel::Start();

return 0;
}
```

21.11 lab9_dynamic_threads/main.cpp

This example demonstrates how to create and destroy threads dynamically at runtime.

[illegible]

```

K_WORD awApp2Stack[PORT_KERNEL_DEFAULT_STACK_SIZE];

//-----
// idle thread -- do nothing
Thread clIdleThread;
K_WORD awIdleStack[PORT_KERNEL_DEFAULT_STACK_SIZE];
void IdleMain(void* /*unused*/)
{
    while (1) {}
}

#define MAX_THREADS (10)
Thread* apclActiveThreads[10];
uint32_t au32ActiveTime[10];

void PrintThreadSlack(void)
{
    Kernel::DebugPrint("Stack Slack");
    for (uint8_t i = 0; i < MAX_THREADS; i++) {
        if (apclActiveThreads[i] != 0) {
            char szStr[10];

            auto u16Slack = apclActiveThreads[i]->GetStackSlack();
            MemUtil::DecimalToHex((K_ADDR)apclActiveThreads[i], szStr);
            Kernel::DebugPrint(szStr);
            Kernel::DebugPrint(" ");
            MemUtil::DecimalToString(u16Slack, szStr);
            Kernel::DebugPrint(szStr);
            Kernel::DebugPrint("\n");
        }
    }
}

void PrintCPUUsage(void)
{
    Kernel::DebugPrint("Cpu usage\n");
    for (int i = 0; i < MAX_THREADS; i++) {
        if (apclActiveThreads[i] != 0) {
            // KernelAware::Trace(0, __LINE__, (K_ADDR)apclActiveThreads[i], au16ActiveTime[i]);
        }
    }
}

void ThreadCreate(Thread* pclThread_)
{
    Kernel::DebugPrint("TC\n");
    CriticalSection::Enter();
    for (uint8_t i = 0; i < MAX_THREADS; i++) {
        if (apclActiveThreads[i] == 0) {
            apclActiveThreads[i] = pclThread_;
            break;
        }
    }
    CriticalSection::Exit();

    PrintThreadSlack();
    PrintCPUUsage();
}

void ThreadExit(Thread* pclThread_)
{
    Kernel::DebugPrint("TX\n");
    CriticalSection::Enter();
    for (uint8_t i = 0; i < MAX_THREADS; i++) {
        if (apclActiveThreads[i] == pclThread_) {
            apclActiveThreads[i] = 0;
            au32ActiveTime[i] = 0;
            break;
        }
    }
    CriticalSection::Exit();

    PrintThreadSlack();
    PrintCPUUsage();
}

void ThreadContextSwitch(Thread* pclThread_)
{
    Kernel::DebugPrint("CS\n");
    static uint32_t u32LastTicks = 0;
    auto u32Ticks = Kernel::GetTicks();

    CriticalSection::Enter();
    for (uint8_t i = 0; i < MAX_THREADS; i++) {
        if (apclActiveThreads[i] == pclThread_) {
            au32ActiveTime[i] += u32Ticks - u32LastTicks;
            break;
        }
    }
}

```



```

    }
}
CriticalSection::Exit();

u32LastTicks = u32Ticks;
}

//-----
void WorkerMain1(void* arg_)
{
    auto*    pclSem    = static_cast<Semaphore*>(arg_);
    uint32_t u32Count = 0;

    // Do some work. Post a semaphore to notify the other thread that the
    // work has been completed.
    while (u32Count < 1000000) { u32Count++; }

    Kernel::DebugPrint("Worker1 -- Done Work\n");
    pclSem->Post();

    // Work is completed, just spin now. Let another thread destroy u16.
    while (1) {}
}

//-----
void WorkerMain2(void* arg_)
{
    uint32_t u32Count = 0;
    while (u32Count < 1000000) { u32Count++; }

    Kernel::DebugPrint("Worker2 -- Done Work\n");

    // A dynamic thread can self-terminate as well:
    Scheduler::GetCurrentThread()->Exit();
}

//-----
void App1Main(void* unused_)
{
    Thread    clMyThread;
    Semaphore clMySem;

    clMySem.Init(0, 1);
    while (1) {
        // Example 1 - create a worker thread at our current priority in order to
        // parallelize some work.
        clMyThread.Init(awApp2Stack, sizeof(awApp2Stack), 1, WorkerMain1, (void*)&clMySem);
        clMyThread.Start();

        // Do some work of our own in parallel, while the other thread works on its project.
        uint32_t u32Count = 0;
        while (u32Count < 1000000) { u32Count++; }

        Kernel::DebugPrint("Thread -- Done Work\n");

        PrintThreadSlack();

        // Wait for the other thread to finish its job.
        clMySem.Pend();

        // Once the thread has signalled u16, we can safely call "Exit" on the thread to
        // remove it from scheduling and recycle it later.
        clMyThread.Exit();

        // Spin the thread up again to do something else in parallel. This time, the thread
        // will run completely asynchronously to this thread.
        clMyThread.Init(awApp2Stack, sizeof(awApp2Stack), 1, WorkerMain2, 0);
        clMyThread.Start();

        u32Count = 0;
        while (u32Count < 1000000) { u32Count++; }

        Kernel::DebugPrint("Thread -- Done Work\n");

        // Check that we're sure the worker thread has terminated before we try running the
        // test loop again.
        while (clMyThread.GetState() != ThreadState::Exit) {}

        Kernel::DebugPrint("  Test Done\n");
        Thread::Sleep(1000);
        PrintThreadSlack();
    }
}

} // anonymous namespace

using namespace Mark3;
//-----
int main(void)

```

```
{  
    // See the annotations in previous labs for details on init.  
    Kernel::Init();  
    Kernel::SetDebugPrintFunction(DebugPrint);  
  
    Kernel::SetThreadCreateCallout(ThreadCreate);  
    Kernel::SetThreadExitCallout(ThreadExit);  
    Kernel::SetThreadContextSwitchCallout(ThreadContextSwitch);  
  
    clIdleThread.Init(awIdleStack, sizeof(awIdleStack), 0, IdleMain, 0);  
    clIdleThread.Start();  
  
    clApp1Thread.Init(awApp1Stack, sizeof(awApp1Stack), 1, App1Main, 0);  
    clApp1Thread.Start();  
    Kernel::Start();  
  
    return 0;  
}
```

Index

/home/moslevin/projects/m3-repo/kernel/lib/mark3c/src/mark3c.↵ cpp, 475, 476
cpp, 295, 346 /home/moslevin/projects/m3-repo/kernel/src/message.↵
/home/moslevin/projects/m3-repo/kernel/lib/mark3c/src/public/fake.↵ cpp, 478, 479
_types.h, 358 /home/moslevin/projects/m3-repo/kernel/src/mutex.cpp,
/home/moslevin/projects/m3-repo/kernel/lib/mark3c/src/public/mark3c.↵ 480, 481
h, 360, 419 /home/moslevin/projects/m3-repo/kernel/src/notify.cpp,
/home/moslevin/projects/m3-repo/kernel/lib/memutil/memutil.↵ 483, 484
cpp, 424, 425 /home/moslevin/projects/m3-repo/kernel/src/profile.cpp,
/home/moslevin/projects/m3-repo/kernel/lib/memutil/public/memutil.↵ 486
h, 431, 432 /home/moslevin/projects/m3-repo/kernel/src/public/atomic.↵
/home/moslevin/projects/m3-repo/kernel/lib/streamer/public/streamer.h, 487, 488
h, 433, 434 /home/moslevin/projects/m3-repo/kernel/src/public/autalloc.↵
/home/moslevin/projects/m3-repo/kernel/lib/streamer/streamer.↵ h, 489, 490
cpp, 434, 435 /home/moslevin/projects/m3-repo/kernel/src/public/blocking.↵
/home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/avr/asm.↵
cpp, 438 /home/moslevin/projects/m3-repo/kernel/src/public/colist.↵
/home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/avr/asm.↵
cpp, 439, 441 /home/moslevin/projects/m3-repo/kernel/src/public/condvar.↵
/home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/avr/asm.↵
h, 442, 447 /home/moslevin/projects/m3-repo/kernel/src/public/coroutine.↵
/home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/avr/asm.↵
h, 448, 450 /home/moslevin/projects/m3-repo/kernel/src/public/cosched.↵
/home/moslevin/projects/m3-repo/kernel/src/arch/avr/atmega1284p/gcc/avr/asm.↵
cpp, 452, 453 /home/moslevin/projects/m3-repo/kernel/src/public/criticalguard.↵
/home/moslevin/projects/m3-repo/kernel/src/atomic.↵ h, 498
cpp, 455 /home/moslevin/projects/m3-repo/kernel/src/public/criticalsection.↵
/home/moslevin/projects/m3-repo/kernel/src/autalloc.↵ h, 499
cpp, 456 /home/moslevin/projects/m3-repo/kernel/src/public/eventflag.↵
/home/moslevin/projects/m3-repo/kernel/src/blocking.↵ h, 500
cpp, 457, 458 /home/moslevin/projects/m3-repo/kernel/src/public/ithreadport.↵
/home/moslevin/projects/m3-repo/kernel/src/colist.cpp, h, 501, 502
459 /home/moslevin/projects/m3-repo/kernel/src/public/kernel.↵
/home/moslevin/projects/m3-repo/kernel/src/condvar.↵ h, 502, 503
cpp, 460 /home/moslevin/projects/m3-repo/kernel/src/public/kerneldebug.↵
/home/moslevin/projects/m3-repo/kernel/src/coroutine.↵ h, 504, 505
cpp, 461, 462 /home/moslevin/projects/m3-repo/kernel/src/public/kernelswi.↵
/home/moslevin/projects/m3-repo/kernel/src/cosched.↵ h, 506
cpp, 463 /home/moslevin/projects/m3-repo/kernel/src/public/kerneltimer.↵
/home/moslevin/projects/m3-repo/kernel/src/eventflag.↵ h, 507
cpp, 464 /home/moslevin/projects/m3-repo/kernel/src/public/kerneltypes.↵
/home/moslevin/projects/m3-repo/kernel/src/kernel.cpp, h, 508
468 /home/moslevin/projects/m3-repo/kernel/src/public/ksemaphore.↵
/home/moslevin/projects/m3-repo/kernel/src/ksemaphore.↵ h, 509, 510
cpp, 469, 470 /home/moslevin/projects/m3-repo/kernel/src/public/ll.h,
/home/moslevin/projects/m3-repo/kernel/src/ll.cpp, 472, 510, 511
473 /home/moslevin/projects/m3-repo/kernel/src/public/lockguard.↵
/home/moslevin/projects/m3-repo/kernel/src/lockguard.↵ h, 513, 514
cpp, 474, 475 /home/moslevin/projects/m3-repo/kernel/src/public/mailbox.↵
/home/moslevin/projects/m3-repo/kernel/src/mailbox.↵ h, 514, 515

/home/moslevin/projects/m3-repo/kernel/src/public/manual.h, 516, 517
 /home/moslevin/projects/m3-repo/kernel/src/public/mark3.h, 517, 518
 /home/moslevin/projects/m3-repo/kernel/src/public/mark3cfg.h, 518, 521
 /home/moslevin/projects/m3-repo/kernel/src/public/message.h, 522, 523
 /home/moslevin/projects/m3-repo/kernel/src/public/mutex.h, 524, 525
 /home/moslevin/projects/m3-repo/kernel/src/public/notify.h, 526, 527
 /home/moslevin/projects/m3-repo/kernel/src/public/paniccodes.h, 527, 530
 /home/moslevin/projects/m3-repo/kernel/src/public/priomap.h, 530, 531
 /home/moslevin/projects/m3-repo/kernel/src/public/priomap1.h, 531, 532
 /home/moslevin/projects/m3-repo/kernel/src/public/priomap2.h, 533, 534
 /home/moslevin/projects/m3-repo/kernel/src/public/profile.h, 535, 536
 /home/moslevin/projects/m3-repo/kernel/src/public/profiling_results.h, 536
 /home/moslevin/projects/m3-repo/kernel/src/public/quantum.h, 537
 /home/moslevin/projects/m3-repo/kernel/src/public/readerwriter.h, 538, 539
 /home/moslevin/projects/m3-repo/kernel/src/public/scheduler.h, 539, 541
 /home/moslevin/projects/m3-repo/kernel/src/public/schedulerguard.h, 541, 542
 /home/moslevin/projects/m3-repo/kernel/src/public/sizeprofile.h, 542
 /home/moslevin/projects/m3-repo/kernel/src/public/thread.h, 543, 544
 /home/moslevin/projects/m3-repo/kernel/src/public/threadlist.h, 546, 547
 /home/moslevin/projects/m3-repo/kernel/src/public/threadlistlist.h, 547, 548
 /home/moslevin/projects/m3-repo/kernel/src/public/timer.h, 548, 549
 /home/moslevin/projects/m3-repo/kernel/src/public/timerlist.h, 550, 551
 /home/moslevin/projects/m3-repo/kernel/src/public/timerscheduler.h, 551, 552
 /home/moslevin/projects/m3-repo/kernel/src/quantum.cpp, 552, 553
 /home/moslevin/projects/m3-repo/kernel/src/readerwriter.cpp, 554
 /home/moslevin/projects/m3-repo/kernel/src/scheduler.cpp, 555, 556
 /home/moslevin/projects/m3-repo/kernel/src/thread.cpp, 557, 558
 /home/moslevin/projects/m3-repo/kernel/src/threadlist.cpp, 563
 /home/moslevin/projects/m3-repo/kernel/src/threadlistlist.cpp, 565

565, 566
 BlockingObject
 Mark3::BlockingObject, 111
 Coroutine
 Mark3::Coroutine, 125
 CriticalGuard
 Mark3::CriticalGuard, 131
 EventFlag
 Mark3::EventFlag, 137
 LockGuard
 Mark3::LockGuard, 159
 Mailbox
 Mark3::Mailbox, 163
 MessagePool
 Mark3::MessagePool, 186
 MessageQueue
 Mark3::MessageQueue, 190
 Mutex
 Mark3::Mutex, 194
 Notify
 Mark3::Notify, 200
 SchedulerGuard
 Mark3::SchedulerGuard, 228
 Semaphore
 Mark3::Semaphore, 230
 Thread
 Mark3::Thread, 245
 Timer
 Mark3::Timer, 275
 ASM
 threadport.h, 449
 AVR
 portcfg.h, 443
 AcquireReader
 Mark3::ReaderWriterLock, 218
 AcquireReader_i
 Mark3::ReaderWriterLock, 218
 AcquireWriter
 Mark3::ReaderWriterLock, 219
 AcquireWriter_i
 Mark3::ReaderWriterLock, 219
 Activate
 Mark3::Coroutine, 125
 Add
 Mark3::Atomic, 103
 Mark3::CircularLinkedList, 116
 Mark3::CoList, 118
 Mark3::DoubleLinkedList, 134
 Mark3::Scheduler, 223
 Mark3::ThreadList, 265
 Mark3::ThreadListList, 270
 Mark3::TimerList, 281
 Mark3::TimerScheduler, 284
 Mark3::TypedCircularLinkedList, 287
 Mark3::TypedDoubleLinkedList, 291

- AddPriority
 - Mark3::ThreadList, [267](#)
- Alloc_ConditionVariable
 - mark3c.cpp, [300](#)
- Alloc_EventFlag
 - mark3c.cpp, [300](#)
- Alloc_Mailbox
 - mark3c.cpp, [300](#)
 - mark3c.h, [375](#)
- Alloc_Memory
 - mark3c.cpp, [300](#)
 - mark3c.h, [375](#)
- Alloc_Message
 - mark3c.cpp, [300](#)
 - mark3c.h, [376](#)
- Alloc_MessagePool
 - mark3c.cpp, [301](#)
 - mark3c.h, [376](#)
- Alloc_MessageQueue
 - mark3c.cpp, [301](#)
 - mark3c.h, [376](#)
- Alloc_Mutex
 - mark3c.cpp, [301](#)
 - mark3c.h, [376](#)
- Alloc_Notify
 - mark3c.cpp, [302](#)
 - mark3c.h, [377](#)
- Alloc_ReaderWriterLock
 - mark3c.cpp, [302](#)
- Alloc_Semaphore
 - mark3c.cpp, [302](#)
 - mark3c.h, [377](#)
- Alloc_Thread
 - mark3c.cpp, [303](#)
 - mark3c.h, [377](#)
- Alloc_Timer
 - mark3c.cpp, [303](#)
 - mark3c.h, [378](#)
- Allocate
 - Mark3::AutoAlloc, [106](#)
- Atomic_Add16
 - mark3c.cpp, [303](#)
 - mark3c.h, [378](#)
- Atomic_Add32
 - mark3c.cpp, [304](#)
 - mark3c.h, [379](#)
- Atomic_Add8
 - mark3c.cpp, [304](#)
 - mark3c.h, [379](#)
- Atomic_Set16
 - mark3c.cpp, [305](#)
 - mark3c.h, [380](#)
- Atomic_Set32
 - mark3c.cpp, [305](#)
 - mark3c.h, [380](#)
- Atomic_Set8
 - mark3c.cpp, [306](#)
 - mark3c.h, [381](#)
- Atomic_Sub16
 - mark3c.cpp, [306](#)
 - mark3c.h, [381](#)
- Atomic_Sub32
 - mark3c.cpp, [307](#)
 - mark3c.h, [382](#)
- Atomic_Sub8
 - mark3c.cpp, [307](#)
 - mark3c.h, [382](#)
- Atomic_TestAndSet
 - mark3c.cpp, [308](#)
 - mark3c.h, [383](#)
- AutoAllocAllocator_t
 - Mark3, [95](#)
- AutoAllocFree_t
 - Mark3, [95](#)
- AutoAllocType
 - Mark3, [97](#)
- Block
 - Mark3::BlockingObject, [111](#)
- BlockPriority
 - Mark3::BlockingObject, [112](#)
- Blocked
 - Mark3, [98](#)
- BlockingObject
 - Mark3::BlockingObject, [111](#)
- Broadcast
 - Mark3::ConditionVariable, [121](#)
- CONDITIONVARIABLE_SIZE
 - mark3c.h, [366](#)
- CanRead
 - Mark3::Streamer, [236](#)
- CanWrite
 - Mark3::Streamer, [236](#)
- Cancel
 - Mark3::Quantum, [214](#)
- Checksum16
 - Mark3::MemUtil, [174](#)
- Checksum8
 - Mark3::MemUtil, [174](#)
- CircularLinkList
 - Mark3::CircularLinkList, [115](#)
 - Mark3::LinkListNode, [157](#)
- Claim
 - Mark3::Mutex, [194](#), [195](#)
 - Mark3::Streamer, [236](#)
- Claim_i
 - Mark3::Mutex, [195](#)
- Clear
 - Mark3::EventFlag, [137](#)
 - Mark3::PriorityMapL1, [203](#)
 - Mark3::PriorityMapL2, [207](#)
- ClearInTimer
 - Mark3::Quantum, [214](#)
- ClearNode
 - Mark3::LinkListNode, [156](#)
- CoPrioMap

- Mark3, [95](#)
- CompareMemory
 - Mark3::MemUtil, [175](#)
- CompareStrings
 - Mark3::MemUtil, [175](#), [176](#)
- CompleteStart
 - Mark3::Kernel, [142](#)
- ConditionVariable_Broadcast
 - mark3c.cpp, [308](#)
 - mark3c.h, [383](#)
- ConditionVariable_Init
 - mark3c.cpp, [309](#)
 - mark3c.h, [384](#)
- ConditionVariable_Signal
 - mark3c.cpp, [309](#)
 - mark3c.h, [384](#)
- ConditionVariable_TimedWait
 - mark3c.cpp, [310](#)
 - mark3c.h, [384](#)
- ConditionVariable_Wait
 - mark3c.cpp, [310](#)
 - mark3c.h, [385](#)
- ConditionVariable_t
 - mark3c.h, [372](#)
- Config
 - Mark3::KernelSWI, [150](#)
 - Mark3::KernelTimer, [151](#)
- ContextSwitchSWI
 - Mark3::Thread, [246](#)
- CoopYield
 - Mark3::Thread, [246](#)
- CopyData
 - Mark3::Mailbox, [163](#)
- CopyMemory
 - Mark3::MemUtil, [176](#)
- CopyString
 - Mark3::MemUtil, [177](#)
- CoroutineHandler
 - Mark3, [96](#)
- CriticalGuard
 - Mark3::CriticalGuard, [131](#)
- DECLARE_CONDITIONVARIABLE
 - mark3c.h, [366](#)
- DECLARE_MAILBOX
 - mark3c.h, [366](#)
- DECLARE_MESSAGEPOOL
 - mark3c.h, [367](#)
- DECLARE_MESSAGEQUEUE
 - mark3c.h, [367](#)
- DECLARE_MESSAGE
 - mark3c.h, [366](#)
- DECLARE_MUTEX
 - mark3c.h, [367](#)
- DECLARE_NOTIFY
 - mark3c.h, [368](#)
- DECLARE_READERWRITERLOCK
 - mark3c.h, [368](#)
- DECLARE_SEMAPHORE
 - mark3c.h, [368](#)
- DECLARE_THREAD
 - mark3c.h, [369](#)
- DECLARE_TIMER
 - mark3c.h, [369](#)
- DebugPrint
 - Mark3::Kernel, [142](#)
- DebugPrintFunction
 - kernel.h, [503](#)
- DecimalToHex
 - Mark3::MemUtil, [177](#), [178](#)
- DecimalToString
 - Mark3::MemUtil, [178](#), [179](#)
- DestroyObject
 - Mark3::AutoAlloc, [106](#)
- DestroyRawData
 - Mark3::AutoAlloc, [106](#)
- DestroyUserTypeAllocation
 - Mark3::AutoAlloc, [107](#)
- DoubleLinkedList
 - Mark3::DoubleLinkedList, [134](#)
 - Mark3::LinkListNode, [157](#)
- Enter
 - Mark3::CriticalSection, [132](#)
- EventFlagOperation
 - Mark3, [98](#)
- Exit
 - Mark3, [98](#)
 - Mark3::CriticalSection, [132](#)
 - Mark3::Thread, [246](#)
- Free
 - Mark3::AutoAlloc, [107](#)
- Free_ConditionVariable
 - mark3c.cpp, [311](#)
- Free_EventFlag
 - mark3c.cpp, [311](#)
- Free_Mailbox
 - mark3c.cpp, [311](#)
 - mark3c.h, [385](#)
- Free_Memory
 - mark3c.cpp, [311](#)
 - mark3c.h, [386](#)
- Free_Message
 - mark3c.cpp, [312](#)
 - mark3c.h, [386](#)
- Free_MessagePool
 - mark3c.cpp, [312](#)
 - mark3c.h, [386](#)
- Free_MessageQueue
 - mark3c.cpp, [312](#)
 - mark3c.h, [386](#)
- Free_Mutex
 - mark3c.cpp, [312](#)
 - mark3c.h, [386](#)
- Free_Notify
 - mark3c.cpp, [312](#)
 - mark3c.h, [387](#)

- Free_ReaderWriterLock
 - mark3c.cpp, [312](#)
- Free_Semaphore
 - mark3c.cpp, [313](#)
 - mark3c.h, [387](#)
- Free_Thread
 - mark3c.cpp, [313](#)
 - mark3c.h, [387](#)
- Free_Timer
 - mark3c.cpp, [313](#)
 - mark3c.h, [387](#)
- g_kwCriticalCount
 - Mark3, [100](#)
 - threadport.cpp, [453](#)
- g_kwSFR
 - Mark3, [100](#)
 - threadport.cpp, [453](#)
- g_pclCurrent
 - scheduler.cpp, [556](#)
 - scheduler.h, [540](#)
- g_pclNext
 - scheduler.cpp, [556](#)
 - scheduler.h, [540](#)
- GetAvailable
 - Mark3::Streamer, [237](#)
- GetAverage
 - Mark3::ProfileTimer, [211](#)
- GetCoList
 - Mark3::CoScheduler, [128](#)
- GetCode
 - Mark3::Message, [183](#)
- GetCount
 - Mark3::MessageQueue, [190](#)
 - Mark3::Semaphore, [230](#)
- GetCurPriority
 - Mark3::Thread, [246](#)
- GetCurrent
 - Mark3::ProfileTimer, [211](#)
 - Mark3::Thread, [247](#)
- GetCurrentThread
 - Mark3::Scheduler, [223](#)
- GetData
 - Mark3::Message, [183](#)
- GetEventFlagMask
 - Mark3::Thread, [247](#)
- GetEventFlagMode
 - Mark3::Thread, [247](#)
- GetExpired
 - Mark3::Thread, [248](#)
- GetExtendedContext
 - Mark3::Thread, [248](#)
- GetFreeSlots
 - Mark3::Mailbox, [163](#)
- GetHead
 - Mark3::LinkList, [153](#)
 - Mark3::MessagePool, [187](#)
 - Mark3::ThreadListList, [271](#)
 - Mark3::TypedCircularLinkList, [288](#)
 - Mark3::TypedDoubleLinkList, [292](#)
- GetHeadPointer
 - Mark3::Mailbox, [163](#)
- GetID
 - Mark3::Thread, [248](#)
- GetMask
 - Mark3::EventFlag, [137](#)
- GetName
 - Mark3::Thread, [248](#)
- GetNext
 - Mark3::LinkListNode, [156](#)
 - Mark3::TypedLinkListNode, [294](#)
- GetNextThread
 - Mark3::Scheduler, [223](#)
- GetOwner
 - Mark3::Thread, [249](#)
- GetPrev
 - Mark3::LinkListNode, [157](#)
 - Mark3::TypedLinkListNode, [294](#)
- GetPrioMap
 - Mark3::CoScheduler, [129](#)
- GetPriority
 - Mark3::Coroutine, [125](#)
 - Mark3::Thread, [249](#)
- GetQuantum
 - Mark3::Thread, [249](#)
- GetStack
 - Mark3::Thread, [250](#)
- GetStackGuardThreshold
 - Mark3::Kernel, [143](#)
- GetStackSize
 - Mark3::Thread, [250](#)
- GetStackSlack
 - Mark3::Thread, [250](#)
- GetState
 - Mark3::Thread, [251](#)
- GetStopList
 - Mark3::CoScheduler, [129](#)
 - Mark3::Scheduler, [223](#)
- GetTail
 - Mark3::LinkList, [153](#)
 - Mark3::TypedCircularLinkList, [288](#)
 - Mark3::TypedDoubleLinkList, [292](#)
- GetTailPointer
 - Mark3::Mailbox, [164](#)
- GetThreadContextSwitchCallout
 - Mark3::Kernel, [143](#)
- GetThreadCreateCallout
 - Mark3::Kernel, [143](#)
- GetThreadExitCallout
 - Mark3::Kernel, [143](#)
- GetThreadList
 - Mark3::Scheduler, [224](#)
- GetTicks
 - Mark3::Kernel, [143](#)
- GetTimer
 - Mark3::Thread, [251](#)
- HighestPriority

- Mark3::PriorityMapL1, [204](#)
- Mark3::PriorityMapL2, [208](#)
- HighestWaiter
 - Mark3::ThreadList, [267](#)
- ISR
 - kerneltimer.cpp, [440](#)
 - Mark3, [99](#)
- IdleFunc
 - Mark3, [96](#)
- InheritPriority
 - Mark3::Thread, [251](#)
- Init
 - Mark3::AutoAlloc, [107](#)
 - Mark3::CoScheduler, [129](#)
 - Mark3::ConditionVariable, [121](#)
 - Mark3::Coroutine, [125](#)
 - Mark3::EventFlag, [137](#)
 - Mark3::Kernel, [144](#)
 - Mark3::LinkList, [153](#)
 - Mark3::Mailbox, [164](#)
 - Mark3::Message, [184](#)
 - Mark3::MessagePool, [187](#)
 - Mark3::MessageQueue, [190](#)
 - Mark3::Mutex, [196](#)
 - Mark3::Notify, [200](#)
 - Mark3::ProfileTimer, [211](#)
 - Mark3::Quantum, [214](#)
 - Mark3::ReaderWriterLock, [220](#)
 - Mark3::Scheduler, [224](#)
 - Mark3::Semaphore, [230](#)
 - Mark3::Streamer, [237](#)
 - Mark3::Thread, [252](#)
 - Mark3::ThreadPort, [272](#)
 - Mark3::Timer, [276](#)
 - Mark3::TimerList, [281](#)
 - Mark3::TimerScheduler, [284](#)
- InitStack
 - Mark3::ThreadPort, [272](#)
- InsertNodeBefore
 - Mark3::CircularLinkList, [116](#)
 - Mark3::TypedCircularLinkList, [288](#)
- Invalid
 - Mark3, [98](#)
- isAcquired
 - Mark3::LockGuard, [160](#)
- IsEmpty
 - Mark3::Mailbox, [165](#)
 - Mark3::Streamer, [237](#)
- IsEnabled
 - Mark3::Scheduler, [224](#)
- IsFull
 - Mark3::Mailbox, [165](#)
- IsInitialized
 - Mark3::BlockingObject, [112](#)
 - Mark3::Thread, [253](#)
 - Mark3::Timer, [276](#)
- IsPanic
 - Mark3::Kernel, [144](#)
- IsStarted
 - Mark3::Kernel, [144](#)
- K_ADDR
 - portcfg.h, [443](#)
- K_INT
 - portcfg.h, [443](#)
- K_WORD
 - portcfg.h, [443](#)
- KERNEL_ASSERT
 - kerneldebug.h, [505](#)
- KERNEL_CONTEXT_SWITCH_CALLOUT
 - mark3cfg.h, [519](#)
- KERNEL_DEBUG
 - mark3cfg.h, [519](#)
- KERNEL_EVENT_FLAGS
 - mark3cfg.h, [519](#)
- KERNEL_EXTENDED_CONTEXT
 - mark3cfg.h, [520](#)
- KERNEL_NAMED_THREADS
 - mark3cfg.h, [520](#)
- KERNEL_NUM_PRIORITIES
 - portcfg.h, [443](#)
- KERNEL_ROUND_ROBIN
 - mark3cfg.h, [520](#)
- KERNEL_STACK_CHECK
 - mark3cfg.h, [520](#)
- KERNEL_STACK_GUARD_DEFAULT
 - portcfg.h, [444](#)
- KERNEL_THREAD_CREATE_CALLOUT
 - mark3cfg.h, [521](#)
- KERNEL_THREAD_EXIT_CALLOUT
 - mark3cfg.h, [521](#)
- KERNEL_TIMERS_THREAD_PRIORITY
 - portcfg.h, [444](#)
- kernel.h
 - DebugPrintFunction, [503](#)
- Kernel_DebugPrint
 - mark3c.cpp, [313](#)
 - mark3c.h, [387](#)
- Kernel_GetStackGuardThreshold
 - mark3c.cpp, [314](#)
- Kernel_GetTicks
 - mark3c.cpp, [314](#)
 - mark3c.h, [388](#)
- Kernel_Init
 - mark3c.cpp, [314](#)
 - mark3c.h, [388](#)
- Kernel_IsPanic
 - mark3c.cpp, [314](#)
 - mark3c.h, [388](#)
- Kernel_IsStarted
 - mark3c.cpp, [315](#)
 - mark3c.h, [389](#)
- Kernel_Panic
 - mark3c.cpp, [315](#)
 - mark3c.h, [389](#)
- Kernel_SetDebugPrintFunction
 - mark3c.cpp, [316](#)

- mark3c.h, [390](#)
- Kernel_SetPanic
 - mark3c.cpp, [316](#)
 - mark3c.h, [390](#)
- Kernel_SetStackGuardThreshold
 - mark3c.cpp, [316](#)
- Kernel_Start
 - mark3c.cpp, [317](#)
 - mark3c.h, [390](#)
- kernel_debug_print_t
 - mark3c.h, [372](#)
- KernelTimer_Task
 - Mark3, [99](#)
- kerneldebug.h
 - KERNEL_ASSERT, [505](#)
- kerneltimer.cpp
 - ISR, [440](#)
 - TCCR1B_INIT, [440](#)
 - TIMER_IFR, [440](#)
 - TIMER_IMSK, [440](#)
- LinkList
 - Mark3::LinkListNode, [157](#)
- LinkListNode
 - Mark3::LinkListNode, [156](#)
- Lock
 - Mark3::Streamer, [238](#)
- LockGuard
 - Mark3::LockGuard, [159](#)
- m_aclPriorities
 - Mark3::CoScheduler, [130](#)
 - Mark3::Scheduler, [226](#)
- m_auxPriorityMap
 - Mark3::PriorityMapL2, [209](#)
- m_bActive
 - Mark3::ProfileTimer, [212](#)
- m_bEnabled
 - Mark3::Scheduler, [226](#)
- m_bExpired
 - Mark3::Thread, [260](#)
- m_blnTimer
 - Mark3::Quantum, [216](#)
- m_blsAcquired
 - Mark3::LockGuard, [160](#)
- m_blsPanic
 - Mark3::Kernel, [148](#)
- m_blsStarted
 - Mark3::Kernel, [148](#)
- m_bPending
 - Mark3::Notify, [202](#)
- m_bQueued
 - Mark3::Coroutine, [126](#)
- m_bQueuedSchedule
 - Mark3::Scheduler, [226](#)
- m_bReady
 - Mark3::Mutex, [197](#)
- m_bRecursive
 - Mark3::Mutex, [197](#)
- m_bSchedState
 - Mark3::SchedulerGuard, [228](#)
- m_bTimerActive
 - Mark3::TimerList, [282](#)
- m_clBlockList
 - Mark3::BlockingObject, [114](#)
- m_clGlobalMutex
 - Mark3::ReaderWriterLock, [221](#)
- m_clLinkList
 - Mark3::MessageQueue, [192](#)
- m_clList
 - Mark3::MessagePool, [188](#)
- m_clMutex
 - Mark3::ConditionVariable, [123](#)
 - Mark3::TimerList, [282](#)
- m_clPrioMap
 - Mark3::CoScheduler, [130](#)
 - Mark3::Scheduler, [226](#)
- m_clReaderMutex
 - Mark3::ReaderWriterLock, [221](#)
- m_clRecvSem
 - Mark3::Mailbox, [171](#)
- m_clSemaphore
 - Mark3::ConditionVariable, [123](#)
 - Mark3::MessageQueue, [193](#)
- m_clSendSem
 - Mark3::Mailbox, [171](#)
- m_clStopList
 - Mark3::CoScheduler, [130](#)
 - Mark3::Scheduler, [227](#)
- m_clThreadListList
 - Mark3::ThreadListList, [271](#)
- m_clTimer
 - Mark3::Thread, [260](#)
- m_clTimerList
 - Mark3::TimerScheduler, [285](#)
- m_eFlagMode
 - Mark3::Thread, [260](#)
- m_eState
 - Mark3::Thread, [260](#)
- m_pau8Buffer
 - Mark3::Streamer, [240](#)
- m_pclActiveThread
 - Mark3::Quantum, [216](#)
- m_pclCurrent
 - Mark3::Thread, [261](#)
- m_pclHead
 - Mark3::LinkList, [154](#)
- m_pclMap
 - Mark3::ThreadList, [269](#)
- m_pclMutex
 - Mark3::LockGuard, [160](#)
- m_pclOwner
 - Mark3::Coroutine, [126](#)
 - Mark3::Mutex, [198](#)
 - Mark3::Thread, [261](#)
 - Mark3::Timer, [278](#)
- m_pclPrioMap

- Mark3::CoList, [120](#)
- m_pclTail
 - Mark3::LinkList, [155](#)
- m_pclTimerThread
 - Mark3::Quantum, [216](#)
- m_pfAllocator
 - Mark3::AutoAlloc, [109](#)
- m_pfCallback
 - Mark3::Timer, [278](#)
- m_pfDebugPrintFunction
 - Mark3::Kernel, [148](#)
- m_pfEntryPoint
 - Mark3::Thread, [261](#)
- m_pfFree
 - Mark3::AutoAlloc, [109](#)
- m_pfHandler
 - Mark3::Coroutine, [127](#)
- m_pfPanic
 - Mark3::Kernel, [148](#)
- m_pfThreadContextCallout
 - Mark3::Kernel, [148](#)
- m_pfThreadCreateCallout
 - Mark3::Kernel, [149](#)
- m_pfThreadExitCallout
 - Mark3::Kernel, [149](#)
- m_pu8LockAddr
 - Mark3::Streamer, [240](#)
- m_pvArg
 - Mark3::Thread, [261](#)
- m_pvBuffer
 - Mark3::Mailbox, [171](#)
- m_pvContext
 - Mark3::Coroutine, [127](#)
- m_pvData
 - Mark3::Message, [185](#)
 - Mark3::Timer, [279](#)
- m_pvExtendedContext
 - Mark3::Thread, [261](#)
- m_pwStack
 - Mark3::Thread, [262](#)
- m_pwStackTop
 - Mark3::Thread, [262](#)
- m_szName
 - Mark3::Thread, [262](#)
- m_u16Avail
 - Mark3::Streamer, [240](#)
- m_u16Code
 - Mark3::Message, [185](#)
- m_u16Count
 - Mark3::Mailbox, [172](#)
- m_u16ElementSize
 - Mark3::Mailbox, [172](#)
- m_u16FlagMask
 - Mark3::Thread, [262](#)
- m_u16Free
 - Mark3::Mailbox, [172](#)
- m_u16GuardThreshold
 - Mark3::Kernel, [149](#)
- m_u16Head
 - Mark3::Mailbox, [172](#)
 - Mark3::Streamer, [241](#)
- m_u16Iterations
 - Mark3::ProfileTimer, [212](#)
- m_u16MaxValue
 - Mark3::Semaphore, [234](#)
- m_u16Quantum
 - Mark3::Thread, [262](#)
- m_u16SetMask
 - Mark3::EventFlag, [140](#)
- m_u16Size
 - Mark3::Streamer, [241](#)
- m_u16StackSize
 - Mark3::Thread, [263](#)
- m_u16Tail
 - Mark3::Mailbox, [172](#)
 - Mark3::Streamer, [241](#)
- m_u16TicksRemain
 - Mark3::Quantum, [216](#)
- m_u16Value
 - Mark3::Semaphore, [234](#)
- m_u32Cumulative
 - Mark3::ProfileTimer, [213](#)
- m_u32CurrentIteration
 - Mark3::ProfileTimer, [213](#)
- m_u32Interval
 - Mark3::Timer, [279](#)
- m_u32NextWakeup
 - Mark3::TimerList, [283](#)
- m_u32StartTicks
 - Mark3::ProfileTimer, [213](#)
- m_u32Ticks
 - Mark3::Kernel, [149](#)
- m_u32TimeLeft
 - Mark3::Timer, [279](#)
- m_u8Flags
 - Mark3::Timer, [279](#)
- m_u8Initialized
 - Mark3::BlockingObject, [114](#)
 - Mark3::Timer, [279](#)
- m_u8ReadCount
 - Mark3::ReaderWriterLock, [221](#)
- m_u8Recurse
 - Mark3::Mutex, [198](#)
- m_u8ThreadID
 - Mark3::Thread, [263](#)
- m_u8Waiters
 - Mark3::ConditionVariable, [123](#)
- m_uBlockingInitCookie
 - Mark3::BlockingObject, [114](#)
- m_uBlockingInvalidCookie
 - Mark3::BlockingObject, [114](#)
- m_uMaxPri
 - Mark3::Mutex, [198](#)
- m_uNumPriorities
 - Mark3::Scheduler, [227](#)
- m_uPriority

- Mark3::CoList, [120](#)
- Mark3::Coroutine, [127](#)
- m_uTimerInitCookie
 - Mark3::Timer, [280](#)
- m_uTimerInvalidCookie
 - Mark3::Timer, [280](#)
- m_uXCurPriority
 - Mark3::Thread, [263](#)
- m_uXPrioMapBitMask
 - Mark3::PriorityMapL1, [205](#)
 - Mark3::PriorityMapL2, [209](#)
- m_uXPrioMapBits
 - Mark3::PriorityMapL1, [205](#)
 - Mark3::PriorityMapL2, [209](#)
- m_uXPrioMapNumWords
 - Mark3::PriorityMapL2, [209](#)
- m_uXPrioMapShiftLUT
 - Mark3::PriorityMapL1, [205](#)
 - Mark3::PriorityMapL2, [209](#)
- m_uXPrioMapWordShift
 - Mark3::PriorityMapL1, [205](#)
 - Mark3::PriorityMapL2, [210](#)
- m_uXPriority
 - Mark3::Thread, [263](#)
 - Mark3::ThreadList, [270](#)
- m_uXPriorityMap
 - Mark3::PriorityMapL1, [205](#)
- m_uXPriorityMapL2
 - Mark3::PriorityMapL2, [210](#)
- MAILBOX_SIZE
 - mark3c.h, [369](#)
- MESSAGE_SIZE
 - mark3c.h, [369](#)
- MESSAGEPOOL_SIZE
 - mark3c.h, [370](#)
- MESSAGEQUEUE_SIZE
 - mark3c.h, [370](#)
- MSecondsToTicks
 - Mark3::Timer, [276](#)
- MUTEX_SIZE
 - mark3c.h, [370](#)
- Mailbox_GetFreeSlots
 - mark3c.cpp, [317](#)
 - mark3c.h, [391](#)
- Mailbox_Init
 - mark3c.cpp, [317](#)
 - mark3c.h, [391](#)
- Mailbox_IsEmpty
 - mark3c.cpp, [318](#)
 - mark3c.h, [392](#)
- Mailbox_IsFull
 - mark3c.cpp, [318](#)
 - mark3c.h, [392](#)
- Mailbox_Receive
 - mark3c.cpp, [319](#)
 - mark3c.h, [393](#)
- Mailbox_ReceiveTail
 - mark3c.cpp, [319](#)
- mark3c.h, [393](#)
- Mailbox_Send
 - mark3c.cpp, [320](#)
 - mark3c.h, [394](#)
- Mailbox_SendTail
 - mark3c.cpp, [320](#)
 - mark3c.h, [394](#)
- Mailbox_TimedReceive
 - mark3c.cpp, [321](#)
 - mark3c.h, [395](#)
- Mailbox_TimedReceiveTail
 - mark3c.cpp, [321](#)
 - mark3c.h, [395](#)
- Mailbox_TimedSend
 - mark3c.cpp, [322](#)
 - mark3c.h, [396](#)
- Mailbox_TimedSendTail
 - mark3c.cpp, [322](#)
 - mark3c.h, [396](#)
- Mailbox_t
 - mark3c.h, [372](#)
- Mark3, [91](#)
 - AutoAllocAllocator_t, [95](#)
 - AutoAllocFree_t, [95](#)
 - AutoAllocType, [97](#)
 - Blocked, [98](#)
 - CoPrioMap, [95](#)
 - CoroutineHandler, [96](#)
 - EventFlagOperation, [98](#)
 - Exit, [98](#)
 - g_kwCriticalCount, [100](#)
 - g_kwSFR, [100](#)
 - ISR, [99](#)
 - IdleFunc, [96](#)
 - Invalid, [98](#)
 - KernelTimer_Task, [99](#)
 - Message, [98](#)
 - Mutex, [98](#)
 - naked, [100](#)
 - Notify, [98](#)
 - PORT_CLZ, [99](#)
 - PORT_CS_ENTER, [99](#)
 - PORT_CS_EXIT, [99](#)
 - PORT_CS_NESTING, [99](#)
 - PORT_IRQ_DISABLE, [100](#)
 - PORT_IRQ_ENABLE, [100](#)
 - PanicFunc, [96](#)
 - PriorityMap, [96](#)
 - Raw, [98](#)
 - Ready, [98](#)
 - SR_, [101](#)
 - Semaphore, [98](#)
 - Stop, [98](#)
 - Thread, [98](#)
 - Thread_Switch, [100](#)
 - ThreadContextCallout, [96](#)
 - ThreadCreateCallout, [96](#)
 - ThreadEntryFunc, [97](#)

- ThreadExitCallout, 97
- ThreadState, 98
- Timer, 98
- TimerCallback, 97
- uMaxTimerTicks, 101
- uTimerFlagActive, 101
- uTimerFlagCallback, 101
- uTimerFlagExpired, 101
- uTimerFlagOneShot, 102
- uTimerTicksInvalid, 102
- User, 98
- Mark3::Atomic, 102
 - Add, 103
 - Set, 103
 - Sub, 104
 - TestAndSet, 104
- Mark3::AutoAlloc, 105
 - Allocate, 106
 - DestroyObject, 106
 - DestroyRawData, 106
 - DestroyUserTypeAllocation, 107
 - Free, 107
 - Init, 107
 - m_pfAllocator, 109
 - m_pfFree, 109
 - NewObject, 107
 - NewRawData, 108
 - NewUserTypeAllocation, 108
 - SetAllocatorFunctions, 108
- Mark3::BlockingObject, 110
 - ~BlockingObject, 111
 - Block, 111
 - BlockPriority, 112
 - BlockingObject, 111
 - IsInitialized, 112
 - m_clBlockList, 114
 - m_u8Initialized, 114
 - m_uBlockingInitCookie, 114
 - m_uBlockingInvalidCookie, 114
 - SetInitialized, 112
 - UnBlock, 112
- Mark3::CircularLinkedList, 115
 - Add, 116
 - CircularLinkedList, 115
 - InsertNodeBefore, 116
 - operator new, 116
 - PivotBackward, 116
 - PivotForward, 117
 - Remove, 117
- Mark3::CoList, 117
 - Add, 118
 - m_pclPrioMap, 120
 - m_uPriority, 120
 - Remove, 119
 - SetPrioMap, 119
 - SetPriority, 119
- Mark3::CoScheduler, 127
 - GetCoList, 128
 - GetPrioMap, 129
 - GetStopList, 129
 - Init, 129
 - m_aclPriorities, 130
 - m_clPrioMap, 130
 - m_clStopList, 130
 - Schedule, 129
- Mark3::ConditionVariable, 120
 - Broadcast, 121
 - Init, 121
 - m_clMutex, 123
 - m_clSemaphore, 123
 - m_u8Waiters, 123
 - operator new, 122
 - Signal, 122
 - Wait, 122
- Mark3::Coroutine, 124
 - ~Coroutine, 125
 - Activate, 125
 - GetPriority, 125
 - Init, 125
 - m_bQueued, 126
 - m_pclOwner, 126
 - m_pfHandler, 127
 - m_pvContext, 127
 - m_uPriority, 127
 - Run, 126
 - SetPriority, 126
- Mark3::CriticalSection, 131
 - ~CriticalSection, 131
 - CriticalGuard, 131
- Mark3::CriticalSection, 132
 - Enter, 132
 - Exit, 132
 - NestingCount, 133
- Mark3::DoubleLinkedList, 133
 - Add, 134
 - DoubleLinkedList, 134
 - operator new, 134
 - Remove, 135
- Mark3::EventFlag, 135
 - ~EventFlag, 137
 - Clear, 137
 - GetMask, 137
 - Init, 137
 - m_u16SetMask, 140
 - operator new, 137
 - Set, 138
 - Wait, 138, 139
 - Wait_i, 139
 - WakeMe, 140
- Mark3::Kernel, 140
 - CompleteStart, 142
 - DebugPrint, 142
 - GetStackGuardThreshold, 143
 - GetThreadContextSwitchCallout, 143
 - GetThreadCreateCallout, 143
 - GetThreadExitCallout, 143

- GetTicks, [143](#)
- Init, [144](#)
- IsPanic, [144](#)
- IsStarted, [144](#)
- m_bIsPanic, [148](#)
- m_bIsStarted, [148](#)
- m_pfDebugPrintFunction, [148](#)
- m_pfPanic, [148](#)
- m_pfThreadContextCallout, [148](#)
- m_pfThreadCreateCallout, [149](#)
- m_pfThreadExitCallout, [149](#)
- m_u16GuardThreshold, [149](#)
- m_u32Ticks, [149](#)
- Panic, [144](#)
- SetDebugPrintFunction, [145](#)
- SetPanic, [145](#)
- SetStackGuardThreshold, [145](#)
- SetThreadContextSwitchCallout, [146](#)
- SetThreadCreateCallout, [146](#)
- SetThreadExitCallout, [147](#)
- Start, [147](#)
- Tick, [147](#)
- Mark3::KernelSWI, [150](#)
 - Config, [150](#)
 - Start, [150](#)
 - Trigger, [150](#)
- Mark3::KernelTimer, [151](#)
 - Config, [151](#)
 - Start, [151](#)
 - Stop, [152](#)
- Mark3::LinkList, [152](#)
 - GetHead, [153](#)
 - GetTail, [153](#)
 - Init, [153](#)
 - m_pclHead, [154](#)
 - m_pclTail, [155](#)
 - SetHead, [154](#)
 - SetTail, [154](#)
- Mark3::LinkListNode, [155](#)
 - CircularLinkList, [157](#)
 - ClearNode, [156](#)
 - DoubleLinkList, [157](#)
 - GetNext, [156](#)
 - GetPrev, [157](#)
 - LinkList, [157](#)
 - LinkListNode, [156](#)
 - next, [158](#)
 - prev, [158](#)
- Mark3::LockGuard, [158](#)
 - ~LockGuard, [159](#)
 - isAcquired, [160](#)
 - LockGuard, [159](#)
 - m_bIsAcquired, [160](#)
 - m_pclMutex, [160](#)
- Mark3::Mailbox, [161](#)
 - ~Mailbox, [163](#)
 - CopyData, [163](#)
 - GetFreeSlots, [163](#)
 - GetHeadPointer, [163](#)
 - GetTailPointer, [164](#)
 - Init, [164](#)
 - IsEmpty, [165](#)
 - IsFull, [165](#)
 - m_clRecvSem, [171](#)
 - m_clSendSem, [171](#)
 - m_pvBuffer, [171](#)
 - m_u16Count, [172](#)
 - m_u16ElementSize, [172](#)
 - m_u16Free, [172](#)
 - m_u16Head, [172](#)
 - m_u16Tail, [172](#)
 - MoveHeadBackward, [165](#)
 - MoveHeadForward, [165](#)
 - MoveTailBackward, [166](#)
 - MoveTailForward, [166](#)
 - operator new, [166](#)
 - Receive, [166](#), [167](#)
 - Receive_i, [167](#)
 - ReceiveTail, [168](#)
 - Send, [168](#), [169](#)
 - Send_i, [169](#)
 - SendTail, [170](#)
- Mark3::MemUtil, [173](#)
 - Checksum16, [174](#)
 - Checksum8, [174](#)
 - CompareMemory, [175](#)
 - CompareStrings, [175](#), [176](#)
 - CopyMemory, [176](#)
 - CopyString, [177](#)
 - DecimalToHex, [177](#), [178](#)
 - DecimalToString, [178](#), [179](#)
 - SetMemory, [179](#)
 - StringLength, [180](#)
 - StringSearch, [180](#)
 - StringToDecimal16, [180](#)
 - StringToDecimal32, [181](#)
 - StringToDecimal64, [181](#)
 - StringToDecimal8, [181](#)
 - Tokenize, [182](#)
- Mark3::Message, [182](#)
 - GetCode, [183](#)
 - GetData, [183](#)
 - Init, [184](#)
 - m_pvData, [185](#)
 - m_u16Code, [185](#)
 - operator new, [184](#)
 - SetCode, [184](#)
 - SetData, [185](#)
- Mark3::MessagePool, [186](#)
 - ~MessagePool, [186](#)
 - GetHead, [187](#)
 - Init, [187](#)
 - m_clList, [188](#)
 - operator new, [187](#)
 - Pop, [187](#)
 - Push, [188](#)

- Mark3::MessageQueue, 189
 - ~MessageQueue, 190
 - GetCount, 190
 - Init, 190
 - m_clLinkedList, 192
 - m_clSemaphore, 193
 - operator new, 190
 - Receive, 191
 - Receive_i, 191
 - Send, 192
- Mark3::Mutex, 193
 - ~Mutex, 194
 - Claim, 194, 195
 - Claim_i, 195
 - Init, 196
 - m_bReady, 197
 - m_bRecursive, 197
 - m_pclOwner, 198
 - m_u8Recurse, 198
 - m_uMaxPri, 198
 - operator new, 196
 - Release, 196
 - WakeMe, 197
 - WakeNext, 197
- Mark3::Notify, 199
 - ~Notify, 200
 - Init, 200
 - m_bPending, 202
 - operator new, 200
 - Signal, 200
 - Wait, 200, 201
 - WakeMe, 201
- Mark3::PriorityMapL1
 - Clear, 203
 - HighestPriority, 204
 - m_uXPrioMapBitMask, 205
 - m_uXPrioMapBits, 205
 - m_uXPrioMapShiftLUT, 205
 - m_uXPrioMapWordShift, 205
 - m_uXPriorityMap, 205
 - PrioBit, 204
 - PrioMapWordIndex, 204
 - PriorityFromBitmap, 204
 - PriorityMapL1, 203
 - Set, 204
- Mark3::PriorityMapL1 < T, C >, 202
- Mark3::PriorityMapL2
 - Clear, 207
 - HighestPriority, 208
 - m_auXPriorityMap, 209
 - m_uXPrioMapBitMask, 209
 - m_uXPrioMapBits, 209
 - m_uXPrioMapNumWords, 209
 - m_uXPrioMapShiftLUT, 209
 - m_uXPrioMapWordShift, 210
 - m_uXPriorityMapL2, 210
 - PrioBit, 208
 - PrioMapWordIndex, 208
 - PriorityFromBitmap, 208
 - PriorityMapL2, 207
 - Set, 208
- Mark3::PriorityMapL2 < T, C >, 206
- Mark3::ProfileTimer, 210
 - GetAverage, 211
 - GetCurrent, 211
 - Init, 211
 - m_bActive, 212
 - m_u16Iterations, 212
 - m_u32Cumulative, 213
 - m_u32CurrentIteration, 213
 - m_u32StartTicks, 213
 - Start, 212
 - Stop, 212
- Mark3::Quantum, 213
 - Cancel, 214
 - ClearInTimer, 214
 - Init, 214
 - m_bInTimer, 216
 - m_pclActiveThread, 216
 - m_pclTimerThread, 216
 - m_u16TicksRemain, 216
 - SetInTimer, 215
 - SetTimerThread, 215
 - Update, 215
- Mark3::ReaderWriterLock, 216
 - AcquireReader, 218
 - AcquireReader_i, 218
 - AcquireWriter, 219
 - AcquireWriter_i, 219
 - Init, 220
 - m_clGlobalMutex, 221
 - m_clReaderMutex, 221
 - m_u8ReadCount, 221
 - operator new, 220
 - ReleaseReader, 220
 - ReleaseWriter, 220
- Mark3::Scheduler, 221
 - Add, 223
 - GetCurrentThread, 223
 - GetNextThread, 223
 - GetStopList, 223
 - GetThreadList, 224
 - Init, 224
 - IsEnabled, 224
 - m_aclPriorities, 226
 - m_bEnabled, 226
 - m_bQueuedSchedule, 226
 - m_clPrioMap, 226
 - m_clStopList, 227
 - m_uNumPriorities, 227
 - QueueScheduler, 225
 - Remove, 225
 - Schedule, 225
 - SetScheduler, 225
- Mark3::SchedulerGuard, 227
 - ~SchedulerGuard, 228

- [m_bSchedState](#), [228](#)
 - [SchedulerGuard](#), [228](#)
- [Mark3::Semaphore](#), [229](#)
 - [~Semaphore](#), [230](#)
 - [GetCount](#), [230](#)
 - [Init](#), [230](#)
 - [m_u16MaxValue](#), [234](#)
 - [m_u16Value](#), [234](#)
 - [operator new](#), [231](#)
 - [Pend](#), [231](#)
 - [Pend_i](#), [231](#)
 - [Post](#), [233](#)
 - [WakeMe](#), [233](#)
 - [WakeNext](#), [233](#)
- [Mark3::Streamer](#), [234](#)
 - [CanRead](#), [236](#)
 - [CanWrite](#), [236](#)
 - [Claim](#), [236](#)
 - [GetAvailable](#), [237](#)
 - [Init](#), [237](#)
 - [IsEmpty](#), [237](#)
 - [Lock](#), [238](#)
 - [m_pau8Buffer](#), [240](#)
 - [m_pu8LockAddr](#), [240](#)
 - [m_u16Avail](#), [240](#)
 - [m_u16Head](#), [241](#)
 - [m_u16Size](#), [241](#)
 - [m_u16Tail](#), [241](#)
 - [Read](#), [238](#)
 - [Unlock](#), [239](#)
 - [Write](#), [239](#), [240](#)
- [Mark3::Thread](#), [242](#)
 - [~Thread](#), [245](#)
 - [ContextSwitchSWI](#), [246](#)
 - [CoopYield](#), [246](#)
 - [Exit](#), [246](#)
 - [GetCurPriority](#), [246](#)
 - [GetCurrent](#), [247](#)
 - [GetEventFlagMask](#), [247](#)
 - [GetEventFlagMode](#), [247](#)
 - [GetExpired](#), [248](#)
 - [GetExtendedContext](#), [248](#)
 - [GetID](#), [248](#)
 - [GetName](#), [248](#)
 - [GetOwner](#), [249](#)
 - [GetPriority](#), [249](#)
 - [GetQuantum](#), [249](#)
 - [GetStack](#), [250](#)
 - [GetStackSize](#), [250](#)
 - [GetStackSlack](#), [250](#)
 - [GetState](#), [251](#)
 - [GetTimer](#), [251](#)
 - [InheritPriority](#), [251](#)
 - [Init](#), [252](#)
 - [IsInitialized](#), [253](#)
 - [m_bExpired](#), [260](#)
 - [m_clTimer](#), [260](#)
 - [m_eFlagMode](#), [260](#)
 - [m_eState](#), [260](#)
 - [m_pclCurrent](#), [261](#)
 - [m_pclOwner](#), [261](#)
 - [m_pfEntryPoint](#), [261](#)
 - [m_pvArg](#), [261](#)
 - [m_pvExtendedContext](#), [261](#)
 - [m_pwStack](#), [262](#)
 - [m_pwStackTop](#), [262](#)
 - [m_szName](#), [262](#)
 - [m_u16FlagMask](#), [262](#)
 - [m_u16Quantum](#), [262](#)
 - [m_u16StackSize](#), [263](#)
 - [m_u8ThreadID](#), [263](#)
 - [m_uXCurPriority](#), [263](#)
 - [m_uXPriority](#), [263](#)
 - [operator new](#), [253](#)
 - [SetCurrent](#), [253](#)
 - [SetEventFlagMask](#), [254](#)
 - [SetEventFlagMode](#), [254](#)
 - [SetExpired](#), [254](#)
 - [SetExtendedContext](#), [255](#)
 - [SetID](#), [255](#)
 - [SetName](#), [255](#)
 - [SetOwner](#), [257](#)
 - [SetPriority](#), [257](#)
 - [SetPriorityBase](#), [257](#)
 - [SetQuantum](#), [258](#)
 - [SetState](#), [258](#)
 - [Sleep](#), [258](#)
 - [Start](#), [259](#)
 - [Stop](#), [259](#)
 - [Thread](#), [245](#)
 - [ThreadPort](#), [260](#)
 - [Yield](#), [259](#)
- [Mark3::ThreadList](#), [264](#)
 - [Add](#), [265](#)
 - [AddPriority](#), [267](#)
 - [HighestWaiter](#), [267](#)
 - [m_pclMap](#), [269](#)
 - [m_uXPriority](#), [270](#)
 - [operator new](#), [267](#)
 - [Remove](#), [267](#)
 - [SetMapPointer](#), [269](#)
 - [SetPriority](#), [269](#)
 - [ThreadList](#), [265](#)
- [Mark3::ThreadListList](#), [270](#)
 - [Add](#), [270](#)
 - [GetHead](#), [271](#)
 - [m_clThreadListList](#), [271](#)
 - [Remove](#), [271](#)
- [Mark3::ThreadPort](#), [272](#)
 - [Init](#), [272](#)
 - [InitStack](#), [272](#)
 - [StartThreads](#), [273](#)
 - [Thread](#), [273](#)
- [Mark3::Timer](#), [273](#)
 - [~Timer](#), [275](#)
 - [Init](#), [276](#)

- IsInitialized, 276
- m_pclOwner, 278
- m_pfCallback, 278
- m_pvData, 279
- m_u32Interval, 279
- m_u32TimeLeft, 279
- m_u8Flags, 279
- m_u8Initialized, 279
- m_uTimerInitCookie, 280
- m_uTimerInvalidCookie, 280
- MSecondsToTicks, 276
- operator new, 276
- SecondsToTicks, 276
- SetInitialized, 277
- Start, 277
- Stop, 278
- Timer, 275
- TimerList, 278
- USecondsToTicks, 278
- Mark3::TimerList, 280
 - Add, 281
 - Init, 281
 - m_bTimerActive, 282
 - m_clMutex, 282
 - m_u32NextWakeup, 283
 - Process, 282
 - Remove, 282
- Mark3::TimerScheduler, 283
 - Add, 284
 - Init, 284
 - m_clTimerList, 285
 - Process, 284
 - Remove, 284
- Mark3::Token_t, 285
 - pcToken, 286
 - u8Len, 286
- Mark3::TypedCircularLinkedList
 - Add, 287
 - GetHead, 288
 - GetTail, 288
 - InsertNodeBefore, 288
 - operator new, 289
 - Remove, 289
 - SetHead, 289
 - SetTail, 290
 - TypedCircularLinkedList, 287
- Mark3::TypedCircularLinkedList< T >, 286
- Mark3::TypedDoubleLinkedList
 - Add, 291
 - GetHead, 292
 - GetTail, 292
 - operator new, 292
 - Remove, 292
 - SetHead, 293
 - SetTail, 293
 - TypedDoubleLinkedList, 291
- Mark3::TypedDoubleLinkedList< T >, 290
- Mark3::TypedLinkedListNode
 - GetNext, 294
 - GetPrev, 294
- Mark3::TypedLinkedListNode< T >, 293
- mark3c.cpp
 - Alloc_ConditionVariable, 300
 - Alloc_EventFlag, 300
 - Alloc_Mailbox, 300
 - Alloc_Memory, 300
 - Alloc_Message, 300
 - Alloc_MessagePool, 301
 - Alloc_MessageQueue, 301
 - Alloc_Mutex, 301
 - Alloc_Notify, 302
 - Alloc_ReaderWriterLock, 302
 - Alloc_Semaphore, 302
 - Alloc_Thread, 303
 - Alloc_Timer, 303
 - Atomic_Add16, 303
 - Atomic_Add32, 304
 - Atomic_Add8, 304
 - Atomic_Set16, 305
 - Atomic_Set32, 305
 - Atomic_Set8, 306
 - Atomic_Sub16, 306
 - Atomic_Sub32, 307
 - Atomic_Sub8, 307
 - Atomic_TestAndSet, 308
 - ConditionVariable_Broadcast, 308
 - ConditionVariable_Init, 309
 - ConditionVariable_Signal, 309
 - ConditionVariable_TimedWait, 310
 - ConditionVariable_Wait, 310
 - Free_ConditionVariable, 311
 - Free_EventFlag, 311
 - Free_Mailbox, 311
 - Free_Memory, 311
 - Free_Message, 312
 - Free_MessagePool, 312
 - Free_MessageQueue, 312
 - Free_Mutex, 312
 - Free_Notify, 312
 - Free_ReaderWriterLock, 312
 - Free_Semaphore, 313
 - Free_Thread, 313
 - Free_Timer, 313
 - Kernel_DebugPrint, 313
 - Kernel_GetStackGuardThreshold, 314
 - Kernel_GetTicks, 314
 - Kernel_Init, 314
 - Kernel_IsPanic, 314
 - Kernel_IsStarted, 315
 - Kernel_Panic, 315
 - Kernel_SetDebugPrintFunction, 316
 - Kernel_SetPanic, 316
 - Kernel_SetStackGuardThreshold, 316
 - Kernel_Start, 317
 - Mailbox_GetFreeSlots, 317
 - Mailbox_Init, 317

- Mailbox_IsEmpty, [318](#)
- Mailbox_IsFull, [318](#)
- Mailbox_Receive, [319](#)
- Mailbox_ReceiveTail, [319](#)
- Mailbox_Send, [320](#)
- Mailbox_SendTail, [320](#)
- Mailbox_TimedReceive, [321](#)
- Mailbox_TimedReceiveTail, [321](#)
- Mailbox_TimedSend, [322](#)
- Mailbox_TimedSendTail, [322](#)
- Message_GetCode, [323](#)
- Message_GetData, [323](#)
- Message_Init, [324](#)
- Message_SetCode, [324](#)
- Message_SetData, [325](#)
- MessagePool_Init, [325](#)
- MessagePool_Pop, [326](#)
- MessagePool_Push, [326](#)
- MessageQueue_GetCount, [327](#)
- MessageQueue_Init, [327](#)
- MessageQueue_Receive, [327](#)
- MessageQueue_Send, [328](#)
- MessageQueue_TimedReceive, [328](#)
- Mutex_Claim, [329](#)
- Mutex_Init, [329](#)
- Mutex_Release, [330](#)
- Mutex_TimedClaim, [330](#)
- Notify_Init, [330](#)
- Notify_Signal, [331](#)
- Notify_TimedWait, [331](#)
- Notify_Wait, [332](#)
- ReaderWriterLock_AcquireReader, [332](#)
- ReaderWriterLock_AcquireWriter, [333](#)
- ReaderWriterLock_Init, [333](#)
- ReaderWriterLock_ReleaseReader, [333](#)
- ReaderWriterLock_ReleaseWriter, [334](#)
- ReaderWriterLock_TimedAcquireReader, [334](#)
- ReaderWriterLock_TimedAcquireWriter, [335](#)
- Scheduler_Enable, [335](#)
- Scheduler_GetCurrentThread, [336](#)
- Scheduler_IsEnabled, [336](#)
- Semaphore_Init, [336](#)
- Semaphore_Pend, [337](#)
- Semaphore_Post, [337](#)
- Semaphore_TimedPend, [337](#)
- Thread_CoopYield, [338](#)
- Thread_Exit, [338](#)
- Thread_GetCurPriority, [339](#)
- Thread_GetID, [339](#)
- Thread_GetName, [340](#)
- Thread_GetPriority, [340](#)
- Thread_GetQuantum, [340](#)
- Thread_GetStackSlack, [340](#)
- Thread_GetState, [341](#)
- Thread_Init, [341](#)
- Thread_SetID, [342](#)
- Thread_SetName, [342](#)
- Thread_SetPriority, [342](#)
- Thread_SetQuantum, [343](#)
- Thread_Sleep, [343](#)
- Thread_Start, [344](#)
- Thread_Stop, [344](#)
- Thread_Yield, [344](#)
- Timer_Init, [345](#)
- Timer_Restart, [345](#)
- Timer_Start, [345](#)
- Timer_Stop, [346](#)
- mark3c.h
 - Alloc_Mailbox, [375](#)
 - Alloc_Memory, [375](#)
 - Alloc_Message, [376](#)
 - Alloc_MessagePool, [376](#)
 - Alloc_MessageQueue, [376](#)
 - Alloc_Mutex, [376](#)
 - Alloc_Notify, [377](#)
 - Alloc_Semaphore, [377](#)
 - Alloc_Thread, [377](#)
 - Alloc_Timer, [378](#)
 - Atomic_Add16, [378](#)
 - Atomic_Add32, [379](#)
 - Atomic_Add8, [379](#)
 - Atomic_Set16, [380](#)
 - Atomic_Set32, [380](#)
 - Atomic_Set8, [381](#)
 - Atomic_Sub16, [381](#)
 - Atomic_Sub32, [382](#)
 - Atomic_Sub8, [382](#)
 - Atomic_TestAndSet, [383](#)
 - CONDITIONVARIABLE_SIZE, [366](#)
 - ConditionVariable_Broadcast, [383](#)
 - ConditionVariable_Init, [384](#)
 - ConditionVariable_Signal, [384](#)
 - ConditionVariable_TimedWait, [384](#)
 - ConditionVariable_Wait, [385](#)
 - ConditionVariable_t, [372](#)
 - DECLARE_CONDITIONVARIABLE, [366](#)
 - DECLARE_MAILBOX, [366](#)
 - DECLARE_MESSAGEPOOL, [367](#)
 - DECLARE_MESSAGEQUEUE, [367](#)
 - DECLARE_MESSAGE, [366](#)
 - DECLARE_MUTEX, [367](#)
 - DECLARE_NOTIFY, [368](#)
 - DECLARE_READERWRITERLOCK, [368](#)
 - DECLARE_SEMAPHORE, [368](#)
 - DECLARE_THREAD, [369](#)
 - DECLARE_TIMER, [369](#)
 - Free_Mailbox, [385](#)
 - Free_Memory, [386](#)
 - Free_Message, [386](#)
 - Free_MessagePool, [386](#)
 - Free_MessageQueue, [386](#)
 - Free_Mutex, [386](#)
 - Free_Notify, [387](#)
 - Free_Semaphore, [387](#)
 - Free_Thread, [387](#)
 - Free_Timer, [387](#)

- Kernel_DebugPrint, 387
- Kernel_GetTicks, 388
- Kernel_Init, 388
- Kernel_IsPanic, 388
- Kernel_IsStarted, 389
- Kernel_Panic, 389
- Kernel_SetDebugPrintFunction, 390
- Kernel_SetPanic, 390
- Kernel_Start, 390
- kernel_debug_print_t, 372
- MAILBOX_SIZE, 369
- MESSAGE_SIZE, 369
- MESSAGEPOOL_SIZE, 370
- MESSAGEQUEUE_SIZE, 370
- MUTEX_SIZE, 370
- Mailbox_GetFreeSlots, 391
- Mailbox_Init, 391
- Mailbox_IsEmpty, 392
- Mailbox_IsFull, 392
- Mailbox_Receive, 393
- Mailbox_ReceiveTail, 393
- Mailbox_Send, 394
- Mailbox_SendTail, 394
- Mailbox_TimedReceive, 395
- Mailbox_TimedReceiveTail, 395
- Mailbox_TimedSend, 396
- Mailbox_TimedSendTail, 396
- Mailbox_t, 372
- Message_GetCode, 397
- Message_GetData, 397
- Message_Init, 398
- Message_SetCode, 398
- Message_SetData, 399
- Message_t, 372
- MessagePool_Init, 399
- MessagePool_Pop, 399
- MessagePool_Push, 400
- MessagePool_t, 372
- MessageQueue_GetCount, 400
- MessageQueue_Init, 401
- MessageQueue_Receive, 401
- MessageQueue_Send, 401
- MessageQueue_TimedReceive, 402
- MessageQueue_t, 372
- Mutex_Claim, 402
- Mutex_Init, 403
- Mutex_Release, 403
- Mutex_TimedClaim, 404
- Mutex_t, 373
- NOTIFY_SIZE, 370
- Notify_Init, 404
- Notify_Signal, 405
- Notify_TimedWait, 405
- Notify_Wait, 405
- Notify_t, 373
- panic_func_t, 373
- READERWRITERLOCK_SIZE, 370
- ReaderWriterLock_AcquireReader, 406
- ReaderWriterLock_AcquireWriter, 406
- ReaderWriterLock_Init, 407
- ReaderWriterLock_ReleaseReader, 407
- ReaderWriterLock_ReleaseWriter, 407
- ReaderWriterLock_TimedAcquireReader, 408
- ReaderWriterLock_TimedAcquireWriter, 408
- ReaderWriterLock_t, 373
- SEMAPHORE_SIZE, 370
- Scheduler_Enable, 409
- Scheduler_GetCurrentThread, 409
- Scheduler_IsEnabled, 410
- Semaphore_Init, 410
- Semaphore_Pend, 410
- Semaphore_Post, 411
- Semaphore_TimedPend, 411
- Semaphore_t, 373
- THREAD_SIZE, 371
- TIMER_SIZE, 371
- TOKEN_1, 371
- TOKEN_2, 371
- Thread_CoopYield, 412
- Thread_Exit, 412
- Thread_GetCurPriority, 412
- Thread_GetID, 413
- Thread_GetPriority, 413
- Thread_GetState, 414
- Thread_Init, 414
- Thread_SetID, 415
- Thread_SetPriority, 415
- Thread_Sleep, 416
- Thread_Start, 416
- Thread_Stop, 416
- Thread_Yield, 417
- thread_entry_func_t, 374
- thread_state_t, 374
- Thread_t, 374
- Timer_Init, 417
- Timer_Restart, 418
- Timer_Start, 418
- Timer_Stop, 419
- timer_callback_t, 374
- Timer_t, 374
- WORD_ROUND, 371
- mark3cfg.h
 - KERNEL_CONTEXT_SWITCH_CALLOUT, 519
 - KERNEL_DEBUG, 519
 - KERNEL_EVENT_FLAGS, 519
 - KERNEL_EXTENDED_CONTEXT, 520
 - KERNEL_NAMED_THREADS, 520
 - KERNEL_ROUND_ROBIN, 520
 - KERNEL_STACK_CHECK, 520
 - KERNEL_THREAD_CREATE_CALLOUT, 521
 - KERNEL_THREAD_EXIT_CALLOUT, 521
- Message
 - Mark3, 98
- Message_GetCode
 - mark3c.cpp, 323
 - mark3c.h, 397

- Message_GetData
 - mark3c.cpp, [323](#)
 - mark3c.h, [397](#)
- Message_Init
 - mark3c.cpp, [324](#)
 - mark3c.h, [398](#)
- Message_SetCode
 - mark3c.cpp, [324](#)
 - mark3c.h, [398](#)
- Message_SetData
 - mark3c.cpp, [325](#)
 - mark3c.h, [399](#)
- Message_t
 - mark3c.h, [372](#)
- MessagePool_Init
 - mark3c.cpp, [325](#)
 - mark3c.h, [399](#)
- MessagePool_Pop
 - mark3c.cpp, [326](#)
 - mark3c.h, [399](#)
- MessagePool_Push
 - mark3c.cpp, [326](#)
 - mark3c.h, [400](#)
- MessagePool_t
 - mark3c.h, [372](#)
- MessageQueue_GetCount
 - mark3c.cpp, [327](#)
 - mark3c.h, [400](#)
- MessageQueue_Init
 - mark3c.cpp, [327](#)
 - mark3c.h, [401](#)
- MessageQueue_Receive
 - mark3c.cpp, [327](#)
 - mark3c.h, [401](#)
- MessageQueue_Send
 - mark3c.cpp, [328](#)
 - mark3c.h, [401](#)
- MessageQueue_TimedReceive
 - mark3c.cpp, [328](#)
 - mark3c.h, [402](#)
- MessageQueue_t
 - mark3c.h, [372](#)
- MoveHeadBackward
 - Mark3::Mailbox, [165](#)
- MoveHeadForward
 - Mark3::Mailbox, [165](#)
- MoveTailBackward
 - Mark3::Mailbox, [166](#)
- MoveTailForward
 - Mark3::Mailbox, [166](#)
- Mutex
 - Mark3, [98](#)
- Mutex_Claim
 - mark3c.cpp, [329](#)
 - mark3c.h, [402](#)
- Mutex_Init
 - mark3c.cpp, [329](#)
 - mark3c.h, [403](#)
- Mutex_Release
 - mark3c.cpp, [330](#)
 - mark3c.h, [403](#)
- Mutex_TimedClaim
 - mark3c.cpp, [330](#)
 - mark3c.h, [404](#)
- Mutex_t
 - mark3c.h, [373](#)
- NOTIFY_SIZE
 - mark3c.h, [370](#)
- naked
 - Mark3, [100](#)
- NestingCount
 - Mark3::CriticalSection, [133](#)
- NewObject
 - Mark3::AutoAlloc, [107](#)
- NewRawData
 - Mark3::AutoAlloc, [108](#)
- NewUserTypeAllocation
 - Mark3::AutoAlloc, [108](#)
- next
 - Mark3::LinkListNode, [158](#)
- Notify
 - Mark3, [98](#)
- Notify_Init
 - mark3c.cpp, [330](#)
 - mark3c.h, [404](#)
- Notify_Signal
 - mark3c.cpp, [331](#)
 - mark3c.h, [405](#)
- Notify_TimedWait
 - mark3c.cpp, [331](#)
 - mark3c.h, [405](#)
- Notify_Wait
 - mark3c.cpp, [332](#)
 - mark3c.h, [405](#)
- Notify_t
 - mark3c.h, [373](#)
- operator new
 - Mark3::CircularLinkList, [116](#)
 - Mark3::ConditionVariable, [122](#)
 - Mark3::DoubleLinkList, [134](#)
 - Mark3::EventFlag, [137](#)
 - Mark3::Mailbox, [166](#)
 - Mark3::Message, [184](#)
 - Mark3::MessagePool, [187](#)
 - Mark3::MessageQueue, [190](#)
 - Mark3::Mutex, [196](#)
 - Mark3::Notify, [200](#)
 - Mark3::ReaderWriterLock, [220](#)
 - Mark3::Semaphore, [231](#)
 - Mark3::Thread, [253](#)
 - Mark3::ThreadList, [267](#)
 - Mark3::Timer, [276](#)
 - Mark3::TypedCircularLinkList, [289](#)
 - Mark3::TypedDoubleLinkList, [292](#)

- PANIC_ACTIVE_COROUTINE_DESCOPED
paniccodes.h, 528
- PANIC_ACTIVE_EVENTFLAG_DESCOPED
paniccodes.h, 528
- PANIC_ACTIVE_MAILBOX_DESCOPED
paniccodes.h, 528
- PANIC_ACTIVE_MUTEX_DESCOPED
paniccodes.h, 528
- PANIC_ACTIVE_NOTIFY_DESCOPED
paniccodes.h, 528
- PANIC_ACTIVE_SEMAPHORE_DESCOPED
paniccodes.h, 528
- PANIC_ACTIVE_TIMER_DESCOPED
paniccodes.h, 529
- PANIC_ASSERT_FAILED
paniccodes.h, 529
- PANIC_AUTO_HEAP_EXHAUSTED
paniccodes.h, 529
- PANIC_LIST_UNLINK_FAILED
paniccodes.h, 529
- PANIC_NO_READY_THREADS
paniccodes.h, 529
- PANIC_POWERMAN_EXHAUSTED
paniccodes.h, 529
- PANIC_RUNNING_THREAD_DESCOPED
paniccodes.h, 530
- PANIC_STACK_SLACK_VIOLATED
paniccodes.h, 530
- PORT_CLZ
Mark3, 99
- PORT_COROUTINE_PRIORITIES
portcfg.h, 444
- PORT_CS_ENTER
Mark3, 99
- PORT_CS_EXIT
Mark3, 99
- PORT_CS_NESTING
Mark3, 99
- PORT_IRQ_DISABLE
Mark3, 100
- PORT_IRQ_ENABLE
Mark3, 100
- PORT_KERNEL_DEFAULT_STACK_SIZE
portcfg.h, 444
- PORT_KERNEL_TIMERS_THREAD_STACK
portcfg.h, 444
- PORT_MIN_TIMER_TICKS
portcfg.h, 445
- PORT_OVERLOAD_NEW
portcfg.h, 445
- PORT_PRIO_MAP_WORD_SIZE
portcfg.h, 445
- PORT_PRIO_TYPE
portcfg.h, 445
- PORT_PUSH_TO_STACK
threadport.h, 449
- PORT_STACK_GROWS_DOWN
portcfg.h, 445
- PORT_SYSTEM_FREQ
portcfg.h, 446
- PORT_TIMER_COUNT_TYPE
portcfg.h, 446
- PORT_TIMER_FREQ
portcfg.h, 446
- PORT_TOP_OF_STACK
threadport.h, 449
- PORT_USE_HW_CLZ
portcfg.h, 446
- Panic
Mark3::Kernel, 144
- panic_func_t
mark3c.h, 373
- PanicFunc
Mark3, 96
- paniccodes.h
PANIC_ACTIVE_COROUTINE_DESCOPED, 528
PANIC_ACTIVE_EVENTFLAG_DESCOPED, 528
PANIC_ACTIVE_MAILBOX_DESCOPED, 528
PANIC_ACTIVE_MUTEX_DESCOPED, 528
PANIC_ACTIVE_NOTIFY_DESCOPED, 528
PANIC_ACTIVE_SEMAPHORE_DESCOPED, 528
PANIC_ACTIVE_TIMER_DESCOPED, 529
PANIC_ASSERT_FAILED, 529
PANIC_AUTO_HEAP_EXHAUSTED, 529
PANIC_LIST_UNLINK_FAILED, 529
PANIC_NO_READY_THREADS, 529
PANIC_POWERMAN_EXHAUSTED, 529
PANIC_RUNNING_THREAD_DESCOPED, 530
PANIC_STACK_SLACK_VIOLATED, 530
- pcToken
Mark3::Token_t, 286
- Pend
Mark3::Semaphore, 231
- Pend_i
Mark3::Semaphore, 231
- PivotBackward
Mark3::CircularLinkList, 116
- PivotForward
Mark3::CircularLinkList, 117
- Pop
Mark3::MessagePool, 187
- portcfg.h
AVR, 443
K_ADDR, 443
K_INT, 443
K_WORD, 443
KERNEL_NUM_PRIORITIES, 443
KERNEL_STACK_GUARD_DEFAULT, 444
KERNEL_TIMERS_THREAD_PRIORITY, 444
PORT_COROUTINE_PRIORITIES, 444
PORT_KERNEL_DEFAULT_STACK_SIZE, 444
PORT_KERNEL_TIMERS_THREAD_STACK, 444
PORT_MIN_TIMER_TICKS, 445
PORT_OVERLOAD_NEW, 445
PORT_PRIO_MAP_WORD_SIZE, 445
PORT_PRIO_TYPE, 445

- PORT_STACK_GROWS_DOWN, [445](#)
- PORT_SYSTEM_FREQ, [446](#)
- PORT_TIMER_COUNT_TYPE, [446](#)
- PORT_TIMER_FREQ, [446](#)
- PORT_USE_HW_CLZ, [446](#)
- THREAD_QUANTUM_DEFAULT, [447](#)
- Post
 - Mark3::Semaphore, [233](#)
- prev
 - Mark3::LinkListNode, [158](#)
- PrioBit
 - Mark3::PriorityMapL1, [204](#)
 - Mark3::PriorityMapL2, [208](#)
- PrioMapWordIndex
 - Mark3::PriorityMapL1, [204](#)
 - Mark3::PriorityMapL2, [208](#)
- PriorityFromBitmap
 - Mark3::PriorityMapL1, [204](#)
 - Mark3::PriorityMapL2, [208](#)
- PriorityMap
 - Mark3, [96](#)
- PriorityMapL1
 - Mark3::PriorityMapL1, [203](#)
- PriorityMapL2
 - Mark3::PriorityMapL2, [207](#)
- Process
 - Mark3::TimerList, [282](#)
 - Mark3::TimerScheduler, [284](#)
- Push
 - Mark3::MessagePool, [188](#)
- QueueScheduler
 - Mark3::Scheduler, [225](#)
- READERWRITERLOCK_SIZE
 - mark3c.h, [370](#)
- Raw
 - Mark3, [98](#)
- Read
 - Mark3::Streamer, [238](#)
- ReaderWriterLock_AcquireReader
 - mark3c.cpp, [332](#)
 - mark3c.h, [406](#)
- ReaderWriterLock_AcquireWriter
 - mark3c.cpp, [333](#)
 - mark3c.h, [406](#)
- ReaderWriterLock_Init
 - mark3c.cpp, [333](#)
 - mark3c.h, [407](#)
- ReaderWriterLock_ReleaseReader
 - mark3c.cpp, [333](#)
 - mark3c.h, [407](#)
- ReaderWriterLock_ReleaseWriter
 - mark3c.cpp, [334](#)
 - mark3c.h, [407](#)
- ReaderWriterLock_TimedAcquireReader
 - mark3c.cpp, [334](#)
 - mark3c.h, [408](#)
- ReaderWriterLock_TimedAcquireWriter
 - mark3c.cpp, [335](#)
 - mark3c.h, [408](#)
- ReaderWriterLock_t
 - mark3c.h, [373](#)
- Ready
 - Mark3, [98](#)
- Receive
 - Mark3::Mailbox, [166](#), [167](#)
 - Mark3::MessageQueue, [191](#)
- Receive_i
 - Mark3::Mailbox, [167](#)
 - Mark3::MessageQueue, [191](#)
- ReceiveTail
 - Mark3::Mailbox, [168](#)
- Release
 - Mark3::Mutex, [196](#)
- ReleaseReader
 - Mark3::ReaderWriterLock, [220](#)
- ReleaseWriter
 - Mark3::ReaderWriterLock, [220](#)
- Remove
 - Mark3::CircularLinkList, [117](#)
 - Mark3::CoList, [119](#)
 - Mark3::DoubleLinkList, [135](#)
 - Mark3::Scheduler, [225](#)
 - Mark3::ThreadList, [267](#)
 - Mark3::ThreadListList, [271](#)
 - Mark3::TimerList, [282](#)
 - Mark3::TimerScheduler, [284](#)
 - Mark3::TypedCircularLinkList, [289](#)
 - Mark3::TypedDoubleLinkList, [292](#)
- Run
 - Mark3::Coroutine, [126](#)
- SEMAPHORE_SIZE
 - mark3c.h, [370](#)
- SR_
 - Mark3, [101](#)
- Schedule
 - Mark3::CoScheduler, [129](#)
 - Mark3::Scheduler, [225](#)
- scheduler.cpp
 - g_pclCurrent, [556](#)
 - g_pclNext, [556](#)
- scheduler.h
 - g_pclCurrent, [540](#)
 - g_pclNext, [540](#)
- Scheduler_Enable
 - mark3c.cpp, [335](#)
 - mark3c.h, [409](#)
- Scheduler_GetCurrentThread
 - mark3c.cpp, [336](#)
 - mark3c.h, [409](#)
- Scheduler_IsEnabled
 - mark3c.cpp, [336](#)
 - mark3c.h, [410](#)
- SchedulerGuard
 - Mark3::SchedulerGuard, [228](#)
- SecondsToTicks

- Mark3::Timer, 276
- Semaphore
 - Mark3, 98
- Semaphore_Init
 - mark3c.cpp, 336
 - mark3c.h, 410
- Semaphore_Pend
 - mark3c.cpp, 337
 - mark3c.h, 410
- Semaphore_Post
 - mark3c.cpp, 337
 - mark3c.h, 411
- Semaphore_TimedPend
 - mark3c.cpp, 337
 - mark3c.h, 411
- Semaphore_t
 - mark3c.h, 373
- Send
 - Mark3::Mailbox, 168, 169
 - Mark3::MessageQueue, 192
- Send_i
 - Mark3::Mailbox, 169
- SendTail
 - Mark3::Mailbox, 170
- Set
 - Mark3::Atomic, 103
 - Mark3::EventFlag, 138
 - Mark3::PriorityMapL1, 204
 - Mark3::PriorityMapL2, 208
- SetAllocatorFunctions
 - Mark3::AutoAlloc, 108
- SetCode
 - Mark3::Message, 184
- SetCurrent
 - Mark3::Thread, 253
- SetData
 - Mark3::Message, 185
- SetDebugPrintFunction
 - Mark3::Kernel, 145
- SetEventFlagMask
 - Mark3::Thread, 254
- SetEventFlagMode
 - Mark3::Thread, 254
- SetExpired
 - Mark3::Thread, 254
- SetExtendedContext
 - Mark3::Thread, 255
- SetHead
 - Mark3::LinkList, 154
 - Mark3::TypedCircularLinkList, 289
 - Mark3::TypedDoubleLinkList, 293
- SetID
 - Mark3::Thread, 255
- SetInTimer
 - Mark3::Quantum, 215
- SetInitialized
 - Mark3::BlockingObject, 112
 - Mark3::Timer, 277
- SetMapPointer
 - Mark3::ThreadList, 269
- SetMemory
 - Mark3::MemUtil, 179
- SetName
 - Mark3::Thread, 255
- SetOwner
 - Mark3::Thread, 257
- SetPanic
 - Mark3::Kernel, 145
- SetPrioMap
 - Mark3::CoList, 119
- SetPriority
 - Mark3::CoList, 119
 - Mark3::Coroutine, 126
 - Mark3::Thread, 257
 - Mark3::ThreadList, 269
- SetPriorityBase
 - Mark3::Thread, 257
- SetQuantum
 - Mark3::Thread, 258
- SetScheduler
 - Mark3::Scheduler, 225
- SetStackGuardThreshold
 - Mark3::Kernel, 145
- SetState
 - Mark3::Thread, 258
- SetTail
 - Mark3::LinkList, 154
 - Mark3::TypedCircularLinkList, 290
 - Mark3::TypedDoubleLinkList, 293
- SetThreadContextSwitchCallout
 - Mark3::Kernel, 146
- SetThreadCreateCallout
 - Mark3::Kernel, 146
- SetThreadExitCallout
 - Mark3::Kernel, 147
- SetTimerThread
 - Mark3::Quantum, 215
- Signal
 - Mark3::ConditionVariable, 122
 - Mark3::Notify, 200
- Sleep
 - Mark3::Thread, 258
- Start
 - Mark3::Kernel, 147
 - Mark3::KernelSWI, 150
 - Mark3::KernelTimer, 151
 - Mark3::ProfileTimer, 212
 - Mark3::Thread, 259
 - Mark3::Timer, 277
- StartThreads
 - Mark3::ThreadPort, 273
- Stop
 - Mark3, 98
 - Mark3::KernelTimer, 152
 - Mark3::ProfileTimer, 212
 - Mark3::Thread, 259

- Mark3::Timer, [278](#)
- StringLength
 - Mark3::MemUtil, [180](#)
- StringSearch
 - Mark3::MemUtil, [180](#)
- StringToDecimal16
 - Mark3::MemUtil, [180](#)
- StringToDecimal32
 - Mark3::MemUtil, [181](#)
- StringToDecimal64
 - Mark3::MemUtil, [181](#)
- StringToDecimal8
 - Mark3::MemUtil, [181](#)
- Sub
 - Mark3::Atomic, [104](#)
- TCCR1B_INIT
 - kerneltimer.cpp, [440](#)
- THREAD_QUANTUM_DEFAULT
 - portcfg.h, [447](#)
- THREAD_SIZE
 - mark3c.h, [371](#)
- TIMER_IFR
 - kerneltimer.cpp, [440](#)
- TIMER_IMSK
 - kerneltimer.cpp, [440](#)
- TIMER_SIZE
 - mark3c.h, [371](#)
- TOKEN_1
 - mark3c.h, [371](#)
- TOKEN_2
 - mark3c.h, [371](#)
- TestAndSet
 - Mark3::Atomic, [104](#)
- Thread
 - Mark3, [98](#)
 - Mark3::Thread, [245](#)
 - Mark3::ThreadPort, [273](#)
- Thread_CoopYield
 - mark3c.cpp, [338](#)
 - mark3c.h, [412](#)
- Thread_Exit
 - mark3c.cpp, [338](#)
 - mark3c.h, [412](#)
- Thread_GetCurPriority
 - mark3c.cpp, [339](#)
 - mark3c.h, [412](#)
- Thread_GetID
 - mark3c.cpp, [339](#)
 - mark3c.h, [413](#)
- Thread_GetName
 - mark3c.cpp, [340](#)
- Thread_GetPriority
 - mark3c.cpp, [340](#)
 - mark3c.h, [413](#)
- Thread_GetQuantum
 - mark3c.cpp, [340](#)
- Thread_GetStackSlack
 - mark3c.cpp, [340](#)
- Thread_GetState
 - mark3c.cpp, [341](#)
 - mark3c.h, [414](#)
- Thread_Init
 - mark3c.cpp, [341](#)
 - mark3c.h, [414](#)
- Thread_RestoreContext
 - threadport.h, [449](#)
- Thread_SaveContext
 - threadport.h, [449](#)
- Thread_SetID
 - mark3c.cpp, [342](#)
 - mark3c.h, [415](#)
- Thread_SetName
 - mark3c.cpp, [342](#)
- Thread_SetPriority
 - mark3c.cpp, [342](#)
 - mark3c.h, [415](#)
- Thread_SetQuantum
 - mark3c.cpp, [343](#)
- Thread_Sleep
 - mark3c.cpp, [343](#)
 - mark3c.h, [416](#)
- Thread_Start
 - mark3c.cpp, [344](#)
 - mark3c.h, [416](#)
- Thread_Stop
 - mark3c.cpp, [344](#)
 - mark3c.h, [416](#)
- Thread_Switch
 - Mark3, [100](#)
- Thread_Yield
 - mark3c.cpp, [344](#)
 - mark3c.h, [417](#)
- thread_entry_func_t
 - mark3c.h, [374](#)
- thread_state_t
 - mark3c.h, [374](#)
- Thread_t
 - mark3c.h, [374](#)
- ThreadContextCallout
 - Mark3, [96](#)
- ThreadCreateCallout
 - Mark3, [96](#)
- ThreadEntryFunc
 - Mark3, [97](#)
- ThreadExitCallout
 - Mark3, [97](#)
- ThreadList
 - Mark3::ThreadList, [265](#)
- ThreadPort
 - Mark3::Thread, [260](#)
- ThreadState
 - Mark3, [98](#)
- threadport.cpp
 - g_kwCriticalCount, [453](#)
 - g_kwSFR, [453](#)
- threadport.h

- ASM, [449](#)
 - PORT_PUSH_TO_STACK, [449](#)
 - PORT_TOP_OF_STACK, [449](#)
 - Thread_RestoreContext, [449](#)
 - Thread_SaveContext, [449](#)
- Tick
 - Mark3::Kernel, [147](#)
- Timer
 - Mark3, [98](#)
 - Mark3::Timer, [275](#)
- Timer_Init
 - mark3c.cpp, [345](#)
 - mark3c.h, [417](#)
- Timer_Restart
 - mark3c.cpp, [345](#)
 - mark3c.h, [418](#)
- Timer_Start
 - mark3c.cpp, [345](#)
 - mark3c.h, [418](#)
- Timer_Stop
 - mark3c.cpp, [346](#)
 - mark3c.h, [419](#)
- timer_callback_t
 - mark3c.h, [374](#)
- Timer_t
 - mark3c.h, [374](#)
- TimerCallback
 - Mark3, [97](#)
- TimerList
 - Mark3::Timer, [278](#)
- Tokenize
 - Mark3::MemUtil, [182](#)
- Trigger
 - Mark3::KernelSWI, [150](#)
- TypedCircularLinkList
 - Mark3::TypedCircularLinkList, [287](#)
- TypedDoubleLinkList
 - Mark3::TypedDoubleLinkList, [291](#)
- u8Len
 - Mark3::Token_t, [286](#)
- uMaxTimerTicks
 - Mark3, [101](#)
- USecondsToTicks
 - Mark3::Timer, [278](#)
- uTimerFlagActive
 - Mark3, [101](#)
- uTimerFlagCallback
 - Mark3, [101](#)
- uTimerFlagExpired
 - Mark3, [101](#)
- uTimerFlagOneShot
 - Mark3, [102](#)
- uTimerTicksInvalid
 - Mark3, [102](#)
- UnBlock
 - Mark3::BlockingObject, [112](#)
- Unlock
 - Mark3::Streamer, [239](#)
- Update
 - Mark3::Quantum, [215](#)
- User
 - Mark3, [98](#)
- WORD_ROUND
 - mark3c.h, [371](#)
- Wait
 - Mark3::ConditionVariable, [122](#)
 - Mark3::EventFlag, [138](#), [139](#)
 - Mark3::Notify, [200](#), [201](#)
- Wait_i
 - Mark3::EventFlag, [139](#)
- WakeMe
 - Mark3::EventFlag, [140](#)
 - Mark3::Mutex, [197](#)
 - Mark3::Notify, [201](#)
 - Mark3::Semaphore, [233](#)
- WakeNext
 - Mark3::Mutex, [197](#)
 - Mark3::Semaphore, [233](#)
- Write
 - Mark3::Streamer, [239](#), [240](#)
- Yield
 - Mark3::Thread, [259](#)